

Face Recognition IV

CSE 40537/60537 Biometrics

Daniel Moreira
Spring 2022



Today you will...

Get to know
Deep-learning-based face recognition.

Feature Extraction

RECAP

Focus

2D-appearance-based methods.



Types

Handcrafted features from Computer Vision.

Data-driven learned features from Machine Learning.

Feature Extraction

RECAP

Focus

2D-appearance-based methods.




Types

Handcrafted features from Computer Vision.

Data-driven learned features from Machine Learning.

Face Recognition

RECAP



Data-Driven Face Recognition

~~Homework~~

Deep Convolutional Neural Networks (CNN)

Data-Driven Face Recognition


~~Non-CAP~~

Deep Convolutional Neural Networks (CNN)

From pixels to classification decision.

Hierarchy of feature extractors.

Each layer extracts features from previous layer.



Data-Driven Face Recognition


~~Non-CAP~~

Deep **Convolutional** Neural Networks (CNN)

Convolutional Layers

E.g., layers 1 and 2.

Feature extractors are convolutional operations which are performed on the output of the previous layer.



Data-Driven Face Recognition

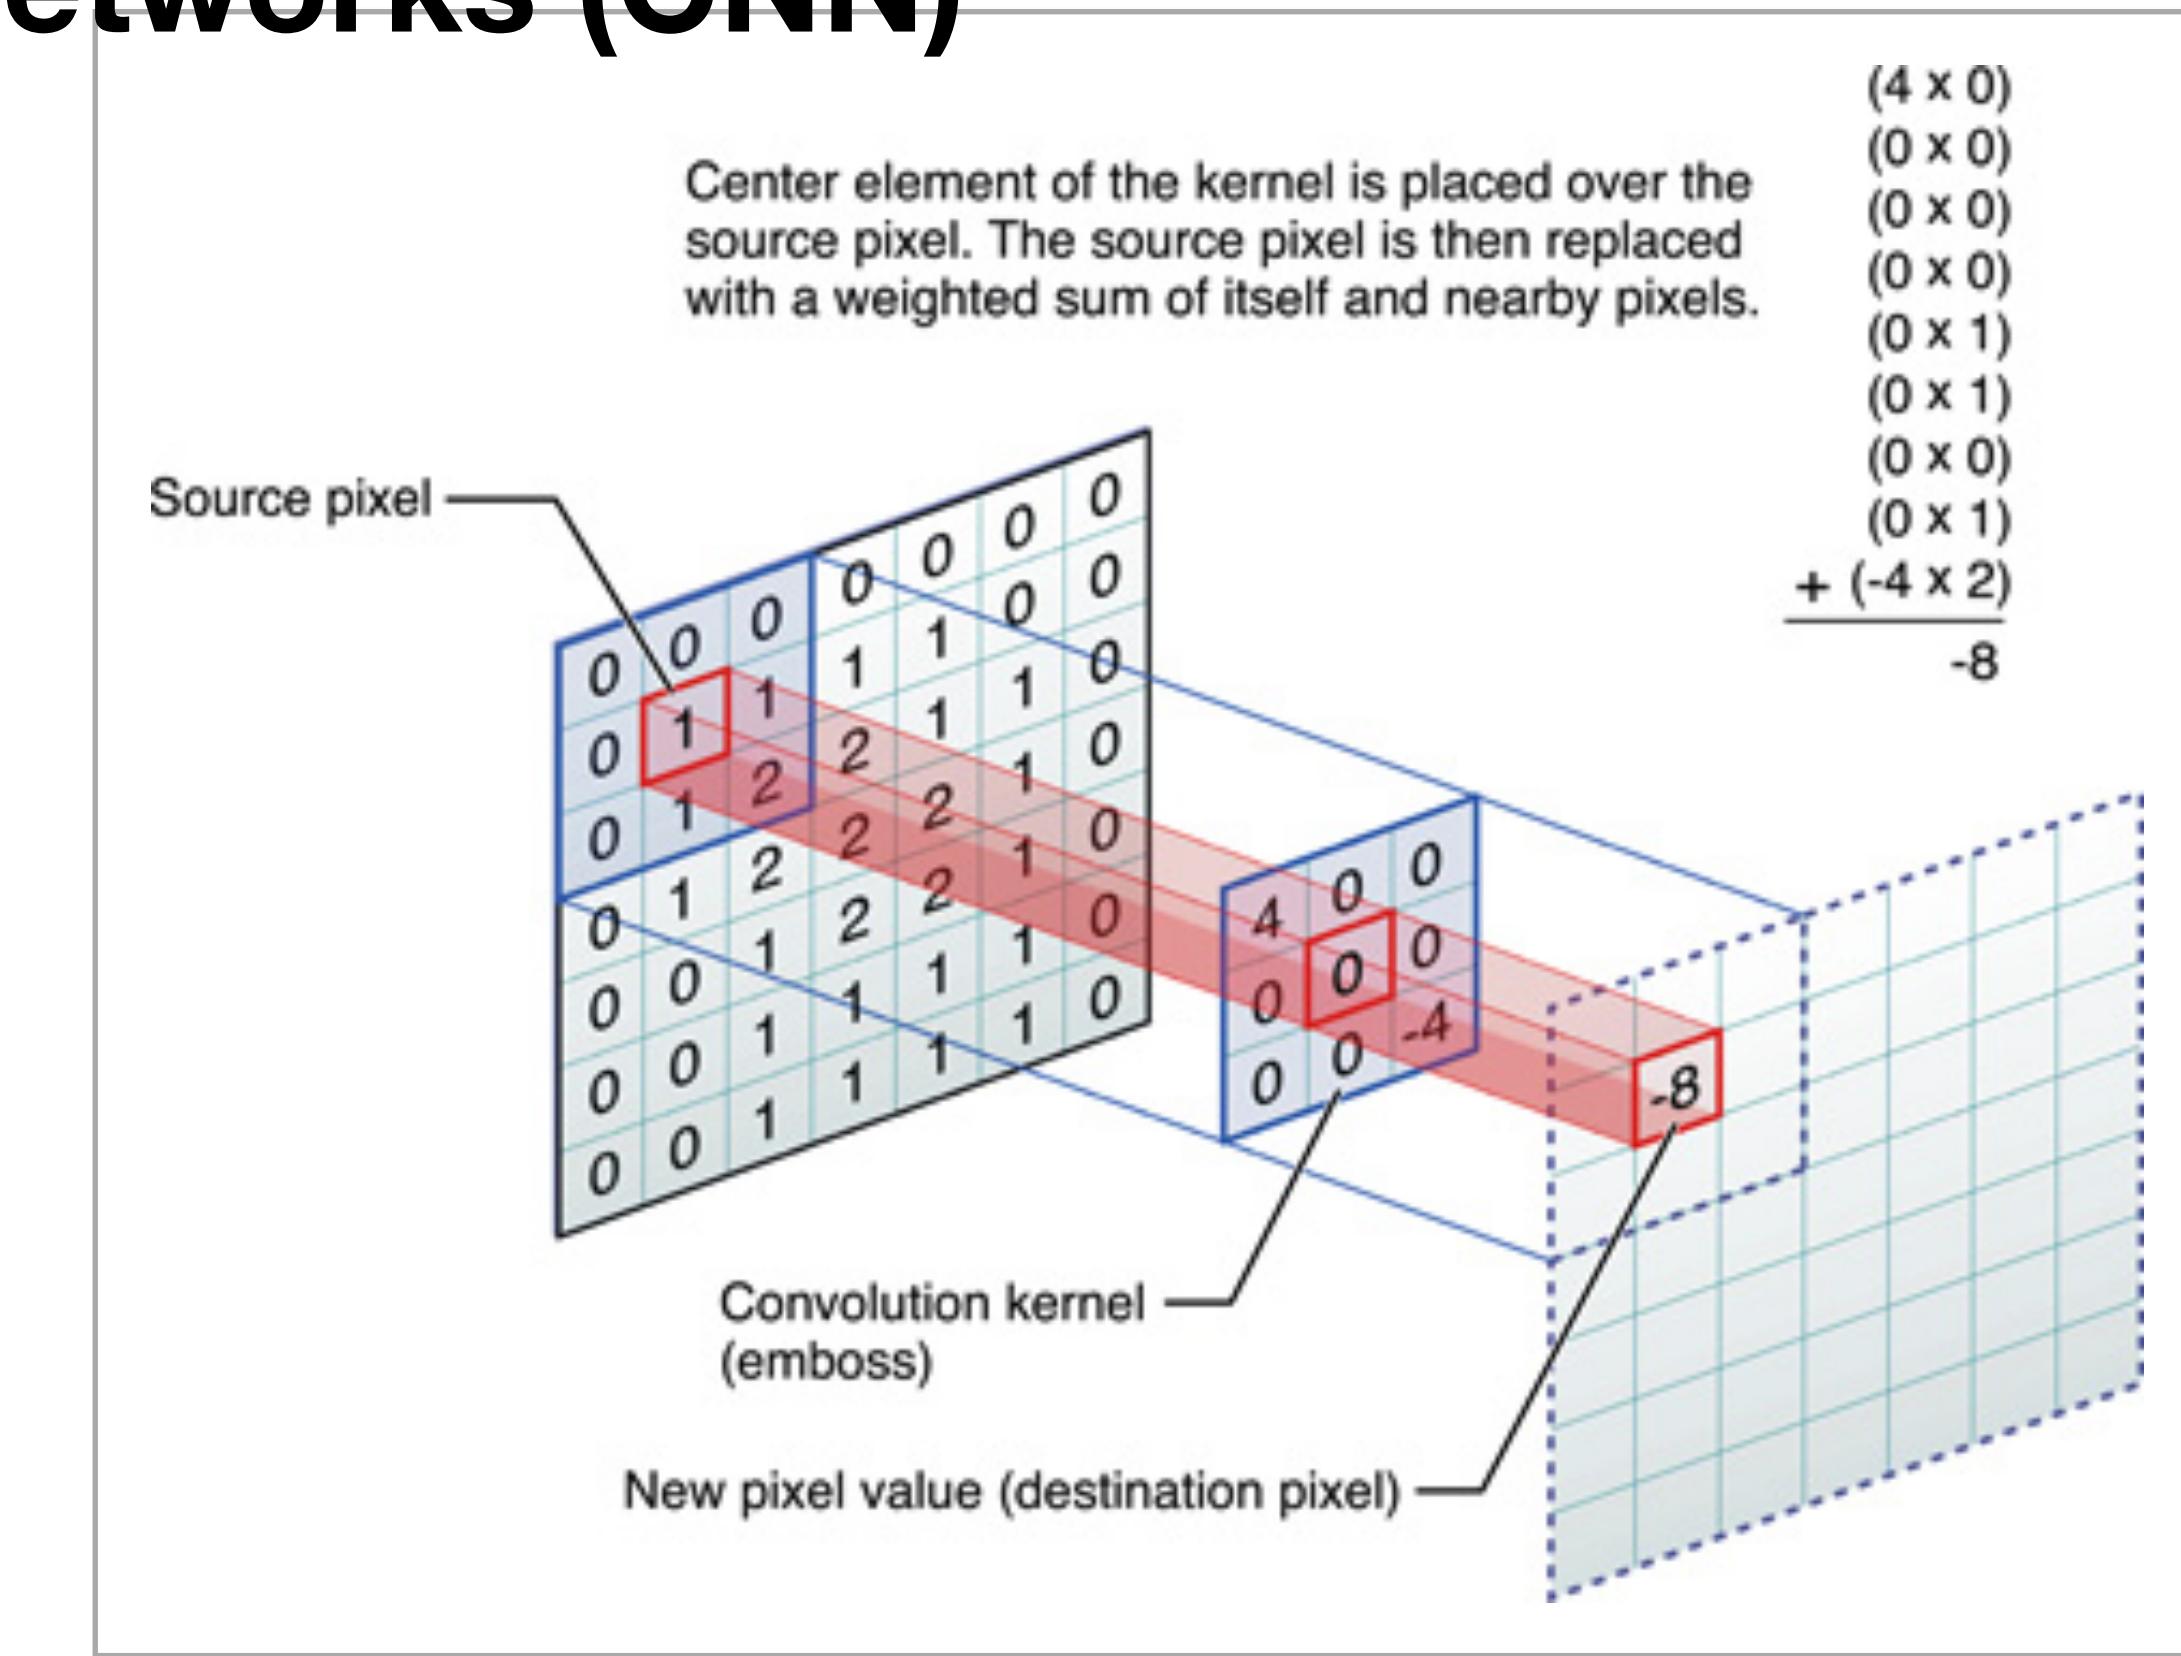
~~Homework CAP~~

Deep **Convolutional** Neural Networks (CNN)

Convolutional Layers

E.g., layers 1 and 2.

Feature extractors are convolutional operations which are performed on the output of the previous layer.



Source:<https://developer.apple.com/library/archive/documentation/Performance/Conceptual/vImage/ConvolutionOperations/ConvolutionOperations.html>

Data-Driven Face Recognition

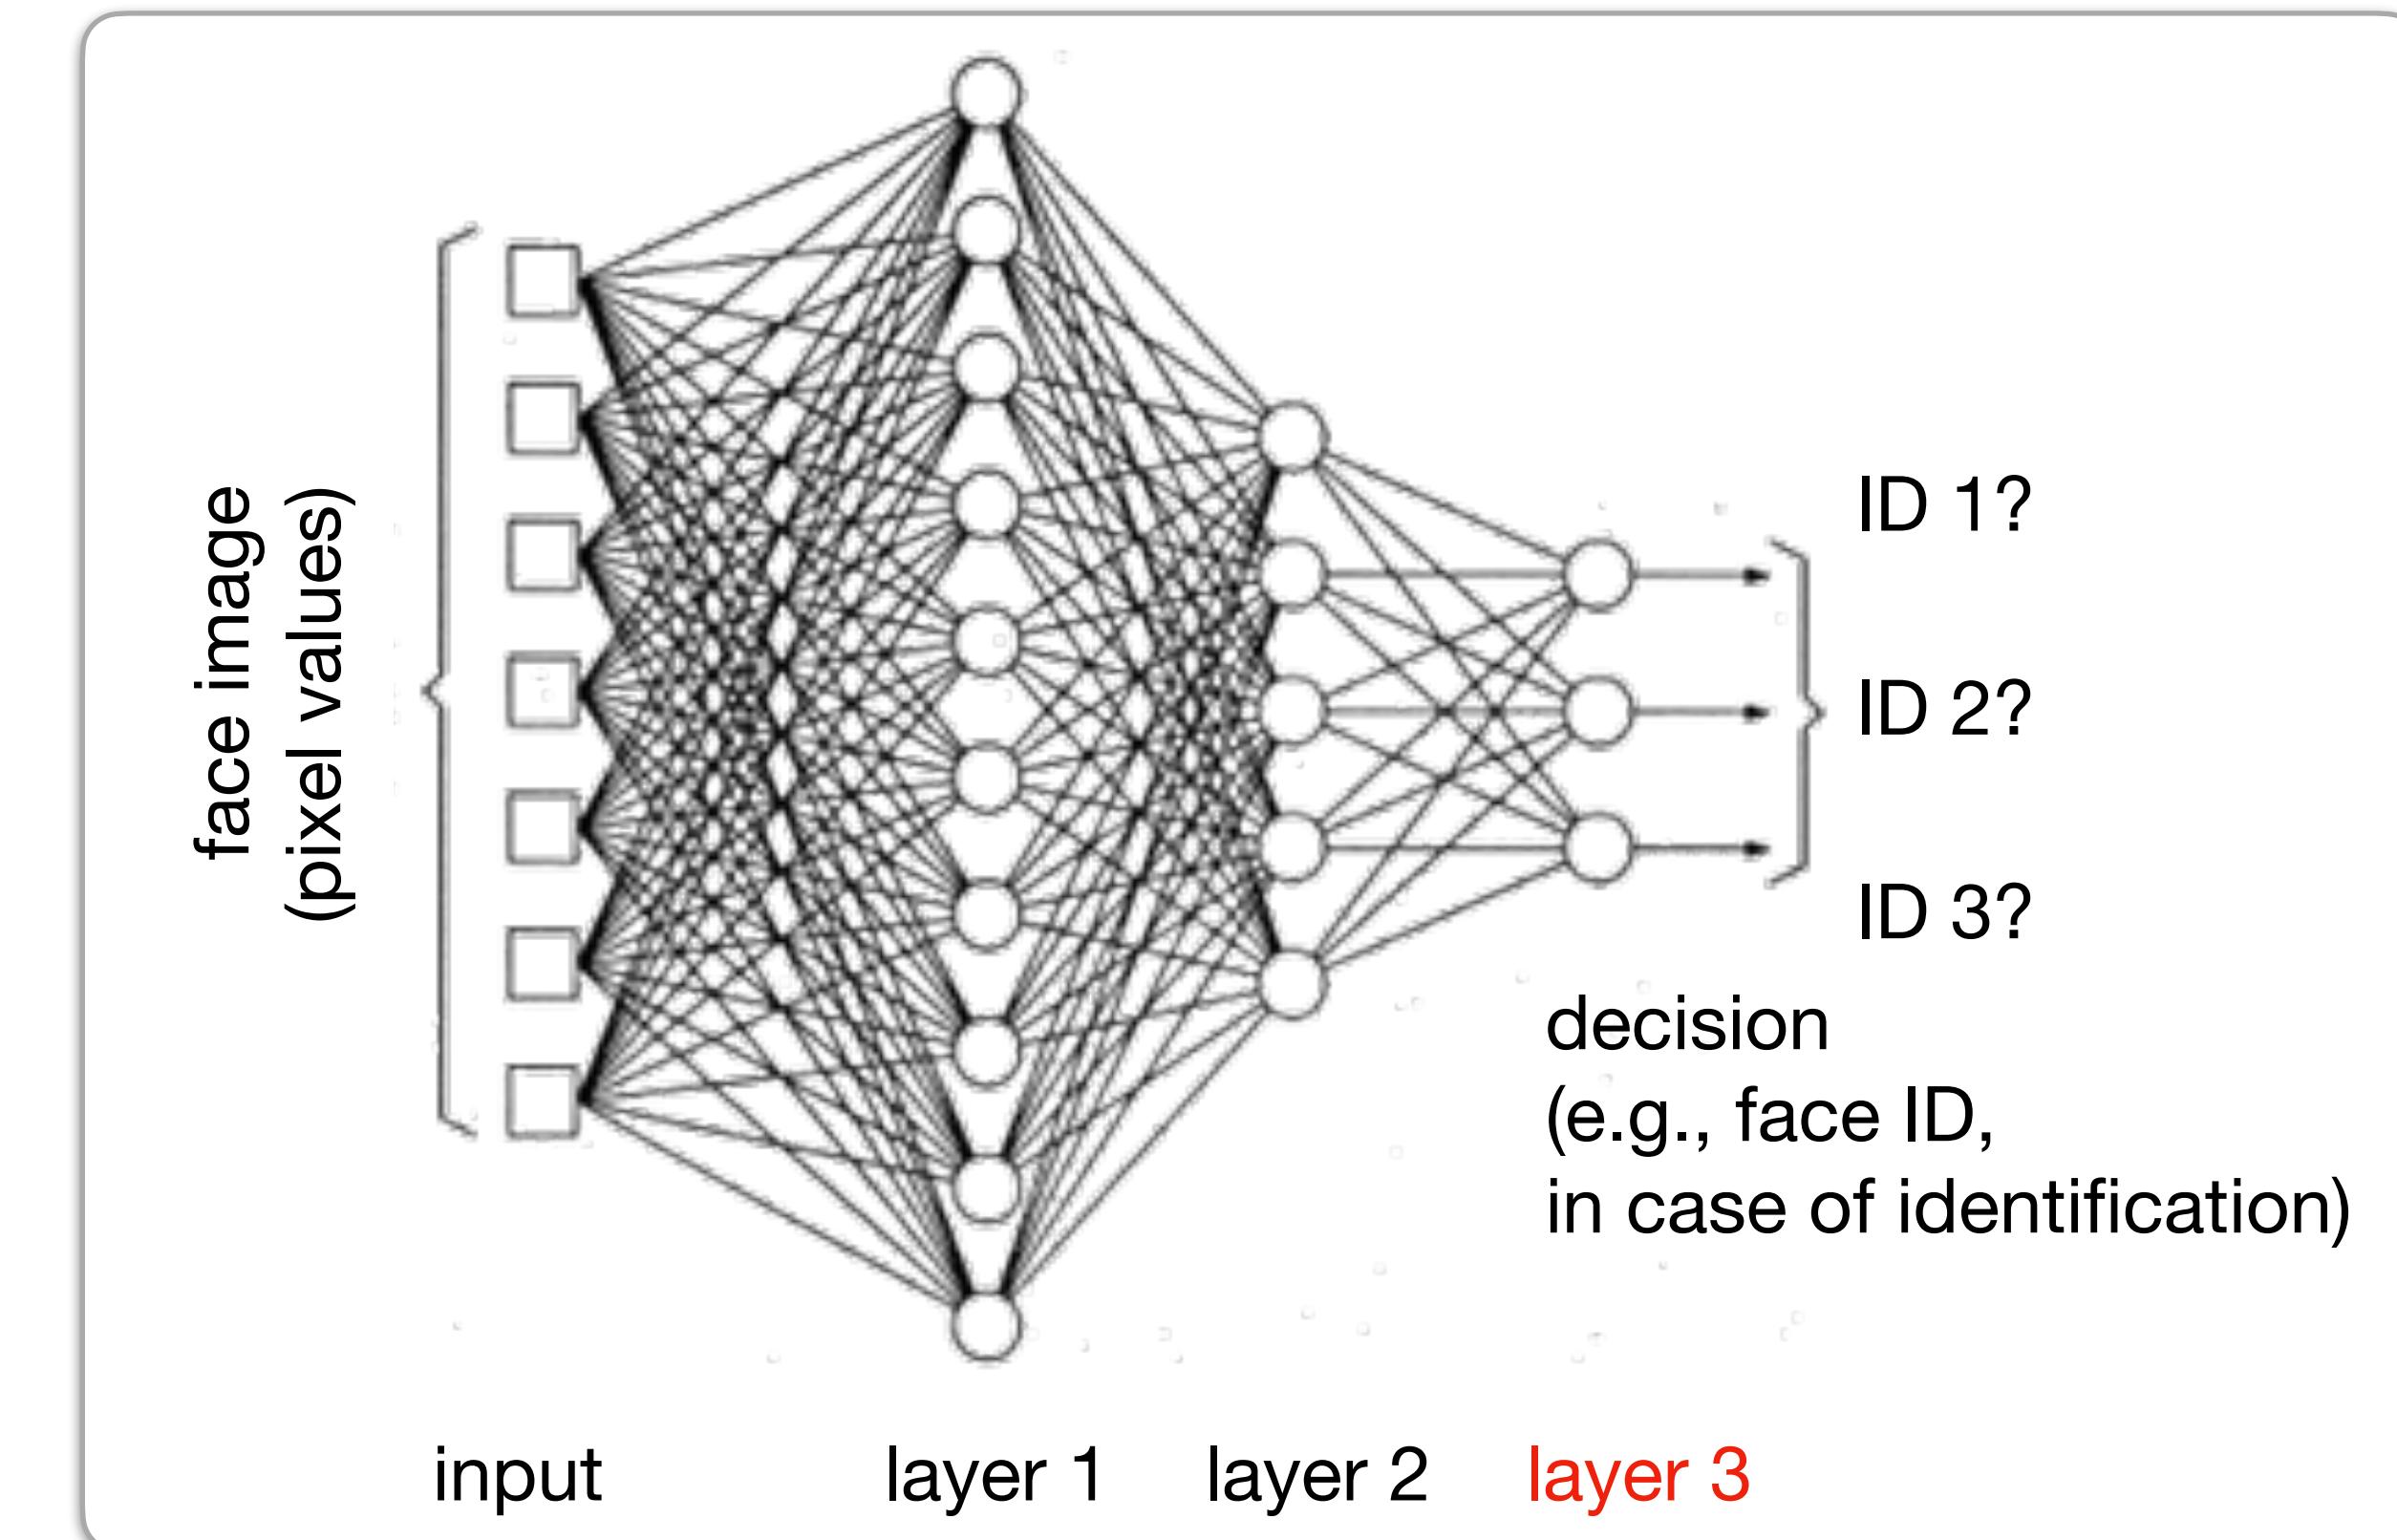
~~Homework~~

Deep **Convolutional** Neural Networks (CNN)

Fully Connected Layer

E.g., layer 3.

It performs the classification, presenting one score output for each class (identity, in the case of Biometrics).



Data-Driven Face Recognition

~~Non-CAP~~

Deep Convolutional Neural Networks (CNN)


How deep can they be?

“Deep” refers to the number of layers.

E.g., VGG16

Simonyan and Zisserman

*Very Deep Convolutional Networks
for Large-Scale Image Recognition*



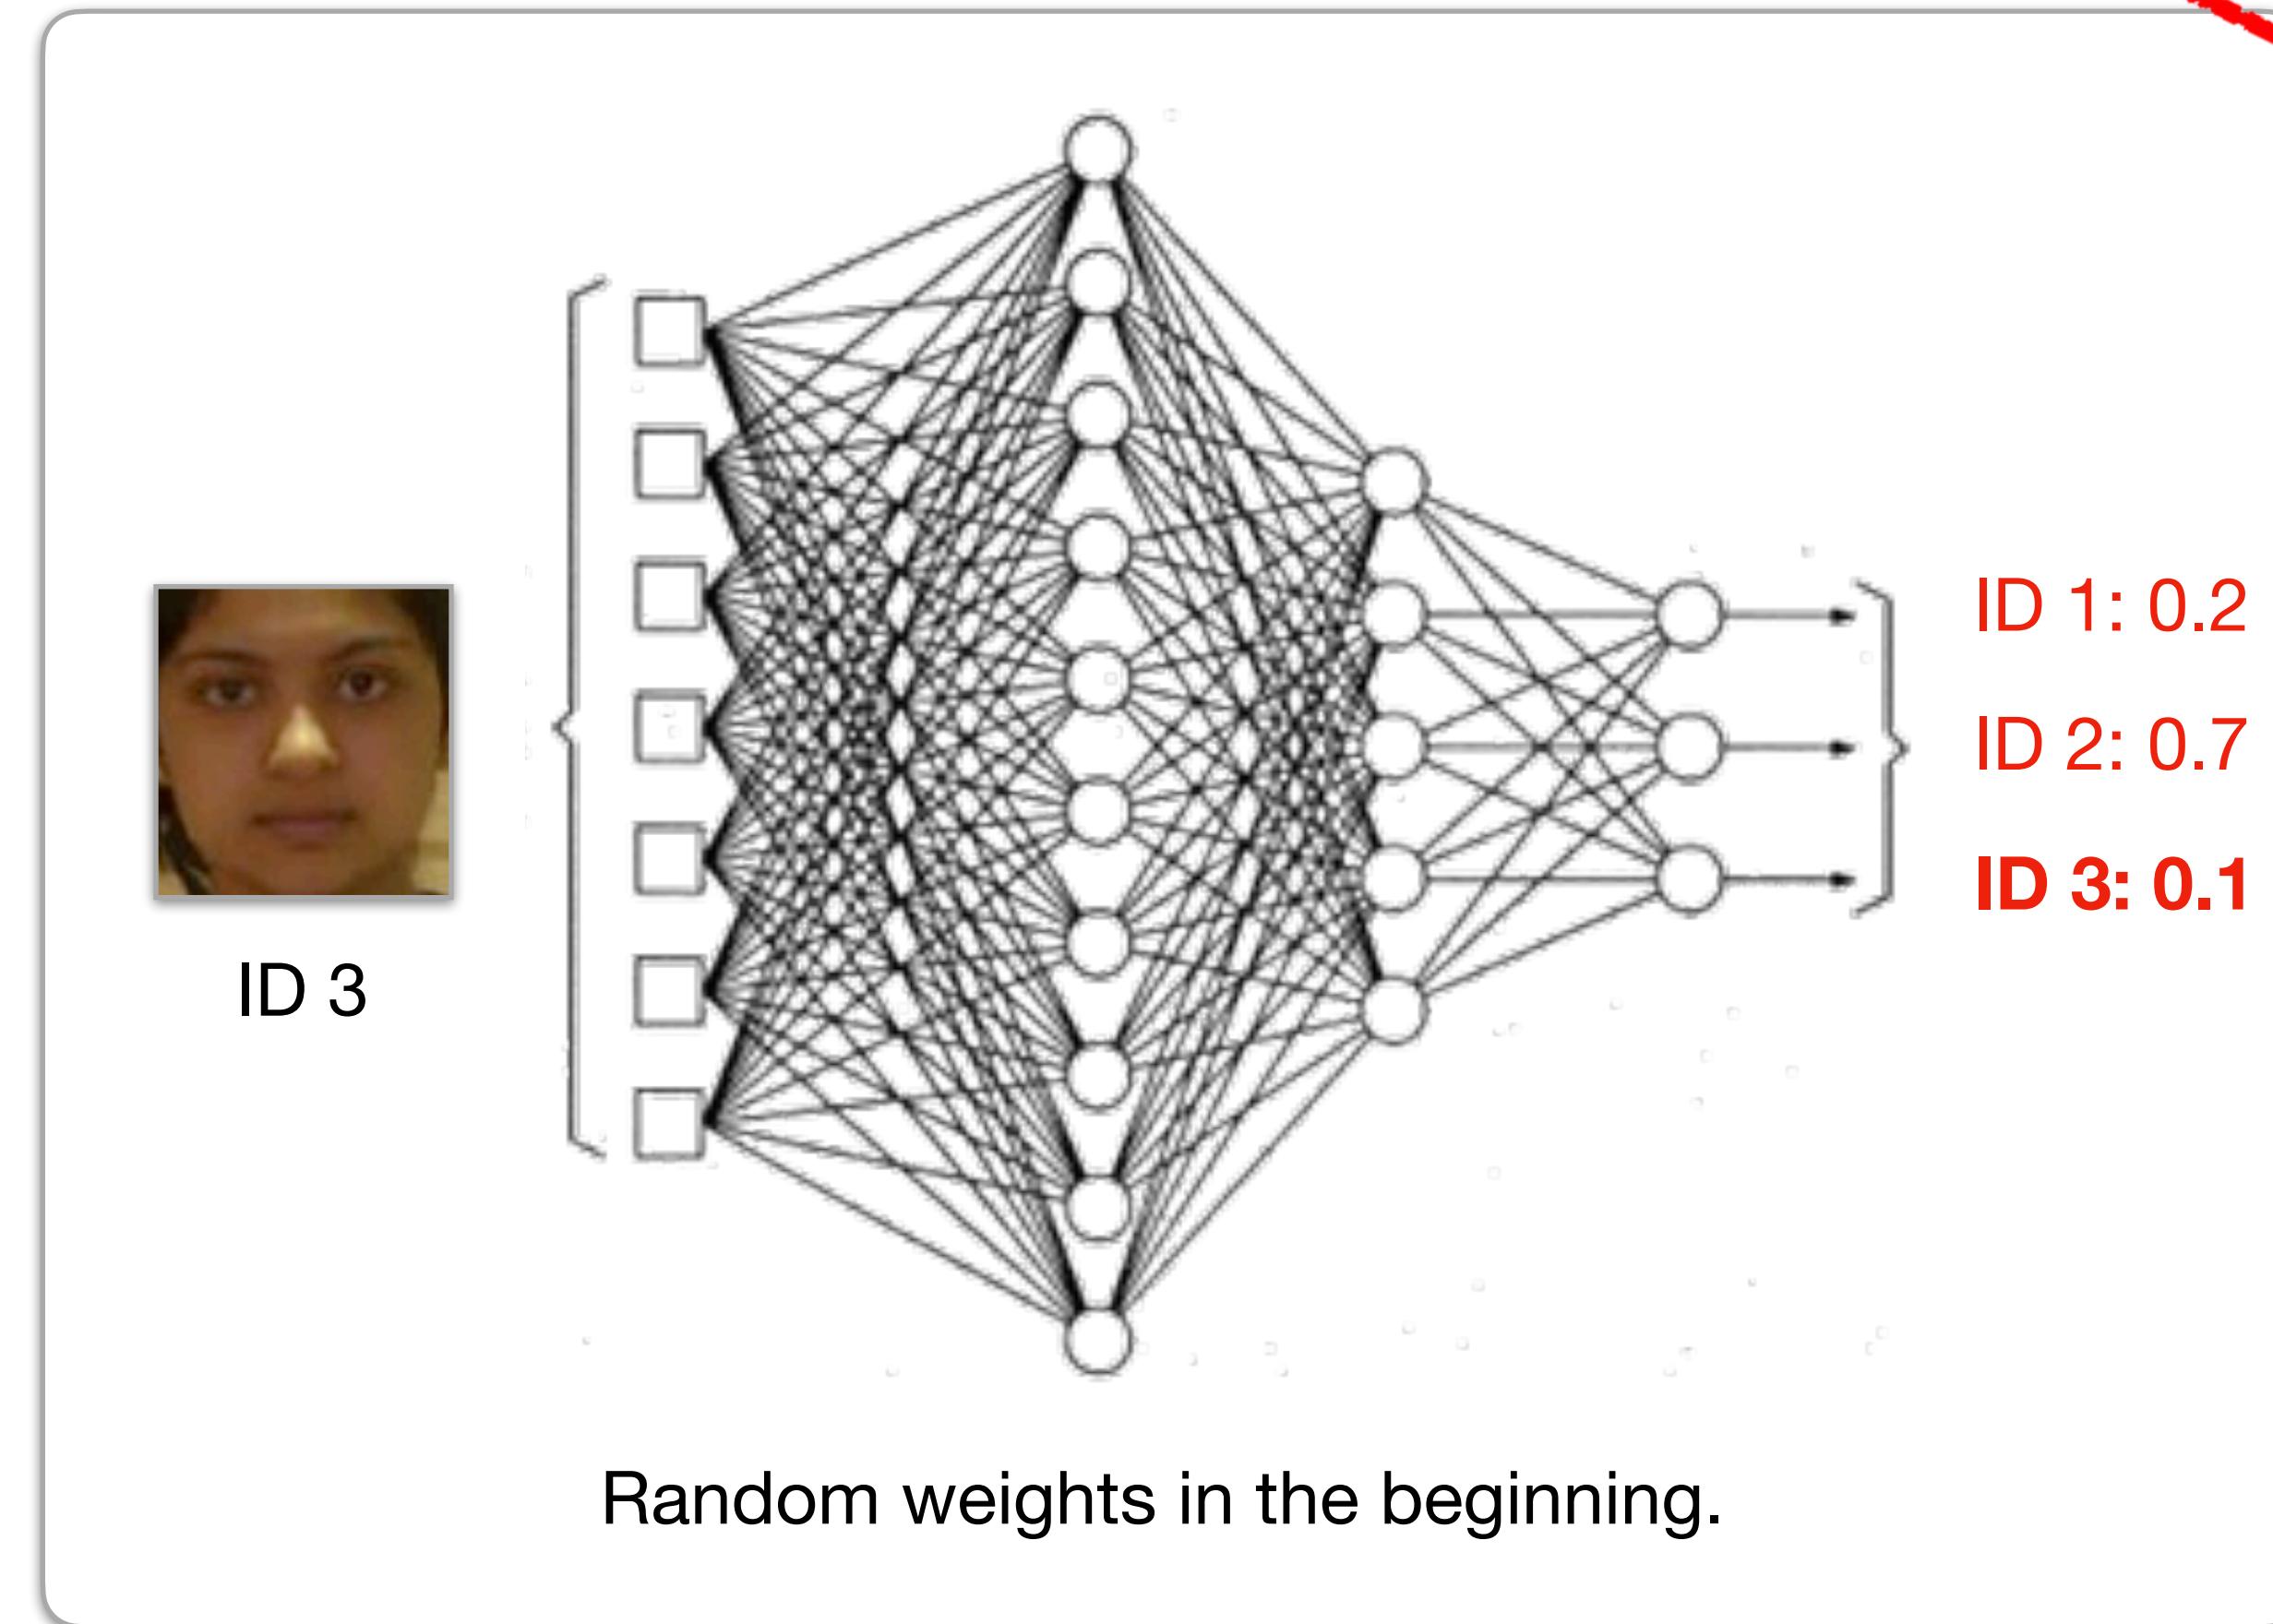
Data-Driven Face Recognition

~~Homework~~

Deep Learning

Training

Labeled examples
(e.g., faces and expected IDs)
are used to teach the network
to classify them correctly.



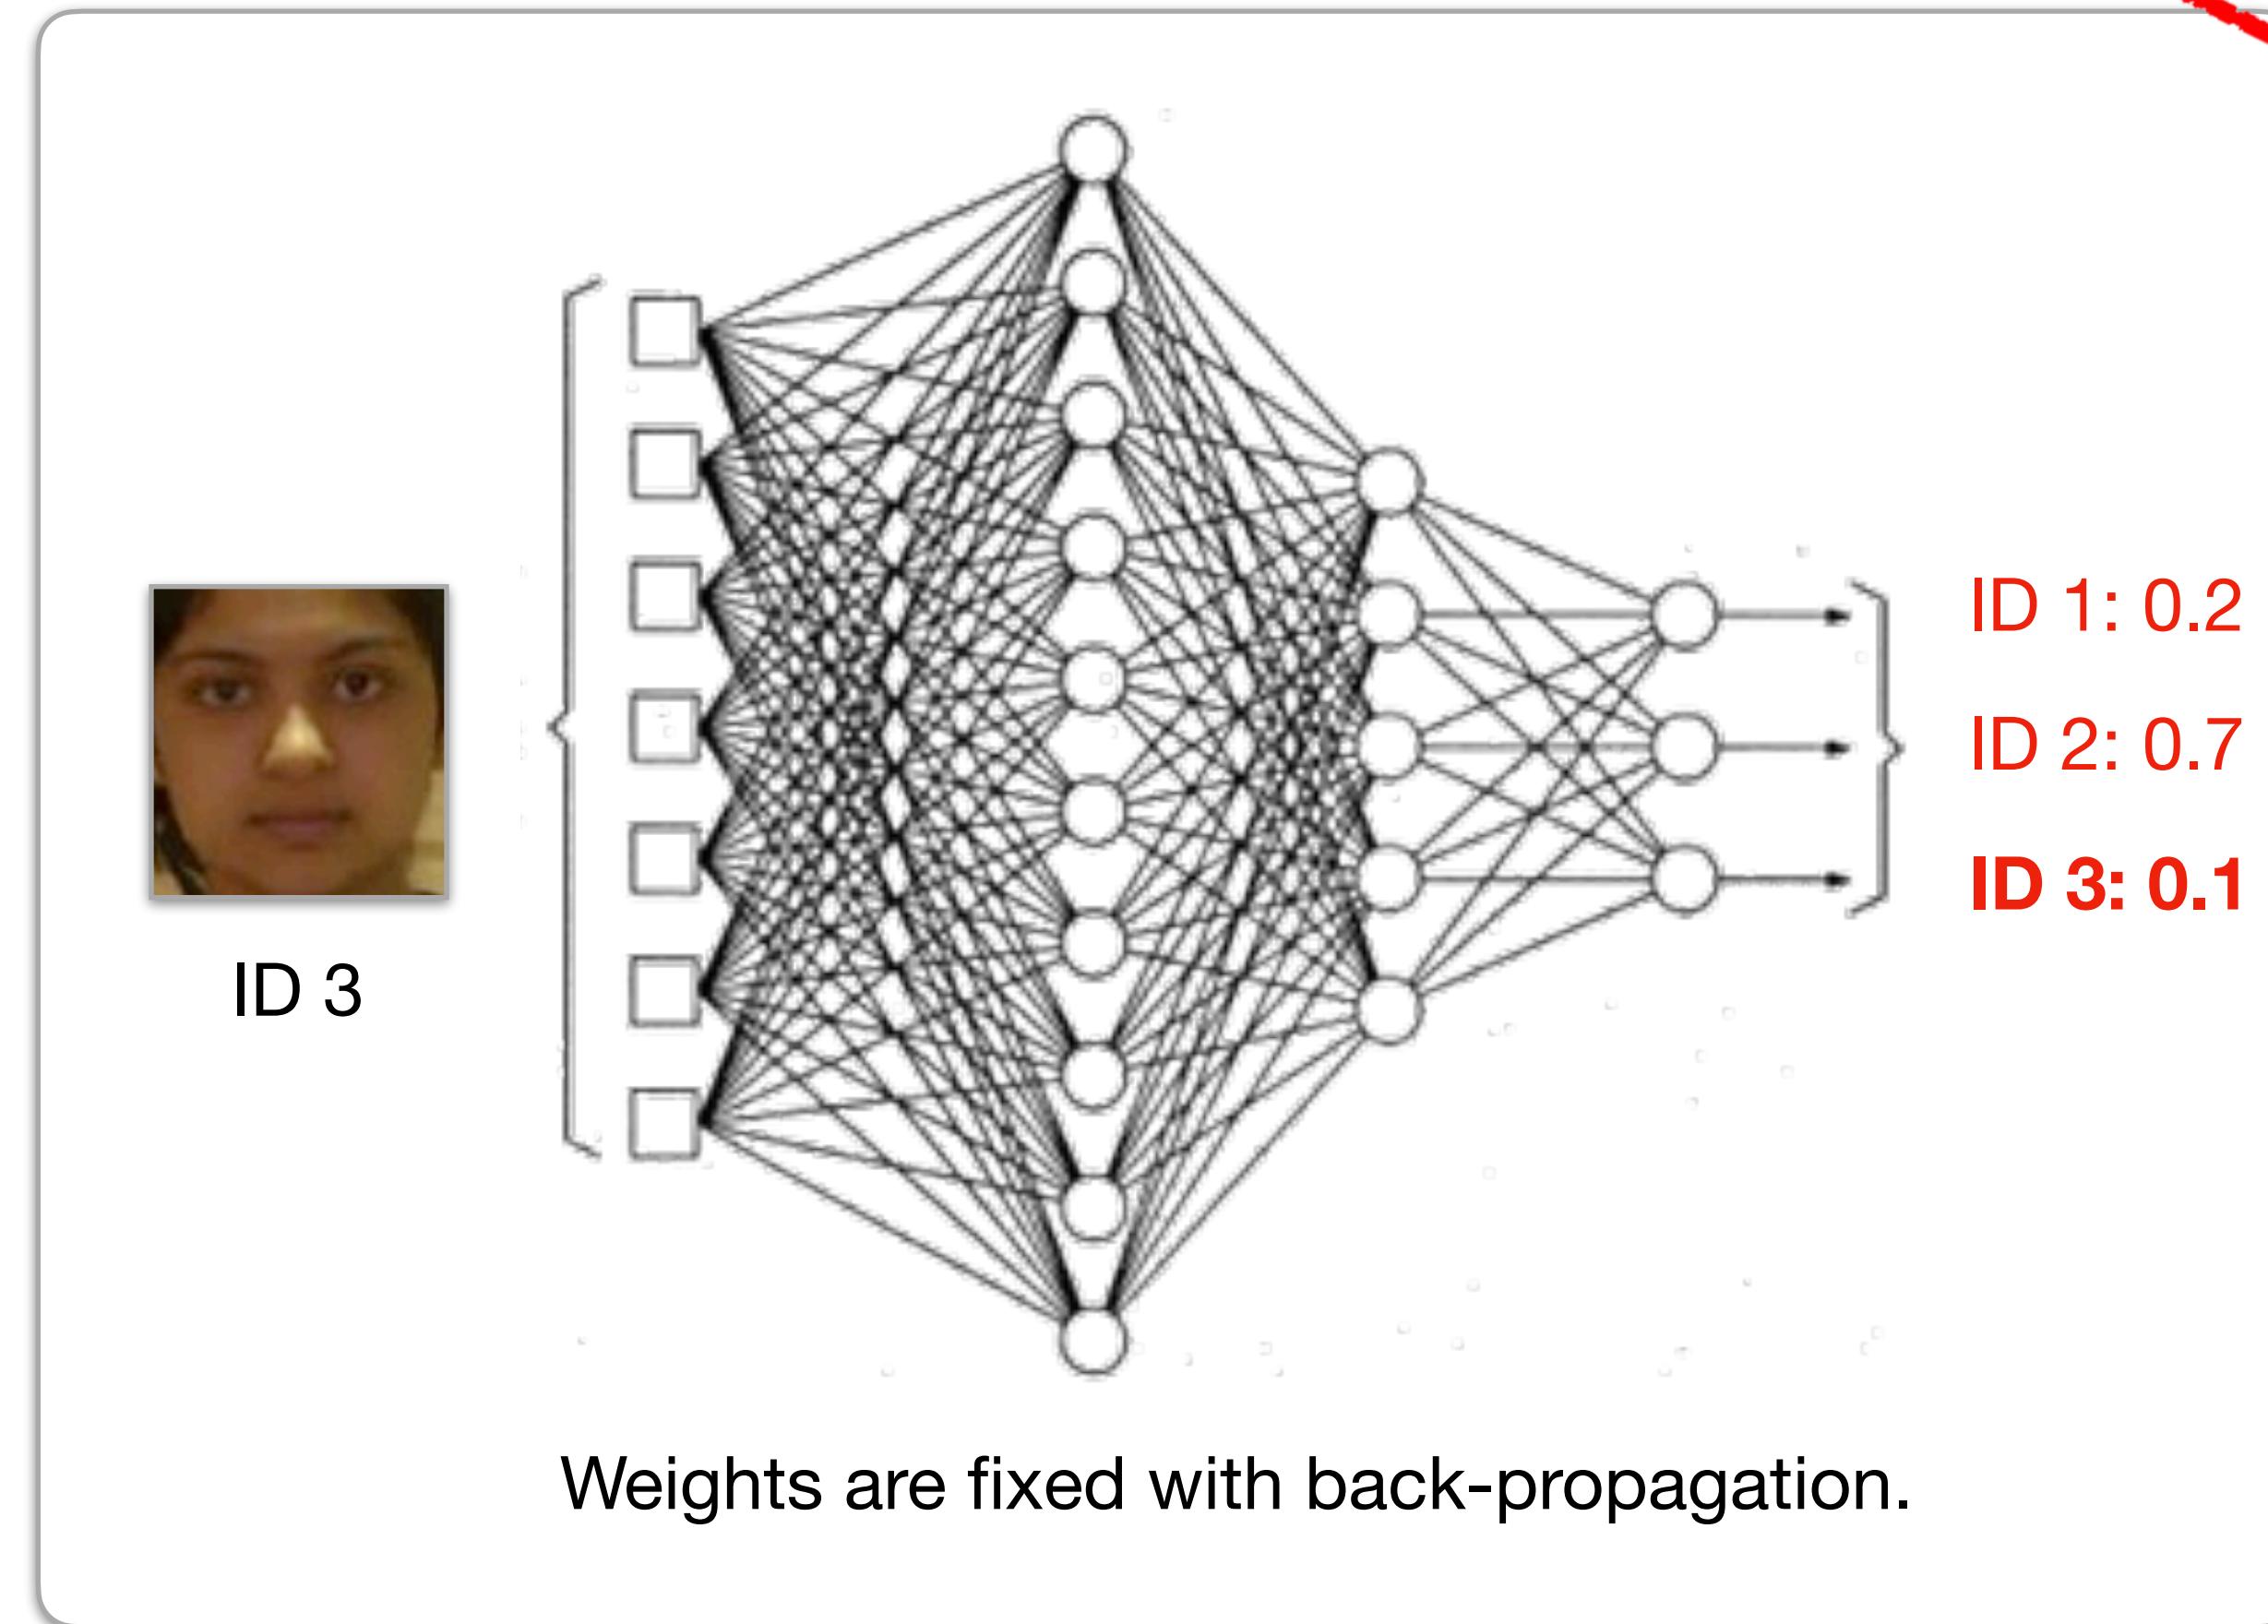
Data-Driven Face Recognition

~~Face CAP~~

Deep Learning

Training

Back-propagation is used to fix the weights of the convolutions within the network.

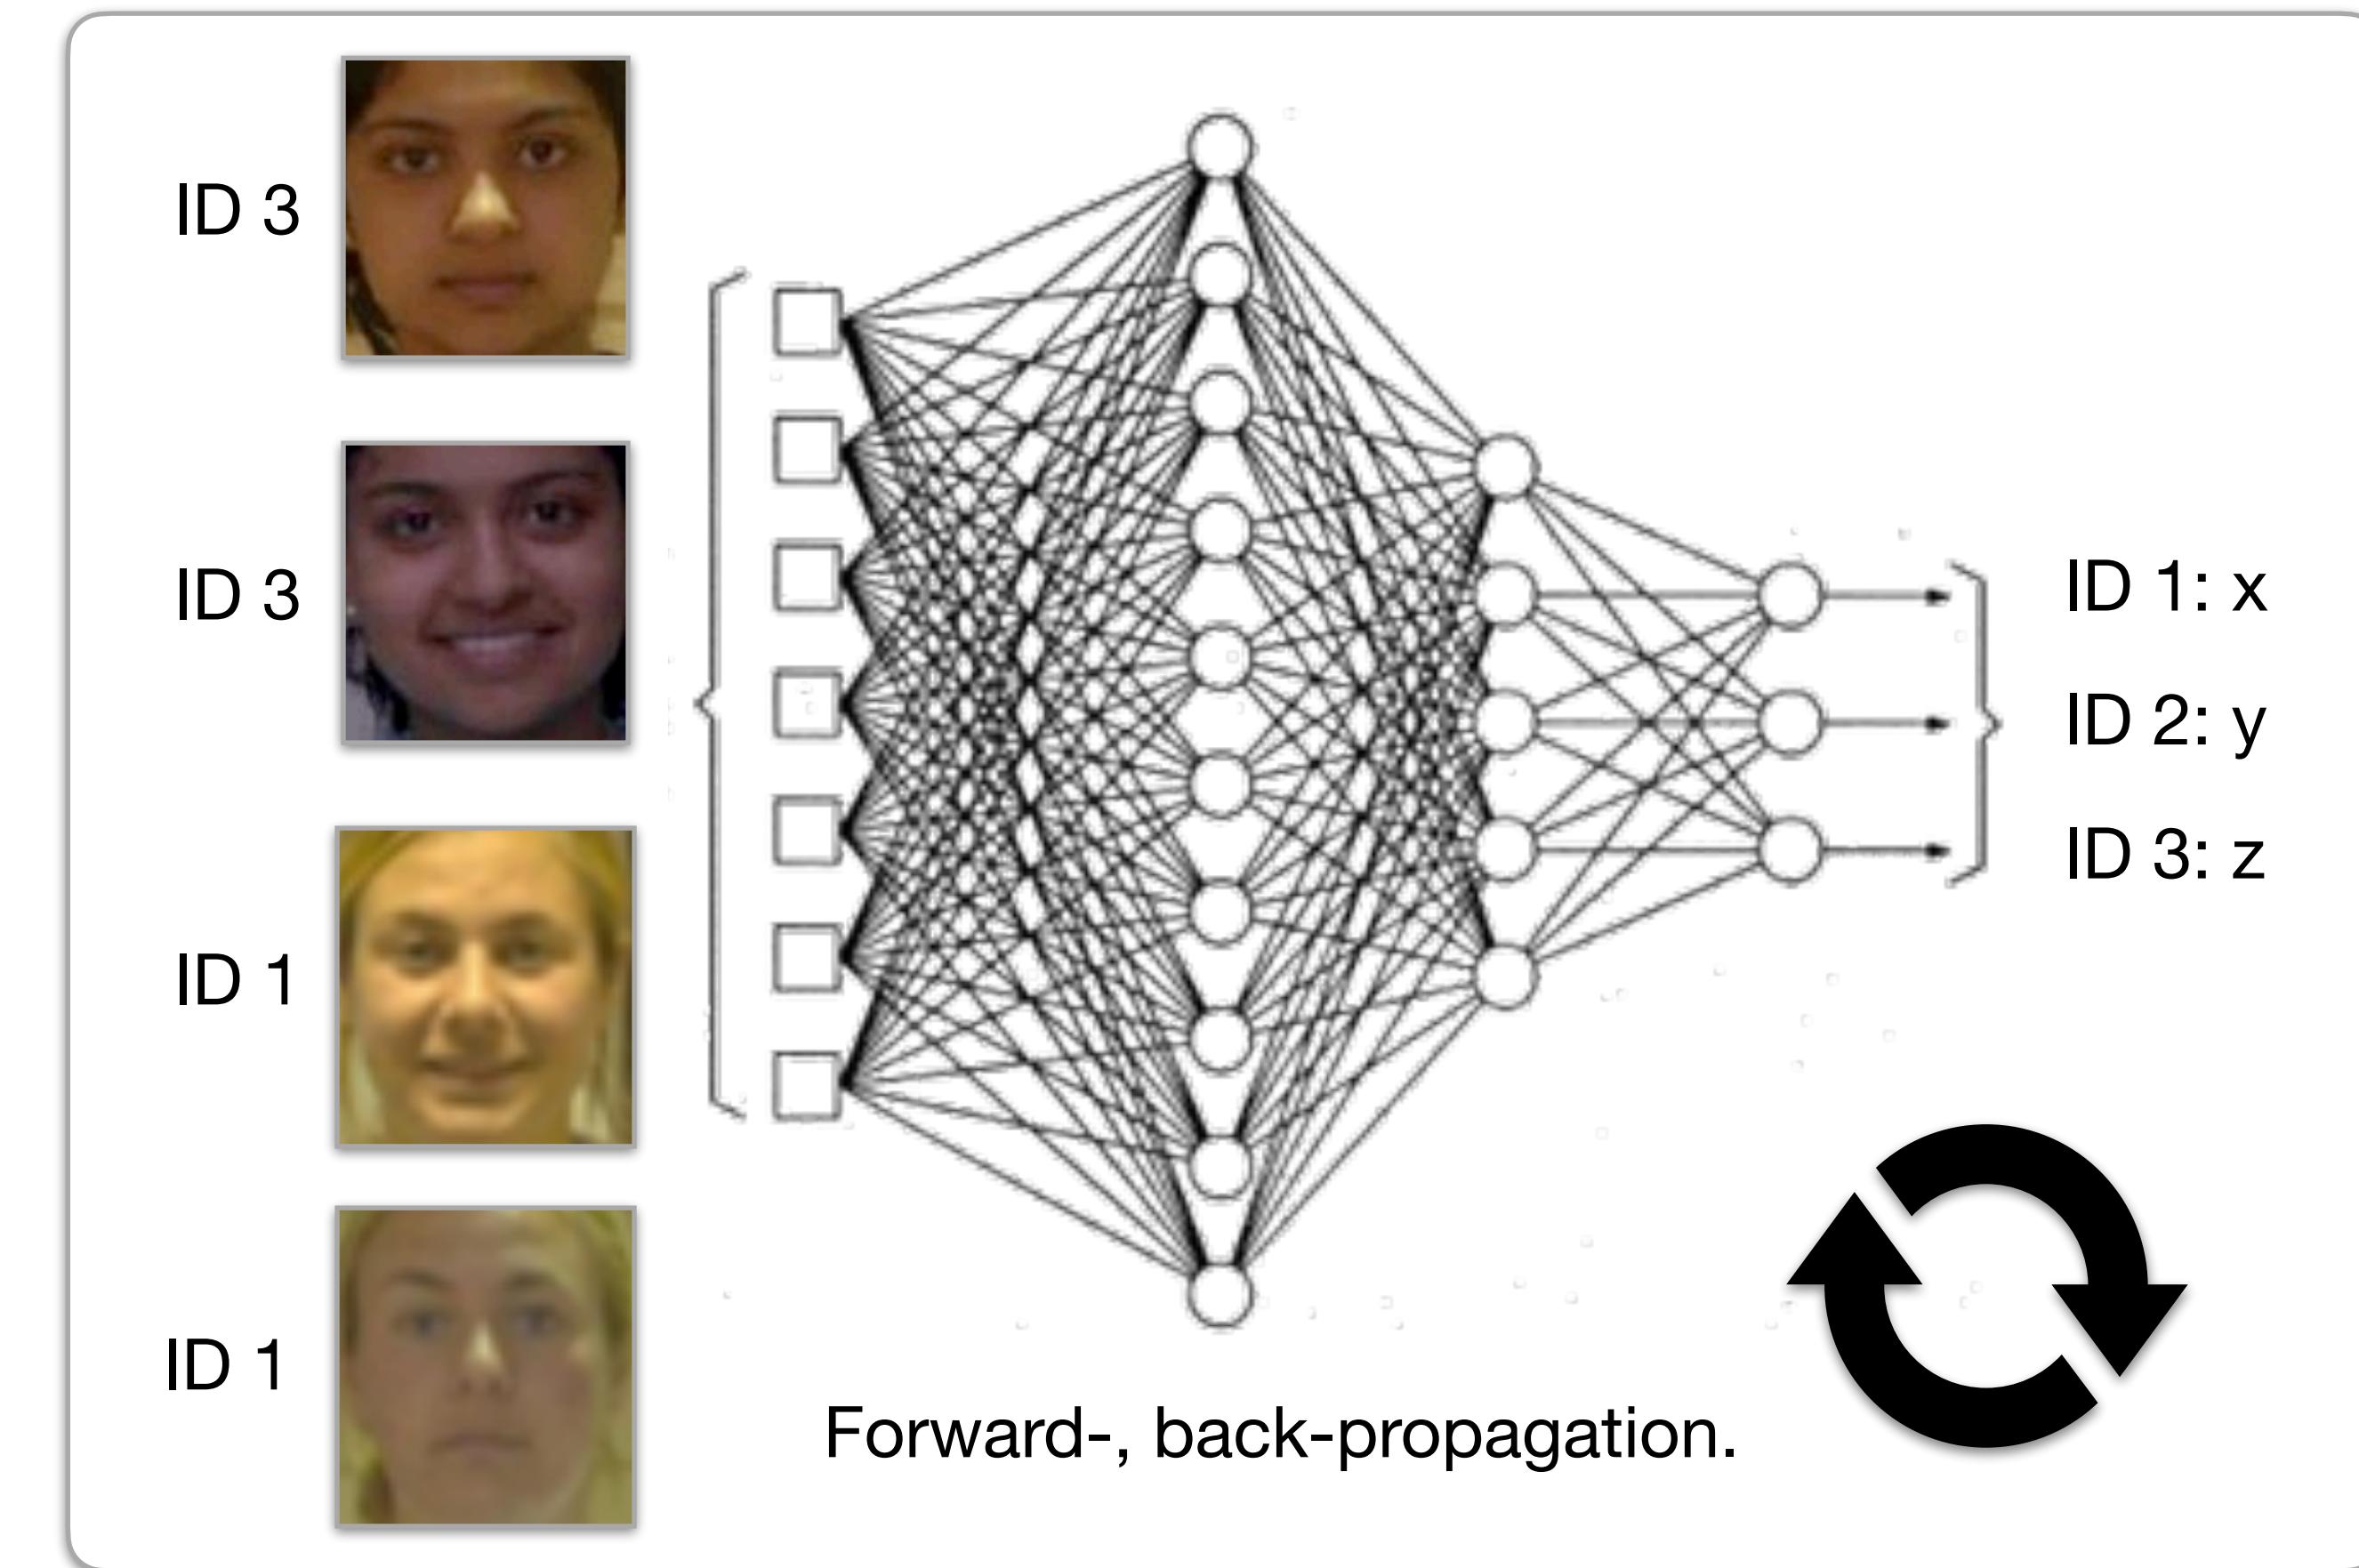


Data-Driven Face Recognition

~~Homework~~

Deep Learning

Present various examples of each class and perform forward-, back-propagation.




Data-Driven Face Recognition

Deep Learning

Optimization target:
minimize classification
error through
loss function.

Popular function:
cross-entropy loss.



Data-Driven Face Recognition


Deep Learning

Cross-entropy Loss (CE)

$$CE = \sum_{face=1}^m \sum_{ID=1}^n (-\log(output(ID)))$$

CNN output for ID

#training faces #people's IDs



Data-Driven Face Recognition

Deep Learning


Cross-entropy Loss (CE)

$$CE = \sum_{face=1}^m \sum_{ID=1}^n (-\log(output(ID)))$$

#training faces

#people's IDs

CNN output for ID



Data-Driven Face Recognition

~~Non-CAP~~


Deep Learning

What are the cons here?

How to enroll a new person?

Fixed number of classes
(i.e., persons).

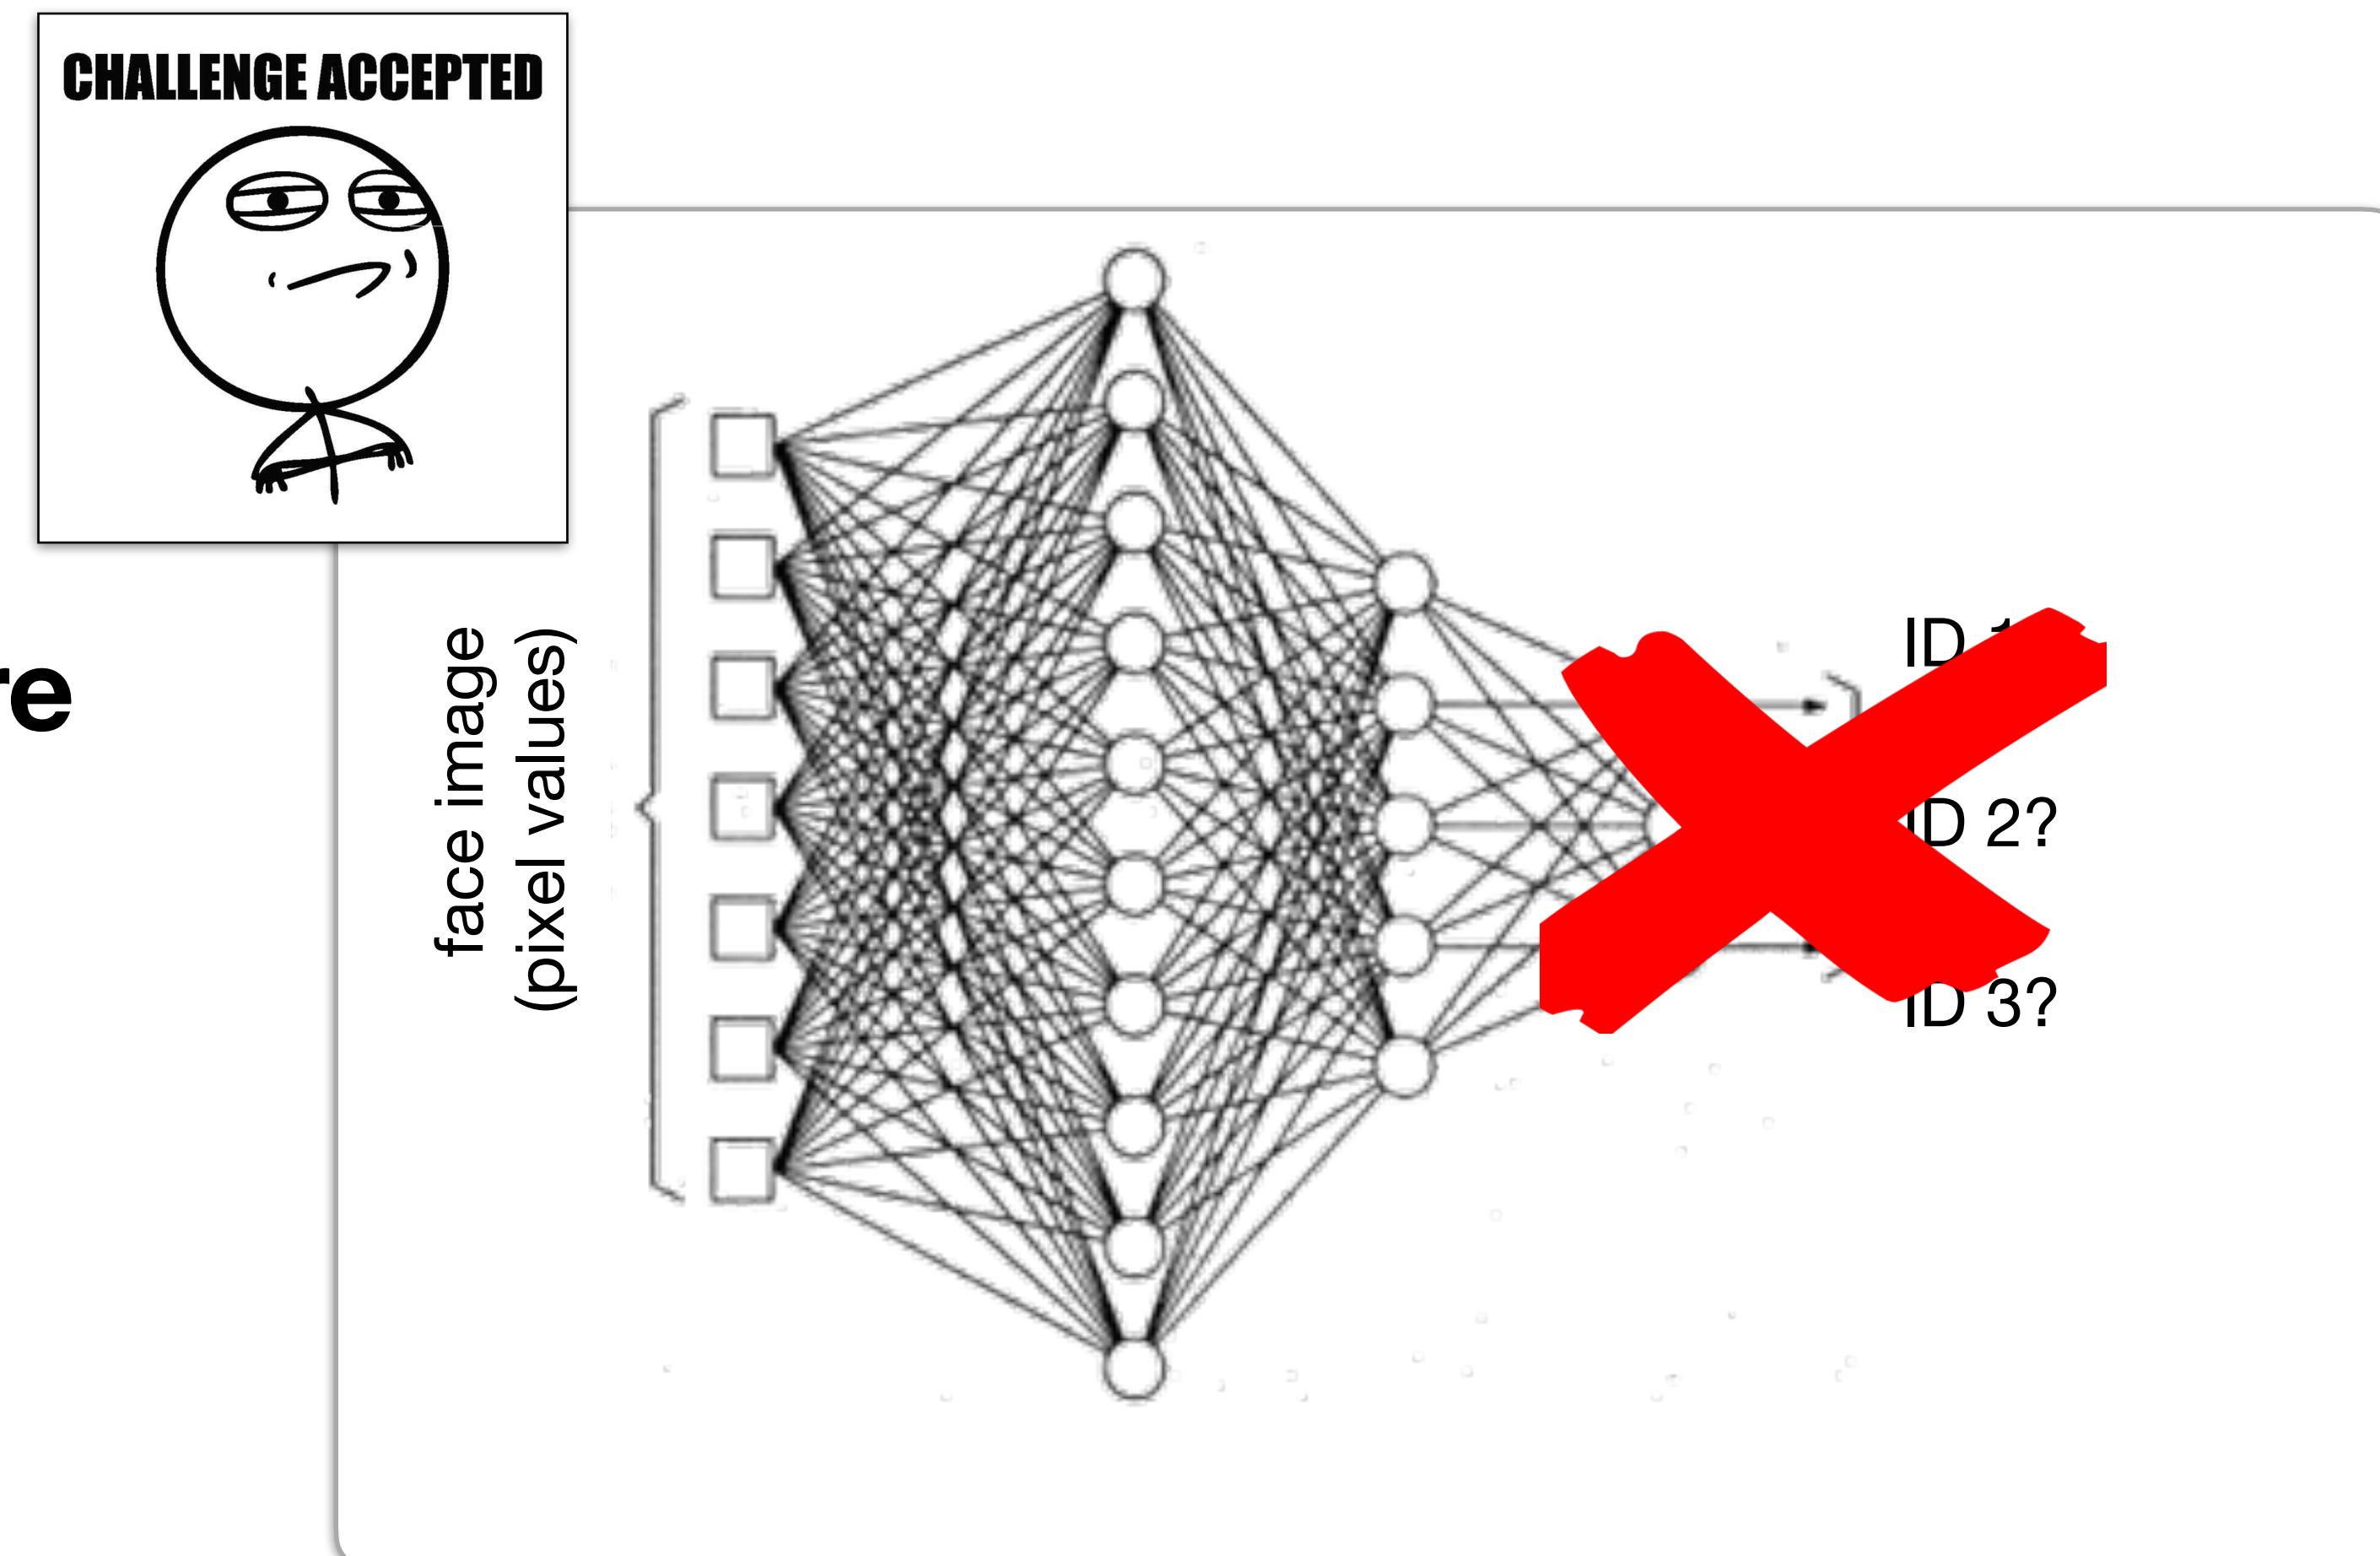
Need for large training dataset (thousands of sample per class).



Data-Driven Face Recognition

Deep Learning


How to make CNN more flexible?
Remove fully connected layer and use last convolutional layers as a feature descriptor.



Data-Driven Face Recognition

Deep Learning

How to make CNN more flexible?
Remove fully connected layer and use last convolutional layers as a feature descriptor.




Data-Driven Face Recognition

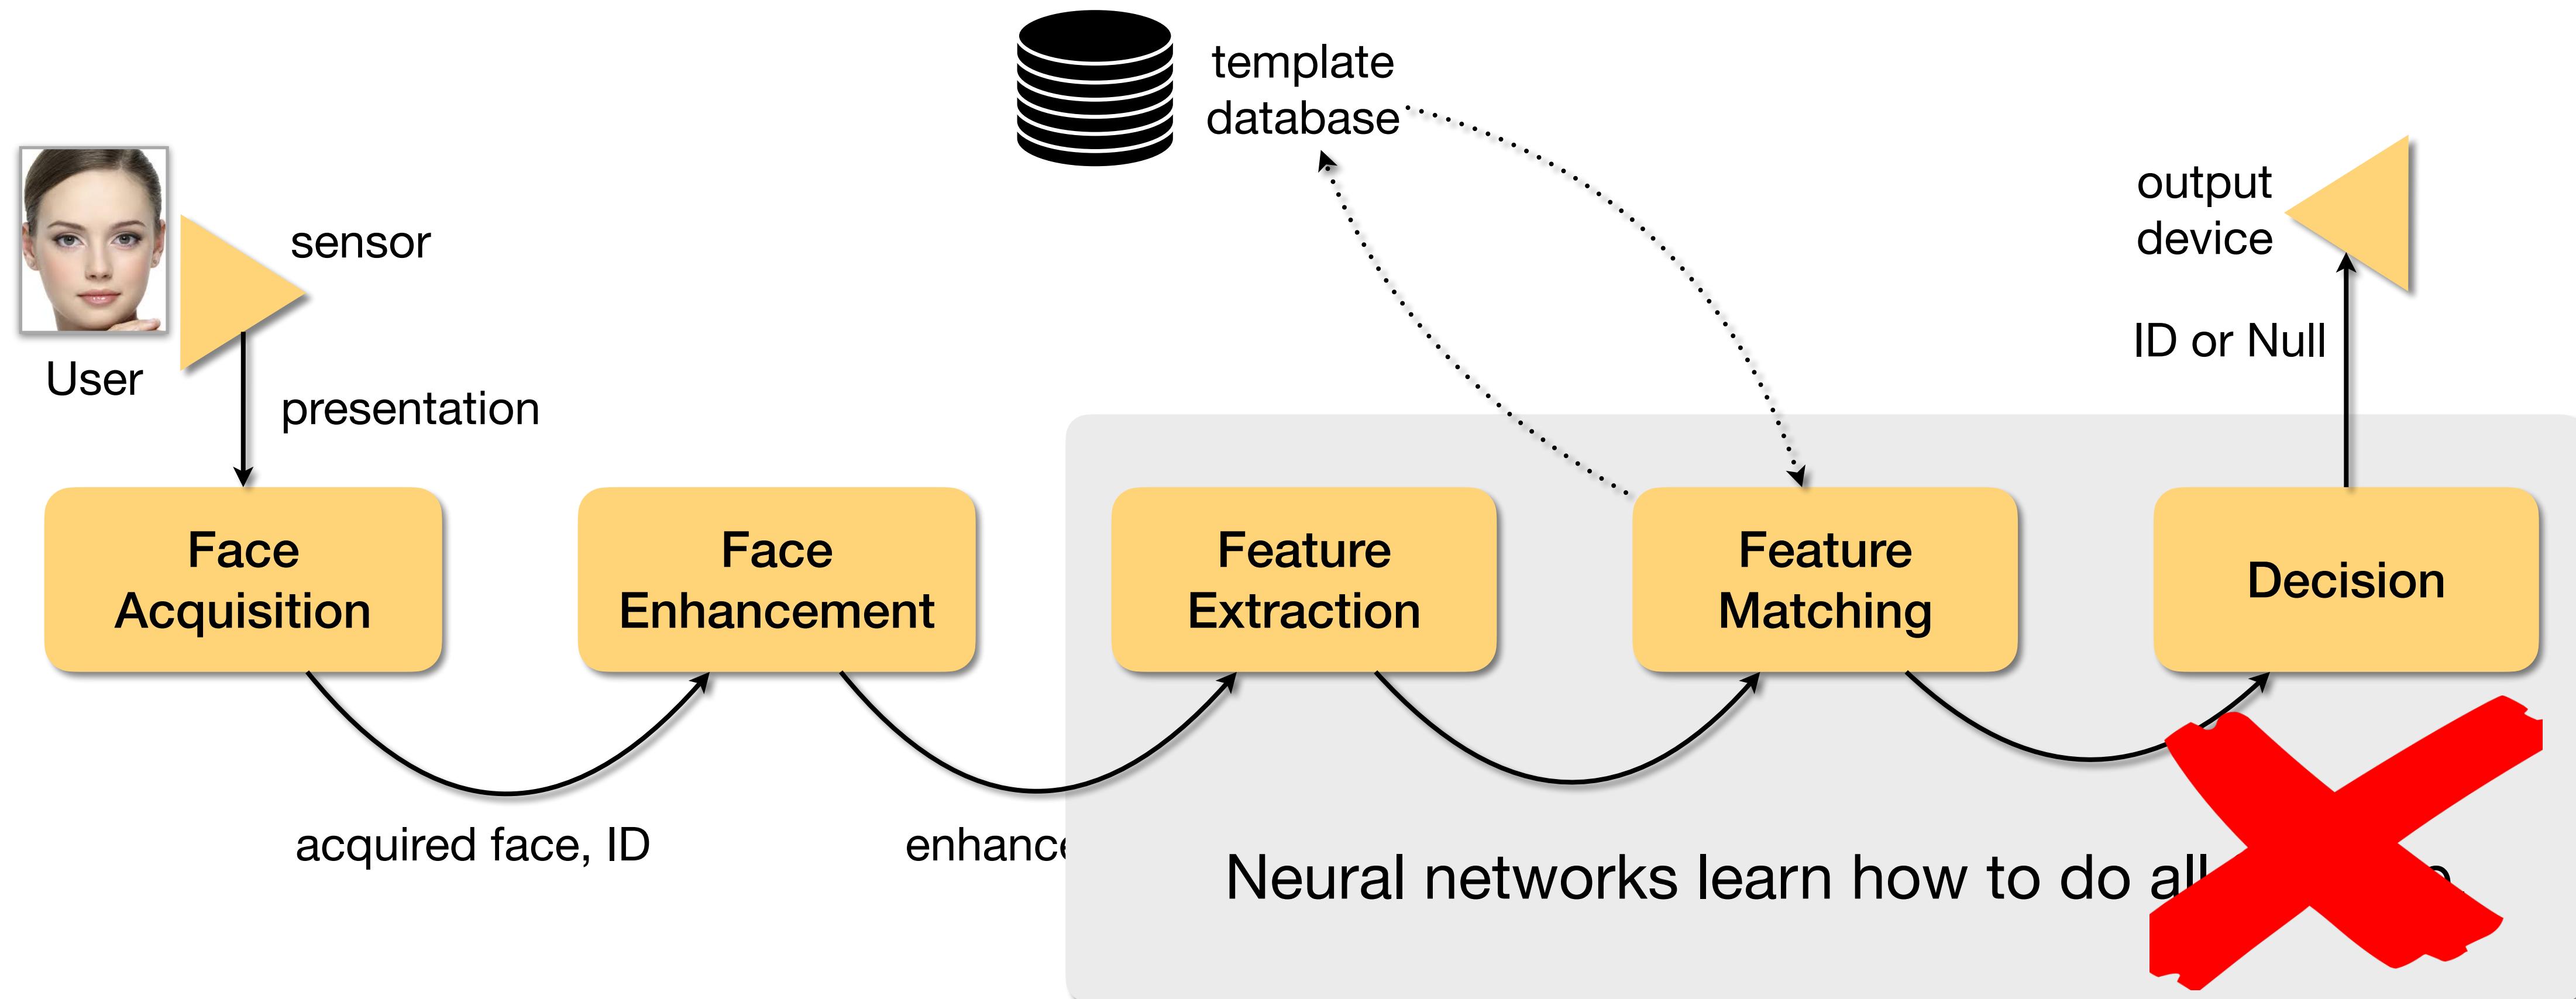
Deep Learning

How to make CNN more flexible?

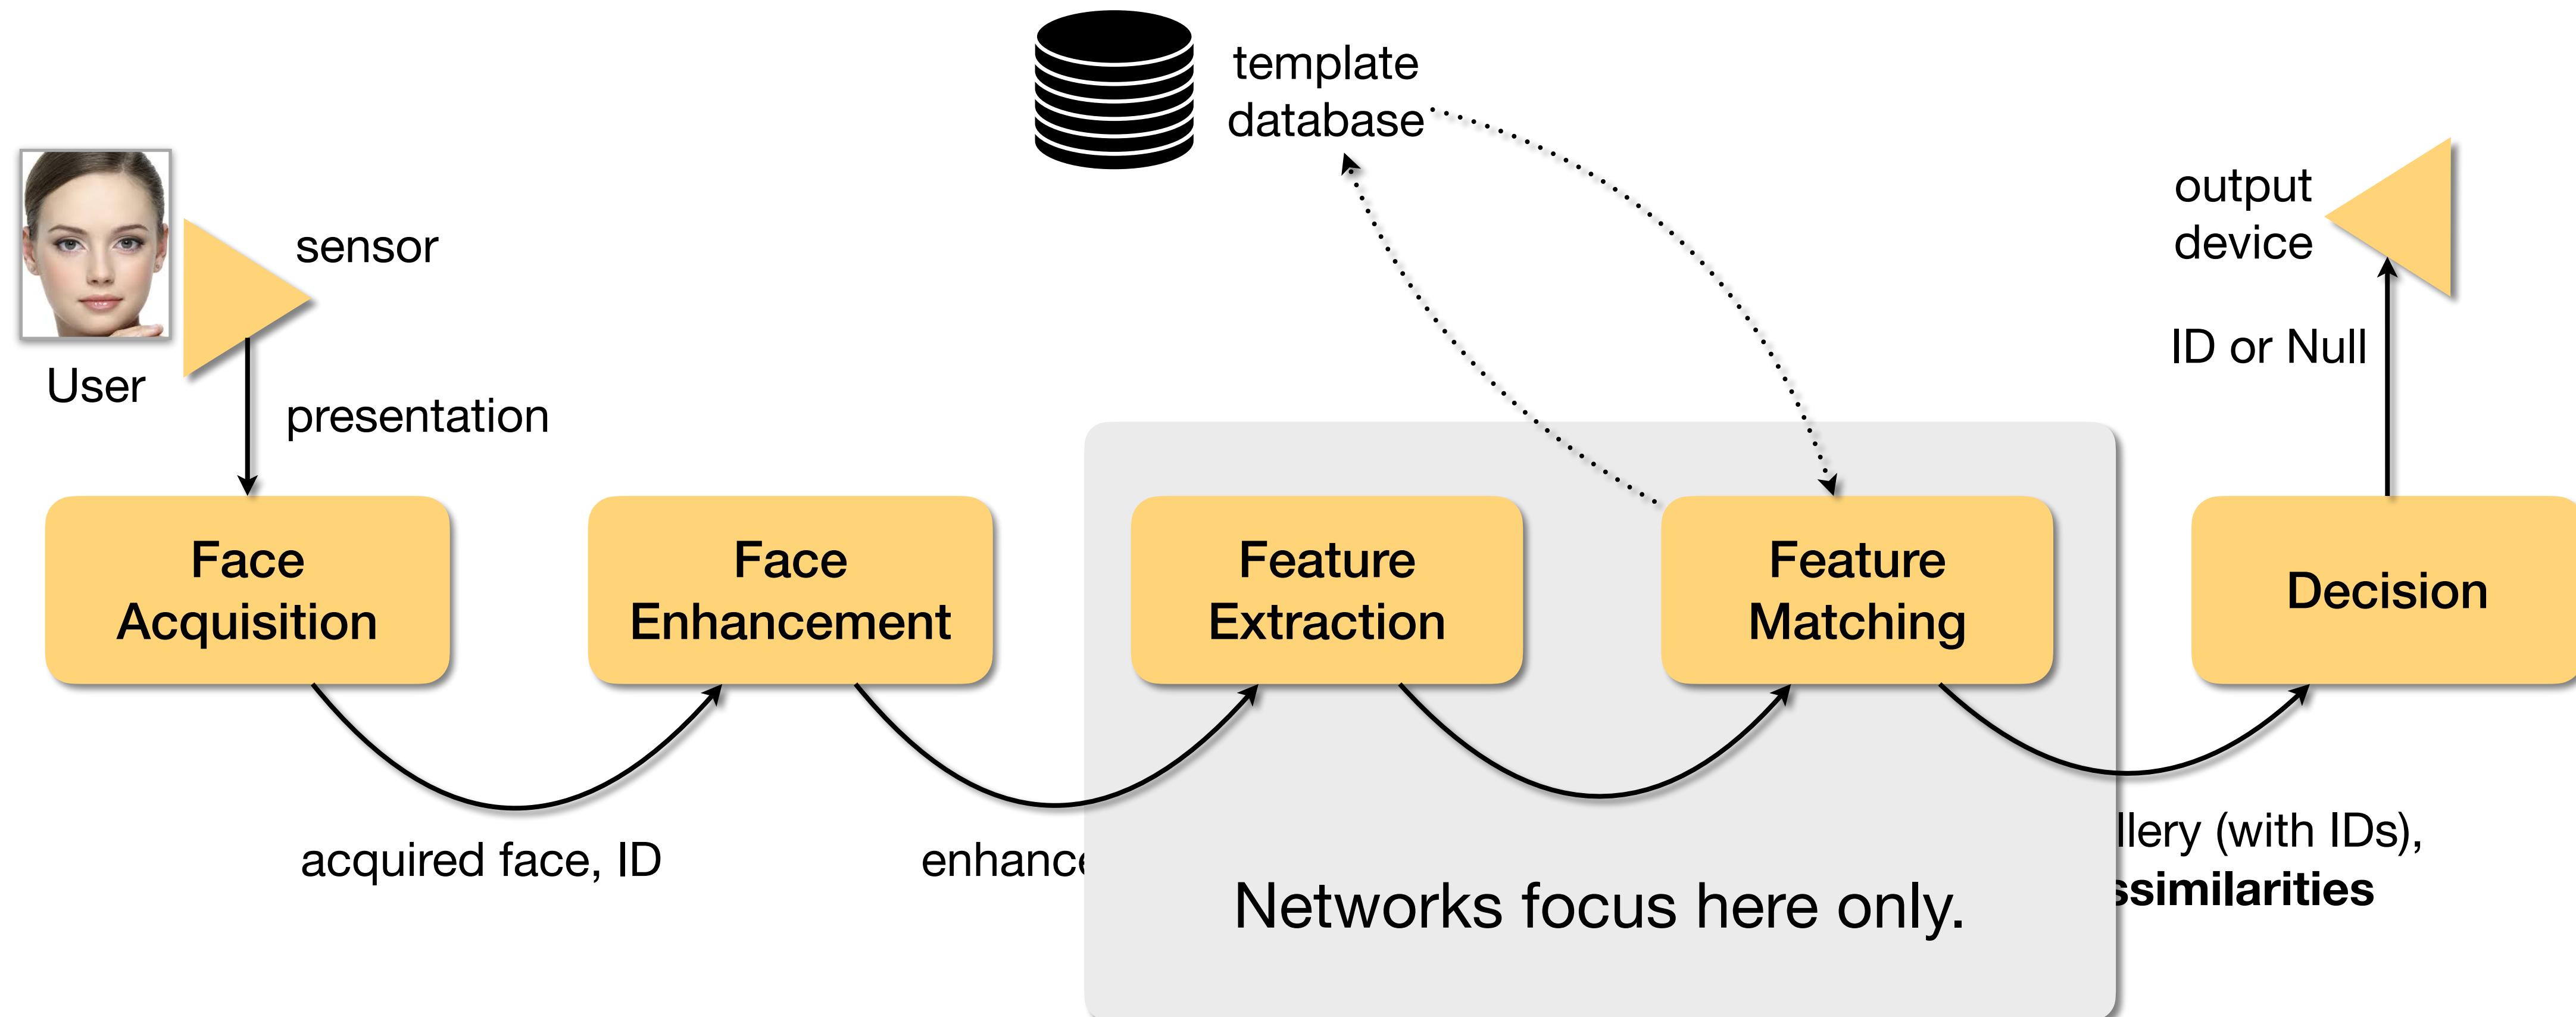
To speed up training, use **siamese networks** (same architecture, same weights).



Data-Driven Face Recognition



Data-Driven Face Recognition



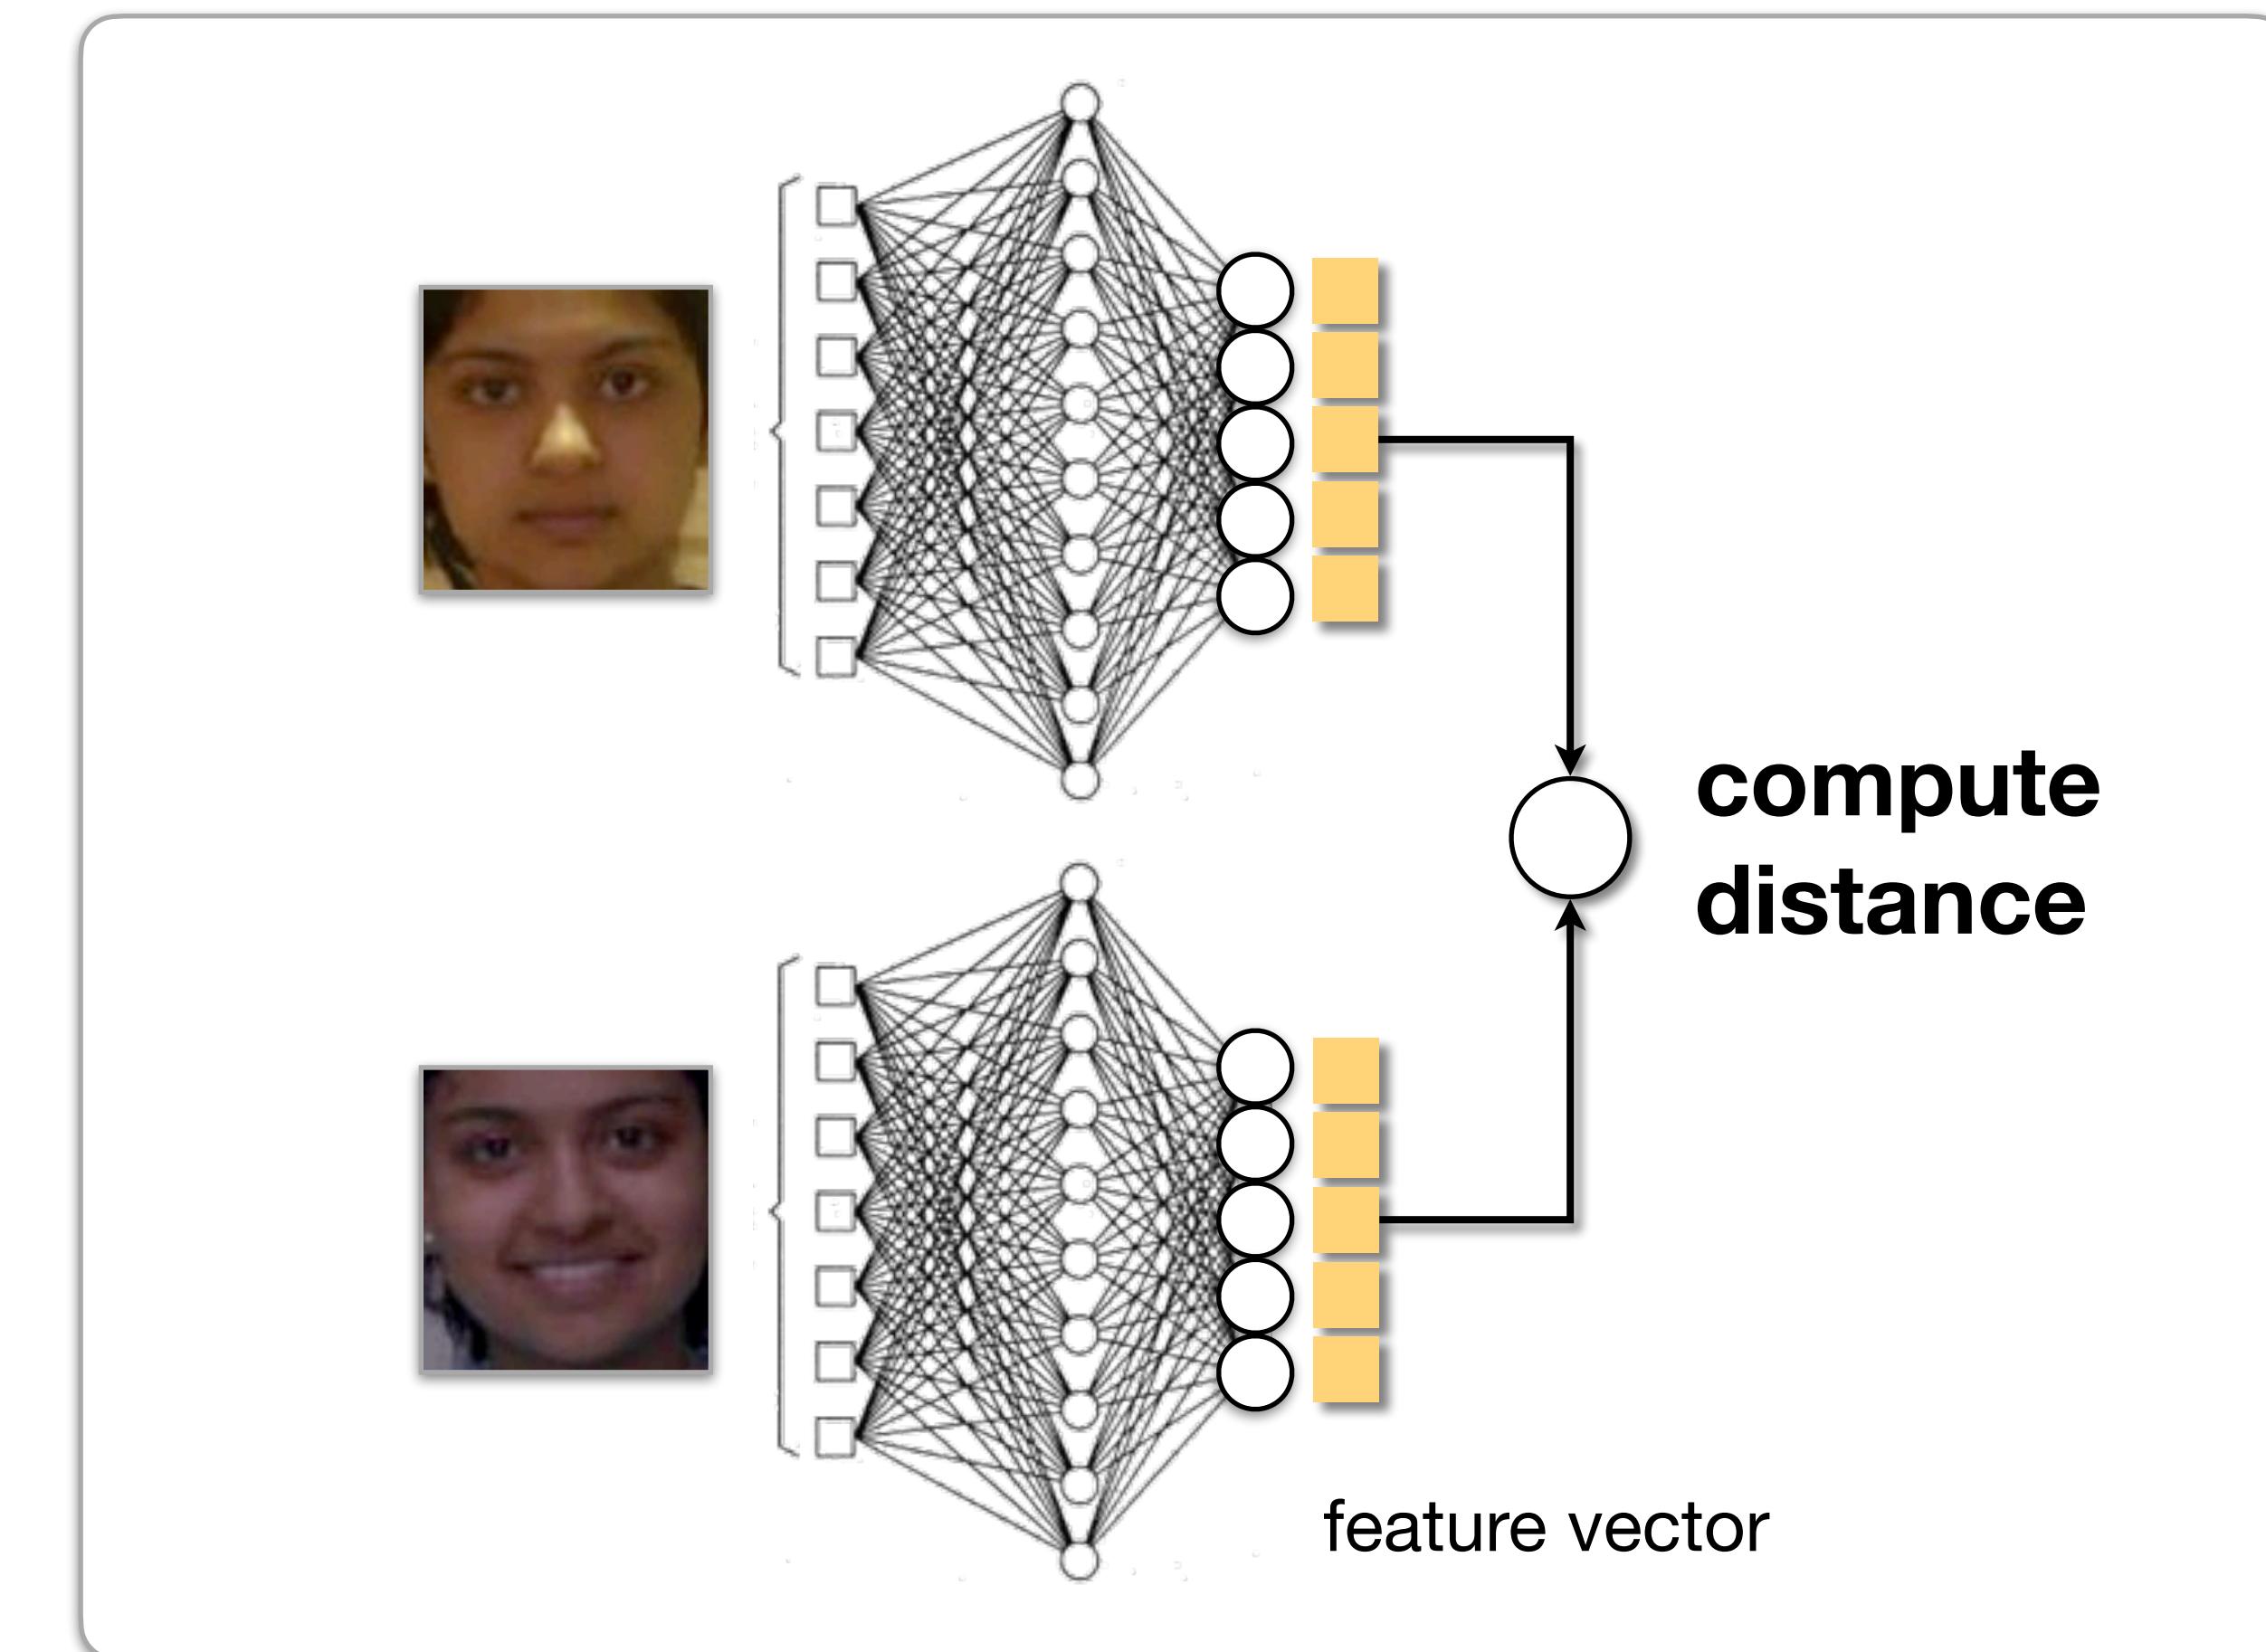
Data-Driven Face Recognition

Deep Learning

Training Approaches

Pairwise-loss-based


Triplet-loss-based



Data-Driven Face Recognition

Deep Learning

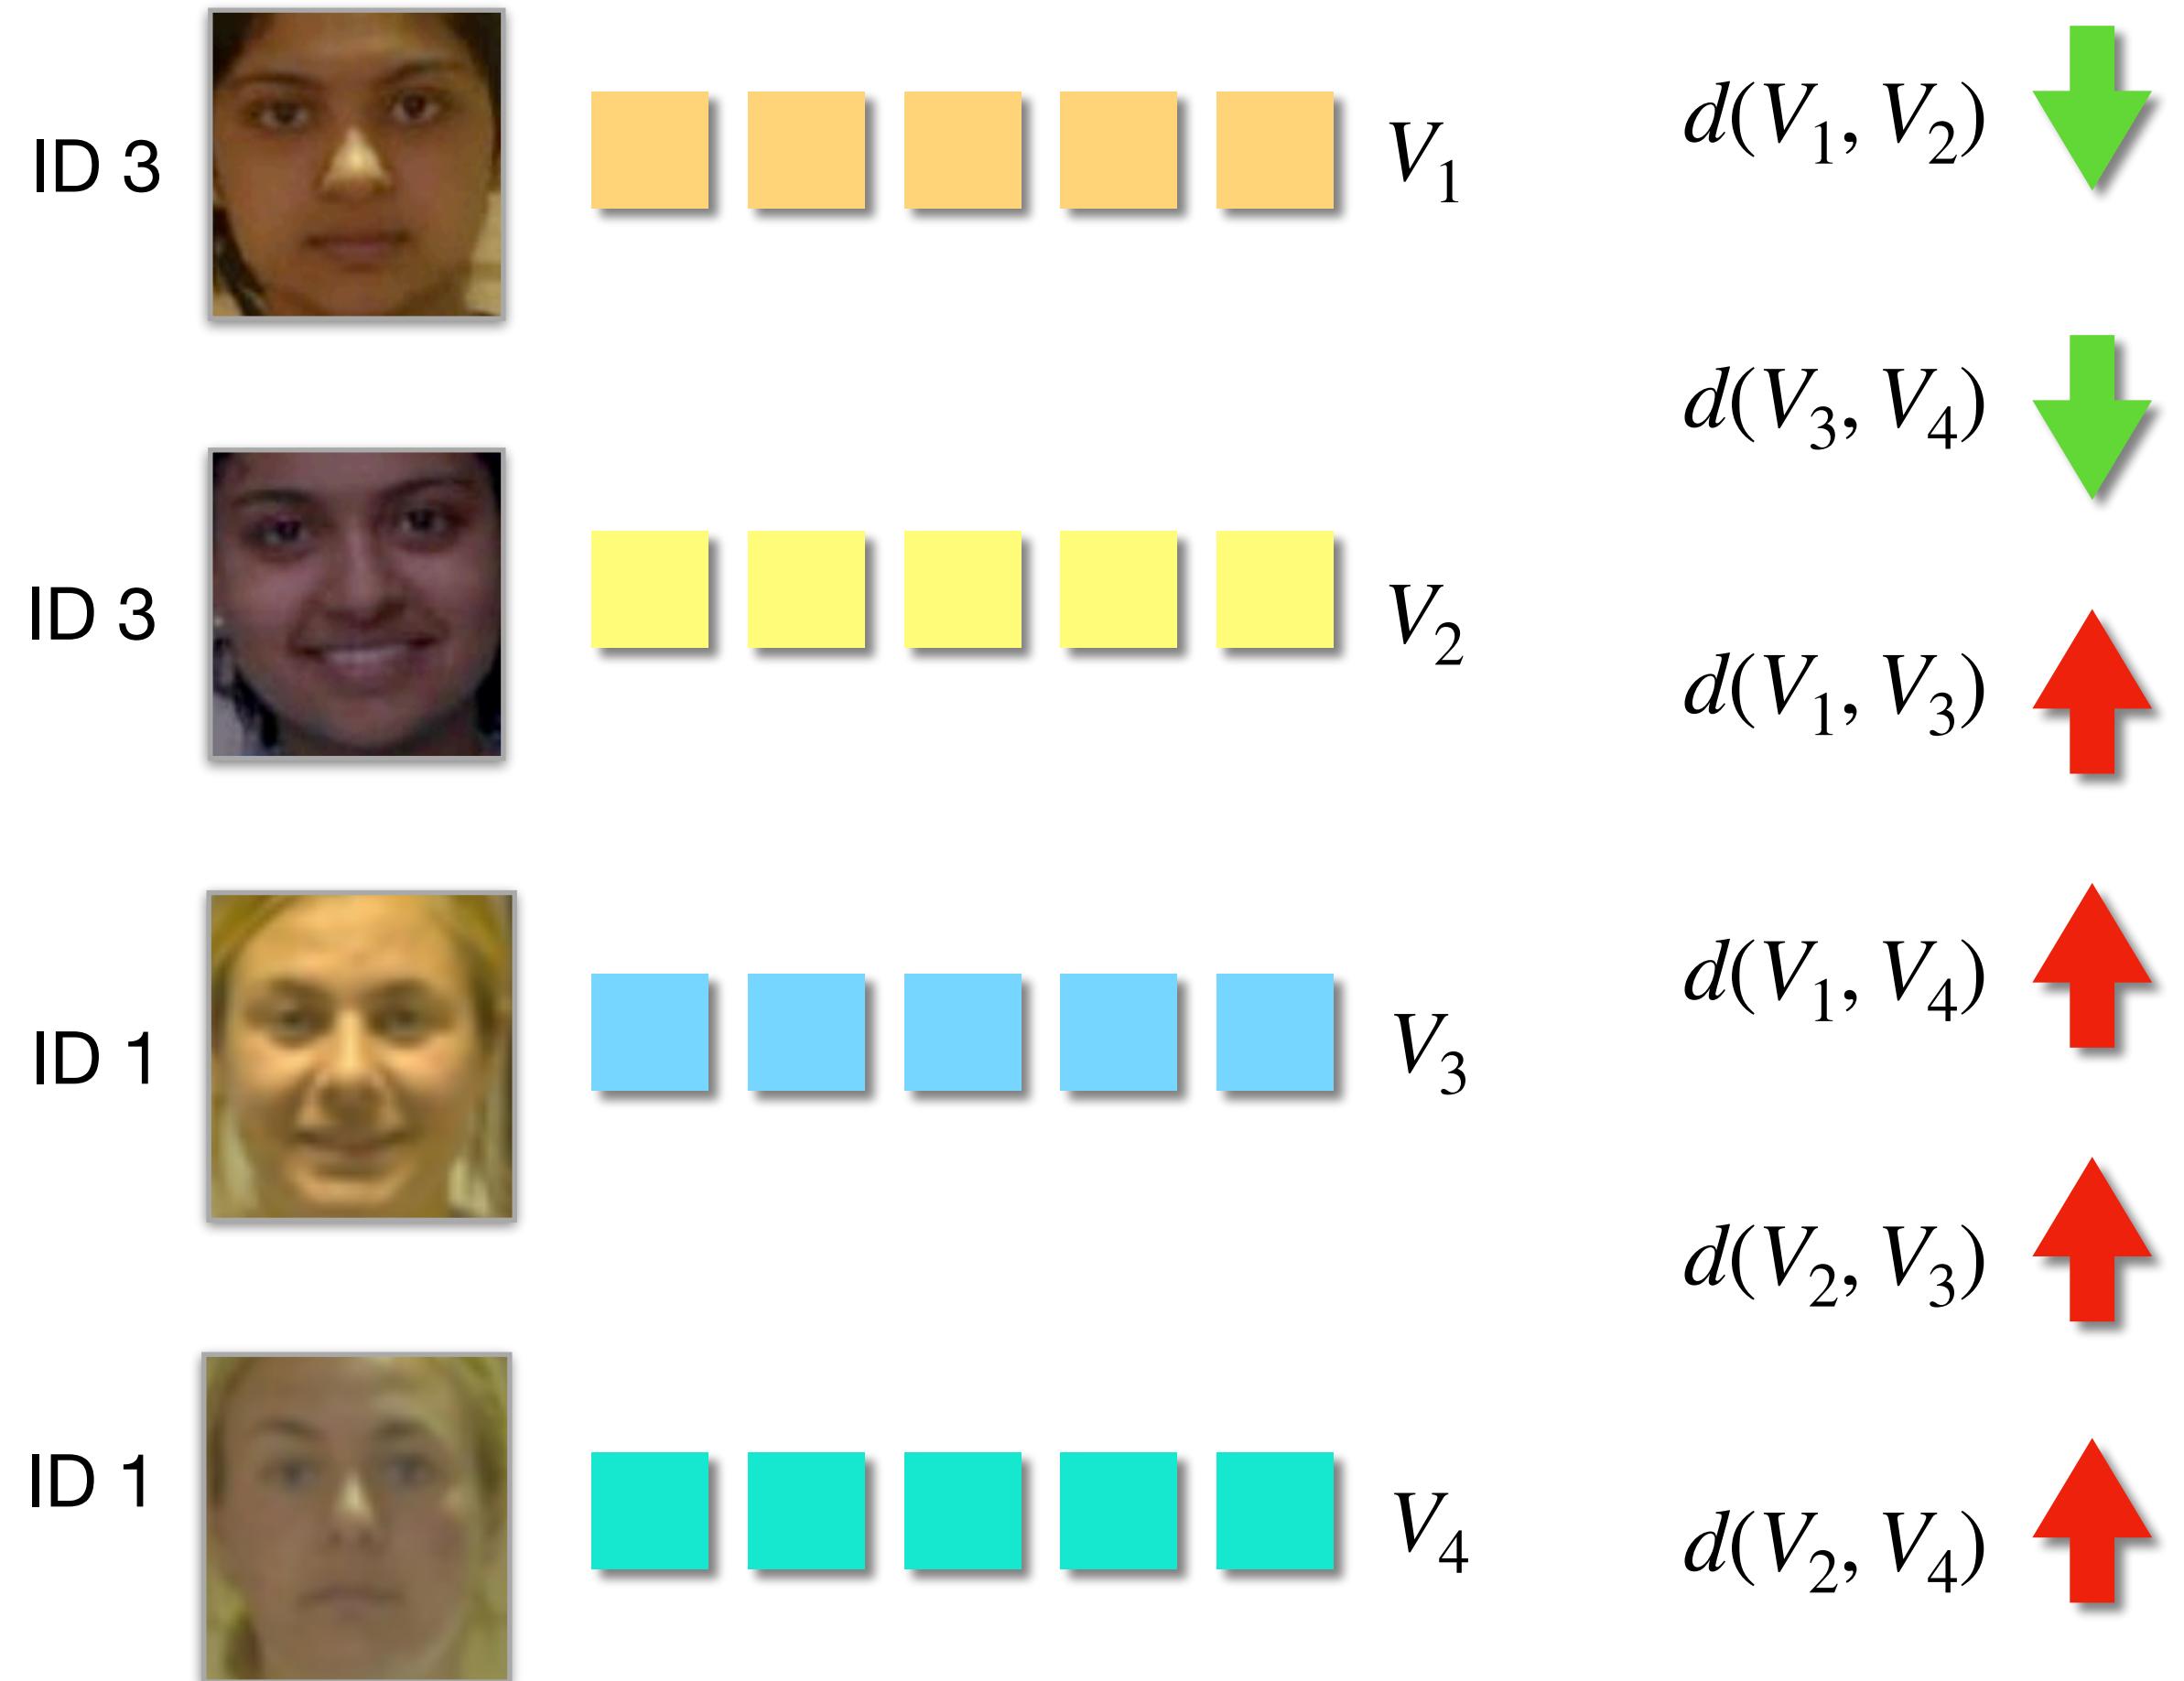
Training Approaches
Pairwise-loss-based
Triplet-loss-based



Pairwise Face Recognition

Pairwise Loss (PL)

Train the network in a way that feature vectors of the same class have small distance, while feature vectors from different classes have large distance.



Pairwise Face Recognition

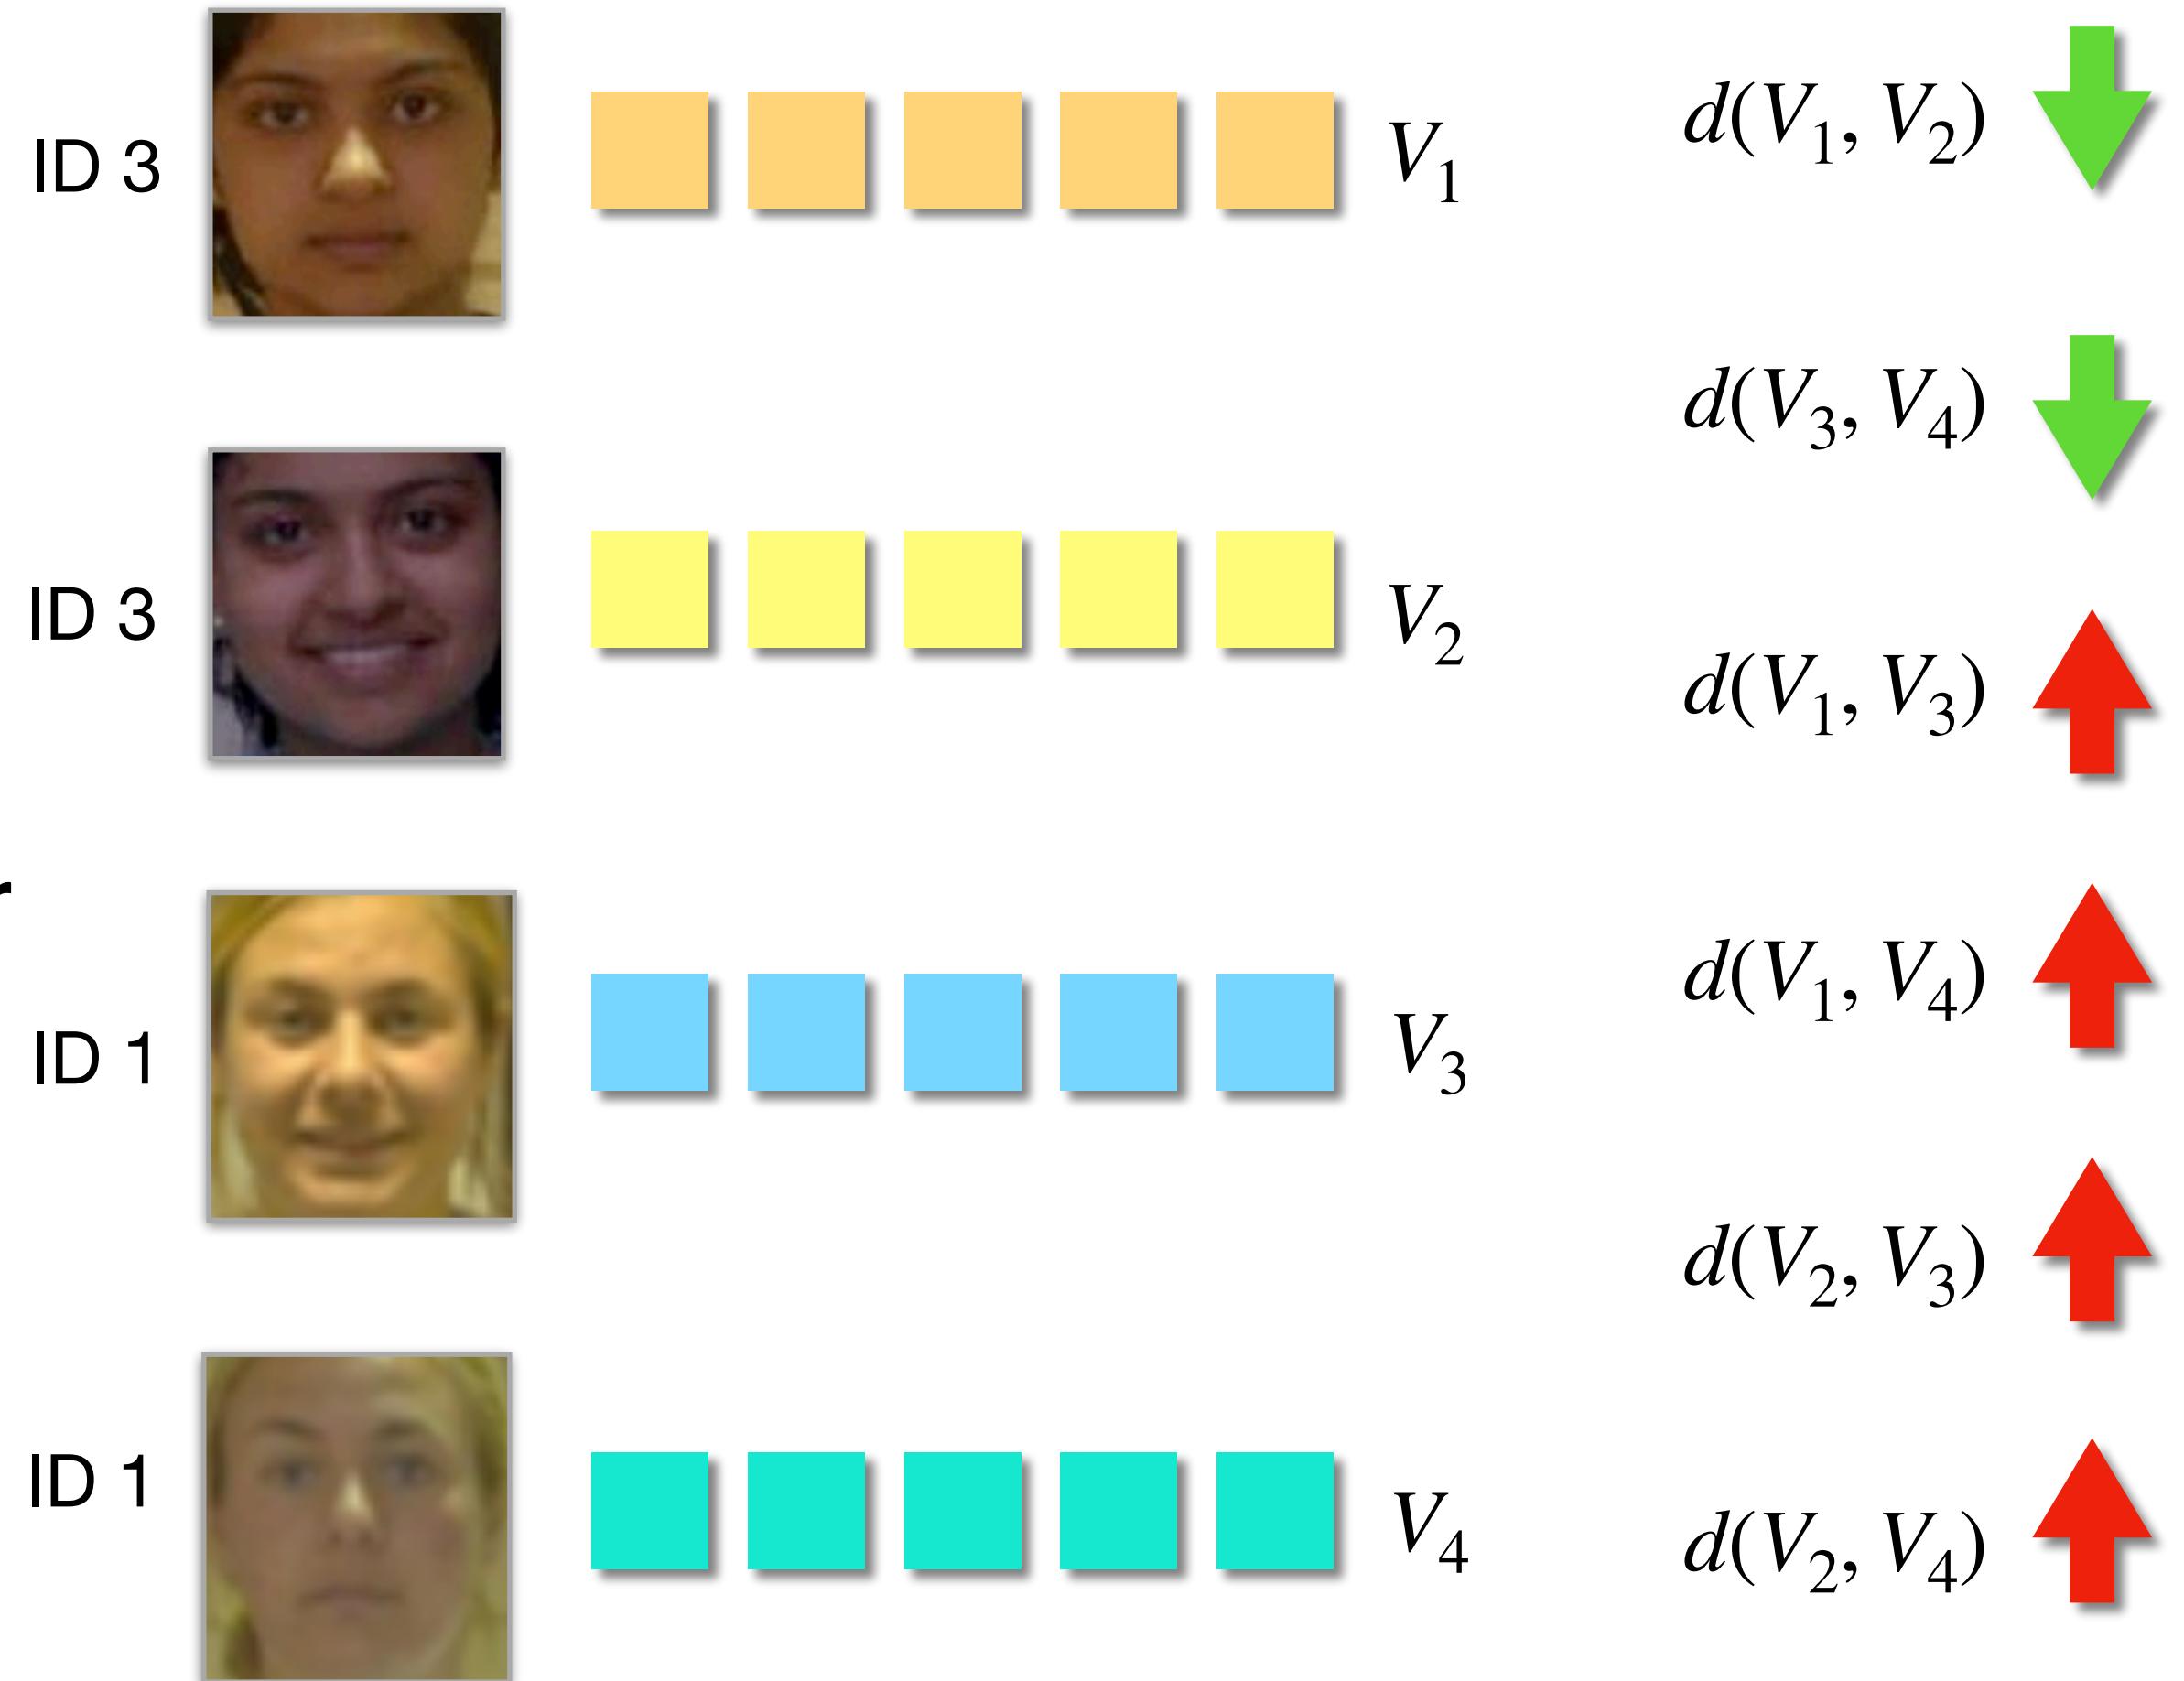
Pairwise Loss (PL)

the smaller, the better


$$PL = \begin{cases} d(V_x, V_y) & \text{if genuine pair} \\ \max(0, m - d(V_x, V_y)) & \text{if impostor pair} \end{cases}$$

enforced margin


it must be larger than m



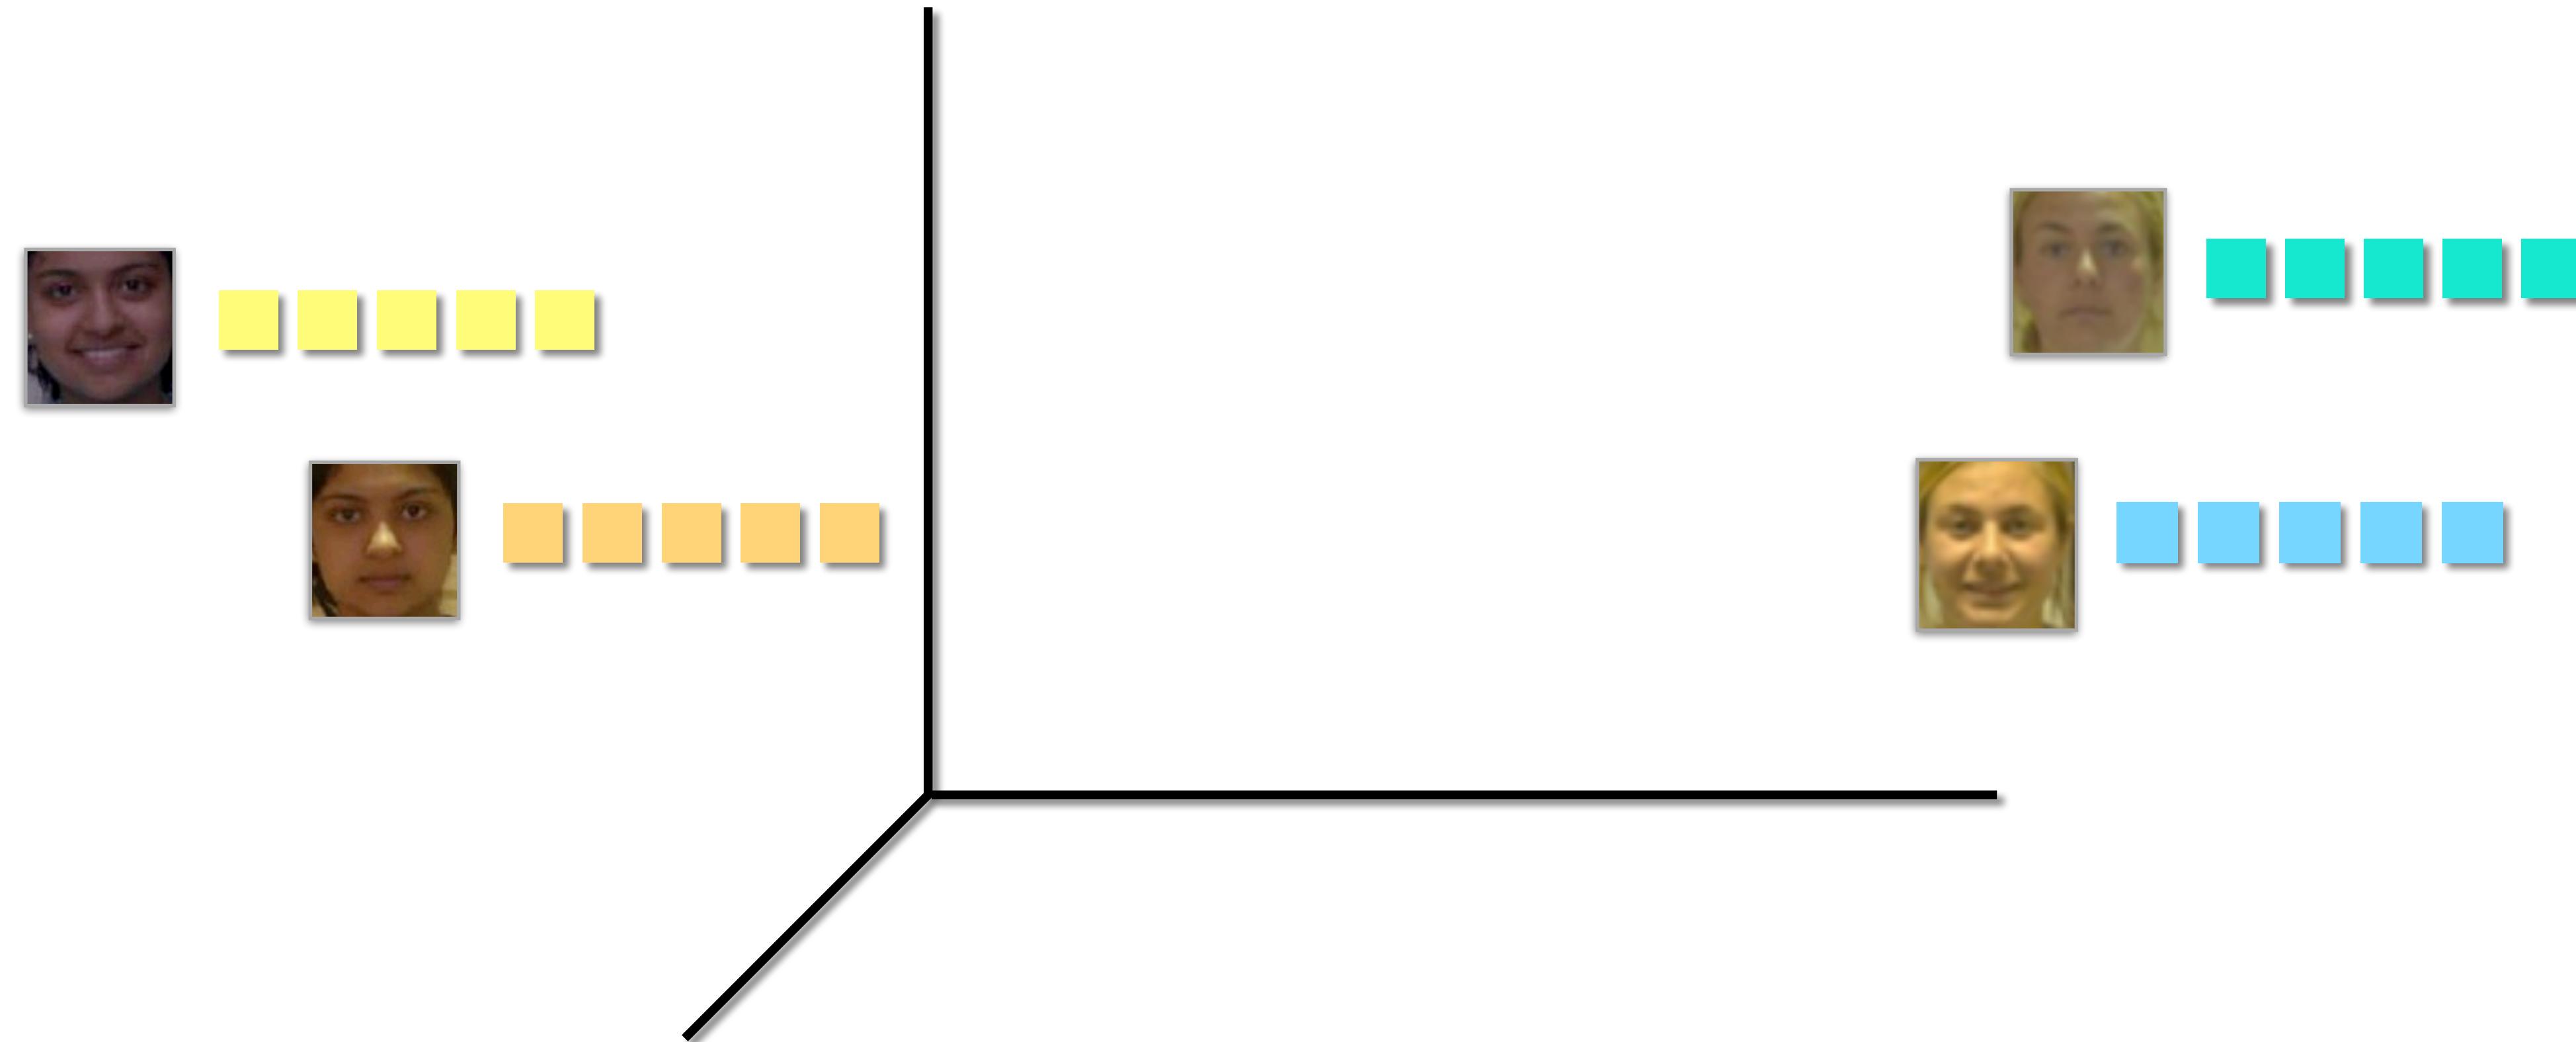
Pairwise Face Recognition



Pairwise Face Recognition



Pairwise Face Recognition



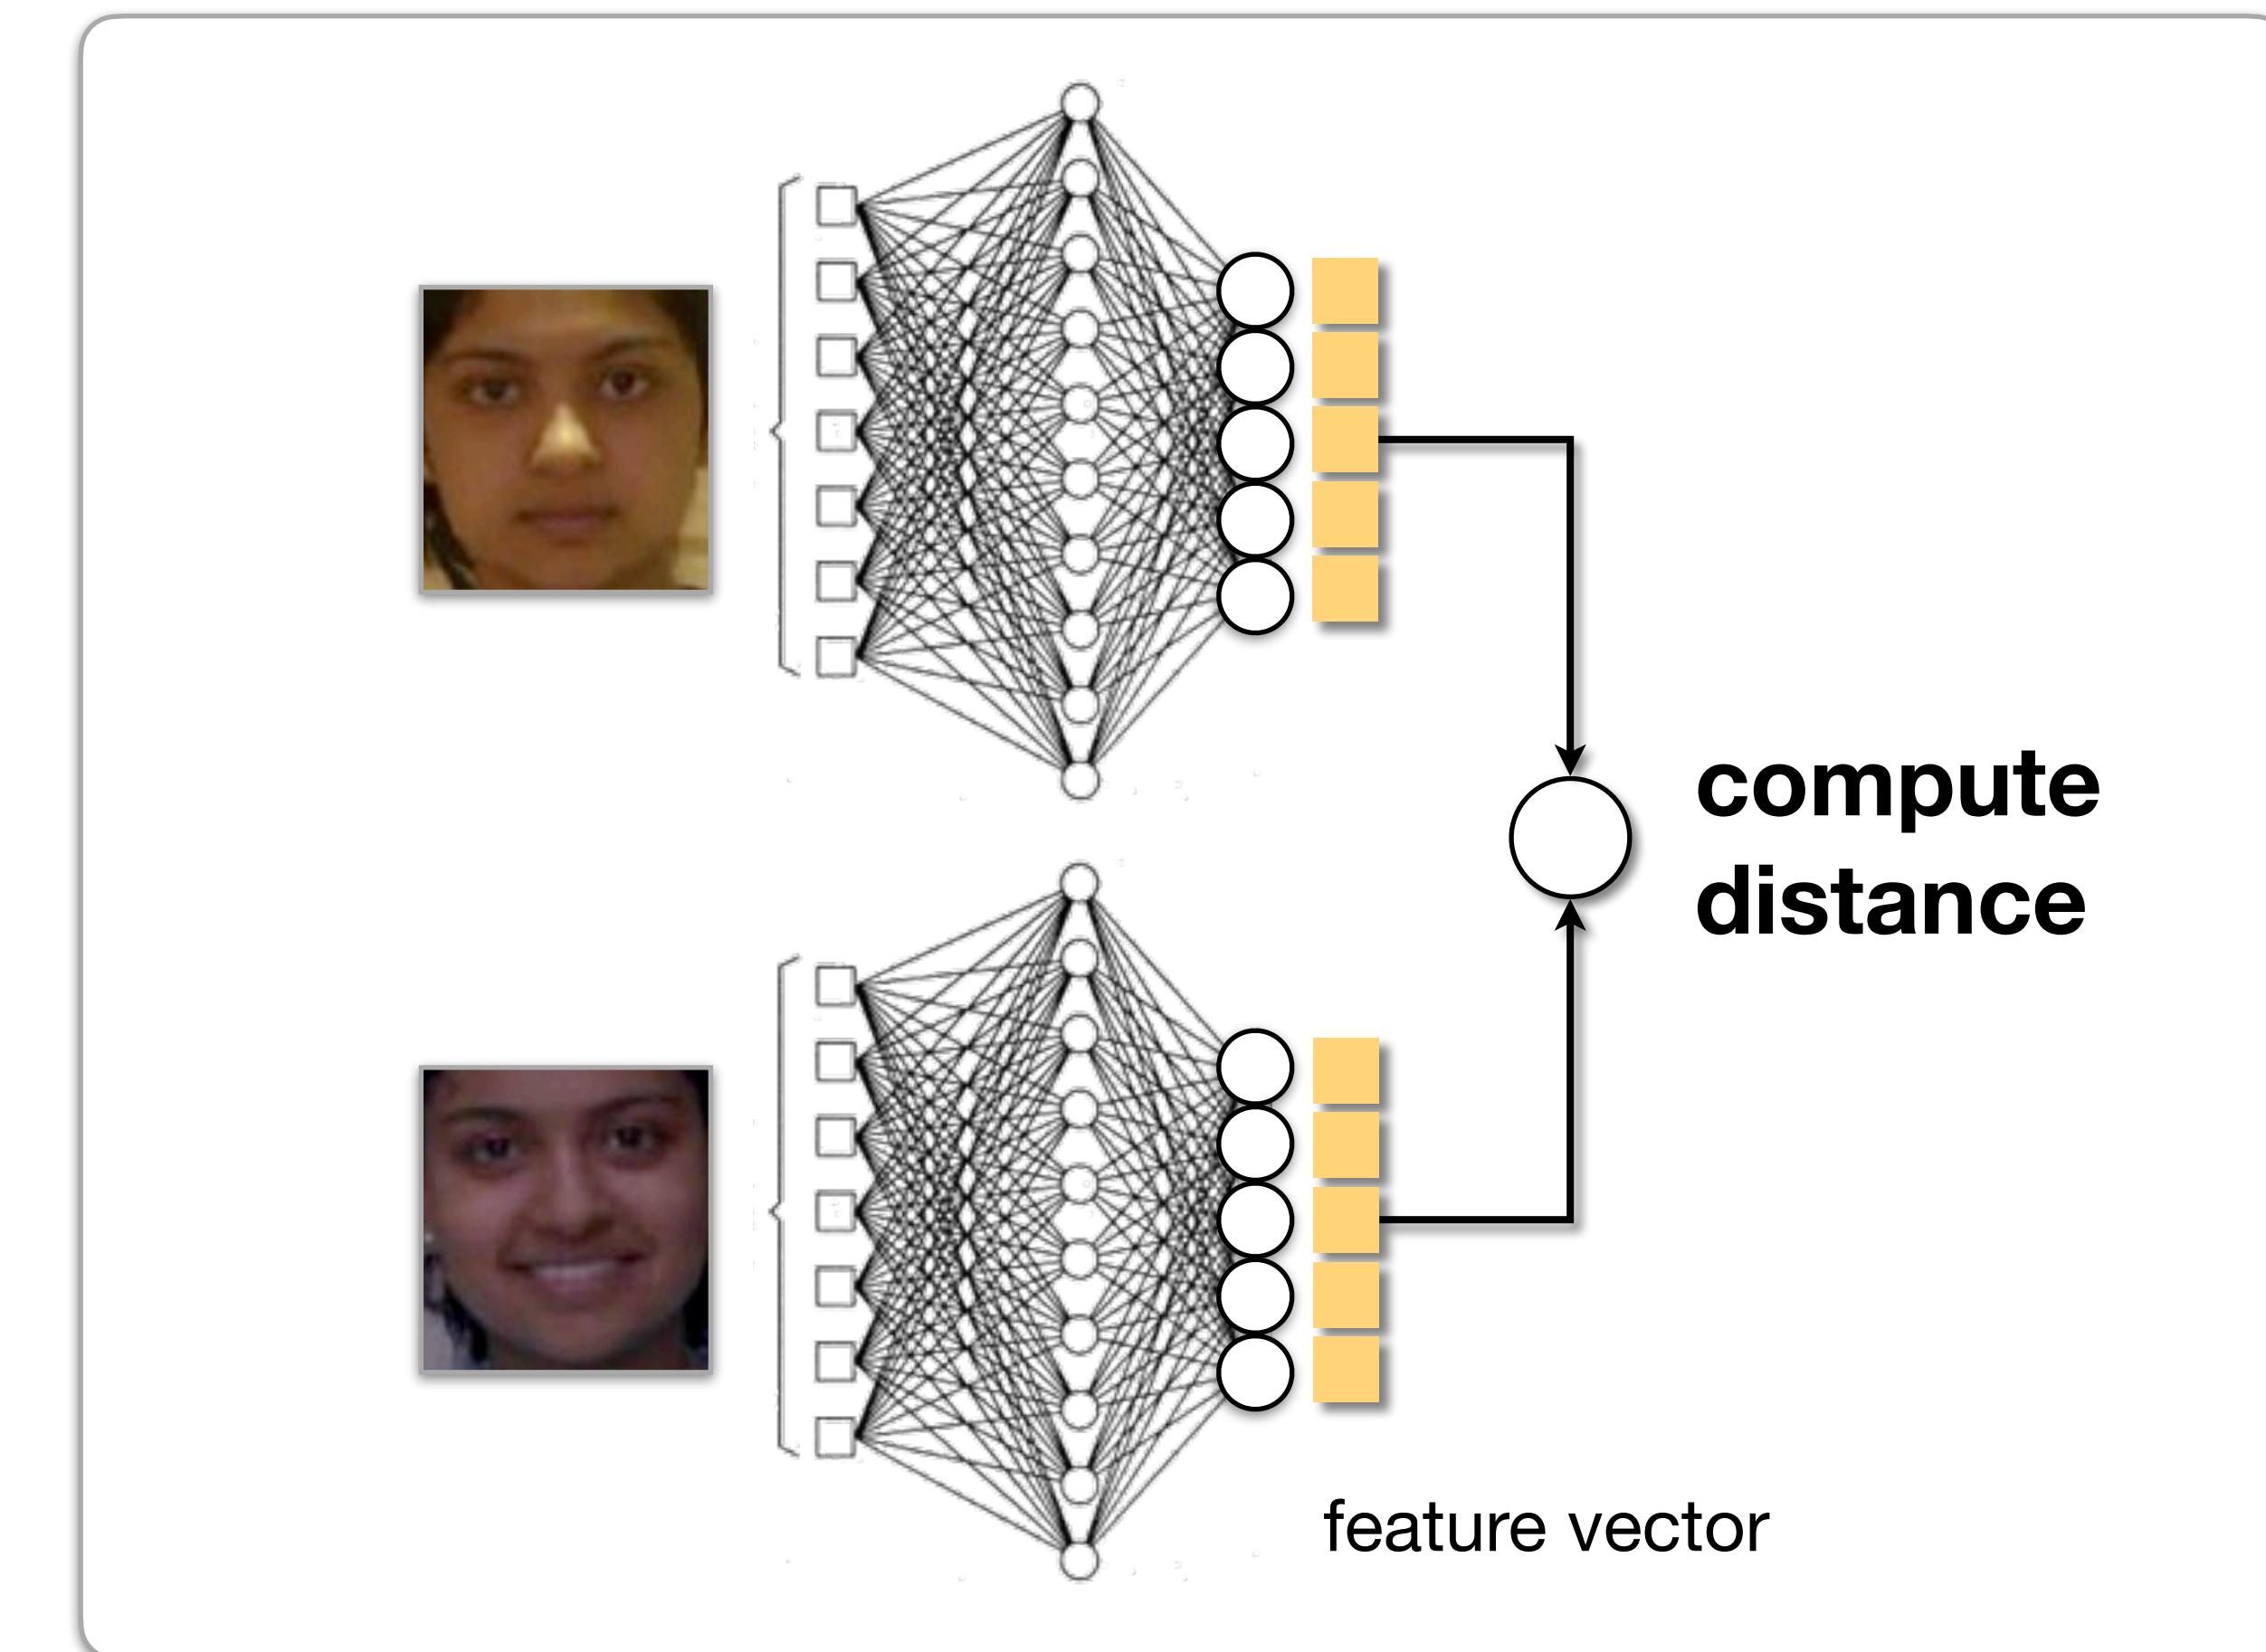
Data-Driven Face Recognition

Deep Learning

Training Approaches

Pairwise-loss-based


Triplet-loss-based



Triplet Face Recognition

Triplet Loss (TL)

Choose a reference data sample (the **anchor**) and a **positive** and a **negative** data samples to optimize their distances.



Triplet Face Recognition

Triplet Loss (TL)


Choose a reference data sample (the **anchor**) and a **positive** and a **negative** data samples to optimize their distances.

Minimize $d(V_1, V_2)$ and maximize $d(V_1, V_3)$.

Schroff et al.

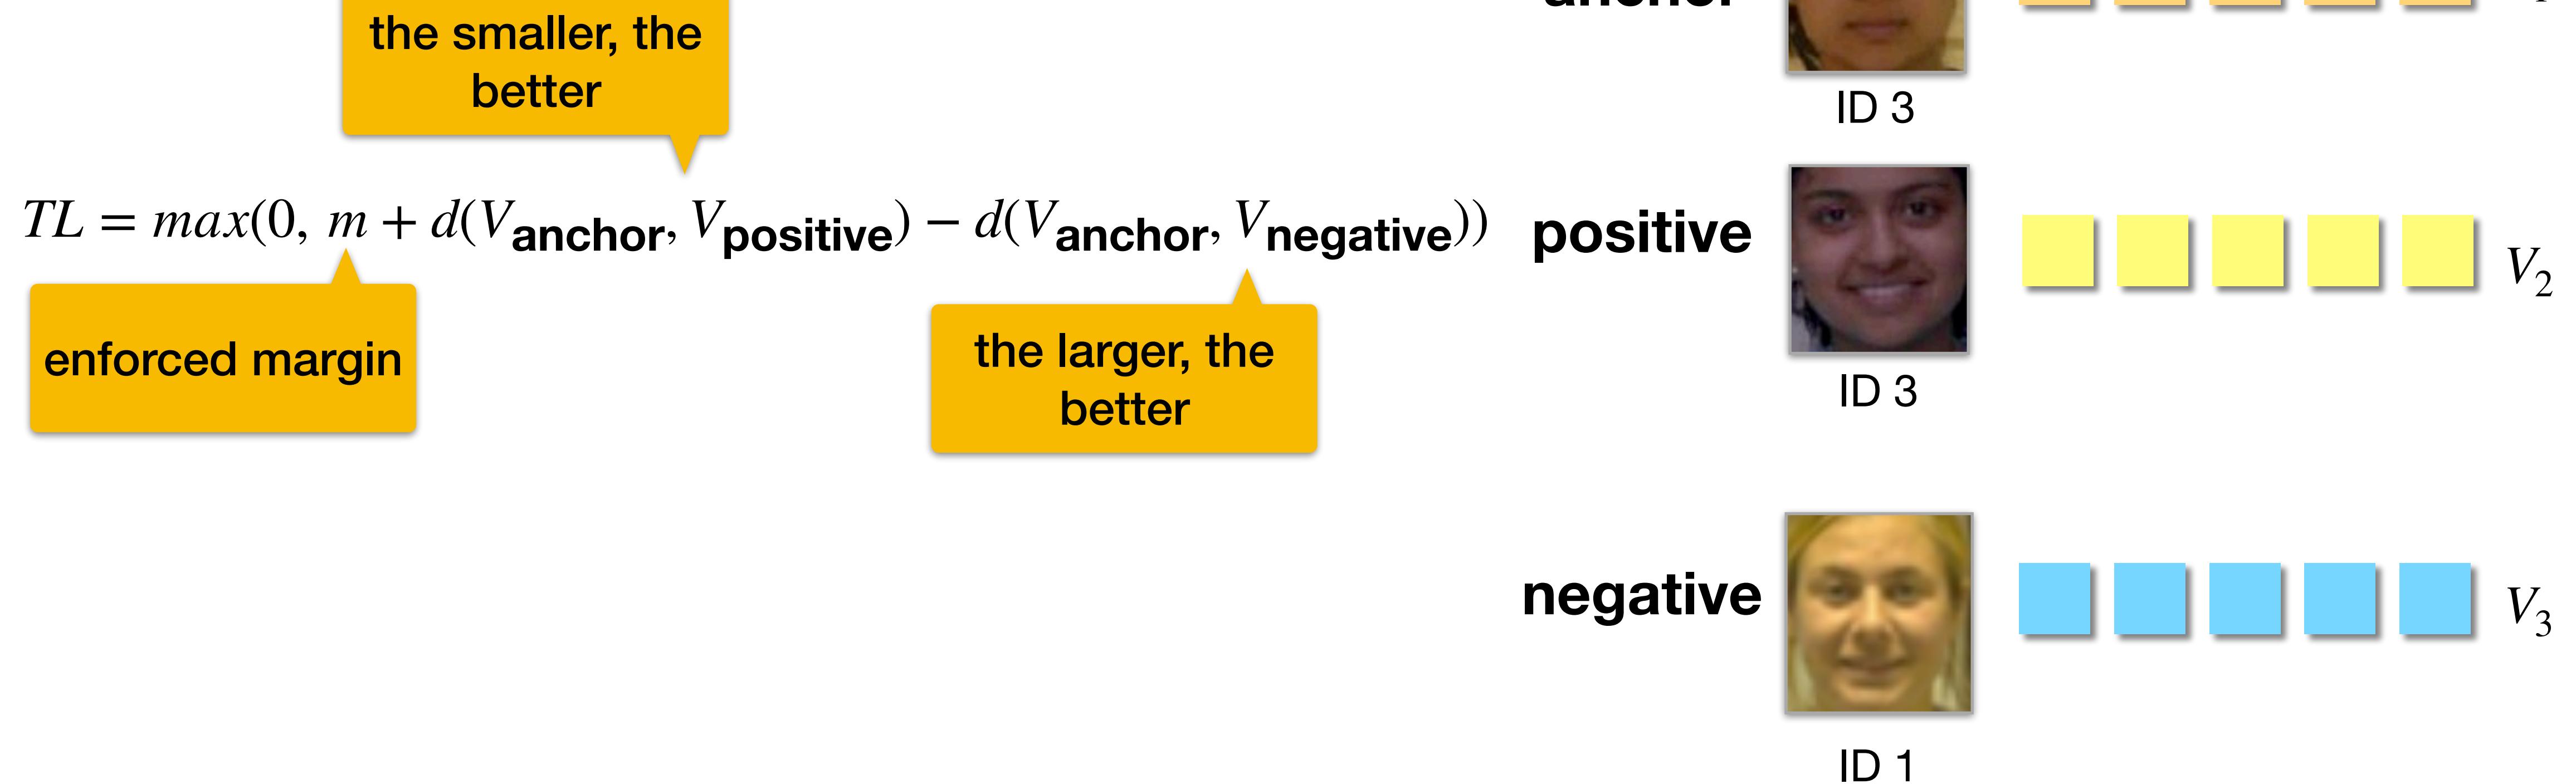
Facenet: A unified embedding for face recognition and clustering.

CVPR 2015



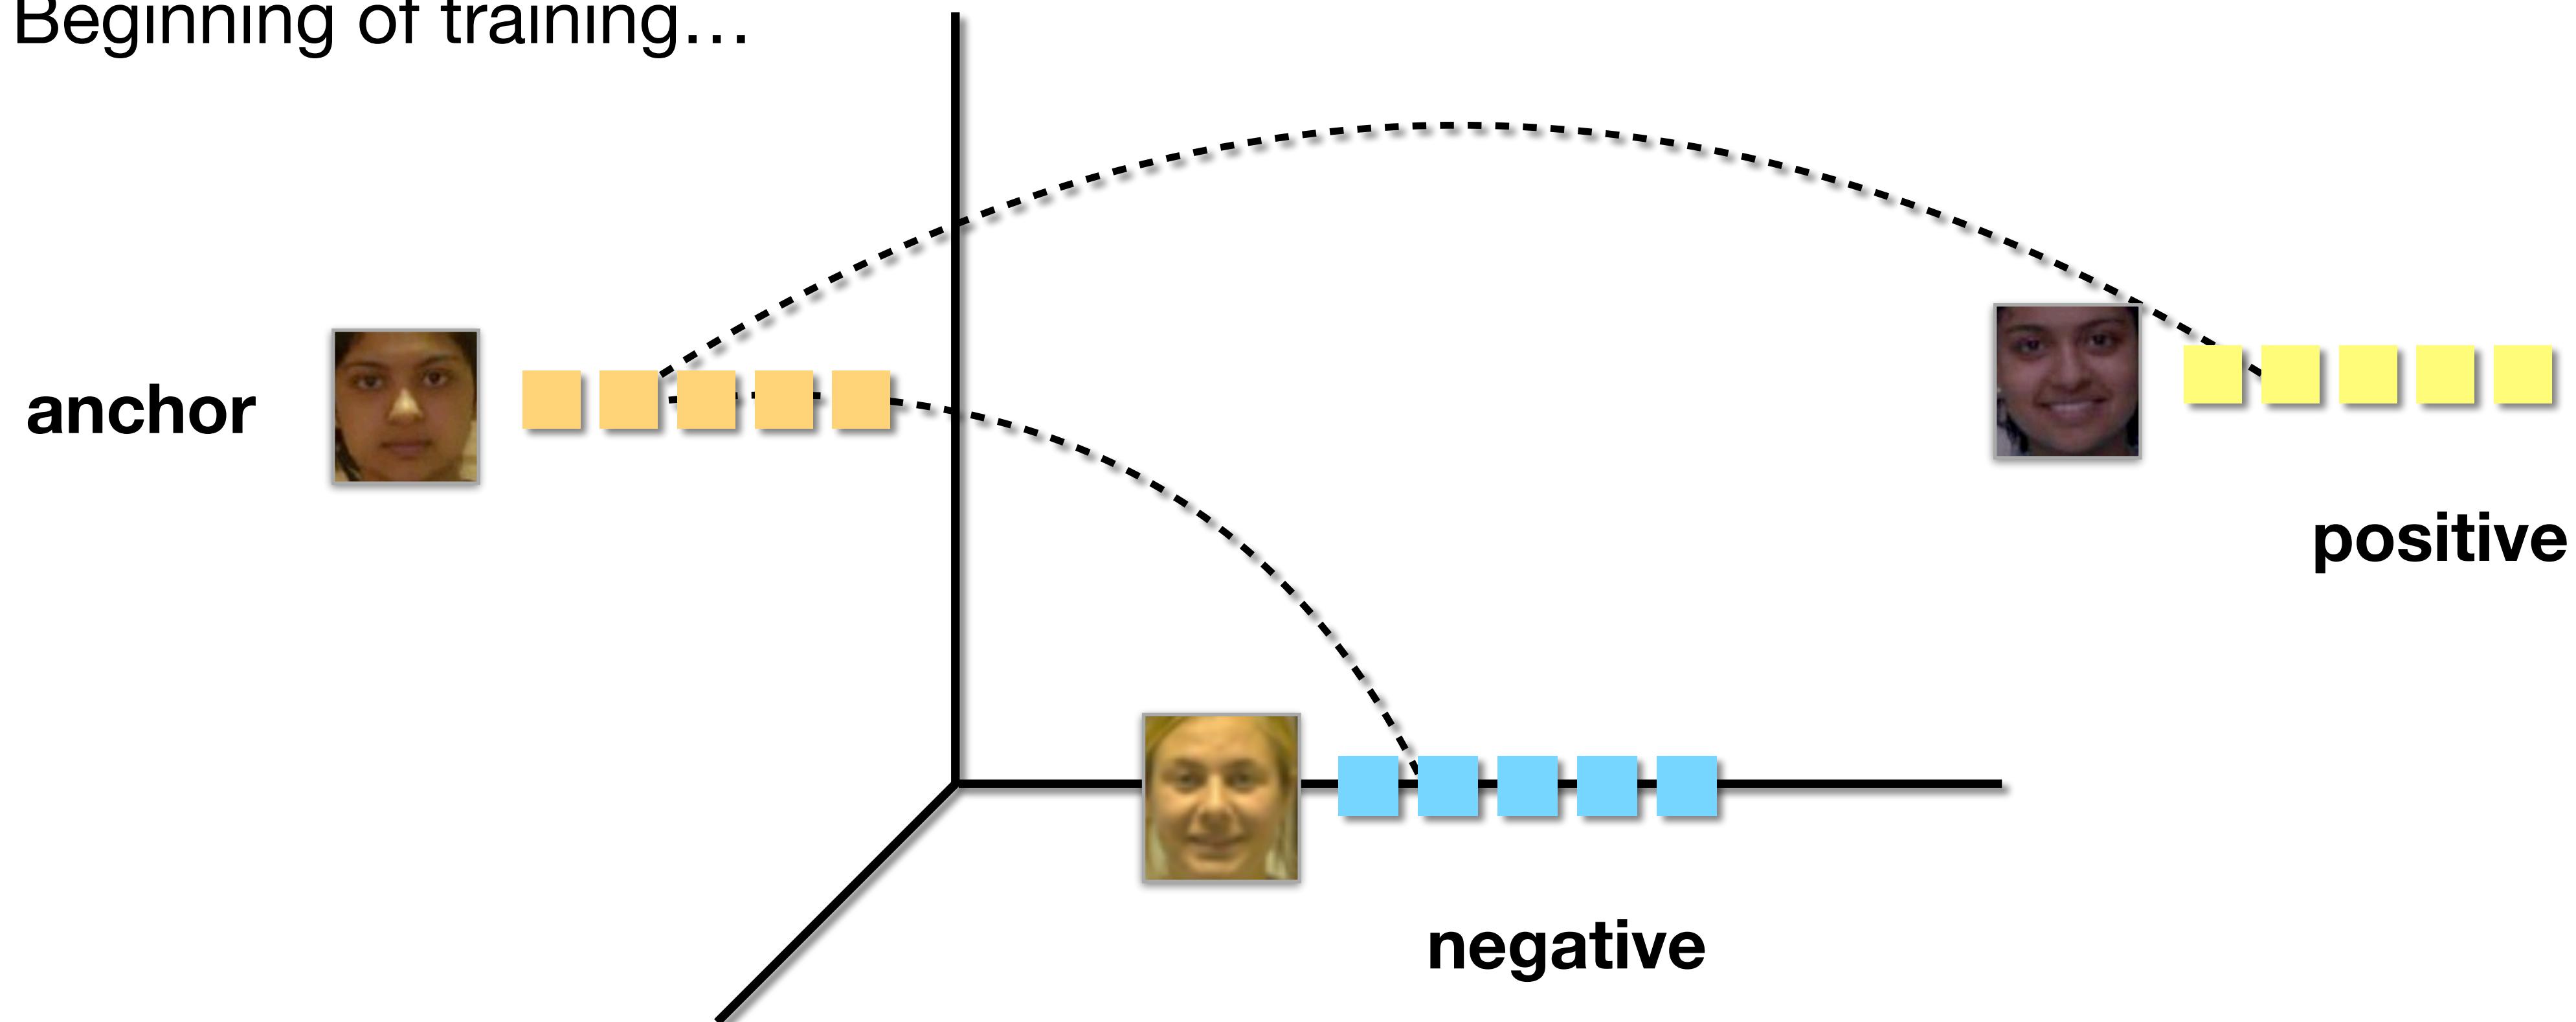
Triplet Face Recognition

Triplet Loss (TL)

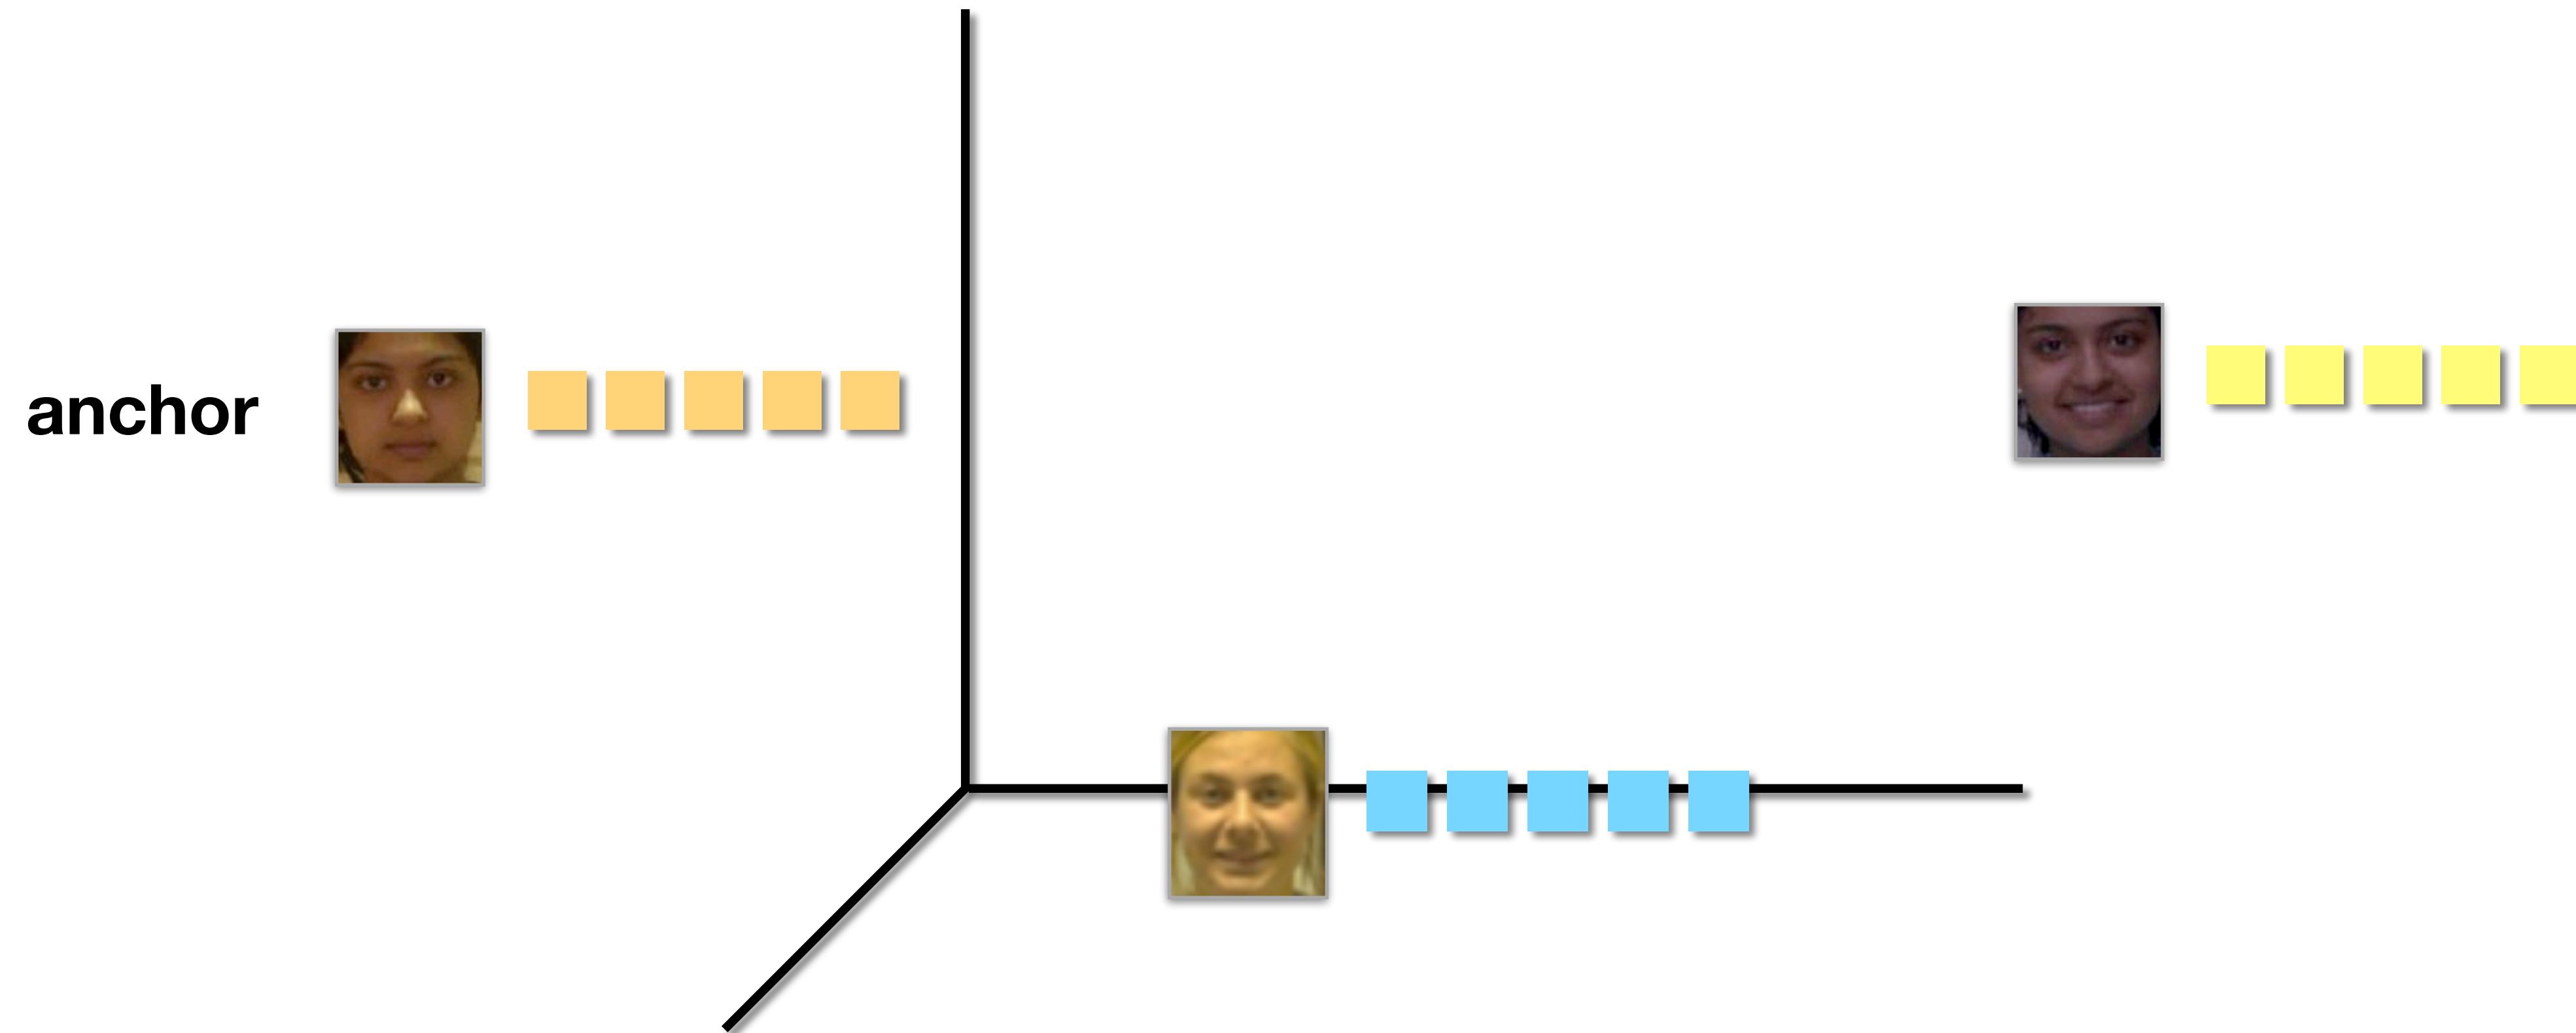


Triplet Face Recognition

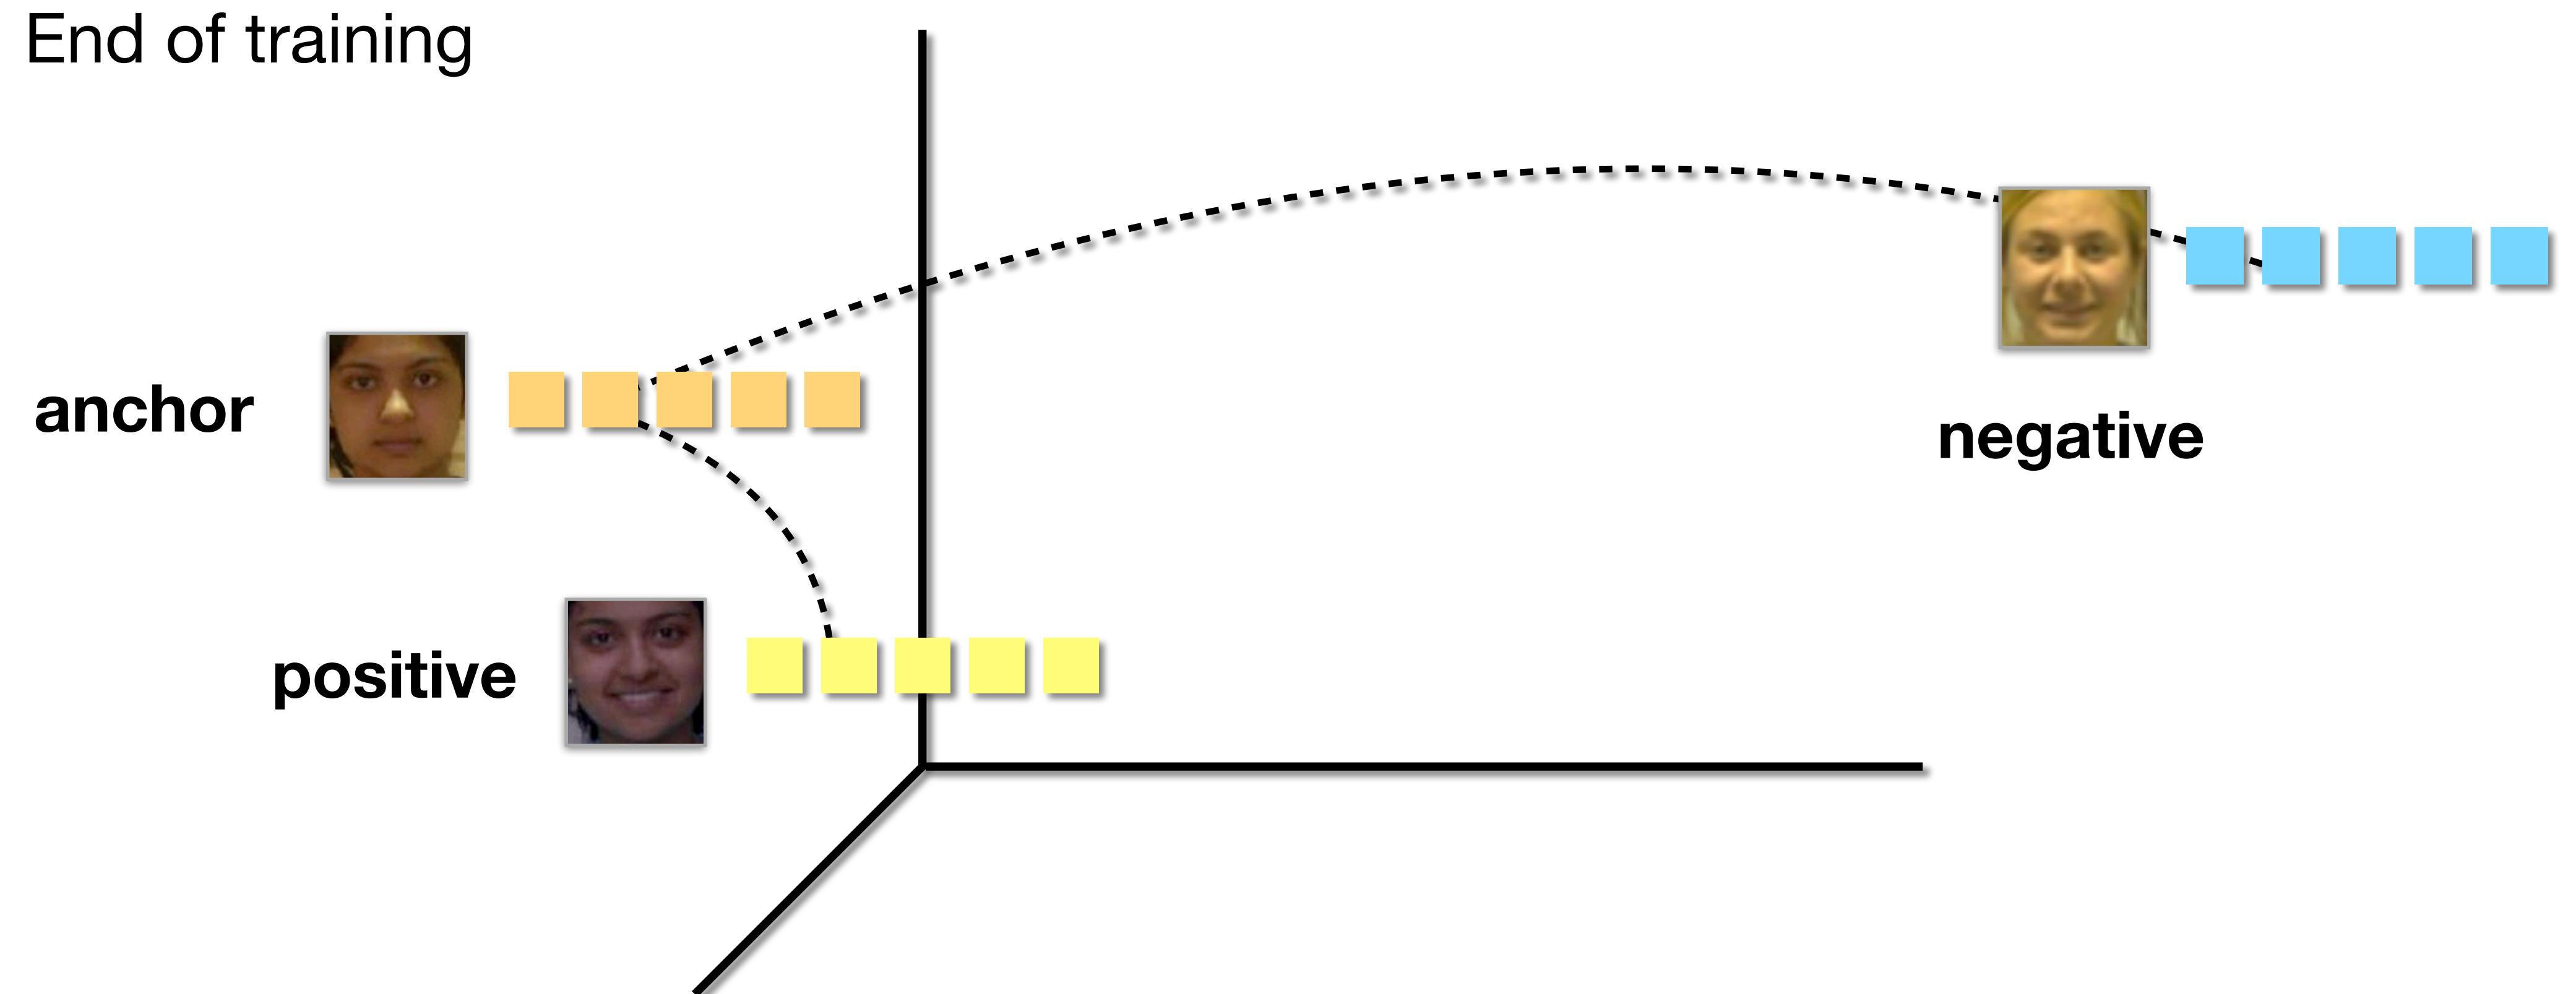
Beginning of training...



Triplet Face Recognition

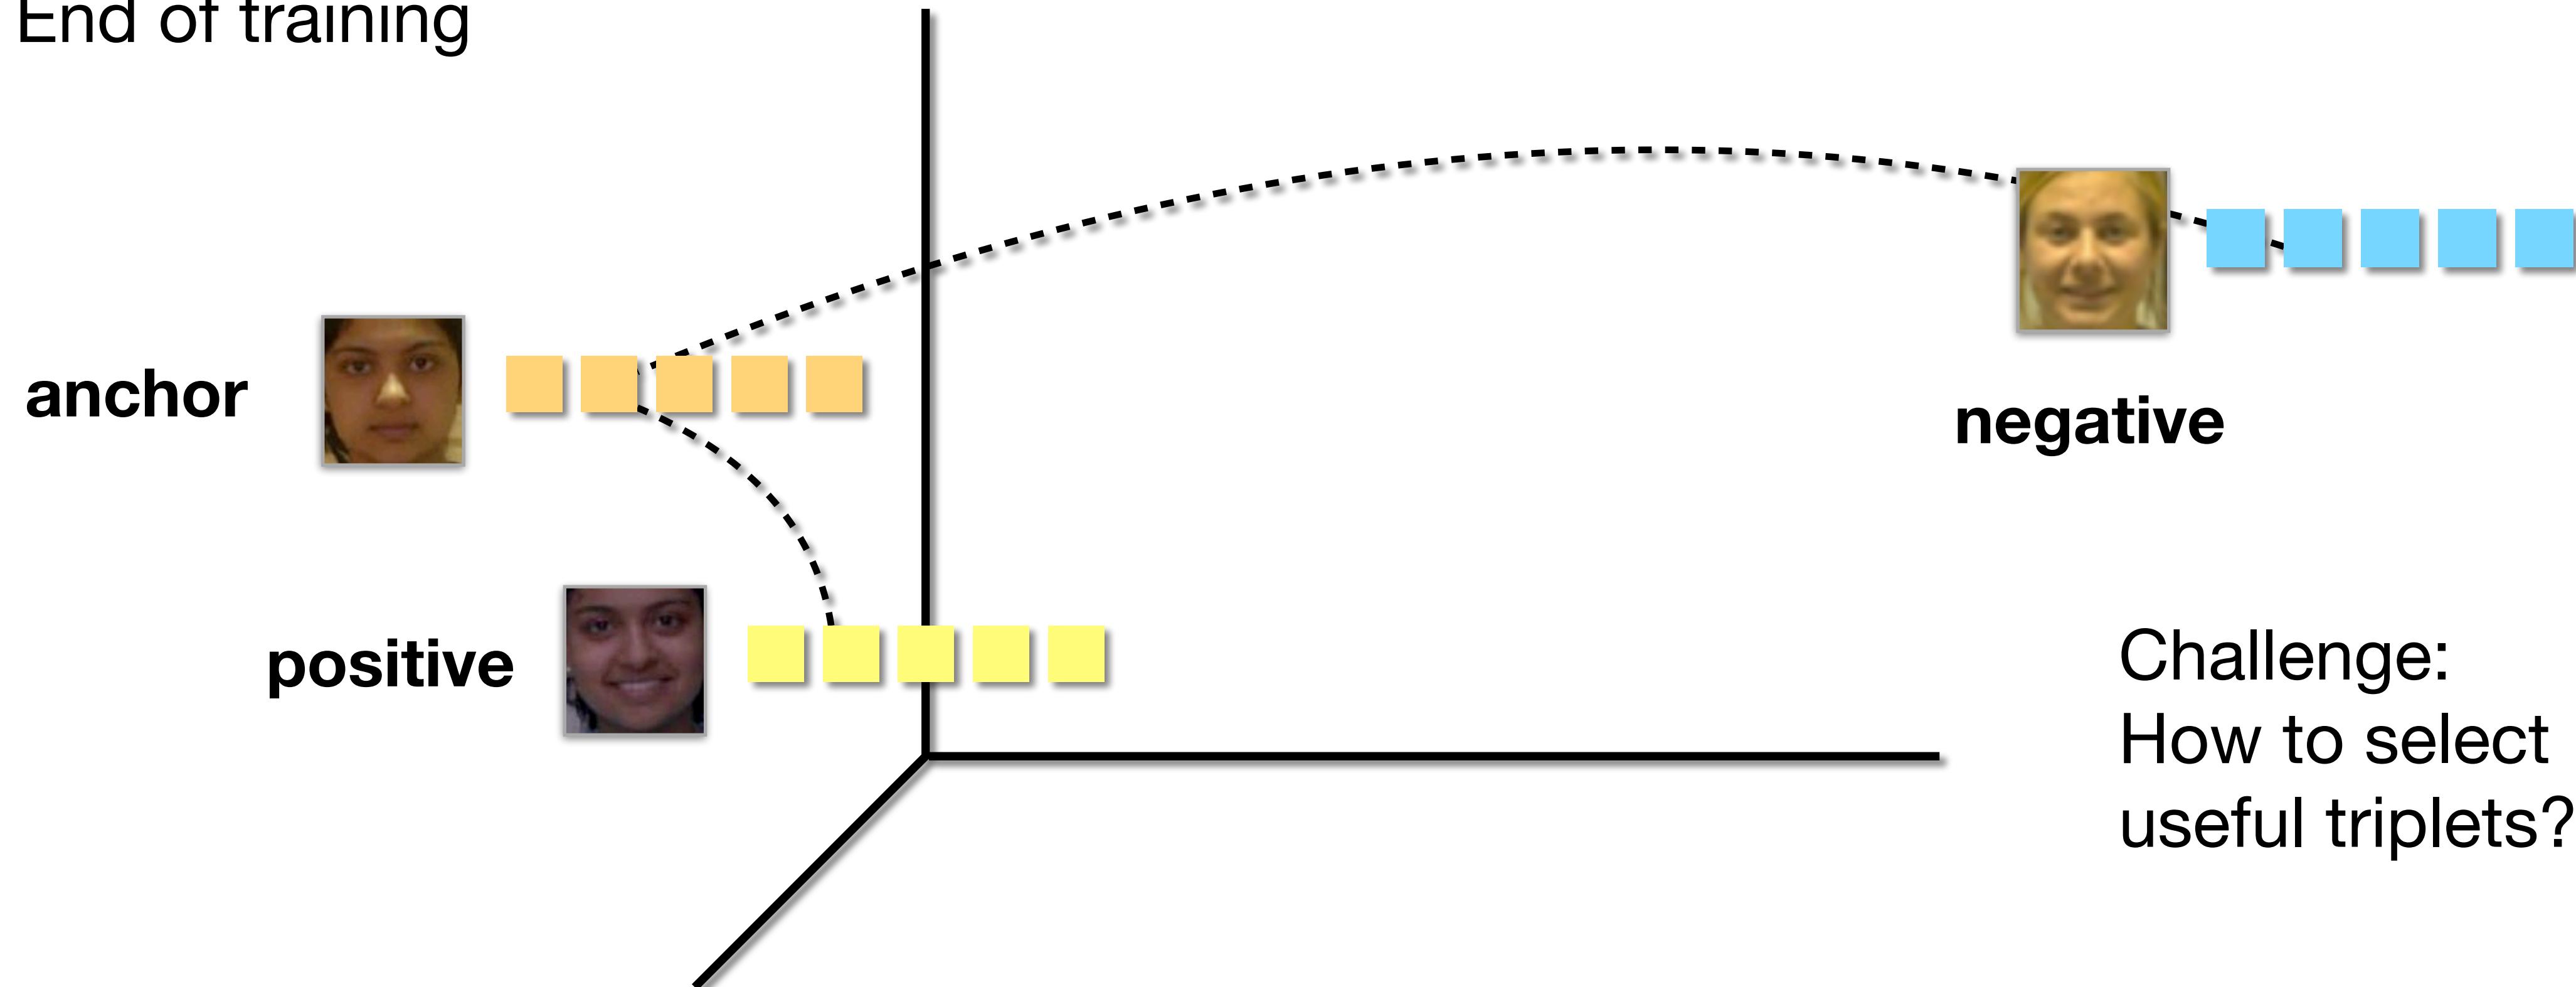


Triplet Face Recognition



Triplet Face Recognition

End of training




Improvements

Source: Deng et al.
*Additive Angular Margin Loss
for Deep Face Recognition.*
CVPR 2019

Centre Loss

Use class clusters' centers to improve the convergence of the learning process.


Liu et al.
Sphereface: Deep hypersphere embedding for face recognition.
CVPR 2017



Improvements

SphereFace

Transform feature space into hypersphere and compute the distances as the **angles** between the feature vectors.




Euclidean
space

Liu et al.
Sphereface: Deep hypersphere embedding for face recognition.
CVPR 2017

Improvements

SphereFace

Transform feature space into hypersphere and compute the distances as the **angles** between the feature vectors.




Liu et al.

Sphereface: Deep hypersphere embedding for face recognition.
CVPR 2017

Improvements

SphereFace

Transform feature space into hypersphere and compute the distances as the **angles** between the feature vectors.



Liu et al.
Sphereface: Deep hypersphere embedding for face recognition.
CVPR 2017

Improvements


ArcFace

Current state of the art.

Deng et al. proposed the **additive angular margin loss** to the problem of face recognition.

Deng et al.

Additive Angular Margin Loss for Deep Face Recognition.
CVPR 2019



**Margin-less
class separation**


Improvements

ArcFace



Current state of the art.

Deng et al. proposed the **additive angular margin loss** to the problem of face recognition.

Deng et al.
Additive Angular Margin Loss for Deep Face Recognition.
CVPR 2019



Margin-less
class separation



Additive angular
margin loss

Improvements


ArcFace

Current state of the art.

Deng et al. proposed the **additive angular margin loss** to the problem of face recognition.

Deng et al.

Additive Angular Margin Loss for Deep Face Recognition.
CVPR 2019



Improvements

Source: Liu et al.
Sphereface: Deep hypersphere embedding for face recognition.
CVPR 2017

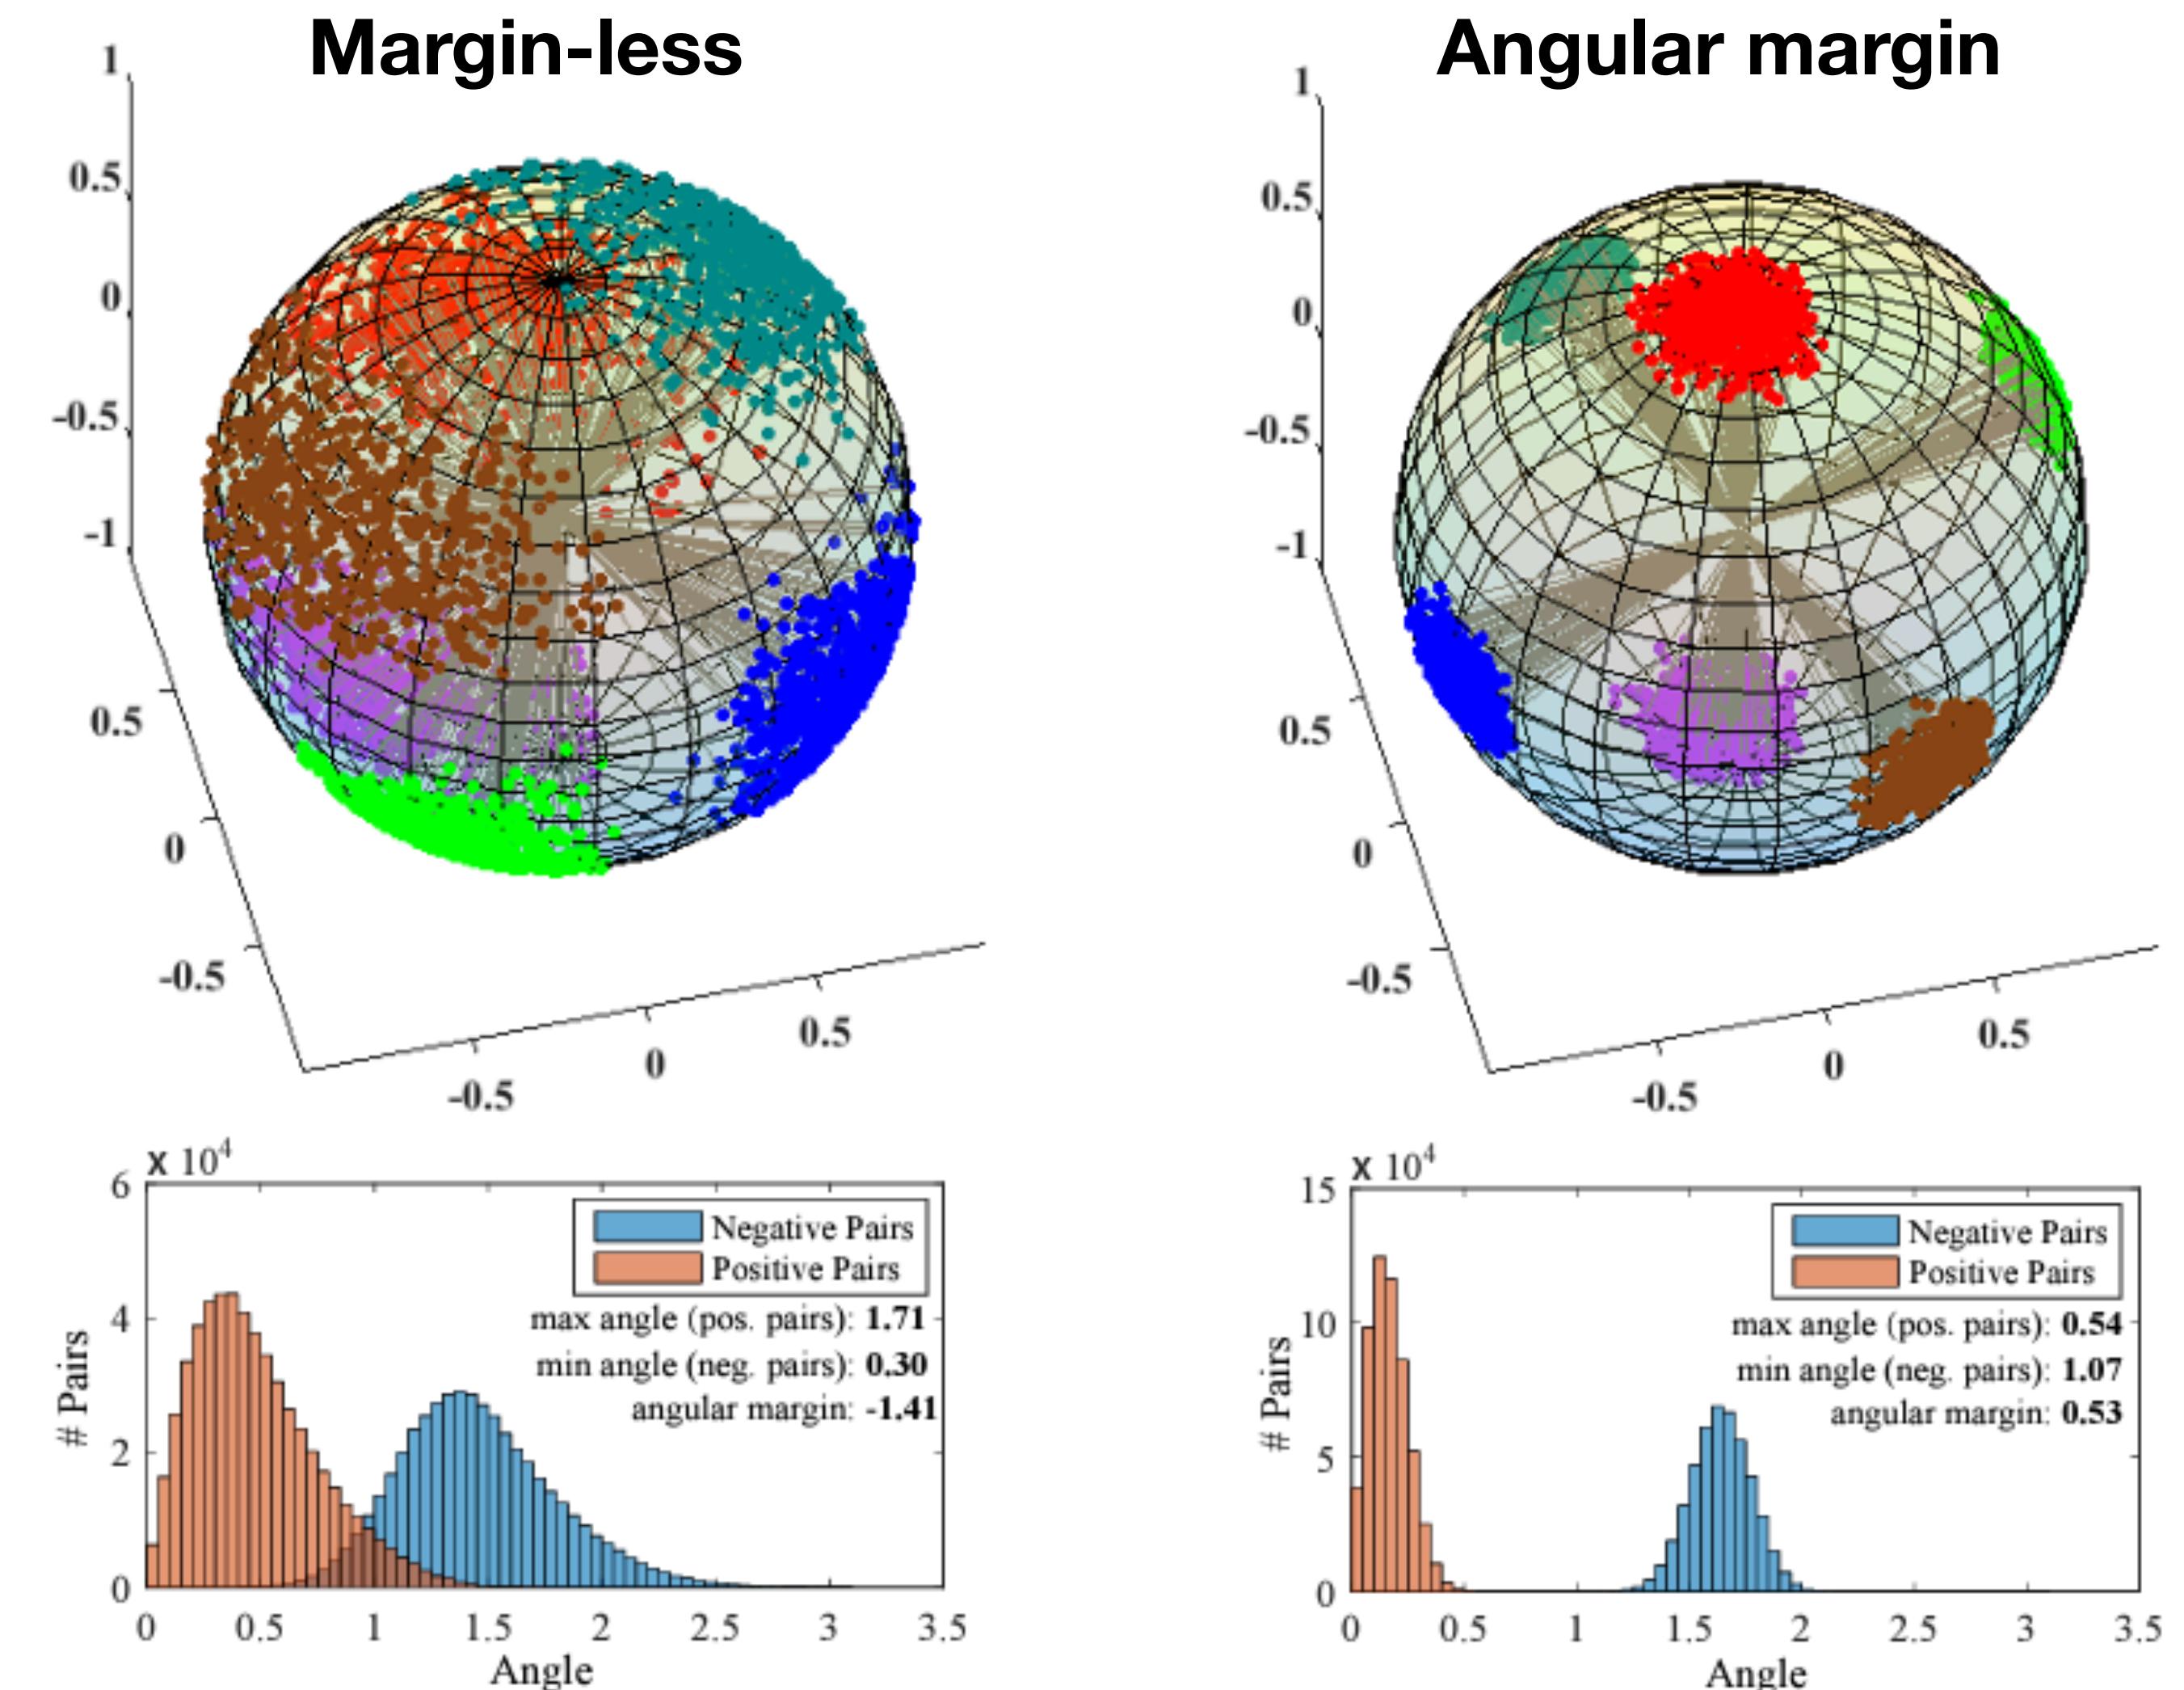
ArcFace

Current state of the art.

Deng et al. proposed the **additive angular margin loss** to the problem of face recognition.

Deng et al.

Additive Angular Margin Loss for Deep Face Recognition.
CVPR 2019



Data-Driven Face Recognition

Problems

Accountability

You must understand what the network is using to classify samples.

You must avoid this
in the case of
Face Recognition!

<https://twitter.com/EricTopol/status/1161657580675985409>

Eric Topol  @EricTopol

How surgical skin markings faked out a deep learning #AI neural net-- a commercially approved product for algorithm-aided melanoma diagnosis. Highly instructive. Machines can be dumb.

jamanetwork.com/journals/jamaderm by @JAMADerm by @UniHeidelberg

JAMA Dermatology | Original Investigation

Association Between Surgical Skin Markings in Dermoscopic Images and Diagnostic Performance of a Deep Learning Convolutional Neural Network for Melanoma Recognition

Julia K. Becker, MD; Christine Fink, MD; Ferdinand Tschann, MD; Alexander Eick, MD; Thomas Gründel, MD; Werner Hofmann-Wohlwend, MD; Luk Thomas, MD; Amélie Lalle, MD; Andreas Blum, MD; Wilhelm Stoeck, MD; Holger A. Haenssle, MD

Editorial

Supplemental content

IMPORTANCE Deep learning convolutional neural networks (CNNs) have shown a performance at the level of dermatologists in the diagnosis of melanoma. Accordingly, further exploring the potential limitations of CNN technology before broadly applying it is of special interest.


OBJECTIVE To investigate the association between gentian violet surgical skin markings in dermoscopic images and the diagnostic performance of a CNN approved for use as a medical device in the European market.

DESIGN AND SETTING A cross-sectional analysis was conducted from August 1, 2018, to November 30, 2018, using a CNN architecture trained with more than 120 000 dermoscopic images of skin neoplasms and corresponding diagnoses. The association of gentian violet skin markings in dermoscopic images with the performance of the CNN was investigated in 3 image sets of 180 melanocytic lesions each (0.07 benign, 0.25 melanocytic).

EXPOSURE The same lesions were sequentially imaged with and without the application of a gentian violet surgical skin marker and then evaluated by the CNNs for their probability of being a melanoma. In addition, the markings were removed by manually cropping the dermoscopic images to focus on the melanocytic lesion.

MAIN OUTCOMES AND MEASURES Sensitivity, specificity, and area under the curve (AUC) of the receiver operating characteristic (ROC) curve for the CNNs diagnostic classification in unmarked, marked, and cropped images.

RESULTS In all, 180 melanocytic lesions (107 benign nevi and 23 melanomas) were imaged. In unmarked lesions, the CNN achieved a sensitivity of 95.7% (95% CI, 79.6–99.2%) and a specificity of 84.7% (95% CI, 76.0%–89.8%). The ROC AUC was 0.969. In marked lesions, an increase in melanoma probability scores was observed that resulted in a sensitivity of 100% (95% CI, 85.7%–100%) and a significantly reduced specificity of 45.8% (95% CI, 36.7%–55.2%; $P < .001$). The ROC AUC was 0.932. Cropping images led to the highest sensitivity of 100% (95% CI, 85.7%–100%) and specificity of 97.2% (95% CI, 92.1%–99.0%), and ROC AUC of 0.993. Heat maps created by vanilla gradient descent backpropagation indicated that the blue markings were associated with the increased false positive rate.



Data-Driven Face Recognition

Problems

Accountability

You must understand what the network is using to classify samples.

Comments on:

<https://www.youtube.com/watch?v=rga2-d1oi30>

Automated Inference on Criminality using Face Images

Xiaolin Wu
Shanghai Jiao Tong University
xwu510@gmail.com

Xi Zhang
Shanghai Jiao Tong University
zhangxi.19930818@sjtu.edu.cn

Abstract

We study, for the first time, automated inference on criminality based solely on still face images. Via supervised machine learning, we build four classifiers (logistic regression, KNN, SVM, CNN) using facial images of 1856 real persons controlled for race, gender, age and facial expressions, nearly half of whom were convicted criminals, for discriminating between criminals and non-criminals. All four classifiers perform consistently well and produce evidence for the validity of automated face-induced inference on criminality based on facial images.

people share the belief that the face alone suffices to reveal innate traits of a person. Aristotle in his famous work Prior Analytics asserted, "It is possible to infer character from features, if it is granted that the body and the soul are changed together by the natural affections". Psychologists have known, for as long as a millennium, the human tendency of inferring innate traits and social attributes (e.g., the trustworthiness, dominance) of a person from his/her facial appearance, and a robust consensus of individuals' inferences . These are the facts found through numerous studies [2, 32, 4, 5, 9, 20, 21, 27, 25].


Data-Driven Face Recognition

Problems

Accountability

You must understand what the network is using to classify samples.

Comments on:
<https://www.youtube.com/watch?v=rga2-d1oi30>



Data-Driven Face Recognition

Problems

Accountability

You must understand what the network is using to classify samples.

Comments on:

<https://www.youtube.com/watch?v=rga2-d1oi30>




Data-Driven Face Recognition

Problems

Bias

What happens if you train the network only with one type of faces (e.g., with only young caucasians)?



Data-Driven Face Recognition

Problems


Avoid Bias

Diversify the training dataset.

There are synthetic ways to do it...
(FaceGen demonstration)



S'up Next?



S'up Next?

Face Recognition Coding Class

Please bring your computers.

