

Computer Vision Applications

COMP 388-002/488-002 Computer Science Topics

Daniel Moreira
Fall 2022



LOYOLA
UNIVERSITY CHICAGO

Letter Soup

**COMP 388-002/488-002 Computer Science Topics
Computer Vision Applications**

Daniel Moreira
Fall 2022



LOYOLA
UNIVERSITY CHICAGO

Today you will...

Get an overview of
CV, AI, ML, PR, SVM,
CNN, DL, GPU, PCA, etc.,
all in favor of the
upcoming seminars.



Course Overview

RECAP

Content

Date	Topic	Leader	Assignment
08/29	Introduction to CV	Instructor	N.A.
09/05	<i>Labor Day</i>	N.A.	A01, due on 09/13
09/12	Letter Soup: AI, ML, NN, and DL	Instructor	A02, due on 09/20
09/19	Local and Global Descriptors	Instructor	A03, due on 09/27
09/26	CBIR and Indexing	TBD (students)	A04, due on 10/04
10/03	Image Classification	TBD (students)	A05, due on 10/18
10/10	<i>Fall Break</i>	N.A.	N.A.

Date	Topic	Leader	Assignment
10/17	Object Detection	TBD (students)	A06, due on 10/25
10/24	Image Segmentation	TBD (students)	A07, due on 11/01
10/31	Face Detection	TBD (students)	A08, due on 11/08
11/07	Face Verification	TBD (students)	A09, due on 11/15
11/14	GANs and Generative DL	TBD (students)	A10, due on 11/29
11/21	Deep and Cheap Fakes	Instructor	N.A.
11/28	Sensitive Media Analysis	Instructor	N.A.
12/05	Provenance Analysis	TBD (students)	N.A.
12/12	<i>Final Exam</i>	N.A.	N.A.

Course Overview

RECAP

Content

Date	Topic	Leader	Assignment
08/29	Introduction to CV	Instructor	N.A.
09/05	<i>Labor Day</i>	N.A.	A01, due on 09/13
09/12	Letter Soup: AI, ML, NN, and DL	Instructor	A02, due on 09/20
09/19	Local and Global Descriptors	Instructor	A03, due on 09/27
09/26	CBIR and Indexing	TBD (students)	A04, due on 10/04
10/03	Image Classification	TBD (students)	A05, due on 10/18
10/10	<i>Fall Break</i>	N.A.	N.A.



Please answer the form at:
<https://forms.gle/wsNjWG3MDiZPsEzA9>

Date	Topic	Leader	Assignment
10/17	Object Detection	TBD (students)	A06, due on 10/25
10/24	Image Segmentation	TBD (students)	A07, due on 11/01
10/31	Face Detection	TBD (students)	A08, due on 11/08
11/07	Face Verification	TBD (students)	A09, due on 11/15
11/14	GANs and Generative DL	TBD (students)	A10, due on 11/29
11/21	Deep and Cheap Fakes	Instructor	N.A.
11/28	Sensitive Media Analysis	Instructor	N.A.
12/05	Provenance Analysis	TBD (students)	N.A.
12/12	<i>Final Exam</i>	N.A.	N.A.



Computer Vision (CV)

RECAP

Face recognition
Reverse imaging
generating synthetic visuals
Deep Fake multimodal (ie: text->img)
different understand computer
Processing present VQA
formats images
Product



Computer Vision (CV)

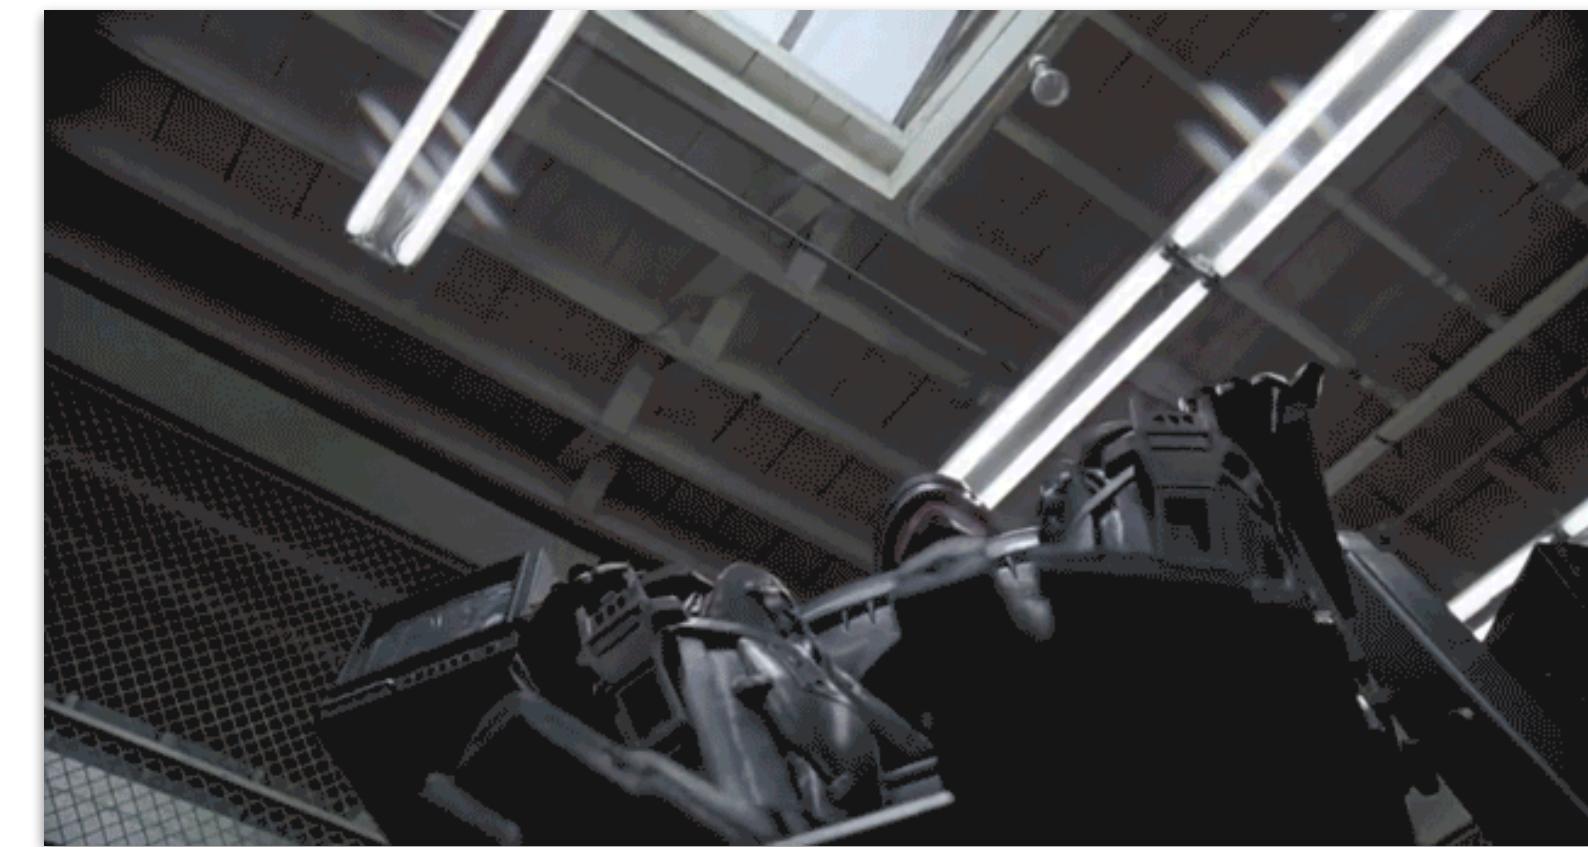
RECAP

Computer Science Subfield

It aims at developing computer systems
that mimic the human visual system.



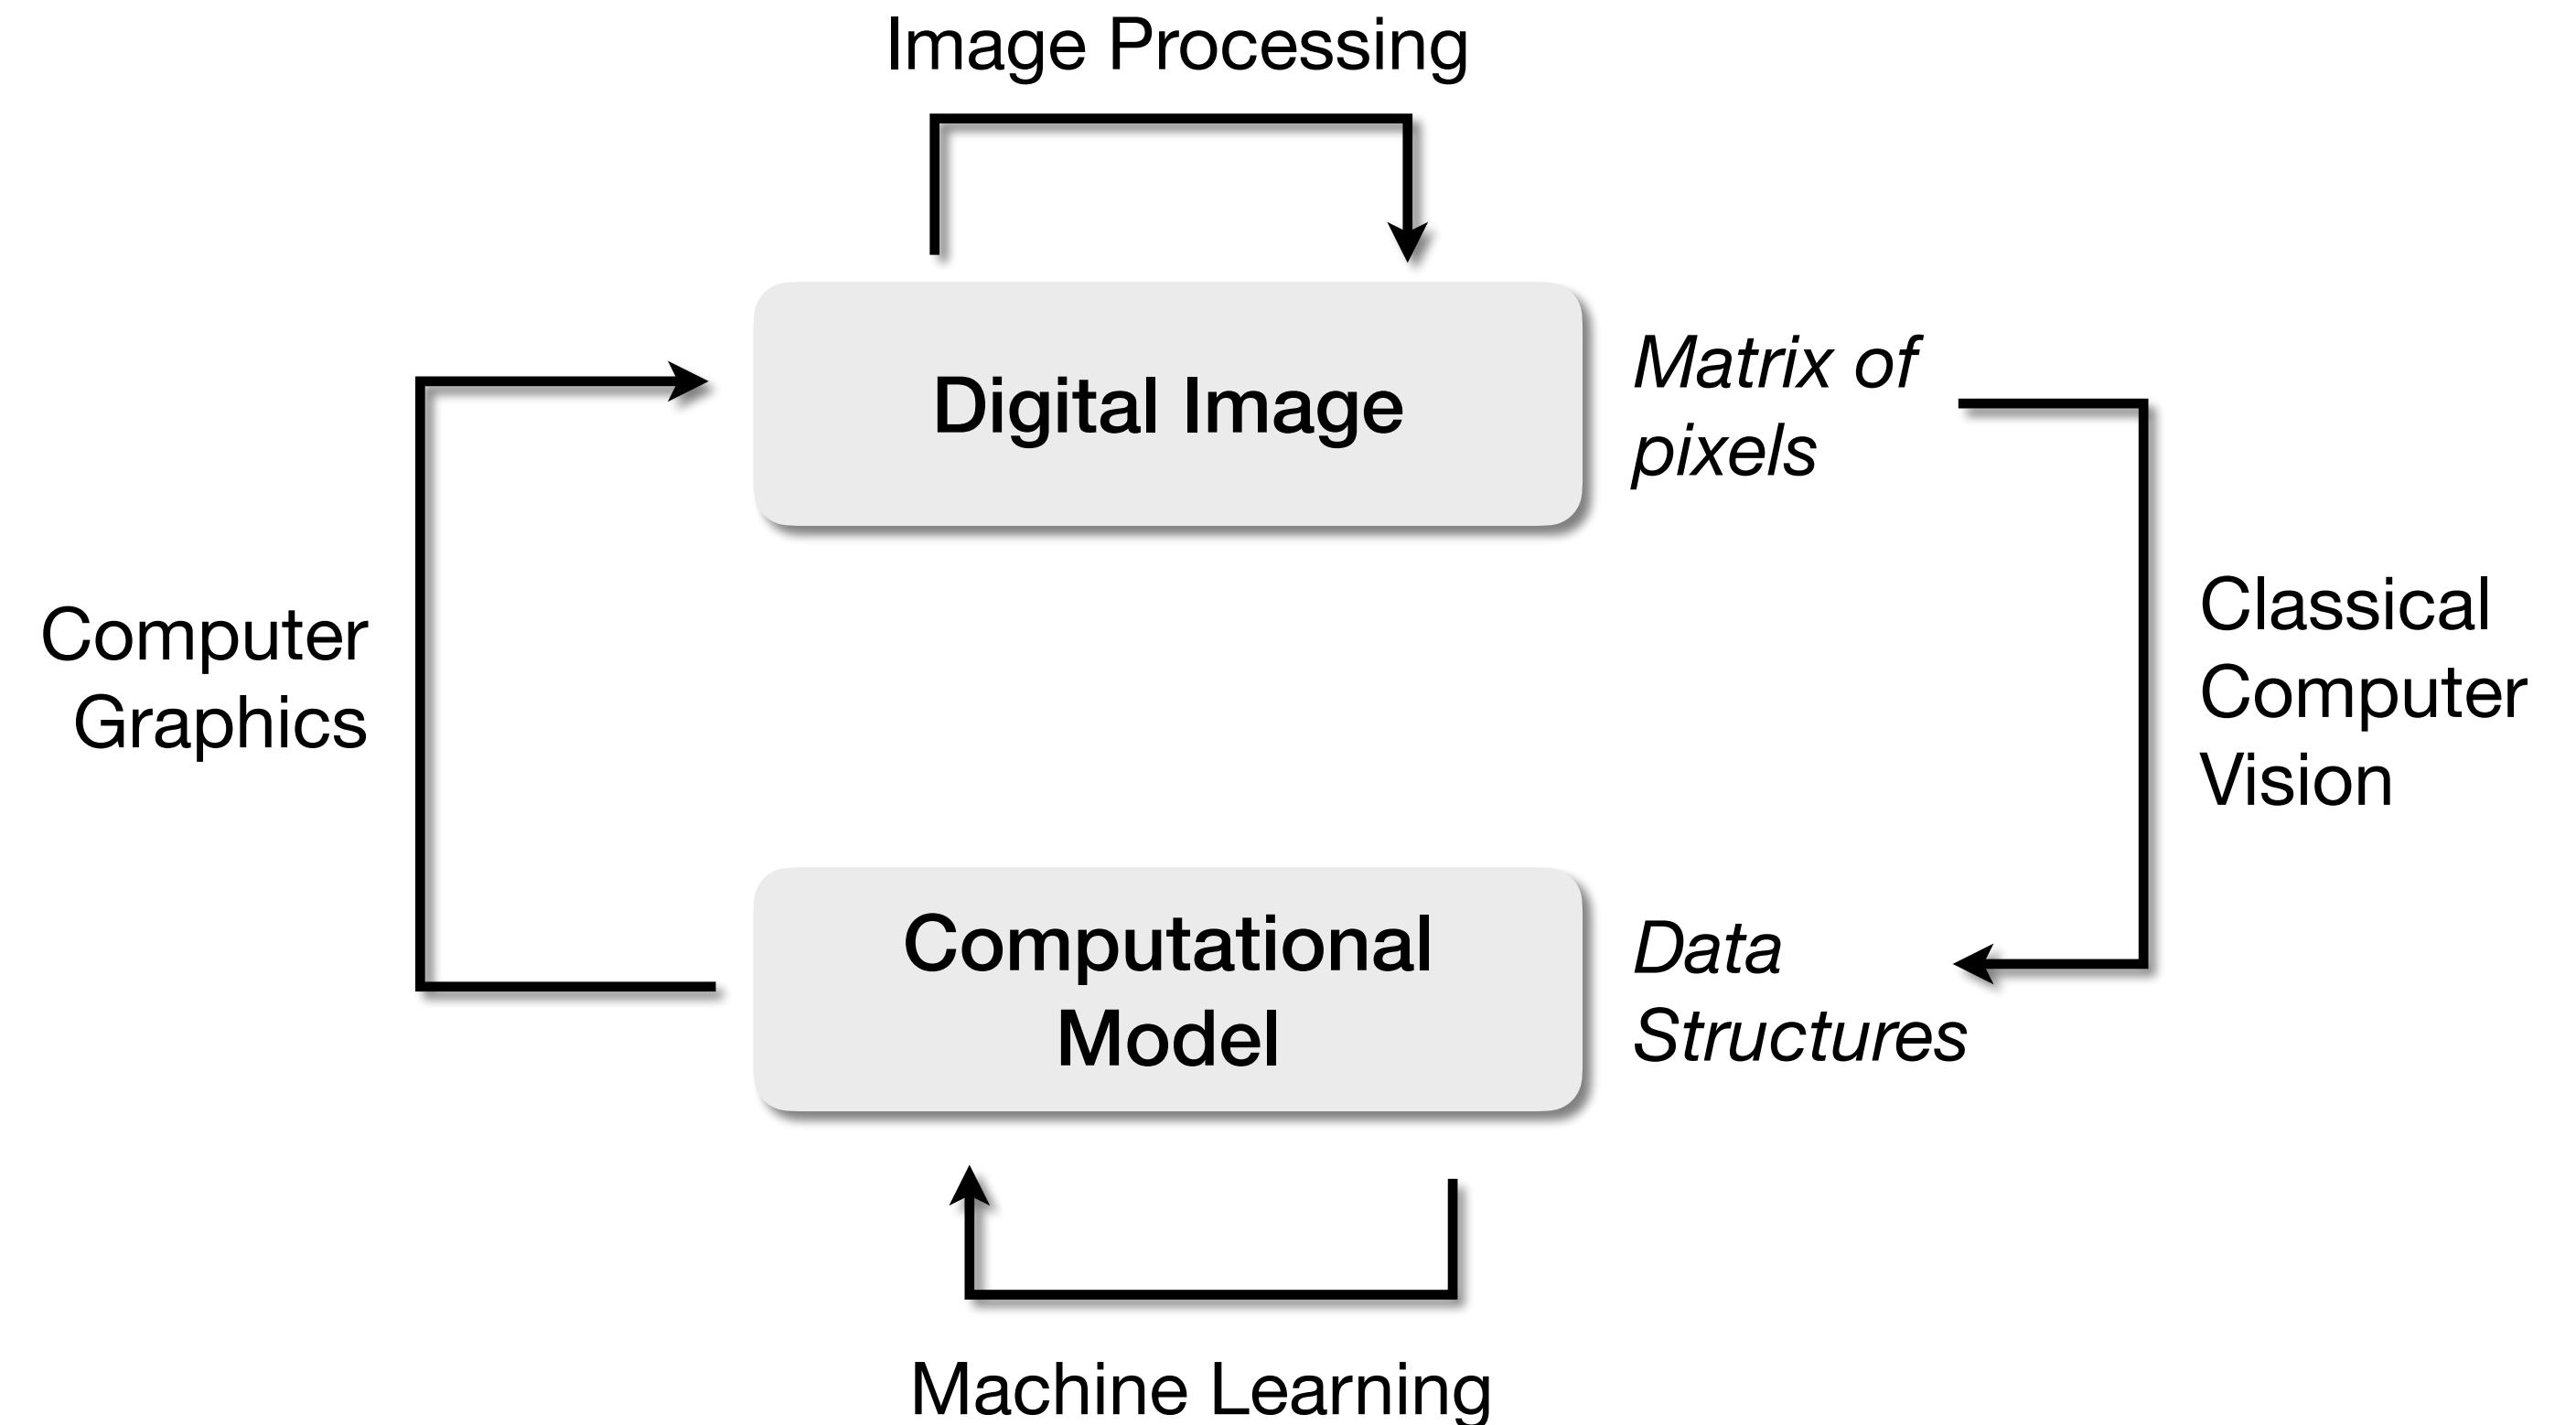
Reference



Objective

Computer Vision (CV)

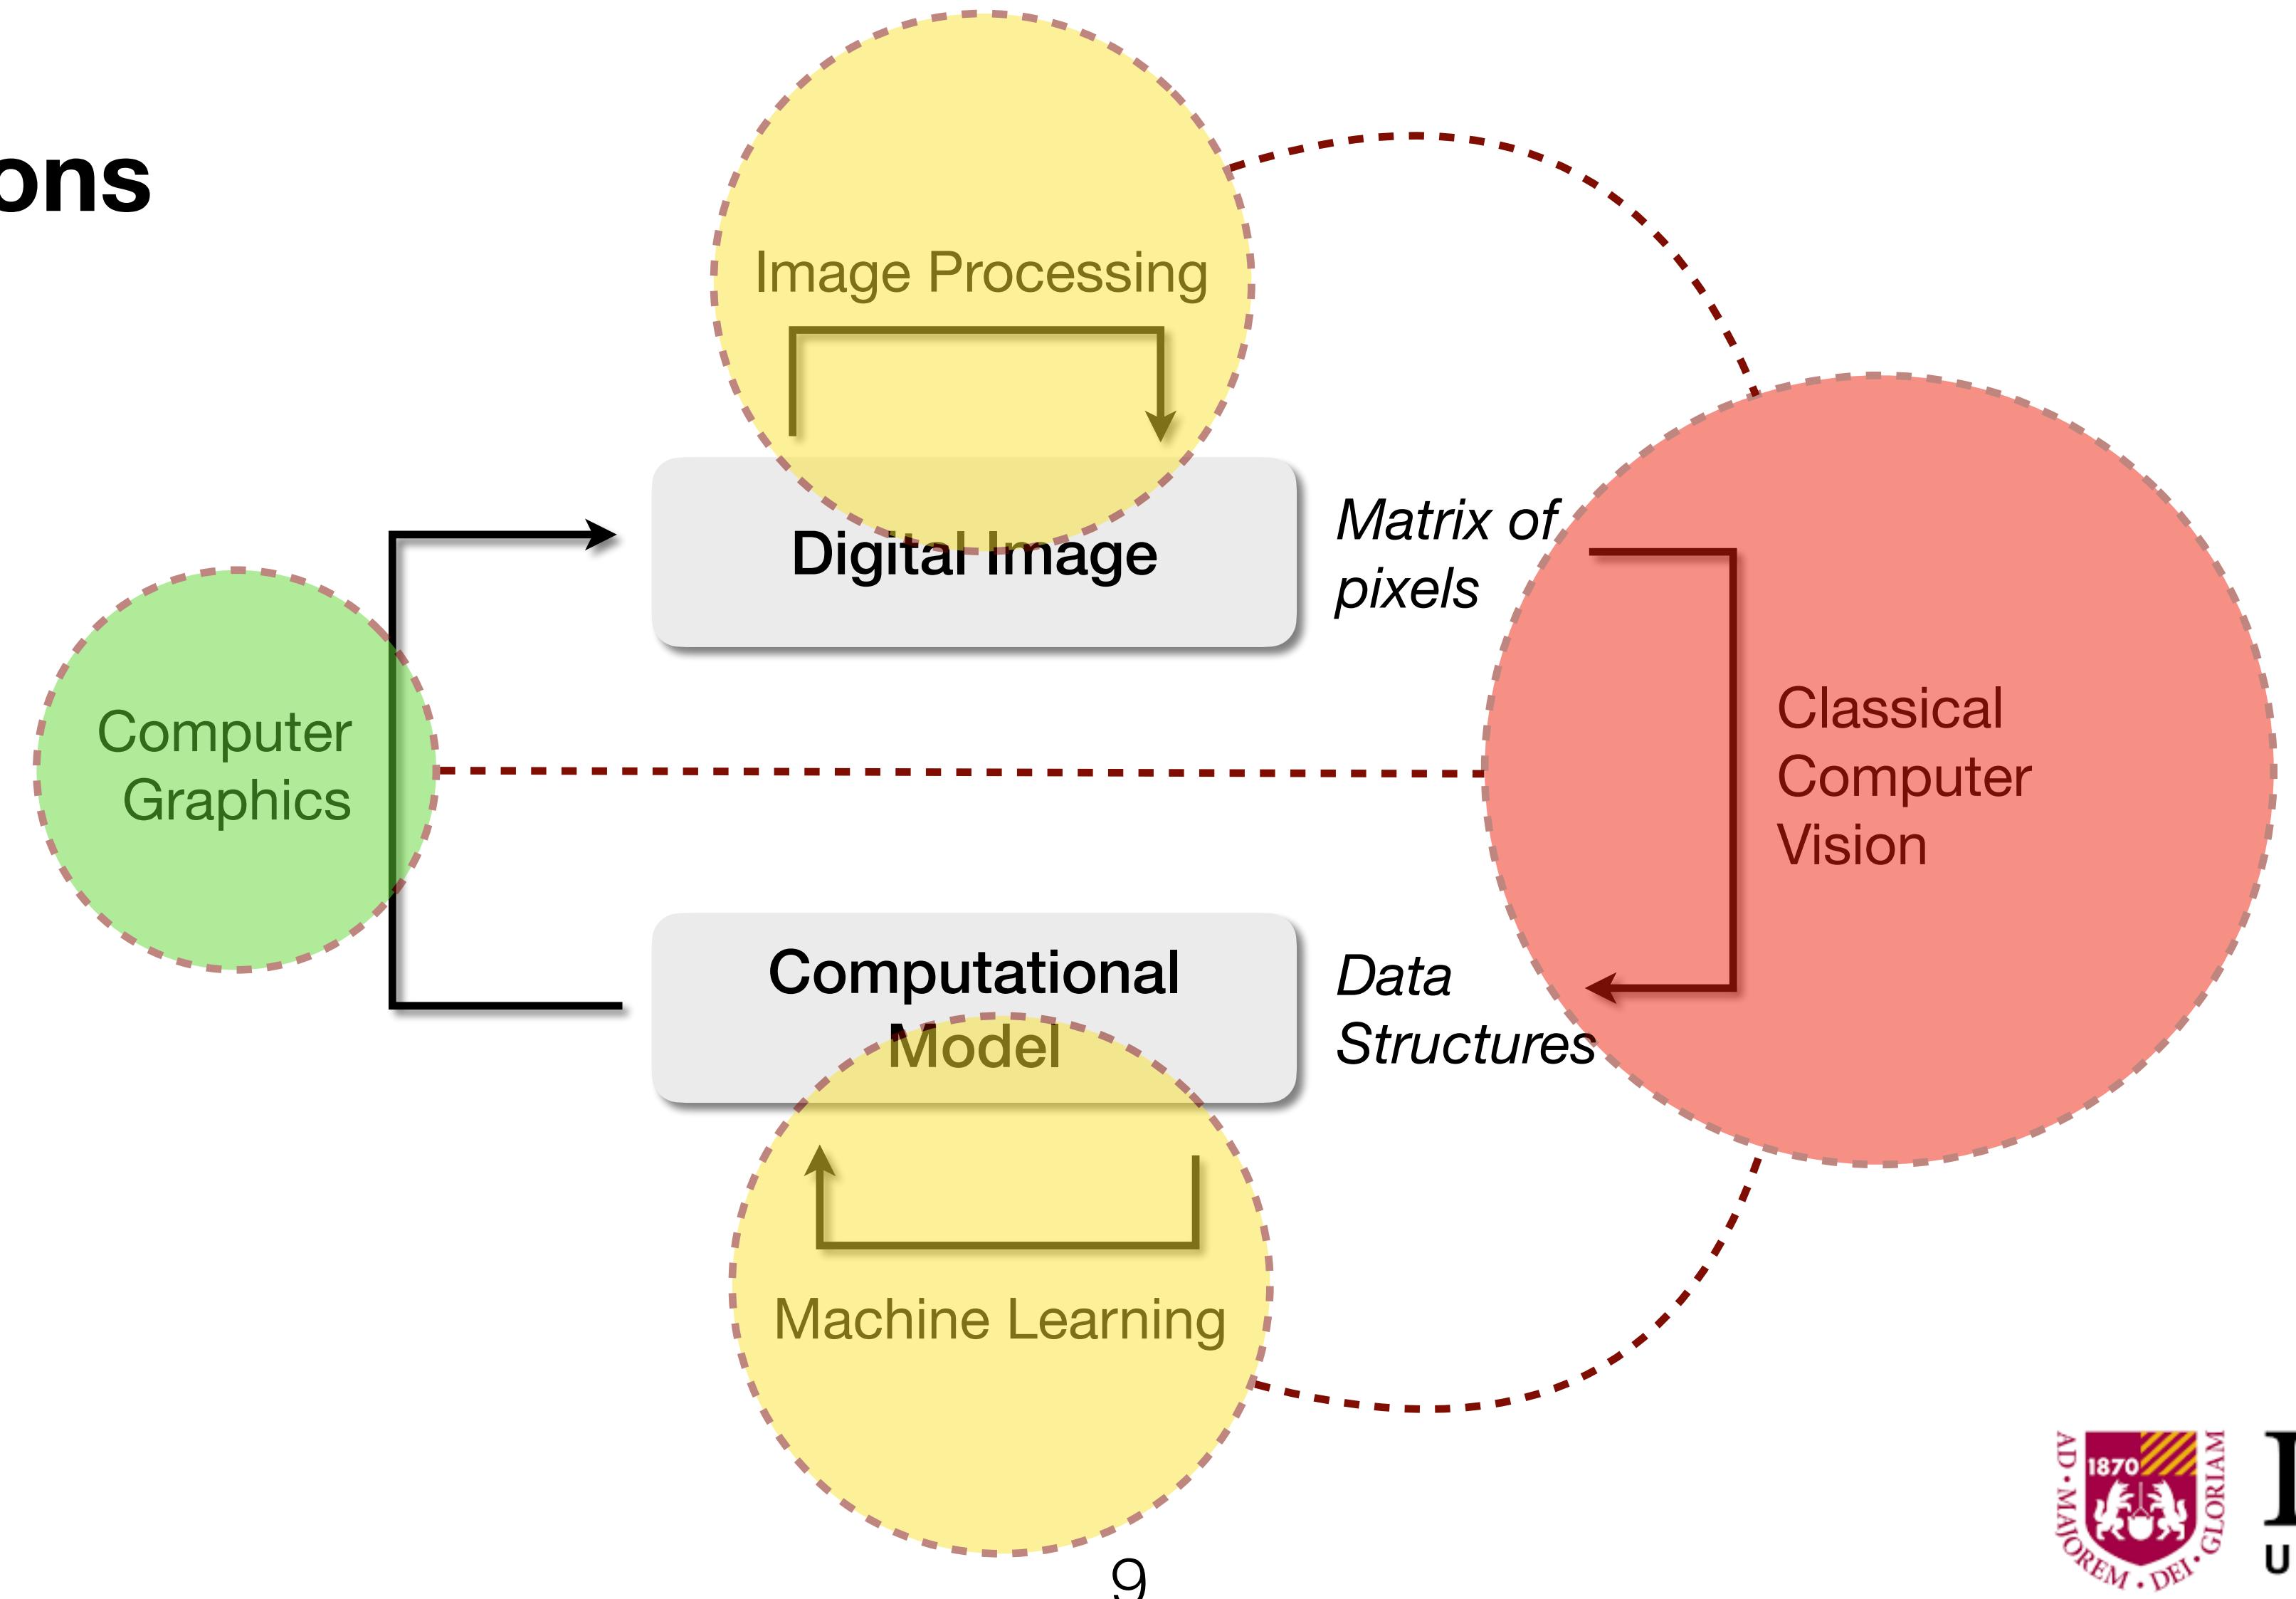
RECAP



Computer Vision (CV)

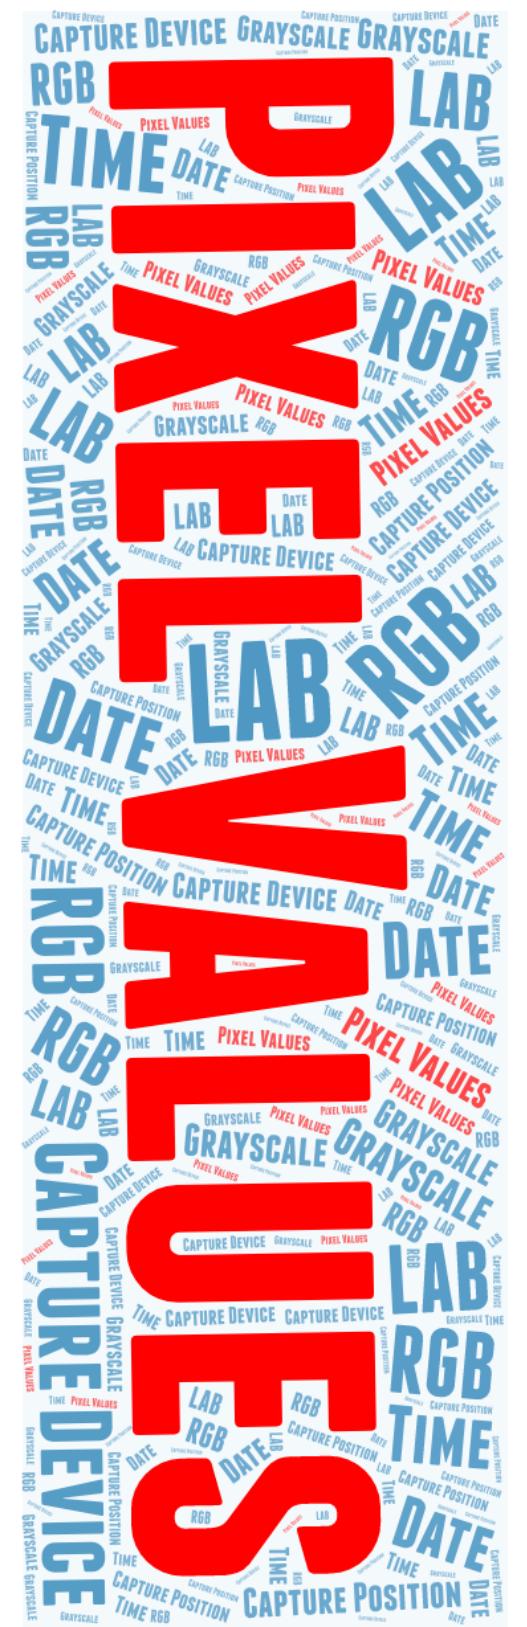
RECAP

Applications



Computer Vision (CV)

RECAP



Level 0



Semantic Gap



Task

Computer Vision (CV)

RECAP

Then



Level 0



Level 1



Level 2



Level 3



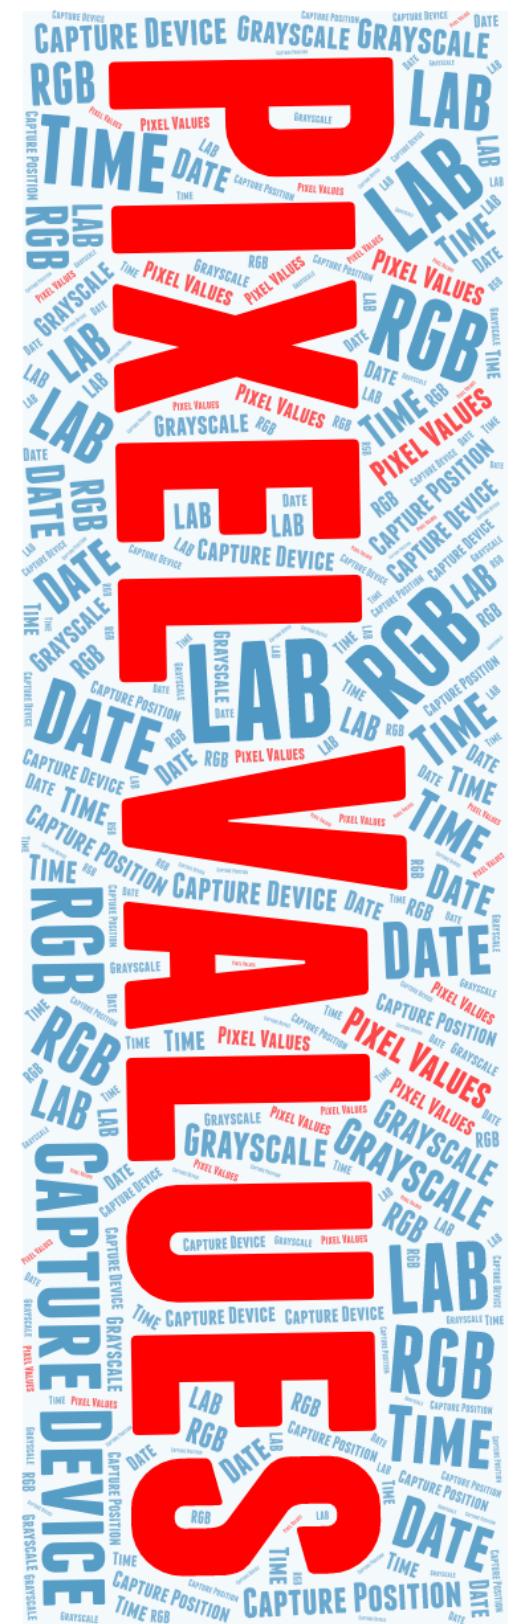
Task

Problem Domain Specialization

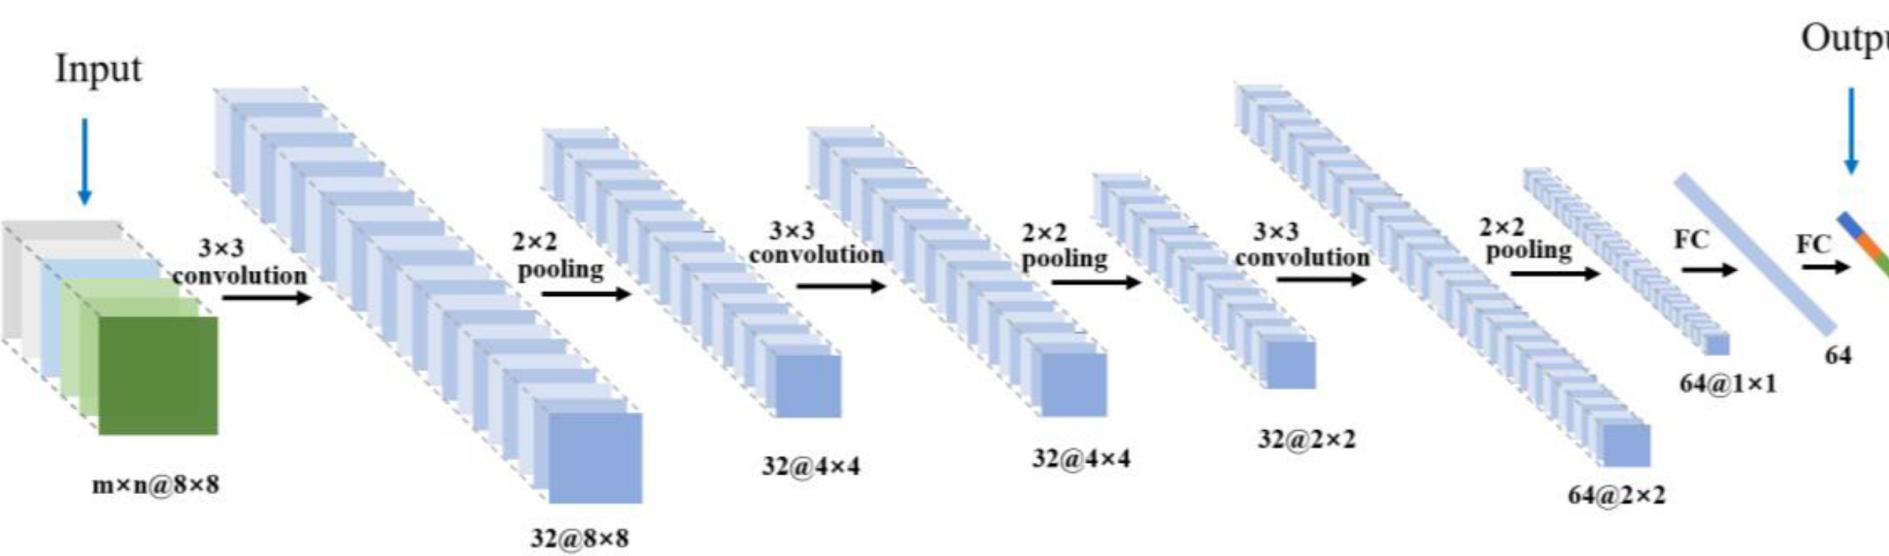
Computer Vision (CV)

RECAP

Now



Level 0



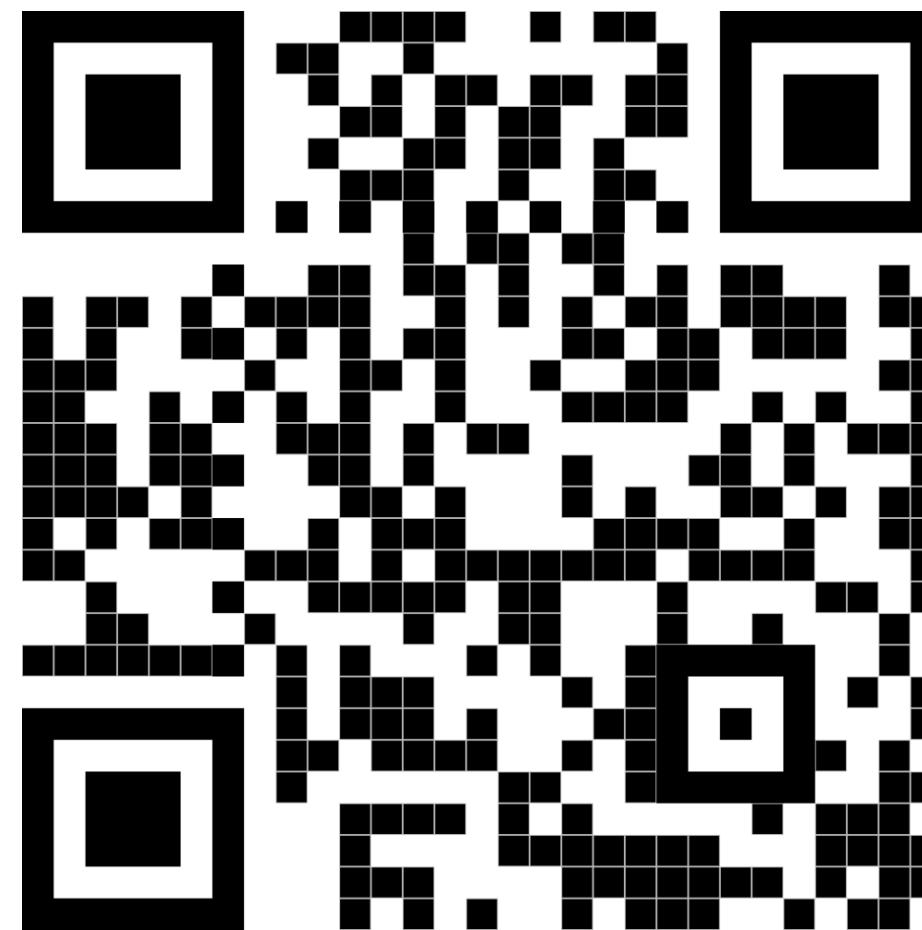
Deep Learning
Machine Learning?
Artificial Intelligence?



Task

Artificial Intelligence (AI)

What comes to your mind?



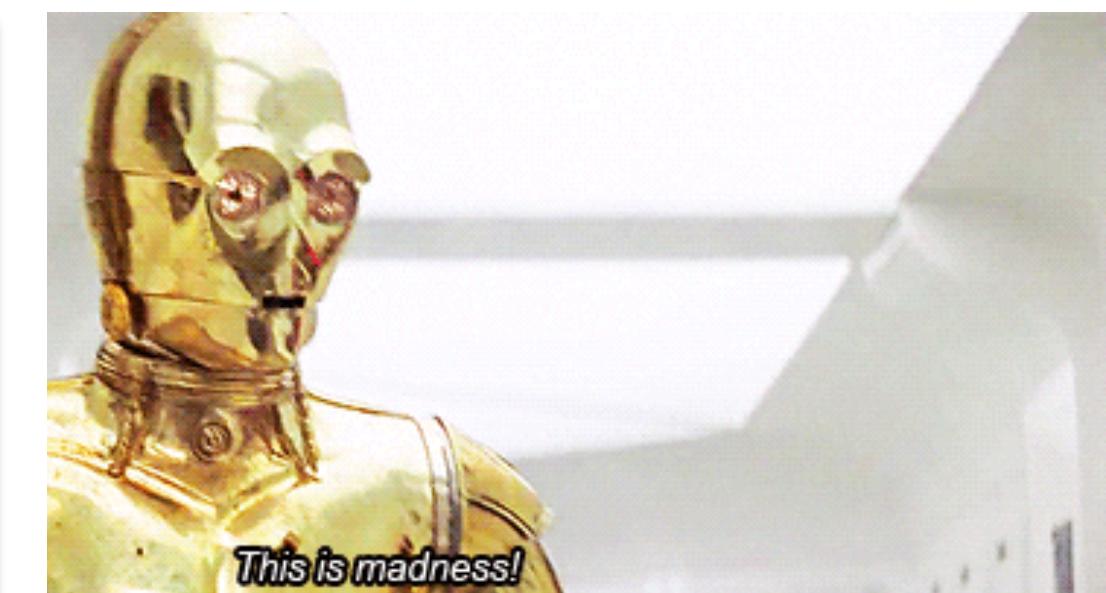
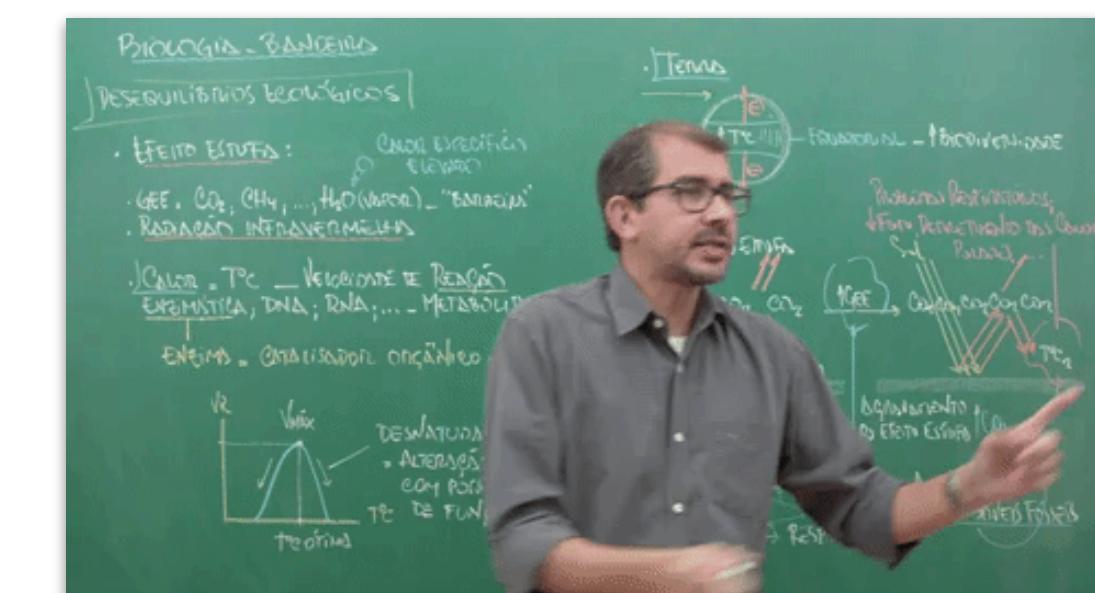
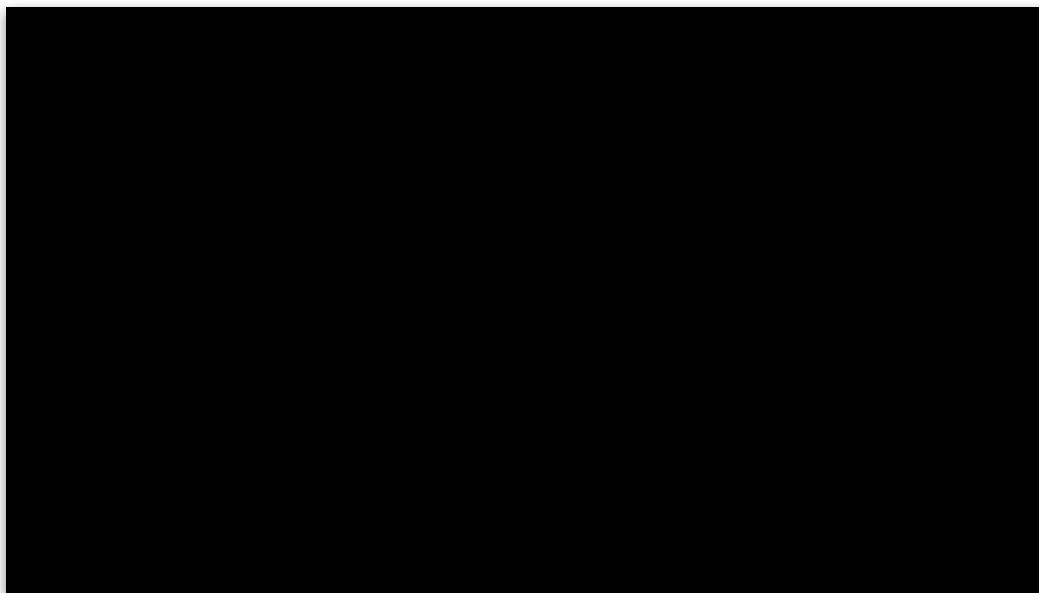
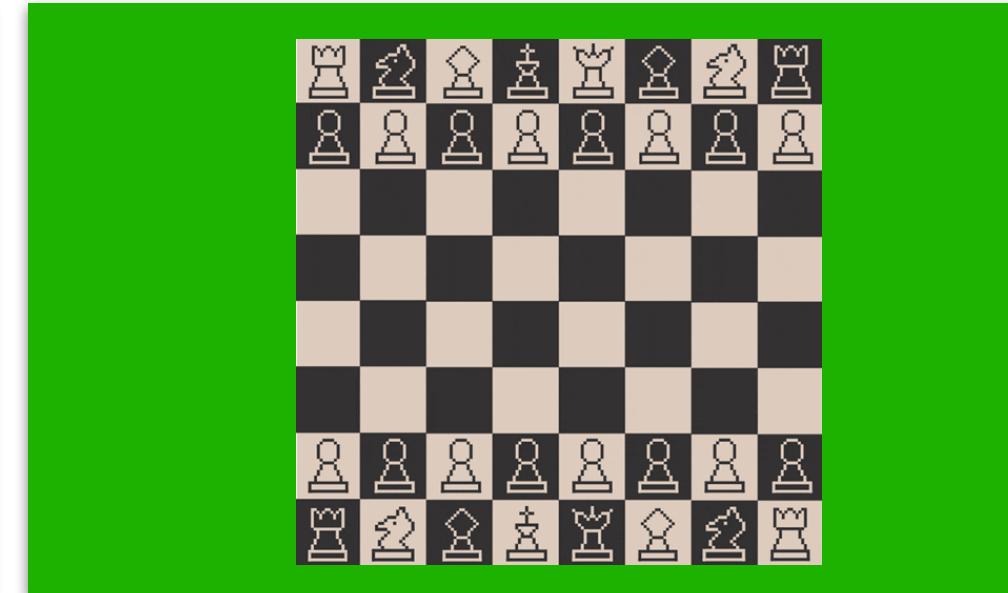
<https://bit.ly/3QFaq5G>



Artificial Intelligence (AI)

Computer Science PoV

It aims at developing computer systems that mimic (or overcome) humans' intelligence (or other living entities').



Humans

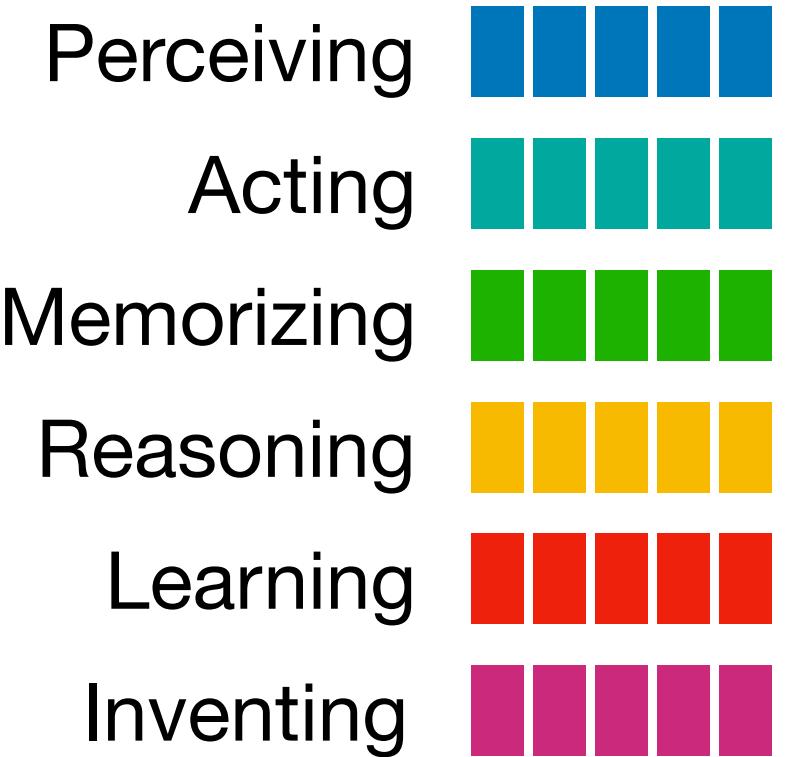
AI

Humans

AI

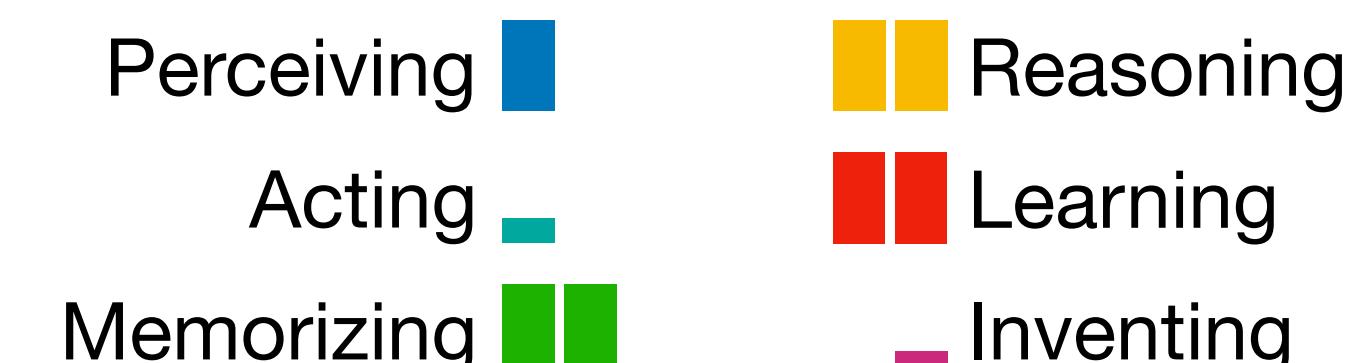
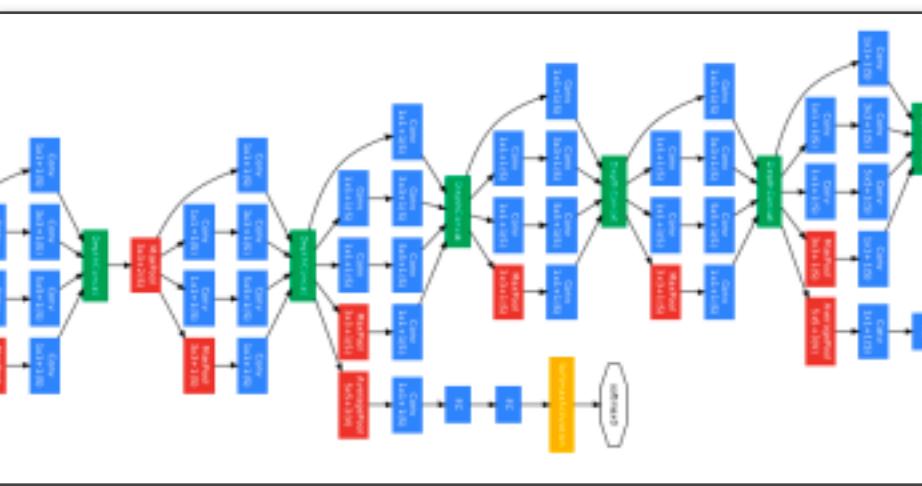
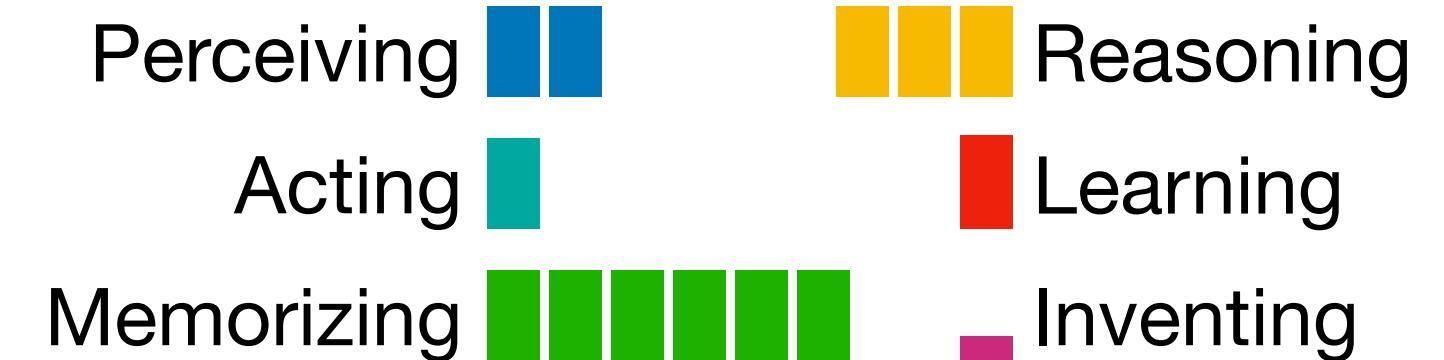
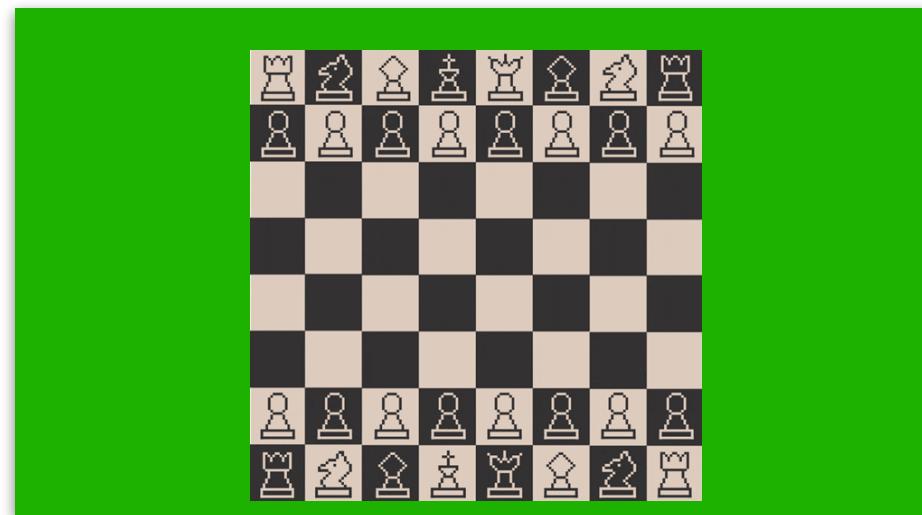
Artificial Intelligence (AI)

Humans (ref.)

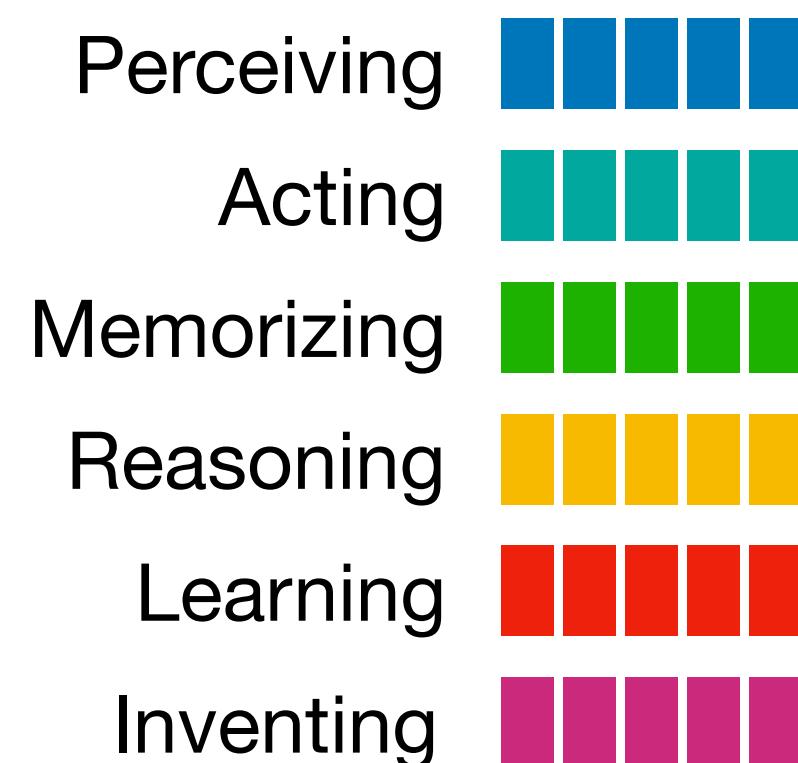
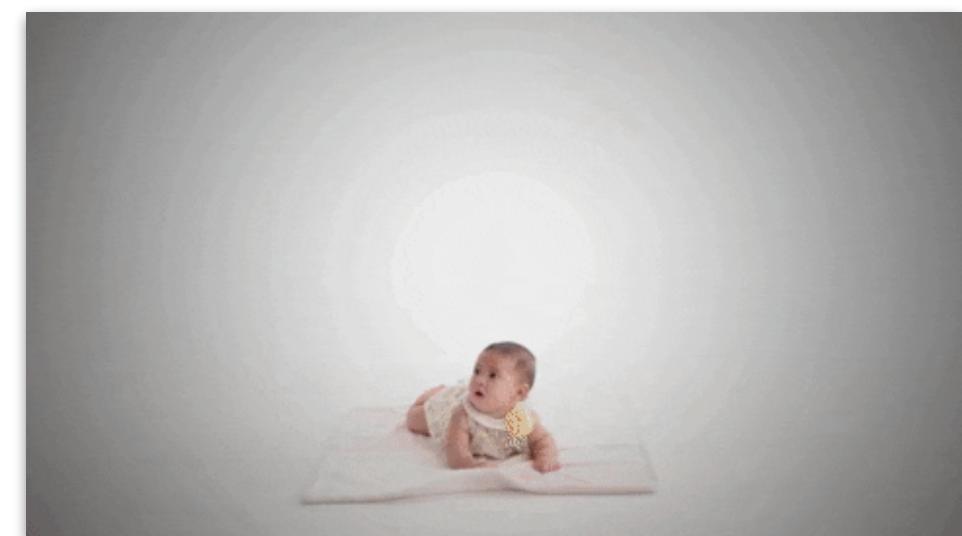


Artificial Intelligence (AI)

Specialized Systems



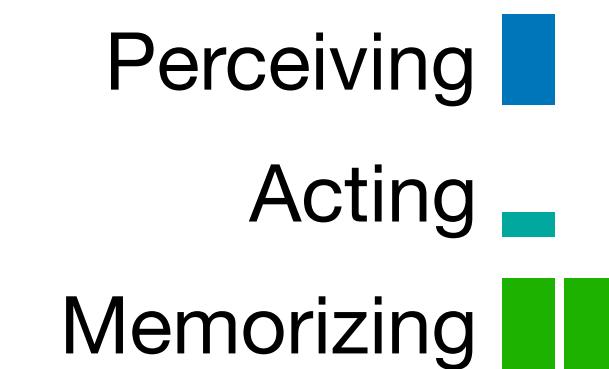
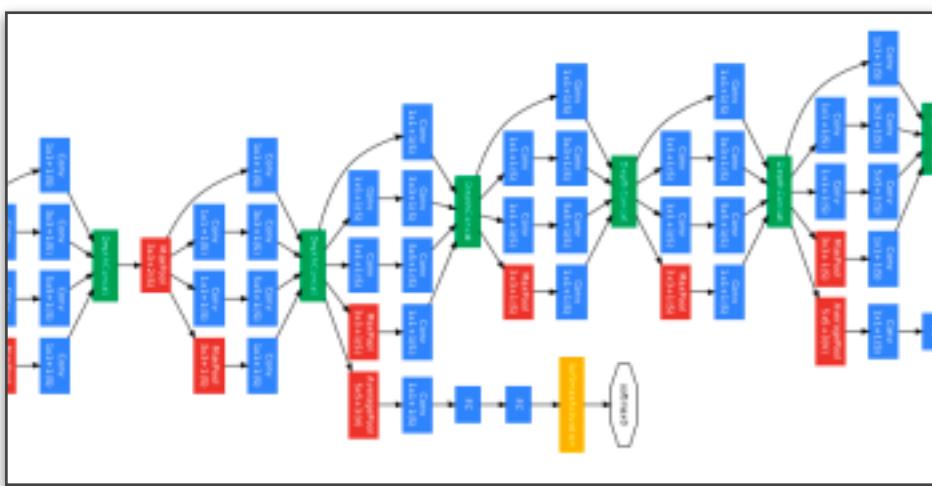
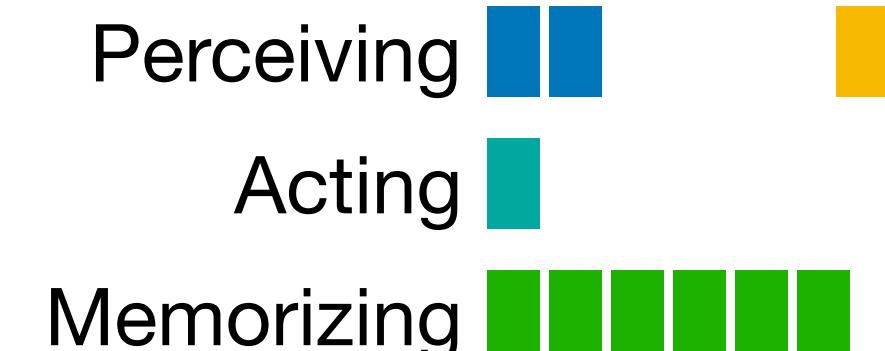
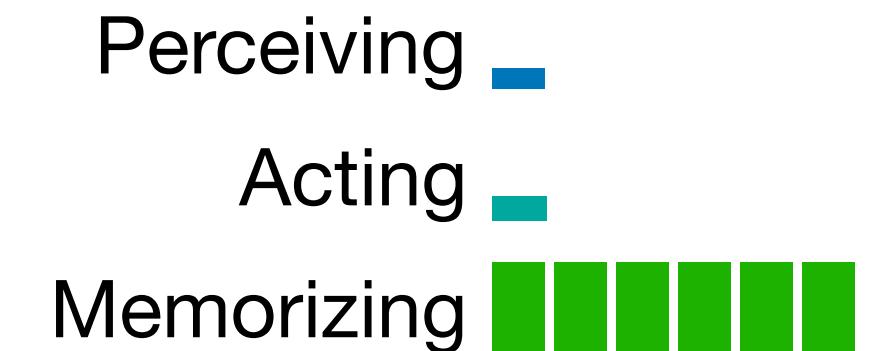
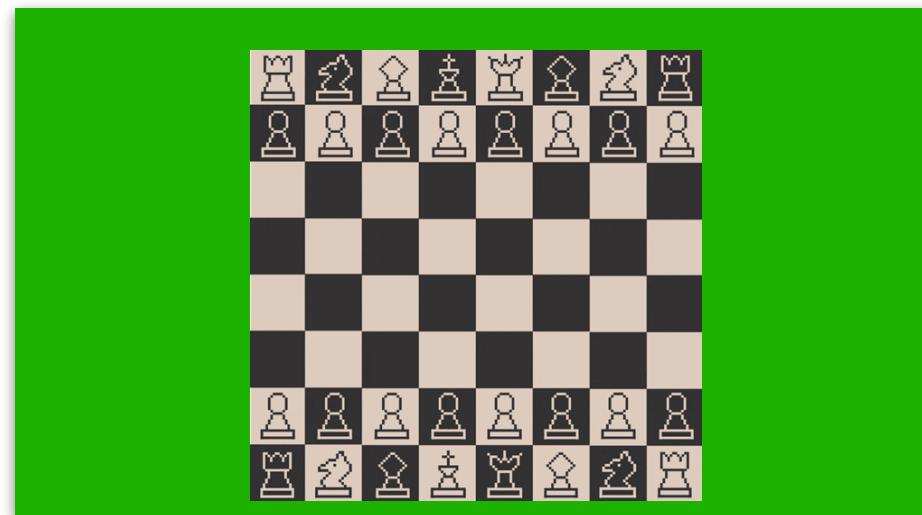
Humans (ref.)



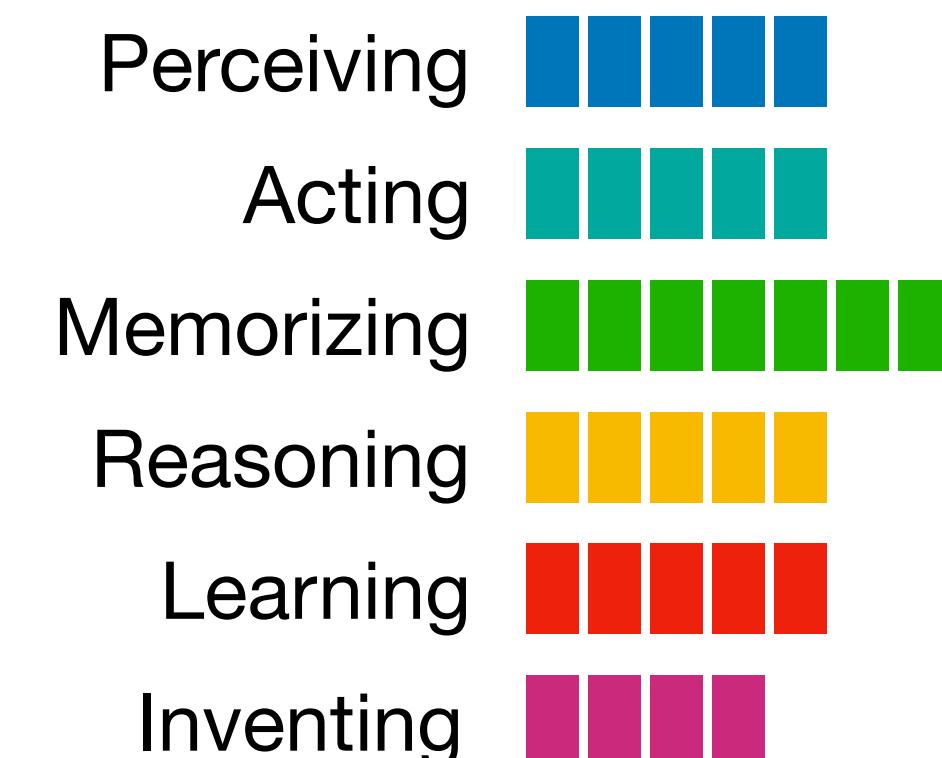
LOYOLA
UNIVERSITY CHICAGO

Artificial Intelligence (AI)

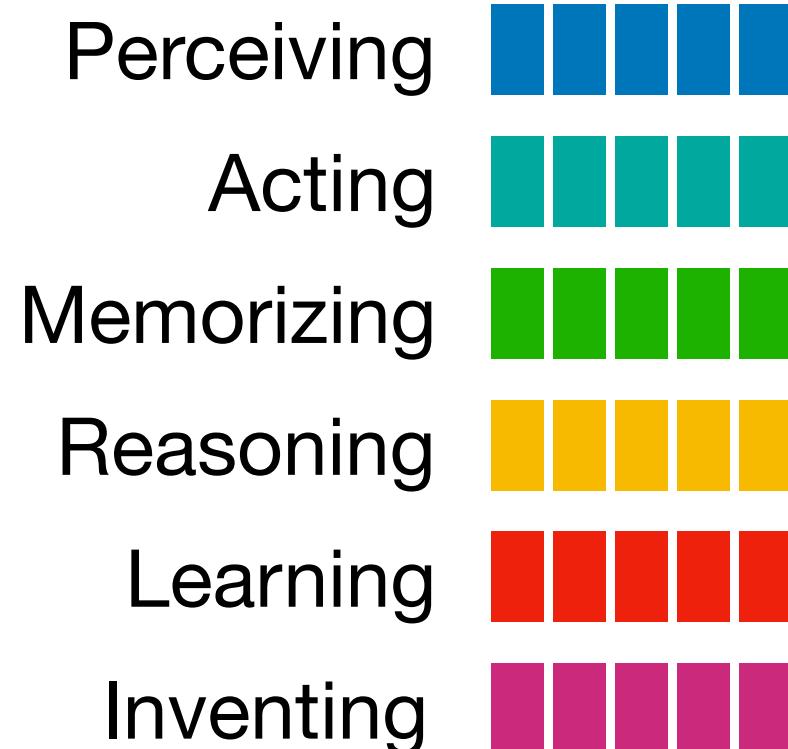
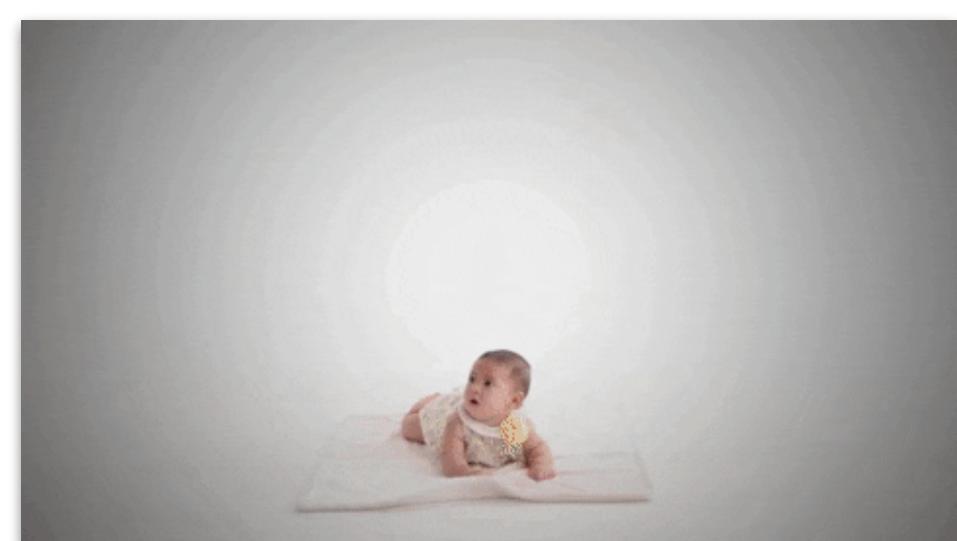
Specialized Systems



General-purpose Systems

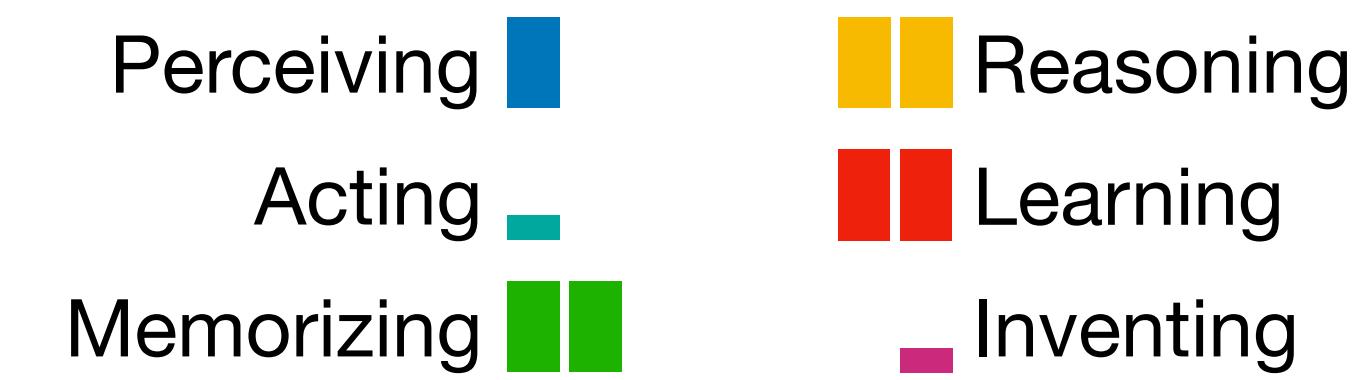
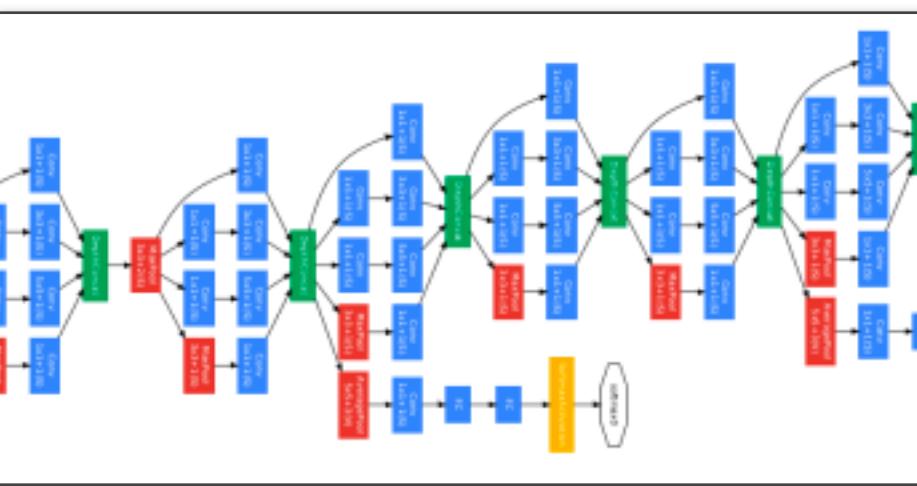
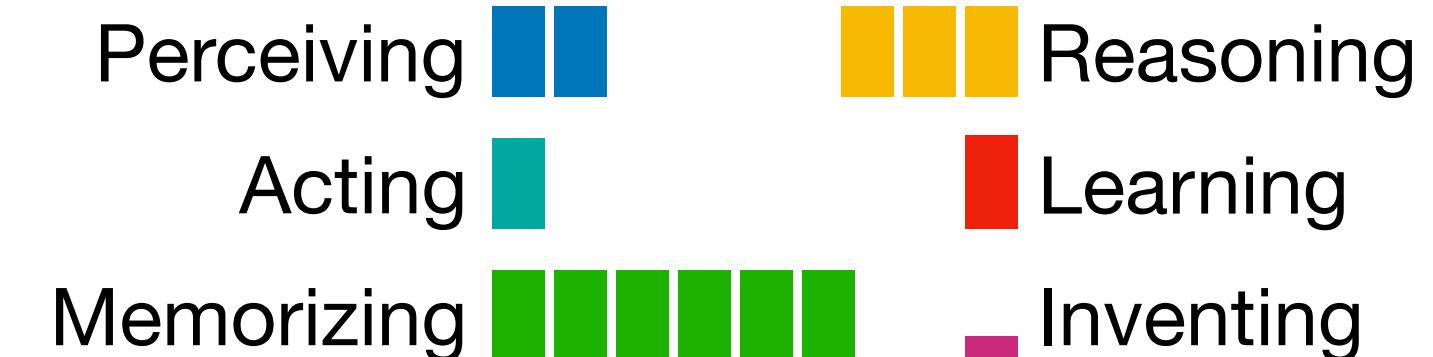
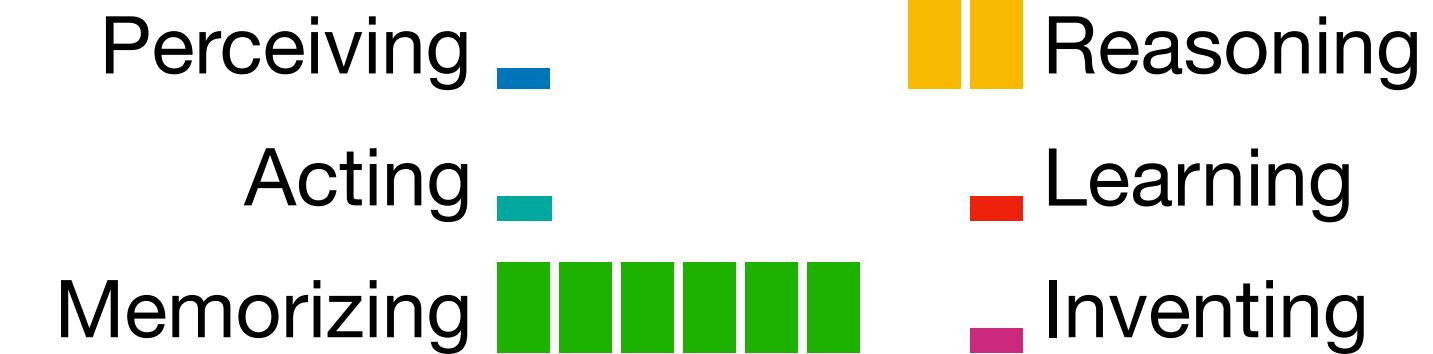
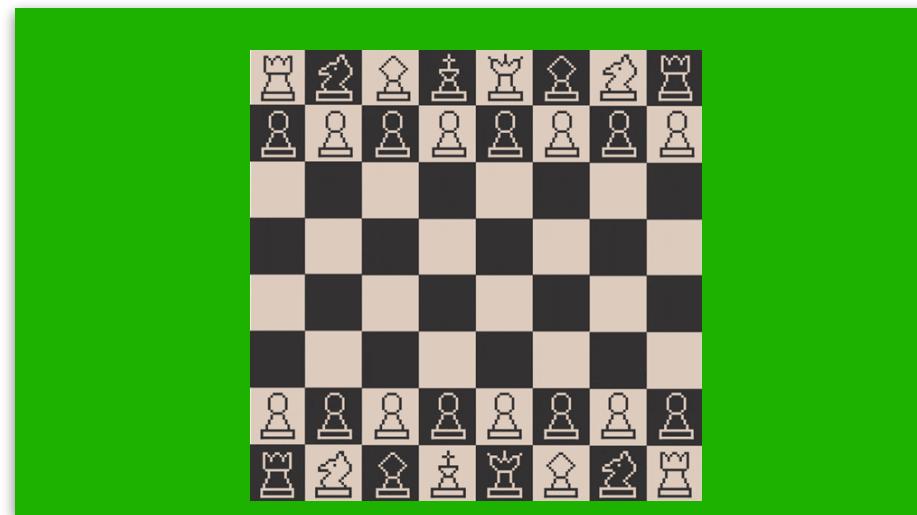


Humans (ref.)

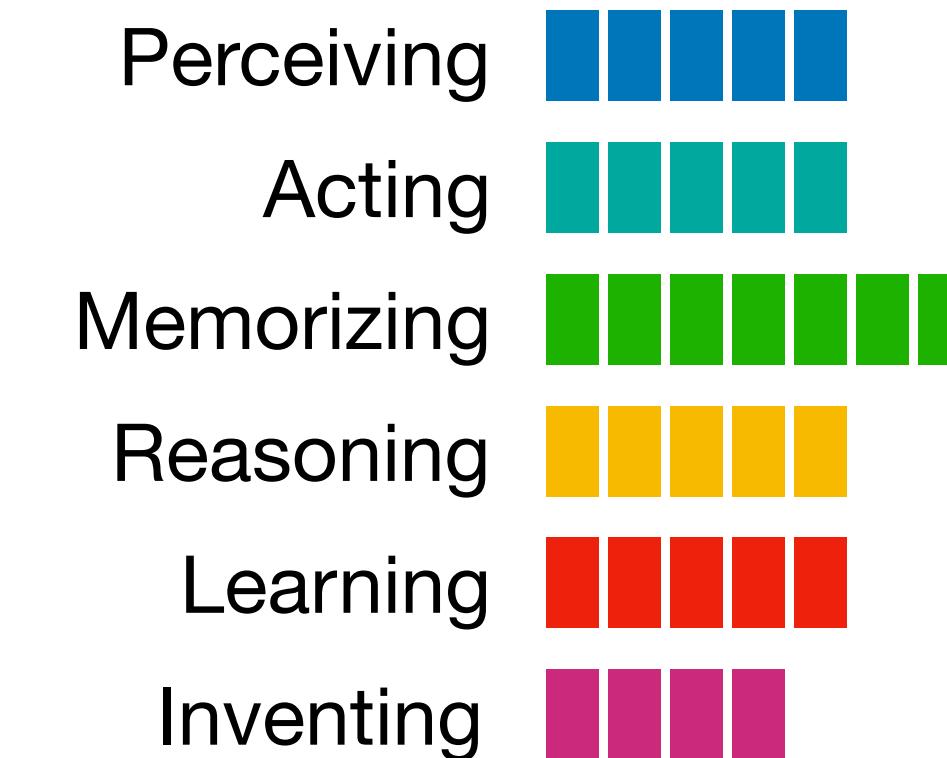


Artificial Intelligence (AI)

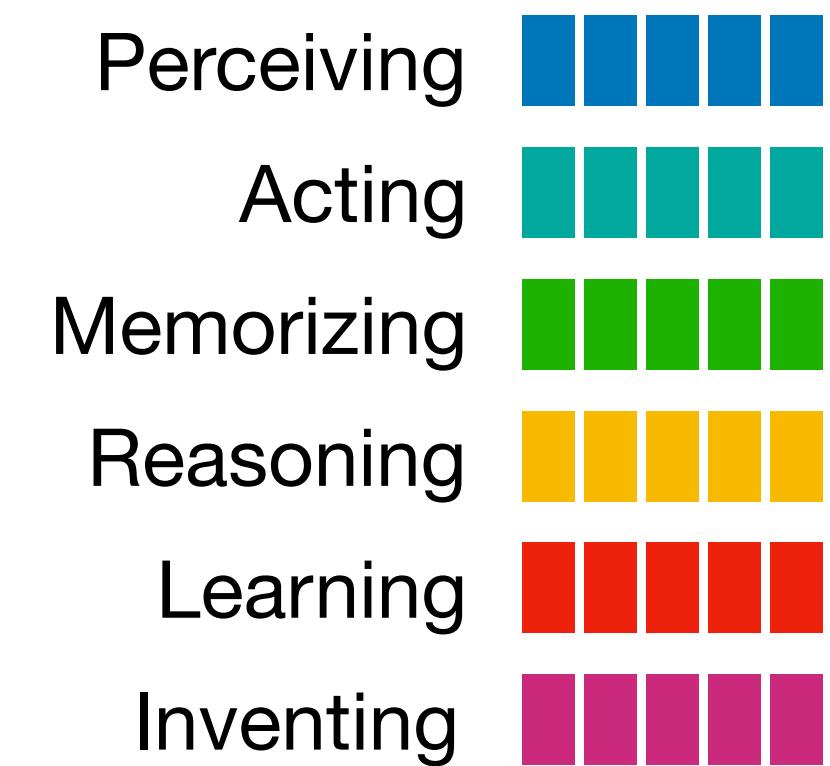
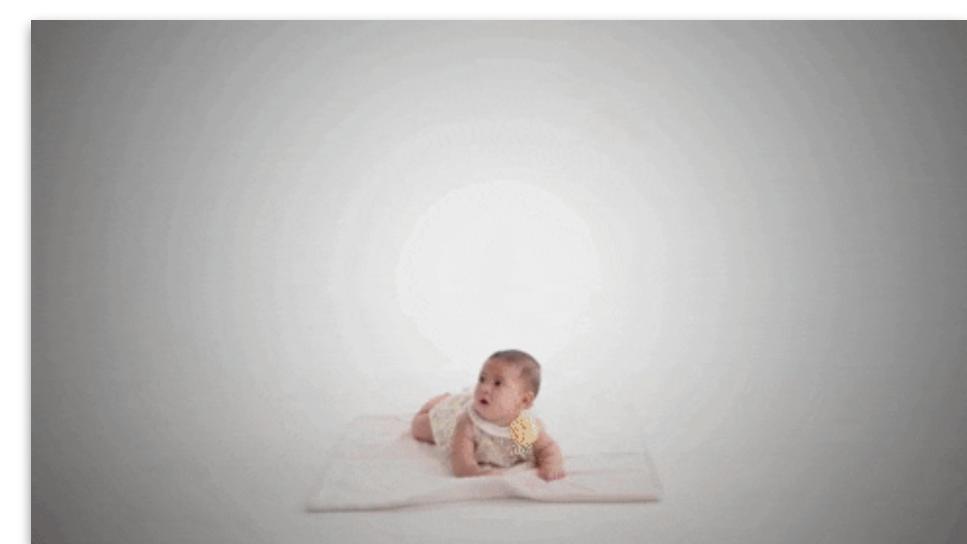
Weak AI



Strong AI



Humans (ref.)

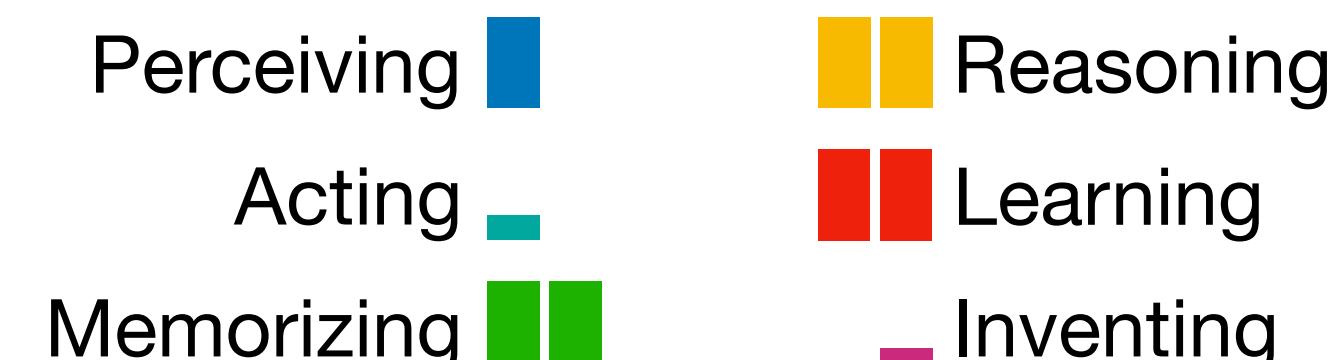
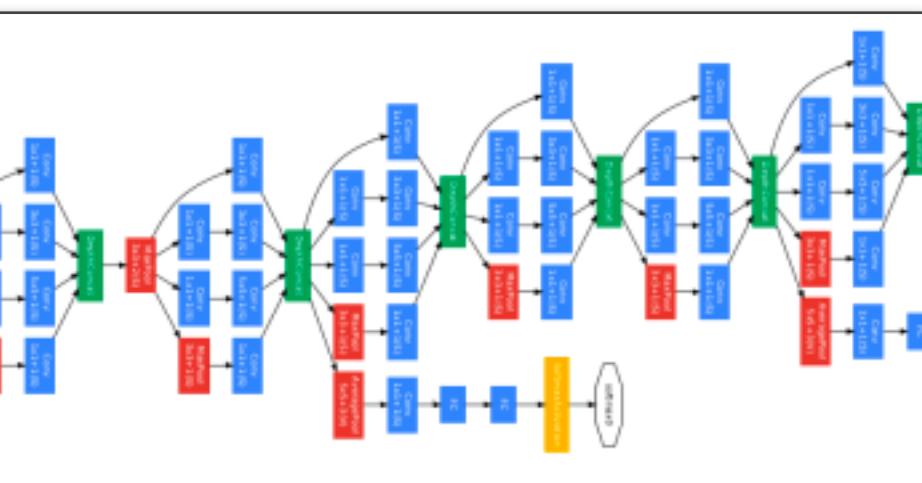
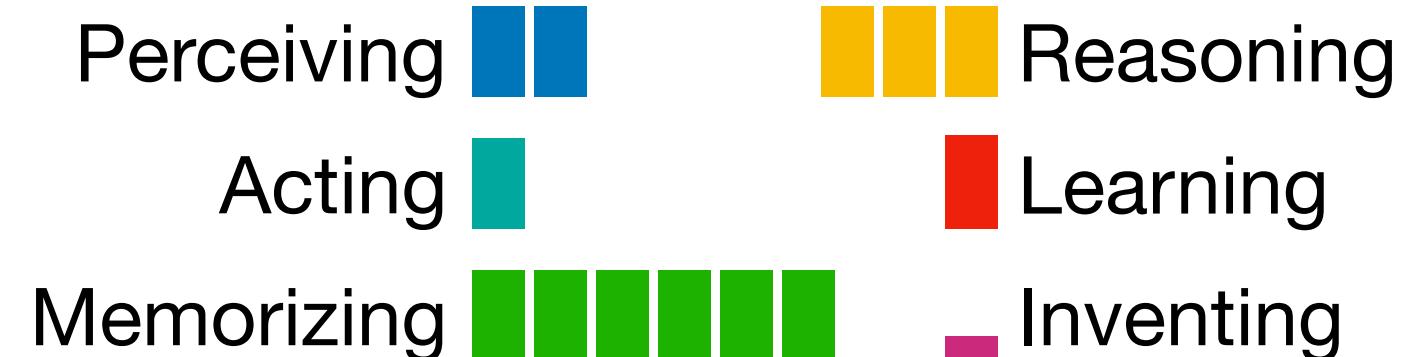
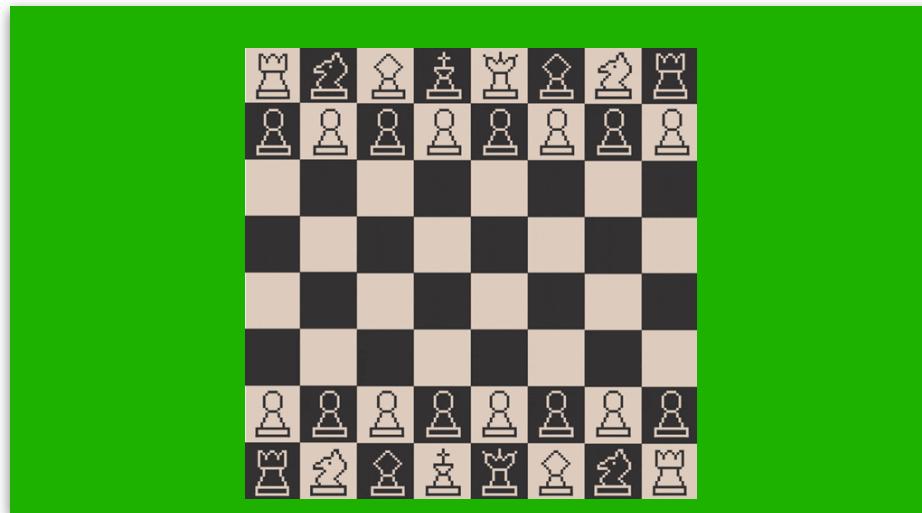


Artificial Intelligence (AI)

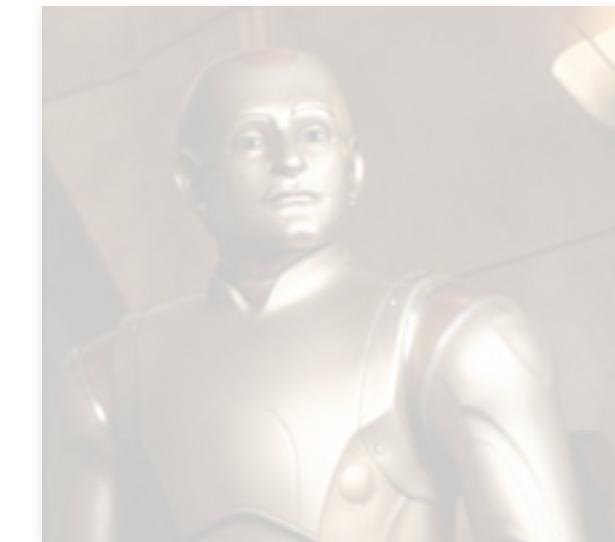
Weak AI



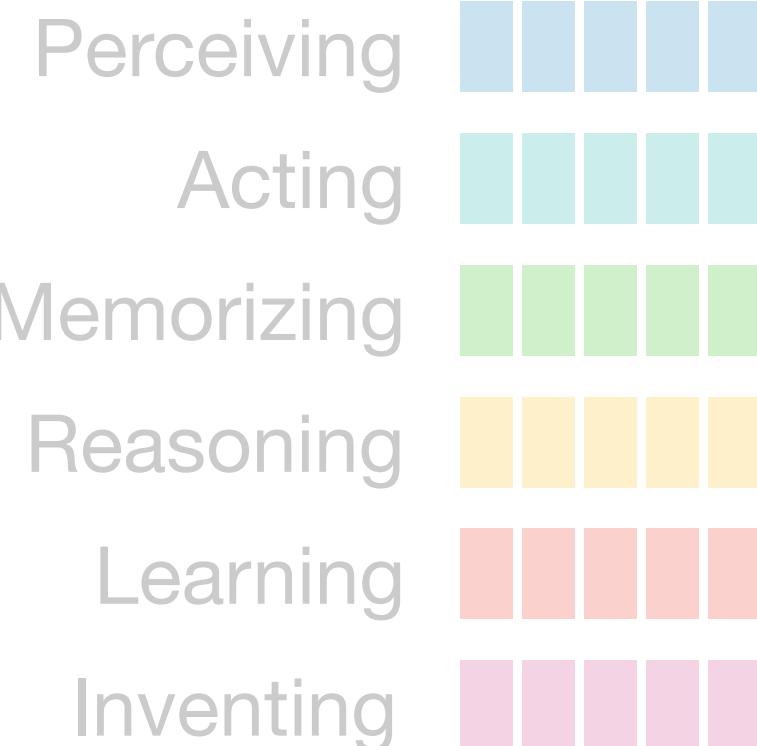
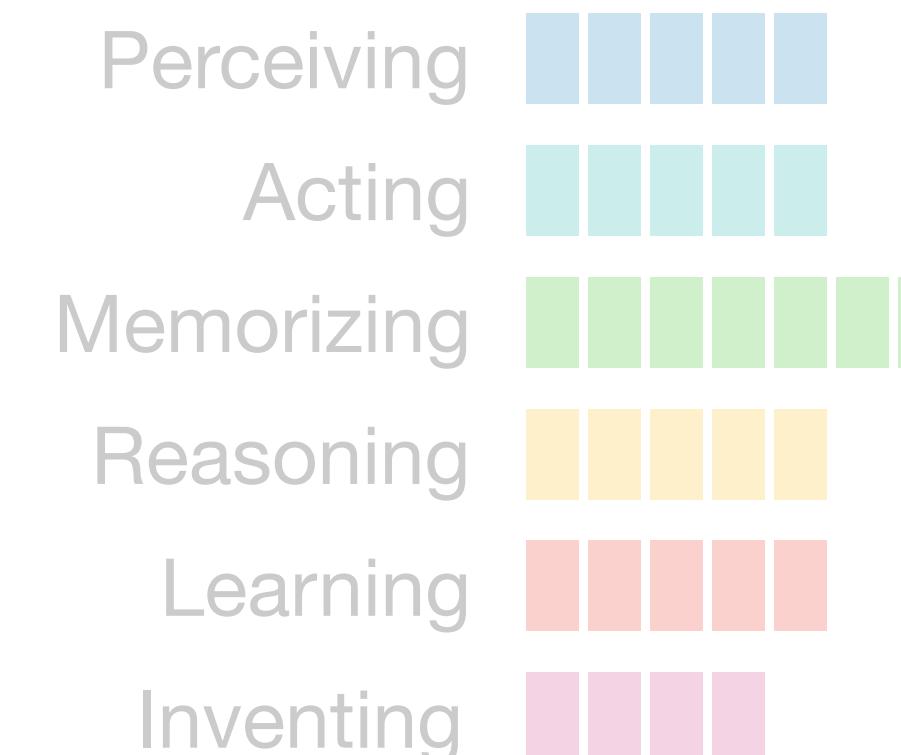
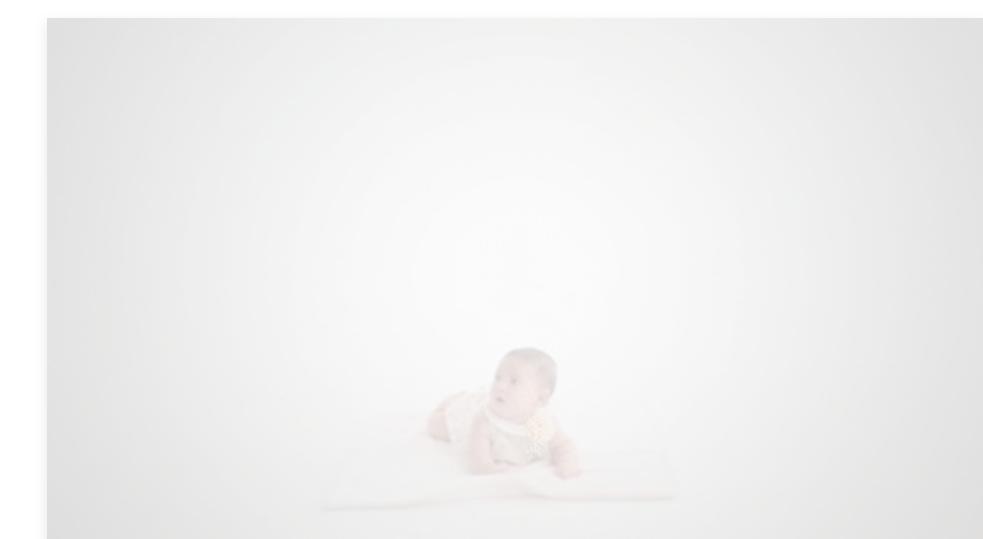
We are here!



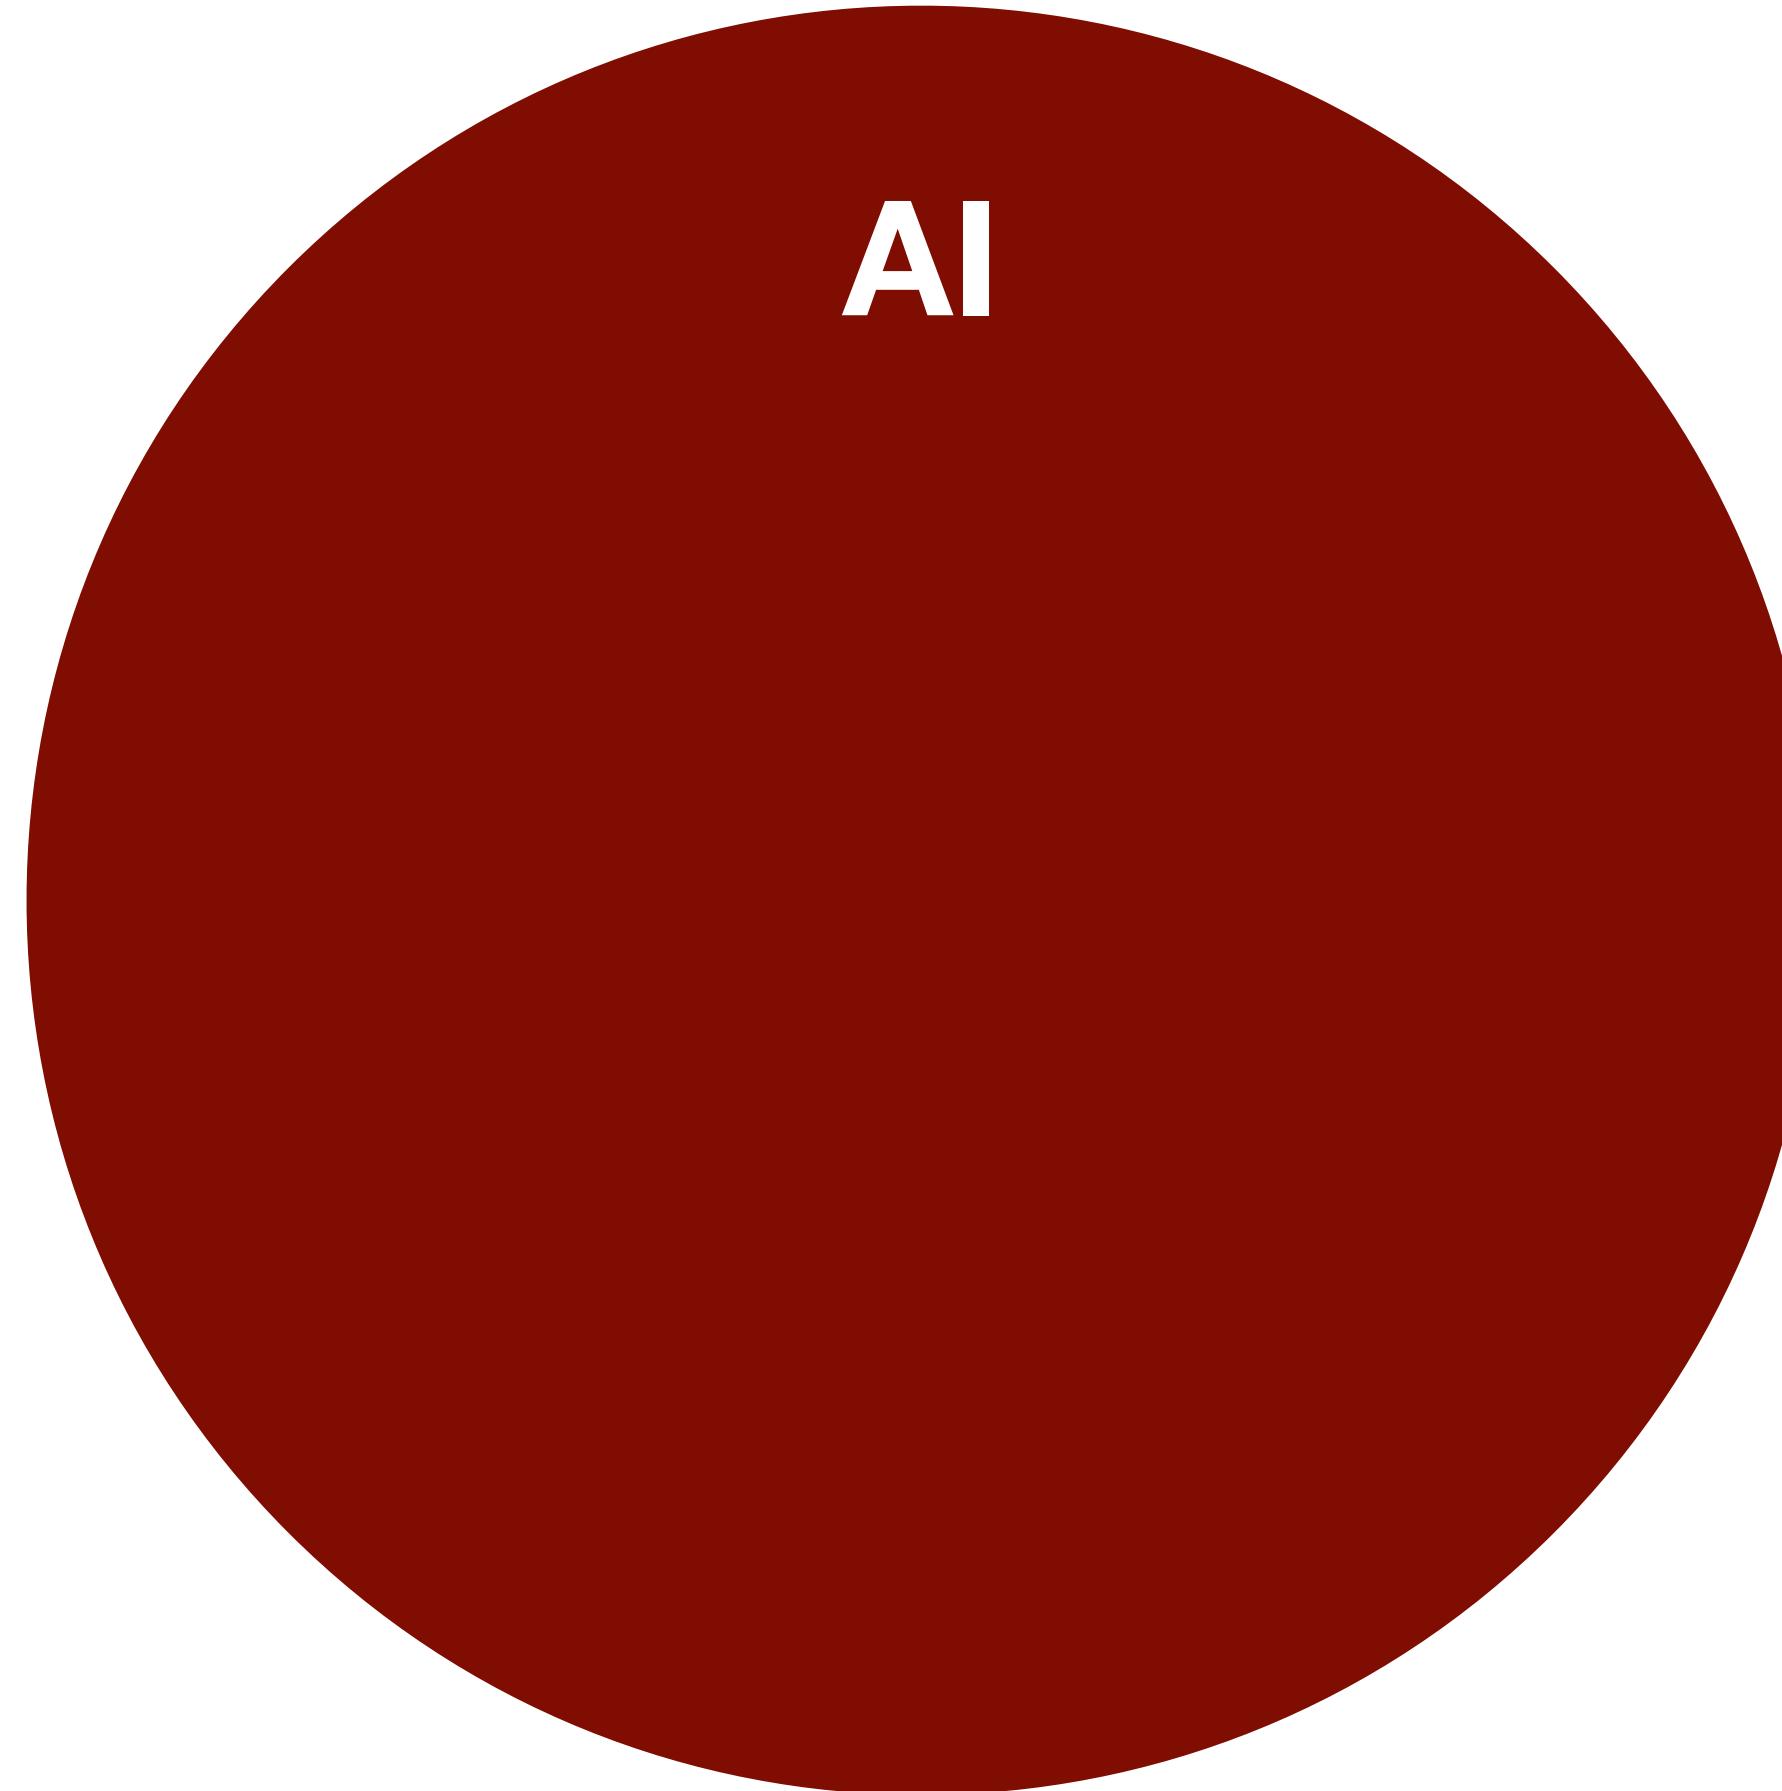
Strong AI



Humans (ref.)



Artificial Intelligence (AI)



Symbolic Logic

Fuzzy Logic

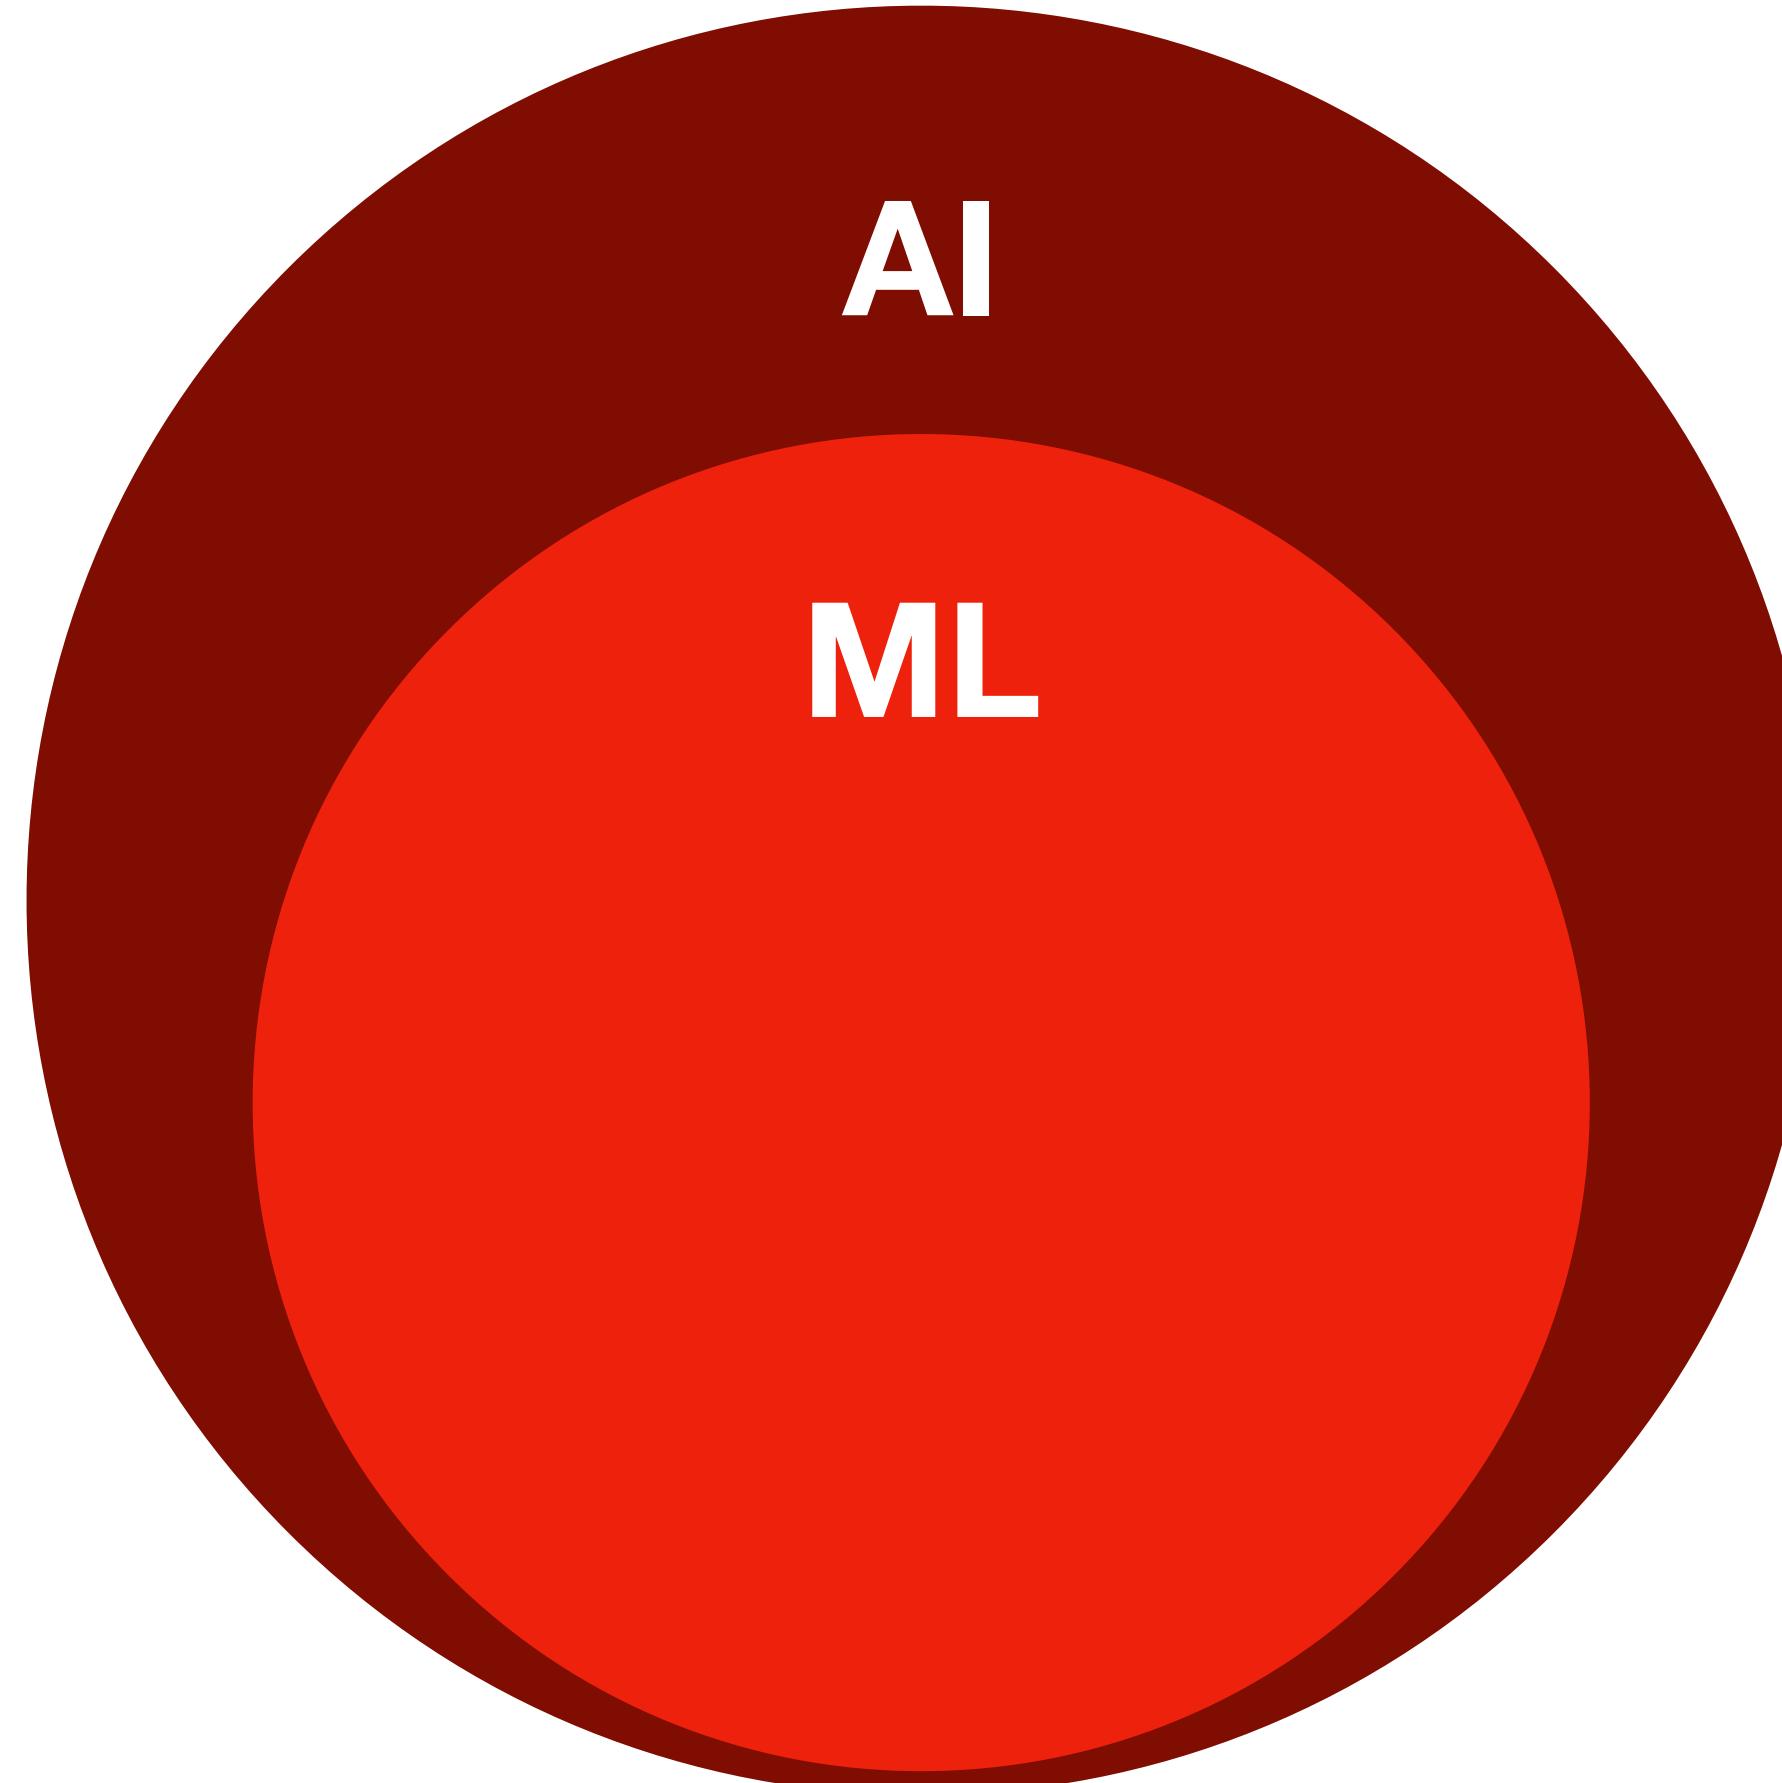
Semantic Networks

Evolutionary Models

Swarm Intelligence

Machine Learning

Artificial Intelligence (AI)



Symbolic Logic

Fuzzy Logic

Semantic Networks

Evolutionary Models

Swarm Intelligence

Machine Learning

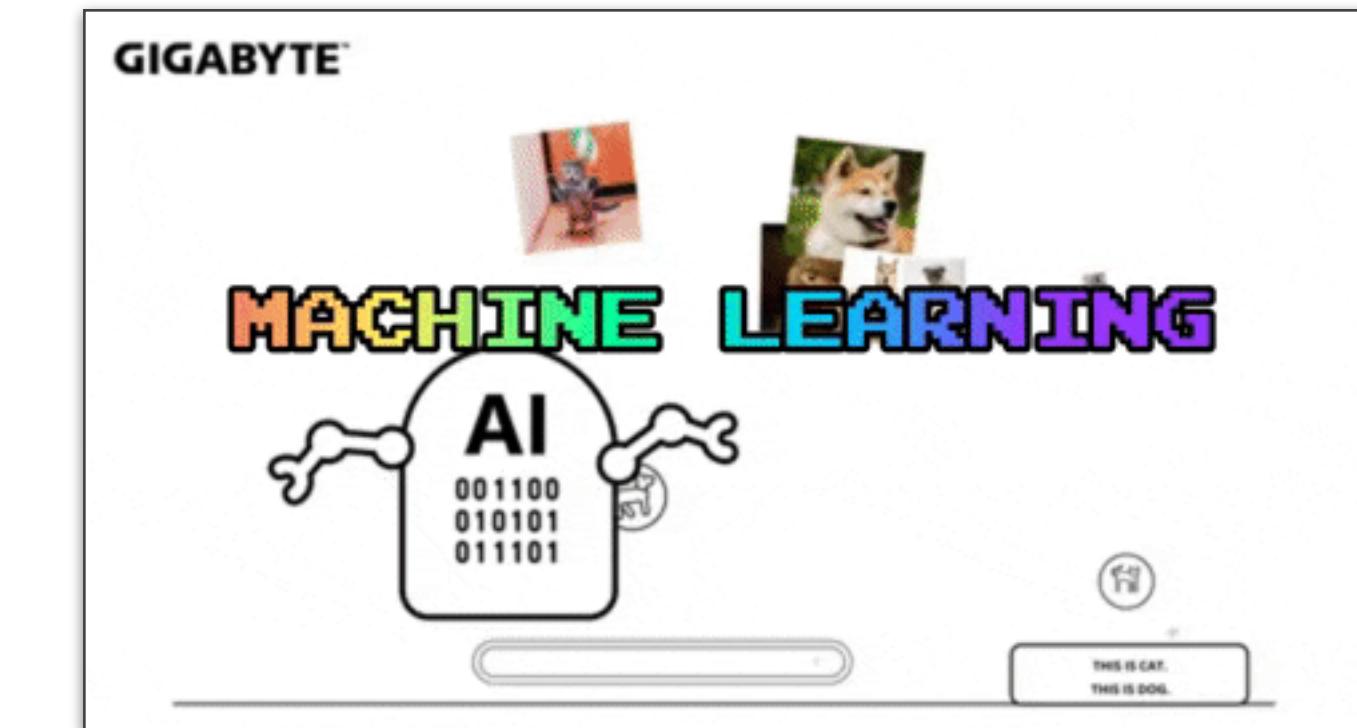
We are
here!

Machine Learning (ML)

What comes to your mind?



<https://bit.ly/3xbkegQ>



Machine Learning (ML)

It is AI

It aims at developing computer systems that mimic humans' intelligence.

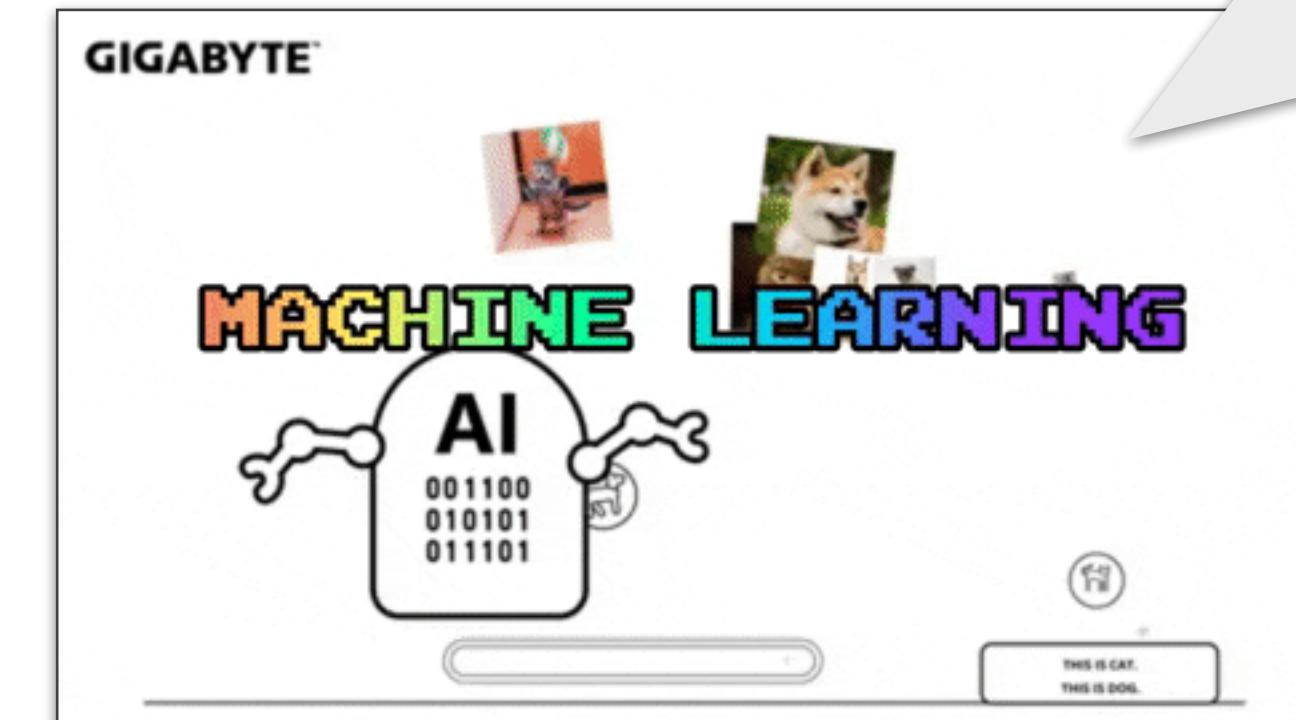
More Specifically

It aims at developing computer systems that mimic the learning-ability of humans.

It is
data-driven.

In Practice

Leverage data examples to improve the performance on a target task.



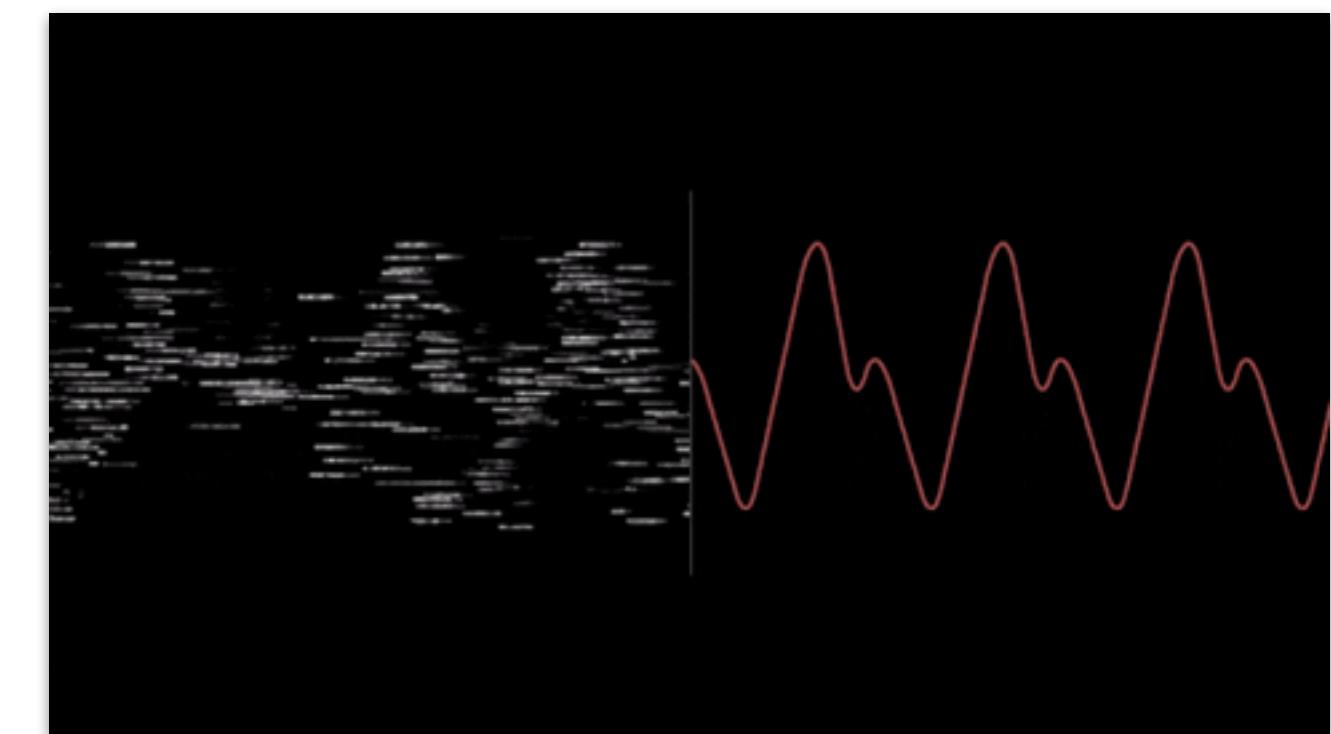
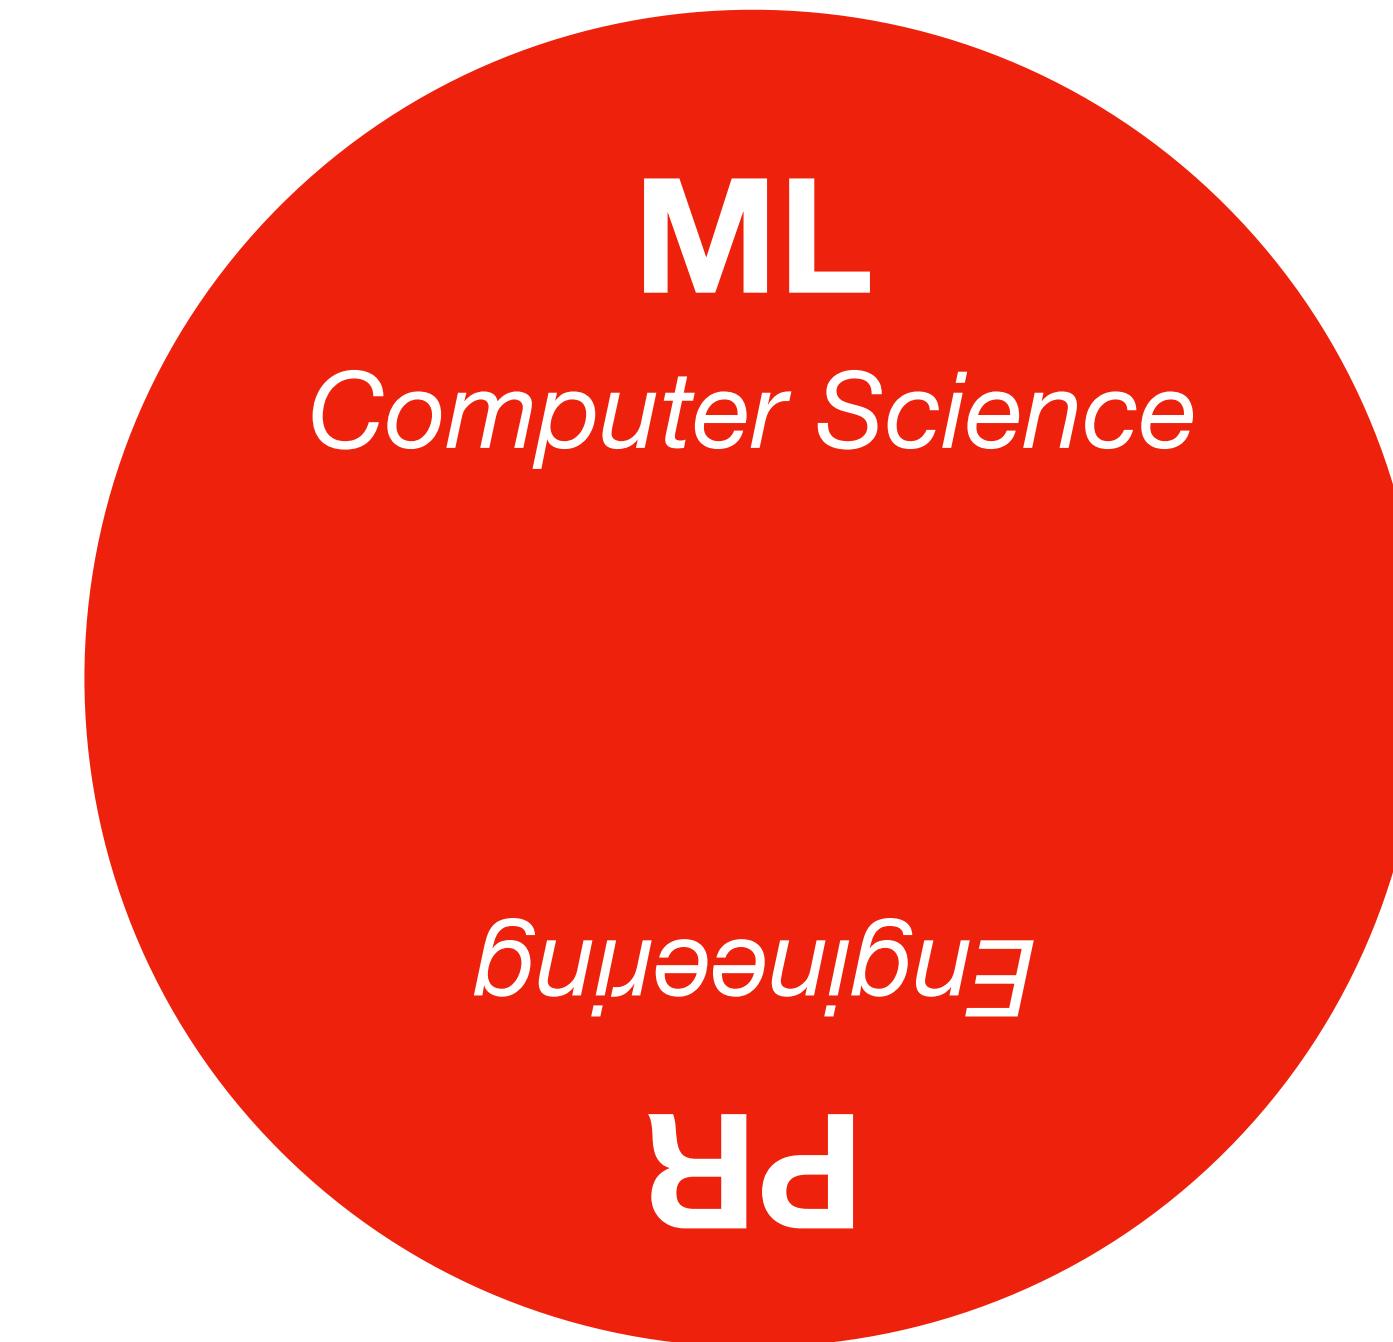
Machine Learning (ML)

ML and Pattern Recognition (PR)

Same field.



Cluster or label data.



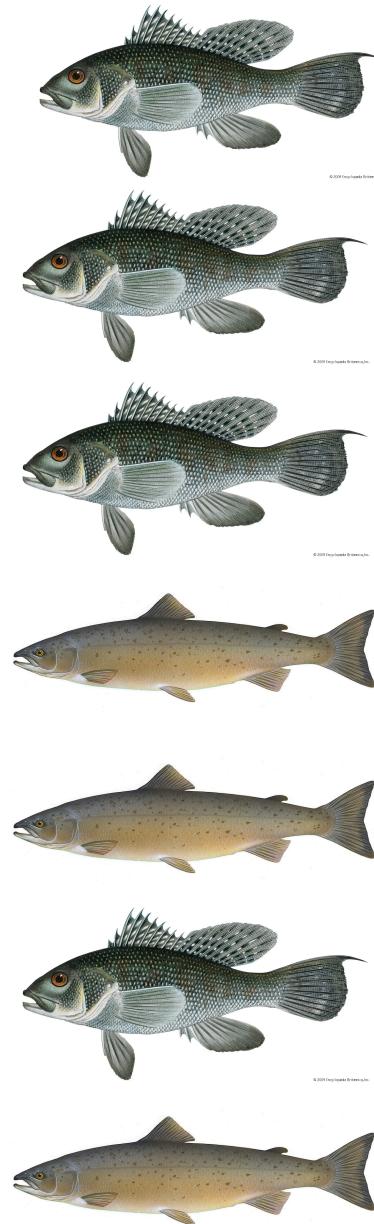
Find patterns on data.

Machine Learning (ML)

What data (structure) are we typically talking about?

Example problem: fish classification

Sea Bass or
Salmon?

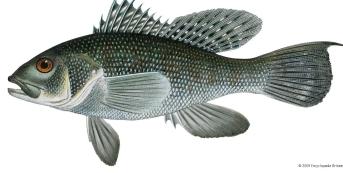
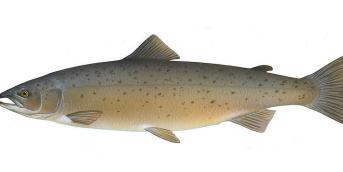
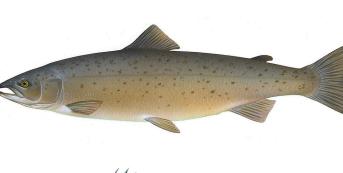
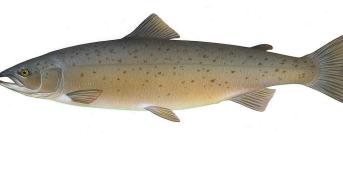


LOYOLA
UNIVERSITY CHICAGO

Machine Learning (ML)

What data (structure) are we typically talking about?

Example problem: fish classification

Sea Bass or Salmon?	Fish	Length	Weight	Width	Fin
		1.12	14.11	0.31	3.0
		0.95	11.22	0.28	3.0
		1.08	12.02	0.31	3.0
		1.45	45.03	0.37	2.0
		1.09	31.01	0.38	2.0
		1.08	11.09	0.29	2.0
		1.51	46.00	0.37	2.0

Machine Learning (ML)

What data (structure) are we typically talking about?

Example problem: fish classification

Sea Bass or
Salmon?

Fish
- length: float
- weight: float
- width: float
- fin_count: int



Length	Weight	Width	Fin	
1.12	14.11	0.31	3.0	
0.95	11.22	0.28	3.0	
1.08	12.02	0.31	3.0	
1.45	45.03	0.37	2.0	
1.09	31.01	0.38	2.0	
1.08	11.09	0.29	2.0	
1.51	46.00	0.37	2.0	

Each data sample is a **Tensor**.

Machine Learning (ML)

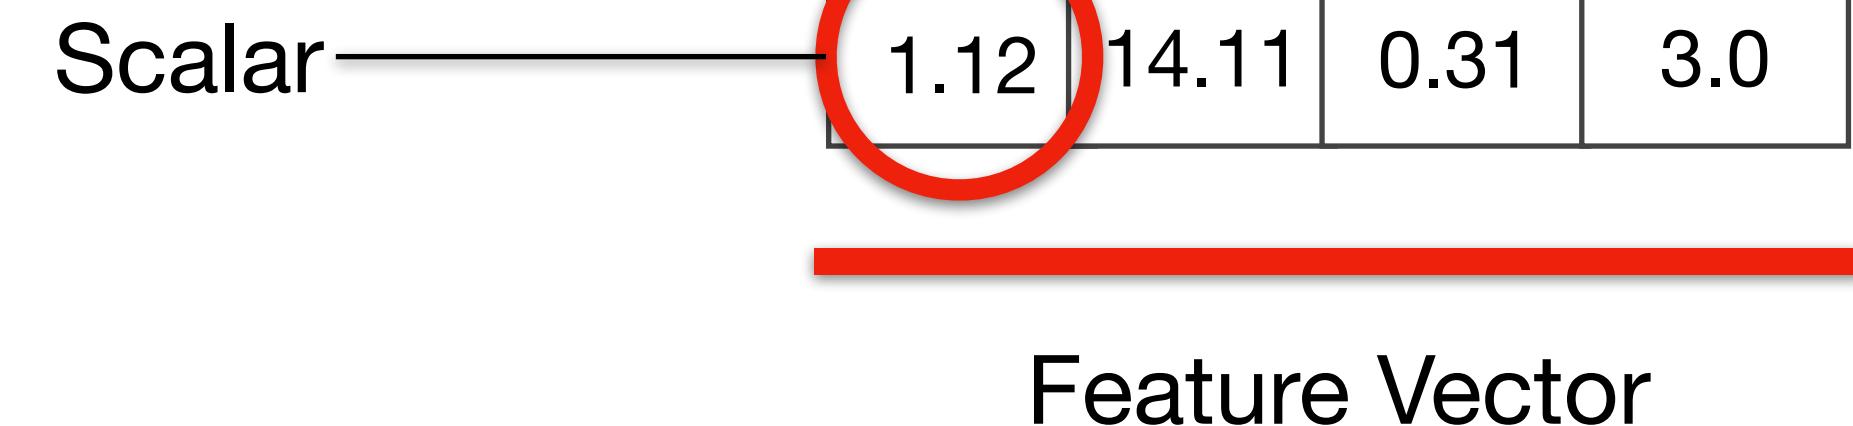
What data (structure) are we typically talking about?

0-order: **Scalar**

1st-order tensor: **Vector**
(feature vectors)

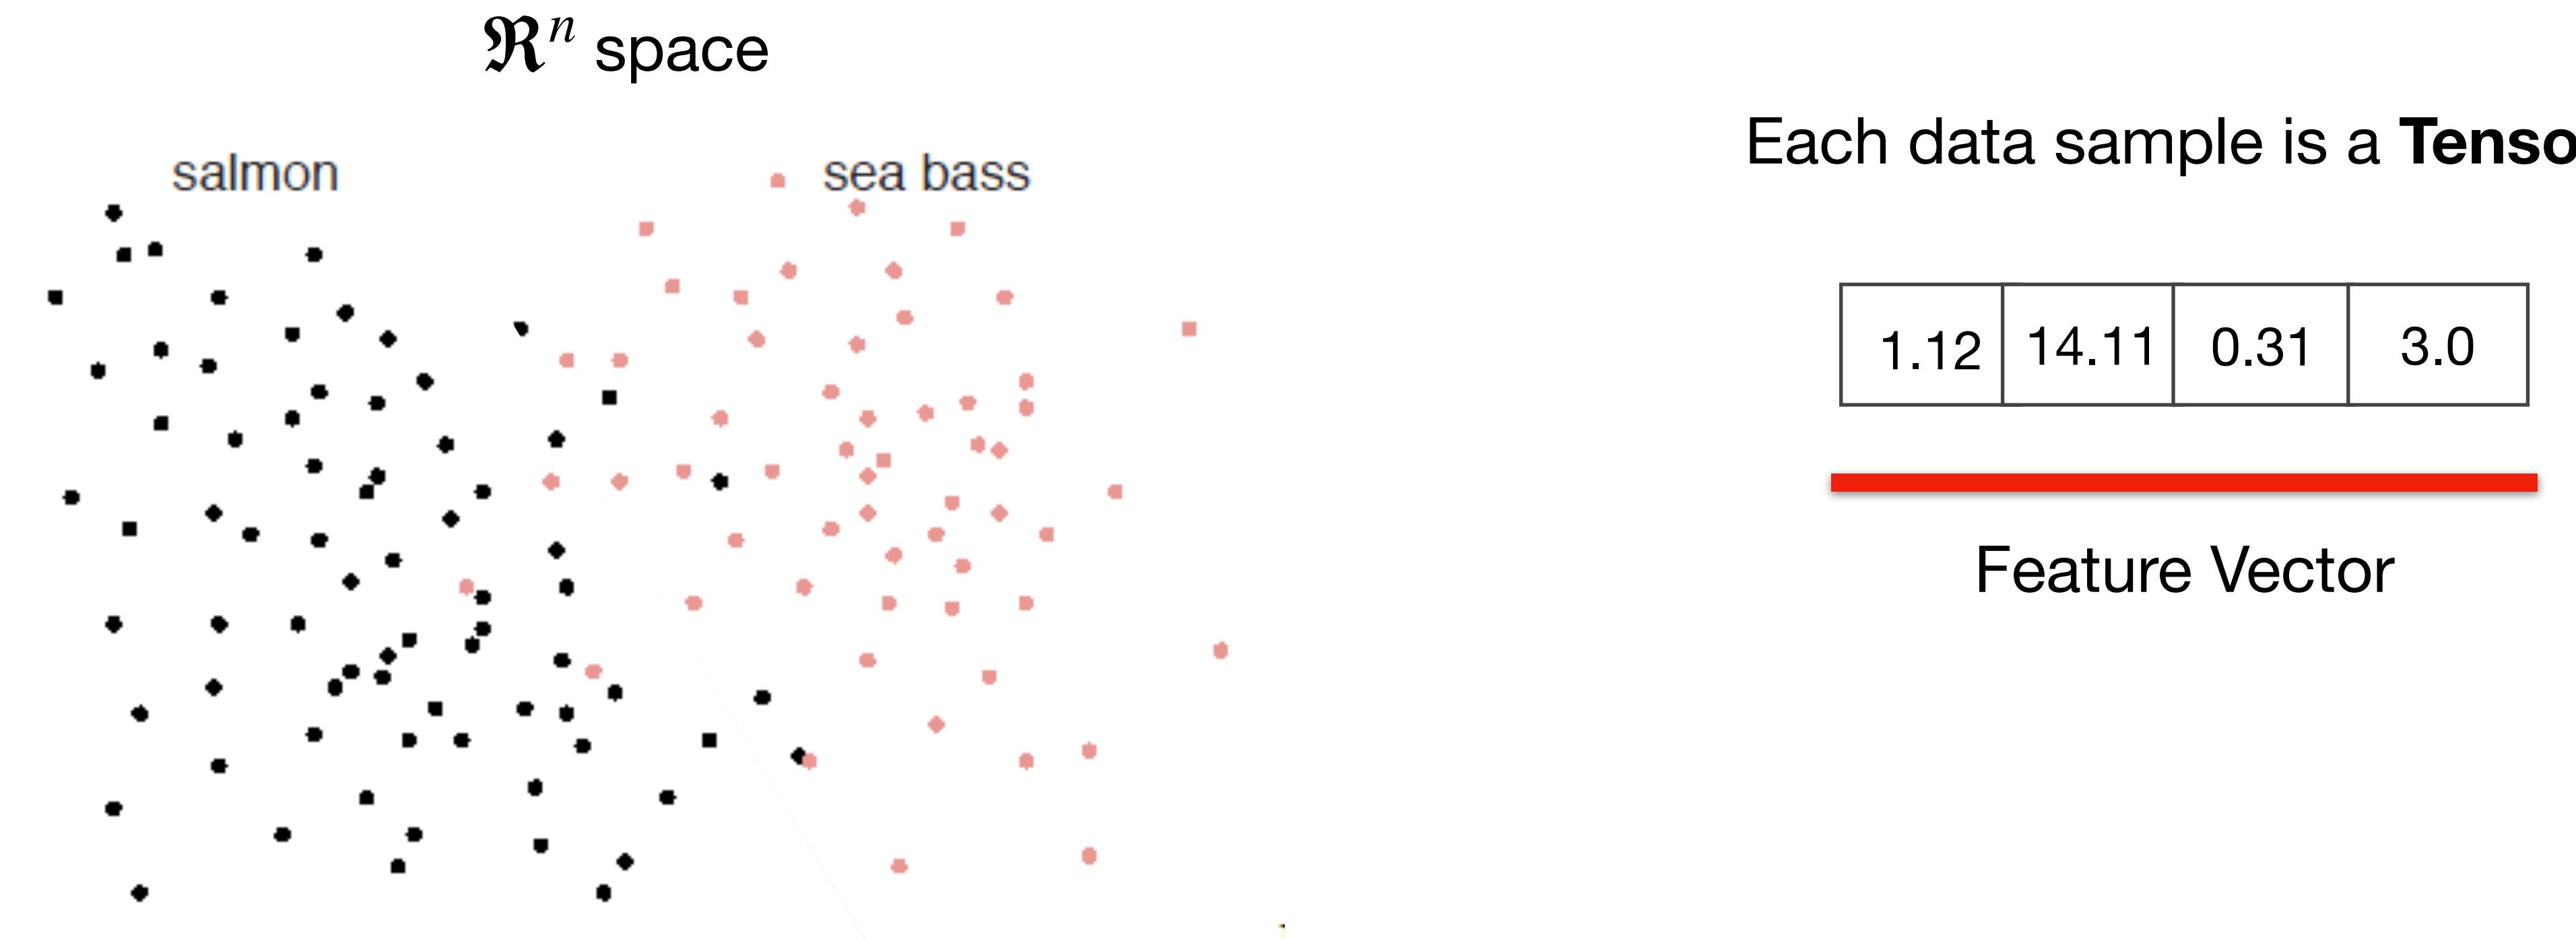
2nd-order tensor: **Matrix**
(images)

Each data sample is a **Tensor**.



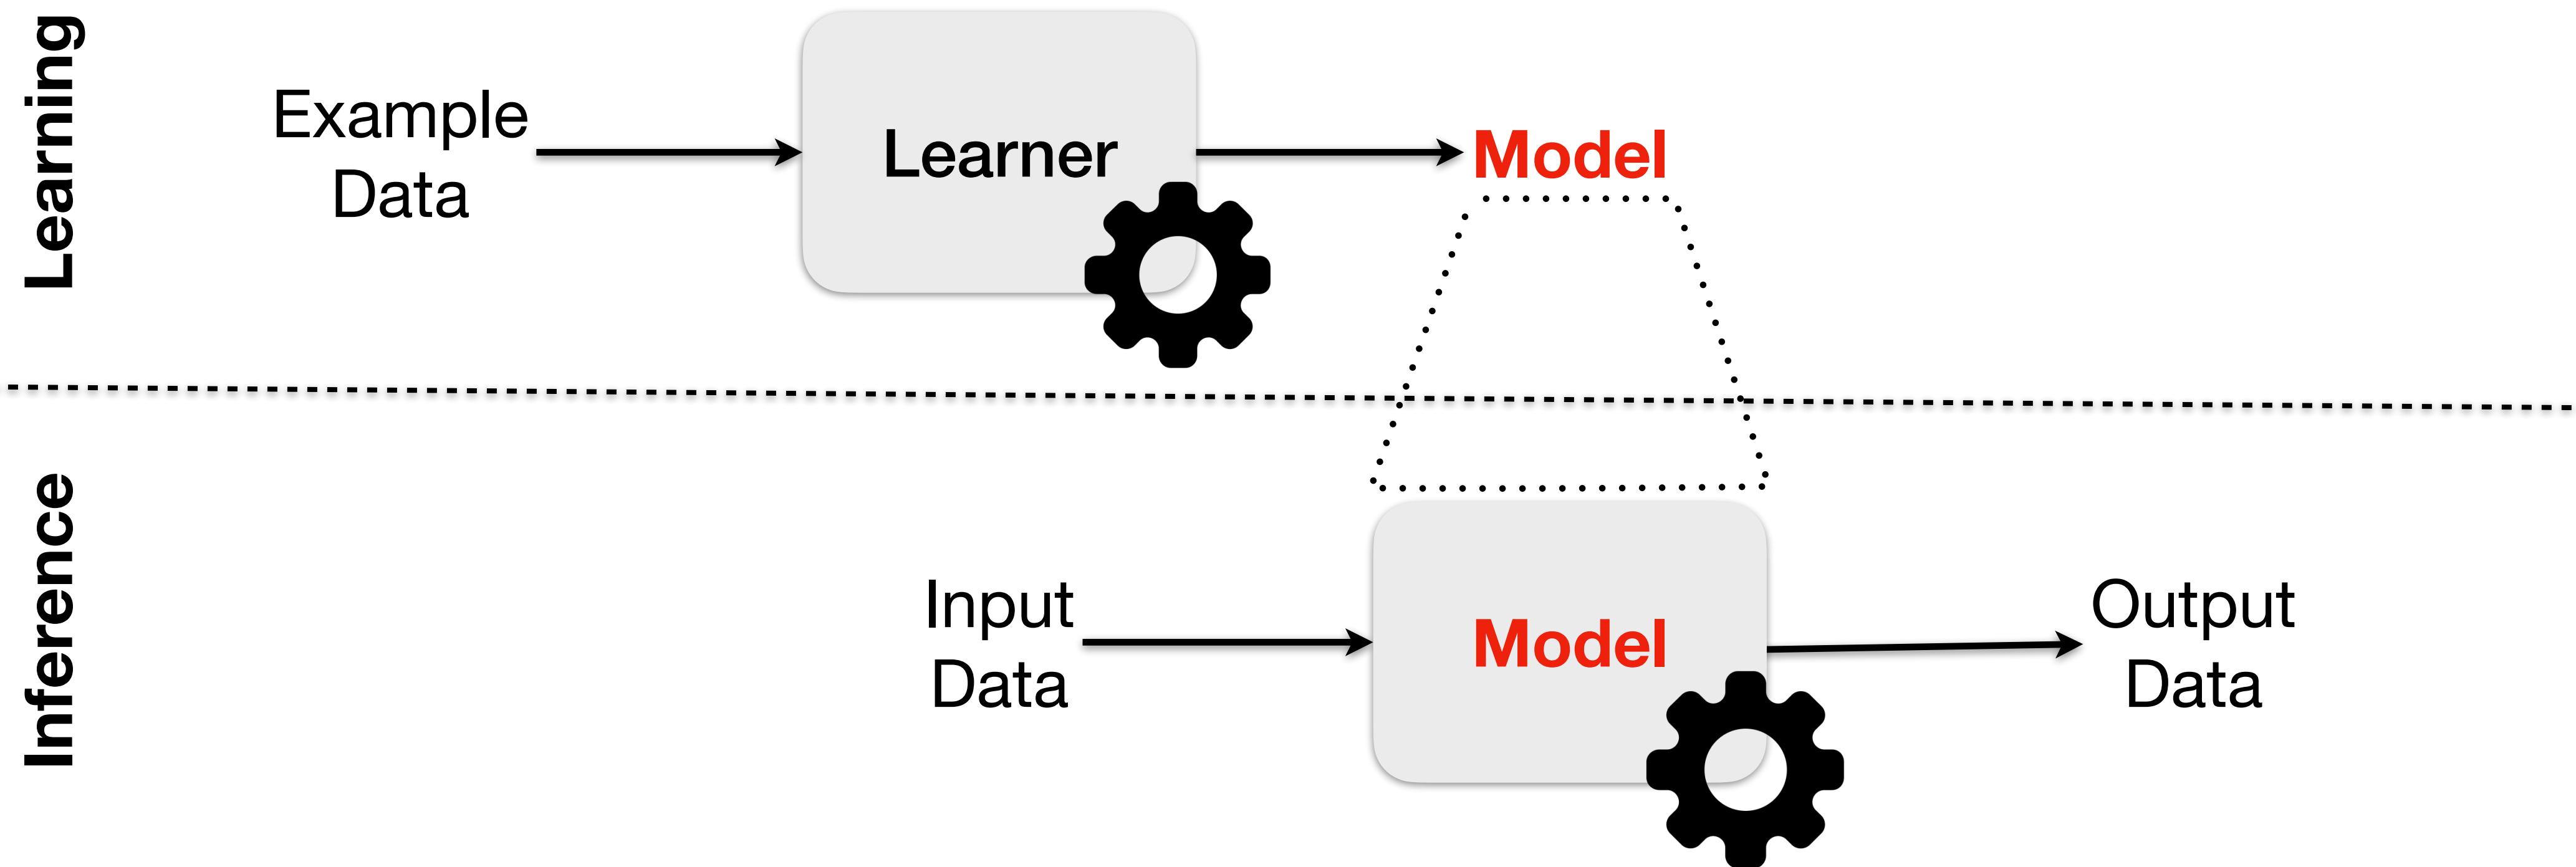
Machine Learning (ML)

What data (structure) are we typically talking about?



Machine Learning (ML)

ML Stages

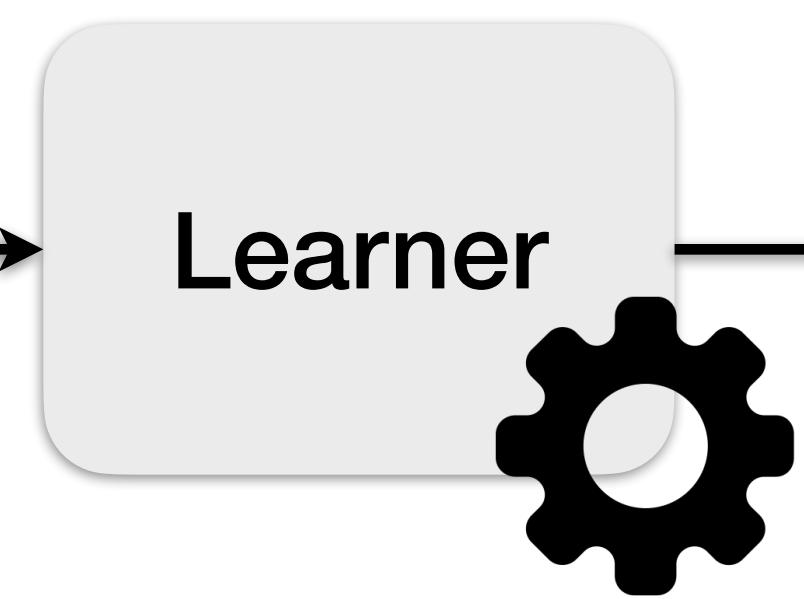


Machine Learning (ML)

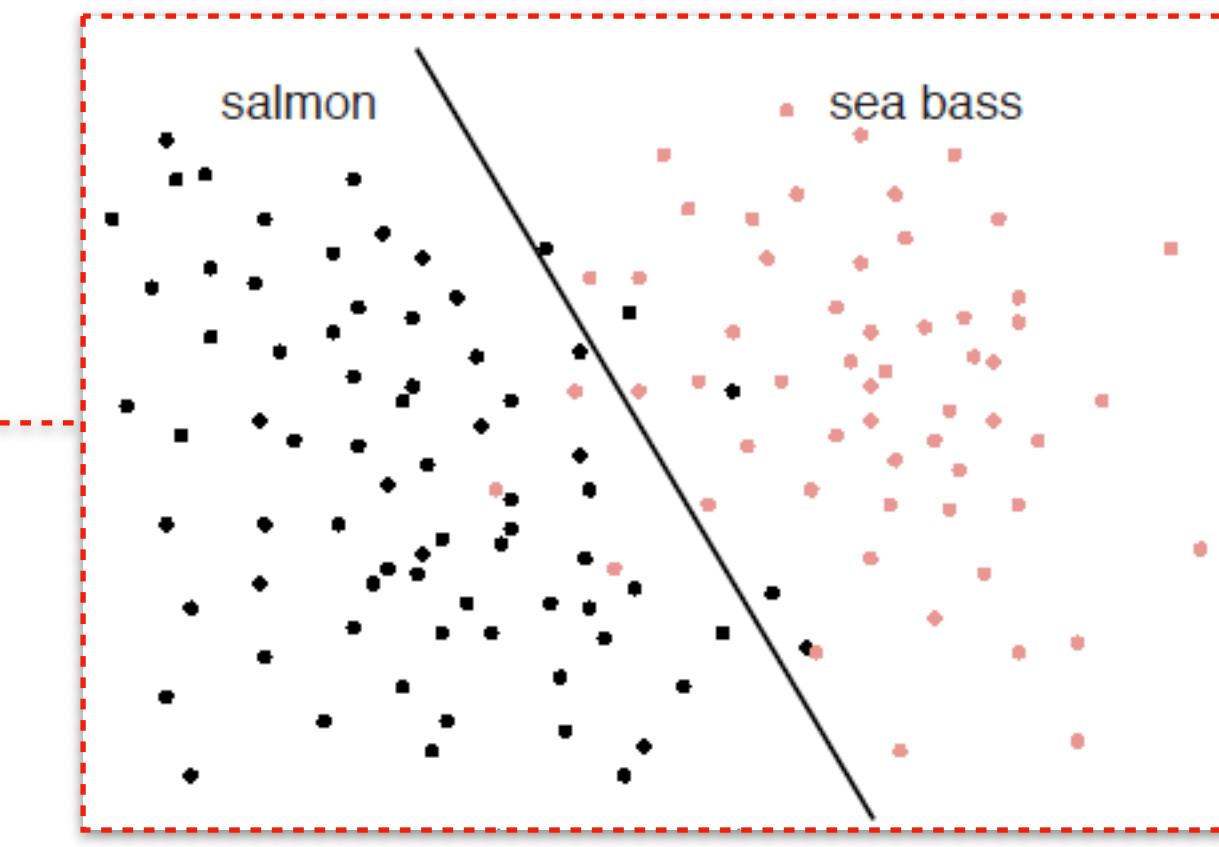
ML Stages

Learning

Example Data



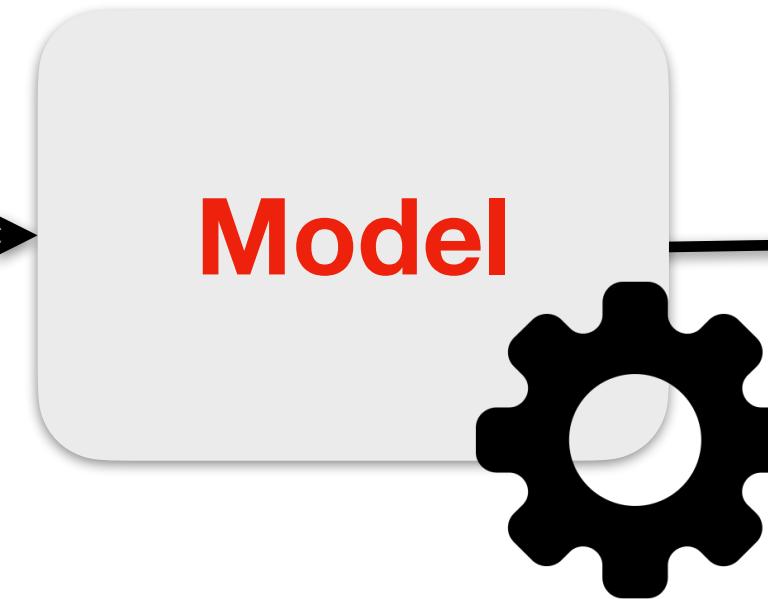
Model



E.g., separation hyperplane.

Inference

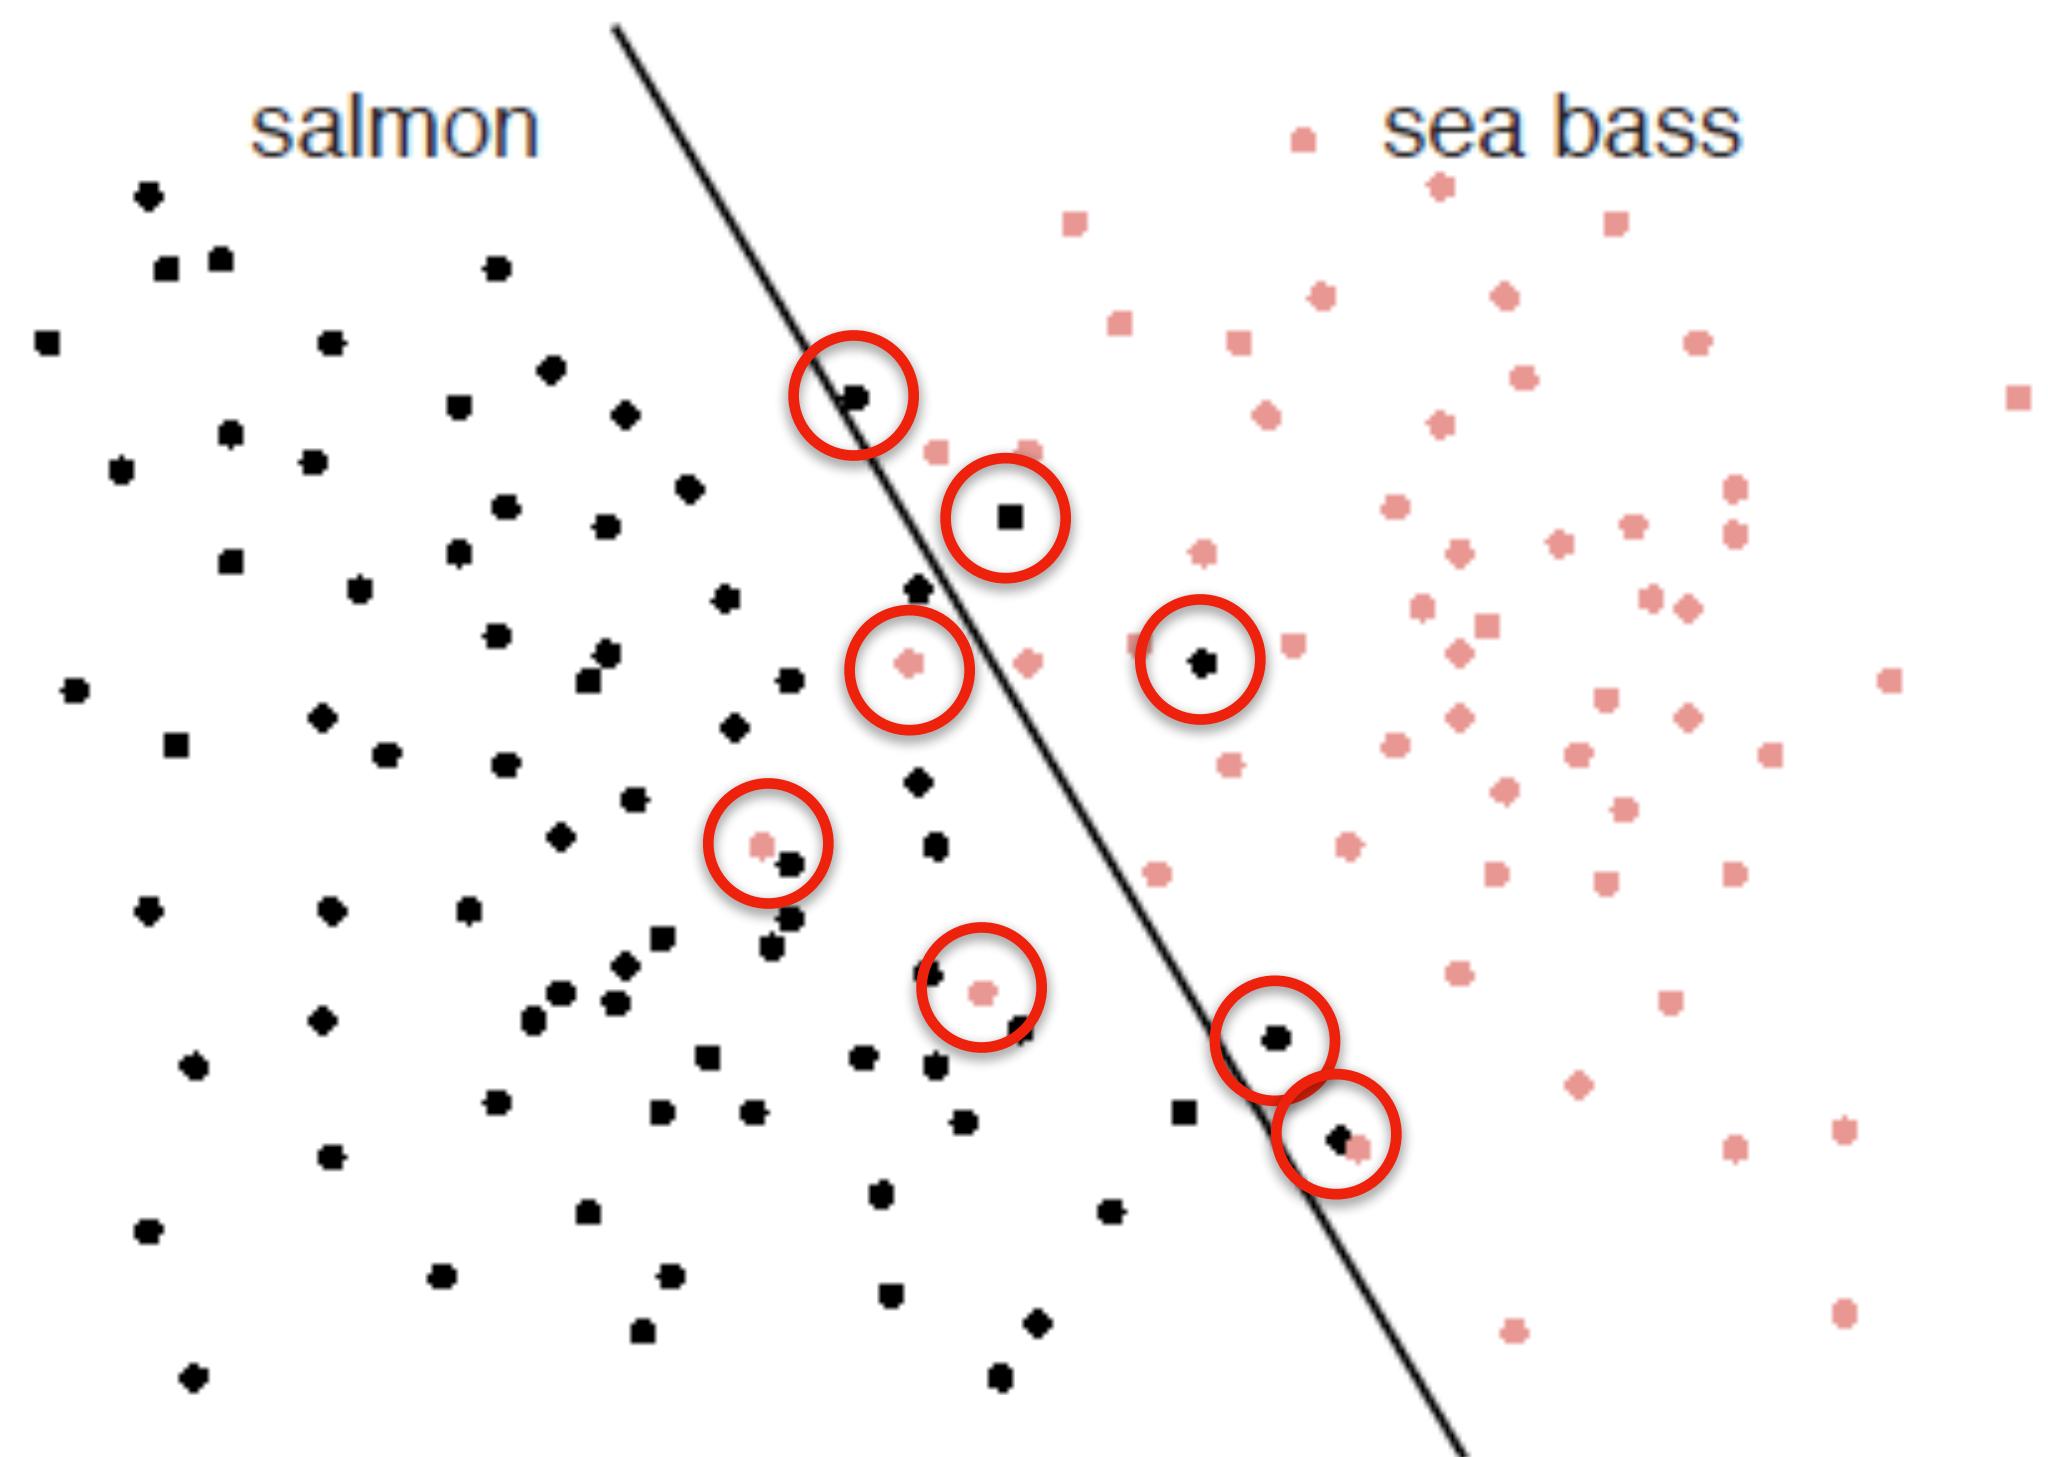
Input Data



LOYOLA
UNIVERSITY CHICAGO

Machine Learning (ML)

Metrics - How to measure the performance of the model?



Actual versus Predicted Labels

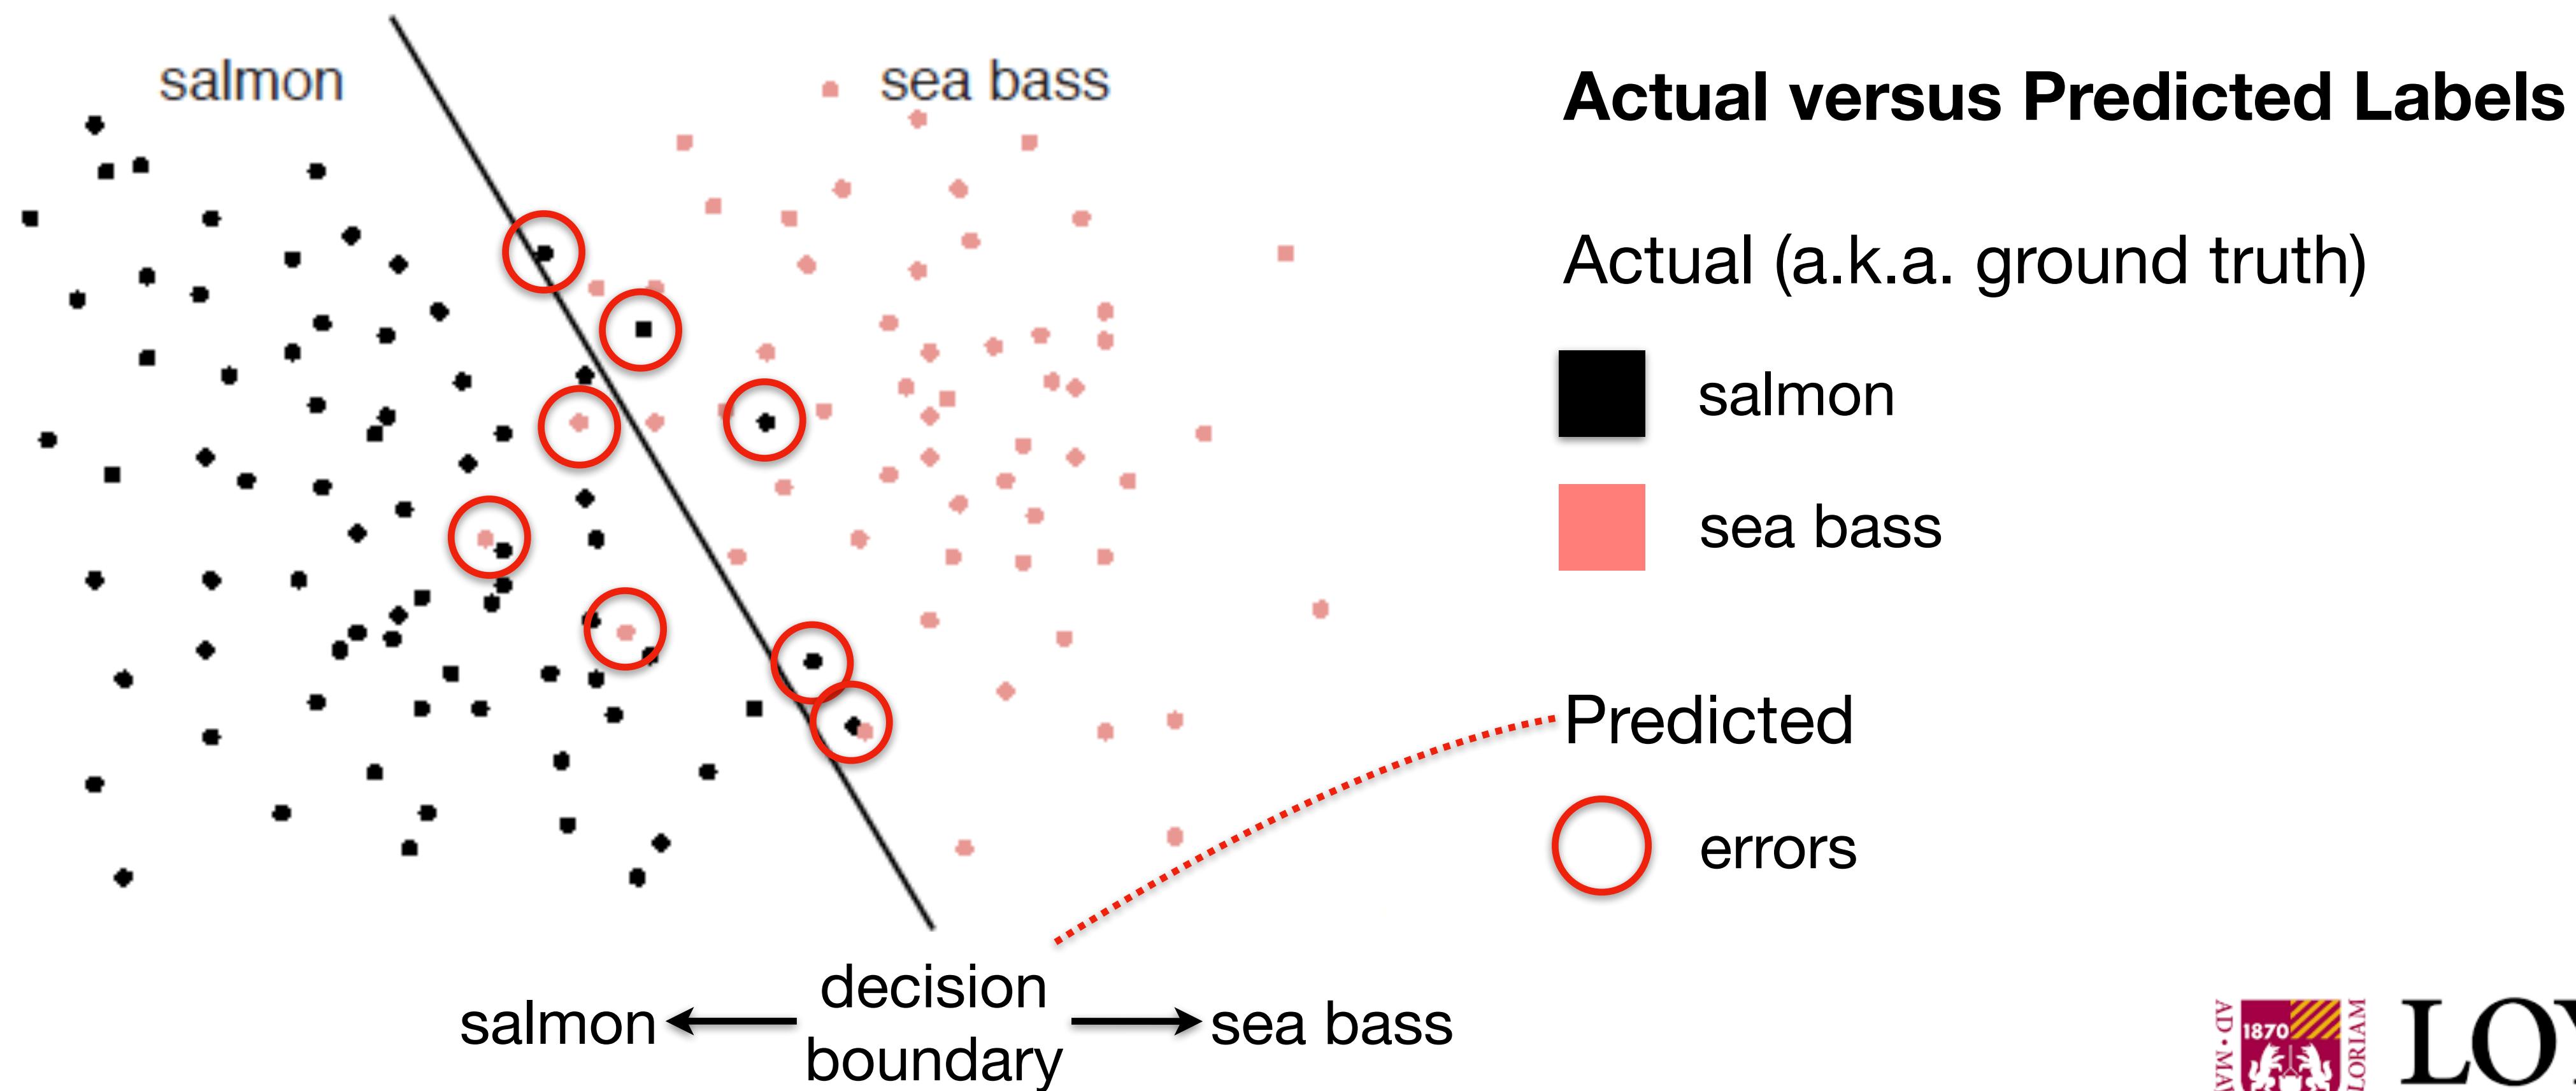
Actual (a.k.a. ground truth)

■ salmon

■ sea bass

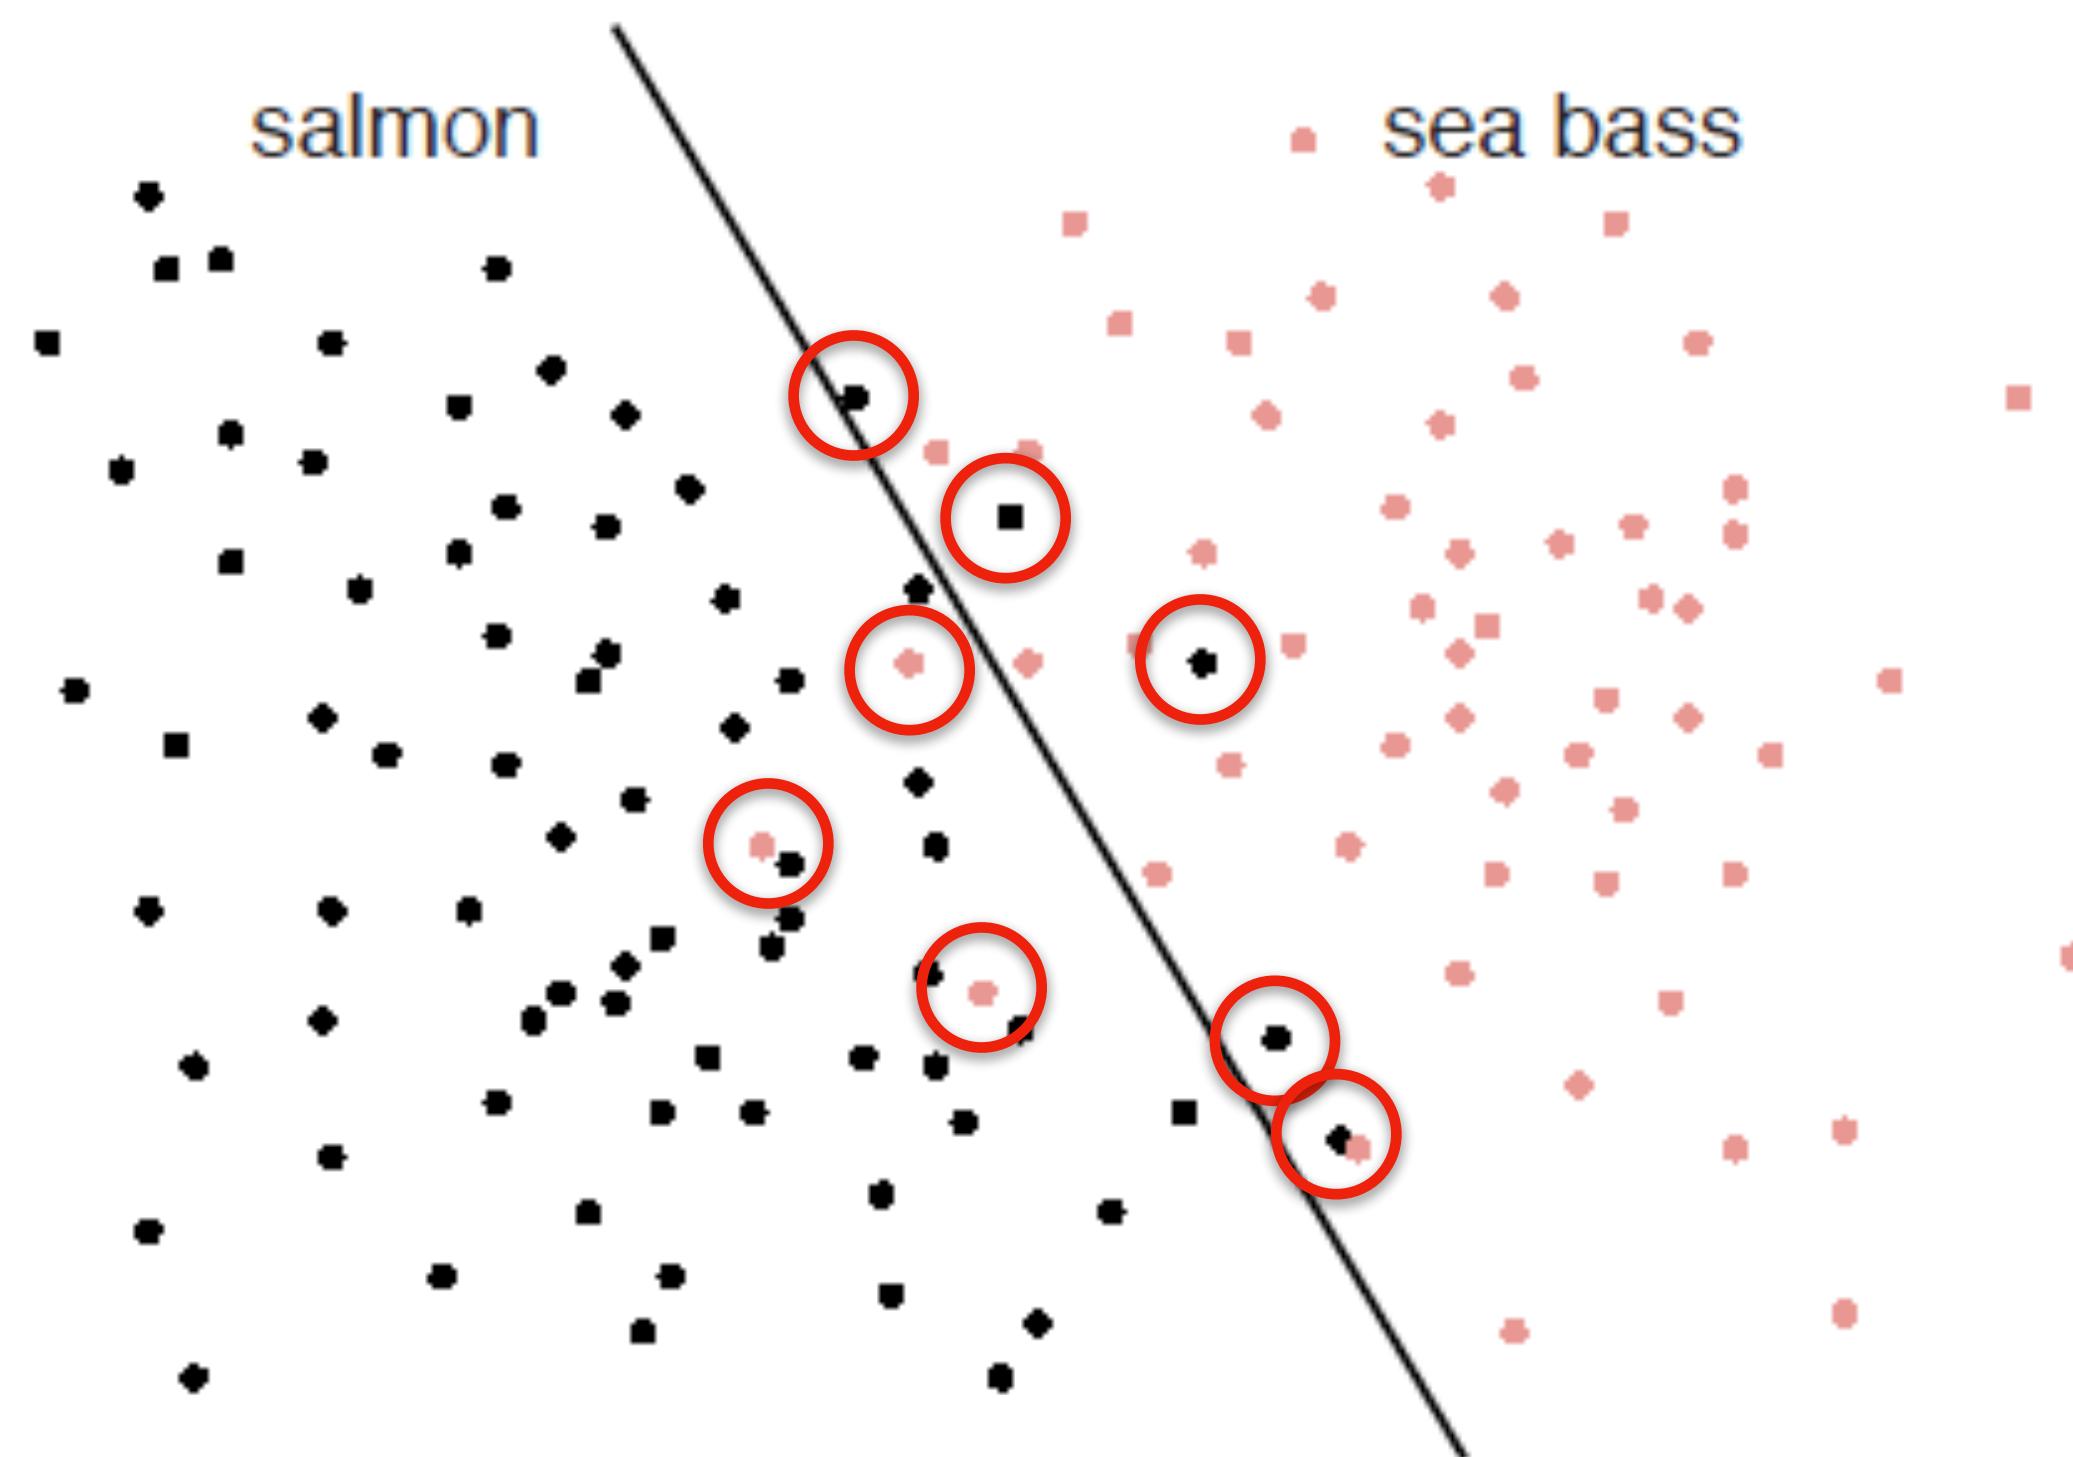
Machine Learning (ML)

Metrics - How to measure the performance of the model?



Machine Learning (ML)

Metrics

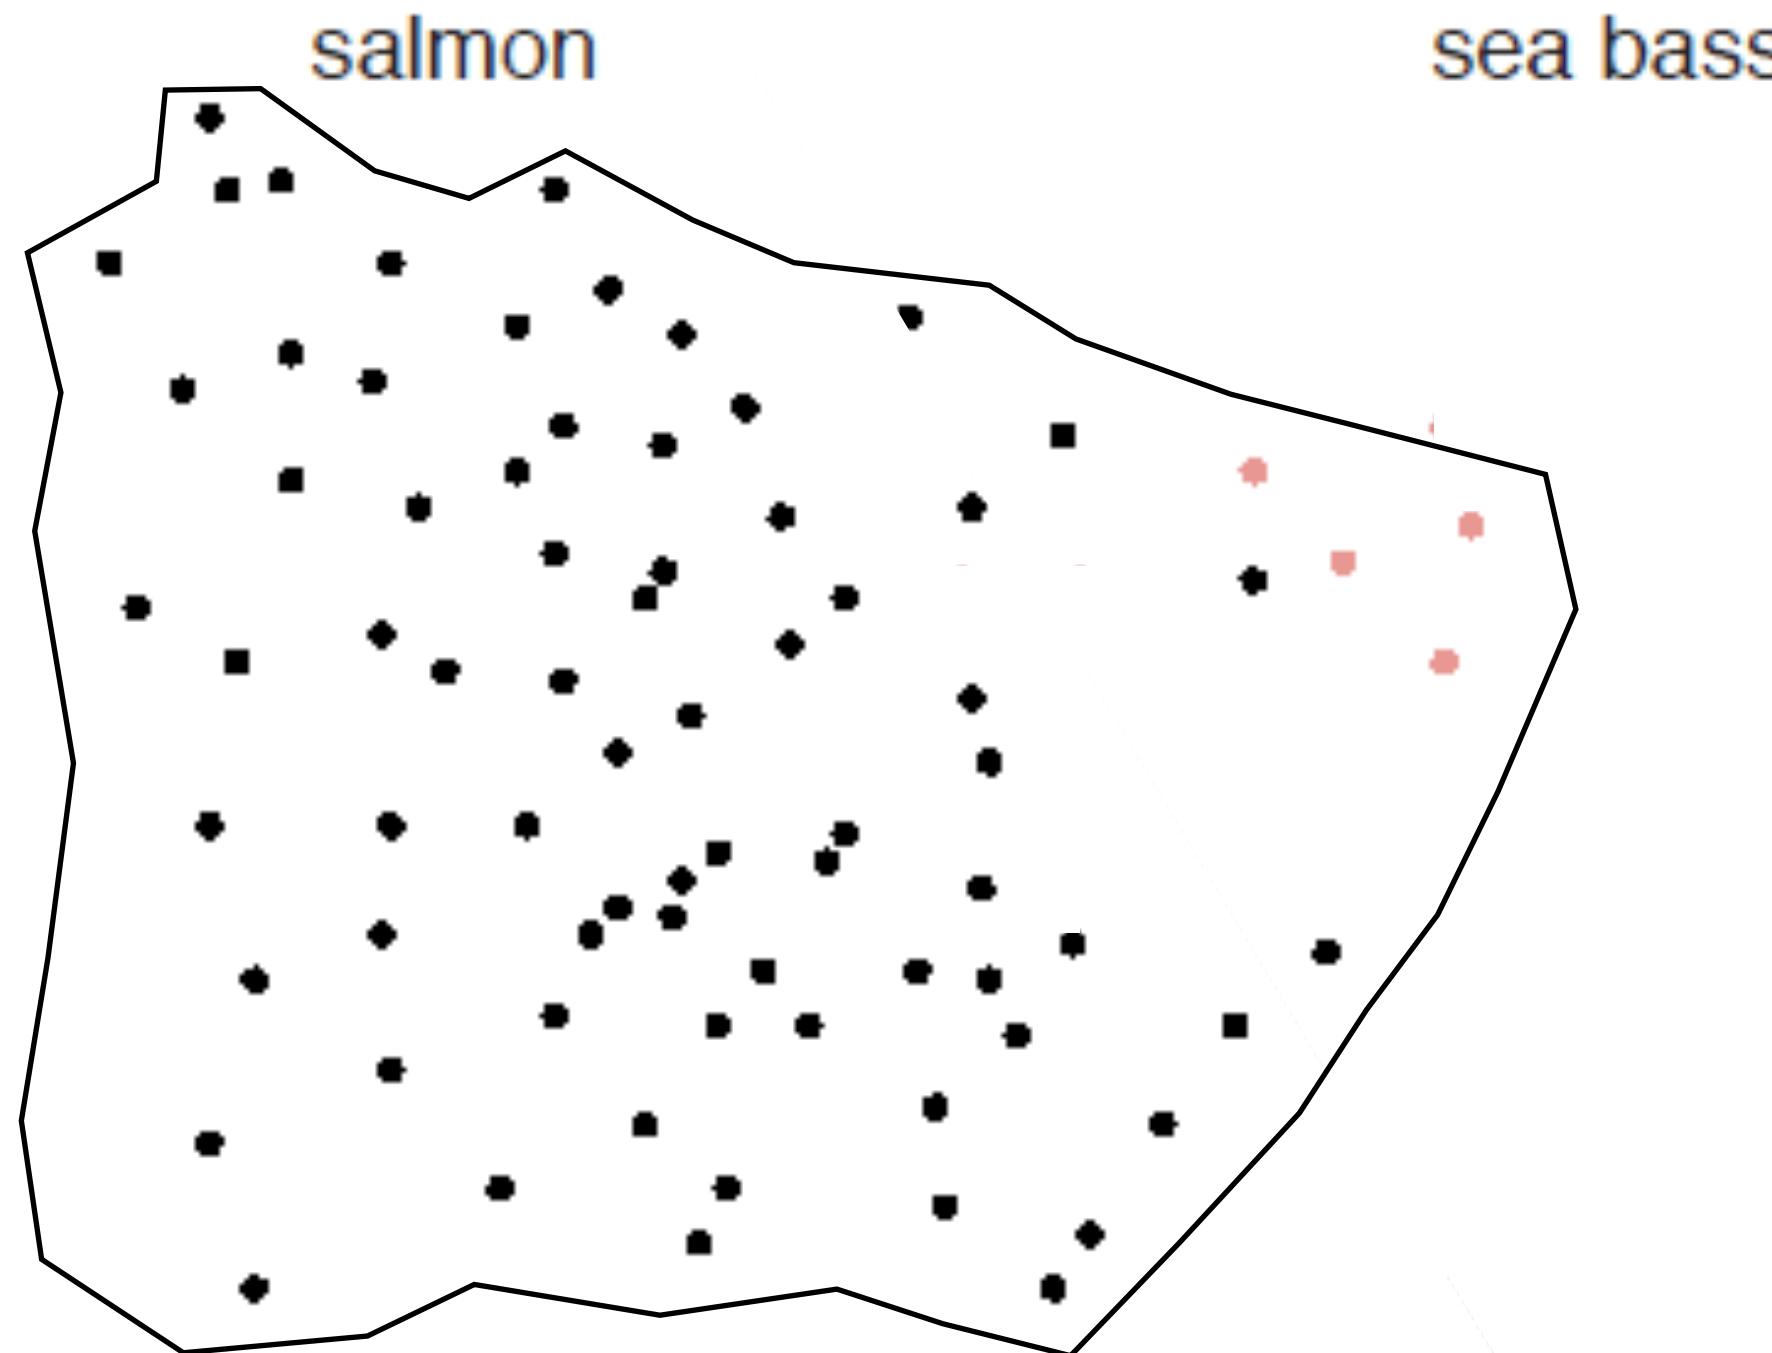


Accuracy (Acc)

$$Acc = \frac{Correct\ Predictions}{Total\ Predictions}$$

Machine Learning (ML)

Metrics



Example model: everything is salmon!

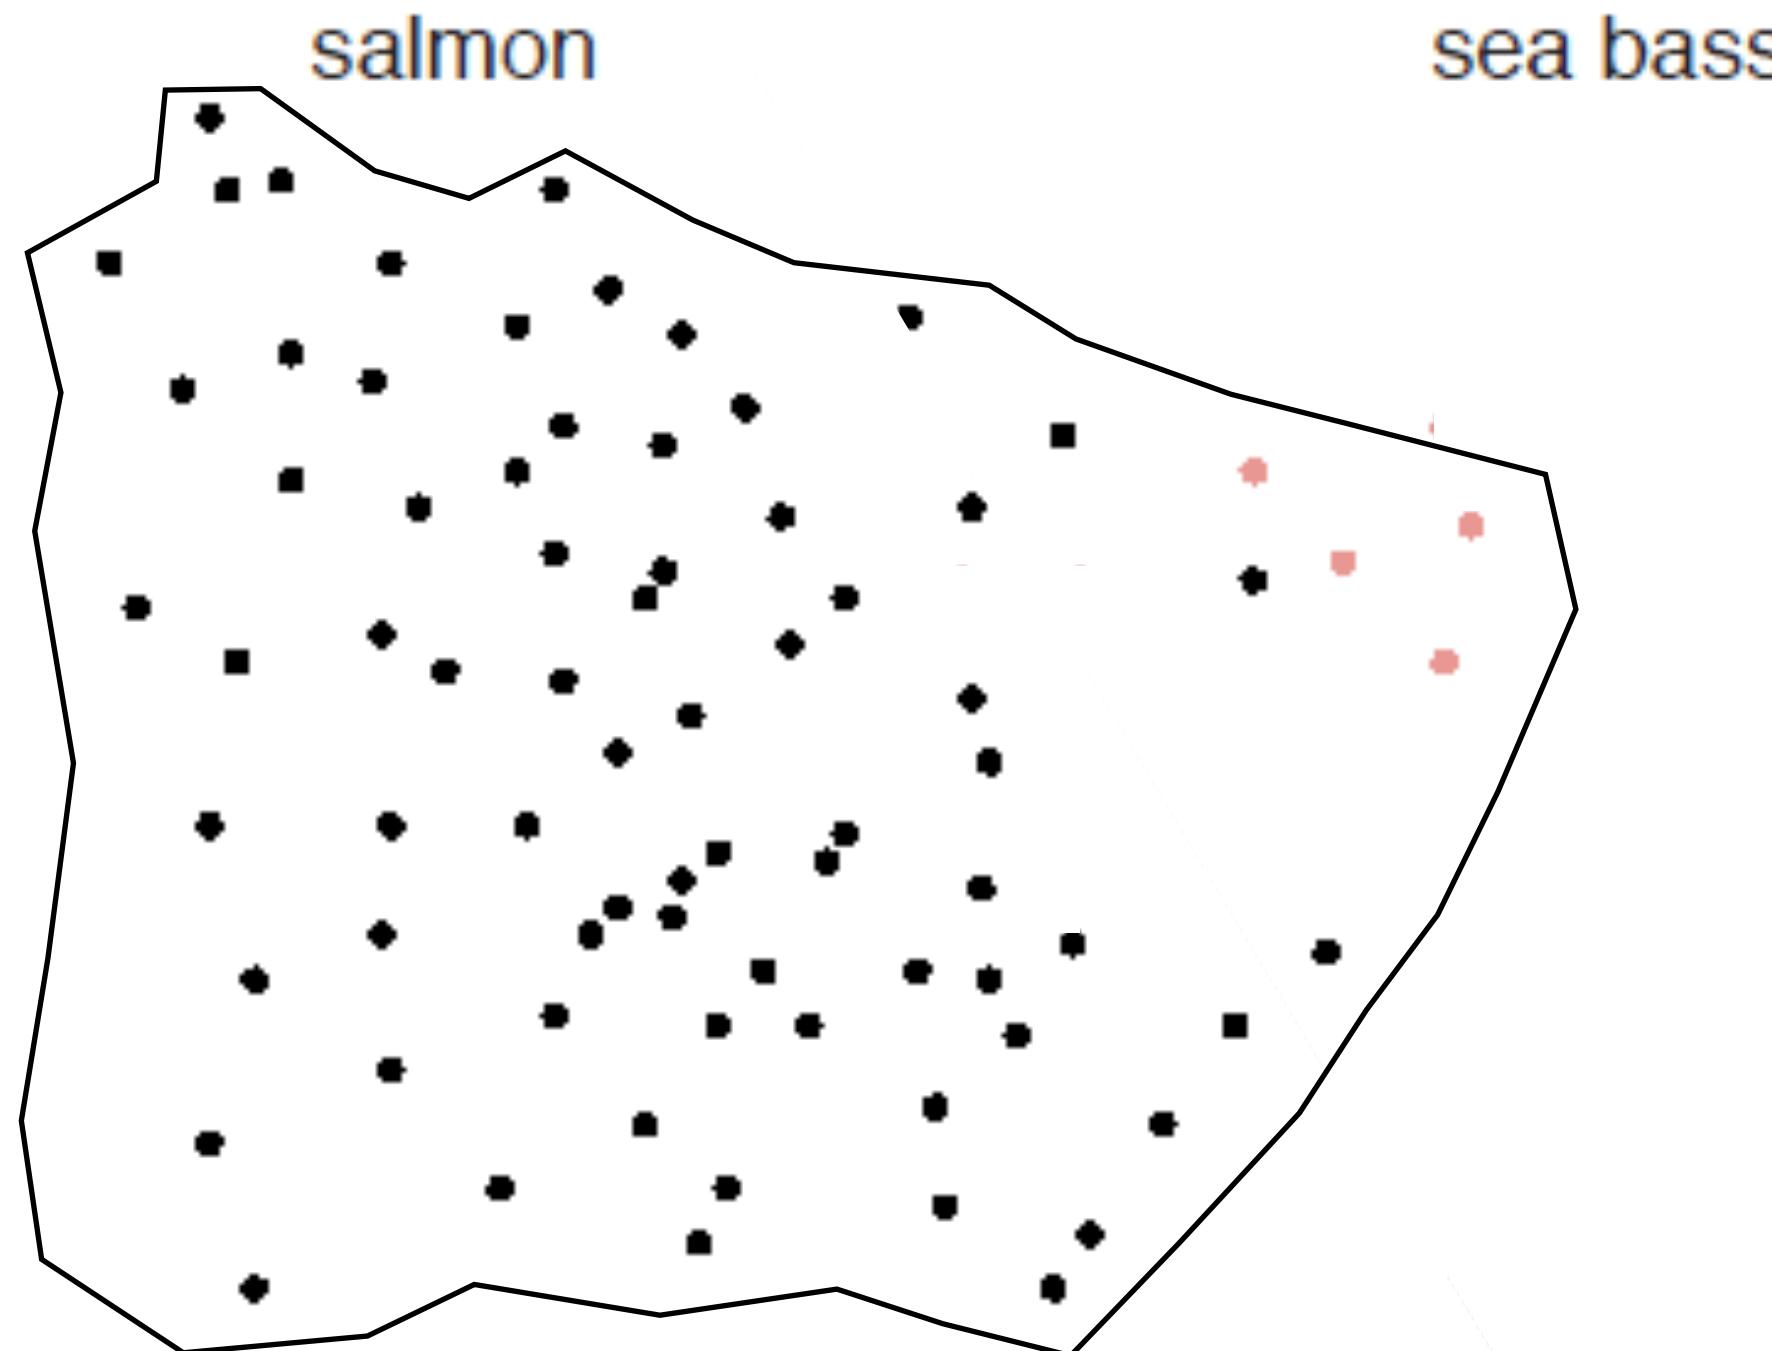
Accuracy (Acc)

$$Acc = \frac{Correct\ Predictions}{Total\ Predictions}$$

Limitation: what happens when we have unbalanced data?

Machine Learning (ML)

Metrics



Example model: everything is salmon!

Balanced Accuracy (BAcc)

$$BAcc = \frac{1}{C} \sum_{i=1}^C \frac{\text{Correct Predictions}_i}{\text{Total Predictions}_i}$$

C is the number of classes.

Average of class-wise accuracy.

Machine Learning (ML)

Metrics

		Ground truth
		salmon
		not salmon
Predicted salmon	True Positive (TP)	False Positive (FP)
	False Negative (FN)	True Negative (TN)

i=salmon

Precision (P) and Recall (R)

$$P = \frac{\sum_{i=1}^C TP_i}{\sum_{i=1}^C (TP_i + FP_i)}$$

How precise is the model when it classifies as i ?
(focus on prediction)

$$R = \frac{\sum_{i=1}^C TP_i}{\sum_{i=1}^C (TP_i + FN_i)}$$

How good is the model in retrieving samples from class i ?
(focus on ground truth)

C is the number of classes.



LOYOLA
UNIVERSITY CHICAGO

Machine Learning (ML)

Metrics

		Ground truth
		salmon
		not salmon
Predicted salmon	True Positive (TP)	False Positive (FP)
	False Negative (FN)	True Negative (TN)

i=salmon

Fscore

$$F_\beta = \frac{(\beta^2 + 1) \times P \times R}{\beta^2 \times P + R}$$

Harmonic mean of P and R
(when we care about both).

$F_1 score$: $\beta = 1$, equal weight to P and to R.

$F_2 score$: $\beta = 2$, more weight to R.

Question: can you think of an application where R is more important than P?

Machine Learning (ML)

Metrics

		Ground truth
		salmon
		not salmon
Predicted salmon	True Positive (TP)	False Positive (FP)
	False Negative (FN)	True Negative (TN)

i=salmon

Fscore

$$F_{\beta} = \frac{(\beta^2 + 1) \times P \times R}{\beta^2 \times P + R}$$

Harmonic mean of P and R
(when we care about both).

$F_1 score$: $\beta = 1$, equal weight to P and to R.

$F_2 score$: $\beta = 2$, more weight to R.

$F_{0.5} score$: $\beta = 0.5$, more weight to P.



Machine Learning (ML)

Metrics

		Ground truth			
		salmon	sea bass	tuna	pirarucu
Predicted	salmon	215	32	0	3
	sea bass	10	701	0	0
tuna	1	0	370	2	
pirarucu	1	0	0	102	

Confusion Matrix

Visualization tool: *what can you see?*

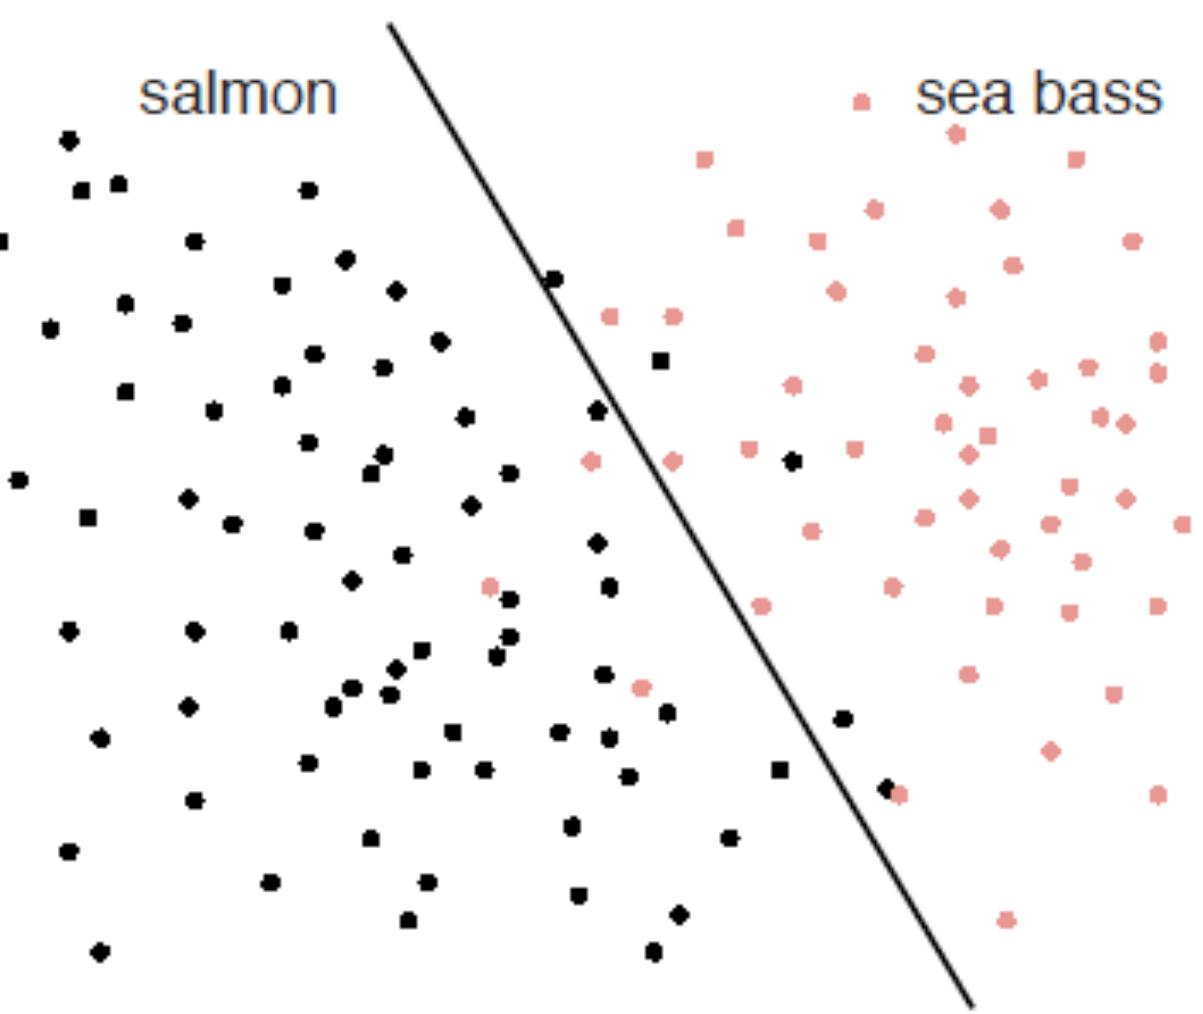
*Why is it called **confusion** matrix?*

Number of samples.

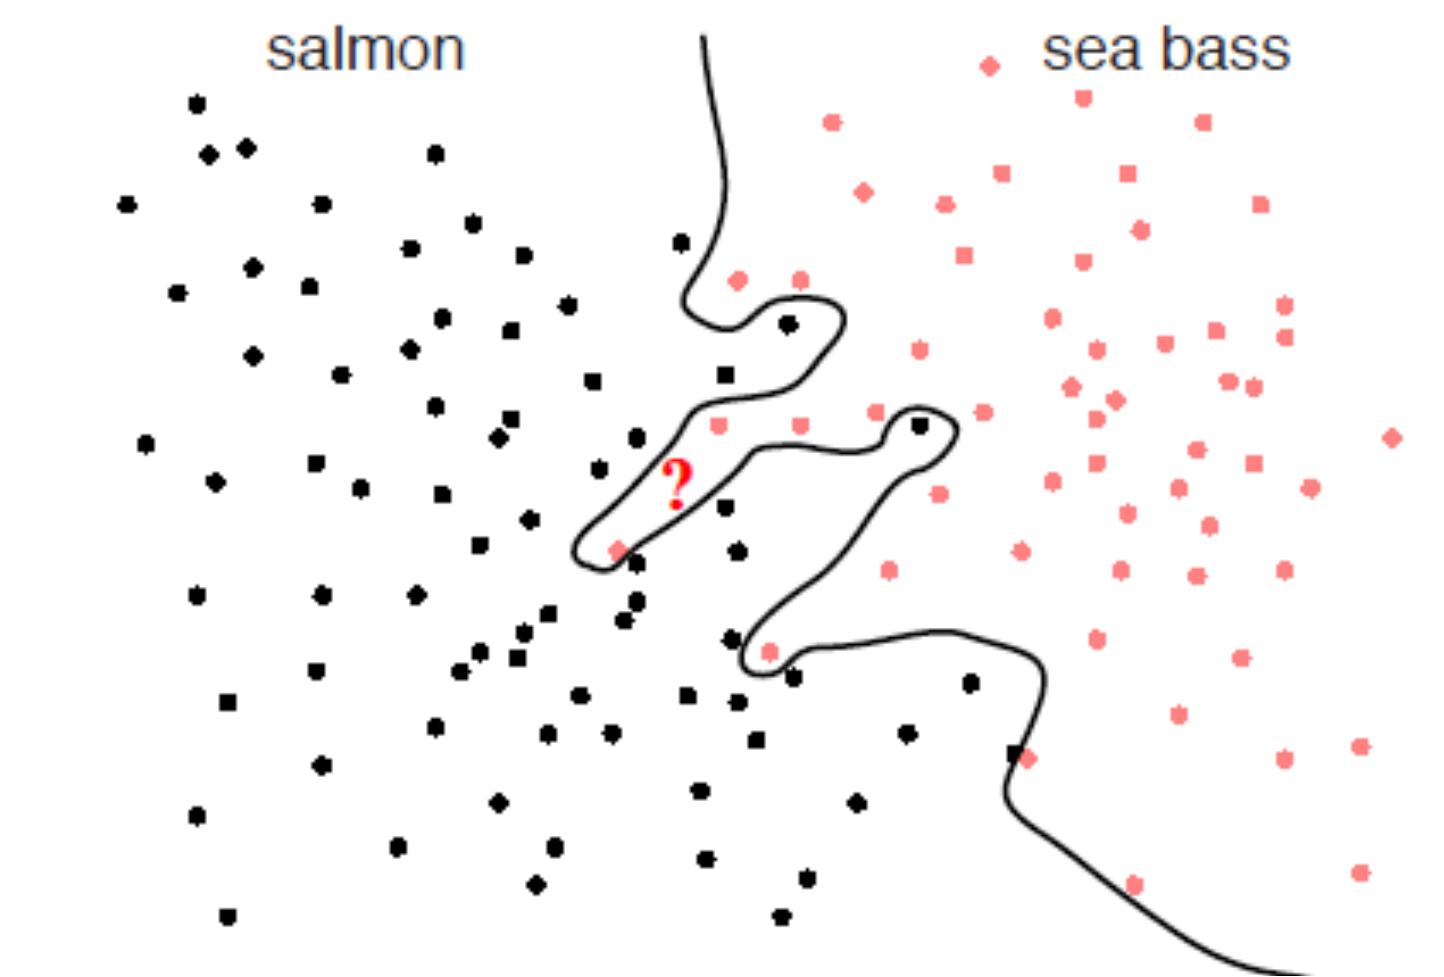
Machine Learning (ML)

Data-driven Learning Issues

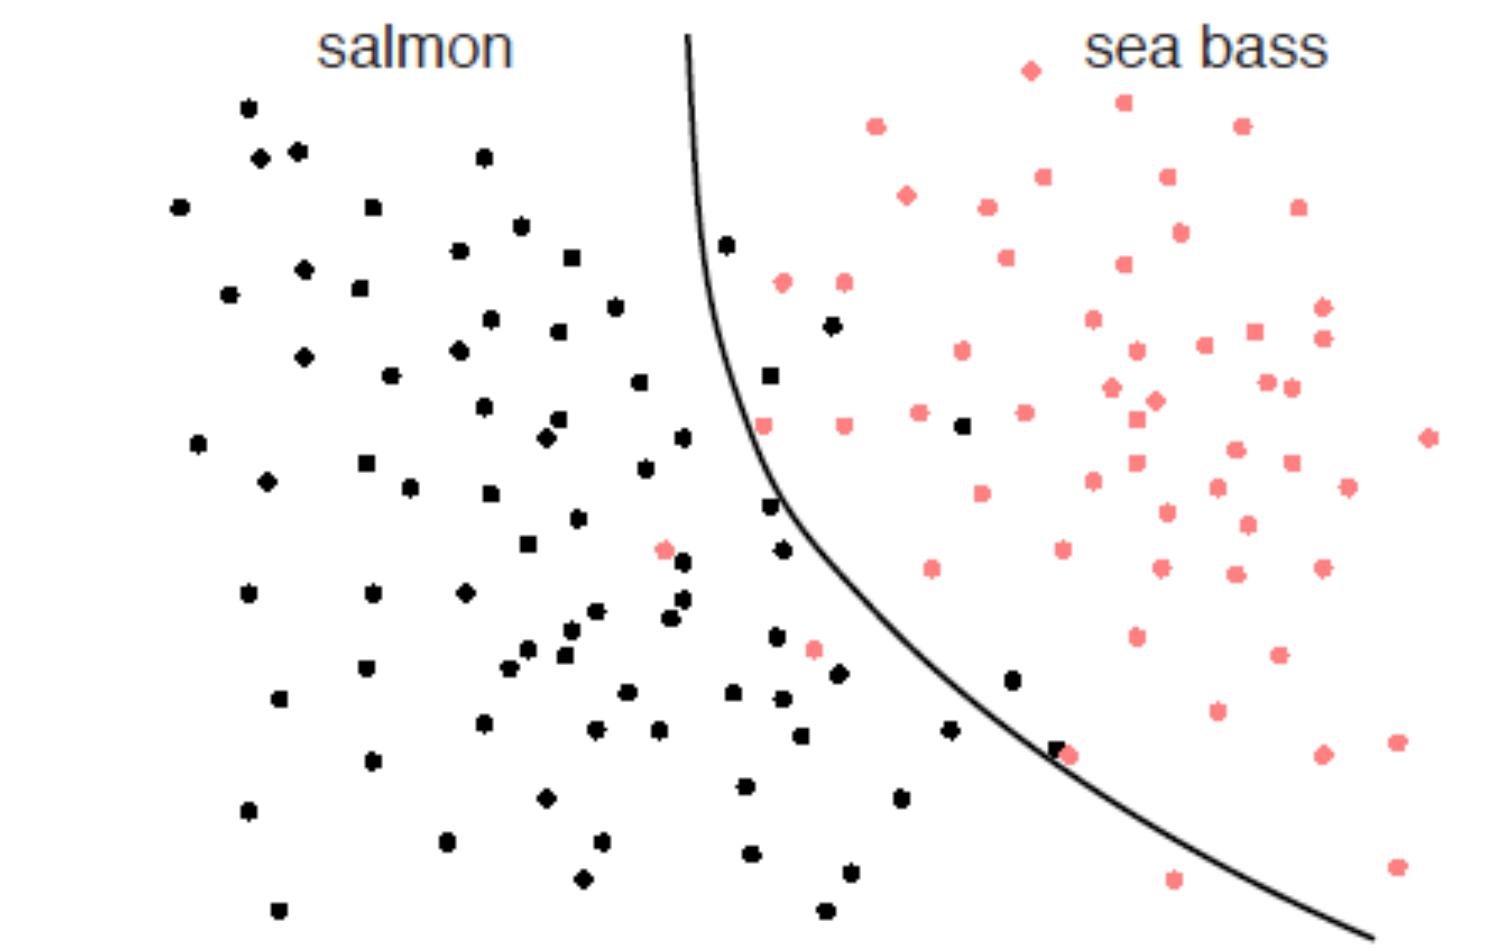
What happens in face of **unseen** data (normal system operation)?



Under-fitting
(*too-simple model*)



Over-fitting
(*too-complex model*)



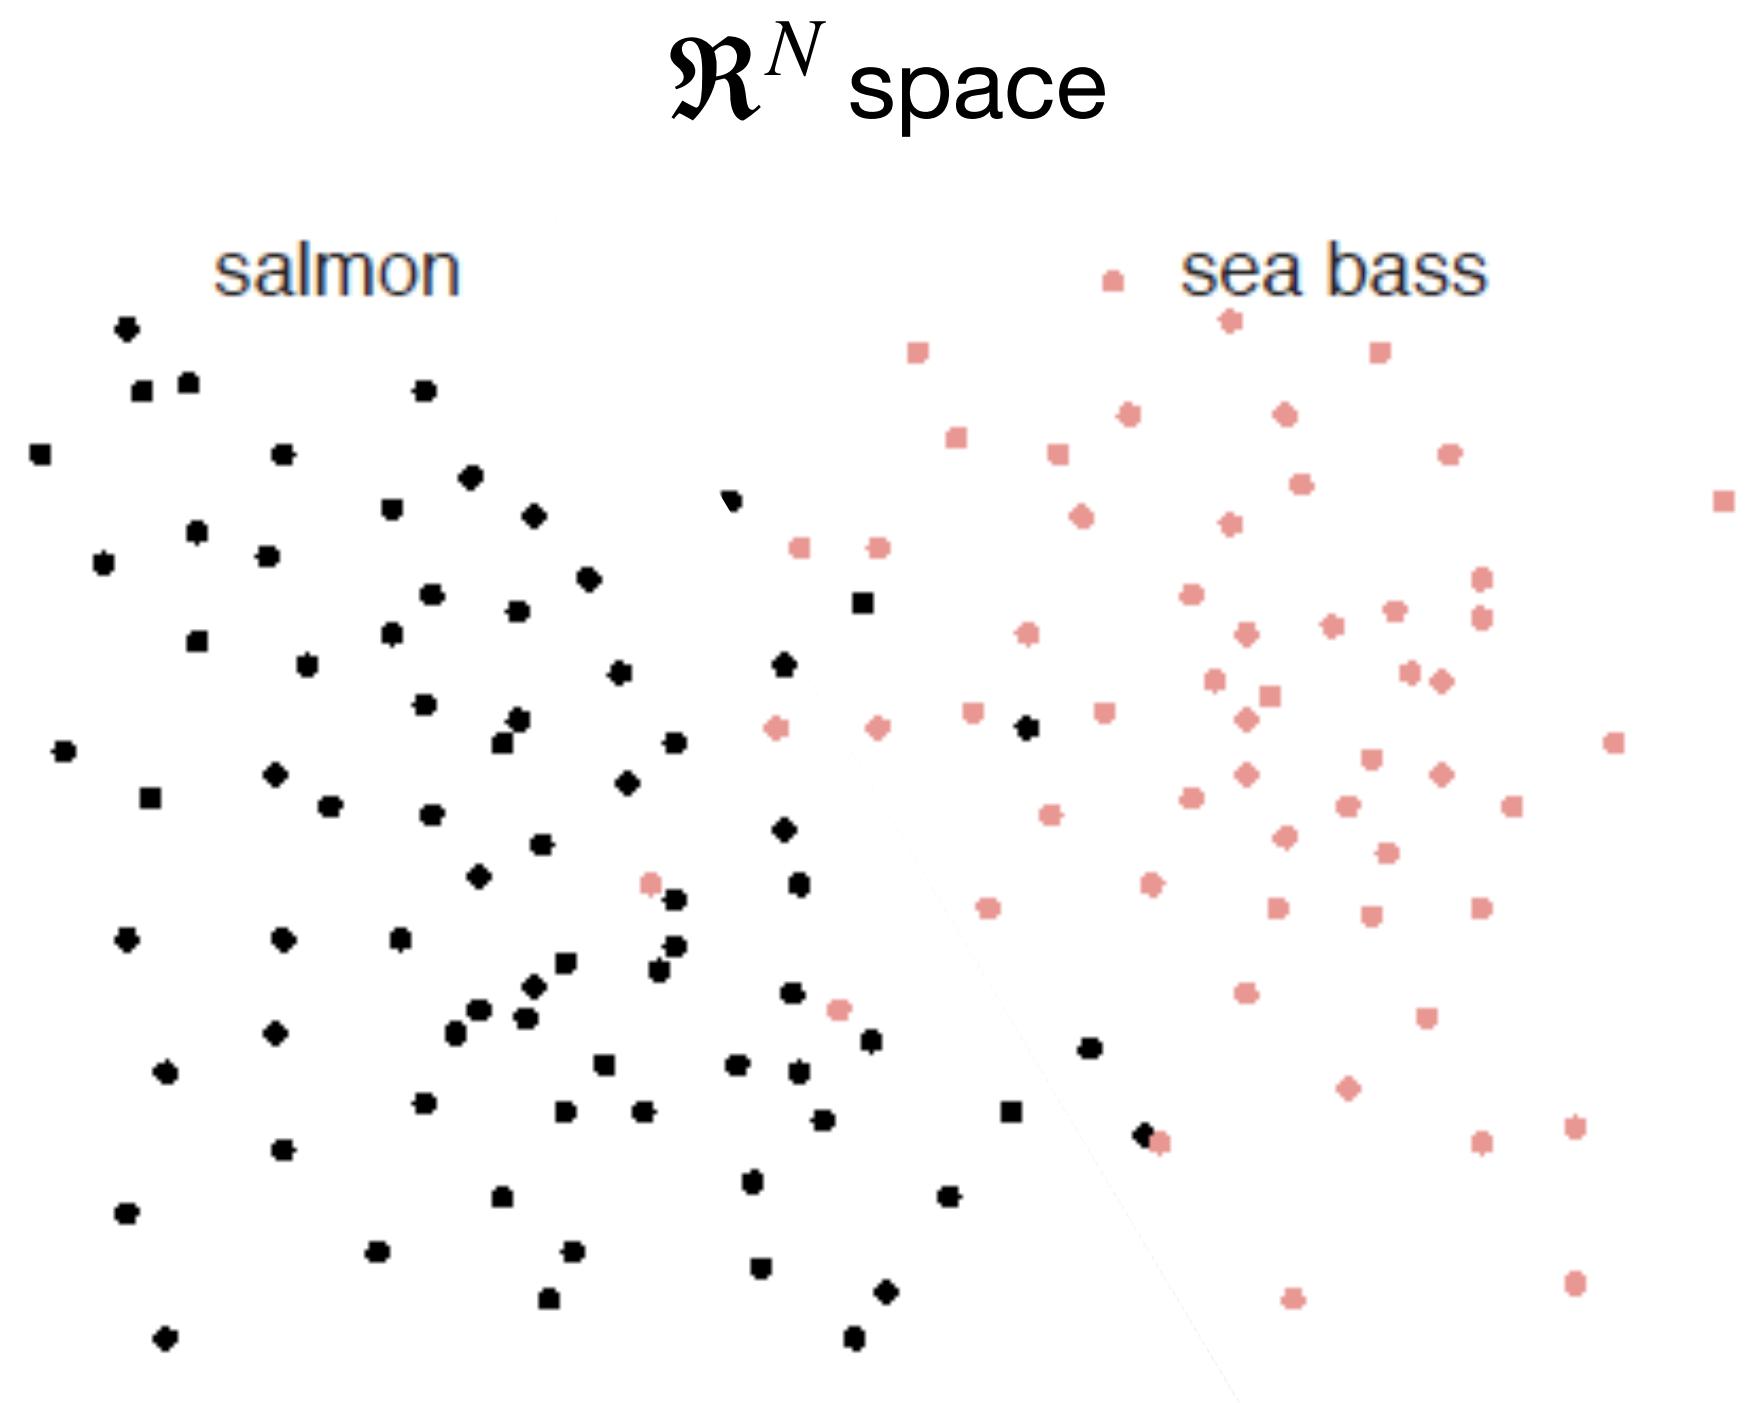
Okay-ish



LOYOLA
UNIVERSITY CHICAGO

Machine Learning (ML)

Data Split



How to estimate the model's performance in face of unseen data?

Random split

$X\%$

$100-X\%$

Training Set

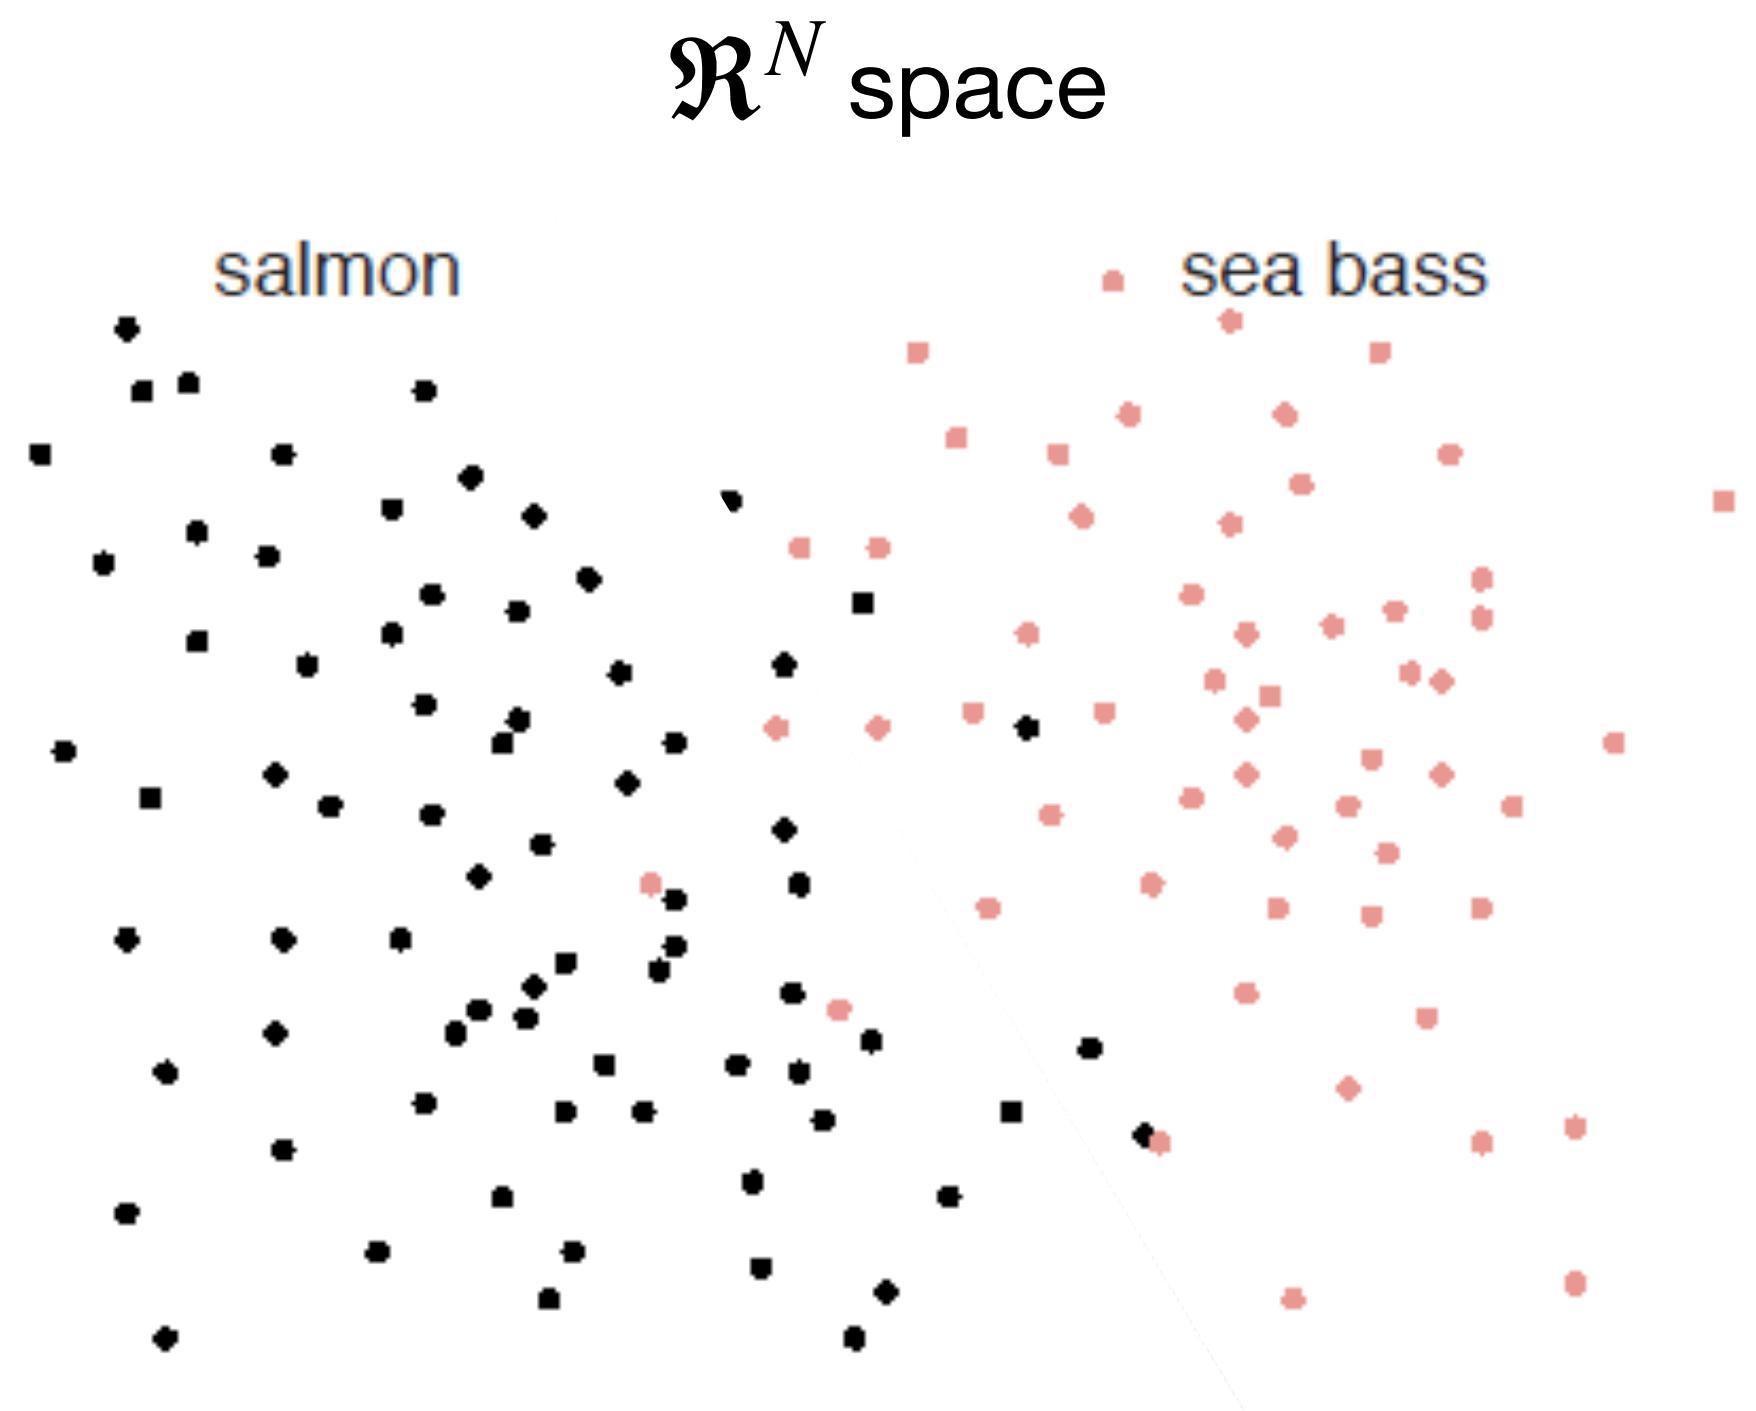
Test Set

Present to the learner
(learning stage).

Reserve to assess
performance in the end
(inference stage).

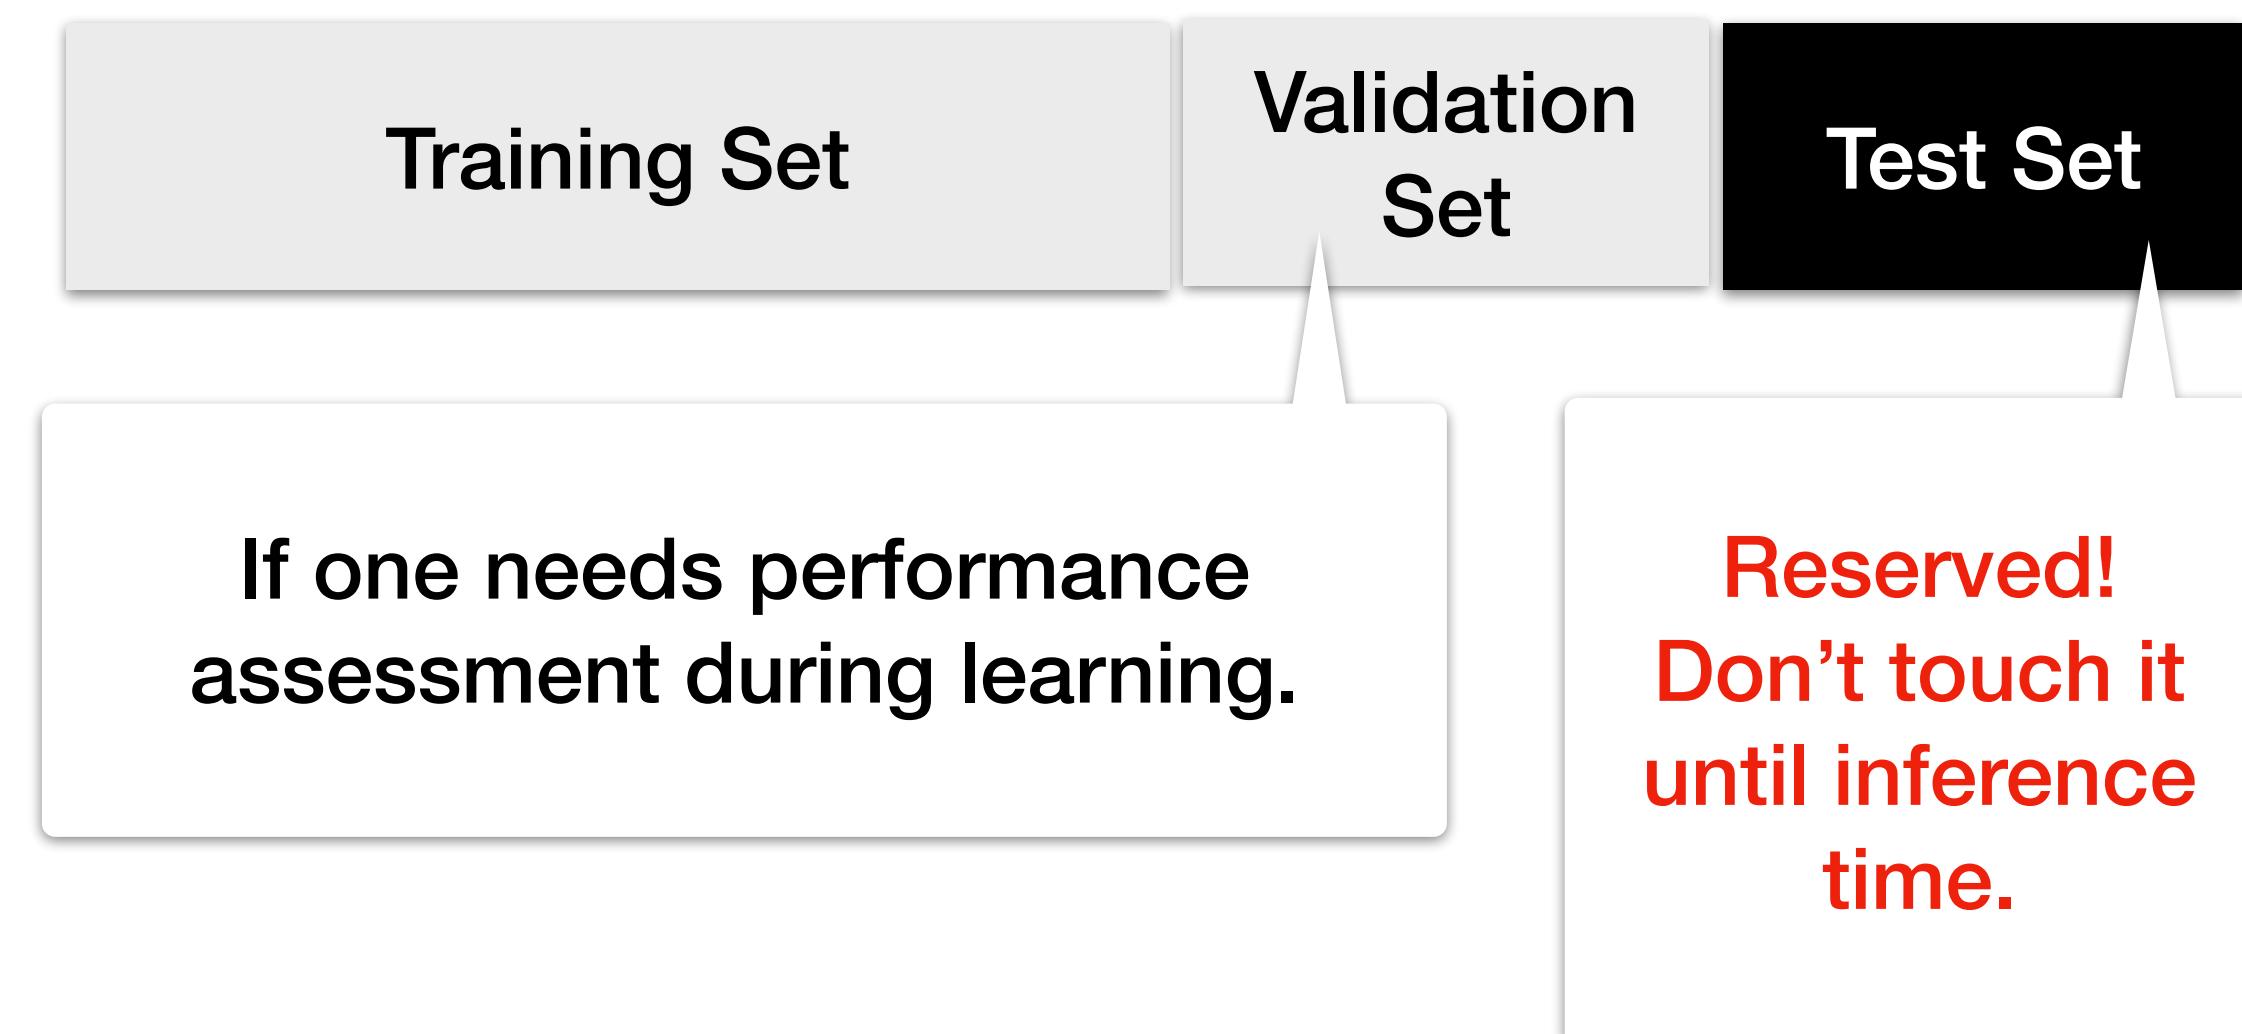
Machine Learning (ML)

Data Split



How to estimate the model's performance in face of unseen data?

Alternative rand. split X % Y % 100-(X+Y)%



Machine Learning (ML)

Random Data Split

What if the estimated performance (P) is a matter of chance?

K-fold Cross Validation

Example: K=5, 5 random splits.



Report μ_P and σ_P .

Machine Learning (ML)

Random Data Split

What if the estimated performance (P) is a matter of chance?

K-fold Cross Validation

Example: K=5, 5 random splits.



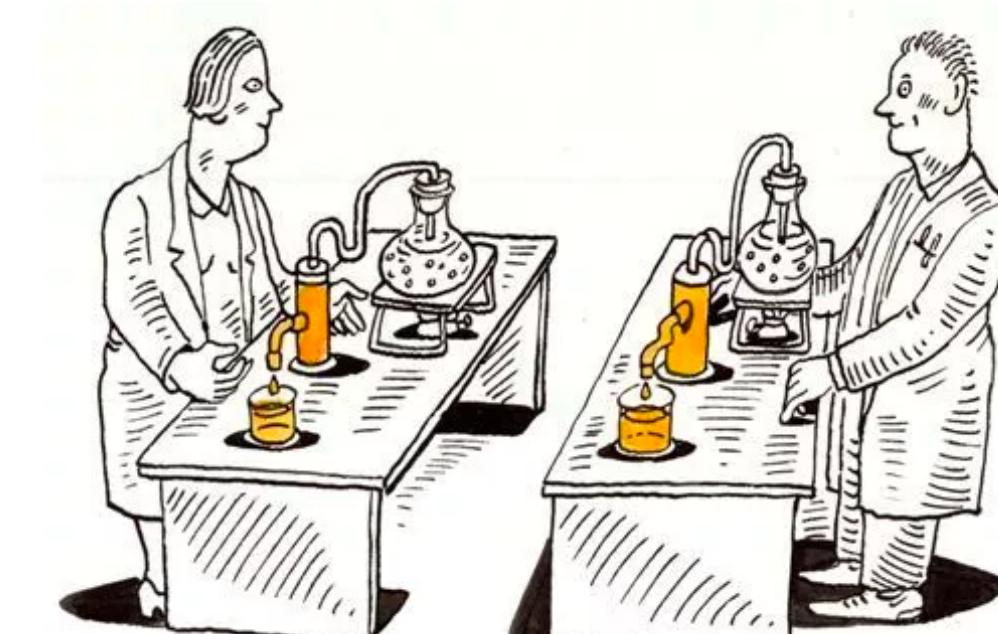
Report μ_P and σ_P .

Public Random Split

No time to train the solution
K times?

Make random split publicly
available, so others can:

1. Reproduce your results.
2. Use the same split to train and compare their solutions.



[https://www.displayr.com/
what-is-reproducible-research/](https://www.displayr.com/what-is-reproducible-research/)



LOYOLA
UNIVERSITY CHICAGO

Machine Learning (ML)

Data-driven Learning Types

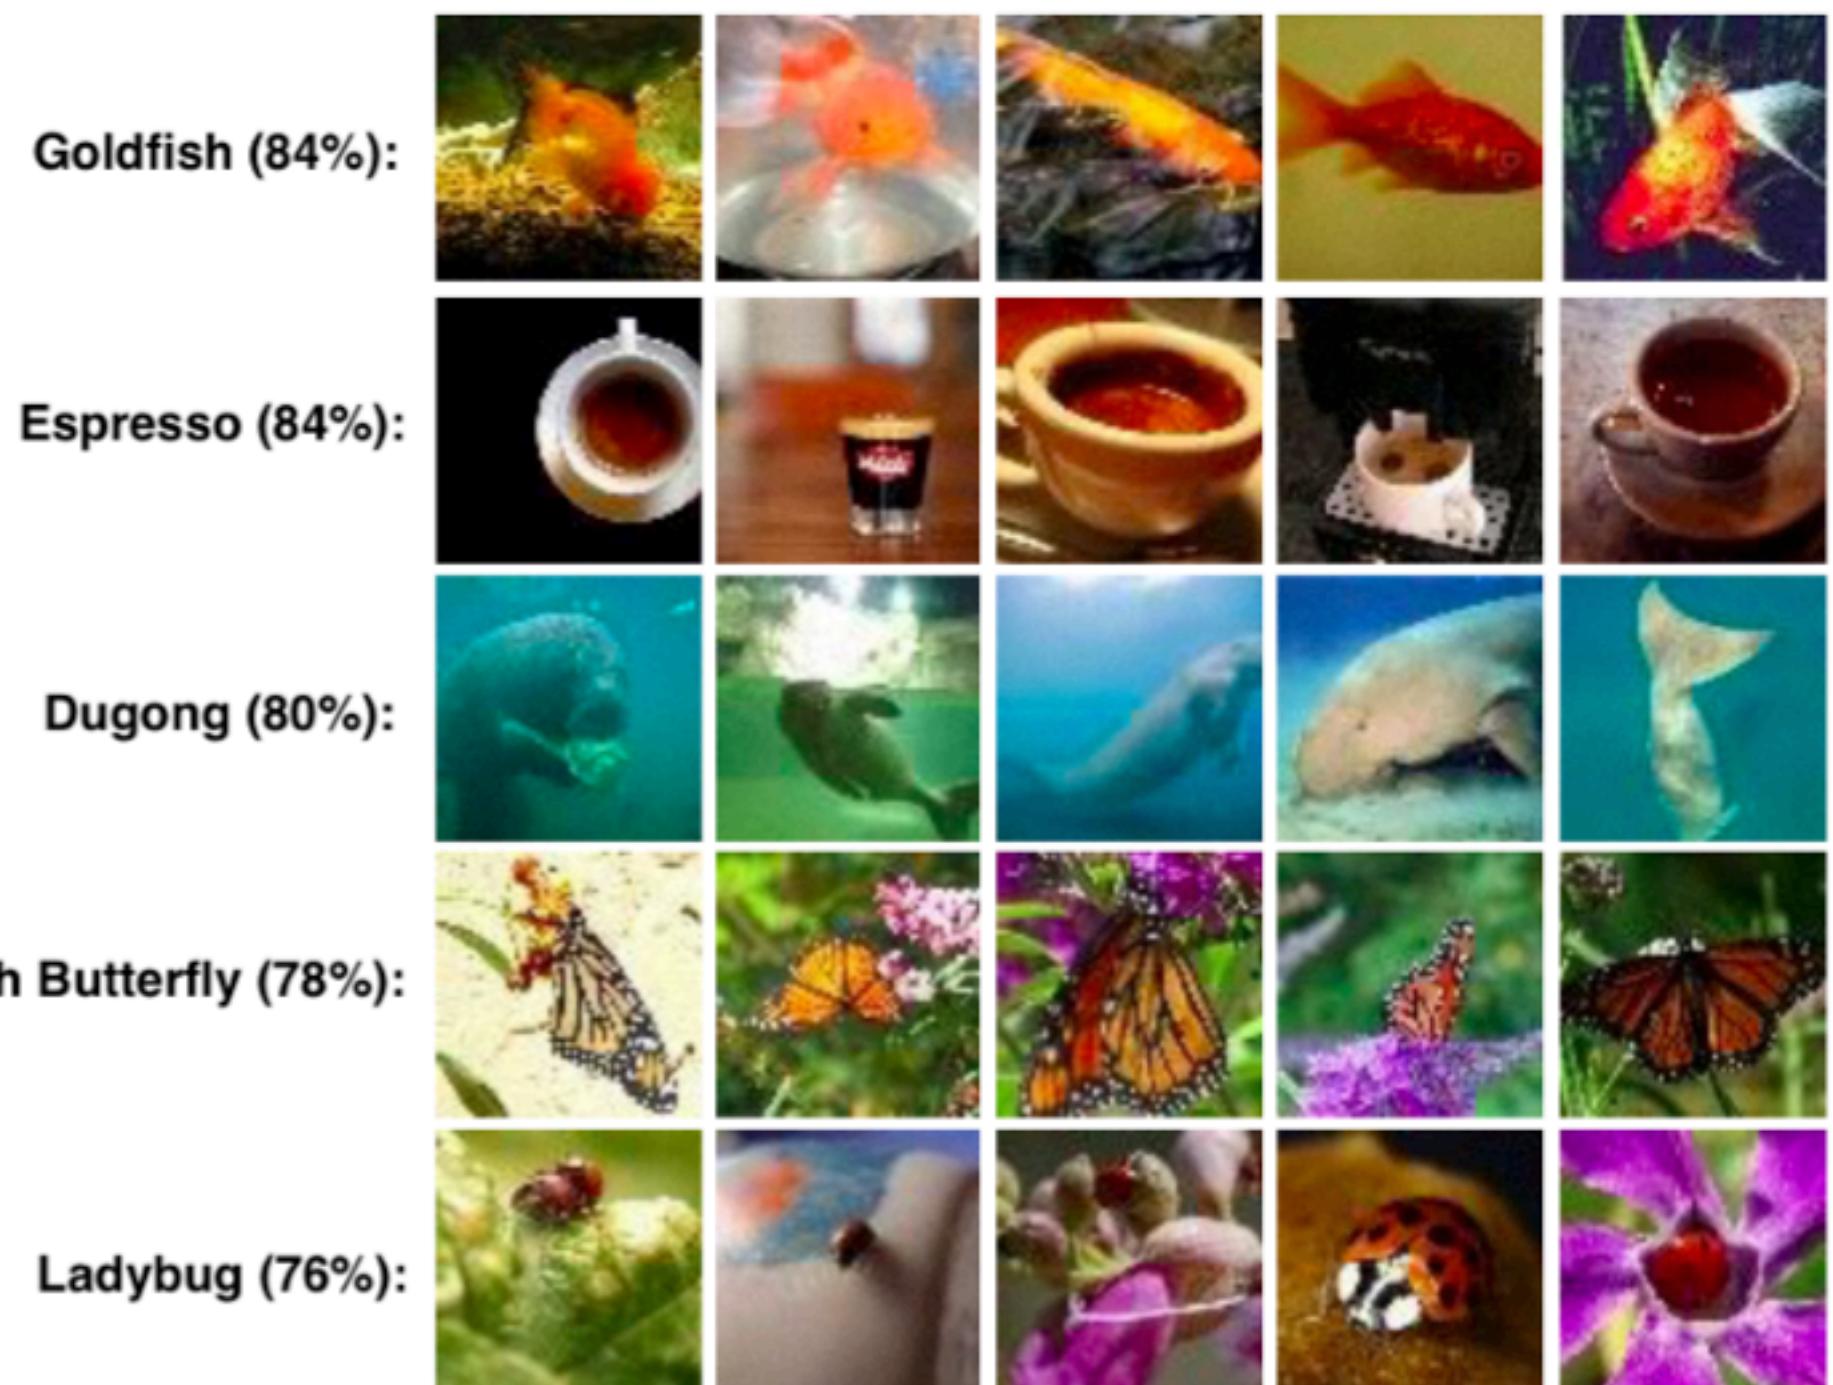
Supervised Learning (1/2)

The target problem has well-defined classes.

There are annotated data to train the learner.

Annotation: each sample (feature vector) has a class.

Yao, L., Miller, J.
Tiny ImageNet with CNNs. CS Stanford, 2015.



Machine Learning (ML)

Data-driven Learning Types

Supervised Learning (1/2) Closed Set versus Open Set

E.g., Face Recognition

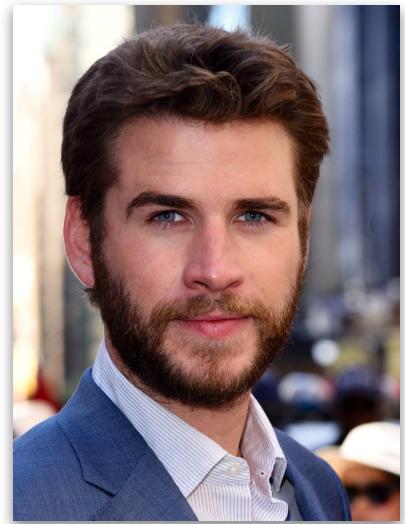


query

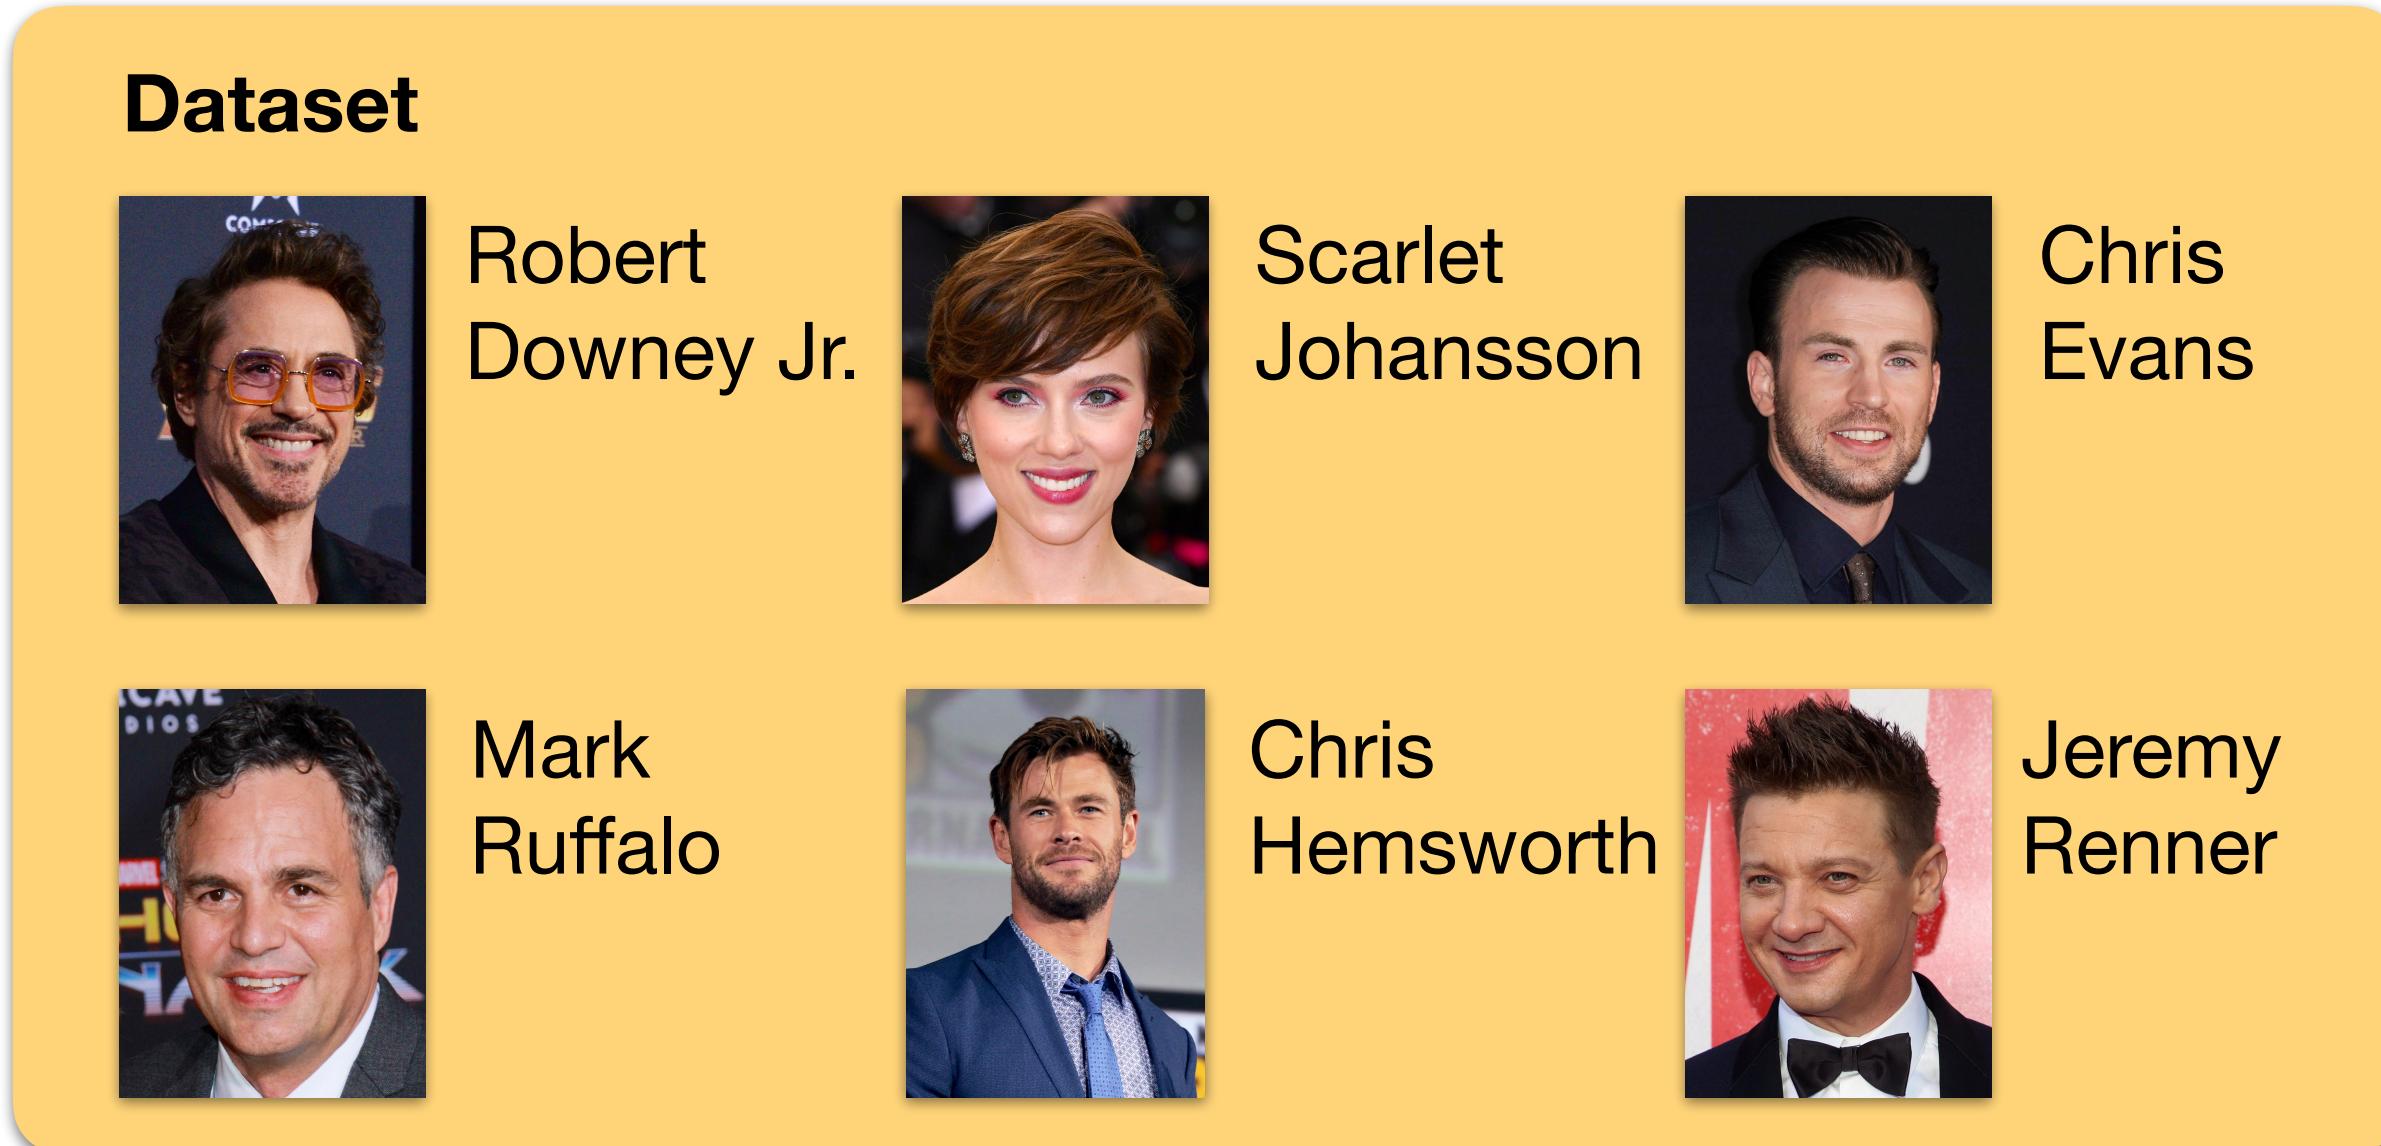


Machine Learning (ML)

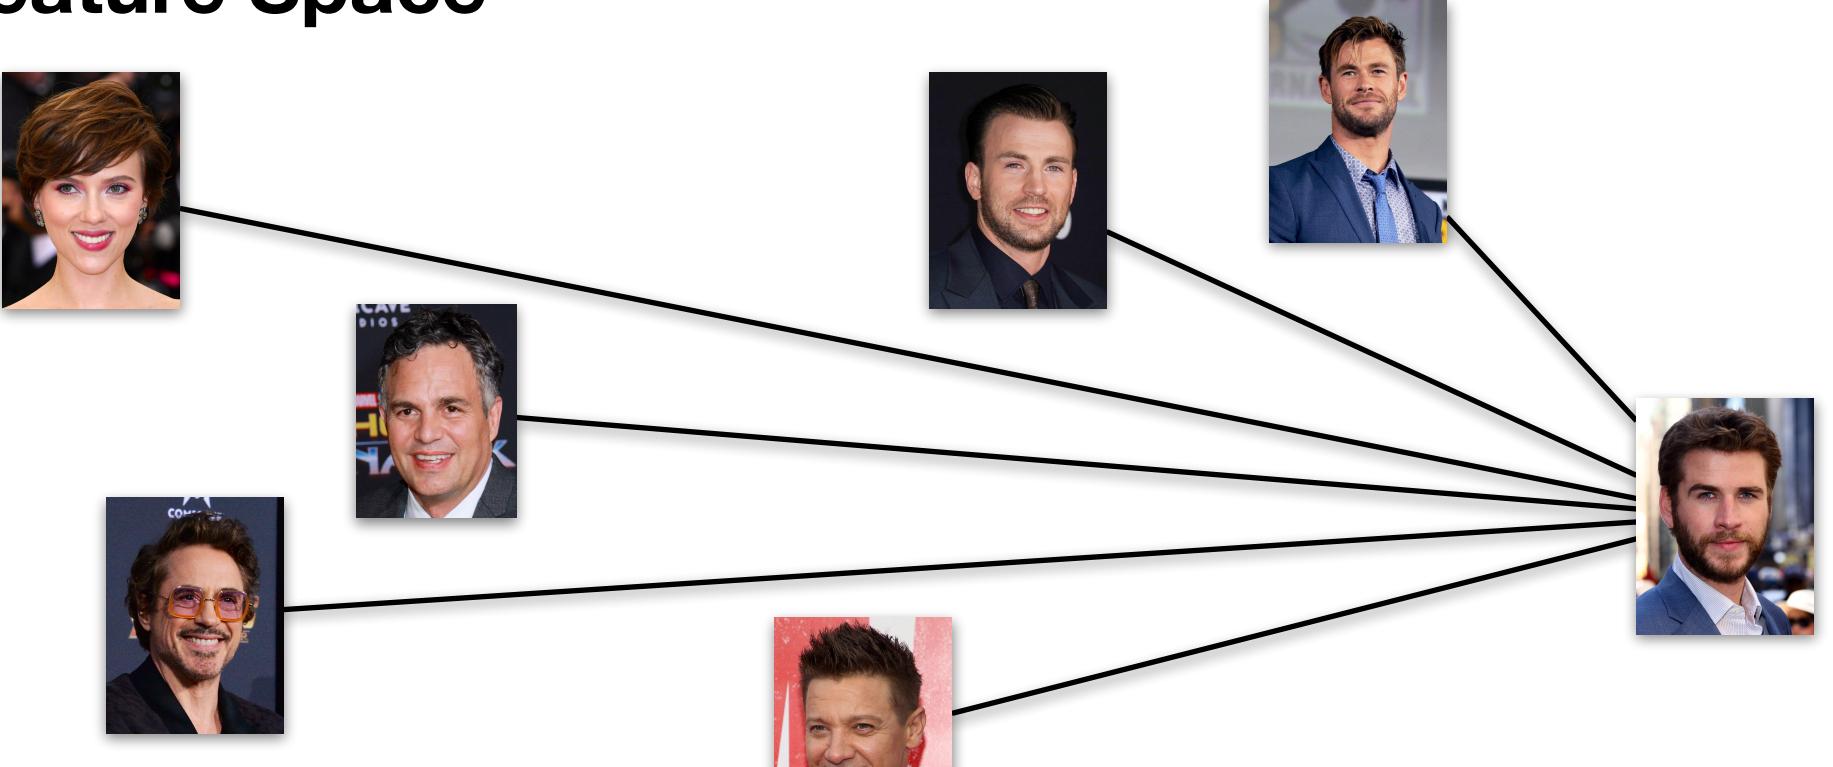
Closed Set versus Open Set



Query
(Liam Hemsworth)



Feature Space



Closed Set

Output
This is
Chris Hemsworth!



Open Set

Output
I don't know
this person!



LOYOLA
UNIVERSITY CHICAGO

Machine Learning (ML)

Data-driven Learning Types

Unsupervised Learning (2/2)

Either the target problem has no well-defined classes.

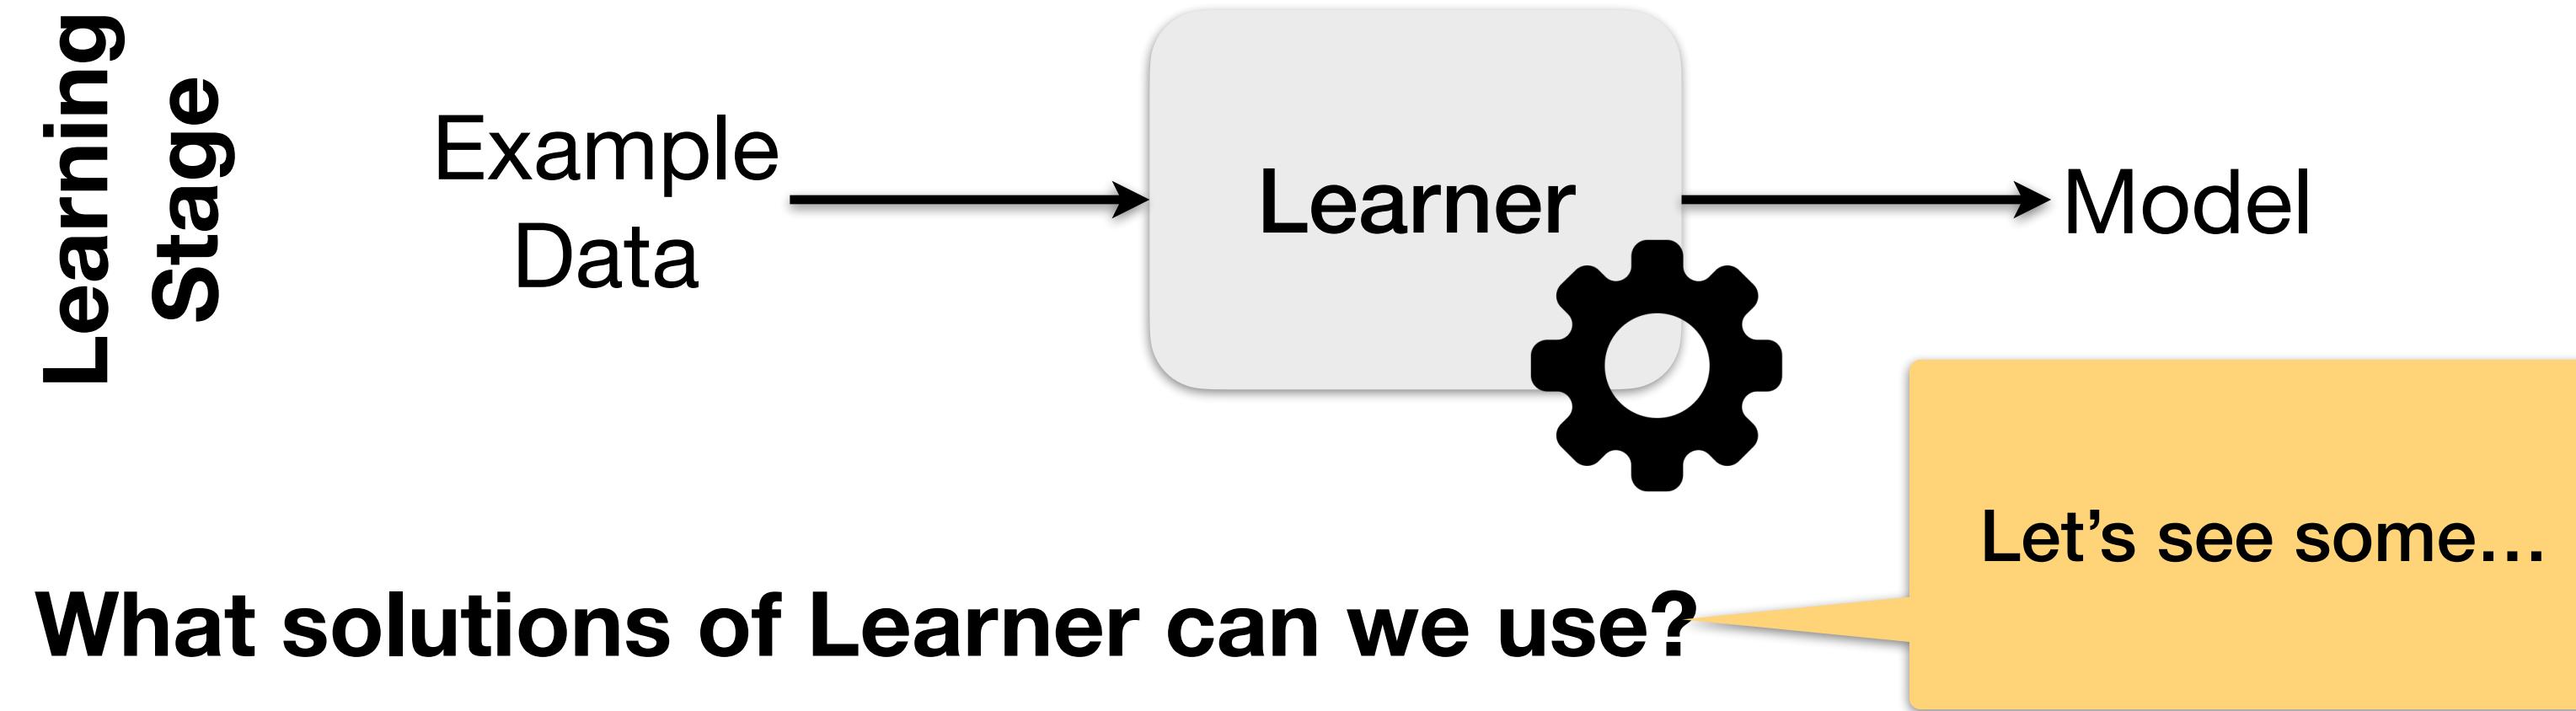
OR

There are not annotated data to train the learner.



Machine Learning (ML)

Learner and Model Solutions



What solutions of Learner can we use?

Unsupervised

Clustering methods such as k-means, k-medoids, etc.

Supervised

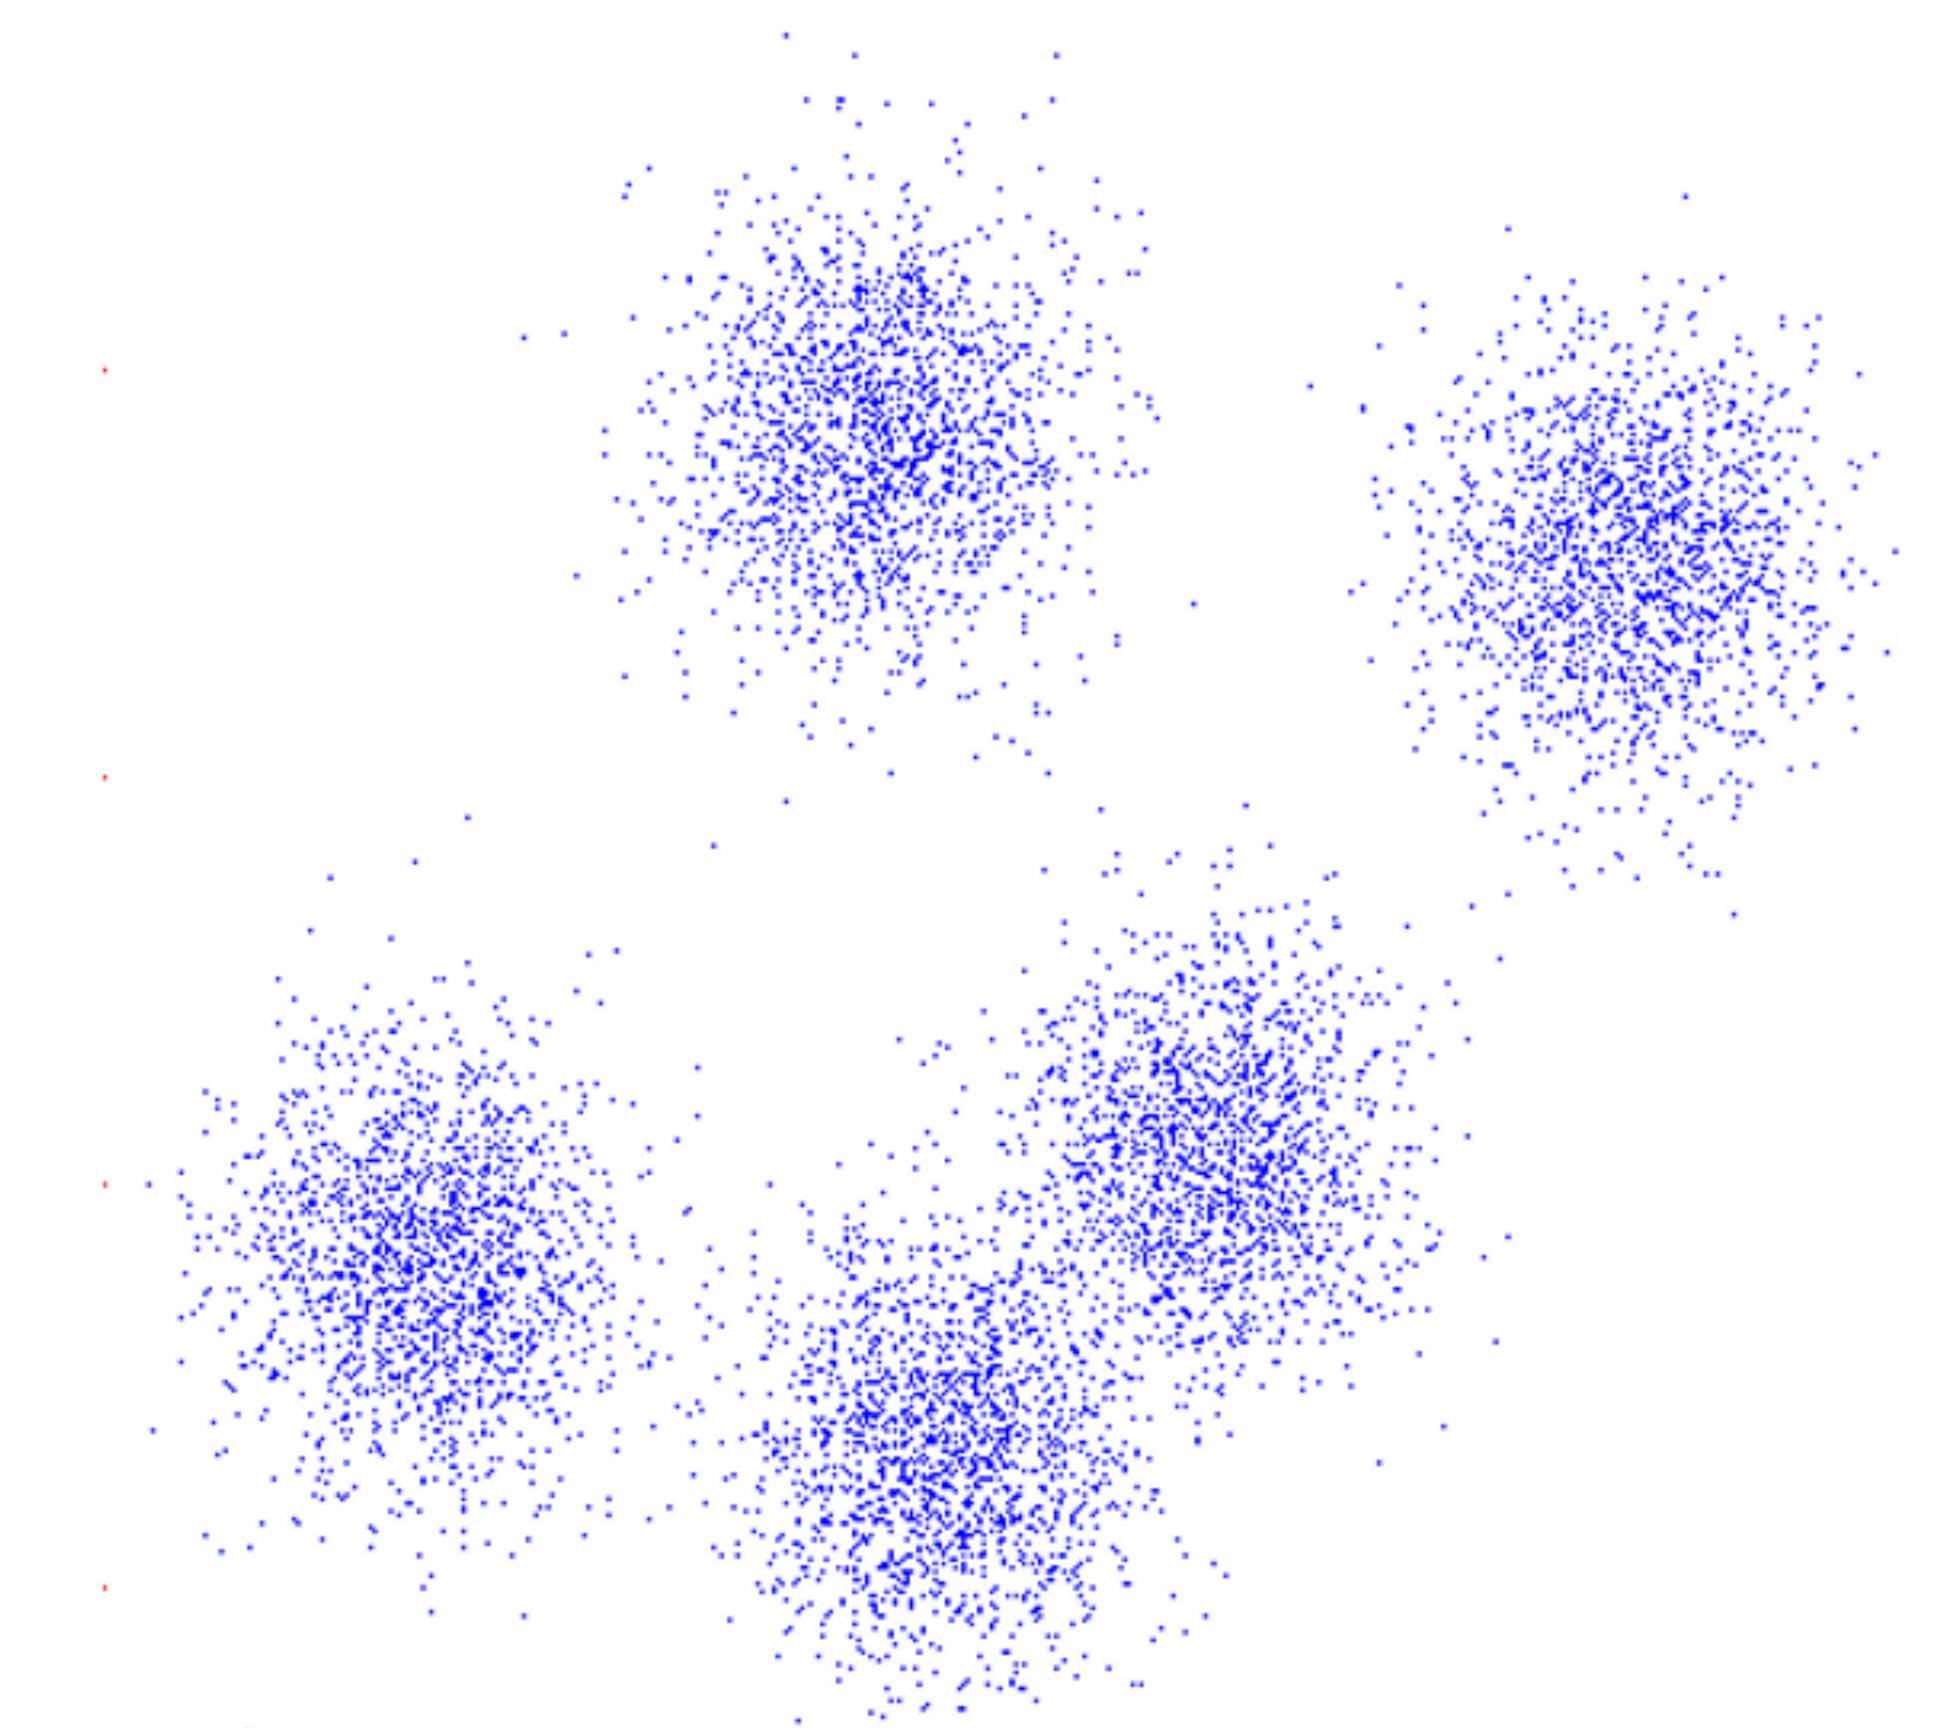
Decision trees, random forests, Support Vector Machines (SVM), typical Neural Networks (NN), etc.

K-Means

How to reduce
data complexity?

Cluster the features and limit
the k-nearest search to one or
a couple of clusters.

There will be less elements to
consider.



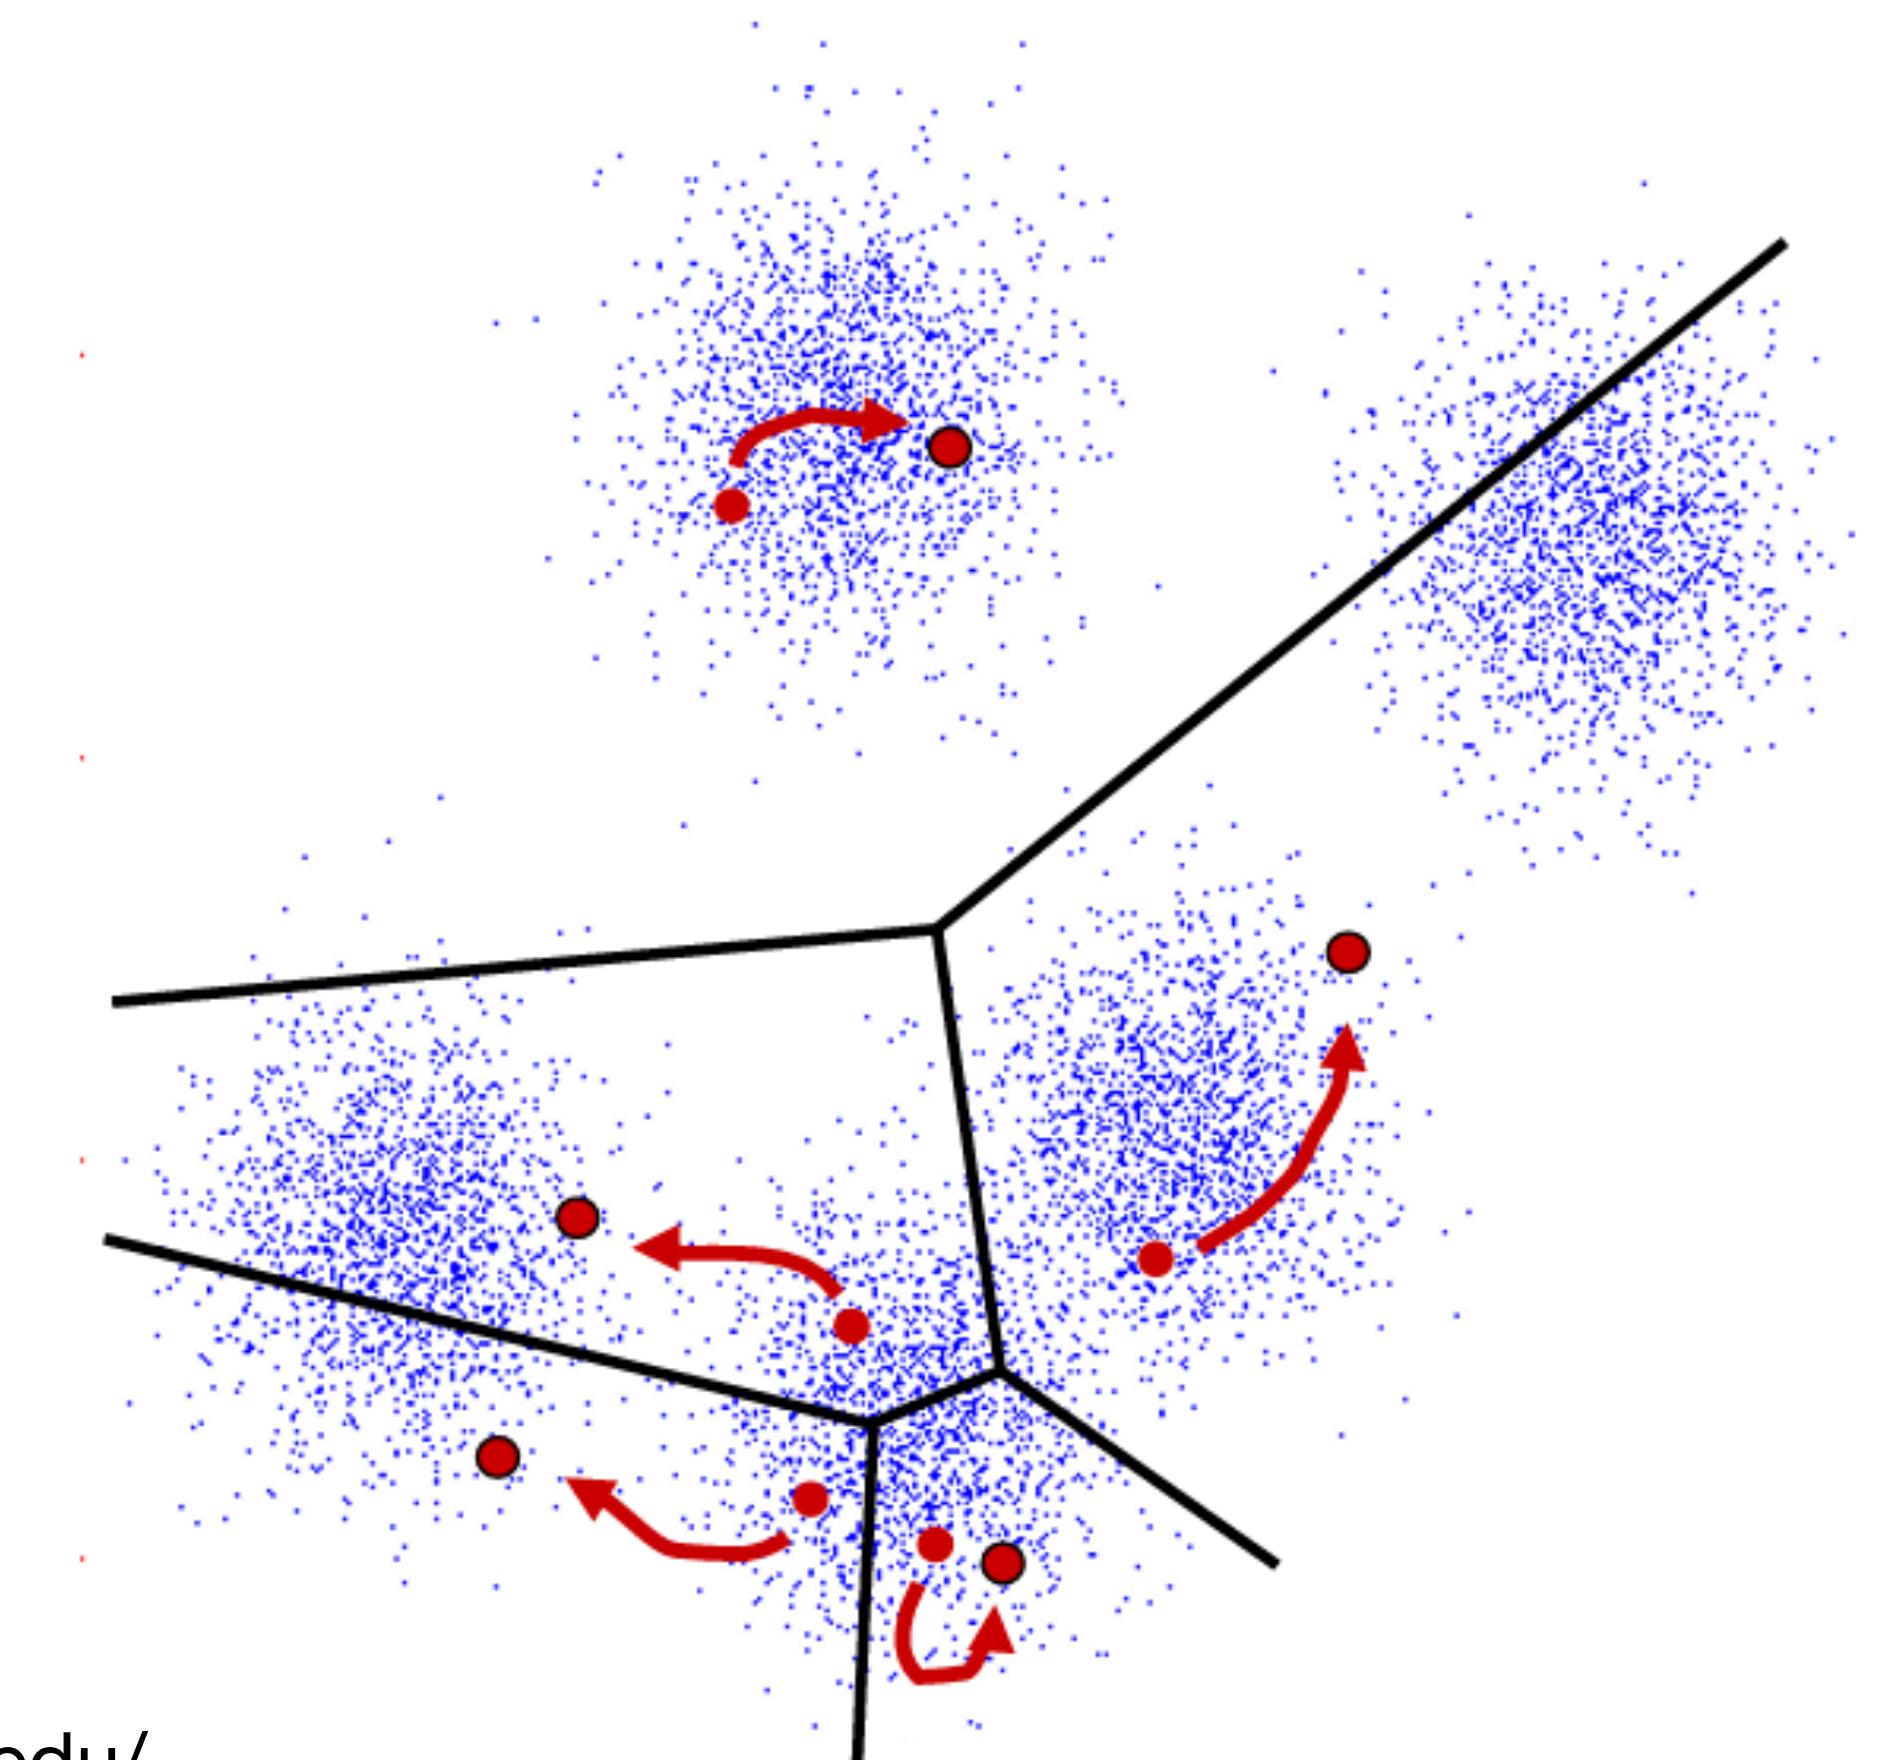
Source: <https://people.csail.mit.edu/dsontag/courses/ml12/slides/lecture14.pdf>

K-Means

How to reduce data complexity?

Cluster the features and limit the k-nearest search to one or a couple of clusters.

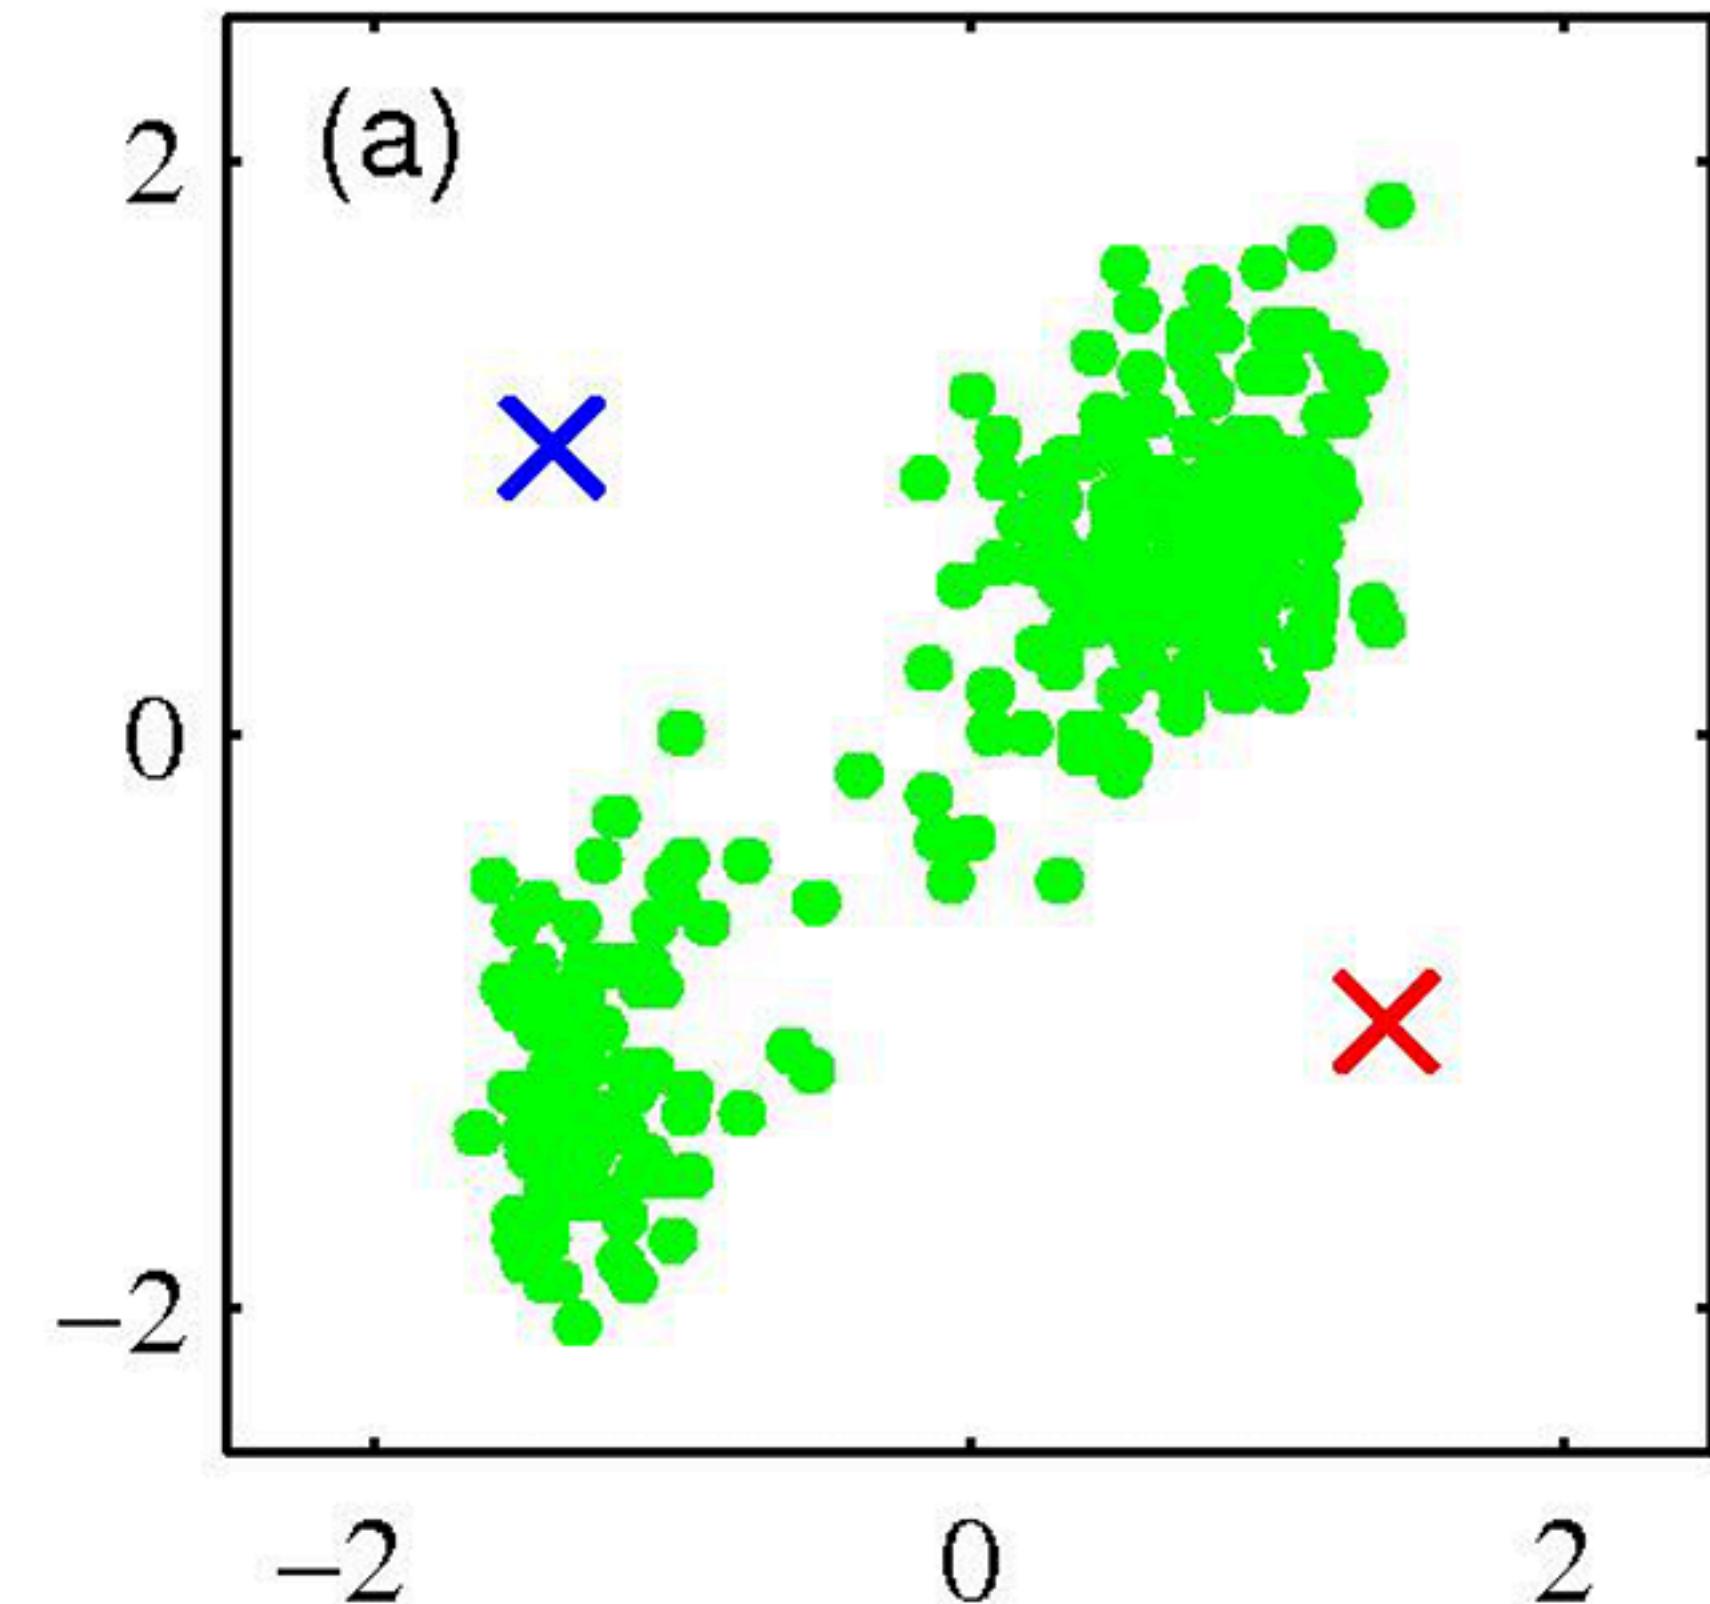
There will be less elements to consider.



Source: <https://people.csail.mit.edu/dsontag/courses/ml12/slides/lecture14.pdf>

K-Means

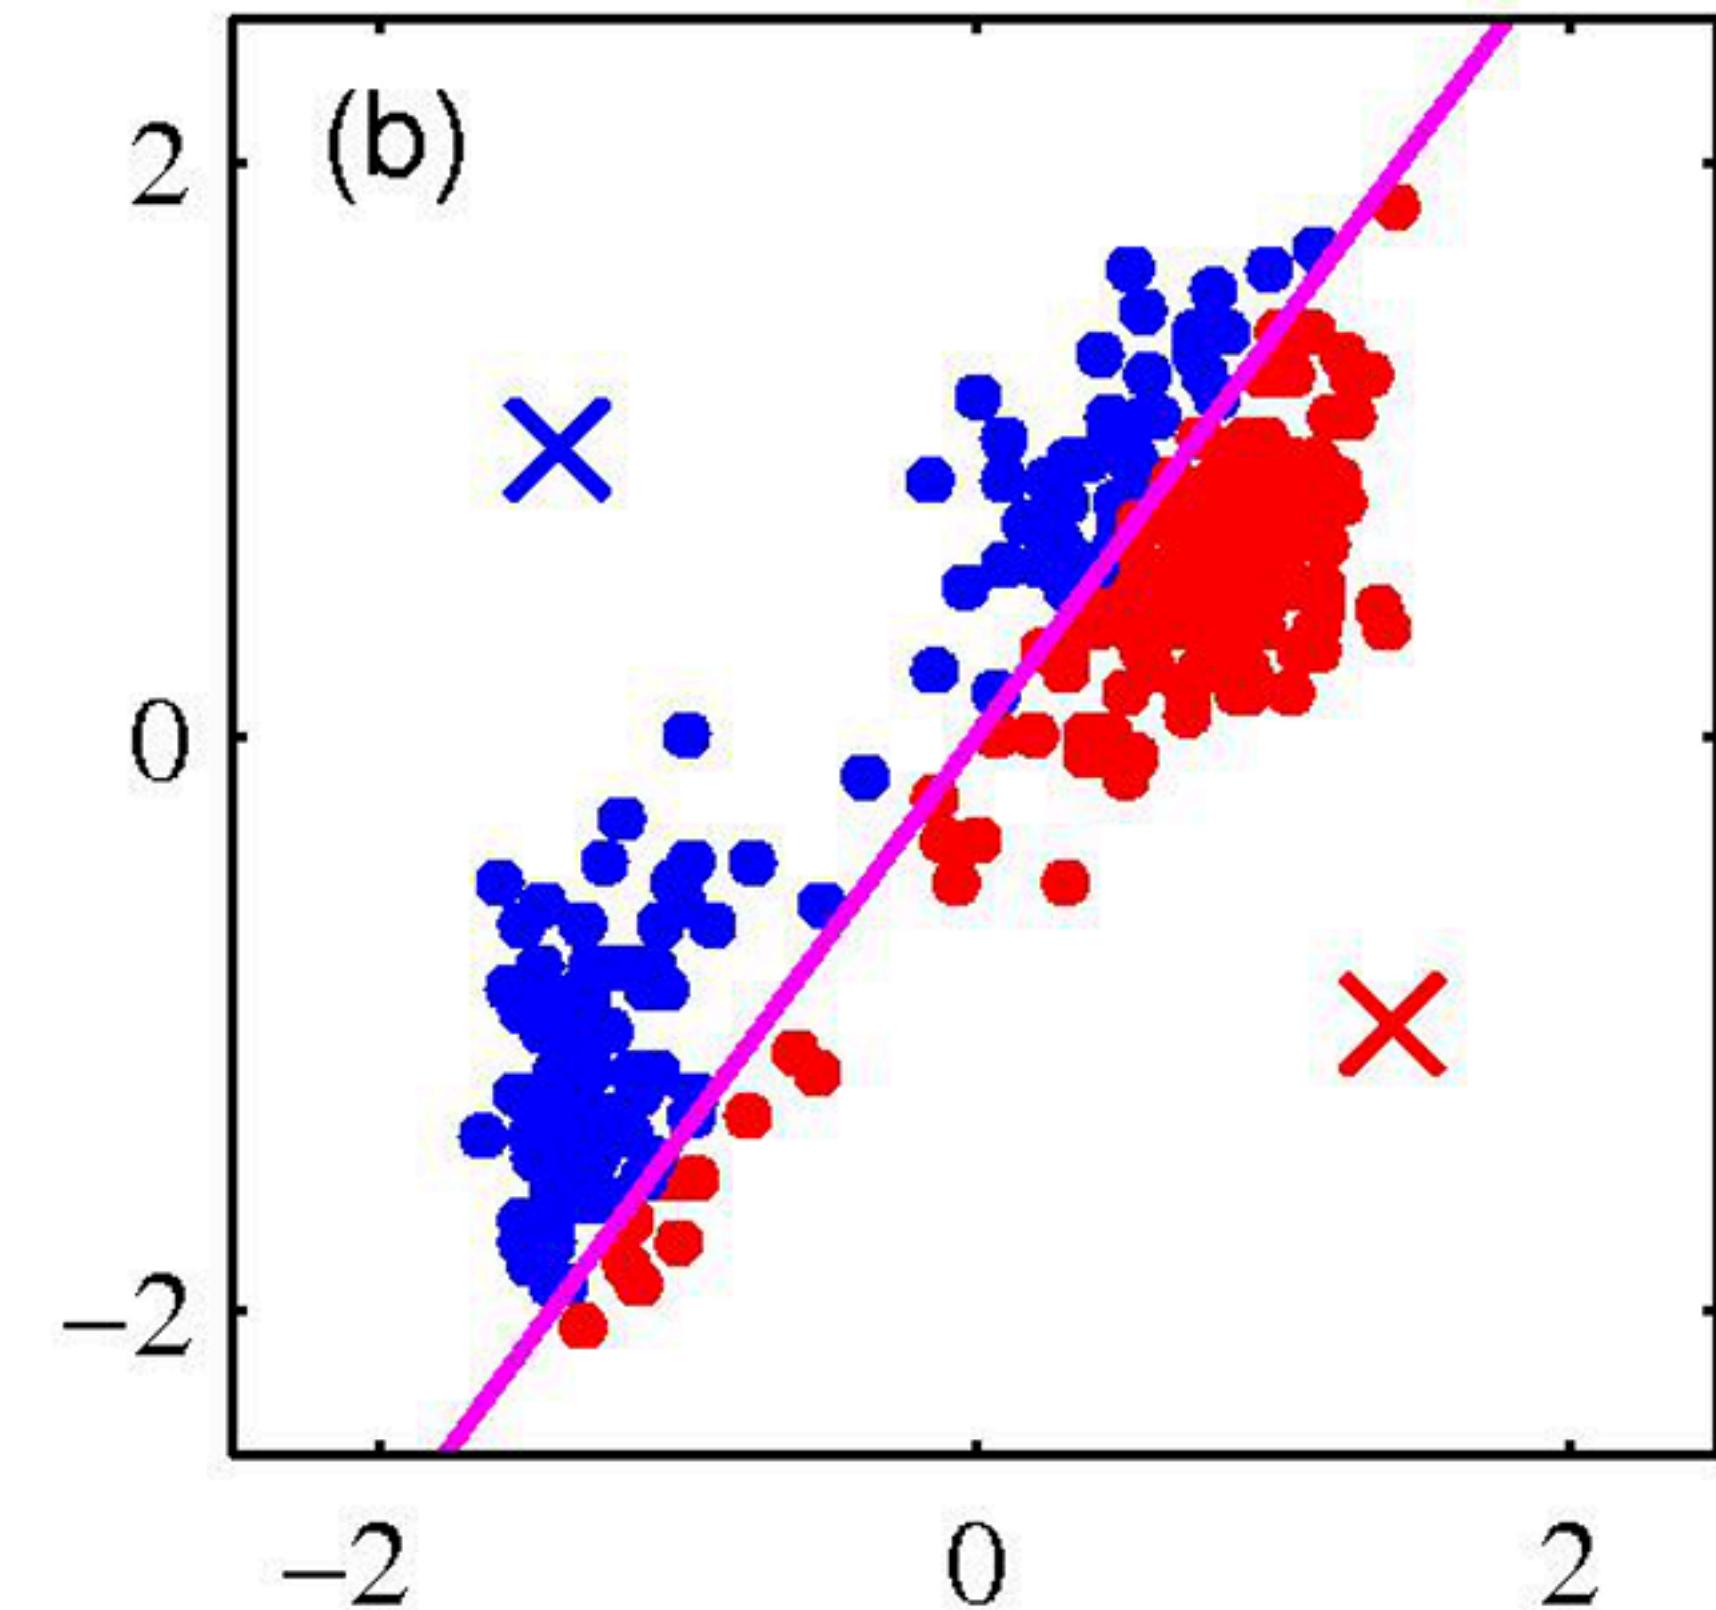
Select K random features as cluster centers.



Source: <https://people.csail.mit.edu/dsontag/courses/ml12/slides/lecture14.pdf>

K-Means

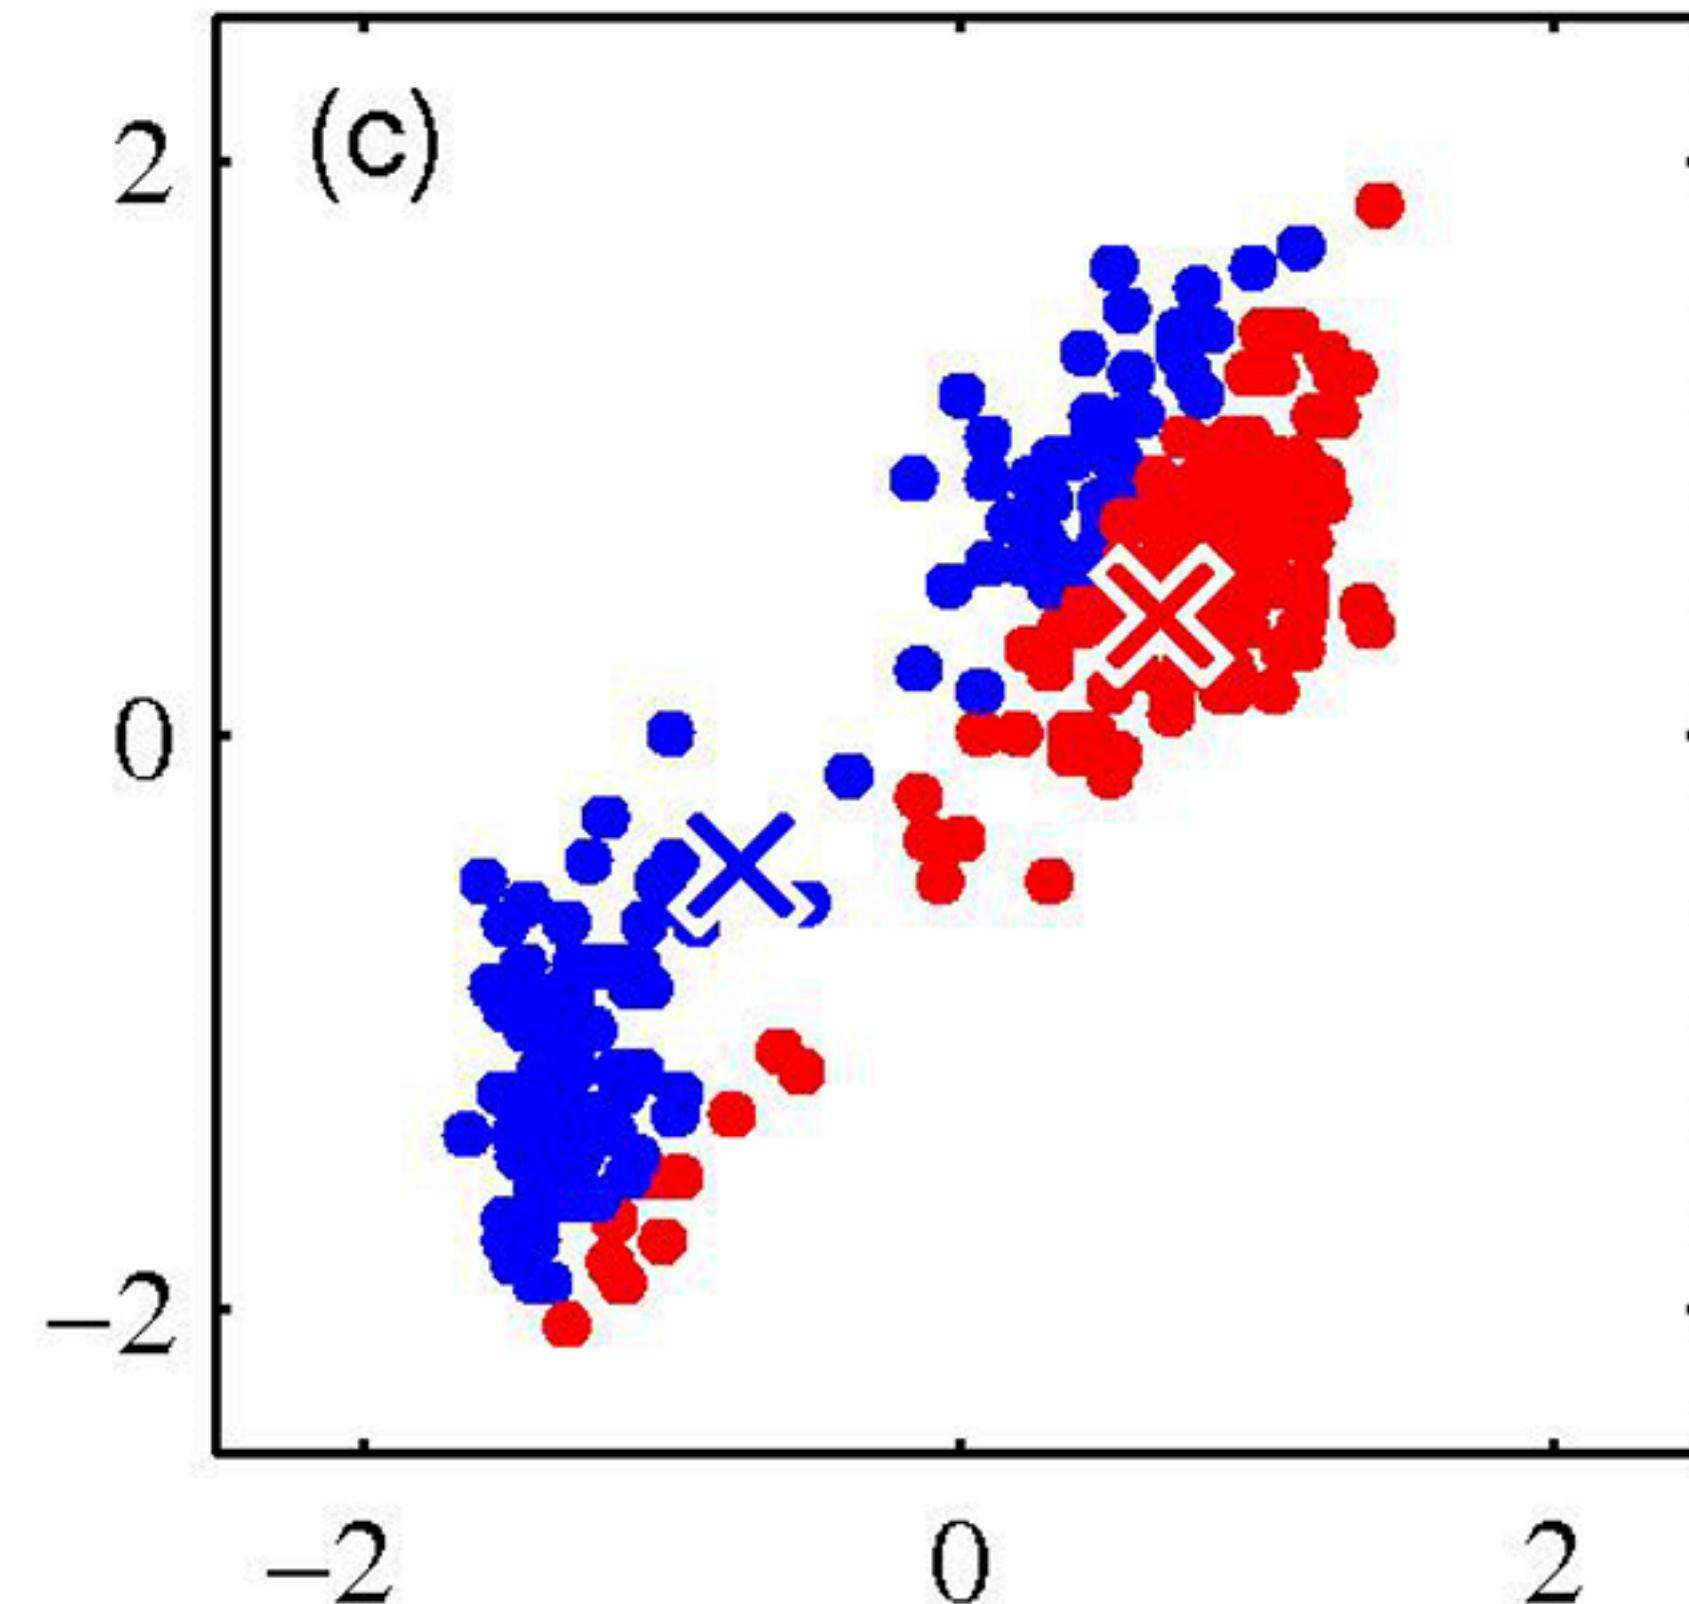
Assign features to the closest cluster centers.



Source: <https://people.csail.mit.edu/dsontag/courses/ml12/slides/lecture14.pdf>

K-Means

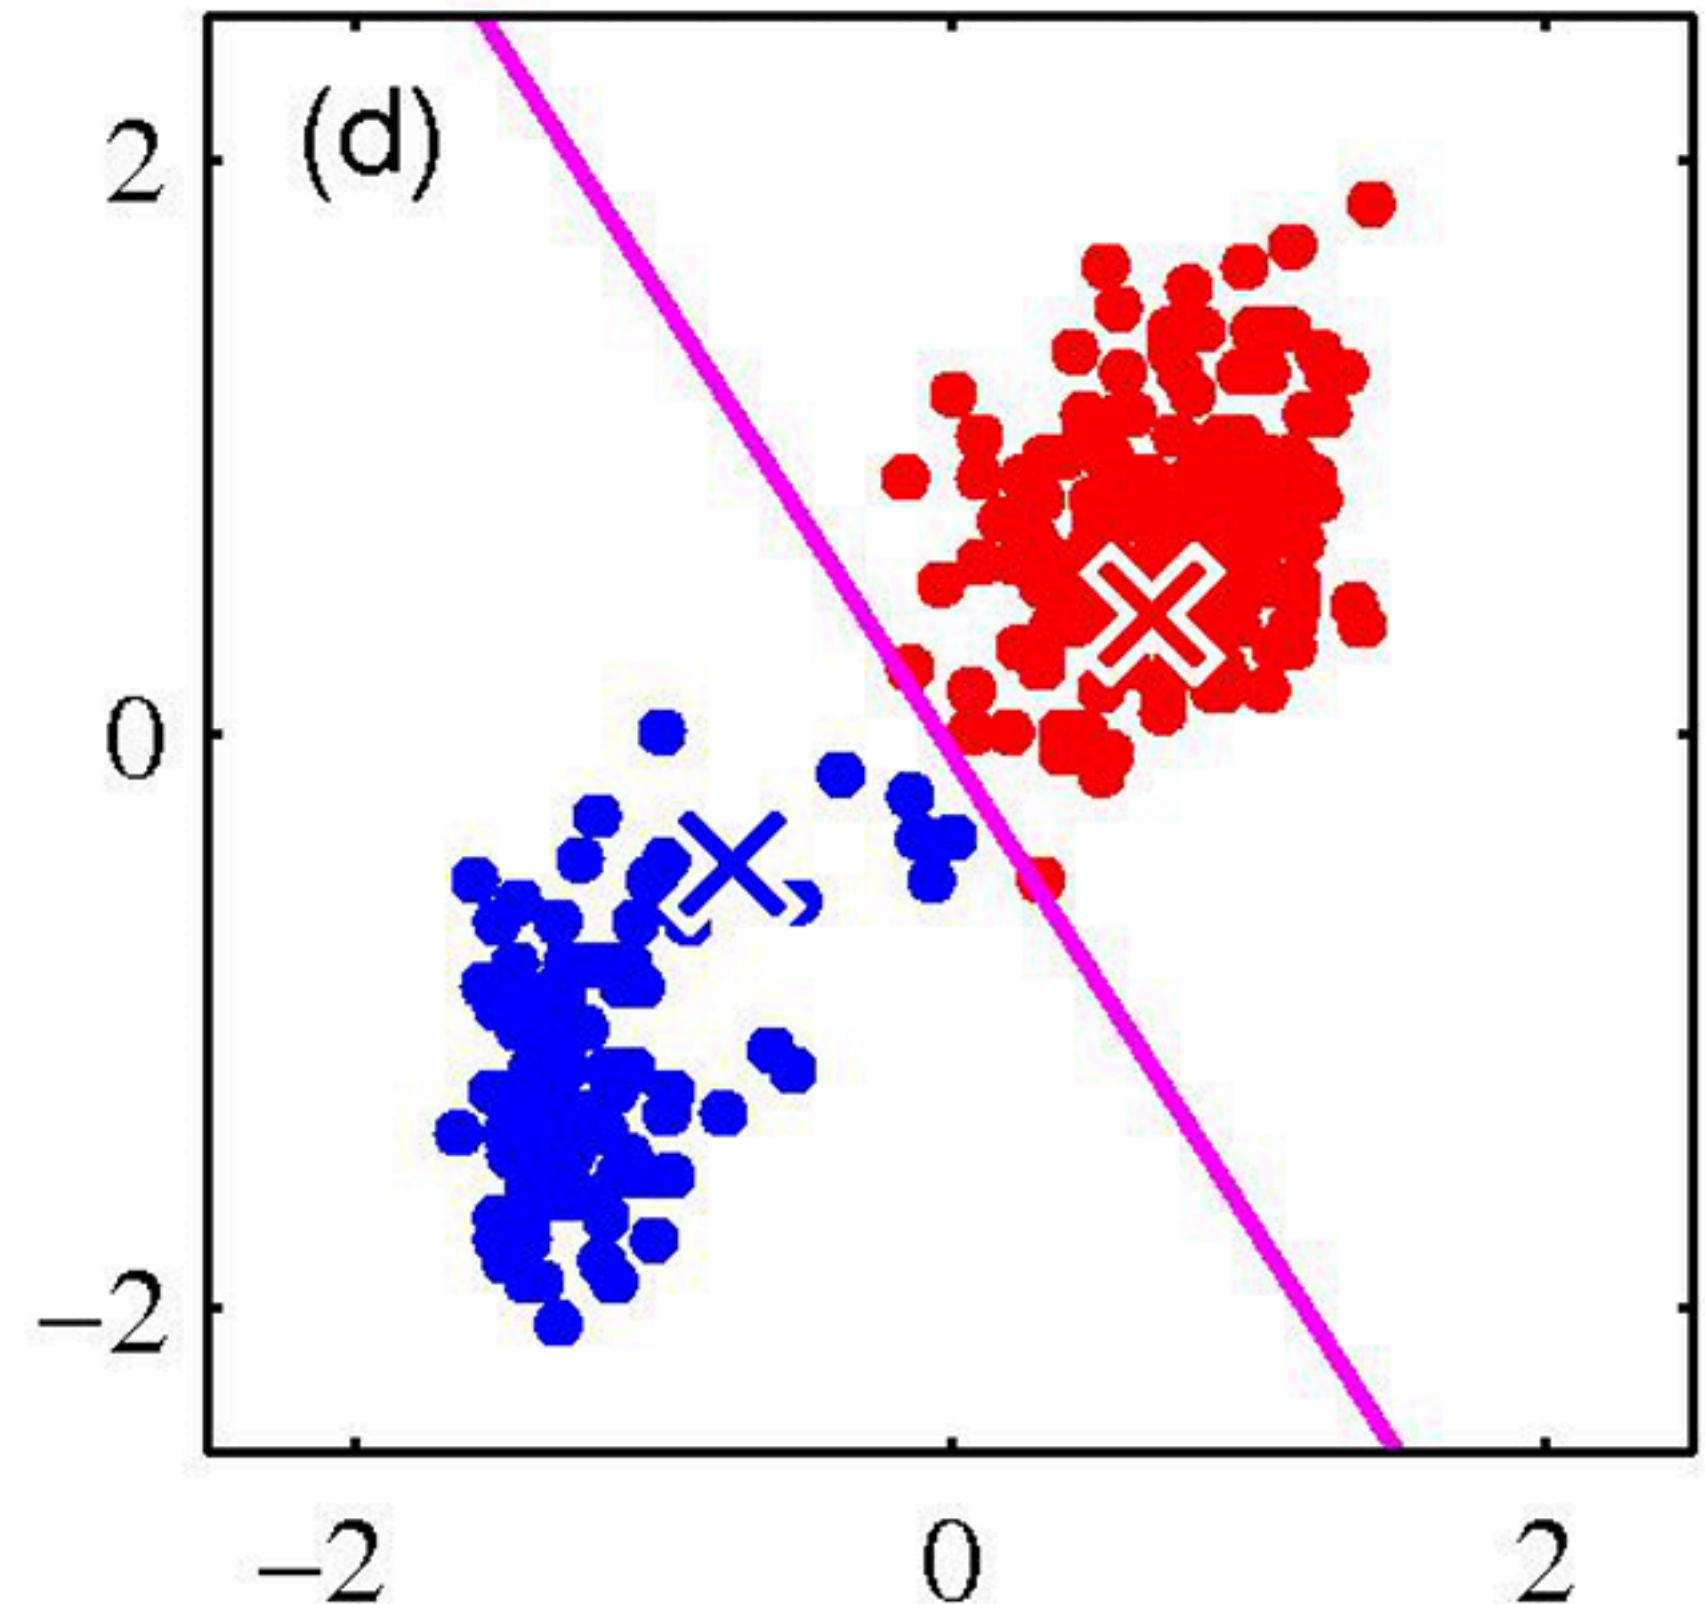
Update the cluster centers by taking the **means** of each cluster.



Source: <https://people.csail.mit.edu/dsontag/courses/ml12/slides/lecture14.pdf>

K-Means

Repeat until convergence.



Source: <https://people.csail.mit.edu/dsontag/courses/ml12/slides/lecture14.pdf>

K-Means

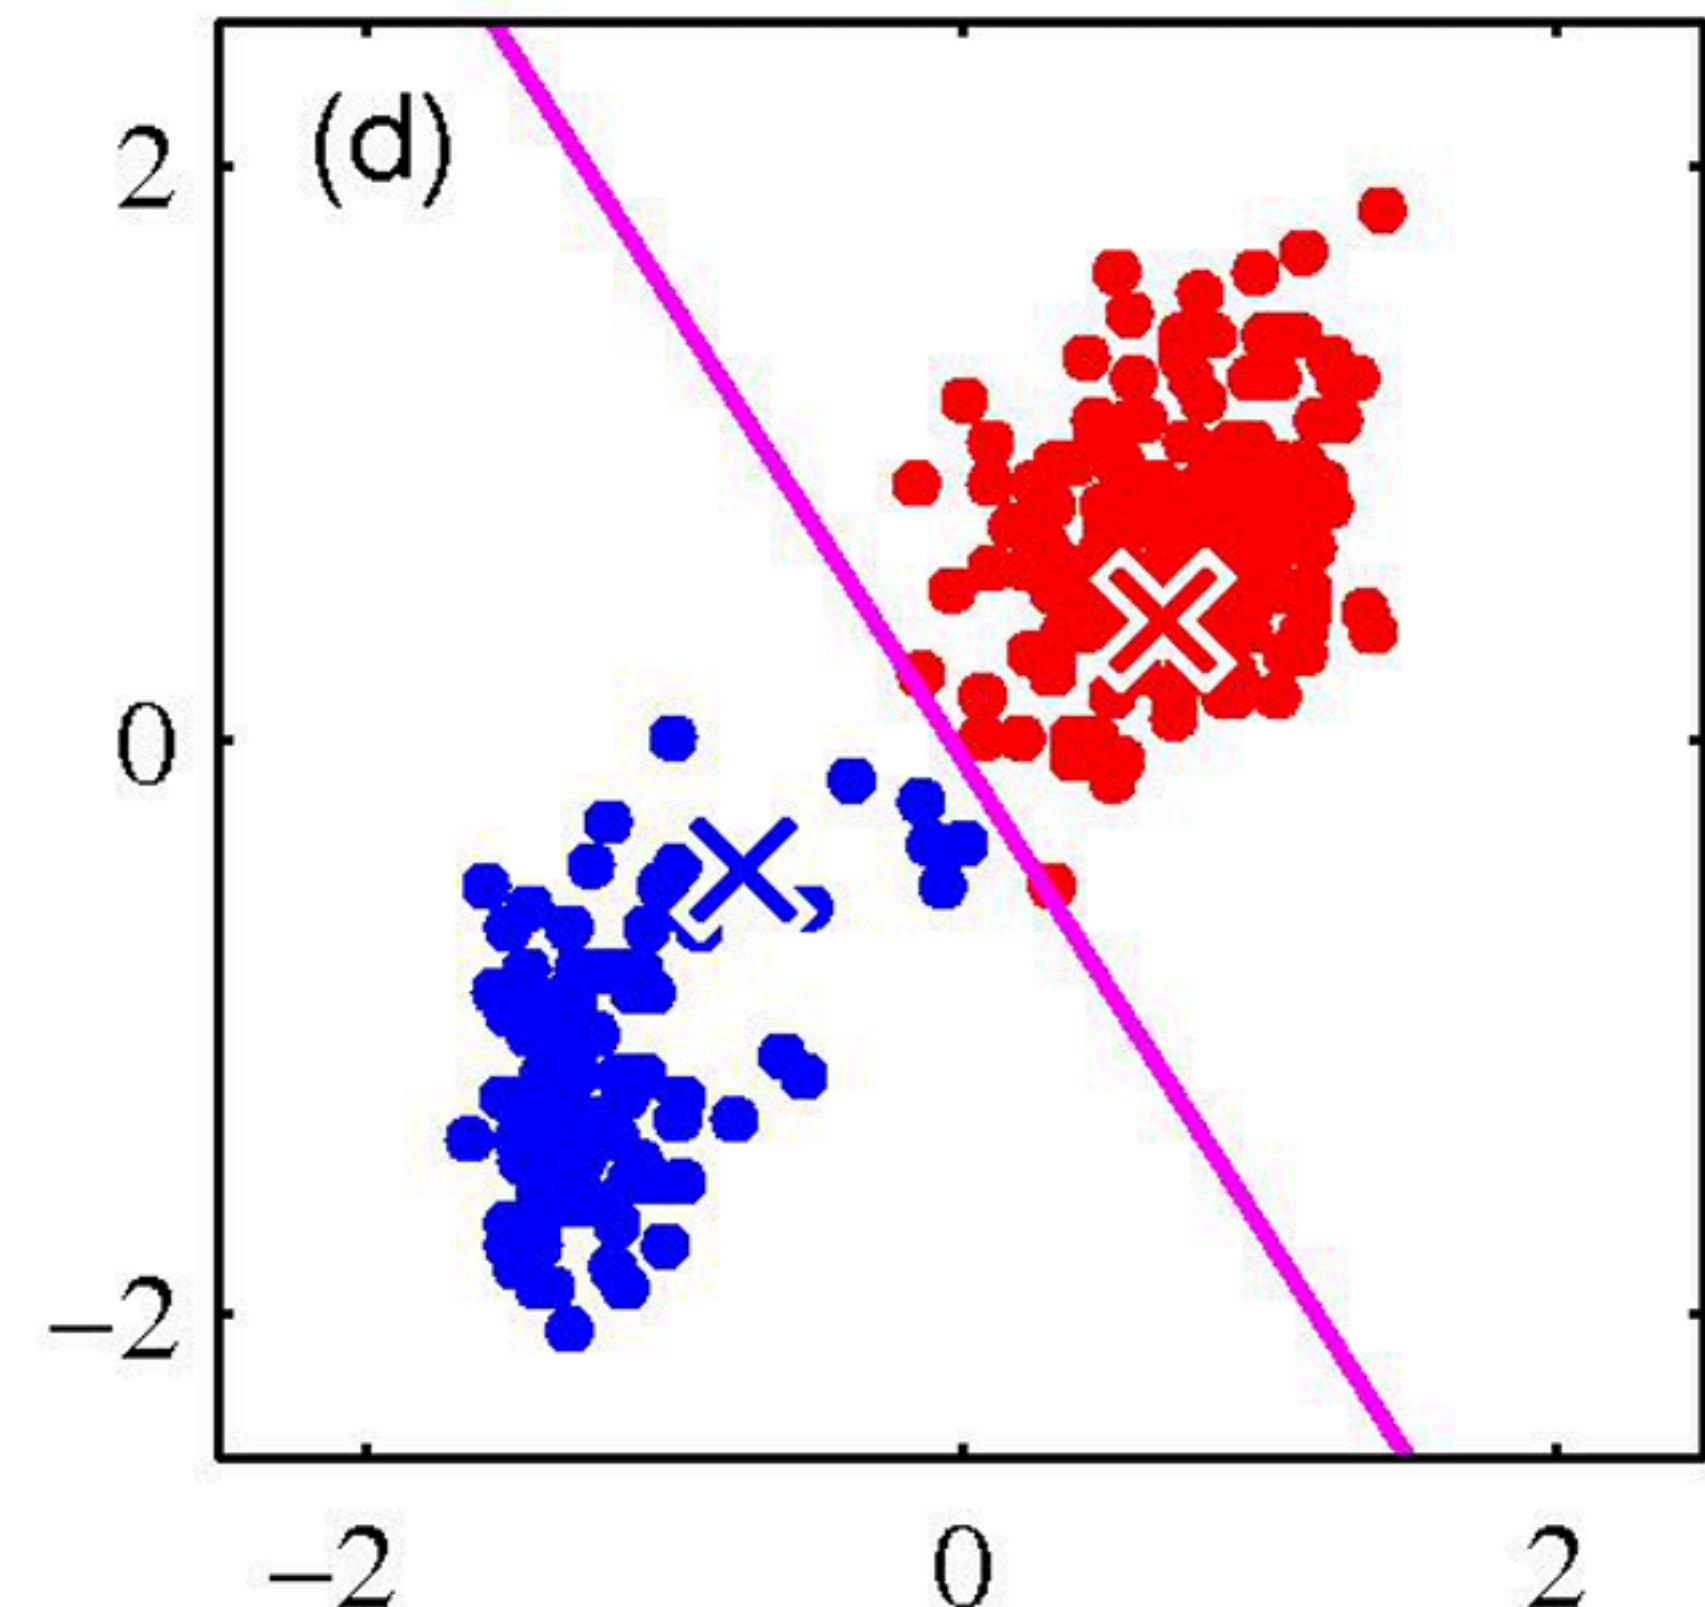
What are the limitations of this approach?

What is the ideal number of clusters?

Complexity of building clusters:
 $O(Kn)$ in each step until convergence.

K: #clusters
n: #features

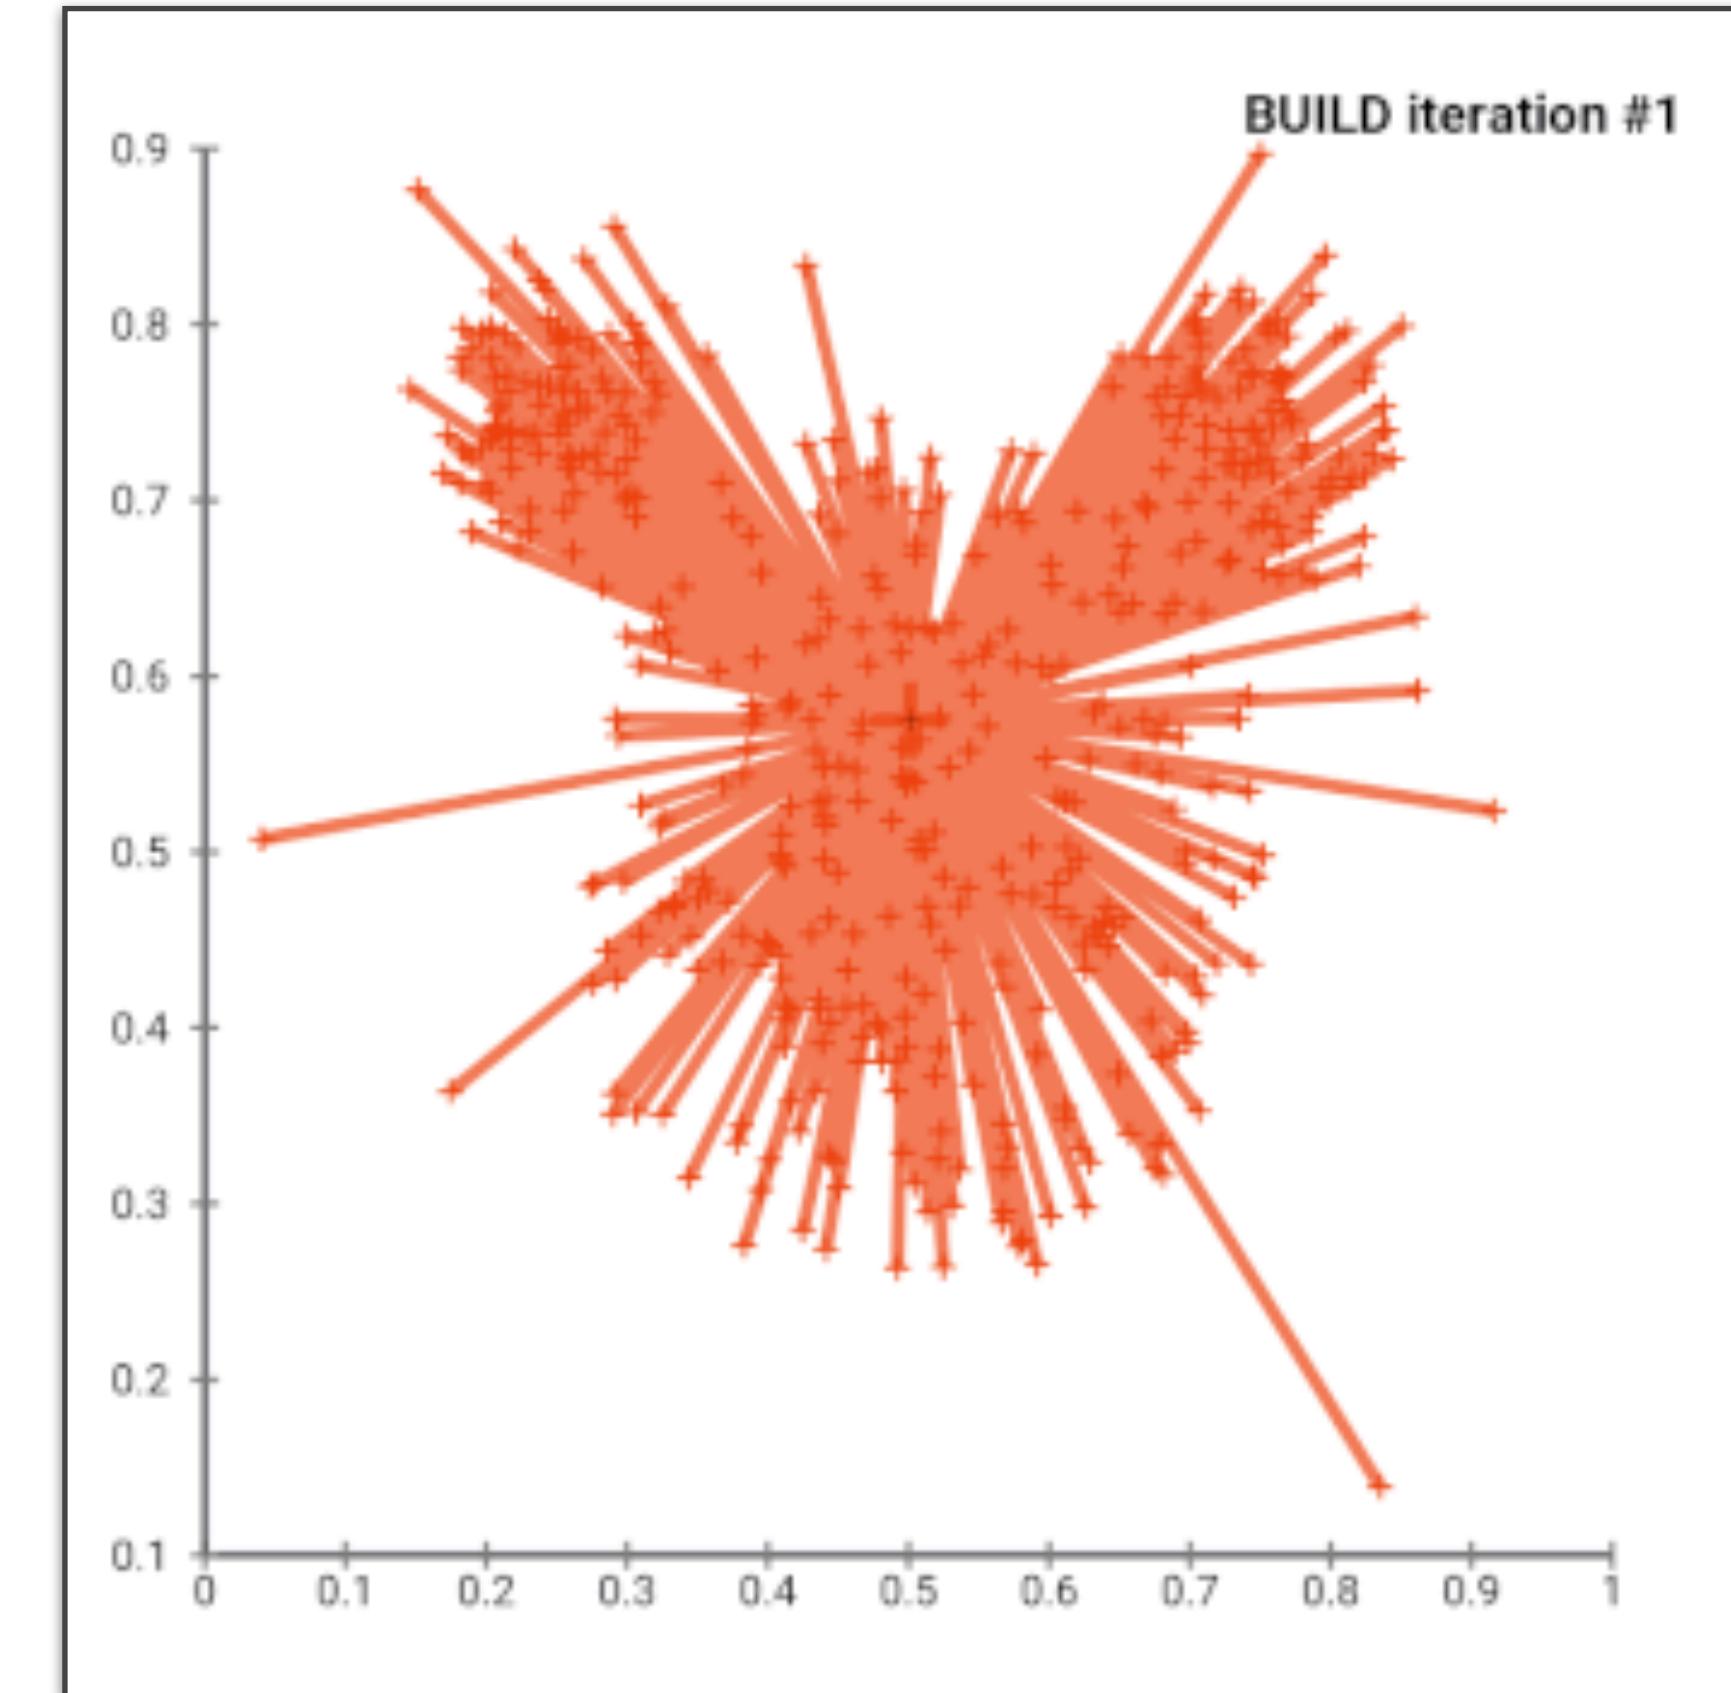
Clustering is *offline*: i.e., it does not happen at feature querying time.



K-Means

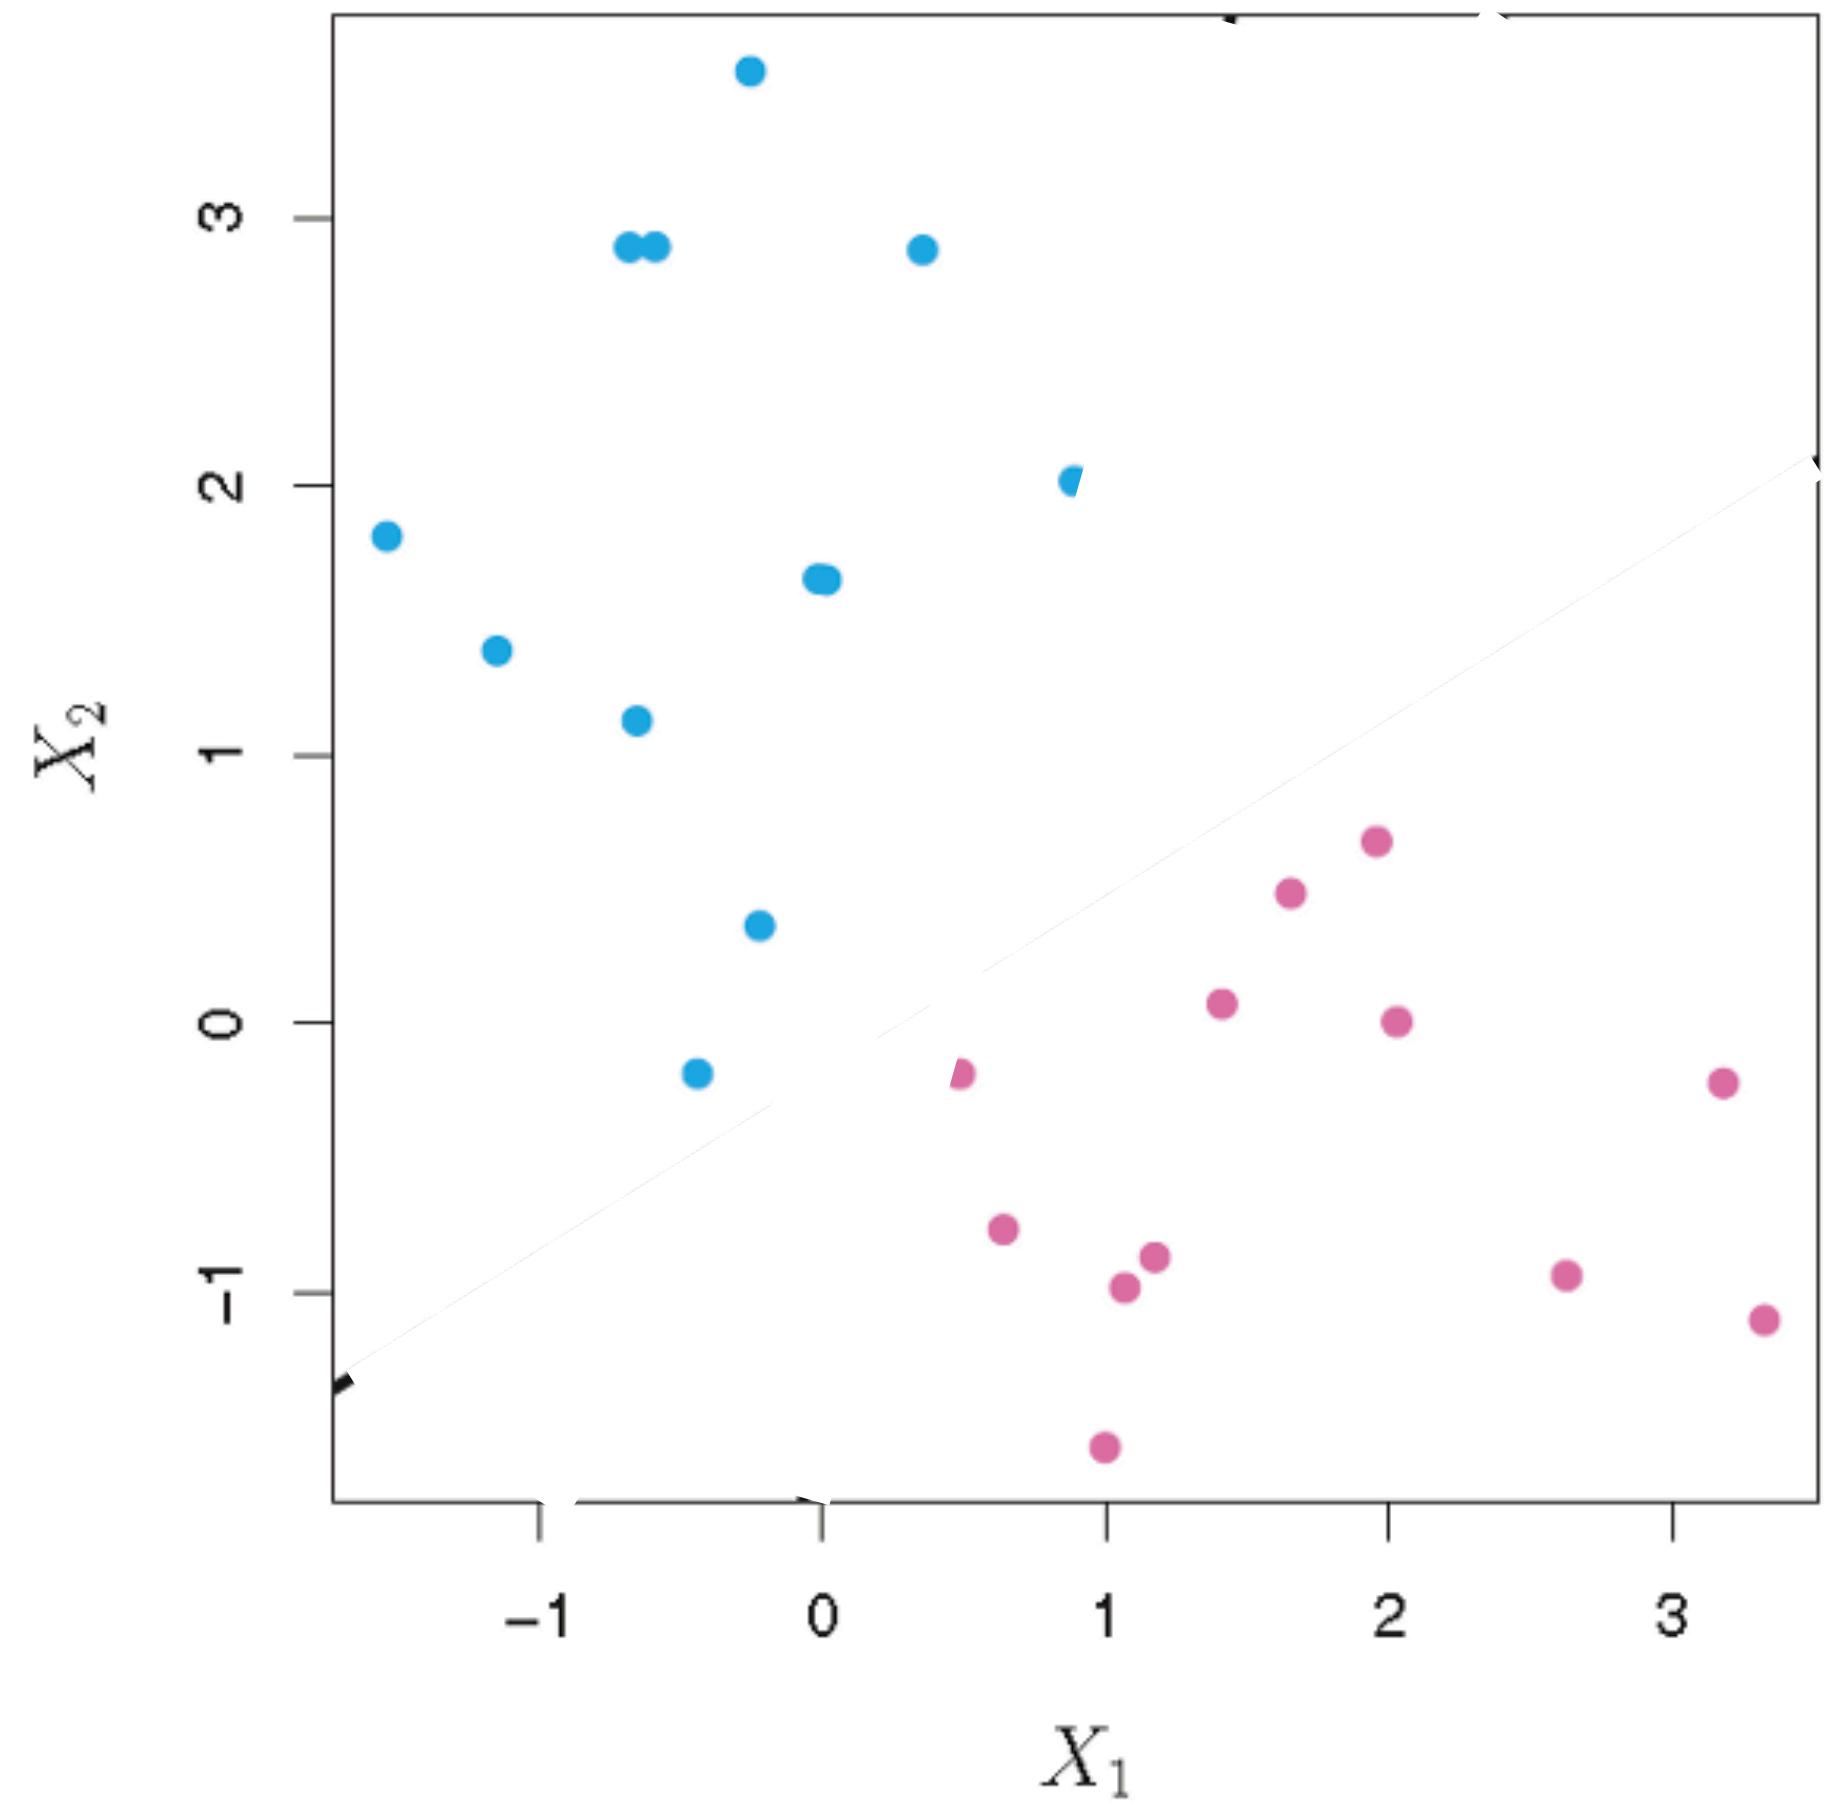
Variation: K-medoids

Instead of using *means* as the cluster centers, use *medians*, which are actual existing features.



Support Vector Machines (SVM)

**How to Separate
these Features?**

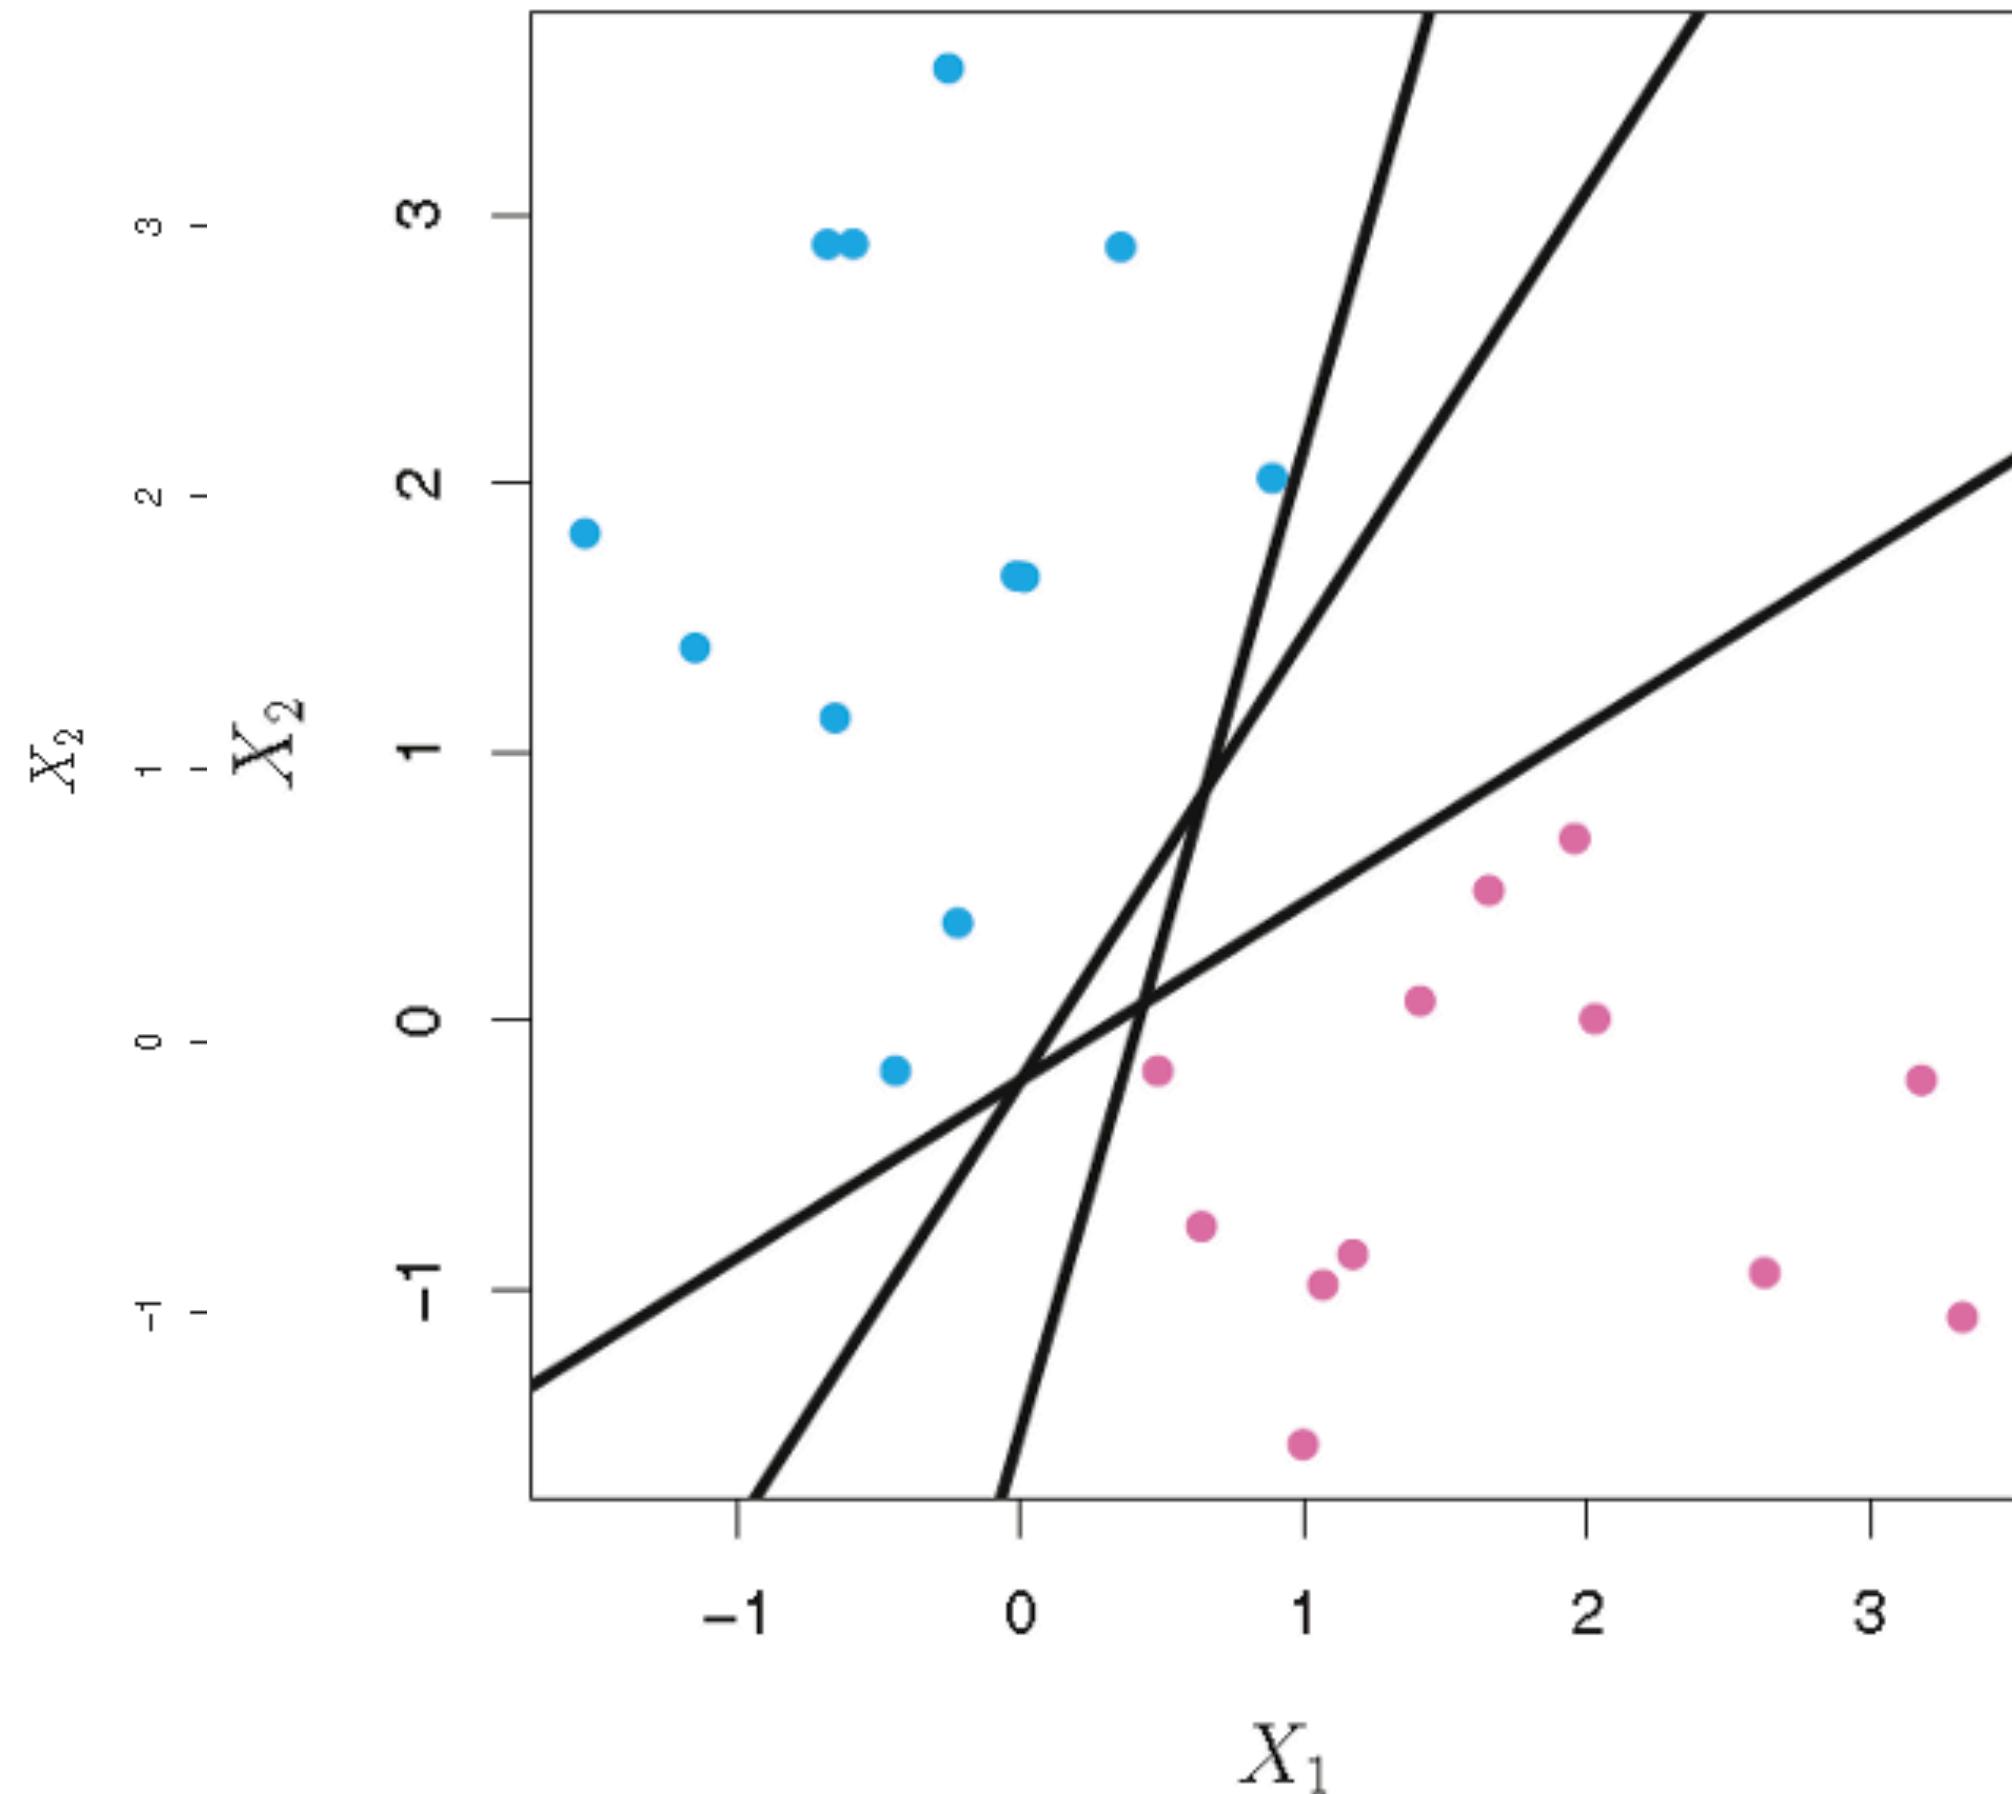


[https://towardsdatascience.com/
the-kernel-trick-c98cdbcaeb3f](https://towardsdatascience.com/the-kernel-trick-c98cdbcaeb3f)

Support Vector Machines (SVM)

**How to Separate
these Features?**

They're linearly separable,
but what is the best
separation?



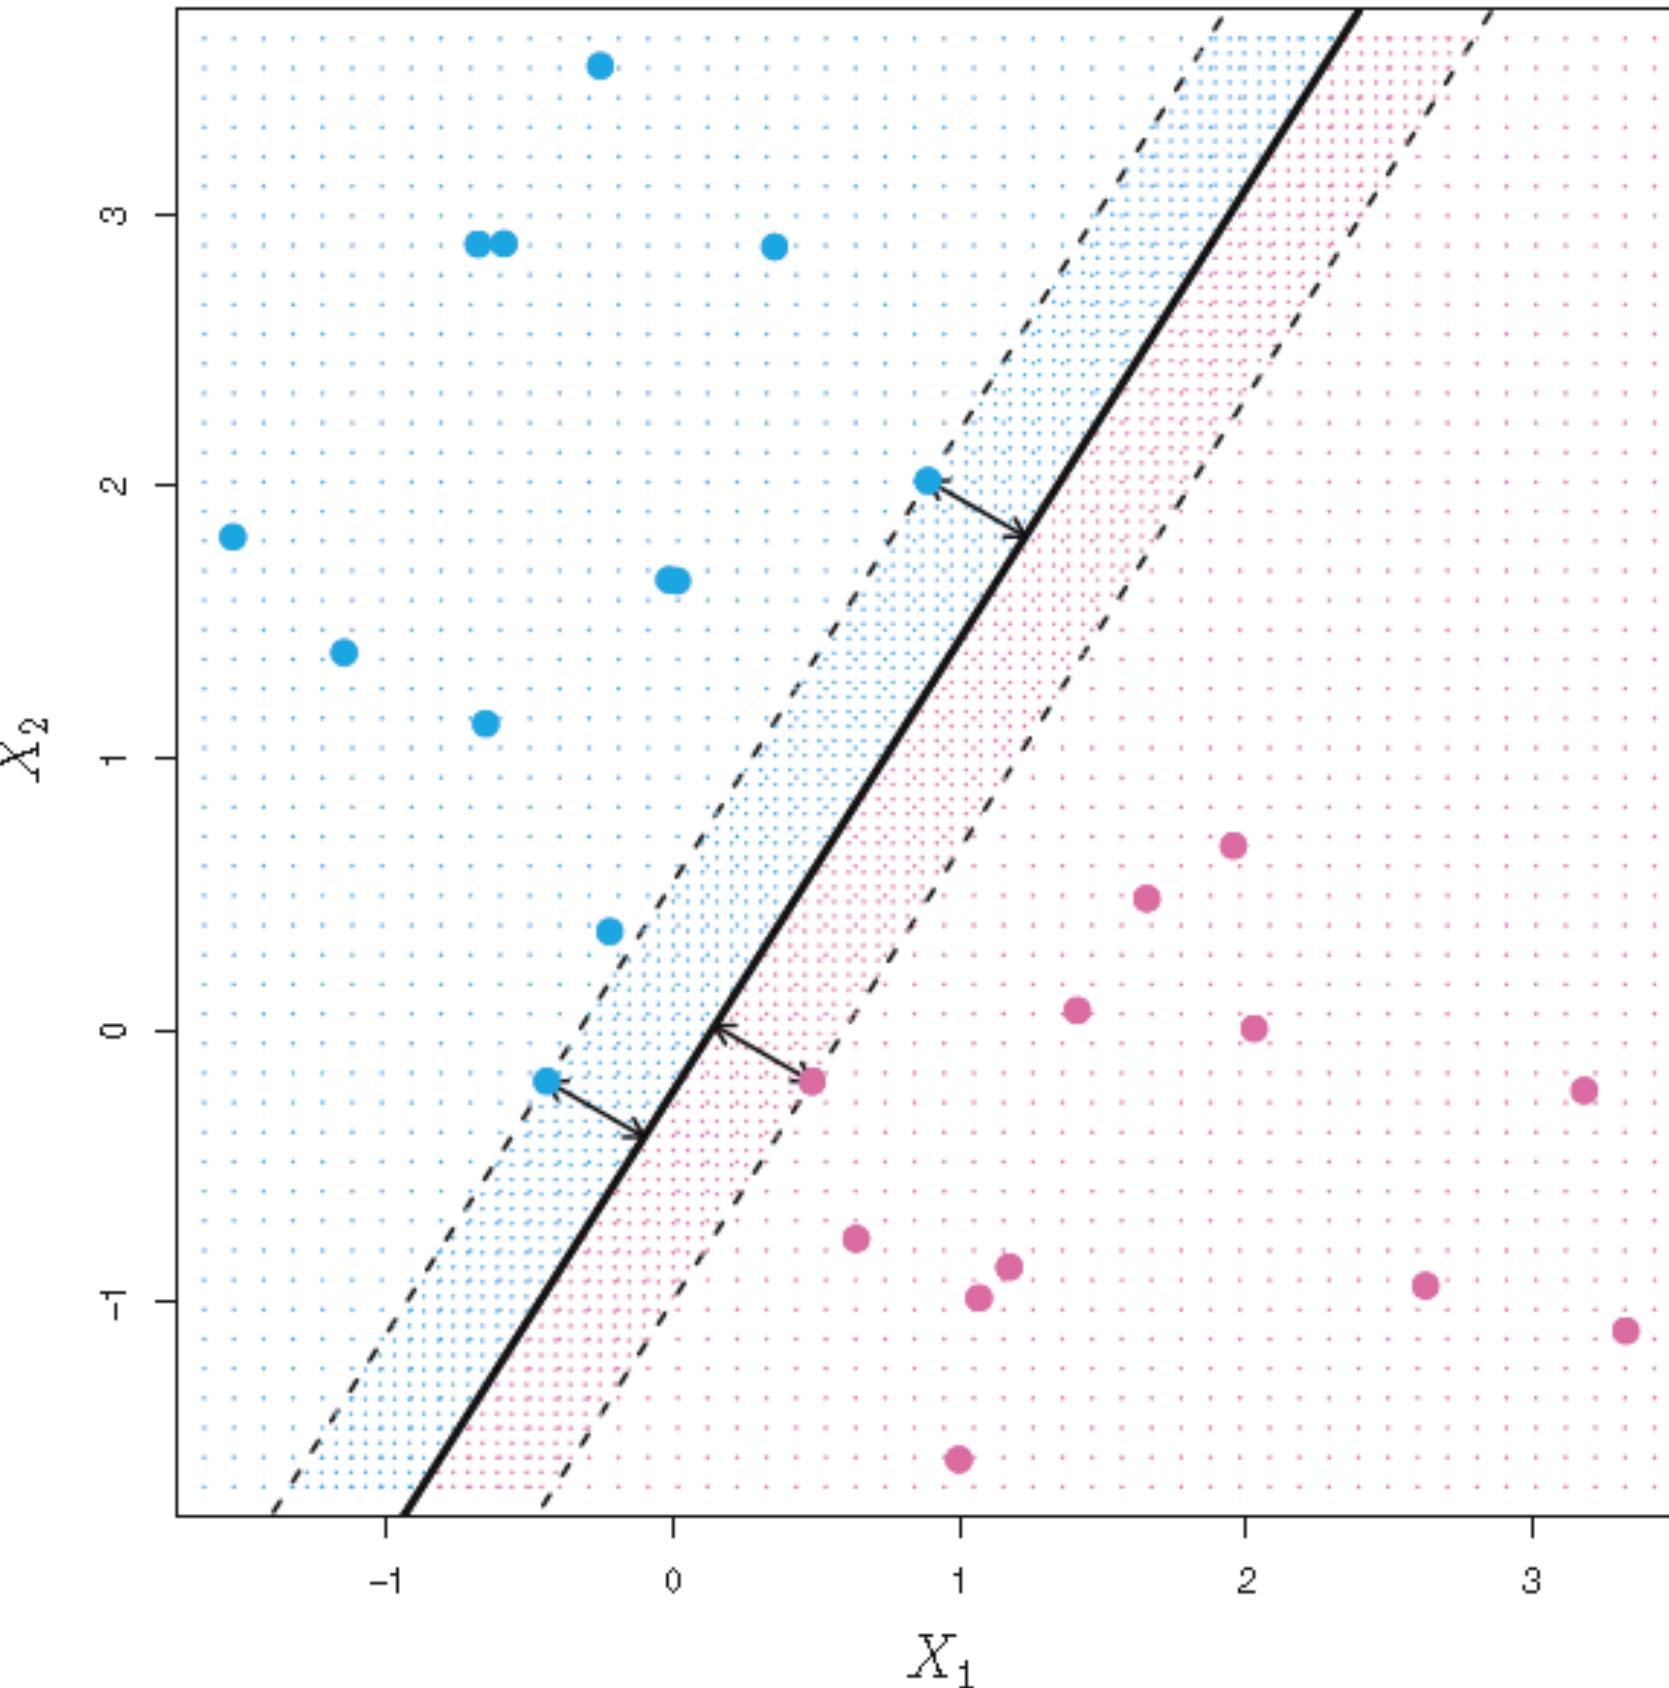
[https://towardsdatascience.com/
the-kernel-trick-c98cdbcaeb3f](https://towardsdatascience.com/the-kernel-trick-c98cdbcaeb3f)

Support Vector Machines (SVM)

**How to Separate
these Features?**

They're linearly separable,
but what is the best
separation?

Solution:
Find the hyperplane that
maximizes the margin
between the classes.



[https://towardsdatascience.com/
the-kernel-trick-c98cdbcaeb3f](https://towardsdatascience.com/the-kernel-trick-c98cdbcaeb3f)

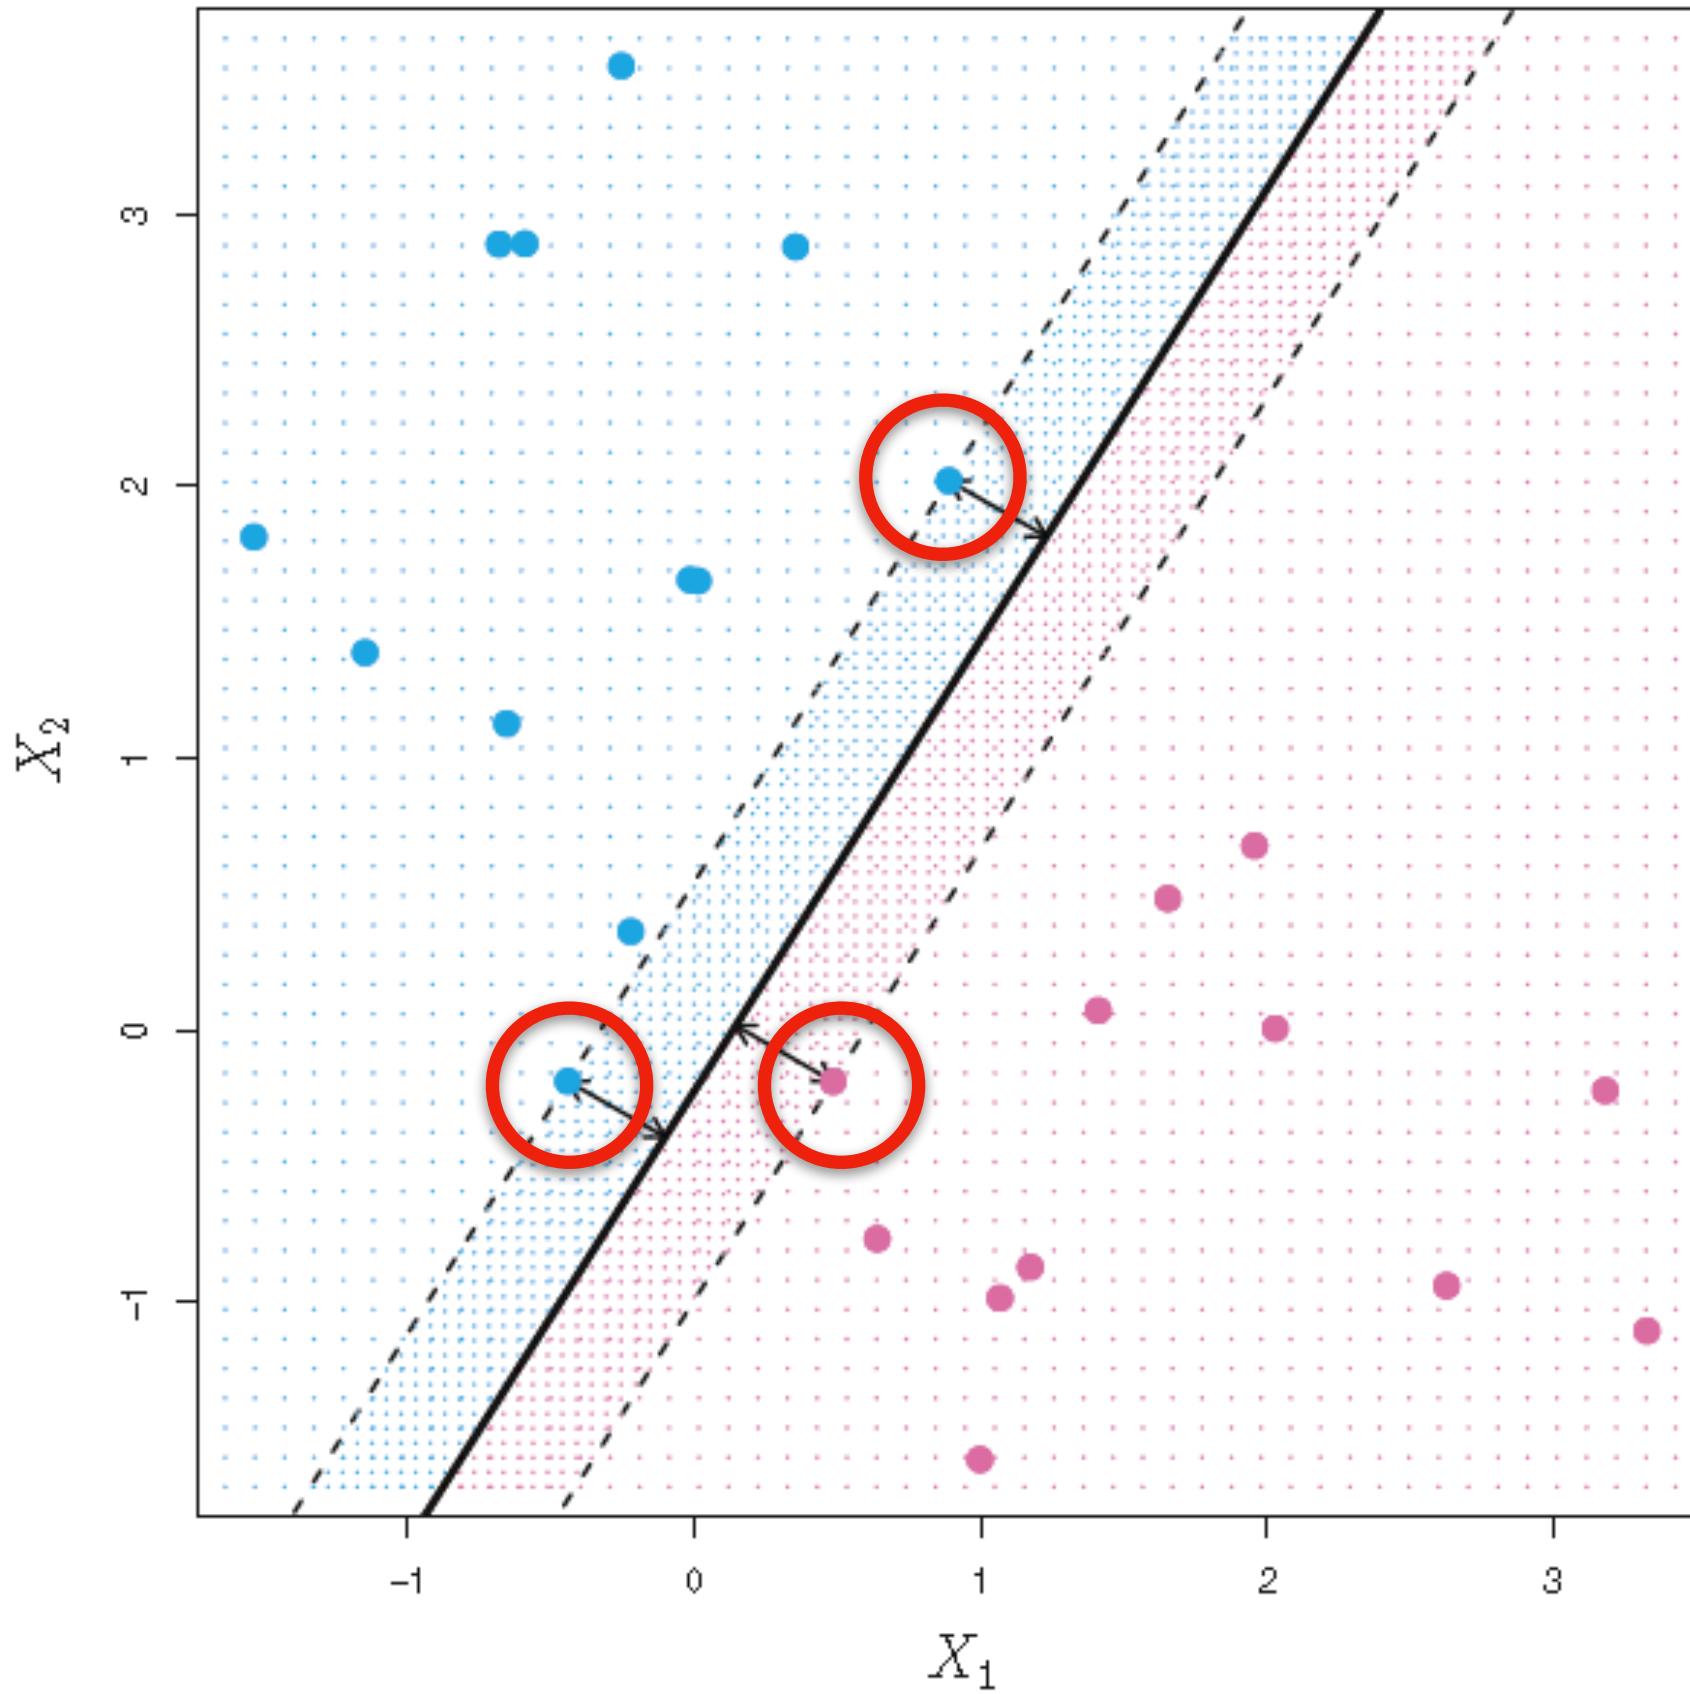


LOYOLA
UNIVERSITY CHICAGO

Support Vector Machines (SVM)

How to Separate these Features?

The feature vectors serving as reference to the margin of the separation hyperplane are called **support vectors**.

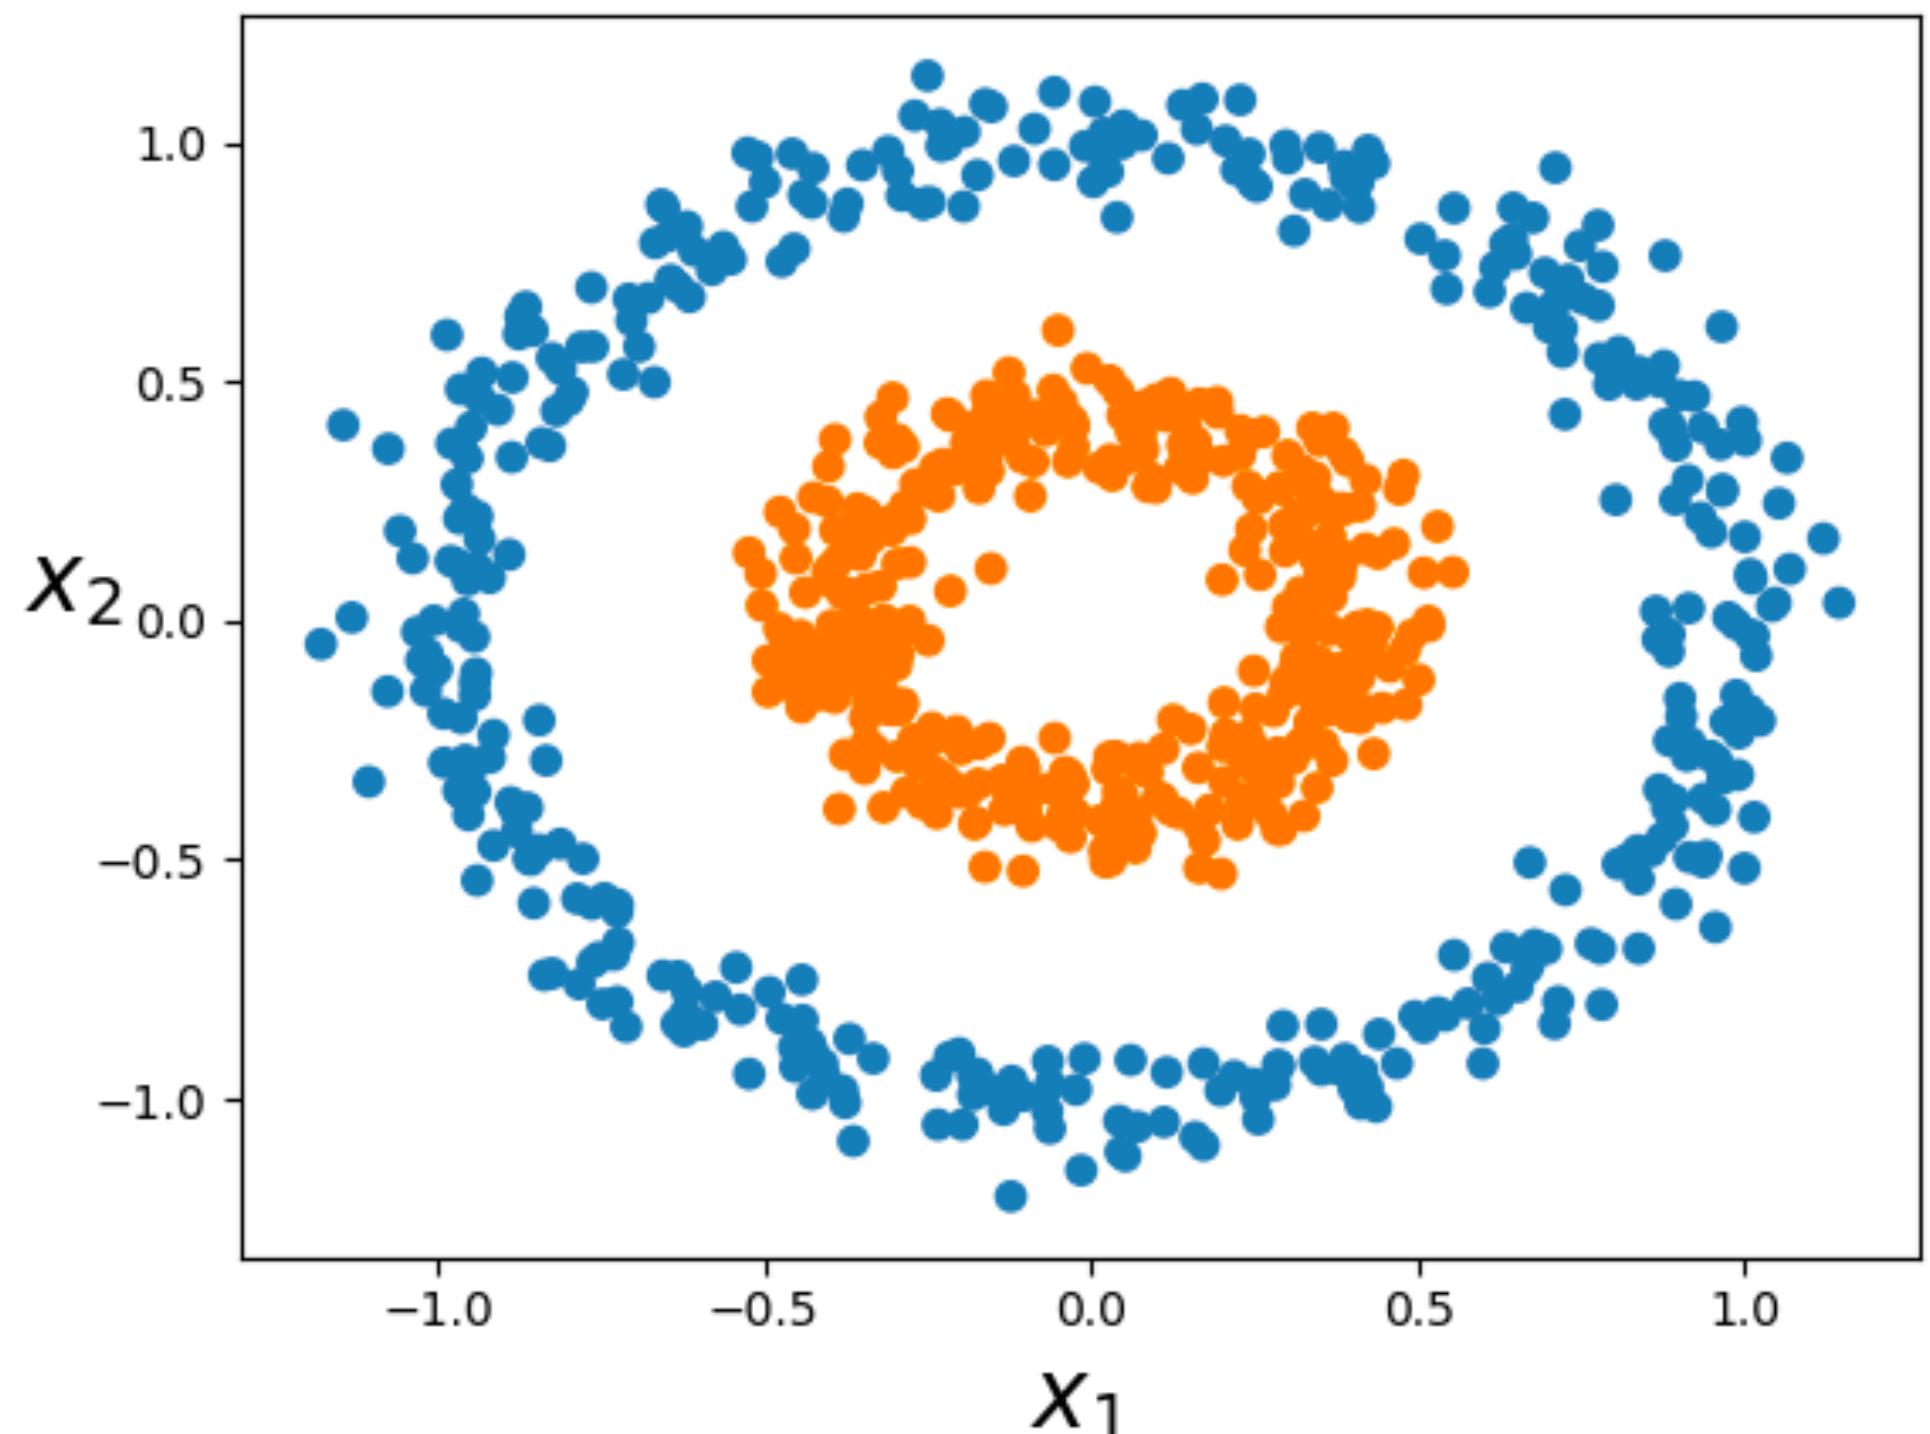


<https://towardsdatascience.com/the-kernel-trick-c98cdbcaeb3f>

Support Vector Machines (SVM)

**How to Separate
these Features?**

How to deal with non-linearly
separable spaces?



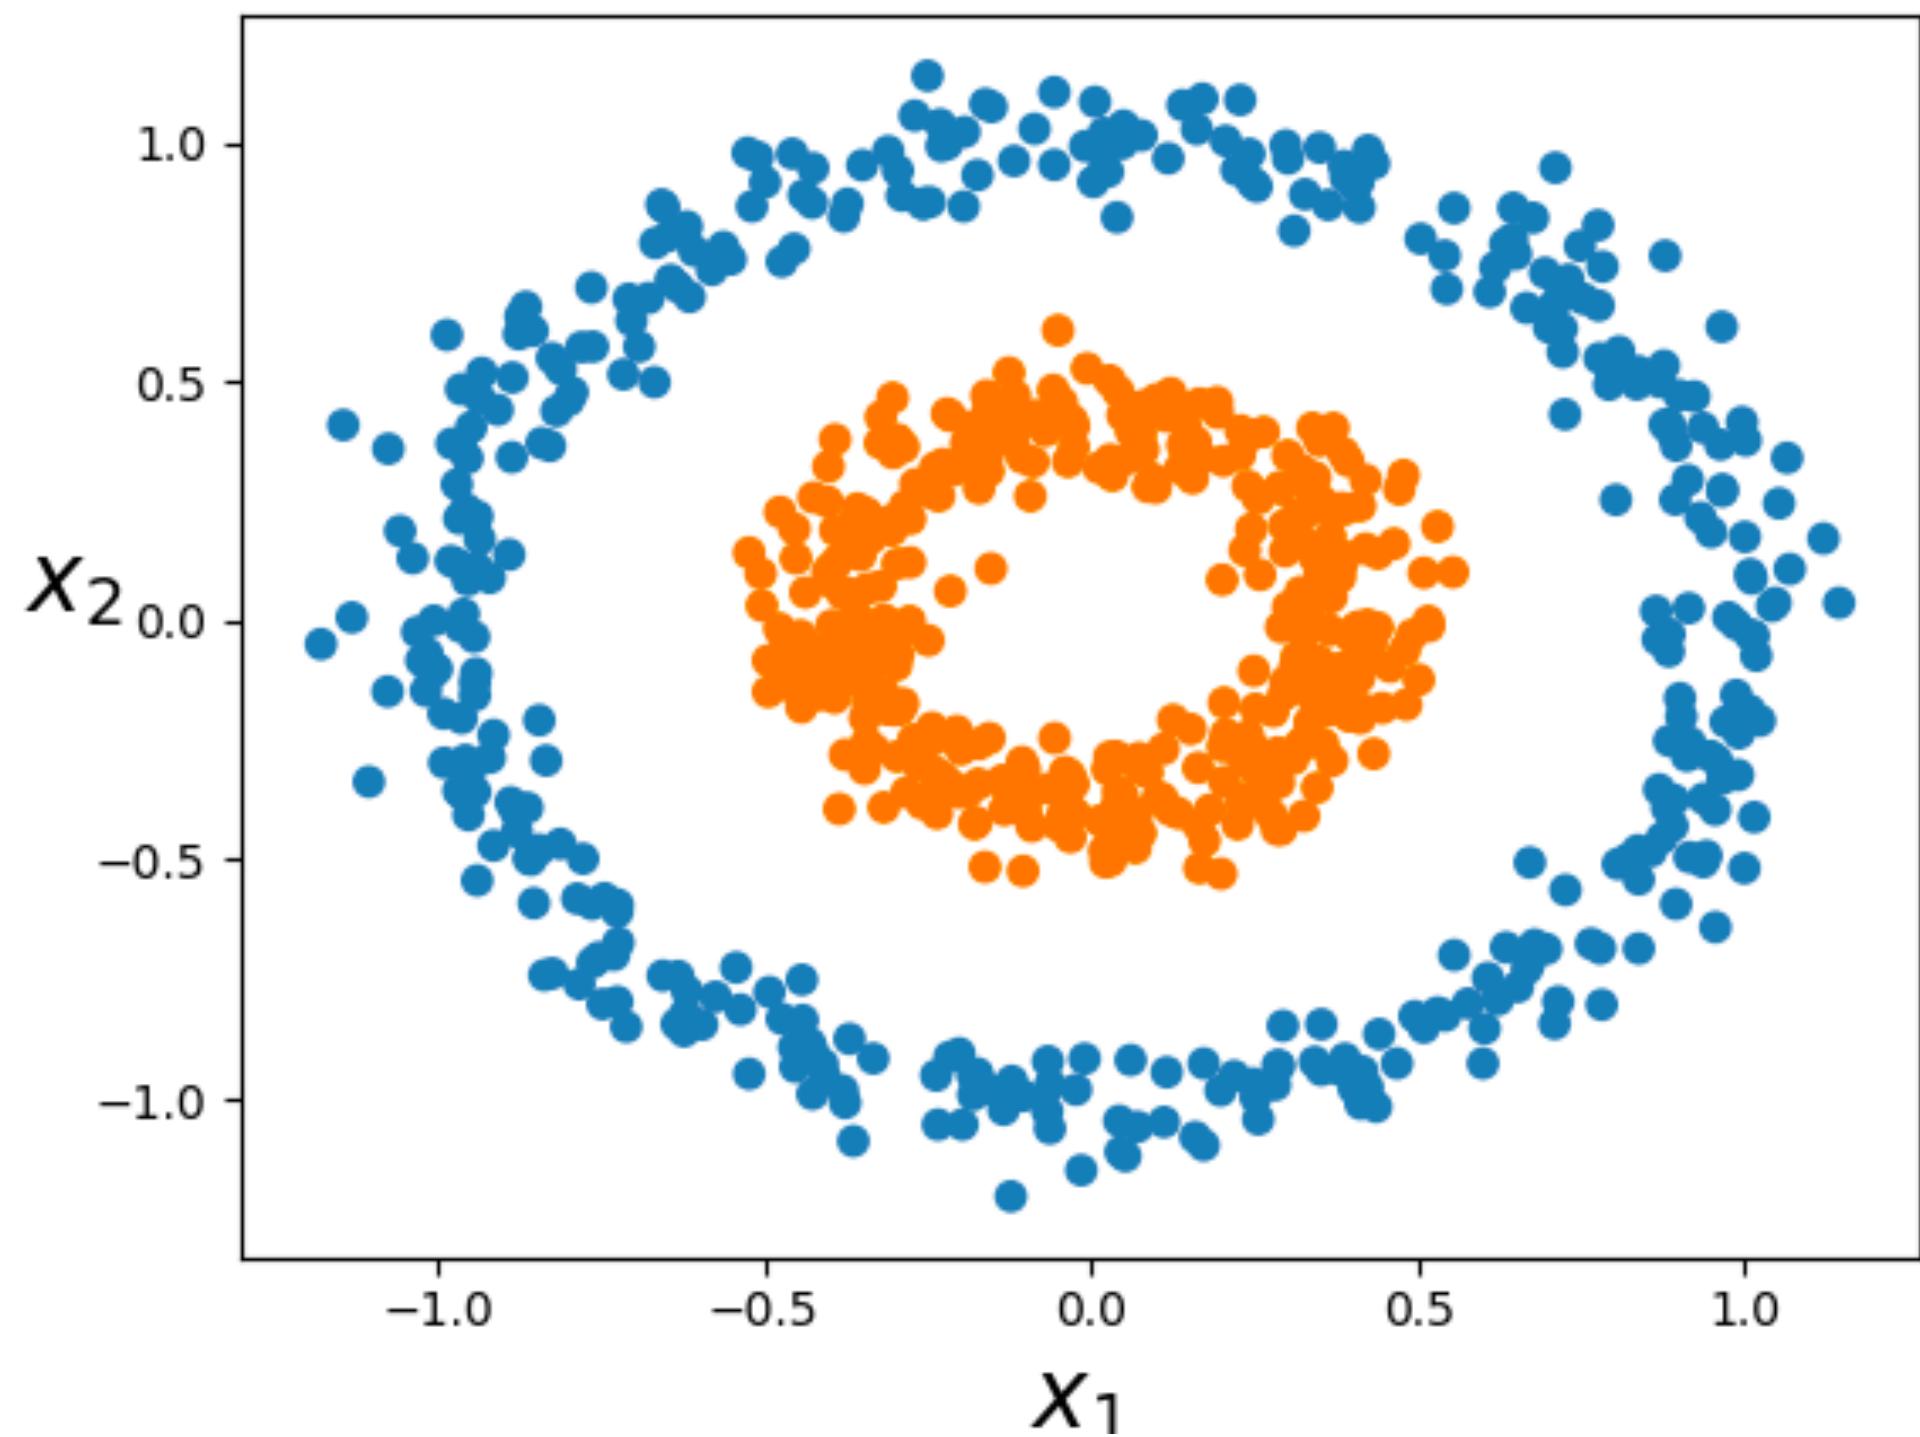
[https://towardsdatascience.com/
the-kernel-trick-c98cdbcaeb3f](https://towardsdatascience.com/the-kernel-trick-c98cdbcaeb3f)

Support Vector Machines (SVM)

**How to Separate
these Features?**

How to deal with non-linearly
separable spaces?

Solution: **kernel trick**.
Transform the data to
higher-dimensional spaces where
they are linearly separable.



[https://towardsdatascience.com/
the-kernel-trick-c98cdbcaeb3f](https://towardsdatascience.com/the-kernel-trick-c98cdbcaeb3f)

Use a kernel function for that (e.g., radial basis).

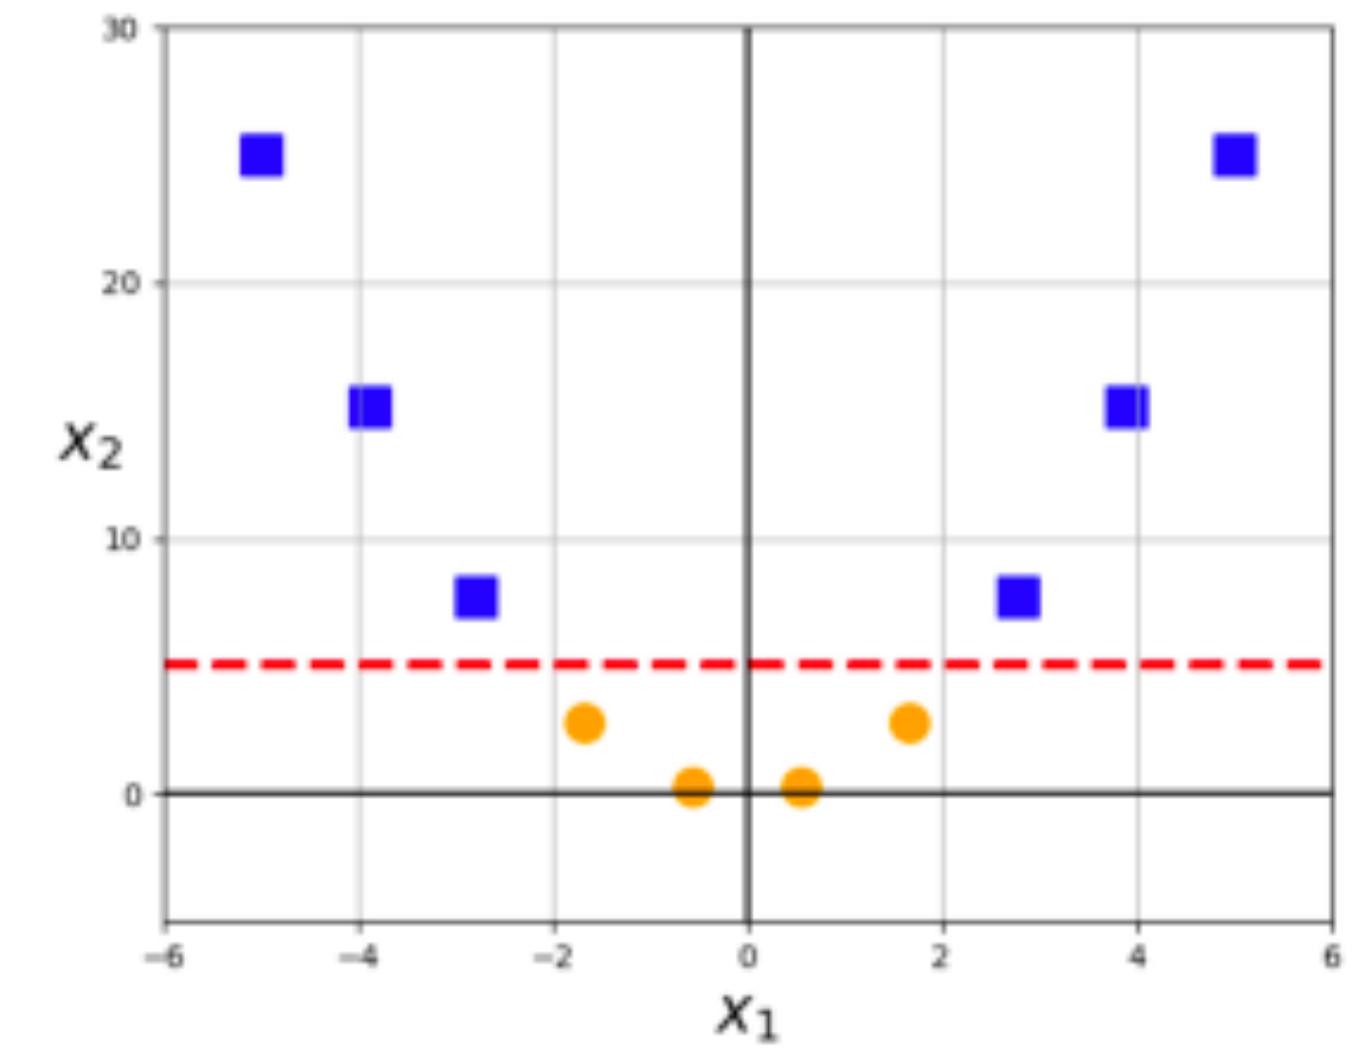
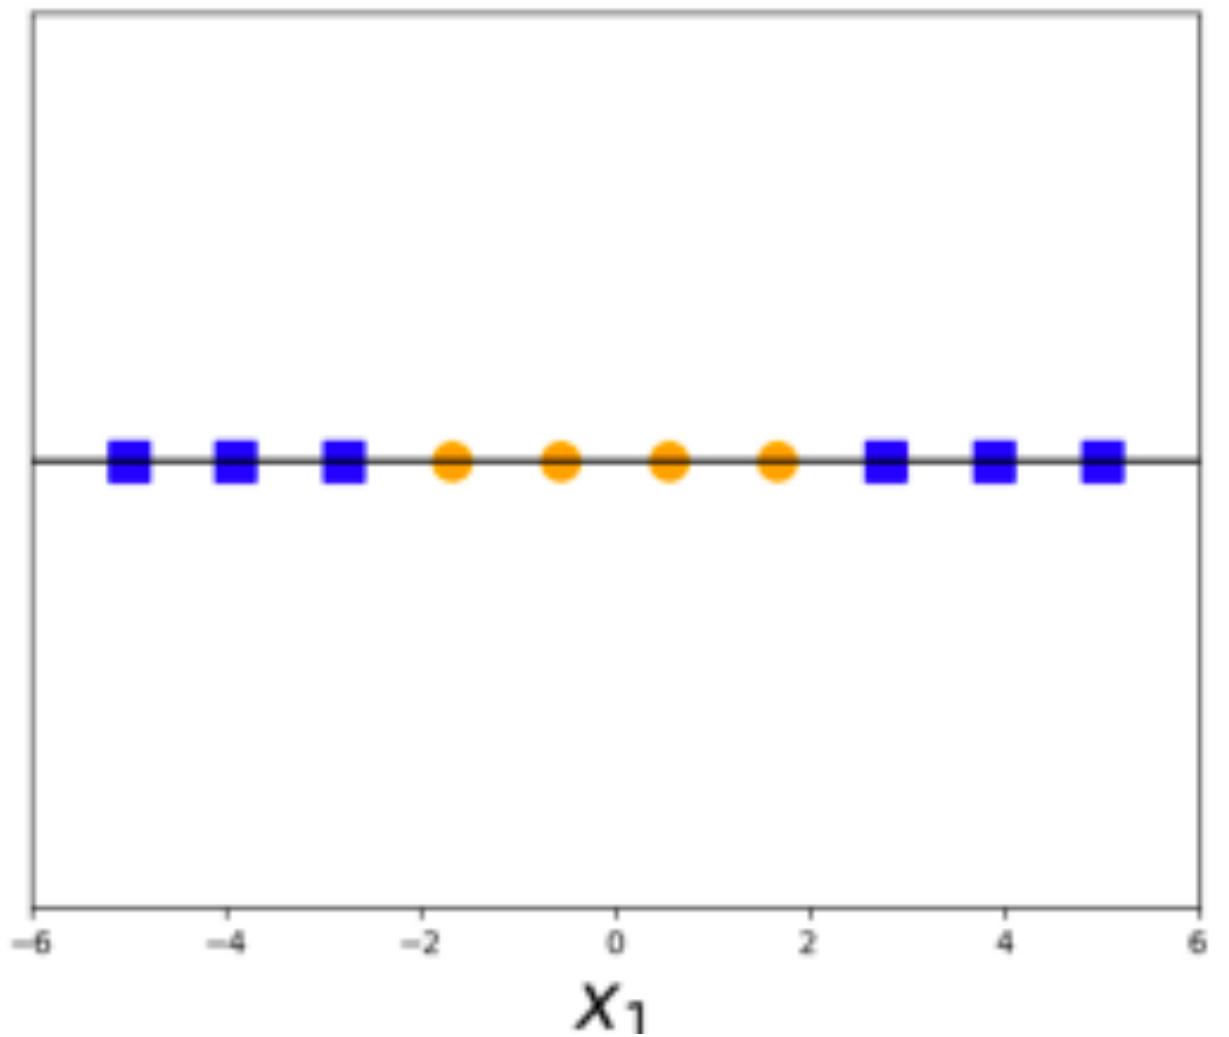


LOYOLA
UNIVERSITY CHICAGO

Support Vector Machines (SVM)

**How to Separate
these Features?**

Kernel-trick
examples.

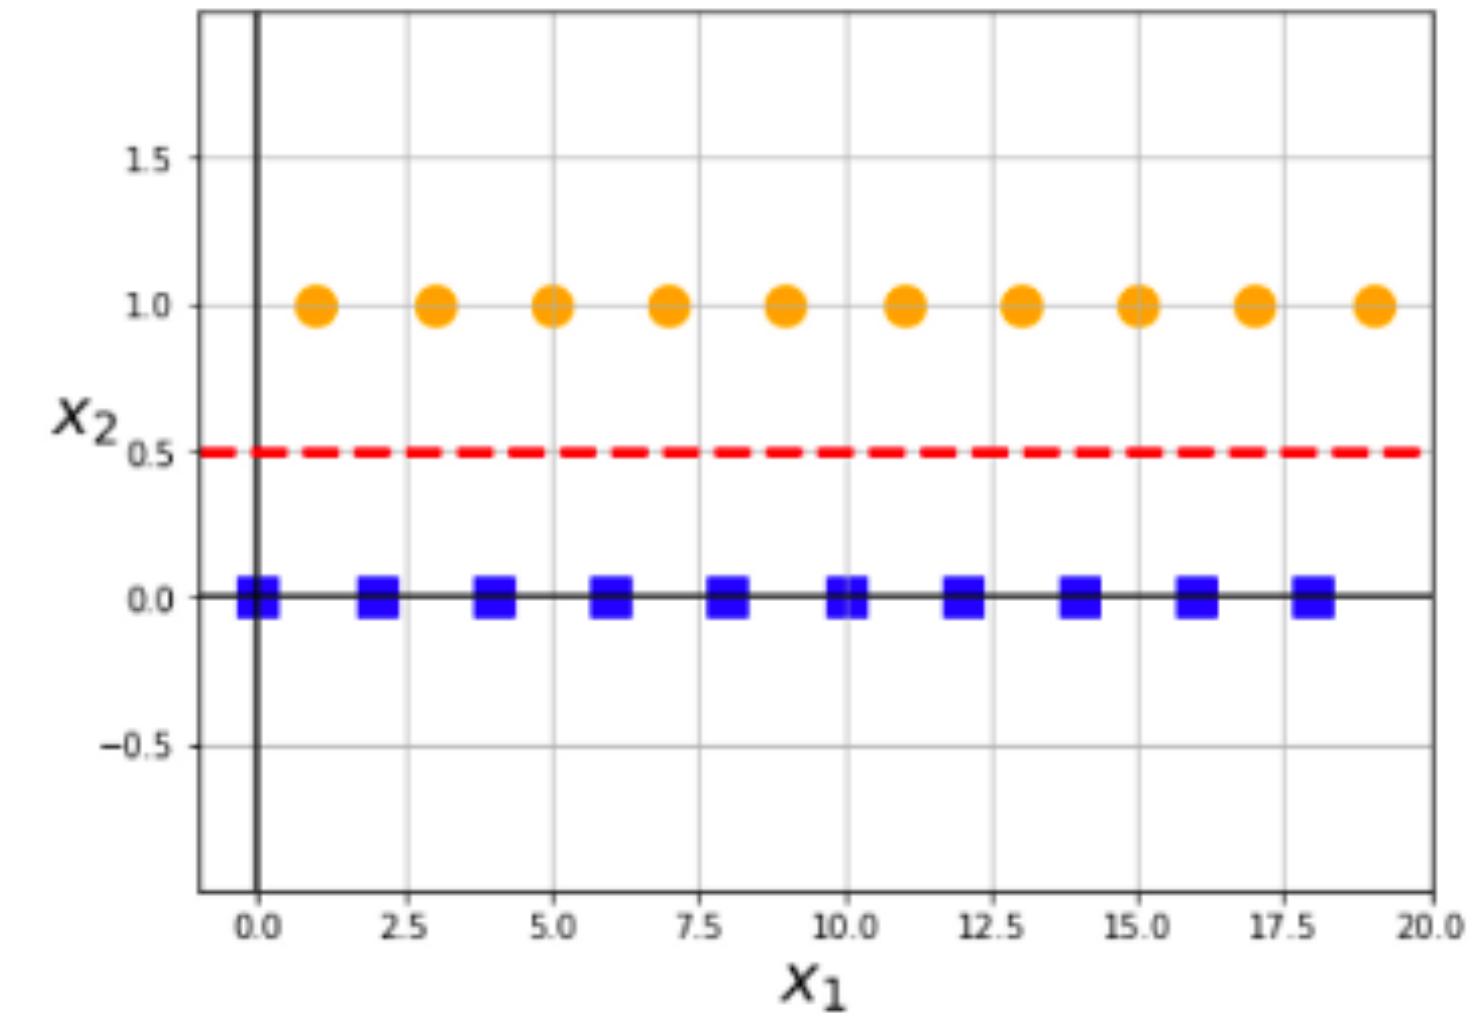
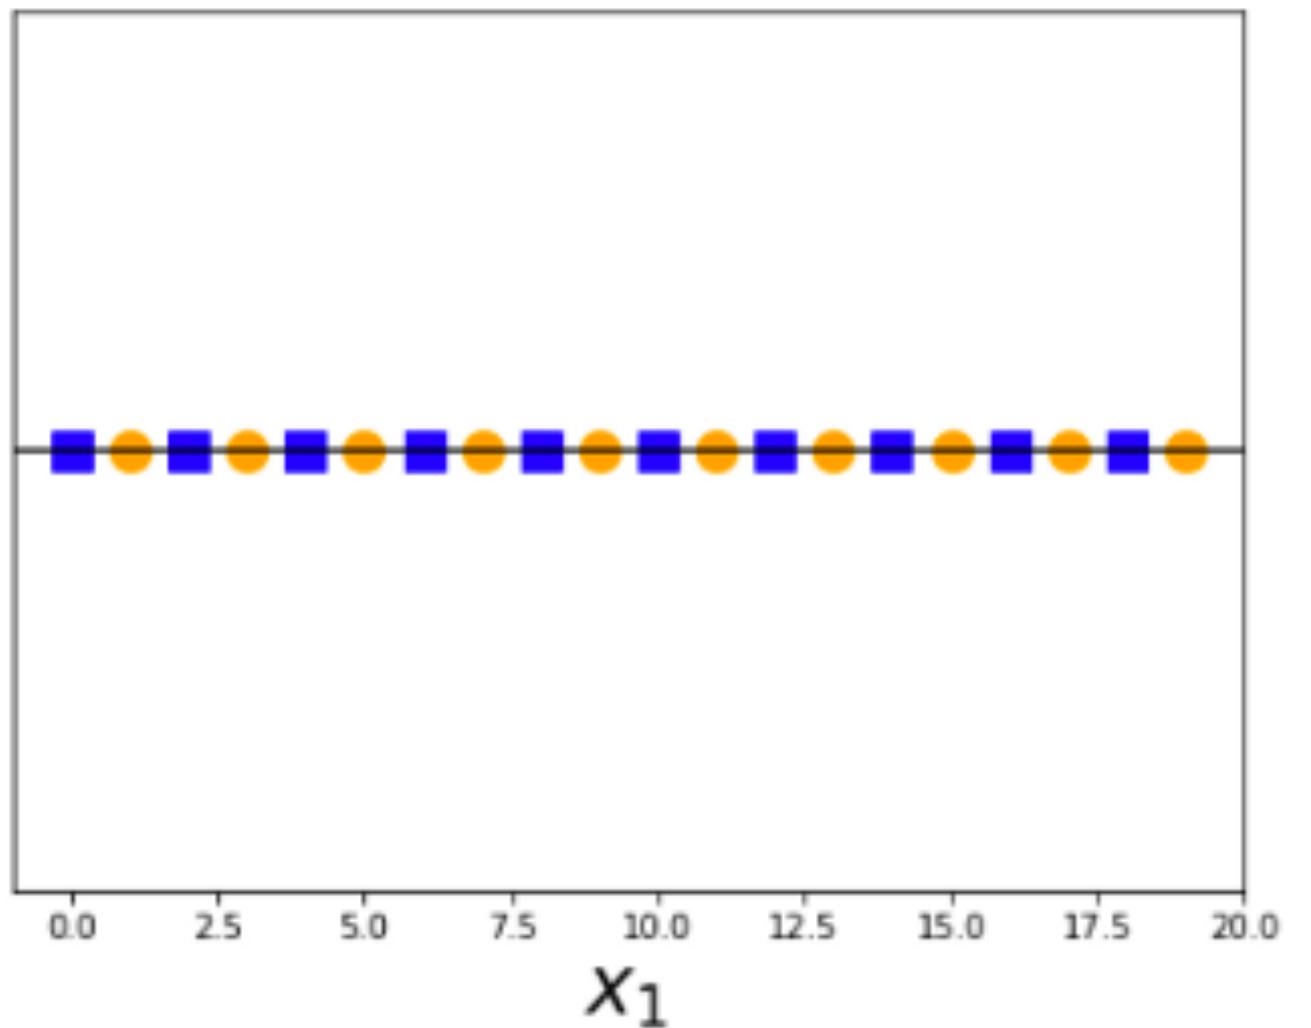


[https://towardsdatascience.com/
the-kernel-trick-c98cdbcaeb3f](https://towardsdatascience.com/the-kernel-trick-c98cdbcaeb3f)

Support Vector Machines (SVM)

**How to Separate
these Features?**

Kernel-trick
examples.

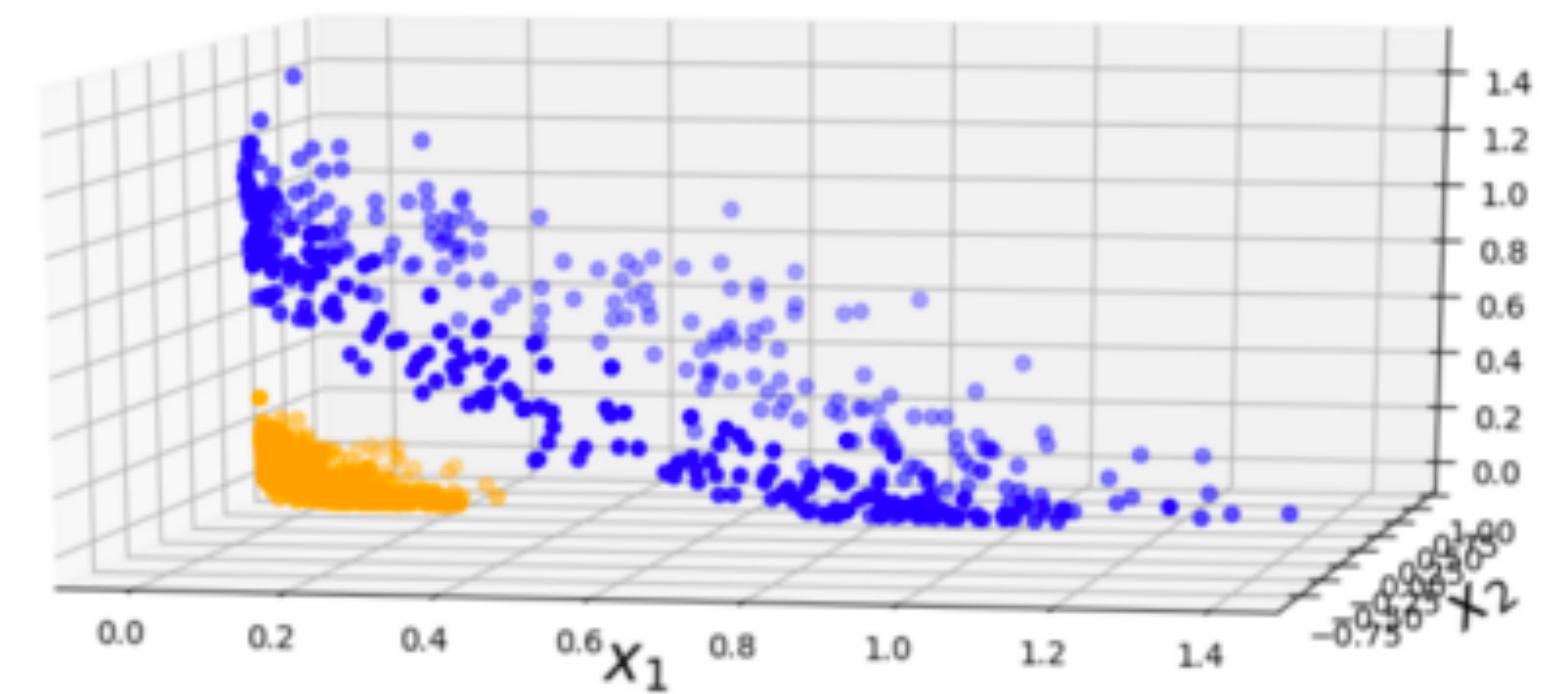
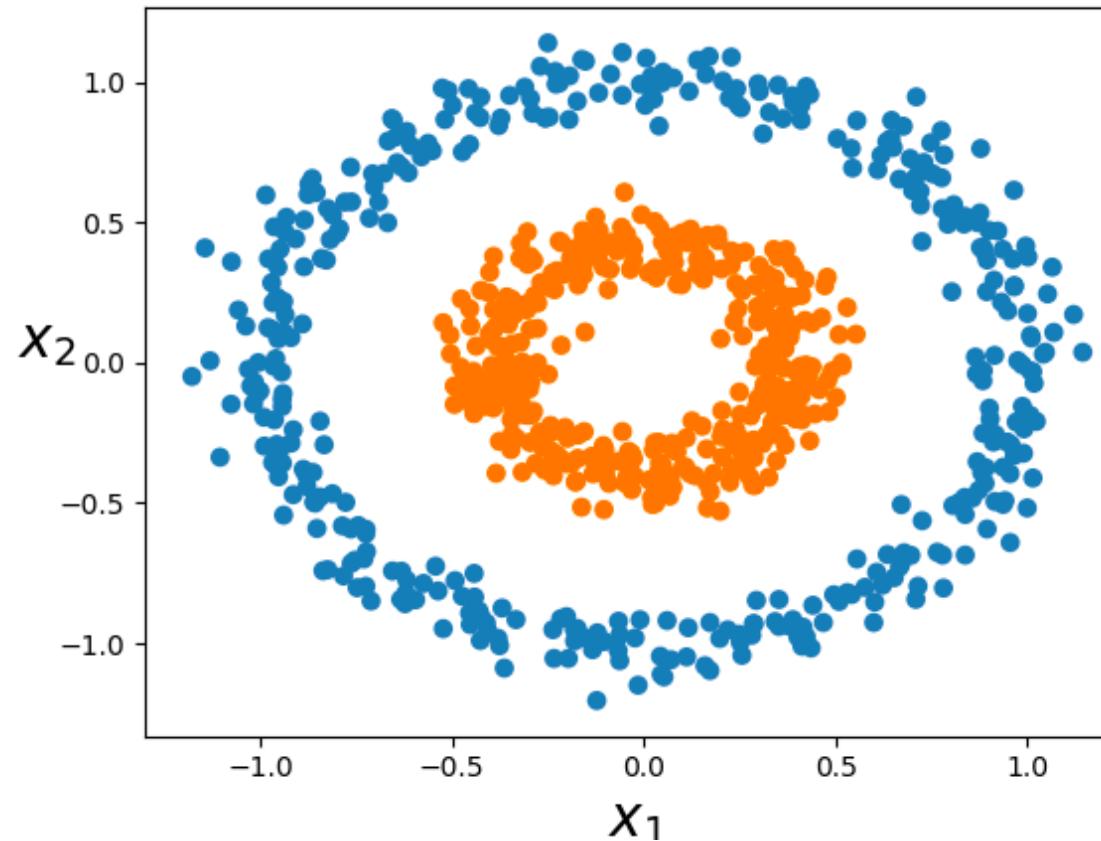


[https://towardsdatascience.com/
the-kernel-trick-c98cdcaeb3f](https://towardsdatascience.com/the-kernel-trick-c98cdcaeb3f)

Support Vector Machines (SVM)

**How to Separate
these Features?**

Kernel-trick
examples.



Implementation
available at
<https://scikit-learn.org/stable/modules/svm.html>

[https://towardsdatascience.com/
the-kernel-trick-c98ccdbcaeb3f](https://towardsdatascience.com/the-kernel-trick-c98ccdbcaeb3f)

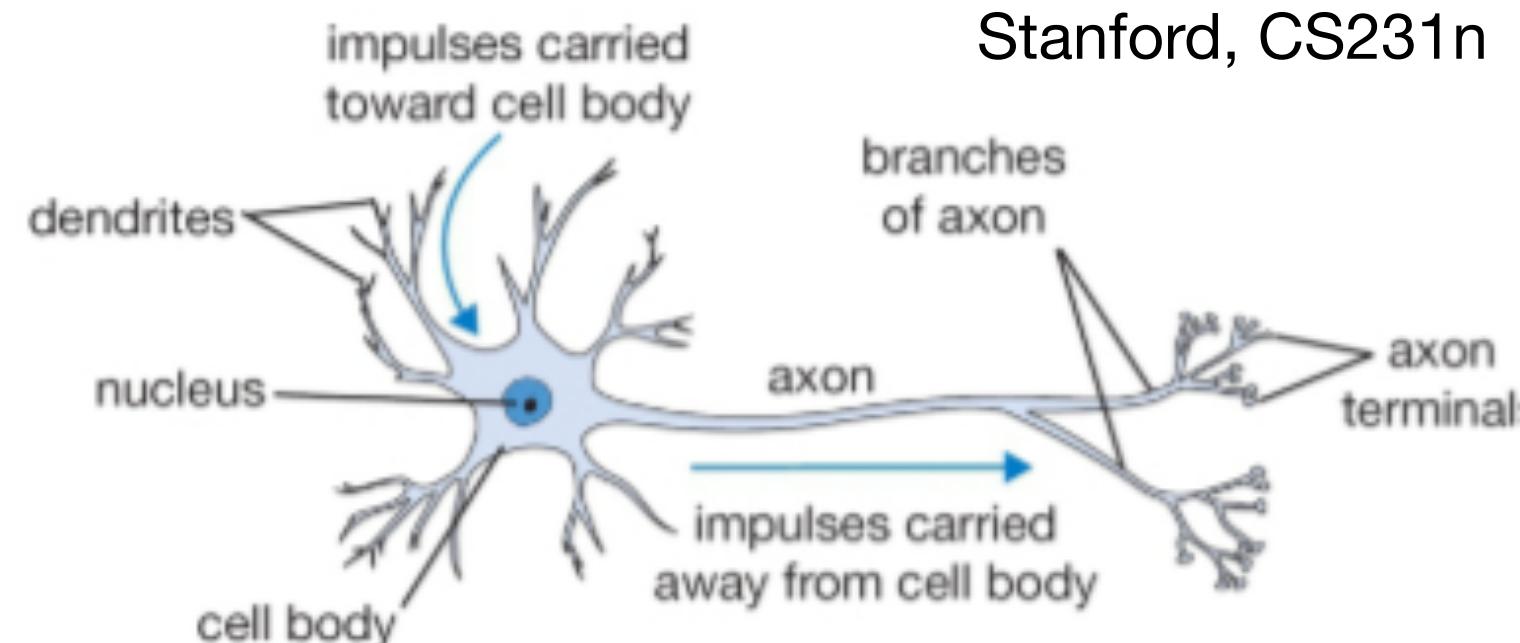


LOYOLA
UNIVERSITY CHICAGO

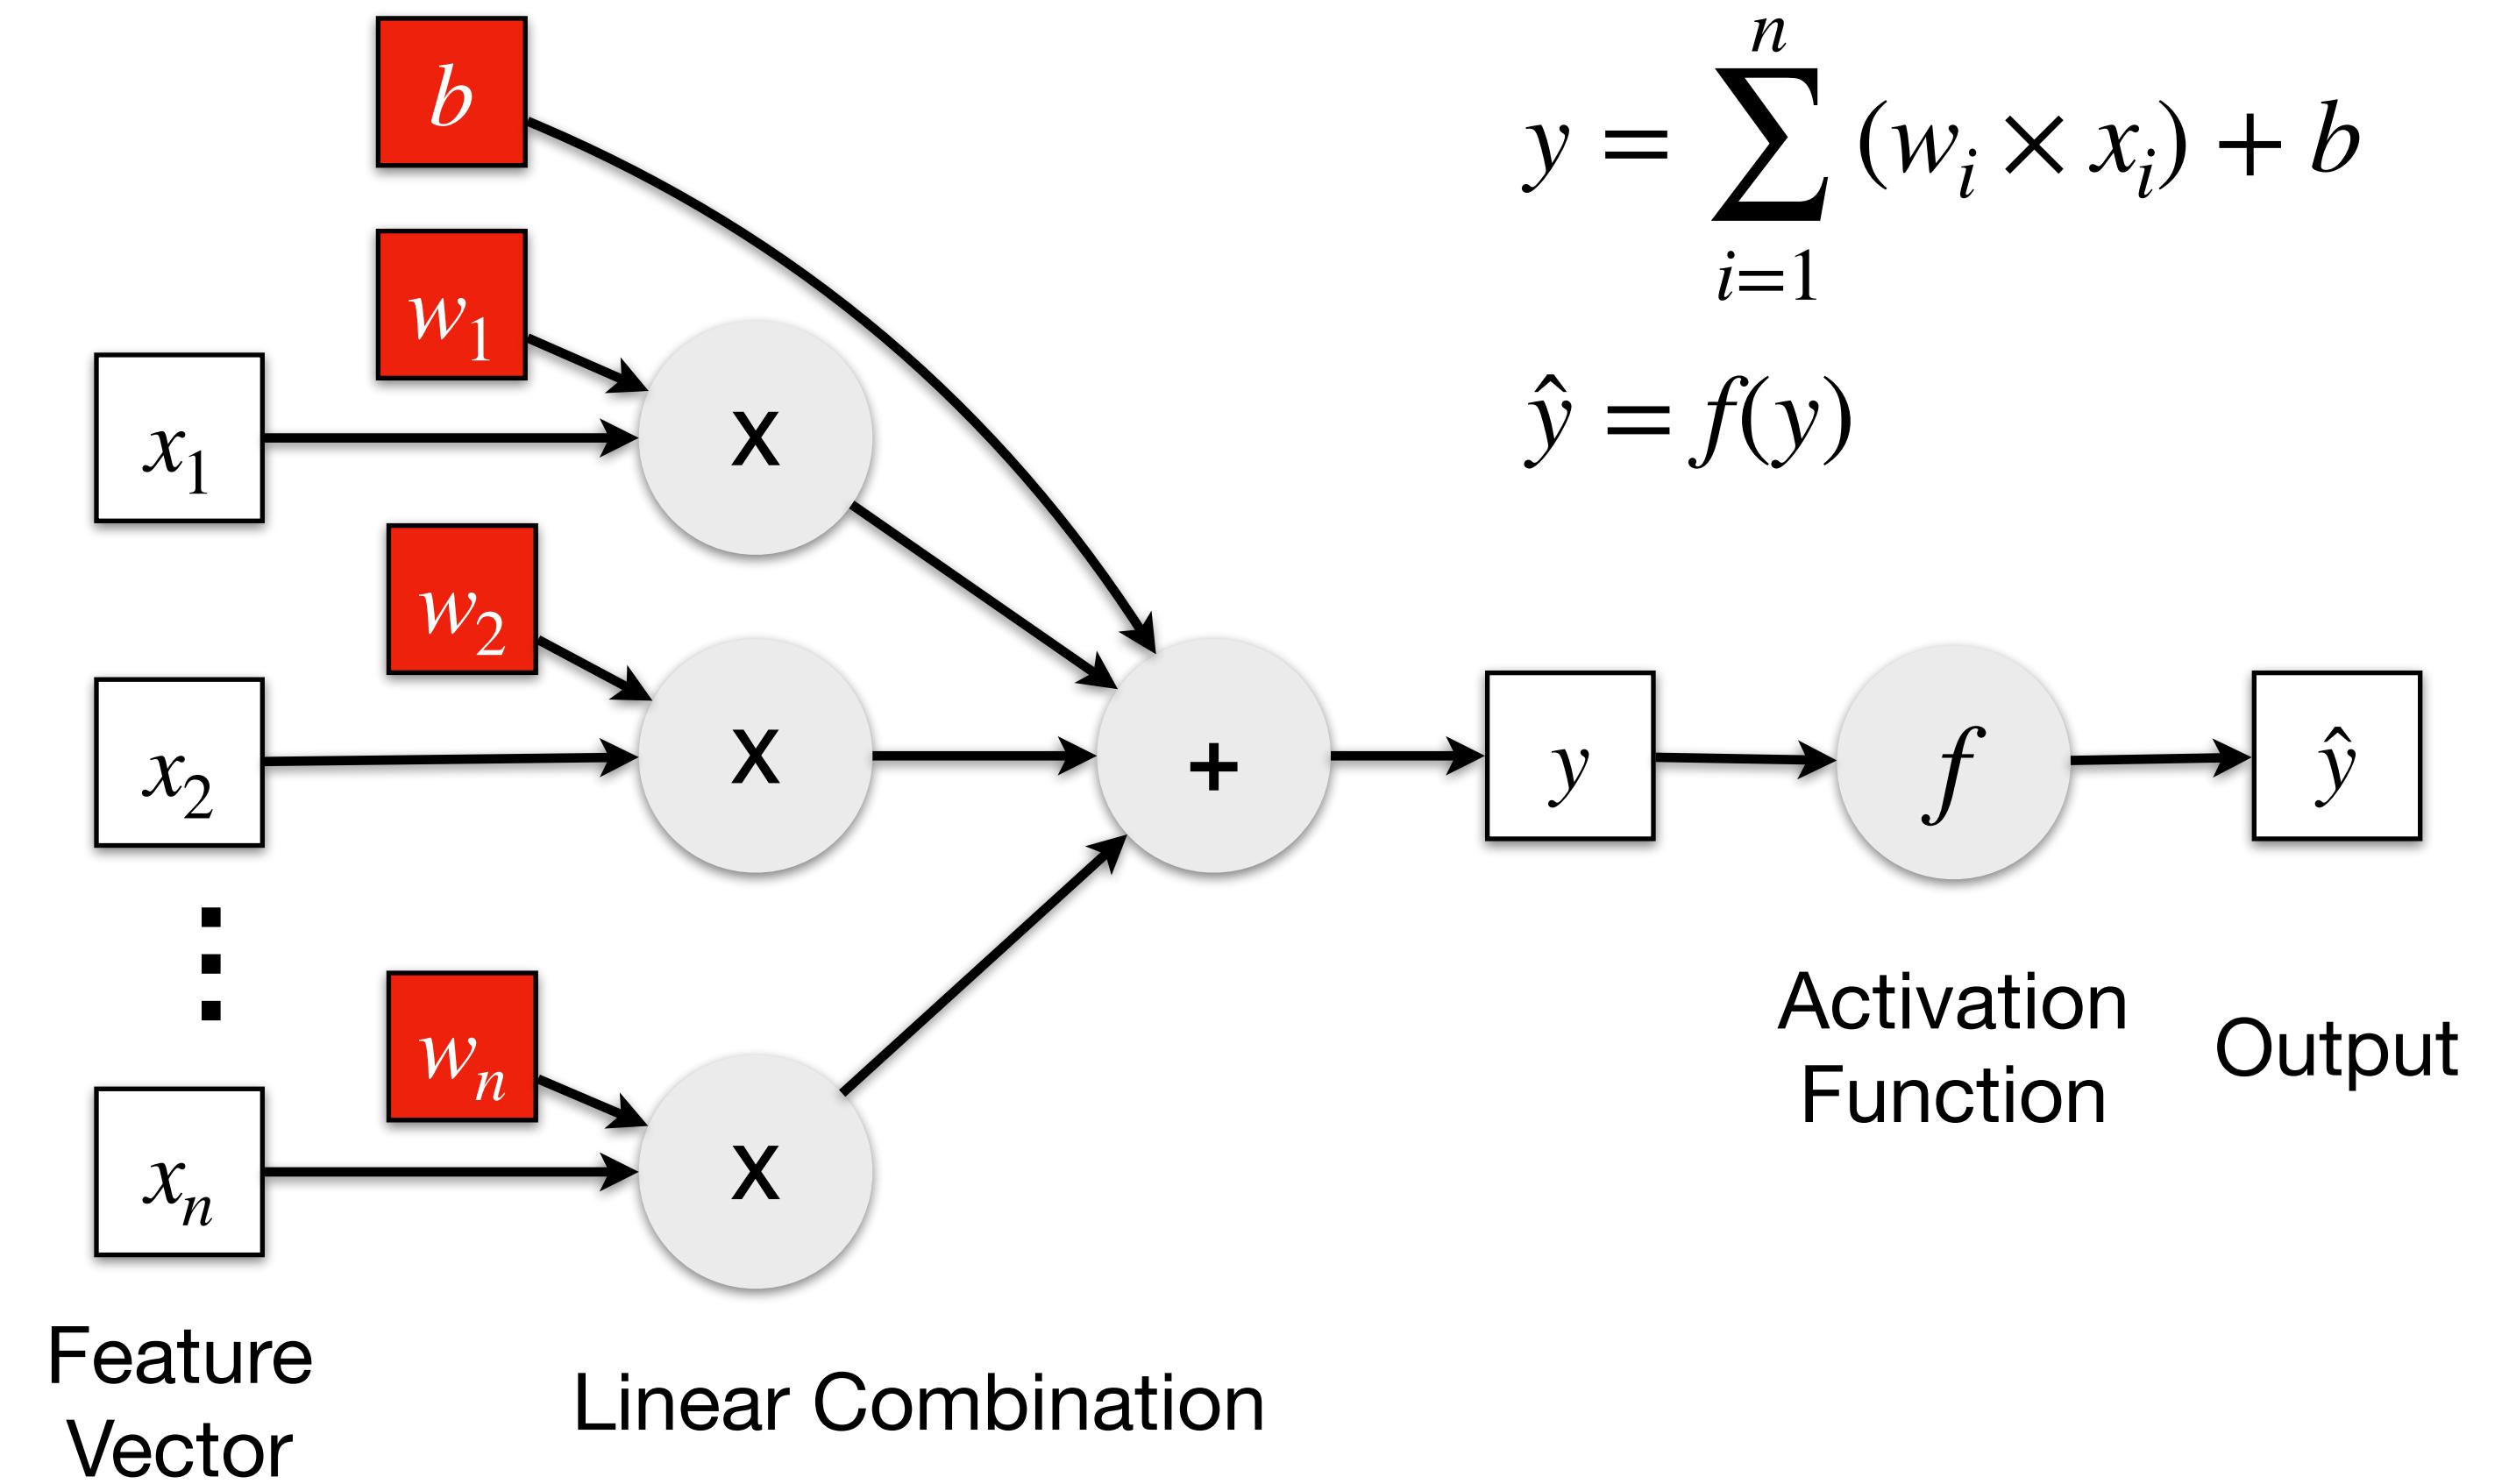
Neural Networks (NN)

Artificial Neuron

Building block of NNs.



Bioinspiration

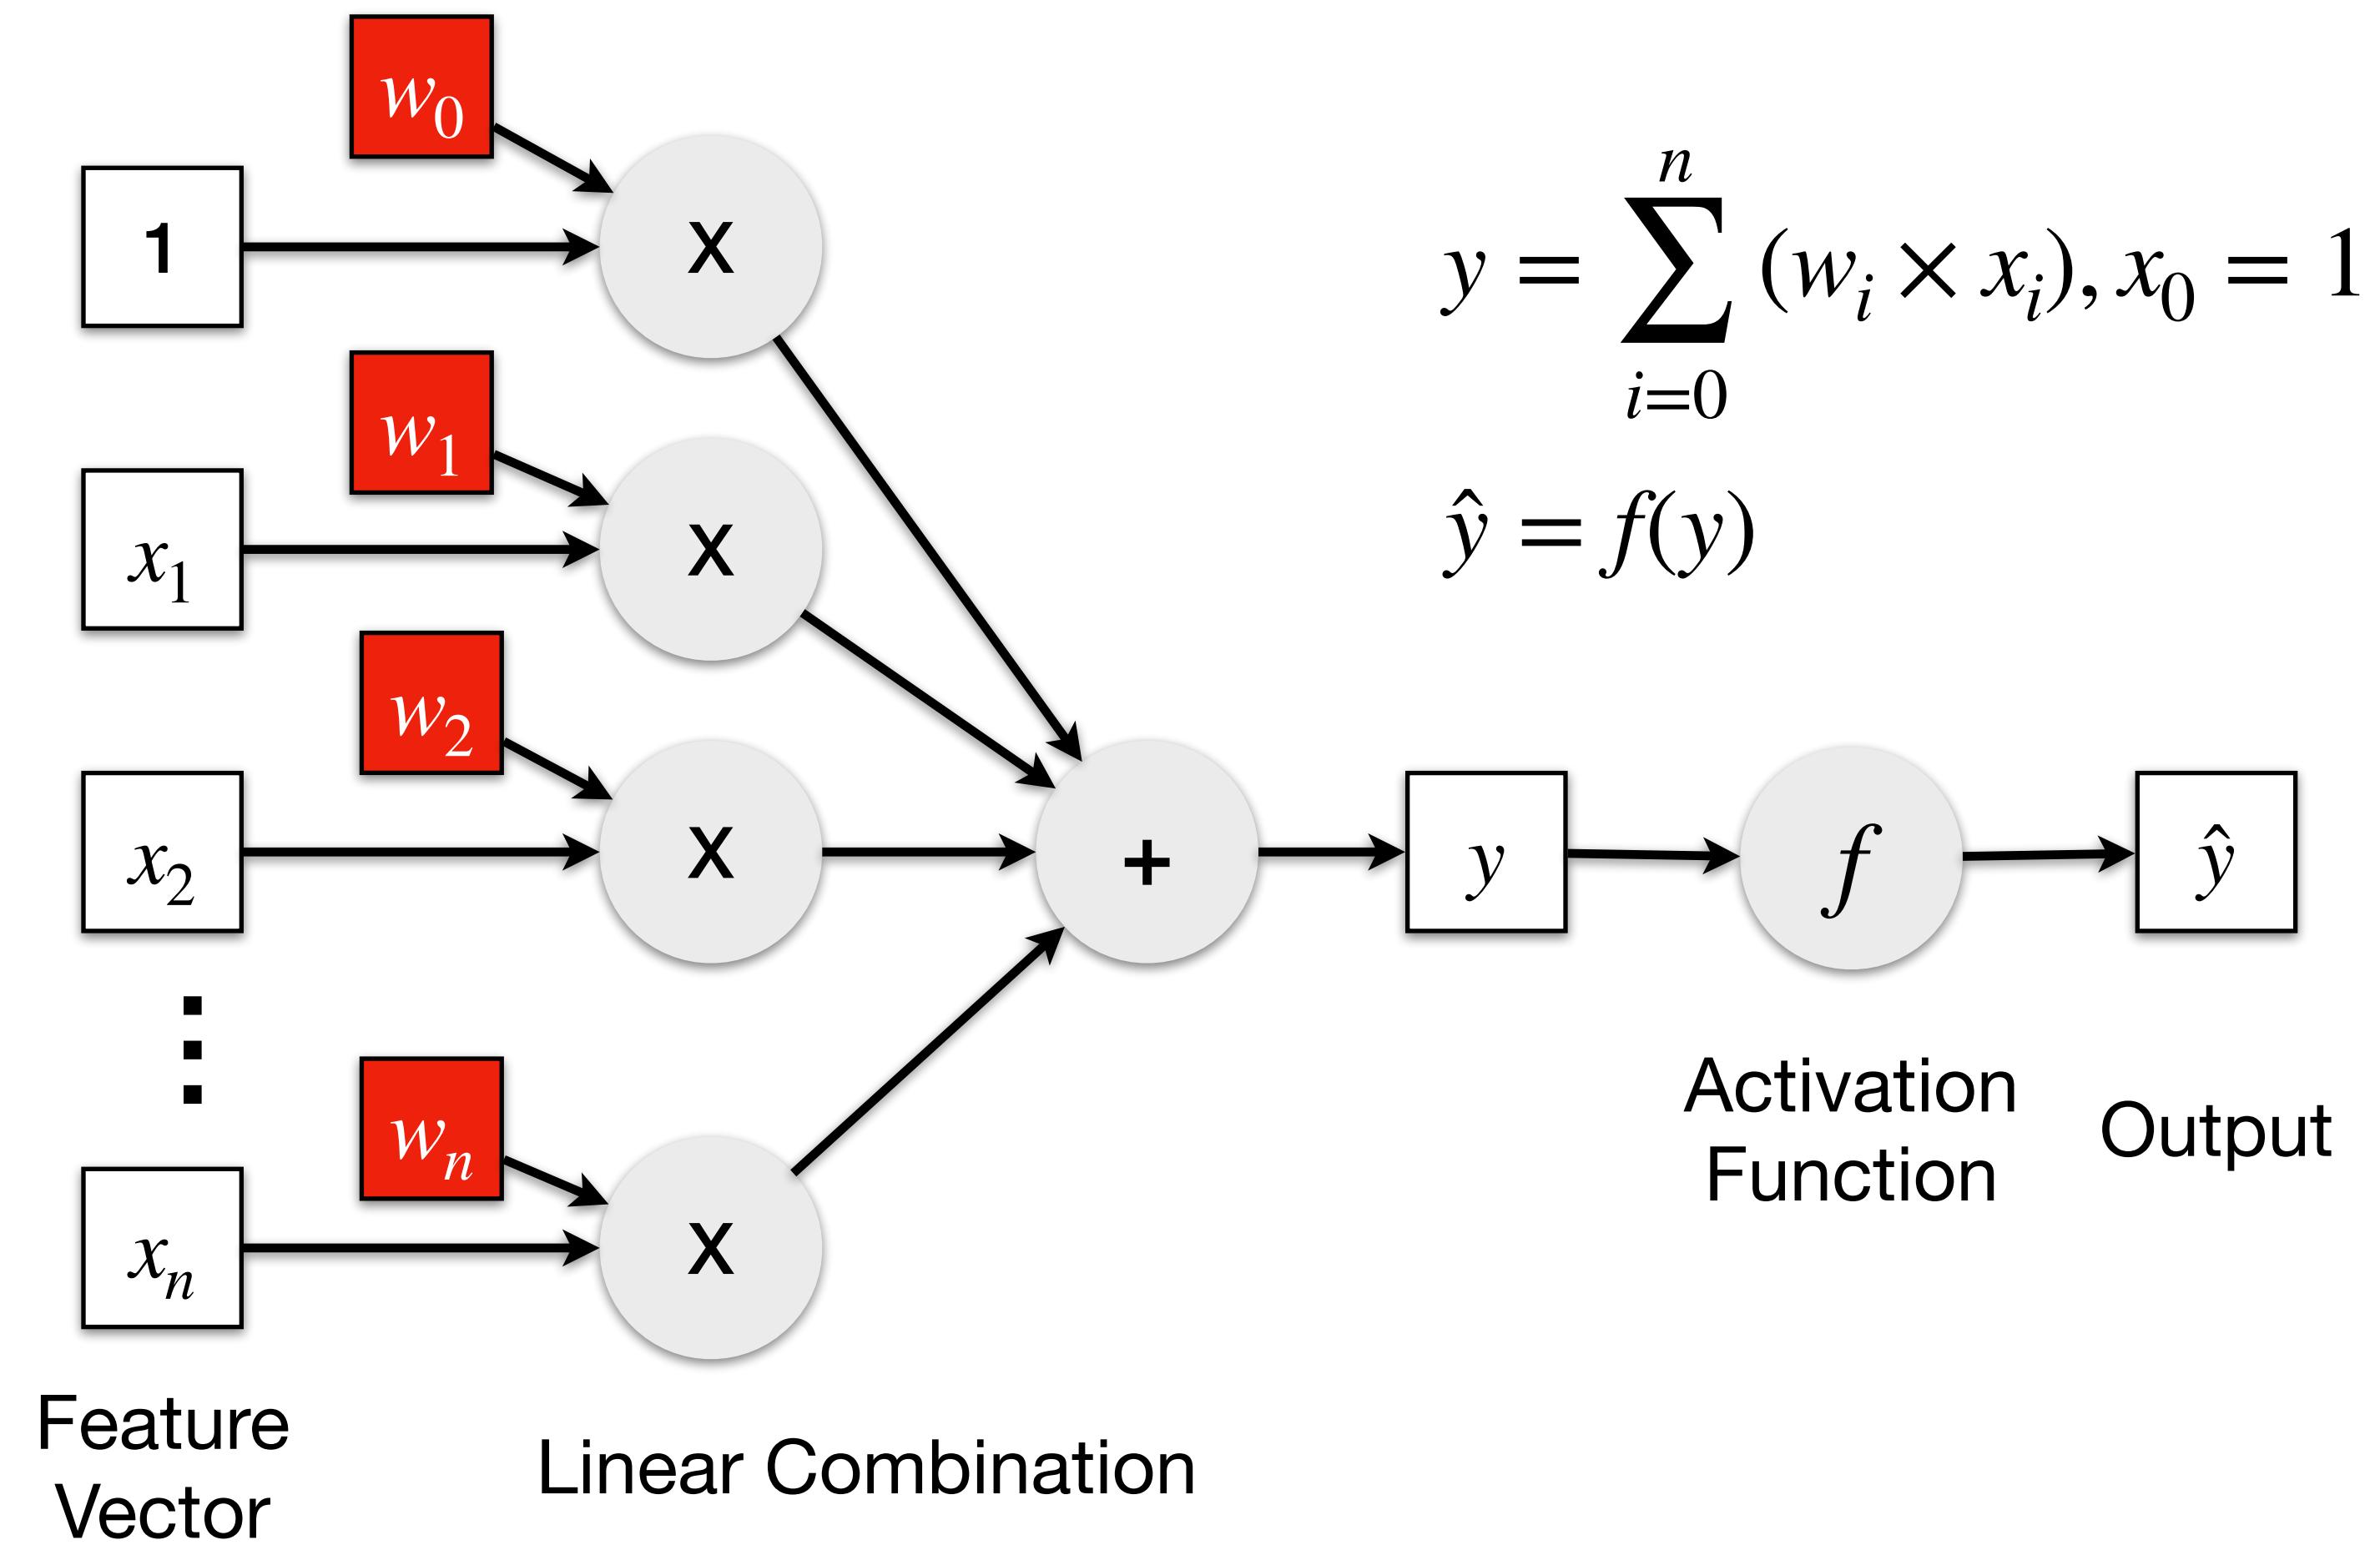


Neural Networks (NN)

Artificial Neuron

Building block of NNs.

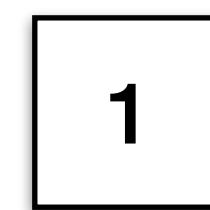
Notation
adjustments.



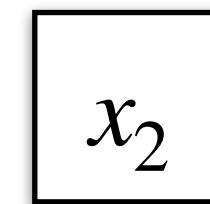
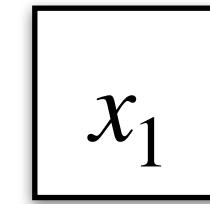
Neural Networks (NN)

Artificial Neuron

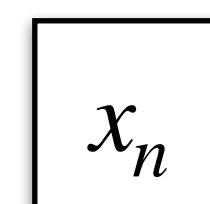
Building block of NNs.



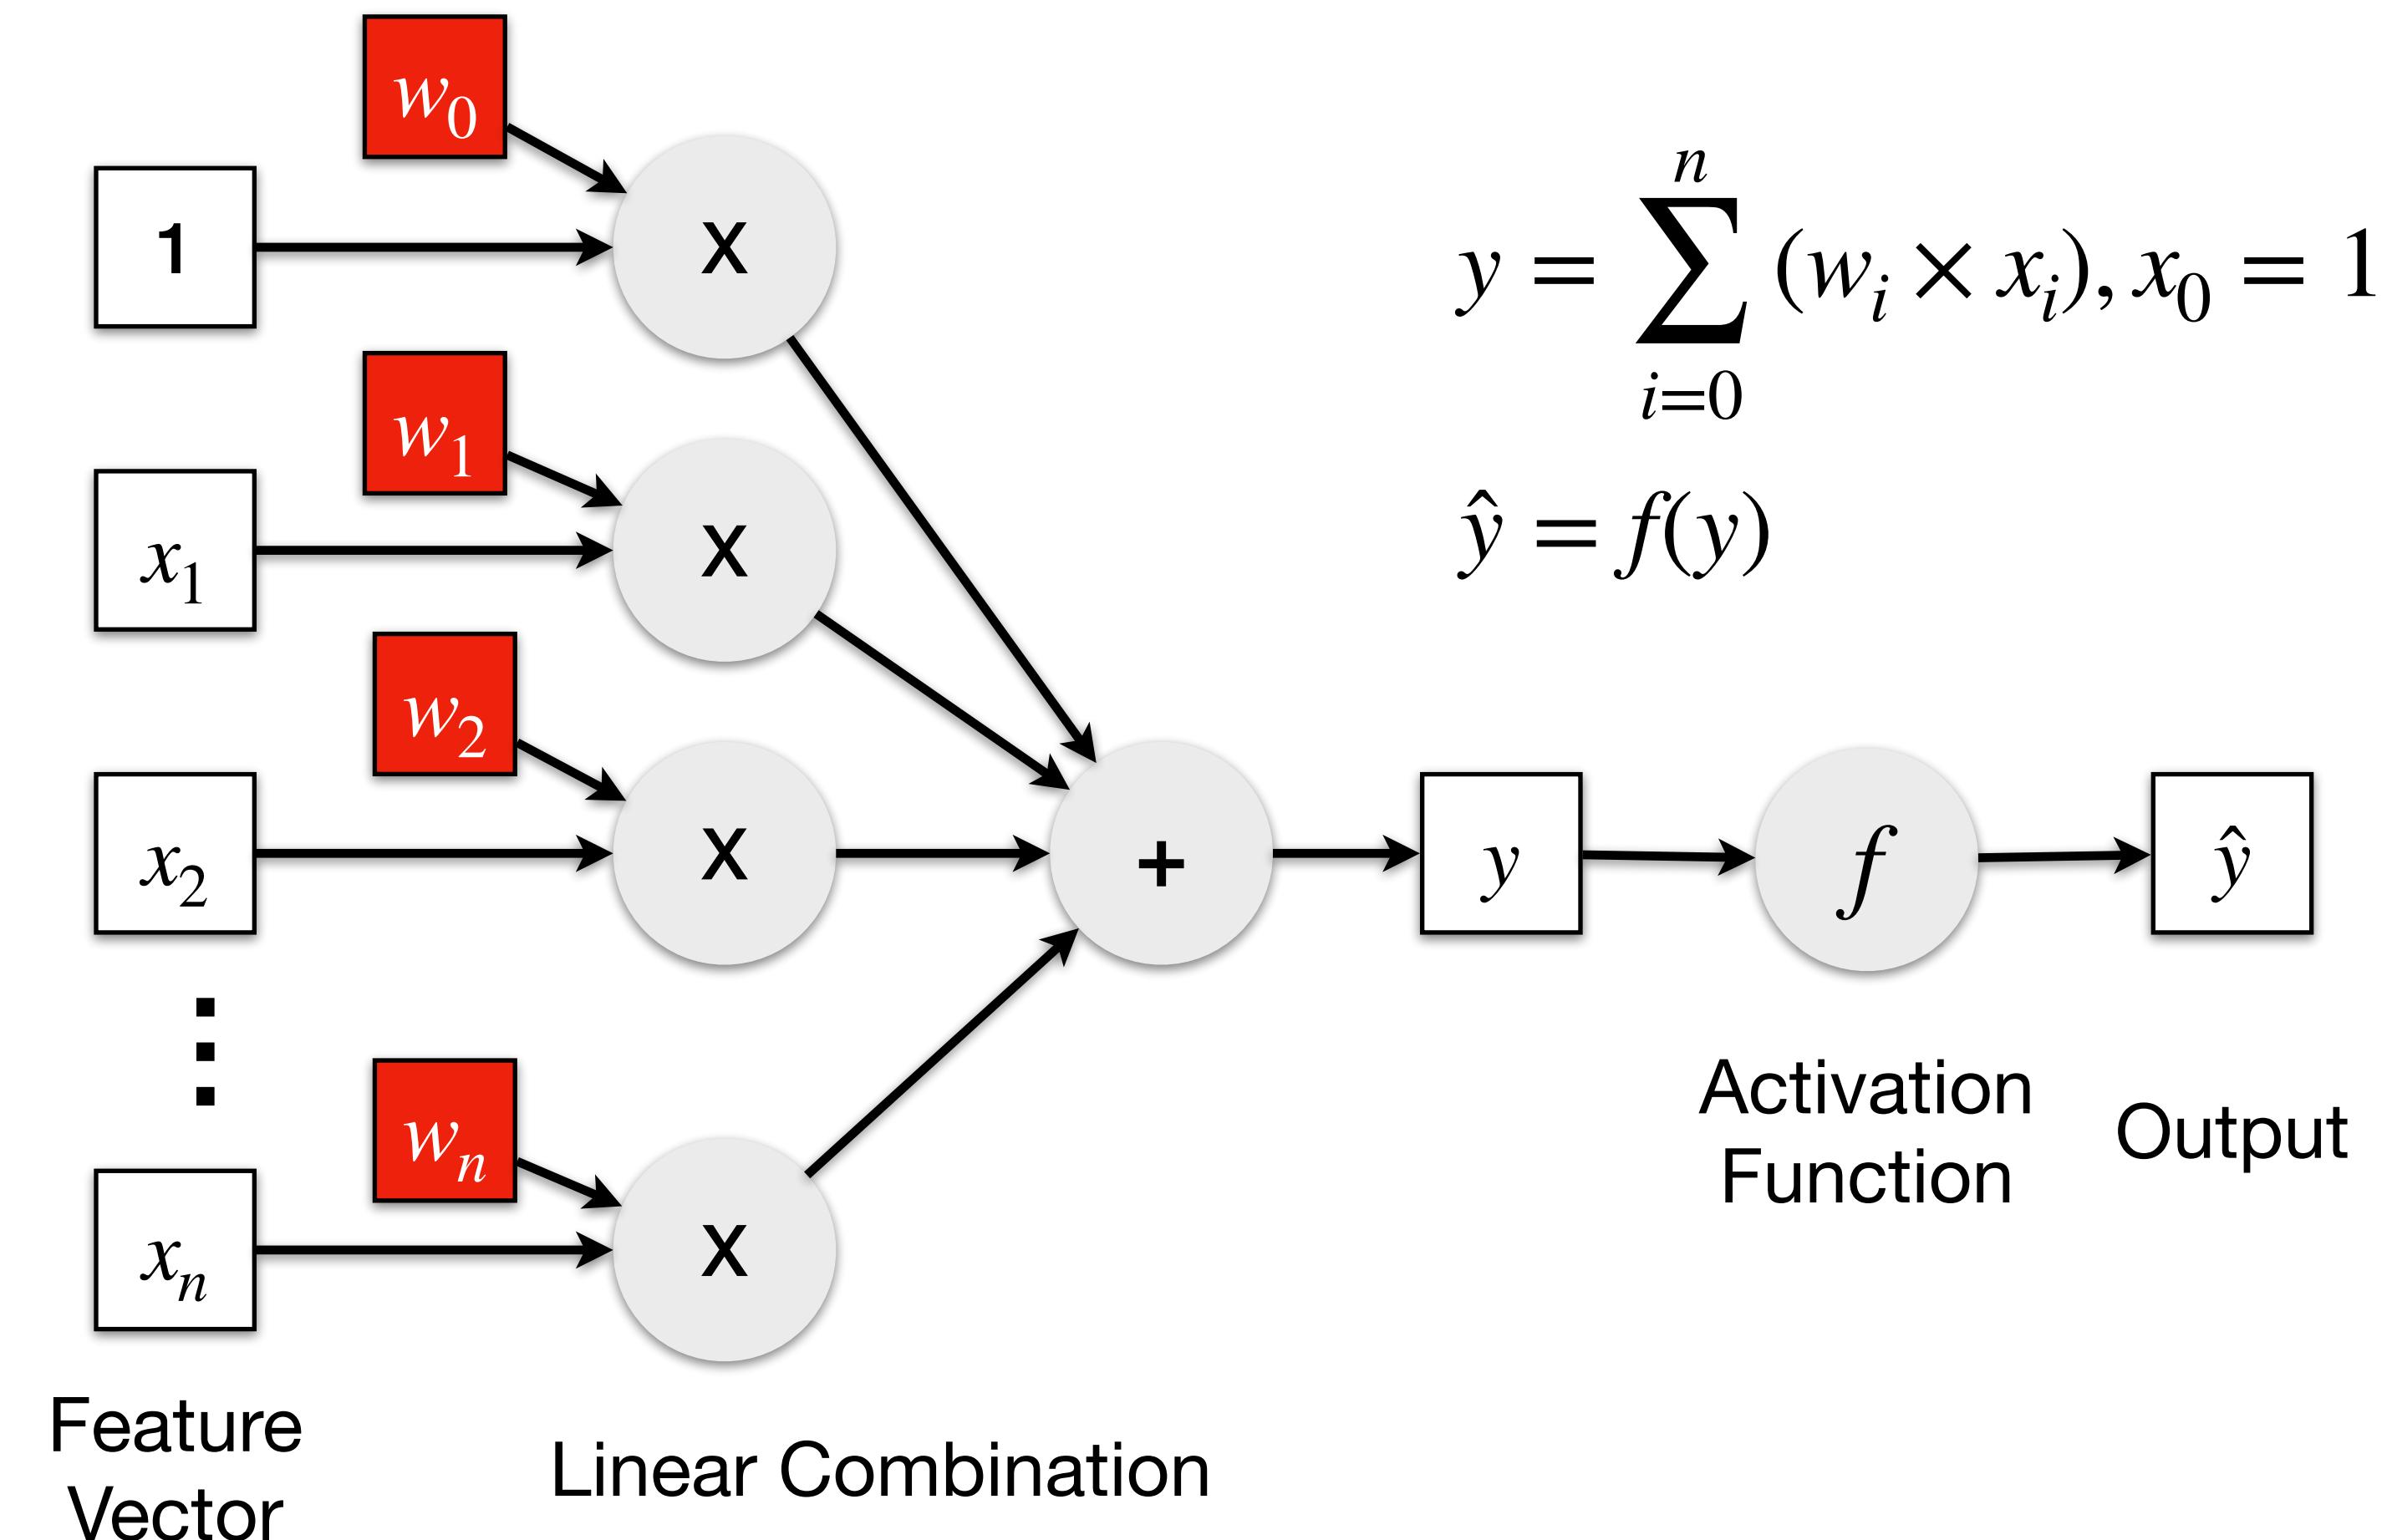
Data sample \vec{x}
with expected
(known) label z .



⋮



Feature
Vector



Neural Networks (NN)

Artificial Neuron

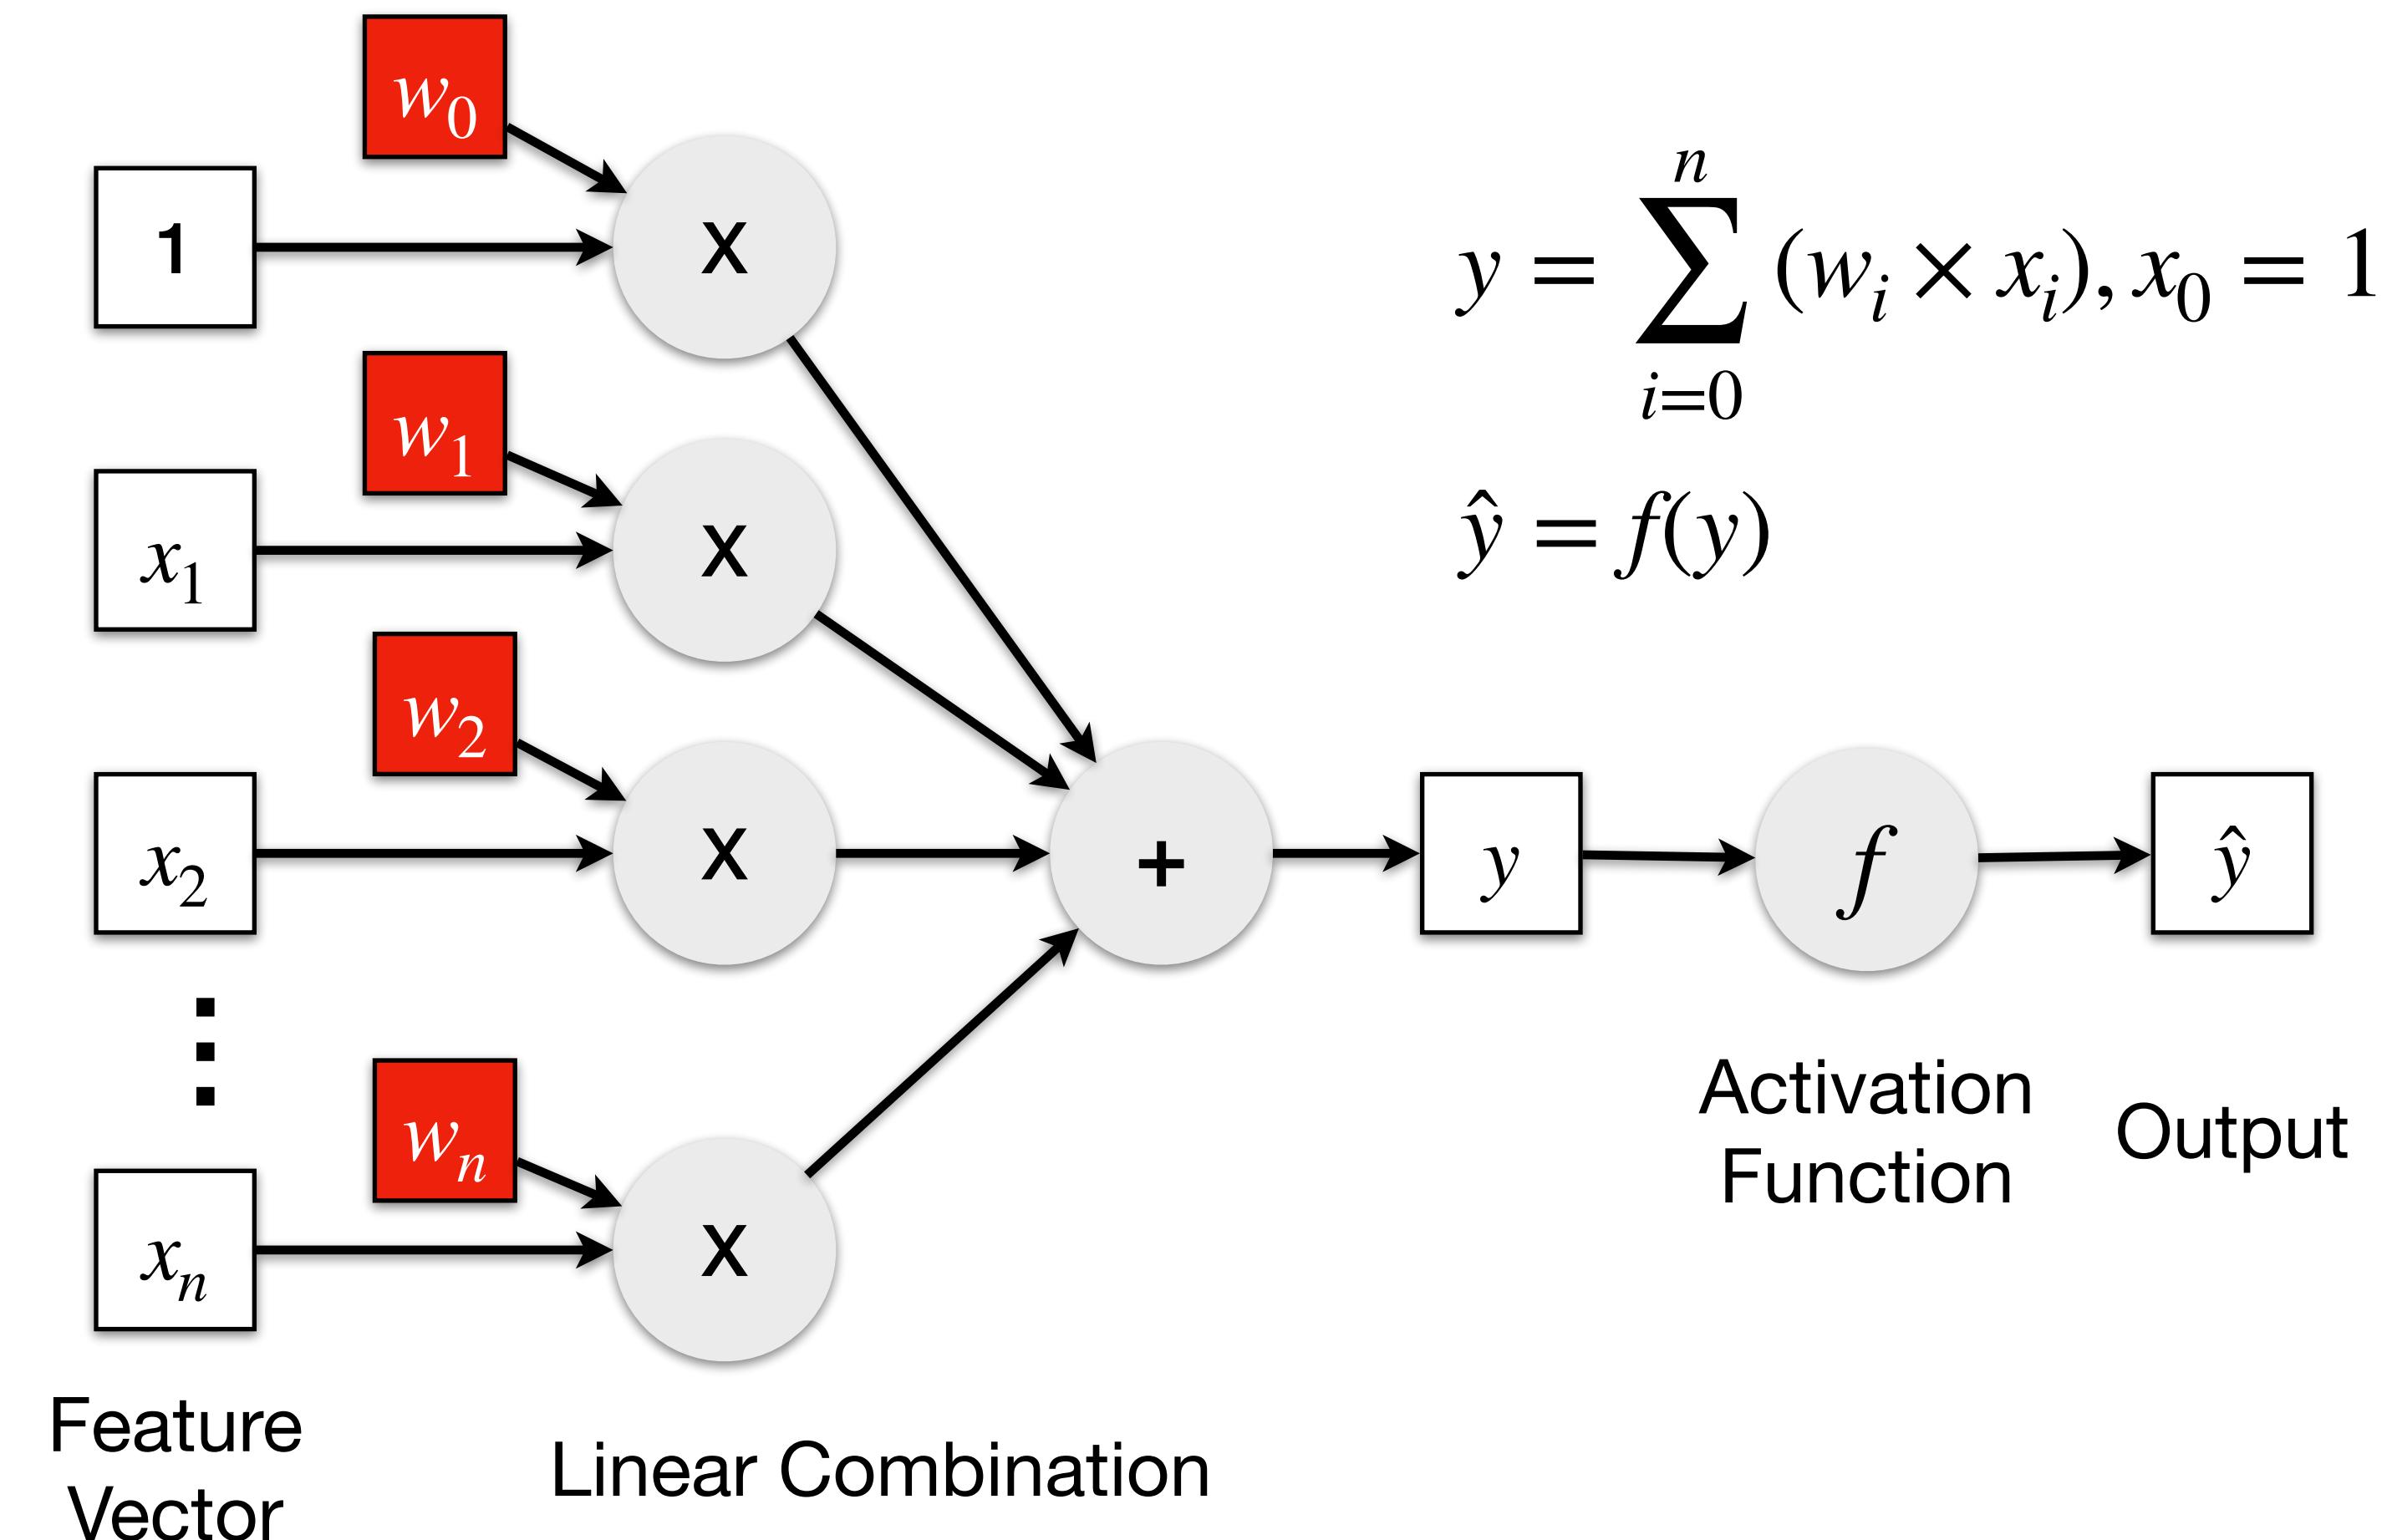
Building block of NNs.

w_0
 w_1
 w_2
⋮
 w_n

Learnable weights \vec{w}
(a.k.a. neuron's parameters).

They may start with random values.

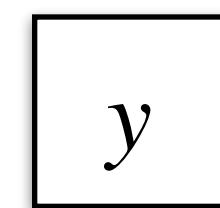
They may start with weights previously learned on other data (transfer learning).

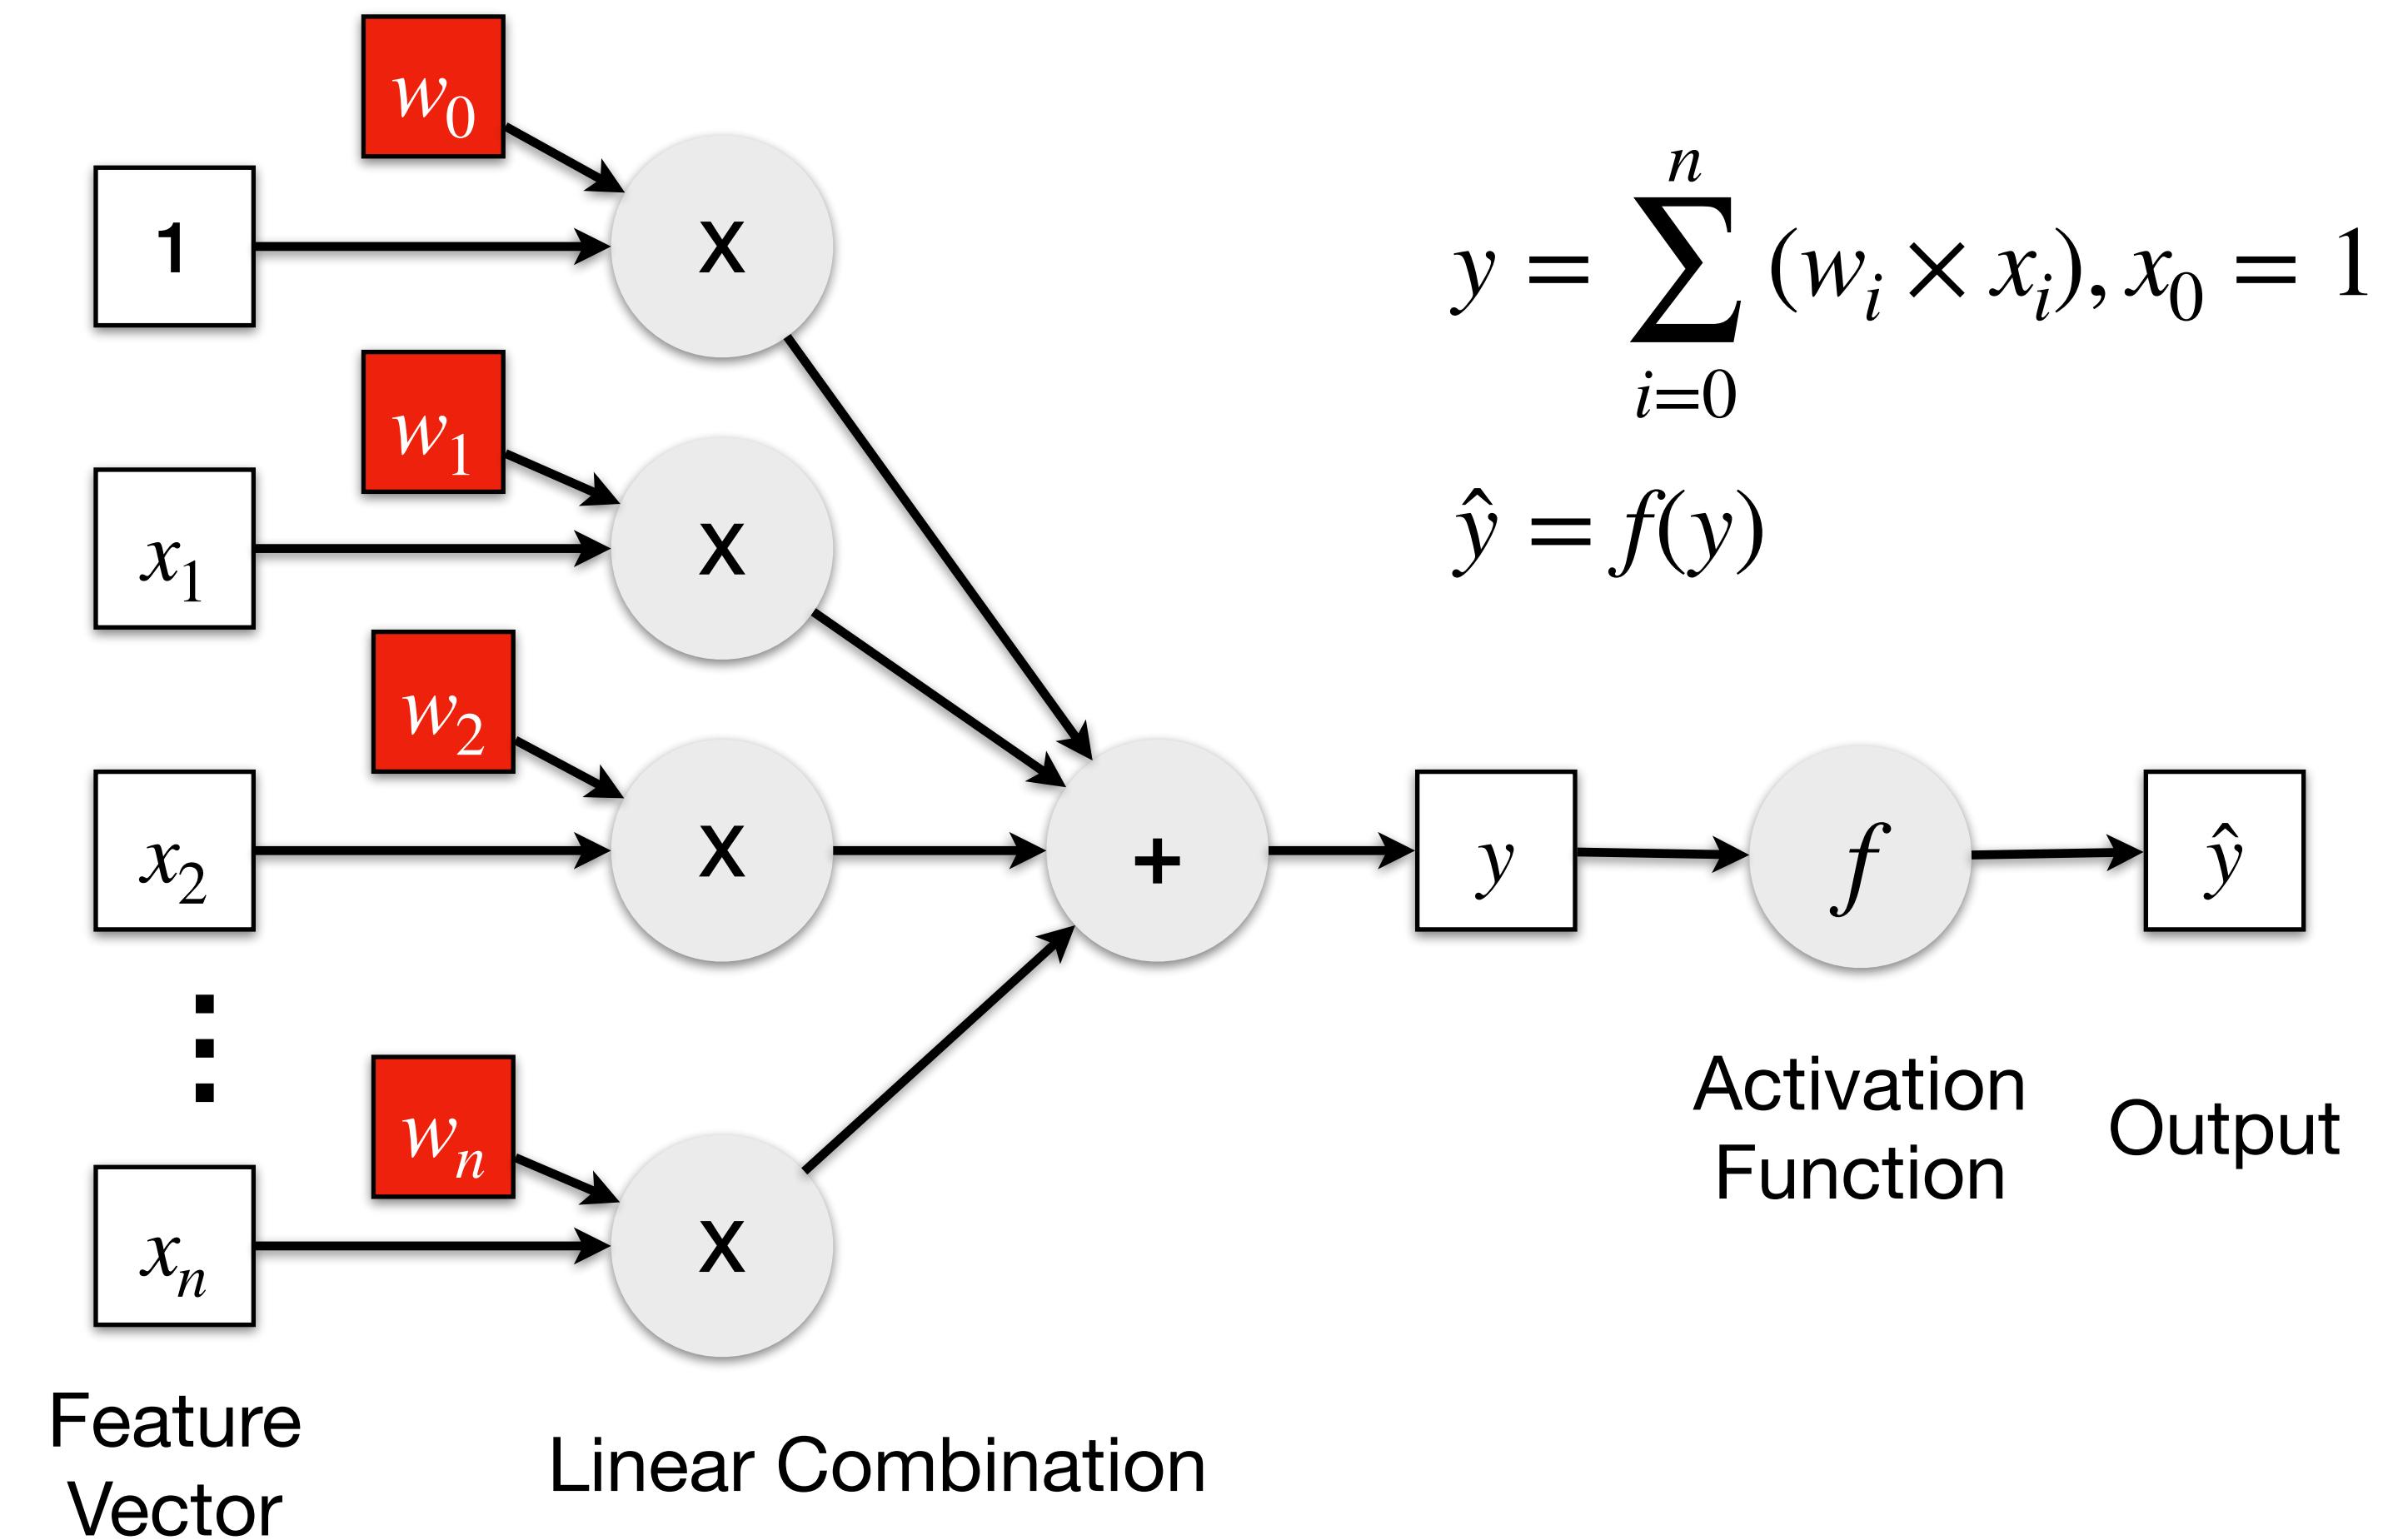


Neural Networks (NN)

Artificial Neuron

Building block of NNs.

 Linear combination
of feature vector's
components and neuron's
weights.



Neural Networks (NN)

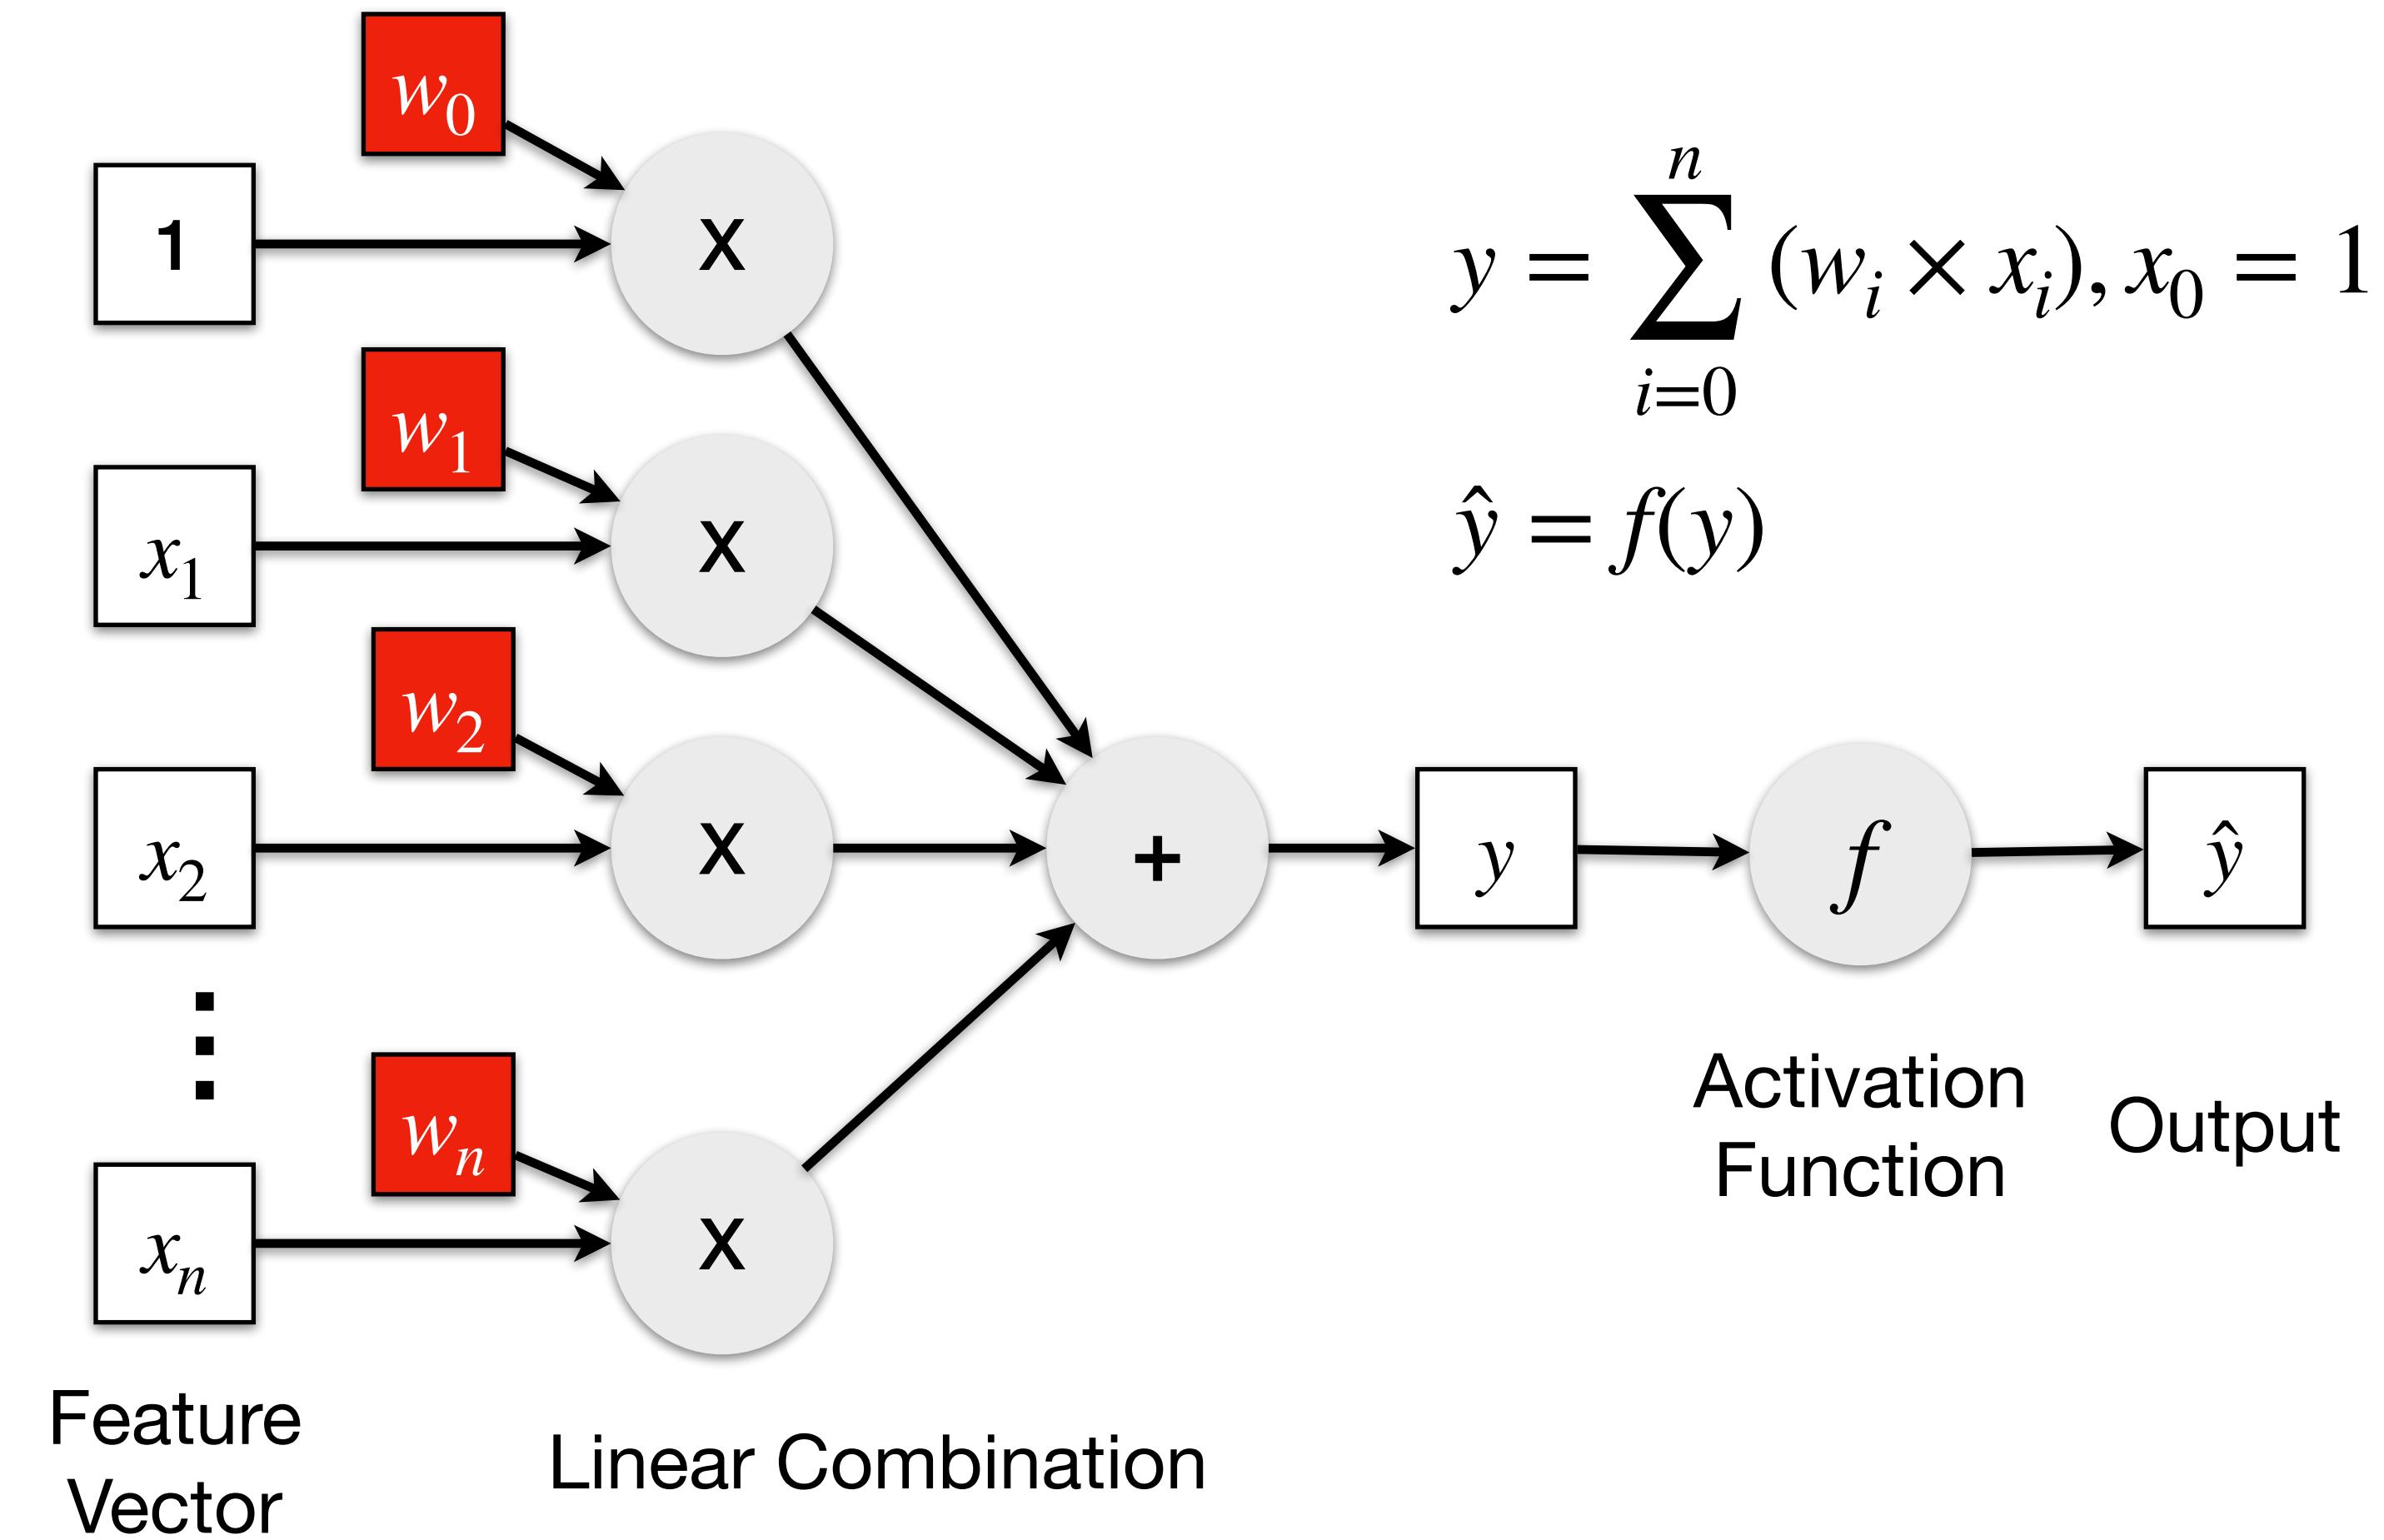
Artificial Neuron

Building block of NNs.

f

Activation Function

Ideally a non-linear
differentiable function
to add non-linearity
to the model.



Neural Networks (NN)

Artificial Neuron

Building block of NNs.

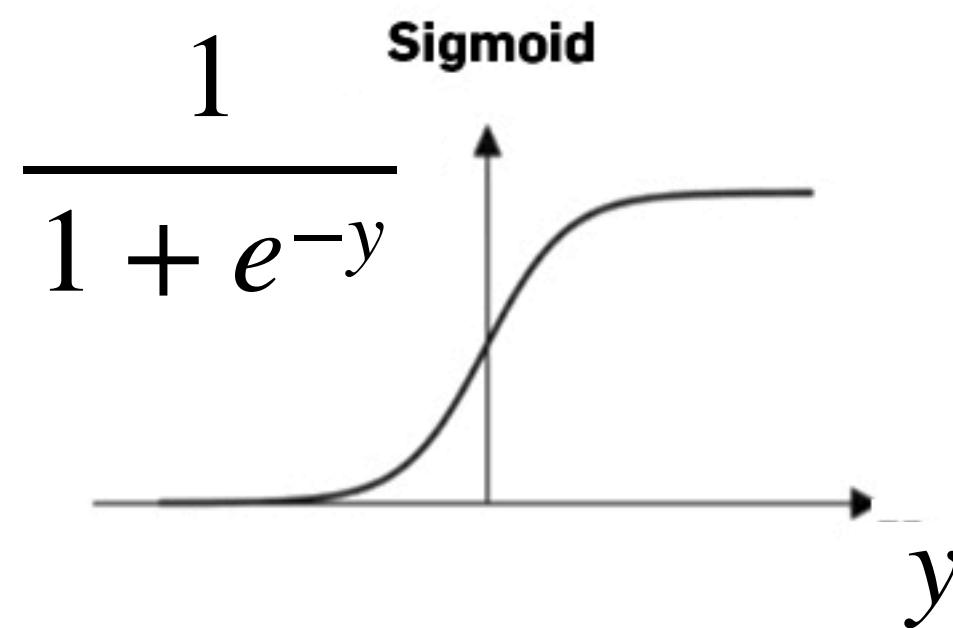
f

Activation Function

Ideally a non-linear differentiable function to add non-linearity to the model.

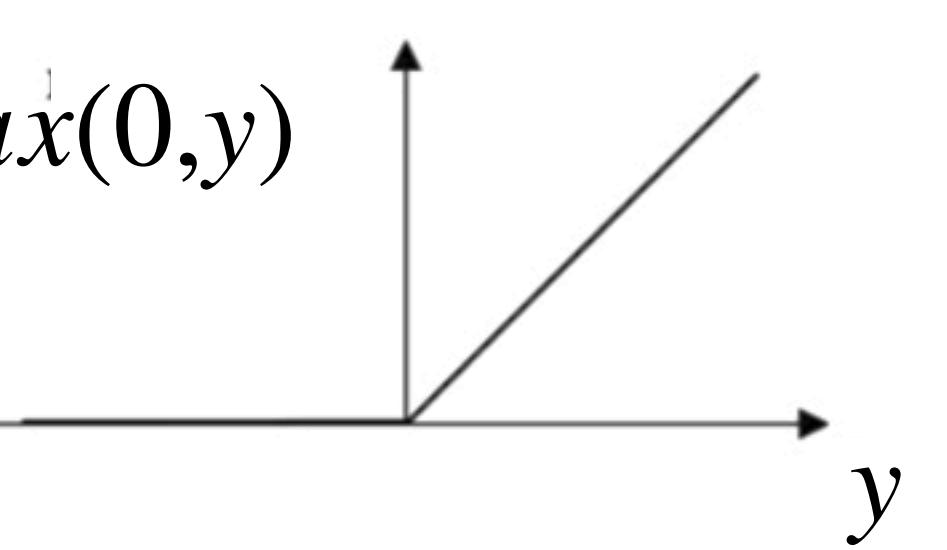
Necessary to allow NNs to learn non-linear functions.

Examples

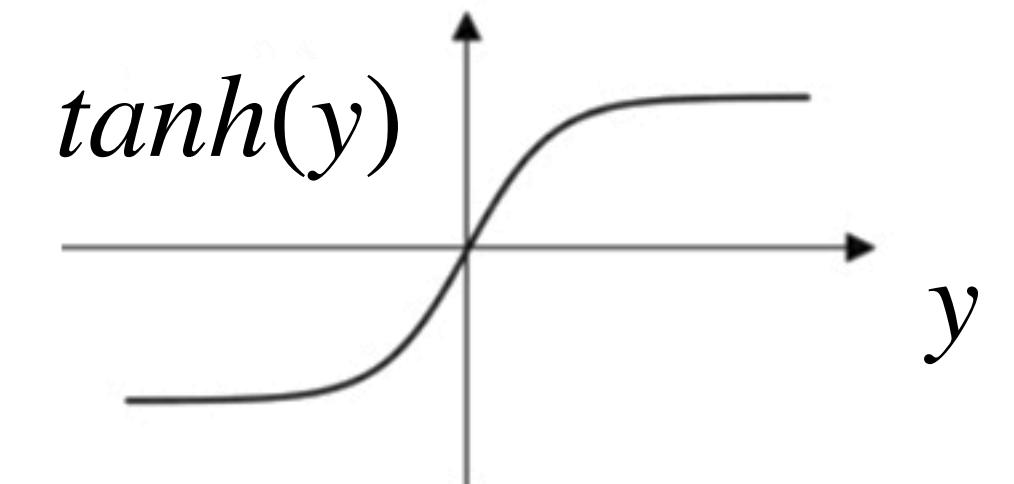


ReLU

$$\max(0, y)$$



Tanh



Neural Networks (NN)

Artificial Neuron

Building block of NNs.

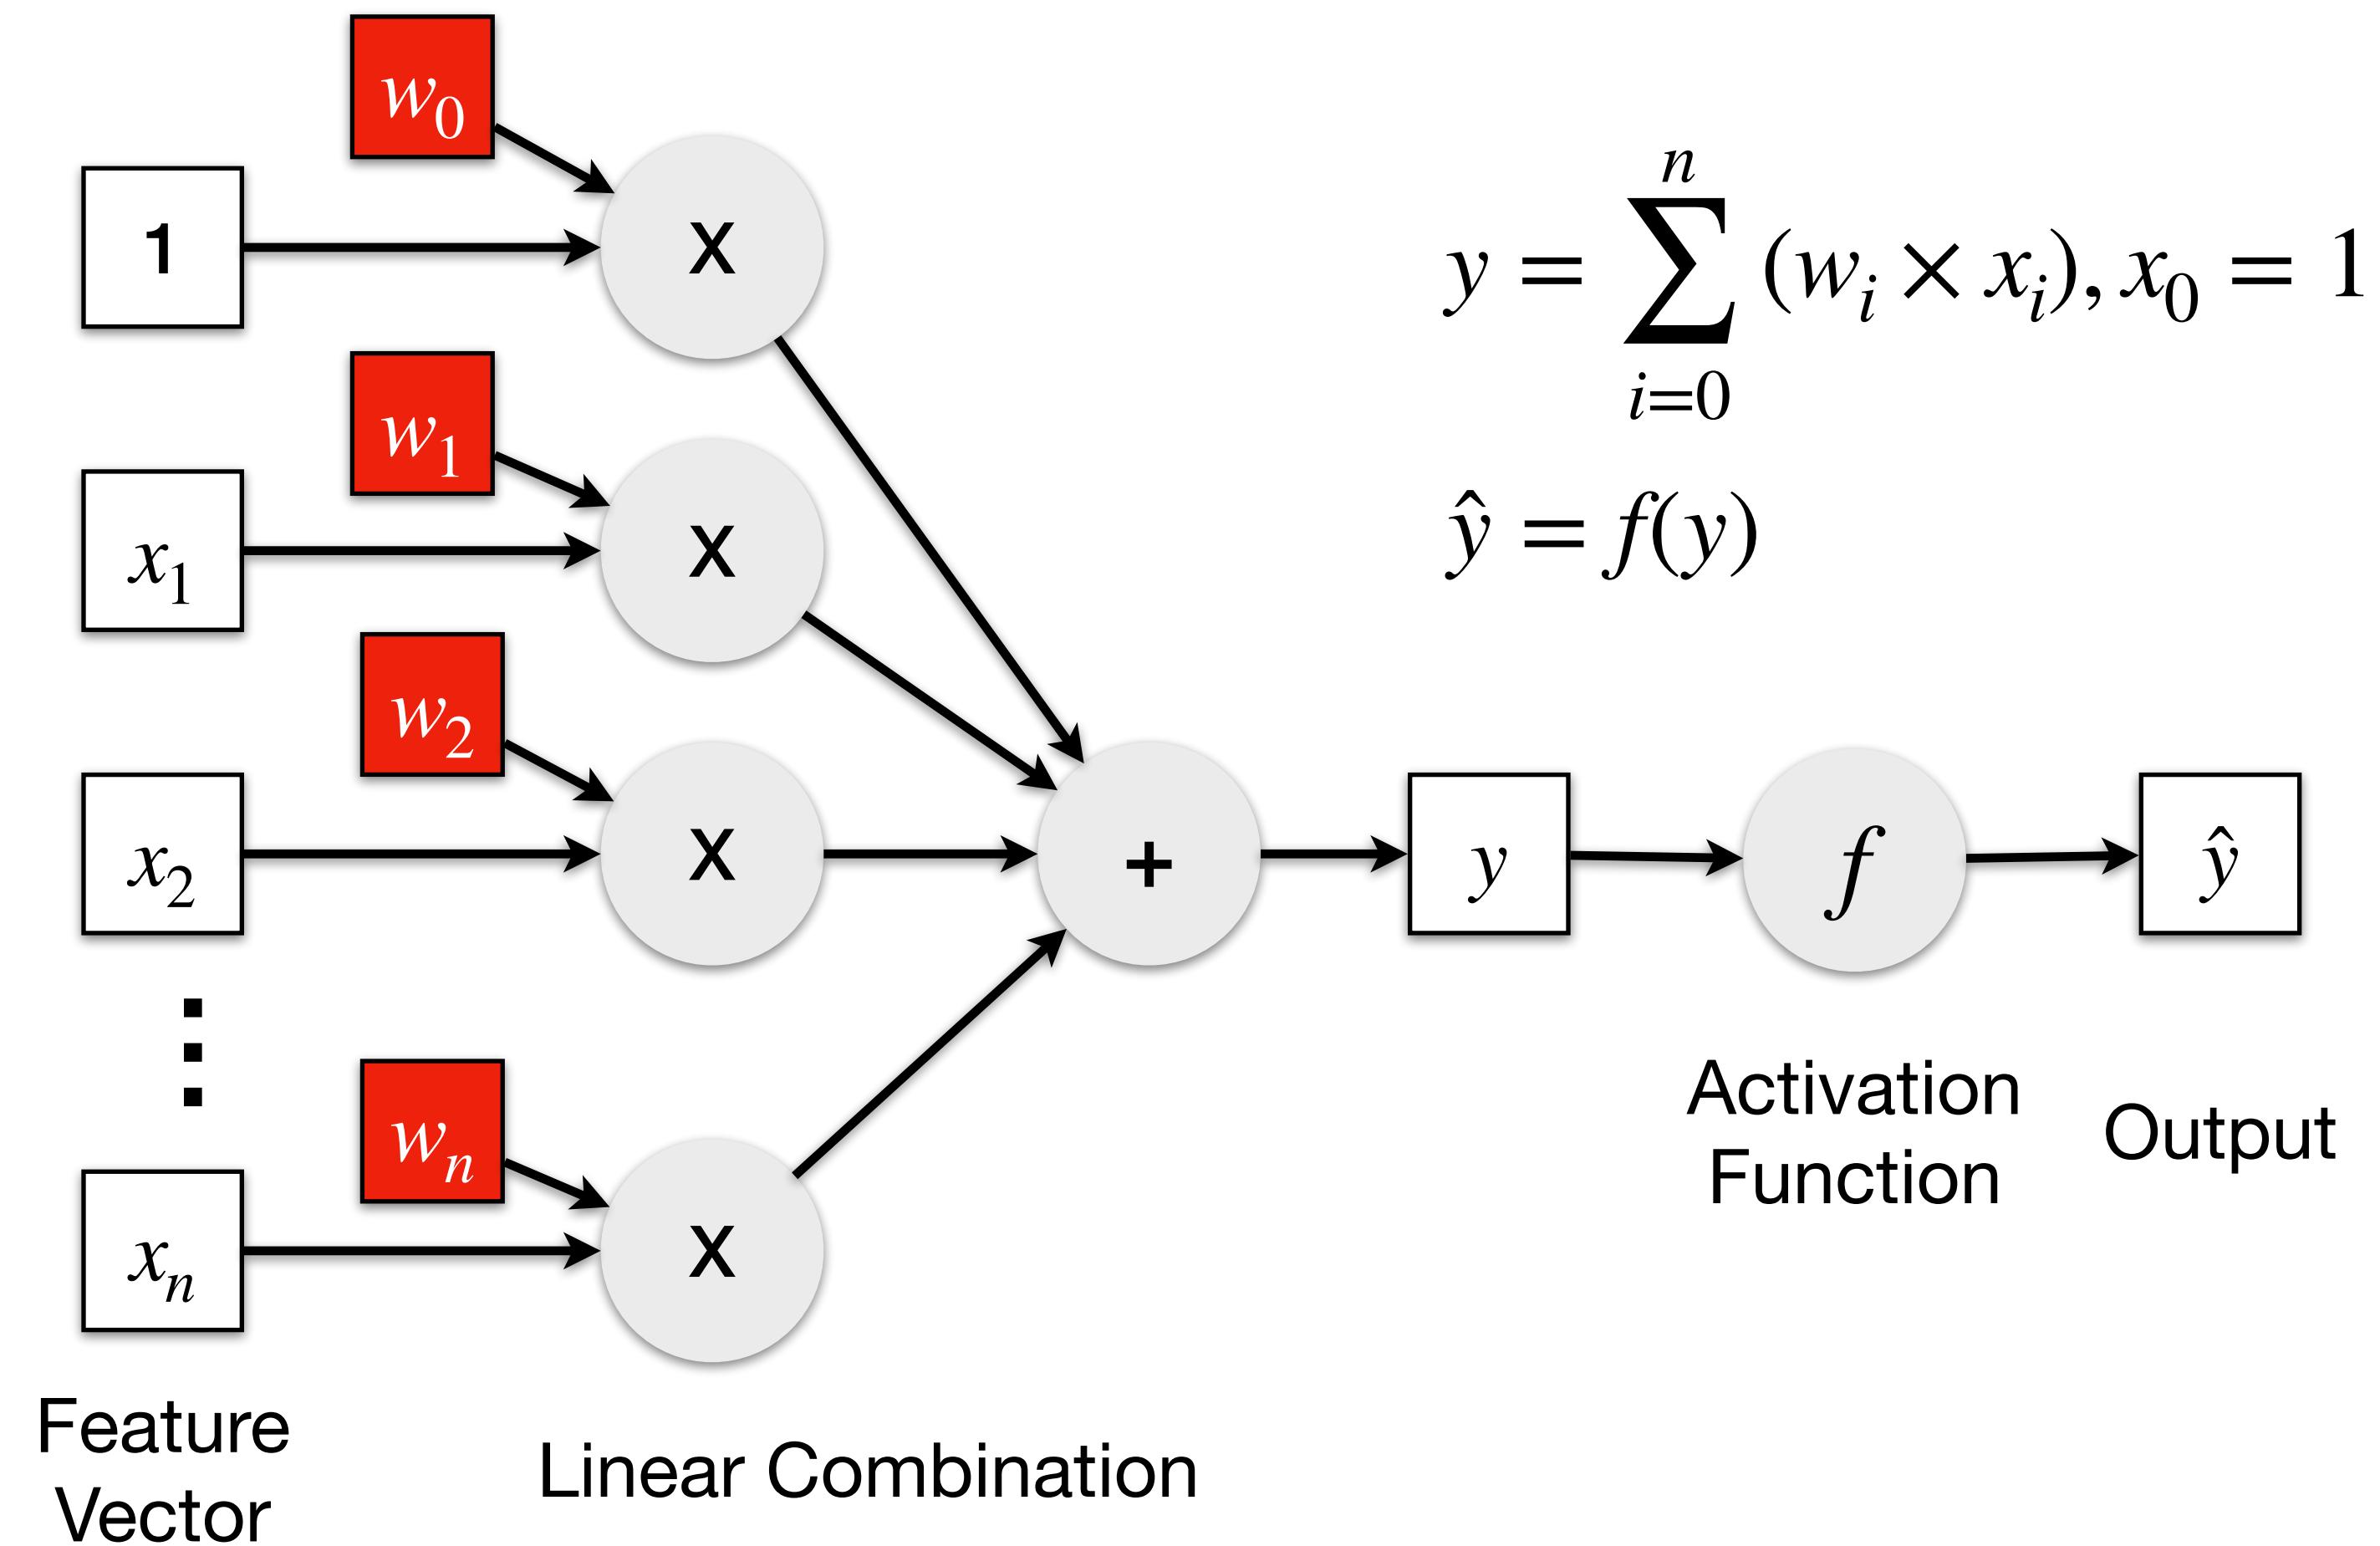
Delta Rule

Supervised learning process

Input: $\vec{x}_1, \vec{x}_2, \dots, \vec{x}_m$, m samples

Output: z_1, z_2, \dots, z_m , m labels

$$Loss(\vec{w}) = \sum_{k=1}^m (z_k - f(\vec{x}_k^T \cdot \vec{w}))^2$$



Neural Networks (NN)

Artificial Neuron

Building block of NNs.

Delta Rule

Supervised learning process

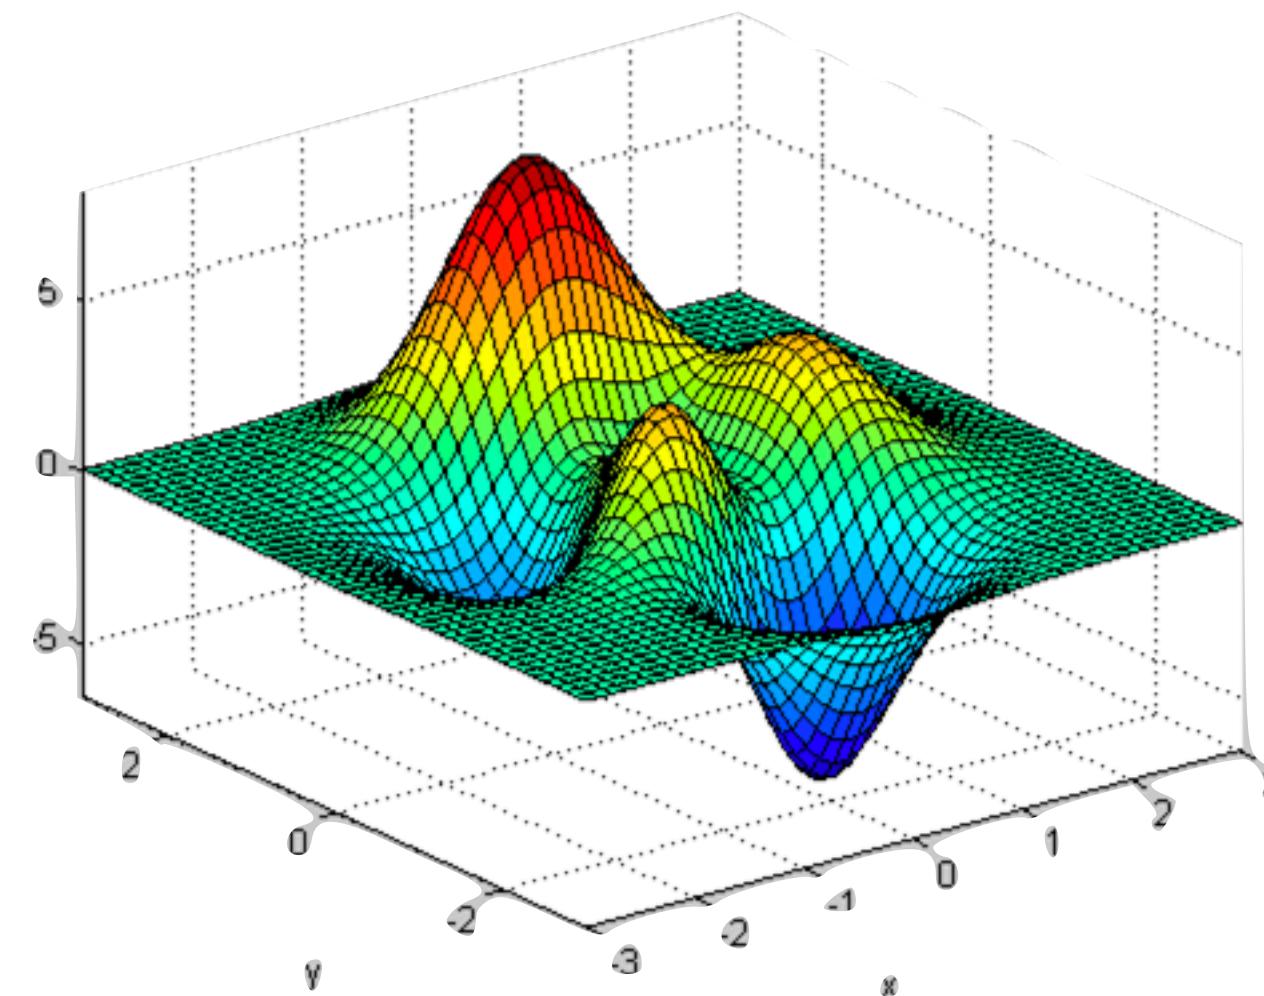
Input: $\vec{x}_1, \vec{x}_2, \dots, \vec{x}_m$, m samples

Output: z_1, z_2, \dots, z_m , m labels

$$Loss(\vec{w}) = \sum_{k=1}^m (z_k - f(\vec{x}_k^T \cdot \vec{w}))^2$$

Partial derivative:

$$\frac{\delta Loss(\vec{w})}{\delta w_i} = - \sum_{k=1}^m 2(z_k - f(\vec{x}_k^T \cdot \vec{w})) \times x_{ki} \times f'(\vec{x}_k^T \cdot \vec{w}) = 0$$



Loss Surface

$$y = \sum_{i=0}^n (w_i \times x_i), x_0 = 1$$
$$\hat{y} = f(y)$$



LOYOLA
UNIVERSITY CHICAGO

Neural Networks (NN)

Artificial Neuron

Building block of NNs.

Delta Rule

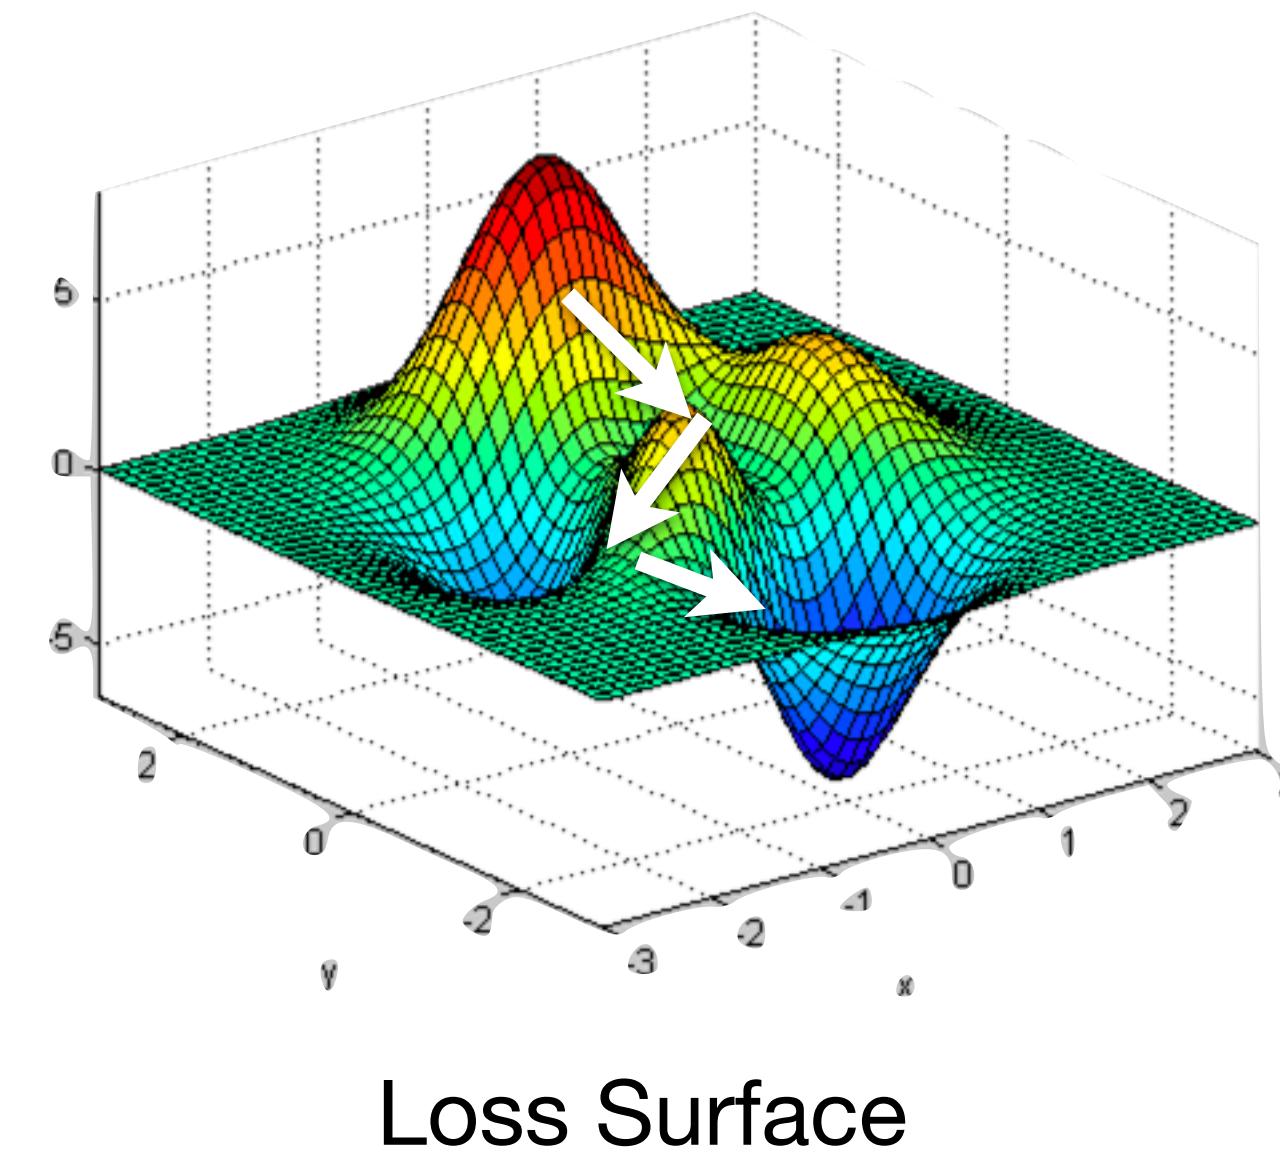
Supervised learning process

Input: $\vec{x}_1, \vec{x}_2, \dots, \vec{x}_m$, m samples

Output: z_1, z_2, \dots, z_m , m labels

$$Loss(\vec{w}) = \sum_{k=1}^m (z_k - f(\vec{x}_k^T \cdot \vec{w}))^2$$

$$\Delta w_i = \sum_{k=1}^m \alpha(z_k - f(\vec{x}_k^T \cdot \vec{w})) \times x_{ki} \times f'(\vec{x}_k^T \cdot \vec{w}) = 0, \quad \alpha \text{ is step size, one } \Delta \text{ for each weight.}$$

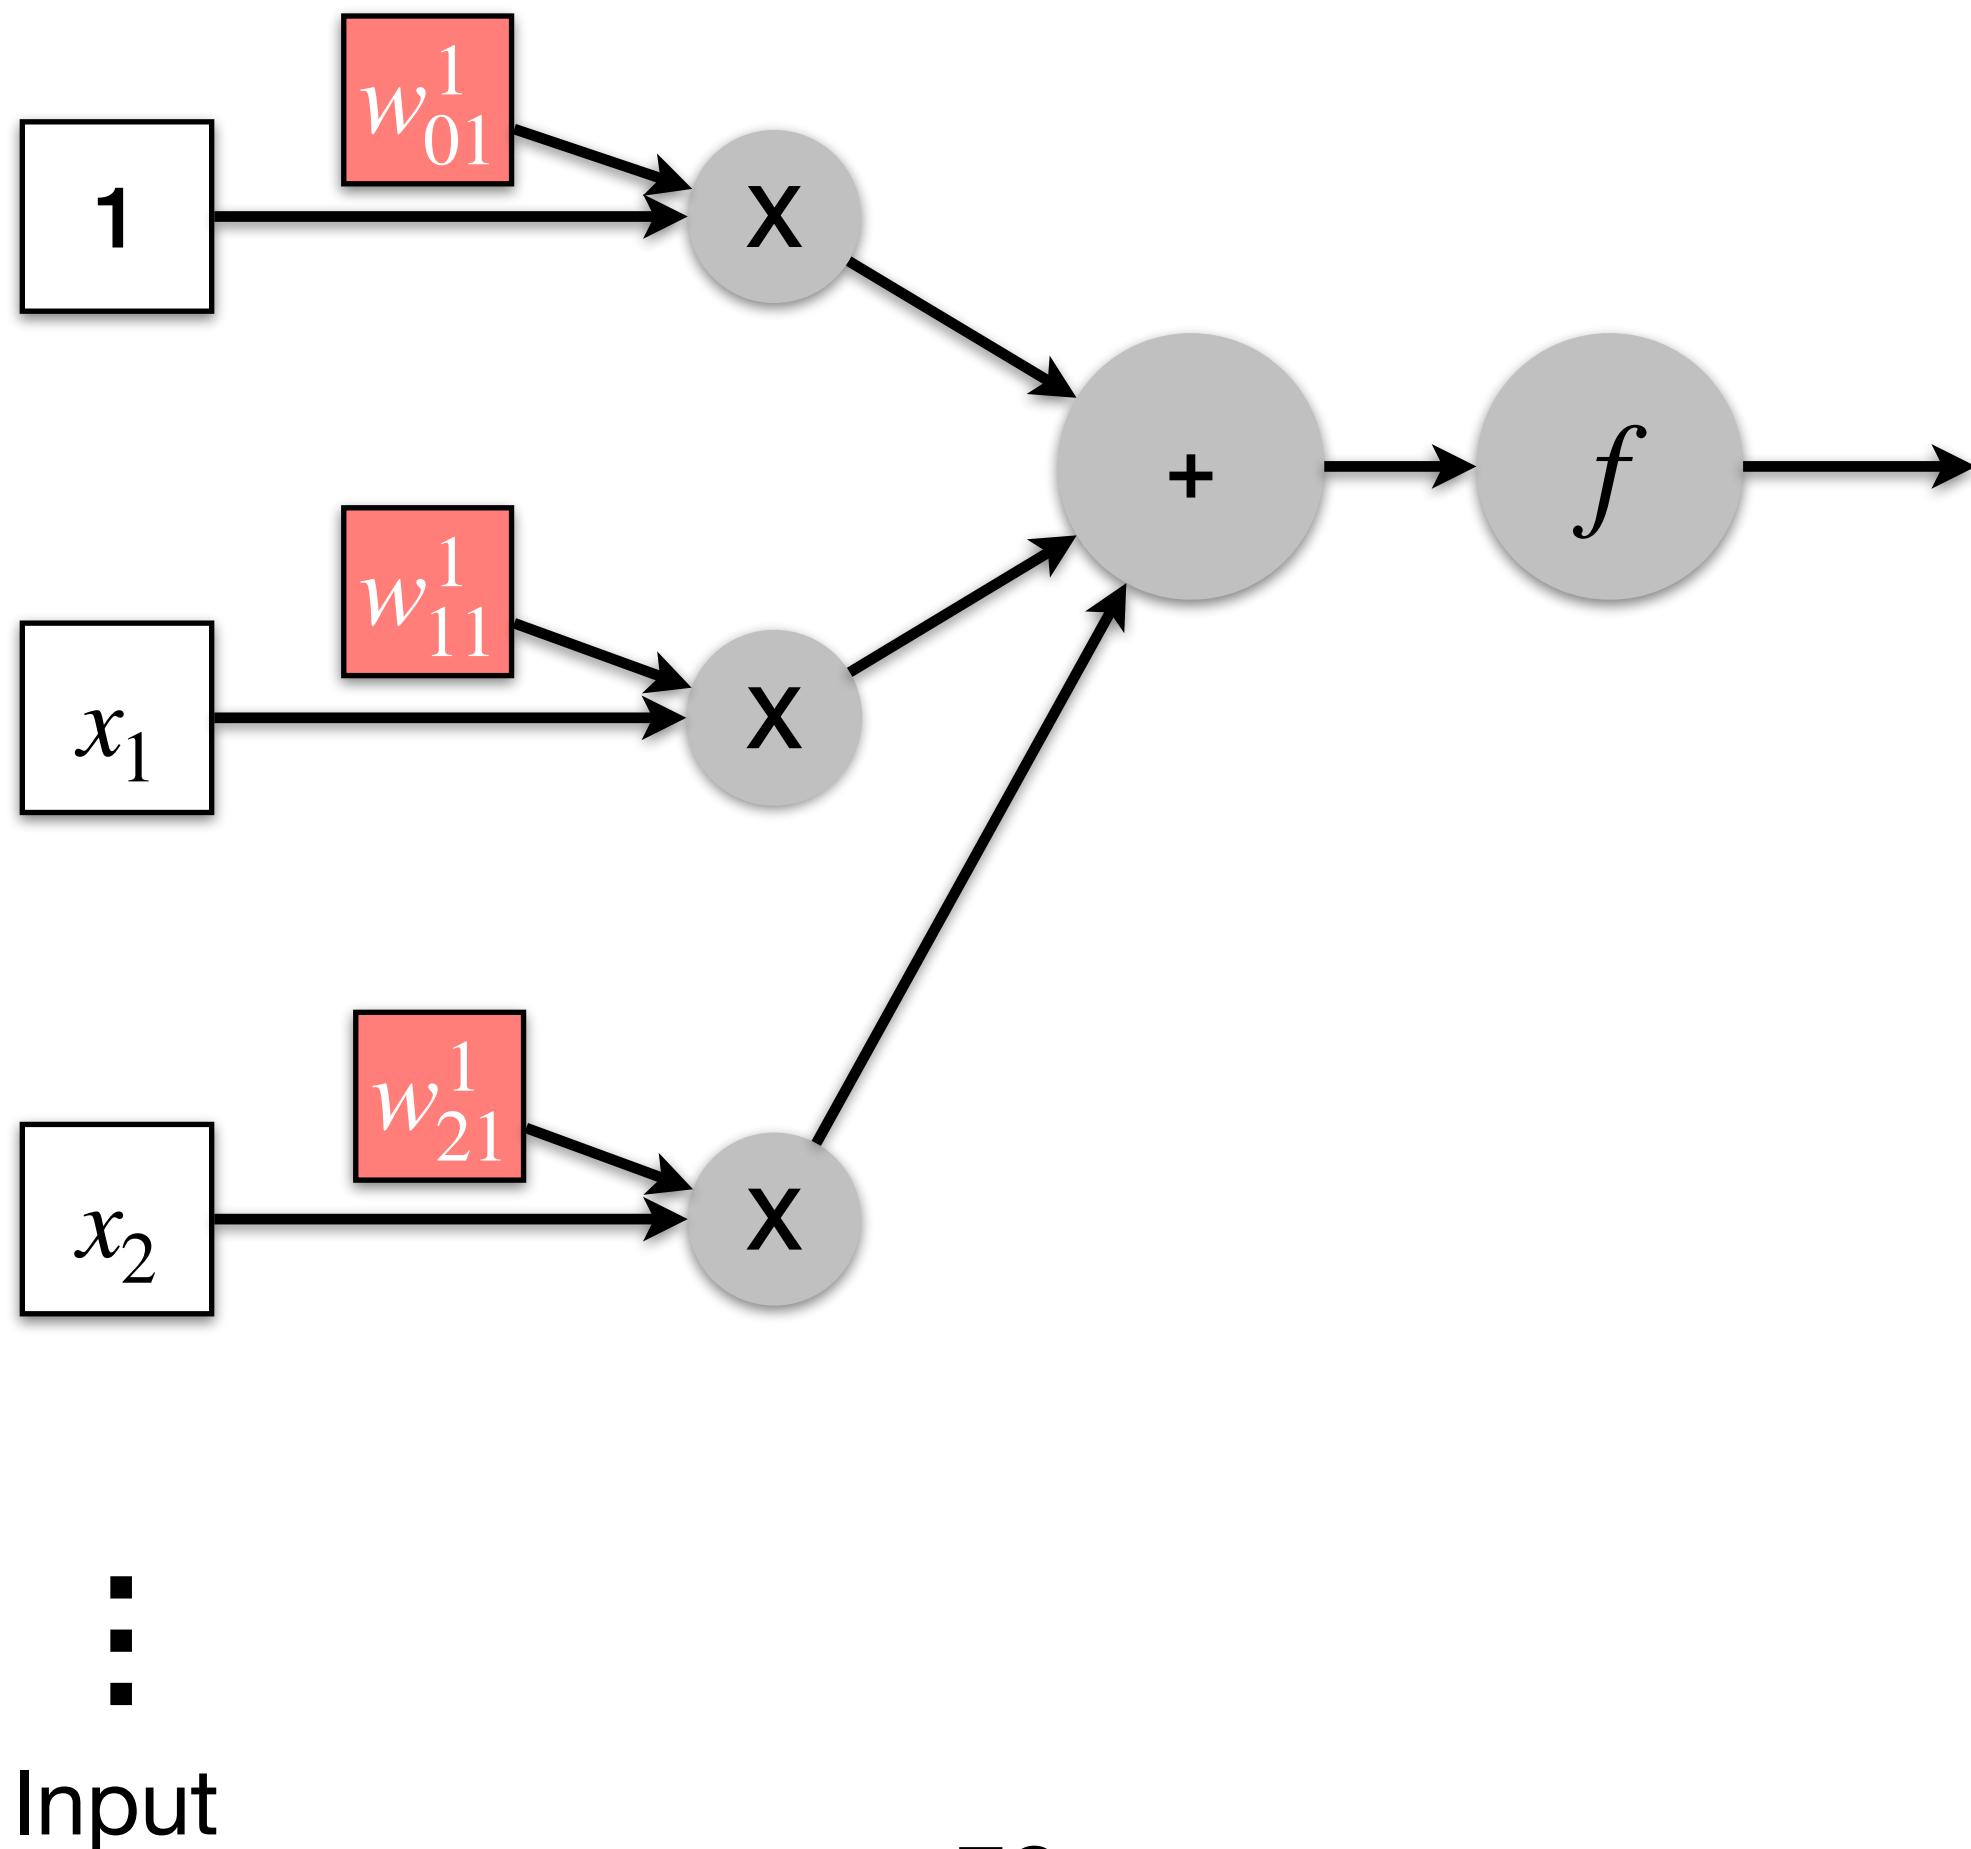


$$y = \sum_{i=0}^n (w_i \times x_i), x_0 = 1$$
$$\hat{y} = f(y)$$

Neural Networks (NN)

Adding Hidden Layers

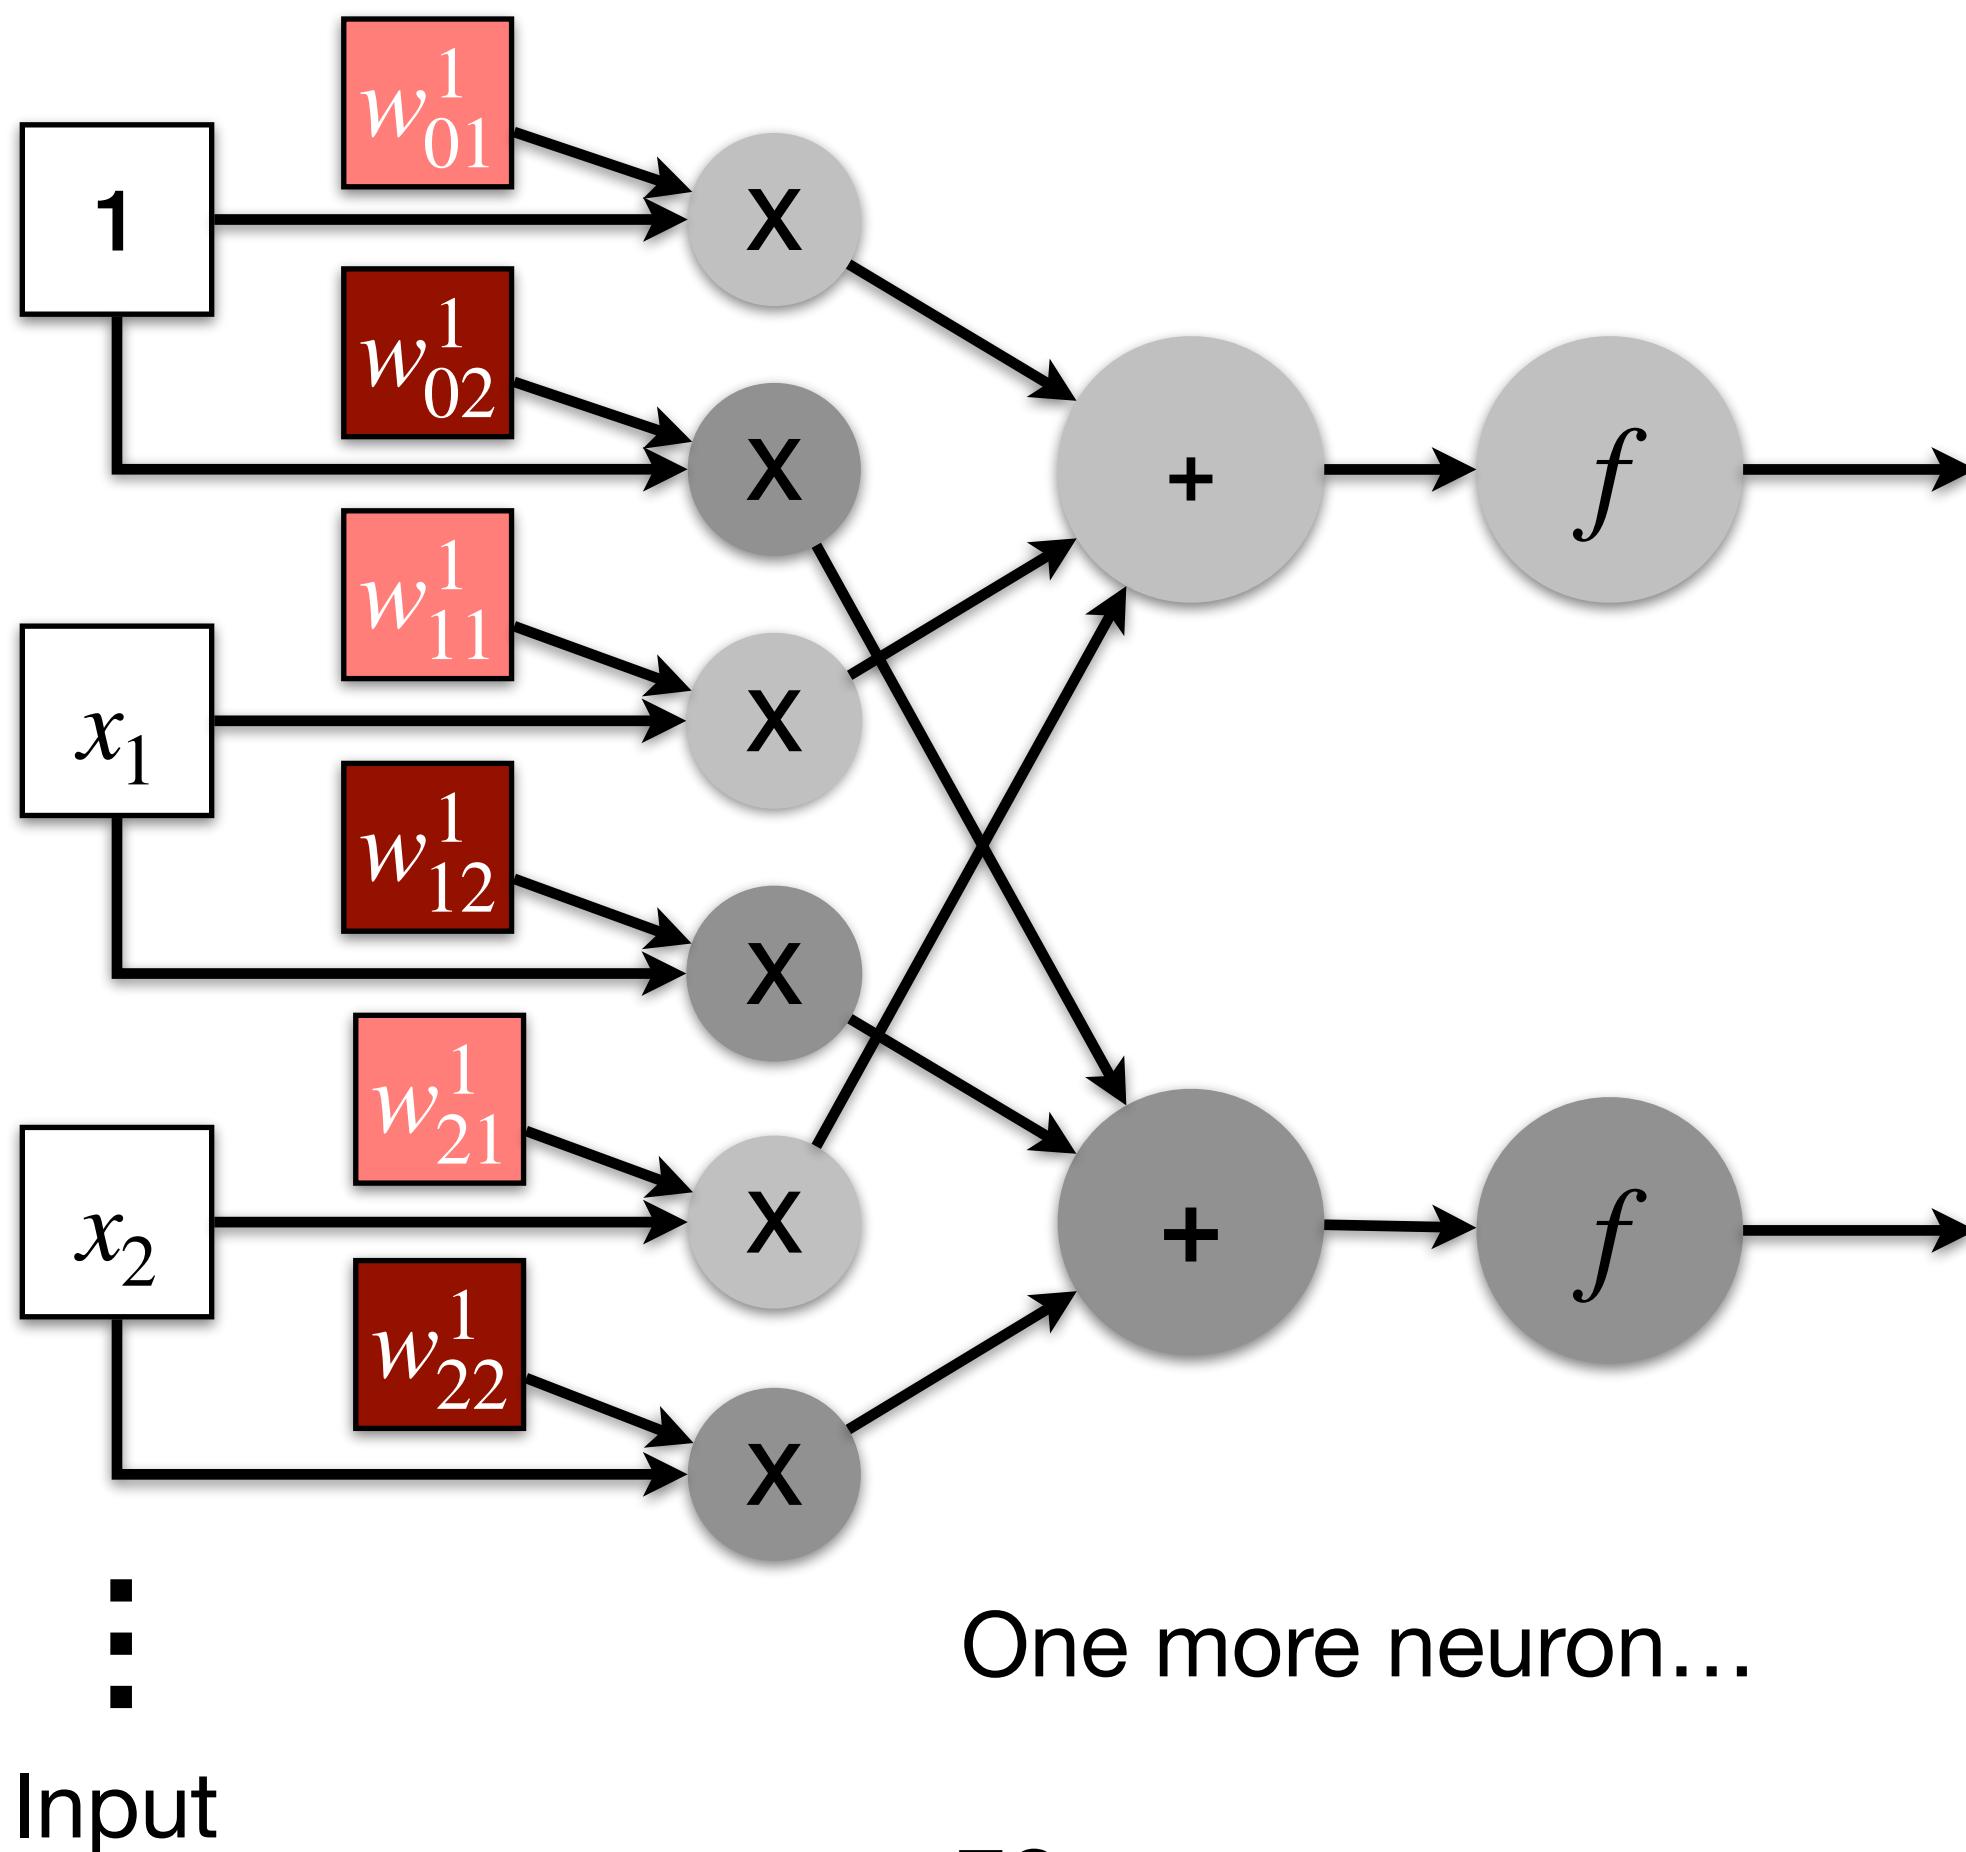
To increase
computing
power.



Neural Networks (NN)

Adding Hidden Layers

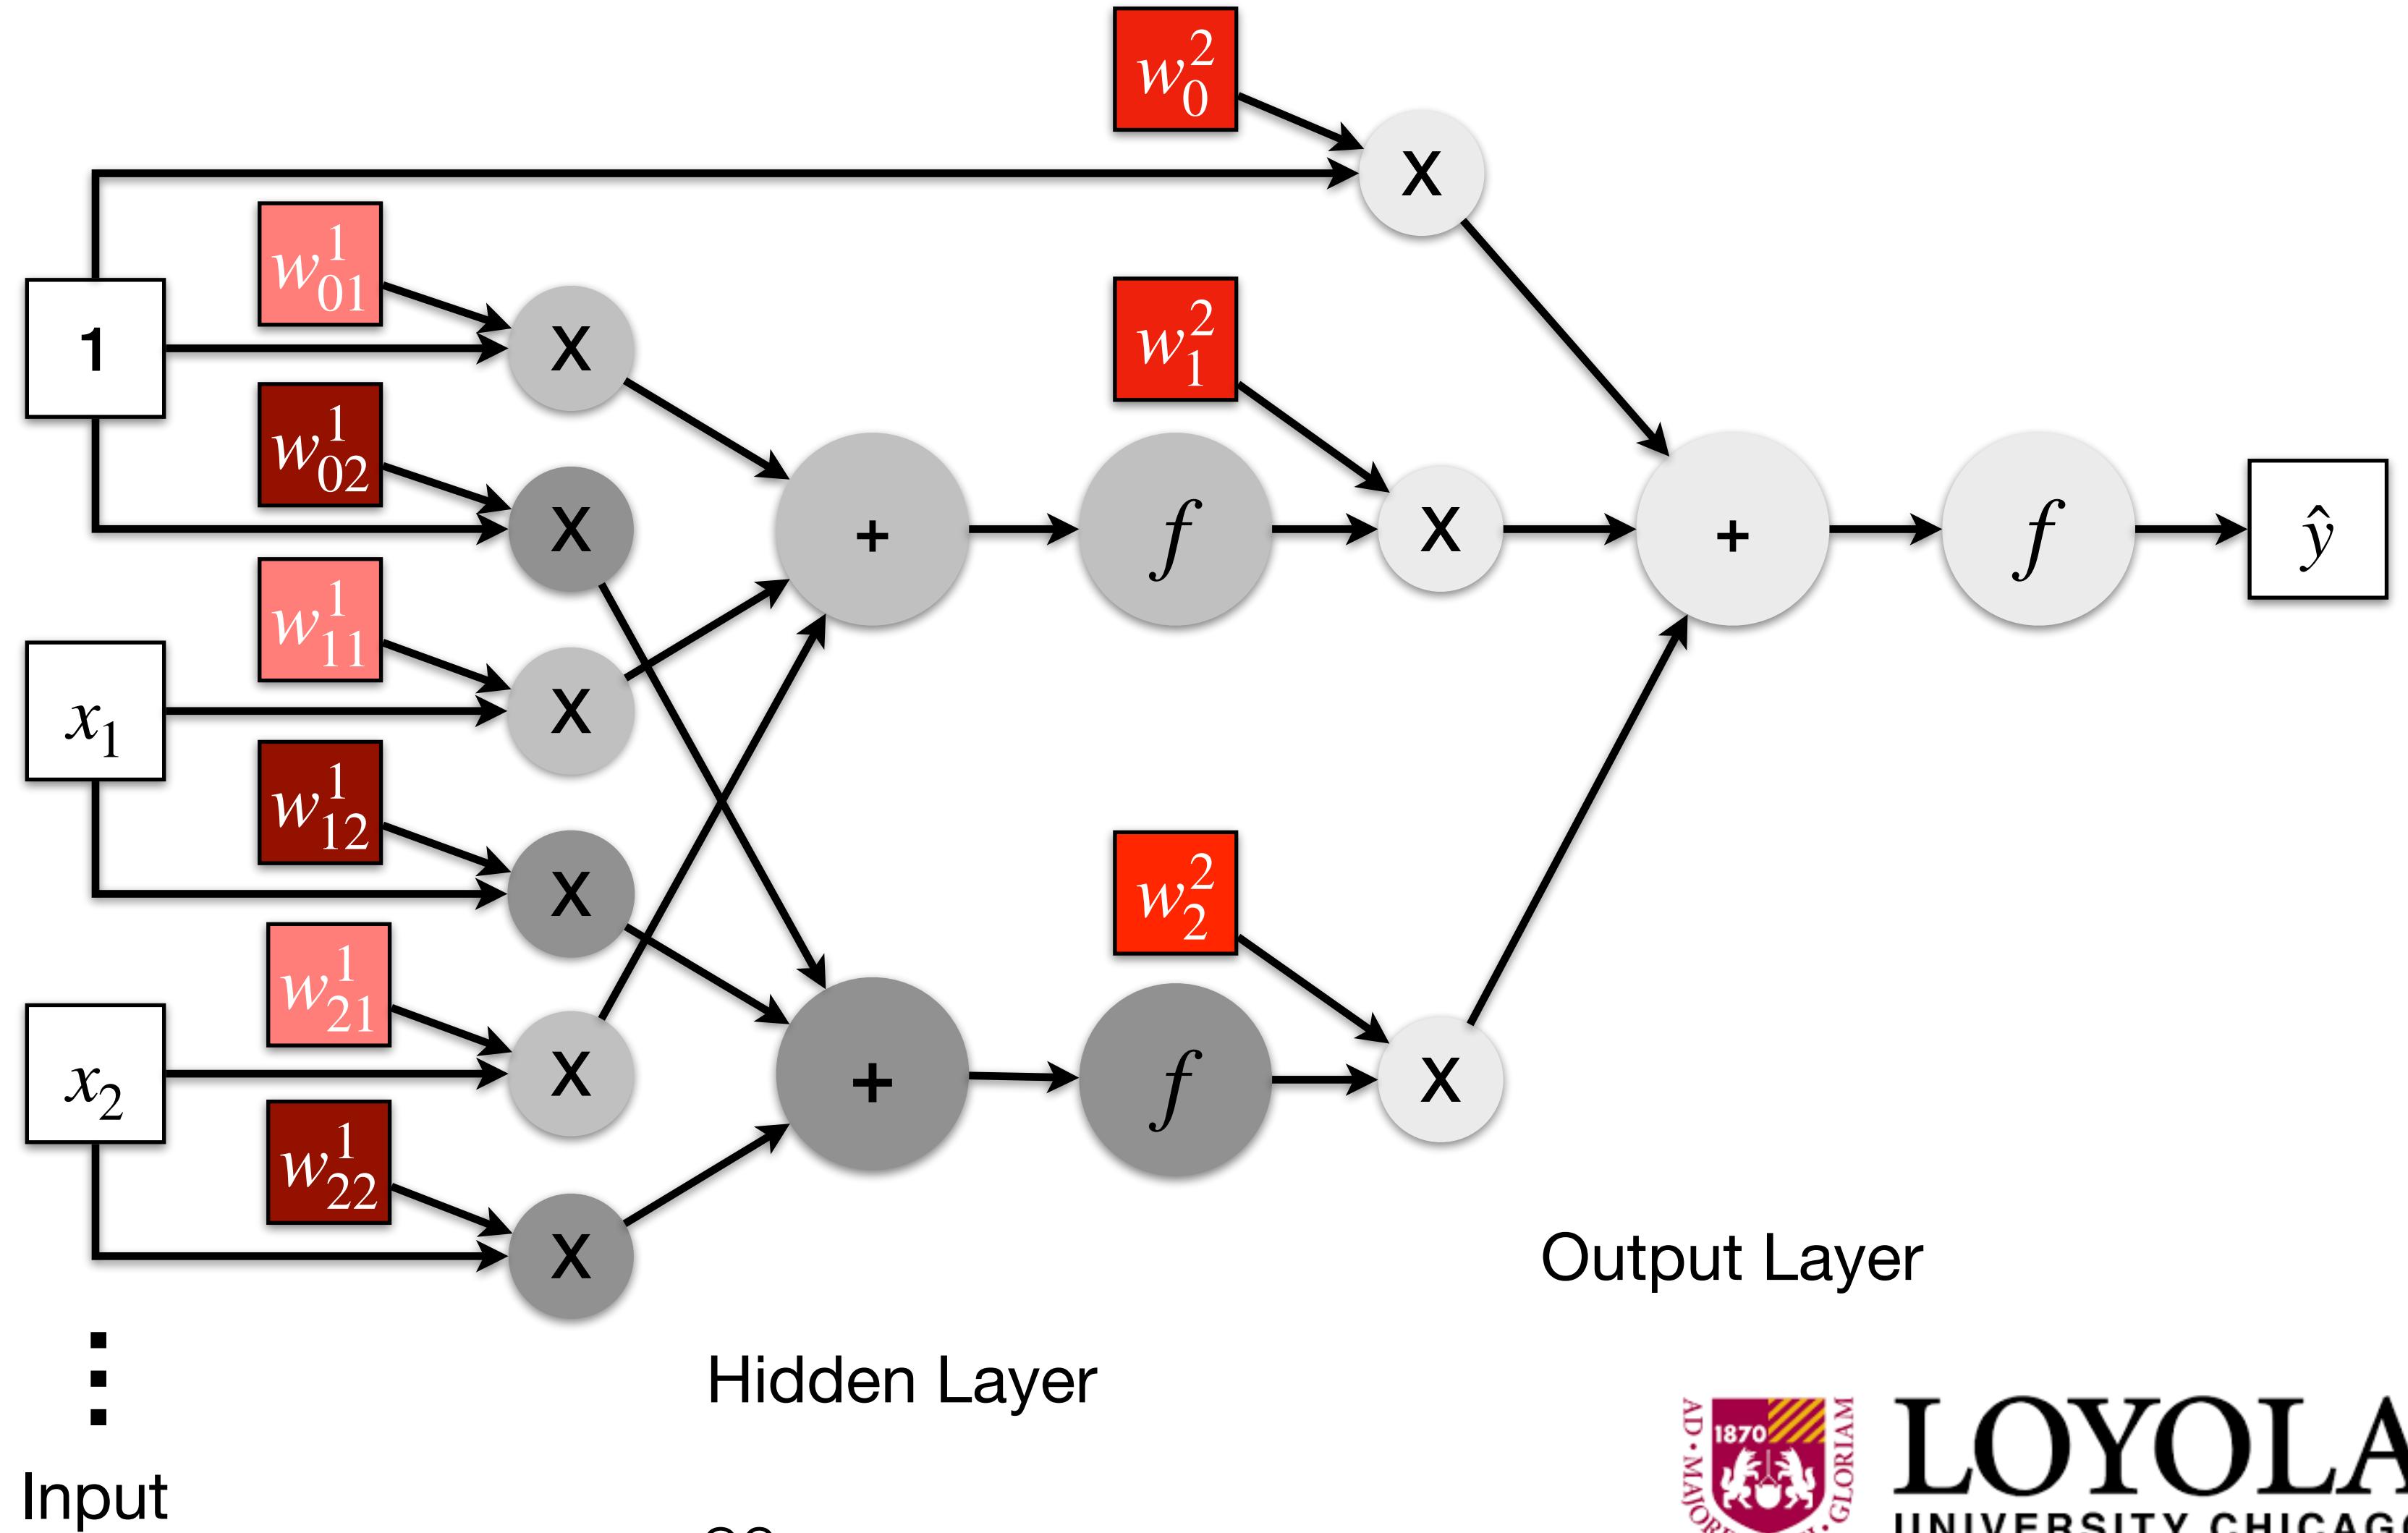
To increase
computing
power.



Neural Networks (NN)

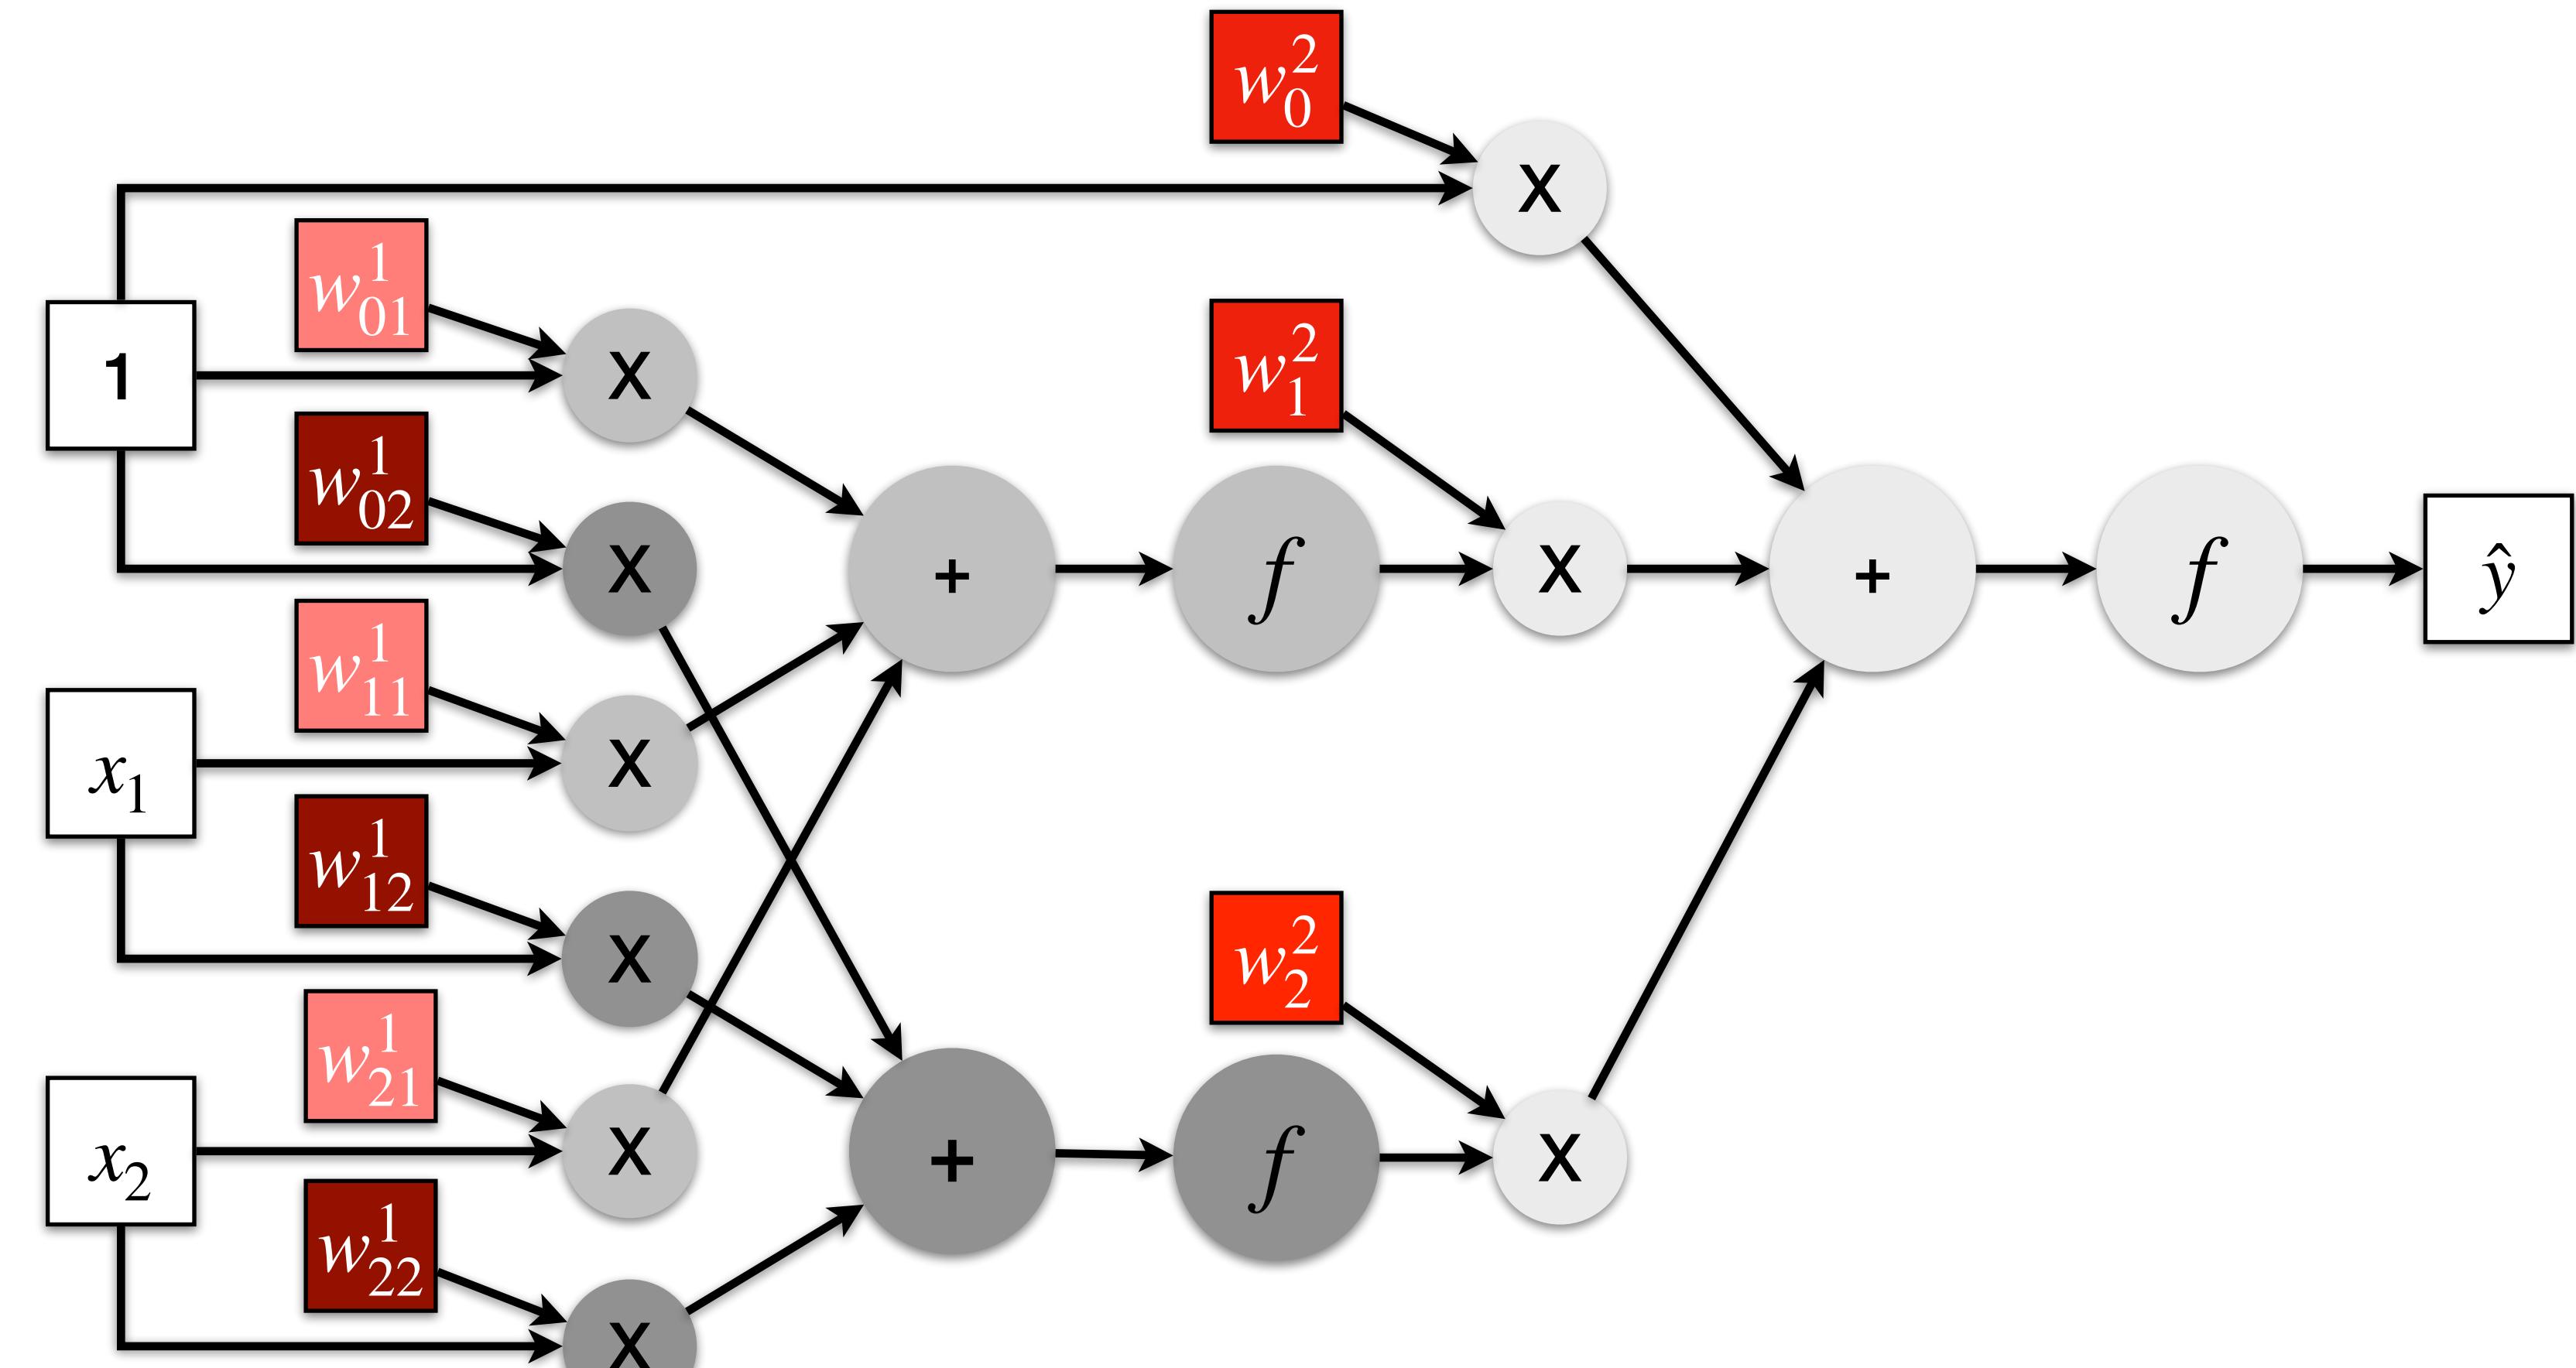
**Adding
Hidden Layers**

To increase
computing
power.



Adding Hidden Layers

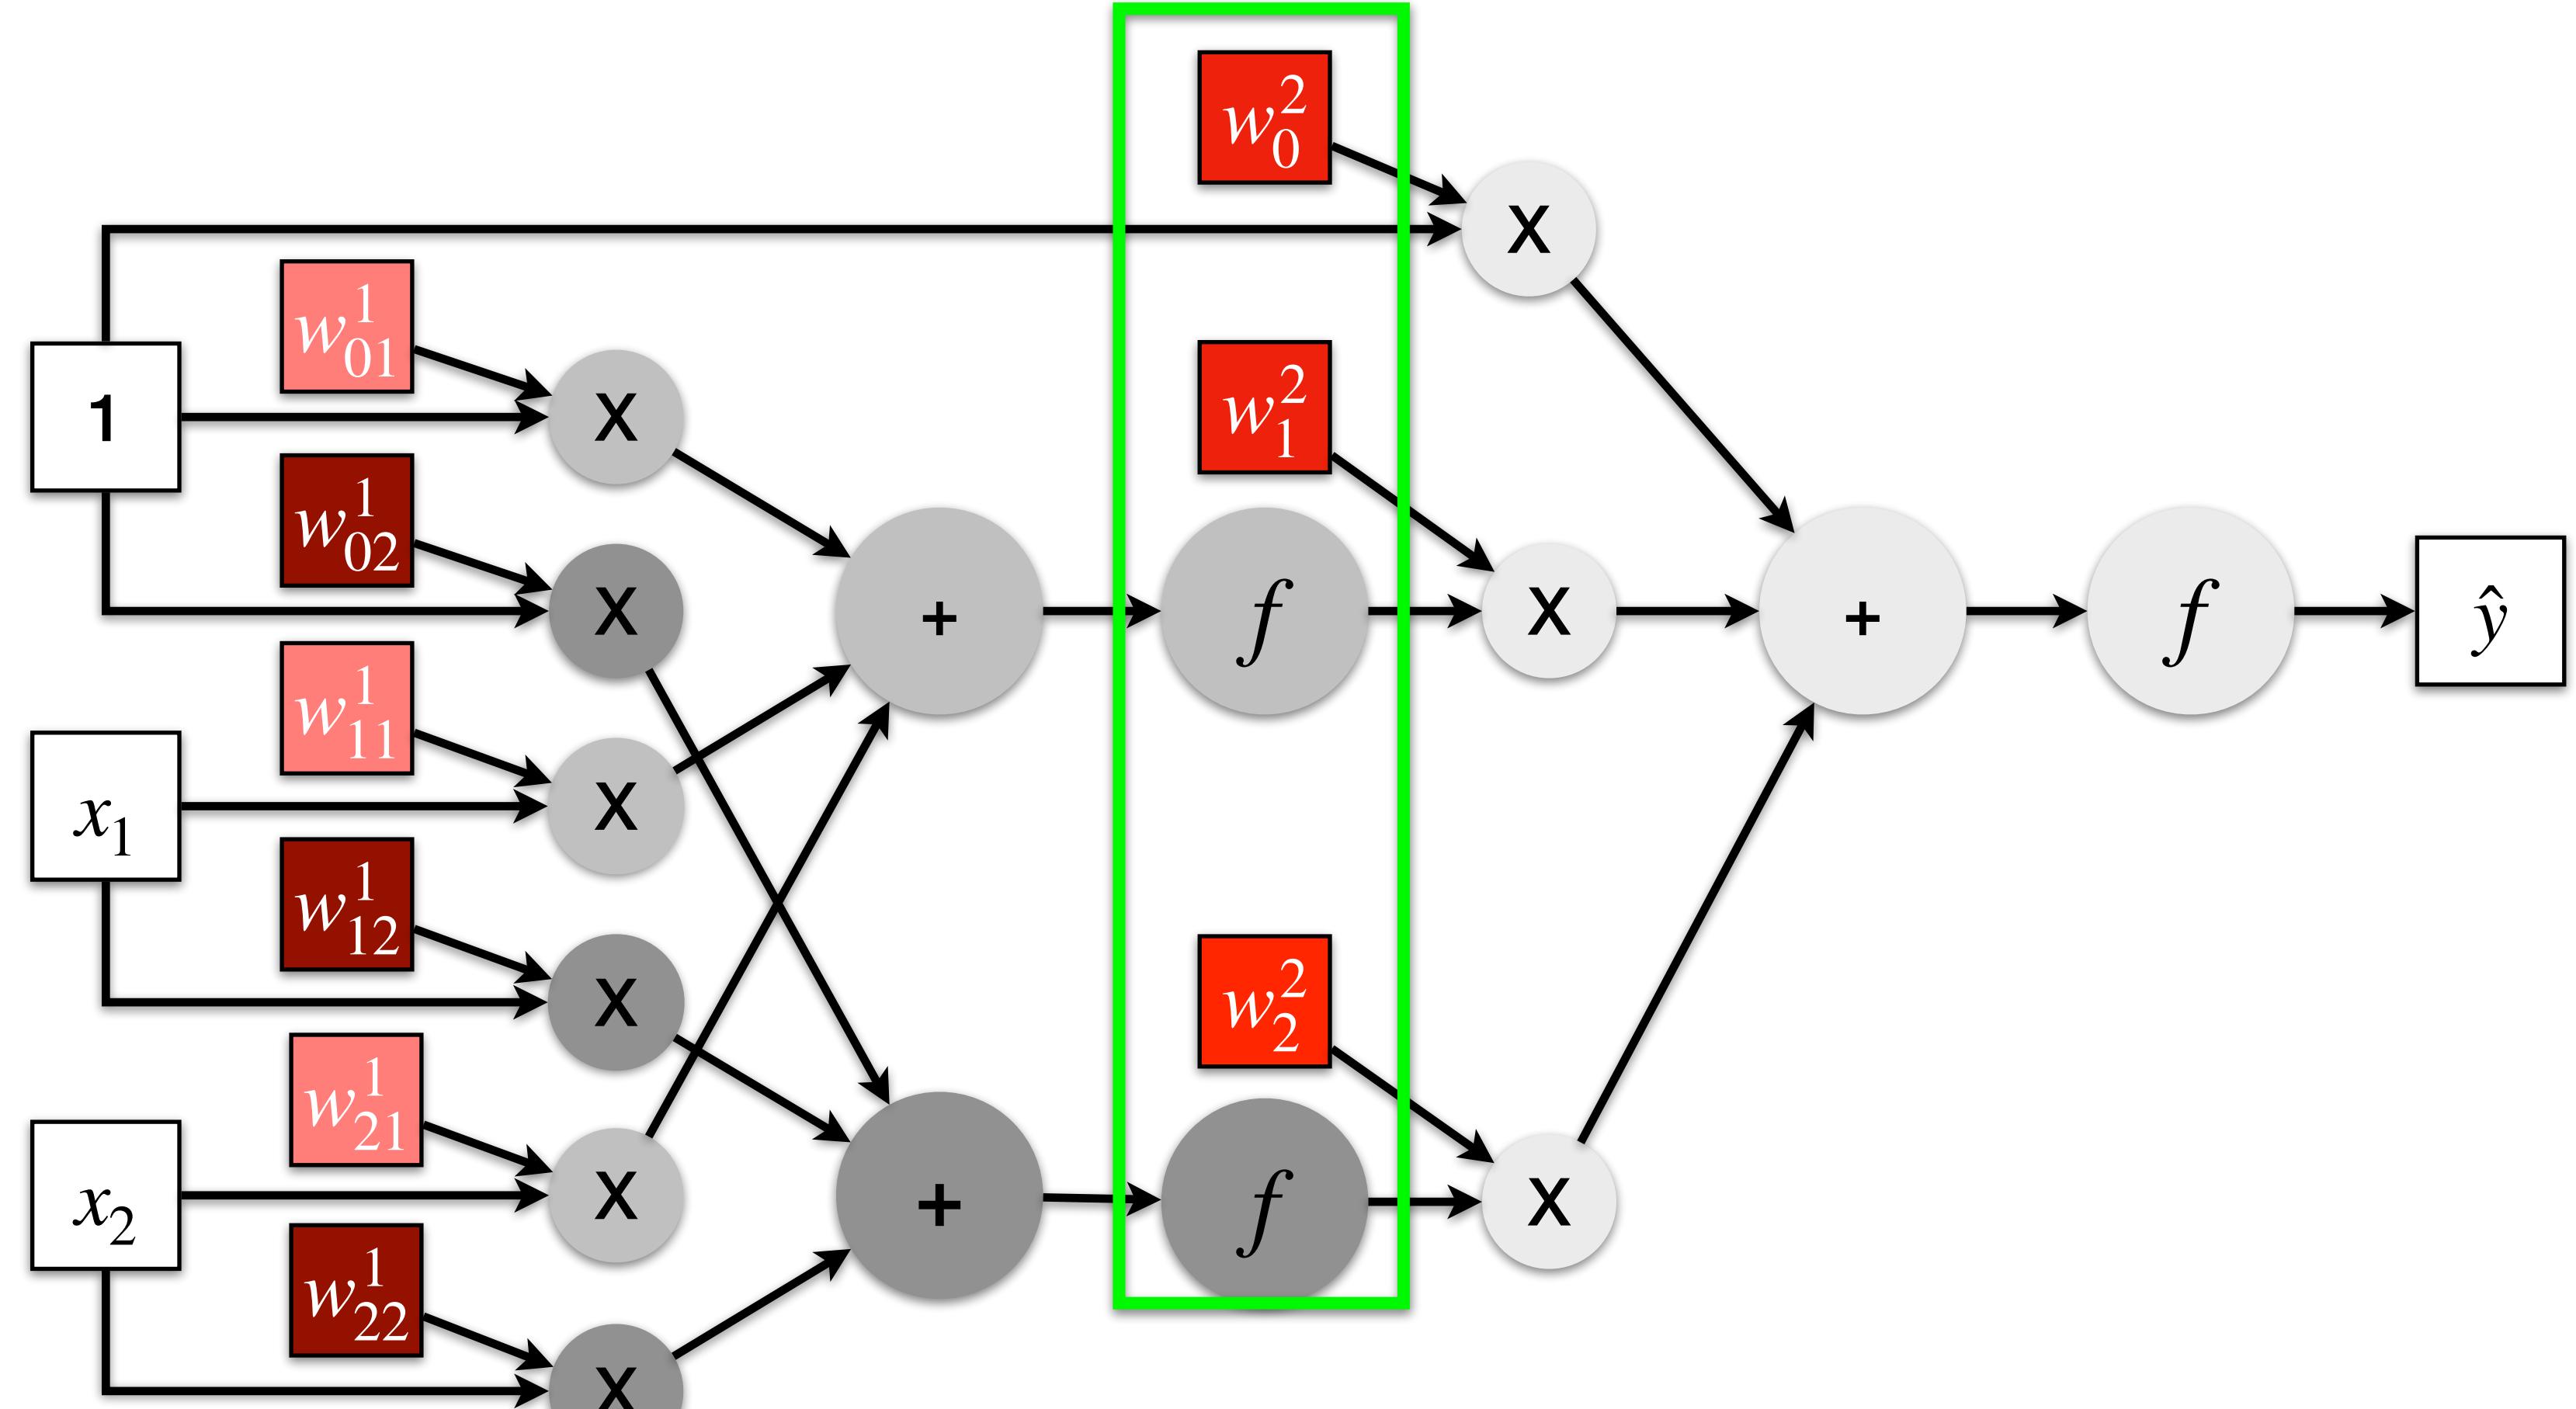
To increase computing power.



$$\hat{y} = f \left(\begin{pmatrix} w_1^2 & w_2^2 \end{pmatrix} f \left(\begin{pmatrix} w_{11}^1 & w_{21}^1 \\ w_{12}^1 & w_{22}^1 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} w_{01}^1 \\ w_{02}^1 \end{pmatrix} \right) + w_0^2 \right)$$

Adding Hidden Layers

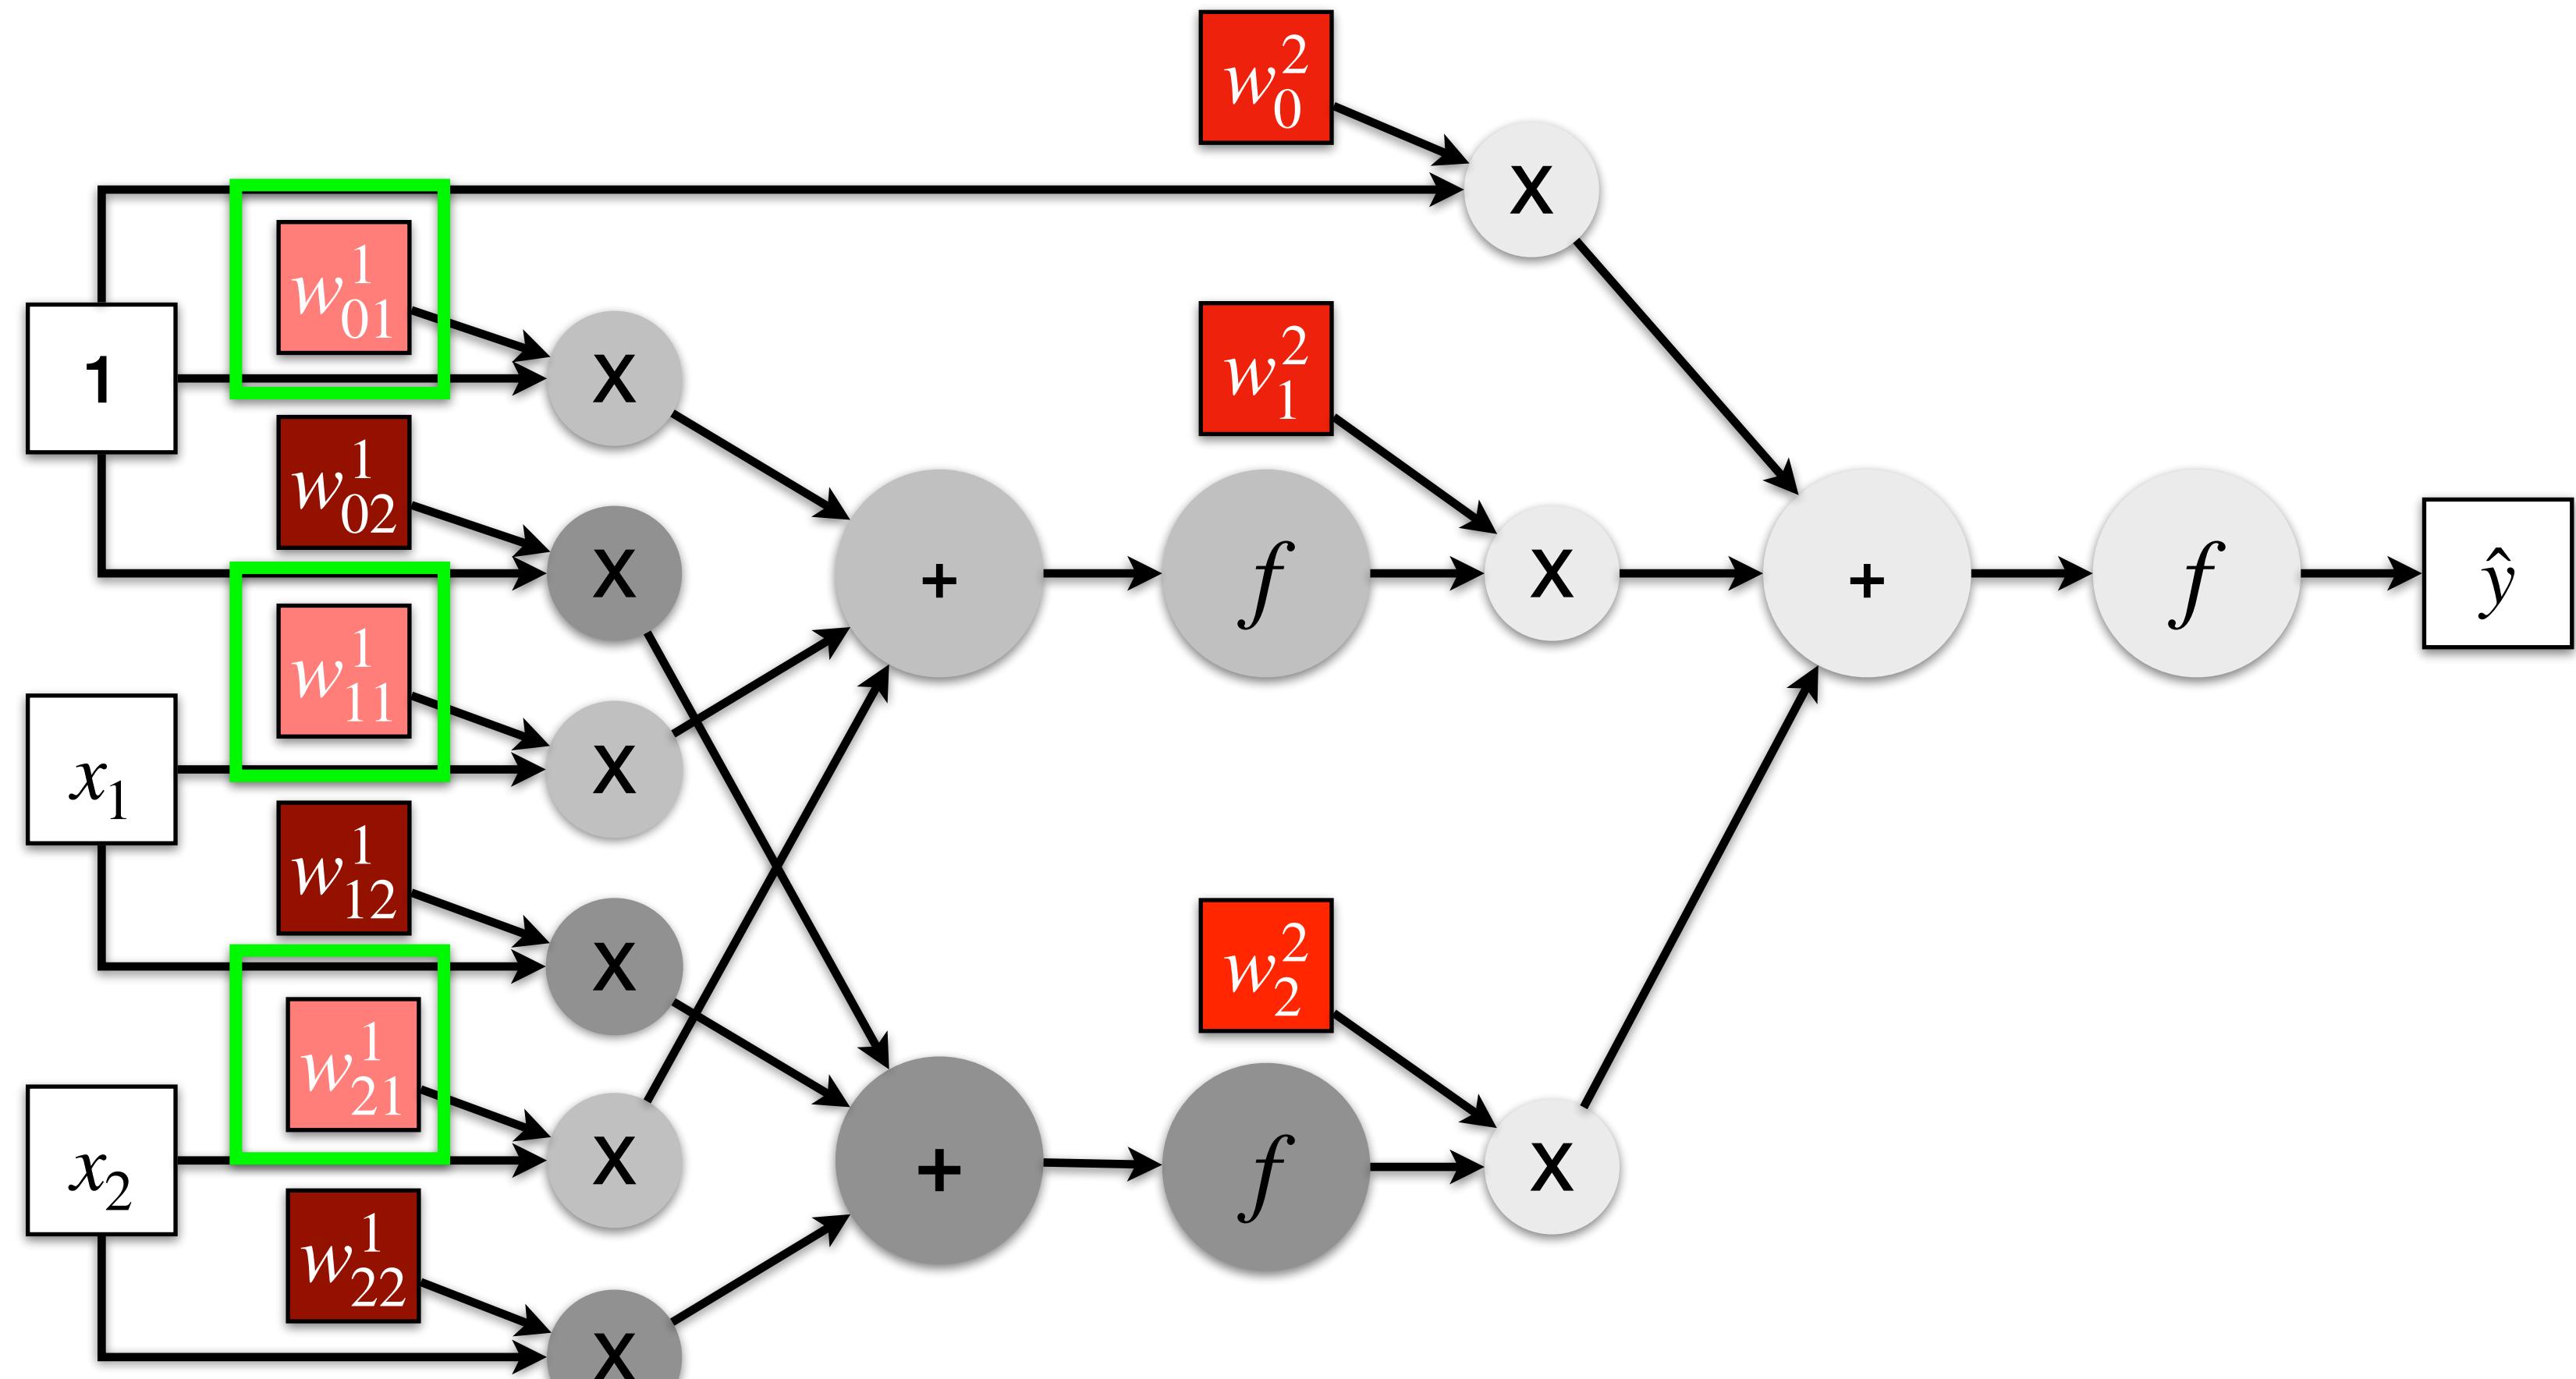
To increase computing power.



$$\hat{y} = f \left(\begin{pmatrix} w_1^2 & w_2^2 \end{pmatrix} f \left(\begin{pmatrix} w_{11}^1 & w_{21}^1 \\ w_{12}^1 & w_{22}^1 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} w_{01}^1 \\ x_{02}^1 \end{pmatrix} \right) + w_0^2 \right)$$

Adding Hidden Layers

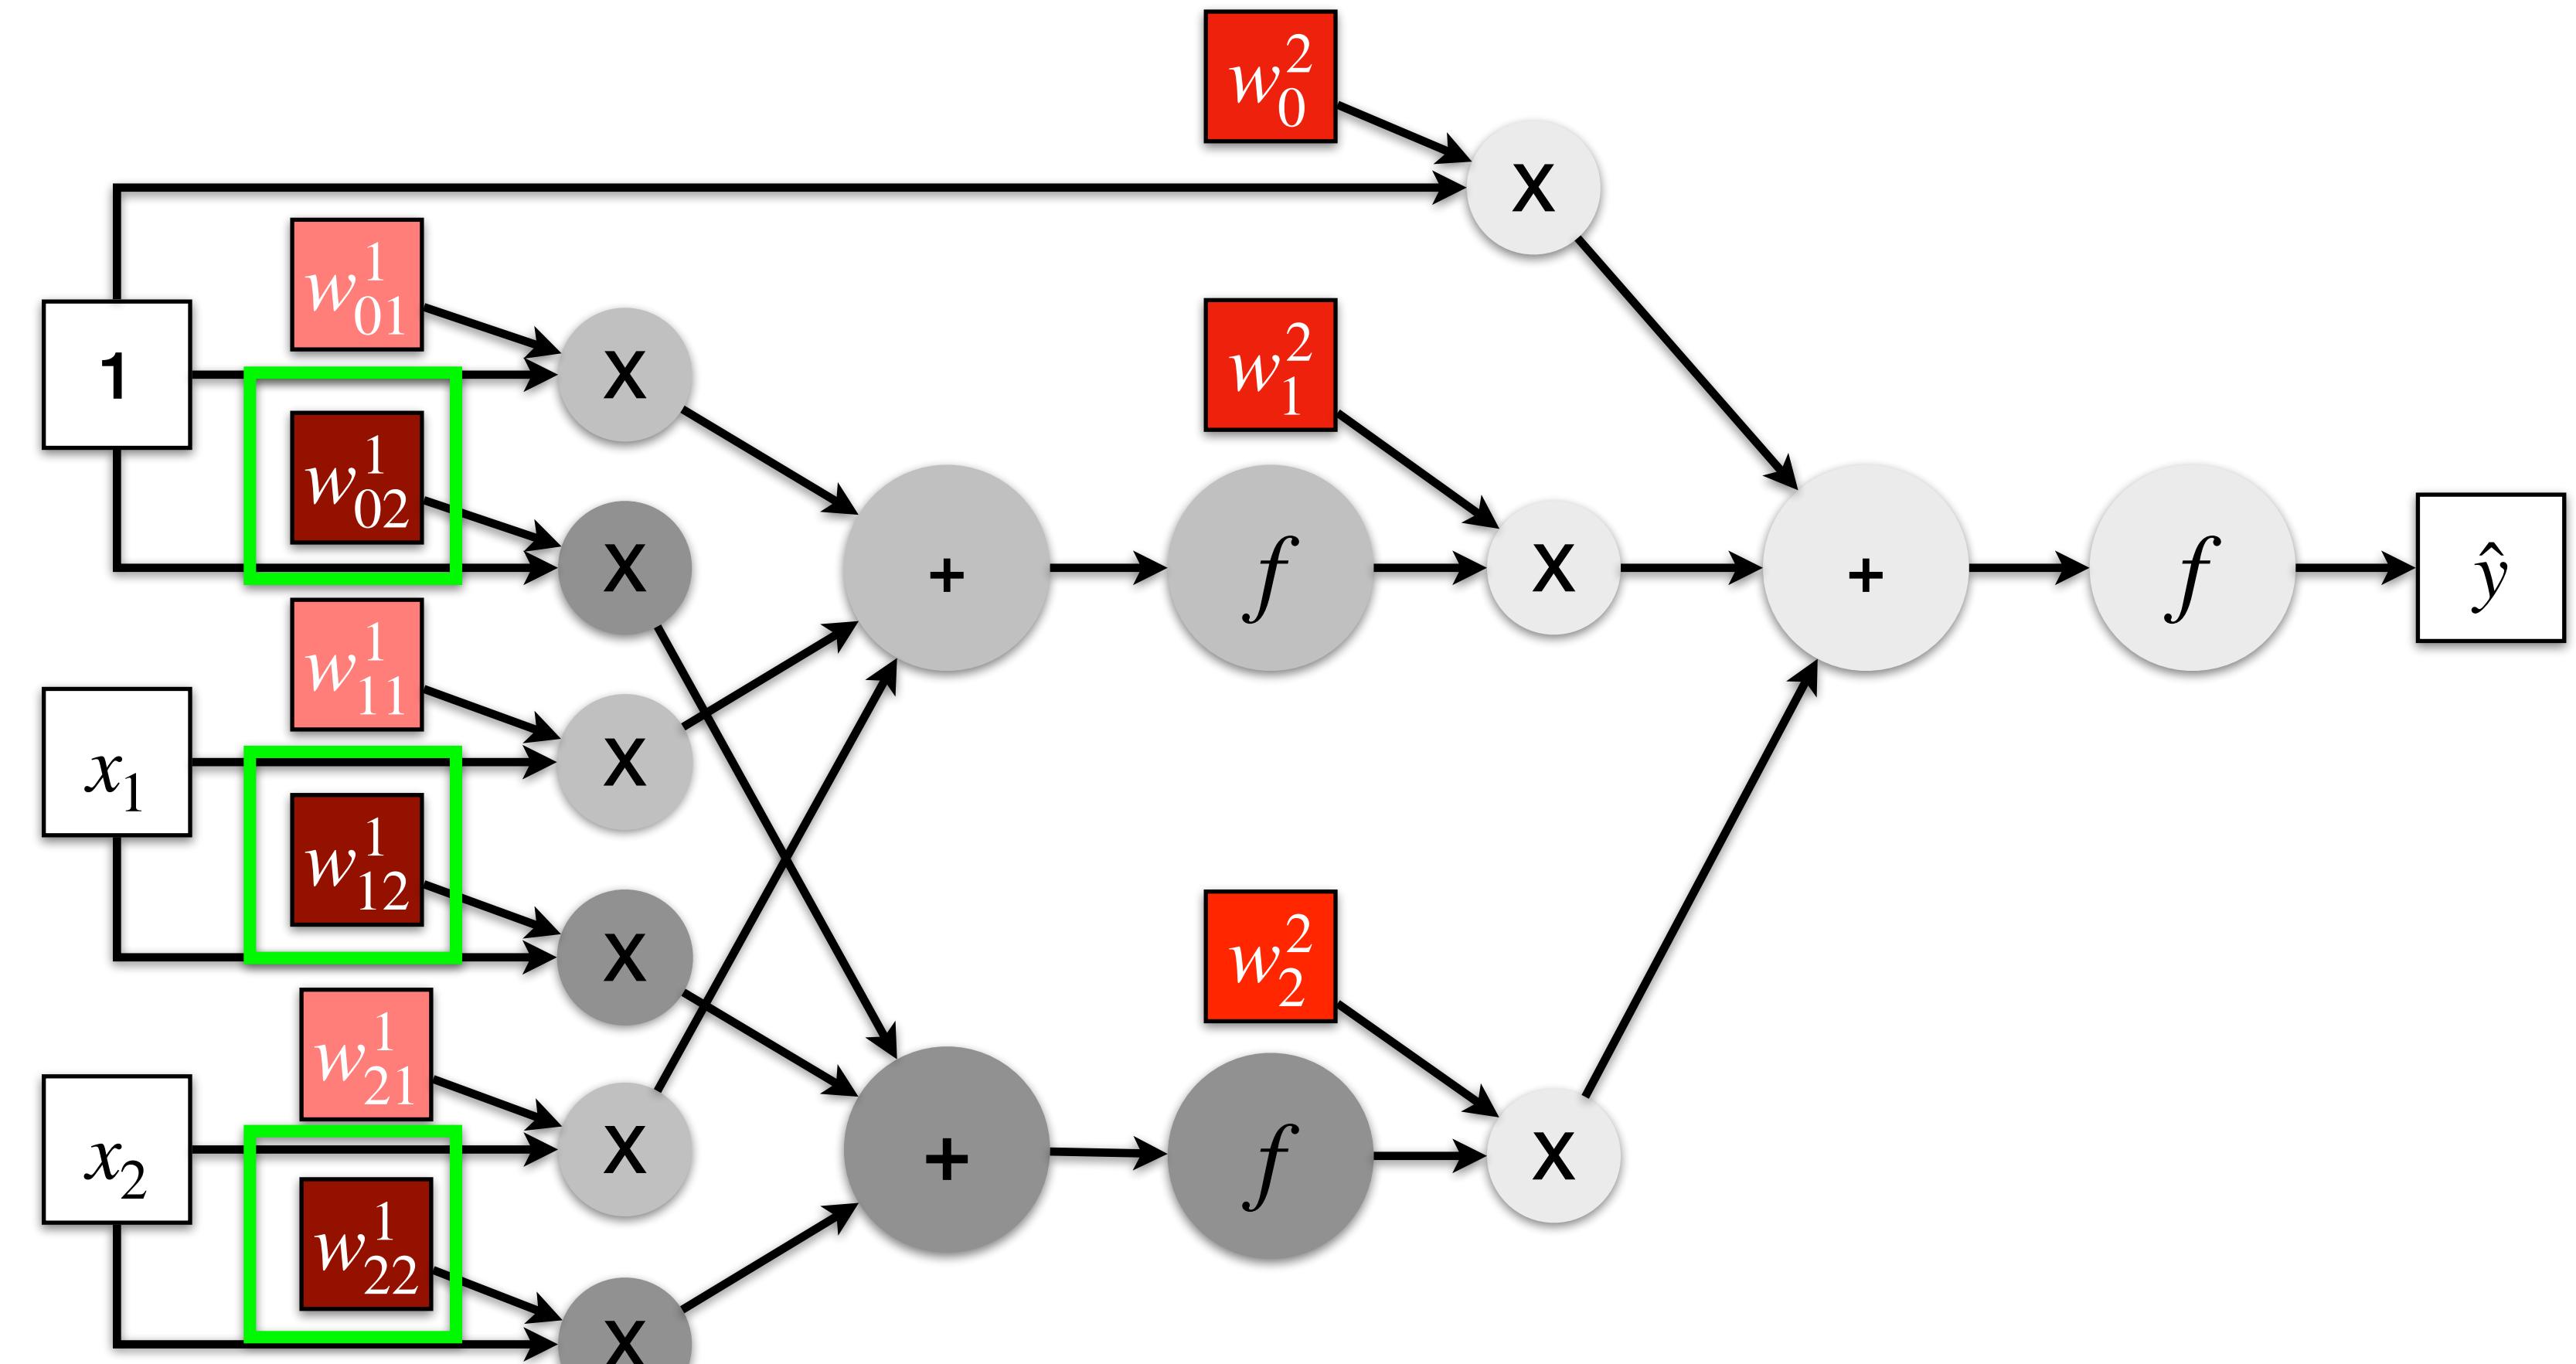
To increase computing power.



$$\hat{y} = f \left(\begin{pmatrix} w_1^2 & w_2^2 \end{pmatrix} f \left(\begin{pmatrix} w_{11}^1 & w_{21}^1 \\ w_{12}^1 & w_{22}^1 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} w_{01}^1 \\ x_{02}^1 \end{pmatrix} \right) + w_0^2 \right)$$

Adding Hidden Layers

To increase computing power.

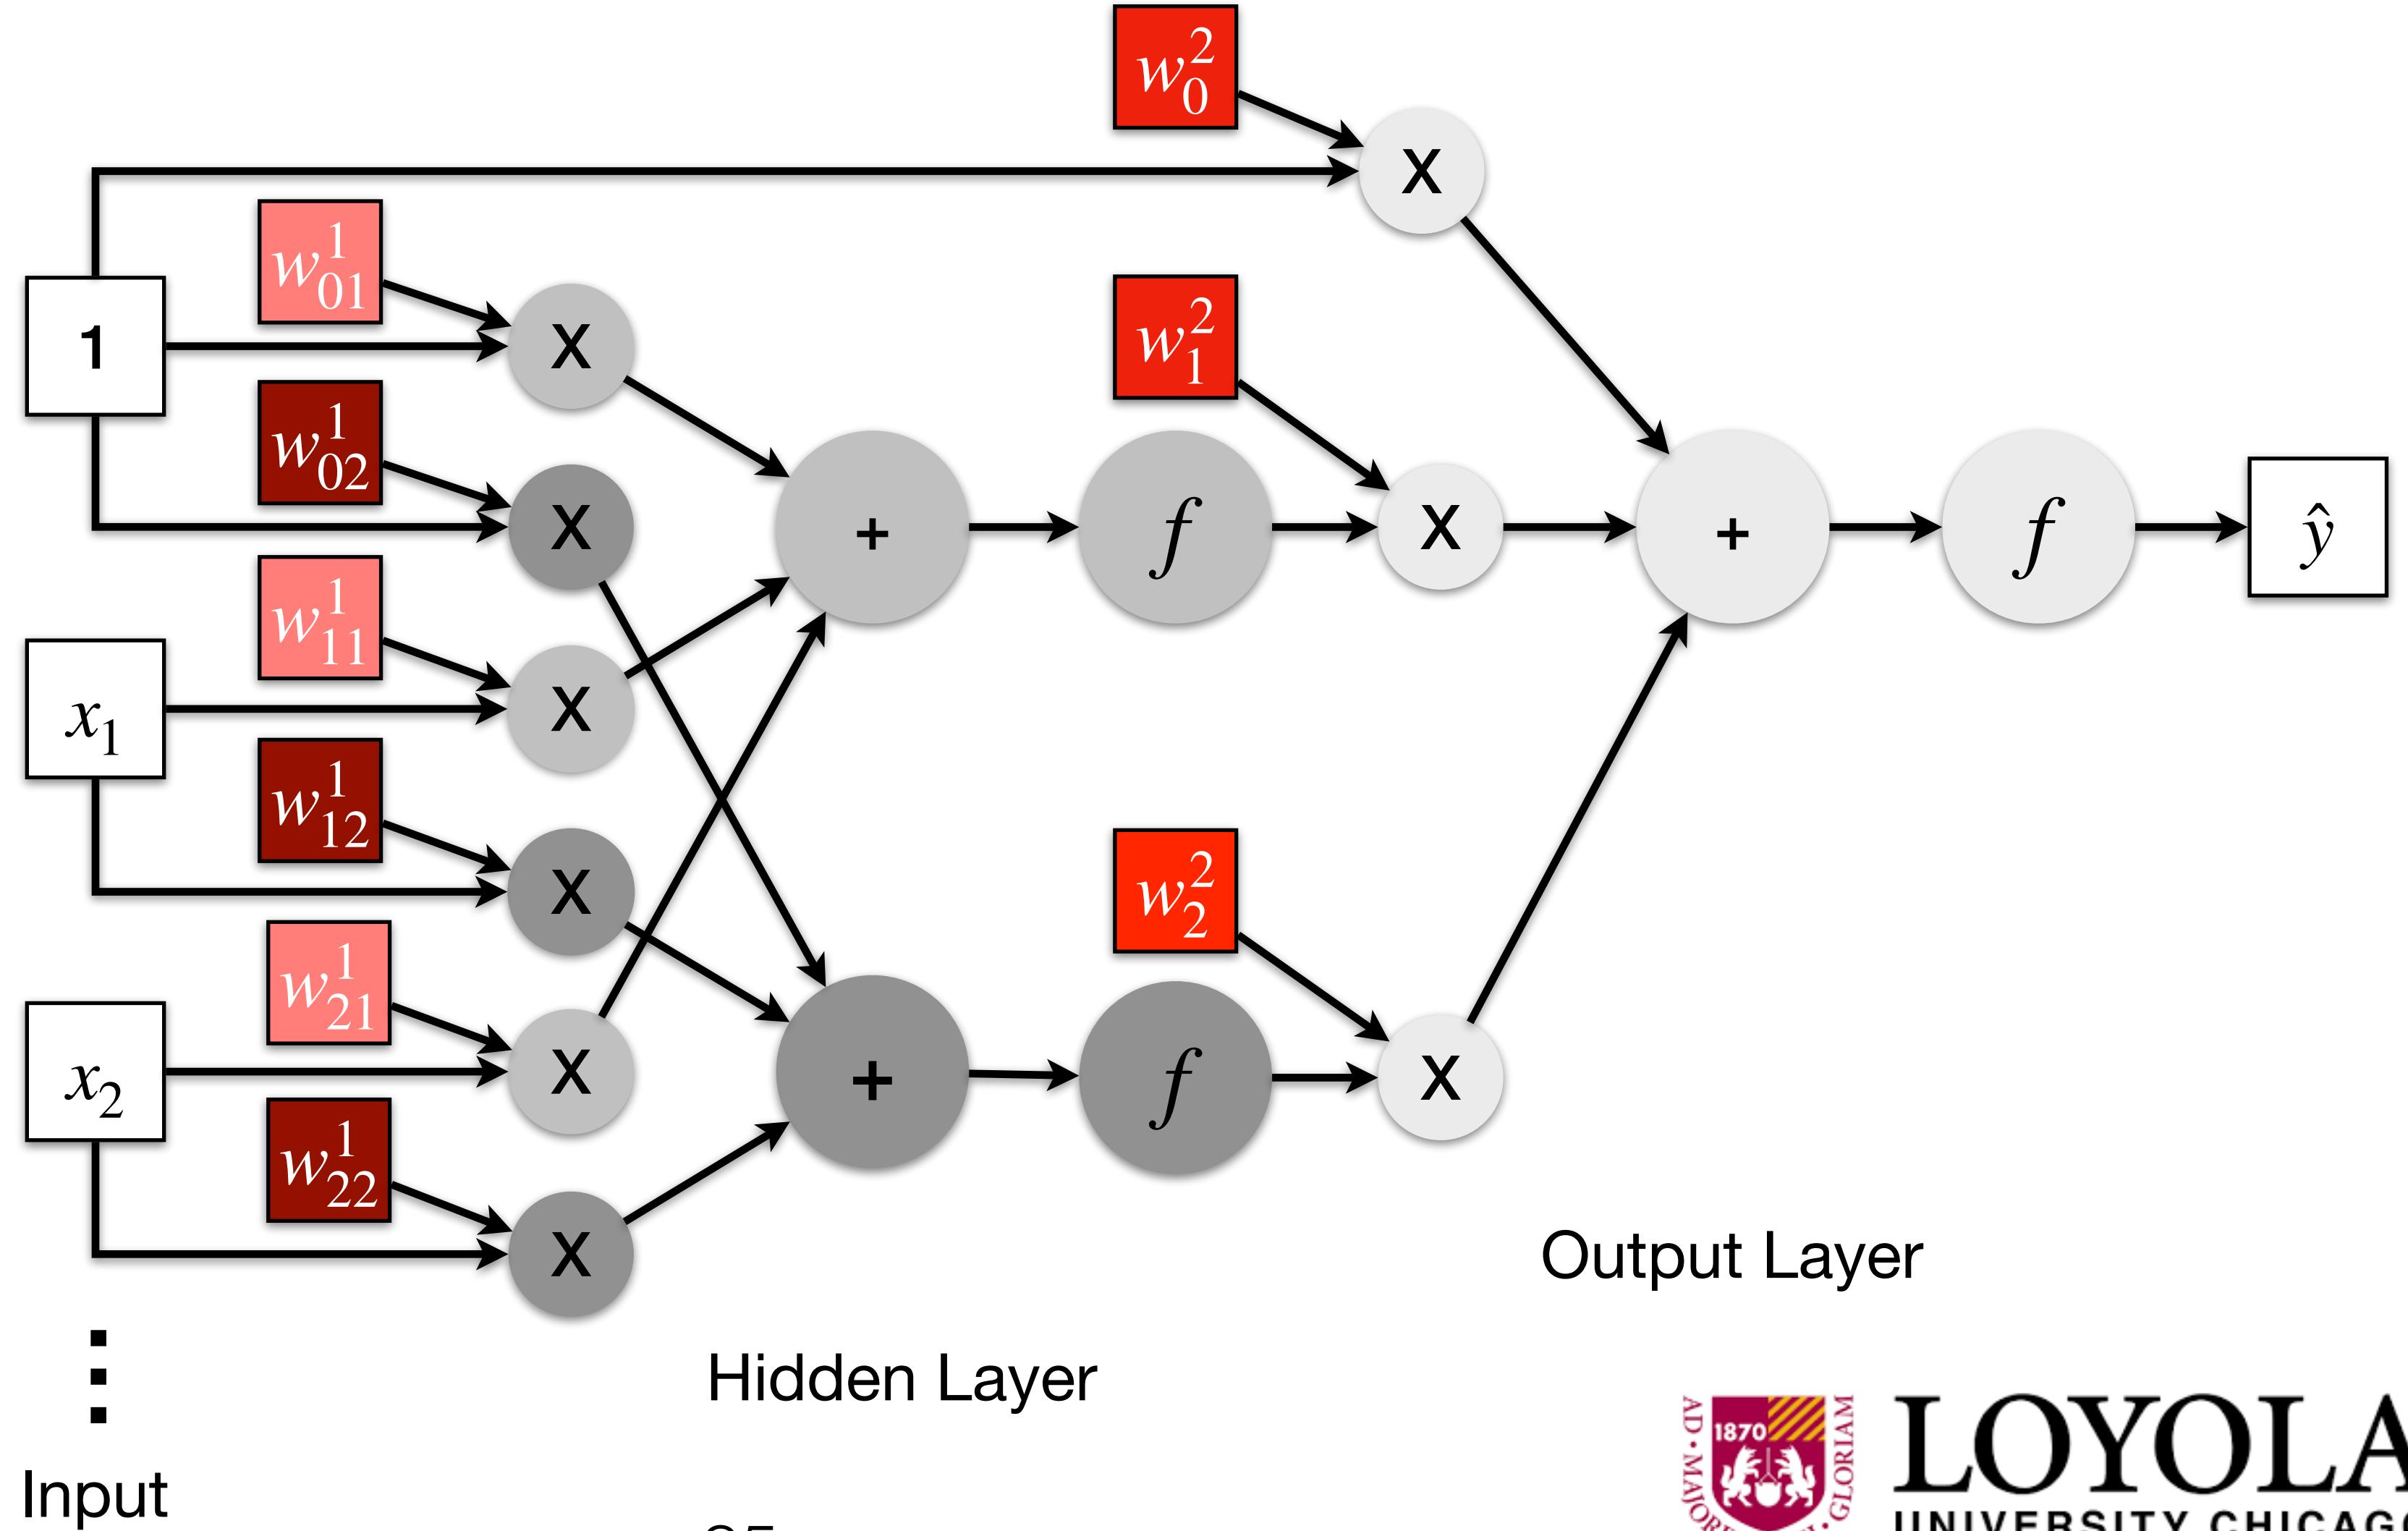


$$\hat{y} = f \left(\begin{pmatrix} w_1^2 & w_2^2 \end{pmatrix} f \left(\begin{pmatrix} w_{11}^1 & w_{21}^1 \\ w_{12}^1 & w_{22}^1 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} w_{01}^1 \\ x_{02}^1 \end{pmatrix} \right) + w_0^2 \right)$$

Neural Networks (NN)

**Adding
Hidden Layers**

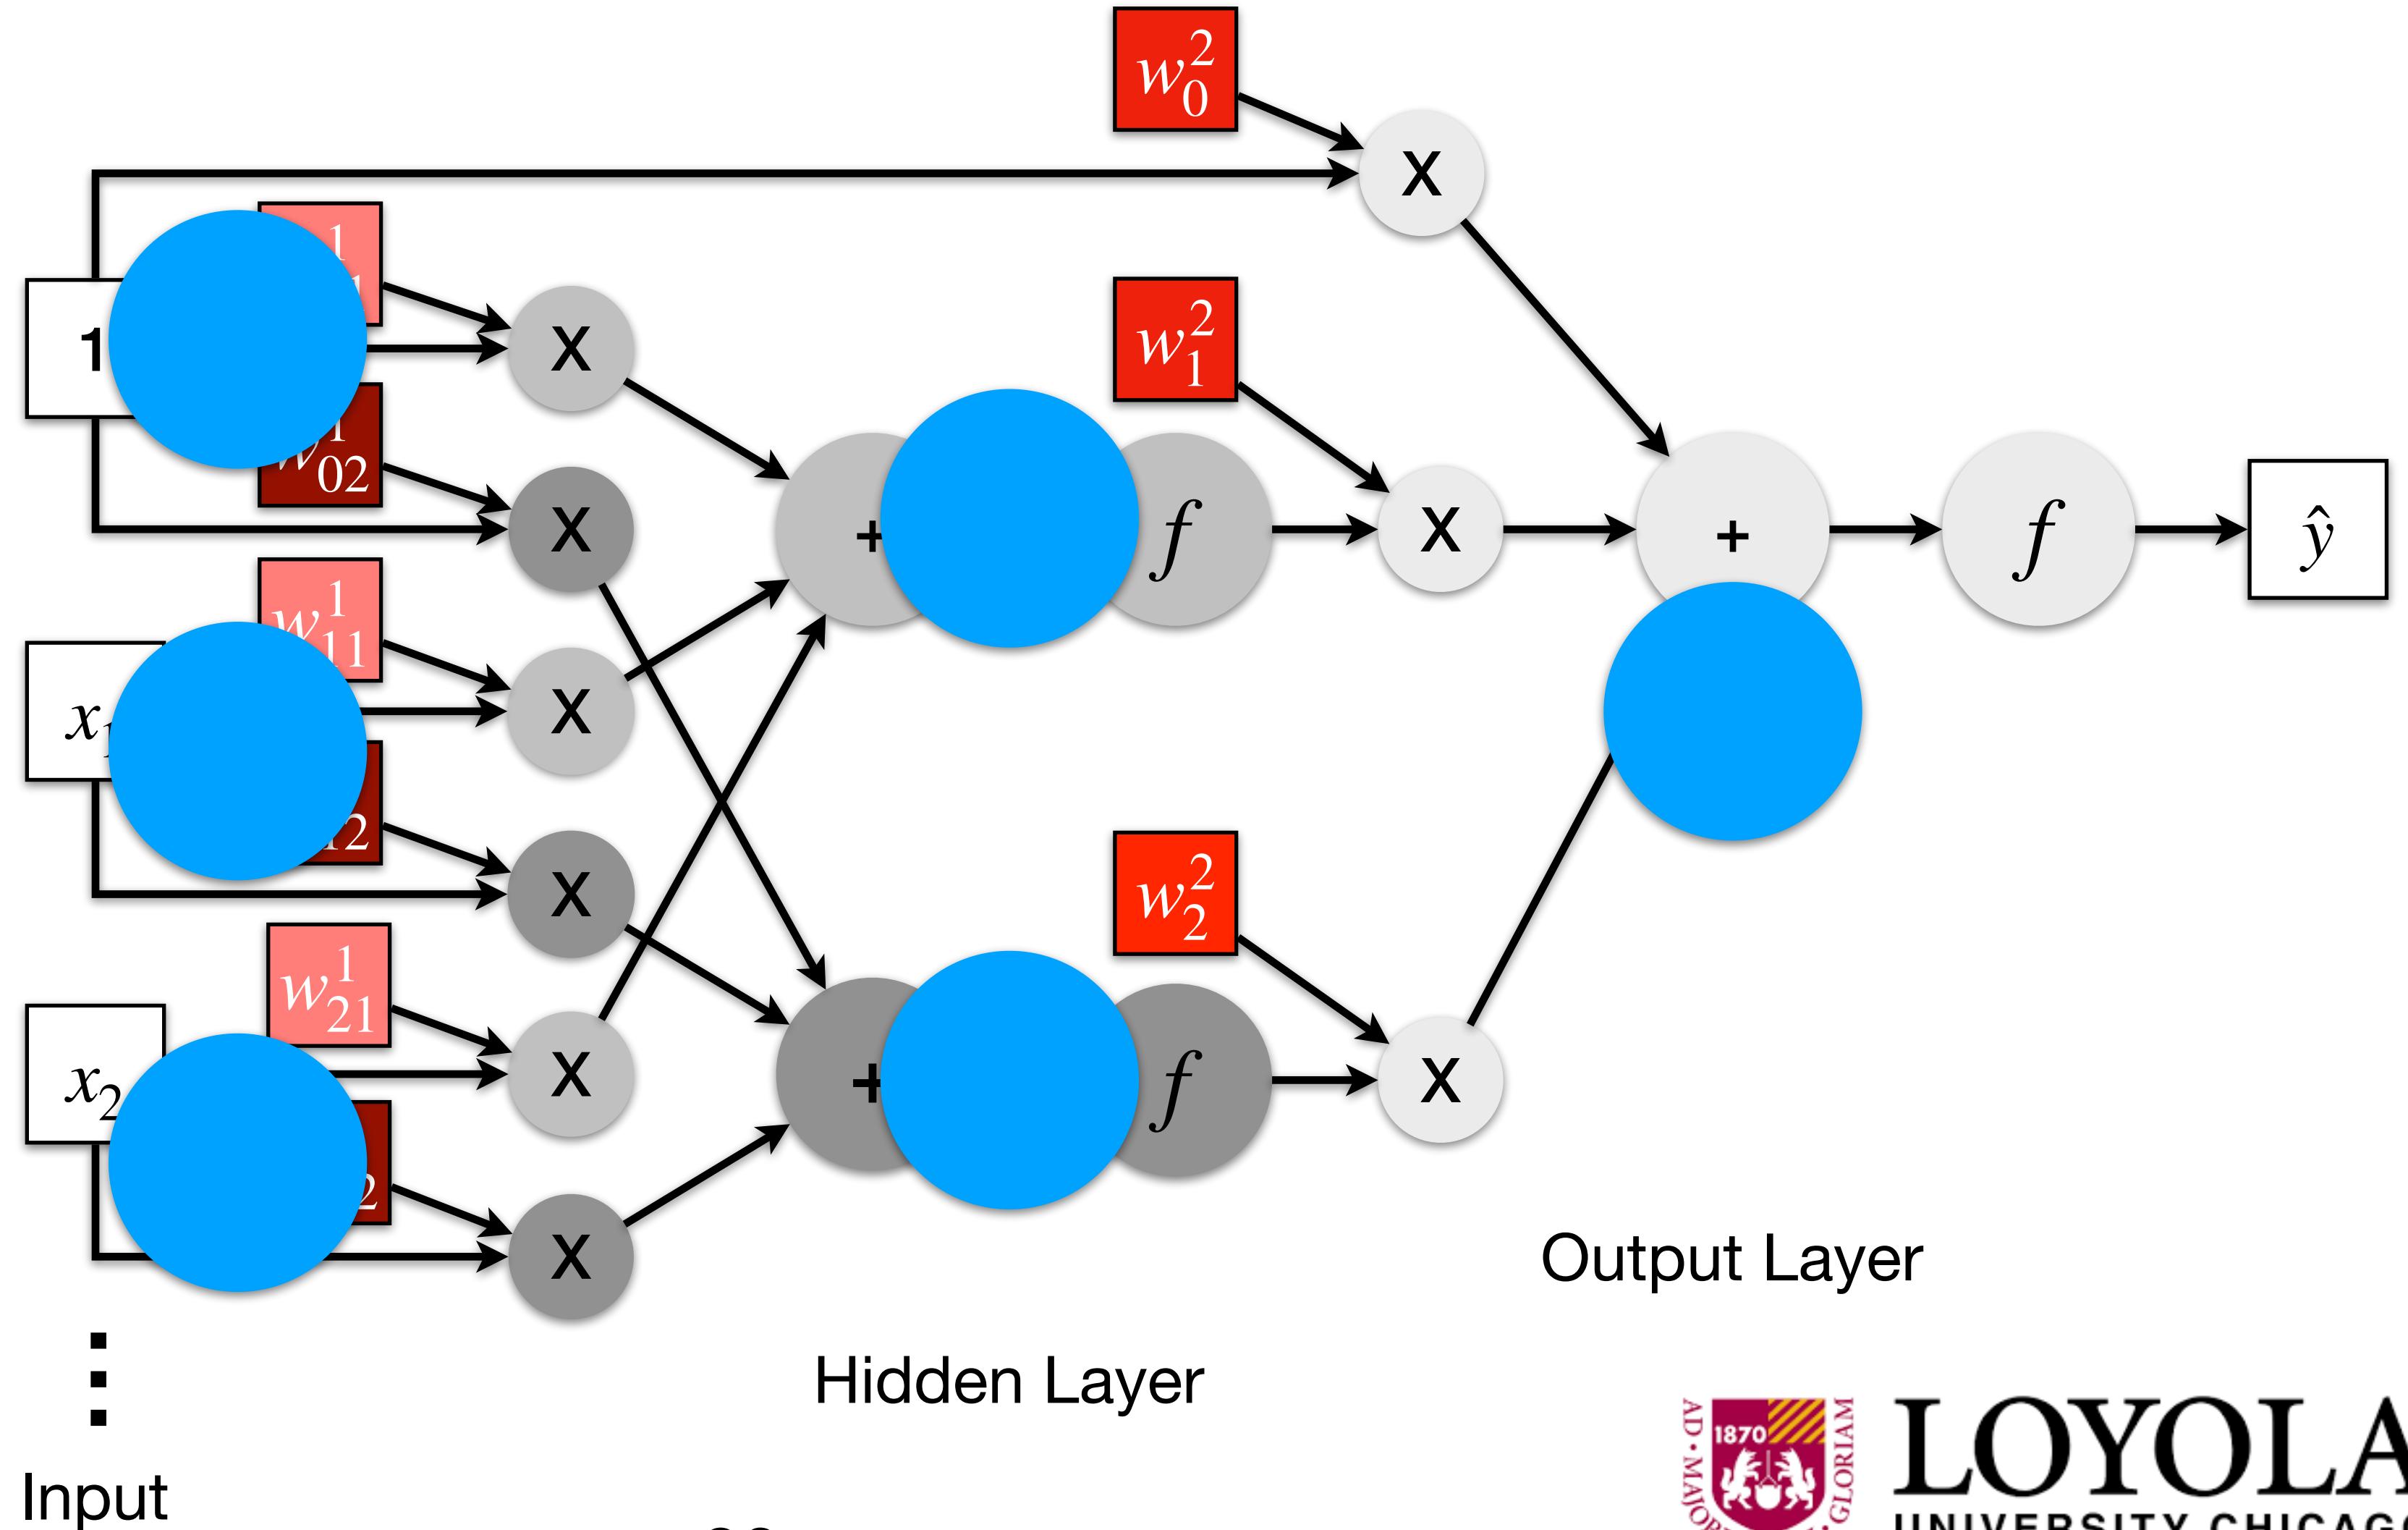
To increase
computing
power.



Neural Networks (NN)

**Adding
Hidden Layers**

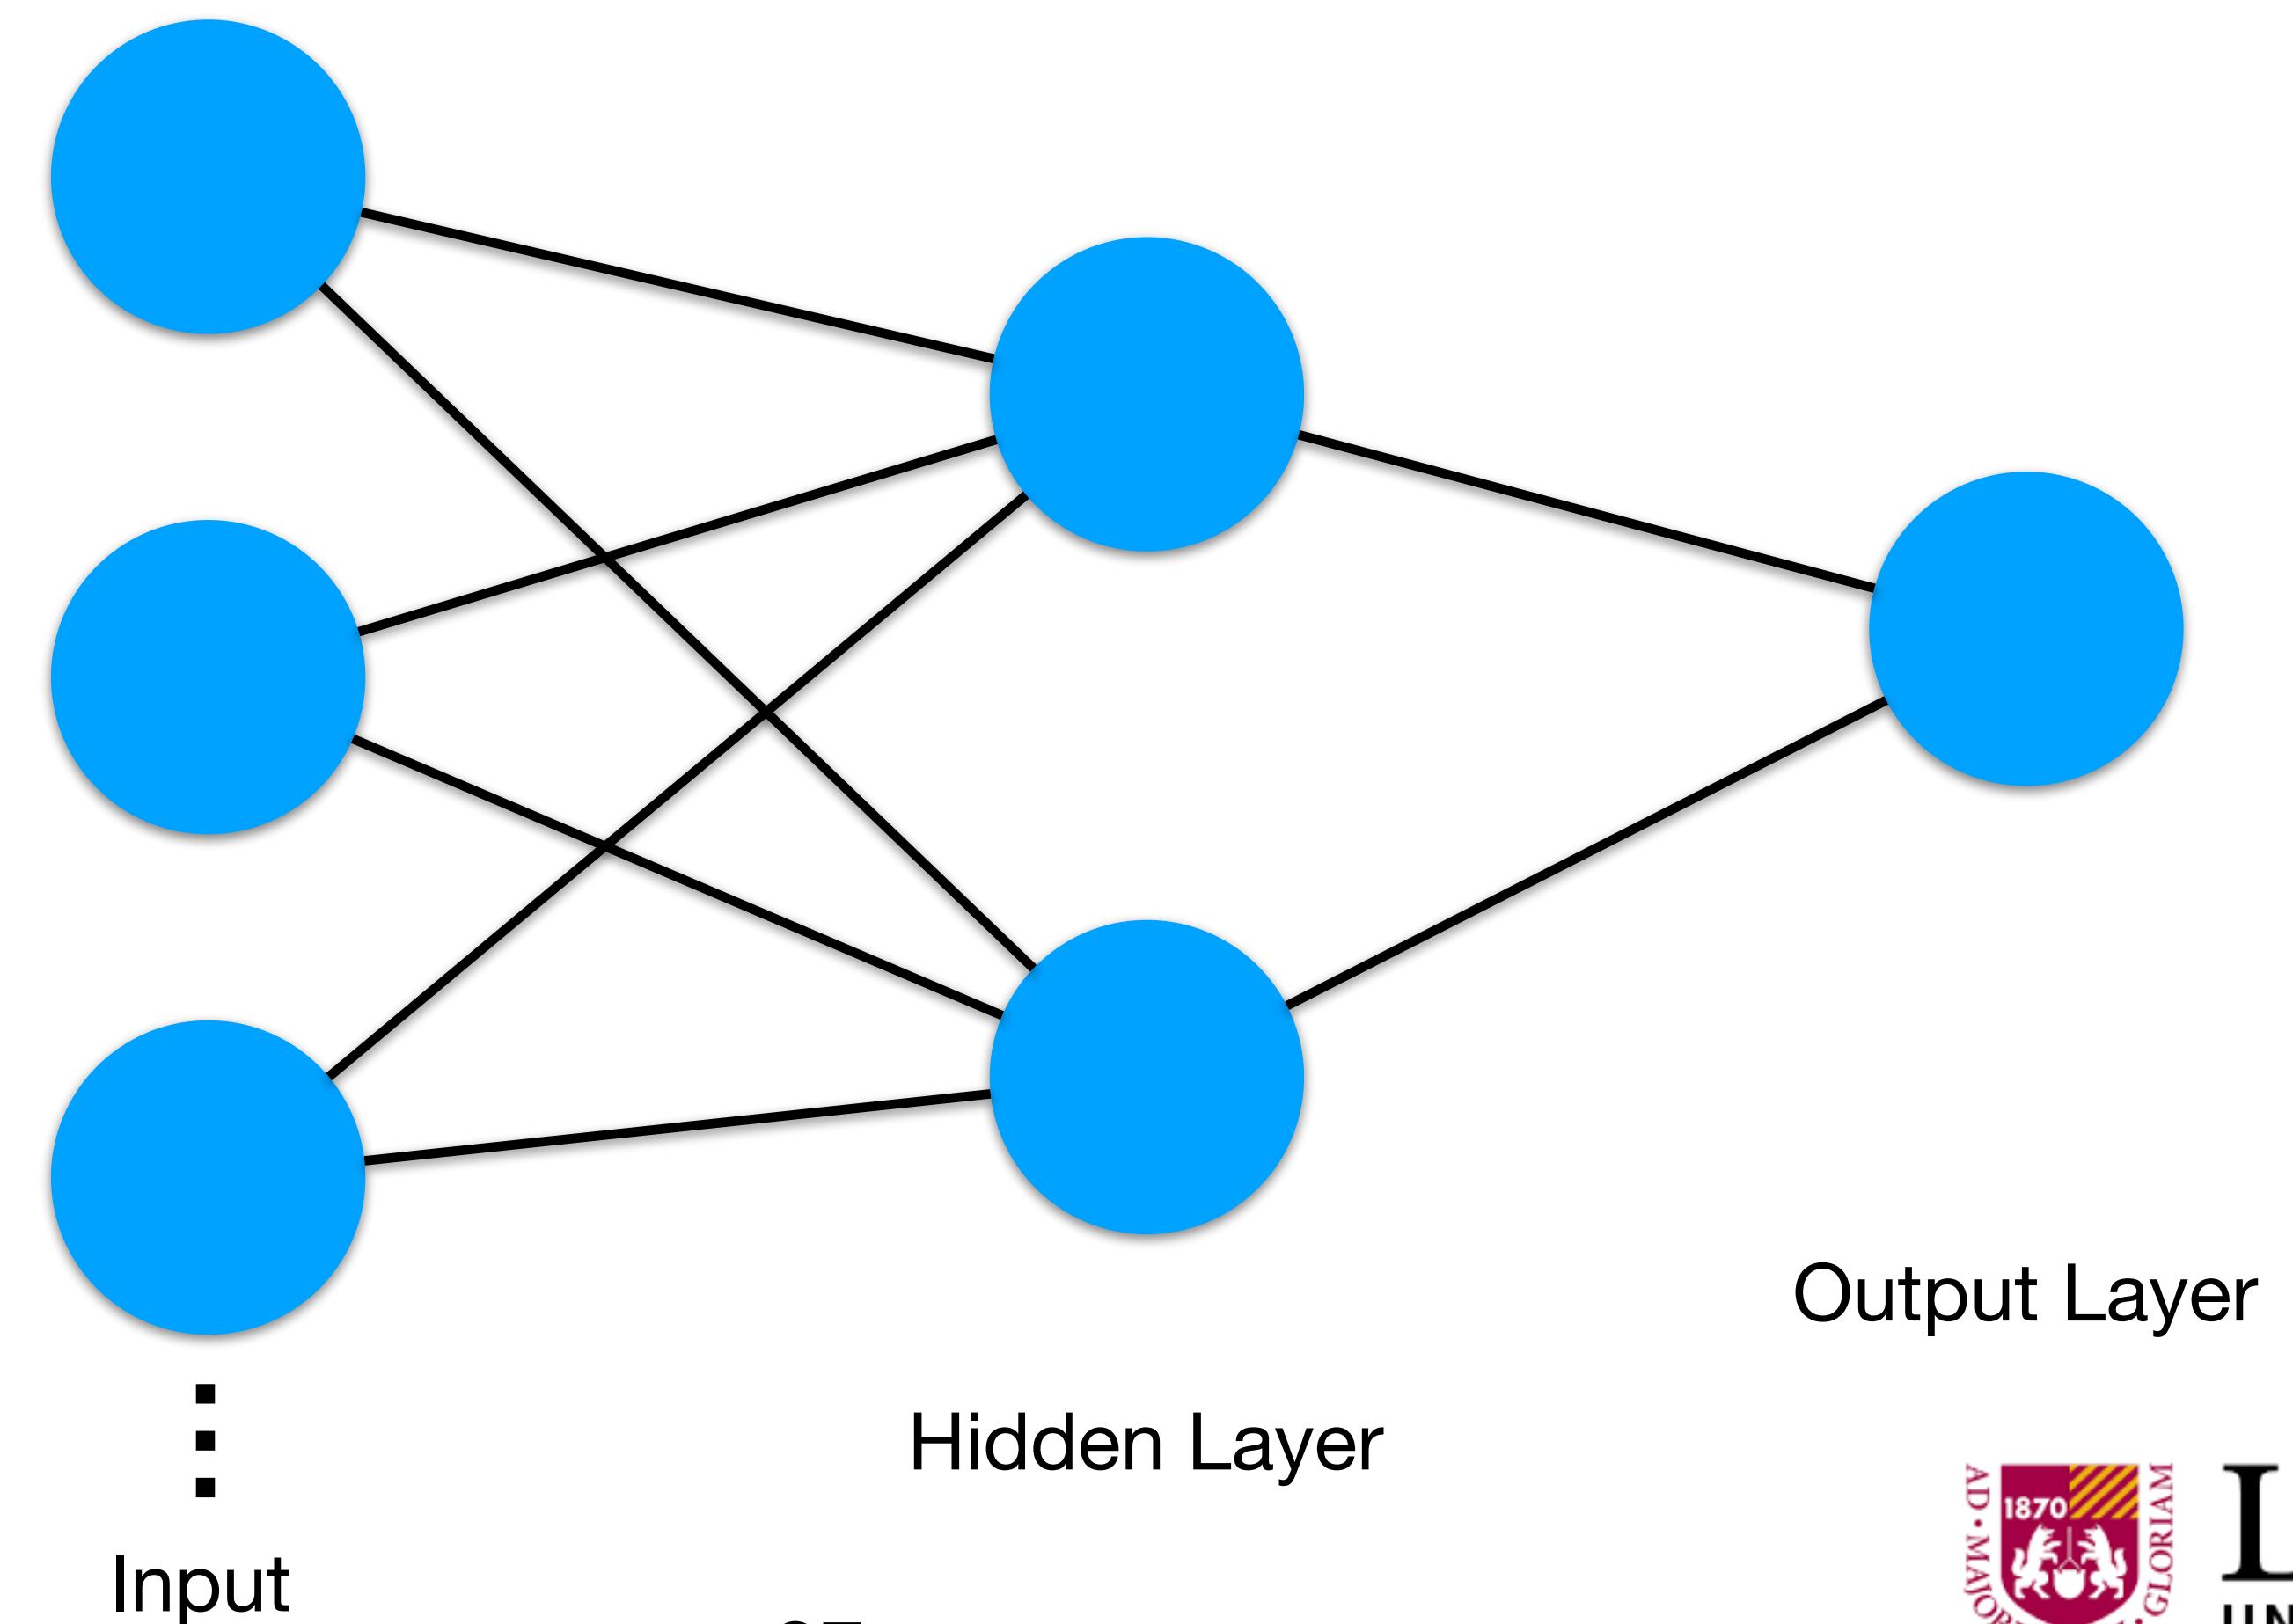
To increase
computing
power.



Neural Networks (NN)

Adding Hidden Layers

To increase
computing
power.

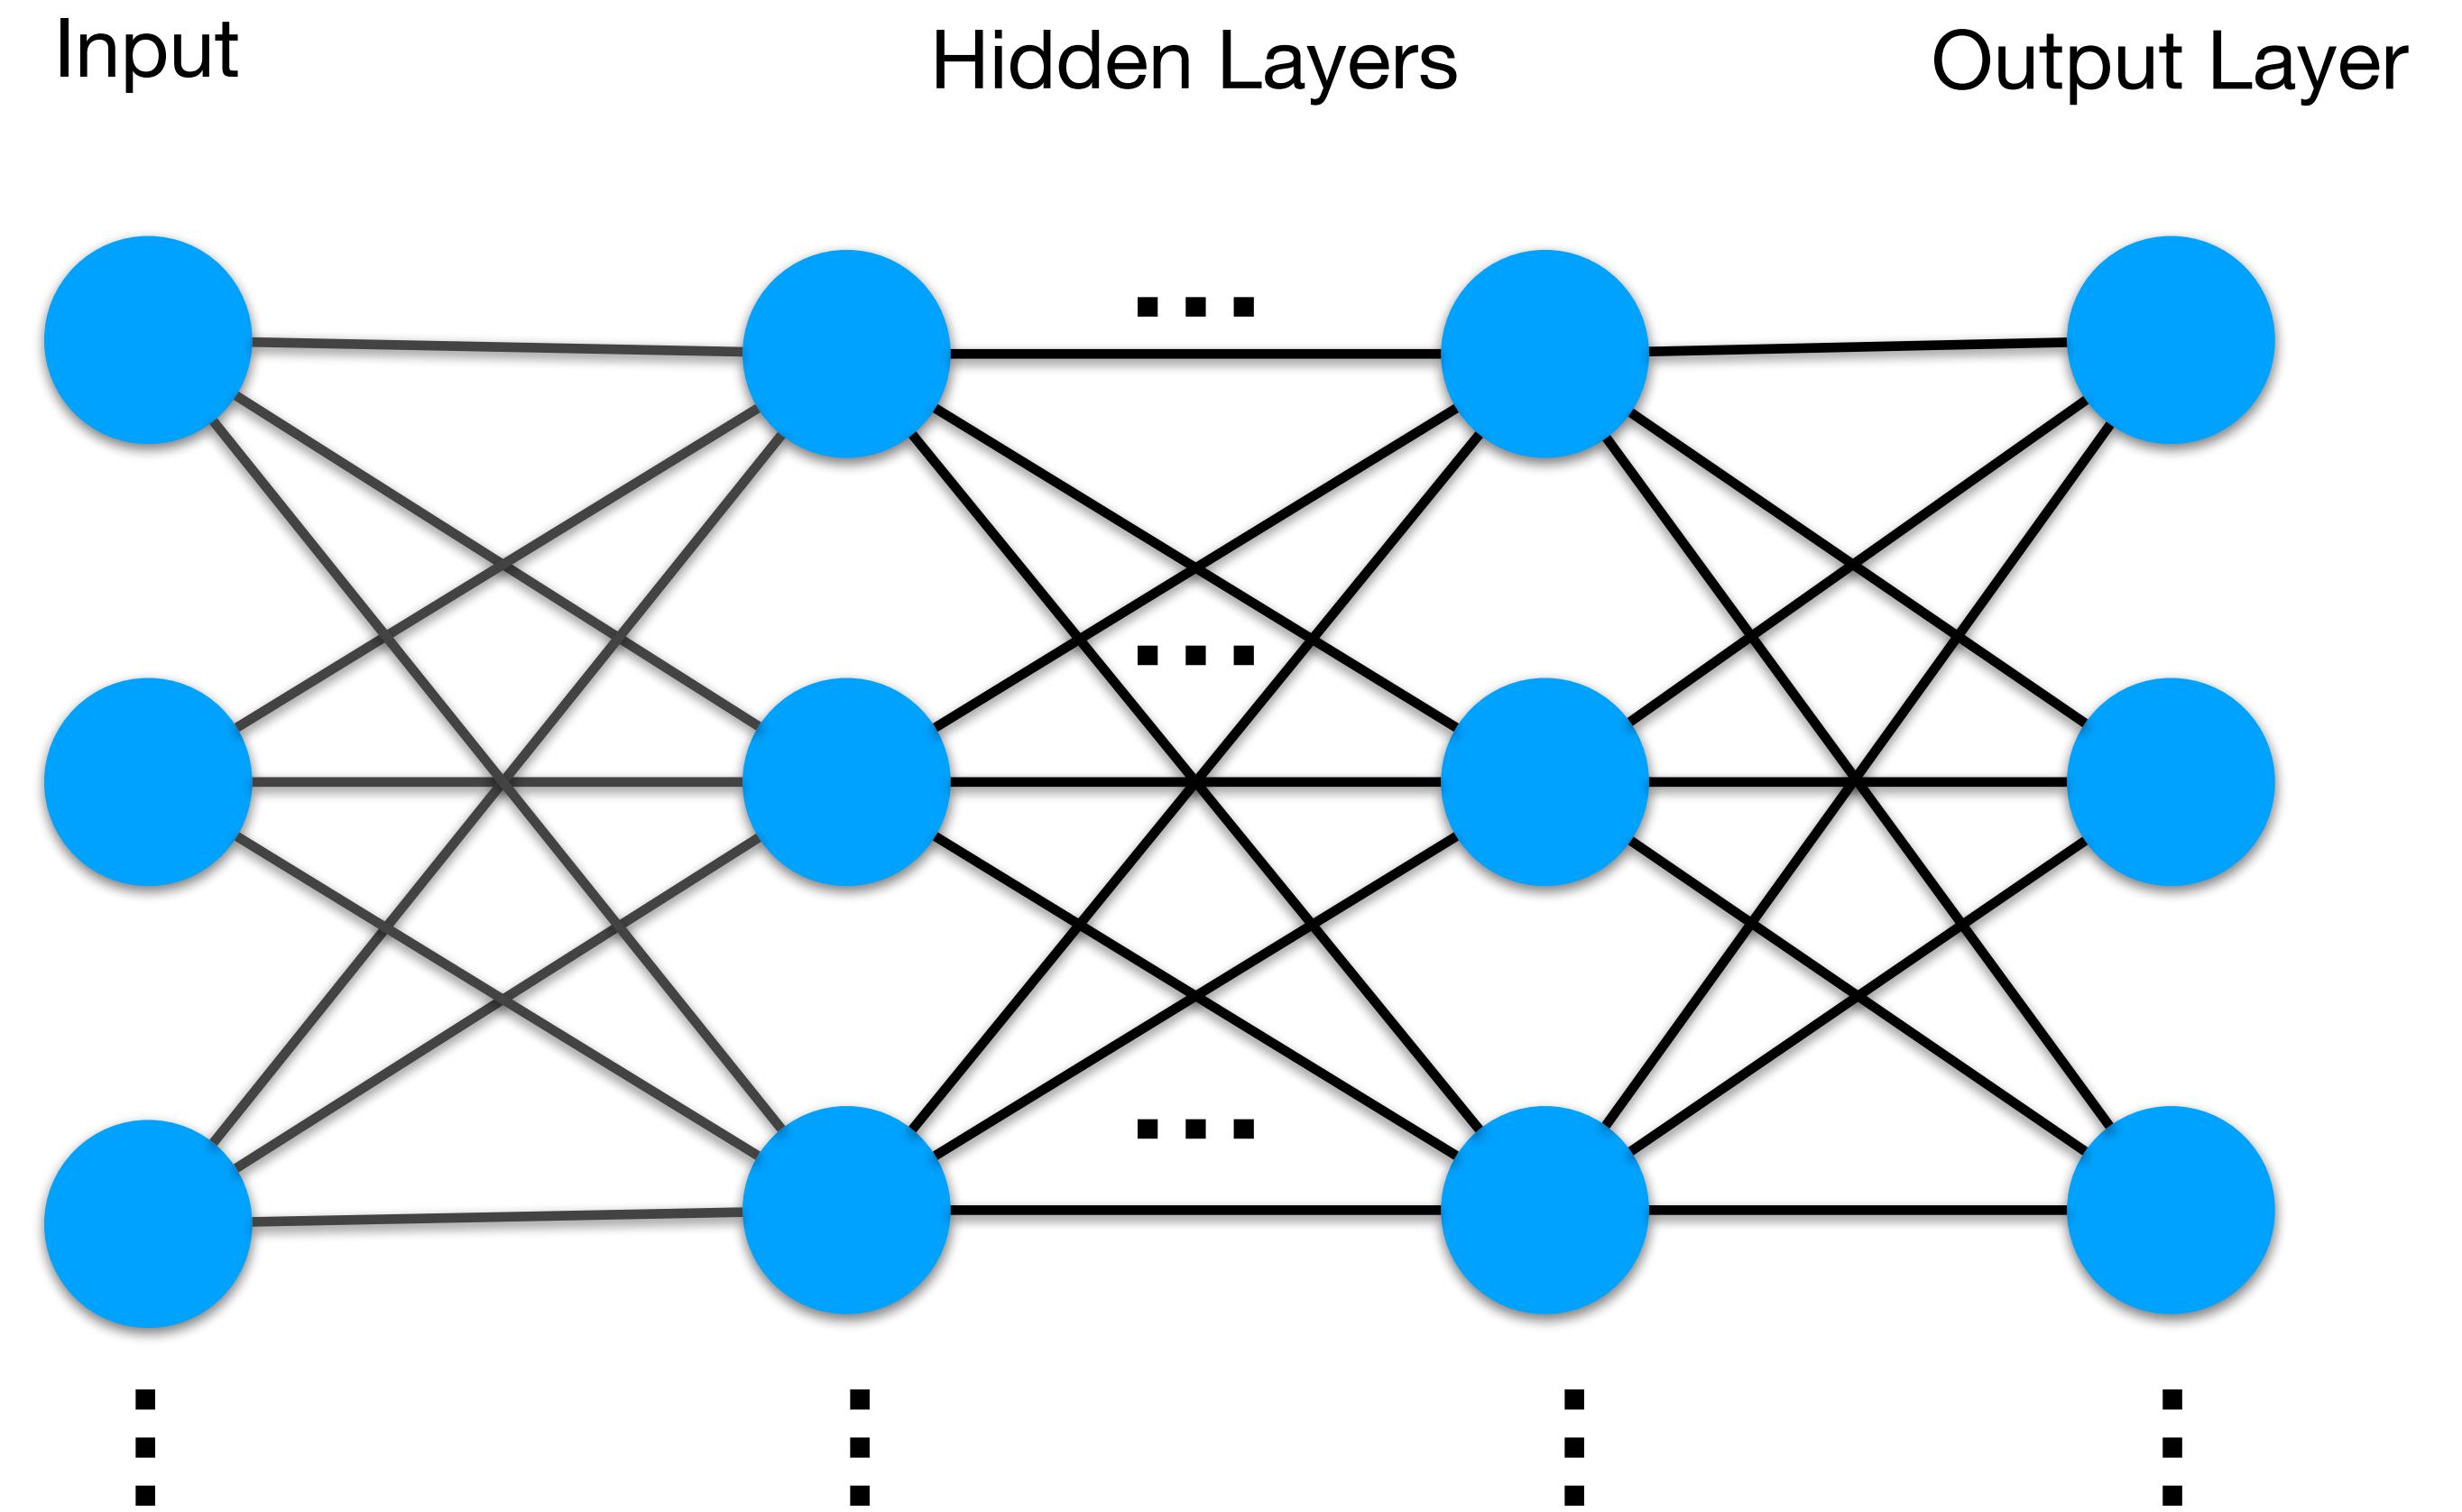


Neural Networks (NN)

Universal Approximation Theorem

An NN with one or more hidden layers with compressive (non-linear) activation function can approximate any continuous function.

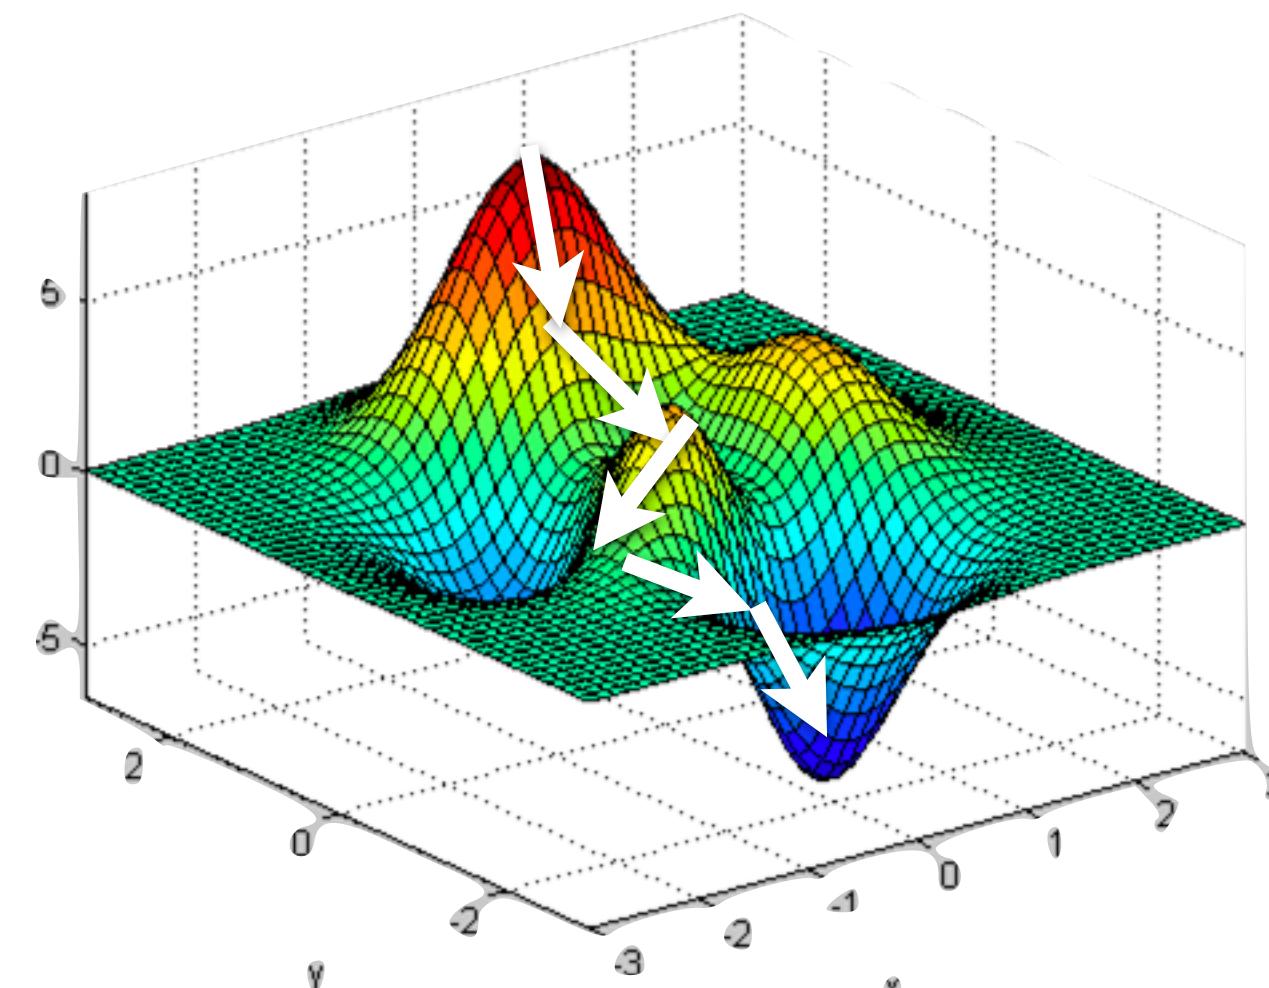
Training this NN might be a nightmare!



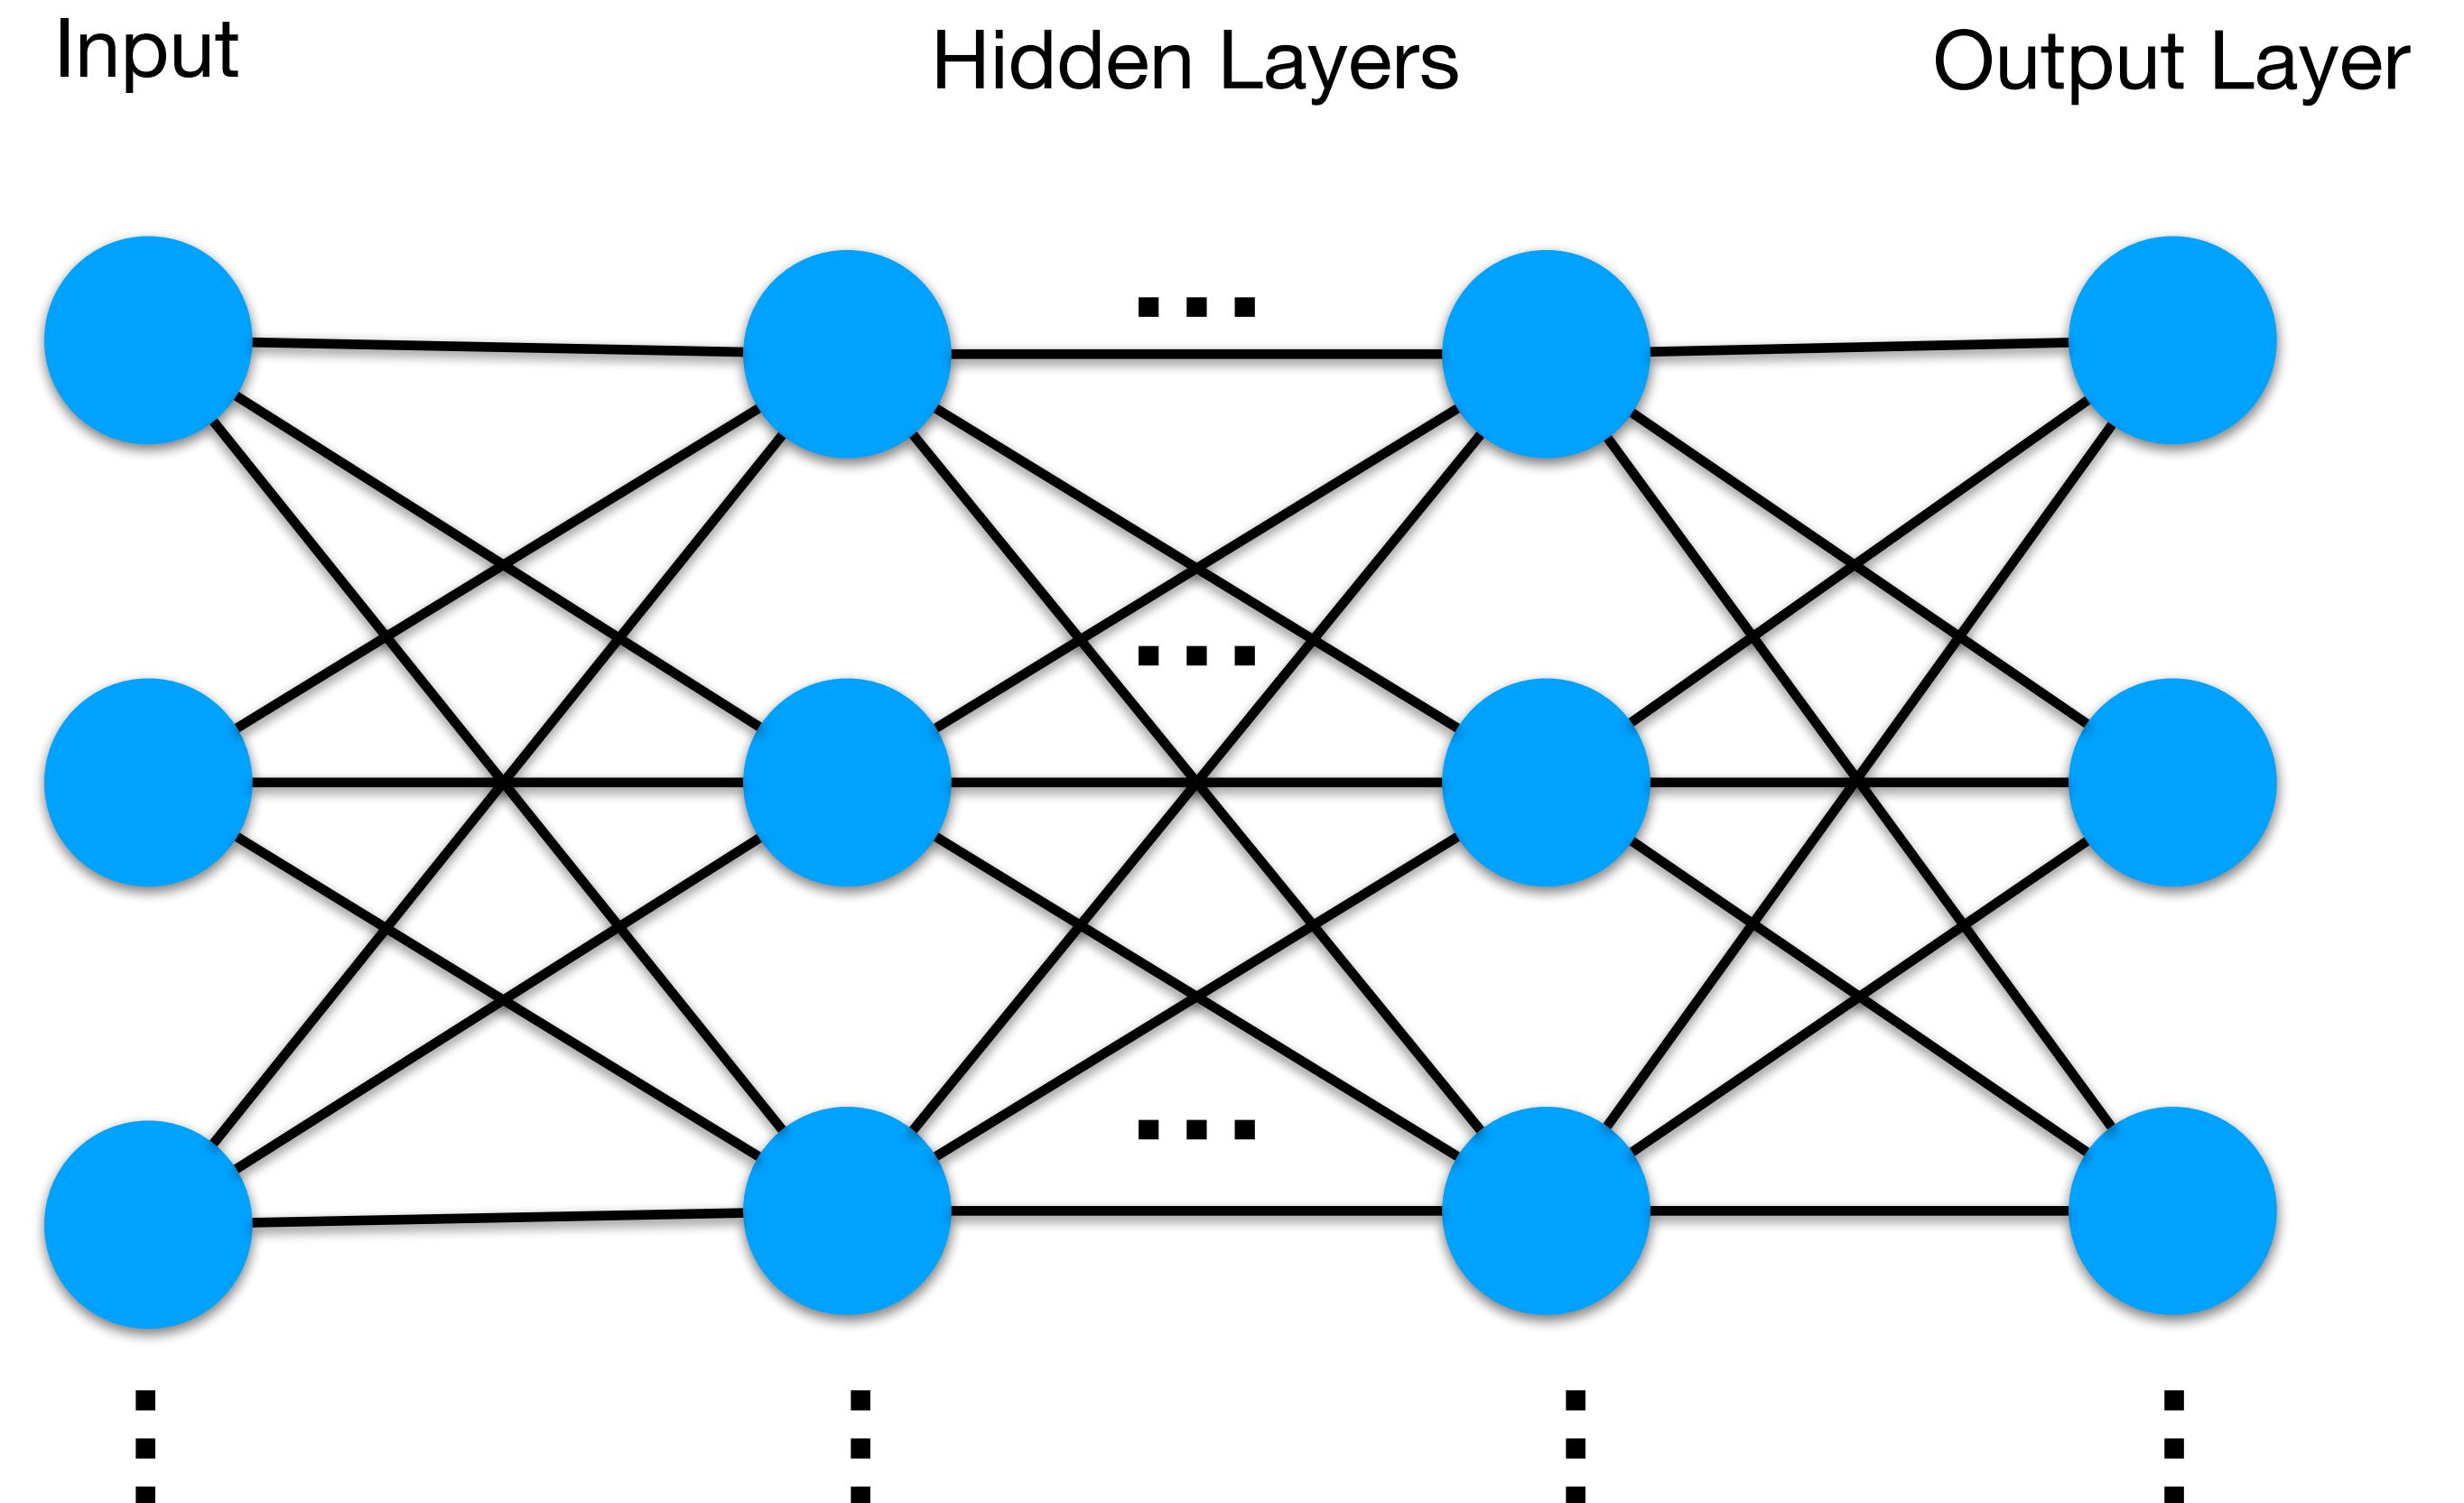
Neural Networks (NN)

Backpropagation

Key algorithm to train the NN.



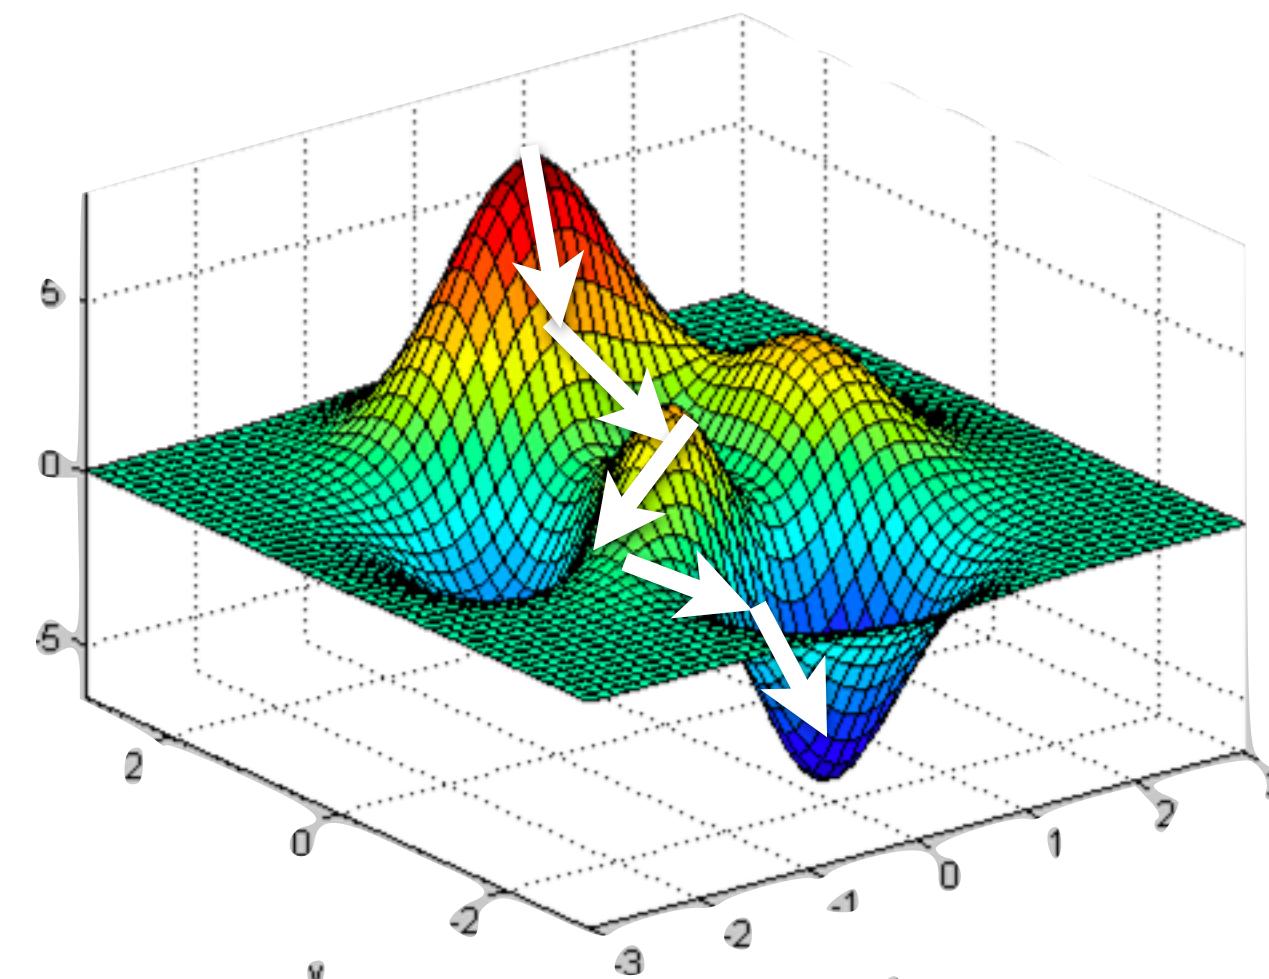
Loss Surface
Gradient Descent



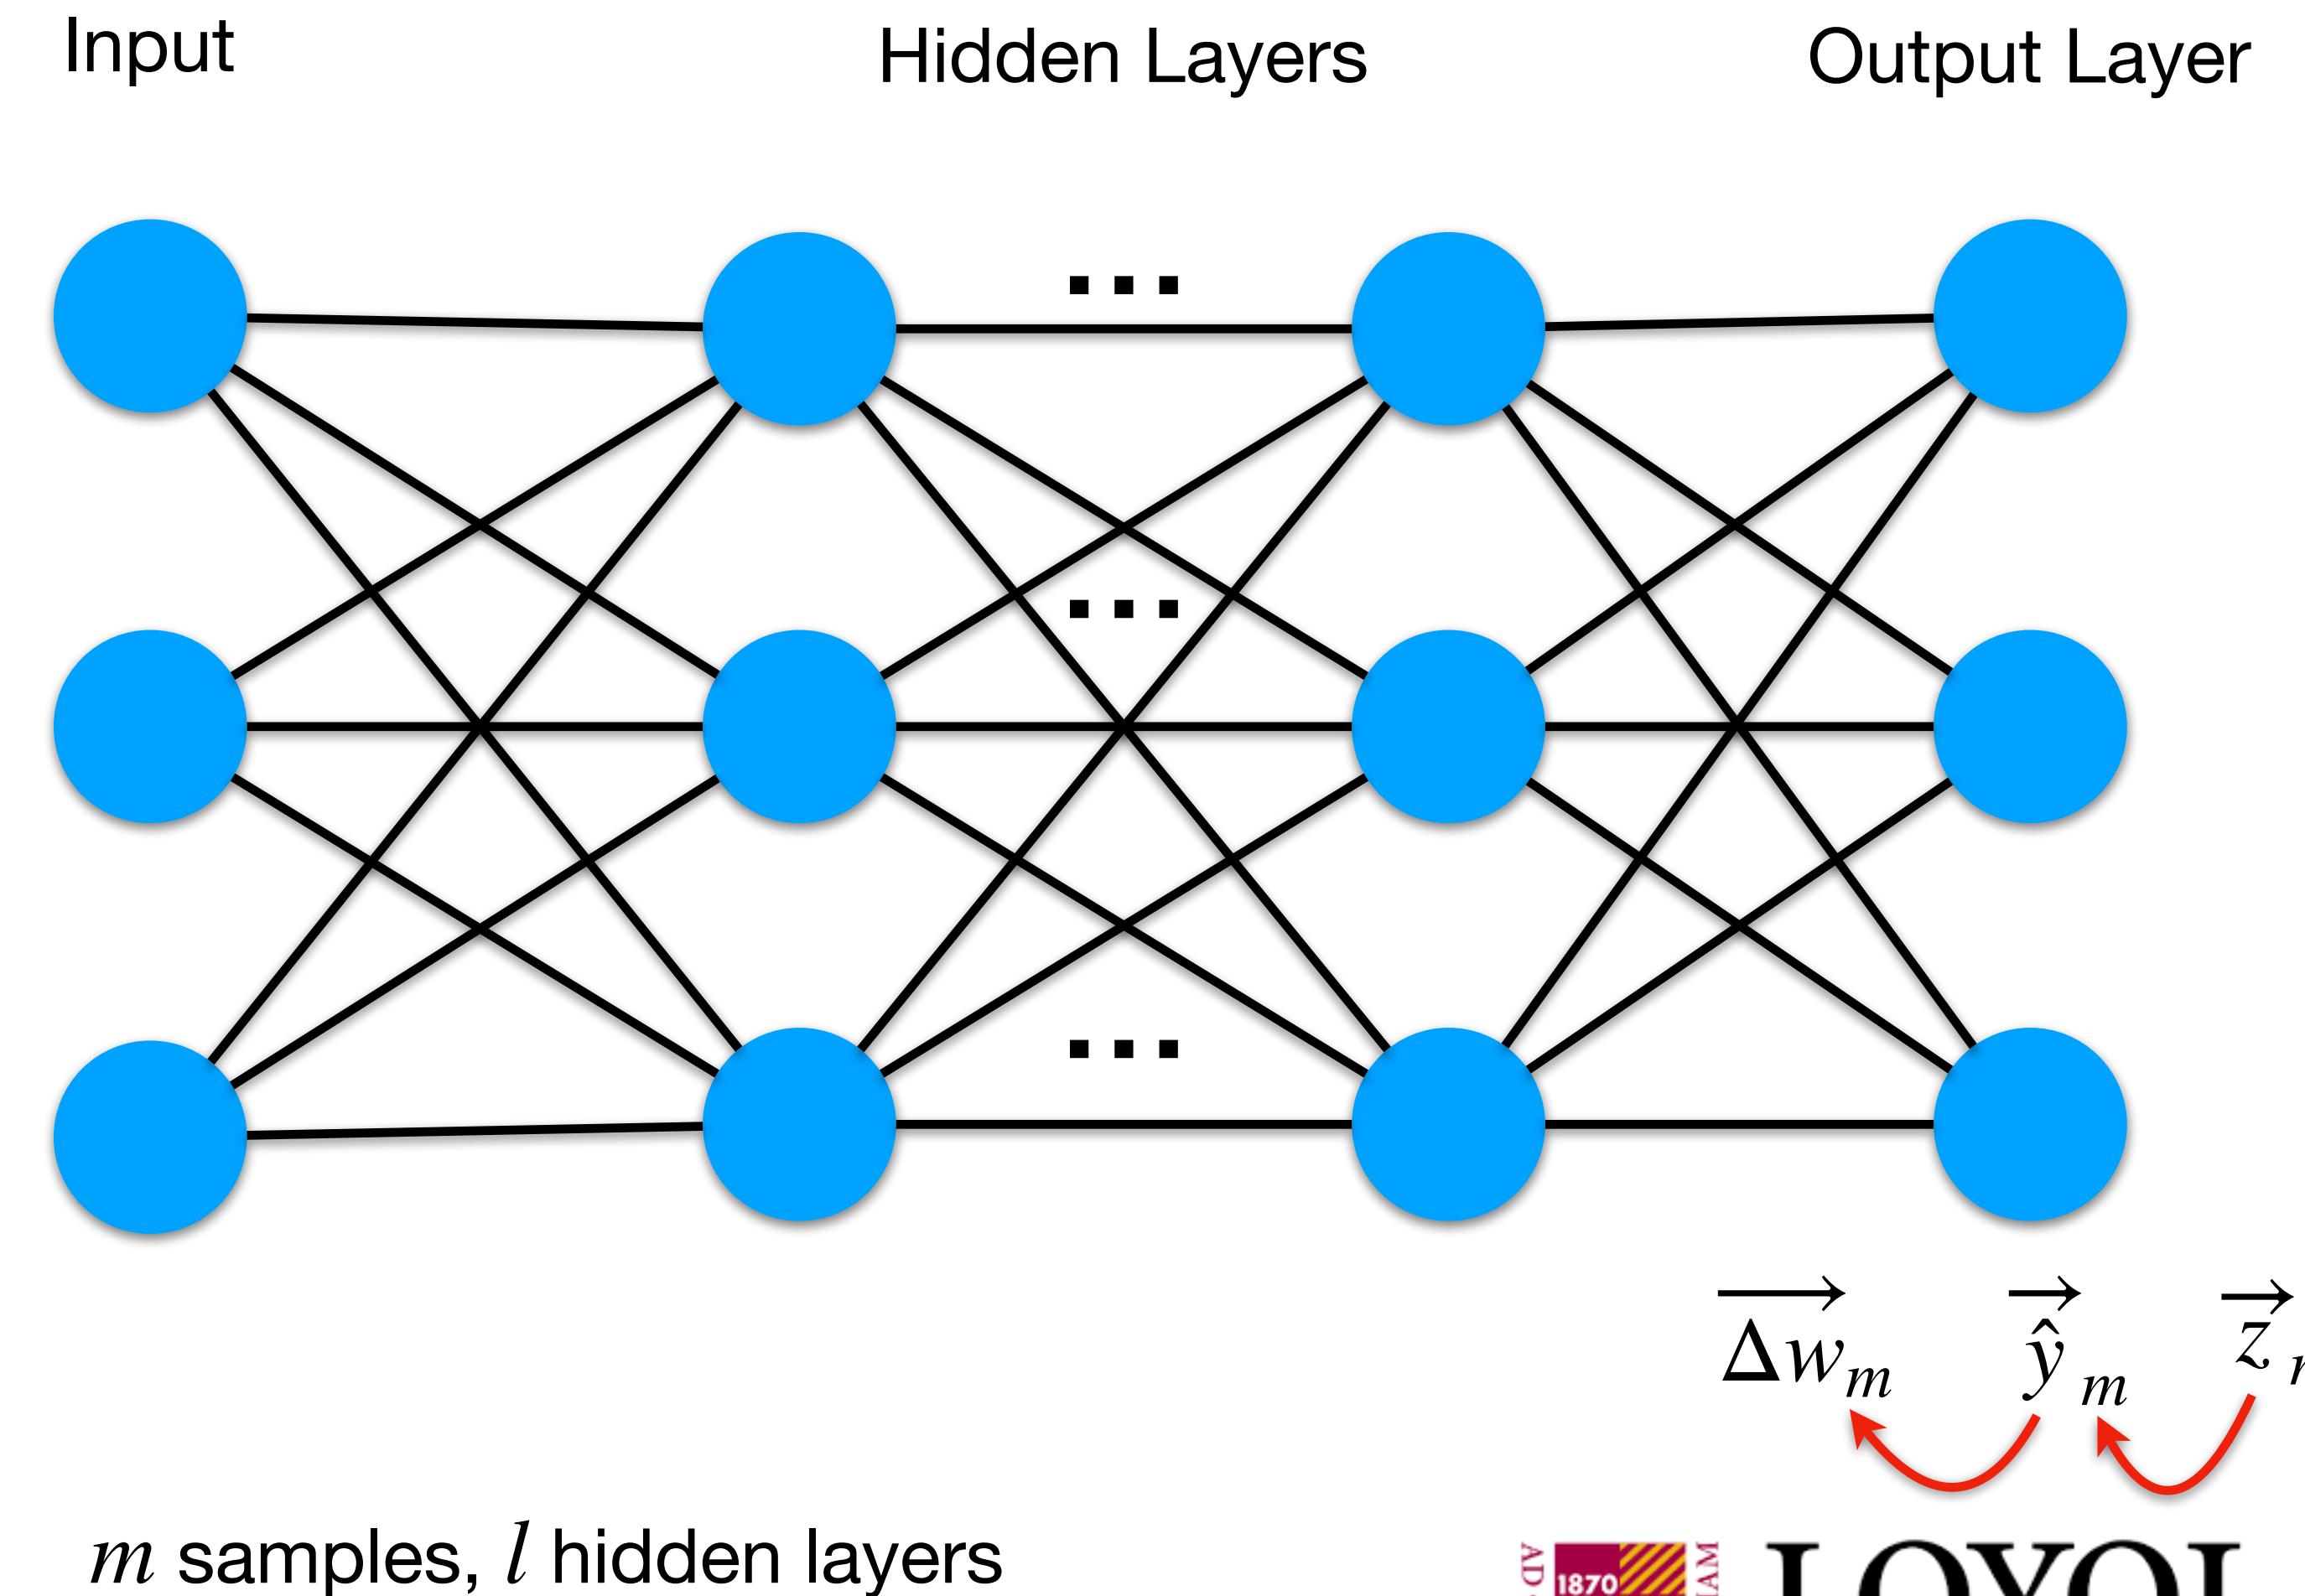
Neural Networks (NN)

Backpropagation

Key algorithm to train the NN.



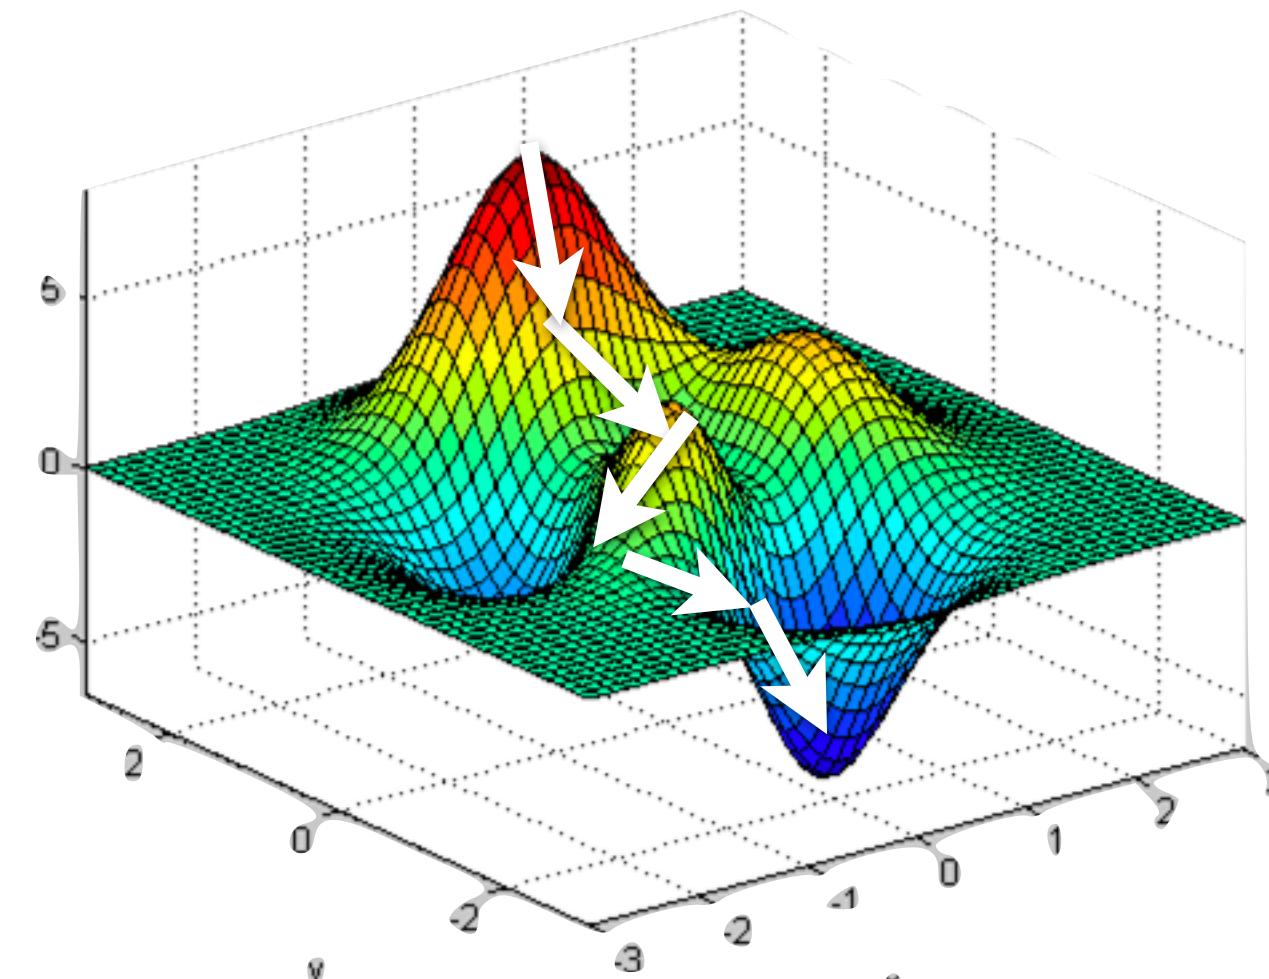
Loss Surface
Gradient Descent



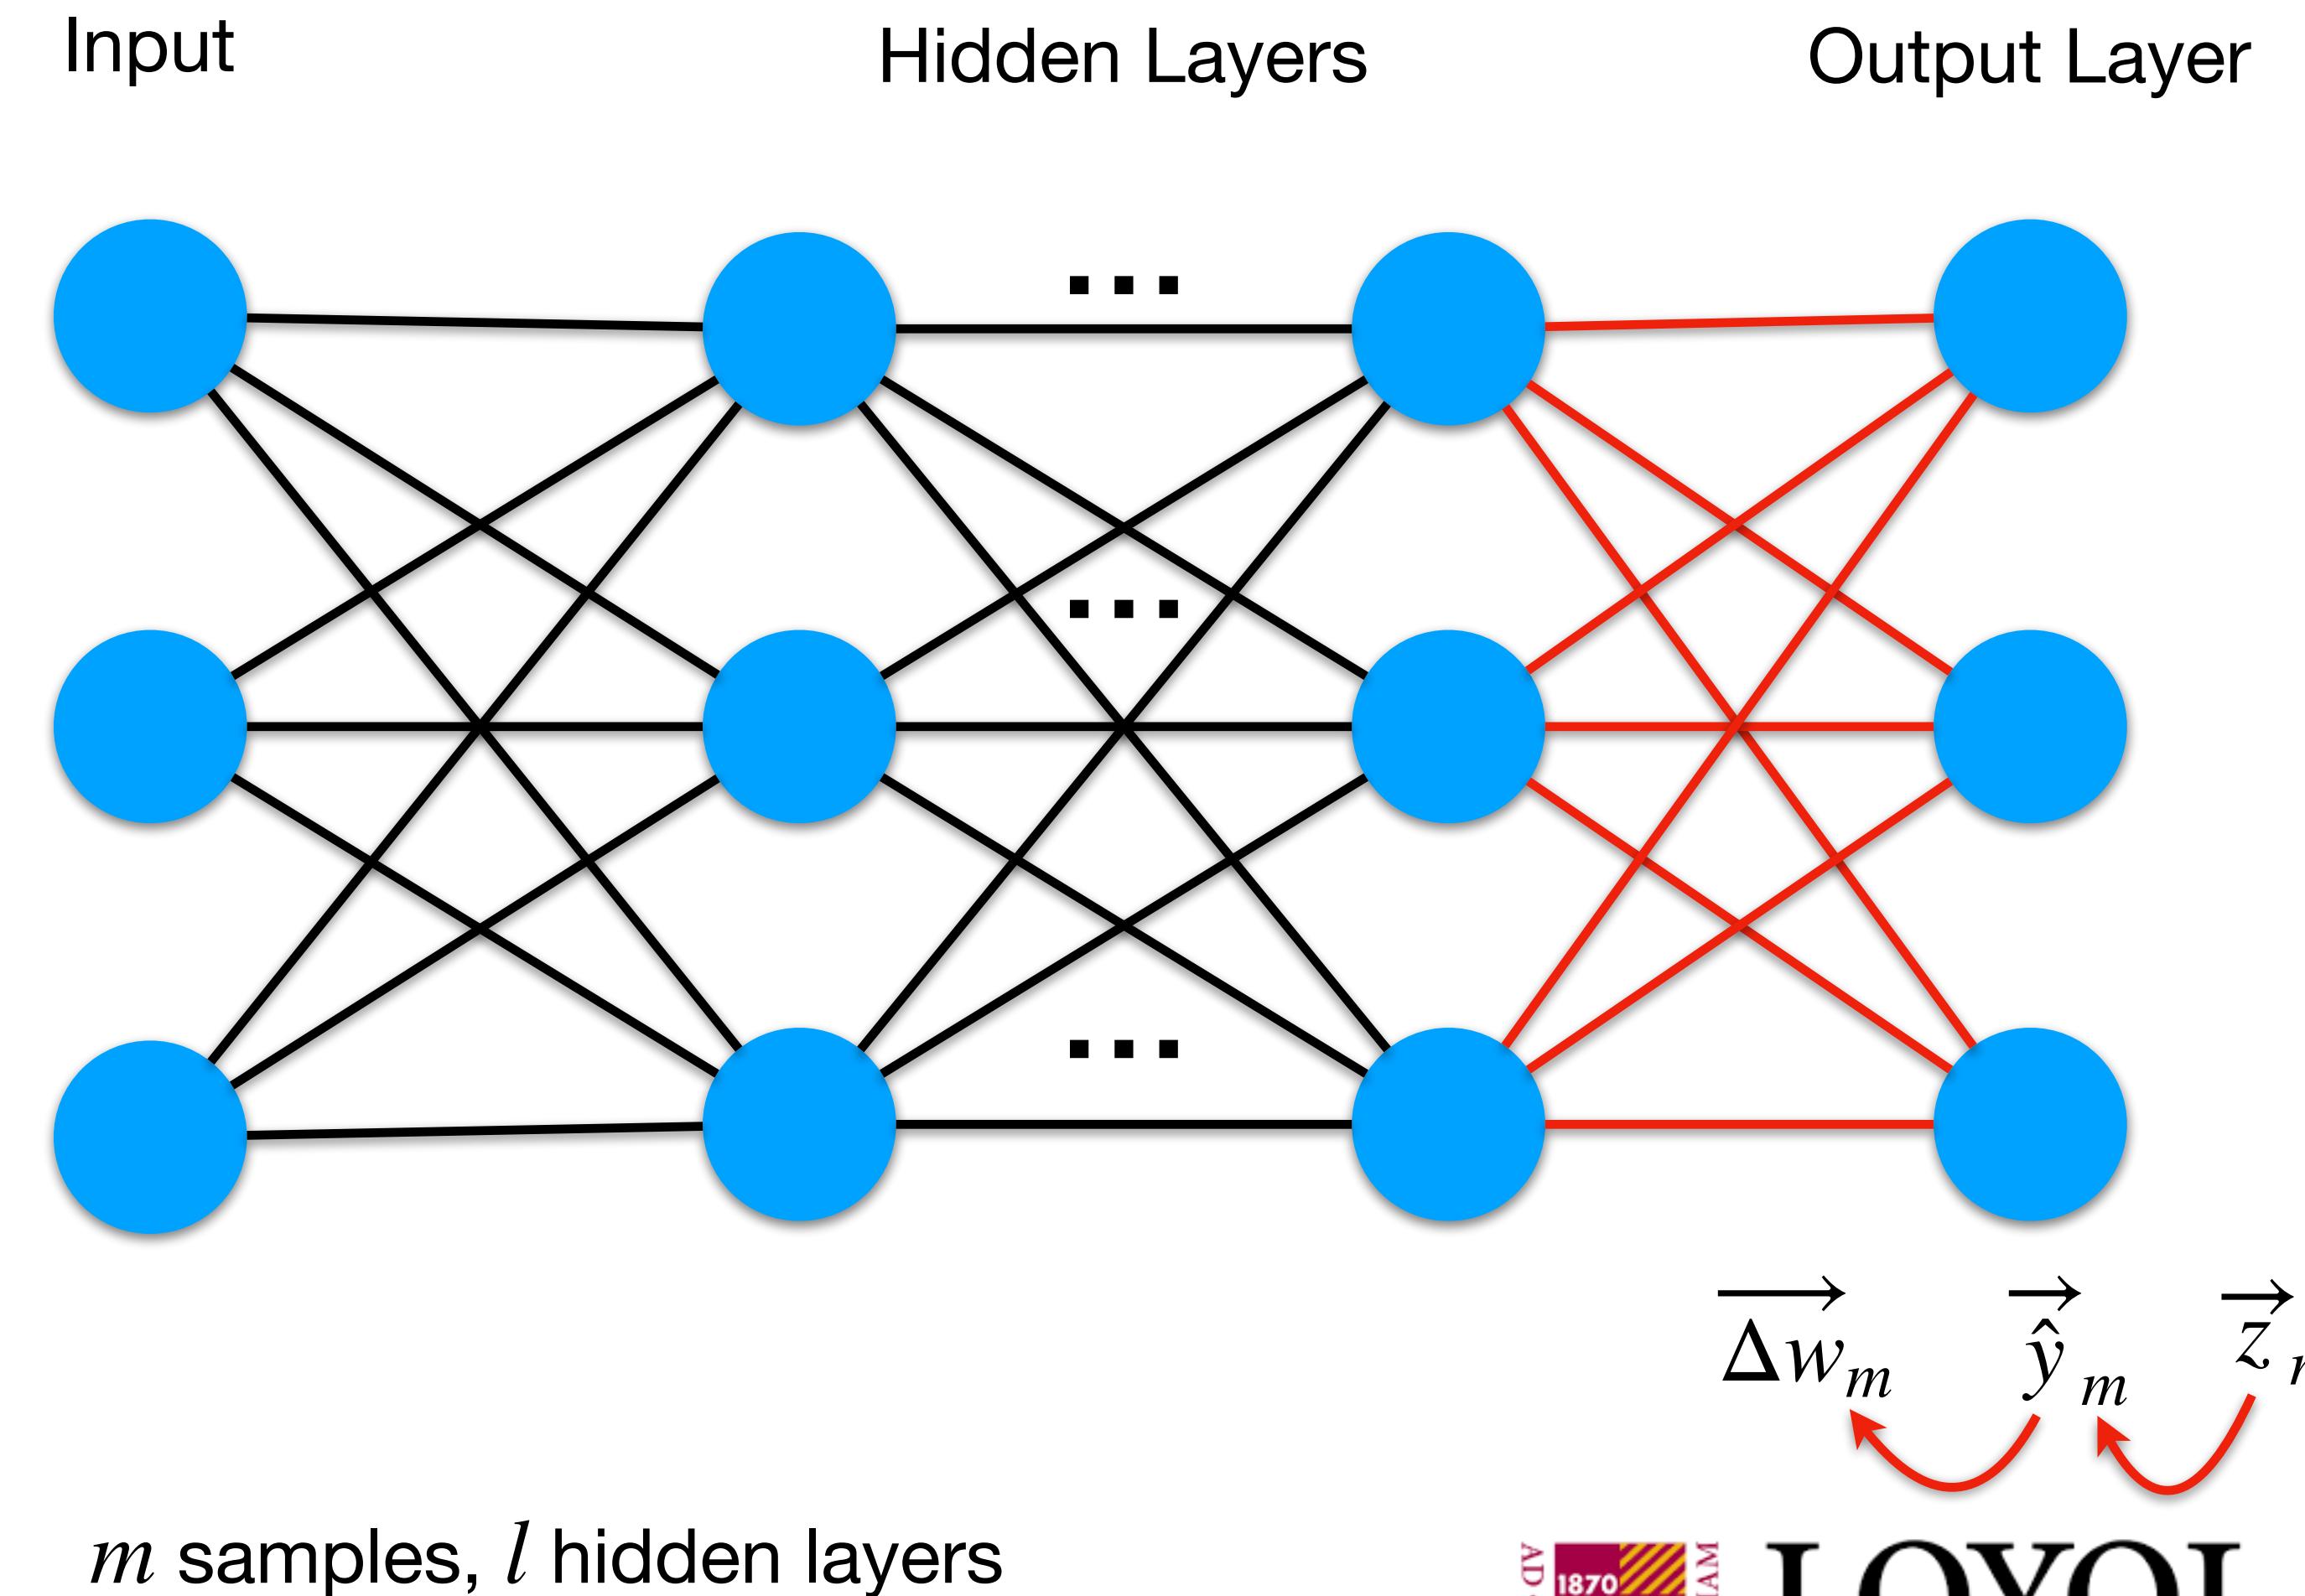
Neural Networks (NN)

Backpropagation

Key algorithm to train the NN.



Loss Surface
Gradient Descent

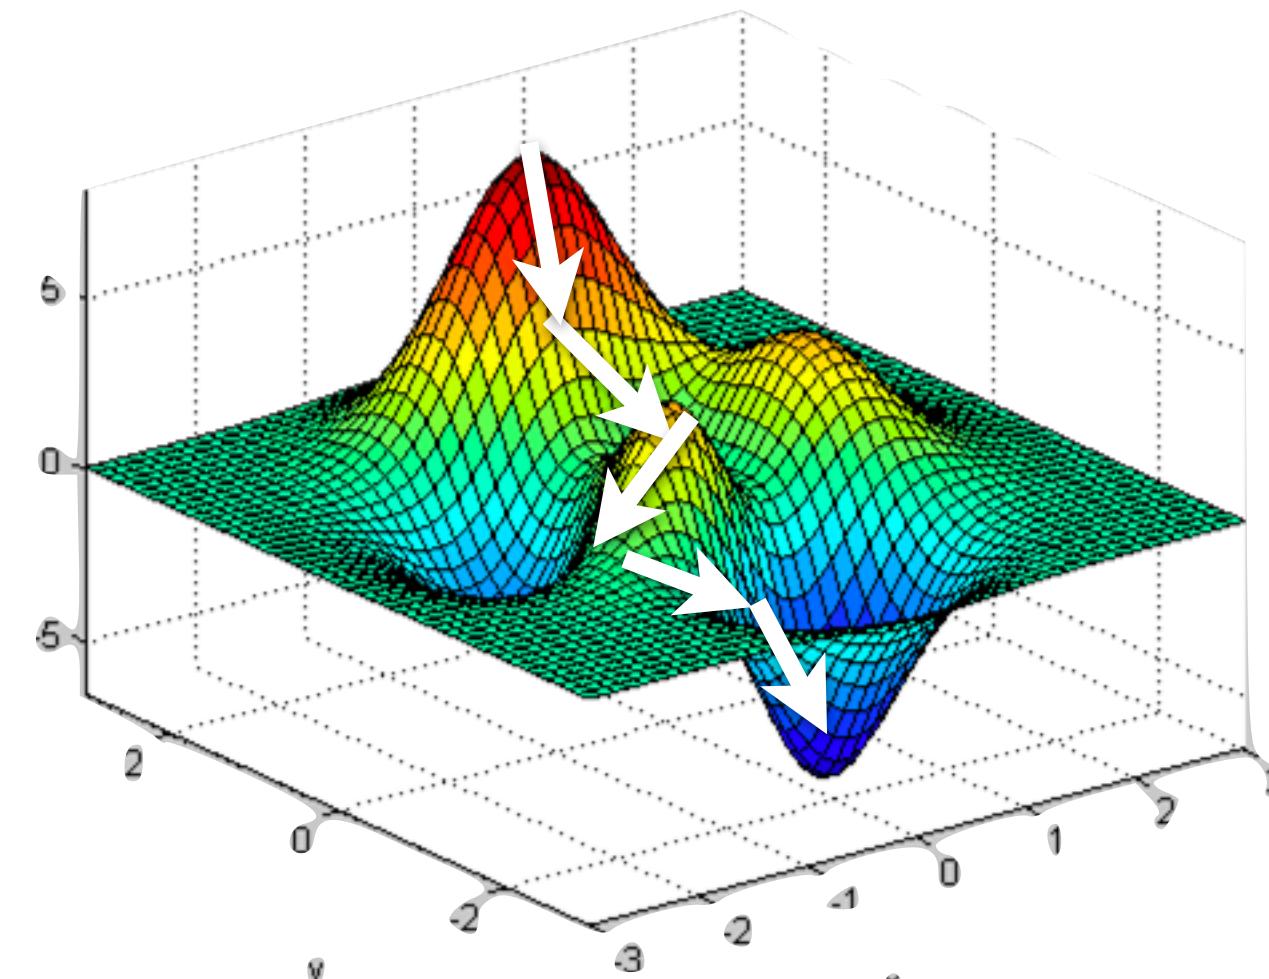


m samples, l hidden layers

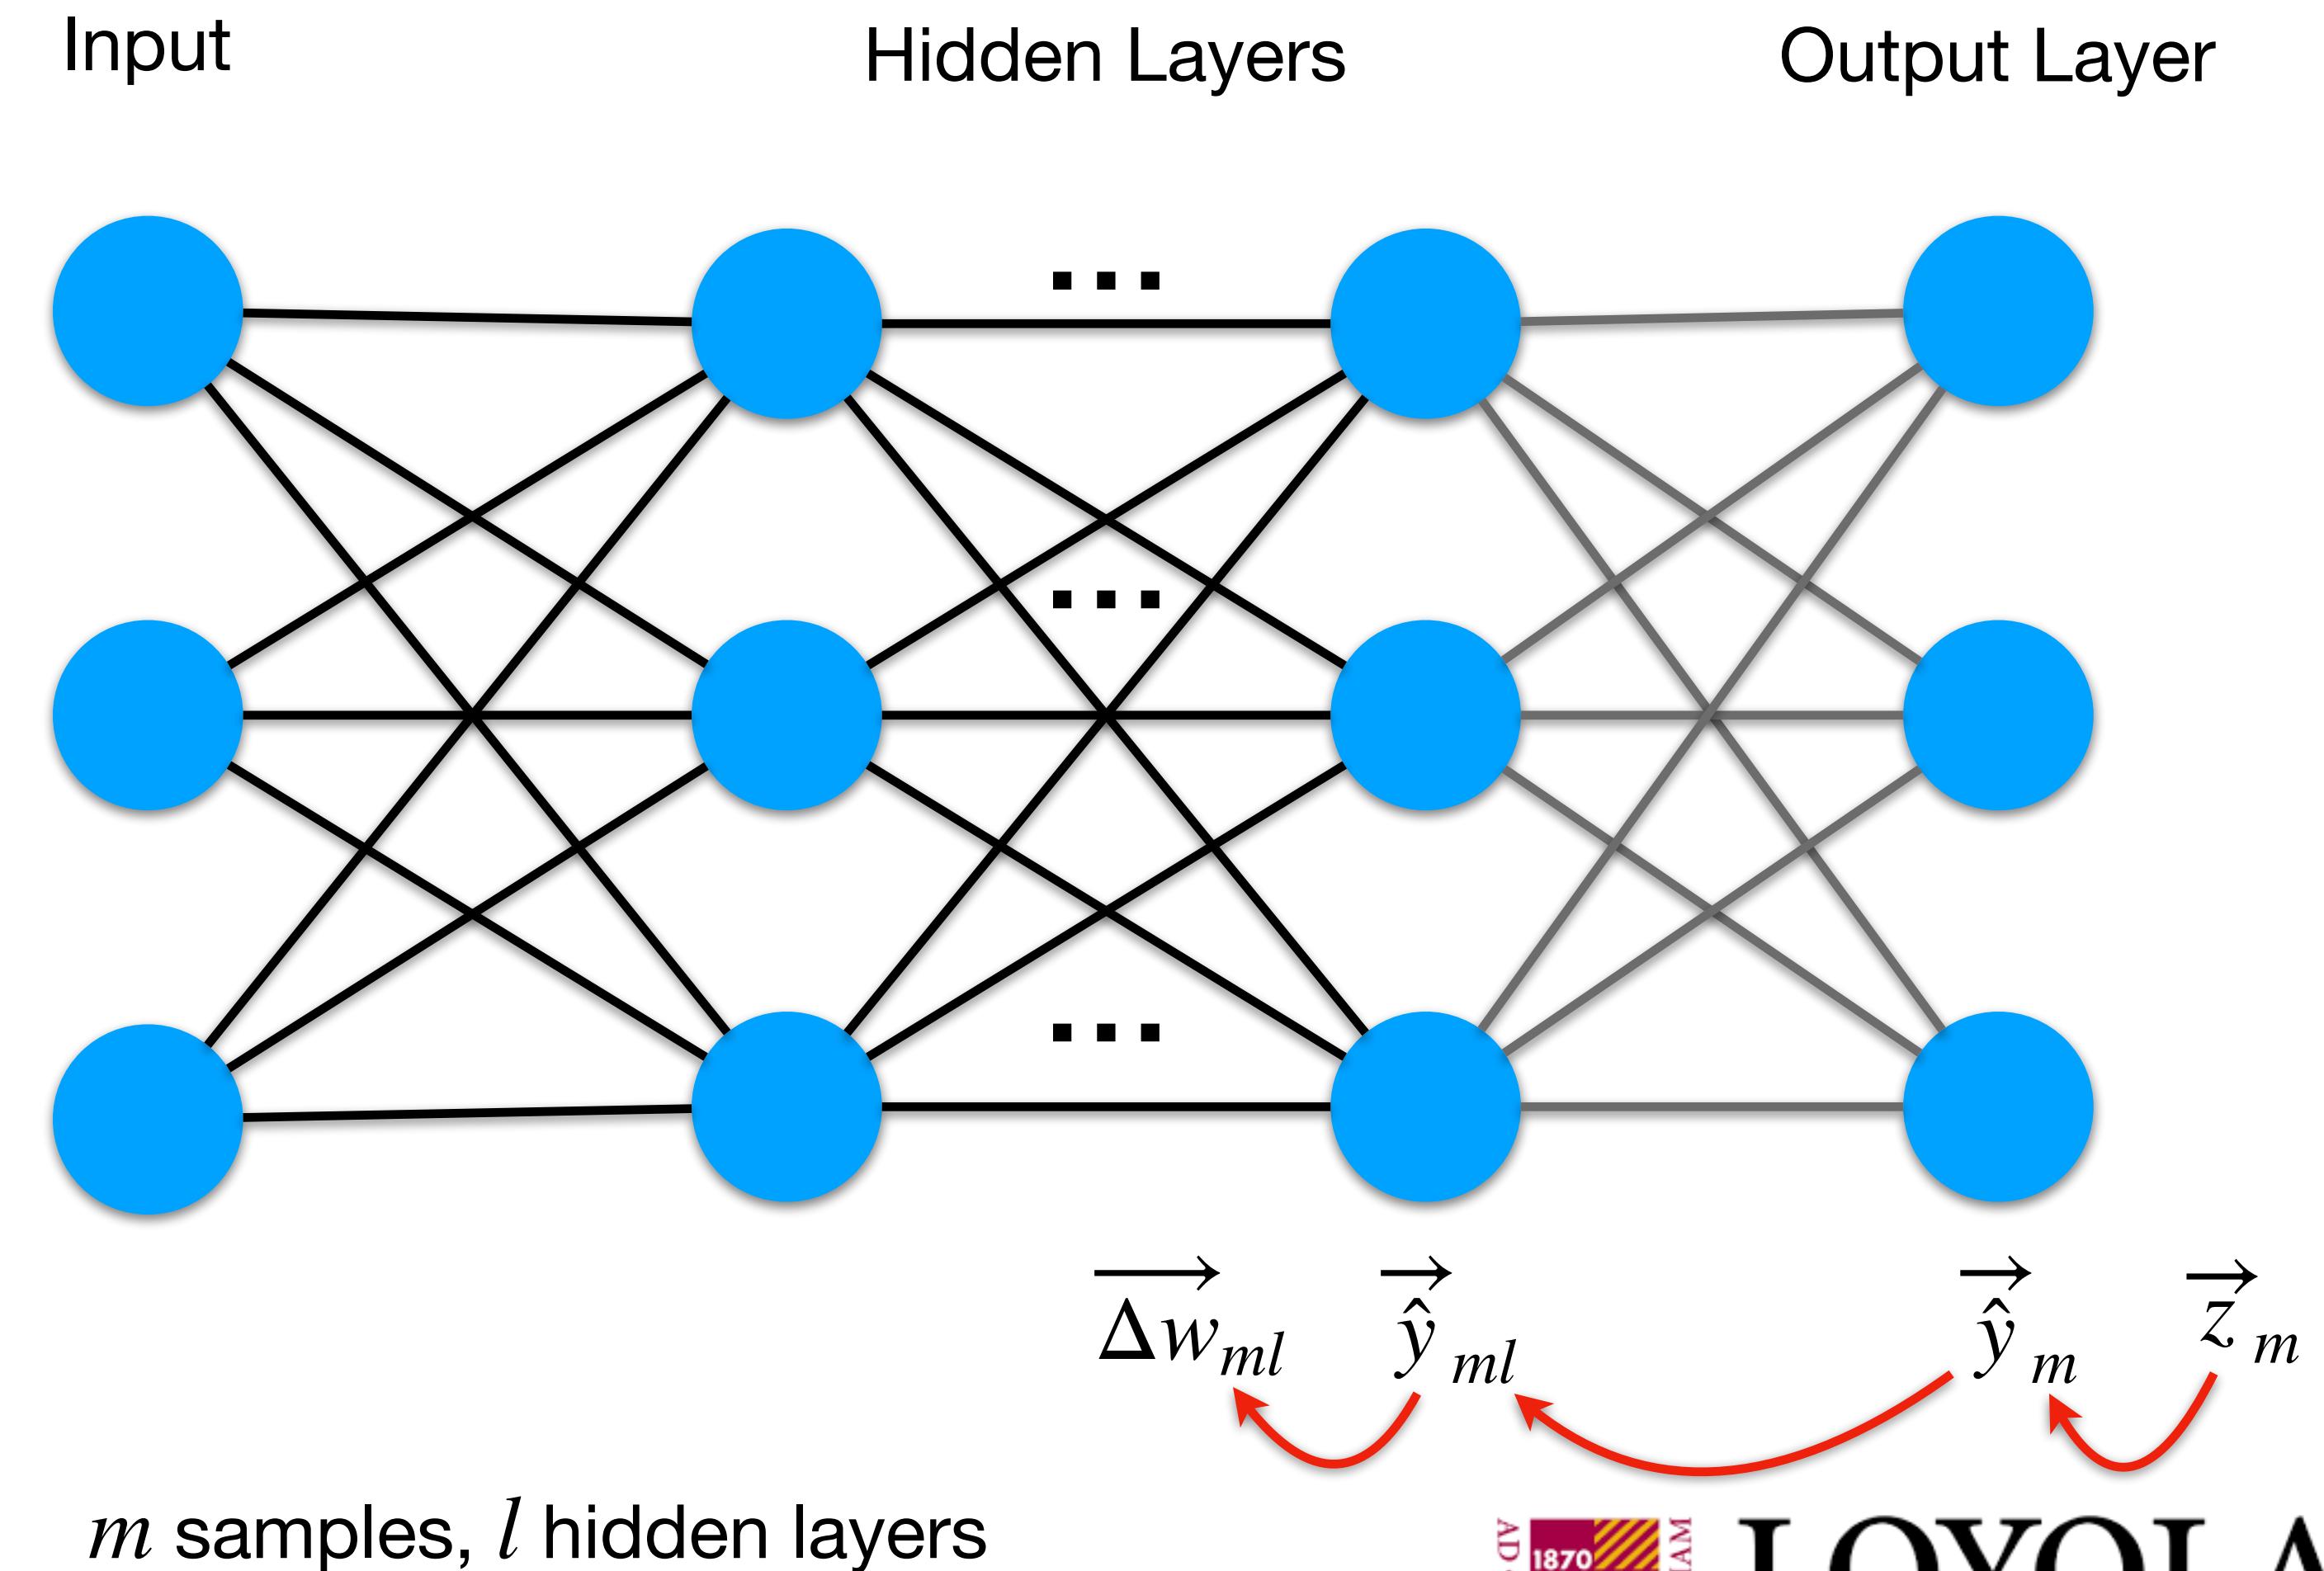
Neural Networks (NN)

Backpropagation

Key algorithm to train the NN.



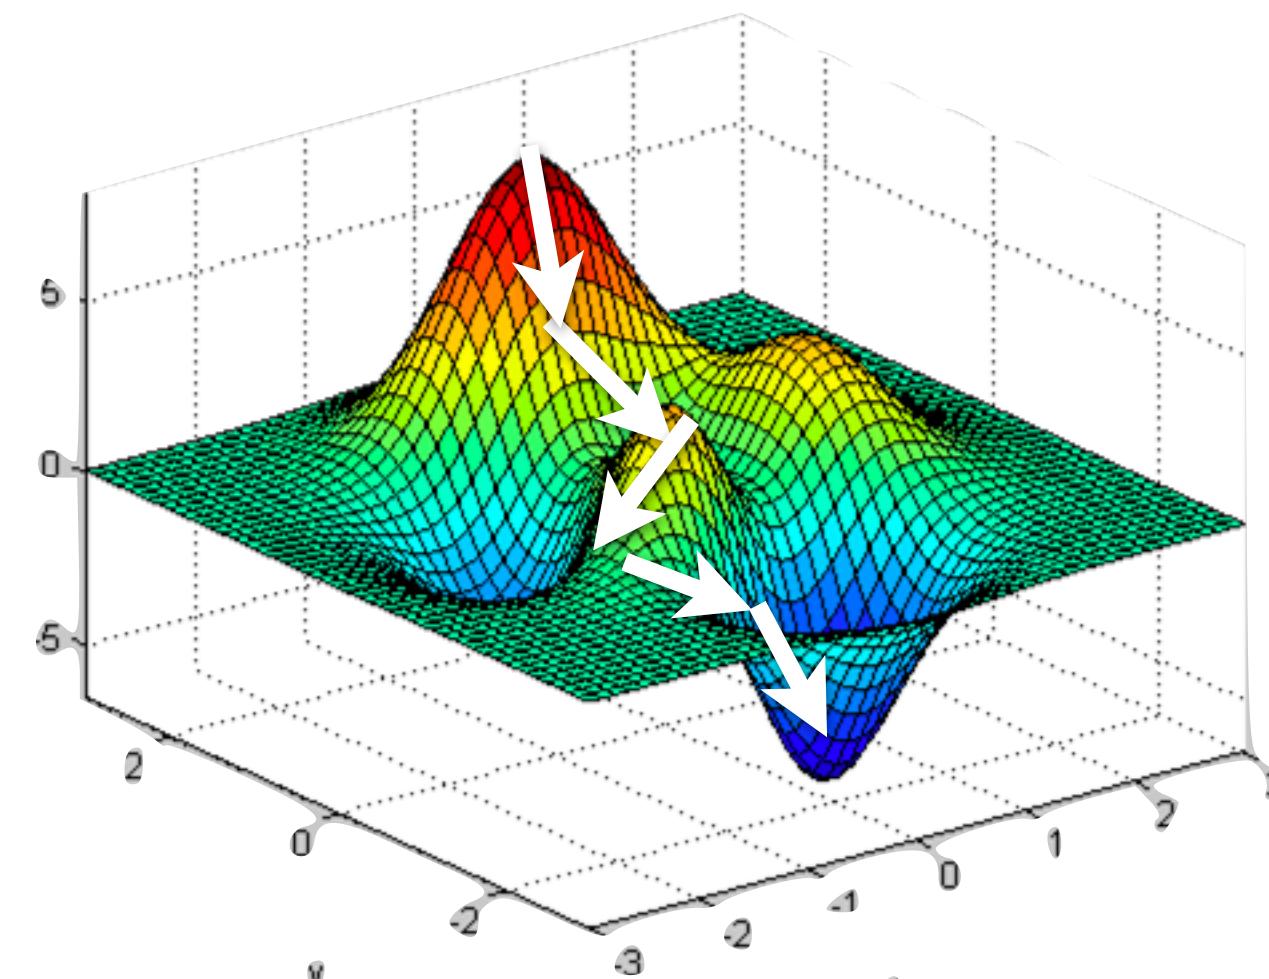
Loss Surface
Gradient Descent



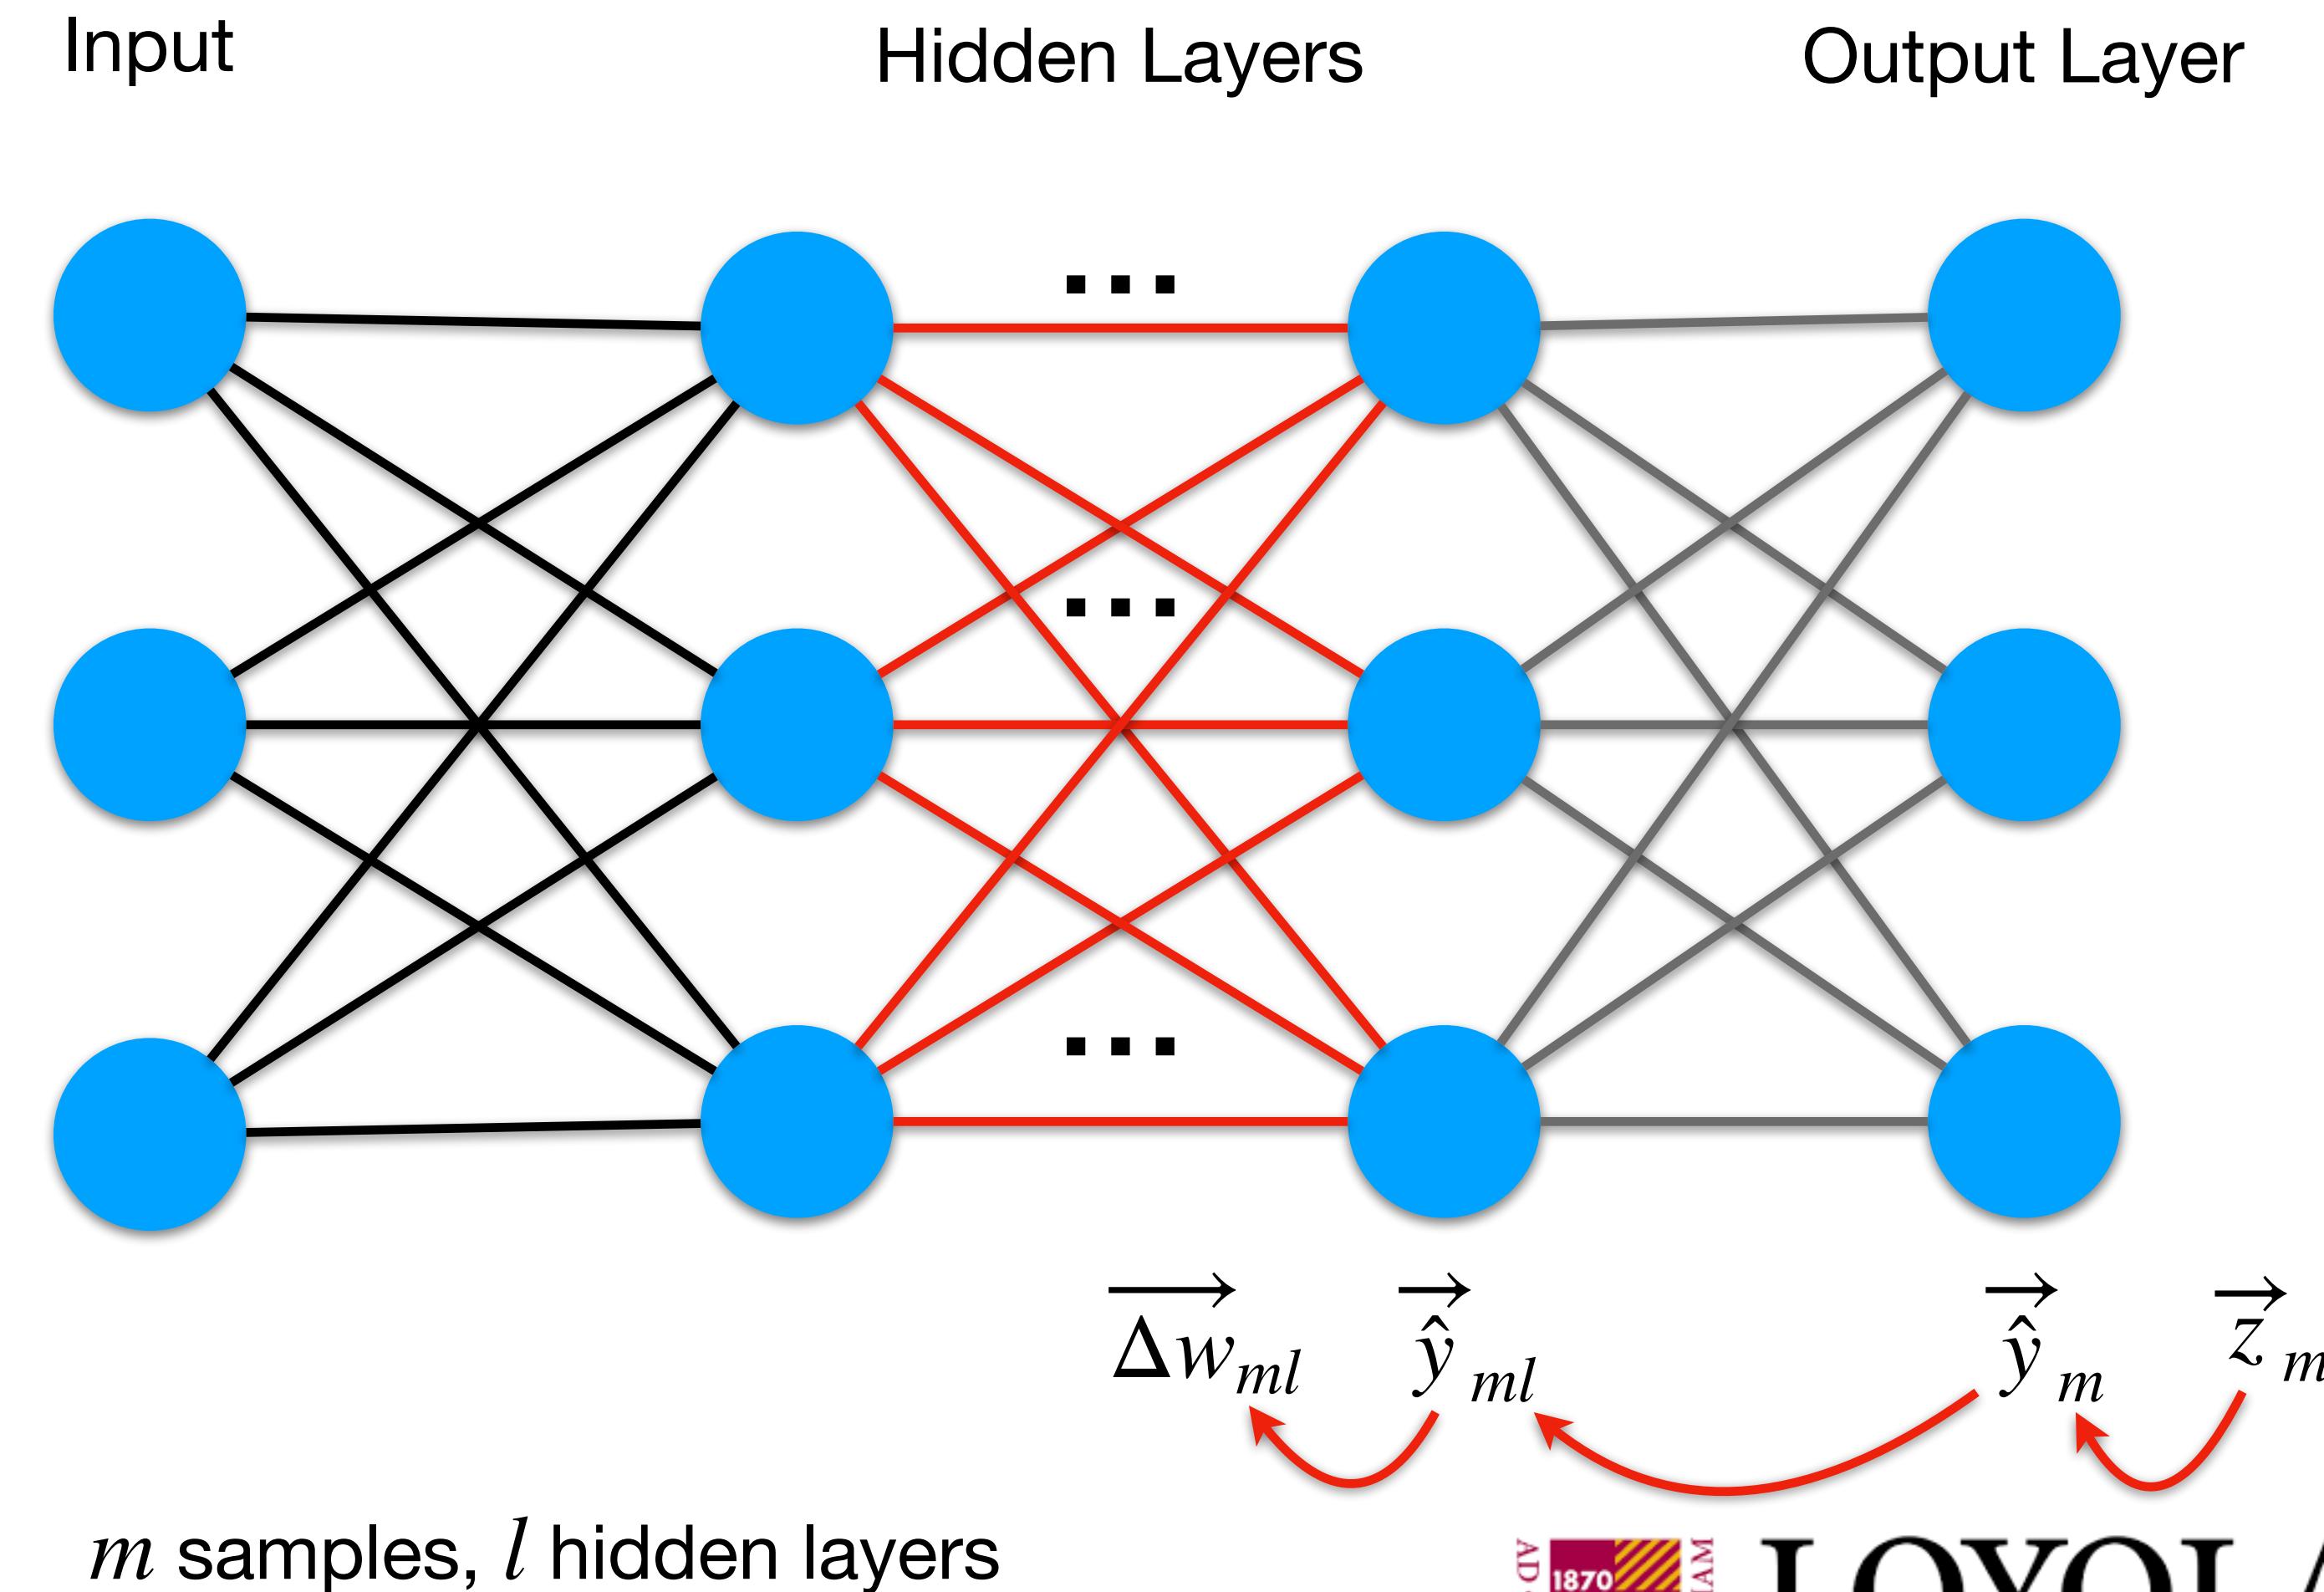
Neural Networks (NN)

Backpropagation

Key algorithm to train the NN.



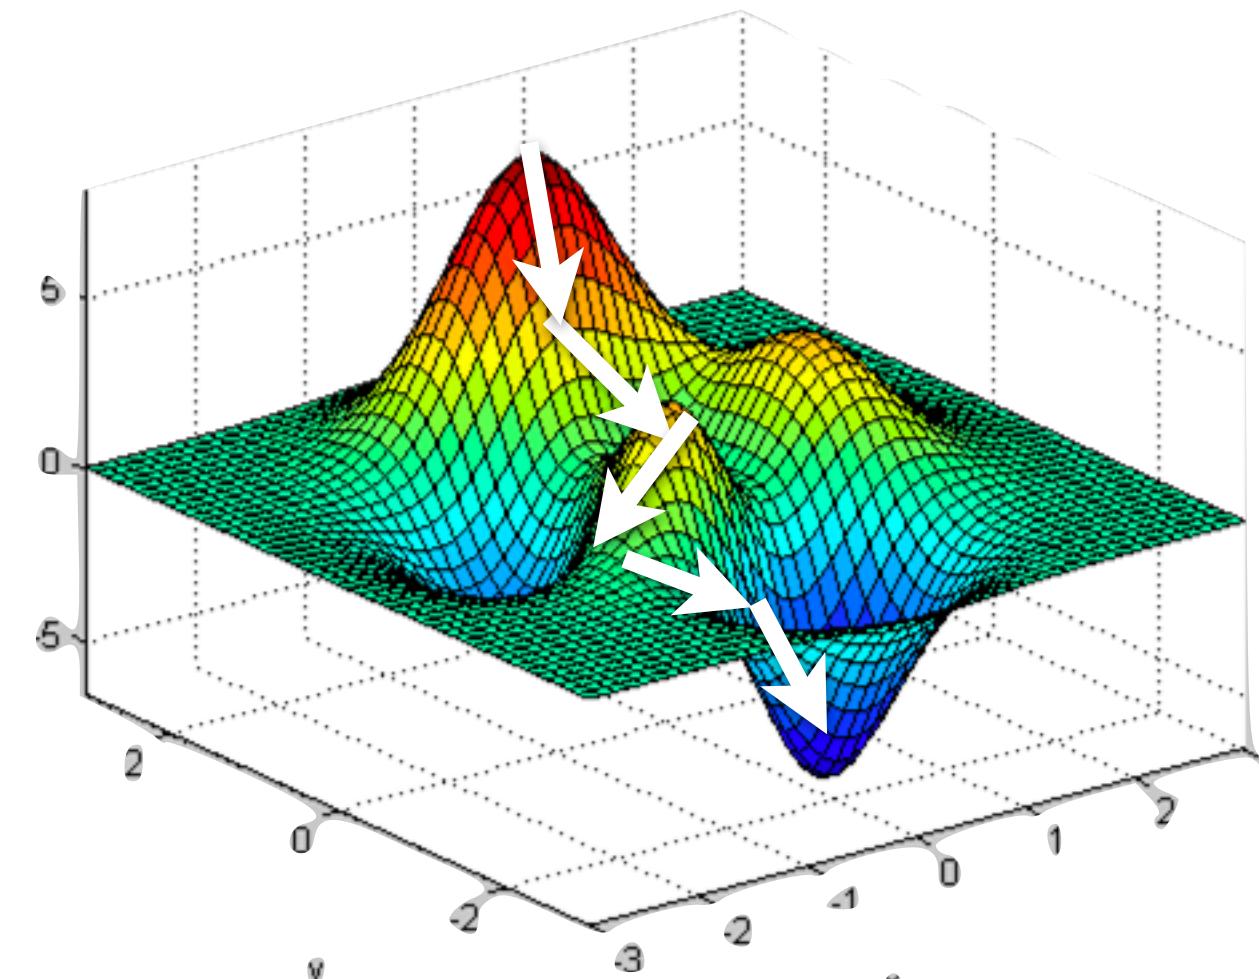
Loss Surface
Gradient Descent



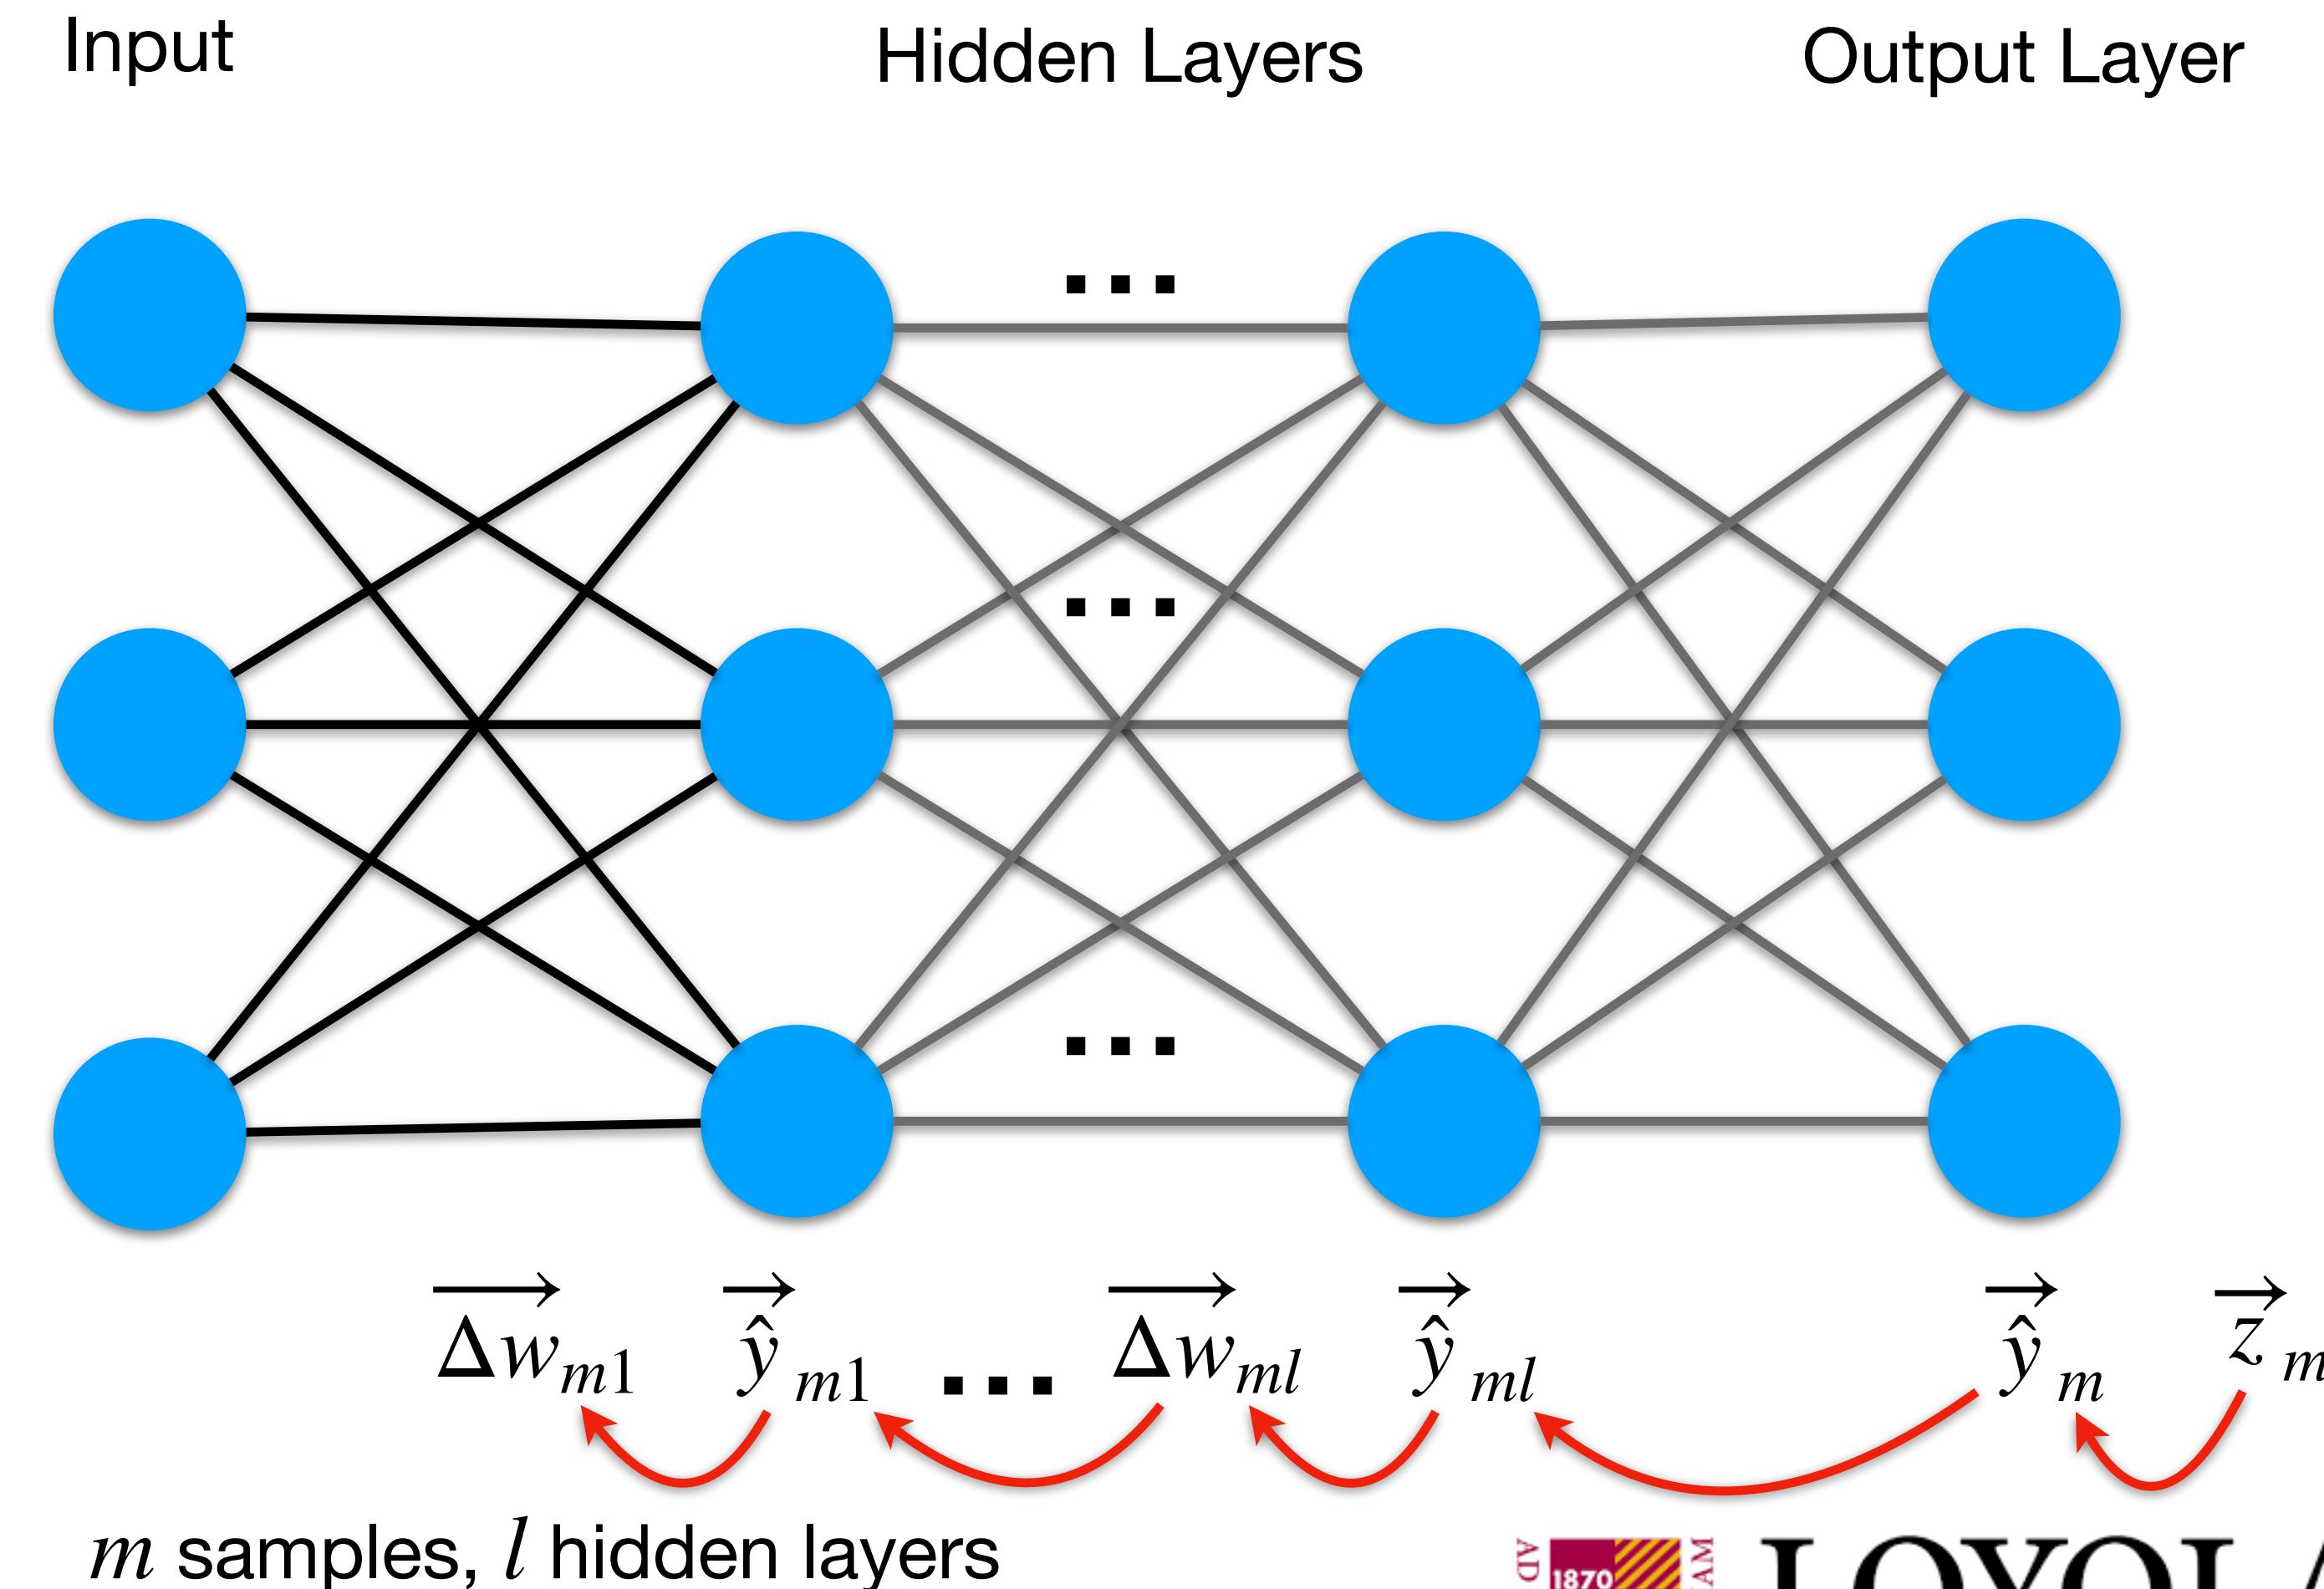
Neural Networks (NN)

Backpropagation

Key algorithm to train the NN.



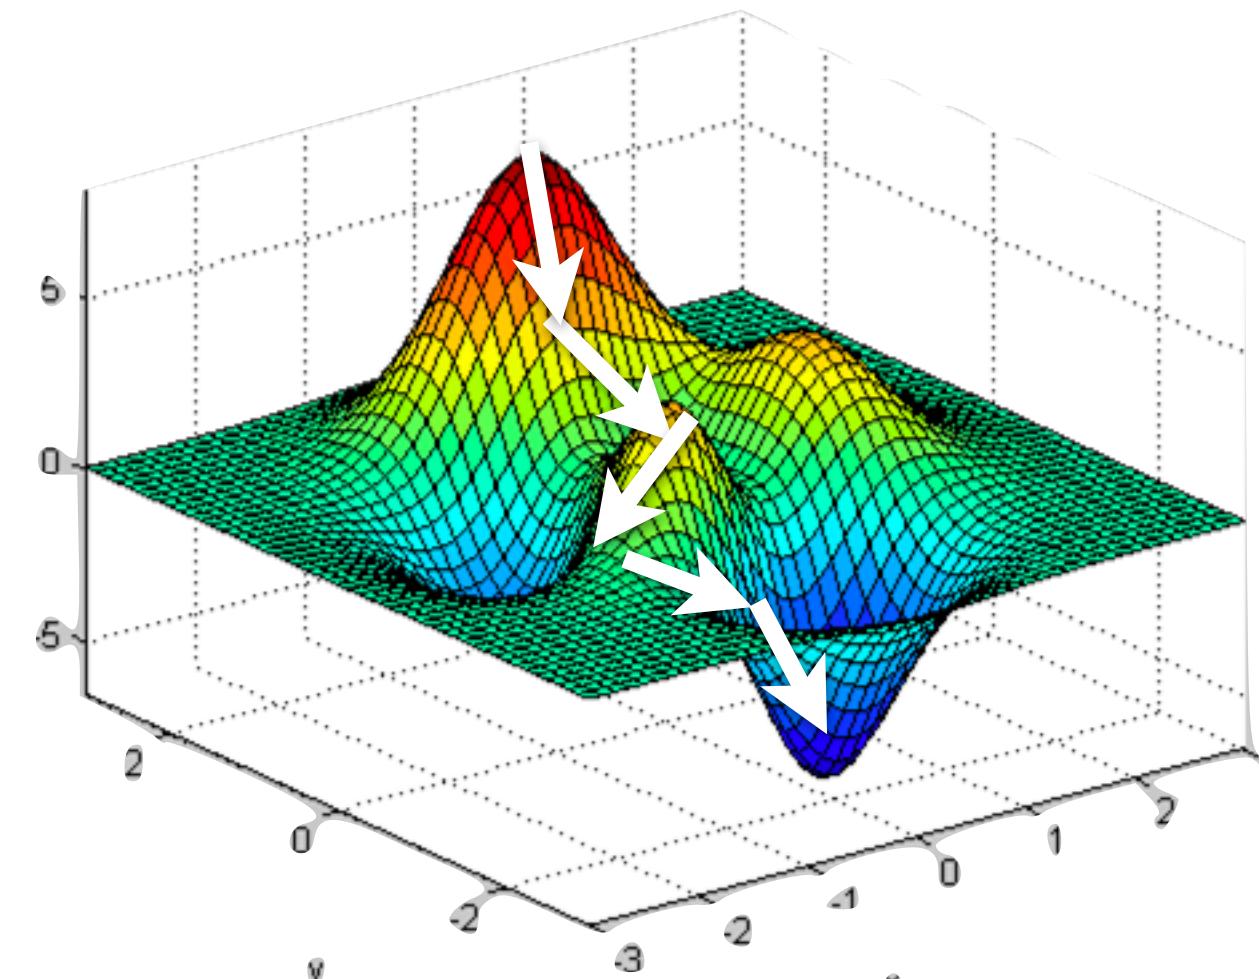
Loss Surface
Gradient Descent



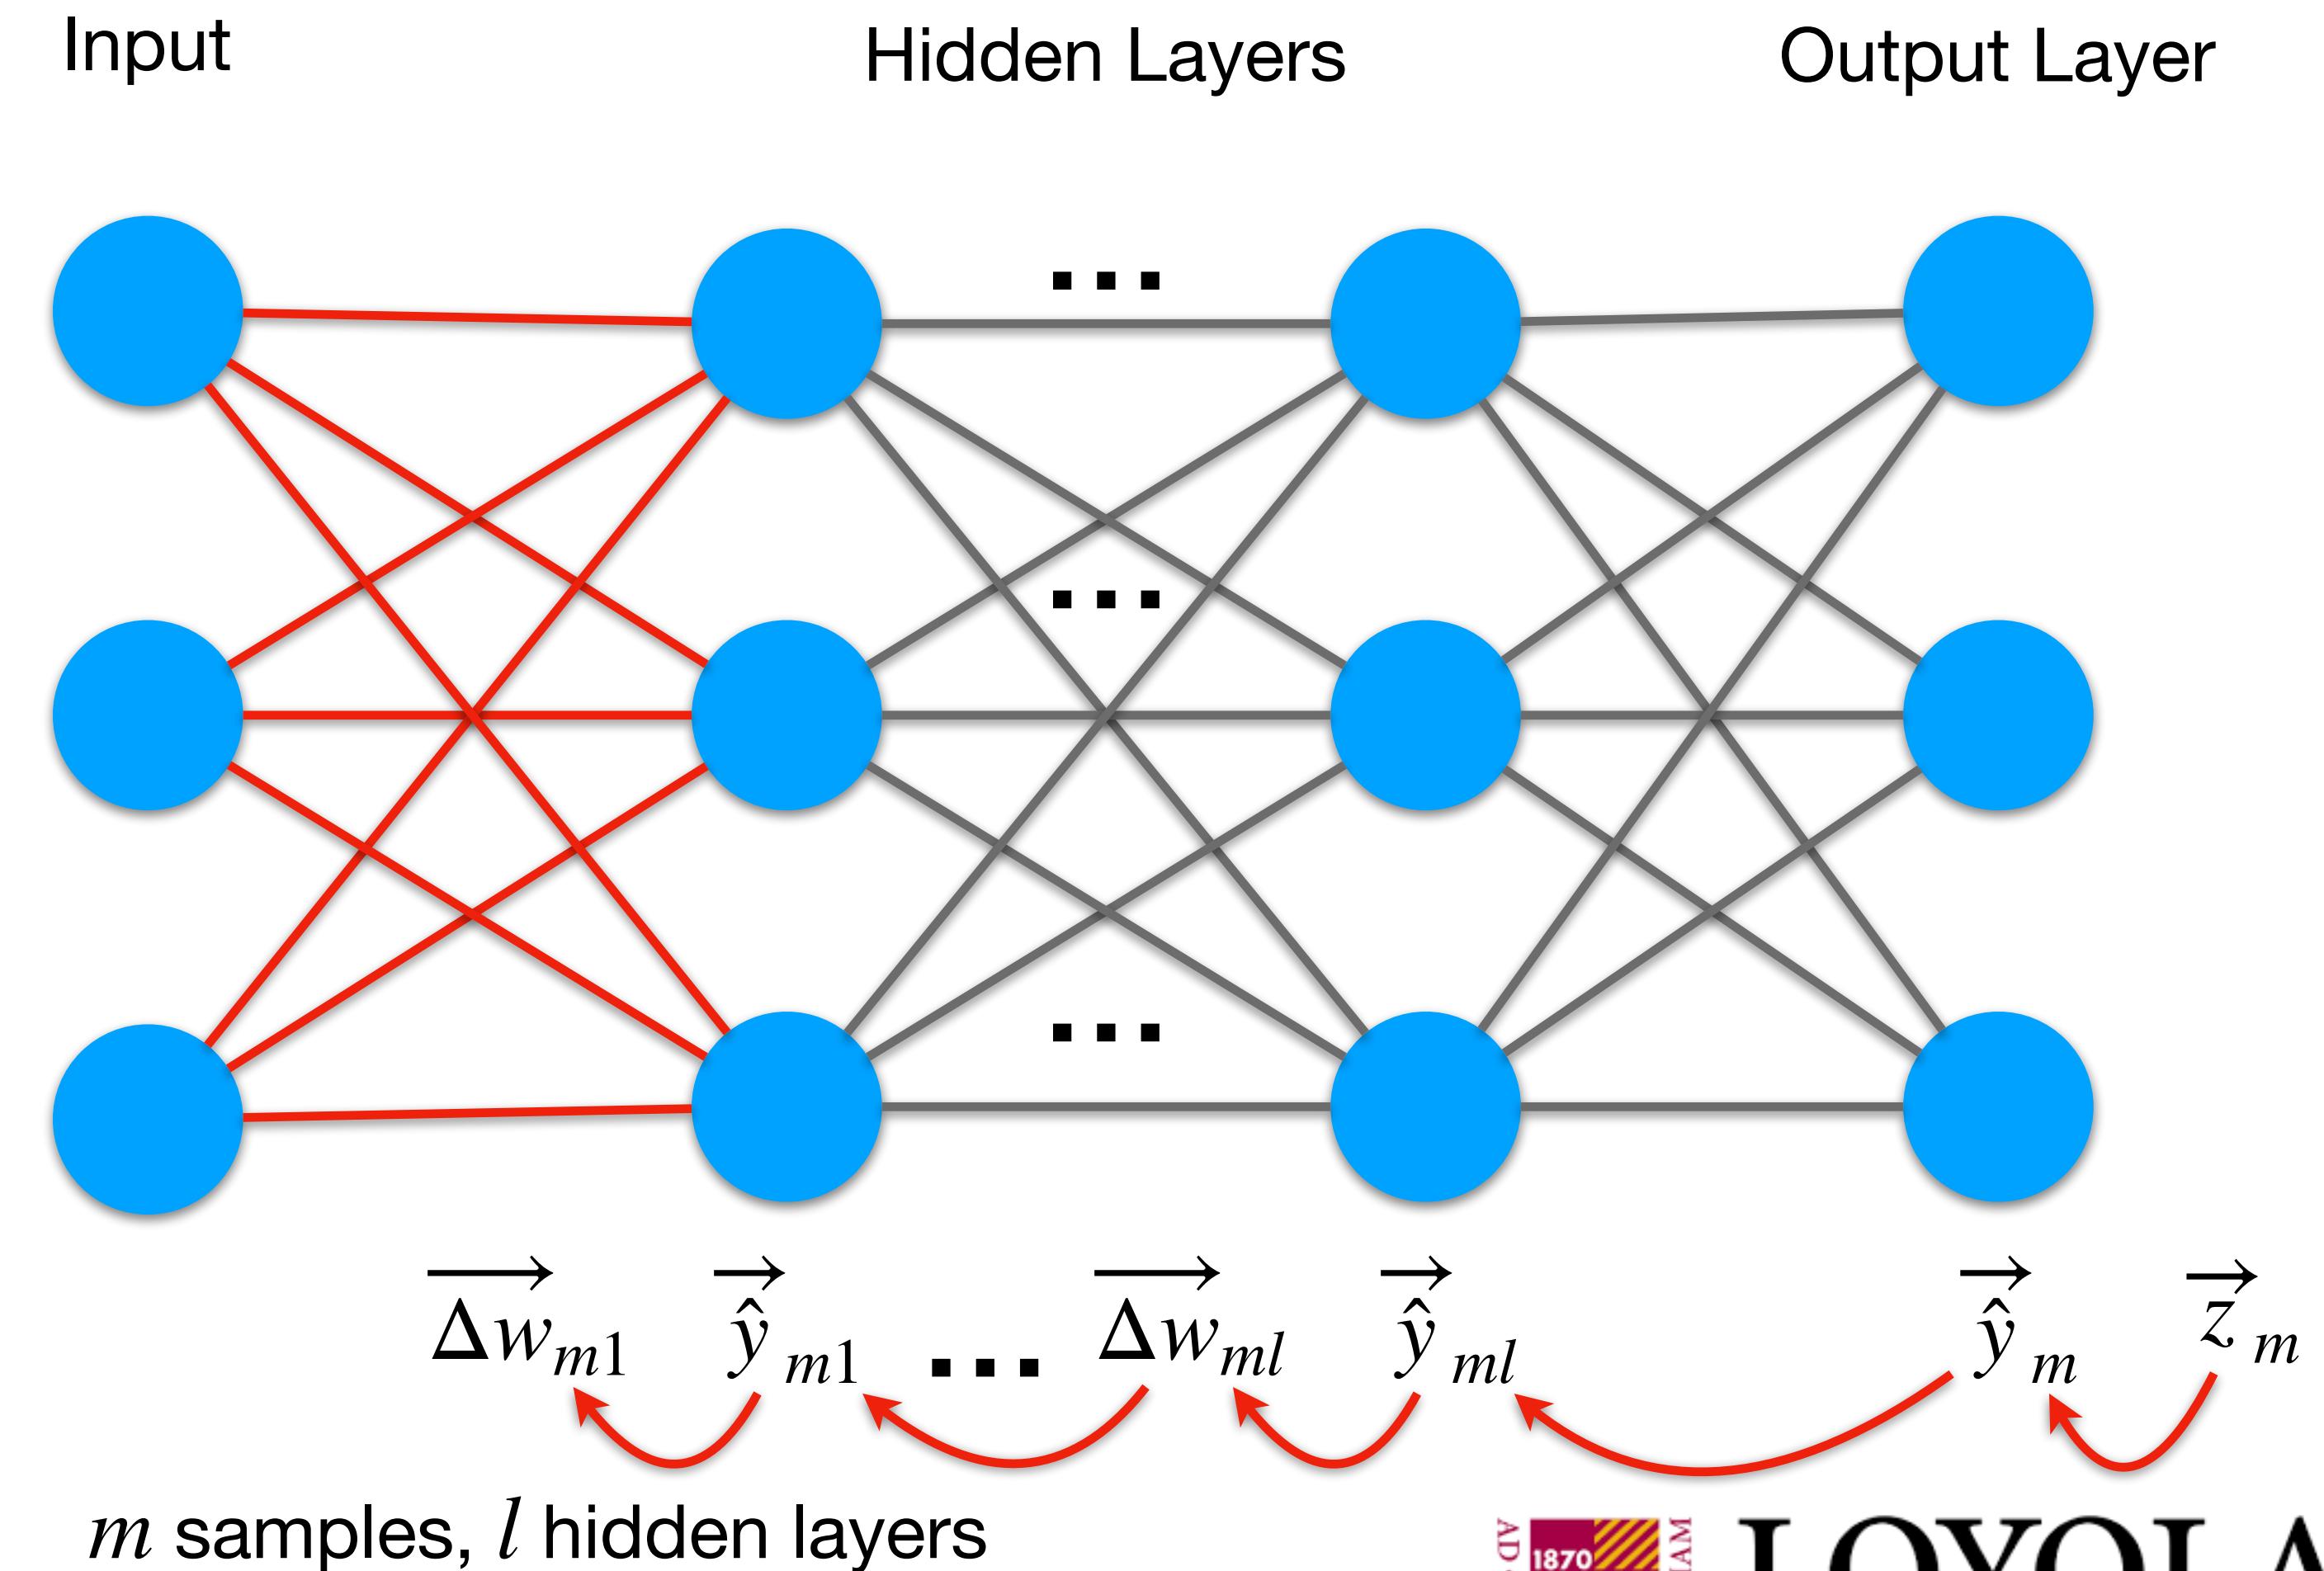
Neural Networks (NN)

Backpropagation

Key algorithm to train the NN.



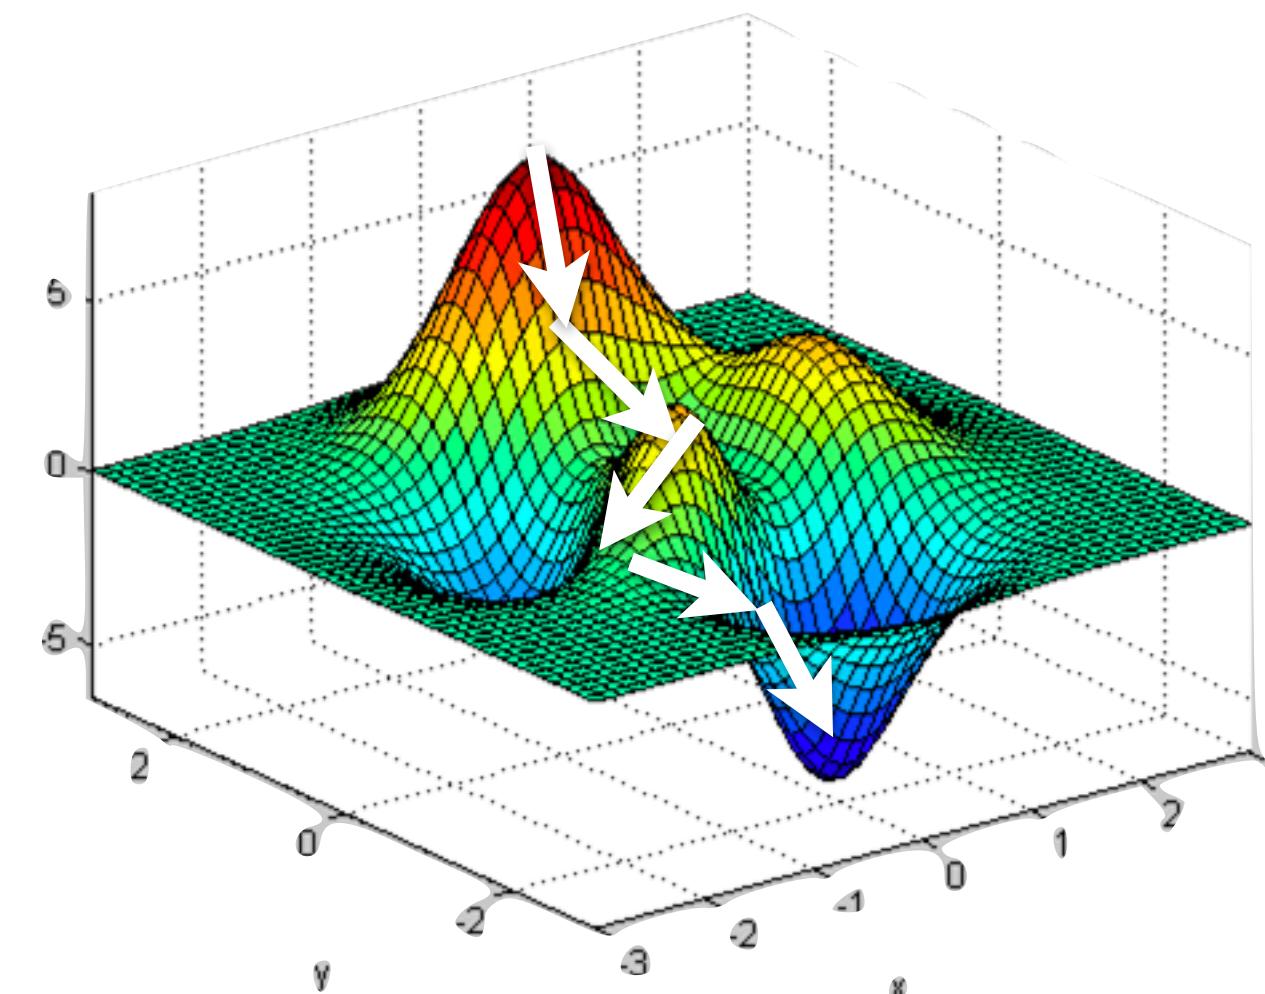
Loss Surface
Gradient Descent



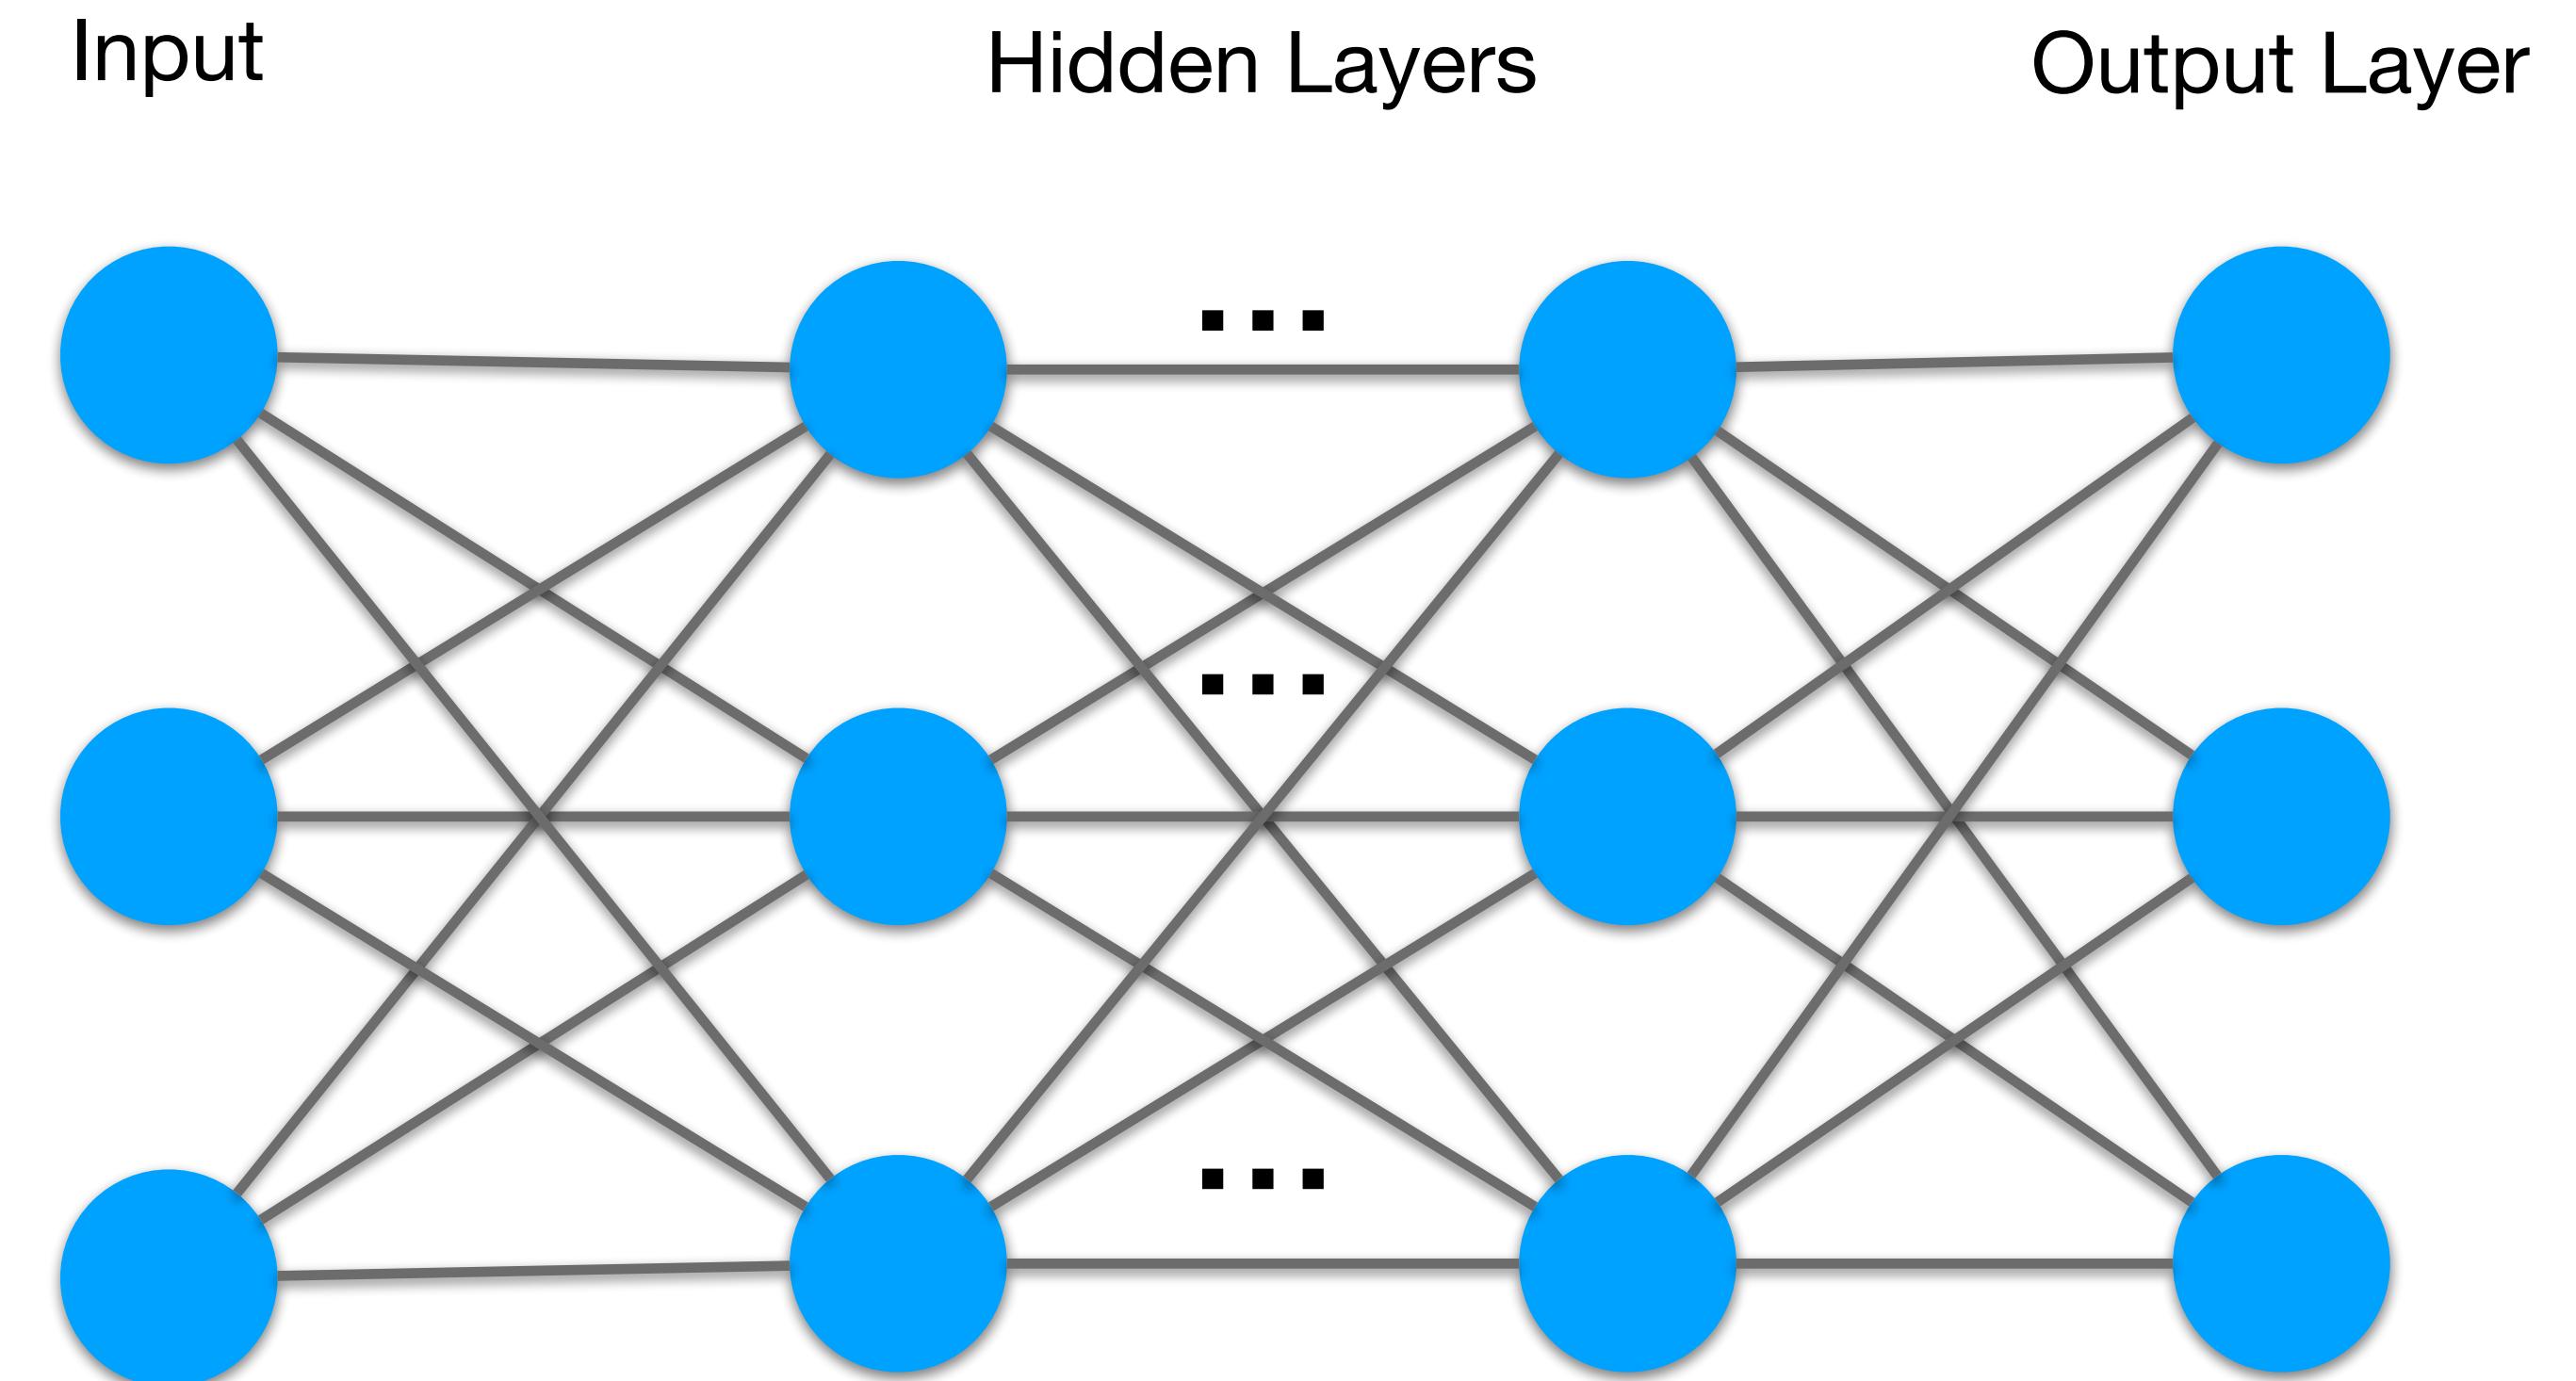
Neural Networks (NN)

Backpropagation

Key algorithm to train the NN.



Loss Surface
Gradient Descent

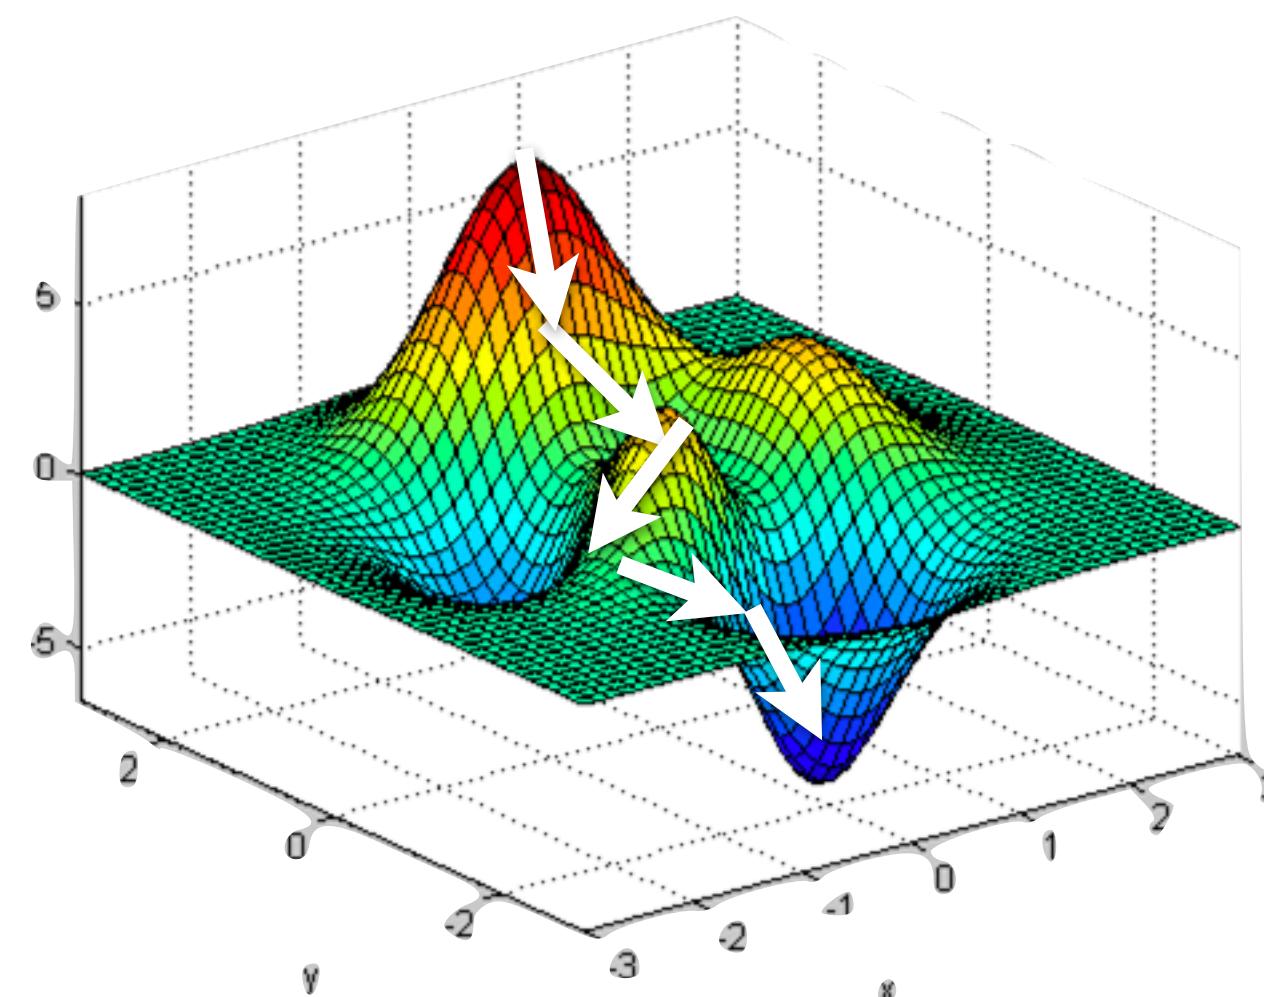


m samples, l hidden layers

Neural Networks (NN)

Backpropagation

Key algorithm to train the NN.



Loss Surface
Gradient Descent

Challenges

The m samples may make back propagation too slow (too much data).

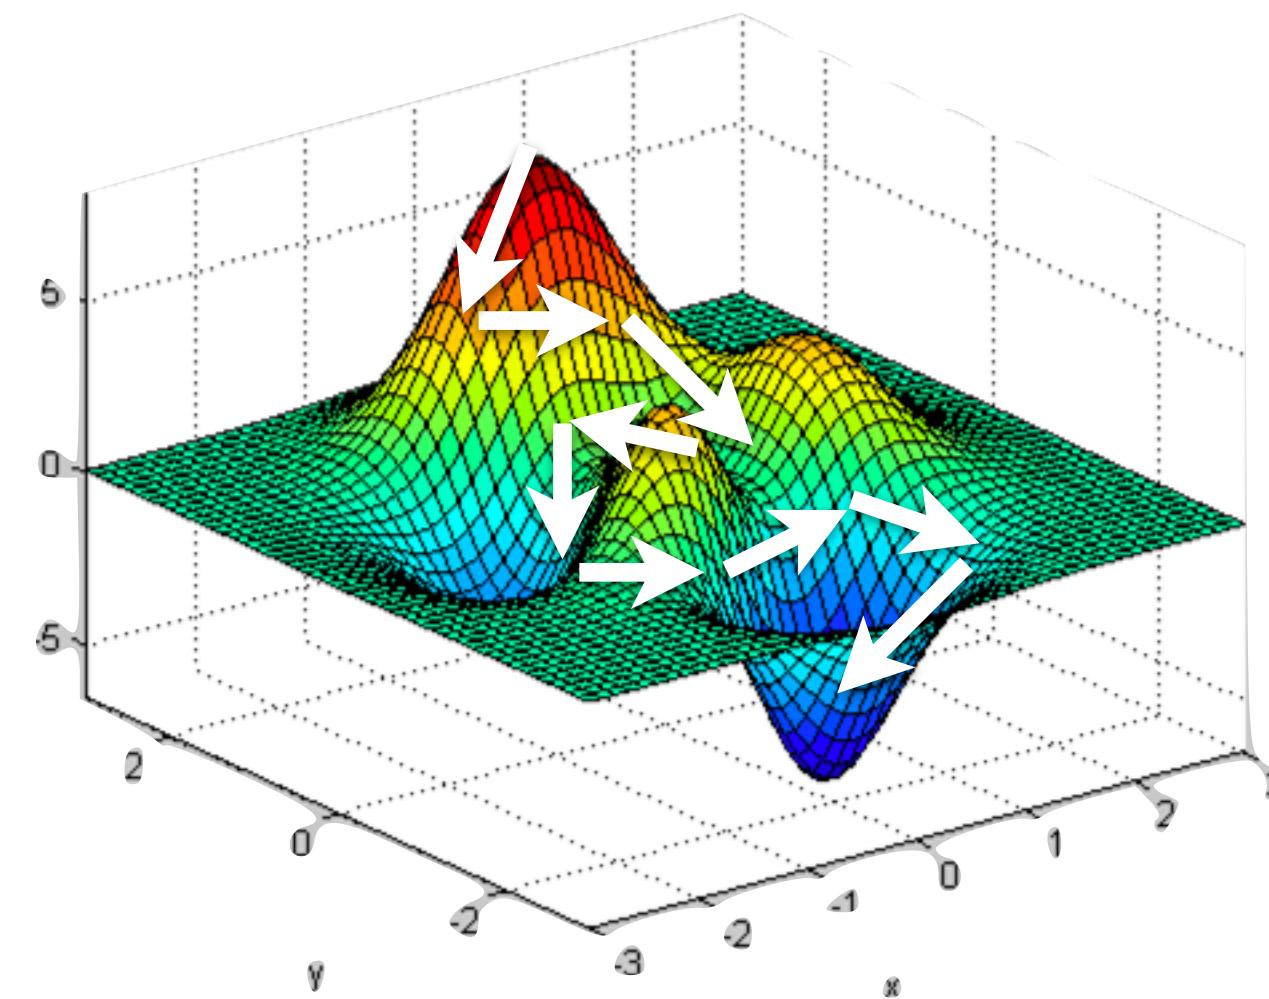
What α (step) should one use?

$$\Delta w_i = \sum_{k=1}^m \alpha(z_k - f(\vec{x}_k^T \cdot \vec{w})) \times x_{ki} \times f'(\vec{x}_k^T \cdot \vec{w}) = 0$$

Neural Networks (NN)

Backpropagation

Key algorithm to train the NN.



Loss Surface
Stochastic Gradient Descent (SGD)

Already implemented
in the NN libraries.

Challenges

The m samples may make back propagation too slow (too much data).

What α (step) should one use?

$$\Delta w_i = \sum_{k=1}^m \alpha(z_k - f(\vec{x}_k^T \cdot \vec{w})) \times x_{ki} \times f'(\vec{x}_k^T \cdot \vec{w}) = 0$$

Solution

Stochastic Gradient Descent

Randomly select multiple smaller subsets of the m samples (mini-batches).

Run more but faster iterations of backpropagation.

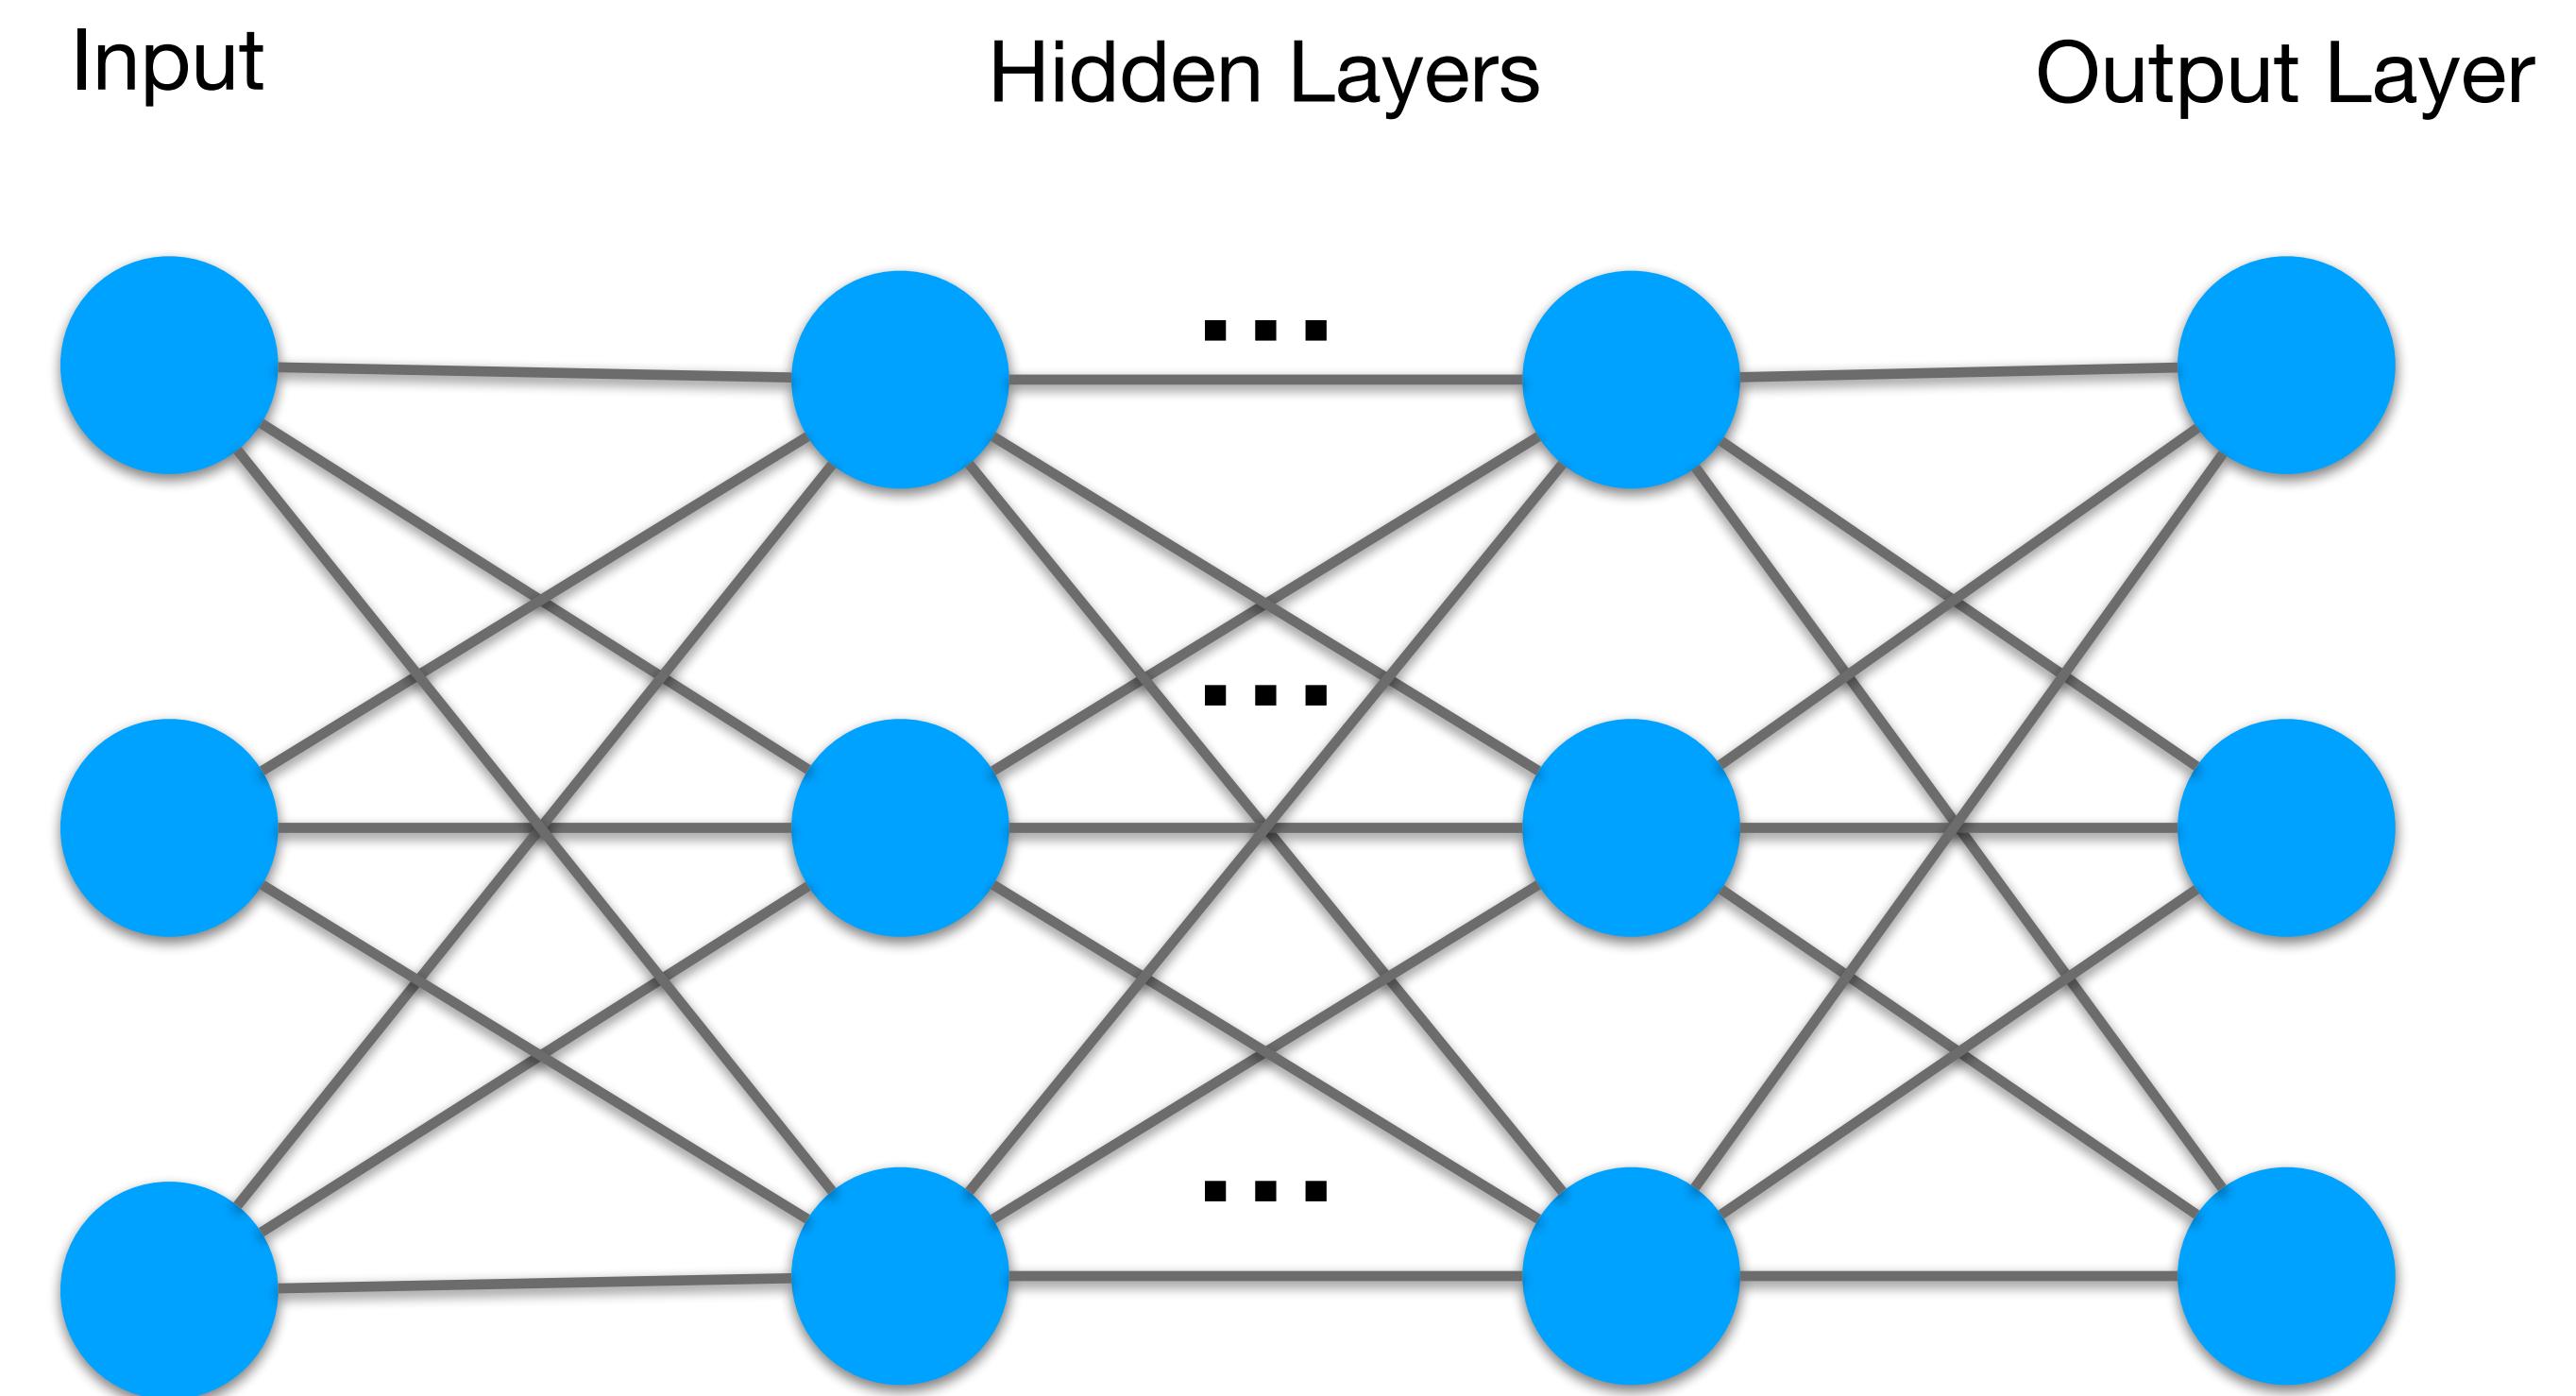
Convolutional Neural Networks (CNN)

How NNs are used in CV?

Flatten the image and feed pixel values to the input?

Cons

Pixel values at homogeneous regions (e.g., texture and blurred regions) have too similar values and may be redundant.

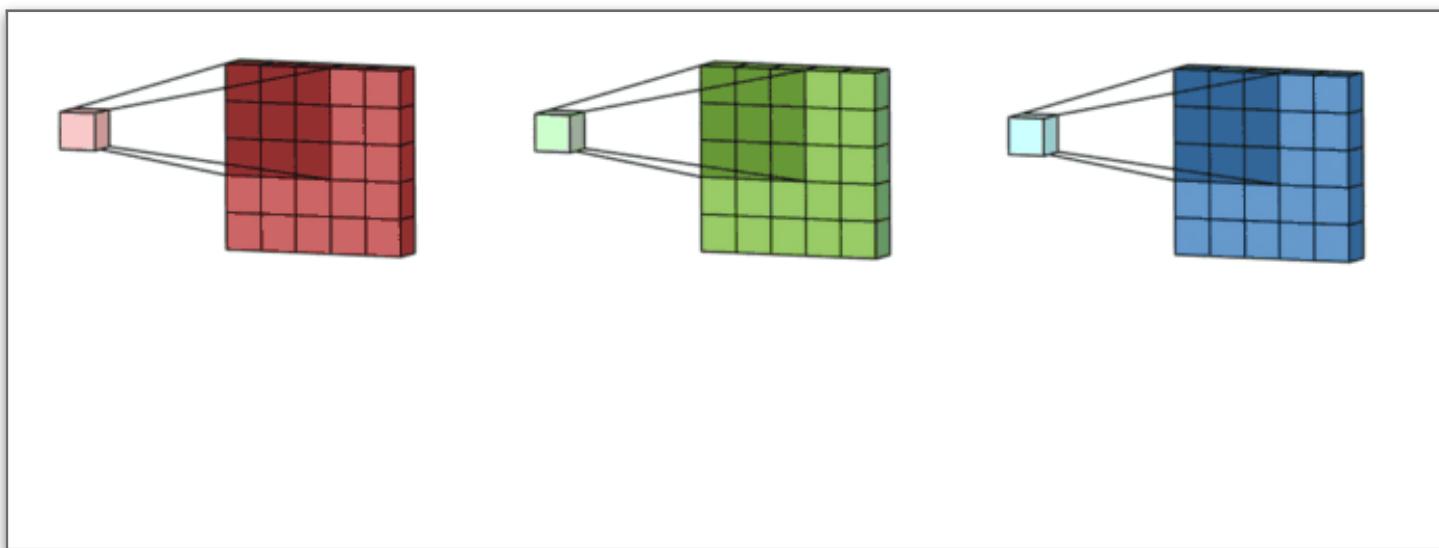


Convolutional Neural Networks (CNN)

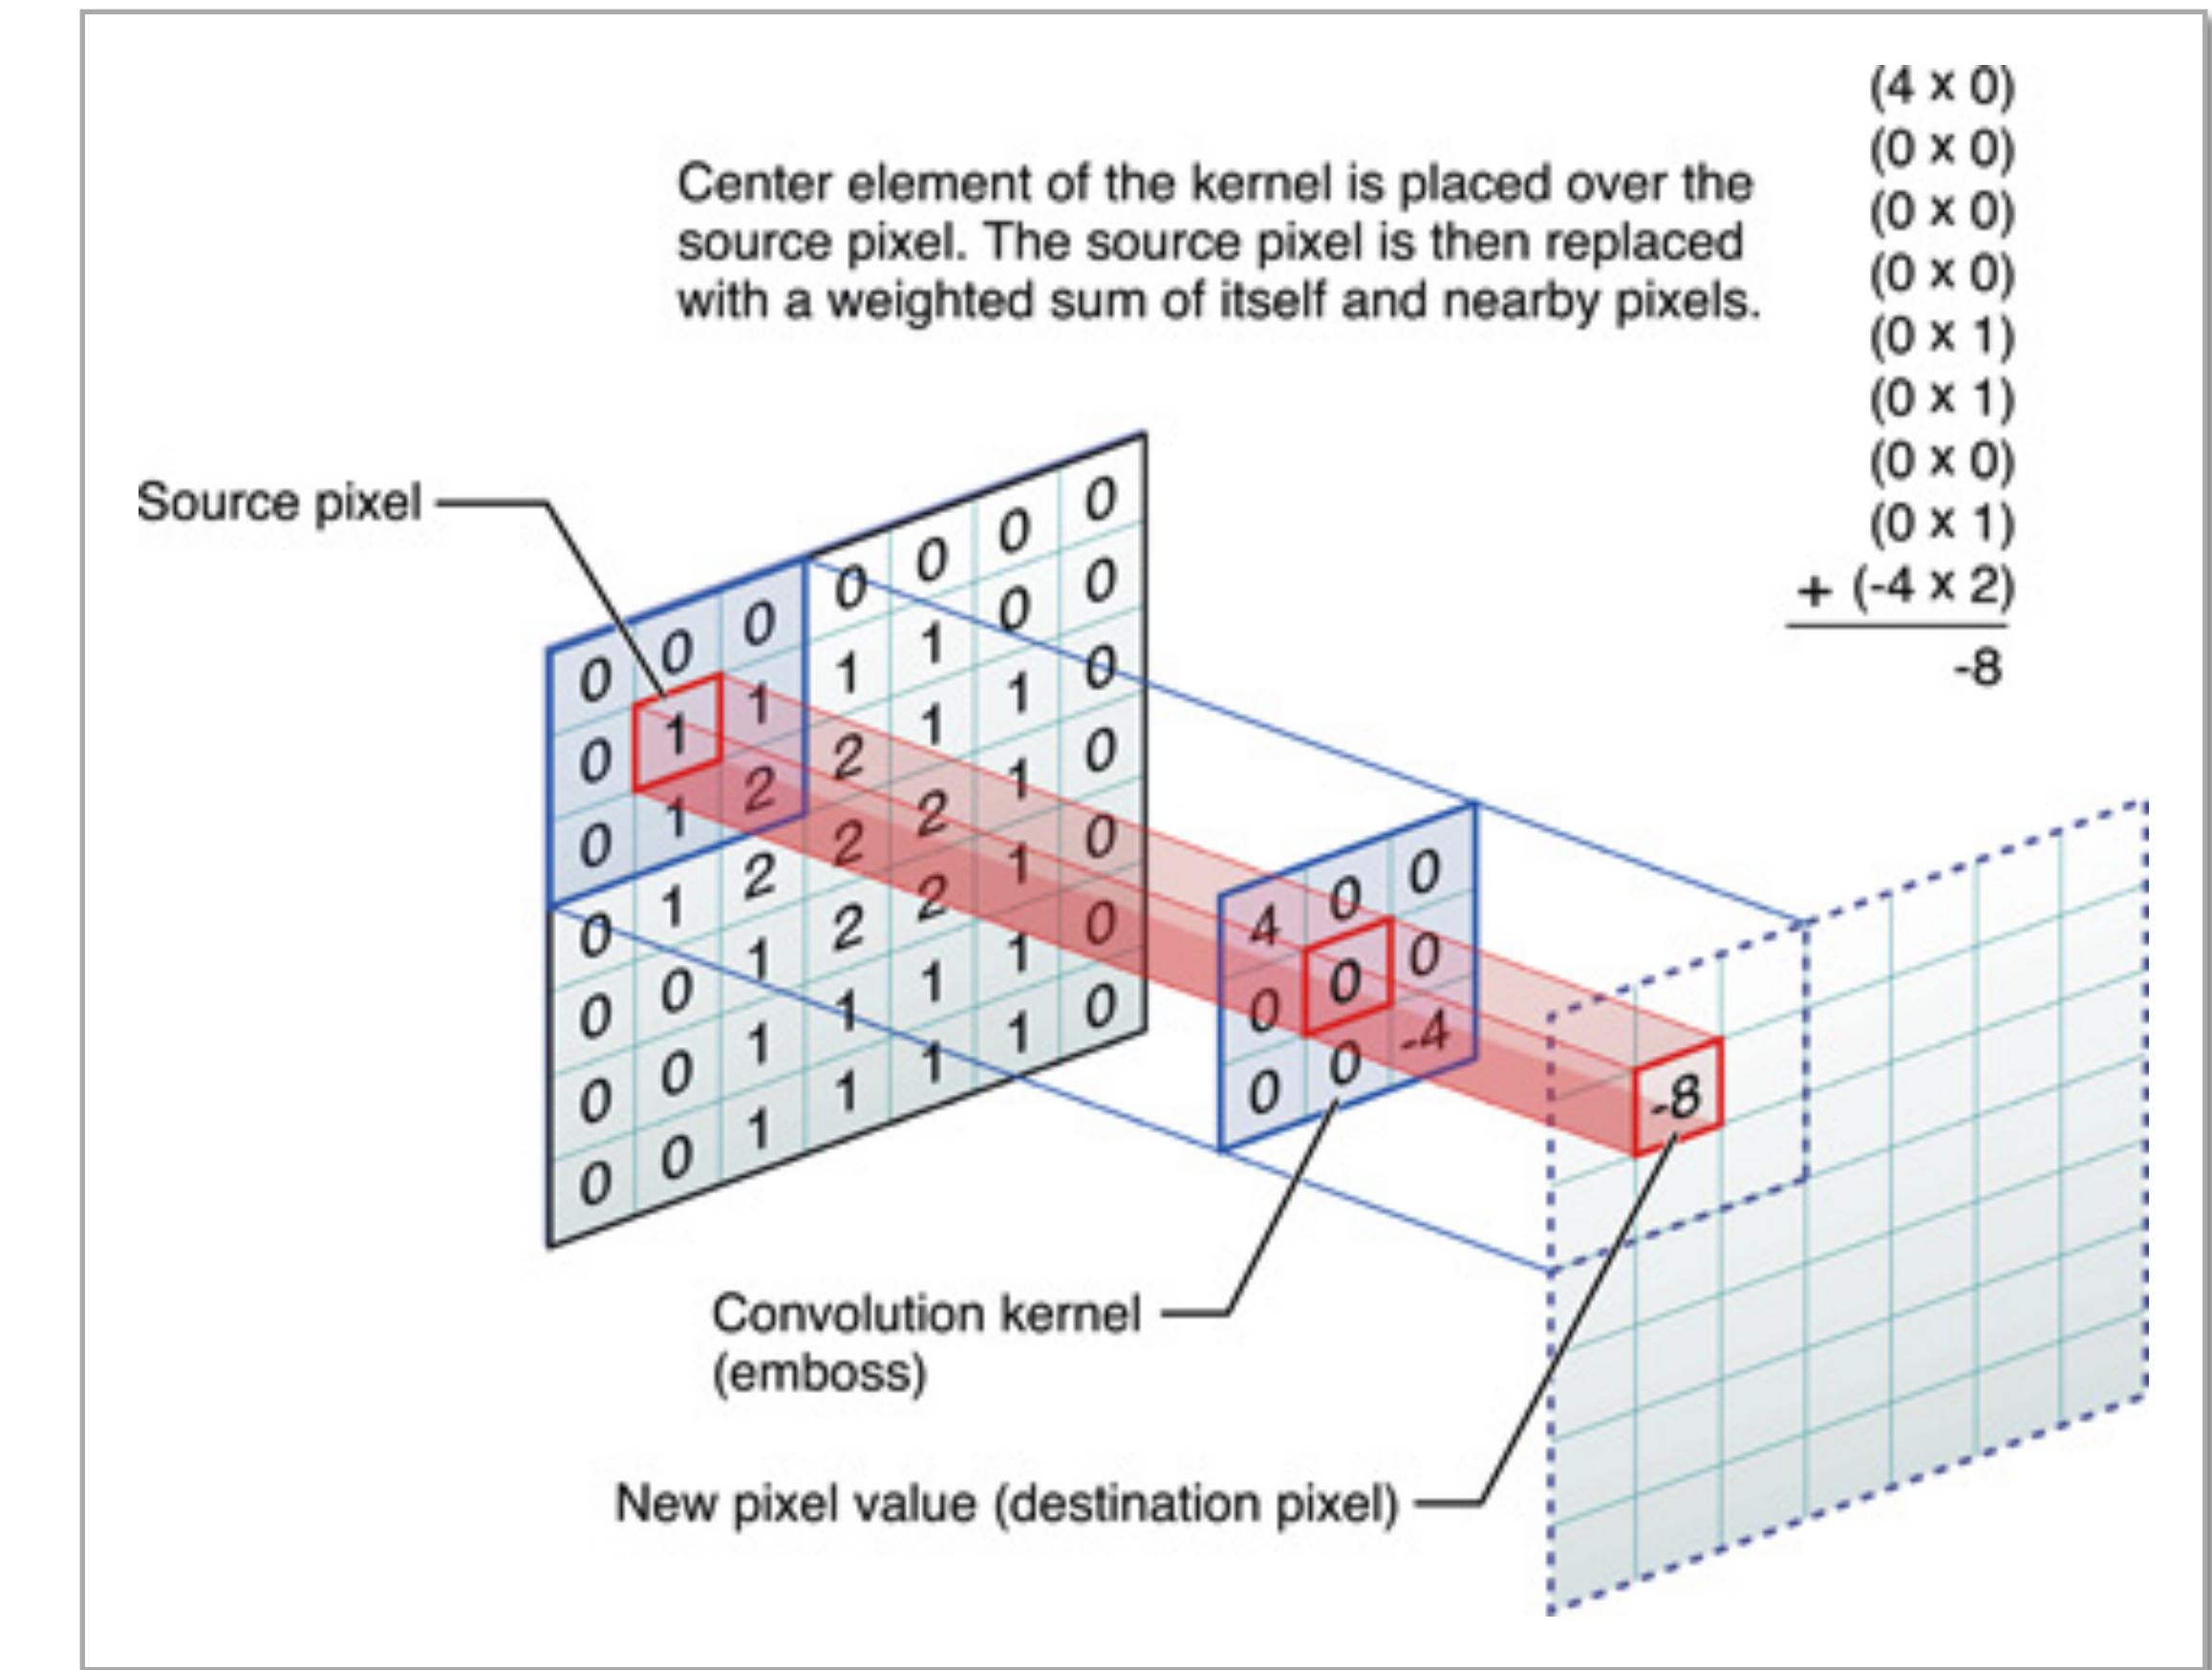
How NNs are used in CV?

Solution

Add convolutions to the NN.



<https://medium.com/analytics-vidhya/understanding-convolution-operations-in-cnn-1914045816d4>



<https://developer.apple.com/library/library/Archive/documentation/Performance/Conceptual/vImage/ConvolutionOperations/ConvolutionOperations.html>

Convolutional Neural Networks (CNN)

How NNs are used in CV?

Solution

Add convolutions to the NN.

Examples of **handcrafted** convolutional filters.

<i>Original</i>	<i>Gaussian Blur</i>	<i>Sharpen</i>	<i>Edge Detection</i>
$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$
			

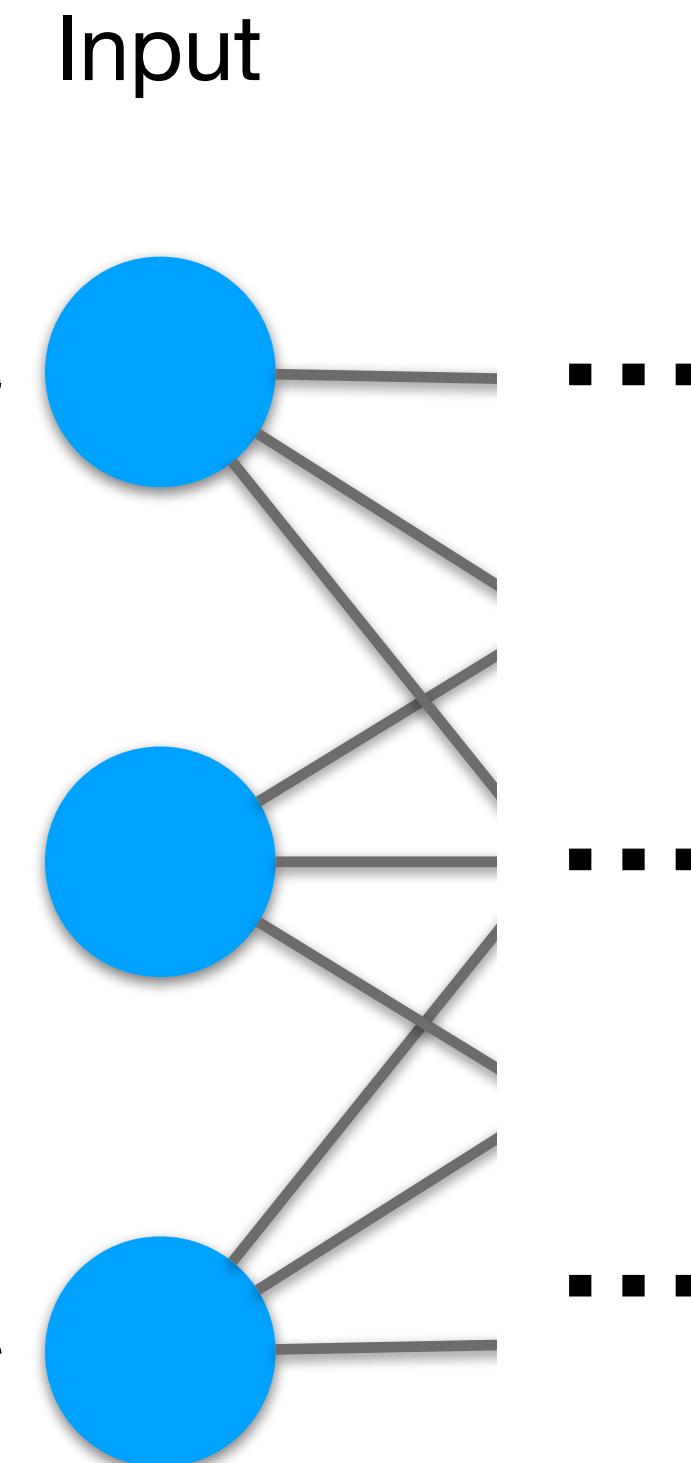
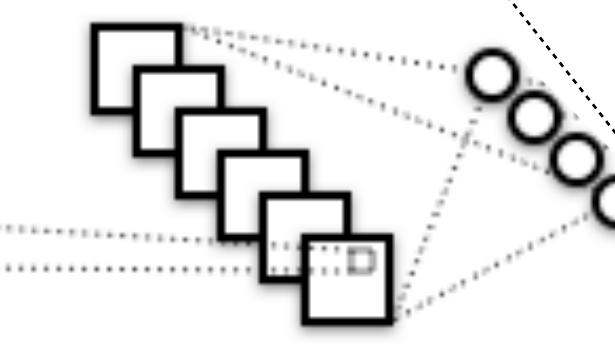
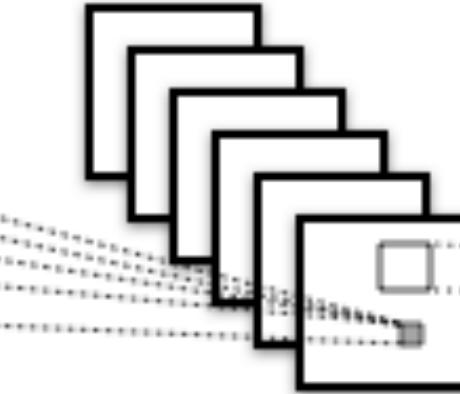
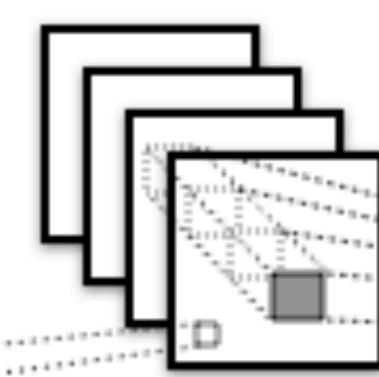
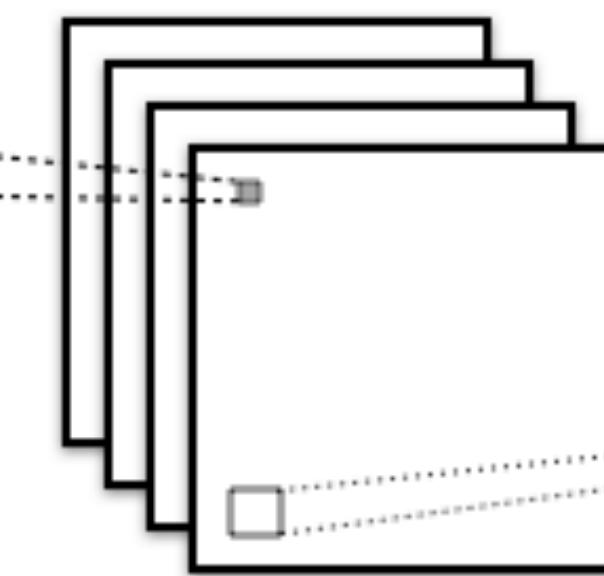
<https://medium.com/analytics-vidhya/understanding-convolution-operations-in-cnn-1914045816d4>

Convolutional Neural Networks (CNN)

How NNs are used in CV?

Solution

Add convolutions to the NN.



<https://towardsdatascience.com/build-your-own-convolution-neural-network-in-5-mins-4217c2cf964f>

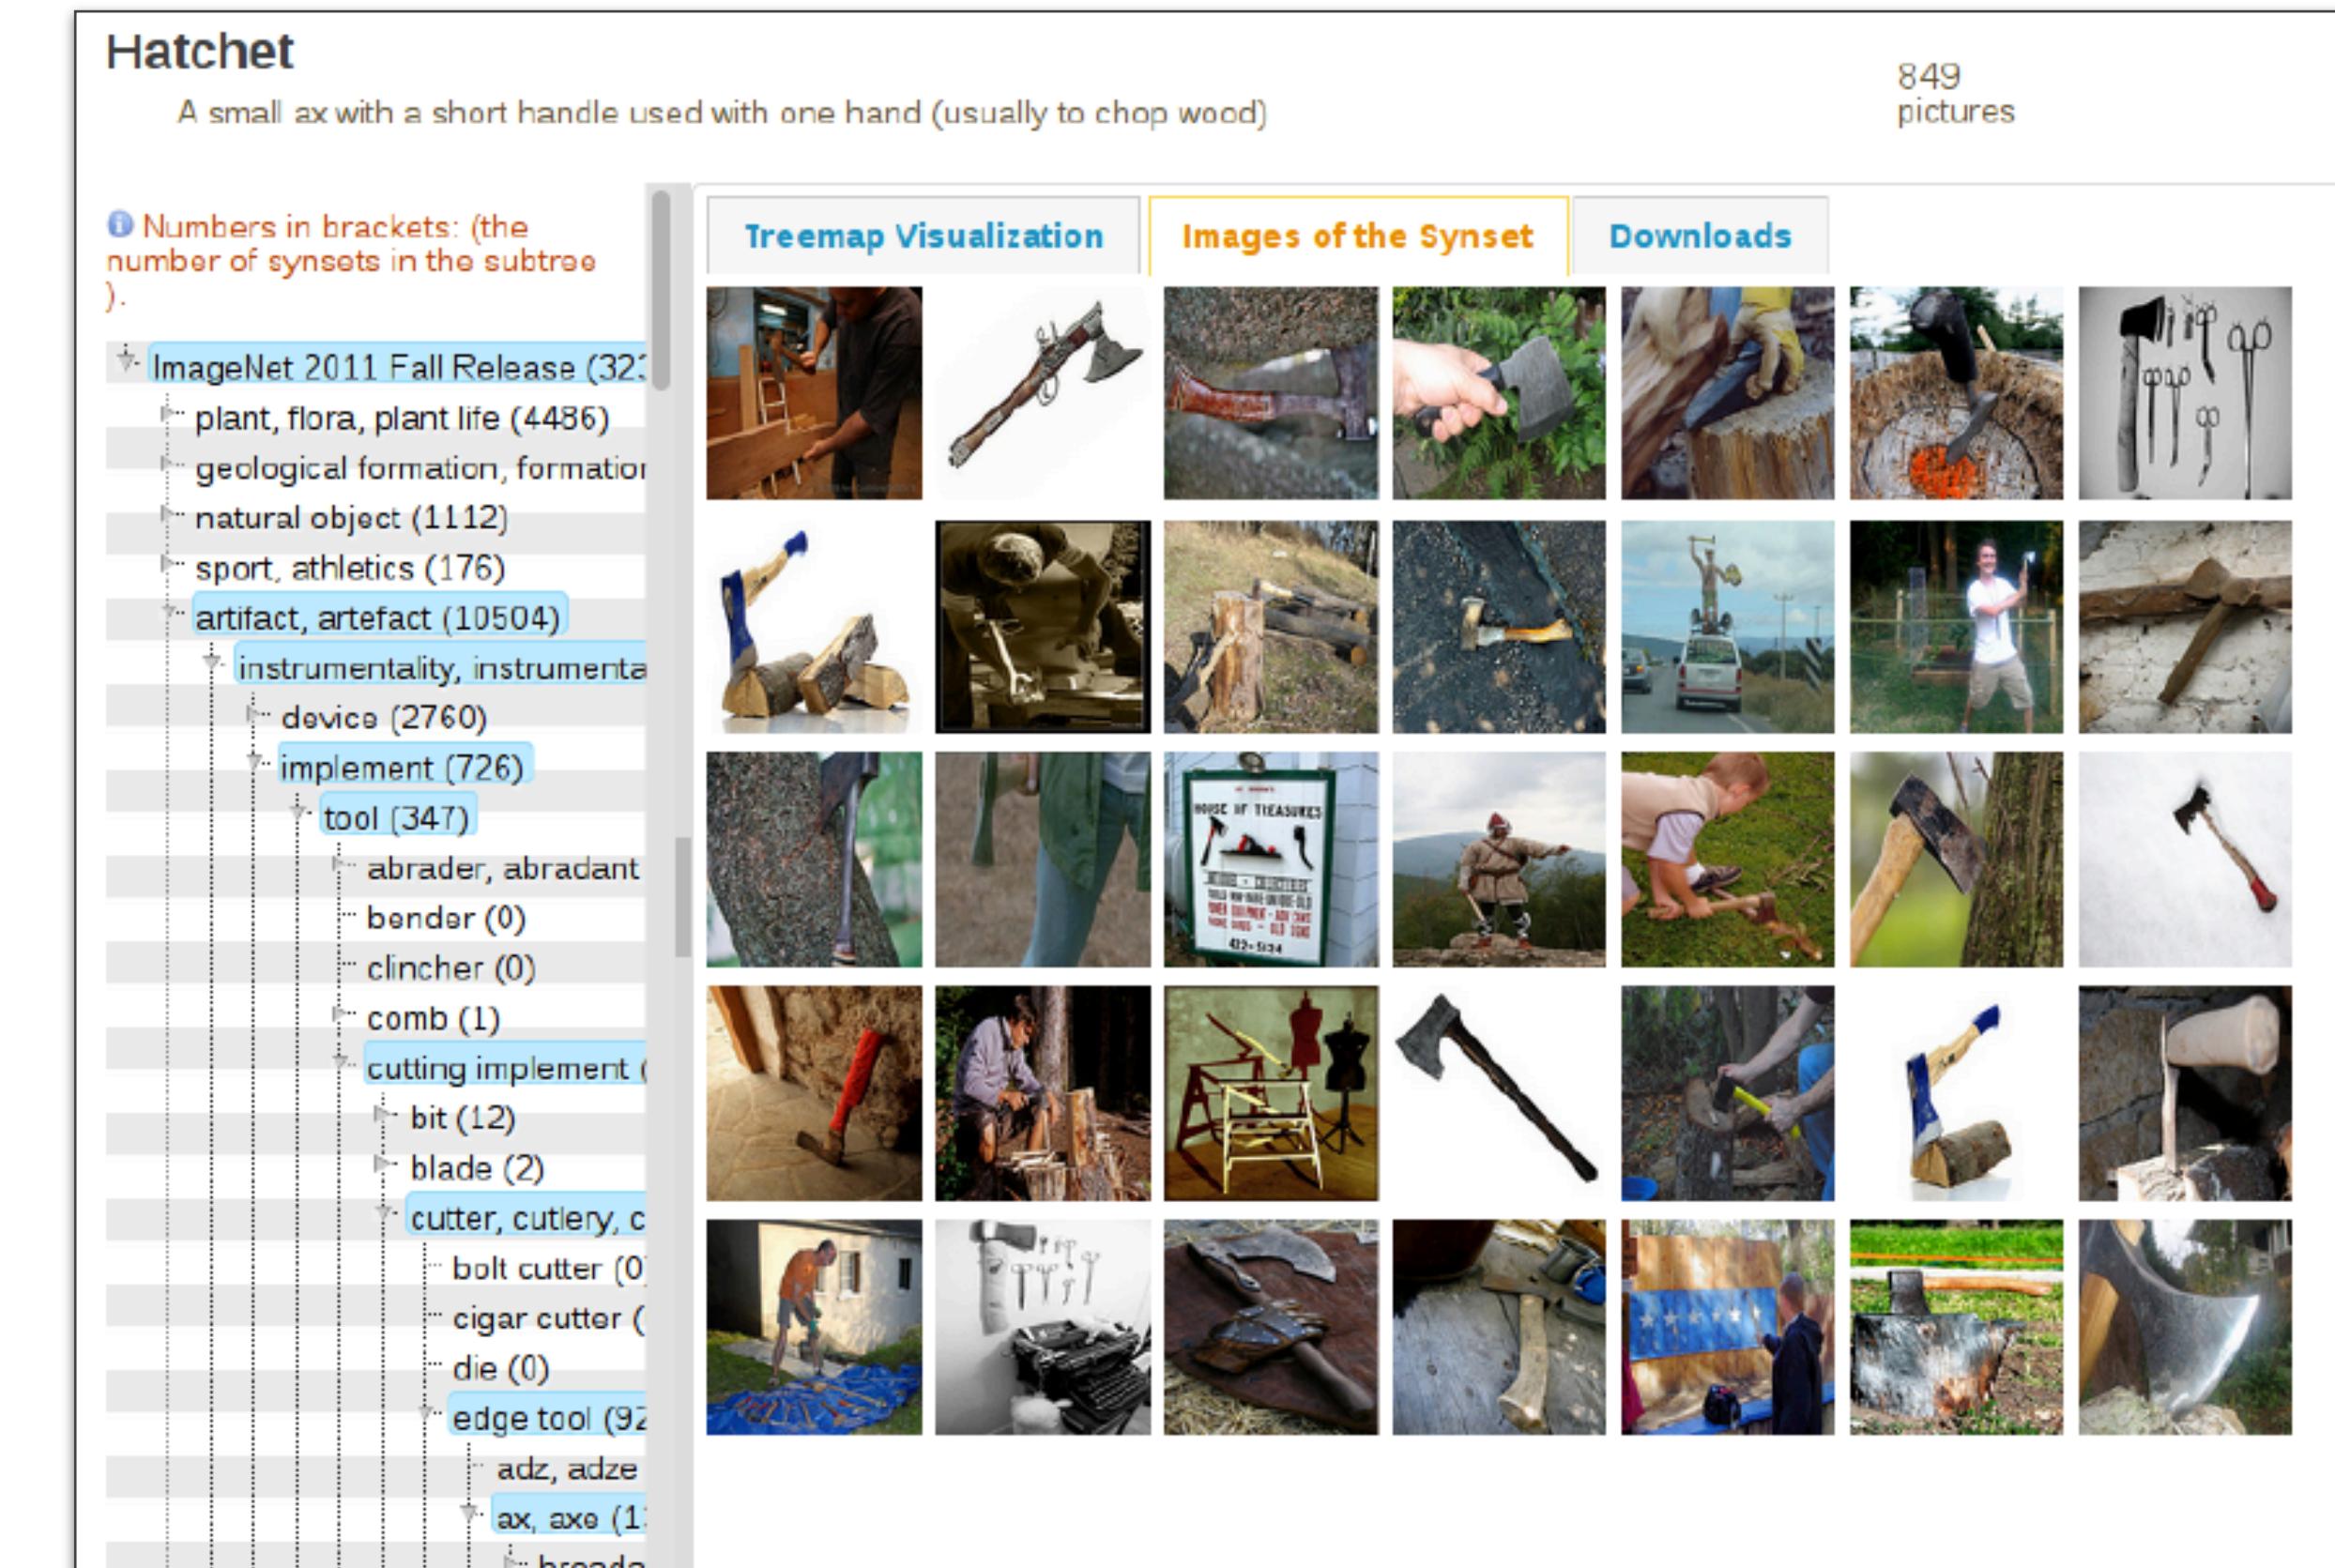
Let the NN also learn the filters.

Deep Learning (DL)

<https://fleuret.org/dlc/#lectures>

ImageNet Large Scale Visual Recognition Challenge

In 2010, 1M images,
1k categories.

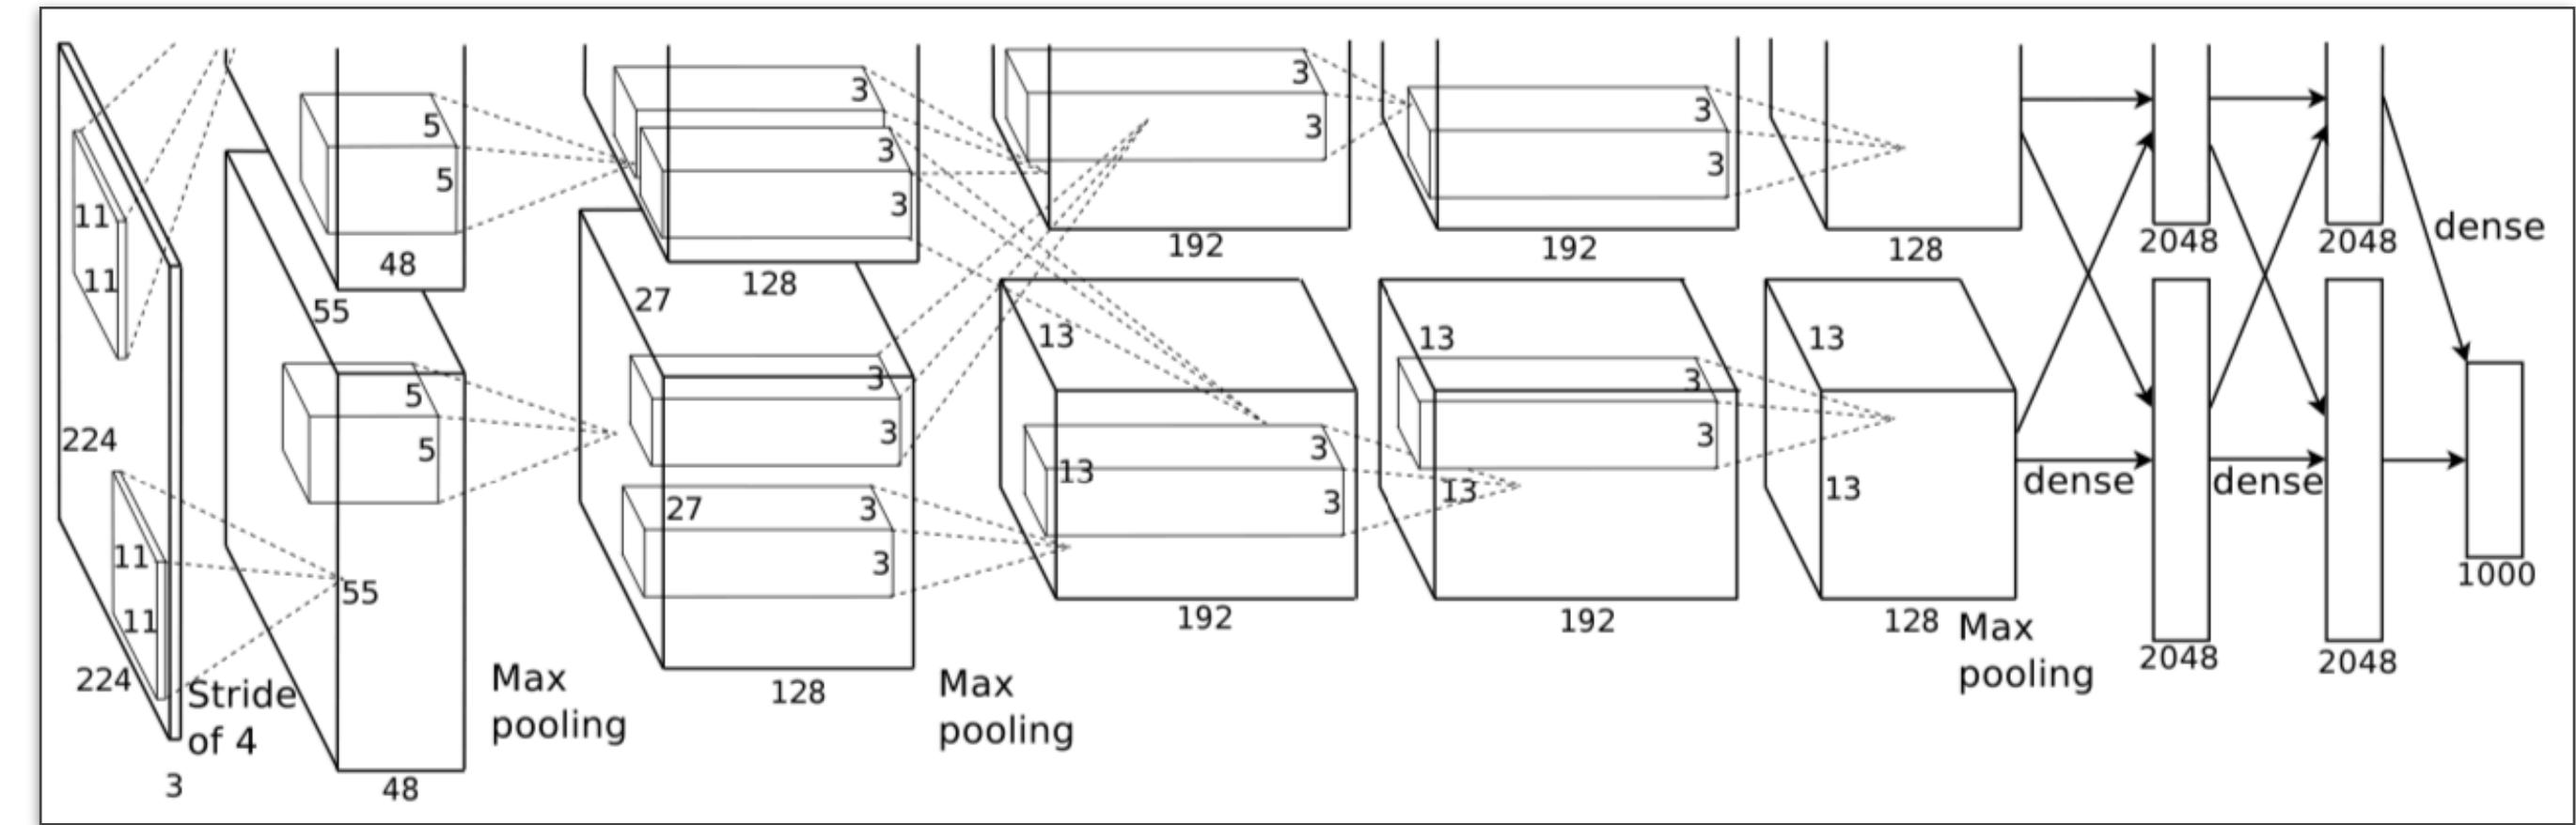


Deep Learning (DL)

AlexNet

In 2012, Krizhevsky et al. employed AlexNet to the challenge.

They used Graphical Processing Units (GPU) in the process.



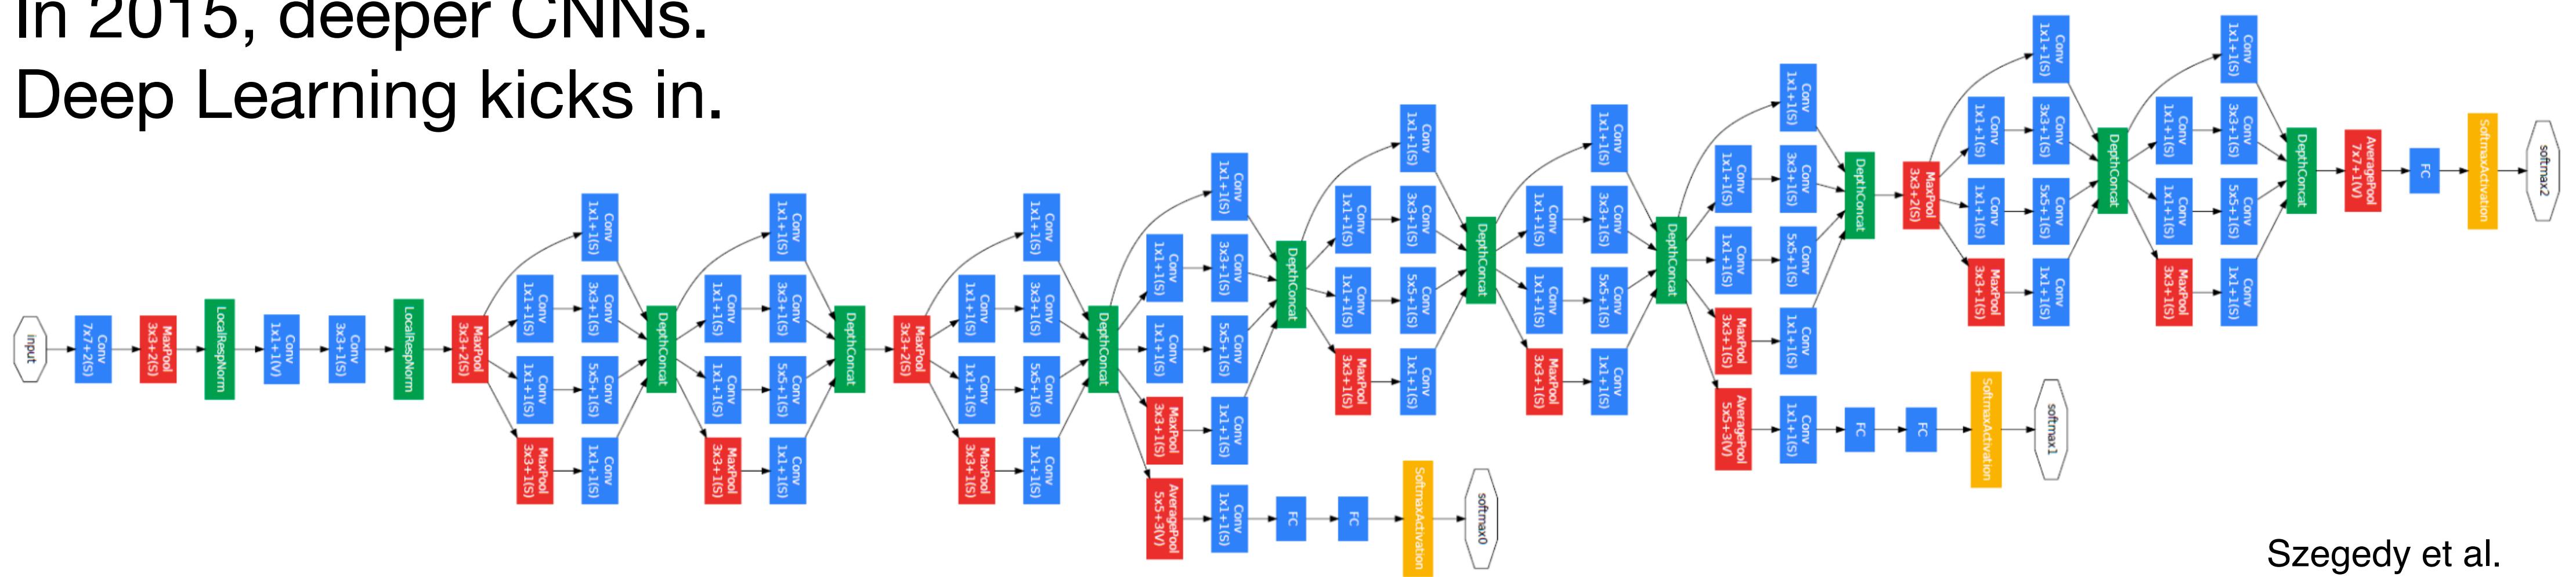
Krizhevsky et al.



Deep Learning (DL)

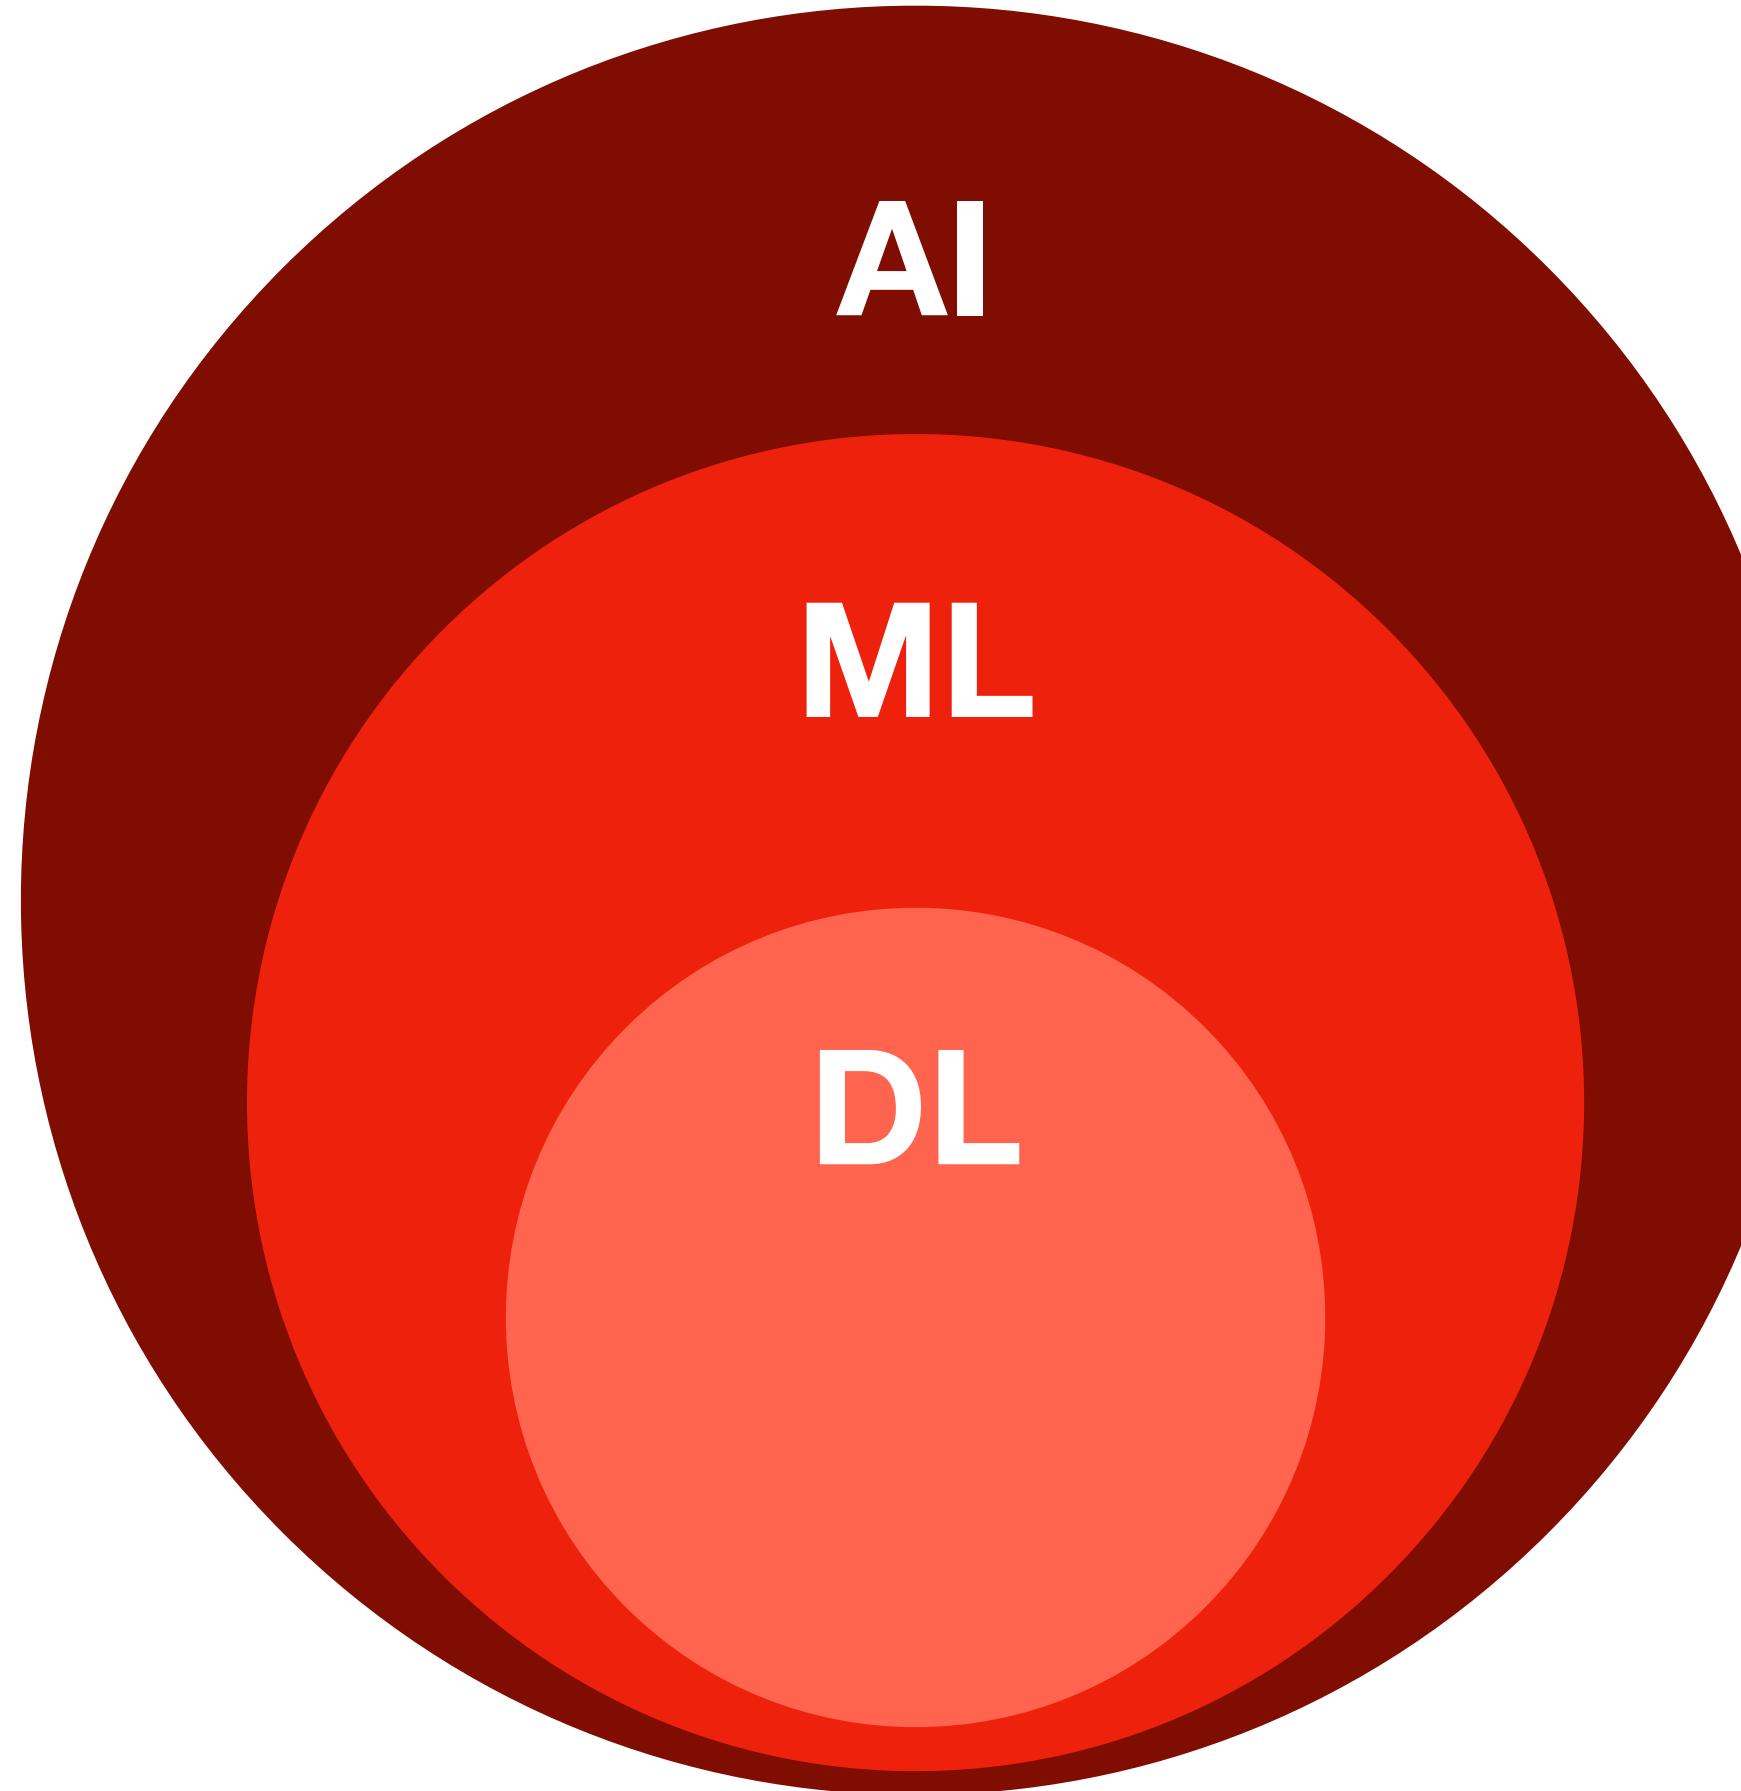
GoogleNet and Others

In 2015, deeper CNNs.
Deep Learning kicks in.

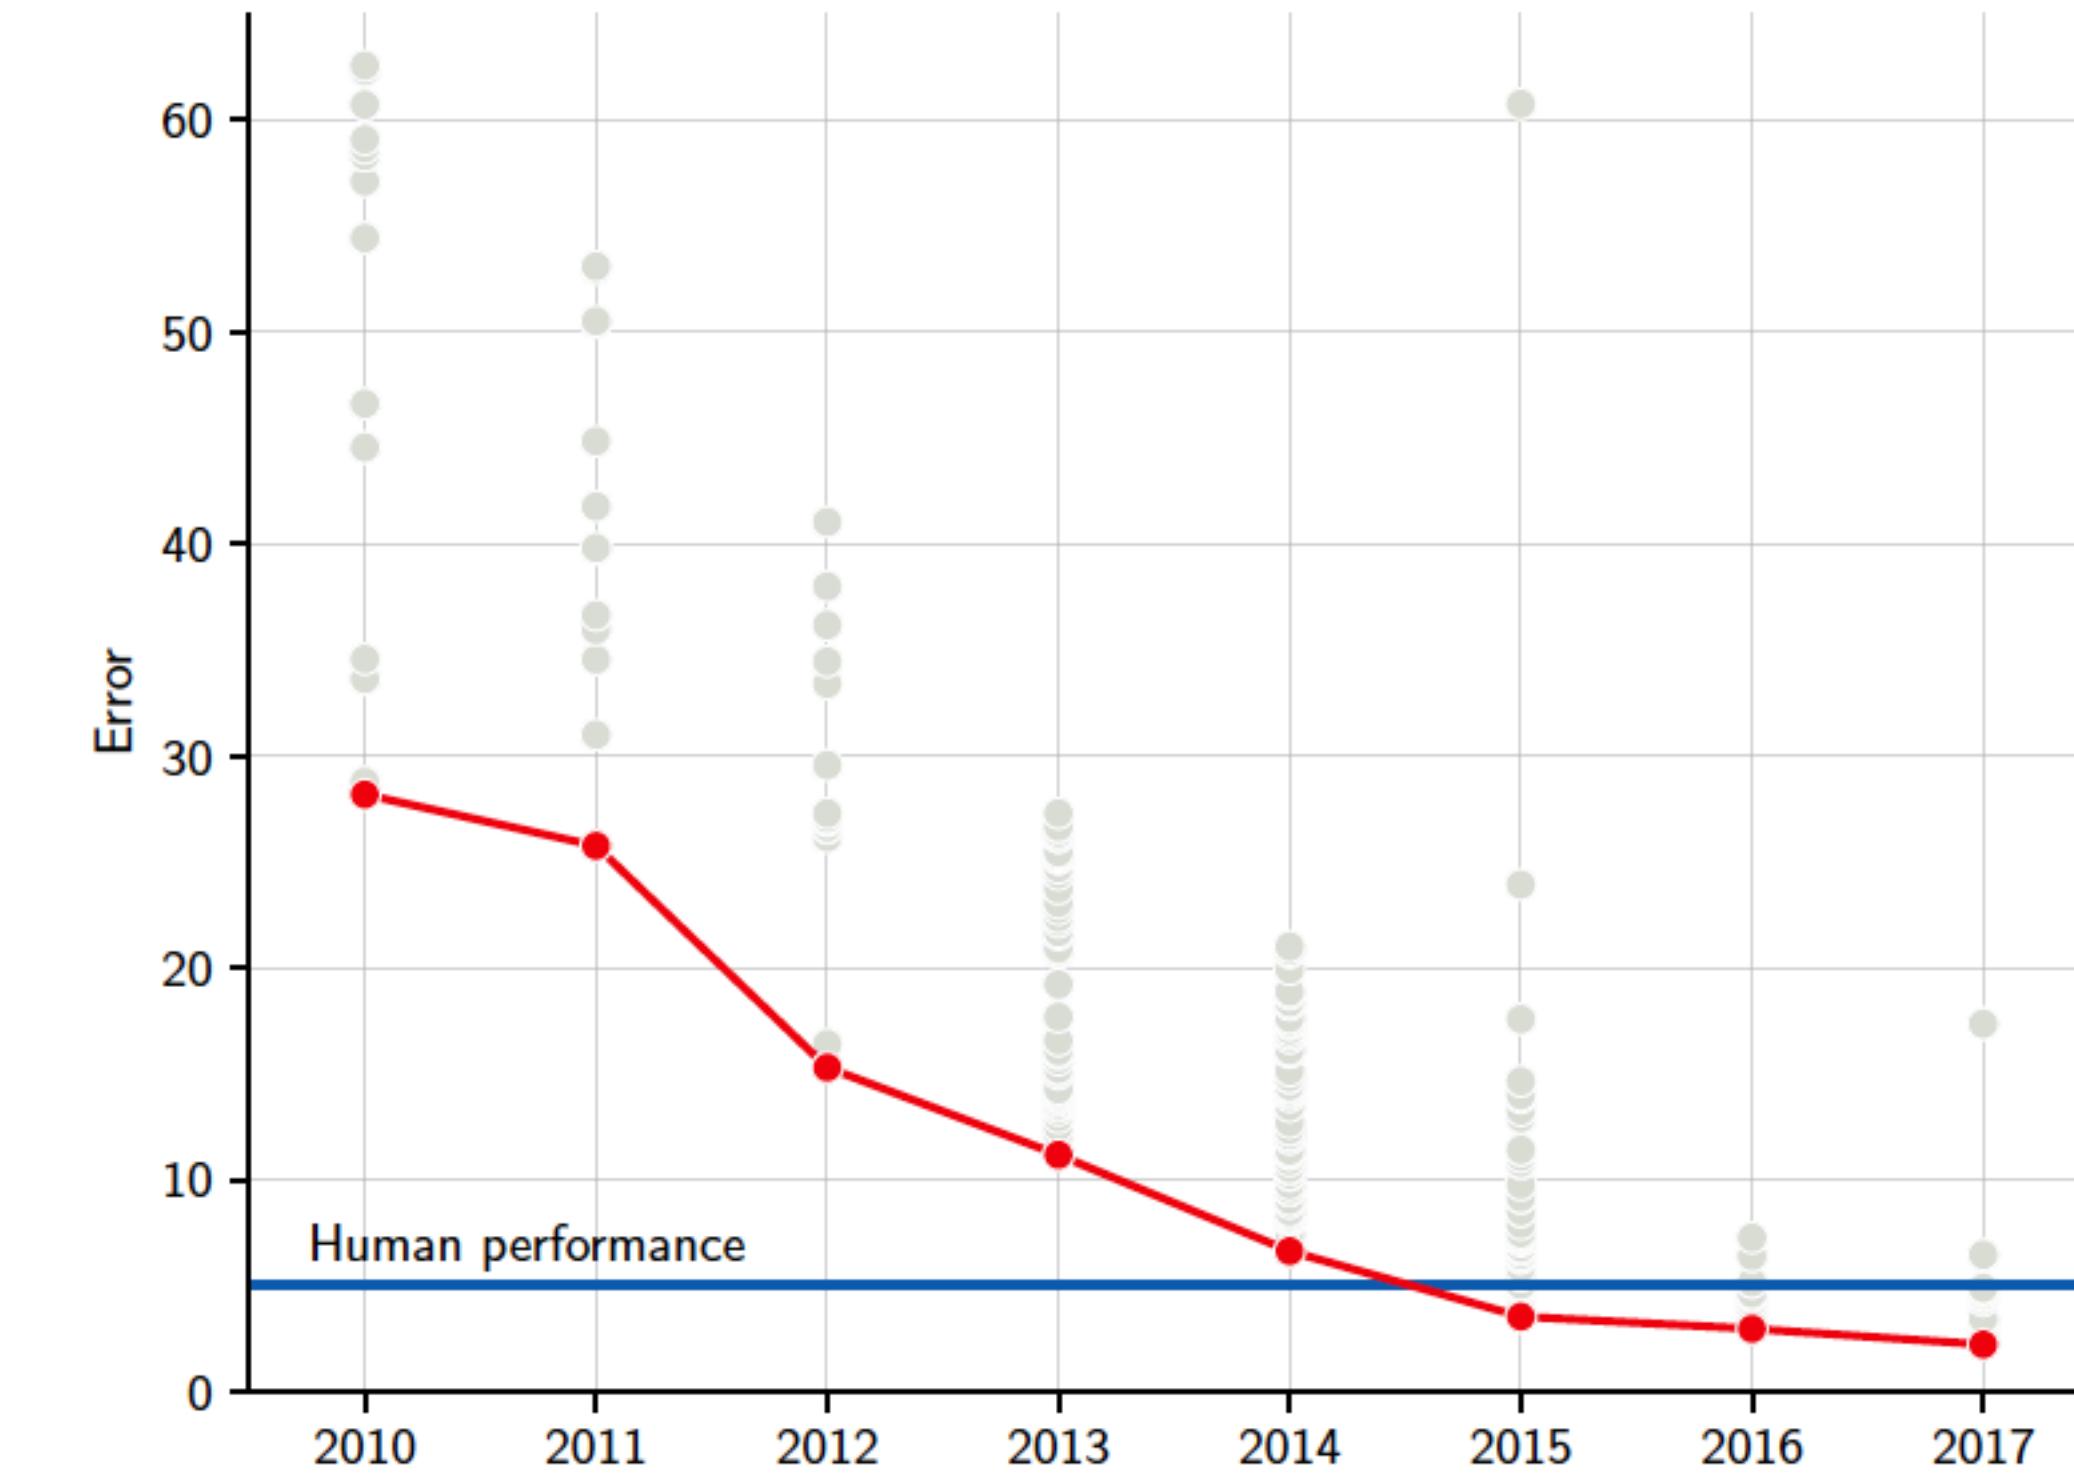


Szegedy et al.

Deep Learning (DL)

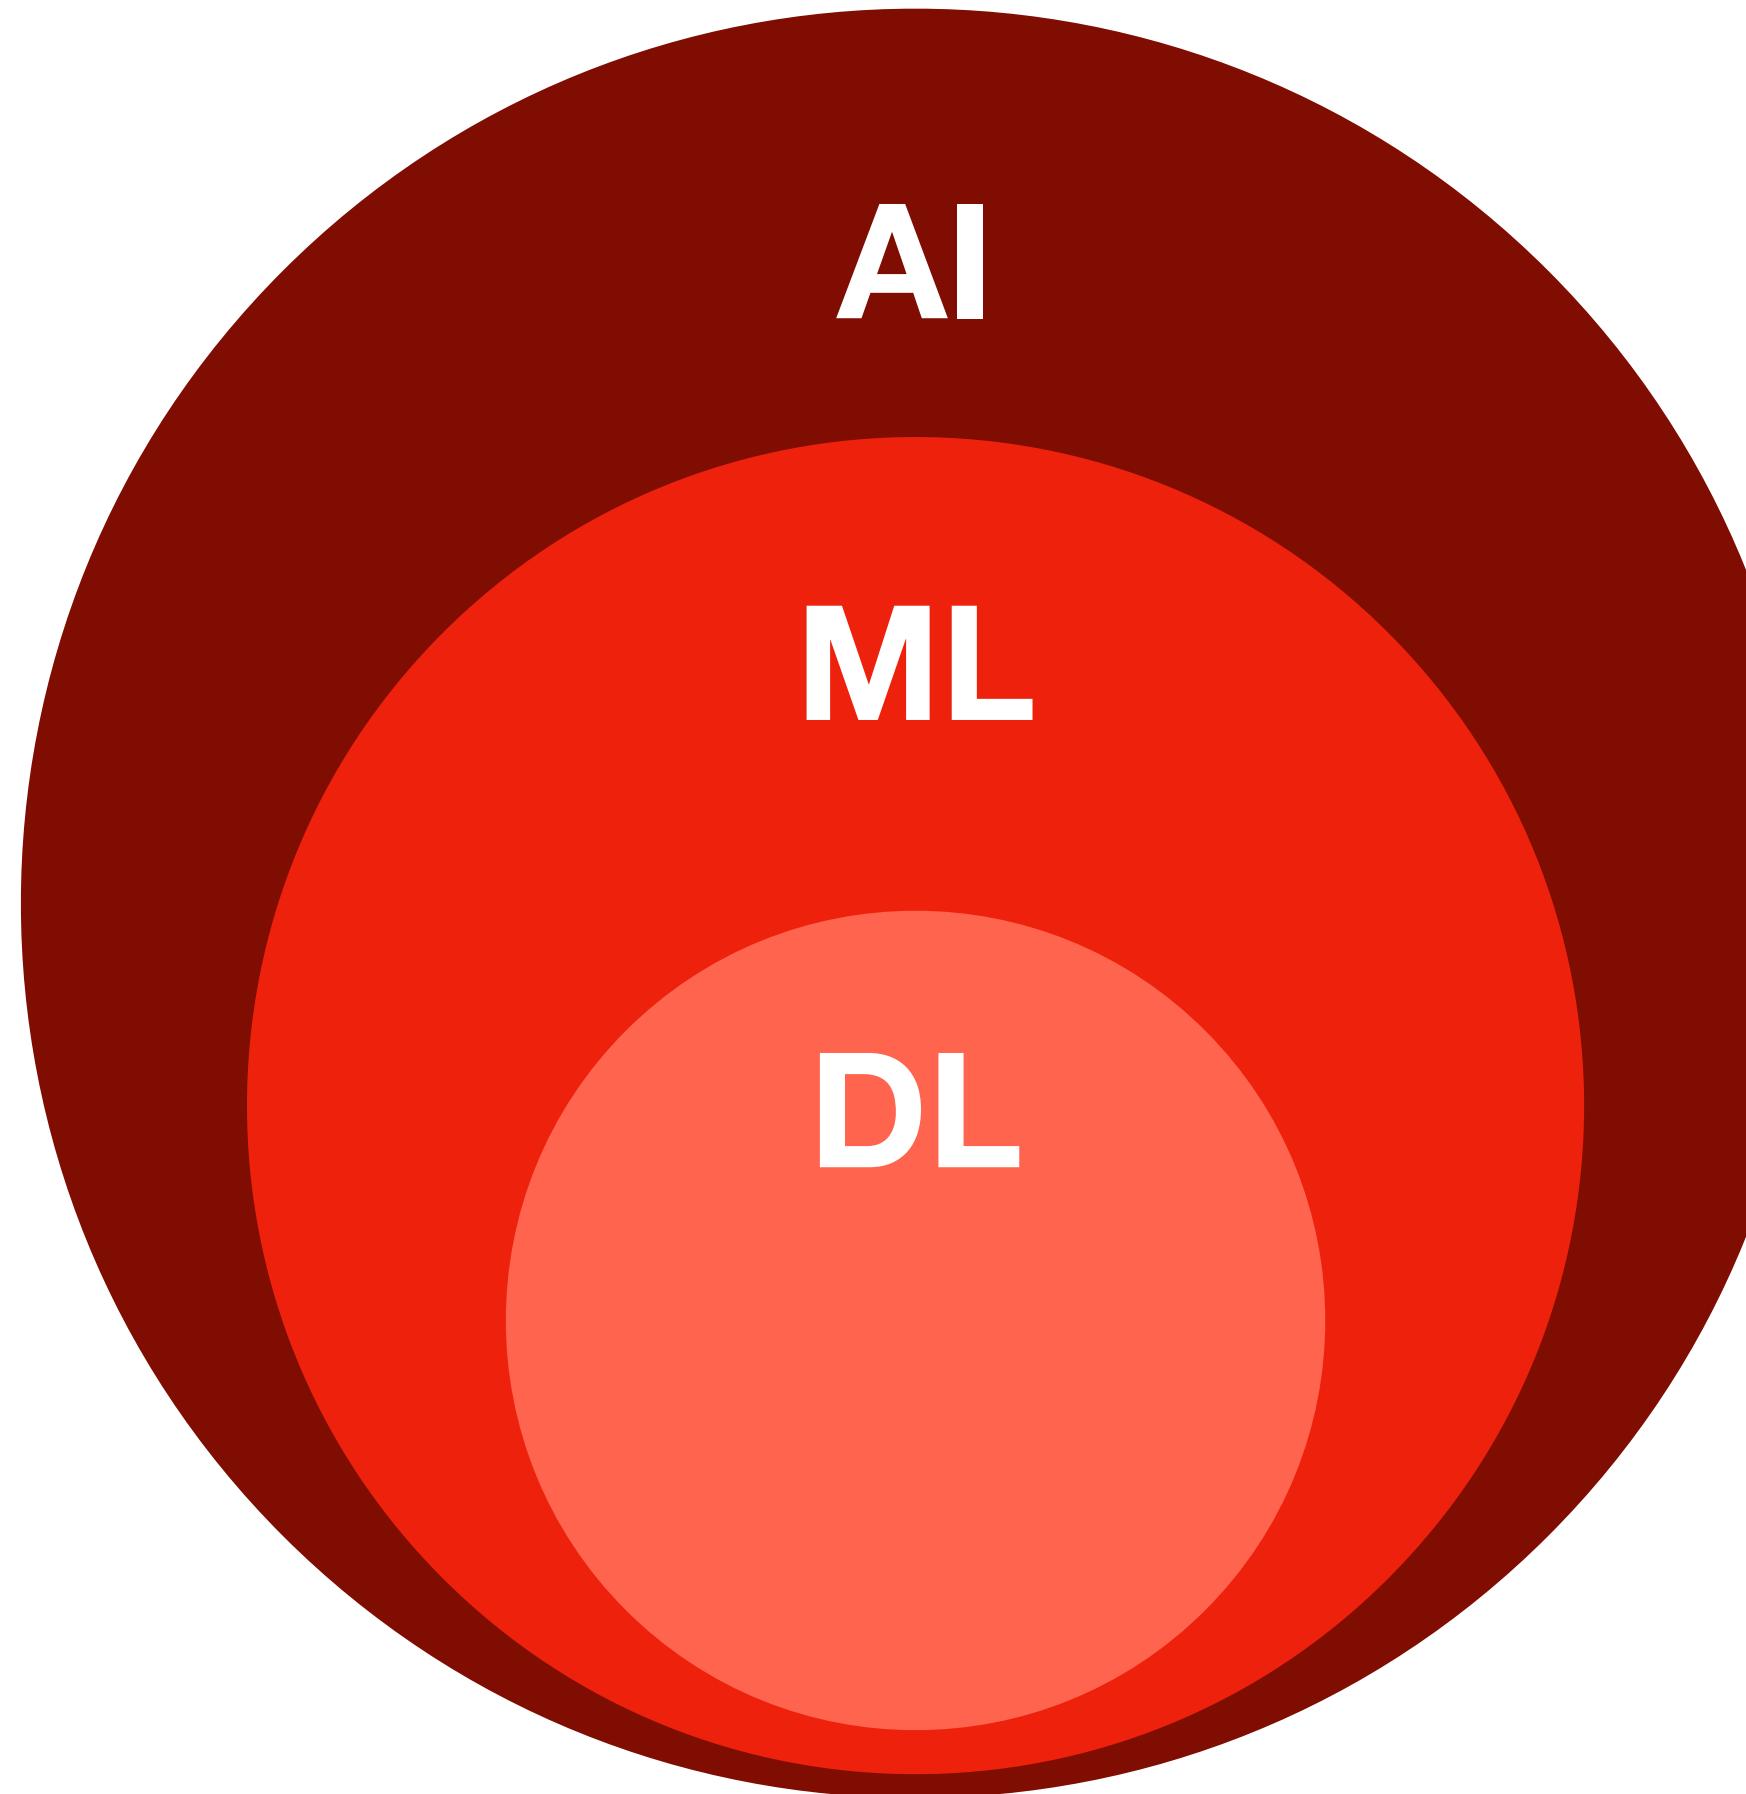


ImageNet Error Rate

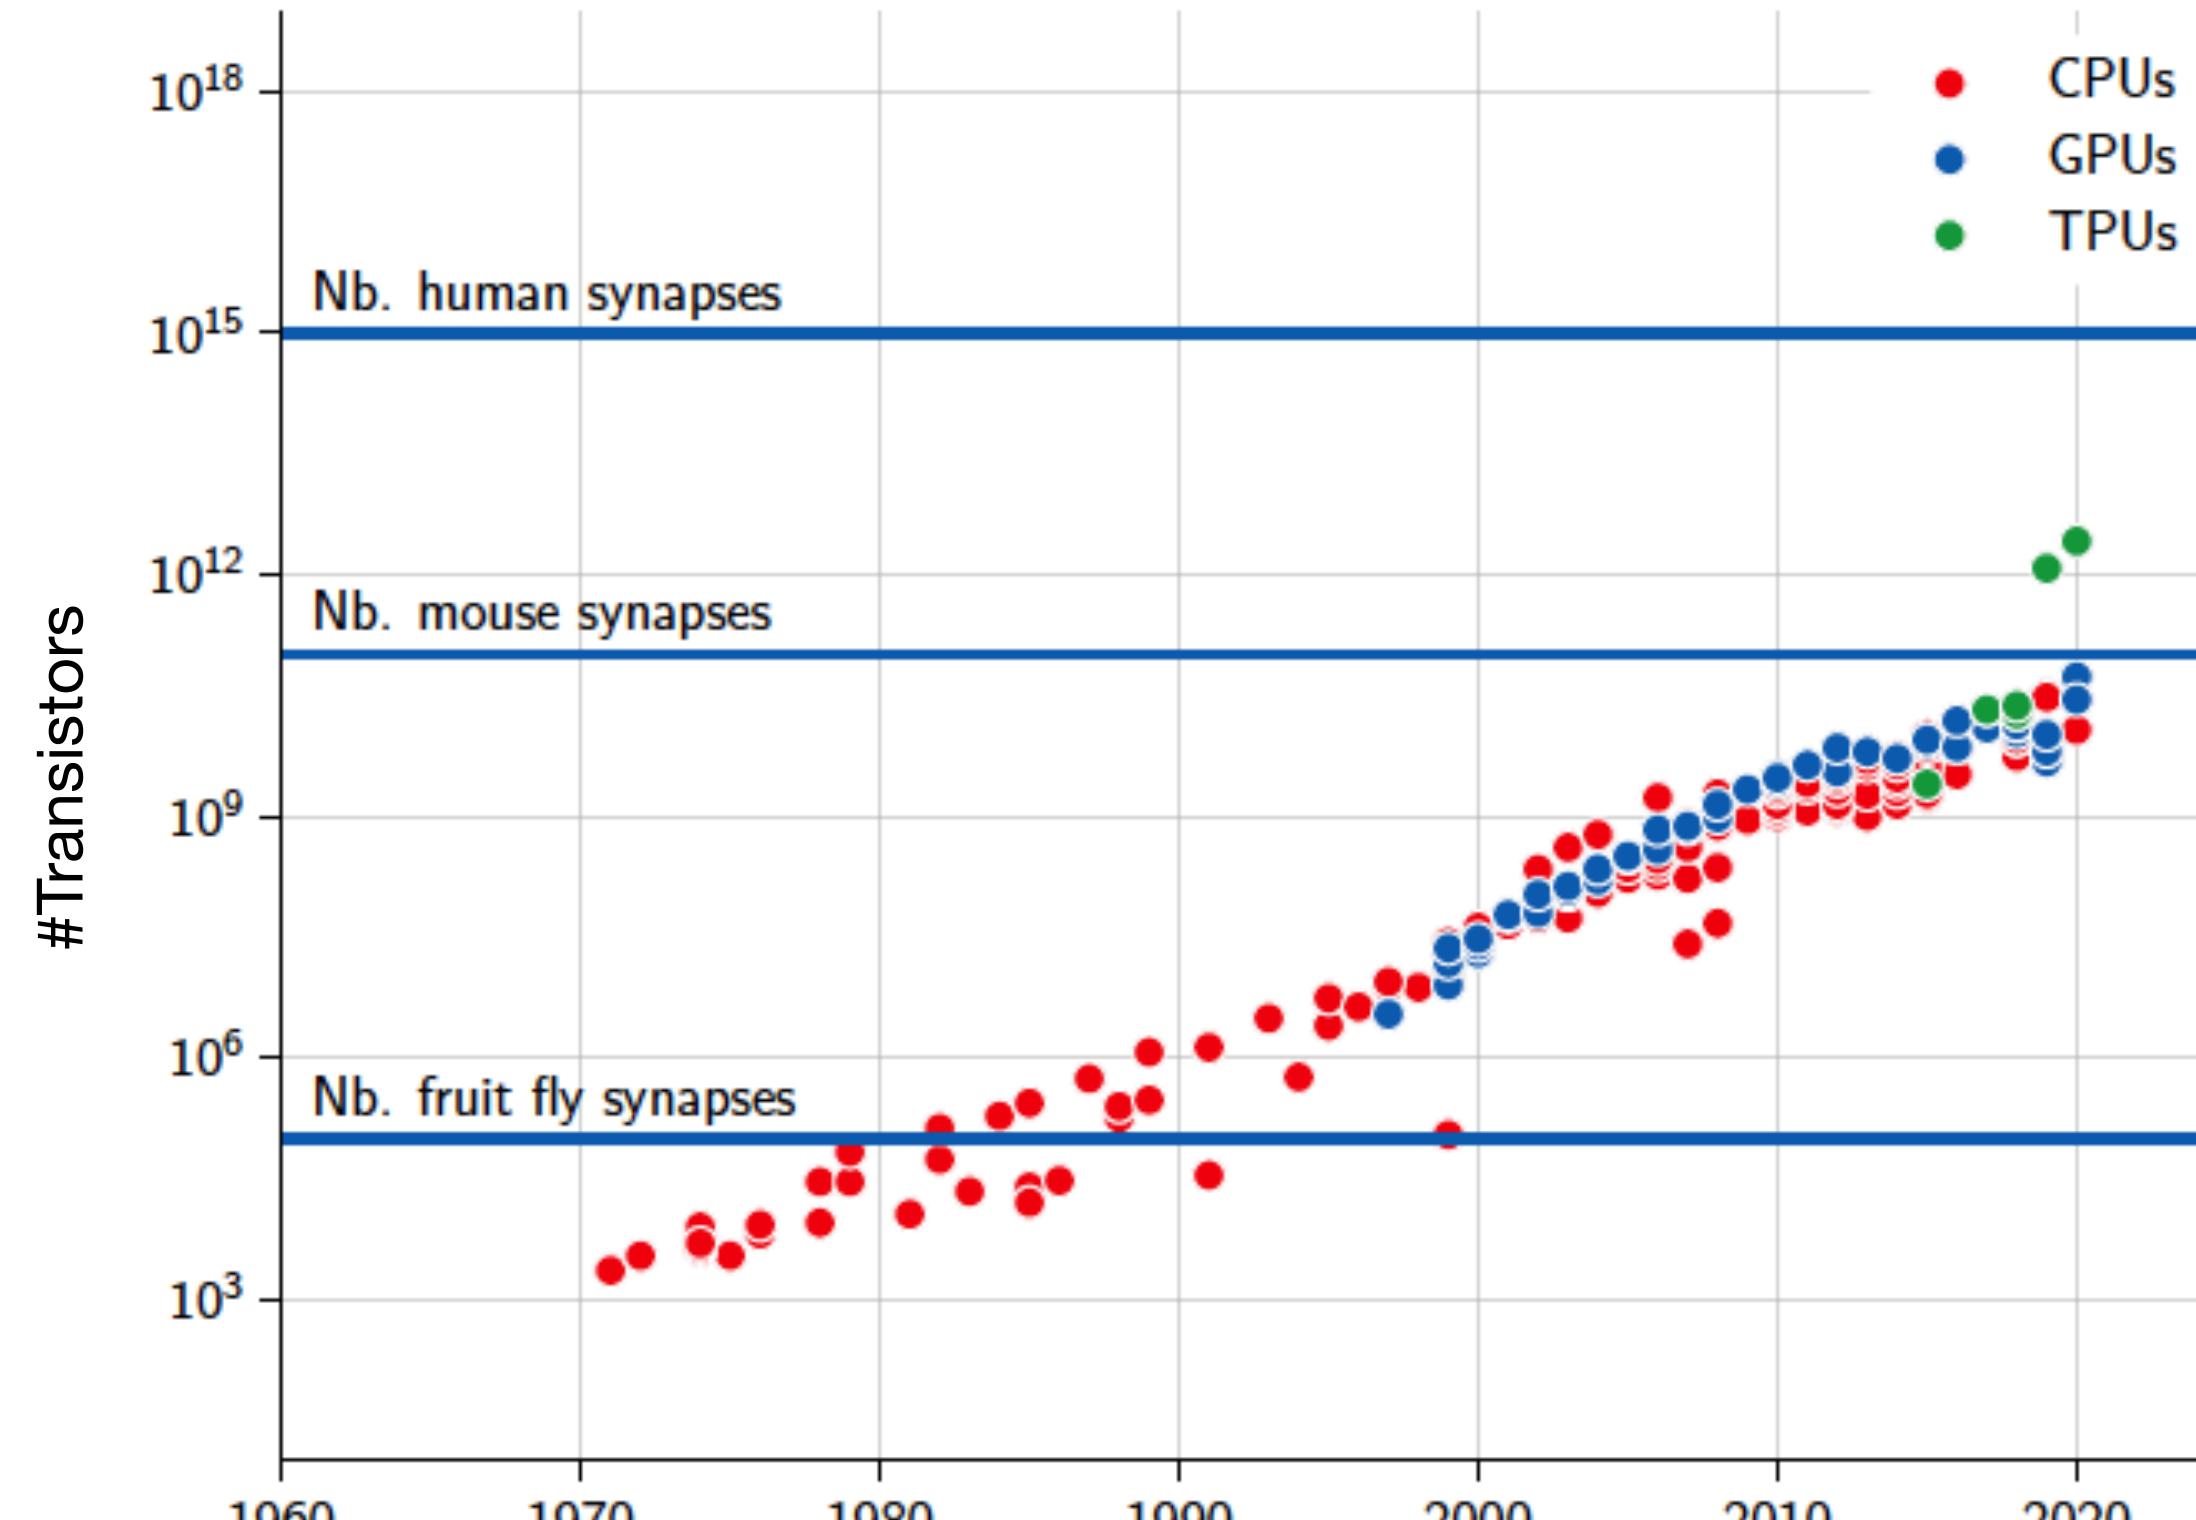


Gershgorn, 2017

Deep Learning (DL)



Number of Transistors Versus Synapses



Deep Learning (DL)

Available Libraries

	Language(s)	License	Main backer
PyTorch	Python, C++	BSD	Facebook
TensorFlow	Python, C++	Apache	Google
JAX	Python	Apache	Google
MXNet	Python, C++, R, Scala	Apache	Amazon
CNTK	Python, C++	MIT	Microsoft
Torch	Lua	BSD	Facebook
Theano	Python	BSD	U. of Montreal
Caffe	C++	BSD 2 clauses	U. of CA, Berkeley

<https://fleuret.org/dlc/#lectures>



Deep Learning (DL)

Pros and Cons

Pros

Deep NNs are powerful tools.
They may have tens of millions
of degrees of freedom.

They can approximate any
continuous function and be fed
with annotated digital images.

Cons

They are data hungry.
They need a massive amount of data
to be trained.

They may become black boxes
with hard inner understanding.
Why are they working or failing?



Be Careful

August 14, 2019

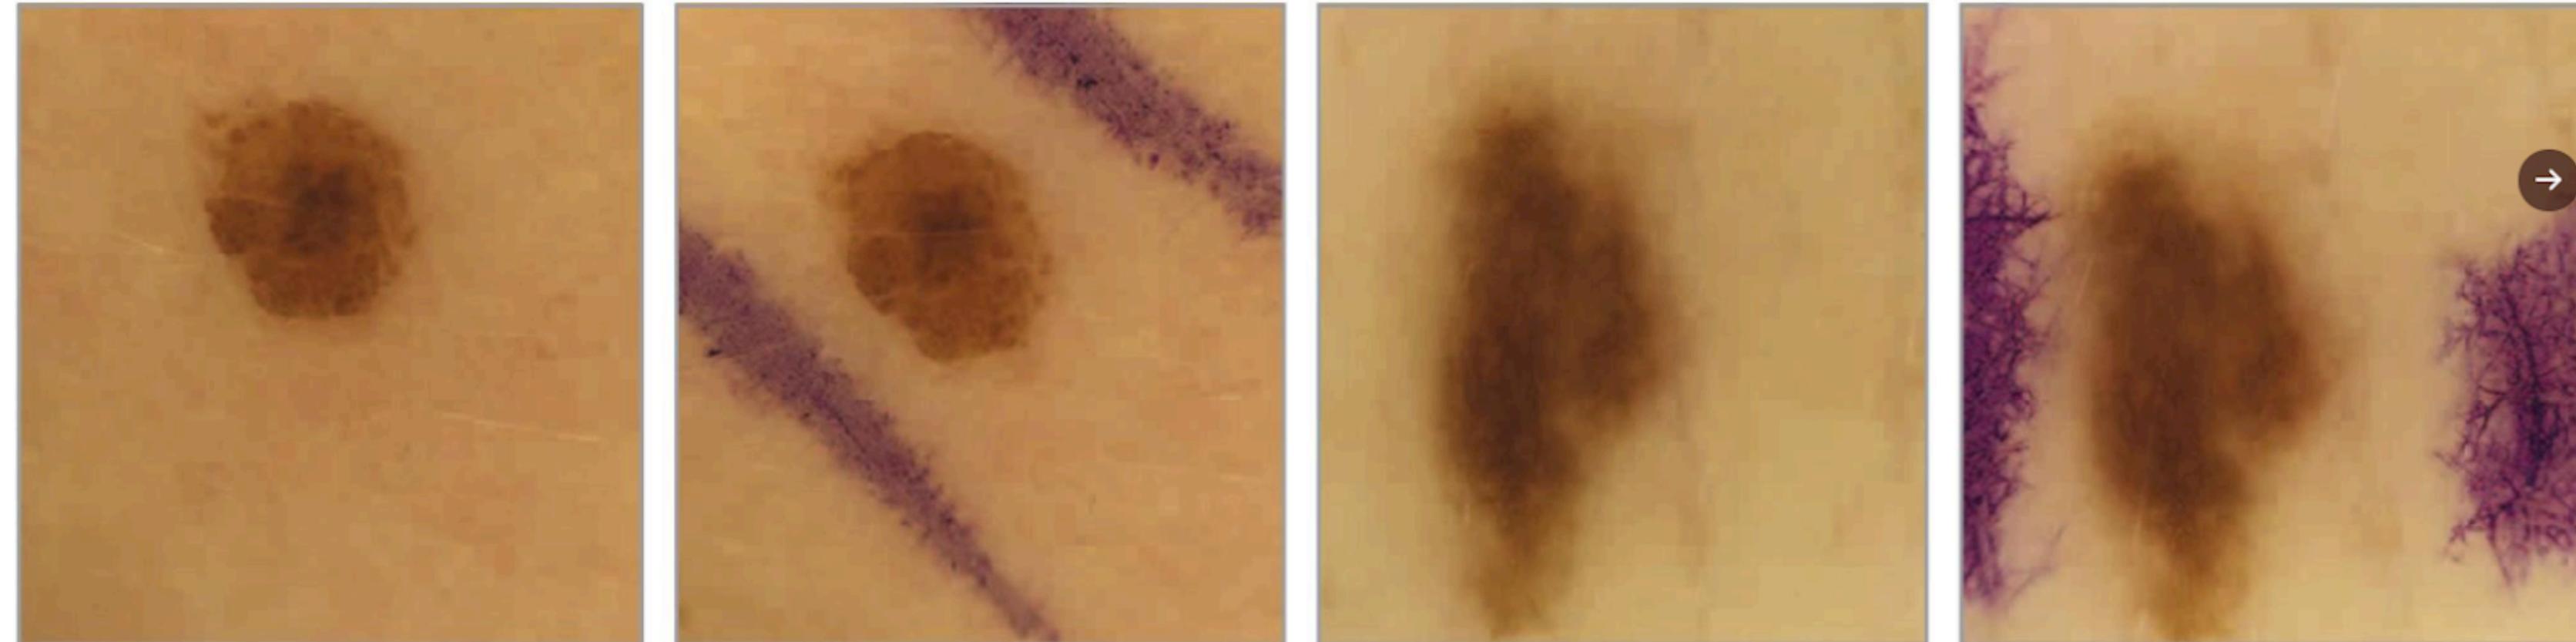
Association Between Surgical Skin Markings in Dermoscopic Images and Diagnostic Performance of a Deep Learning Convolutional Neural Network for Melanoma Recognition

Julia K. Winkler, MD¹; Christine Fink, MD¹; Ferdinand Toberer, MD¹; [et al](#)

[» Author Affiliations](#) | [Article Information](#)

JAMA Dermatol. 2019;155(10):1135-1141. doi:10.1001/jamadermatol.2019.1735

What is the network learning?



What's Next?

Next Class

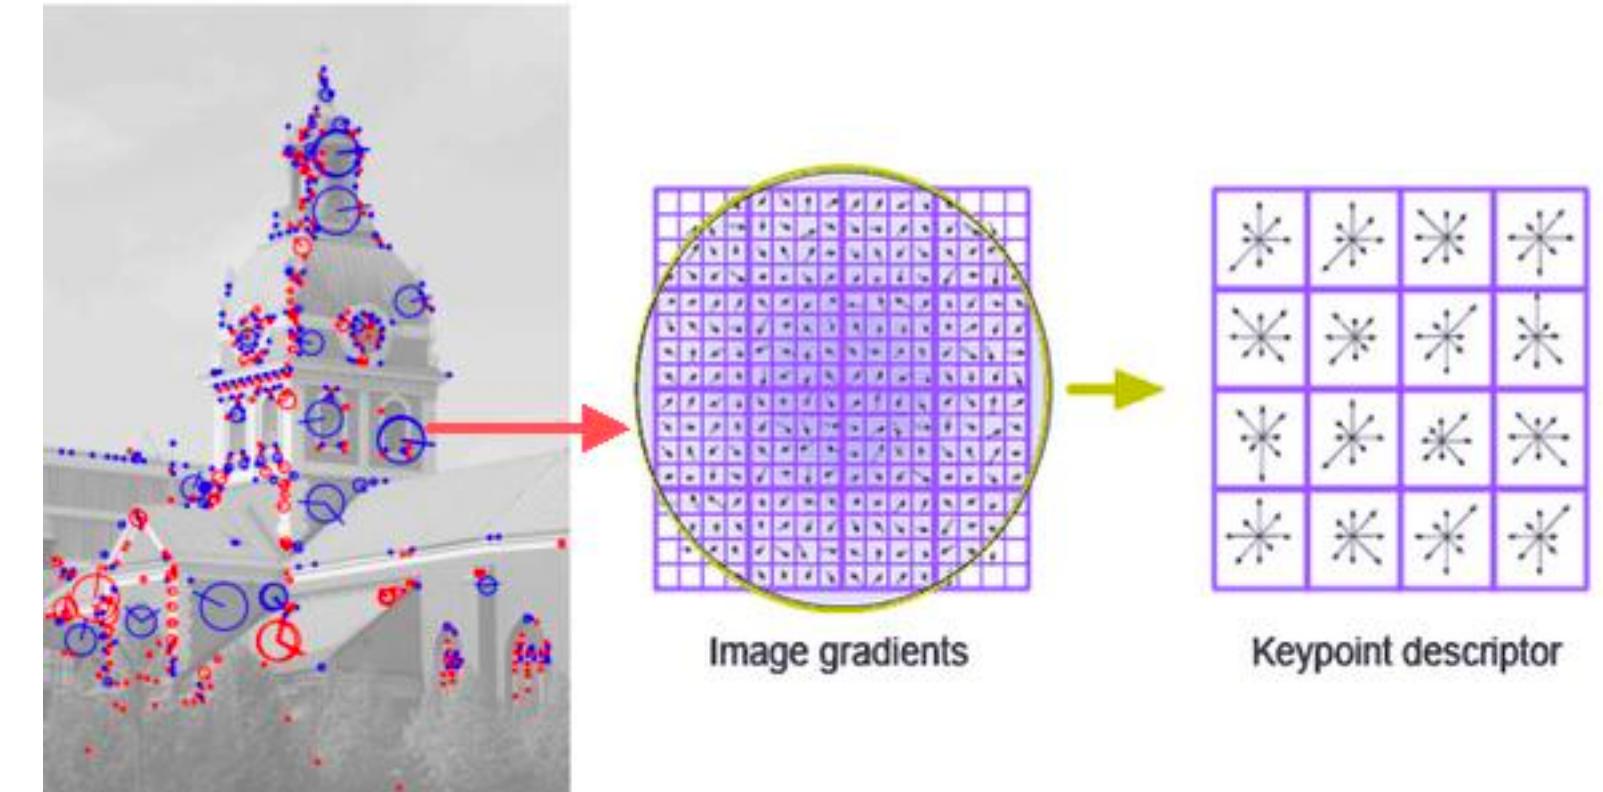
Local and Global Image Descriptors
(Daniel's presentation).

Sakai is up!

The assignments, content of the classes,
and reference papers are being posted there.

Start working on your 2 seminars

I'll announce today the groups based on your answers.
Count on me during office hours (and outside of them)
to prepare your seminars.



[https://www.codeproject.com/Articles/619039/
Bag-of-Features-Descriptor-on-SIFT-Features-with-O](https://www.codeproject.com/Articles/619039/Bag-of-Features-Descriptor-on-SIFT-Features-with-O)