

Computer Vision Applications

COMP 388-002/488-002 Computer Science Topics

Daniel Moreira
Fall 2022



LOYOLA
UNIVERSITY CHICAGO

Image Description

COMP 388-002/488-002 Computer Science Topics
Computer Vision Applications

Daniel Moreira
Fall 2022



LOYOLA
UNIVERSITY CHICAGO

Today you will...

Get to know global and local
image description.

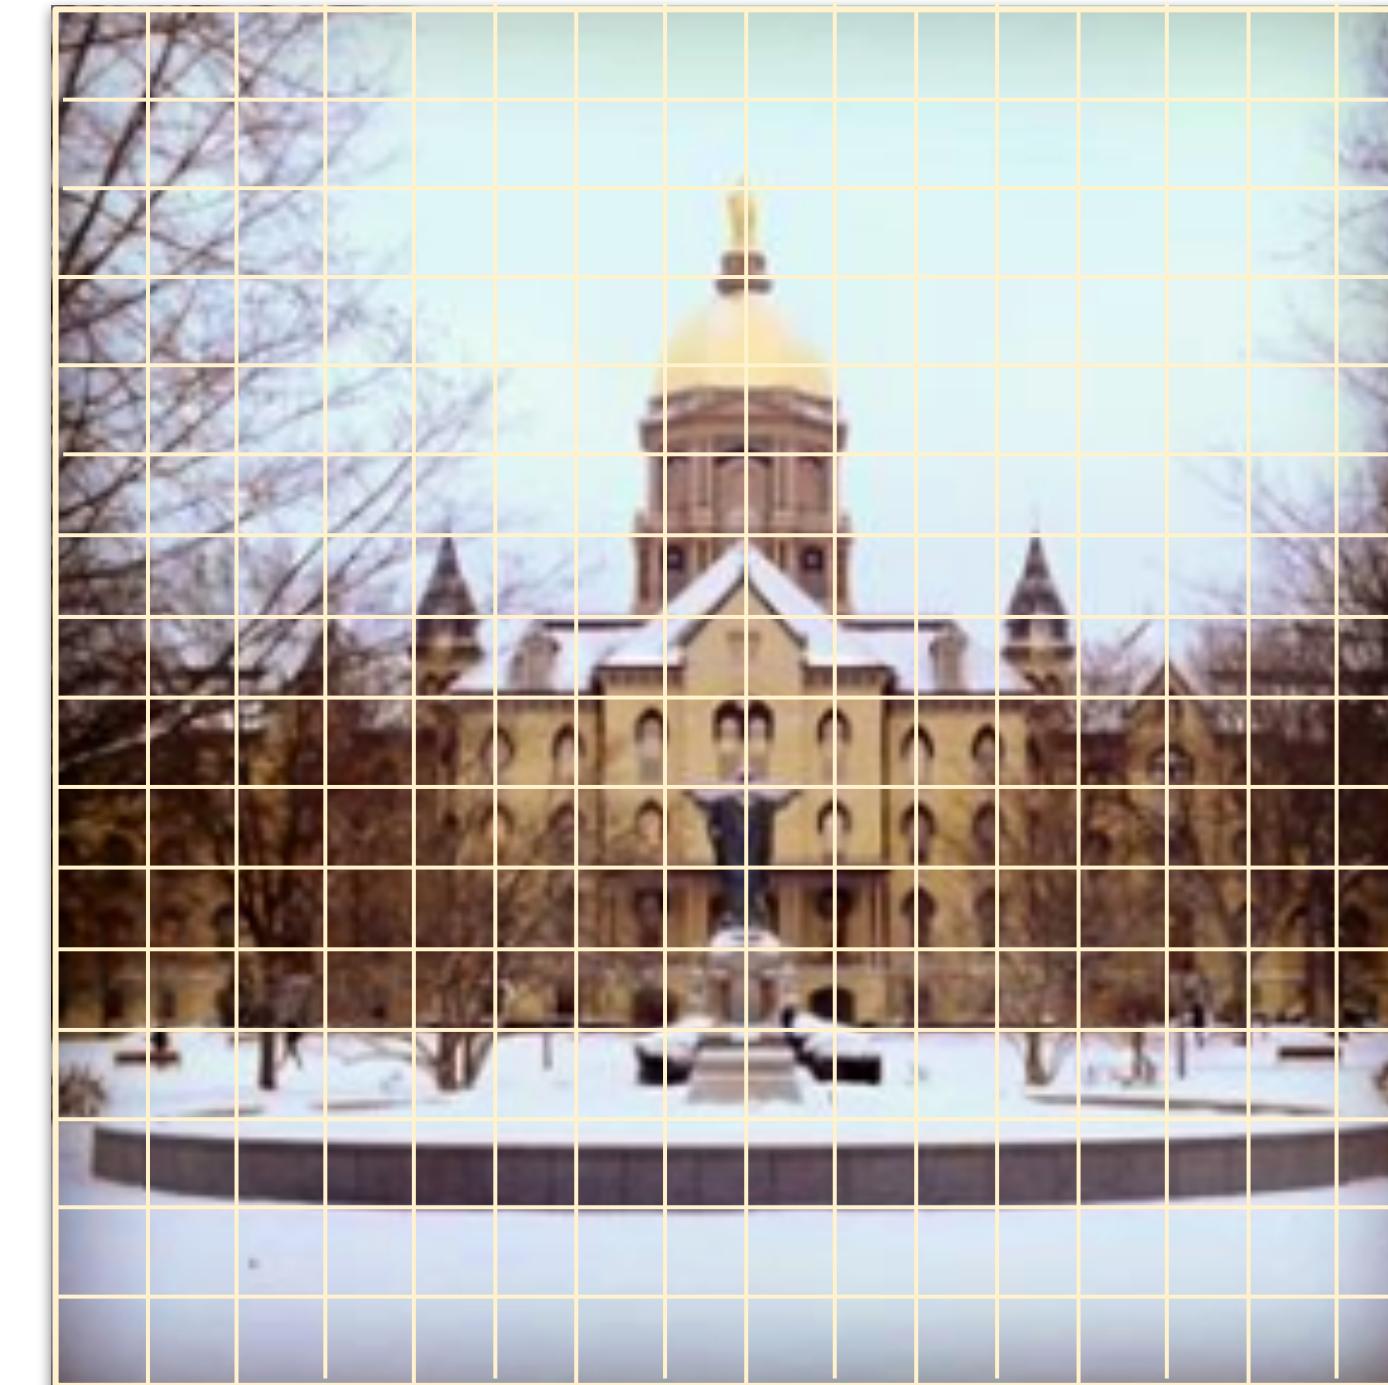
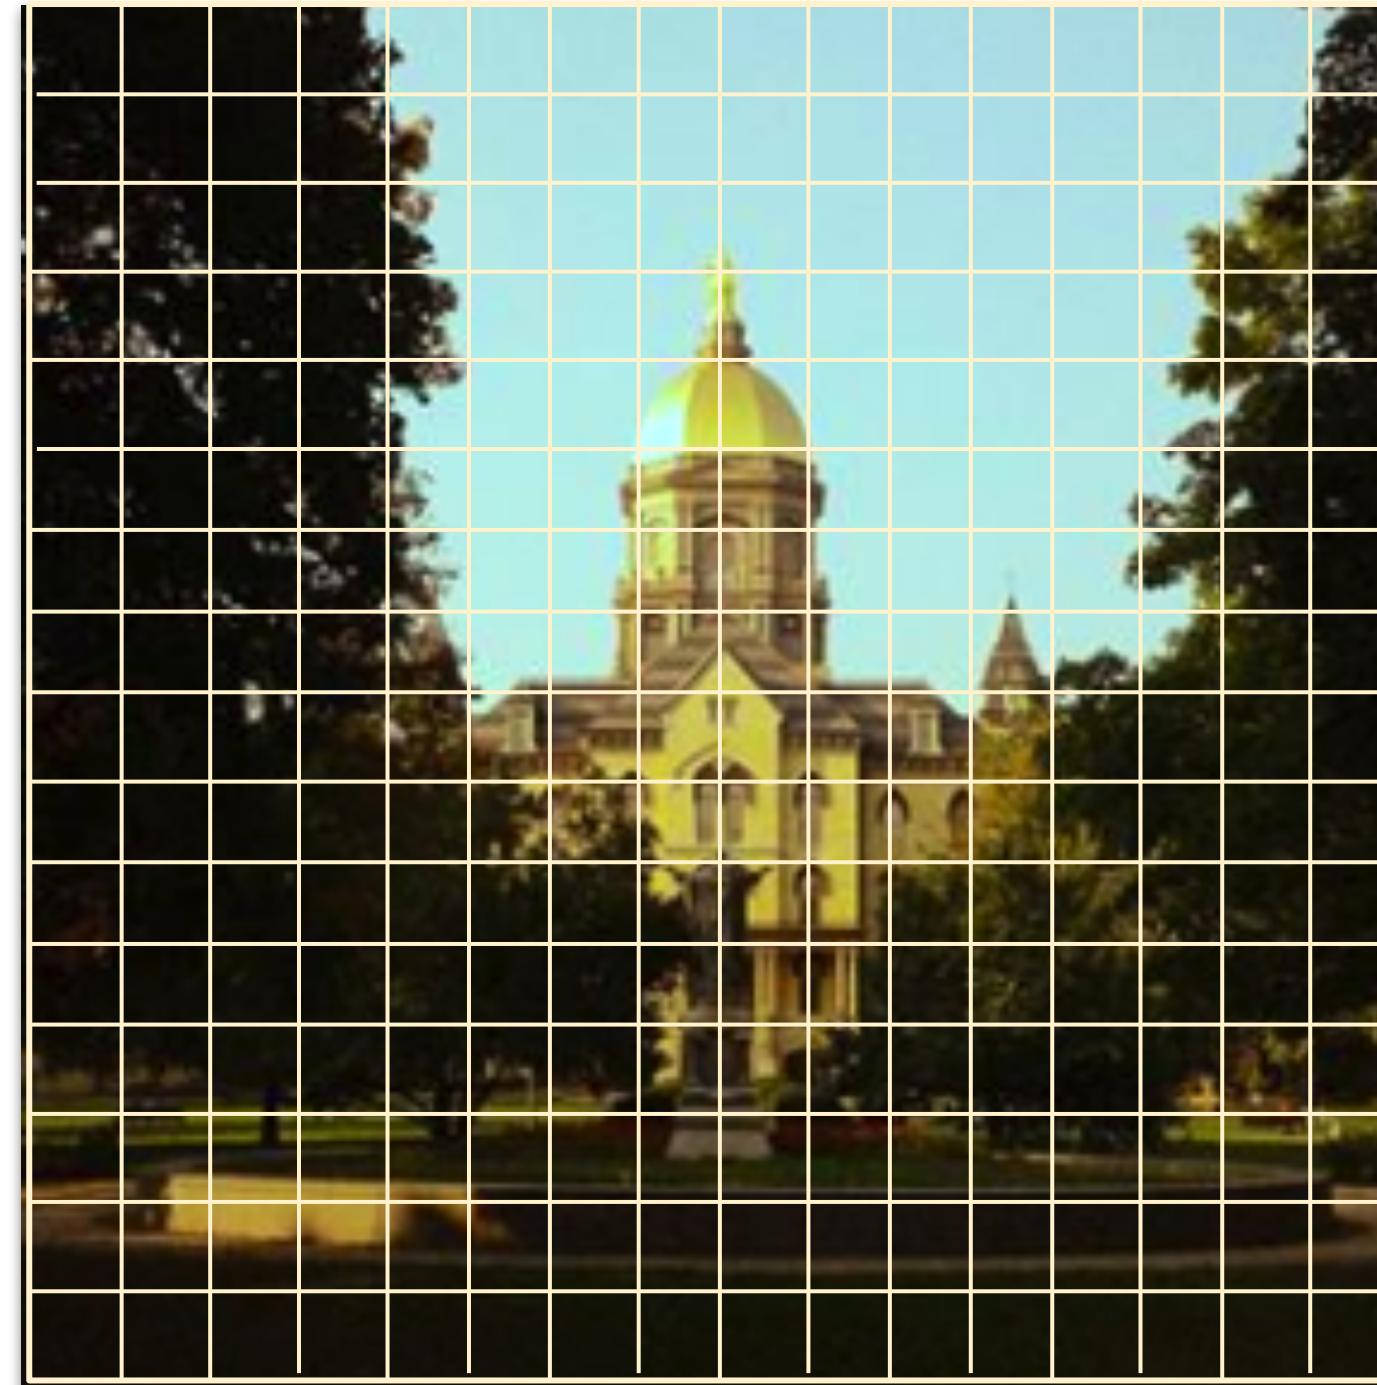


Why do We Need Image Description?



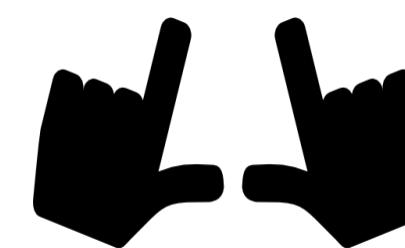
Can a computer system
decide if these images
depict the same building?

Why do We Need Image Description?



Can a computer system
decide if these images
depict the same building?

Yes, but not directly
based on the pixel values.



capture
position



date and
time



device

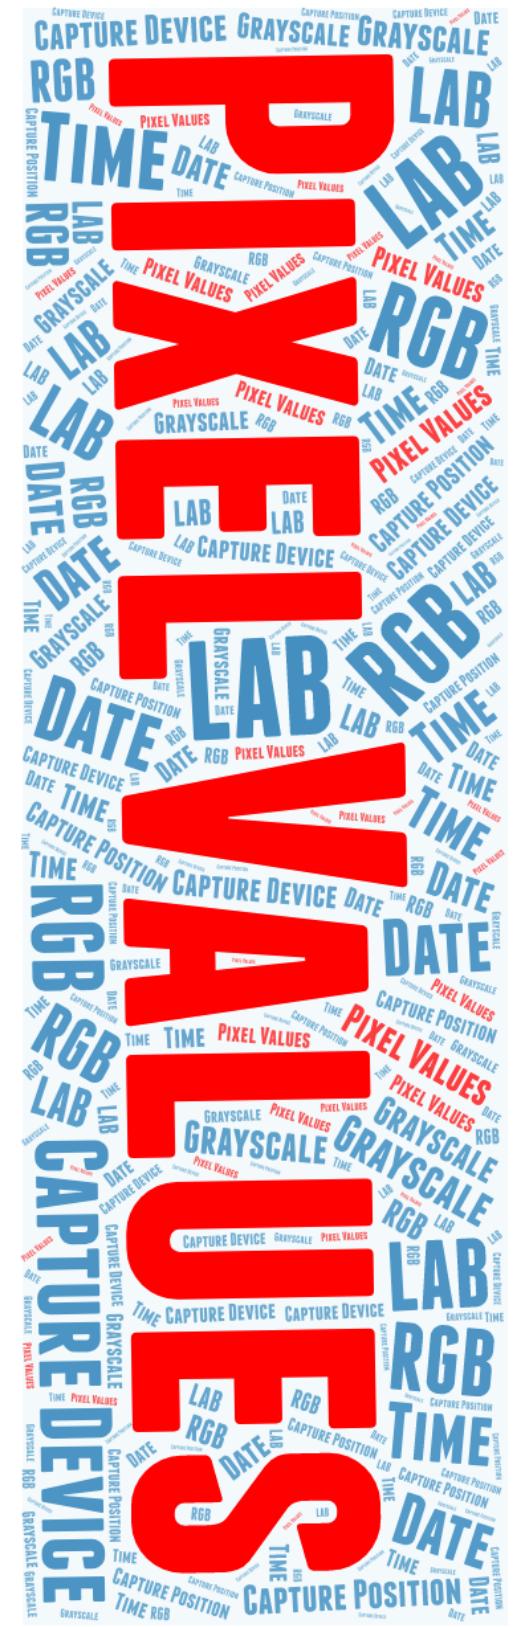
Pixel values depend on complex
settings.



LOYOLA
UNIVERSITY CHICAGO

Semantic Gap

RECAP



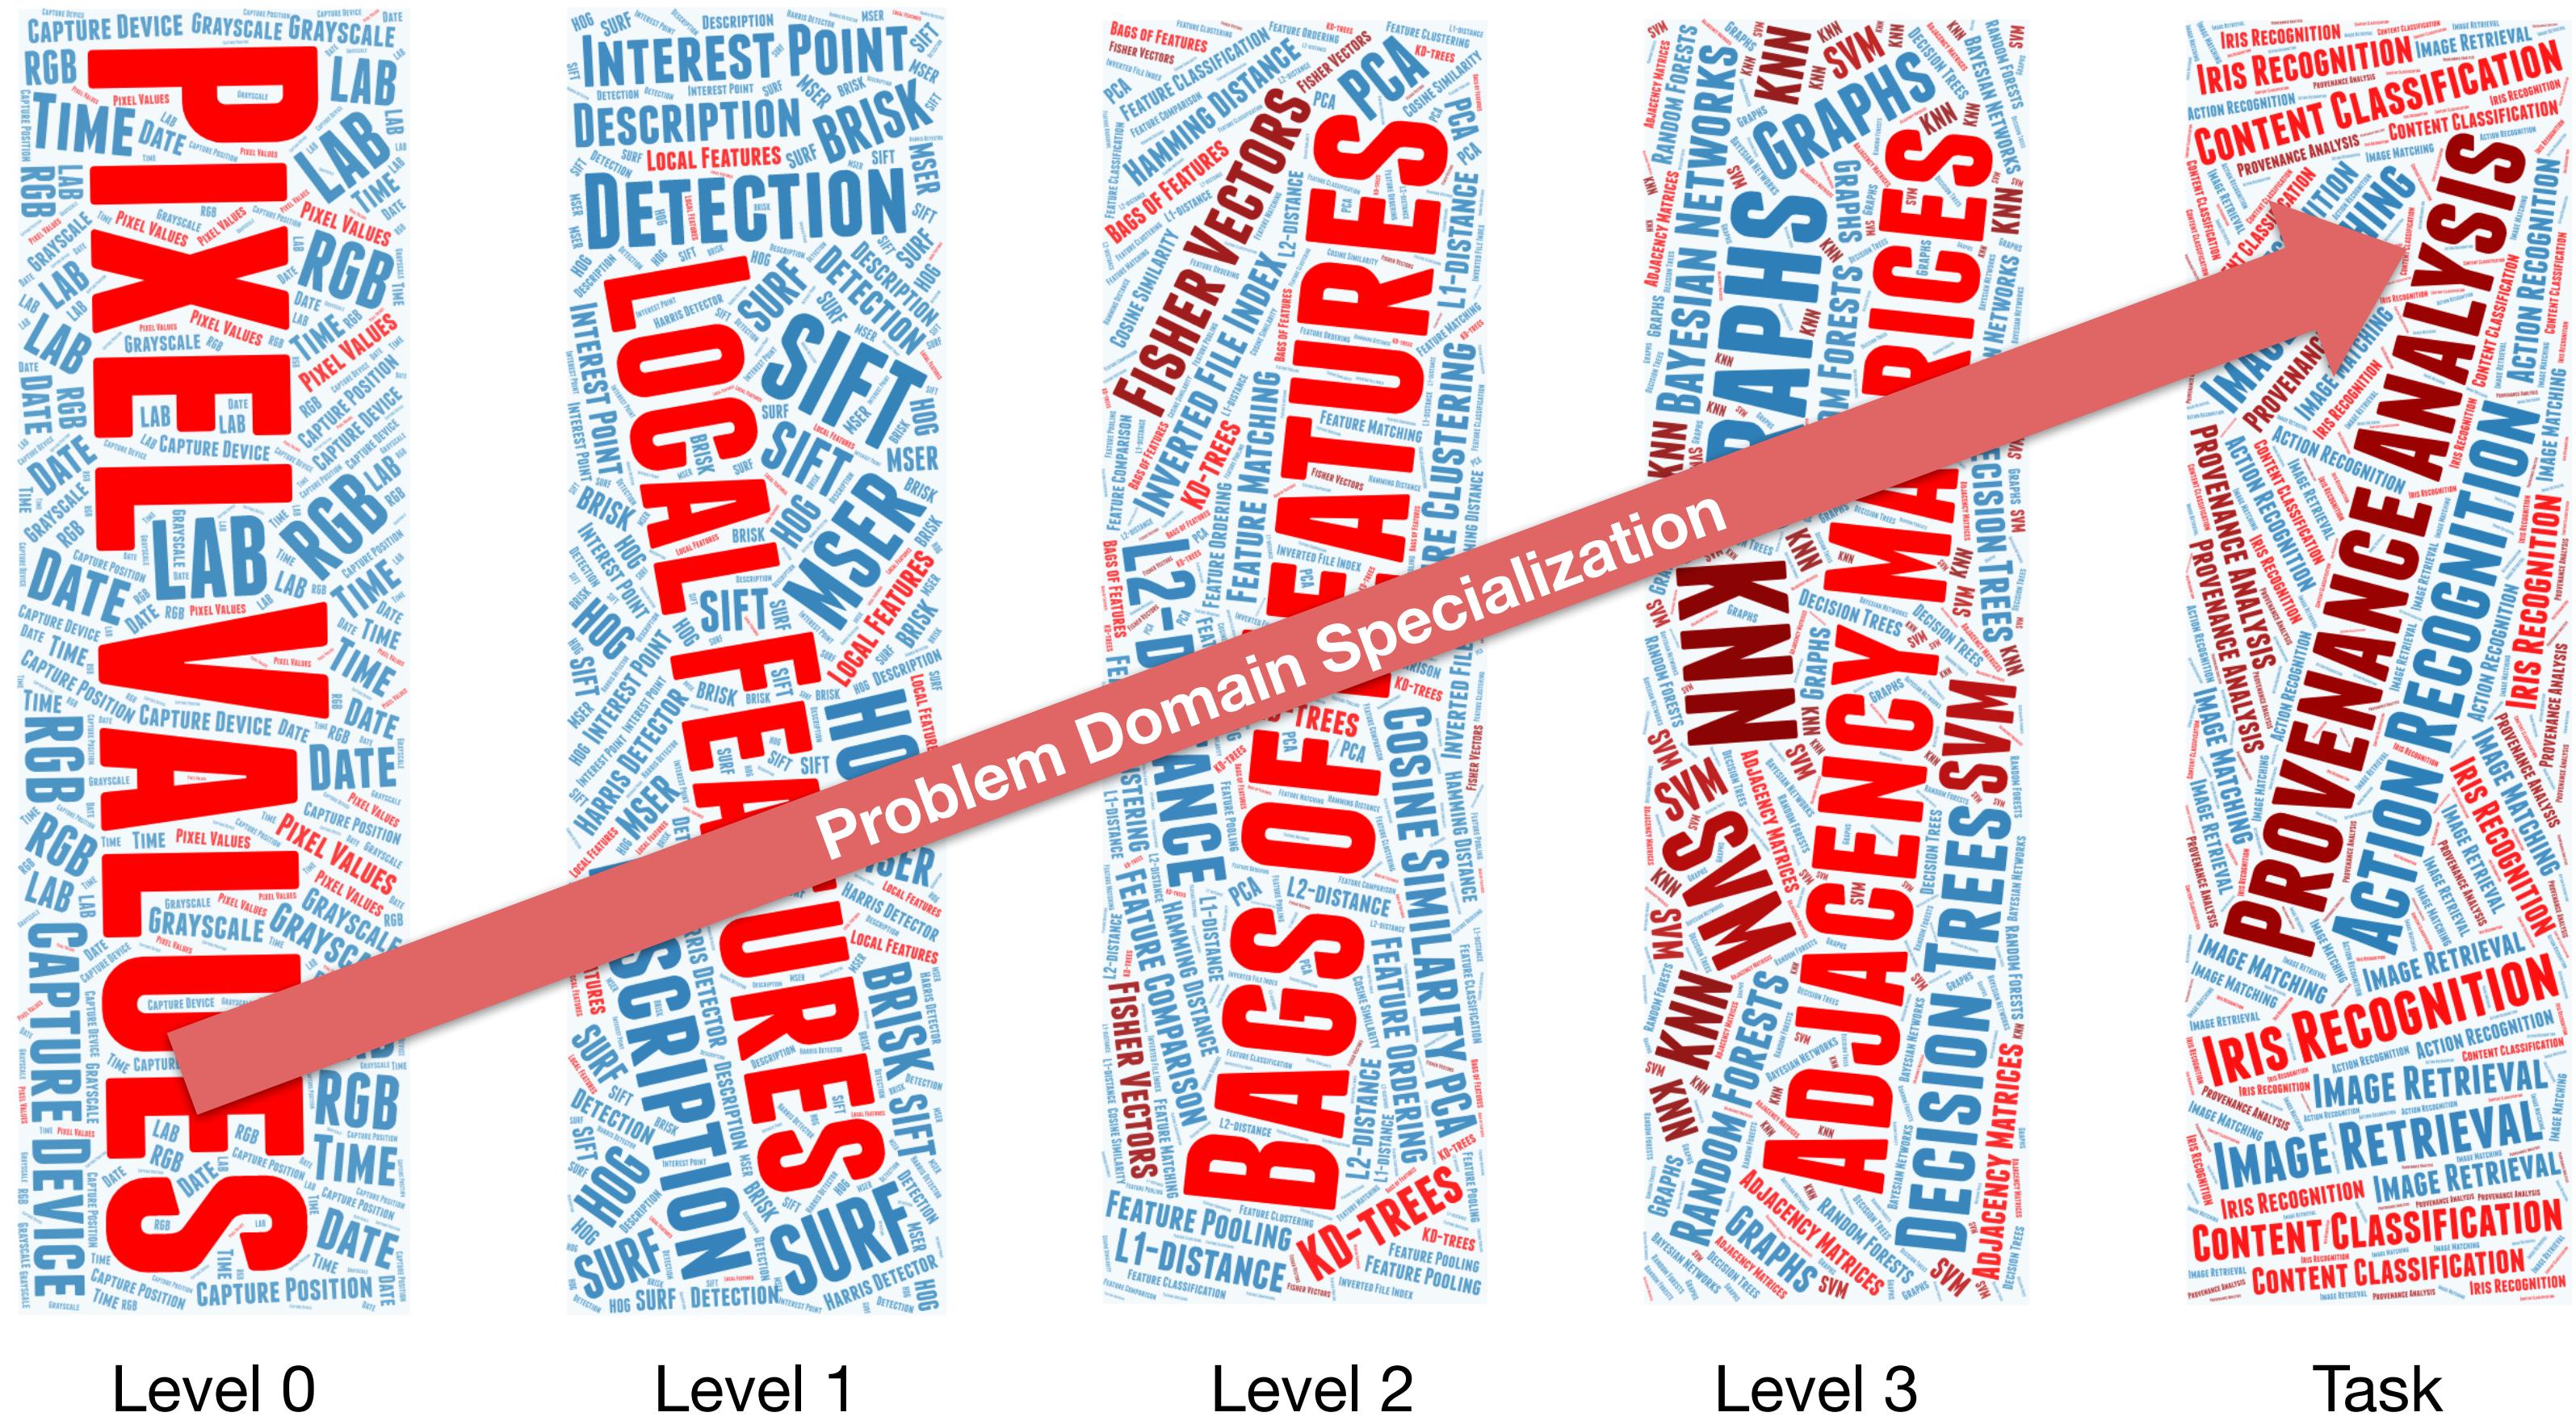
Level 0



Task

Semantic Gap

RECAP

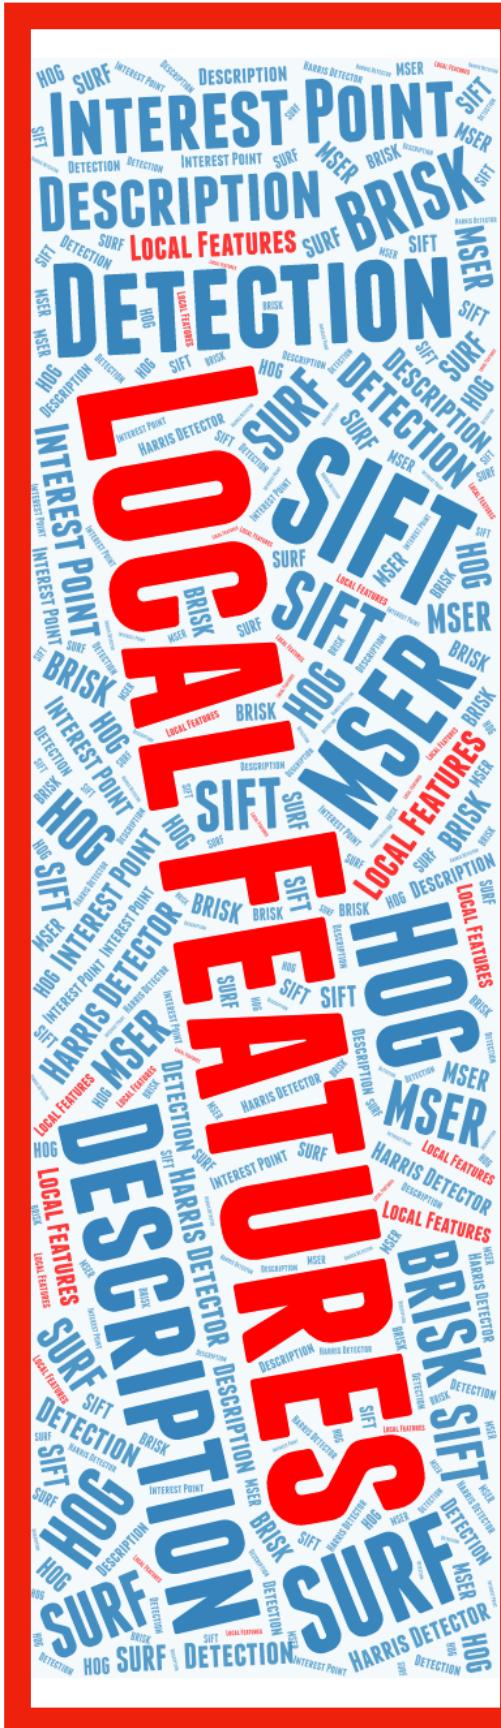


Semantic Gap

A large, red, rectangular stamp with a double-line border containing the word "RECAP" in bold, sans-serif capital letters.



Level 0



Level 1



Level 2



Level 3



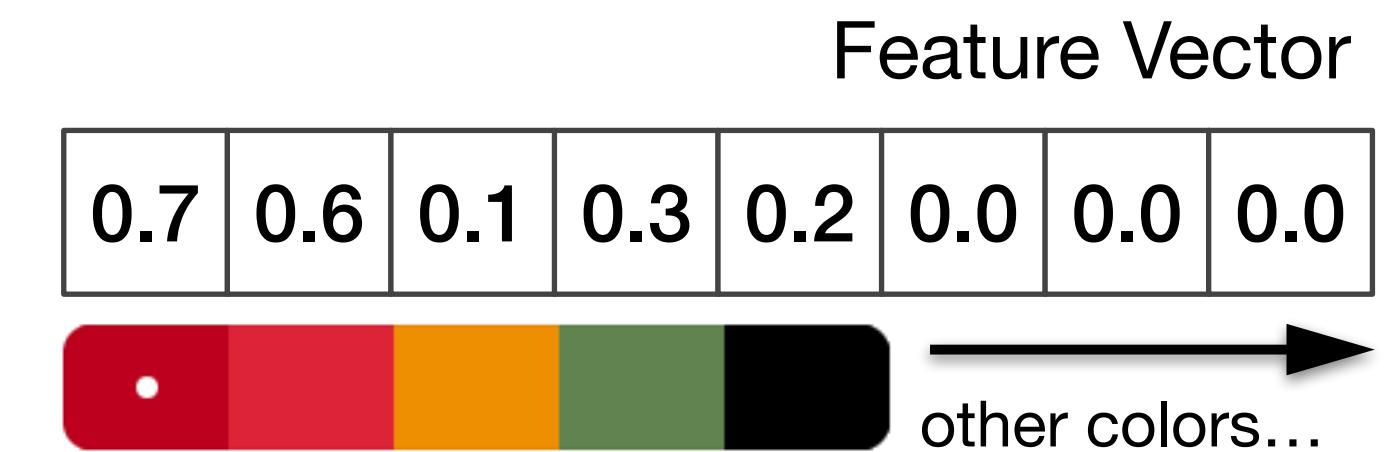
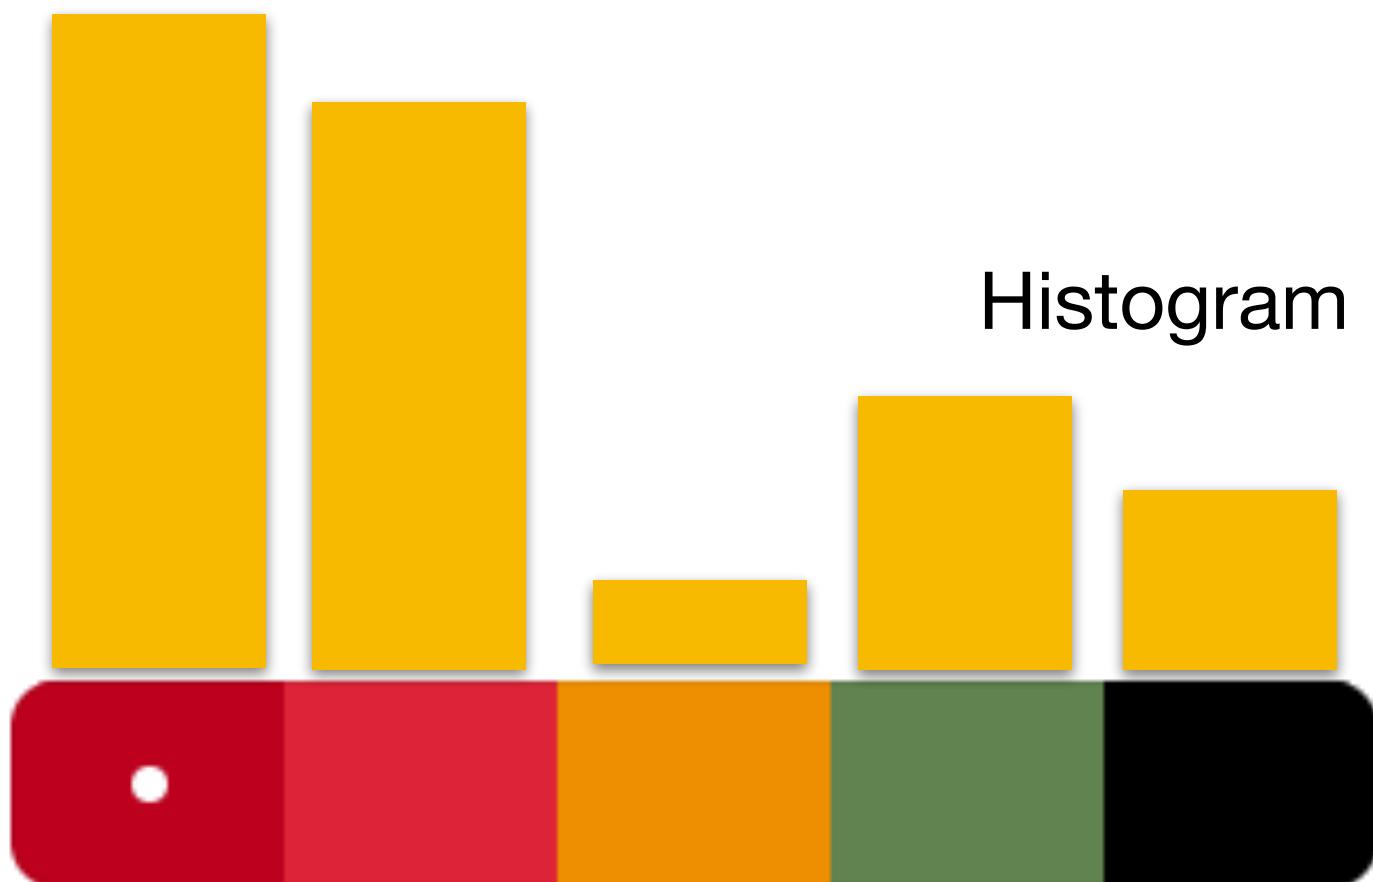
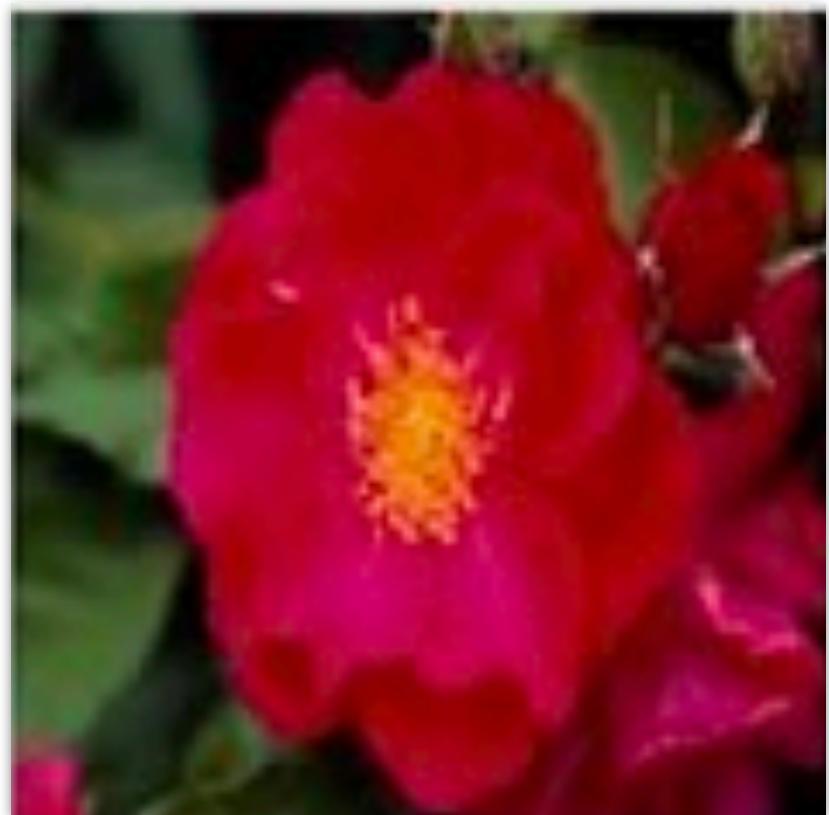
Task



Global Features

The entire image is represented by a single tensor.

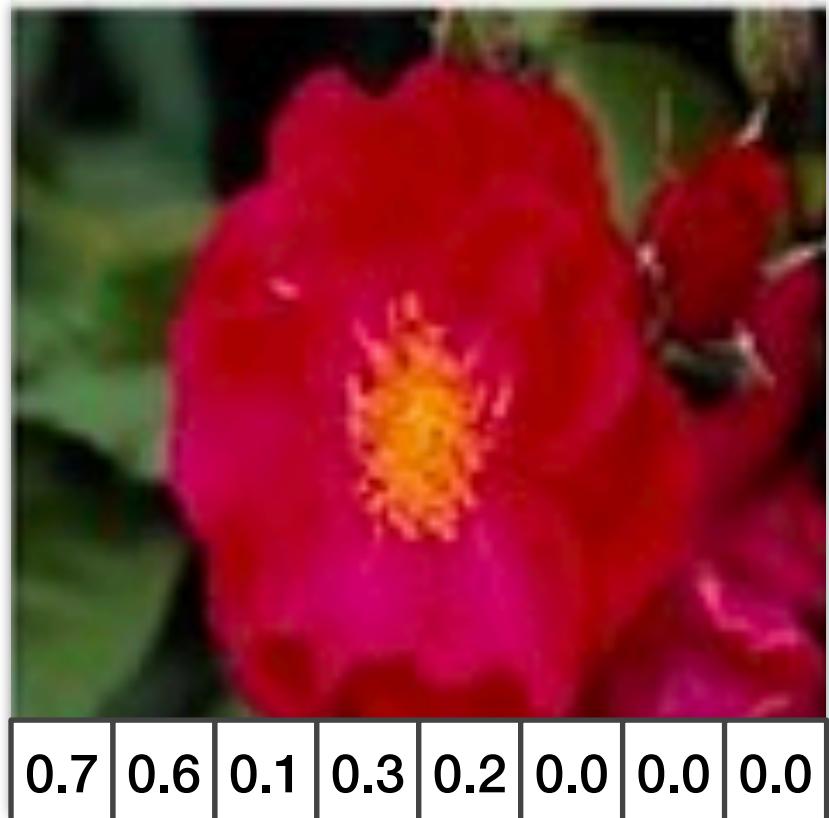
Example: **Color Histogram**



Global Features

The entire image is represented by a single tensor.

Example: **Color Histogram**



query

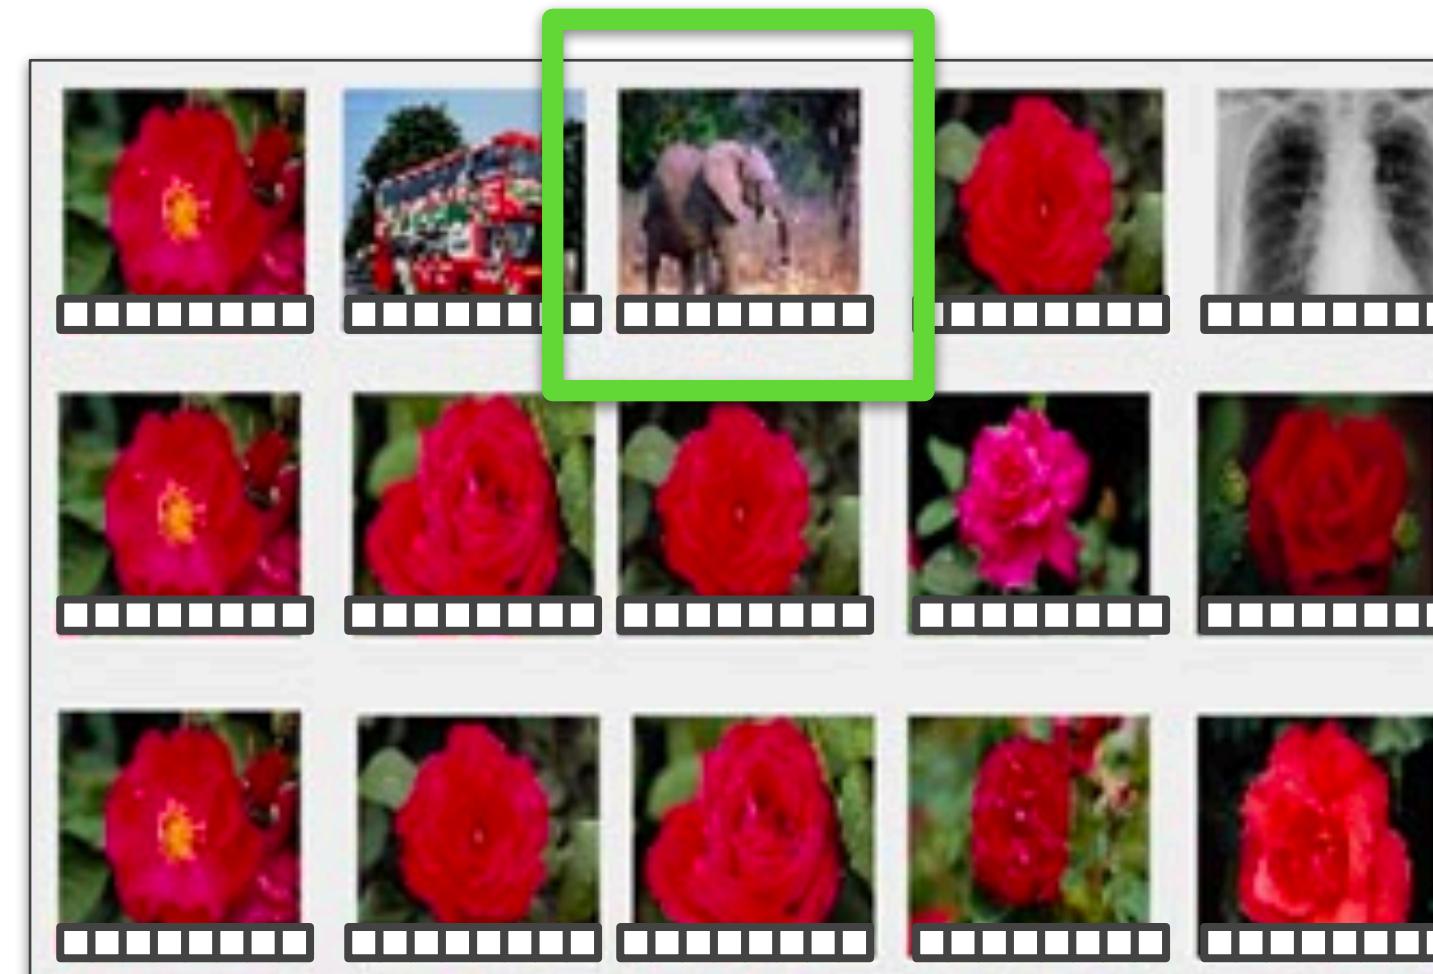
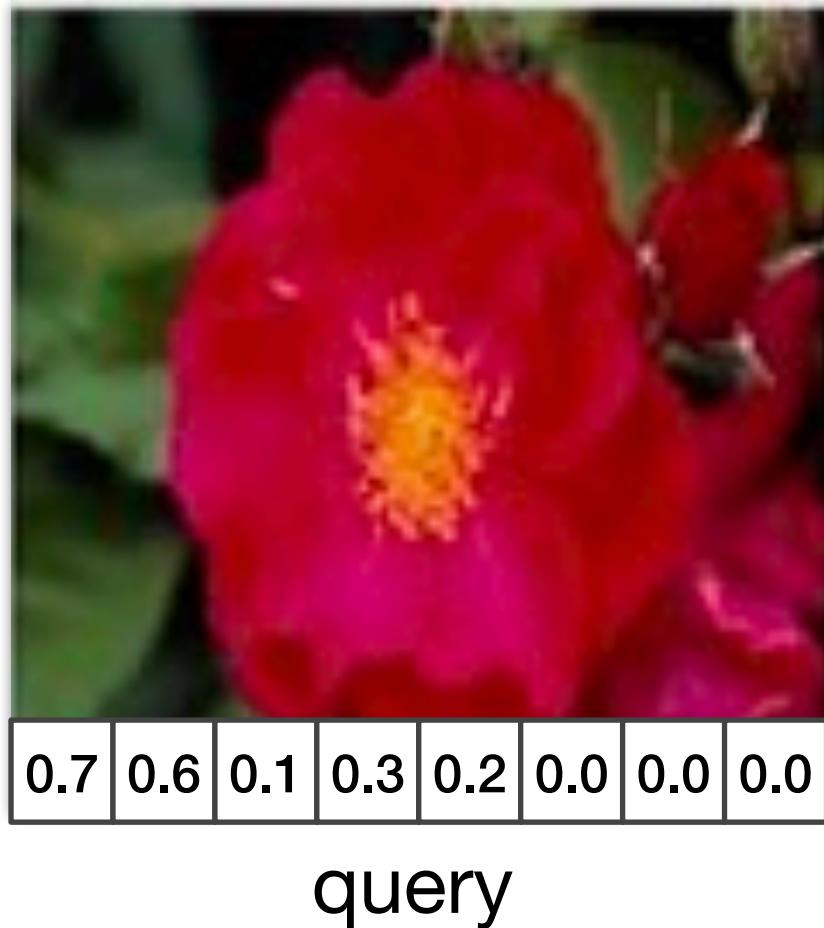


<https://www.pantechsolutions.net/matlab-code-for-image-retrieval>

Global Features

The entire image is represented by a single tensor.

Example: **Color Histogram**



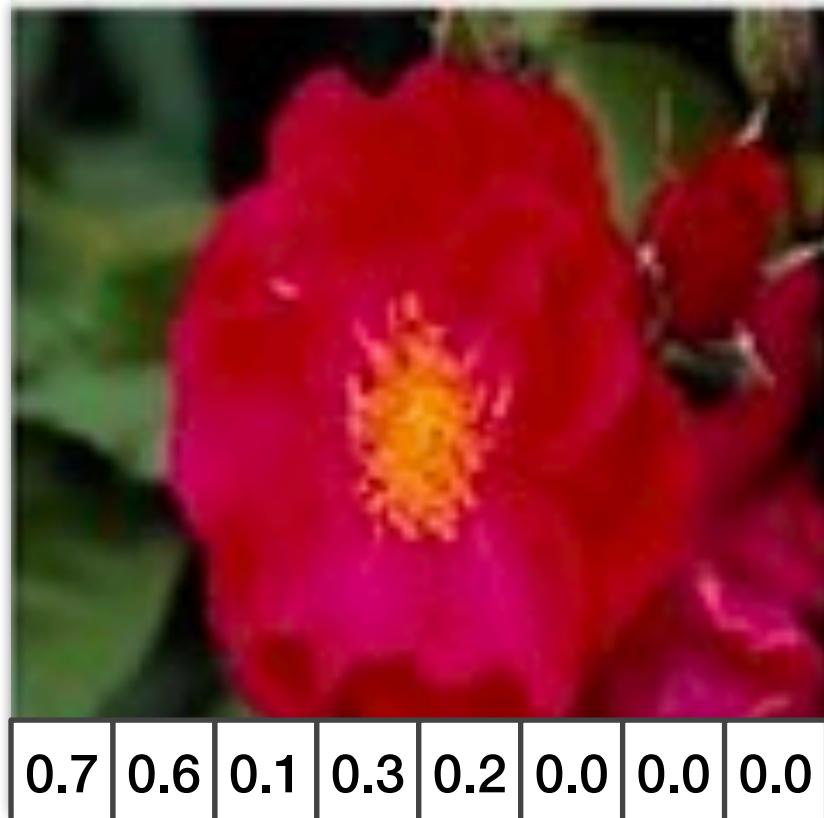
Cons

No distinction between foreground and background.

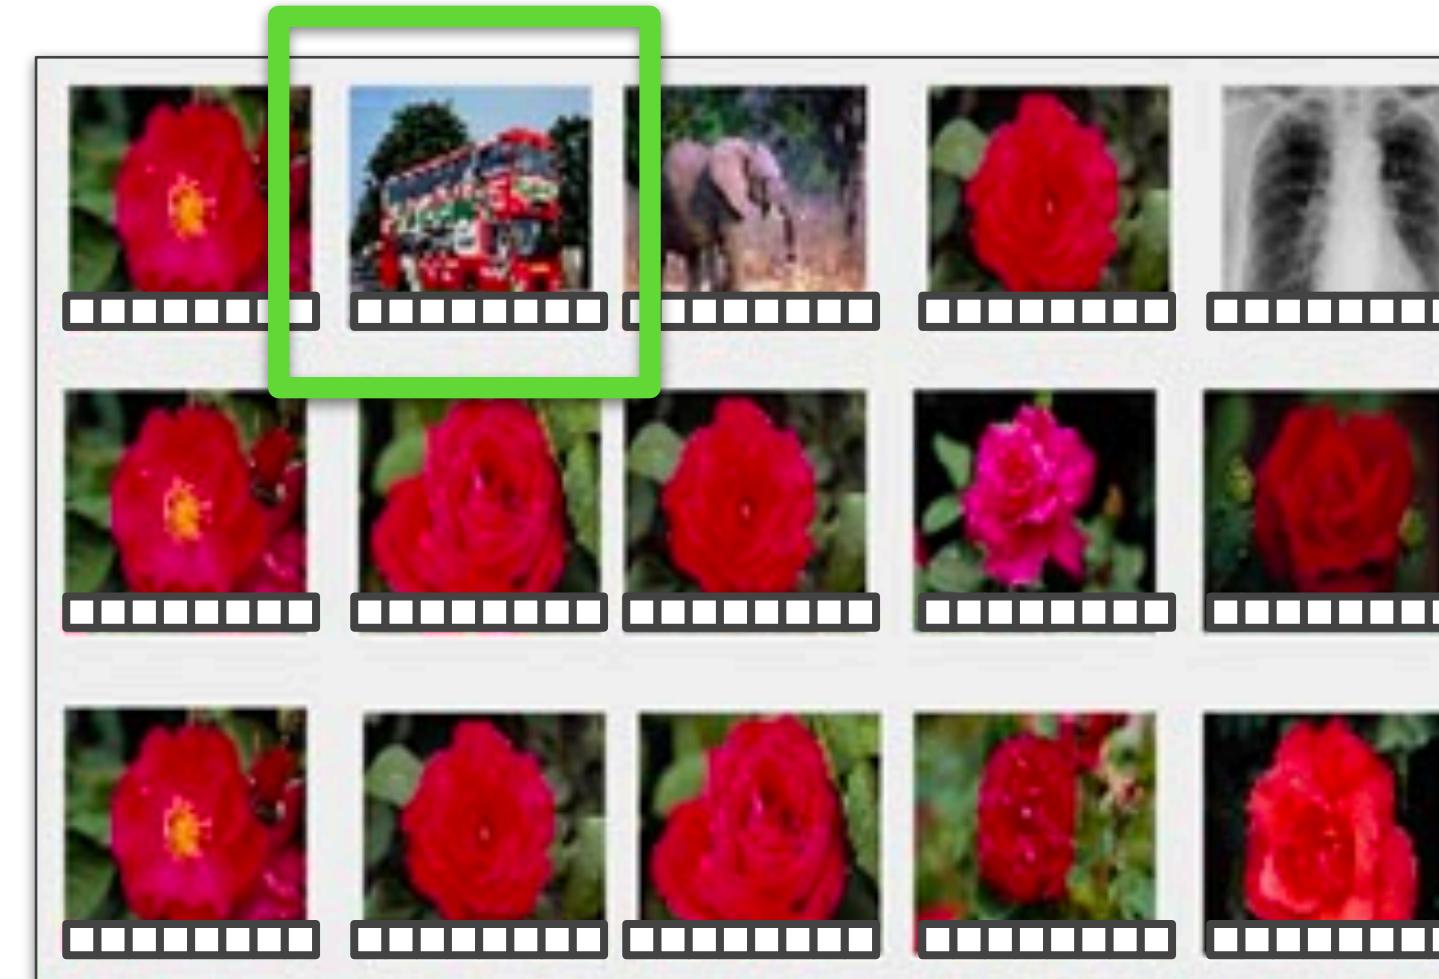
Global Features

The entire image is represented by a single tensor.

Example: **Color Histogram**



query



Cons

No semantics.

Global Features

The entire image is represented by a single tensor.

Example: **Color Histogram**



A lot of dark green.



A lot of white.

Cons

No semantics.

Not robust to occlusions.

No match.

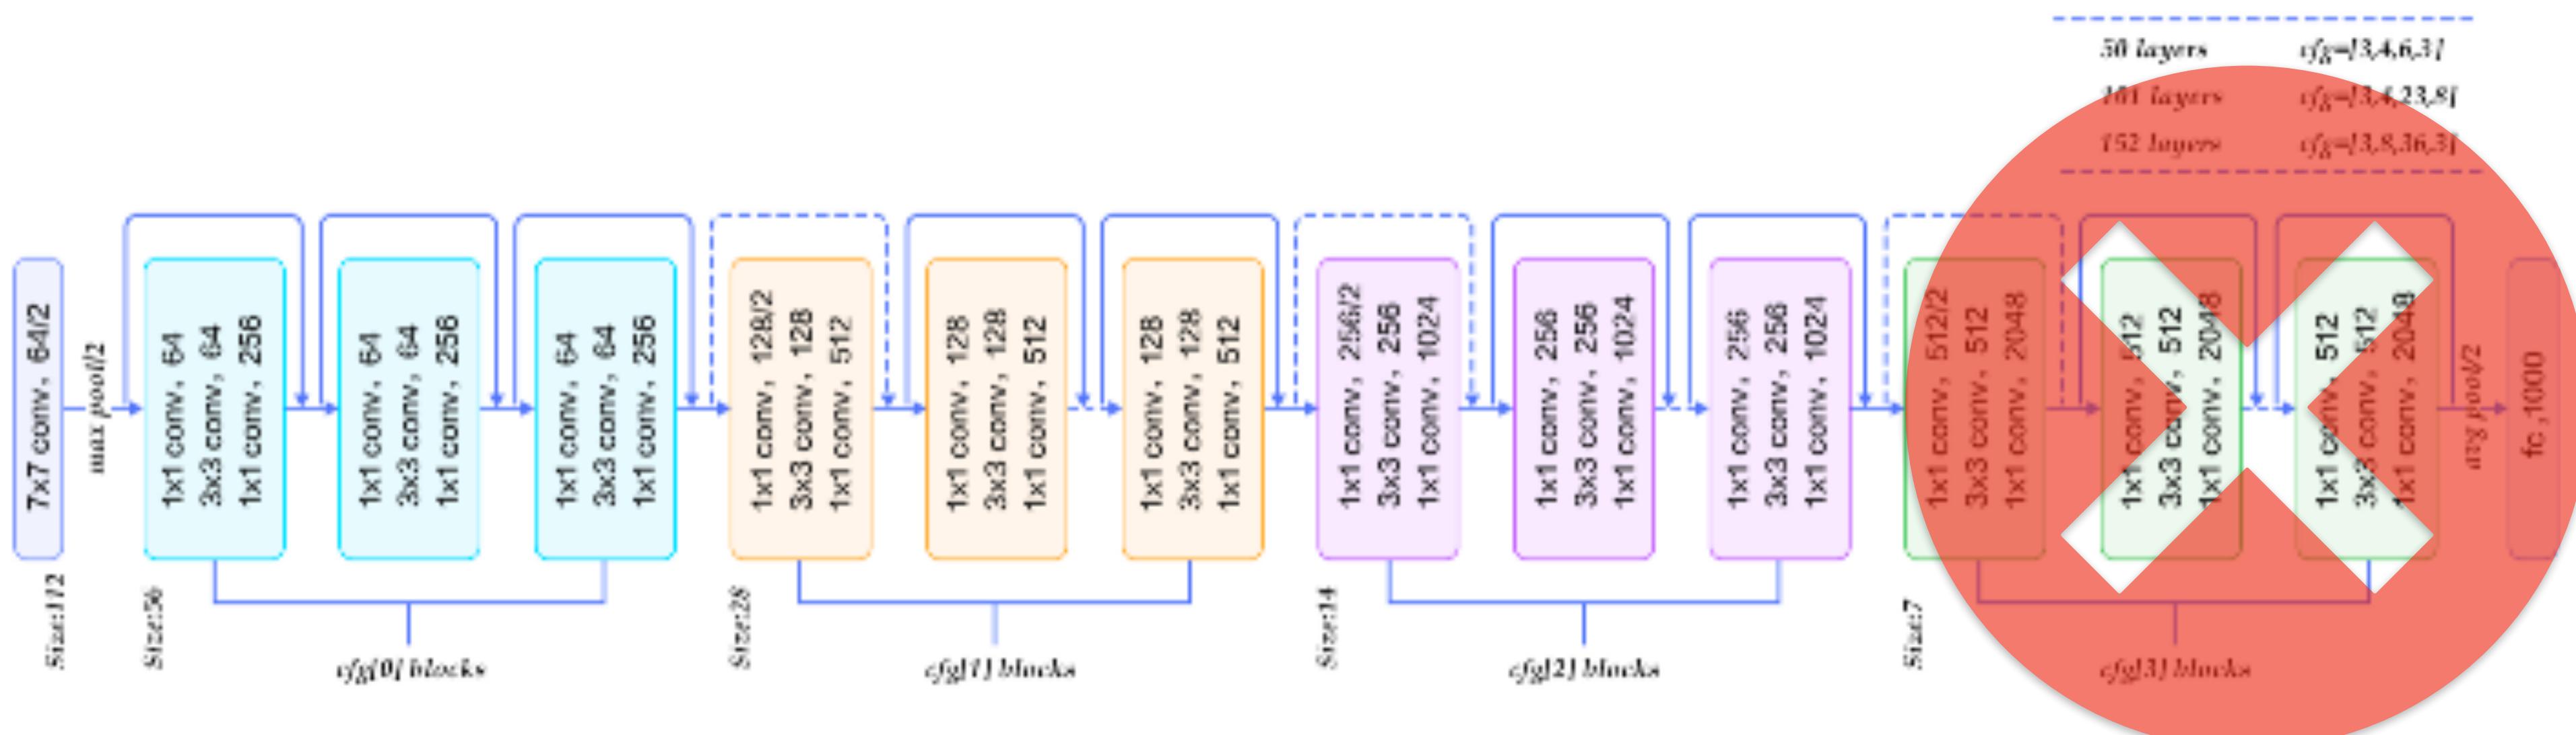


LOYOLA
UNIVERSITY CHICAGO

Global Features

The entire image is represented by a single tensor.

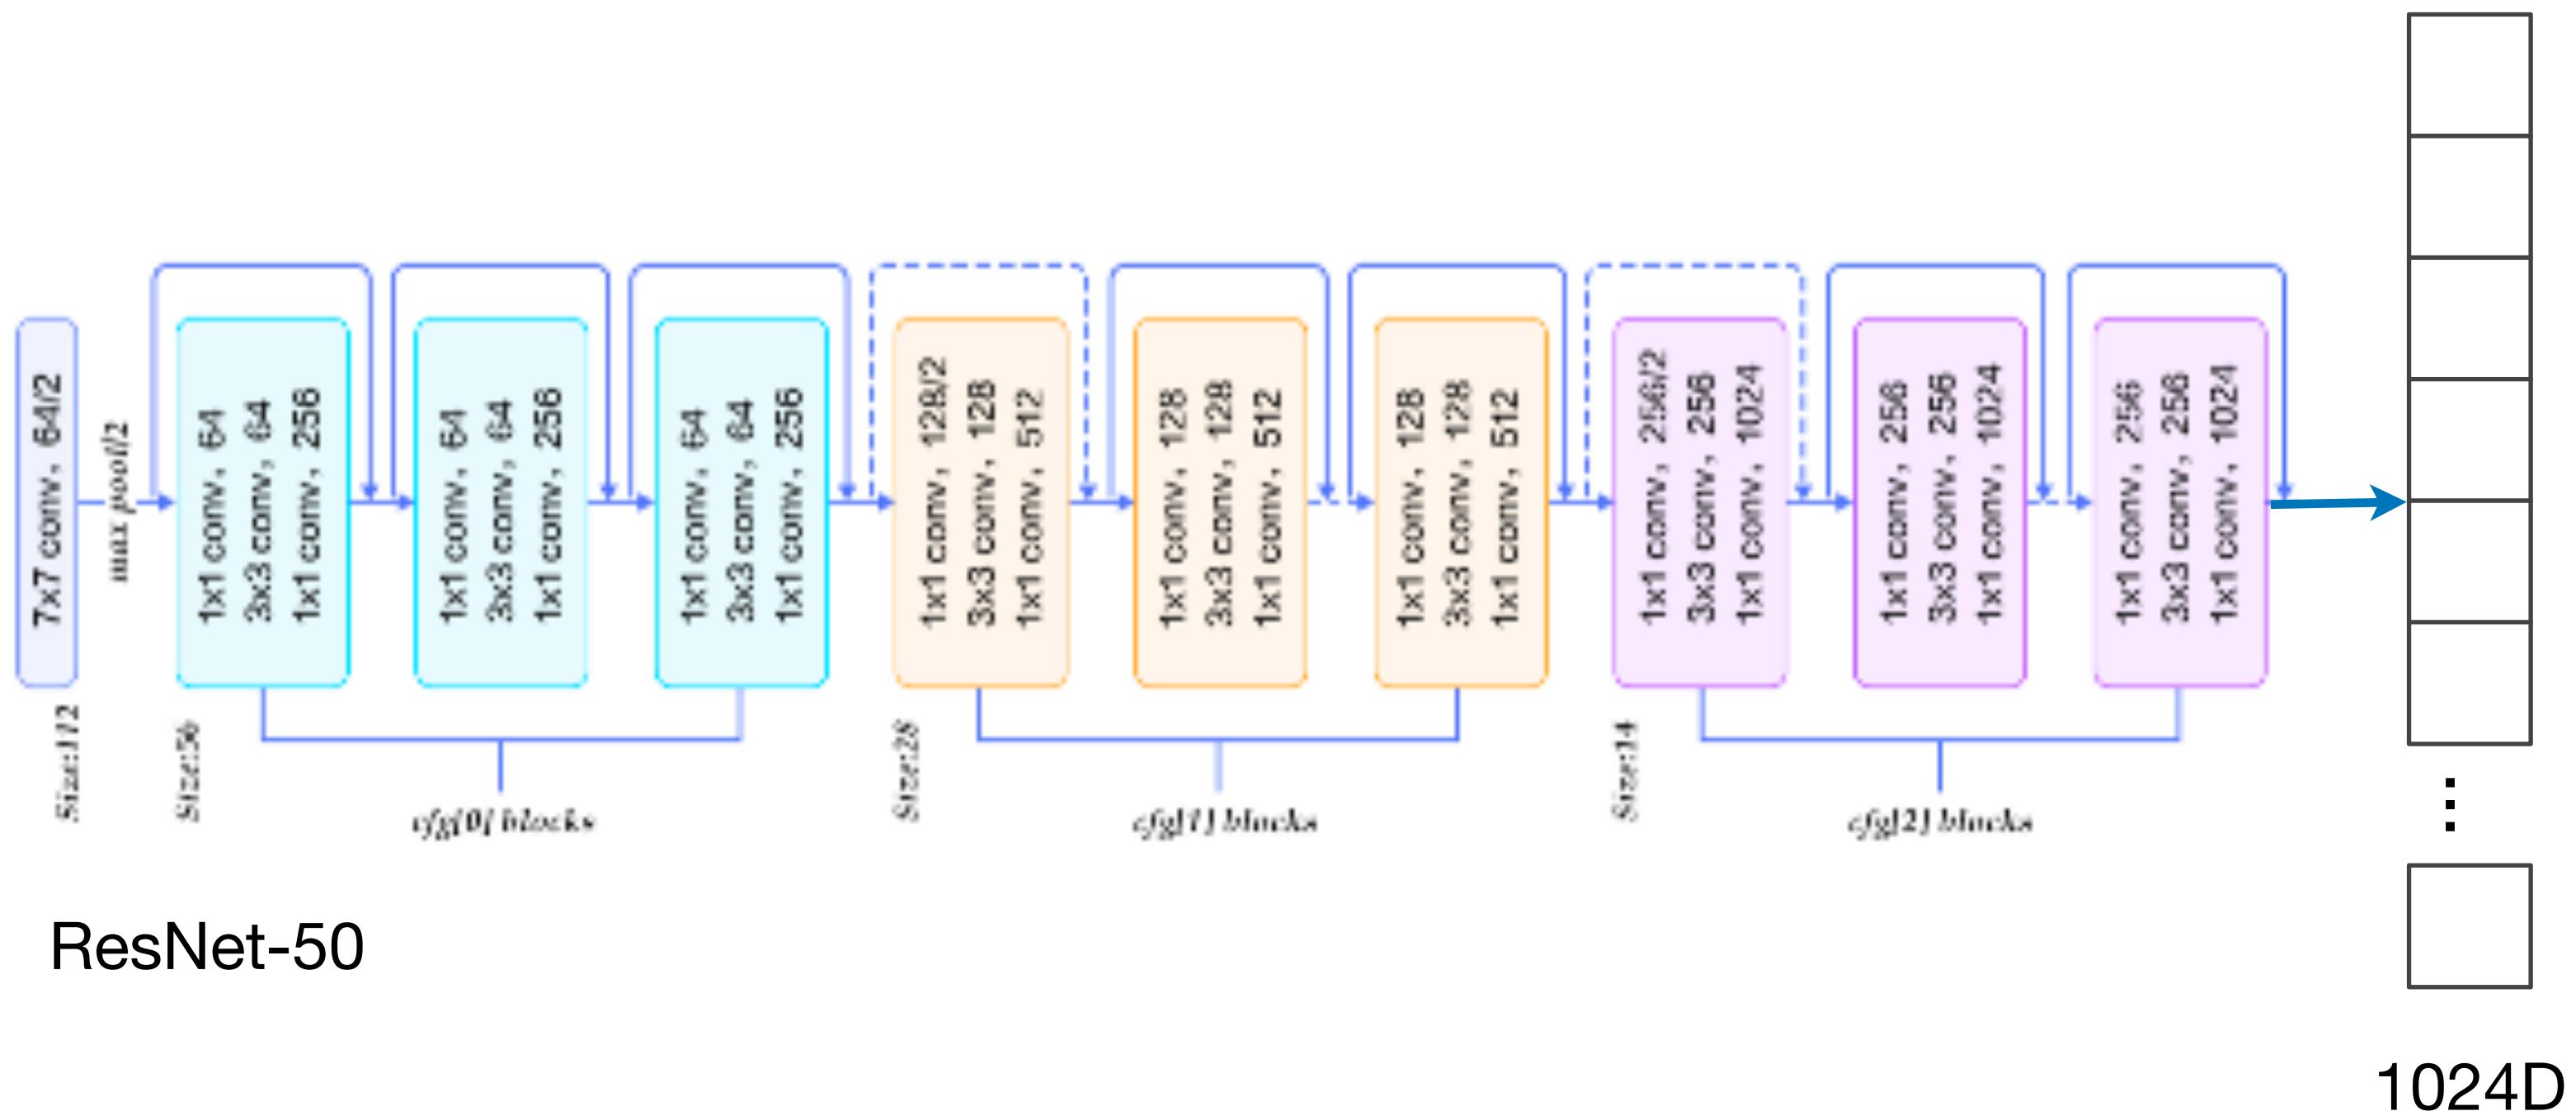
Example: **CNN-based**



Global Features

The entire image is represented by a single tensor.

Example: **CNN-based**



Pros
Inherited semantic awareness.

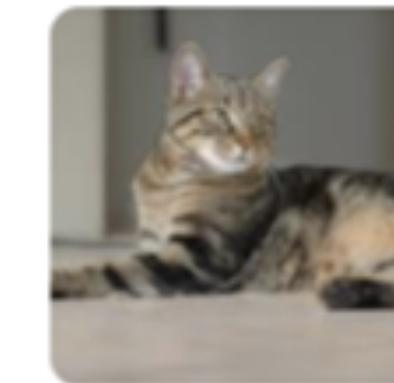
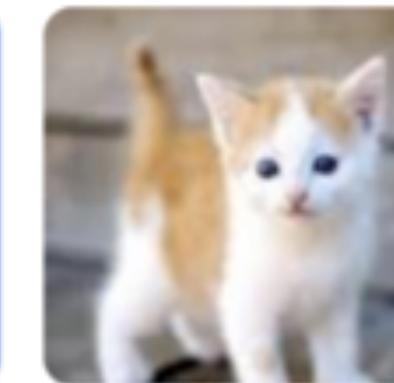
Global Features

The entire image is represented by a single tensor.

Example: **CNN-based**



query



database semantically similar images

Pros

Inherited semantic awareness.

Global Features

The entire image is represented by a single tensor.

Example: **CNN-based**



A lot of occlusion.



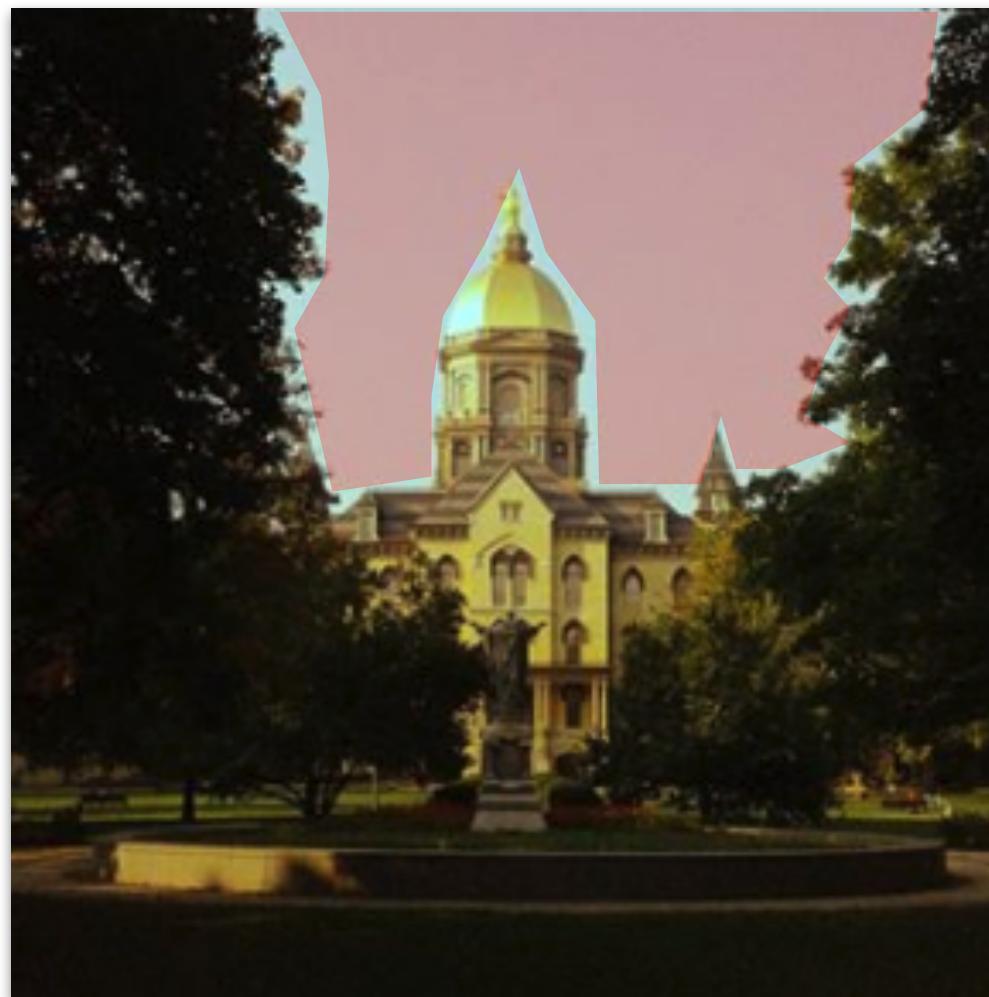
Cons

Can it help to decide if these images depict the same building?

Global Features

The entire image is represented by a single tensor.

Example: **CNN-based**



A lot of occlusion.



Unnecessary regions.

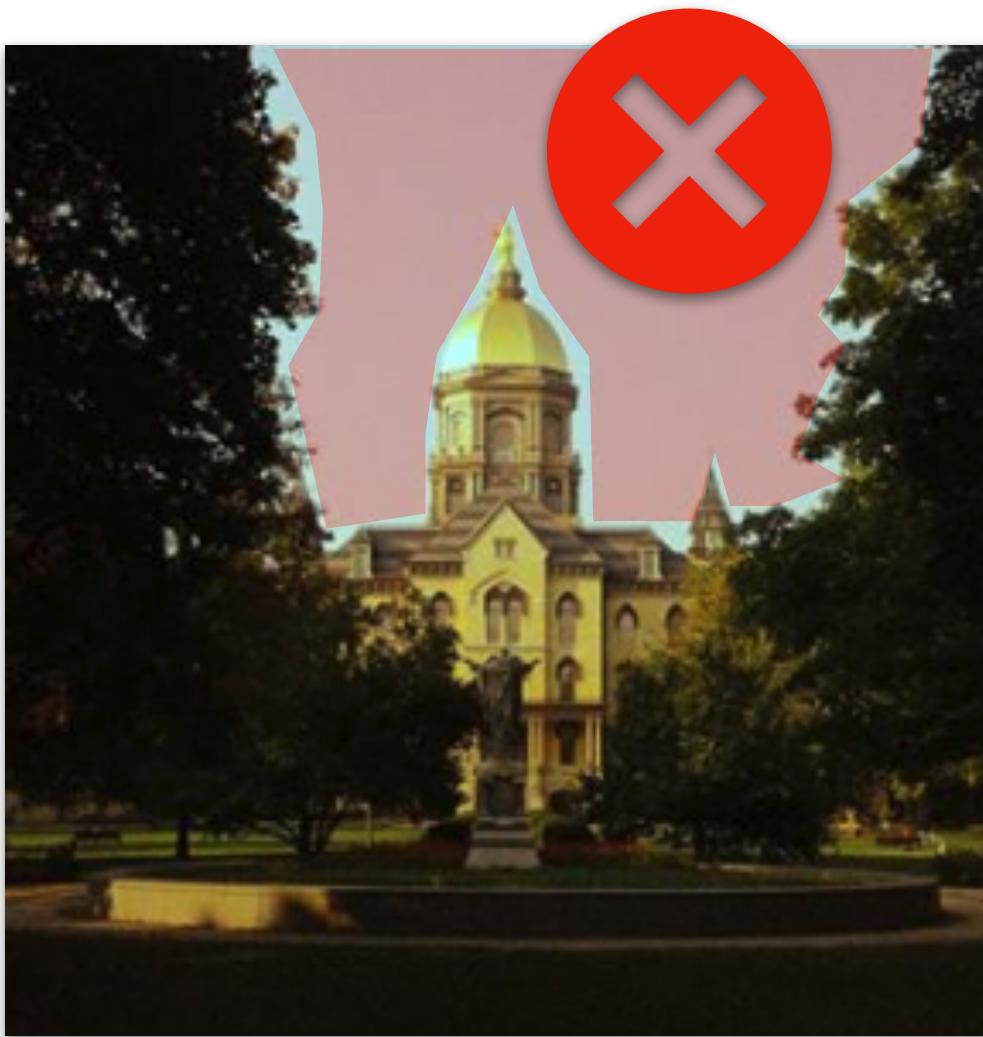
Cons

Can it help to decide if these images depict the same building?

Local Features

What are Local Features?

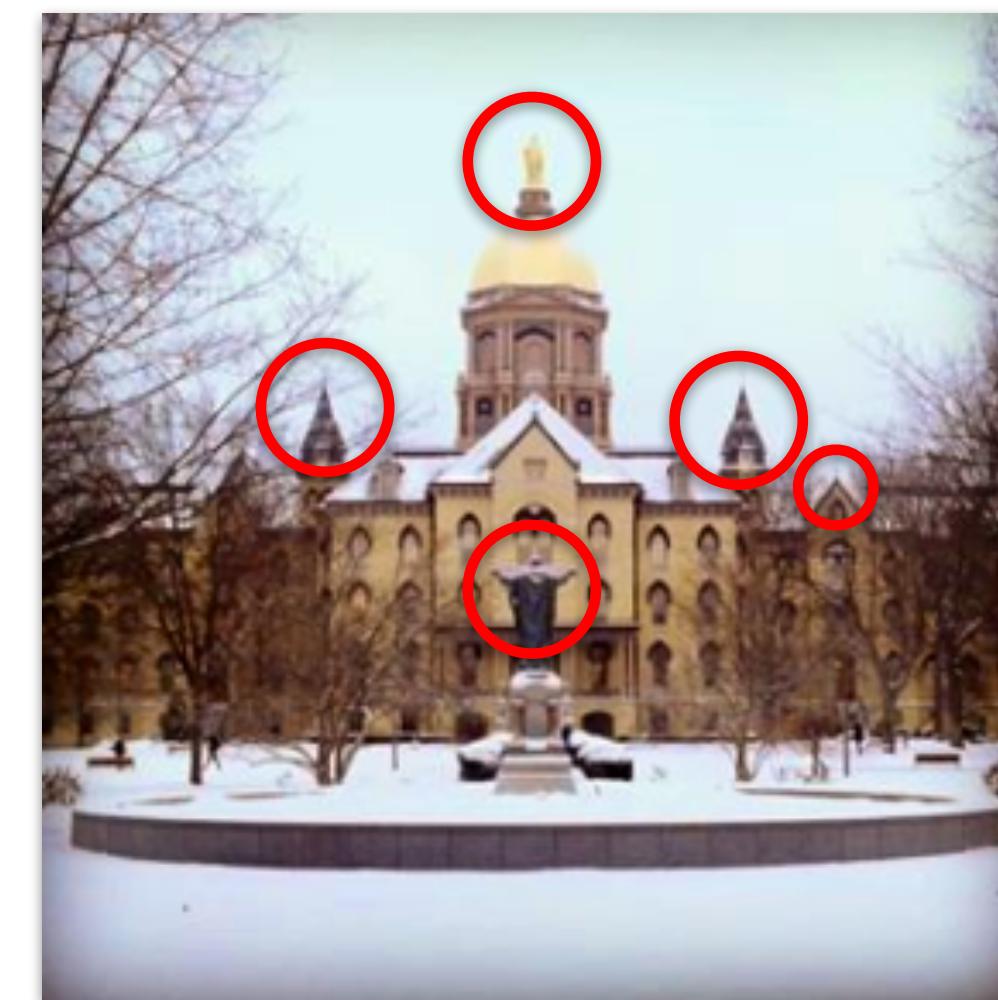
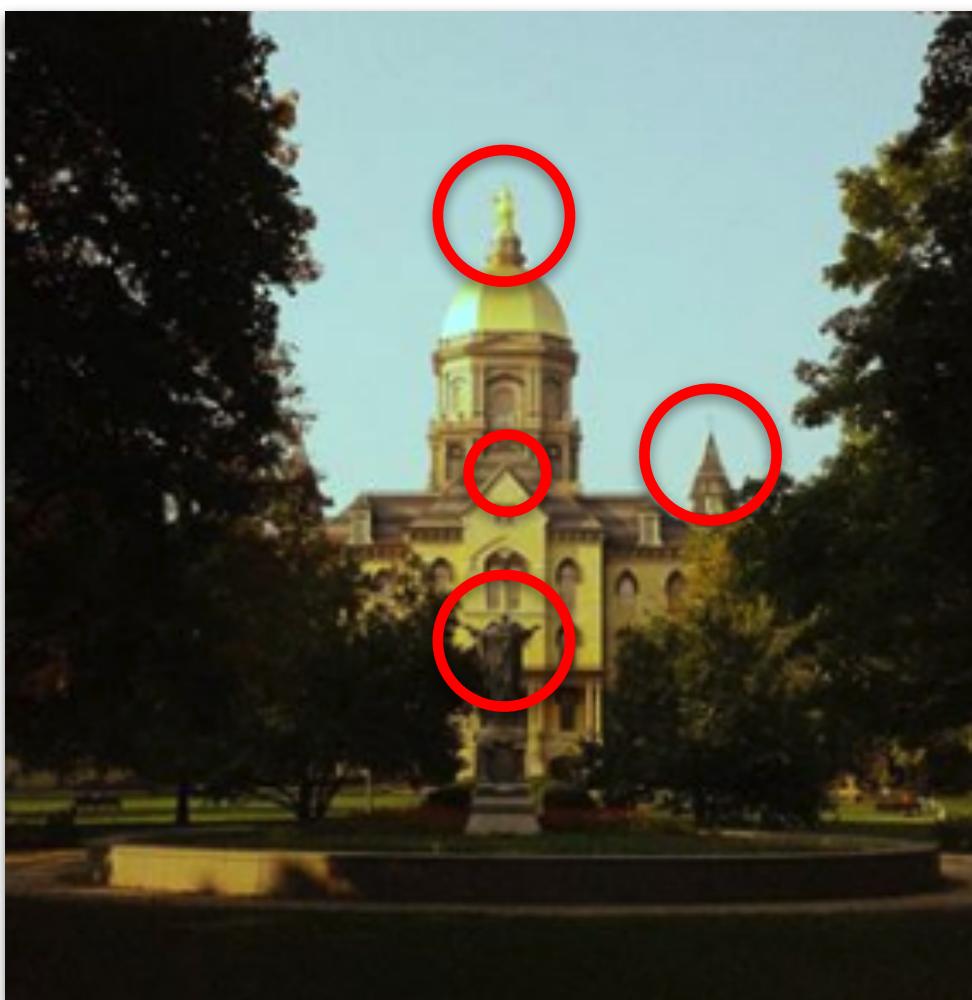
Image patterns that differ from their immediate neighborhood.



Local Features

What are Local Features?

Image patterns that differ from their immediate neighborhood.

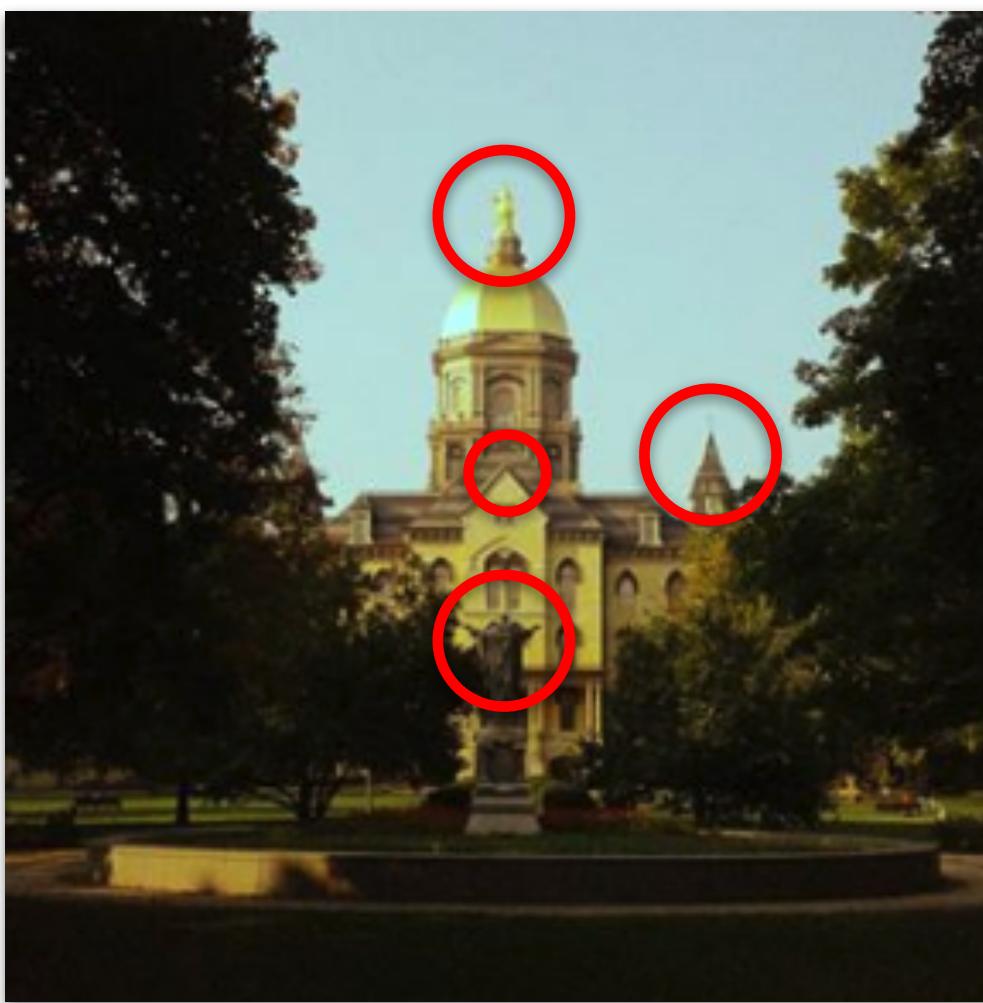


Local Features

What are Local Features?

Image patterns that differ from their immediate neighborhood.

Possible targets: points, edges, corners, junctions, blobs, etc.



Local Features

Why should one use Local Features?

Relevance

- (i) Edges are usually enough for humans' object recognition.
- (ii) Removing the corners hinders humans' abilities.



craiyon.com

Local Features

Why should one use Local Features?

Establish anchor points for image registration.

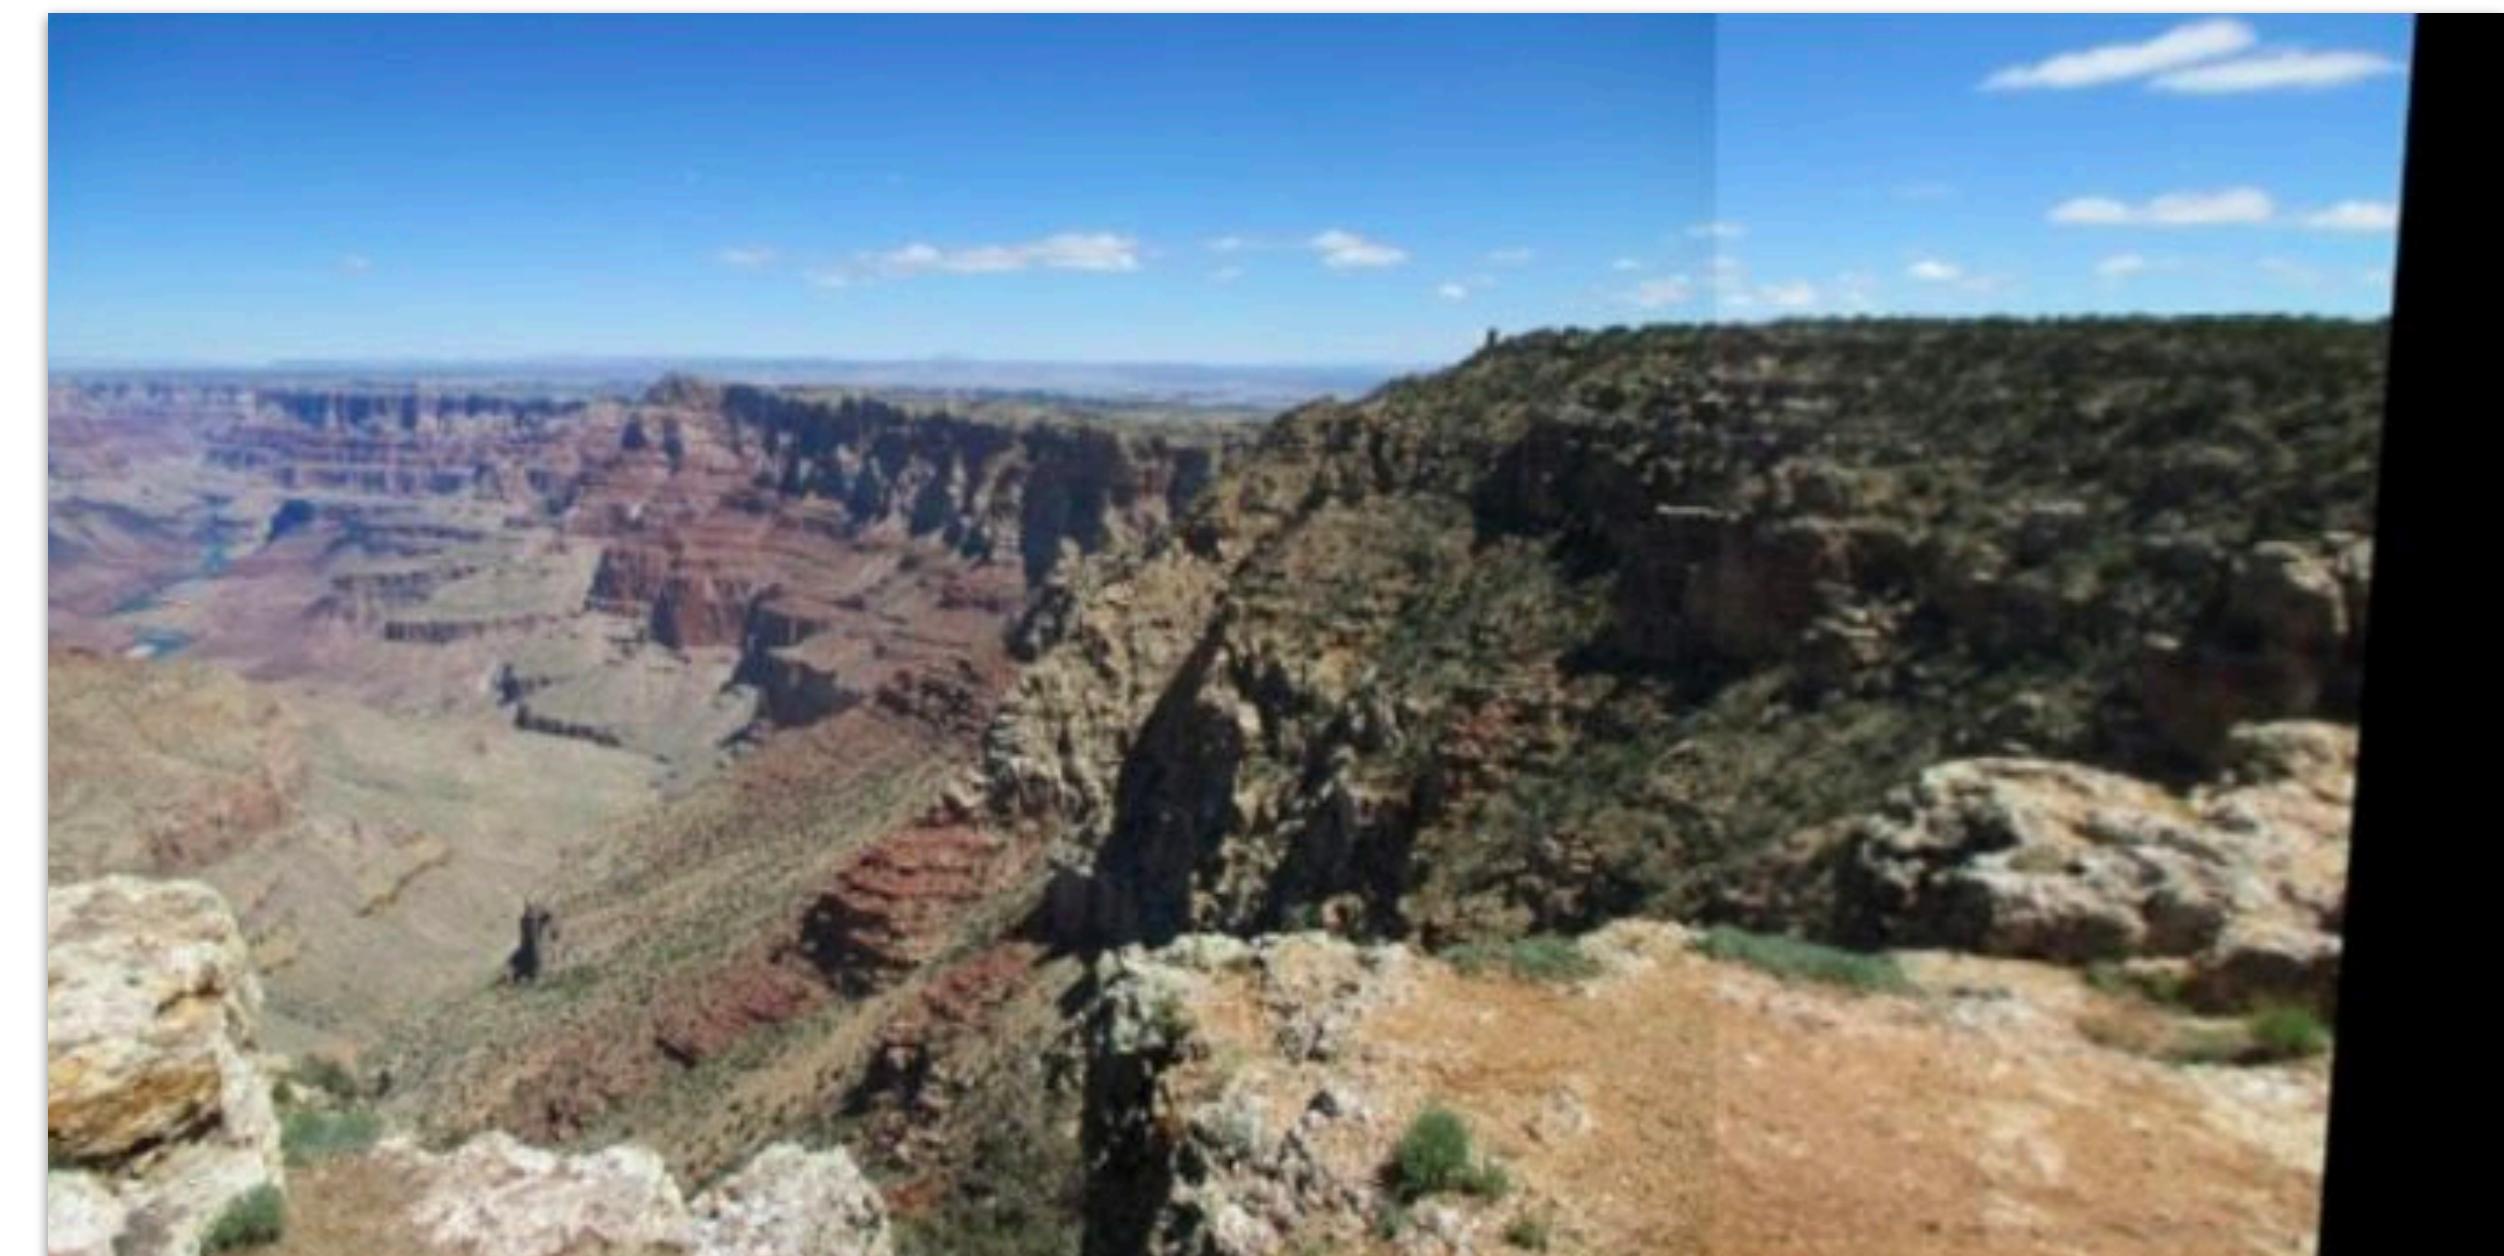
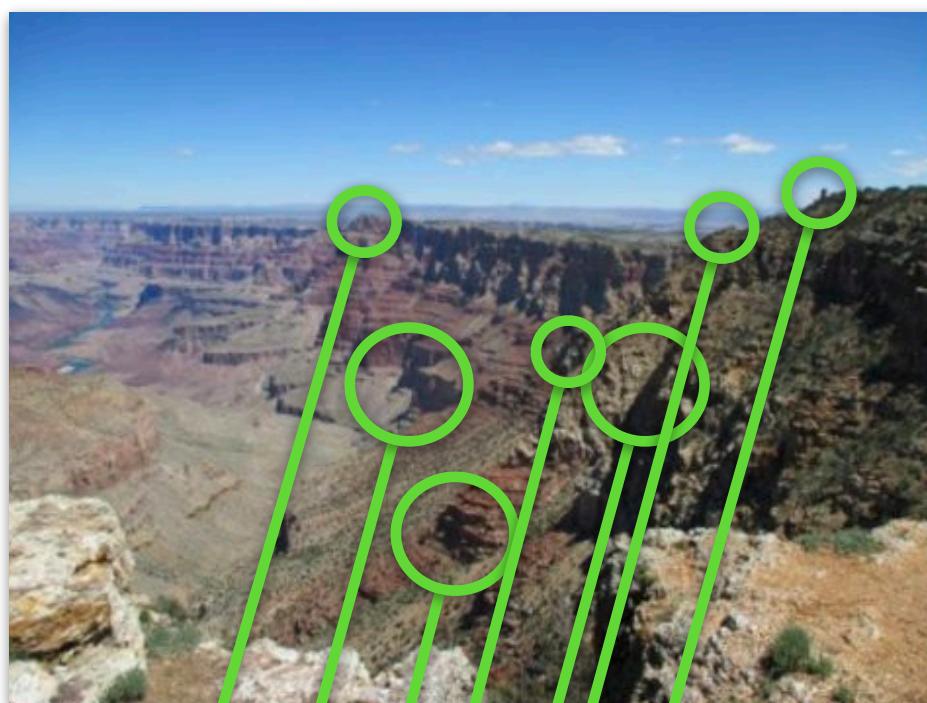


Local Features

Why should one use Local Features?

Establish anchor points for image registration.

pyimagesearch.com

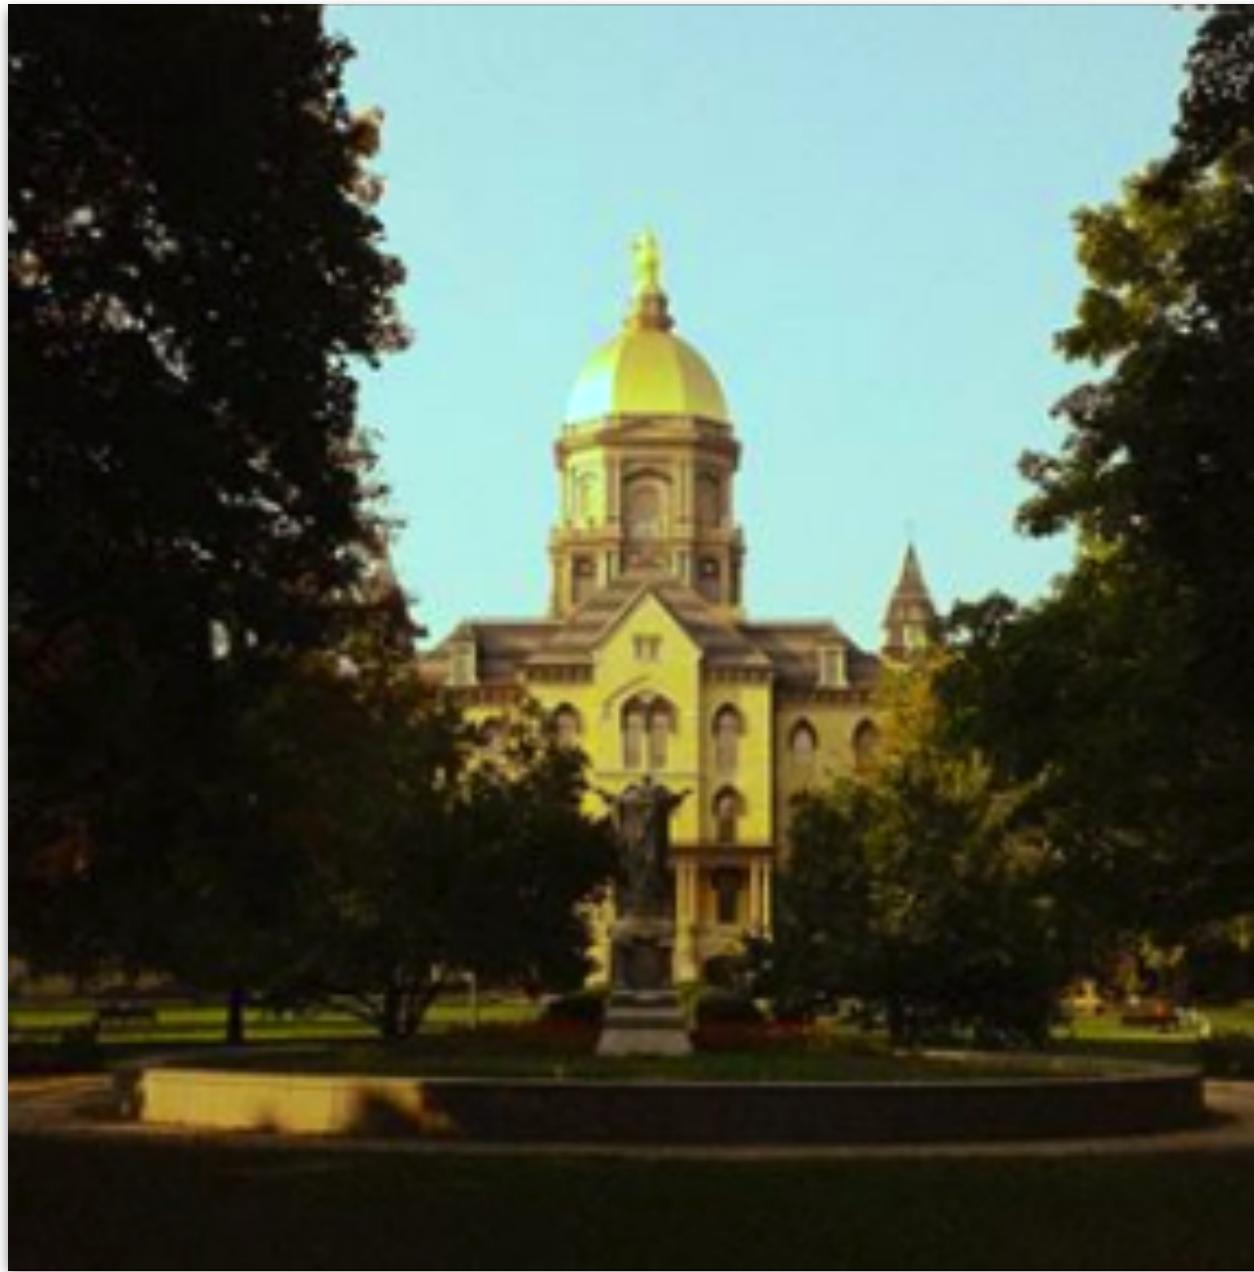


LOYOLA
UNIVERSITY CHICAGO

Local Features

Why should one use Local Features?

Obtain robust and compact image representation for many CV tasks.

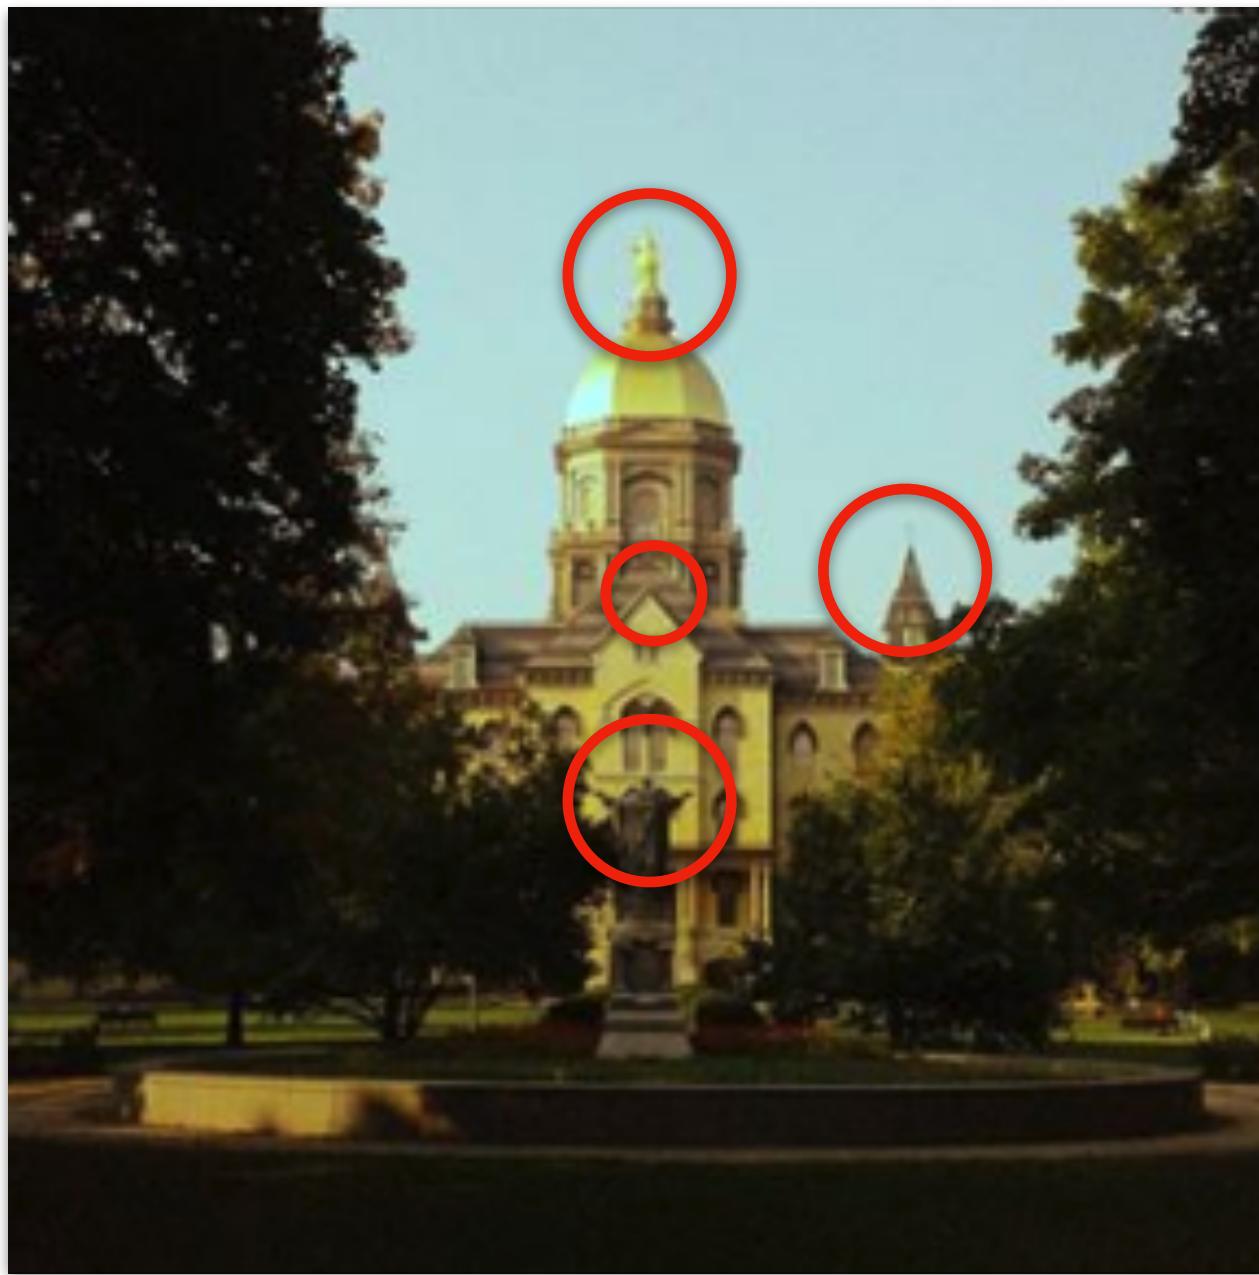


Can a computer system
decide if these images
depict the same building?

Local Features

Why should one use Local Features?

Obtain robust and compact image representation for many CV tasks.

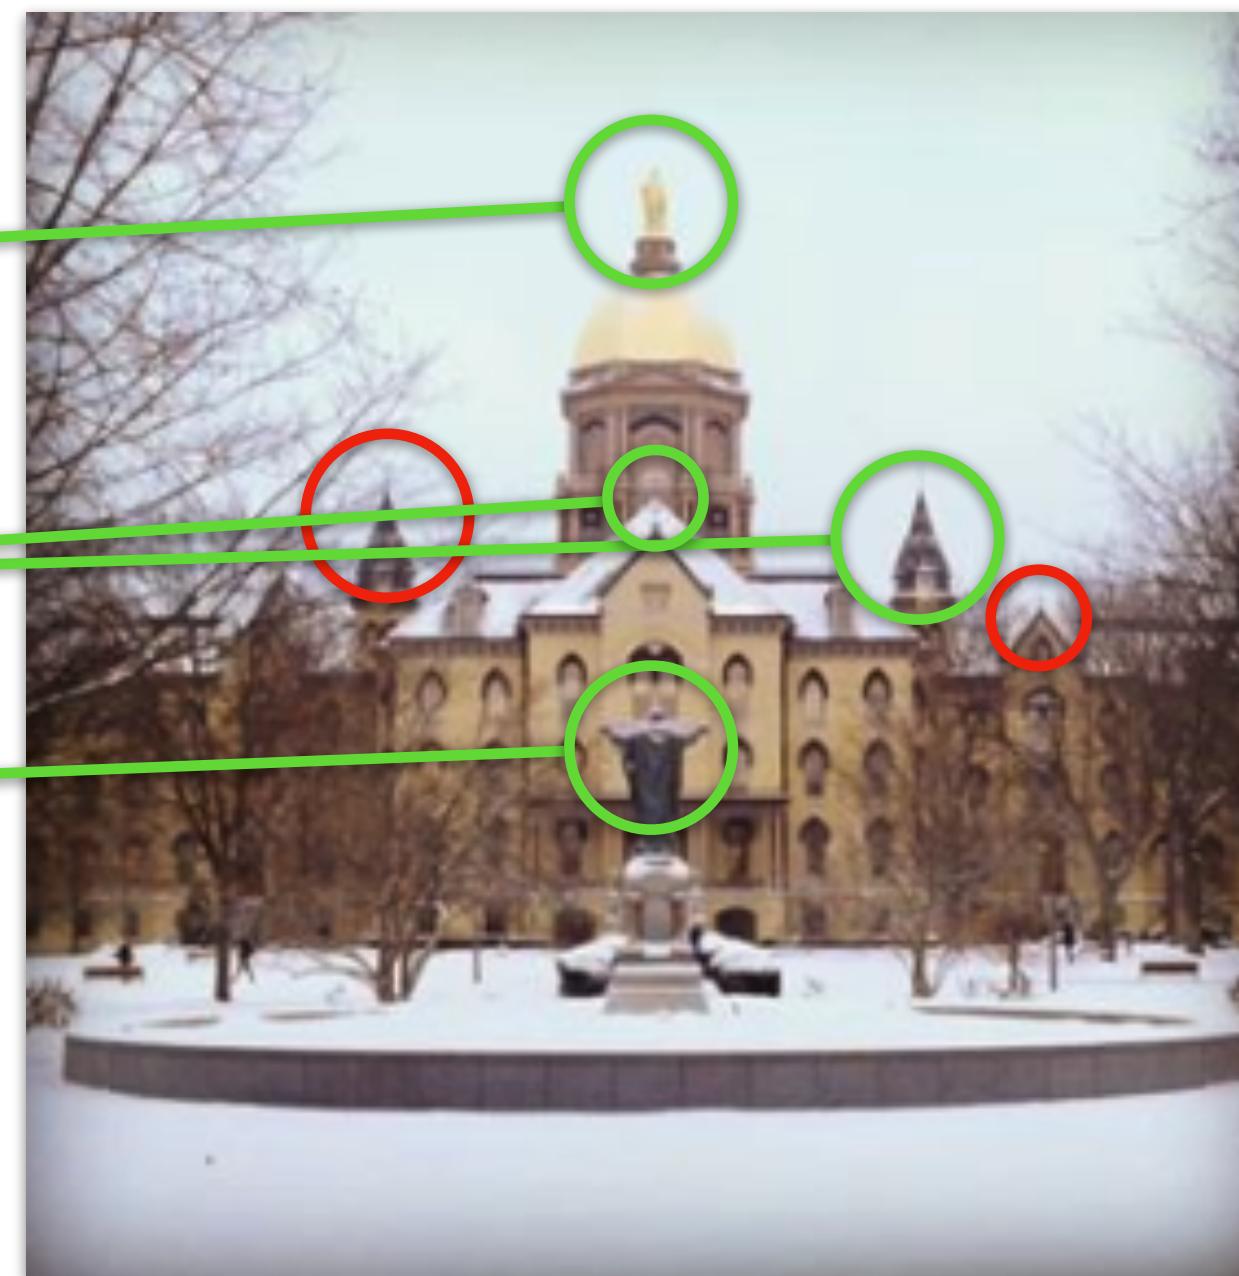
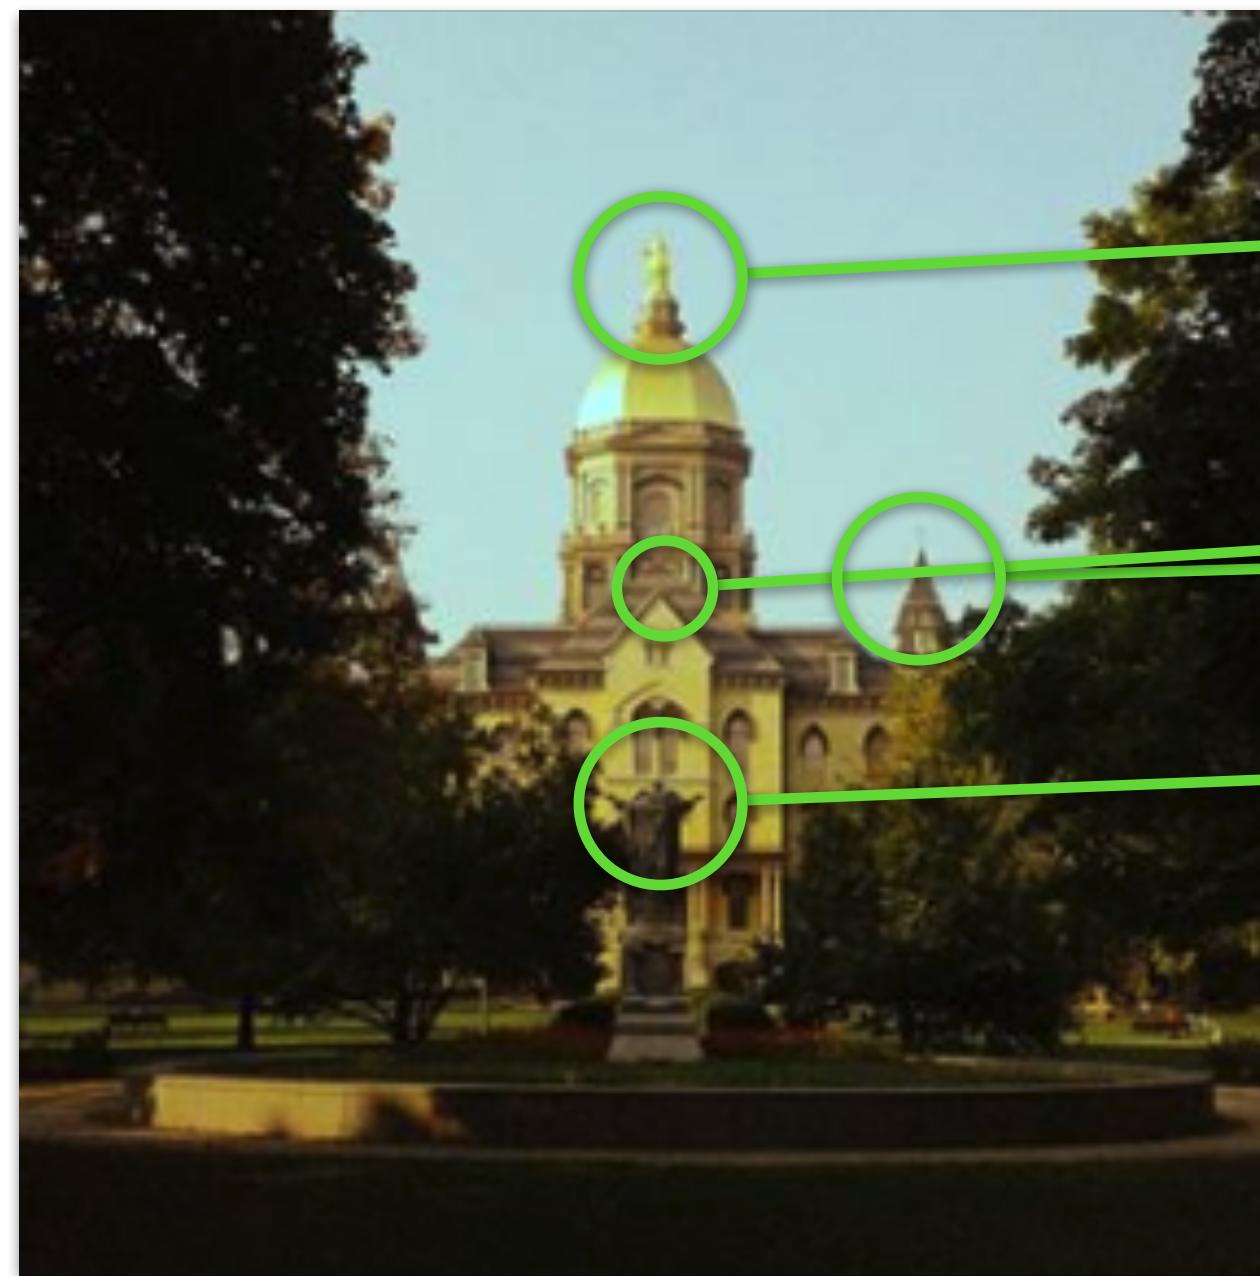


Can a computer system
decide if these images
depict the same building?

Local Features

Why should one use Local Features?

Obtain **robust** and compact image representation for many CV tasks.



Can a computer system
decide if these images
depict the same building?

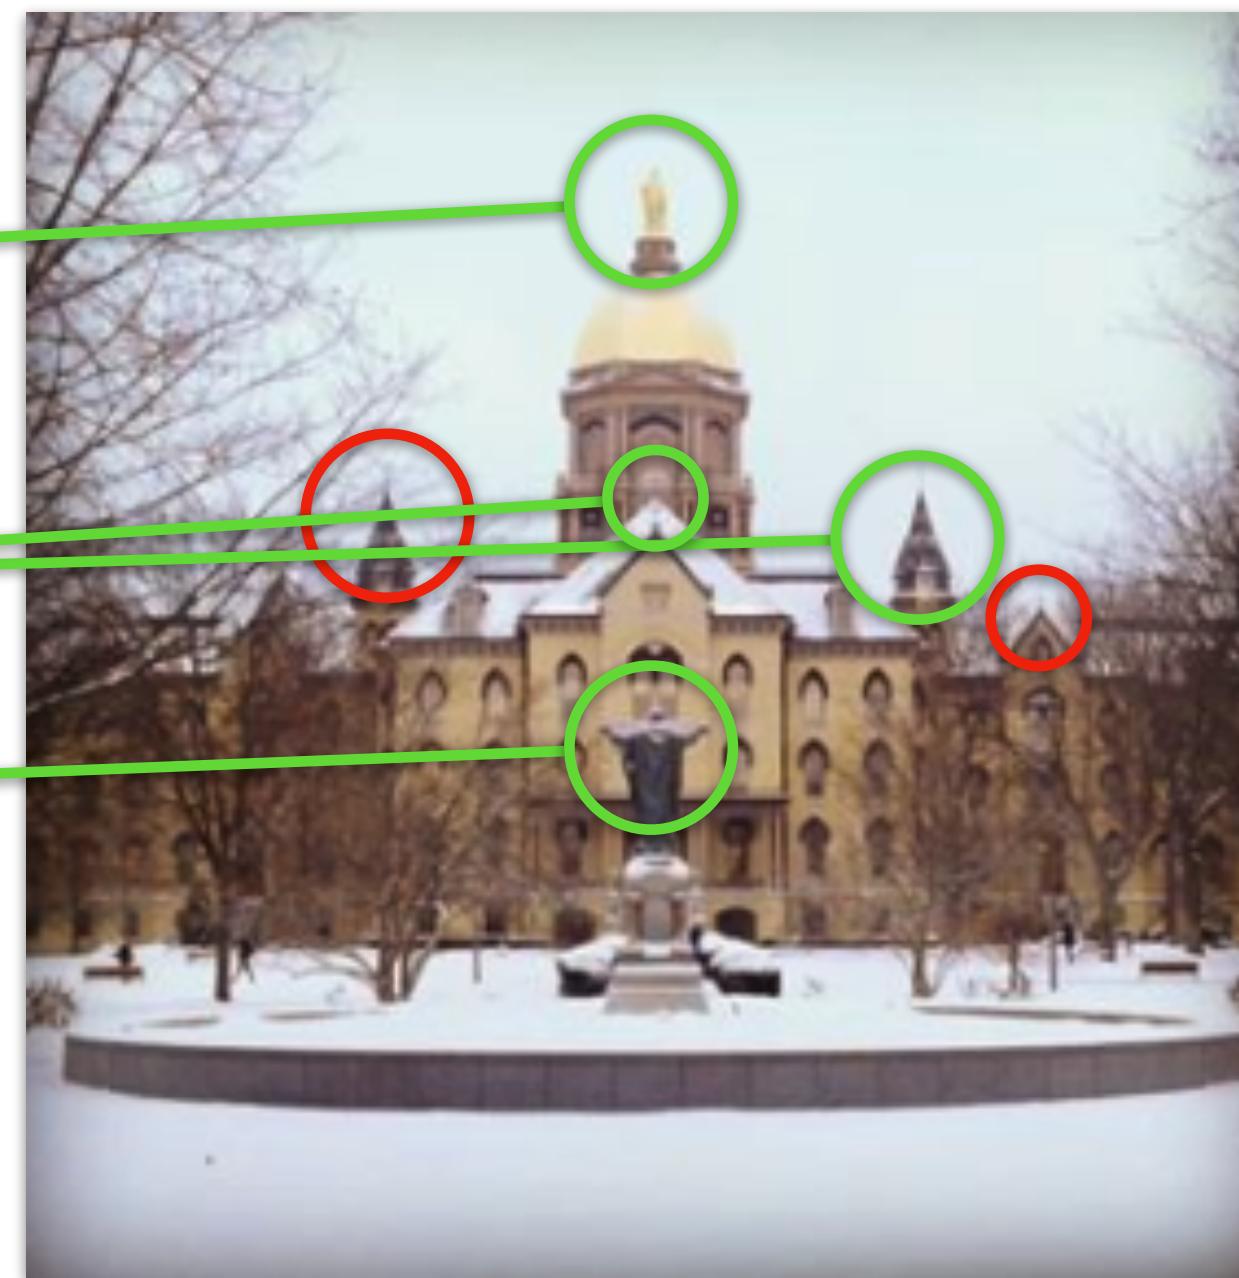
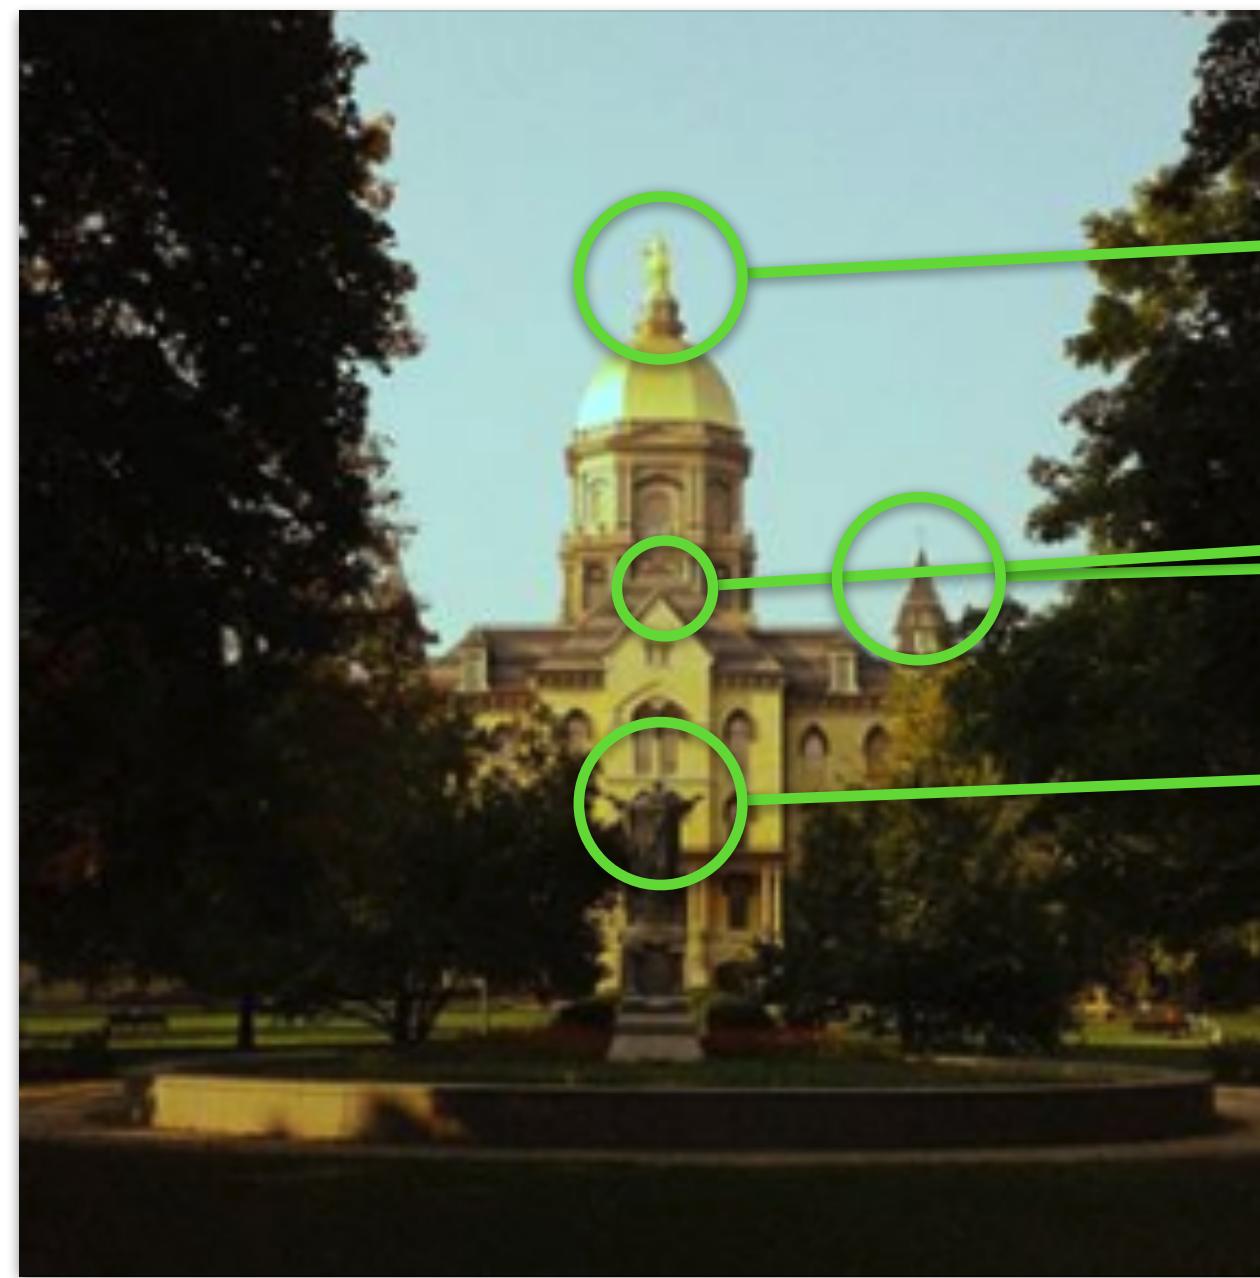
Yes_

In spite of **occlusions**.

Local Features

Why should one use Local Features?

Obtain robust and **compact** image representation for many CV tasks.



Can a computer system
decide if these images
depict the same building?

Yes_

Index the local features
instead of the entire images.

Local Features

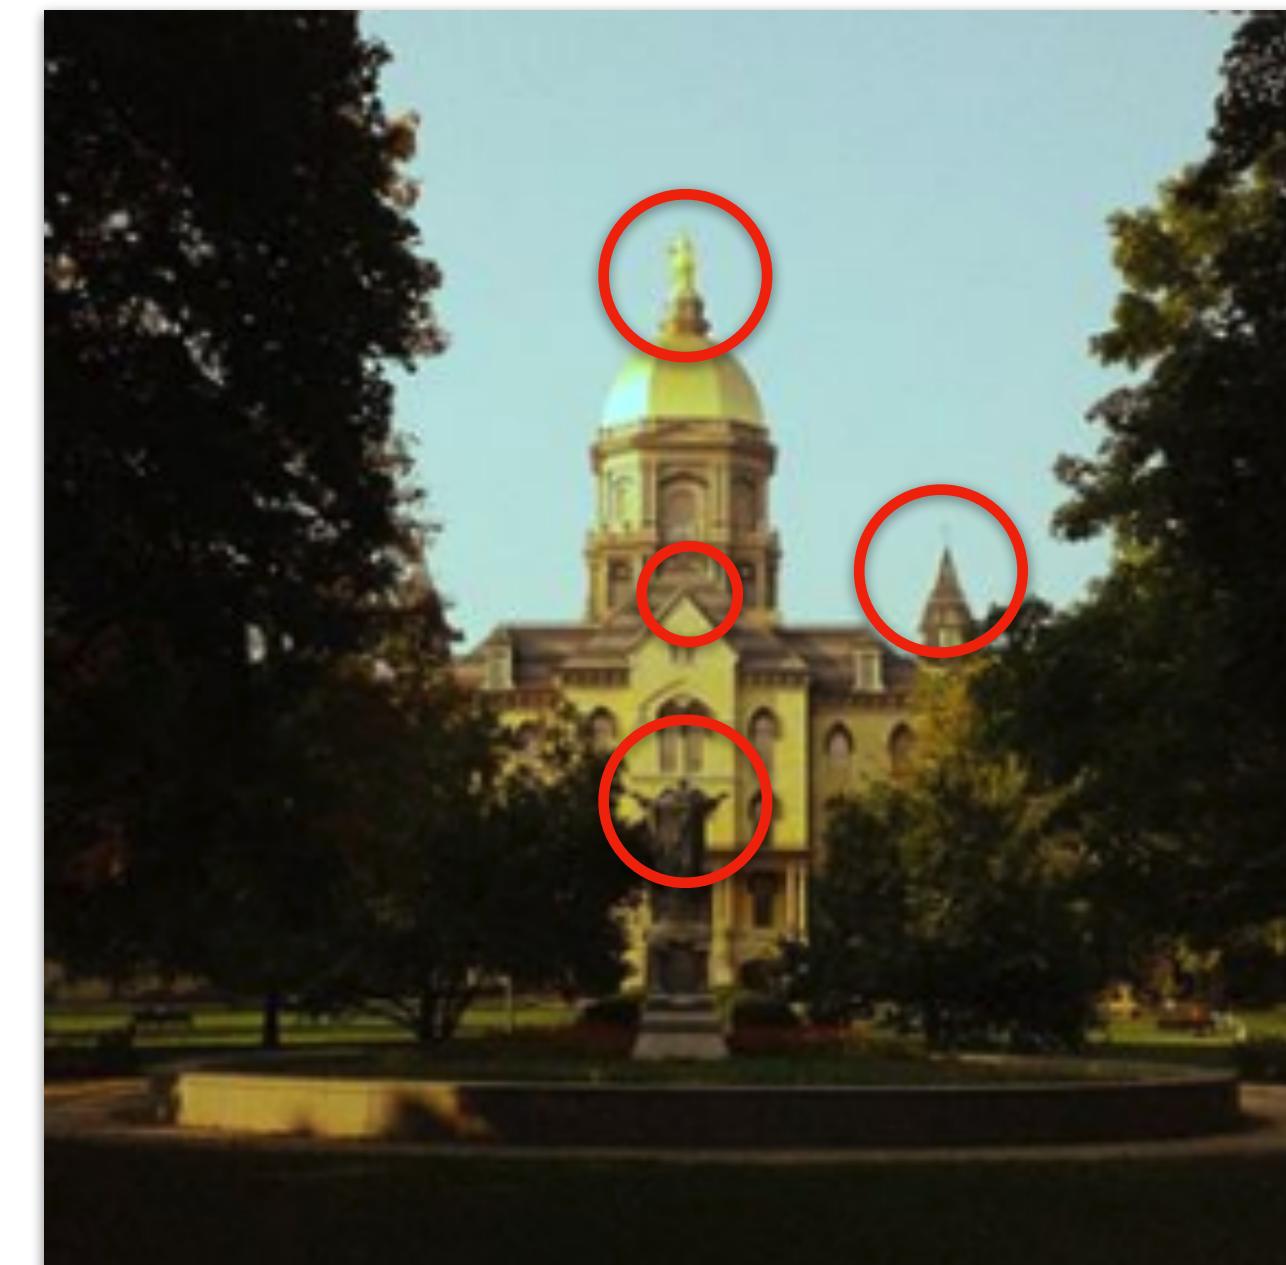
Local Feature Steps

Feature Detection

Interest points, keypoints, or regions of interest are identified within the image.

Desired elements: location (x, y), scale, orientation, strength.

Desired properties: repeatability and distinctiveness.



Local Features

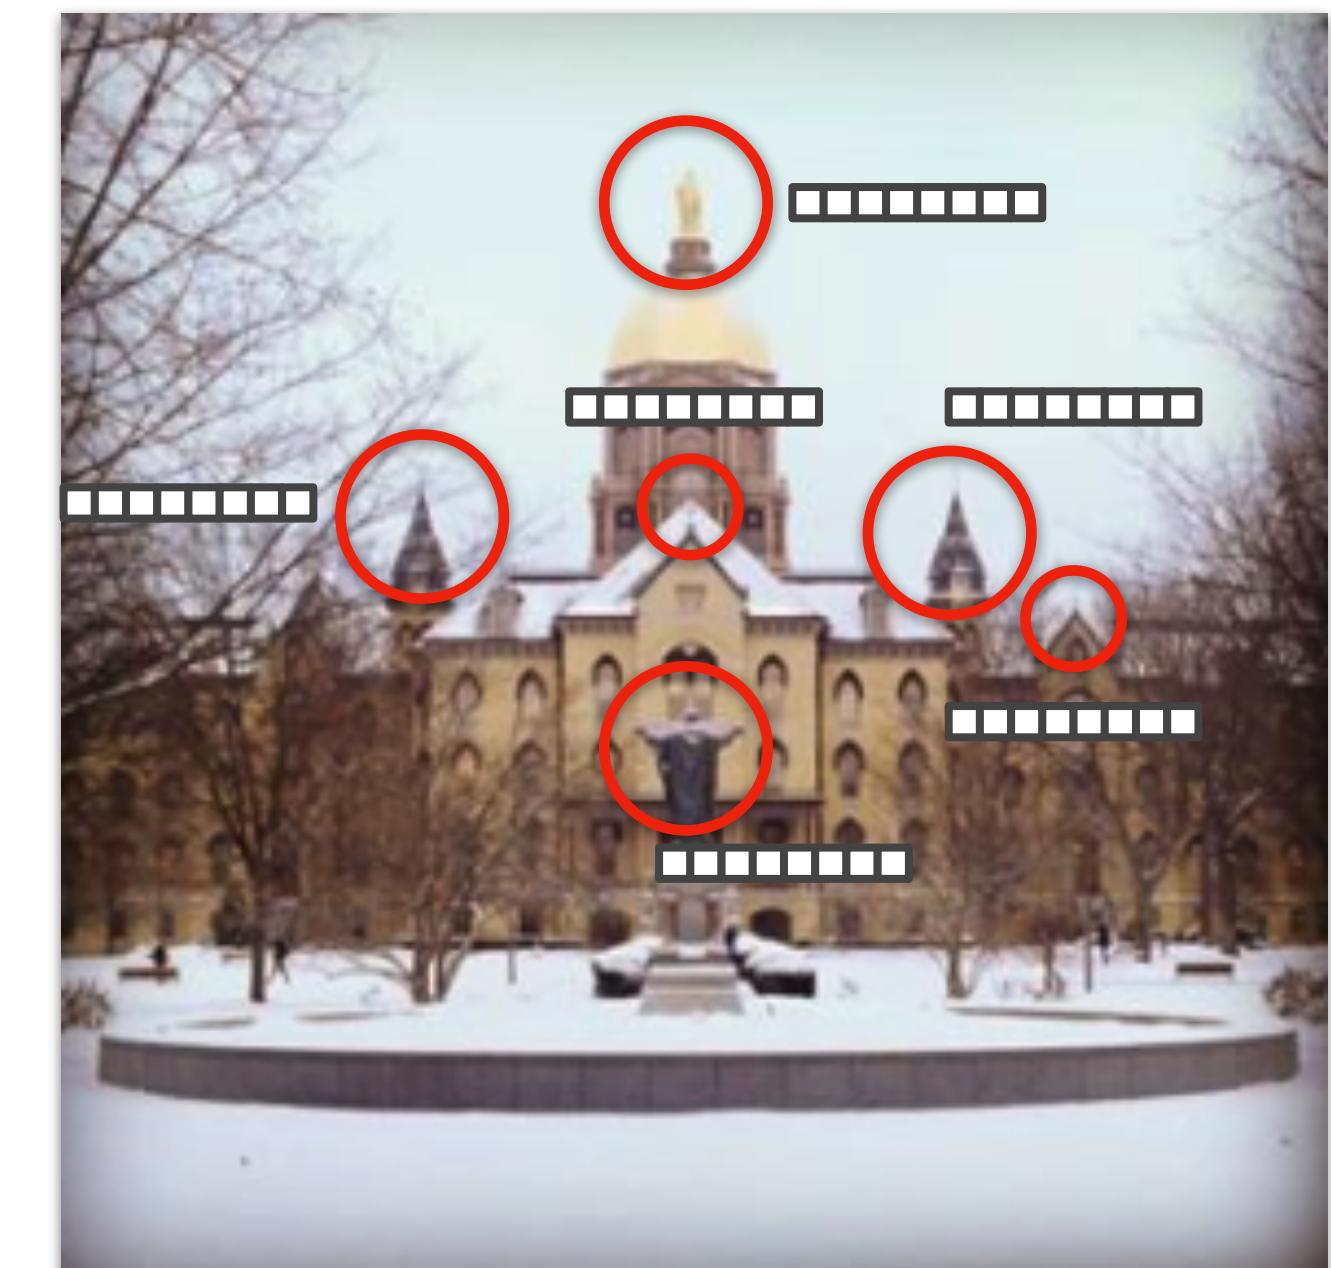
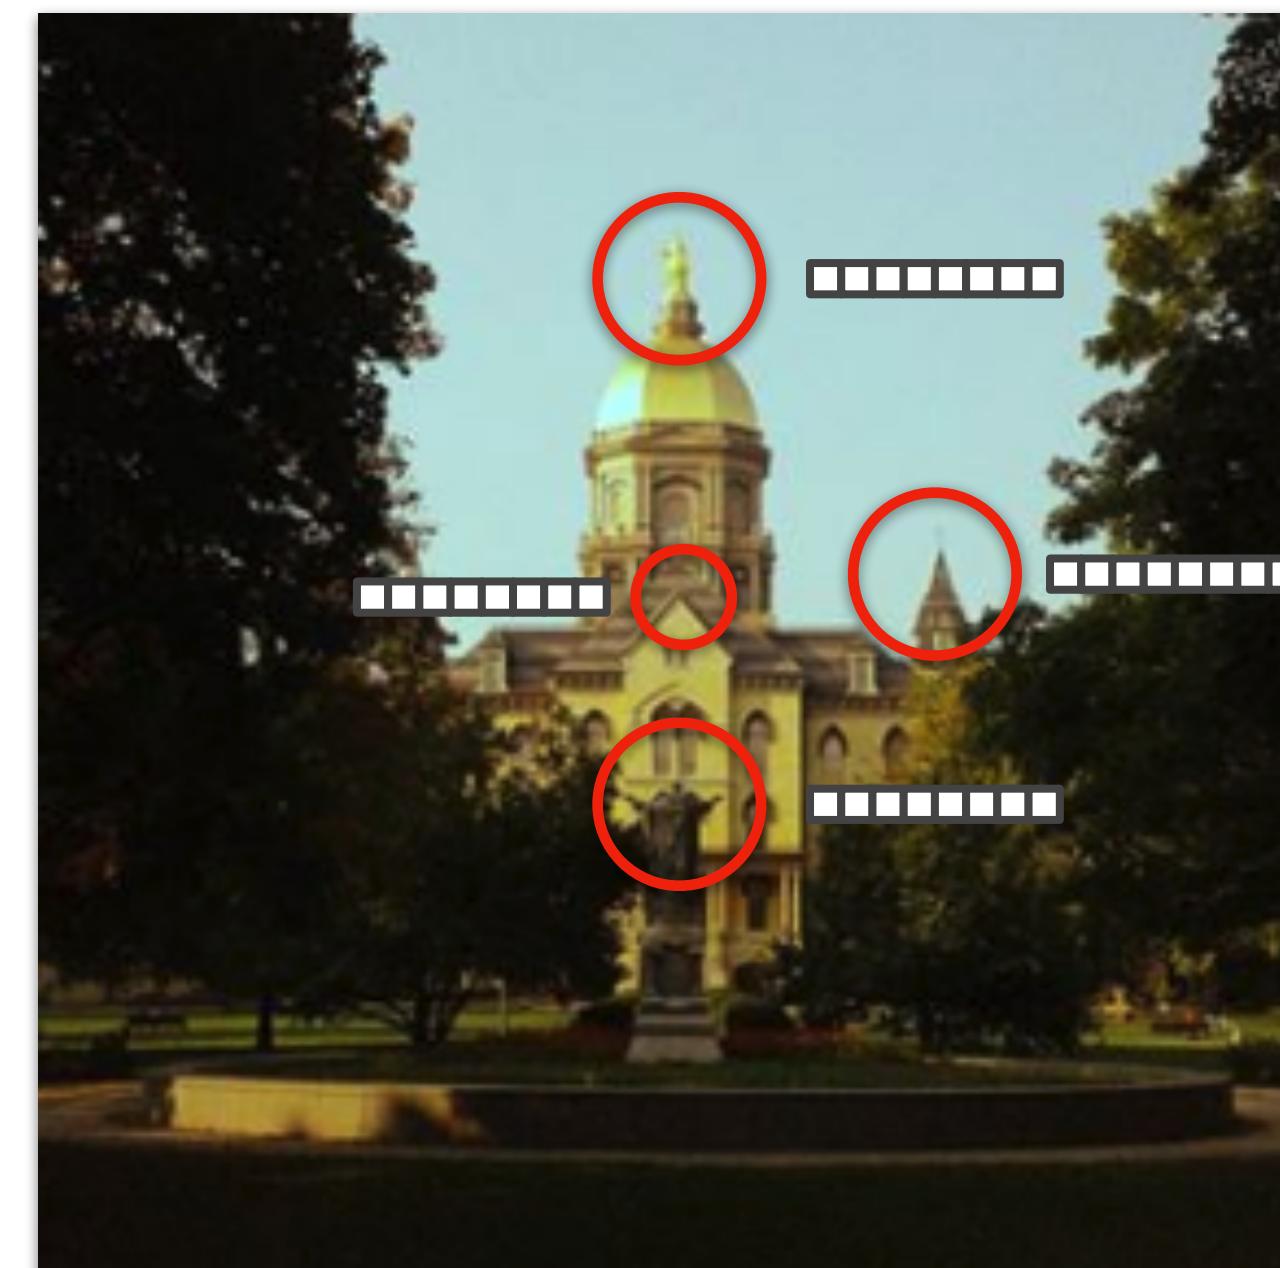
Local Feature Steps

Feature Description

Tensors are computed over the regions of the detected local features.

Desired element: N -Dimensional feature vector.

Desired properties: efficiency and robustness to different capture conditions.



Harris Corner Detector

Focus on Corners

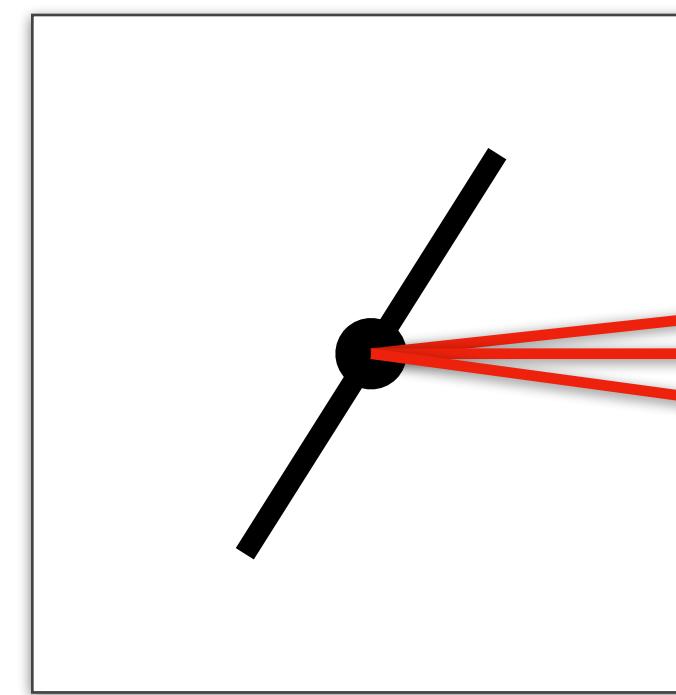


Image 1

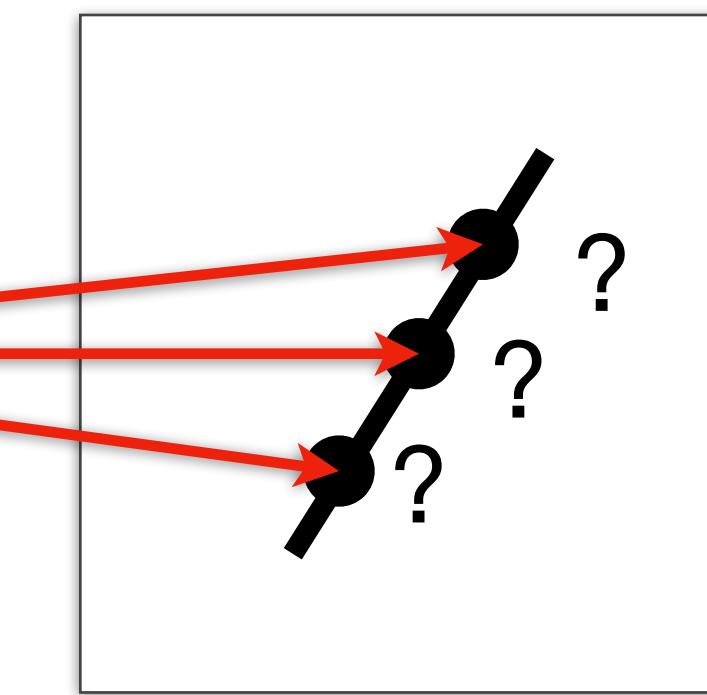
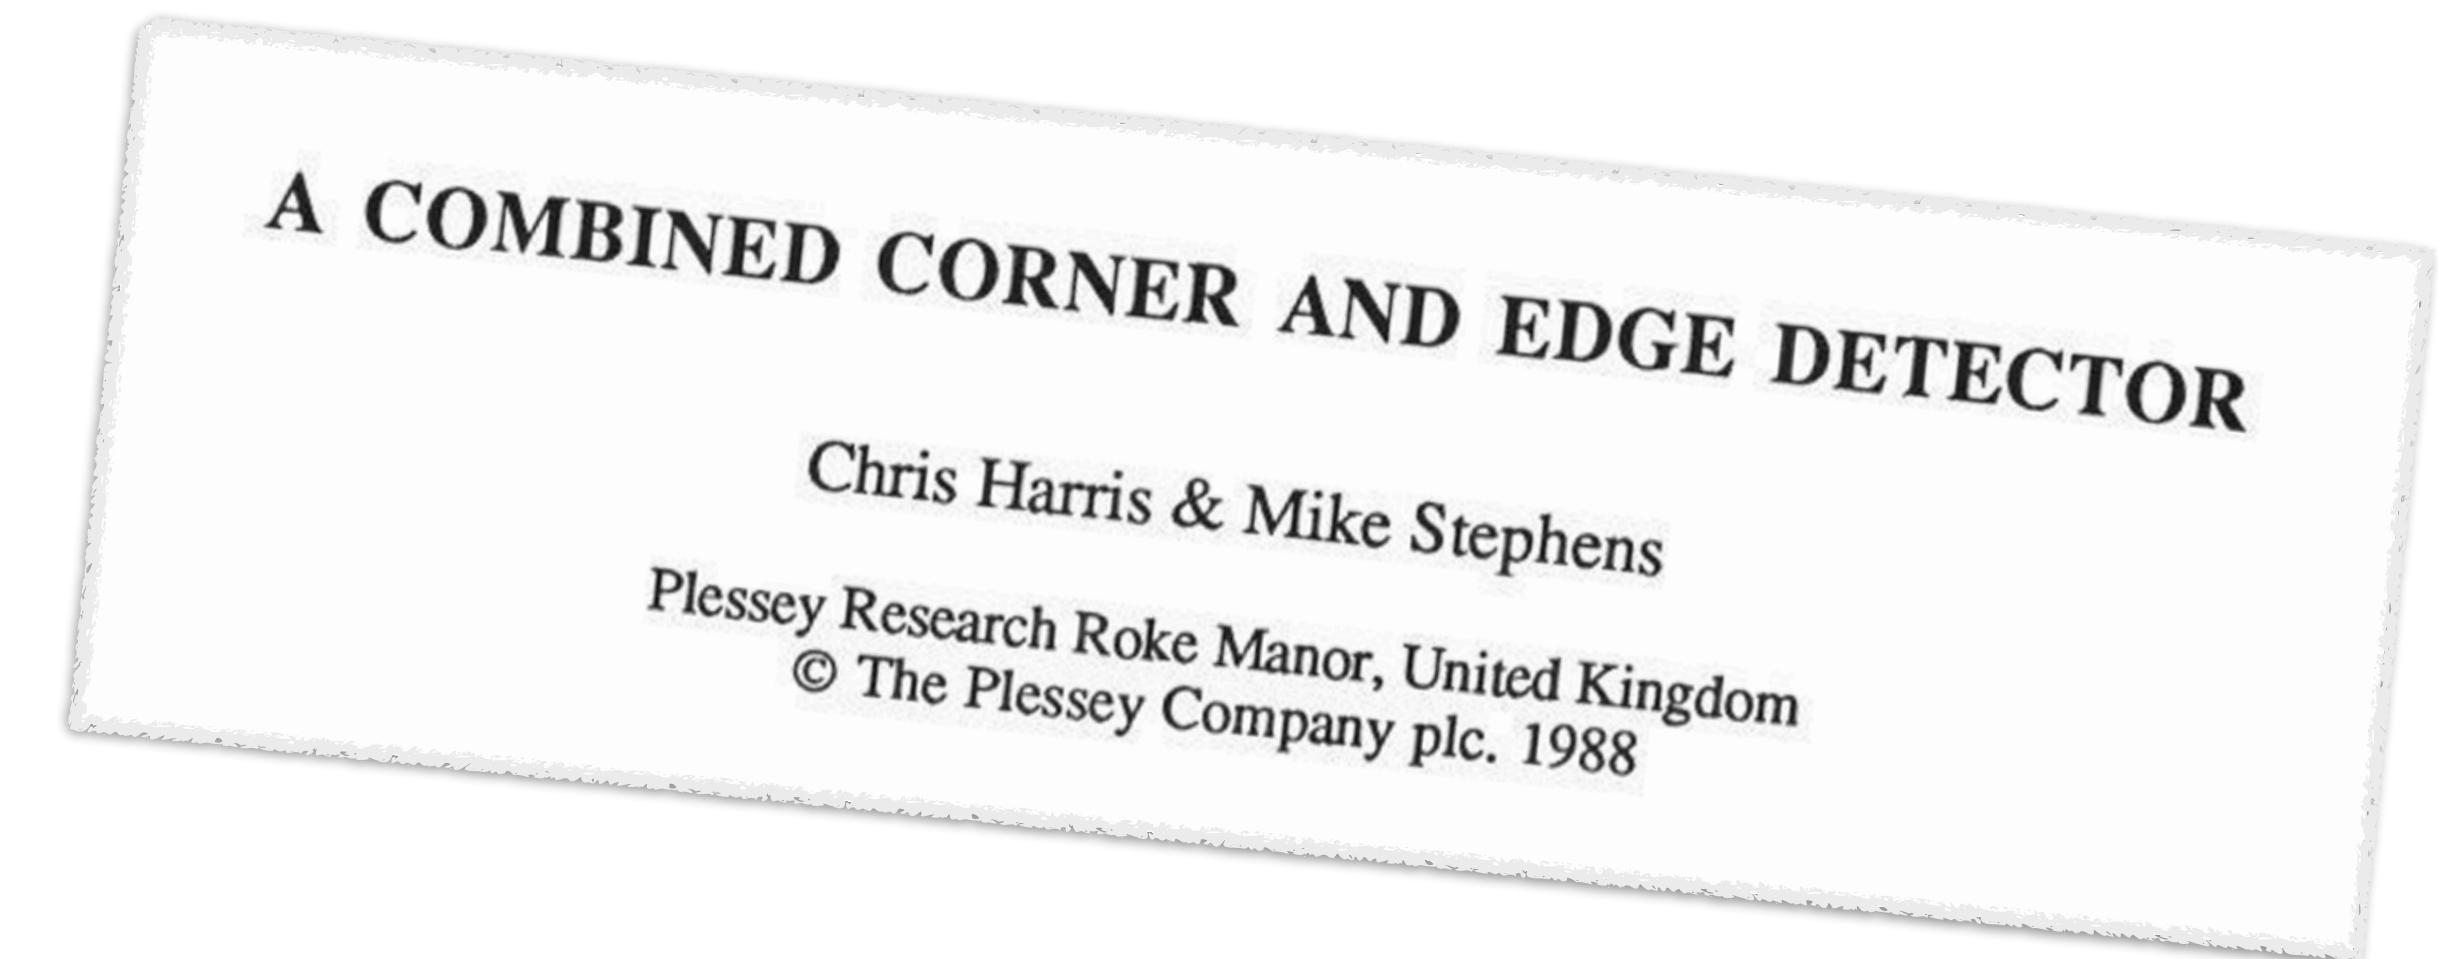


Image 2



Harris Corner Detector

Focus on Corners

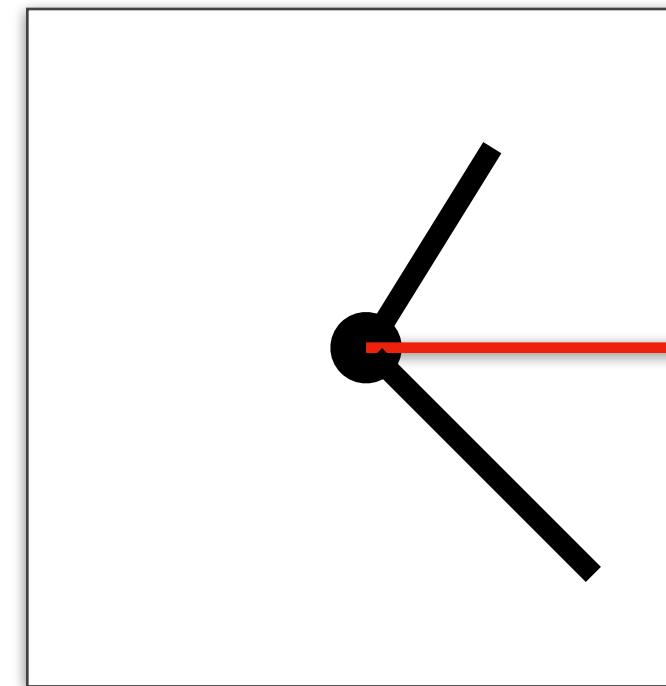


Image 1

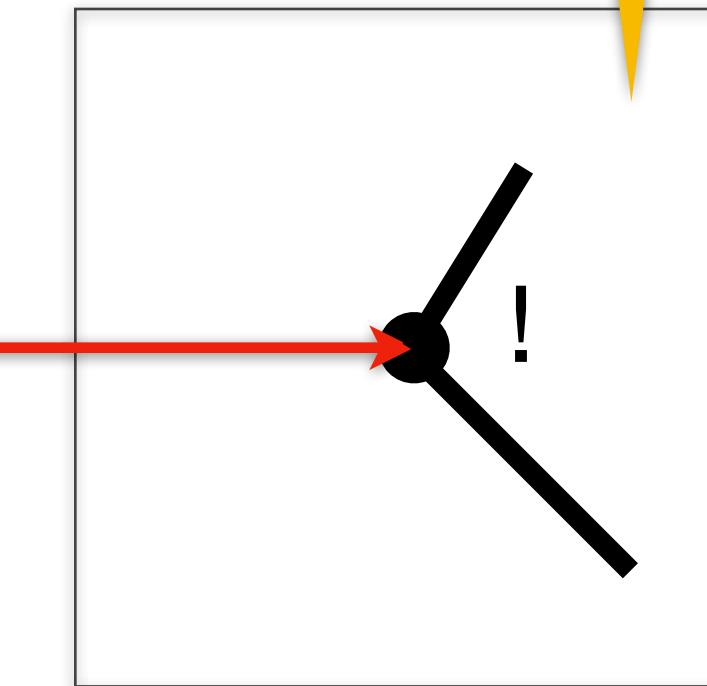


Image 2

Corners are easier to
find and match.

A COMBINED CORNER AND EDGE DETECTOR

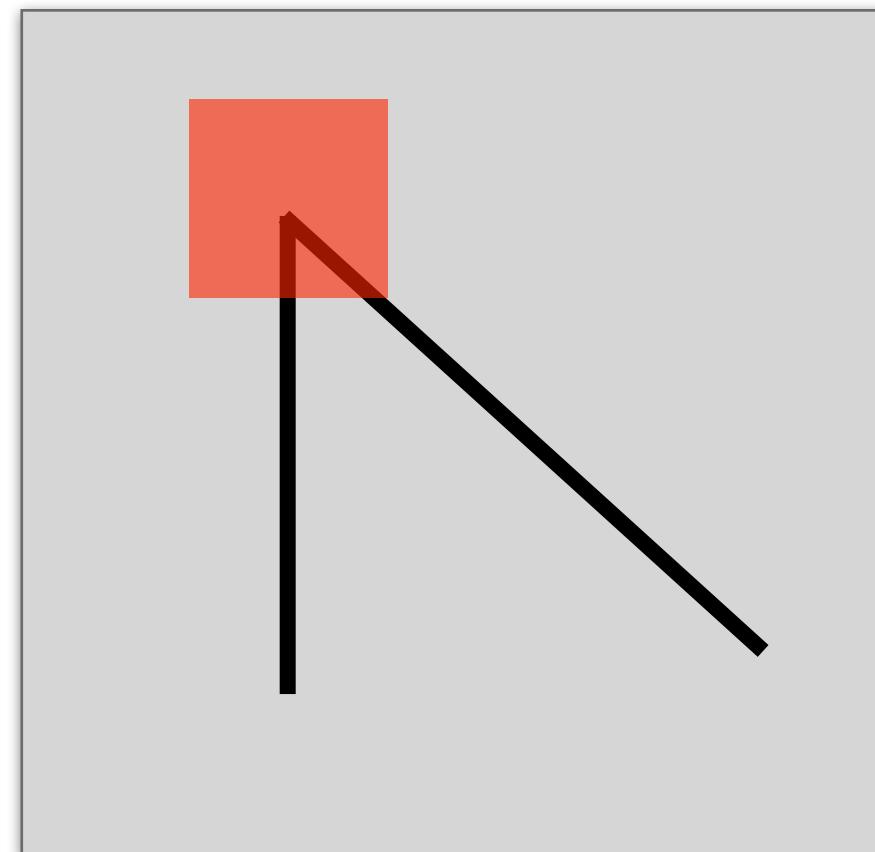
Chris Harris & Mike Stephens
Plessey Research Roke Manor, United Kingdom
© The Plessey Company plc. 1988



LOYOLA
UNIVERSITY CHICAGO

Harris Corner Detector

Focus on Corners

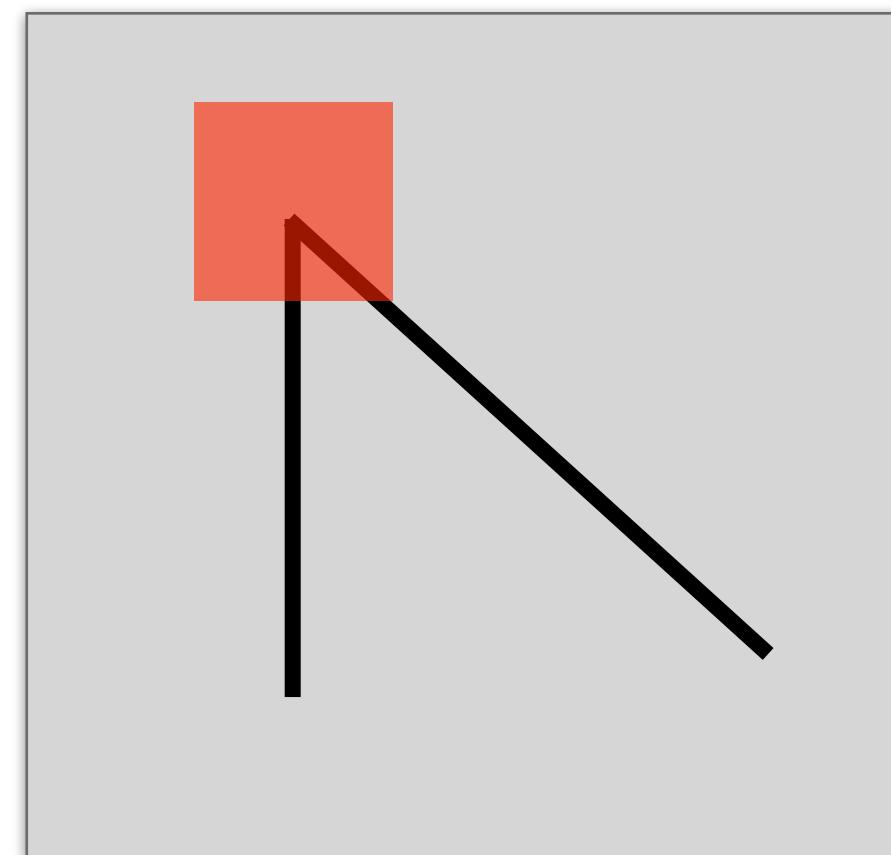


Consider a small $(n \times n)$ -pixel **window** $w(x, y)$ around each pixel $I(x, y)$ of a target image I .

Which of the windows depict corners?

Harris Corner Detector

Focus on Corners



Consider a small $(n \times n)$ -pixel window $w(x, y)$ around each pixel $I(x, y)$ of a target image I .

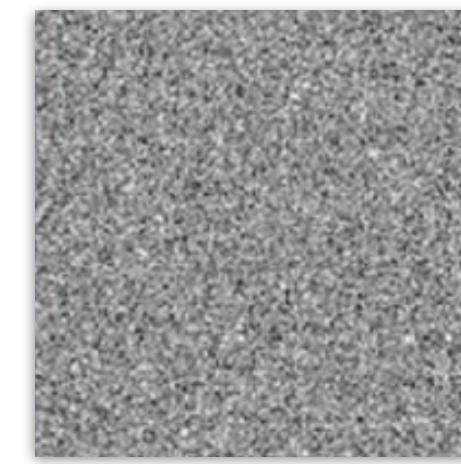
For each window, consider the **image gradients**:

$$\begin{aligned} \text{- In } x \text{ direction: } I_x &= \frac{\delta I}{\delta_x} & \text{- In } y \text{ direction: } I_y &= \frac{\delta I}{\delta_y} \end{aligned}$$

Harris Corner Detector

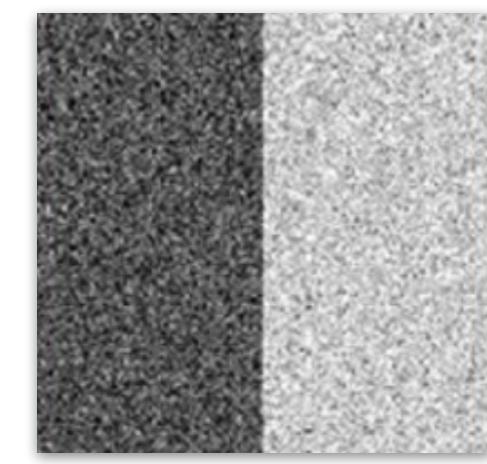
**Image
Gradients**

Image

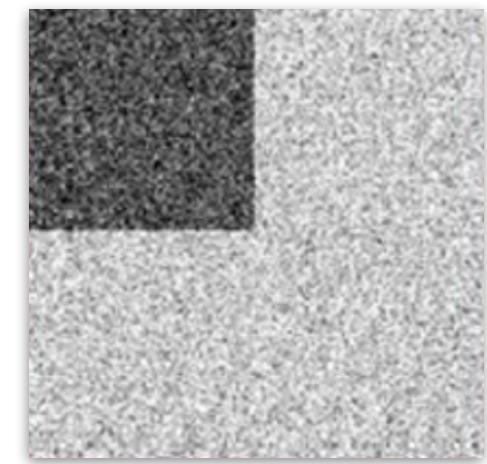


Flat

*Edge
(vertical)*

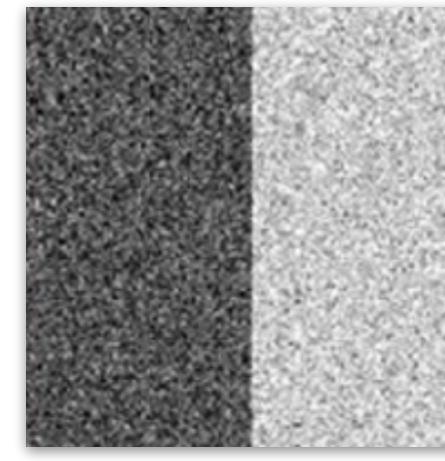
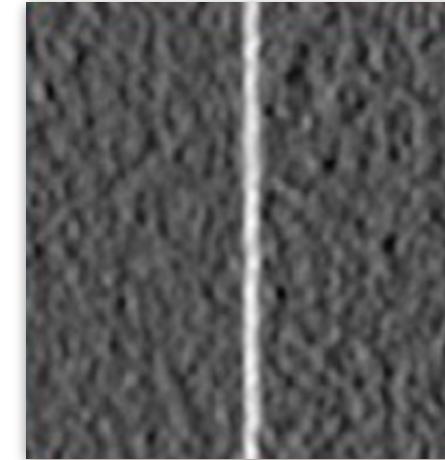
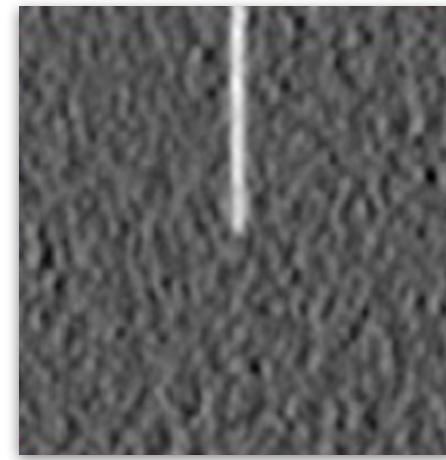


Corner



Harris Corner Detector

**Image
Gradients**

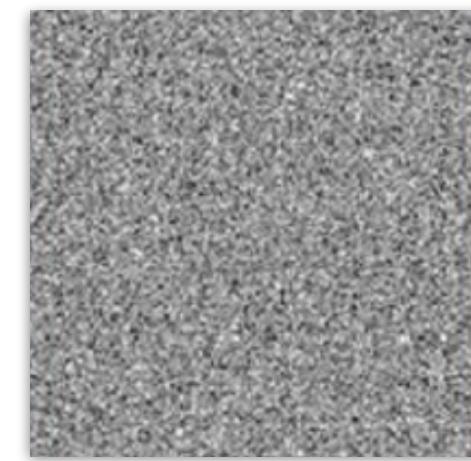
	<i>Flat</i>	<i>Edge (vertical)</i>	<i>Corner</i>	<i>Possible Implementation</i>
Image				
I_x				Image convolution with Sobel filter $\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$

Harris Corner Detector

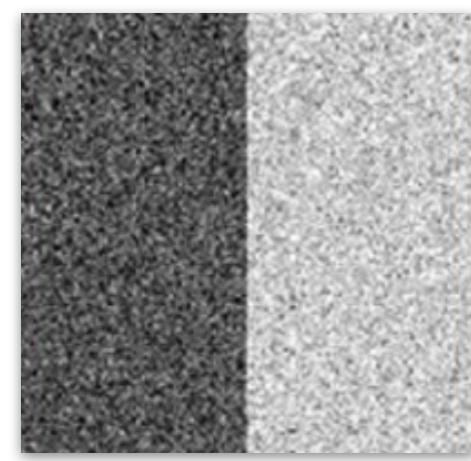
Image Gradients

Flat

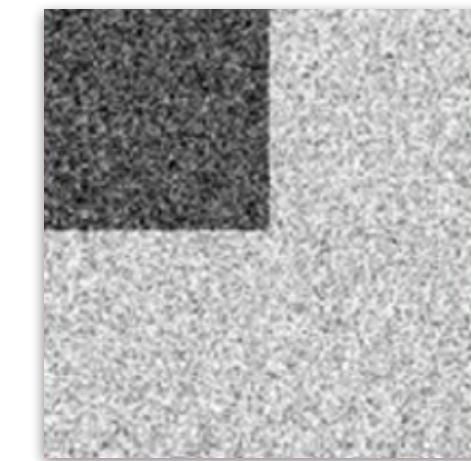
Image



*Edge
(vertical)*



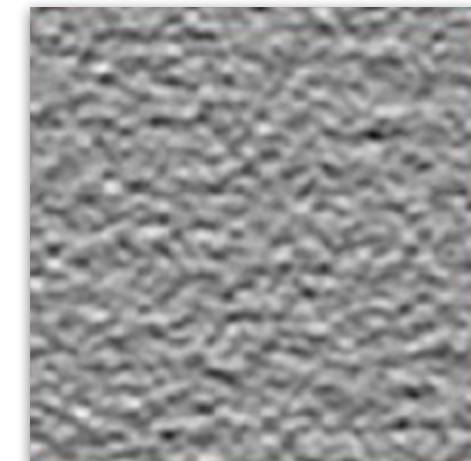
Corner



*Possible
Implementation*

Both I_x and I_y should be large.

I_x



I_y

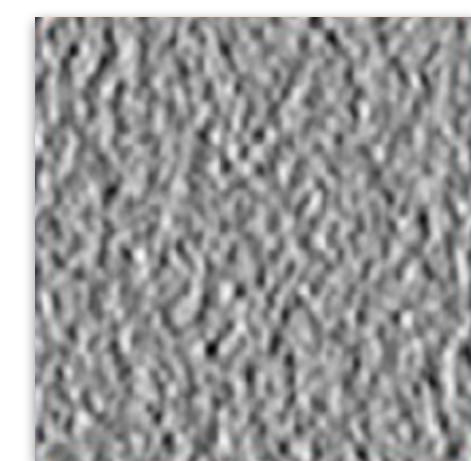


Image convolution with
Sobel filter $\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$

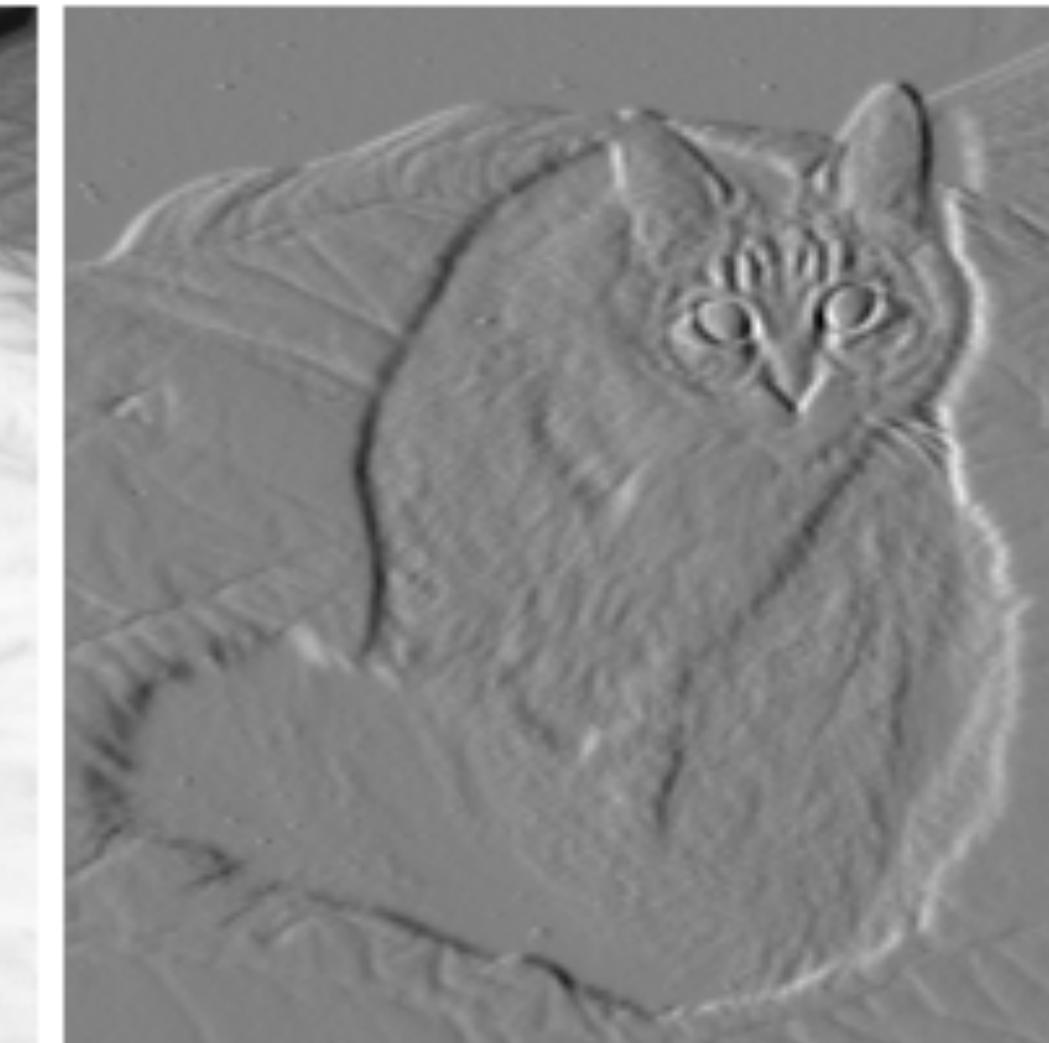
Image convolution with
Sobel filter $\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

Harris Corner Detector

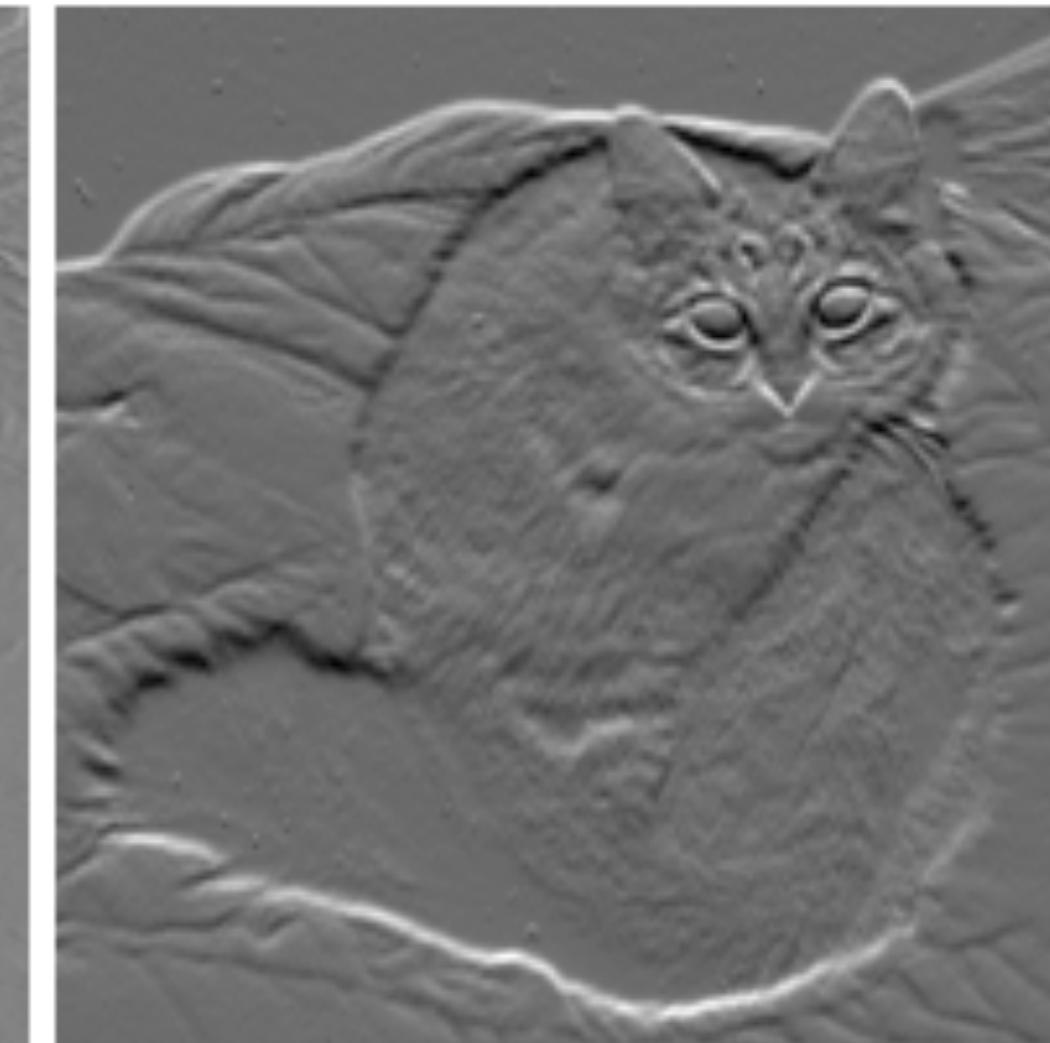
Image
Gradients



Image



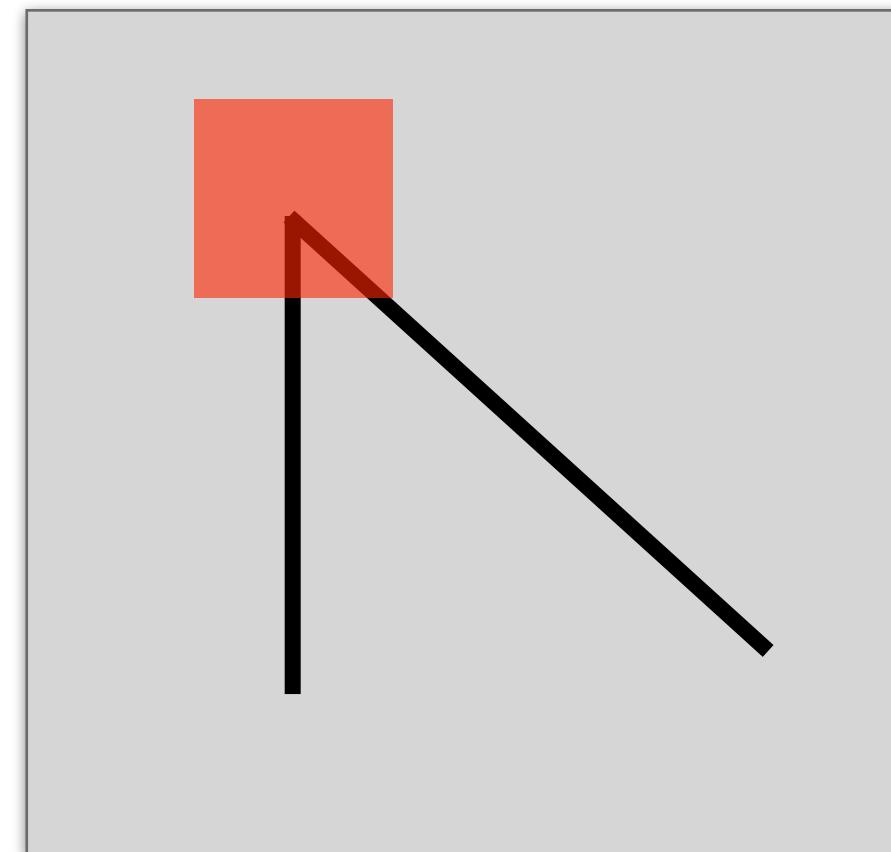
I_x



I_y

Harris Corner Detector

Focus on Corners



Consider a small $(n \times n)$ -pixel window $w(x, y)$ around each pixel $I(x, y)$ of a target image I .

For each window, compute the **structure tensor**:

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix} = \begin{bmatrix} \sum_{x,y} w(x, y) I_x^2 & \sum_{x,y} w(x, y) I_x I_y \\ \sum_{x,y} w(x, y) I_x I_y & \sum_{x,y} w(x, y) I_y^2 \end{bmatrix}$$

Harris Corner Detector

Structure Tensor

(a.k.a. second-moment matrix)

$$M = \begin{bmatrix} \sum_{x,y} w(x,y) I_x^2 & \sum_{x,y} w(x,y) I_x I_y \\ \sum_{x,y} w(x,y) I_x I_y & \sum_{x,y} w(x,y) I_y^2 \end{bmatrix}$$

According to linear algebra principles, the eigenvalues λ_1 and λ_2 of M express the spread of image gradient values in two different directions.

We want both values large, since two different directions define what a corner is.

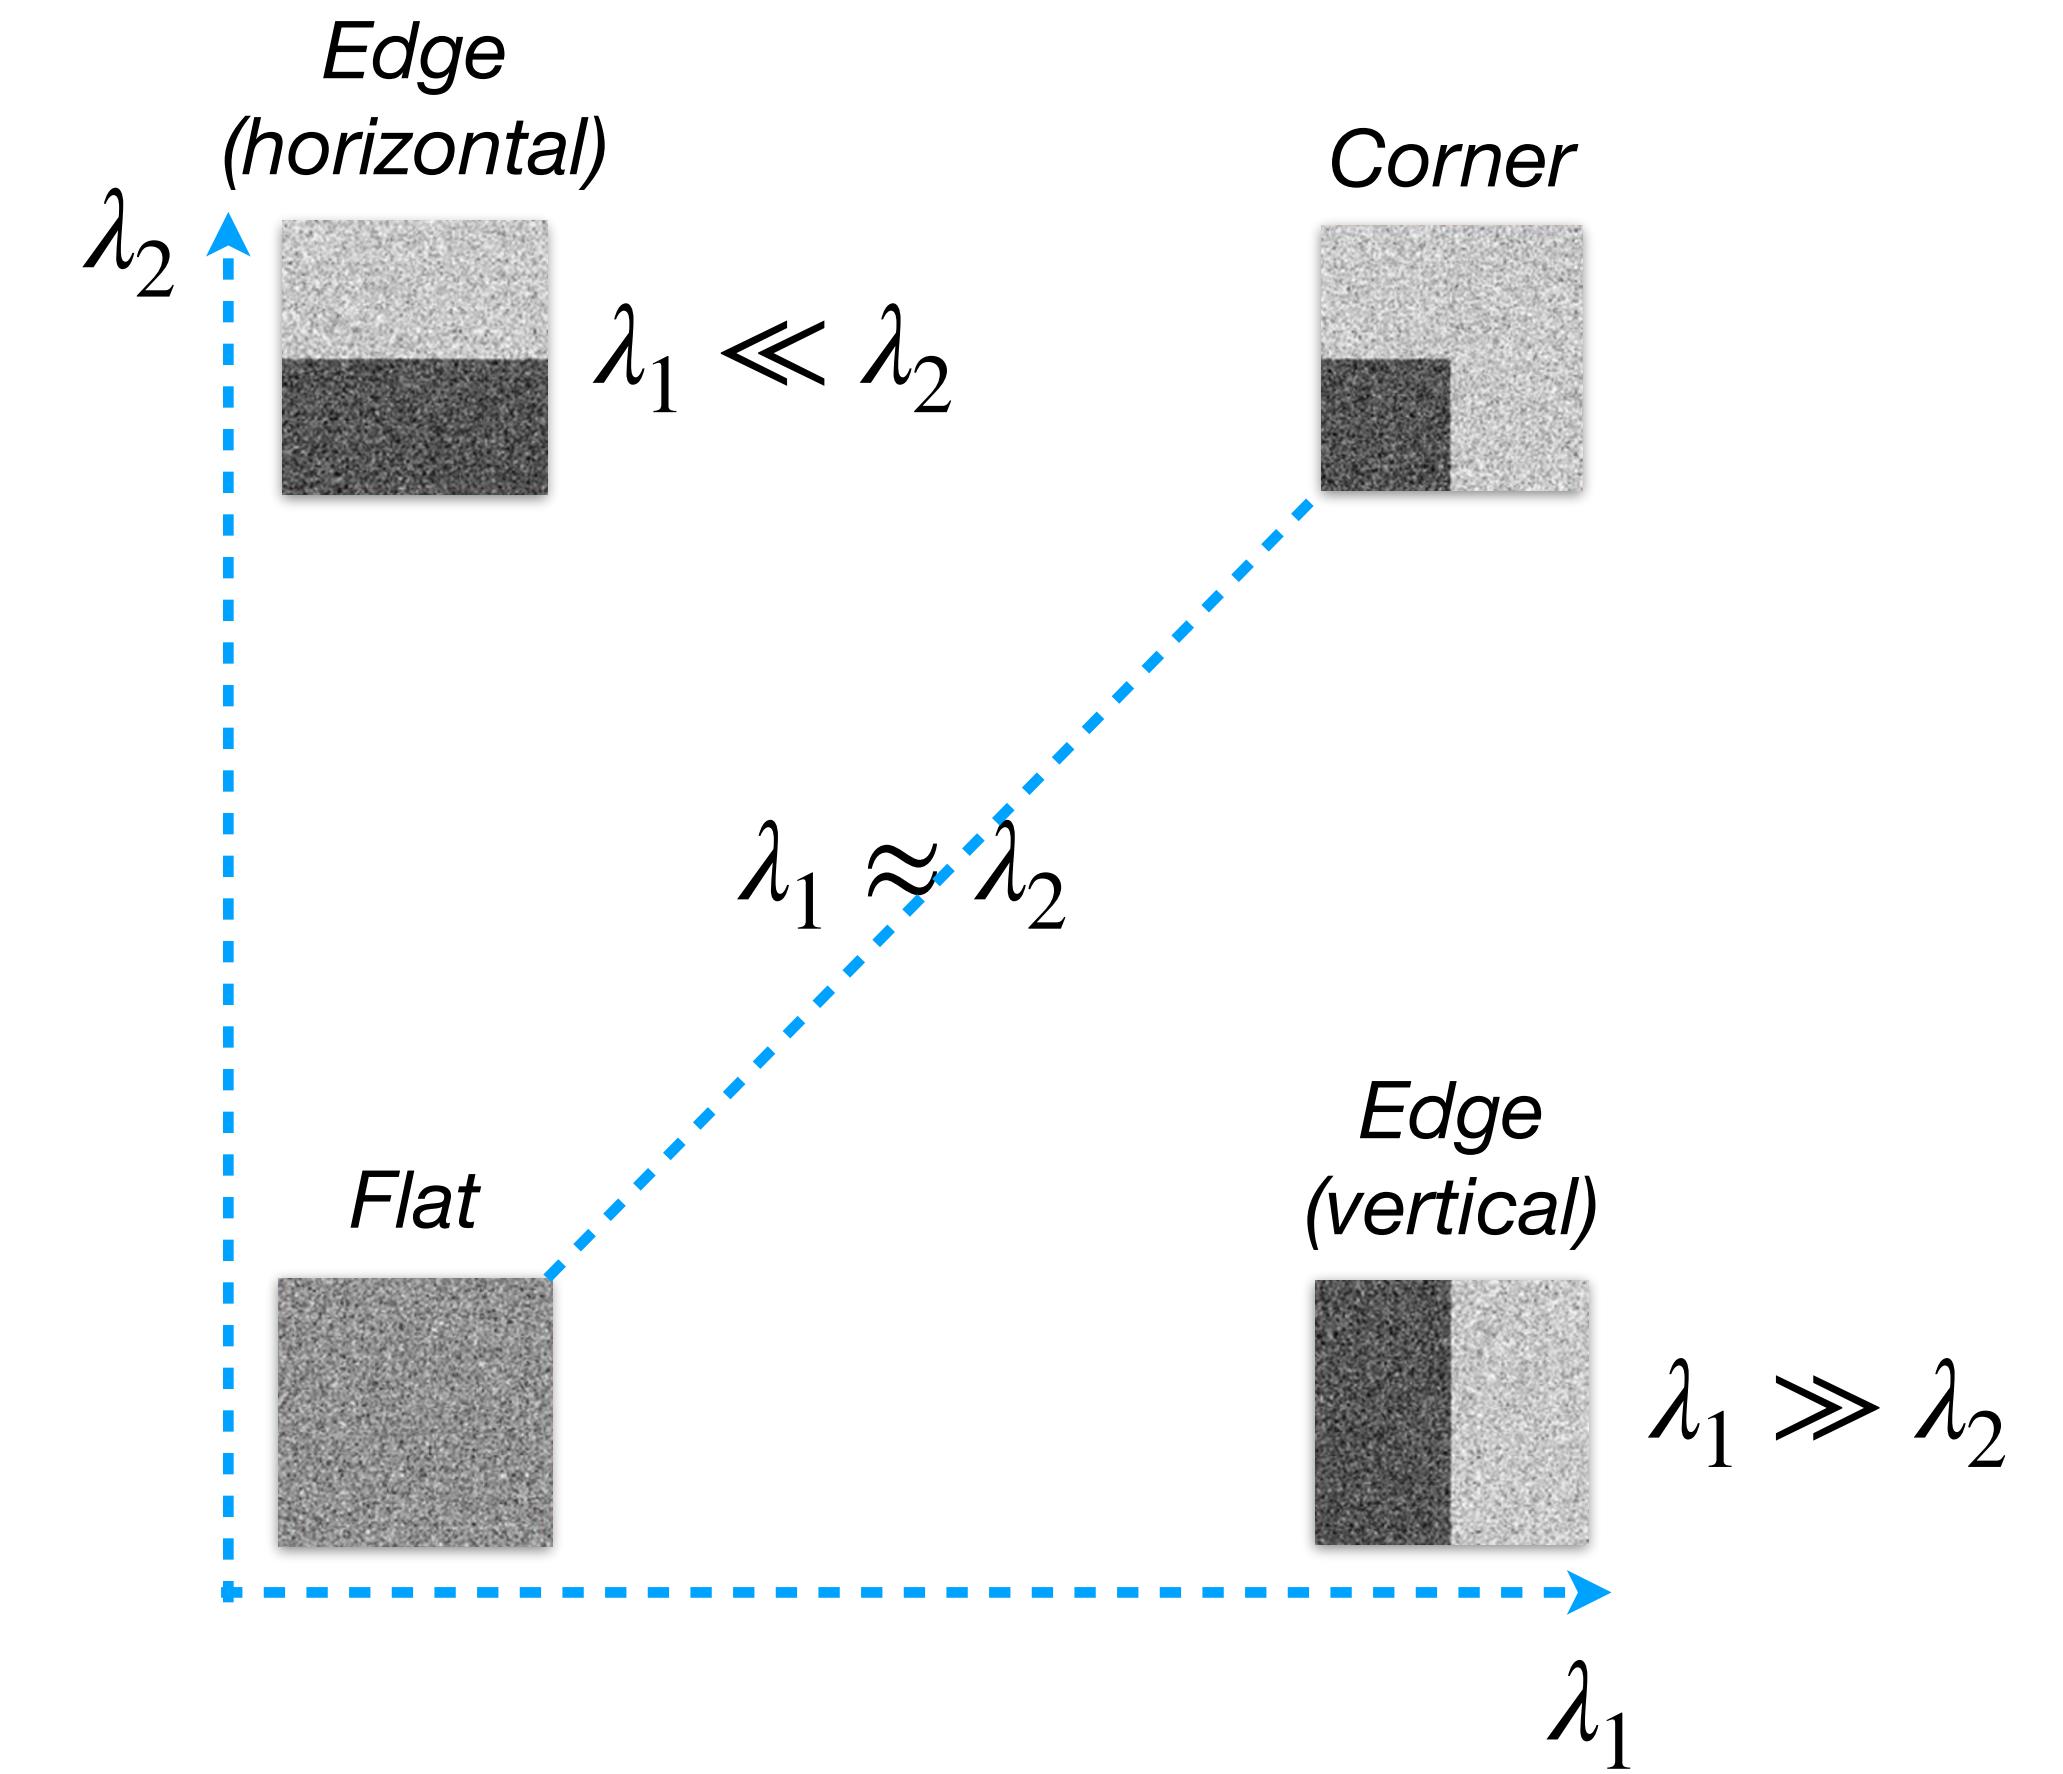
Harris Corner Detector

Structure Tensor

(a.k.a. second-moment matrix)

$$M = \begin{bmatrix} \sum_{x,y} w(x,y) I_x^2 & \sum_{x,y} w(x,y) I_x I_y \\ \sum_{x,y} w(x,y) I_x I_y & \sum_{x,y} w(x,y) I_y^2 \end{bmatrix}$$

Eigenvalues: λ_1 and λ_2



Harris Corner Detector

Structure Tensor

(a.k.a. second-moment matrix)

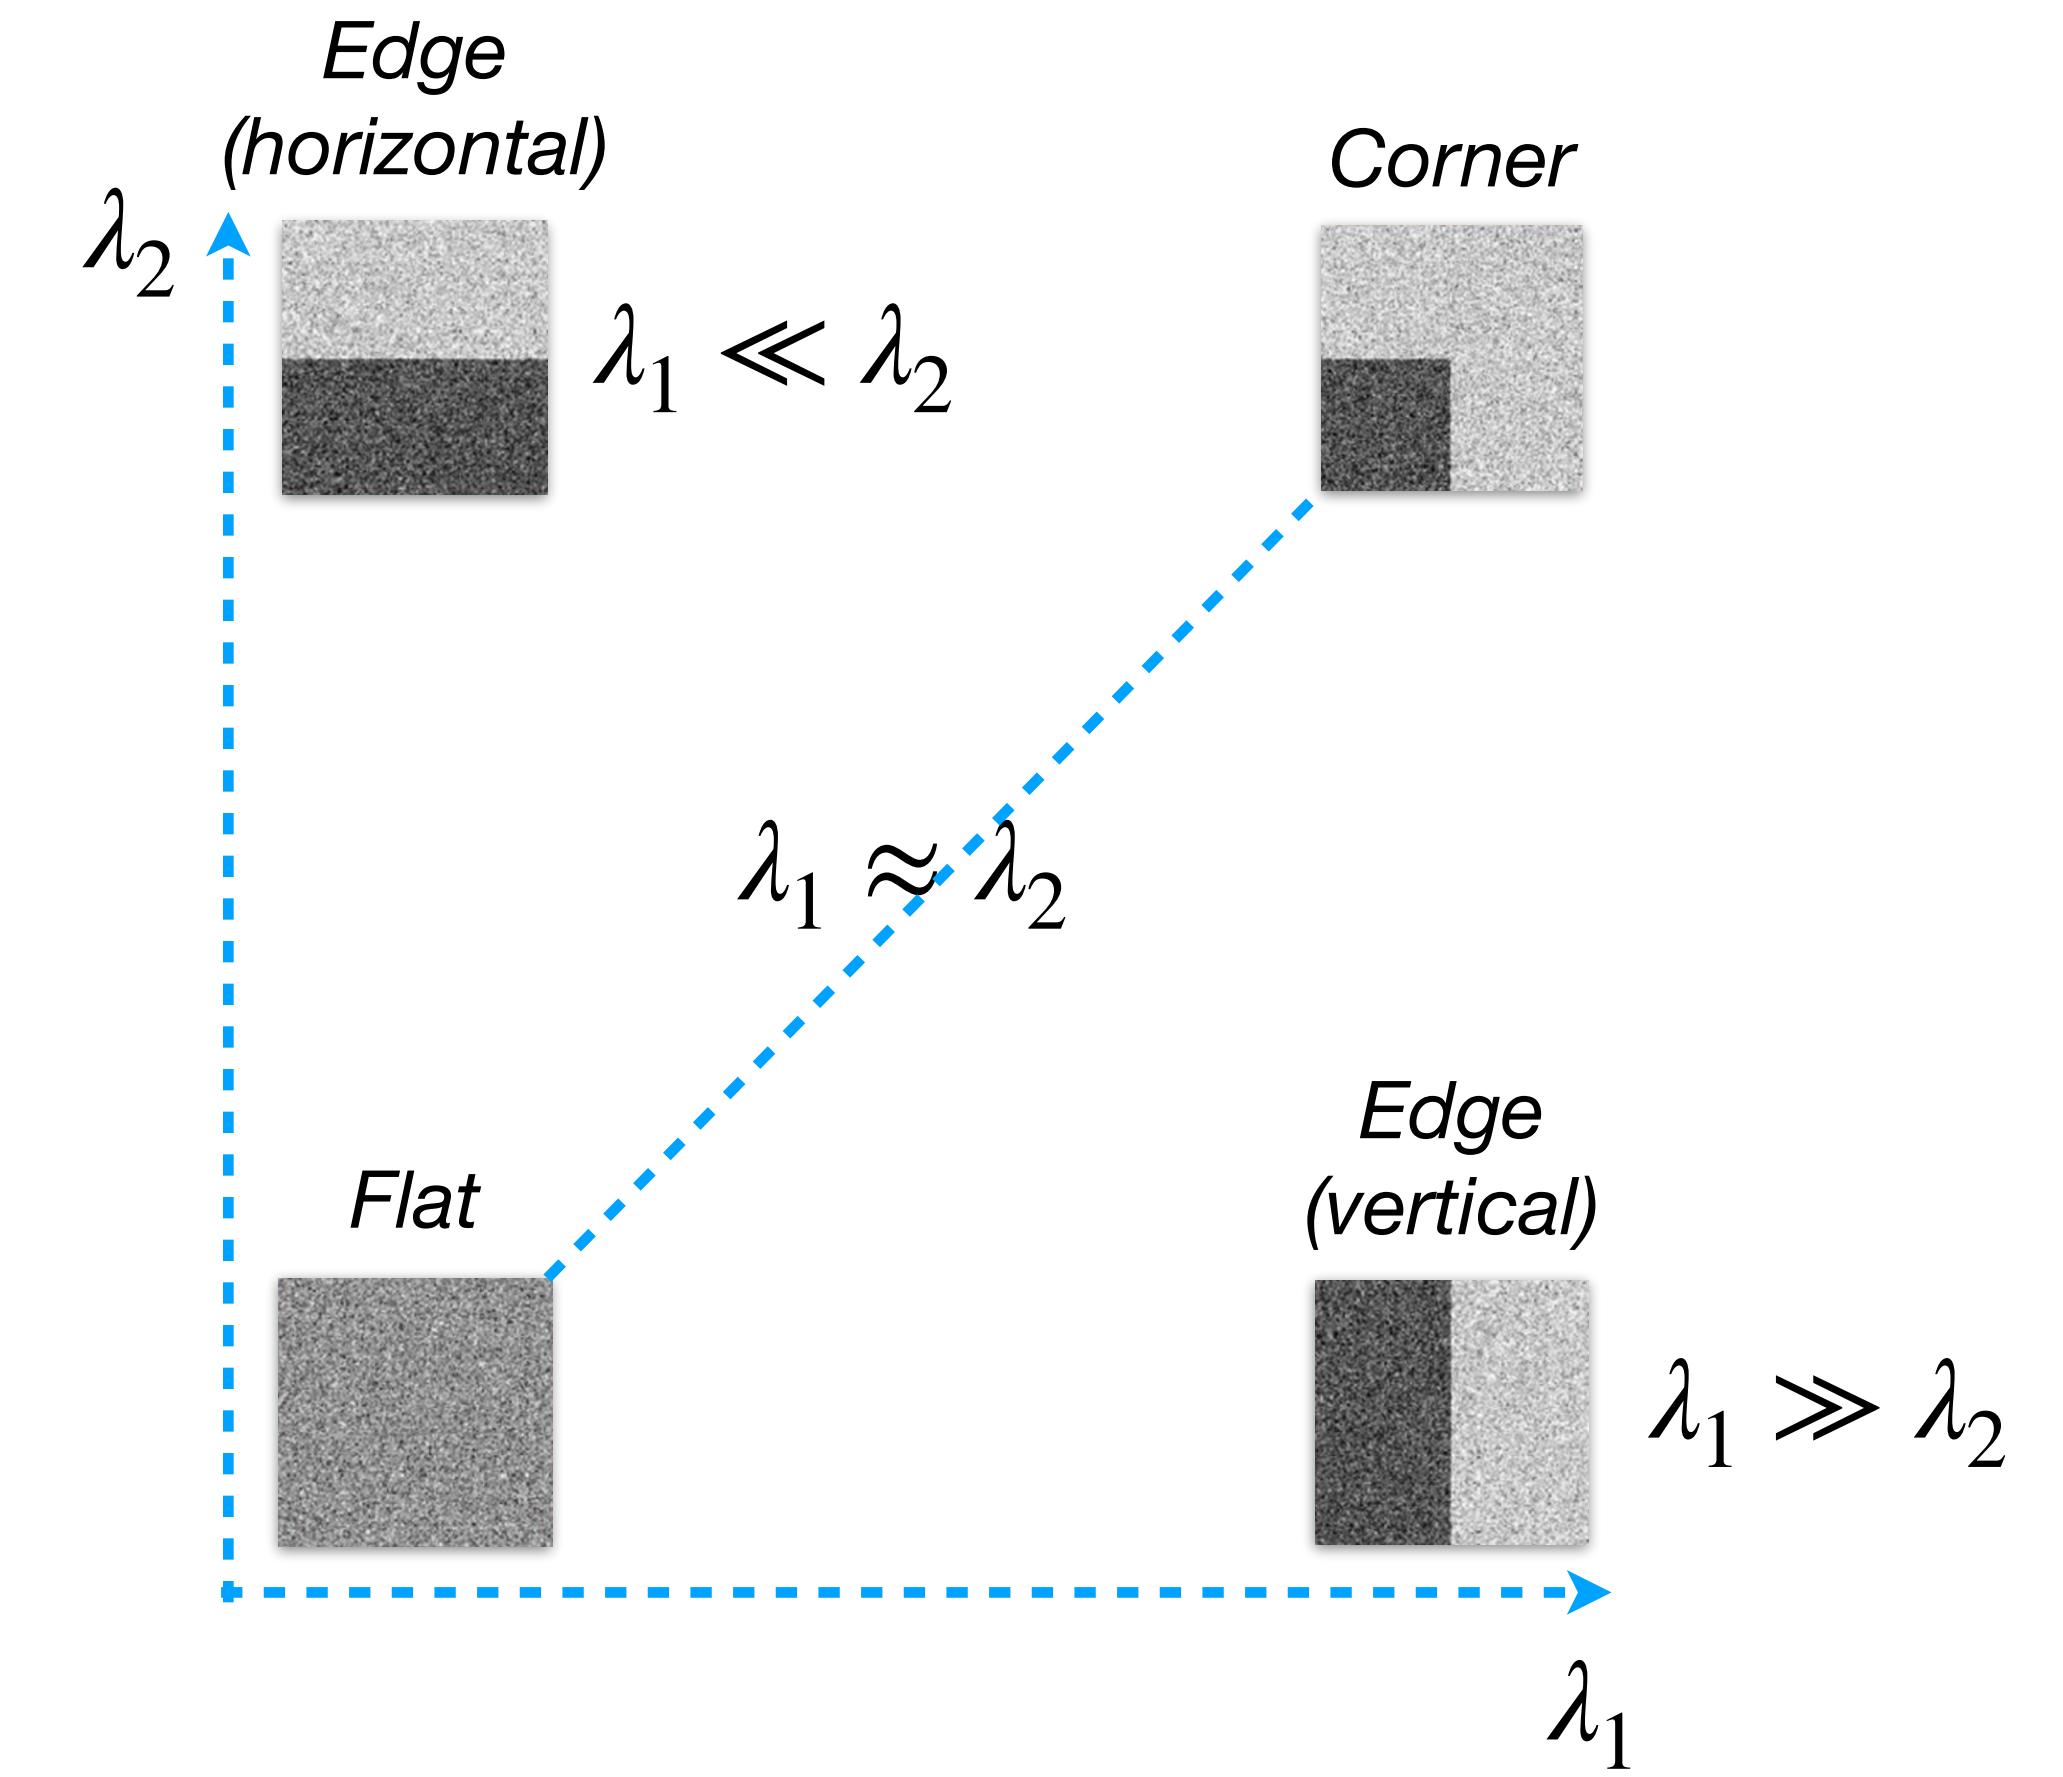
$$M = \begin{bmatrix} \sum_{x,y} w(x,y) I_x^2 & \sum_{x,y} w(x,y) I_x I_y \\ \sum_{x,y} w(x,y) I_x I_y & \sum_{x,y} w(x,y) I_y^2 \end{bmatrix}$$

Eigenvalues: λ_1 and λ_2

Harris and Stephens' idea

Leverage the following properties:

$$\det(M) = \lambda_1 \lambda_2 \text{ and } \text{trace}(M) = \lambda_1 + \lambda_2$$



Harris Corner Detector

Structure Tensor

(a.k.a. second-moment matrix)

$$M = \begin{bmatrix} \sum_{x,y} w(x,y) I_x^2 & \sum_{x,y} w(x,y) I_x I_y \\ \sum_{x,y} w(x,y) I_x I_y & \sum_{x,y} w(x,y) I_y^2 \end{bmatrix}$$

Eigenvalues: λ_1 and λ_2

“Cornerness” Score R

No need to compute λ_1 and λ_2 explicitly
but leverage:

$$\det(M) = \lambda_1 \lambda_2 \text{ and } \text{trace}(M) = \lambda_1 + \lambda_2$$

$$R = \det(M) - k(\text{trace}(M))^2$$



Harris Corner Detector

Structure Tensor

(a.k.a. second-moment matrix)

$$M = \begin{bmatrix} \sum_{x,y} w(x,y)I_x^2 & \sum_{x,y} w(x,y)I_x I_y \\ \sum_{x,y} w(x,y)I_x I_y & \sum_{x,y} w(x,y)I_y^2 \end{bmatrix}$$

$R \gg 0$: we have a corner! (λ_1 and λ_2 are both large)

$R < 0$: we have an edge, so ignore. ($\lambda_1 \gg \lambda_2$ or vice versa)

$|R| \approx 0$: we have a flat region, so ignore.

“Cornerness” Score R

No need to compute λ_1 and λ_2 explicitly
but leverage:

$$\det(M) = \lambda_1 \lambda_2 \text{ and } \text{trace}(M) = \lambda_1 + \lambda_2$$

$$R = \det(M) - k(\text{trace}(M))^2$$

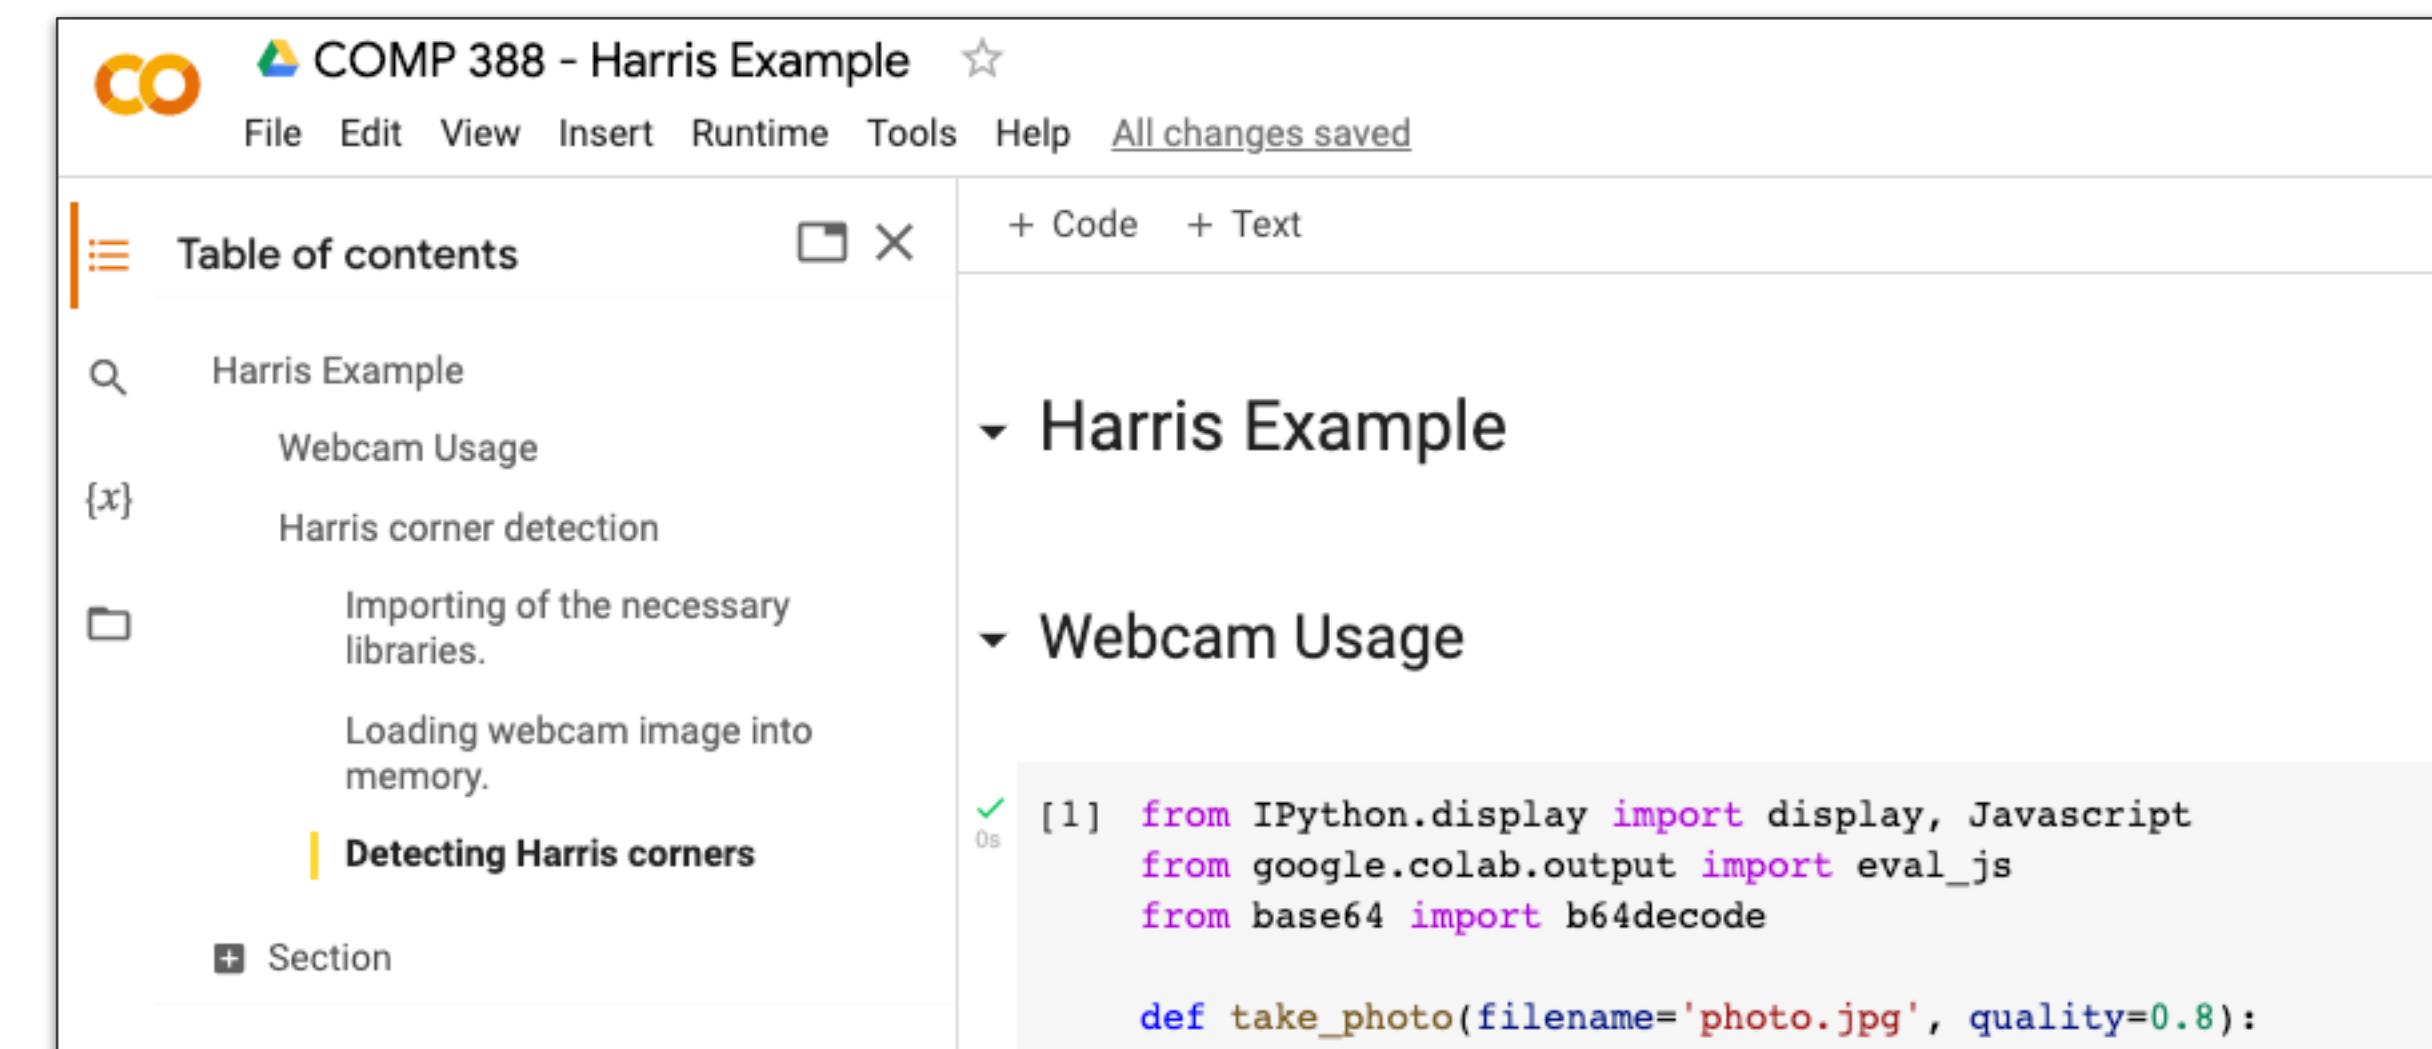
Handcrafted value:
 $k = 0.04$

Harris Corner Detector

Example



<https://bit.ly/3qRBt2U>



The screenshot shows a Google Colab notebook interface. The title bar says "COMP 388 - Harris Example". The left sidebar has a "Table of contents" section with the following items:

- Harris Example
- Webcam Usage
- Harris corner detection
- Importing of the necessary libraries.
- Loading webcam image into memory.
- Detecting Harris corners
- Section

The main area shows a code cell with the following Python code:

```
[1] from IPython.display import display, Javascript
from google.colab.output import eval_js
from base64 import b64decode

def take_photo(filename='photo.jpg', quality=0.8):
```

Scale Invariant Feature Transform (SIFT)

How to match the same content captured with different resolutions?

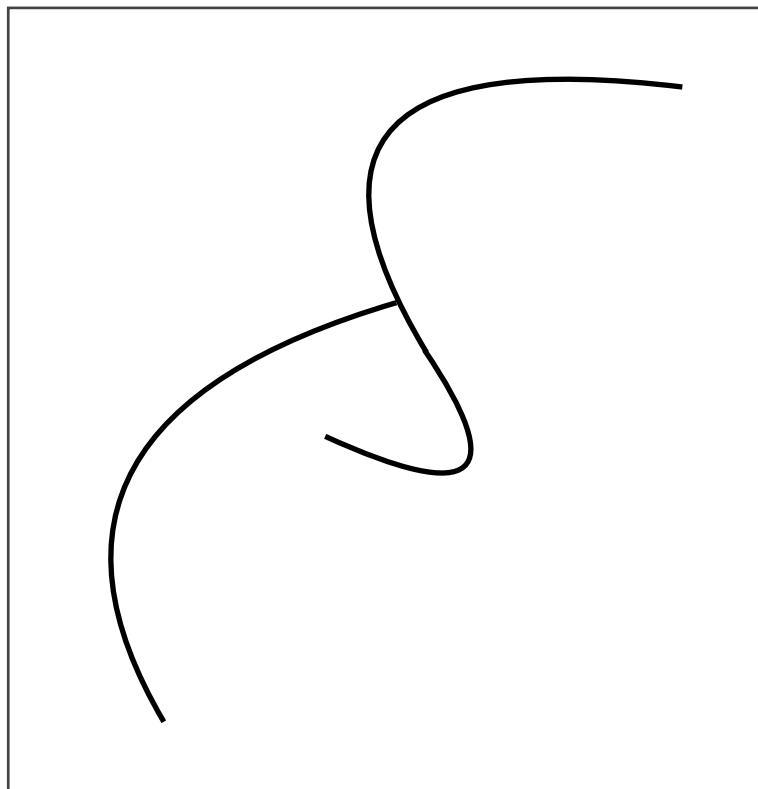


Image 1

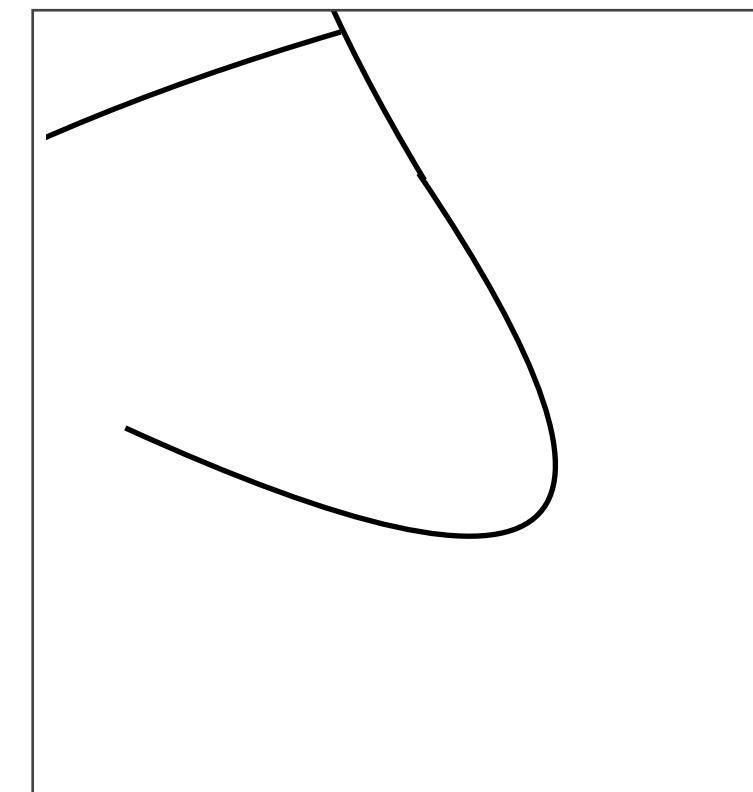
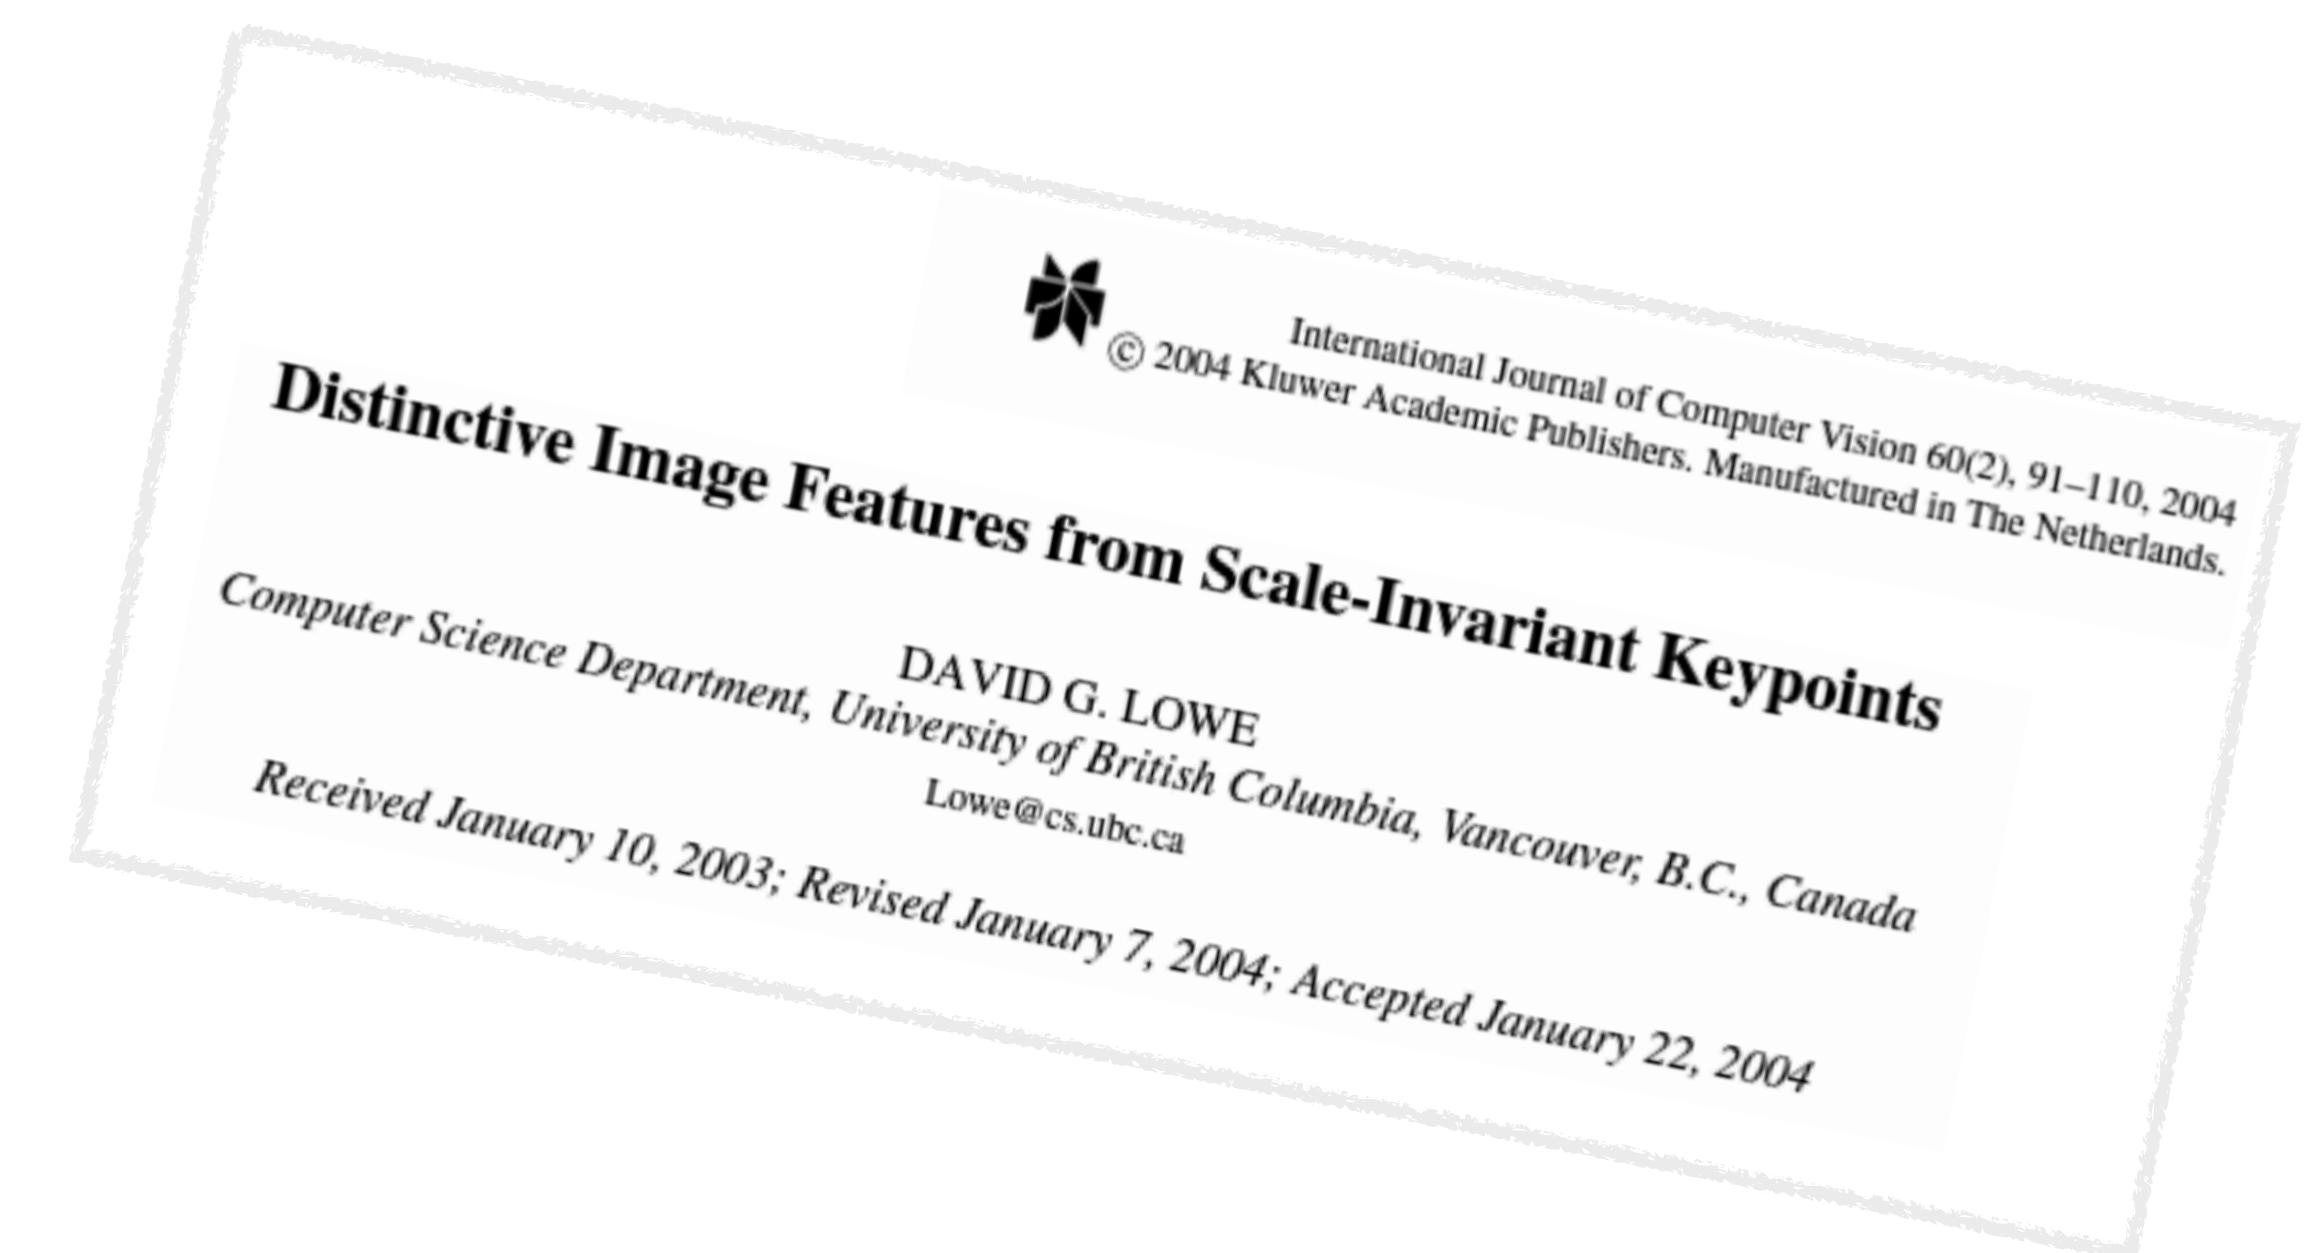


Image 2



Scale Invariant Feature Transform (SIFT)

How to match the same content captured with different resolutions?

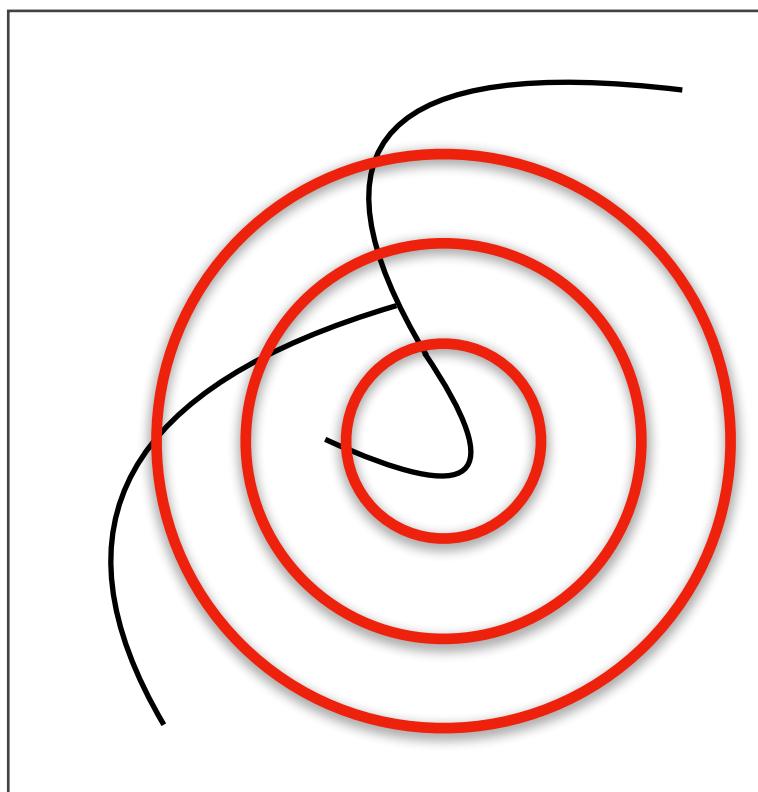


Image 1

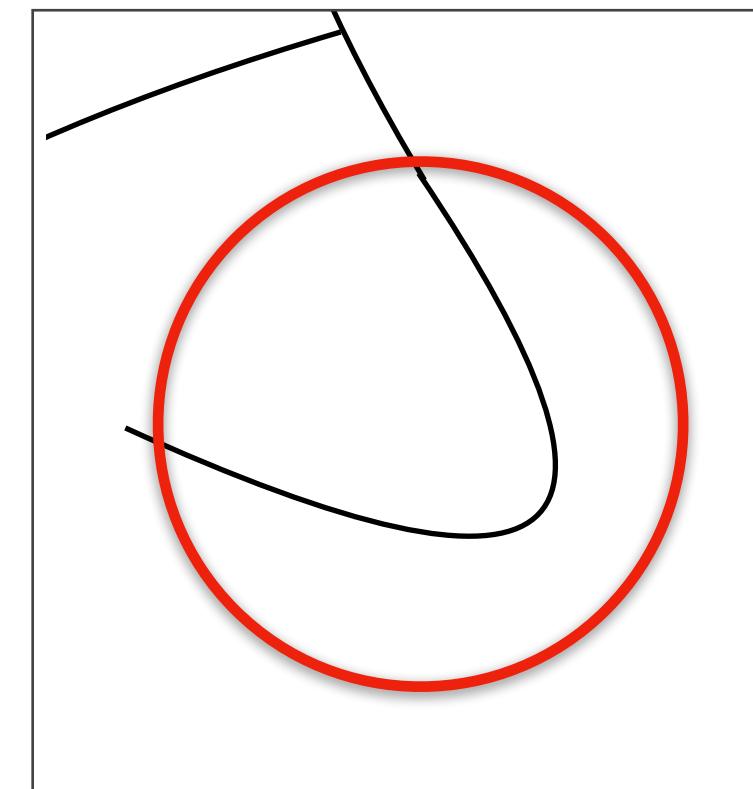
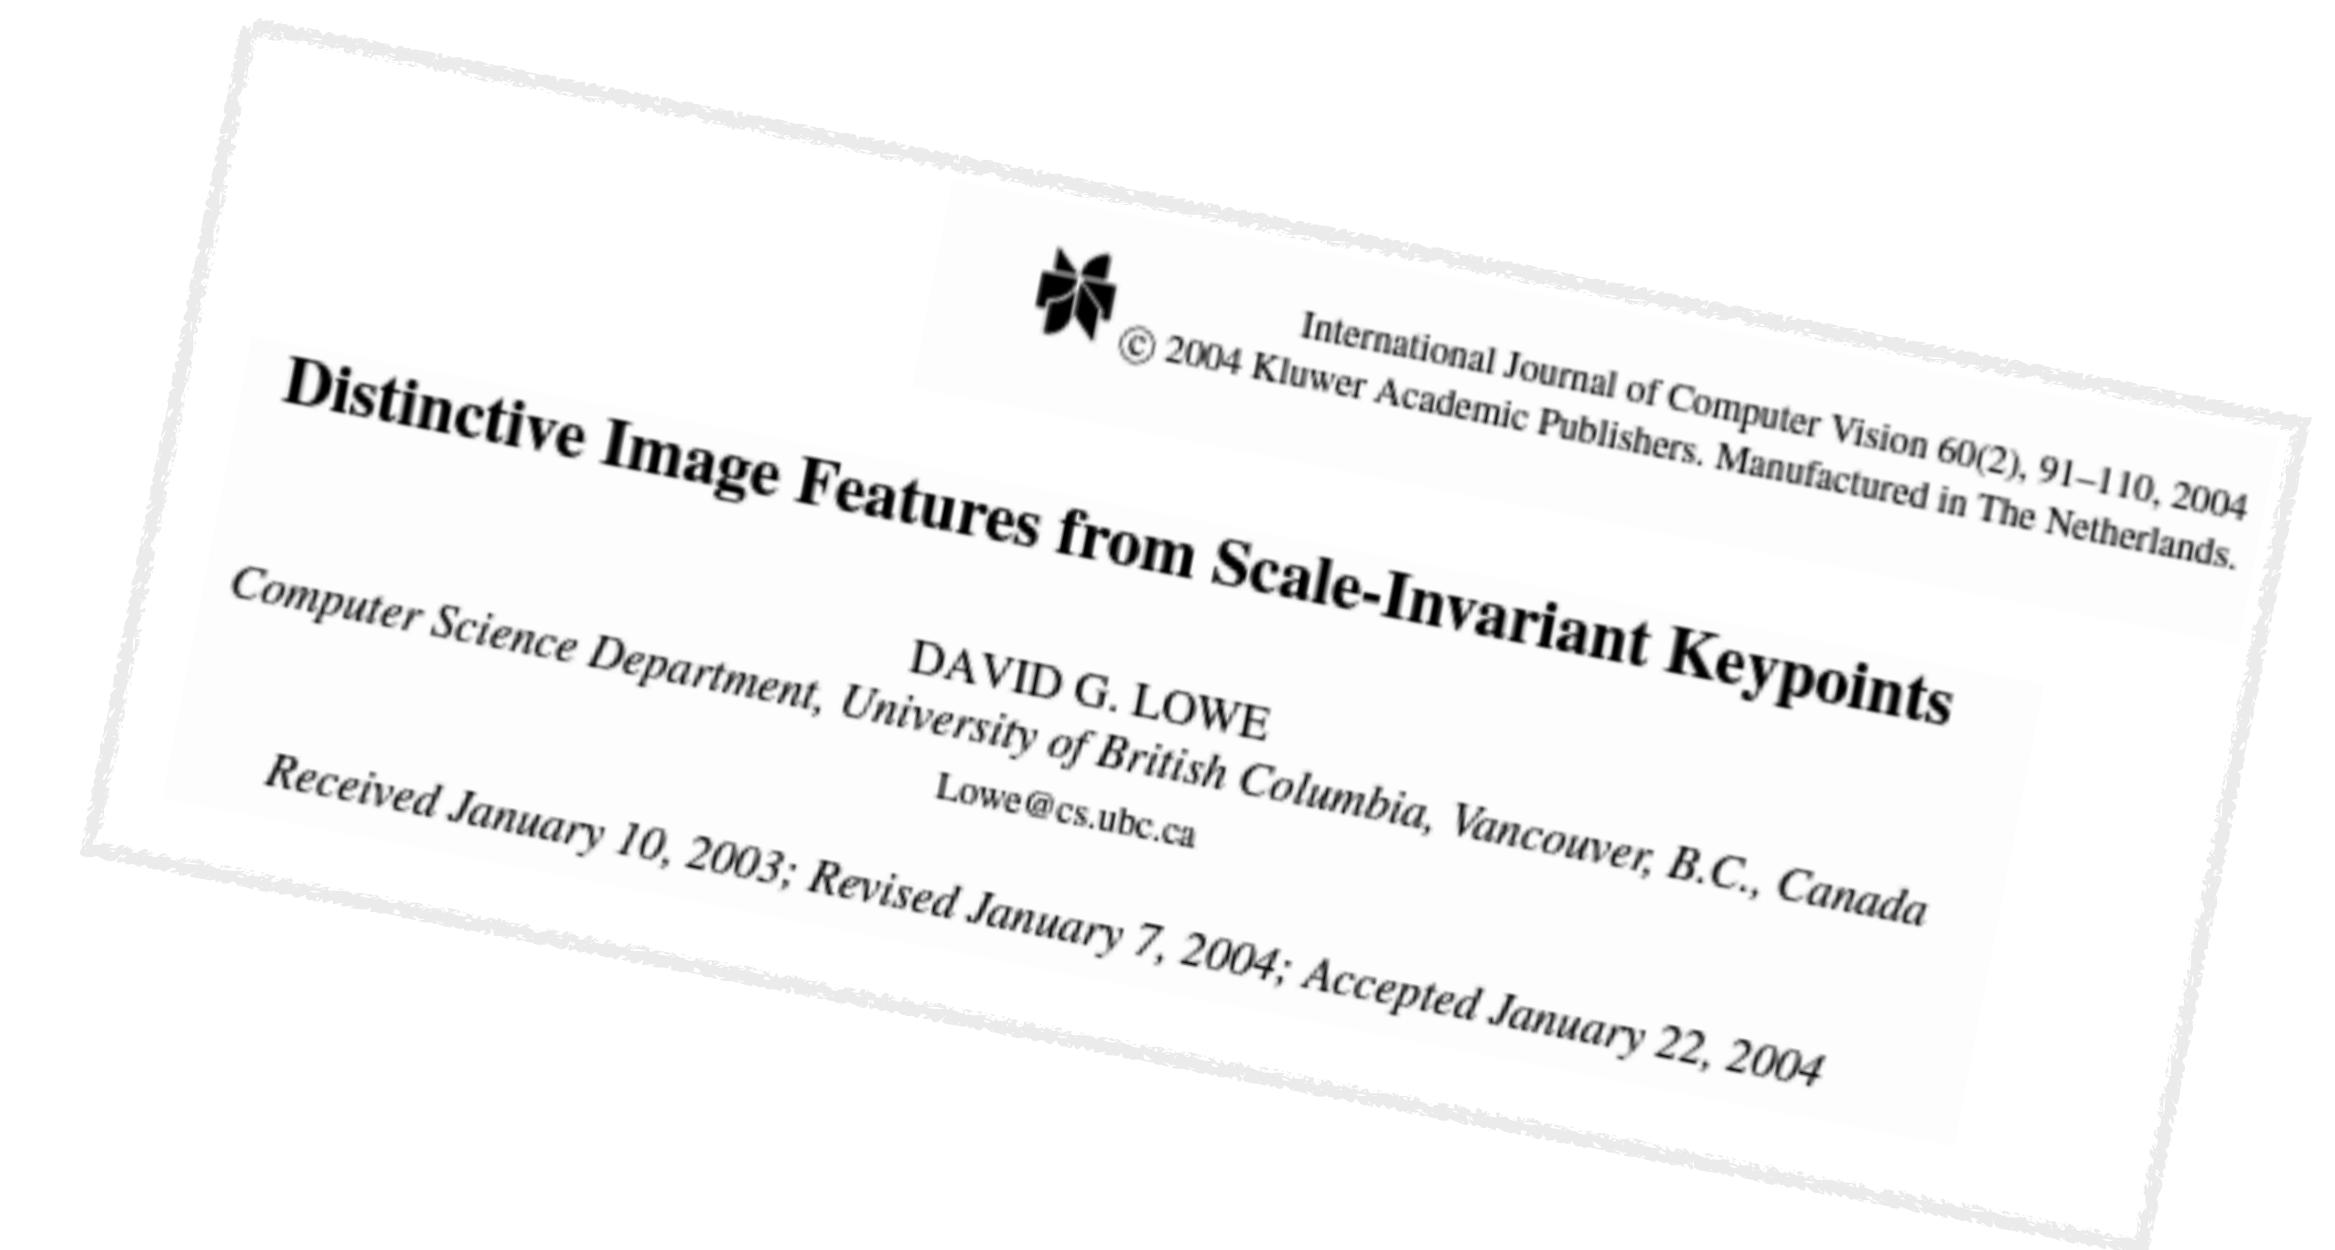


Image 2



Scale Invariant Feature Transform (SIFT)

How to match the same content captured with different resolutions?

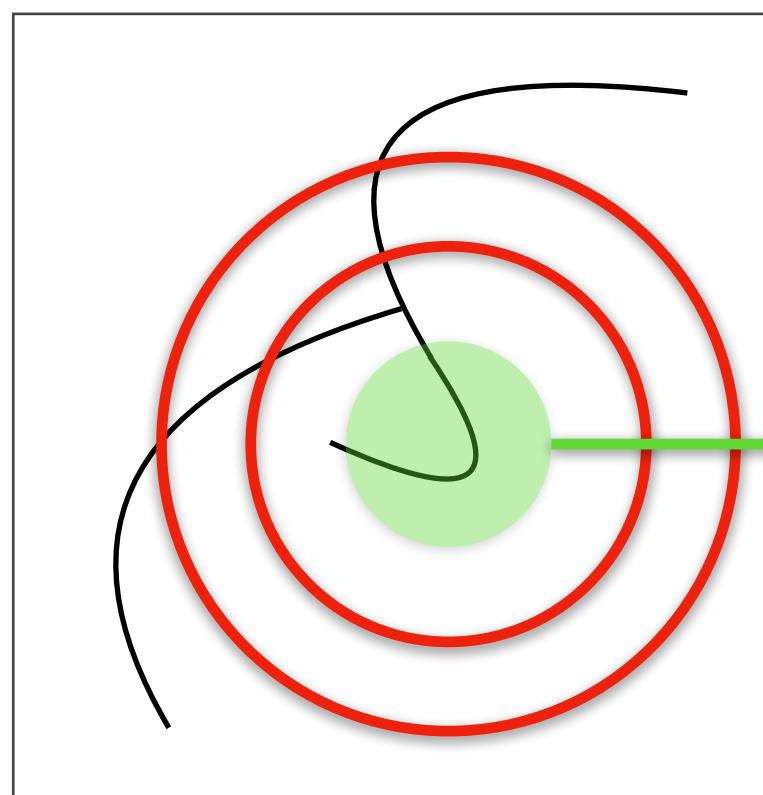


Image 1

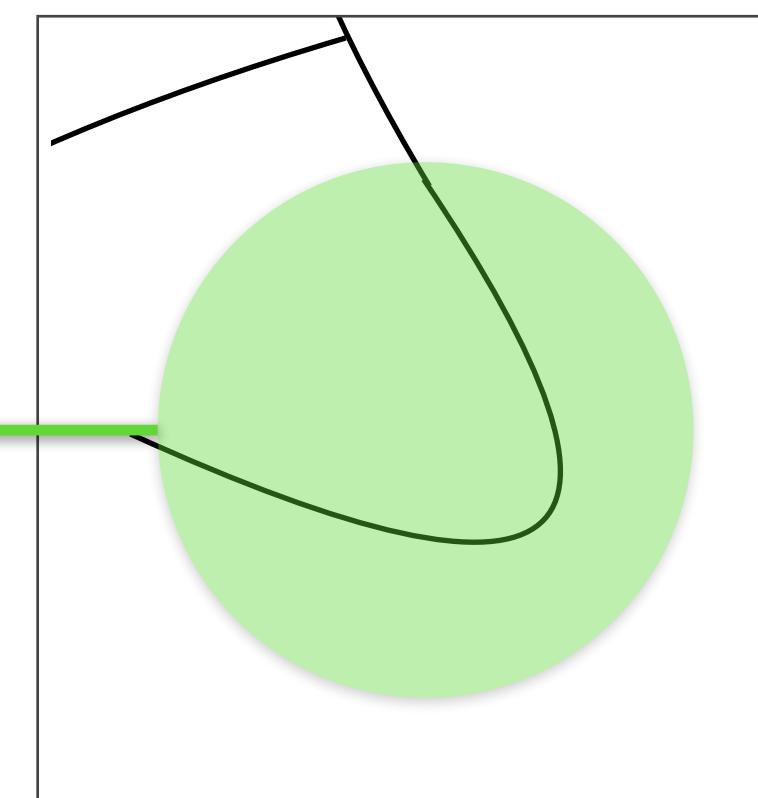
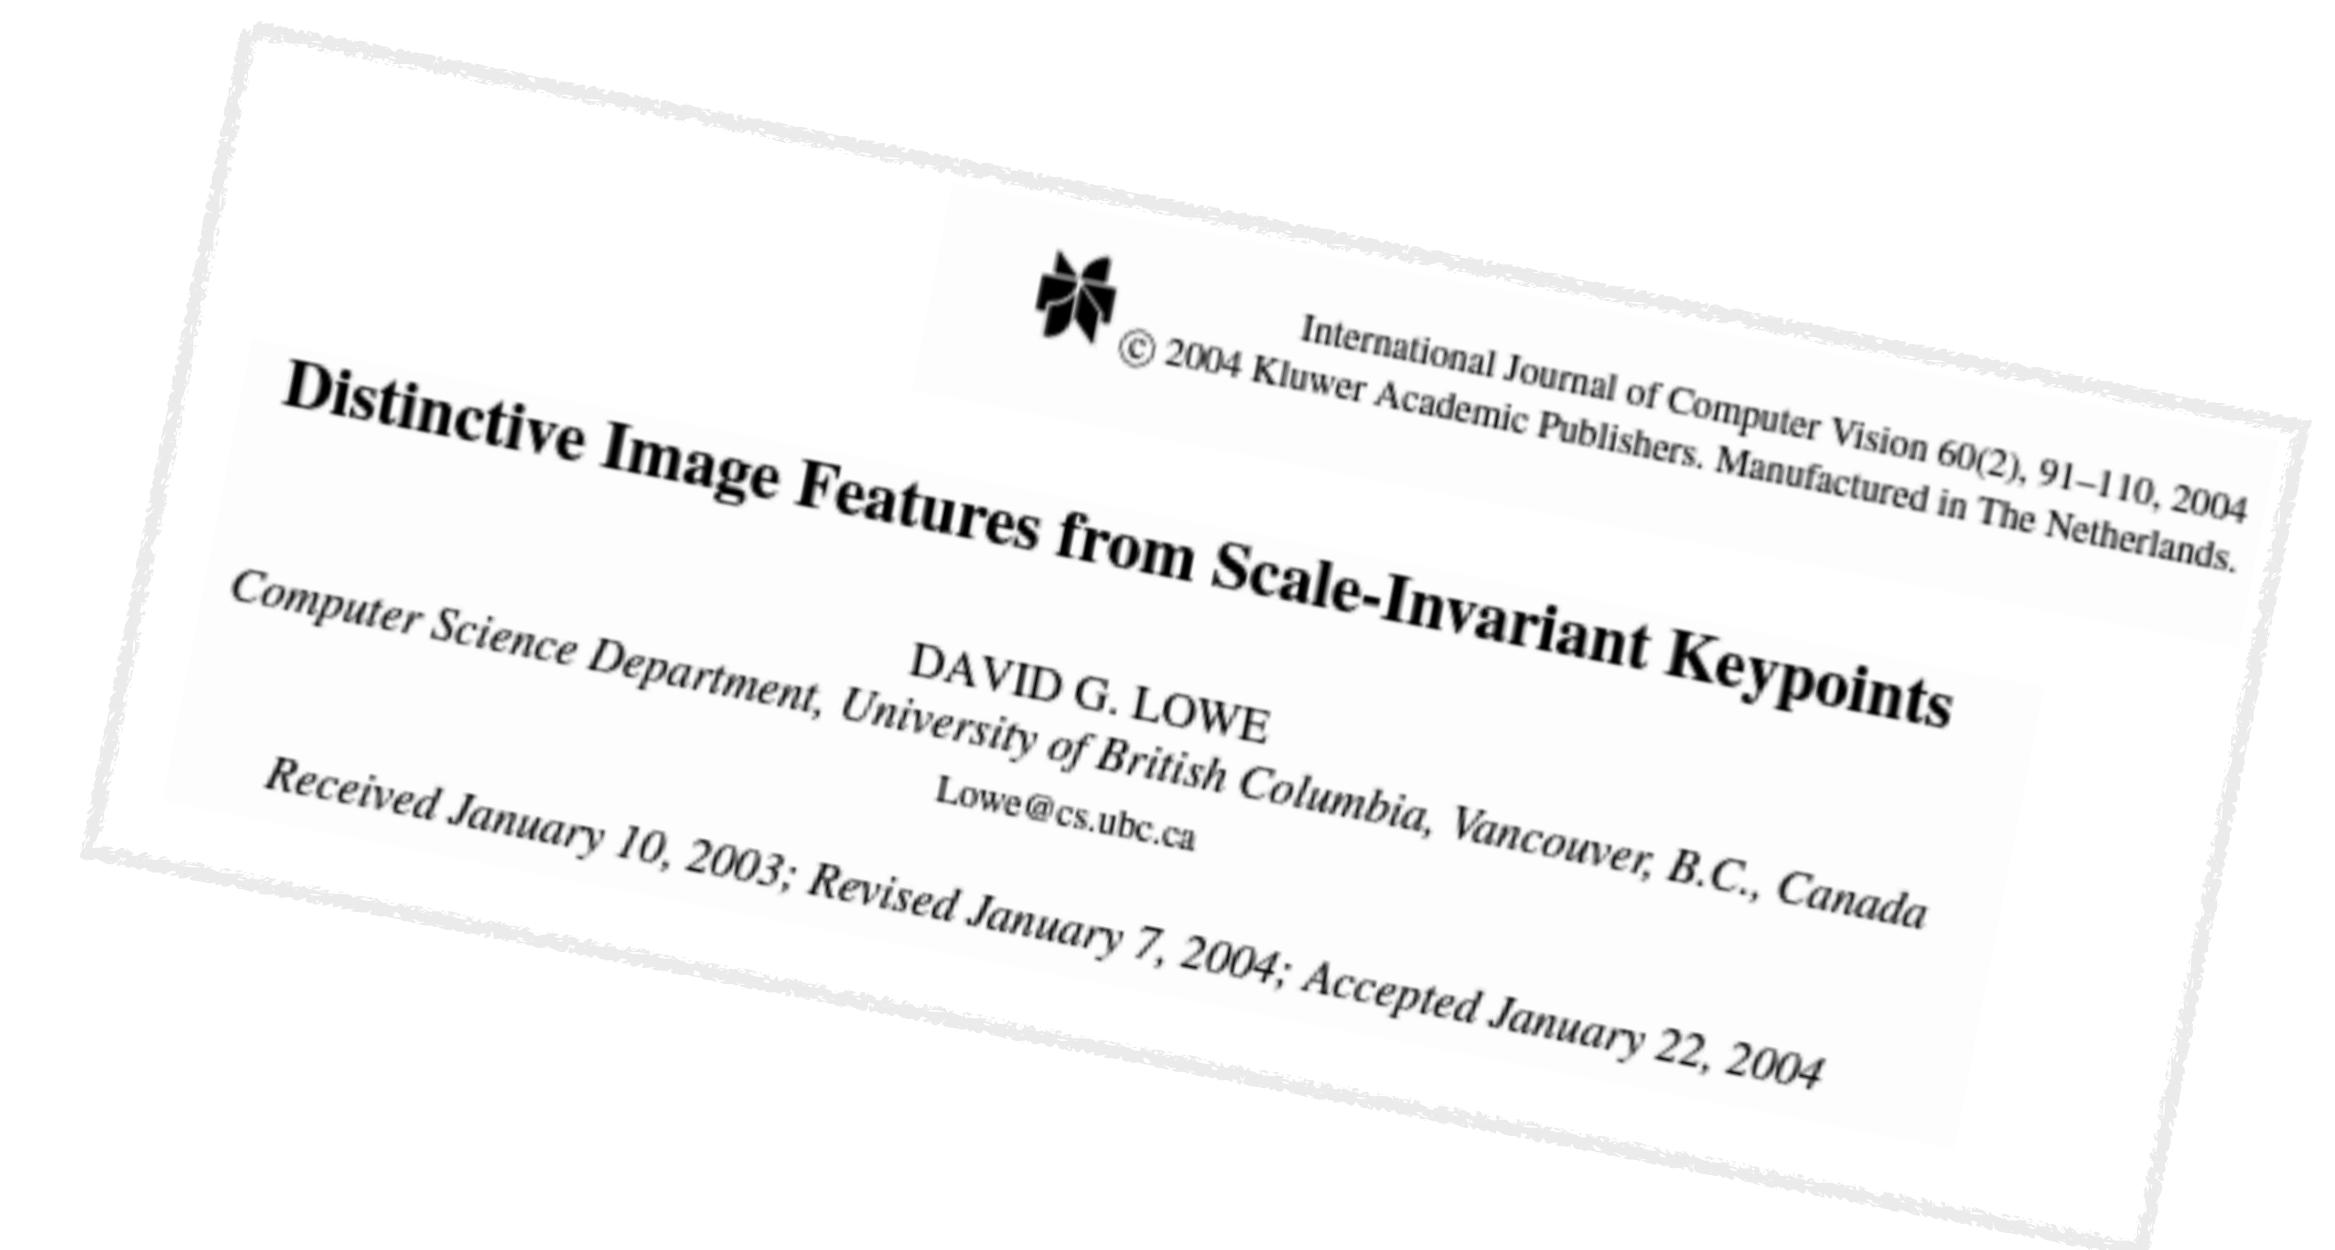
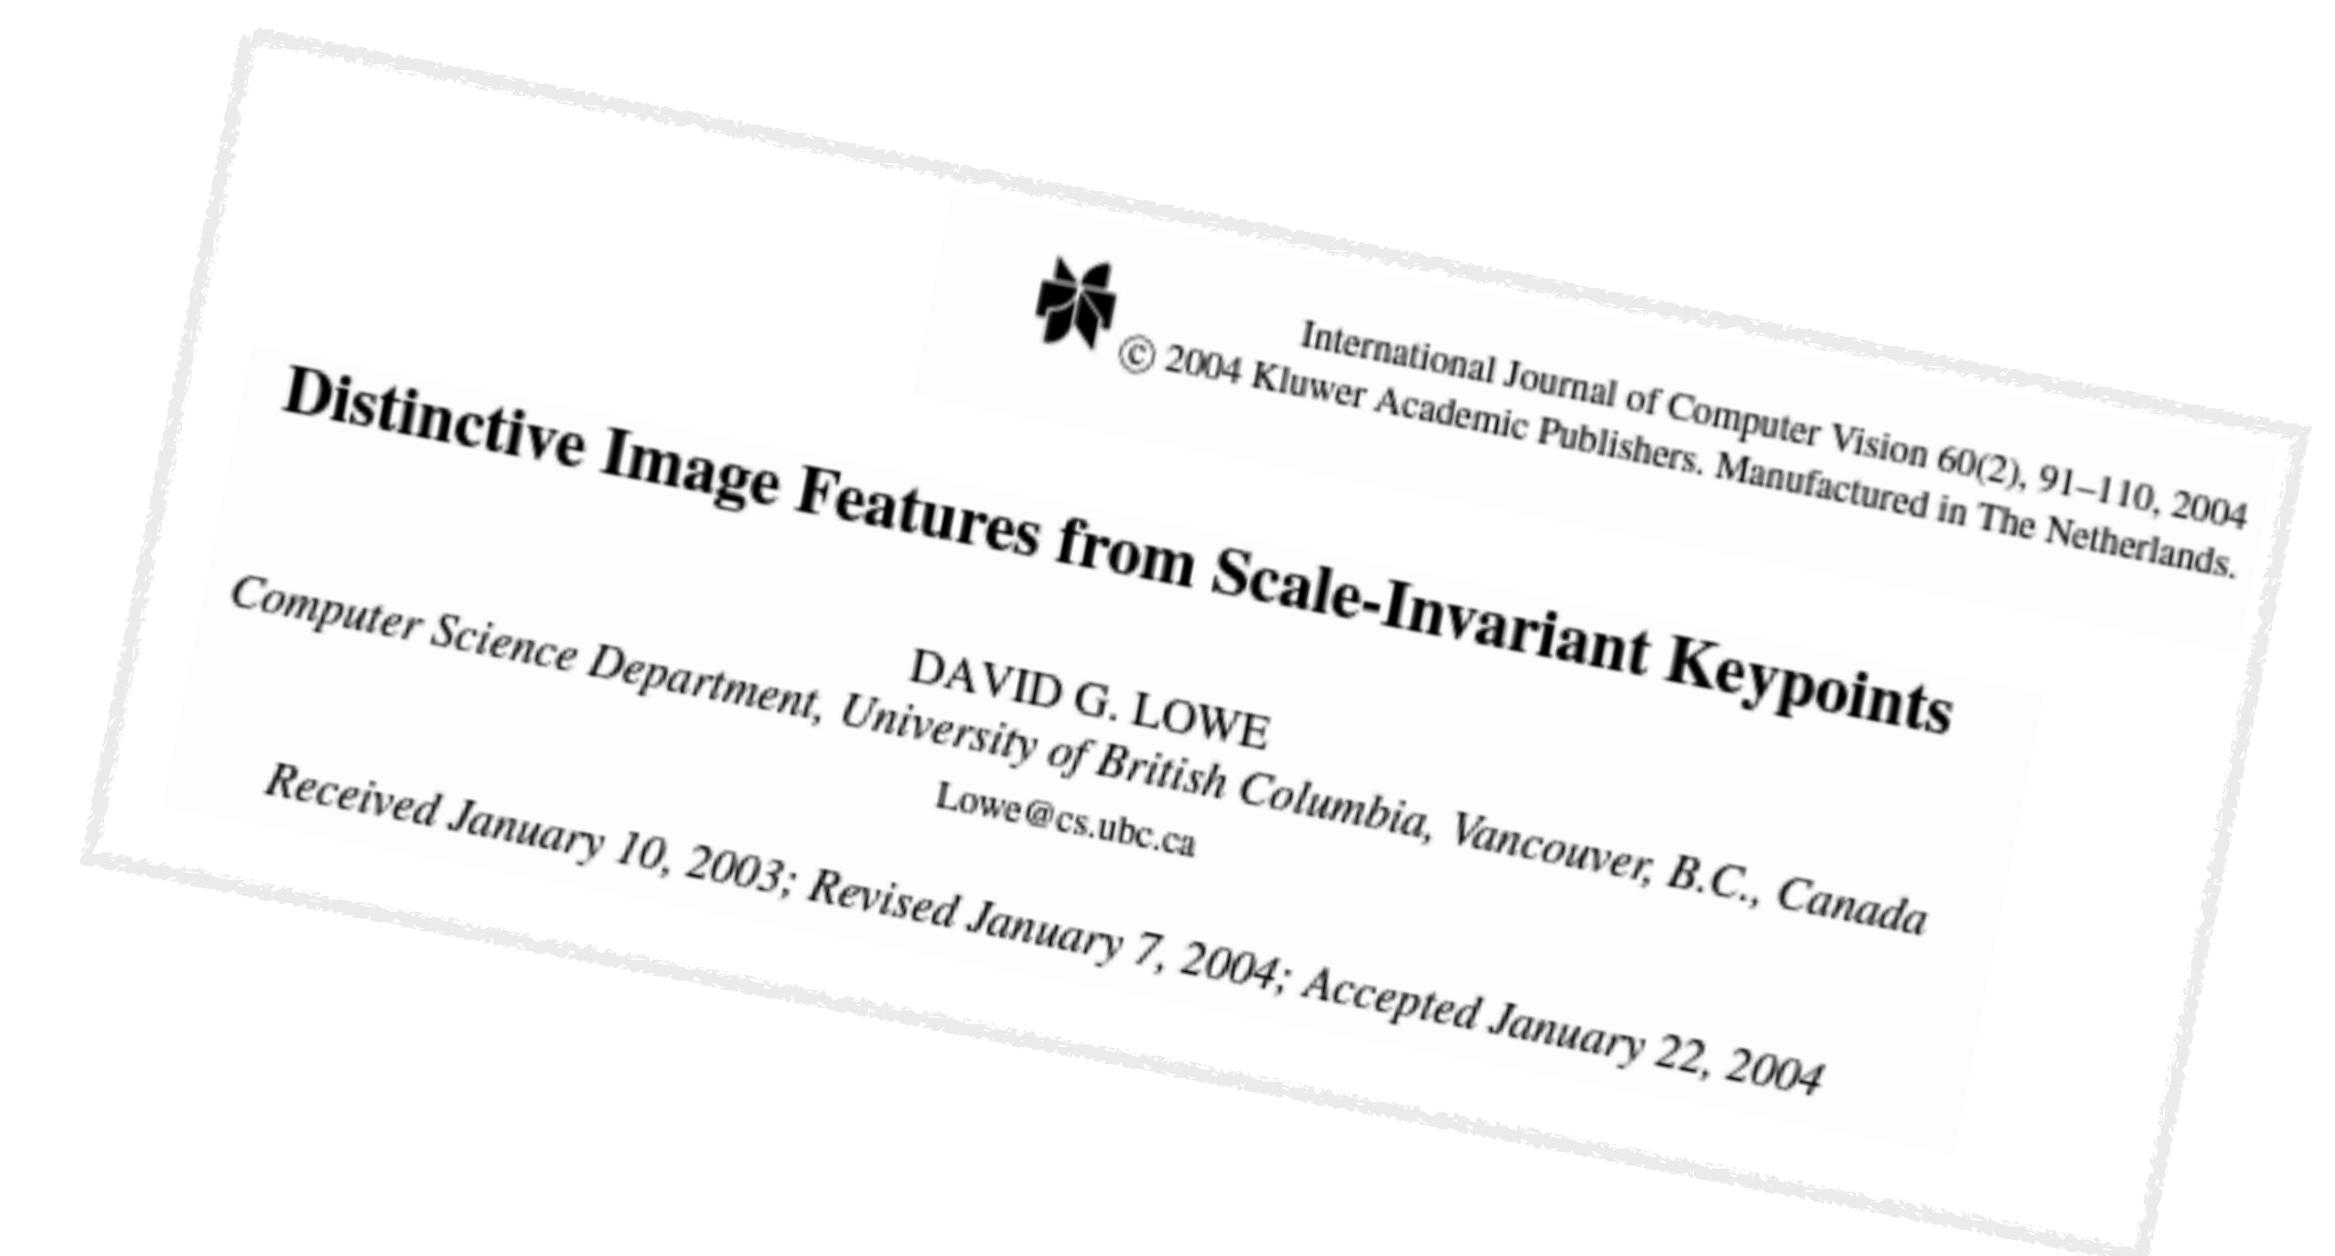
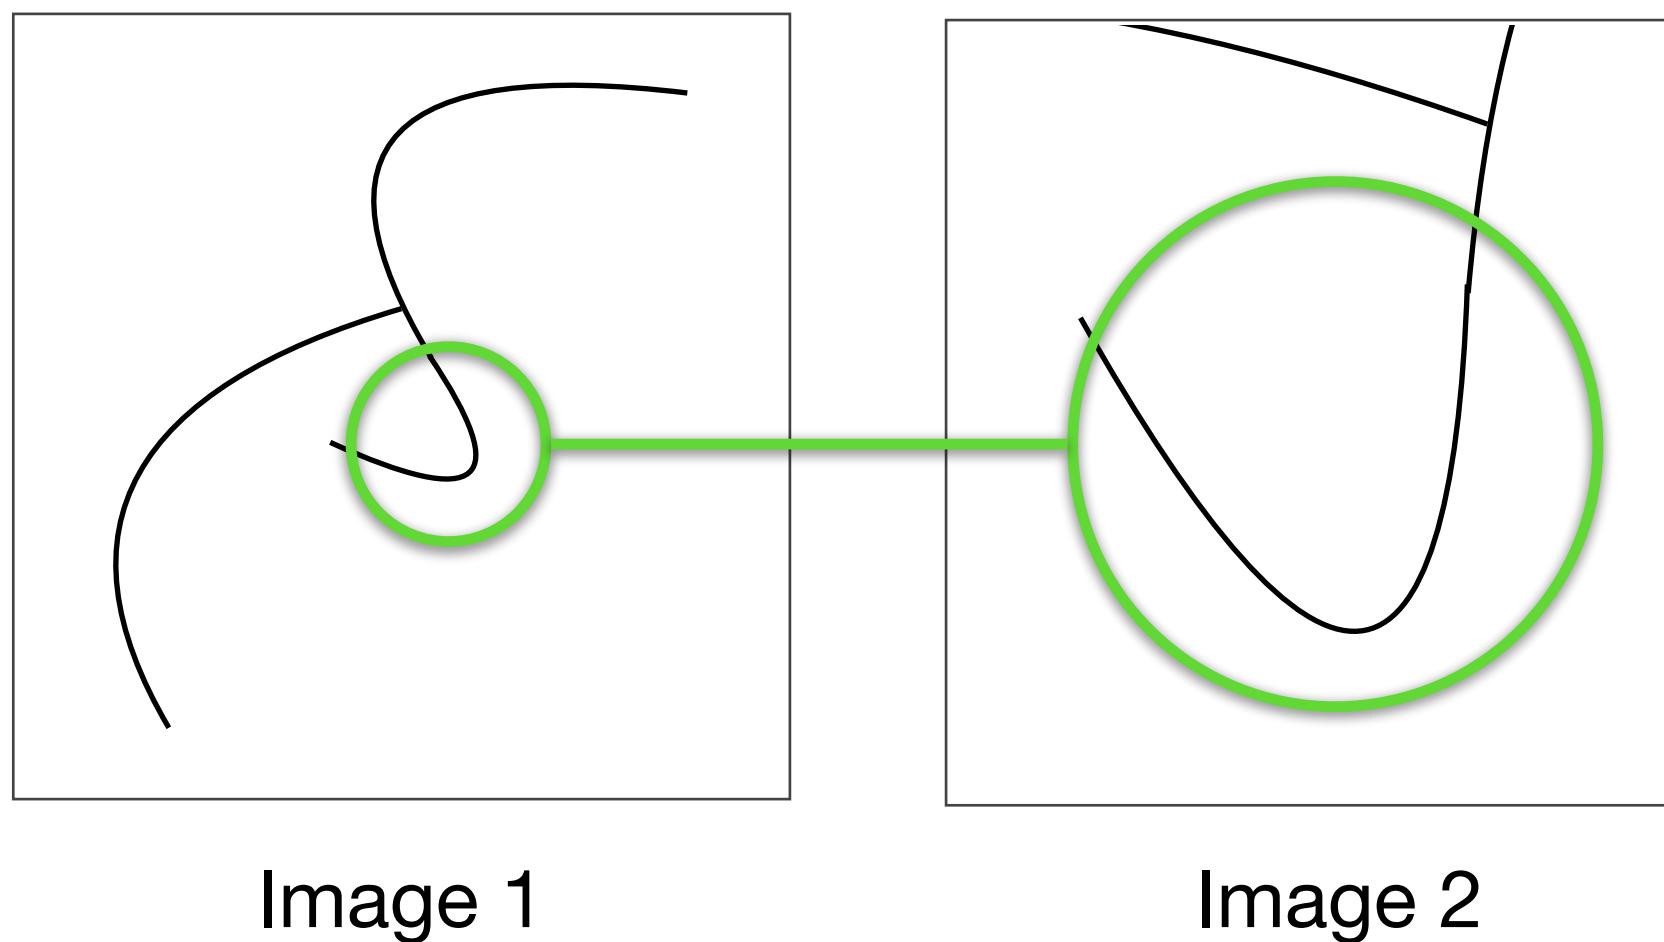


Image 2

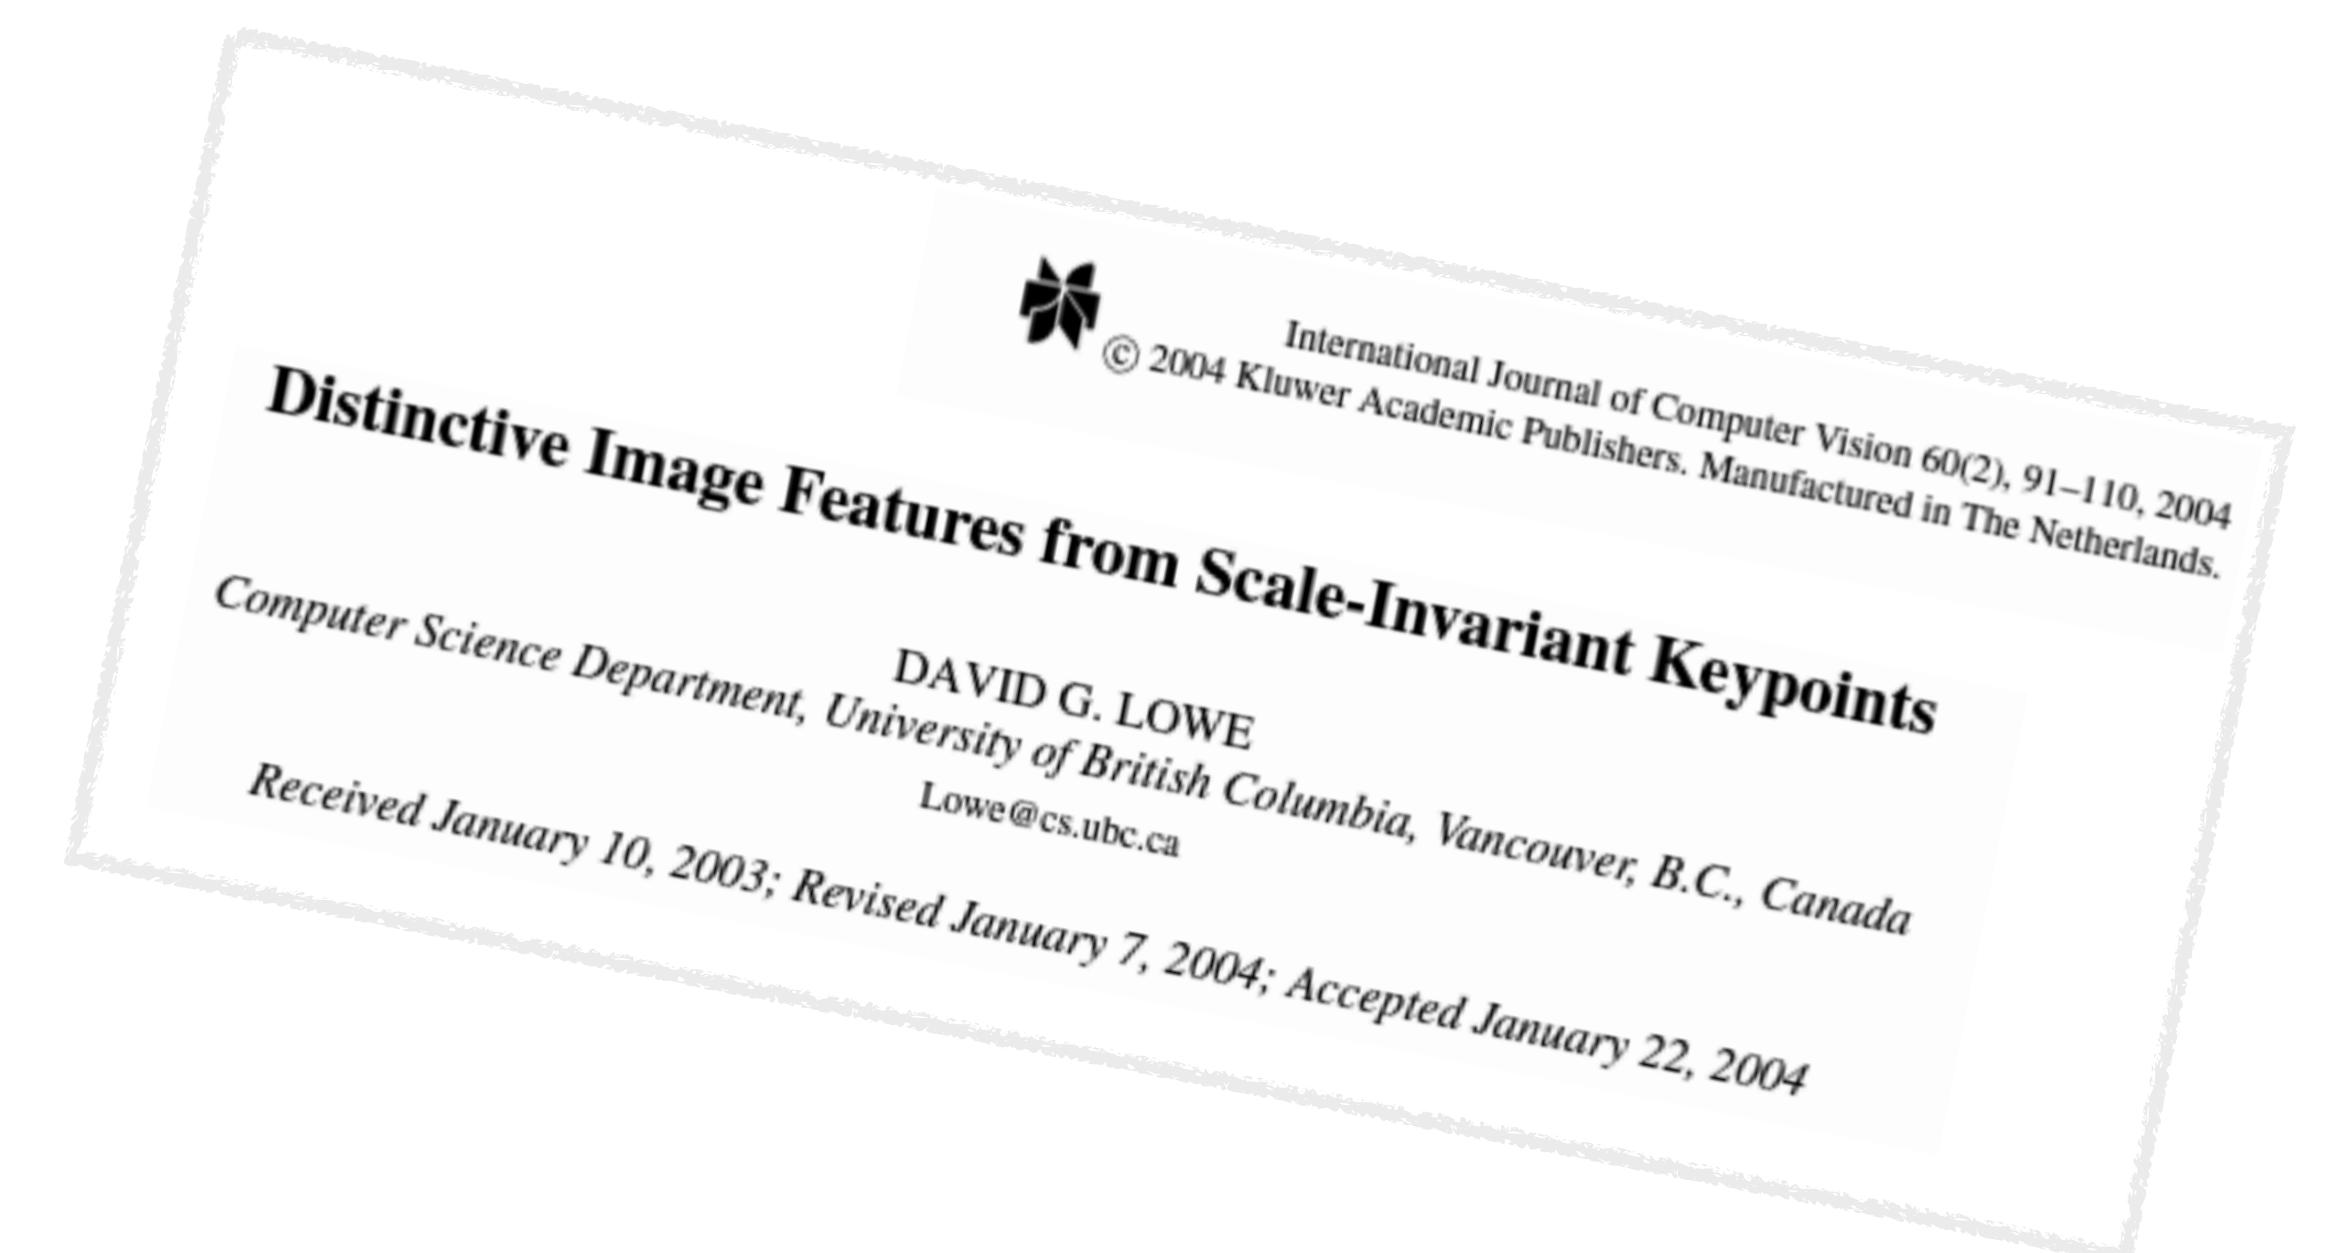
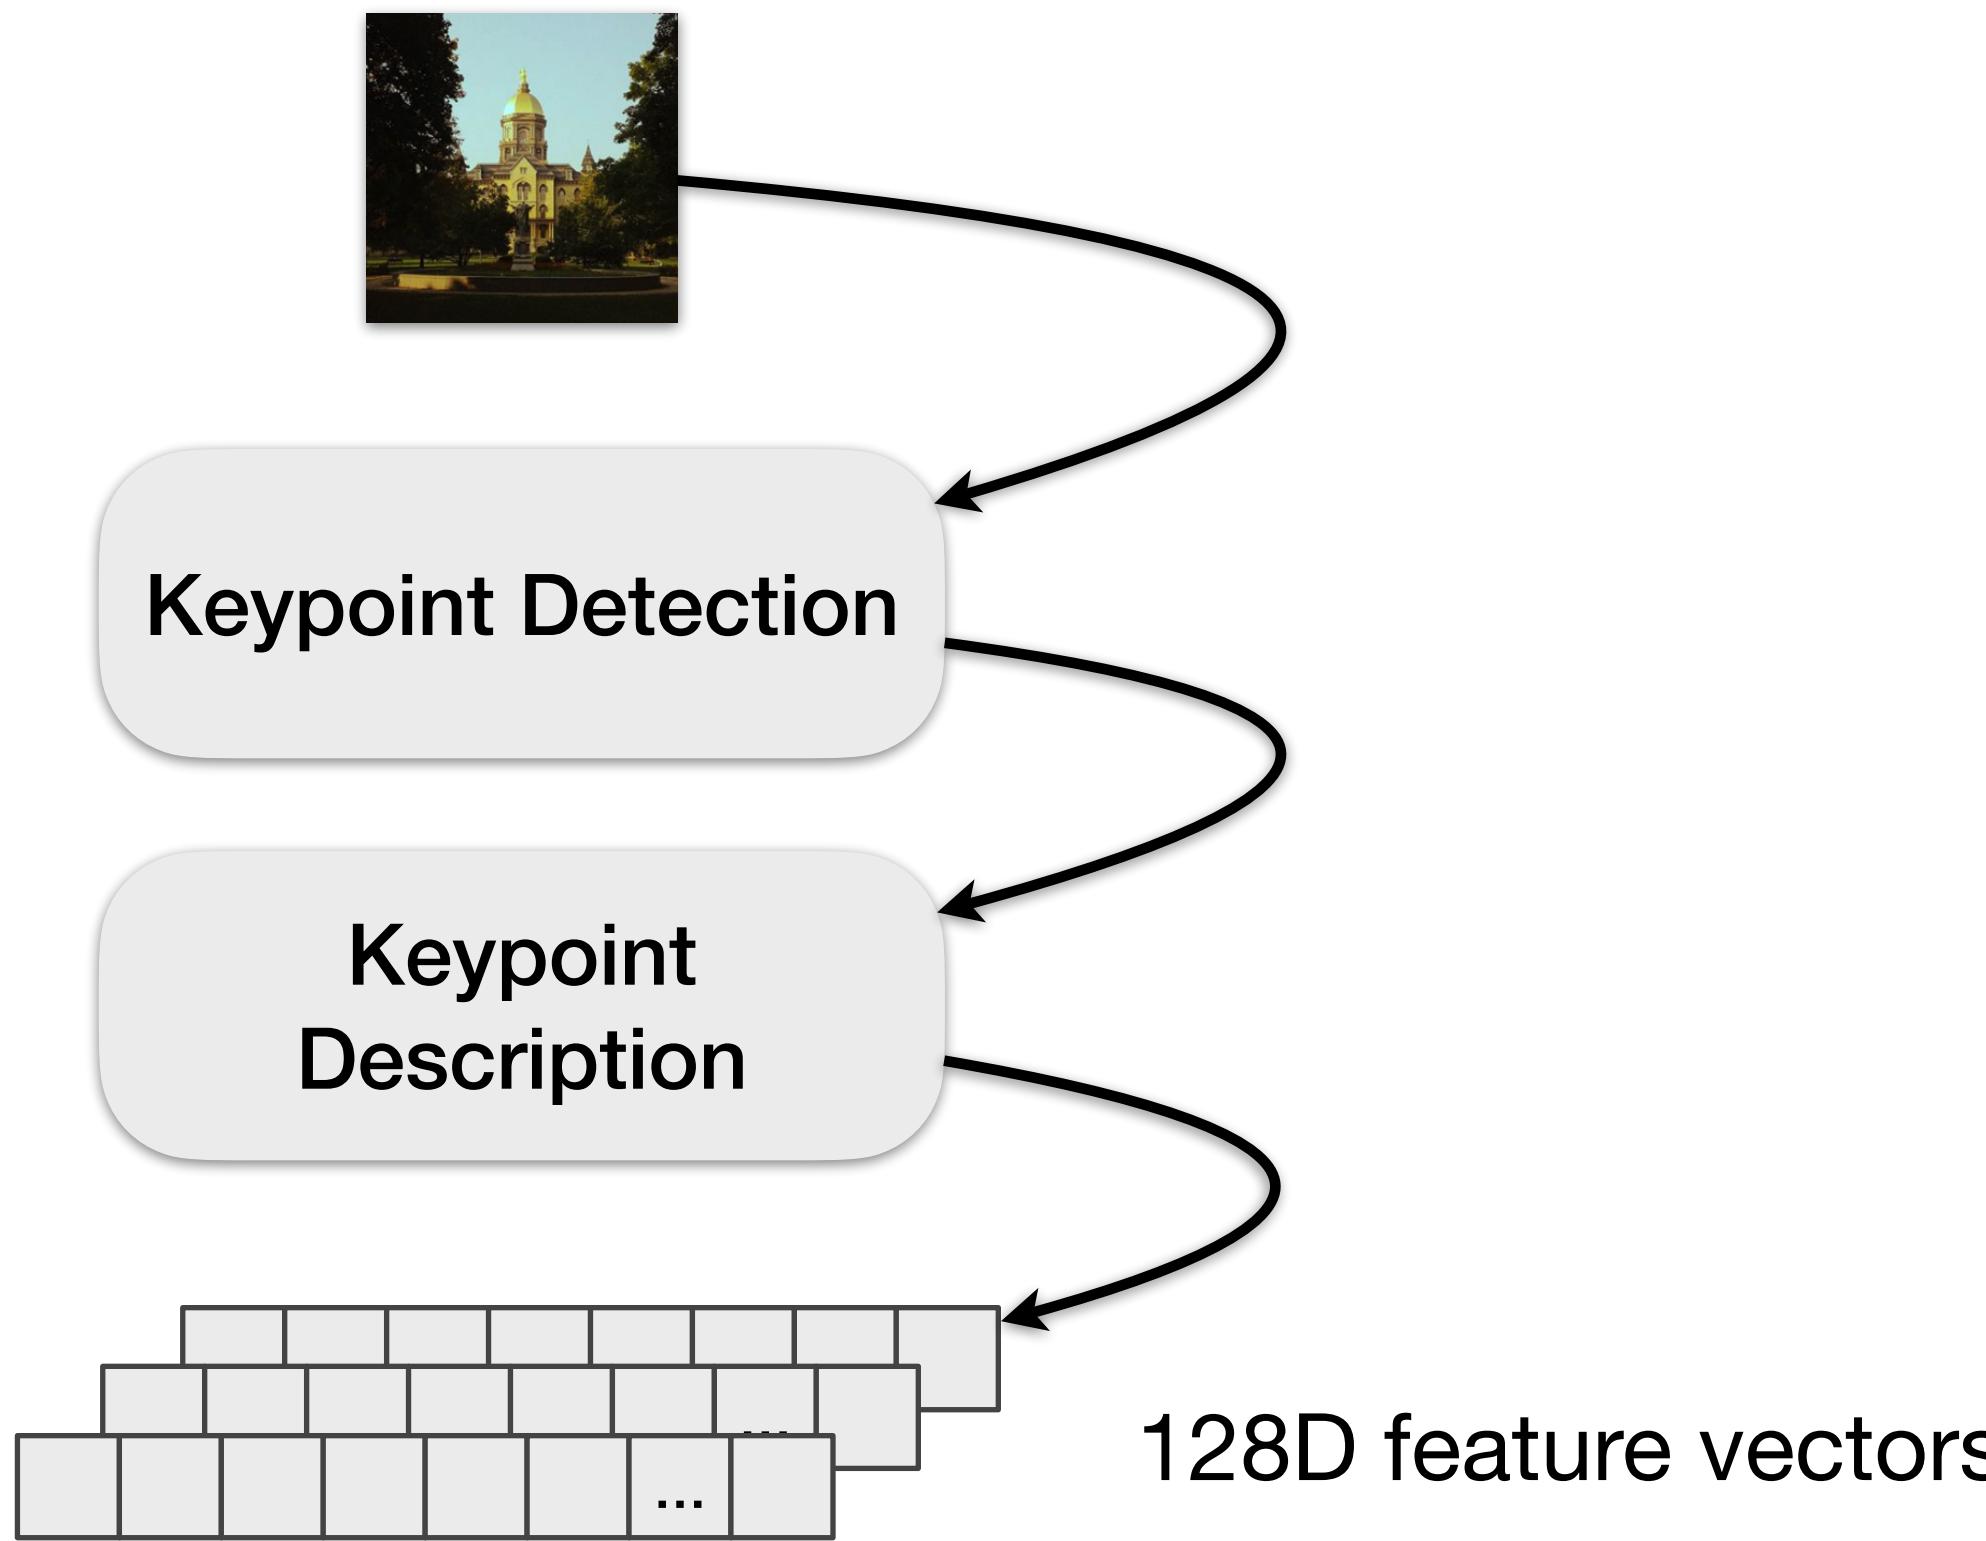


Scale Invariant Feature Transform (SIFT)

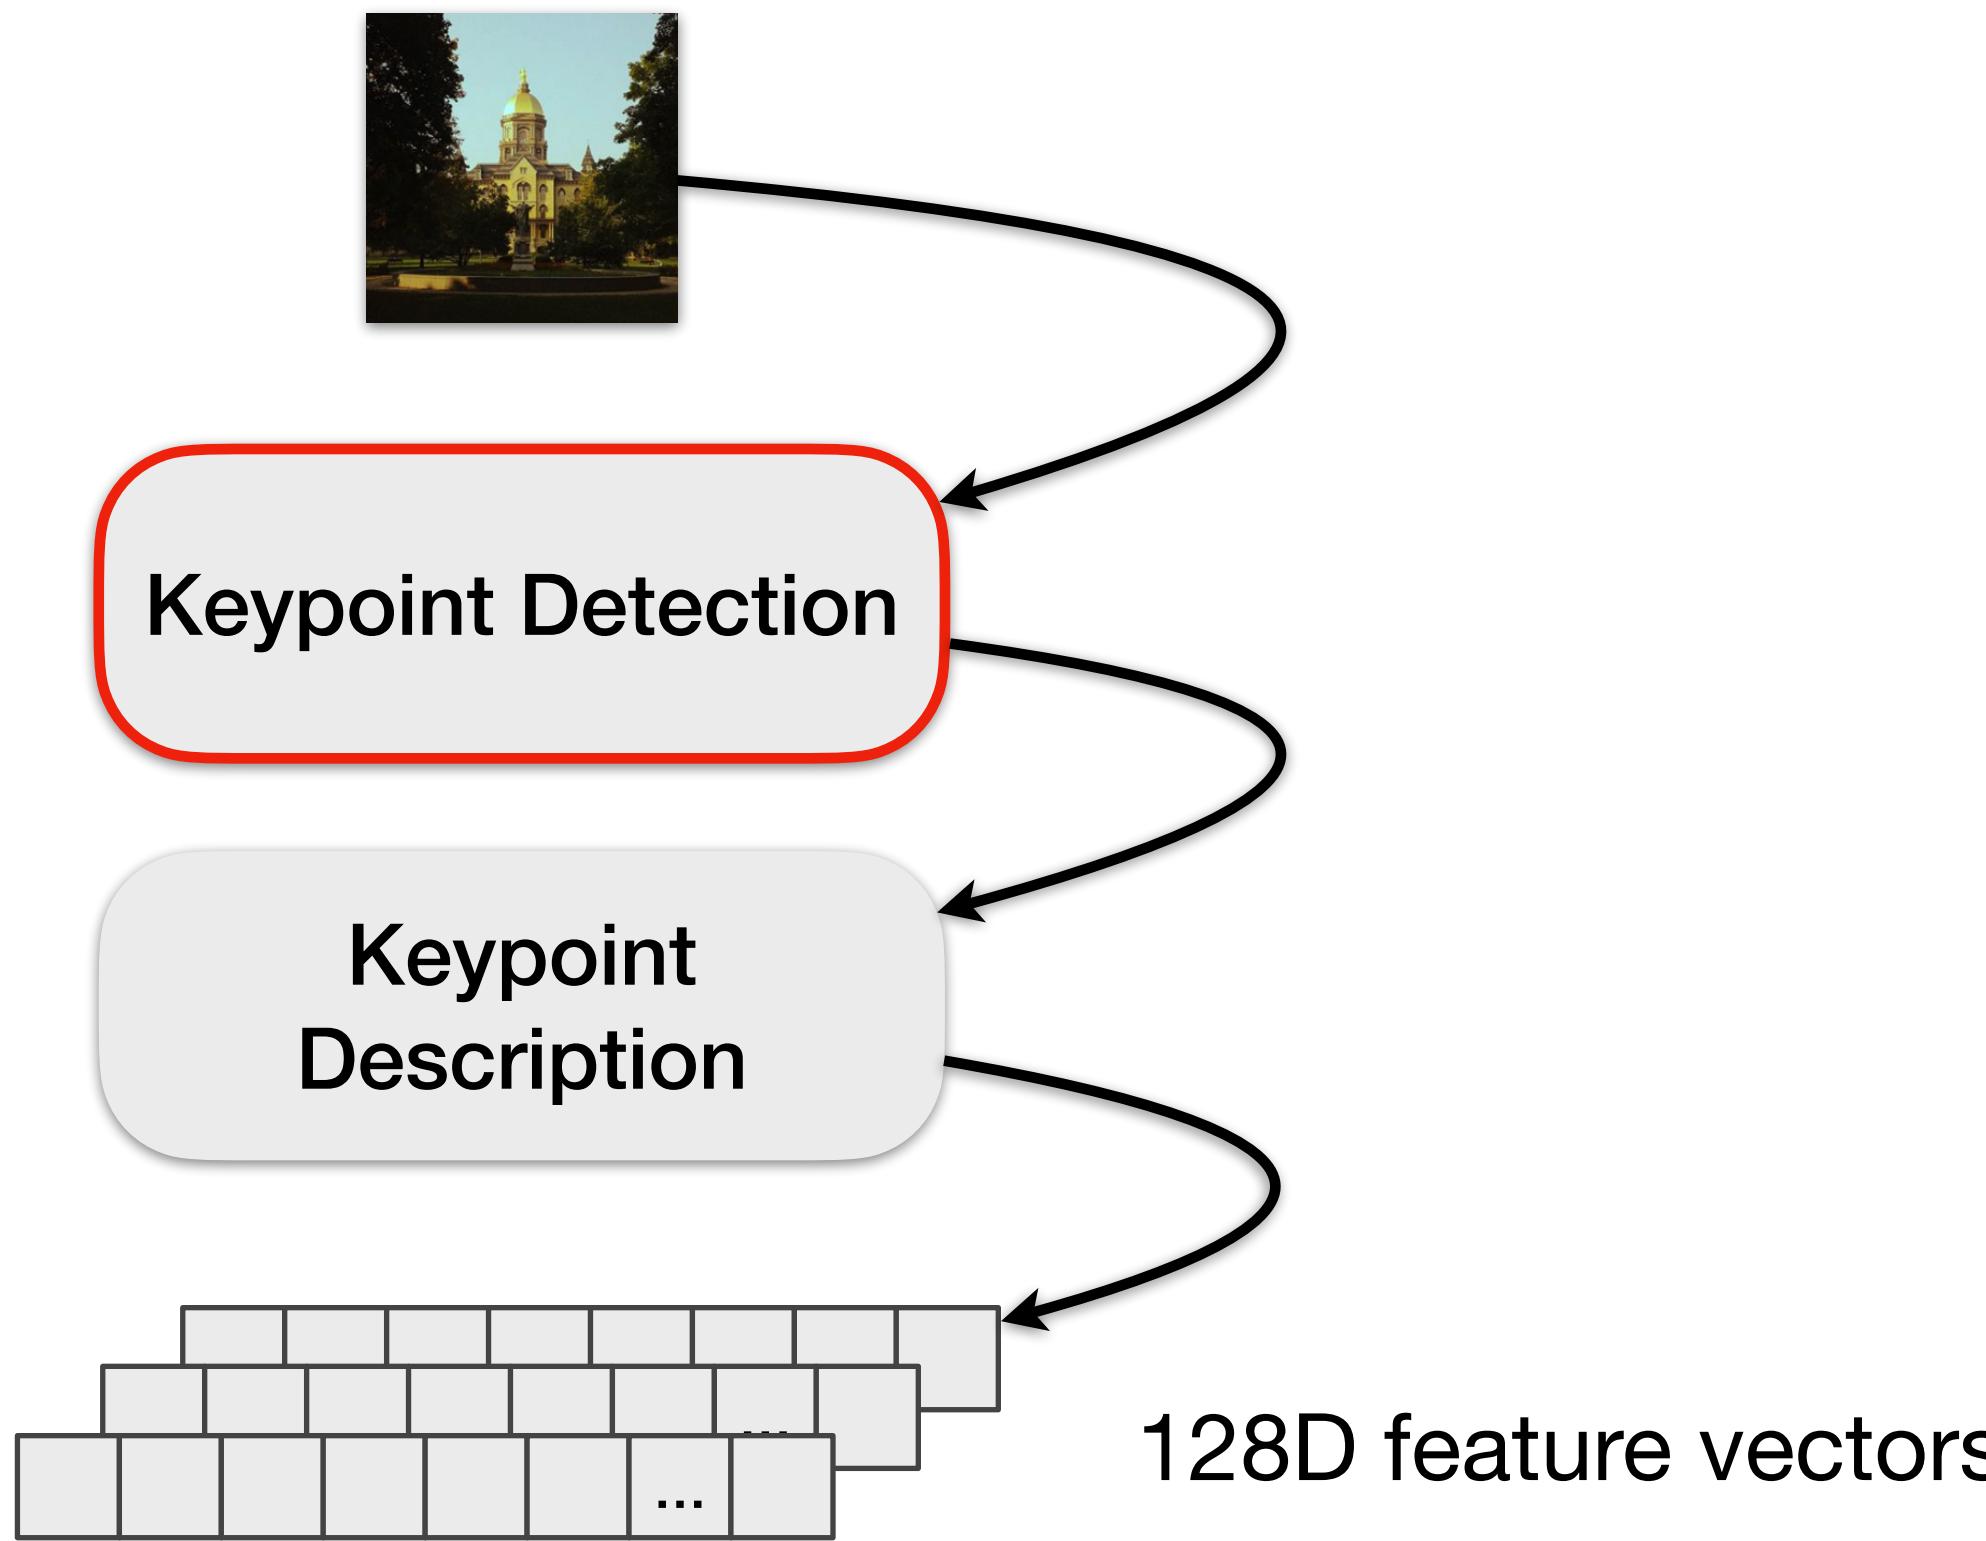
How to match the same content captured with tilt?



Scale Invariant Feature Transform (SIFT)



Scale Invariant Feature Transform (SIFT)



Scale Invariant Feature Transform (SIFT)

Keypoint Detection

How to detect interest points (a.k.a. keypoints) in an image?

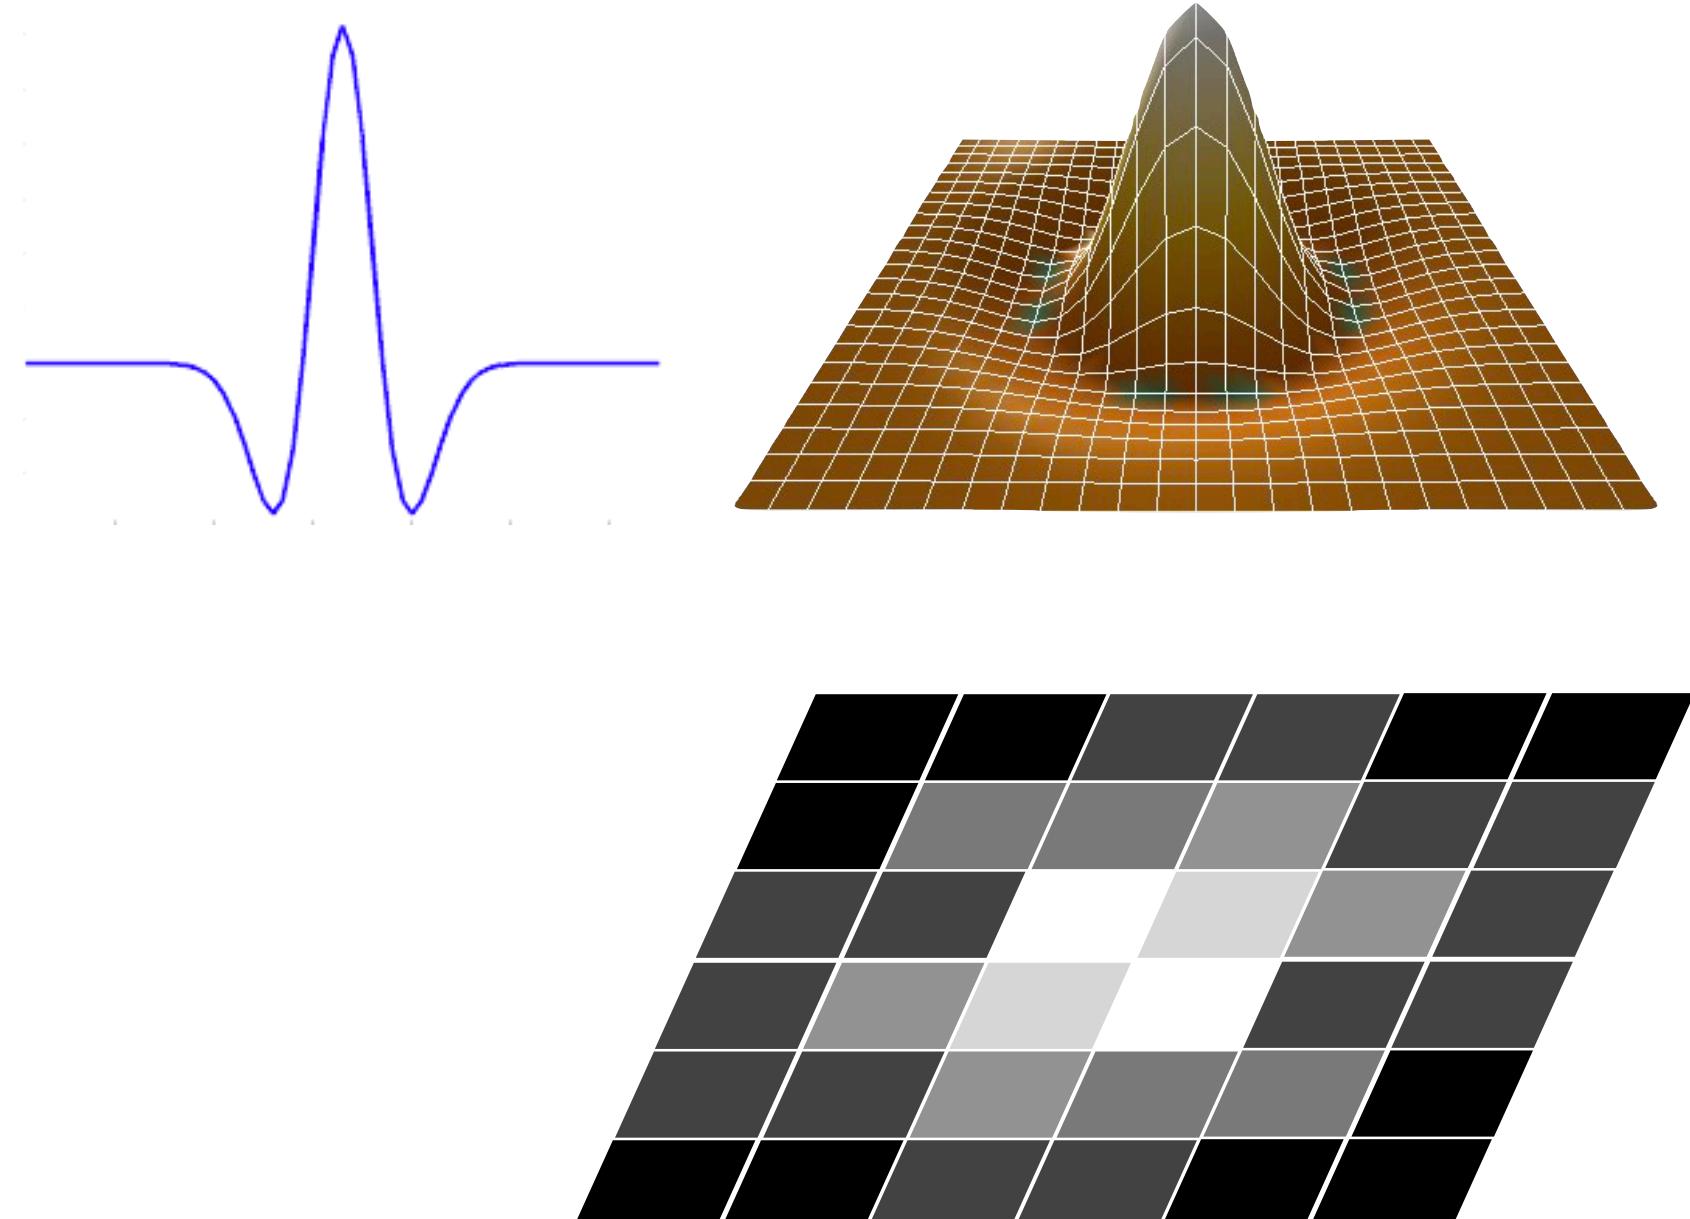
Previous Literature

To focus on blobs:

1. Apply Gaussian to remove noise (blur).
2. Apply Laplacian to detect good regions.

Good regions will have high values after convolution.

This is known as Laplacian of Gaussian (LoG).



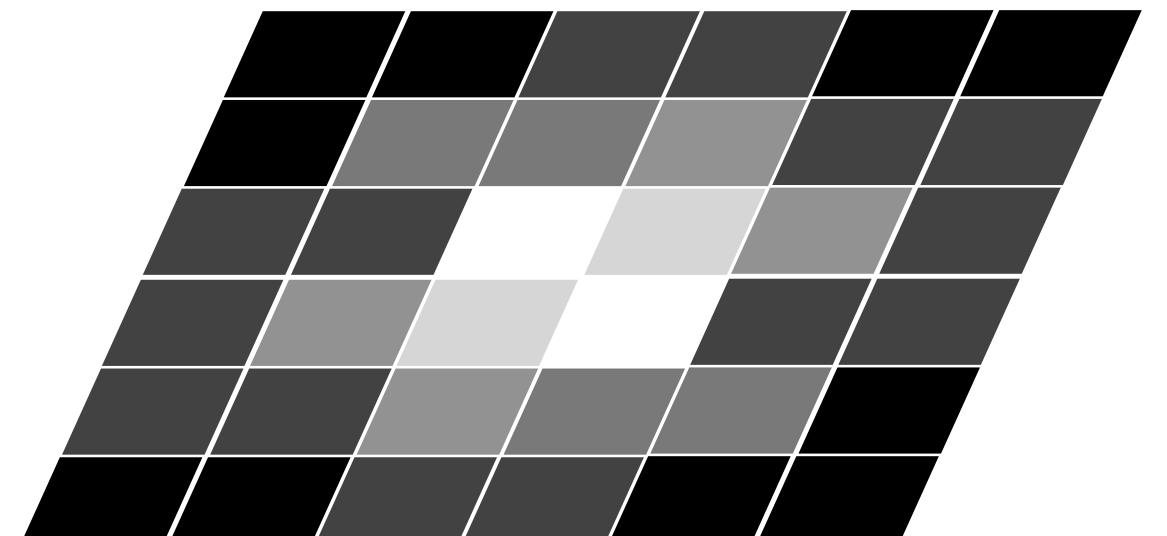
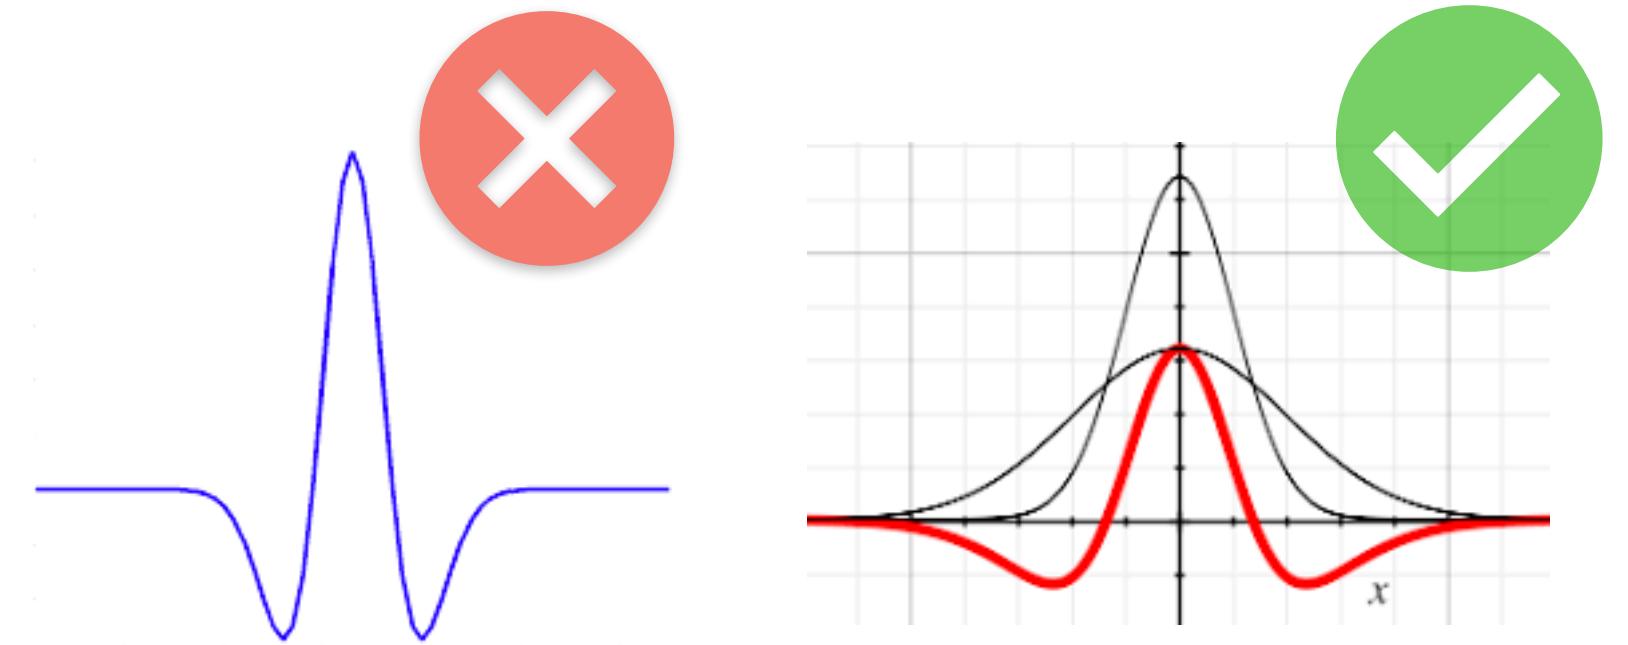
Scale Invariant Feature Transform (SIFT)

Keypoint Detection

How to detect interest points (a.k.a. keypoints) in an image?

Lowe's ideas

Approximate LoG by a Difference of Gaussians (DoG).



Scale Invariant Feature Transform (SIFT)

<https://medium.com/@vad710/cv-for-busy-devs-improving-features-df20c3aa5887>

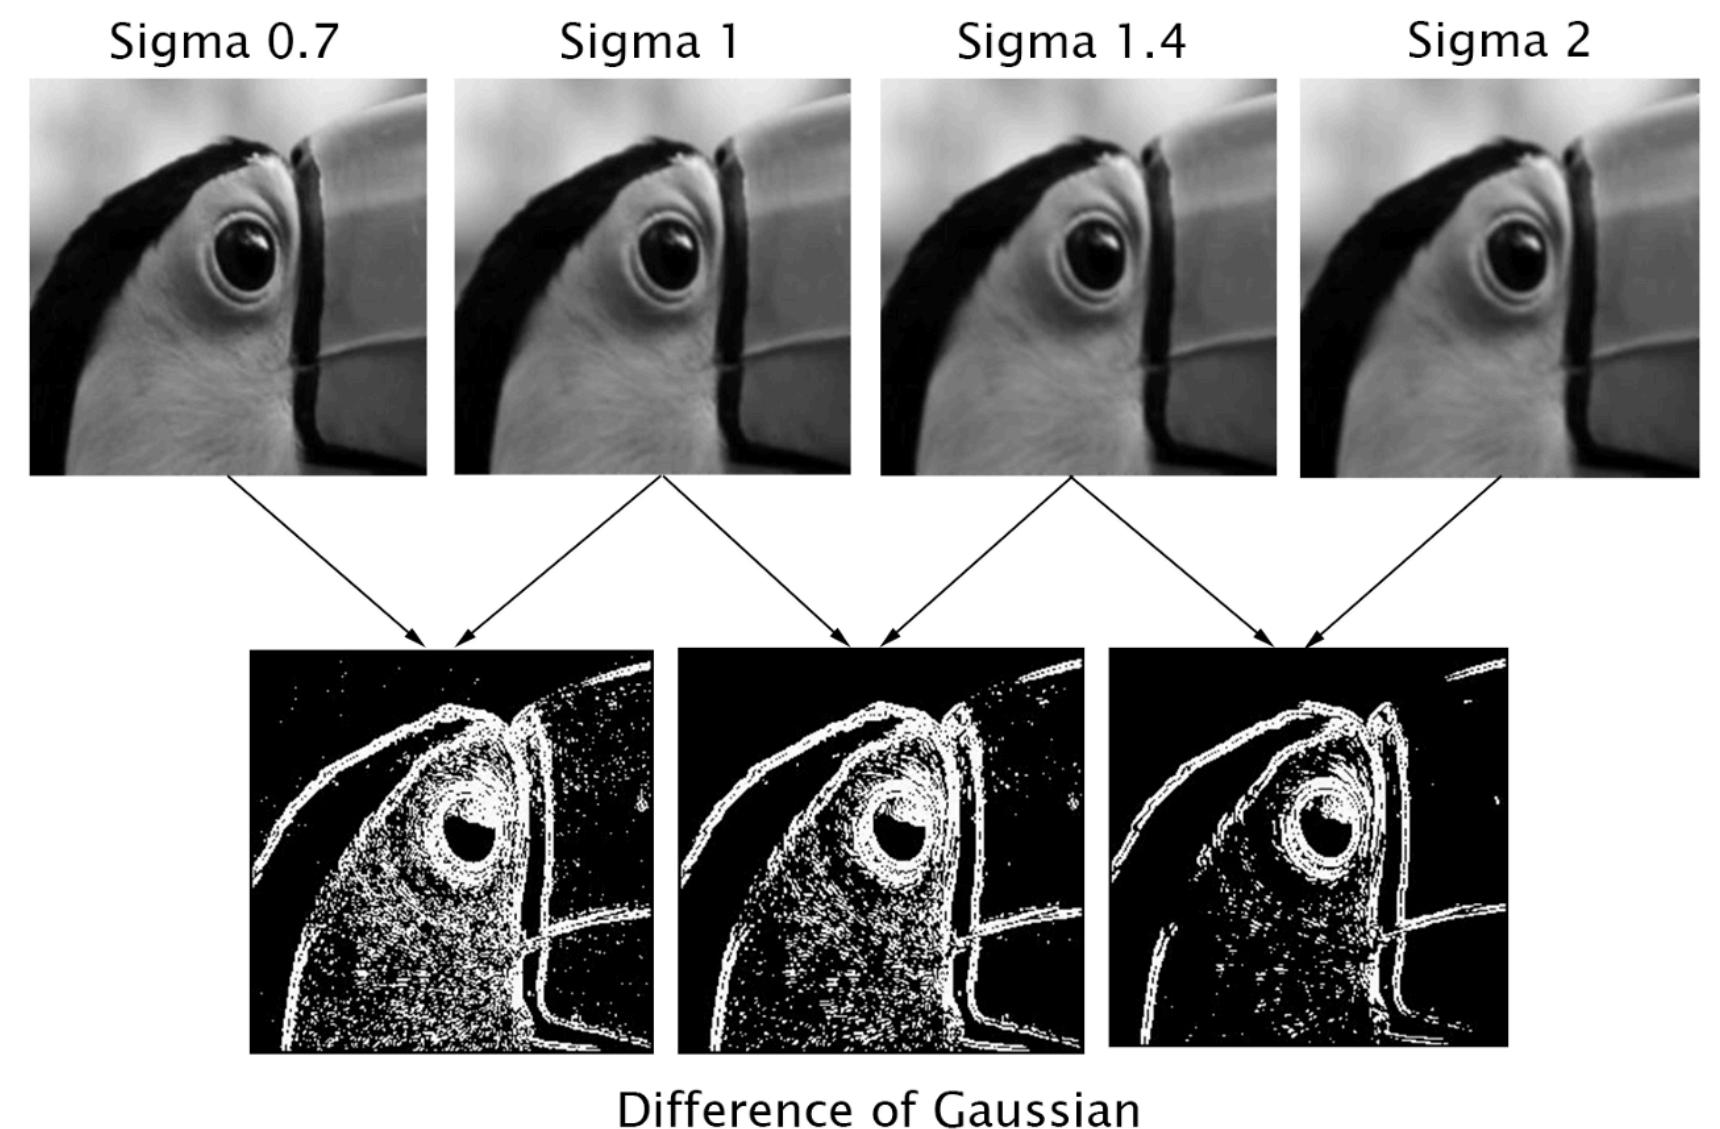
Keypoint Detection

How to detect interest points (a.k.a. keypoints) in an image?

Lowe's ideas

Approximate LoG by a Difference of Gaussians (DoG).

The two subtracted Gaussians come from distinct variances (scales).



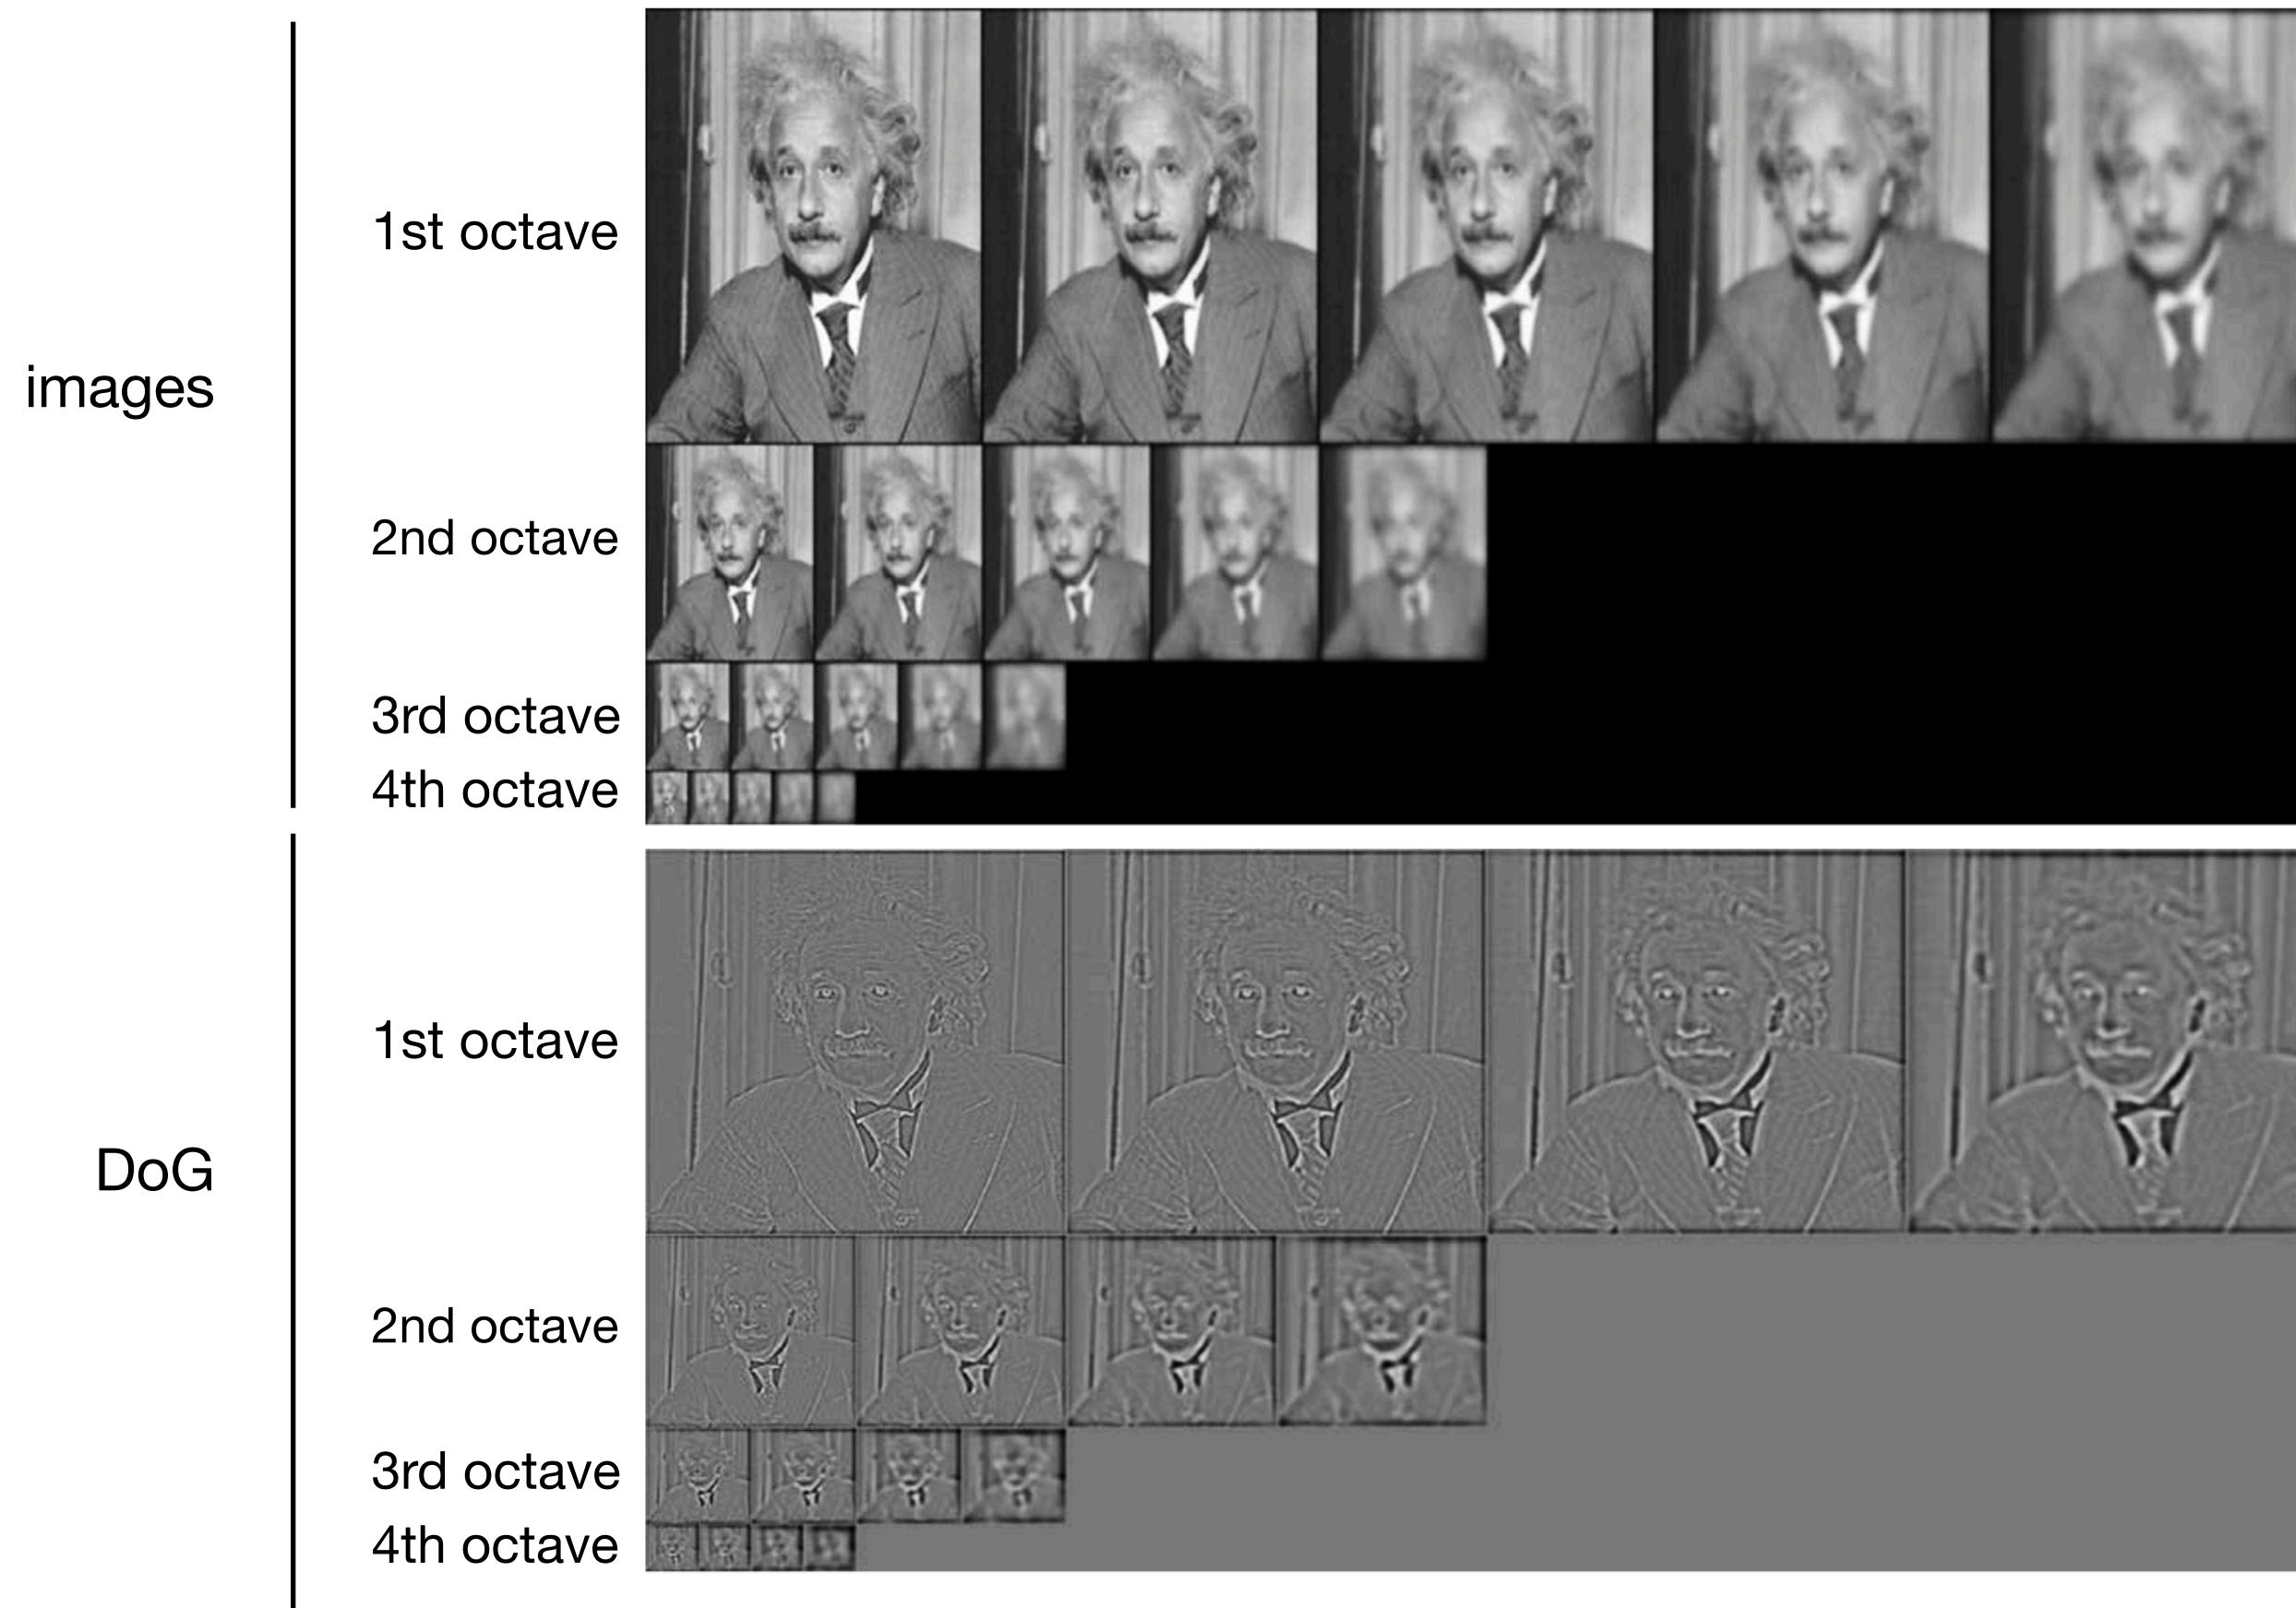
Scale Invariant Feature Transform (SIFT)

<https://faculty.cc.gatech.edu/~afb/classes/CS4495-Fall2013/slides/CS4495-11-Features2.pdf>

Keypoint Detection Scale Invariance

Process the image of interest at different scales.

This is called resolution pyramid.

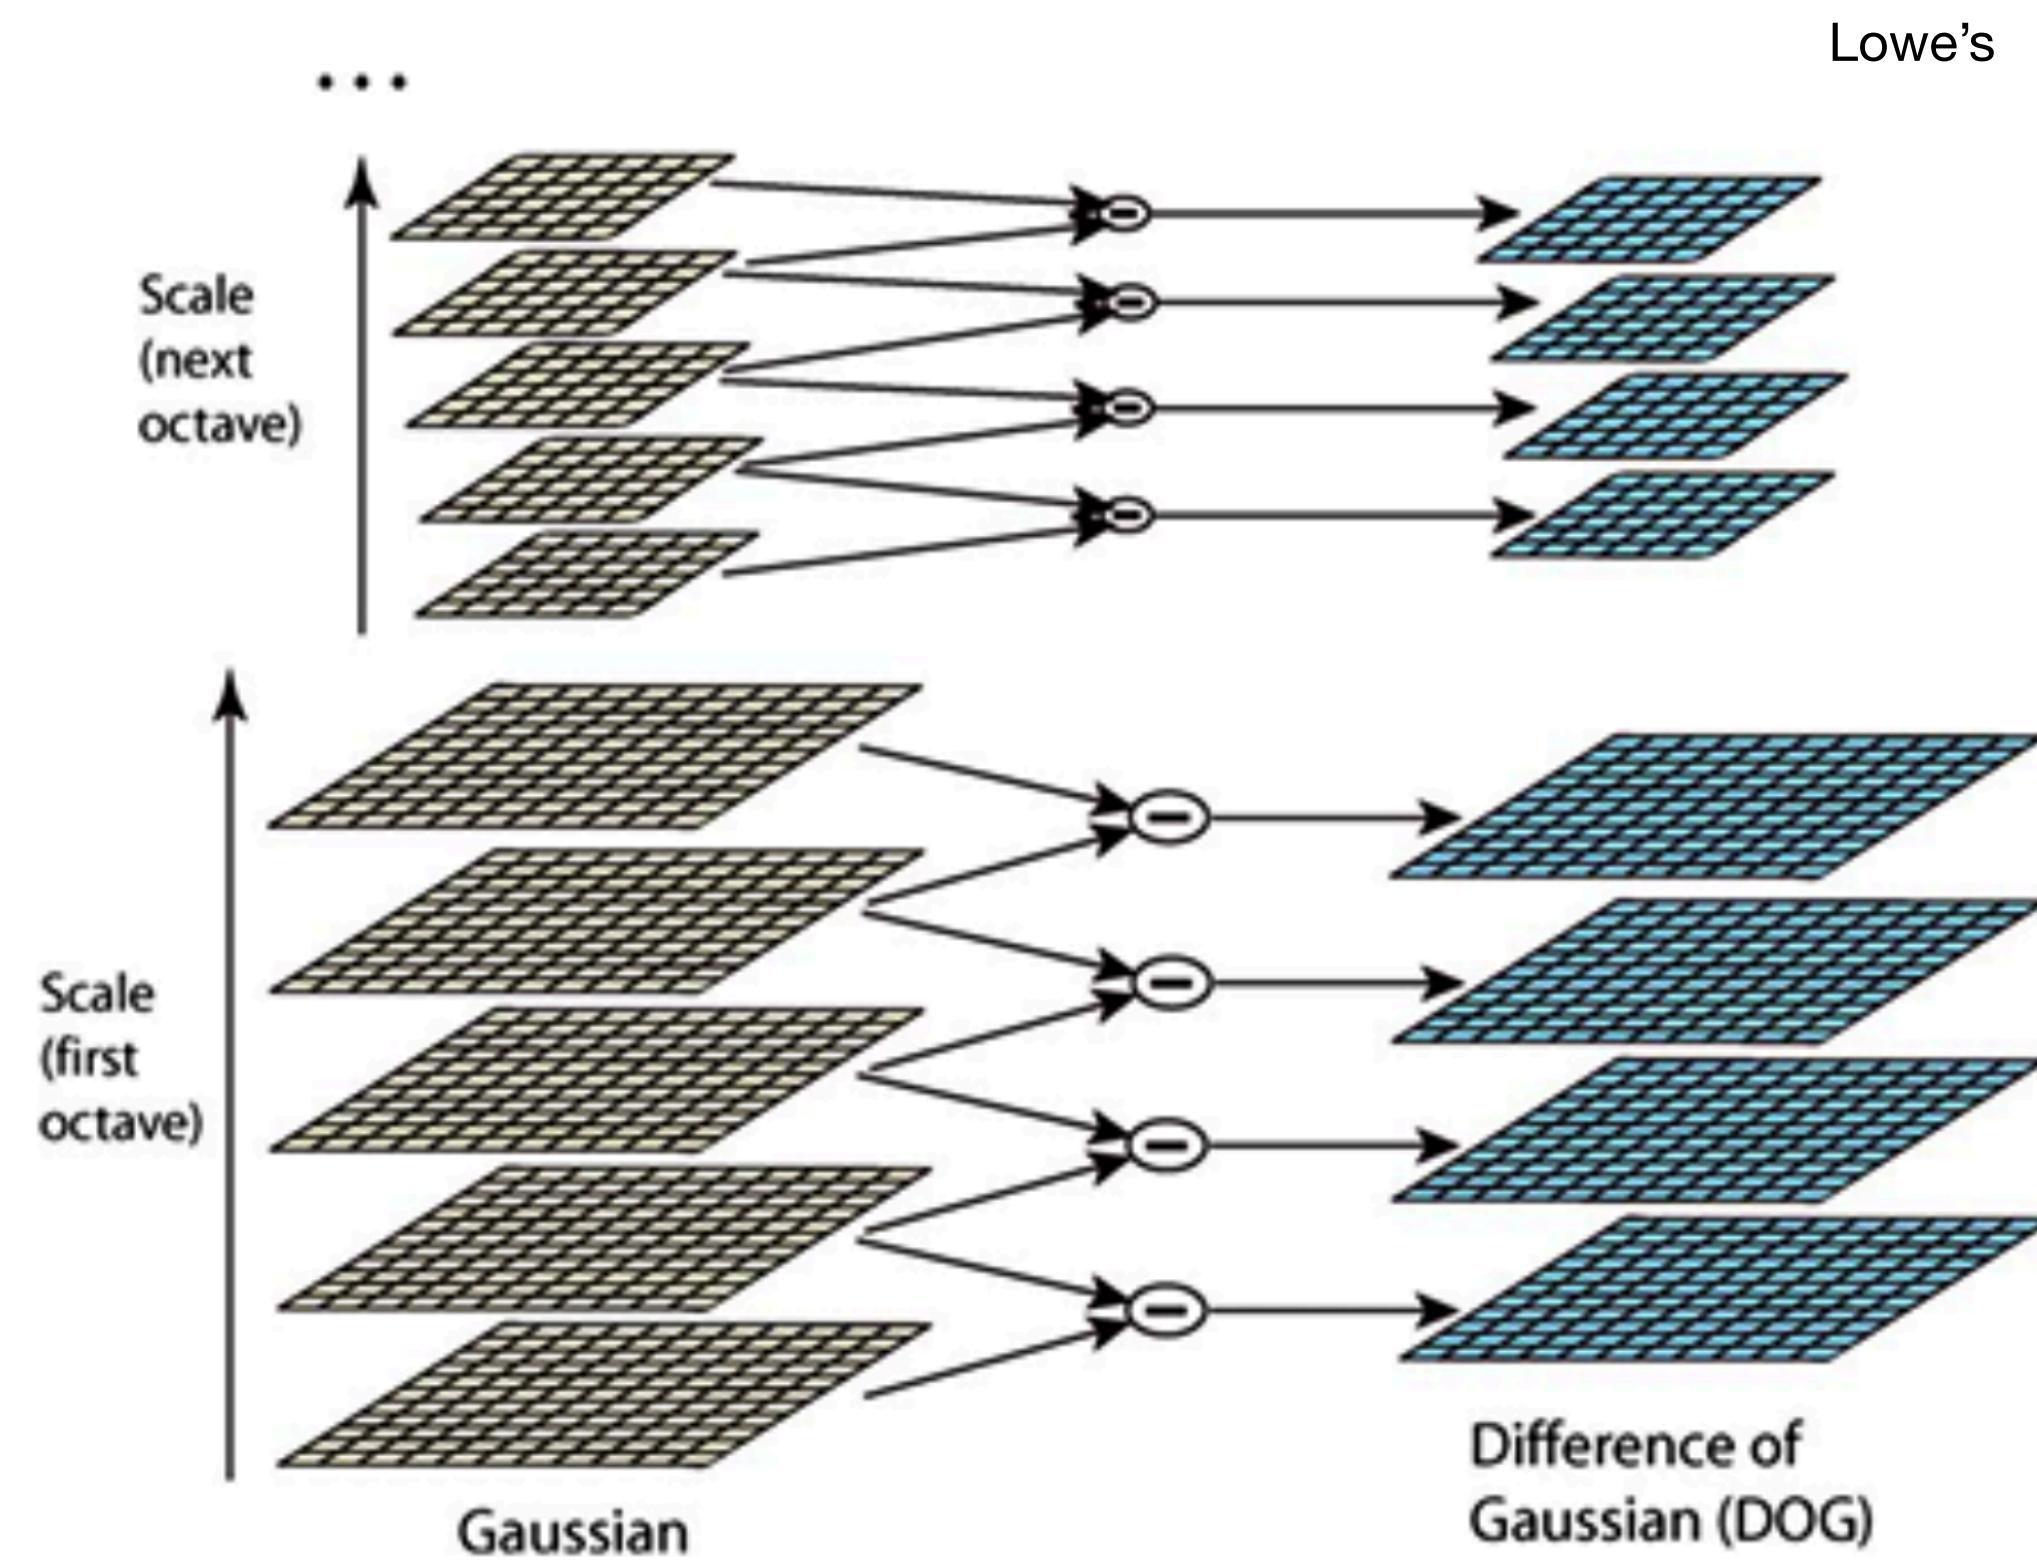


Scale Invariant Feature Transform (SIFT)

Keypoint Detection

Scale Invariance

Process the image of interest
at different scales.



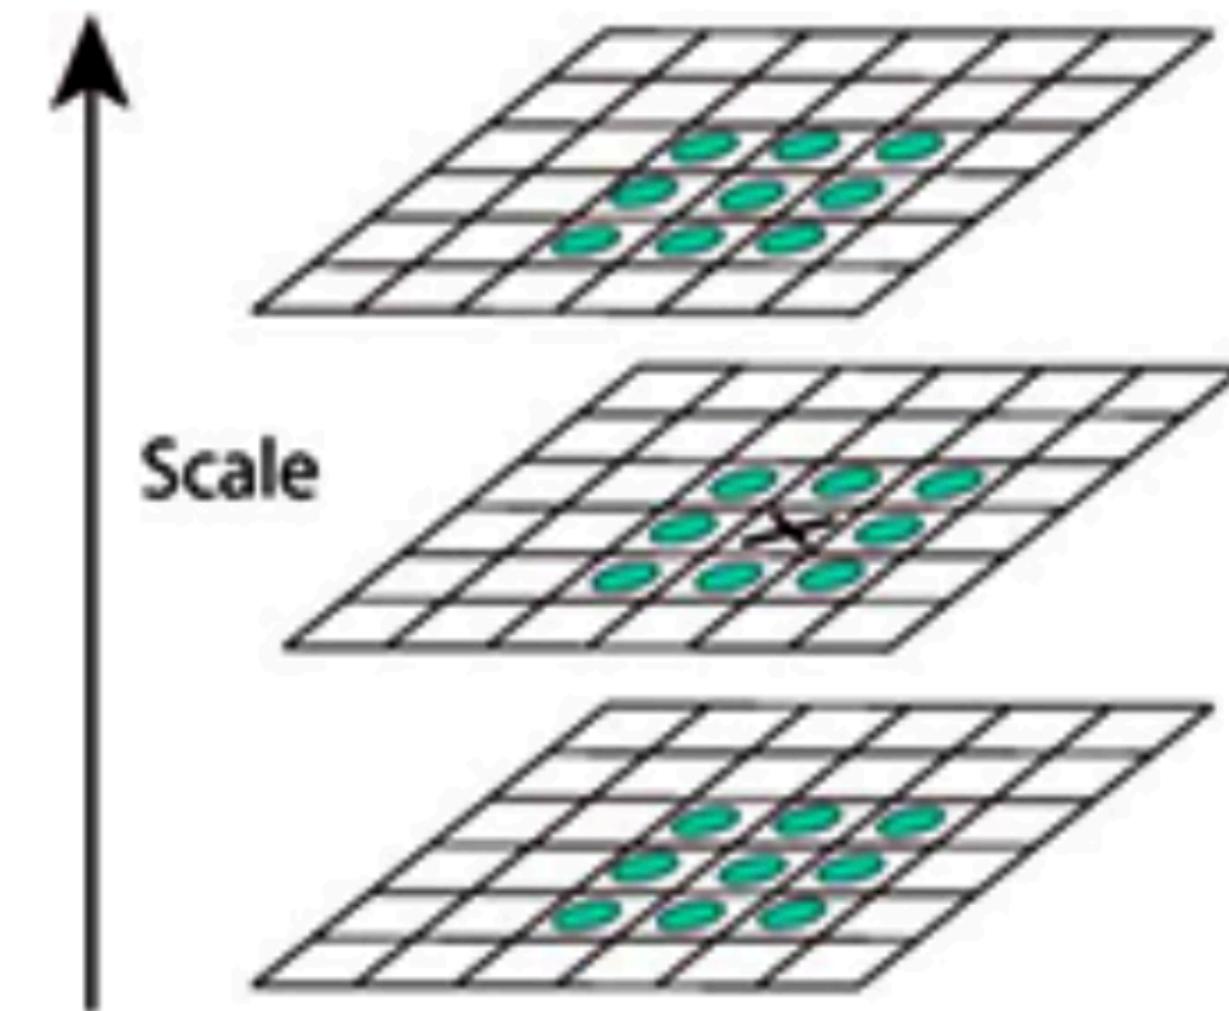
Scale Invariant Feature Transform (SIFT)

Lowe's

Keypoint Detection

Non-maximal Suppression

Good SIFT keypoints present the highest DoG value among their immediate ($3 \times 3 \times 3$) scale-space neighborhood.



Scale Invariant Feature Transform (SIFT)

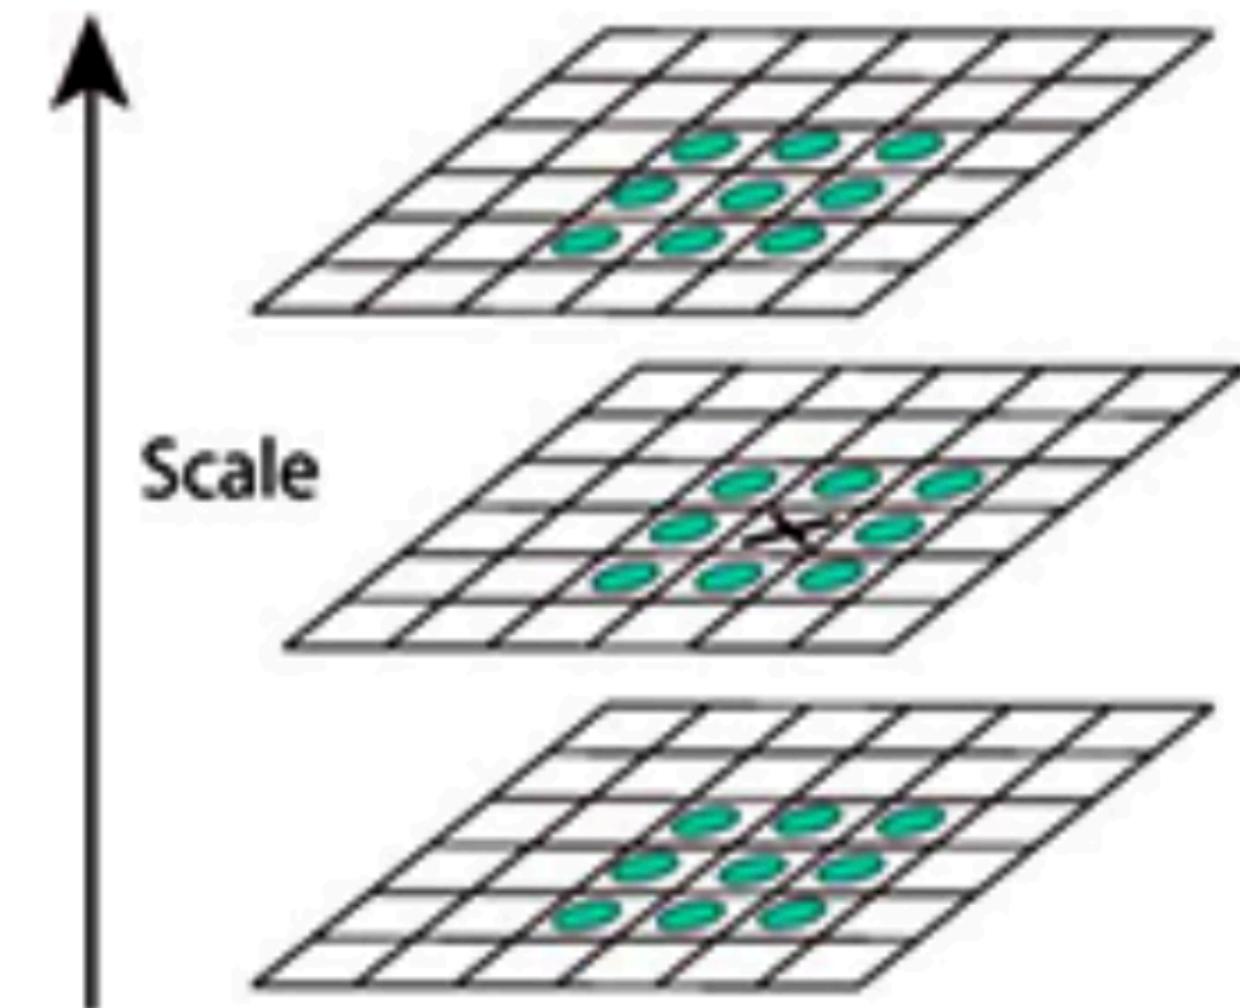
Lowe's

Keypoint Detection

Non-maximal Suppression

Good SIFT keypoints present the highest DoG value among their immediate ($3 \times 3 \times 3$) scale-space neighborhood.

Local scale-space extrema have $I(x, y)$ position and inherit the scale from the level/octave it belongs to within the resolution pyramid.

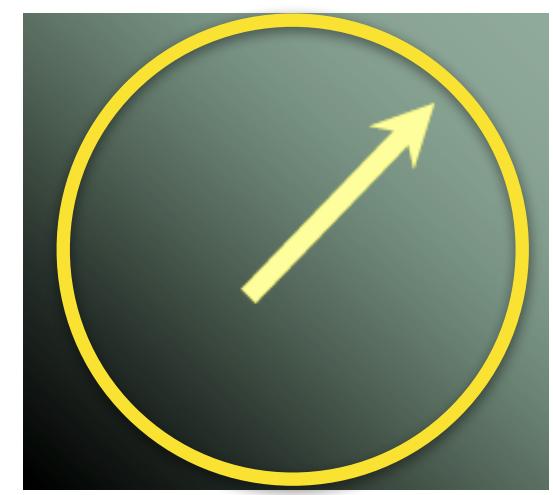
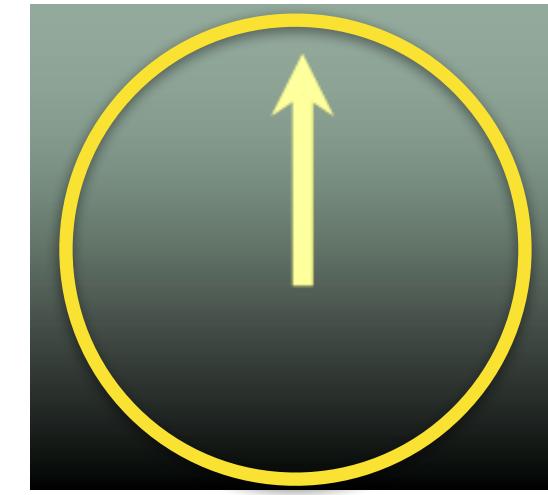
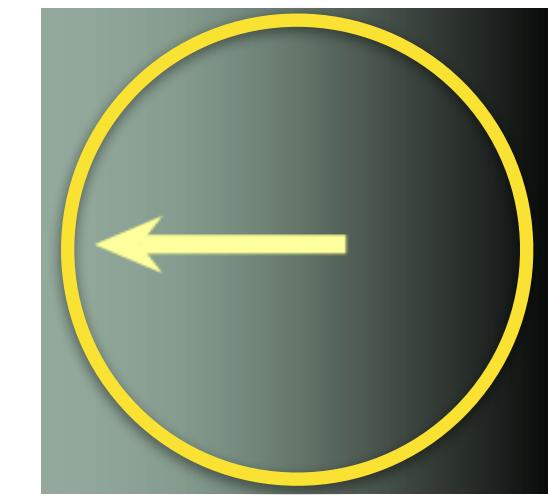


Scale Invariant Feature Transform (SIFT)

Keypoint Detection

Orientation Assignment

To become robust to tilt (rotation), compute the gradient angle for all the pixels within the keypoint neighborhood considering its scale.



Scale Invariant Feature Transform (SIFT)

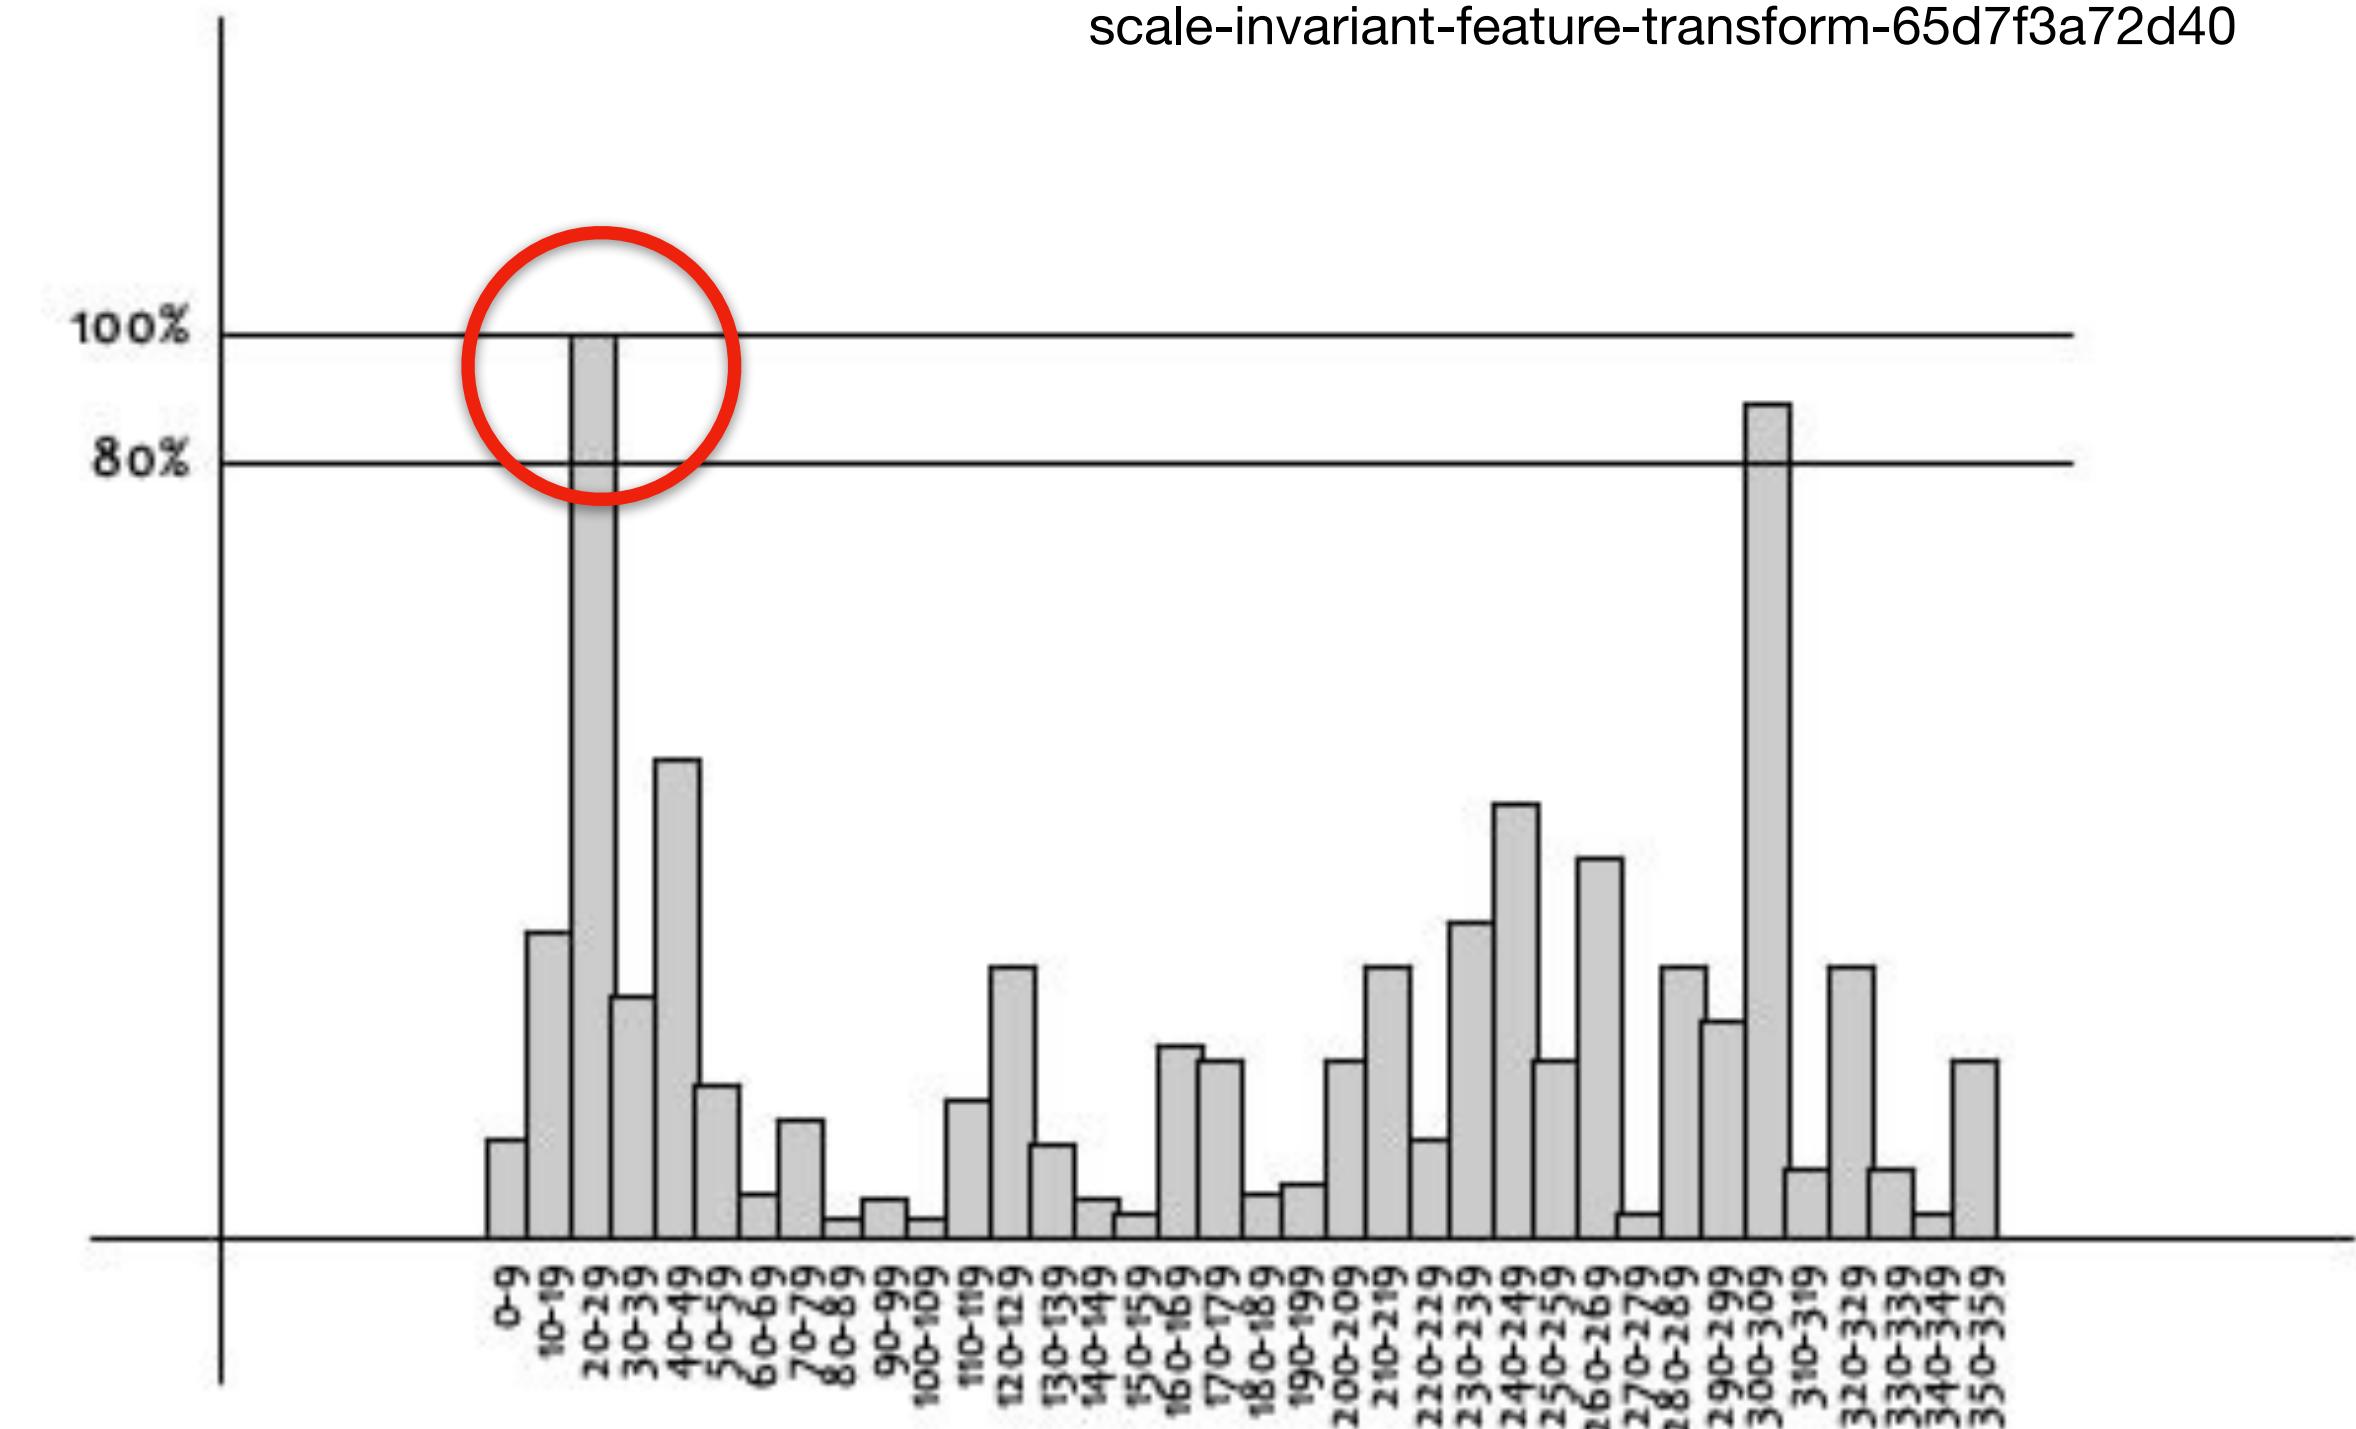
<https://medium.com/data-breach/introduction-to-sift-scale-invariant-feature-transform-65d7f3a72d40>

Keypoint Detection

Orientation Assignment

To become robust to tilt (rotation),
compute the gradient angle for all the pixels
within the keypoint neighborhood
considering its scale.

Create an angle histogram with 36 bins.
Take the dominant angle as the
keypoint orientation.



Scale Invariant Feature Transform (SIFT)

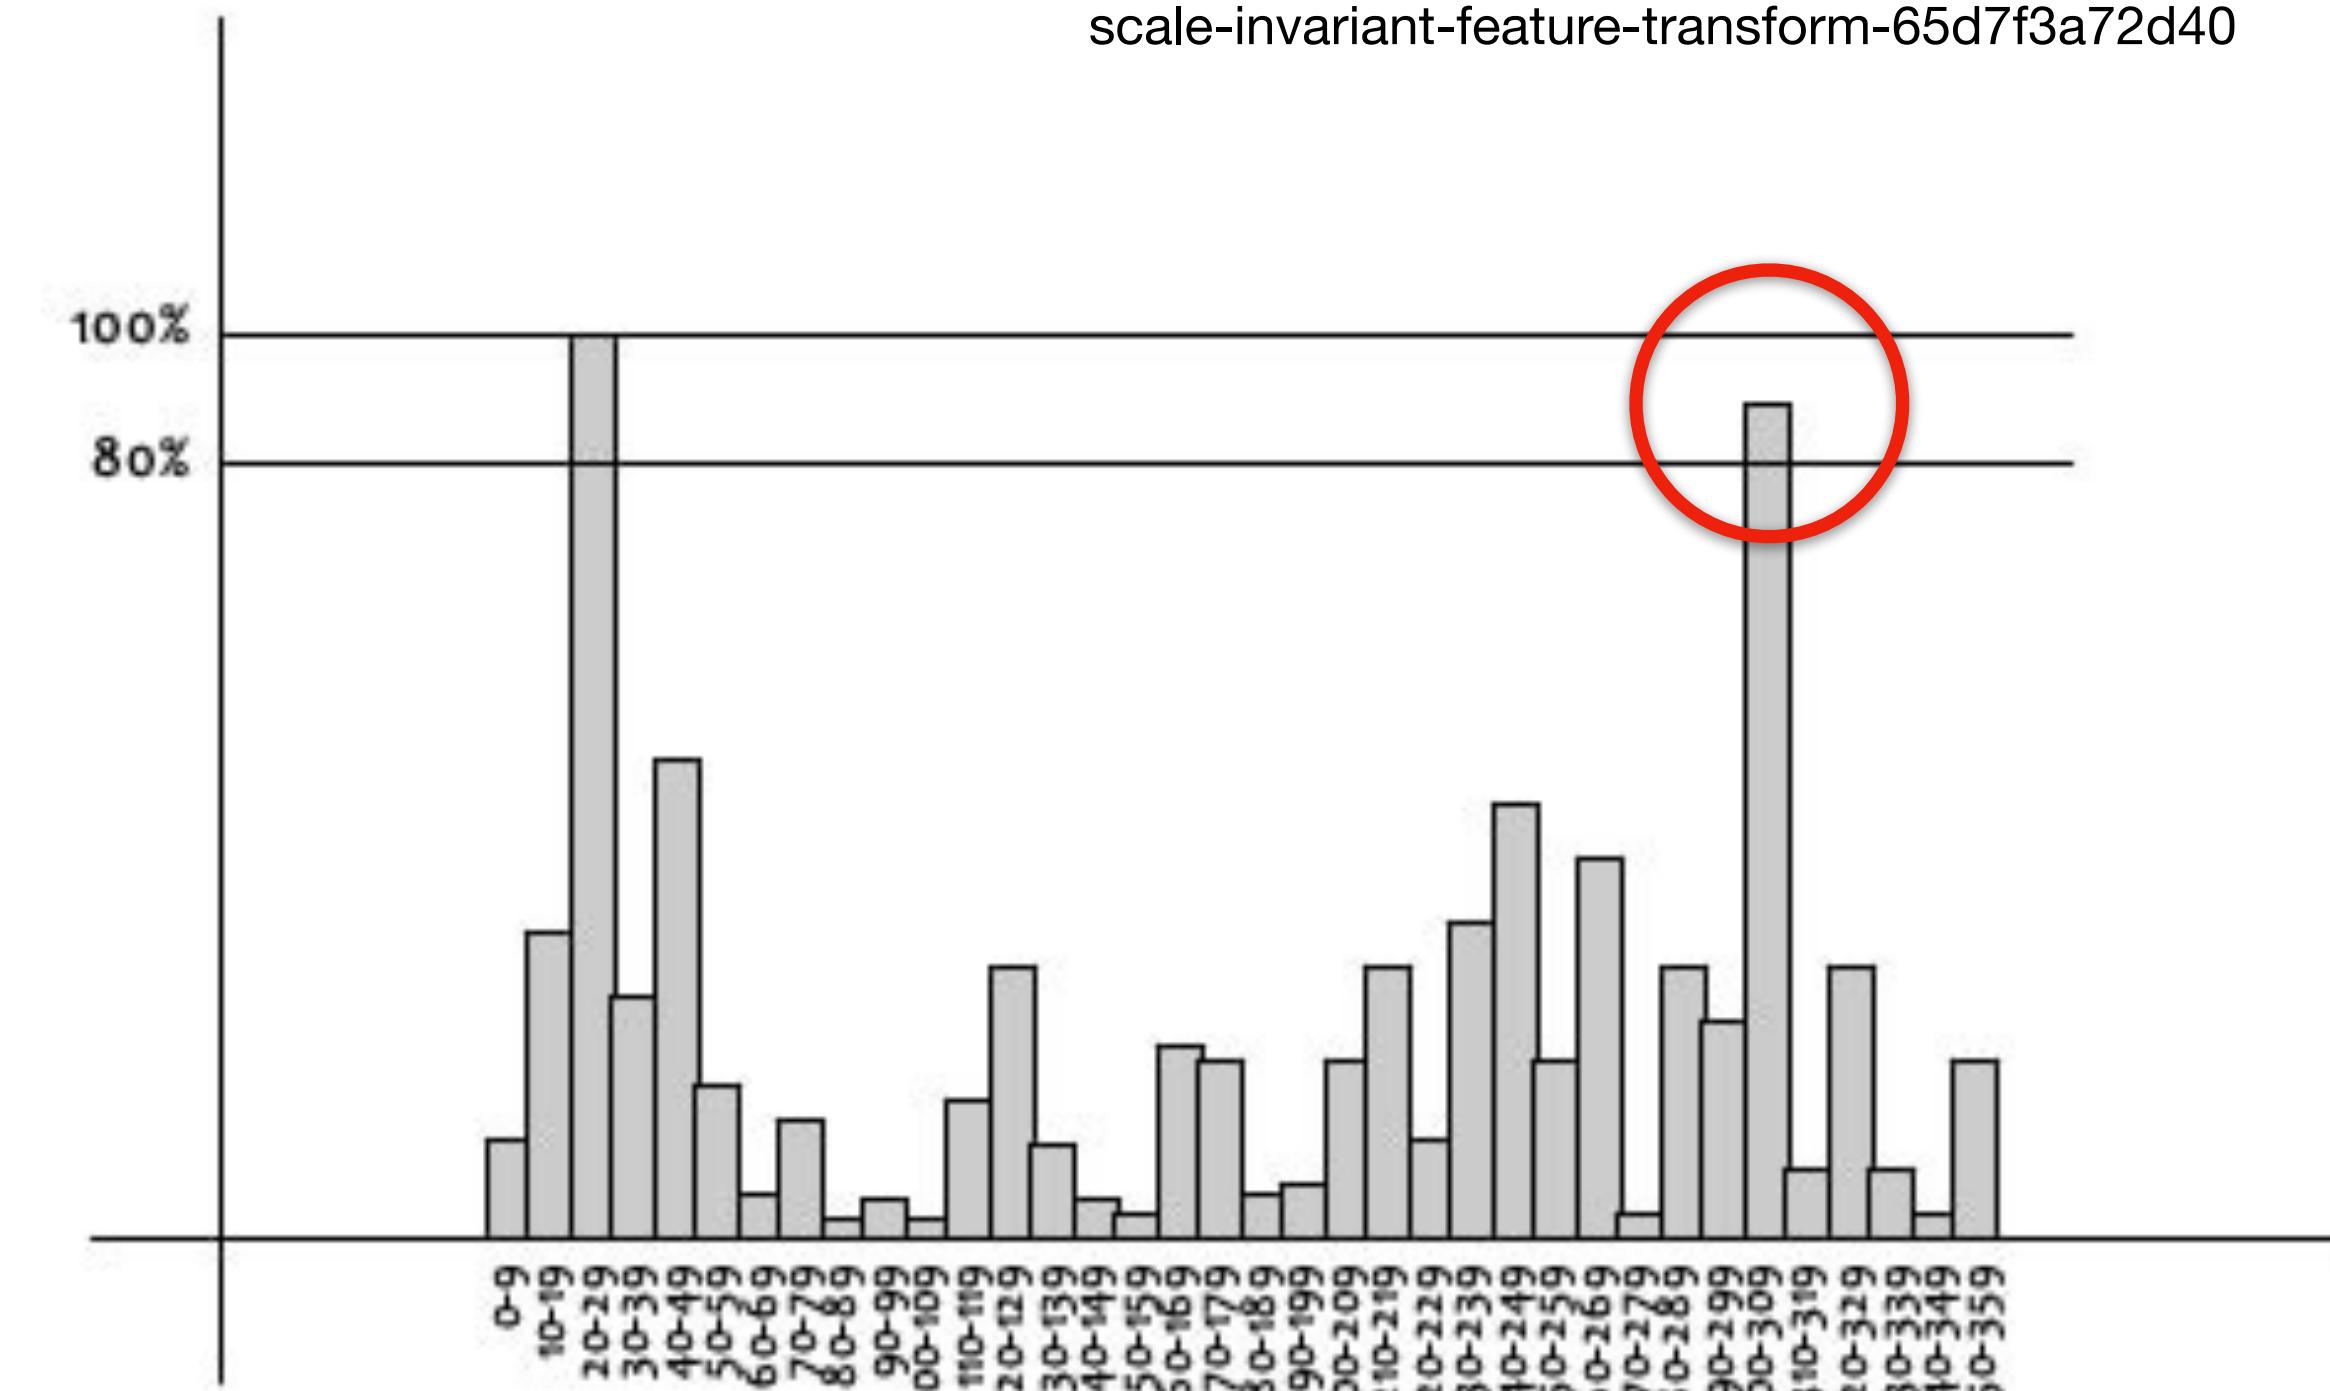
<https://medium.com/data-breach/introduction-to-sift-scale-invariant-feature-transform-65d7f3a72d40>

Keypoint Detection

Orientation Assignment

To become robust to tilt (rotation), compute the gradient angle for all the pixels within the keypoint neighborhood considering its scale.

If other angles have at least 80% of the frequency of the maximum angle, create other keypoints with the same location and scale but different orientation.



Scale Invariant Feature Transform (SIFT)

Keypoint Detection

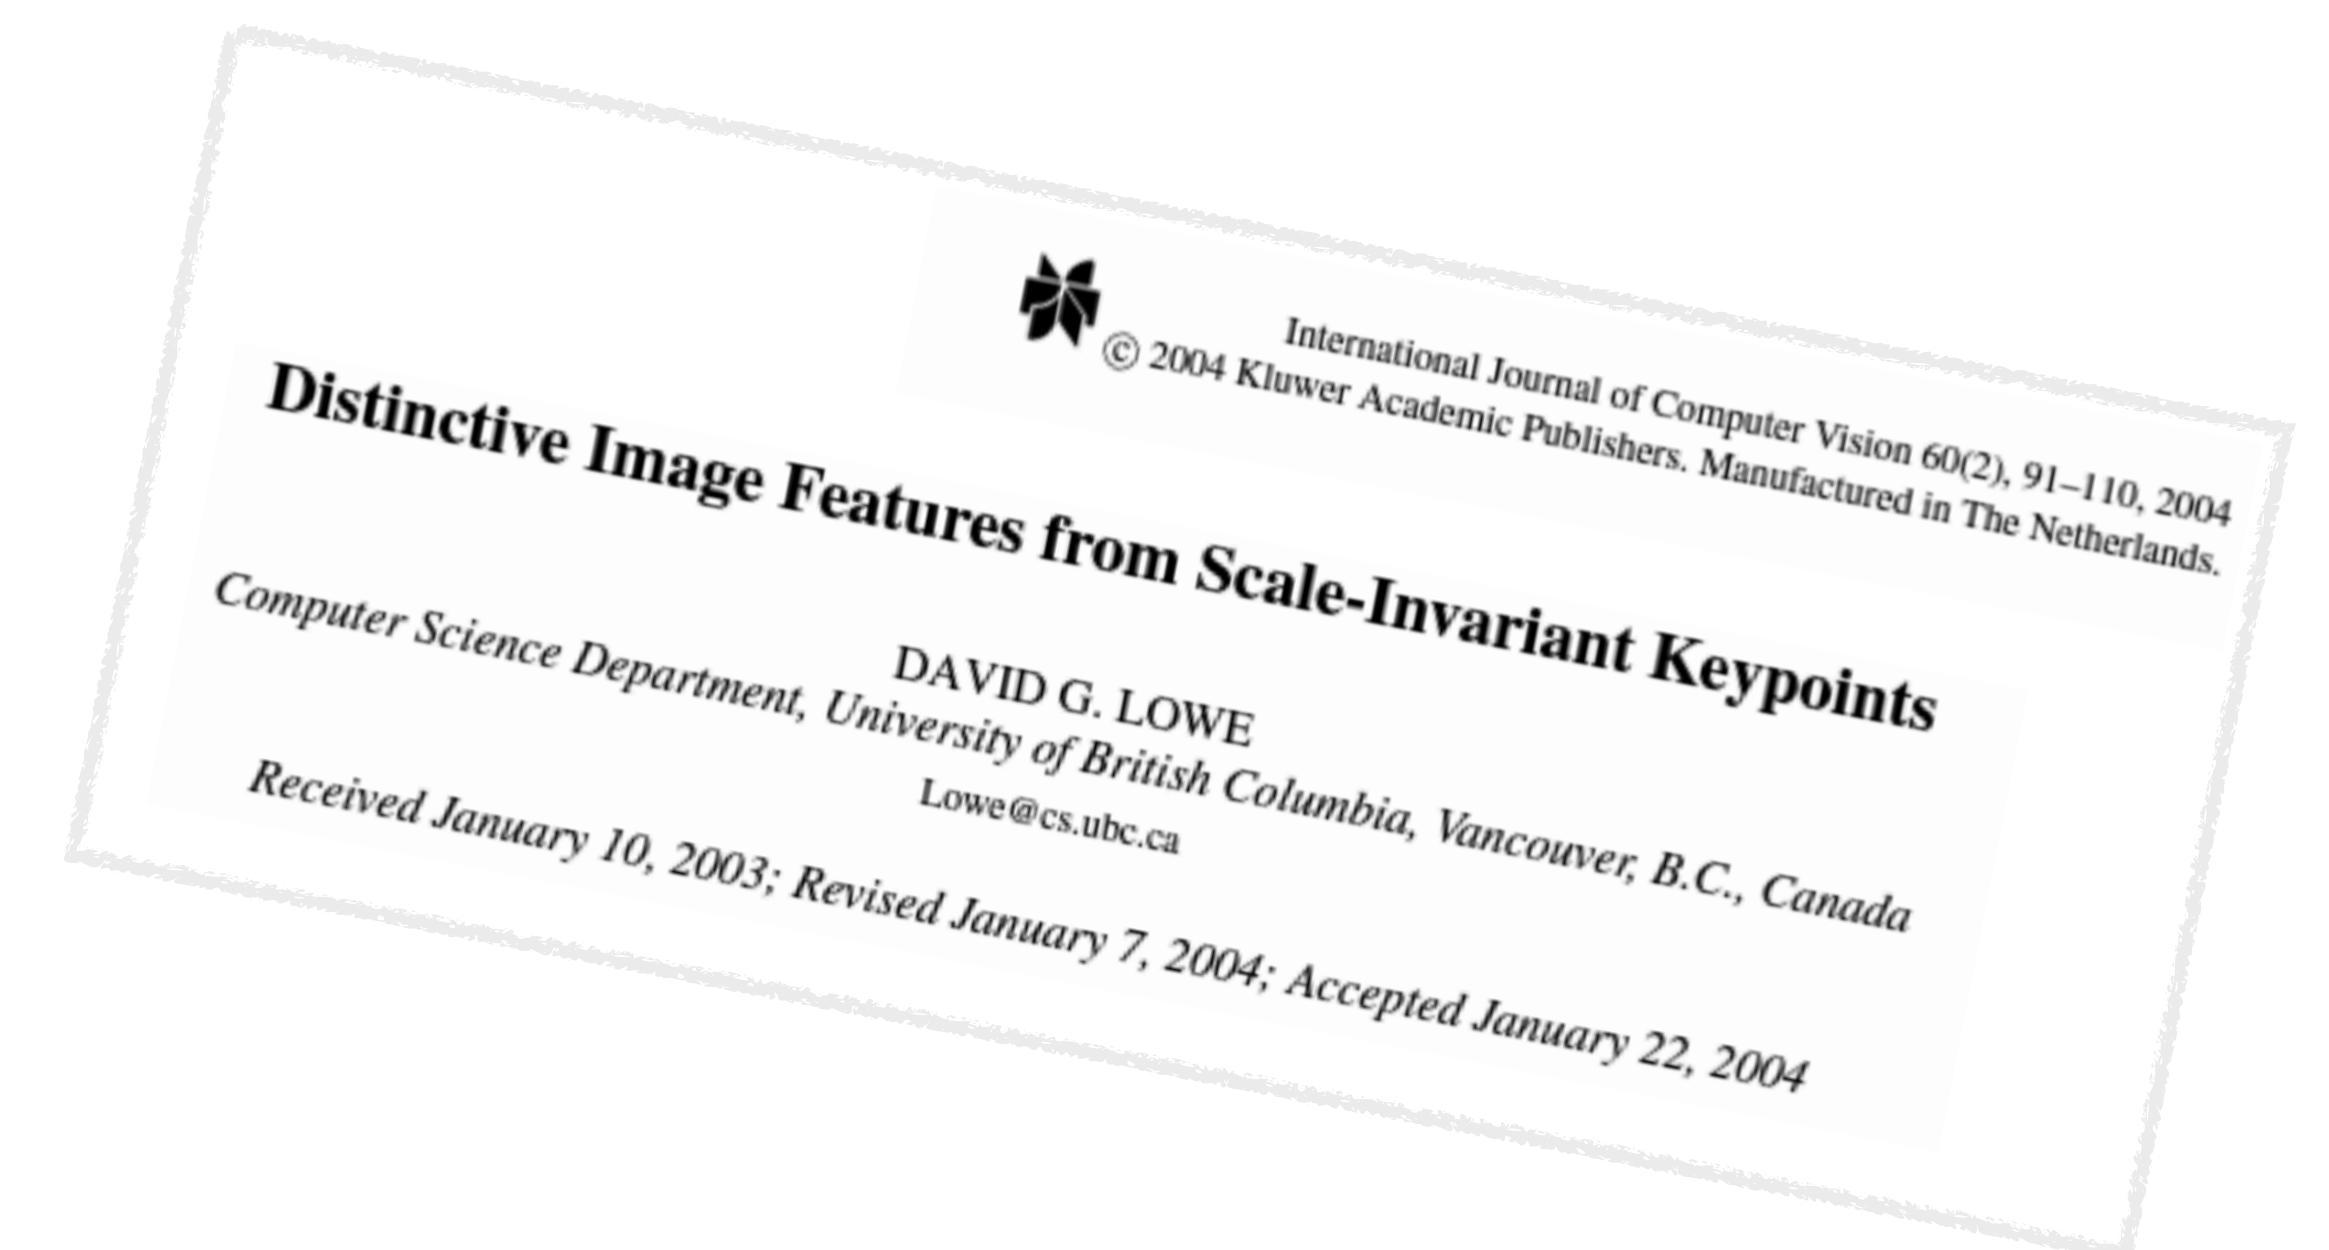
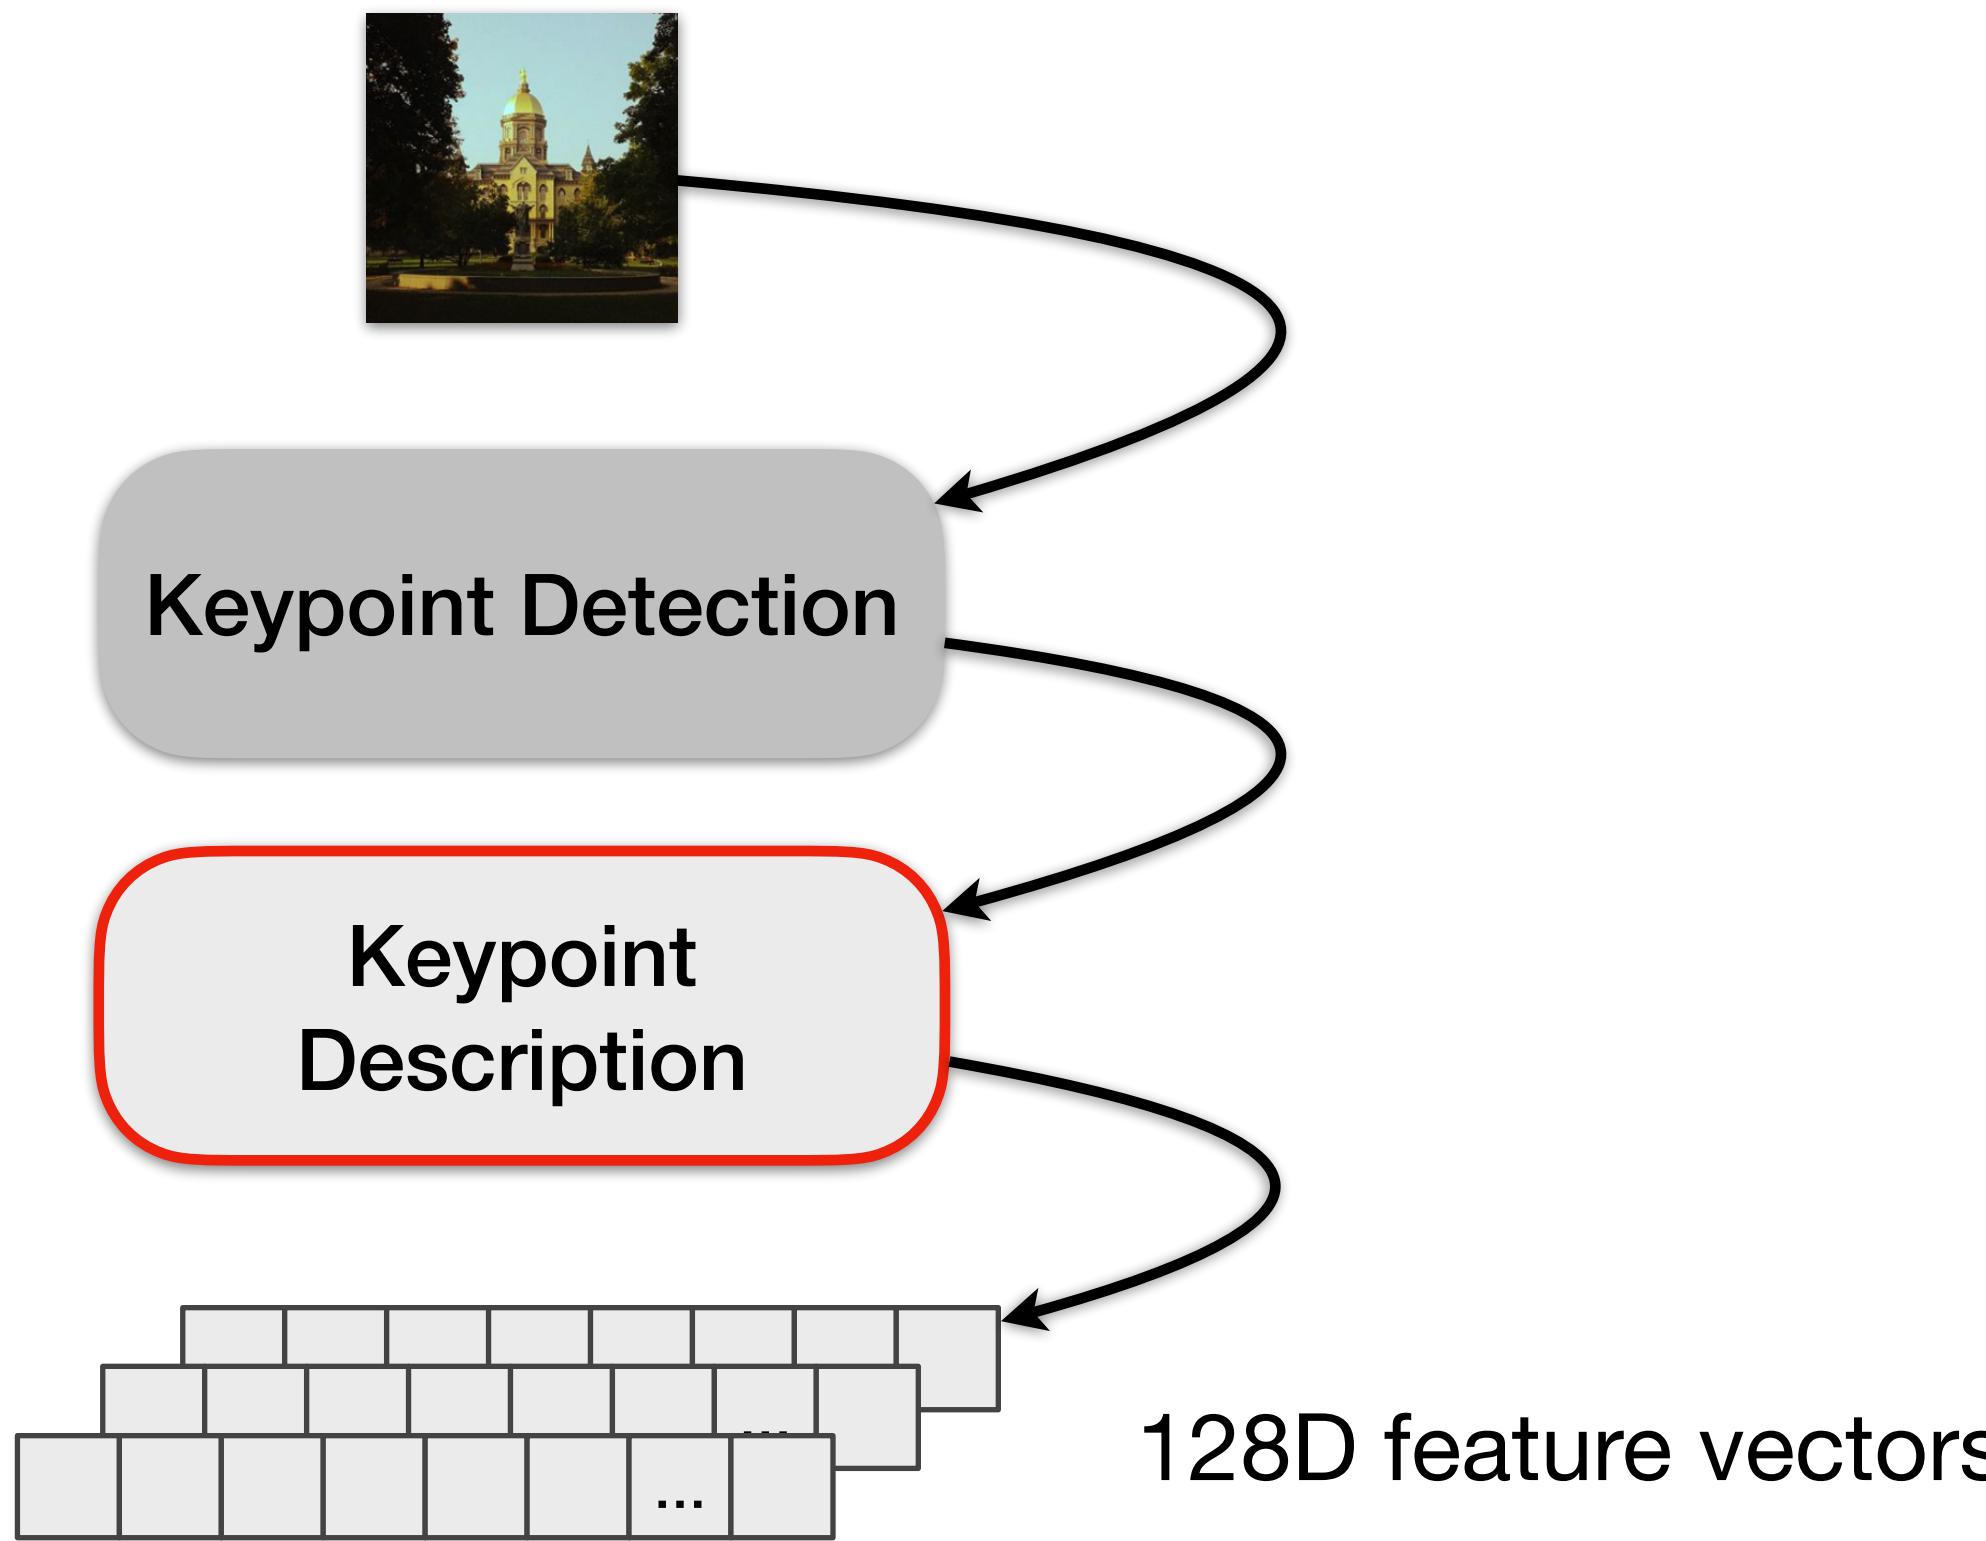
SIFT keypoints have:

1. **Location** (x, y)
2. **Scale** (from resolution pyramid)
3. **Orientation** (dominant local gradient angle)
4. **Strength** (DoG value)



https://docs.opencv.org/4.x/da/df5/tutorial_py_sift_intro.html

Scale Invariant Feature Transform (SIFT)



Scale Invariant Feature Transform (SIFT)

Keypoint Description

For each keypoint, rotate the image according to the keypoint orientation.

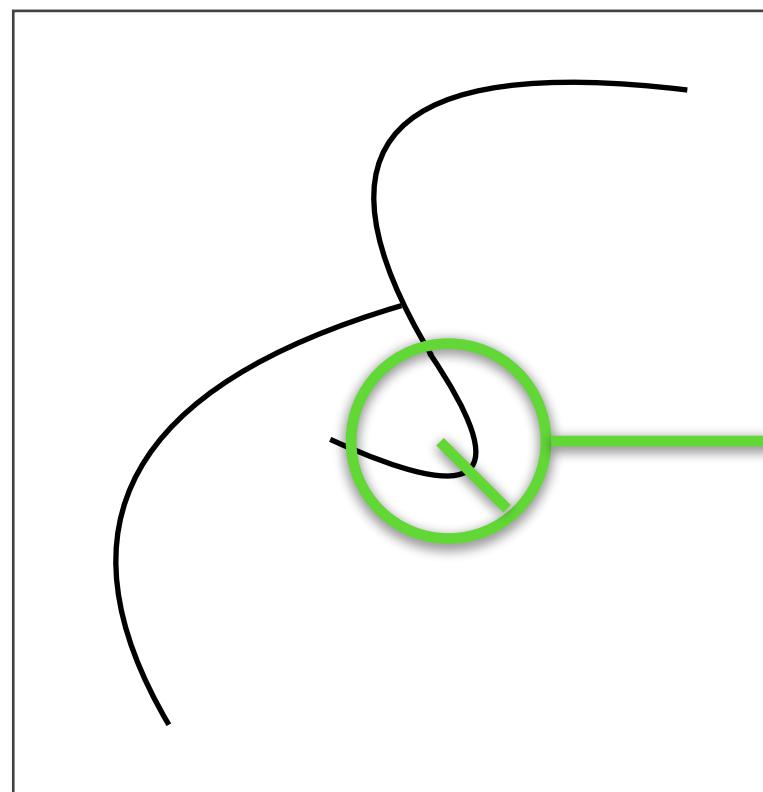


Image 1

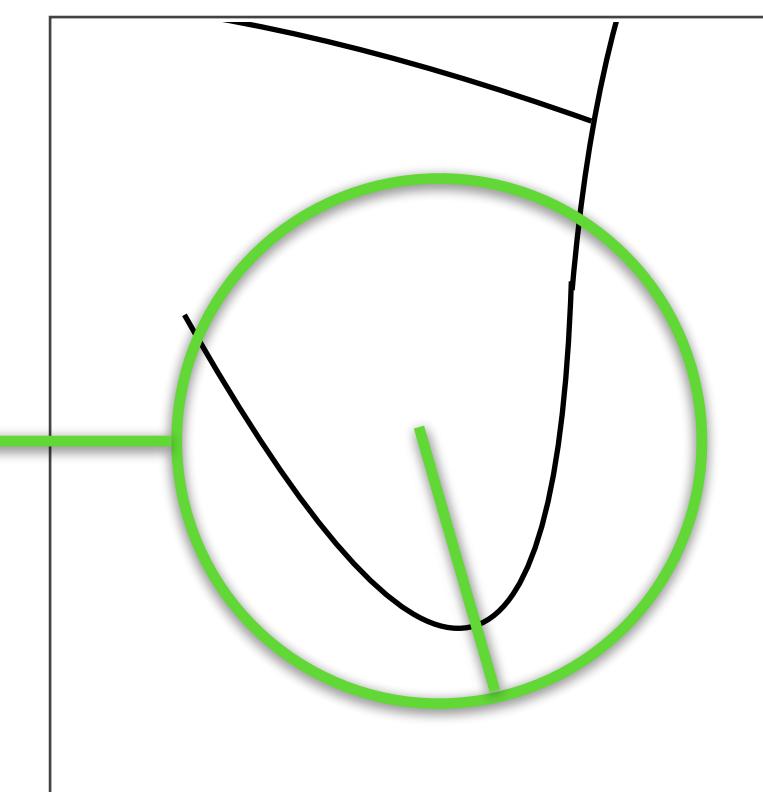


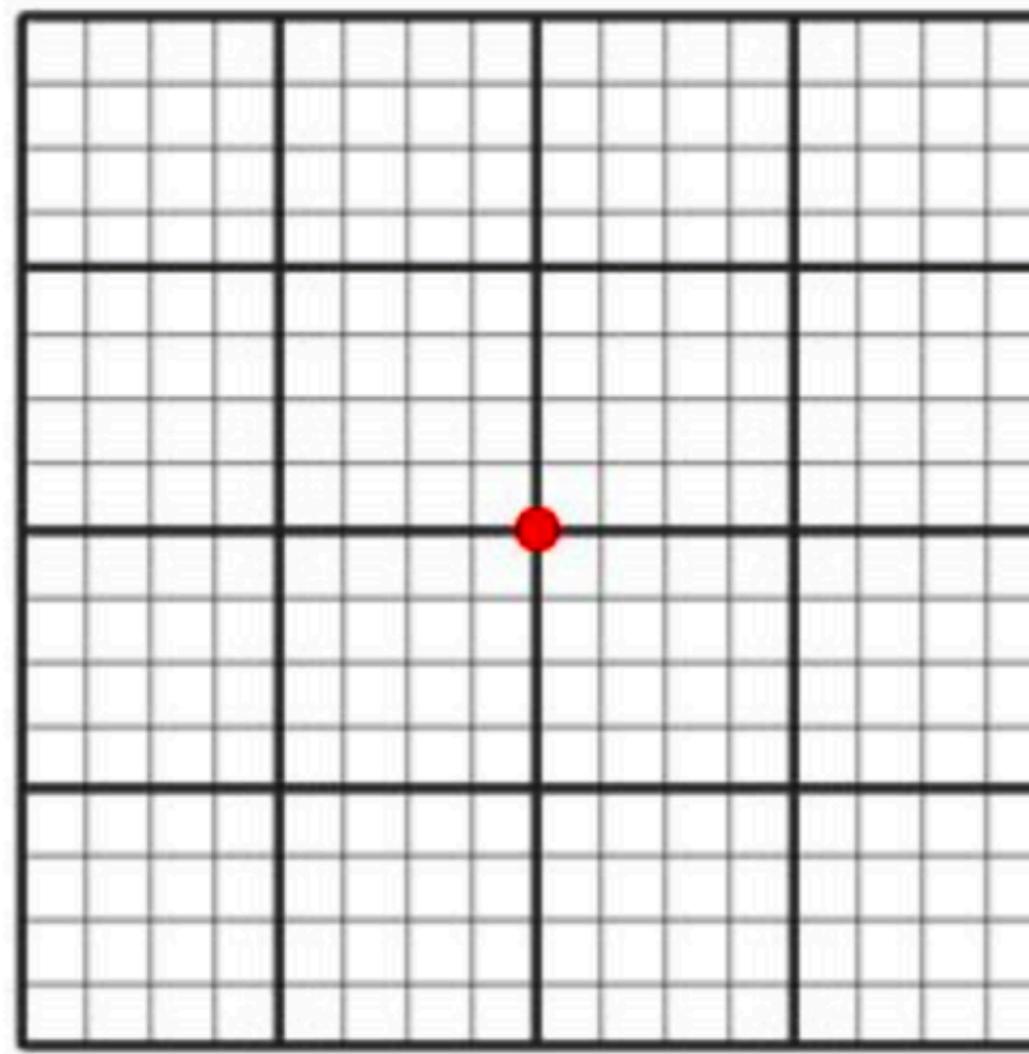
Image 2



Scale Invariant Feature Transform (SIFT)

Keypoint Description

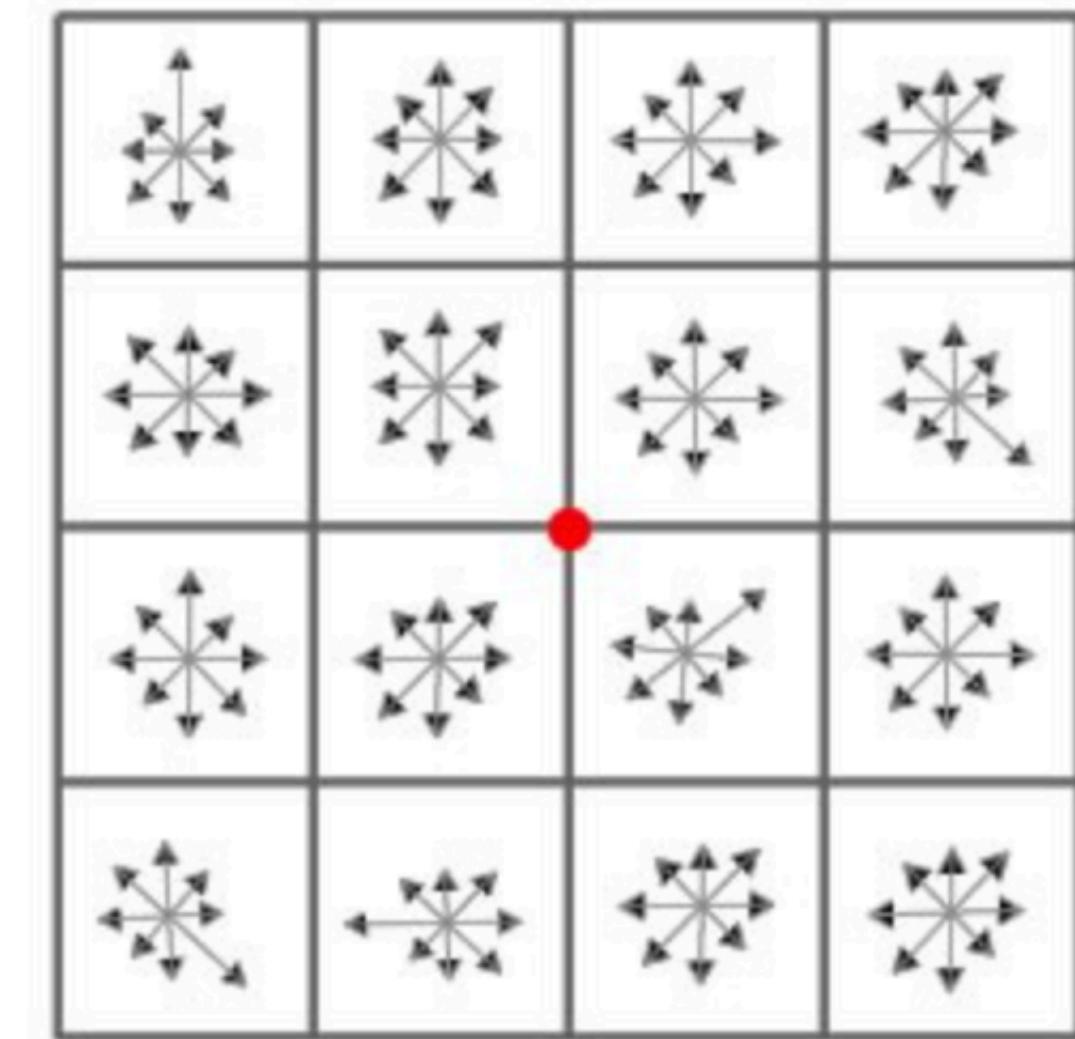
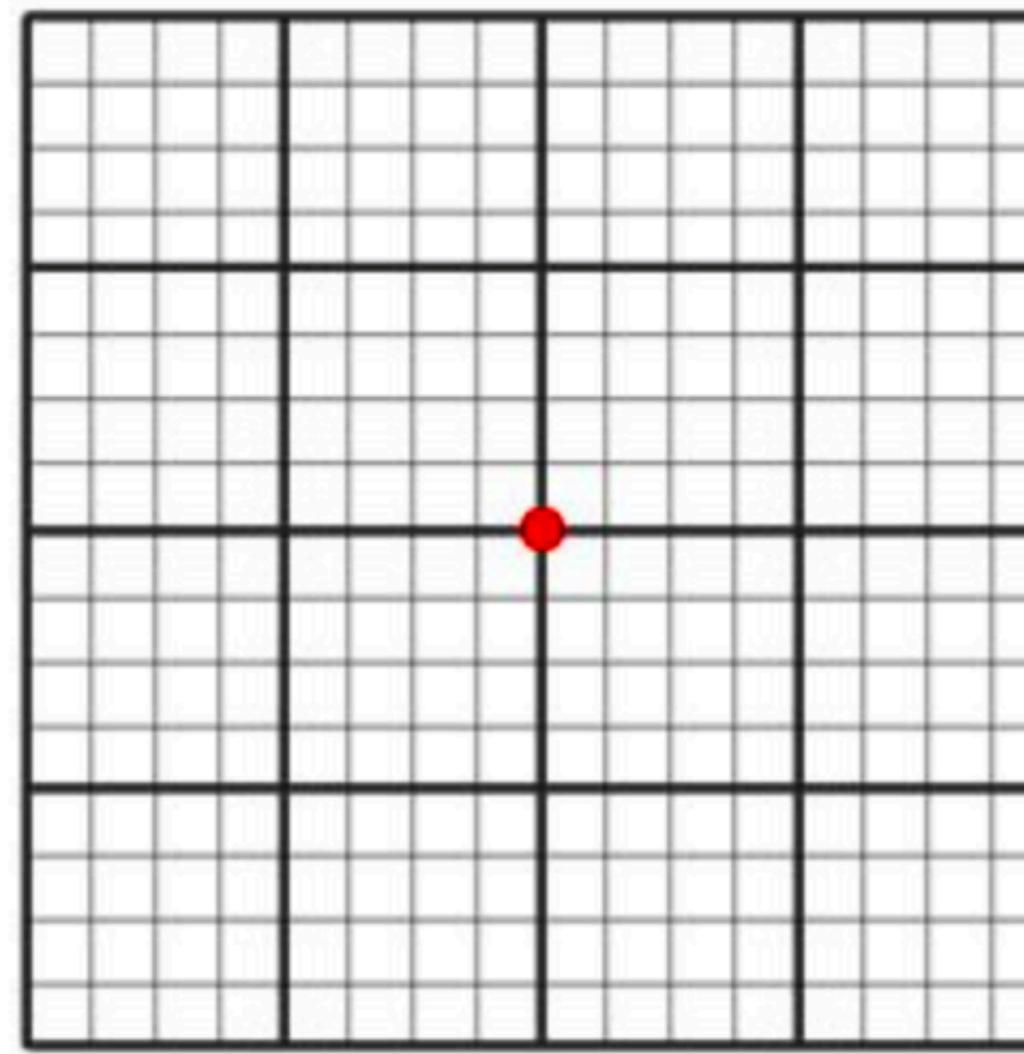
For each rotated keypoint, sample a 4×4 window on its neighborhood, according to the keypoint scale.



Scale Invariant Feature Transform (SIFT)

Keypoint Description

For each rotated keypoint, sample a 4×4 window on its neighborhood, according to the keypoint scale.
For each one of the 4×4 cells, compute a 8-bin histogram of gradient directions.



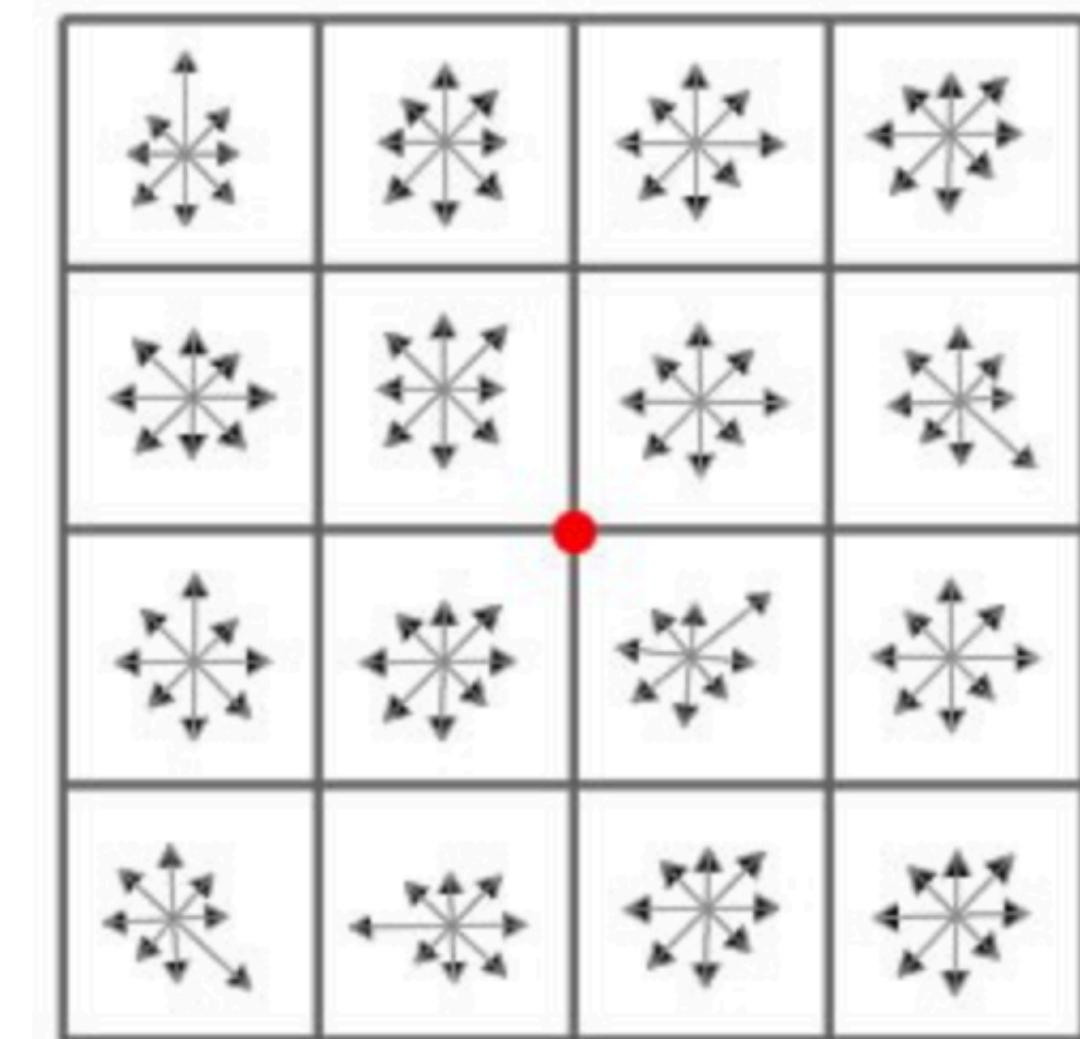
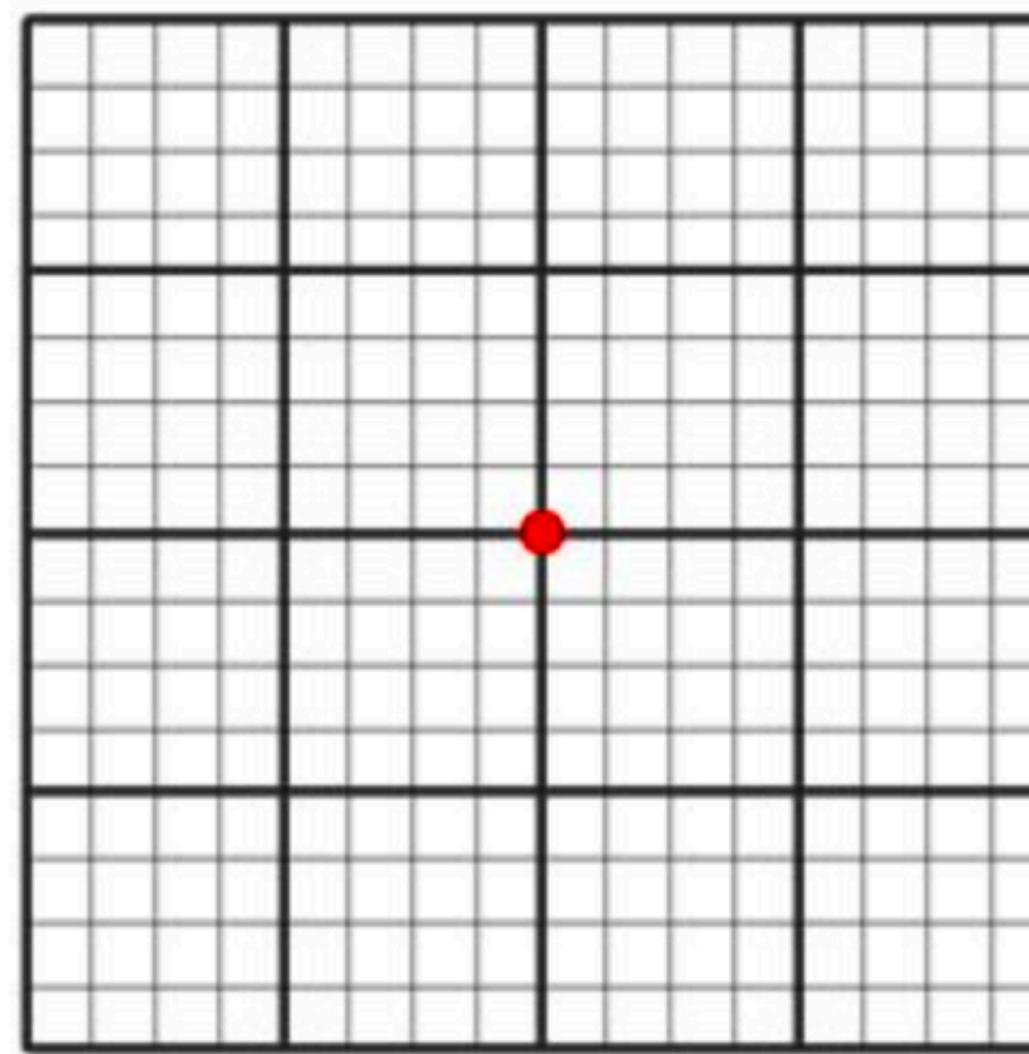
Scale Invariant Feature Transform (SIFT)

Keypoint Description

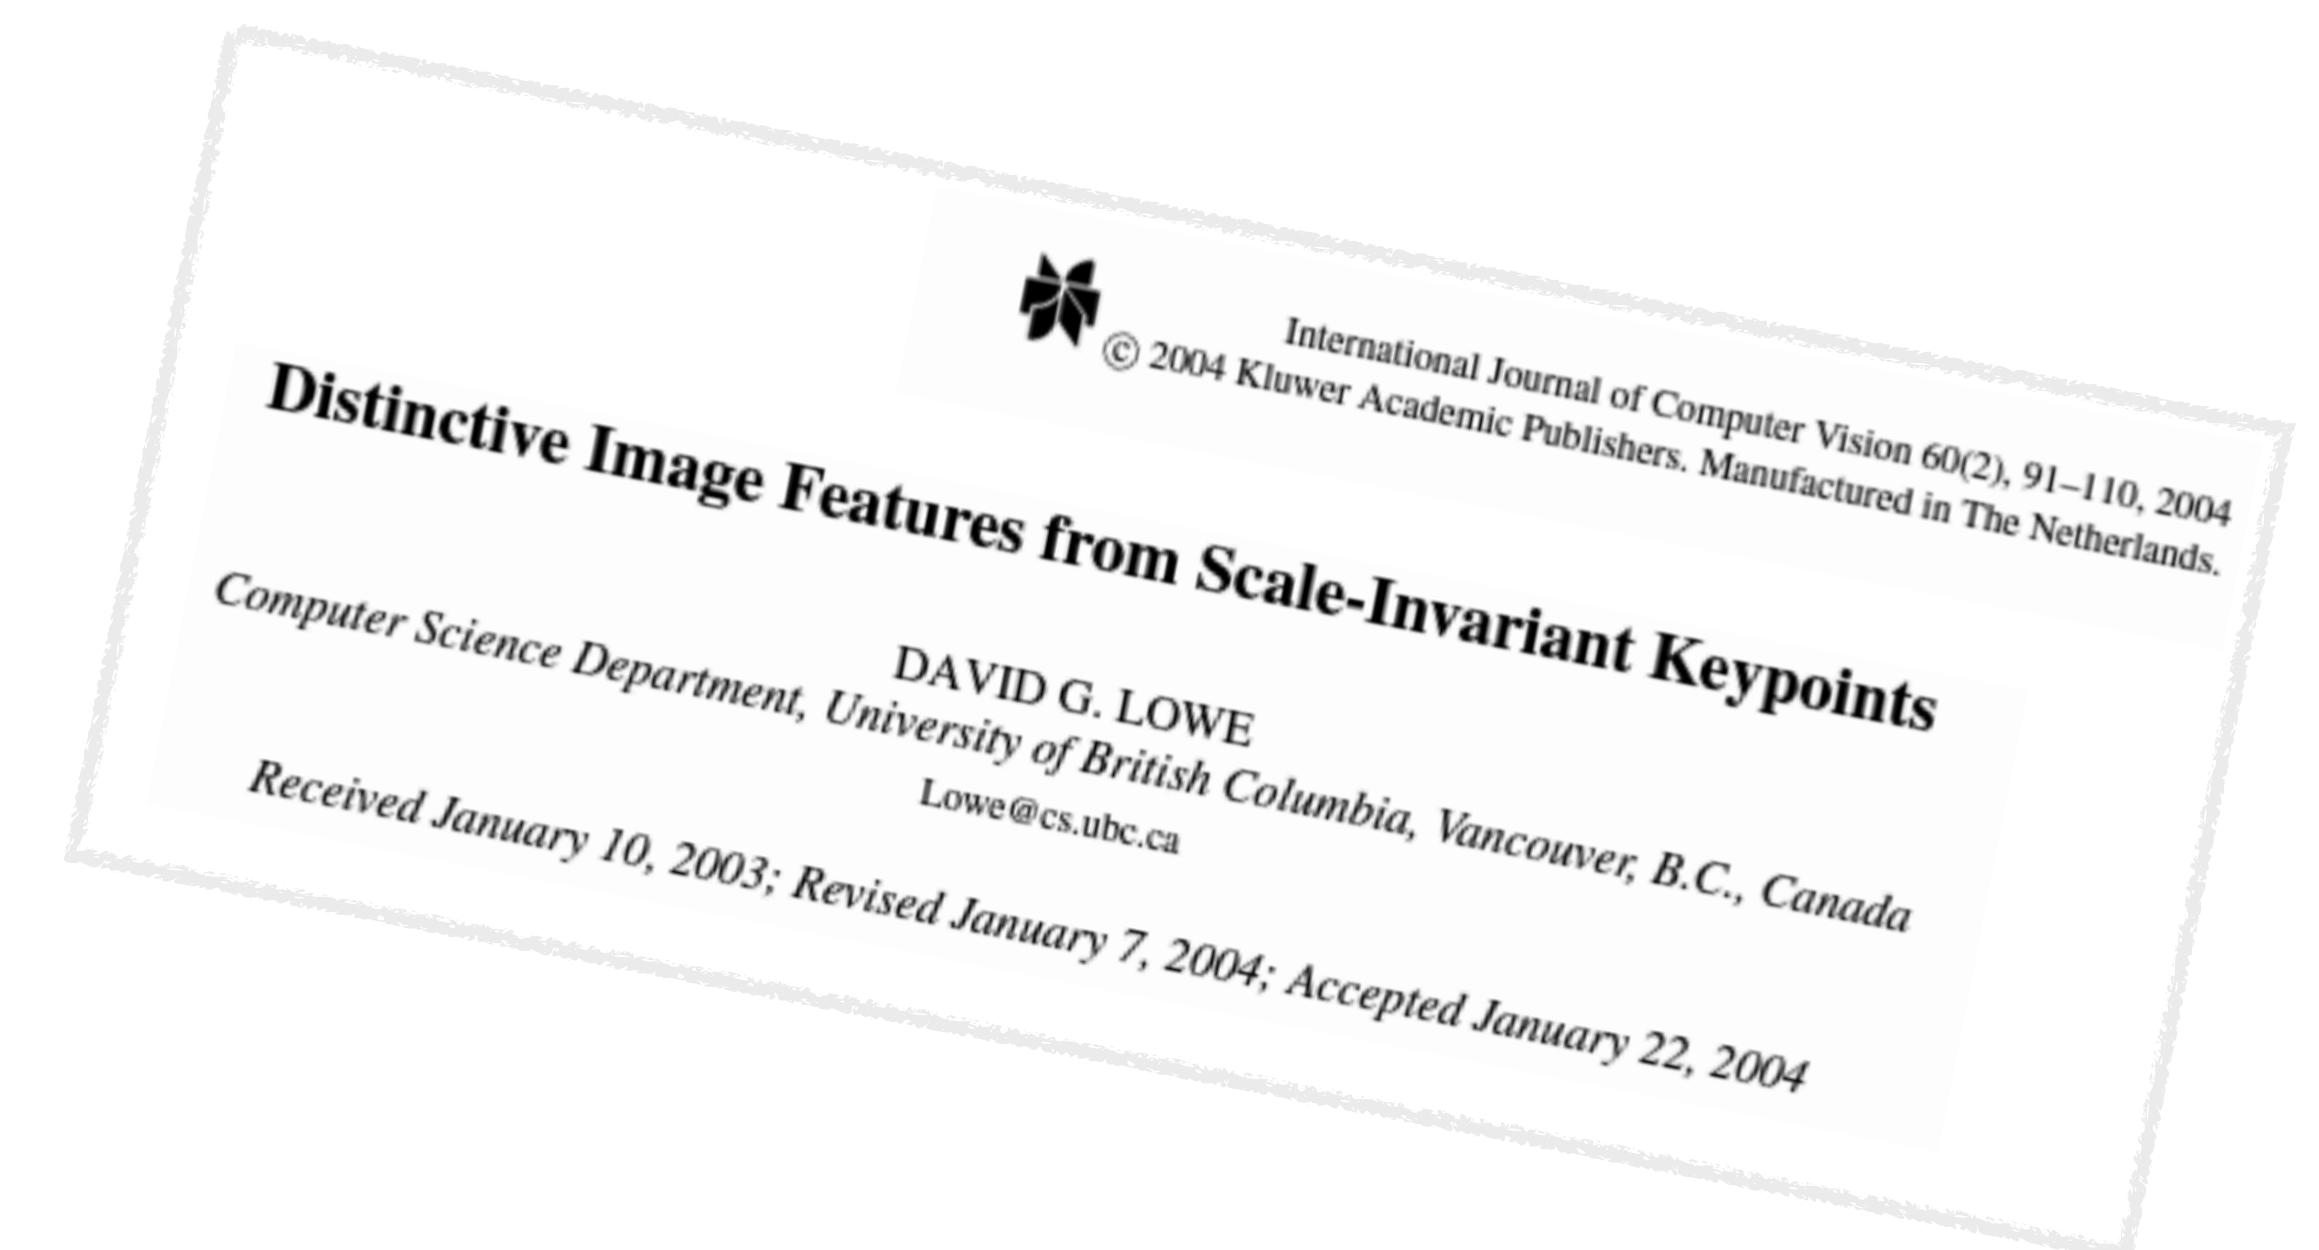
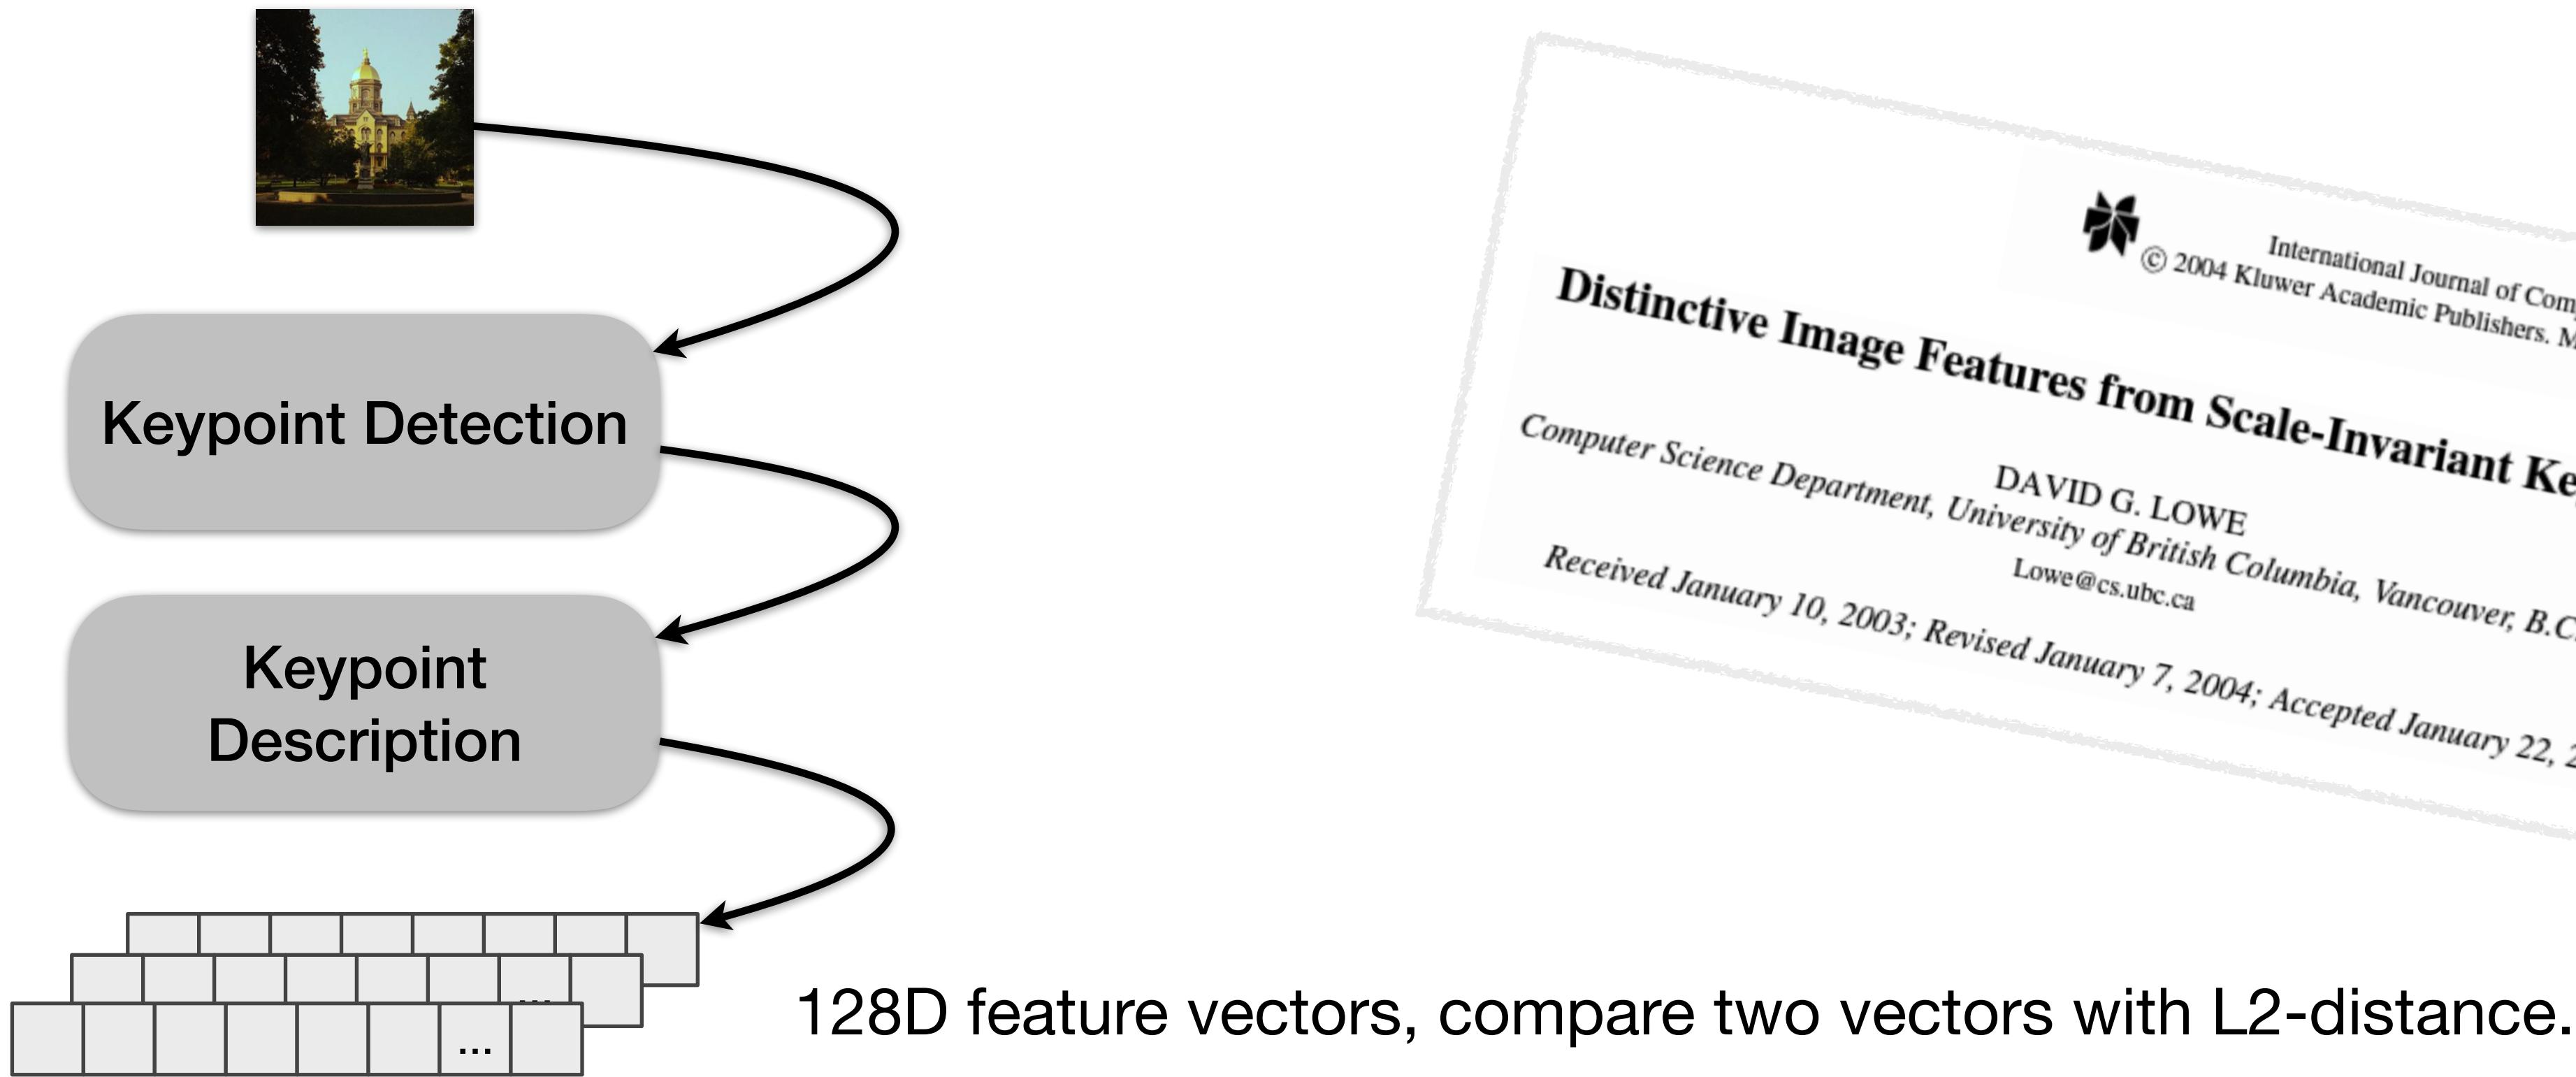
For each rotated keypoint, sample a 4×4 window on its neighborhood, according to the keypoint scale.

For each one of the 4×4 cells, compute a 8-bin histogram of gradient directions.

Fill out a feature vector with the $4 \times 4 \times 8 = 128$ histogram values.



Scale Invariant Feature Transform (SIFT)



Scale Invariant Feature Transform (SIFT)

Example



<https://bit.ly/3QP3EKw>

COMP 388 - SIFT Example

File Edit View Insert Runtime Tools Help All changes saved

Table of contents

- SIFT Example
- Webcam Usage
- SIFT keypoints detection and description
- Importing of the necessary libraries.
- Loading webcam image into memory.
- Detecting and Describing SIFT keypoints

+ Code + Text

▼ SIFT Example

▼ Webcam Usage

```
[37] from IPython.display import display, Javascript
      from google.colab.output import eval_js
      from base64 import b64decode

def take_photo(filename='photo.jpg', quality=0.8):
    js = Javascript('''
        async function takePhoto(quality) {
            const div = document.createElement('div');
            ... ''')
```

Speeded-Up Robust Features (SURF)

How to develop a faster alternative to SIFT?



https://docs.opencv.org/4.x/da/df5/tutorial_py_sift_intro.html

SURF: Speeded Up Robust Features

Herbert Bay¹, Tinne Tuytelaars², and Luc Van Gool^{1,2}

¹ ETH Zurich

{bay, vangool}@vision.ee.ethz.ch

2006

² Katholieke Universiteit Leuven

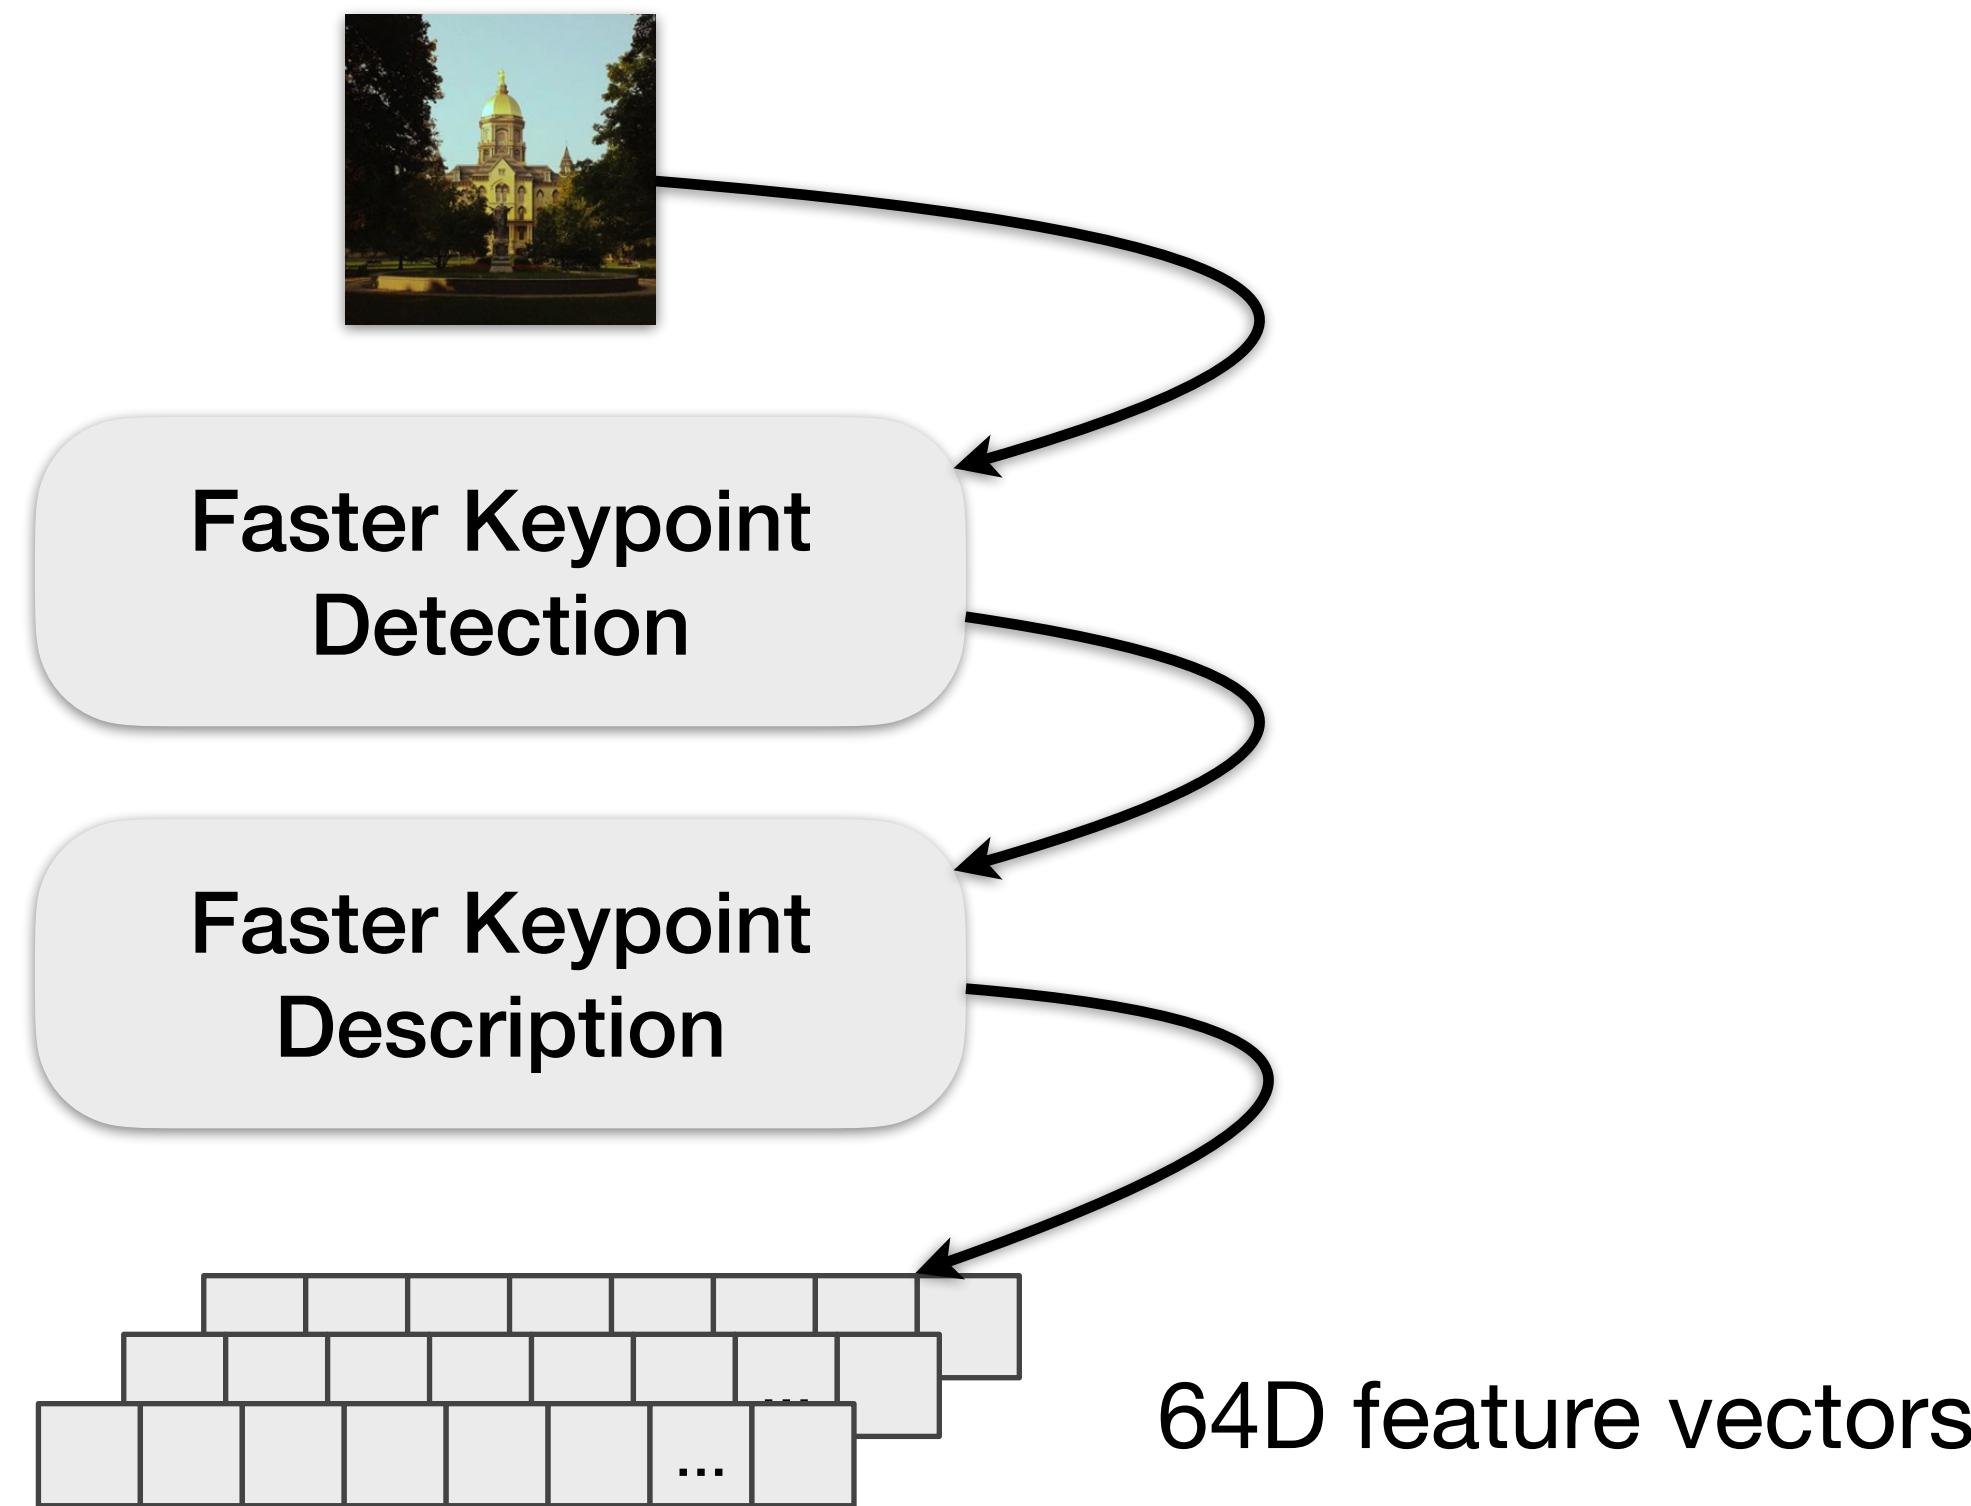
{Tinne.Tuytelaars, Luc.Vangool}@esat.kuleuven.be



LOYOLA
UNIVERSITY CHICAGO

Speeded-Up Robust Features (SURF)

Stages



SURF: Speeded Up Robust Features

Herbert Bay¹, Tinne Tuytelaars², and Luc Van Gool^{1,2}

¹ ETH Zurich

{bay, vangoool}@vision.ee.ethz.ch

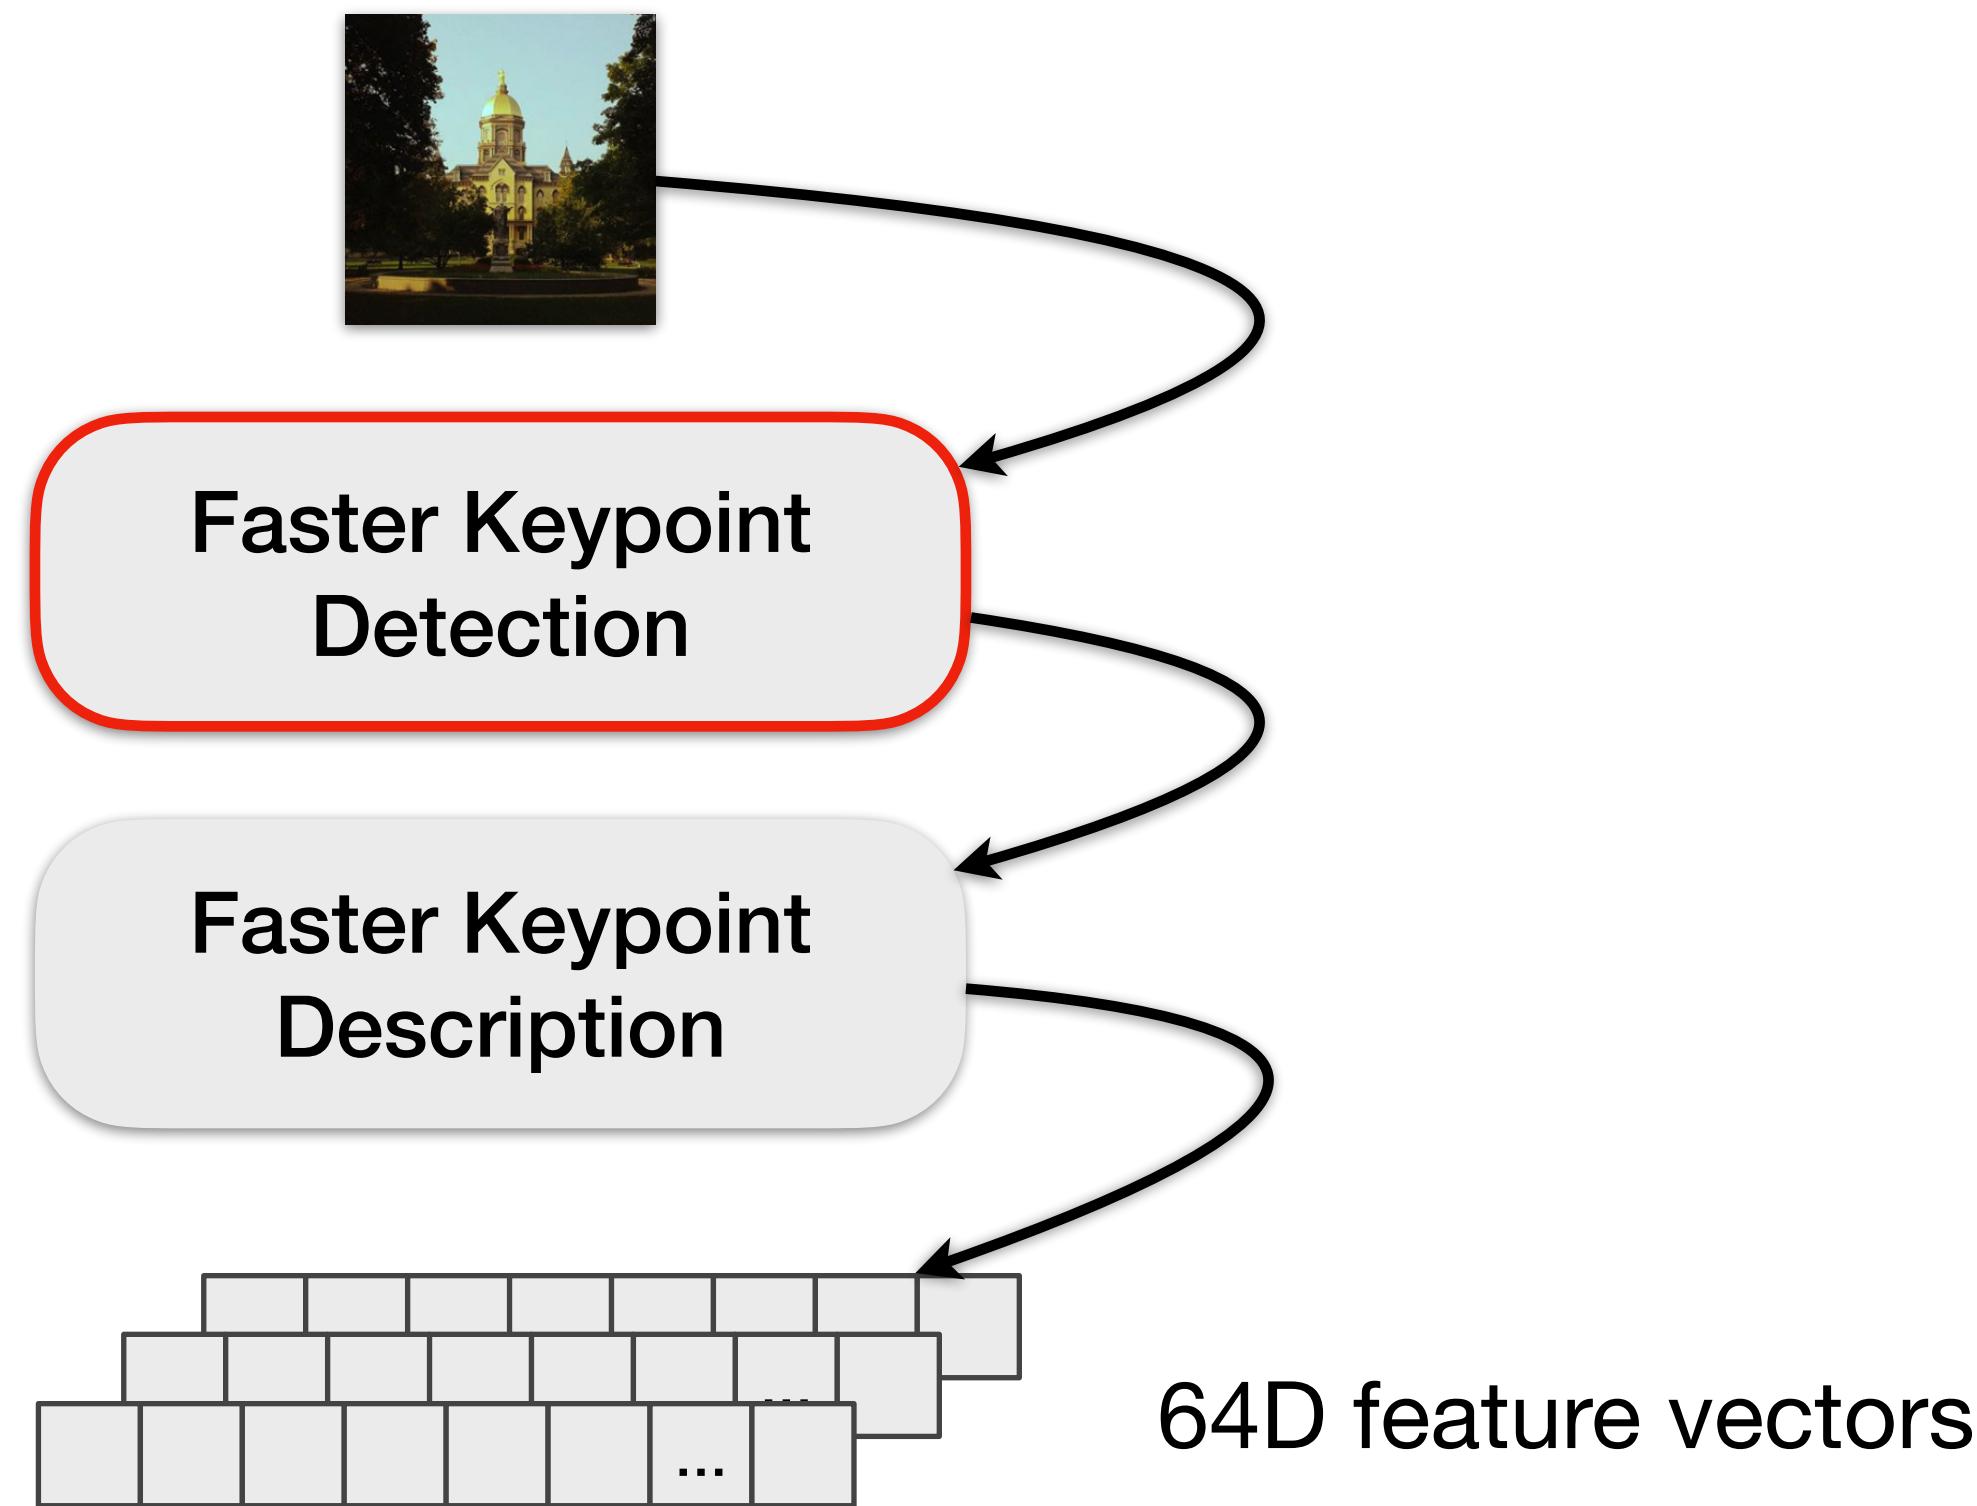
2006

² Katholieke Universiteit Leuven

{Tinne.Tuytelaars, Luc.Vangoool}@esat.kuleuven.be

Speeded-Up Robust Features (SURF)

Stages



SURF: Speeded Up Robust Features

Herbert Bay¹, Tinne Tuytelaars², and Luc Van Gool^{1,2}

¹ ETH Zurich

{bay, vangoool}@vision.ee.ethz.ch

² Katholieke Universiteit Leuven

{Tinne.Tuytelaars, Luc.Vangoool}@esat.kuleuven.be

2006



LOYOLA
UNIVERSITY CHICAGO

Speeded-Up Robust Features (SURF)

Faster Keypoint Detection

Hessian Matrix

Given an image pixel $I(x, y)$, a scale of interest σ ,

and Gaussian second order derivative functions $\frac{\delta^2}{\delta x^2}G(\sigma)$, $\frac{\delta^2}{\delta y^2}G(\sigma)$, and $\frac{\delta^2}{\delta xy}g(\sigma)$,

the Hessian matrix H is given by:

$$H(x, y, \sigma) = \begin{bmatrix} \frac{\delta^2}{\delta x^2}g(\sigma) * I(x, y) & \frac{\delta^2}{\delta xy}g(\sigma) * I(x, y) \\ \frac{\delta^2}{\delta xy}g(\sigma) * I(x, y) & \frac{\delta^2}{\delta y^2}g(\sigma) * I(x, y) \end{bmatrix}$$

Speeded-Up Robust Features (SURF)

Faster Keypoint Detection

Hessian Matrix

Given an image pixel $I(x, y)$, a scale of interest σ ,

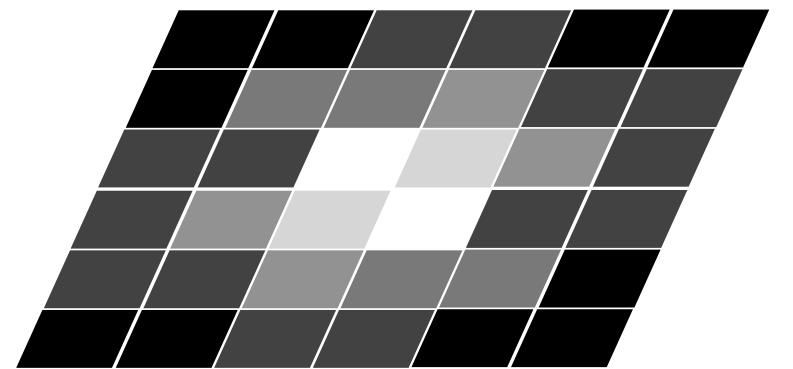
and Gaussian second order derivative functions $\frac{\delta^2}{\delta x^2}G(\sigma)$, $\frac{\delta^2}{\delta y^2}G(\sigma)$, and $\frac{\delta^2}{\delta xy}g(\sigma)$,

the Hessian matrix H is given by:

$$H(x, y, \sigma) = \begin{bmatrix} \frac{\delta^2}{\delta x^2}g(\sigma) * I(x, y) & \frac{\delta^2}{\delta xy}g(\sigma) * I(x, y) \\ \frac{\delta^2}{\delta xy}g(\sigma) * I(x, y) & \frac{\delta^2}{\delta y^2}g(\sigma) * I(x, y) \end{bmatrix}$$

Property: blobs with scale σ and centered at $I(x, y)$ will lead to a large $\det(H)$.

Take the regions with large $\det(H)$ as candidate keypoints.



LOYOLA
UNIVERSITY CHICAGO

Speeded-Up Robust Features (SURF)

Faster Keypoint Detection

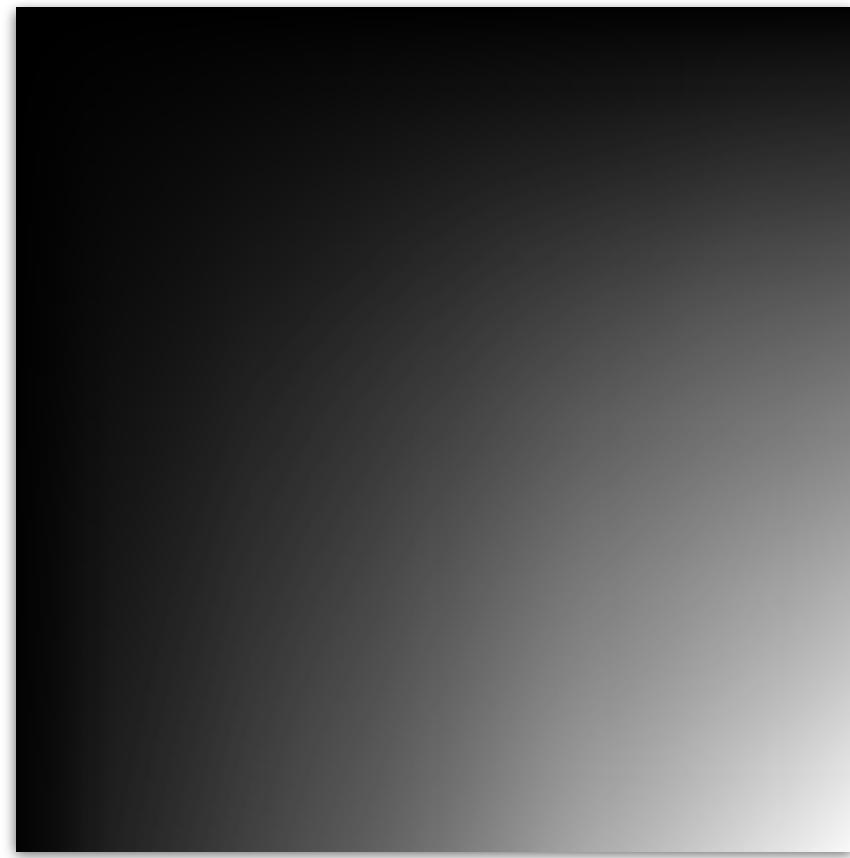
Integral Image

Data structure I_{Σ} computed from a given image I that shares the same resolution (i.e., same number of rows and of columns).

Each “pixel” of I_{Σ} has the following value:

$$I_{\Sigma}(x, y) = \sum_{i=0}^x \sum_{j=0}^y I(i, j)$$

i.e., it holds the sum of all the pixel values of I that spatially precede the position (x, y) .

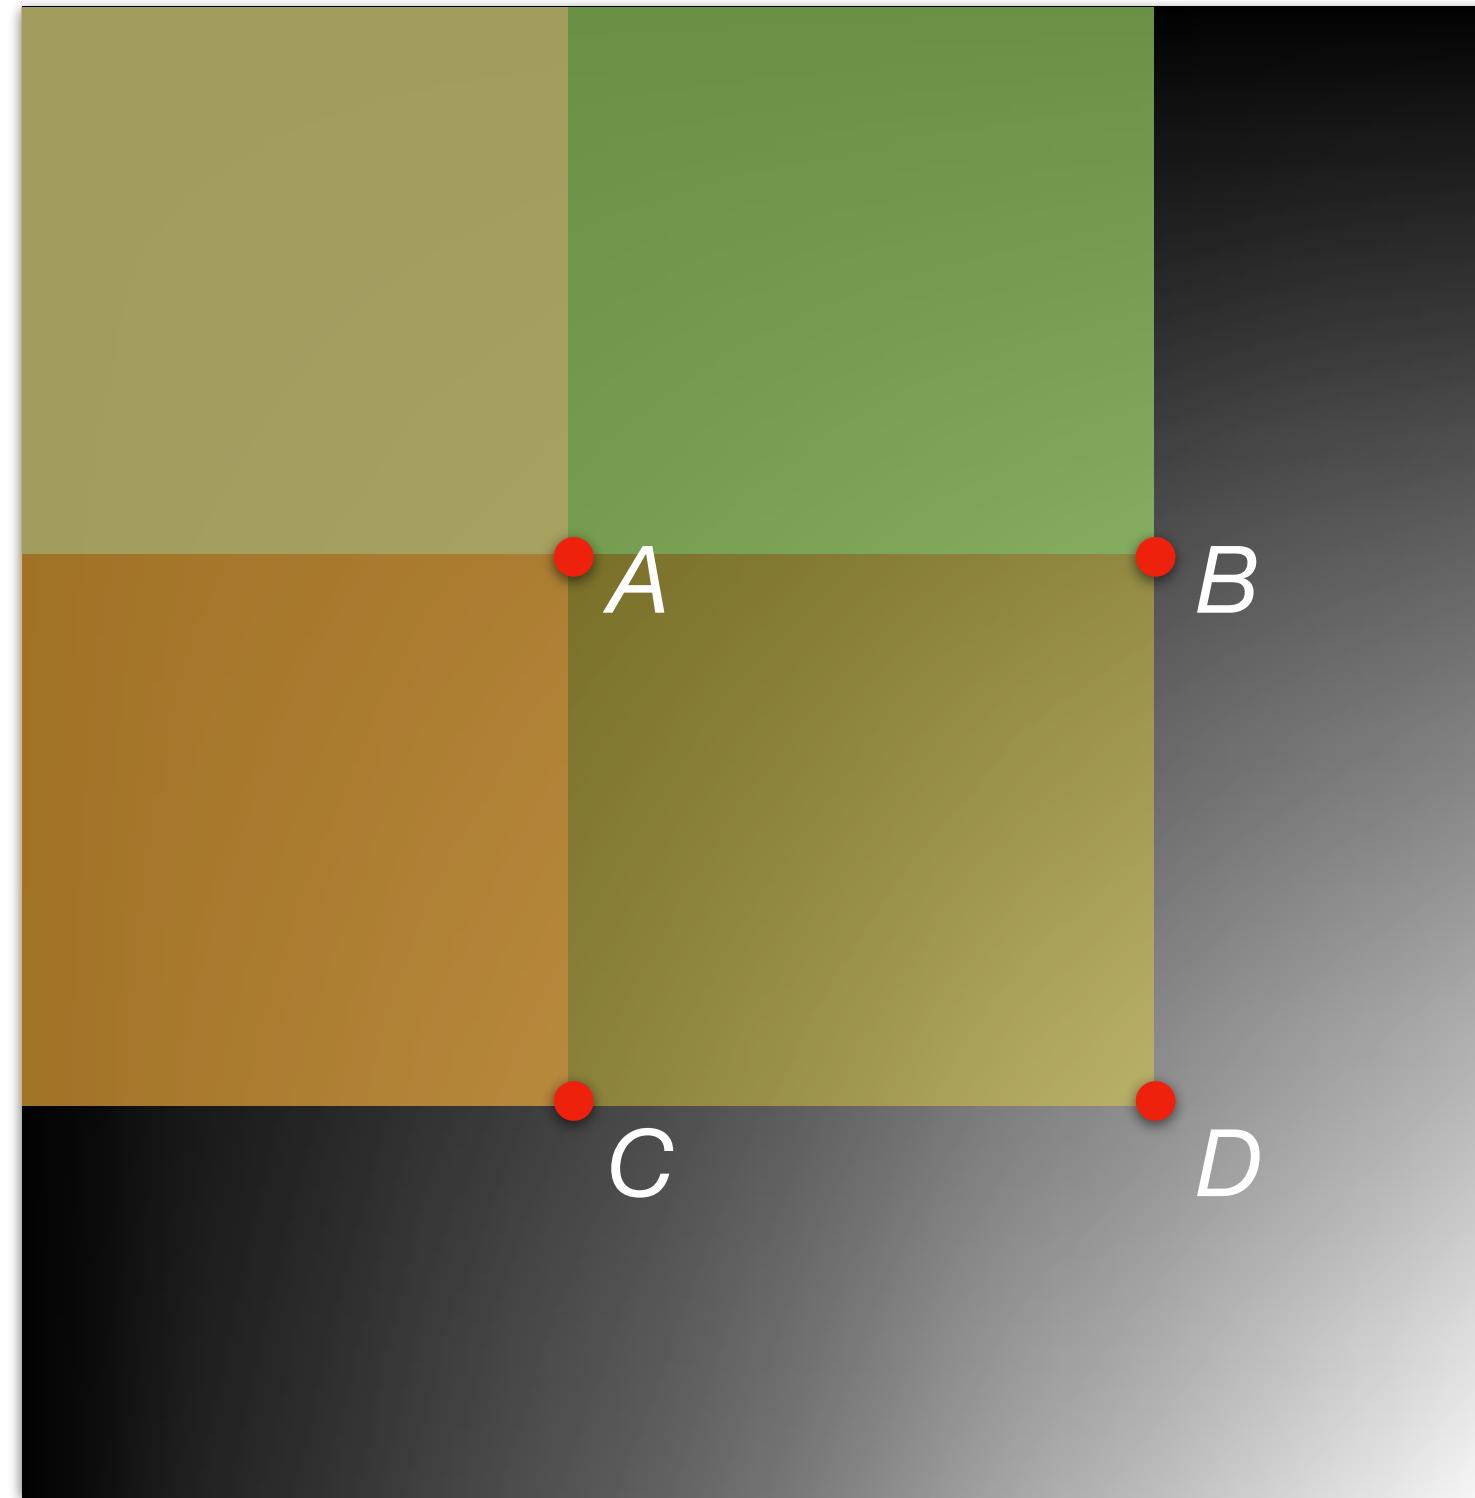


Speeded-Up Robust Features (SURF)

Faster Keypoint Detection $(0,0)$
Integral Image

What is the utility?

It is easy to compute
the sum of pixel values
within any region regardless
of the region size.



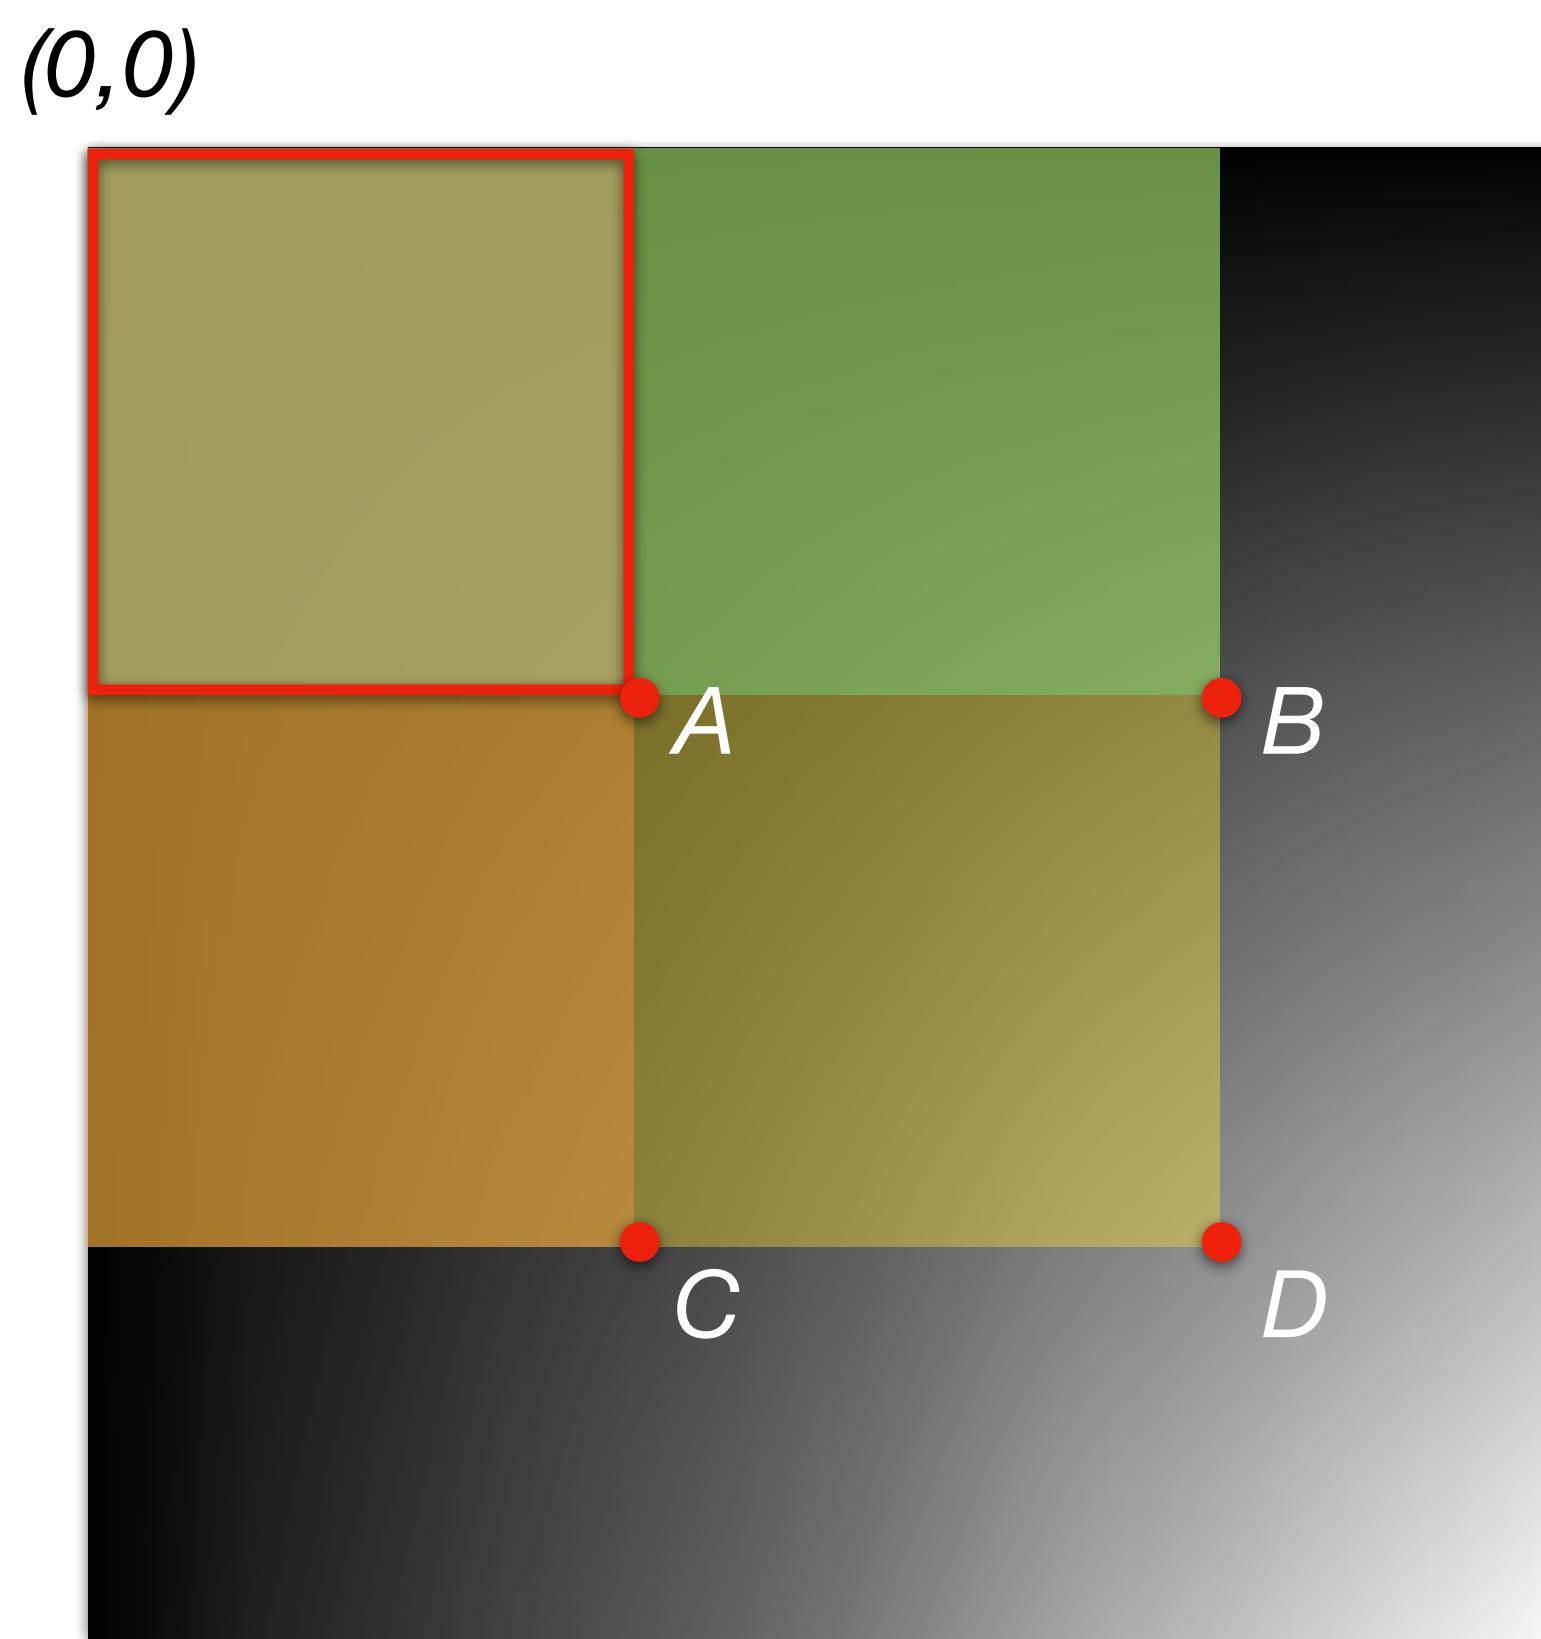
Speeded-Up Robust Features (SURF)

Faster Keypoint Detection

Integral Image

What is the utility?

It is easy to compute the sum of pixel values within any region regardless of the region size.



Sum of region between $(0,0)$ and A ?

Answer: $I_{\Sigma}(A)$

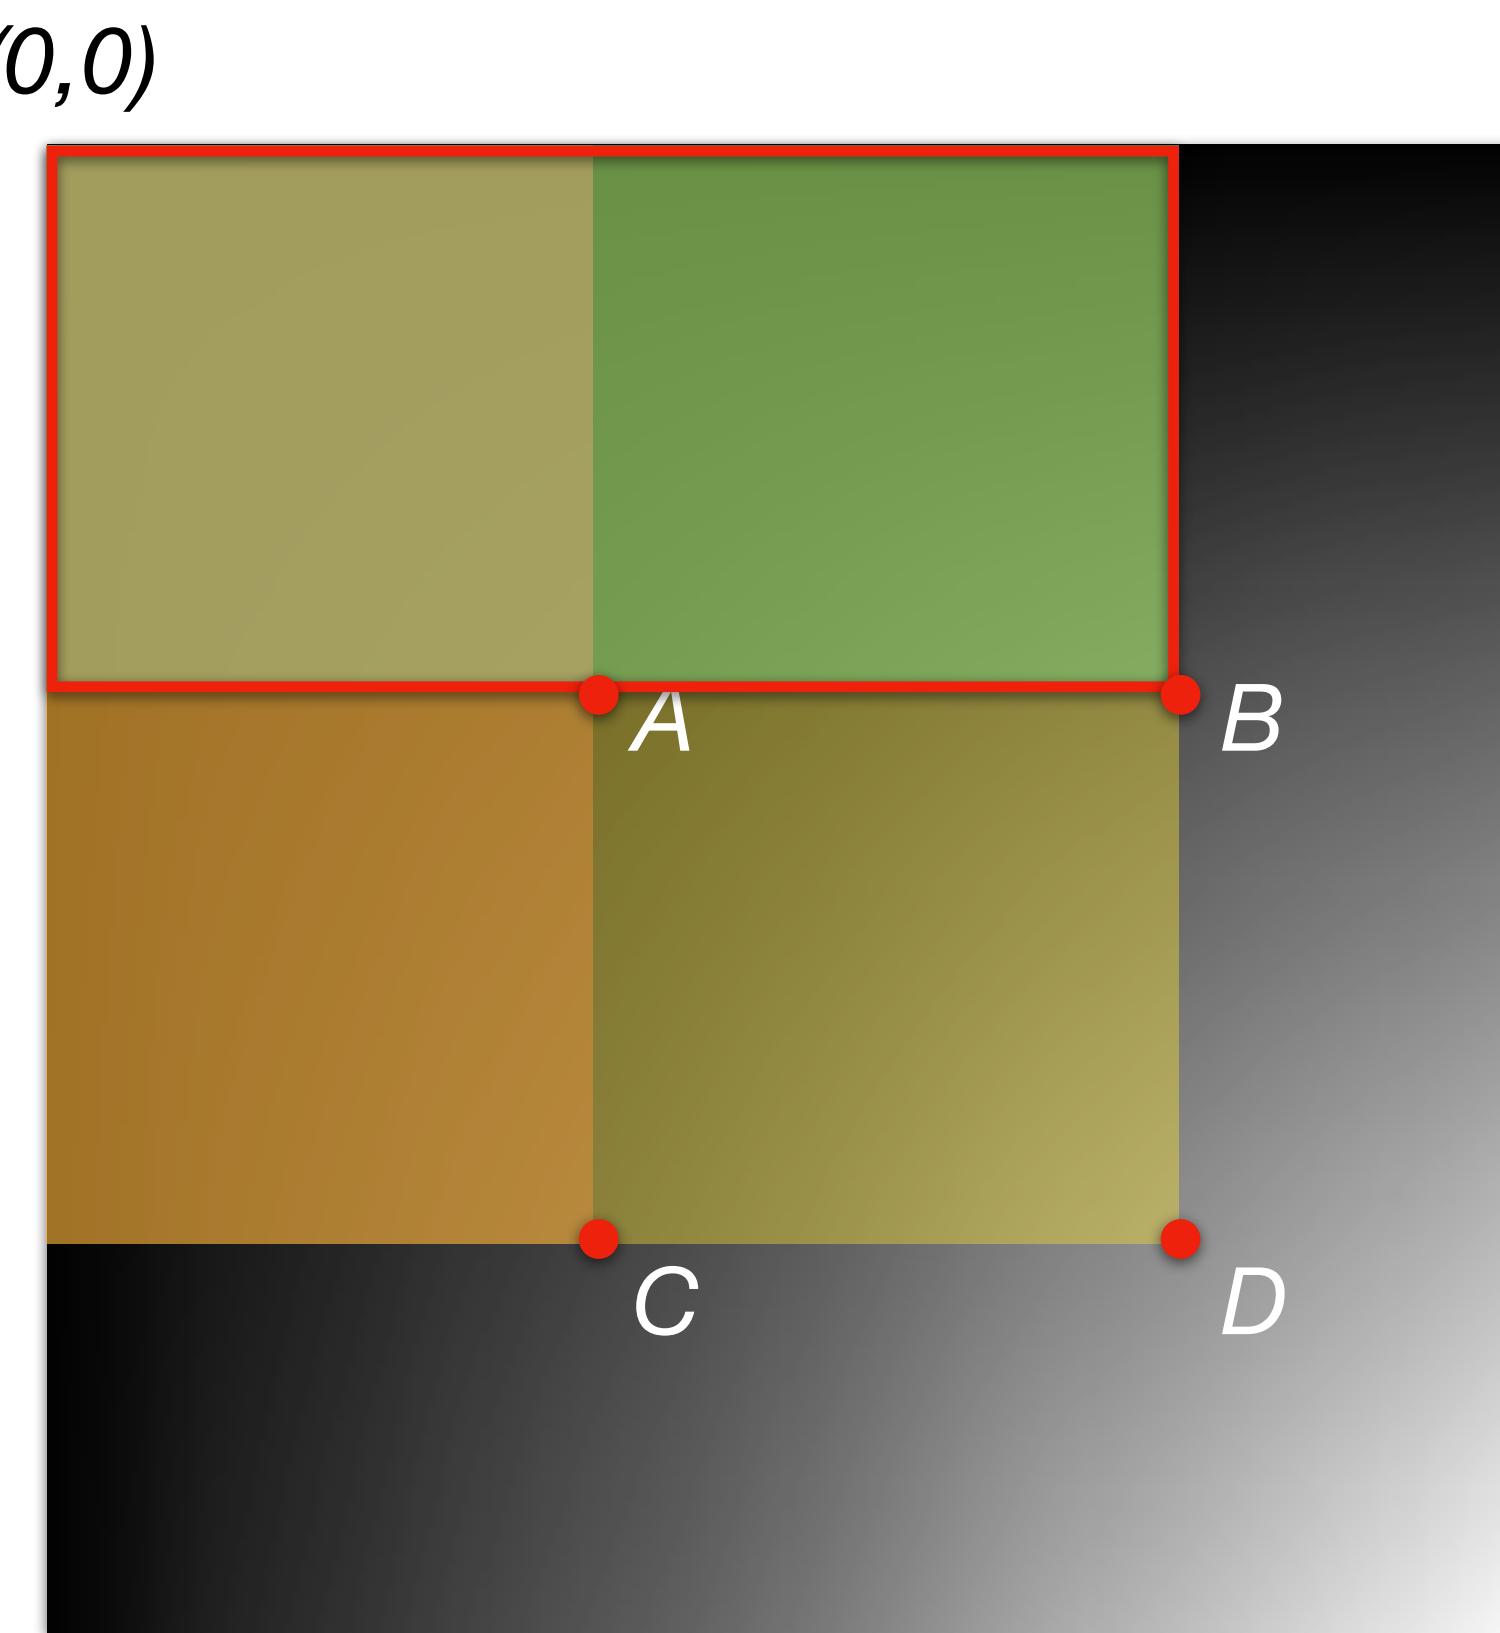
Speeded-Up Robust Features (SURF)

Faster Keypoint Detection

Integral Image

What is the utility?

It is easy to compute the sum of pixel values within any region regardless of the region size.



Sum of region between $(0,0)$ and A ?

Answer: $I_{\Sigma}(A)$

Sum of region between $(0,0)$ and B ?

Answer: $I_{\Sigma}(B)$

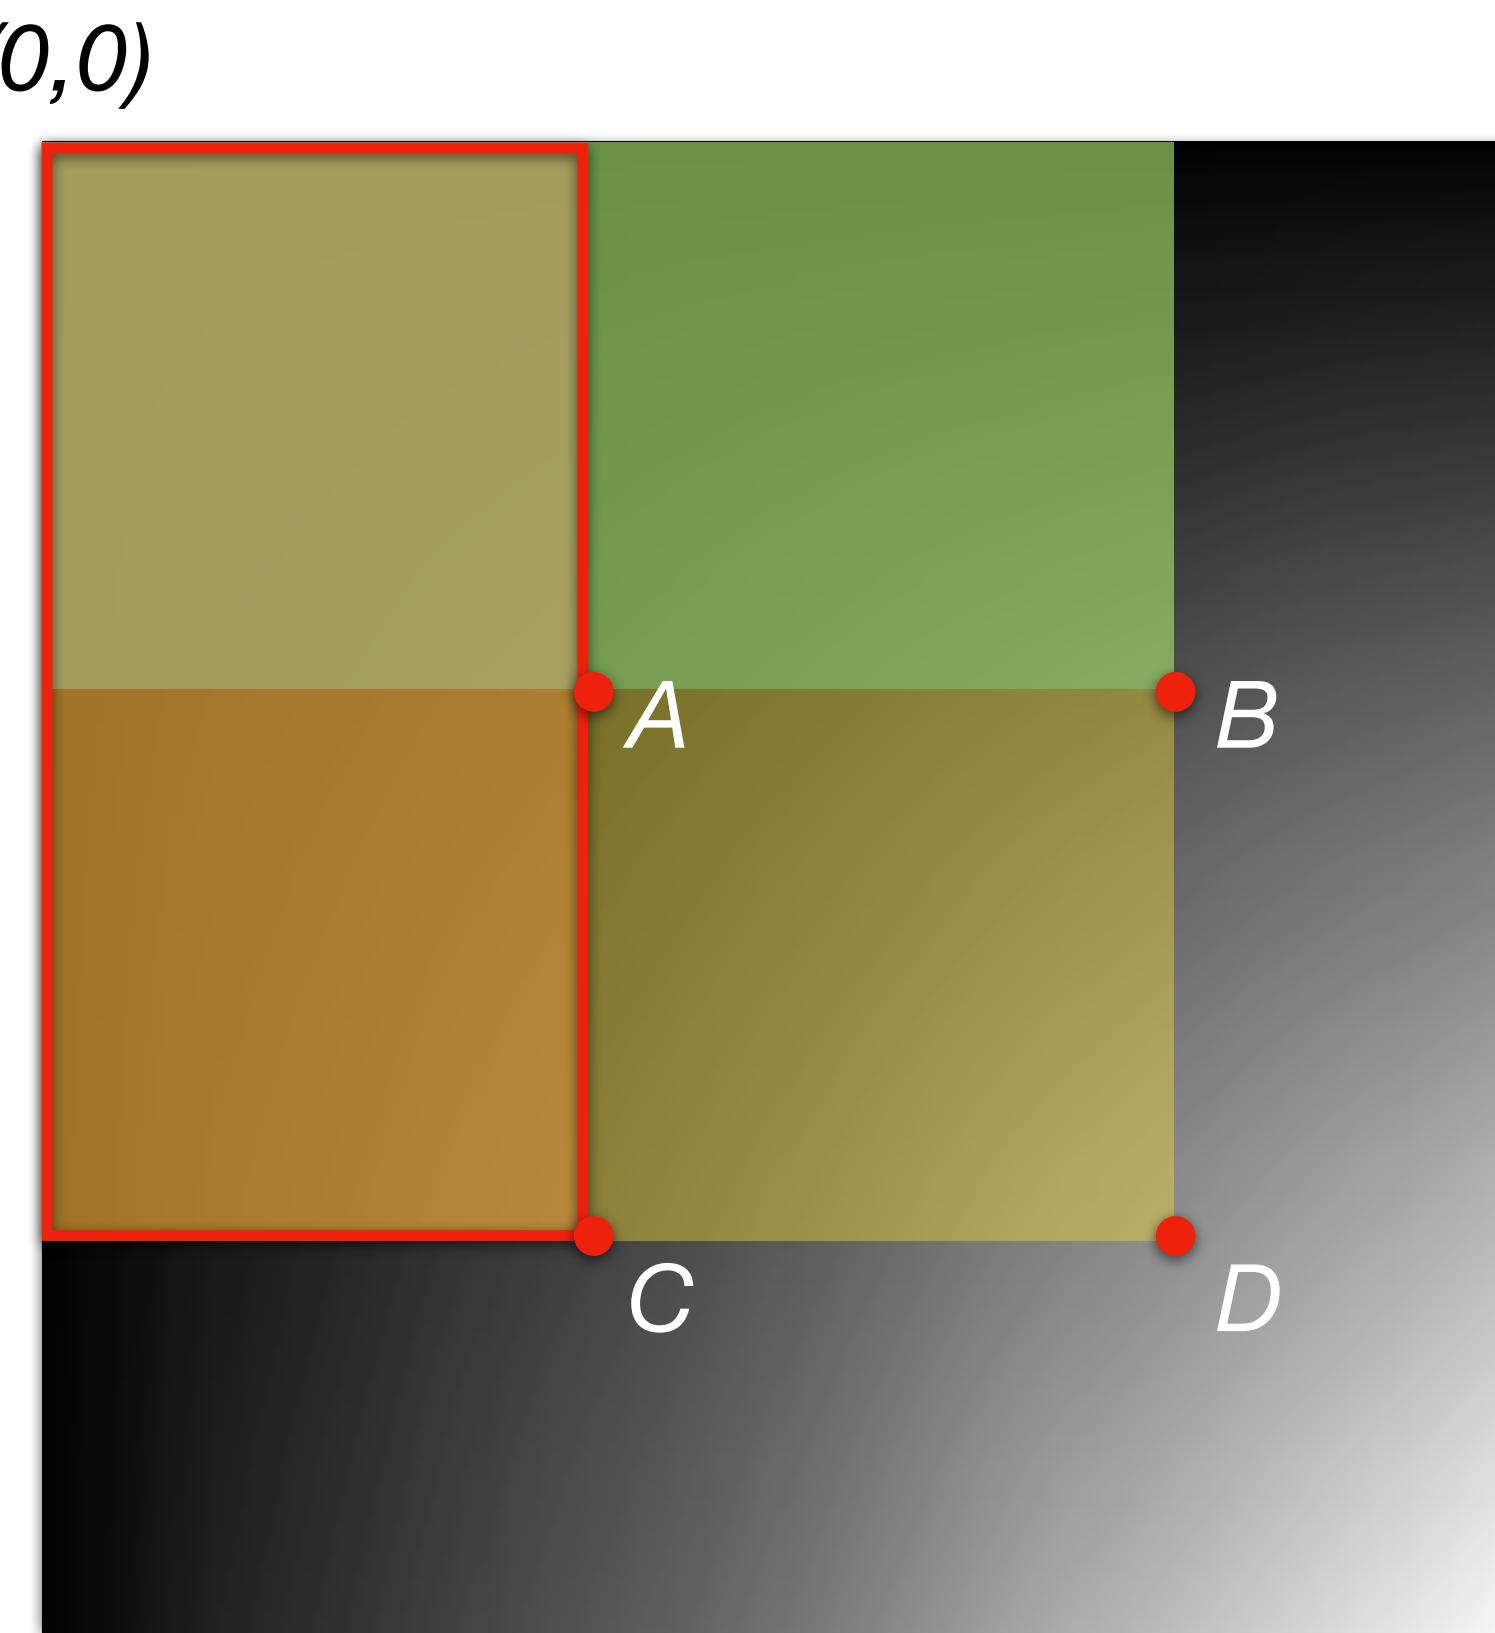
Speeded-Up Robust Features (SURF)

Faster Keypoint Detection

Integral Image

What is the utility?

It is easy to compute the sum of pixel values within any region regardless of the region size.



Sum of region between (0,0) and A?

Answer: $I_{\Sigma}(A)$

Sum of region between (0,0) and B?

Answer: $I_{\Sigma}(B)$

Sum of region between (0,0) and C?

Answer: $I_{\Sigma}(C)$

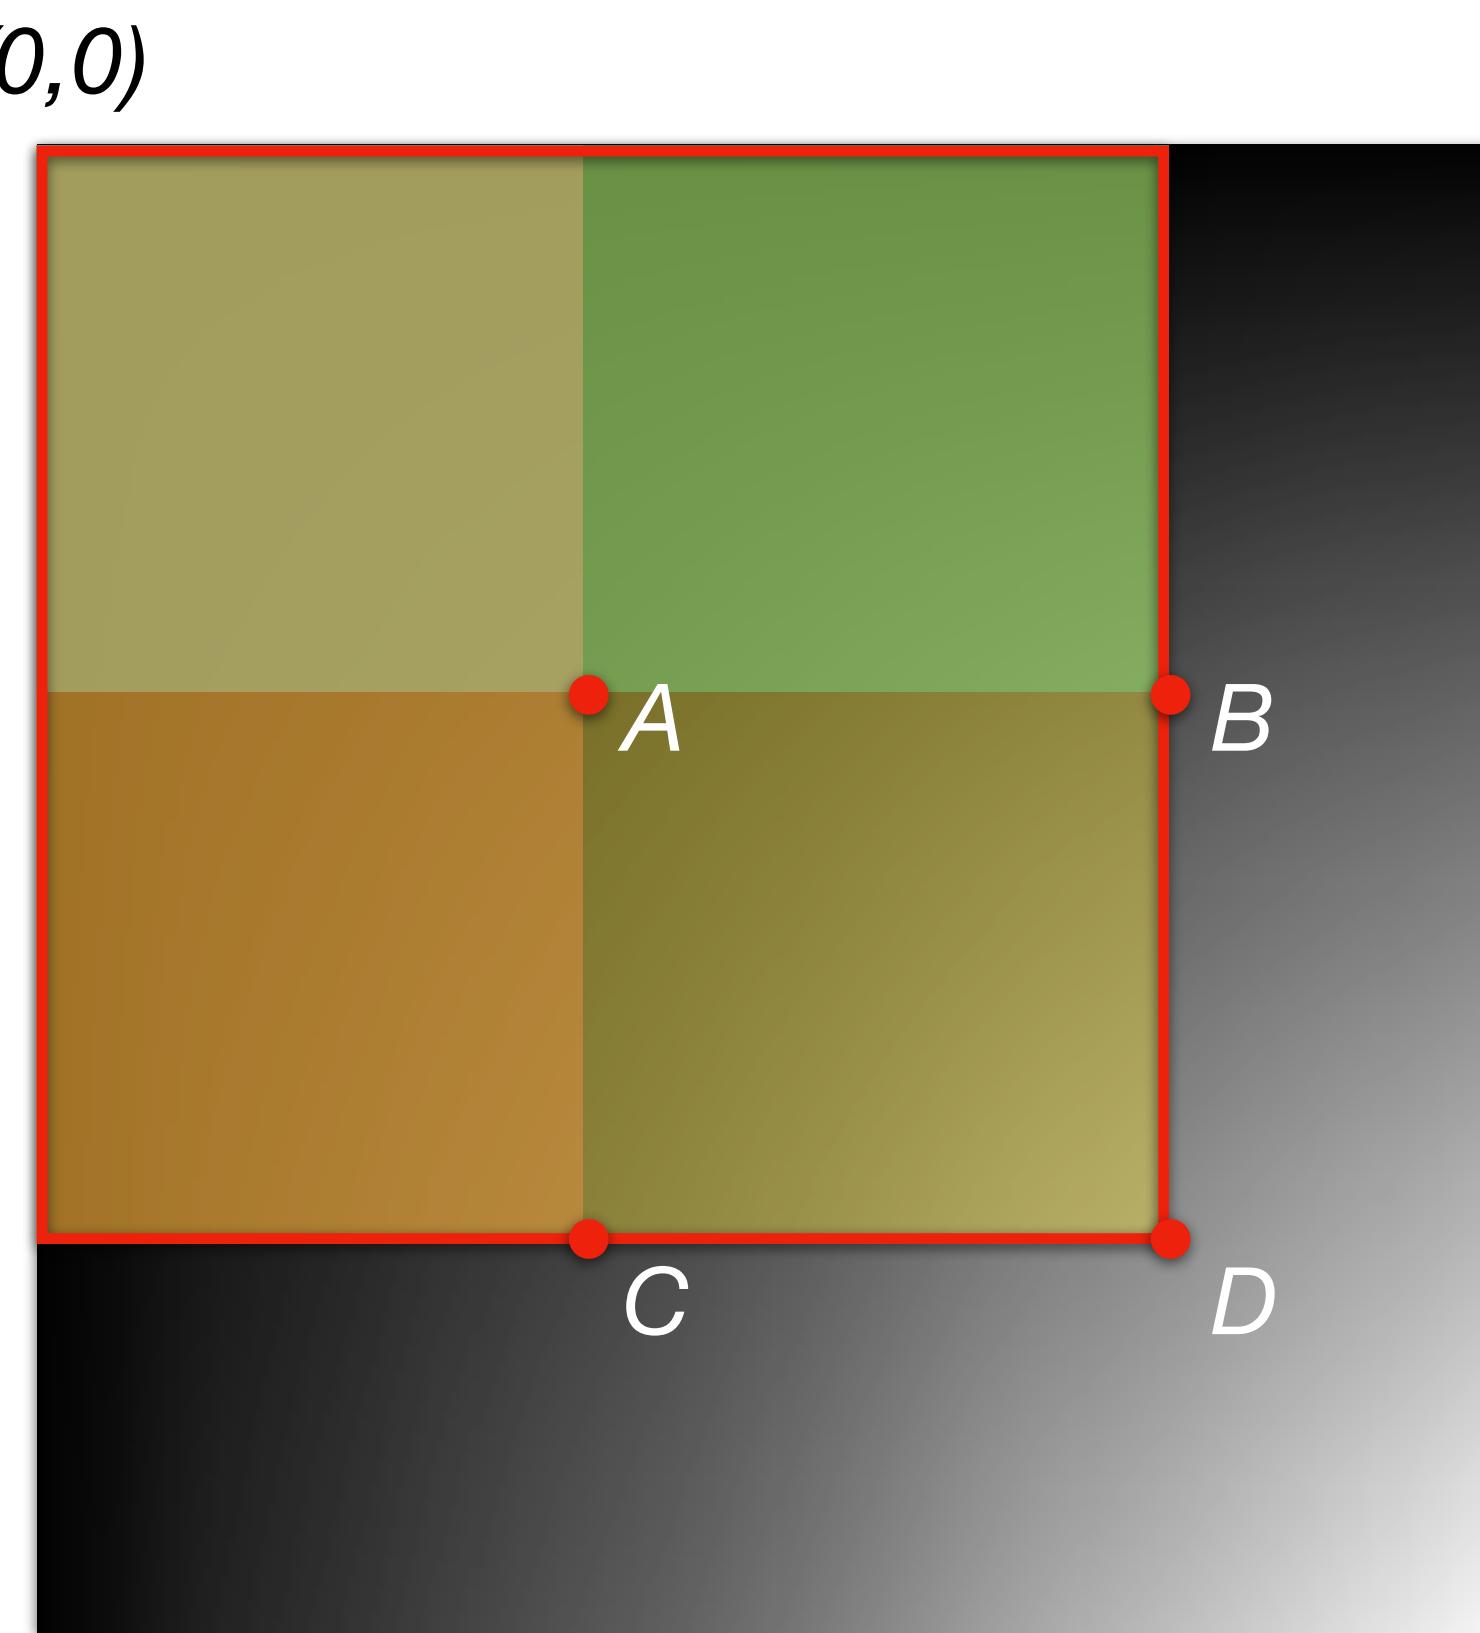
Speeded-Up Robust Features (SURF)

Faster Keypoint Detection

Integral Image

What is the utility?

It is easy to compute the sum of pixel values within any region regardless of the region size.



Sum of region between (0,0) and A?

Answer: $I_{\Sigma}(A)$

Sum of region between (0,0) and B?

Answer: $I_{\Sigma}(B)$

Sum of region between (0,0) and C?

Answer: $I_{\Sigma}(C)$

Sum of region between (0,0) and D?

Answer: $I_{\Sigma}(D)$

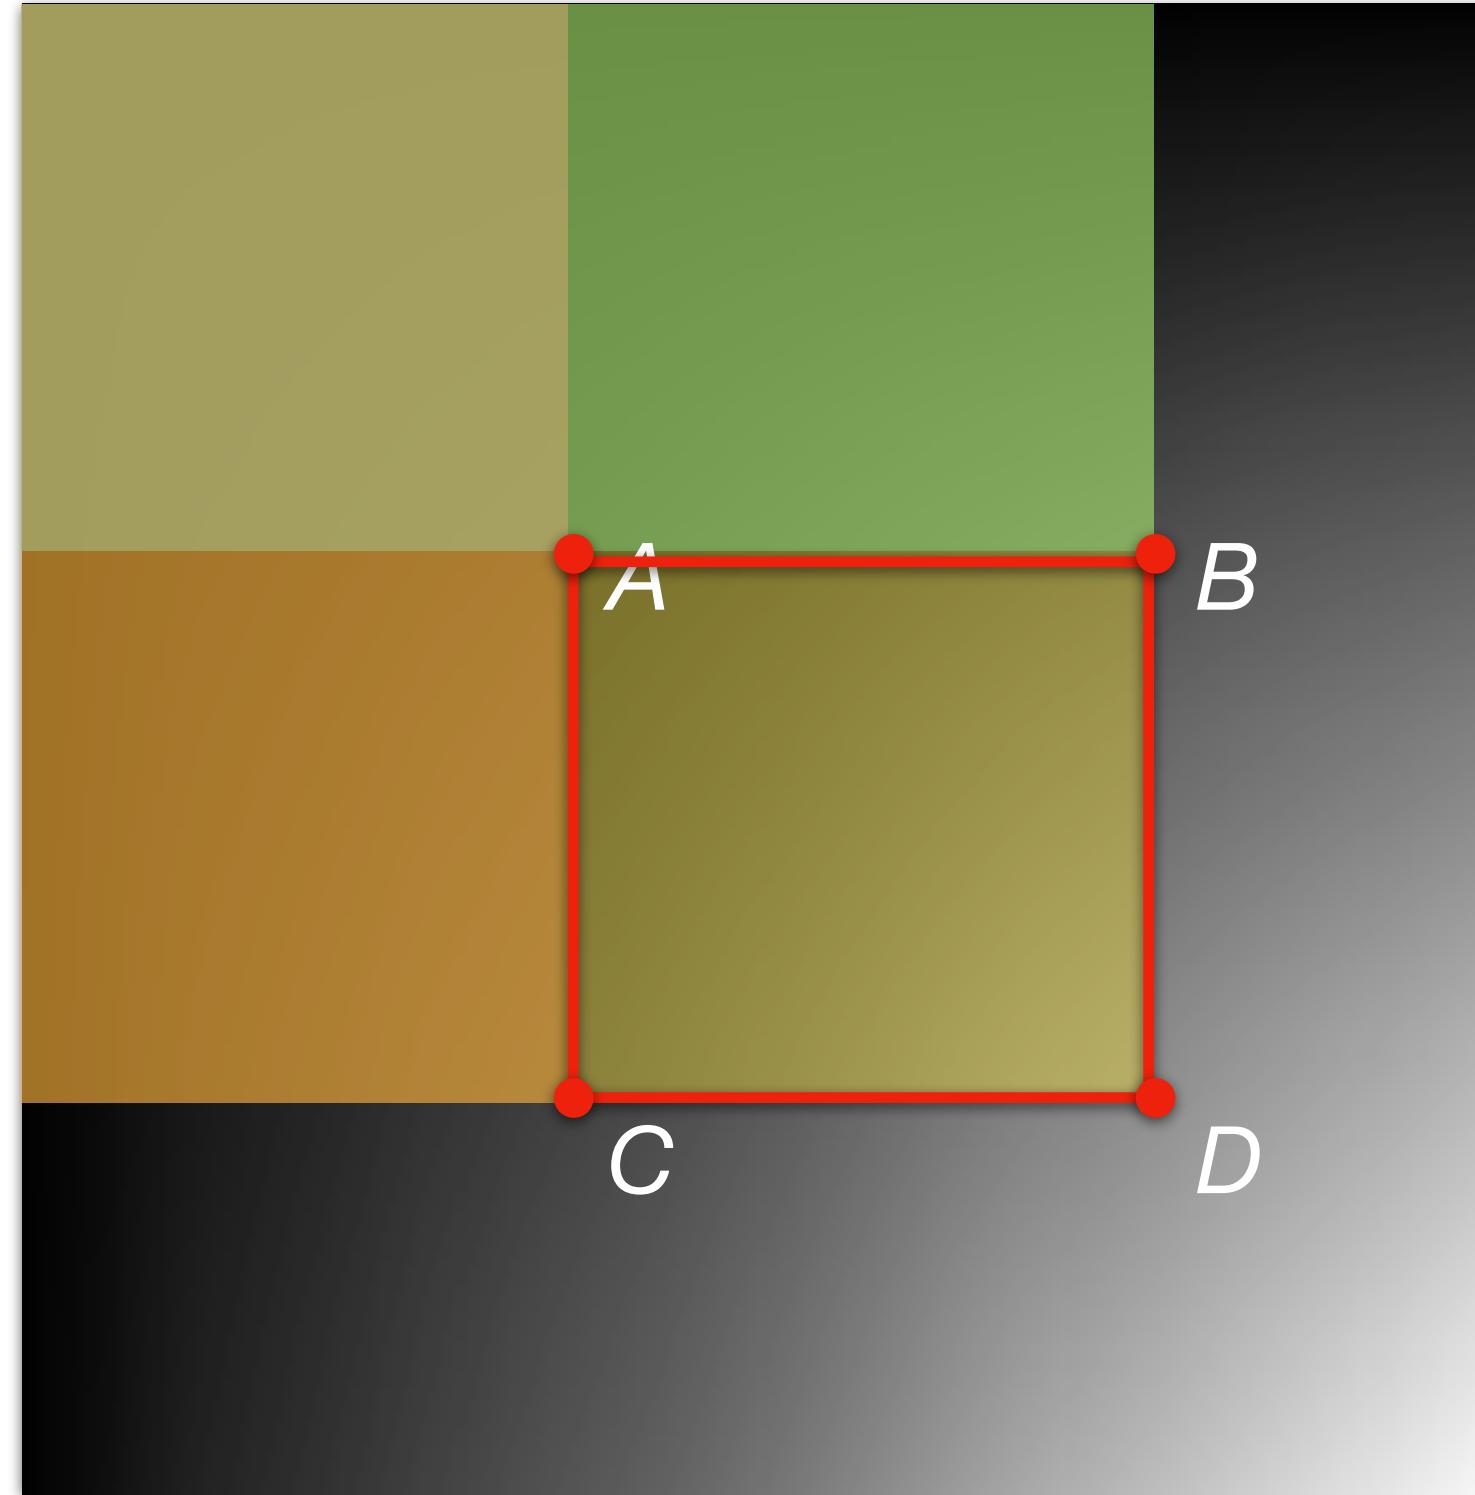
Speeded-Up Robust Features (SURF)

Faster Keypoint Detection $(0,0)$

Integral Image

What is the utility?

It is easy to compute
the sum of pixel values
within any region regardless
of the region size.



Sum of region between A and D ?

Answer:

$$I_{\Sigma}(D) - I_{\Sigma}(B) - I_{\Sigma}(C) + I_{\Sigma}(A)$$

One can get any sum with at most
4 accesses, regardless of the
resolution.

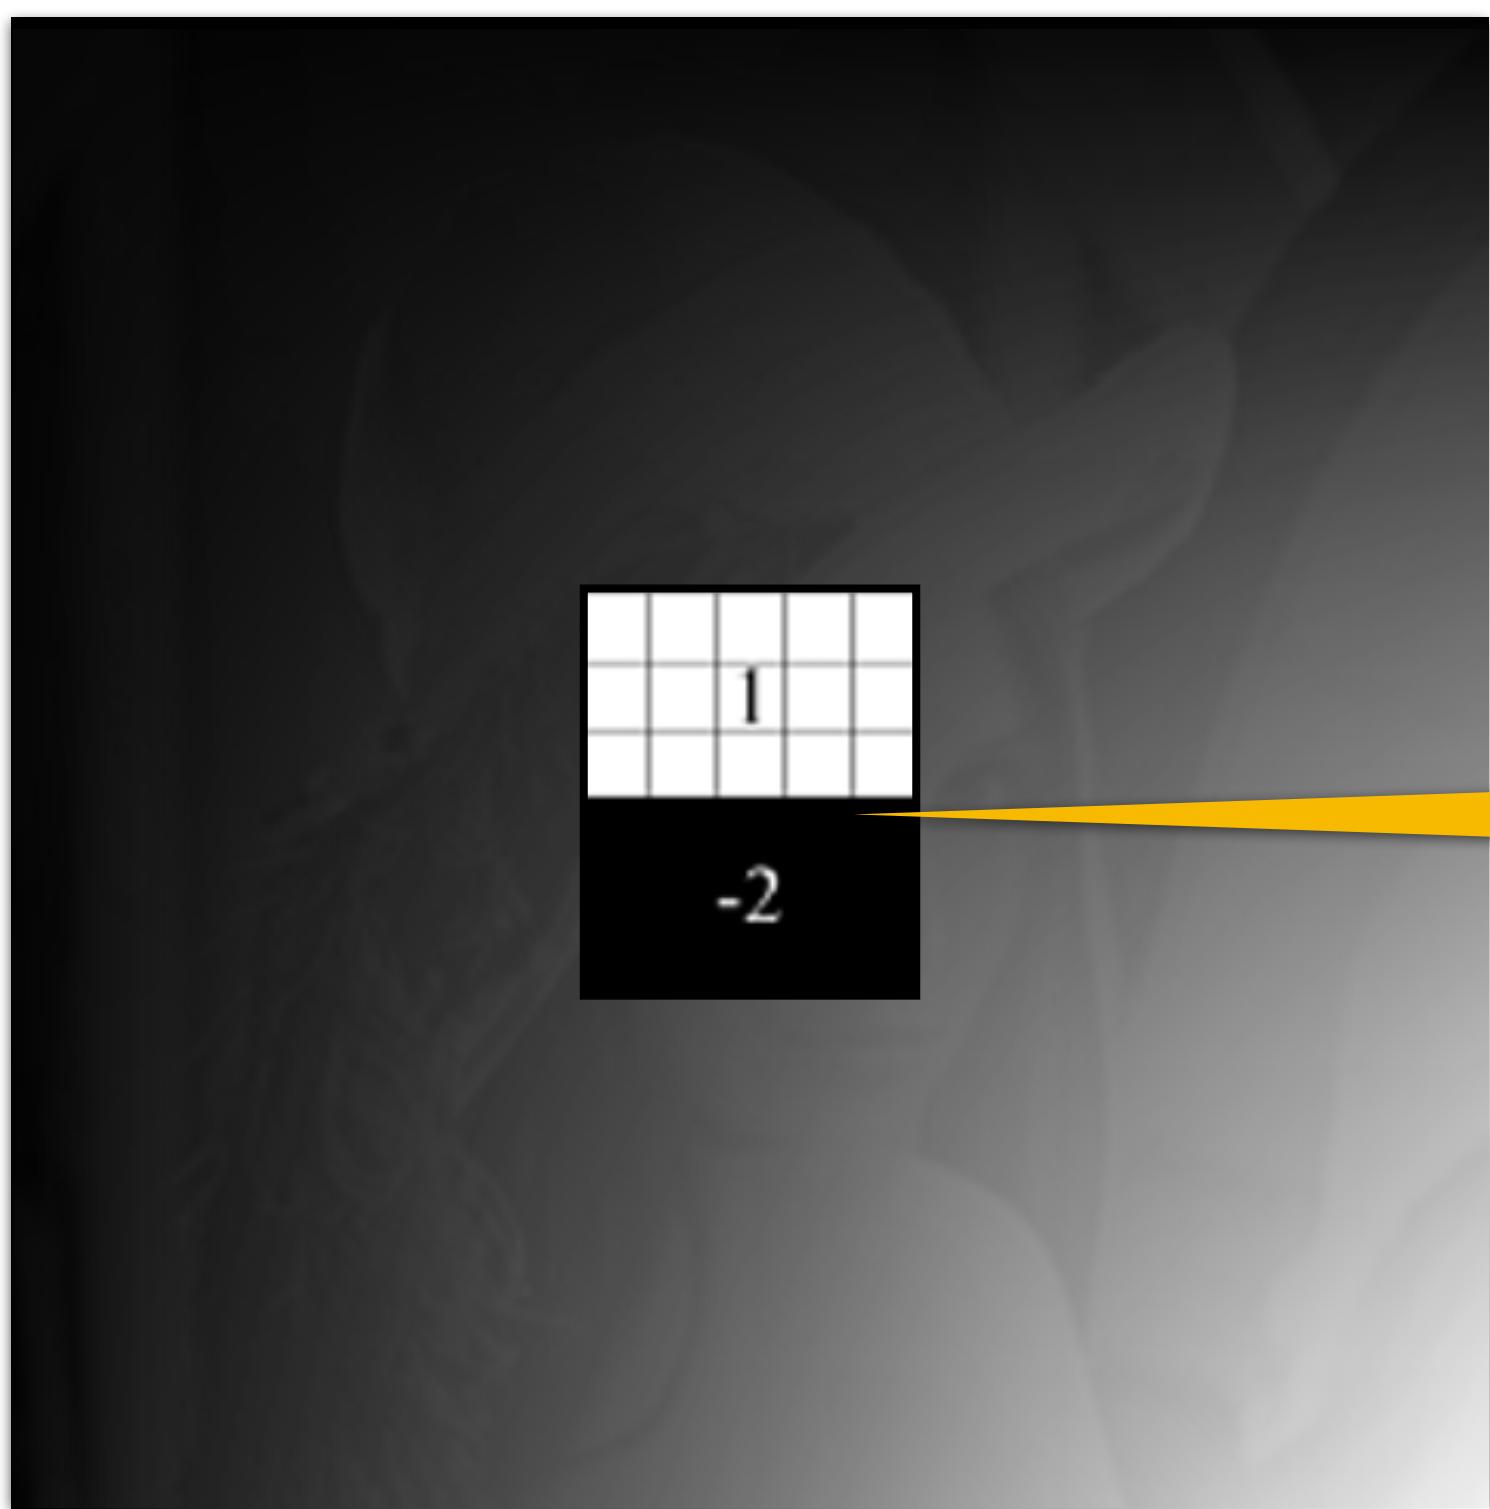
Speeded-Up Robust Features (SURF)

Faster Keypoint Detection $(0,0)$

Integral Image

What is the utility?

It is easy to compute
the sum of pixel values
within any region regardless
of the region size.



Box Filters can be easily convoluted with the integral image.

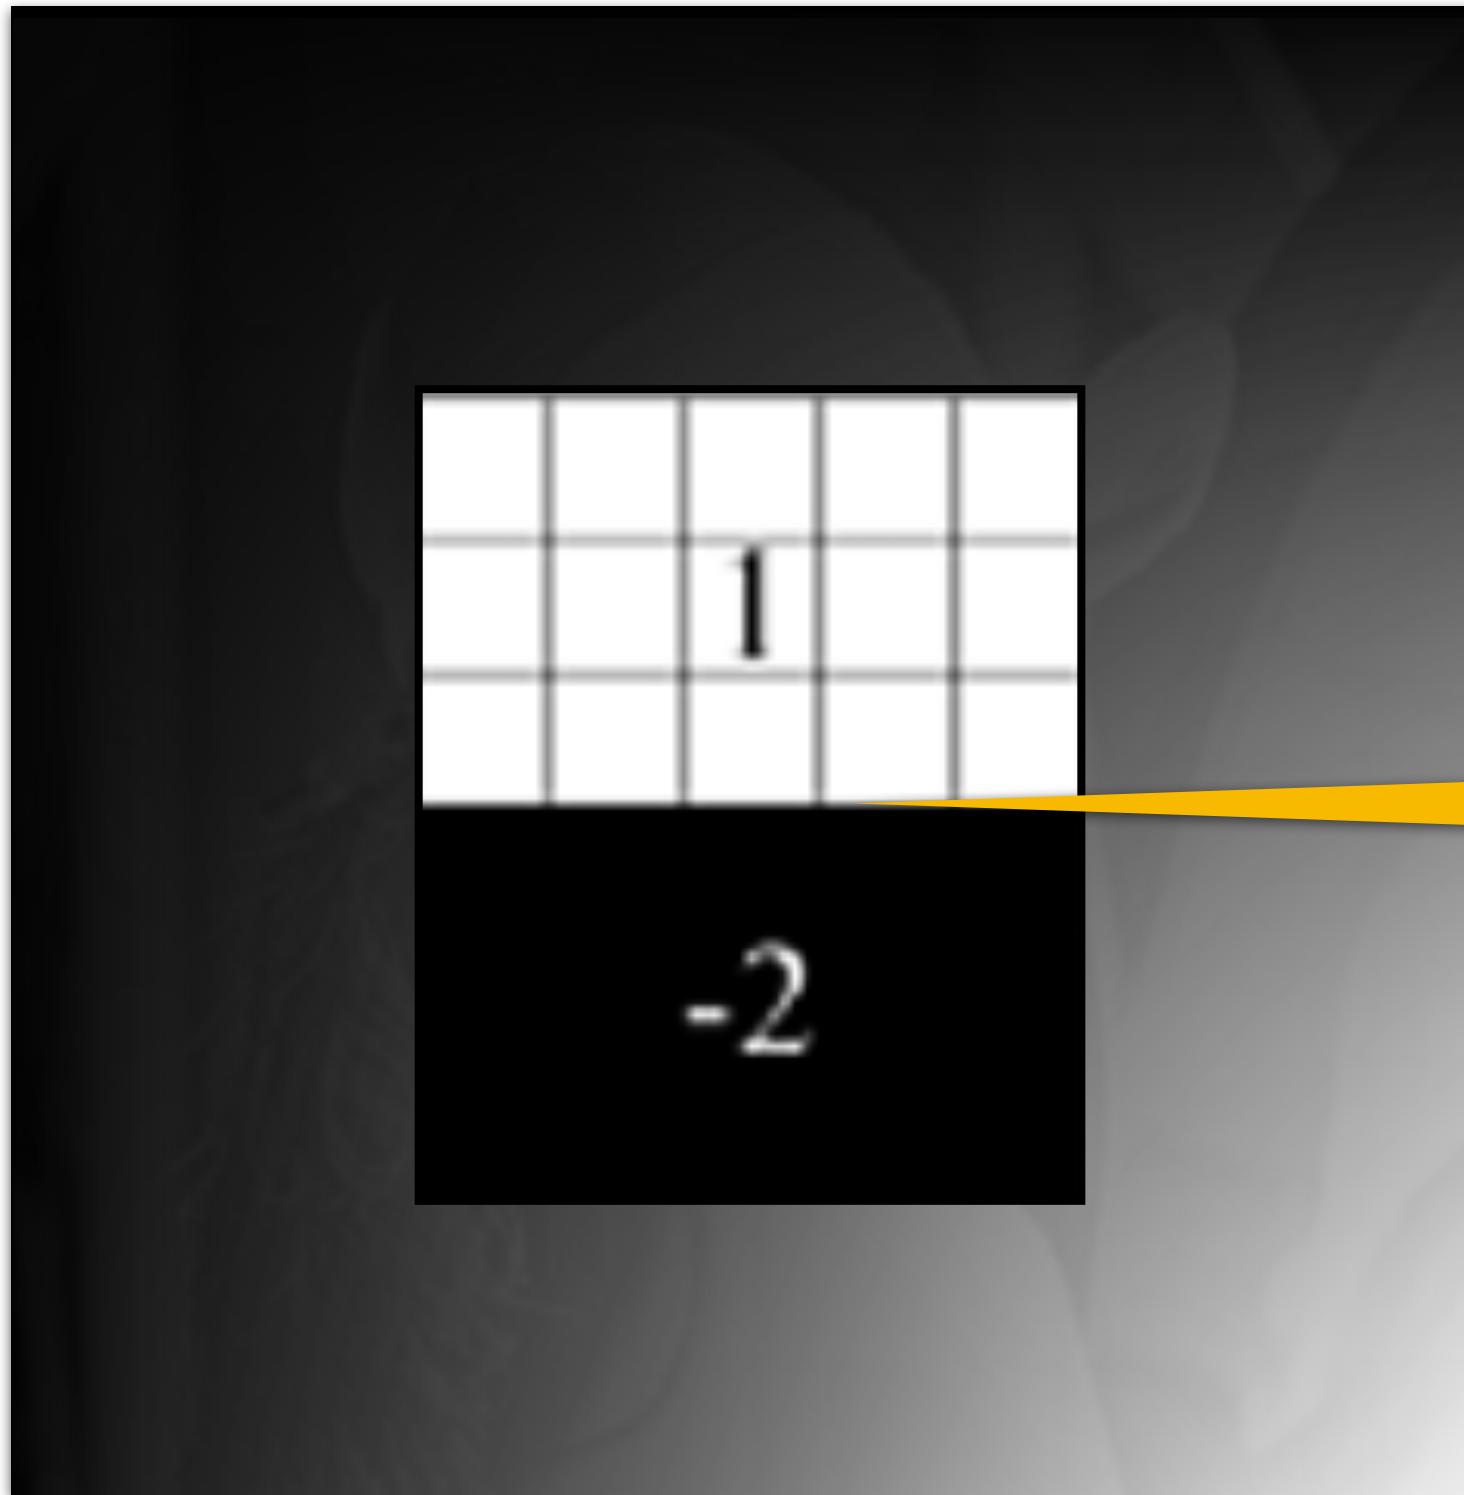
One can quickly test a horizontal edge here (8 accesses)

Speeded-Up Robust Features (SURF)

Faster Keypoint Detection $(0,0)$
Integral Image

What is the utility?

It is easy to compute
the sum of pixel values
within any region regardless
of the region size.



Box Filters can be easily
convolved with the integral
image.

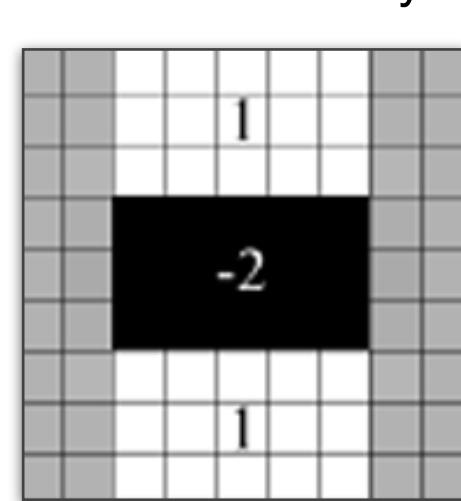
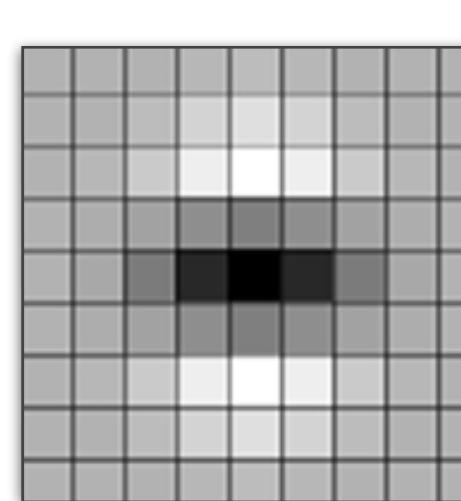
One can quickly test a larger
scale horizontal edge here
(still 8 accesses)

Speeded-Up Robust Features (SURF)

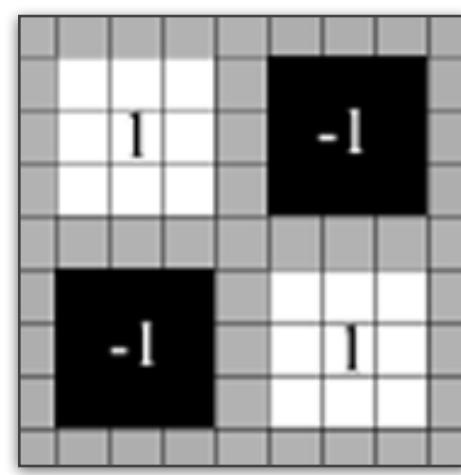
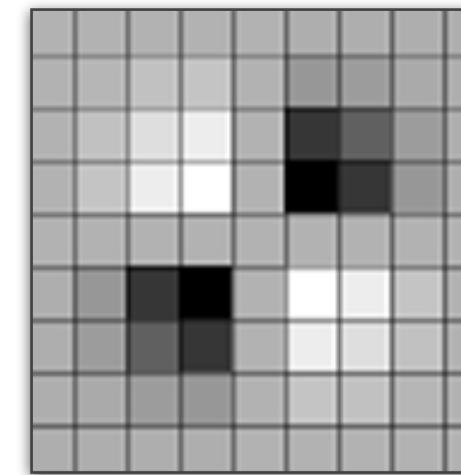
Faster Keypoint Detection

Box Filters

The Gaussian second order derivative functions $\frac{\delta^2}{\delta x^2}G(\sigma)$, $\frac{\delta^2}{\delta y^2}G(\sigma)$, and $\frac{\delta^2}{\delta xy}g(\sigma)$ can be approximated by box filters.



$$\frac{\delta^2}{\delta y^2}G(\sigma)$$



$$\frac{\delta^2}{\delta xy}g(\sigma)$$

Speeded-Up Robust Features (SURF)

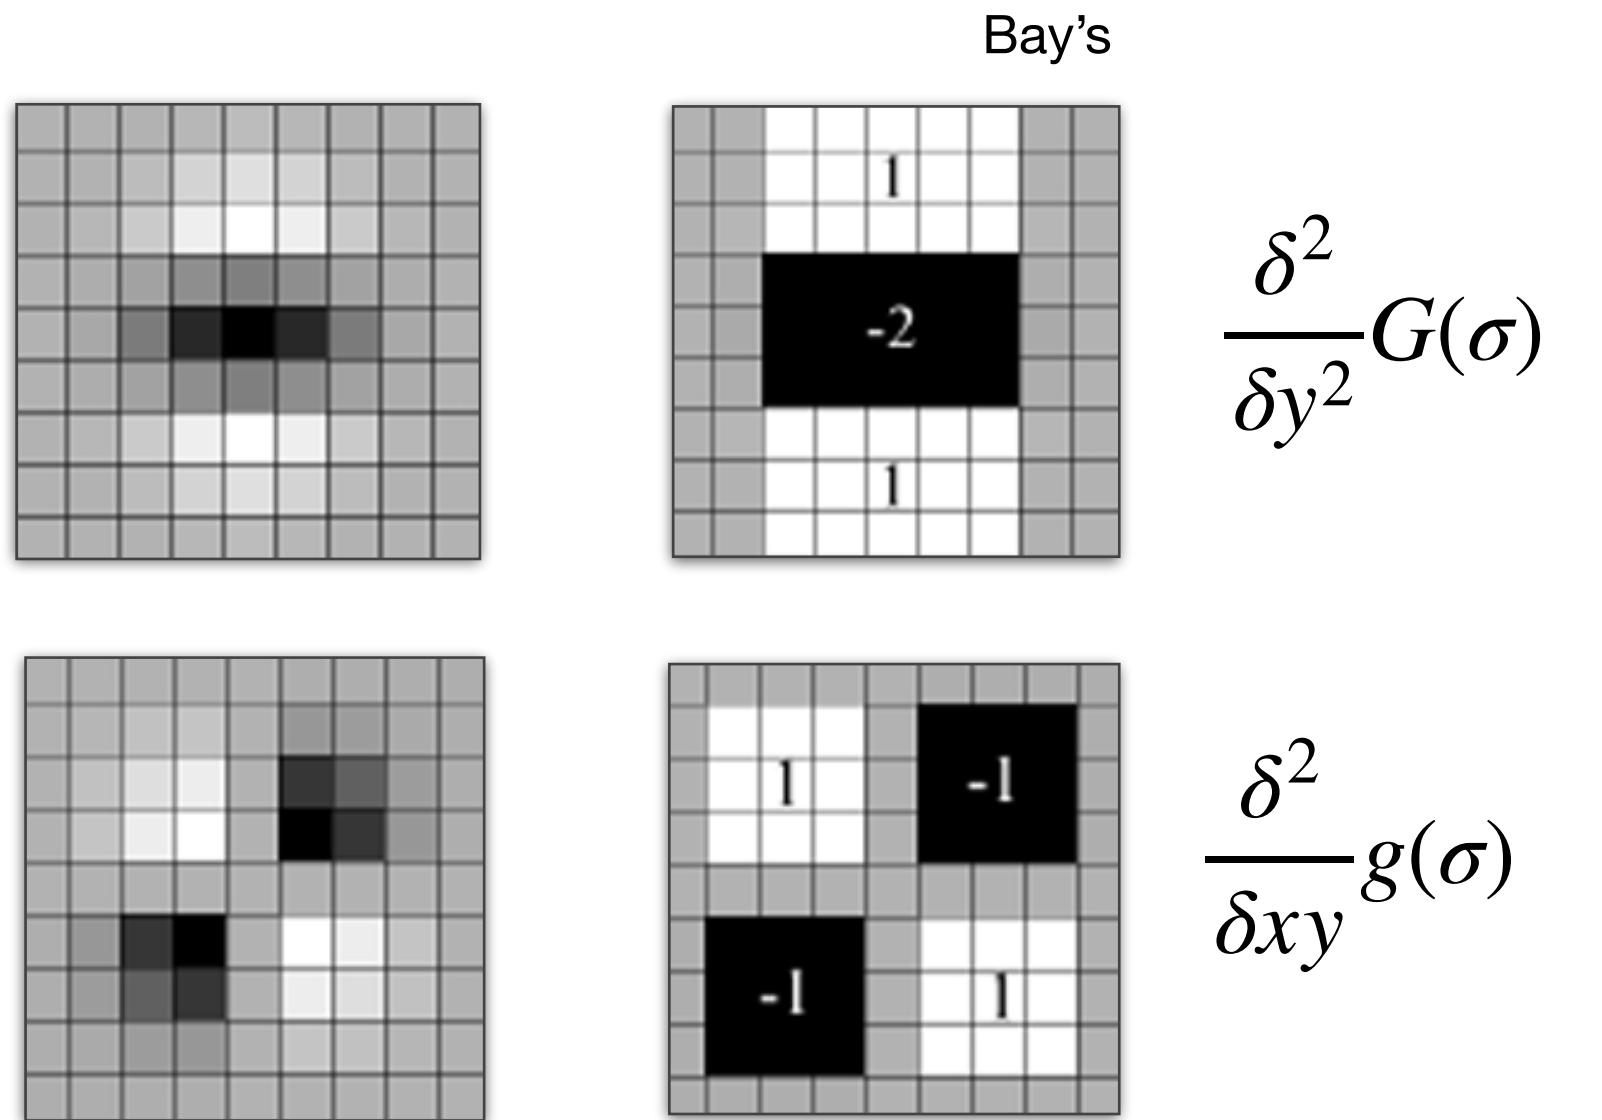
Faster Keypoint Detection

Box Filters

The Gaussian second order derivative functions $\frac{\delta^2}{\delta x^2}G(\sigma)$, $\frac{\delta^2}{\delta y^2}G(\sigma)$, and $\frac{\delta^2}{\delta xy}g(\sigma)$ can be approximated by box filters.

Compute the $\det(H)$ quickly by using the box filters and the integral image!

$$H(x, y, \sigma) = \begin{bmatrix} \frac{\delta^2}{\delta x^2}g(\sigma) * I(x, y) & \frac{\delta^2}{\delta xy}g(\sigma) * I(x, y) \\ \frac{\delta^2}{\delta xy}g(\sigma) * I(x, y) & \frac{\delta^2}{\delta y^2}g(\sigma) * I(x, y) \end{bmatrix}$$



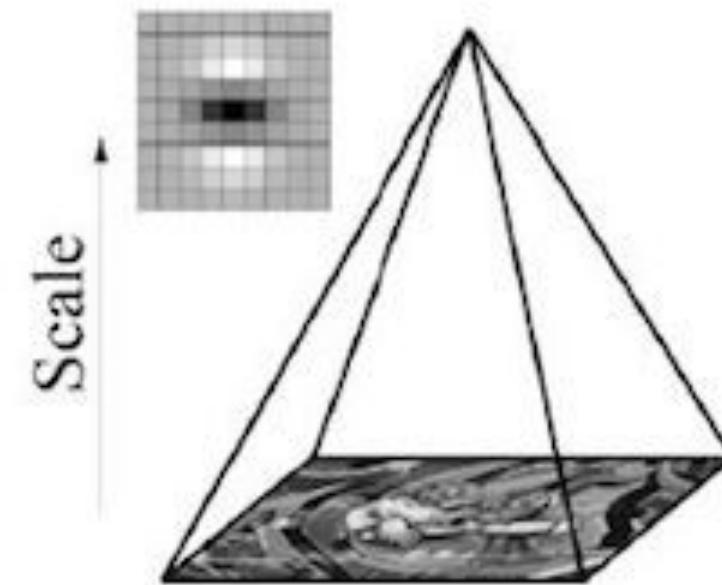
Speeded-Up Robust Features (SURF)

Faster Keypoint Detection

Scale Invariance

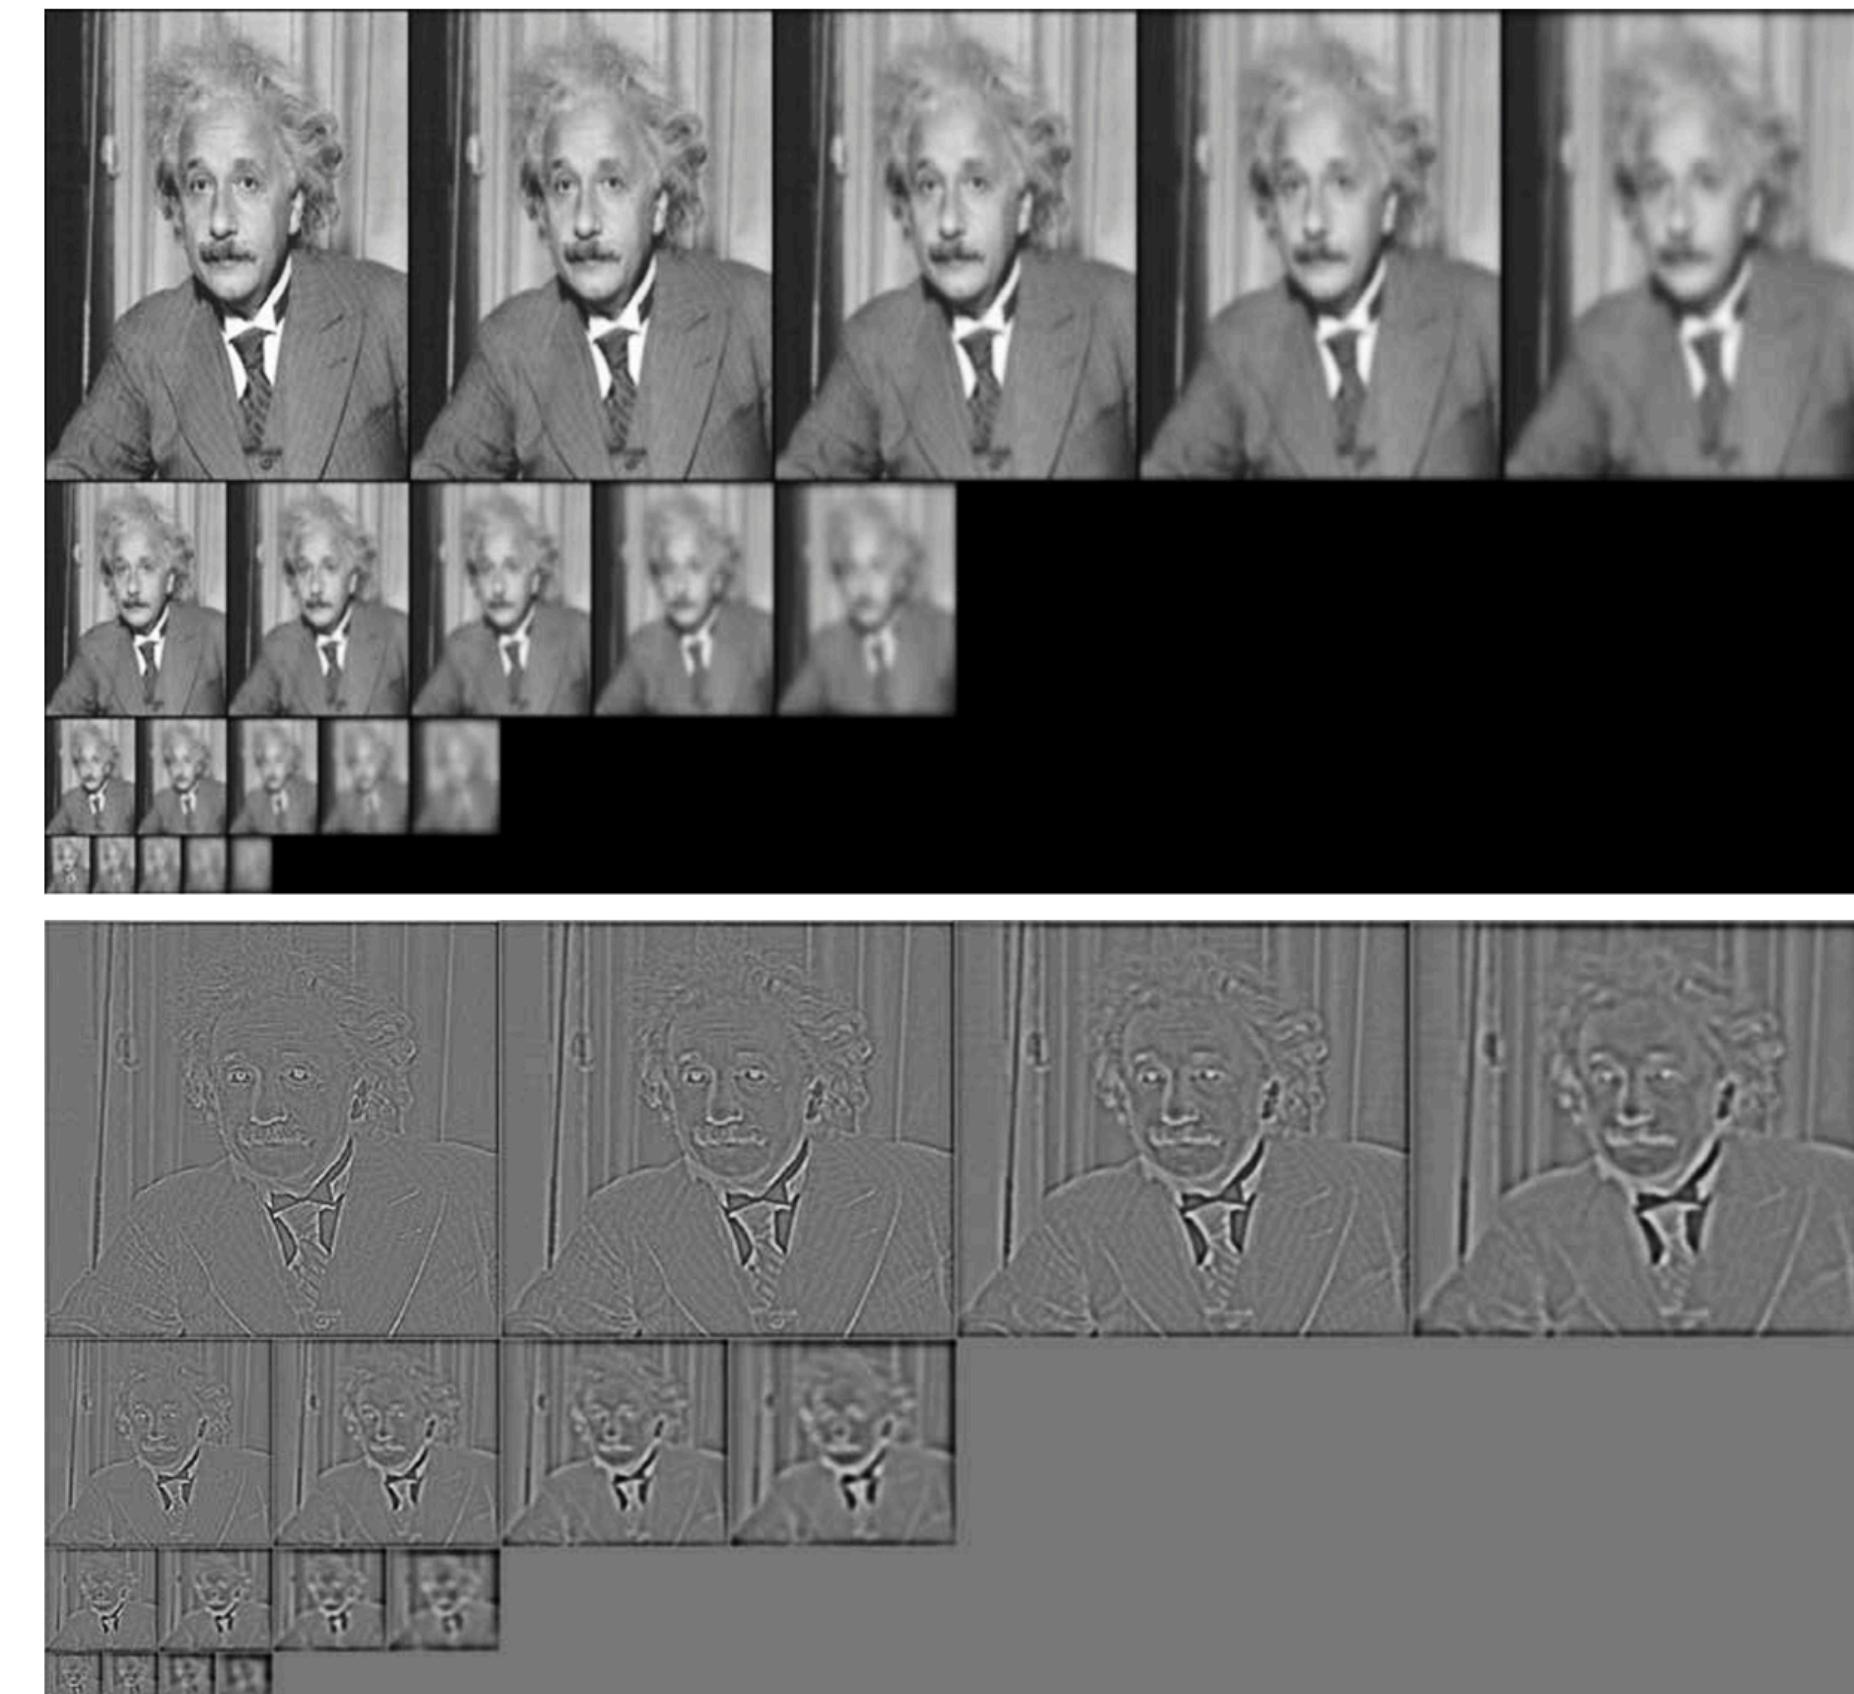
How to obtain scale invariance?

From SIFT: Use resolution pyramid.



<https://medium.com/@deepanshut041/introduction-to-surf-speeded-up-robust-features-c7396d6e7c4e>

<https://faculty.cc.gatech.edu/~afb/classes/CS4495-Fall2013/slides/CS4495-11-Features2.pdf>



Speeded-Up Robust Features (SURF)

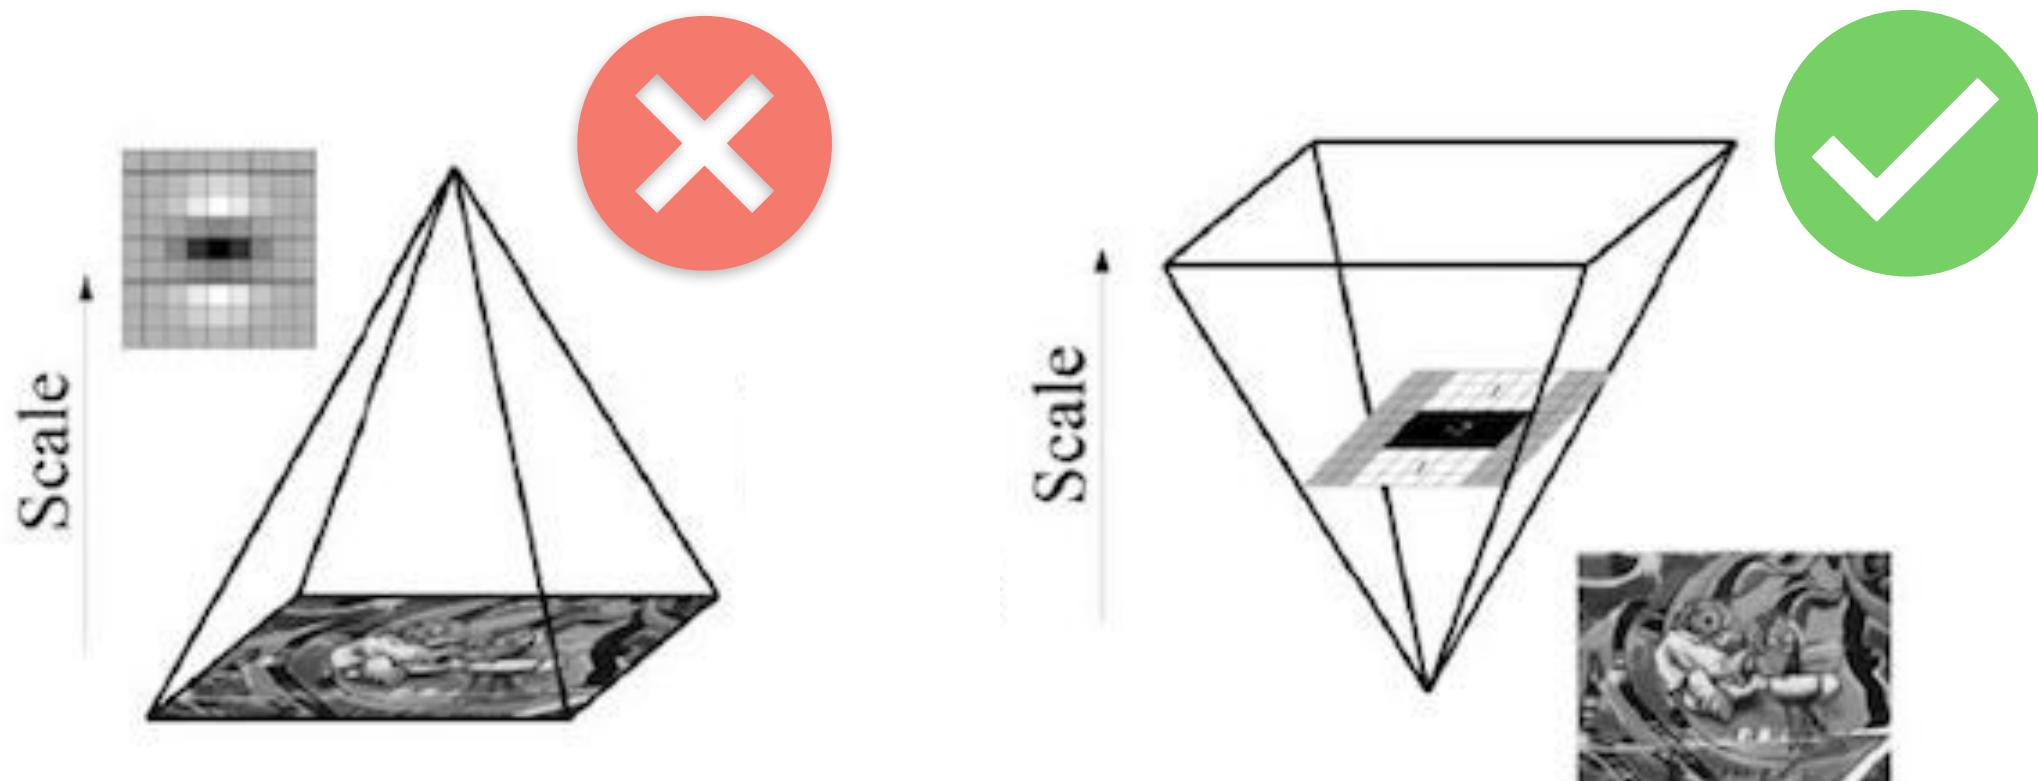
Faster Keypoint Detection

Scale Invariance

How to obtain scale invariance?

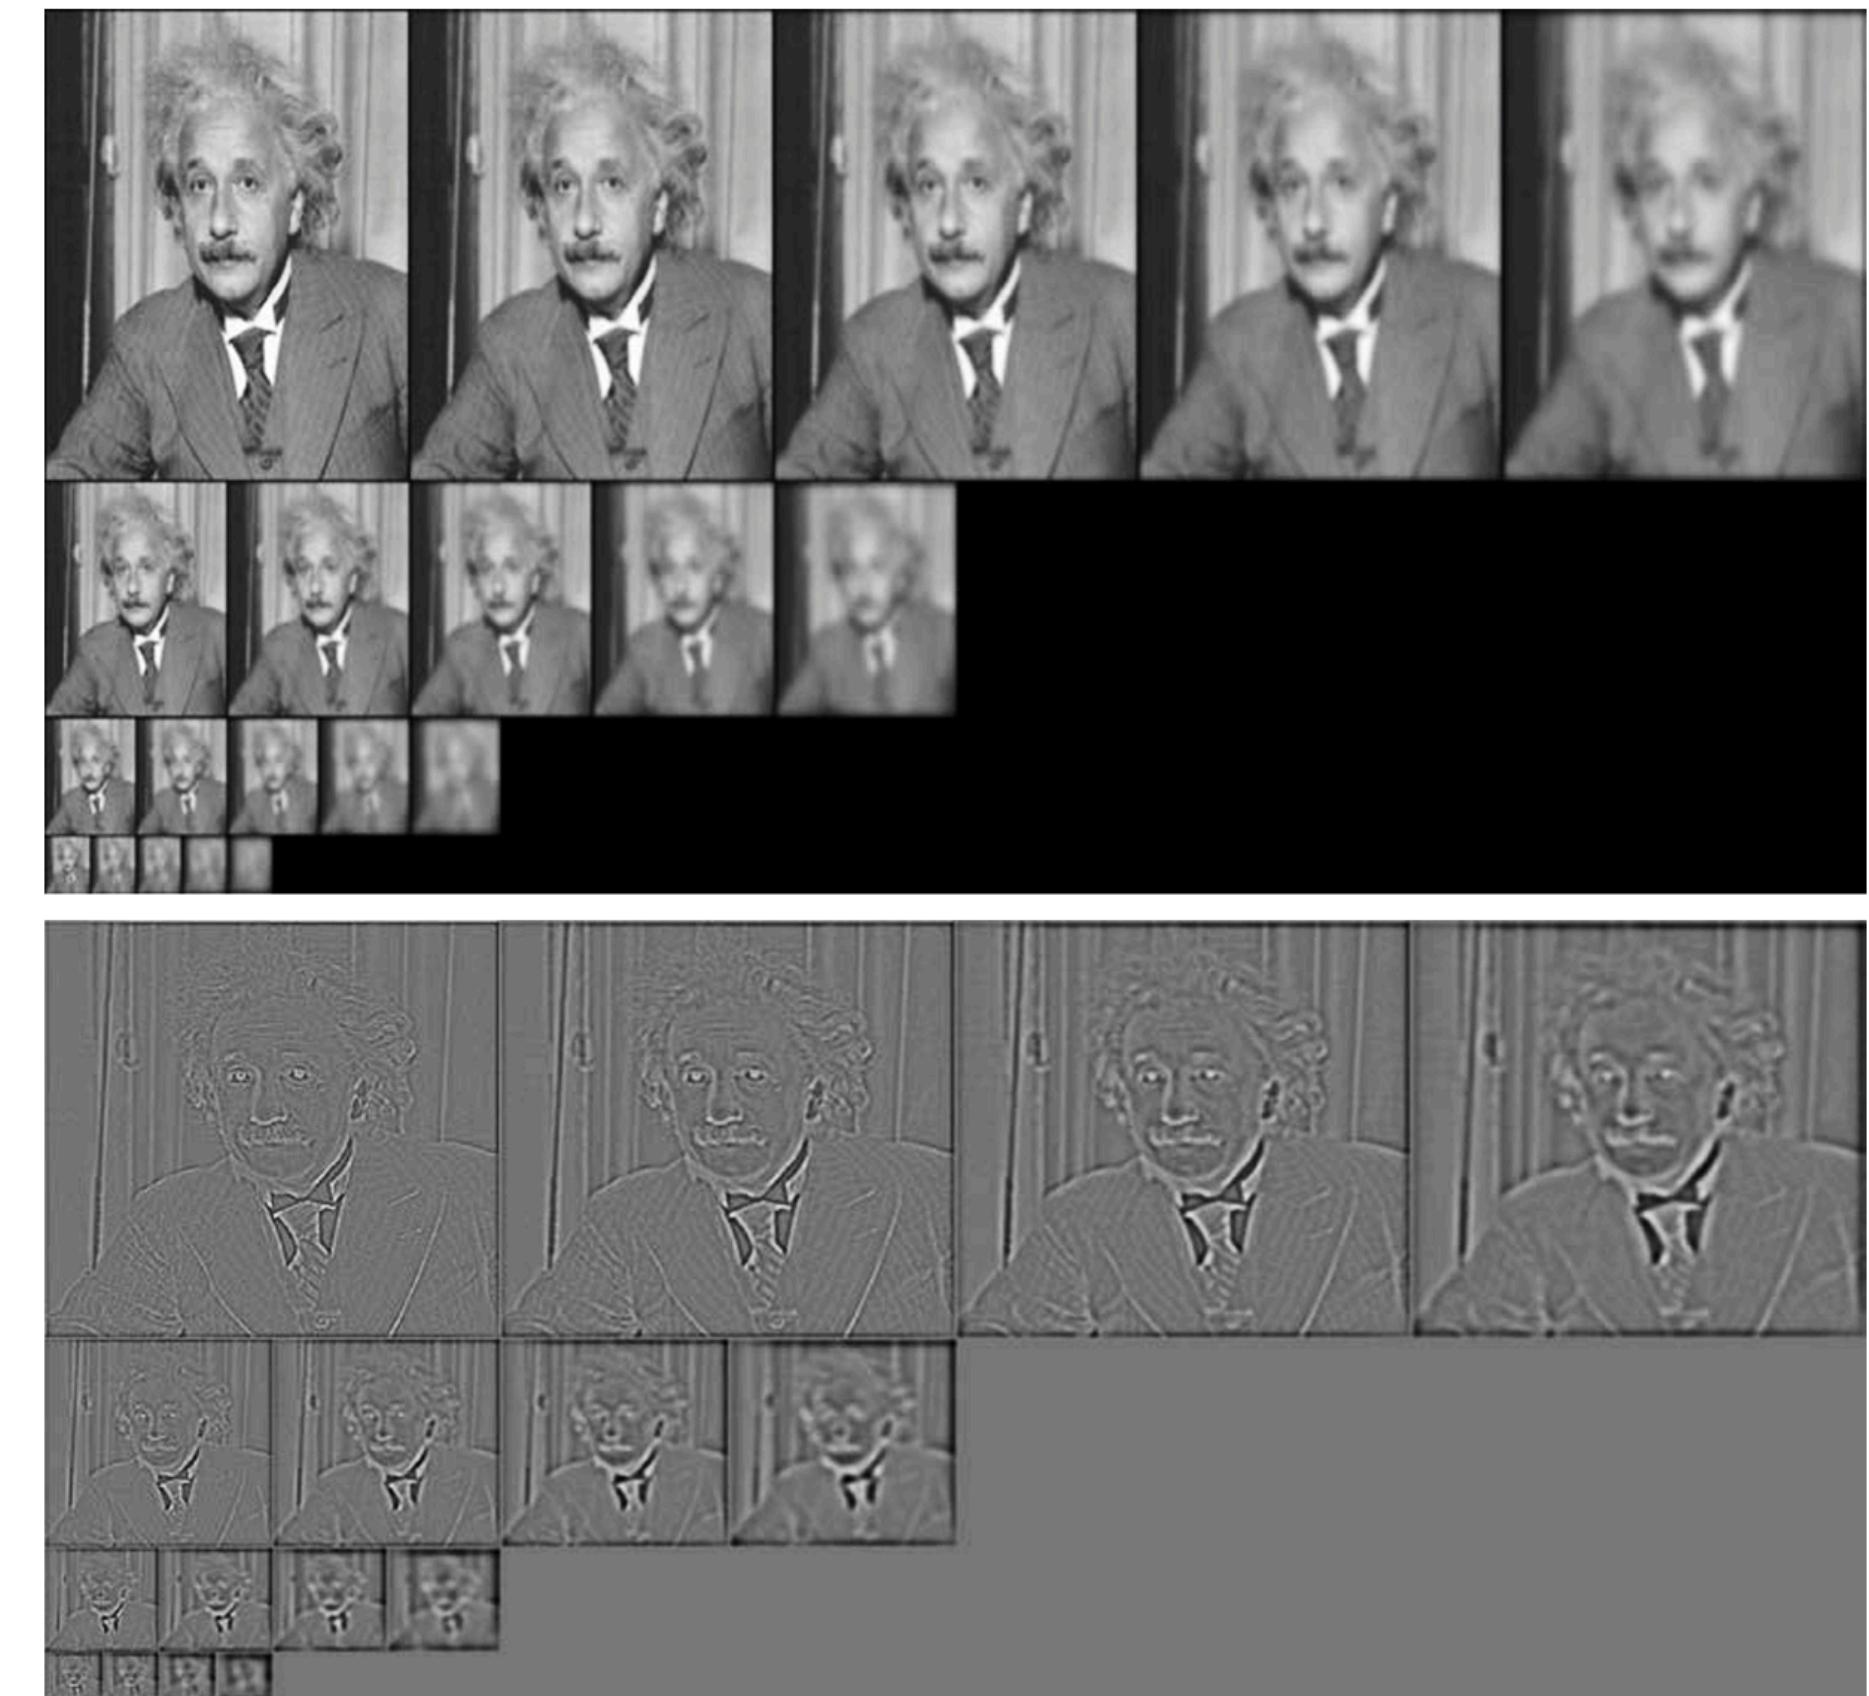
From SIFT: Use resolution pyramid.

But instead of reducing the images,
reduce the box filters!



<https://medium.com/@deepanshut041/introduction-to-surf-speeded-up-robust-features-c7396d6e7c4e>

<https://faculty.cc.gatech.edu/~afb/classes/CS4495-Fall2013/slides/CS4495-11-Features2.pdf>

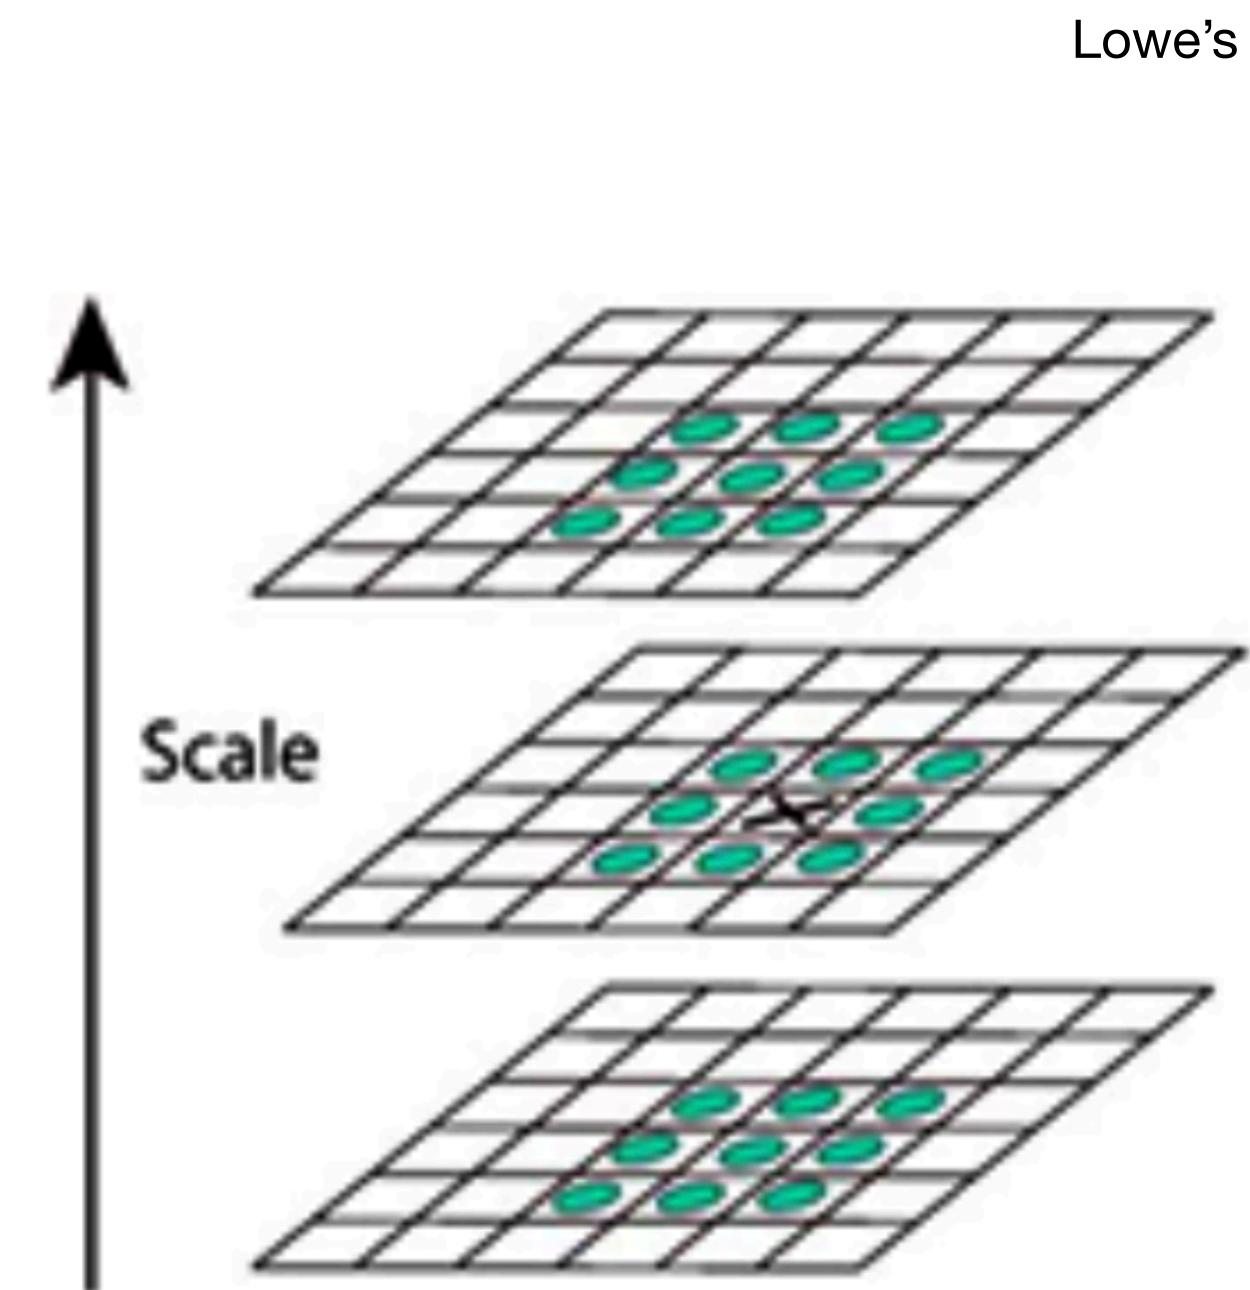


Speeded-Up Robust Features (SURF)

Faster Keypoint Detection

Non-maximal Suppression

Good SURF keypoints present the highest $det(H)$ values among their immediate $(3 \times 3 \times 3)$ scale-space neighborhood.

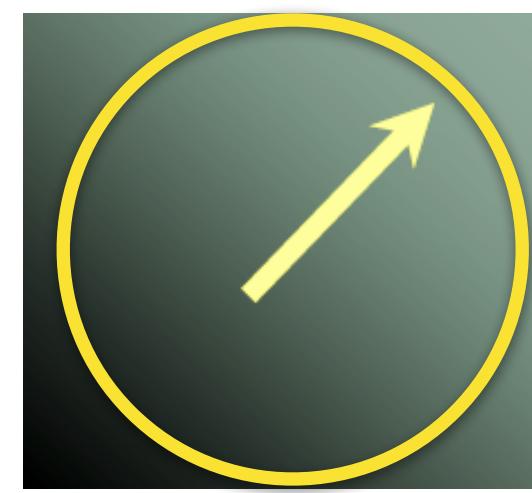
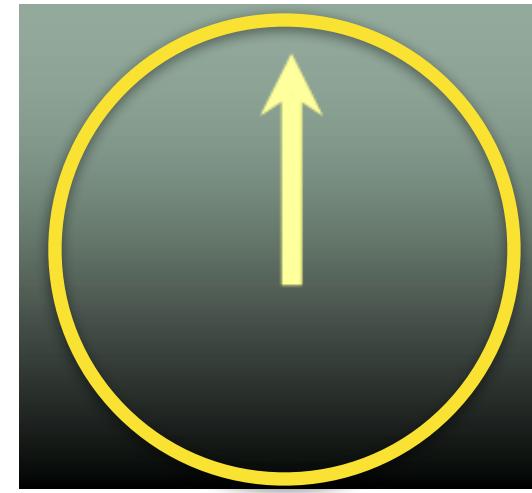
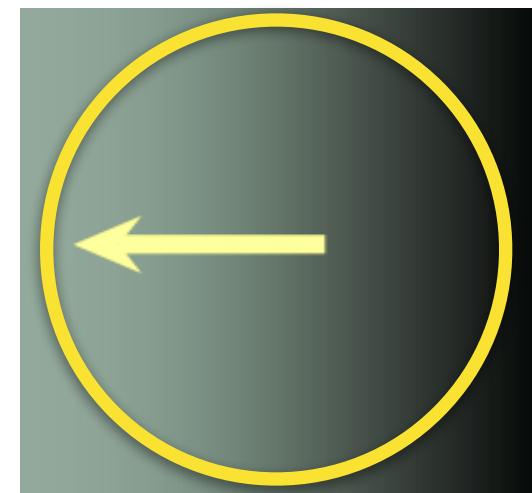


Speeded-Up Robust Features (SURF)

Faster Keypoint Detection

Orientation Assignment

Compute the gradient angle for all the pixels within the keypoint neighborhood considering its scale σ (from the Hessian matrix).



Speeded-Up Robust Features (SURF)

Faster Keypoint Detection

Orientation Assignment

Compute the gradient angle for all the pixels
within the keypoint neighborhood
considering its scale σ (from the Hessian matrix).



d_x



d_y

Use box filters d_x and d_y with scale proportional to σ .

Speeded-Up Robust Features (SURF)

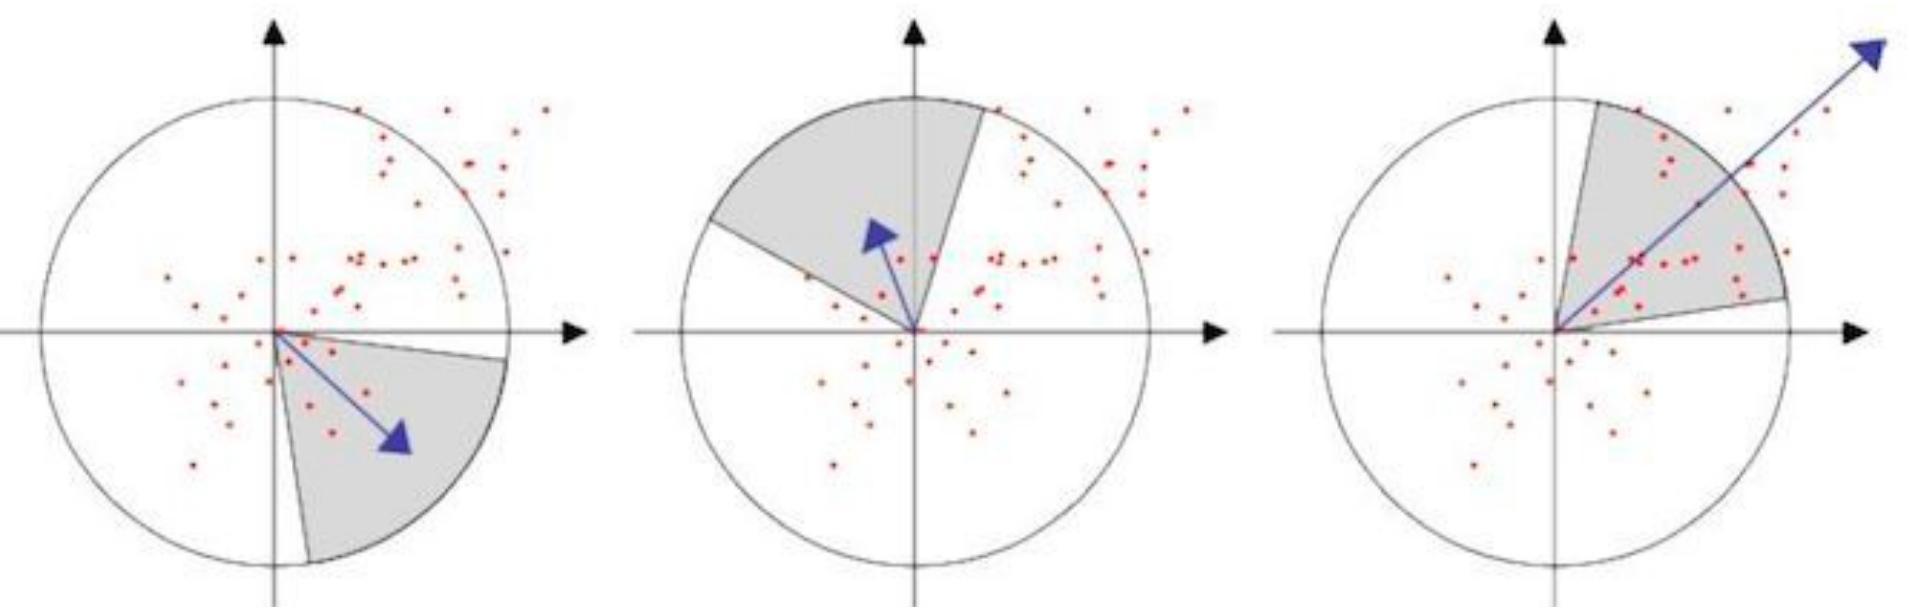
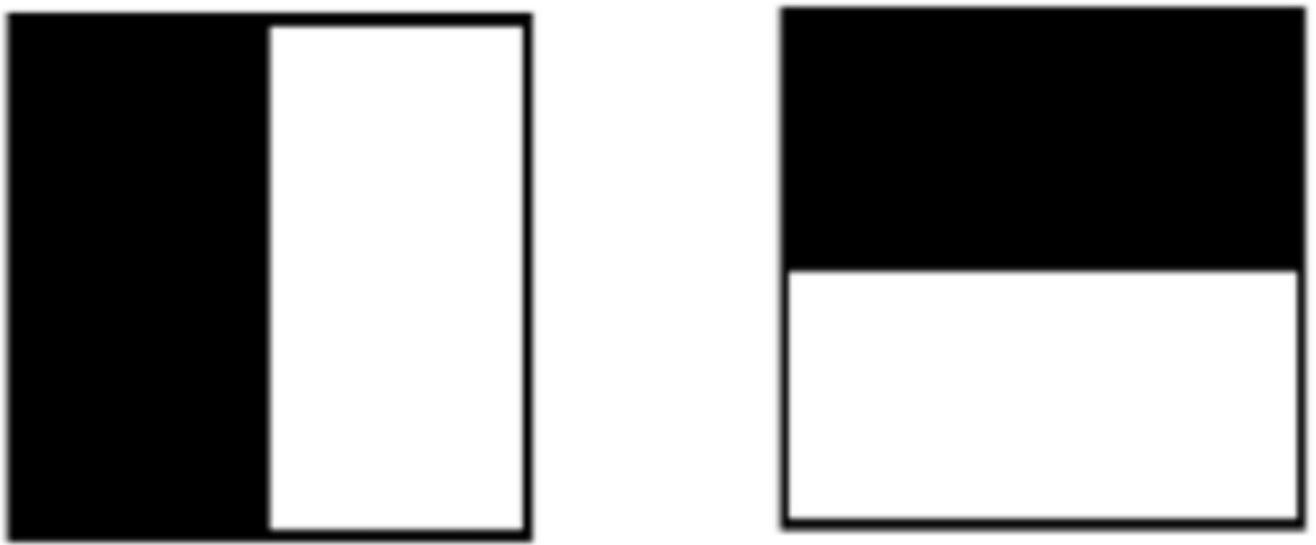
Faster Keypoint Detection

Orientation Assignment

Compute the gradient angle for all the pixels within the keypoint neighborhood considering its scale σ (from the Hessian matrix).

Use box filters d_x and d_y with scale proportional to σ .

Create an angle histogram with 6 bins.
Take the dominant angle as the keypoint orientation.



<https://medium.com/@deepanshut041/introduction-to-surf-speeded-up-robust-features-c7396d6e7c4e>

Speeded-Up Robust Features (SURF)

Faster Keypoint Detection

SURF keypoints have:

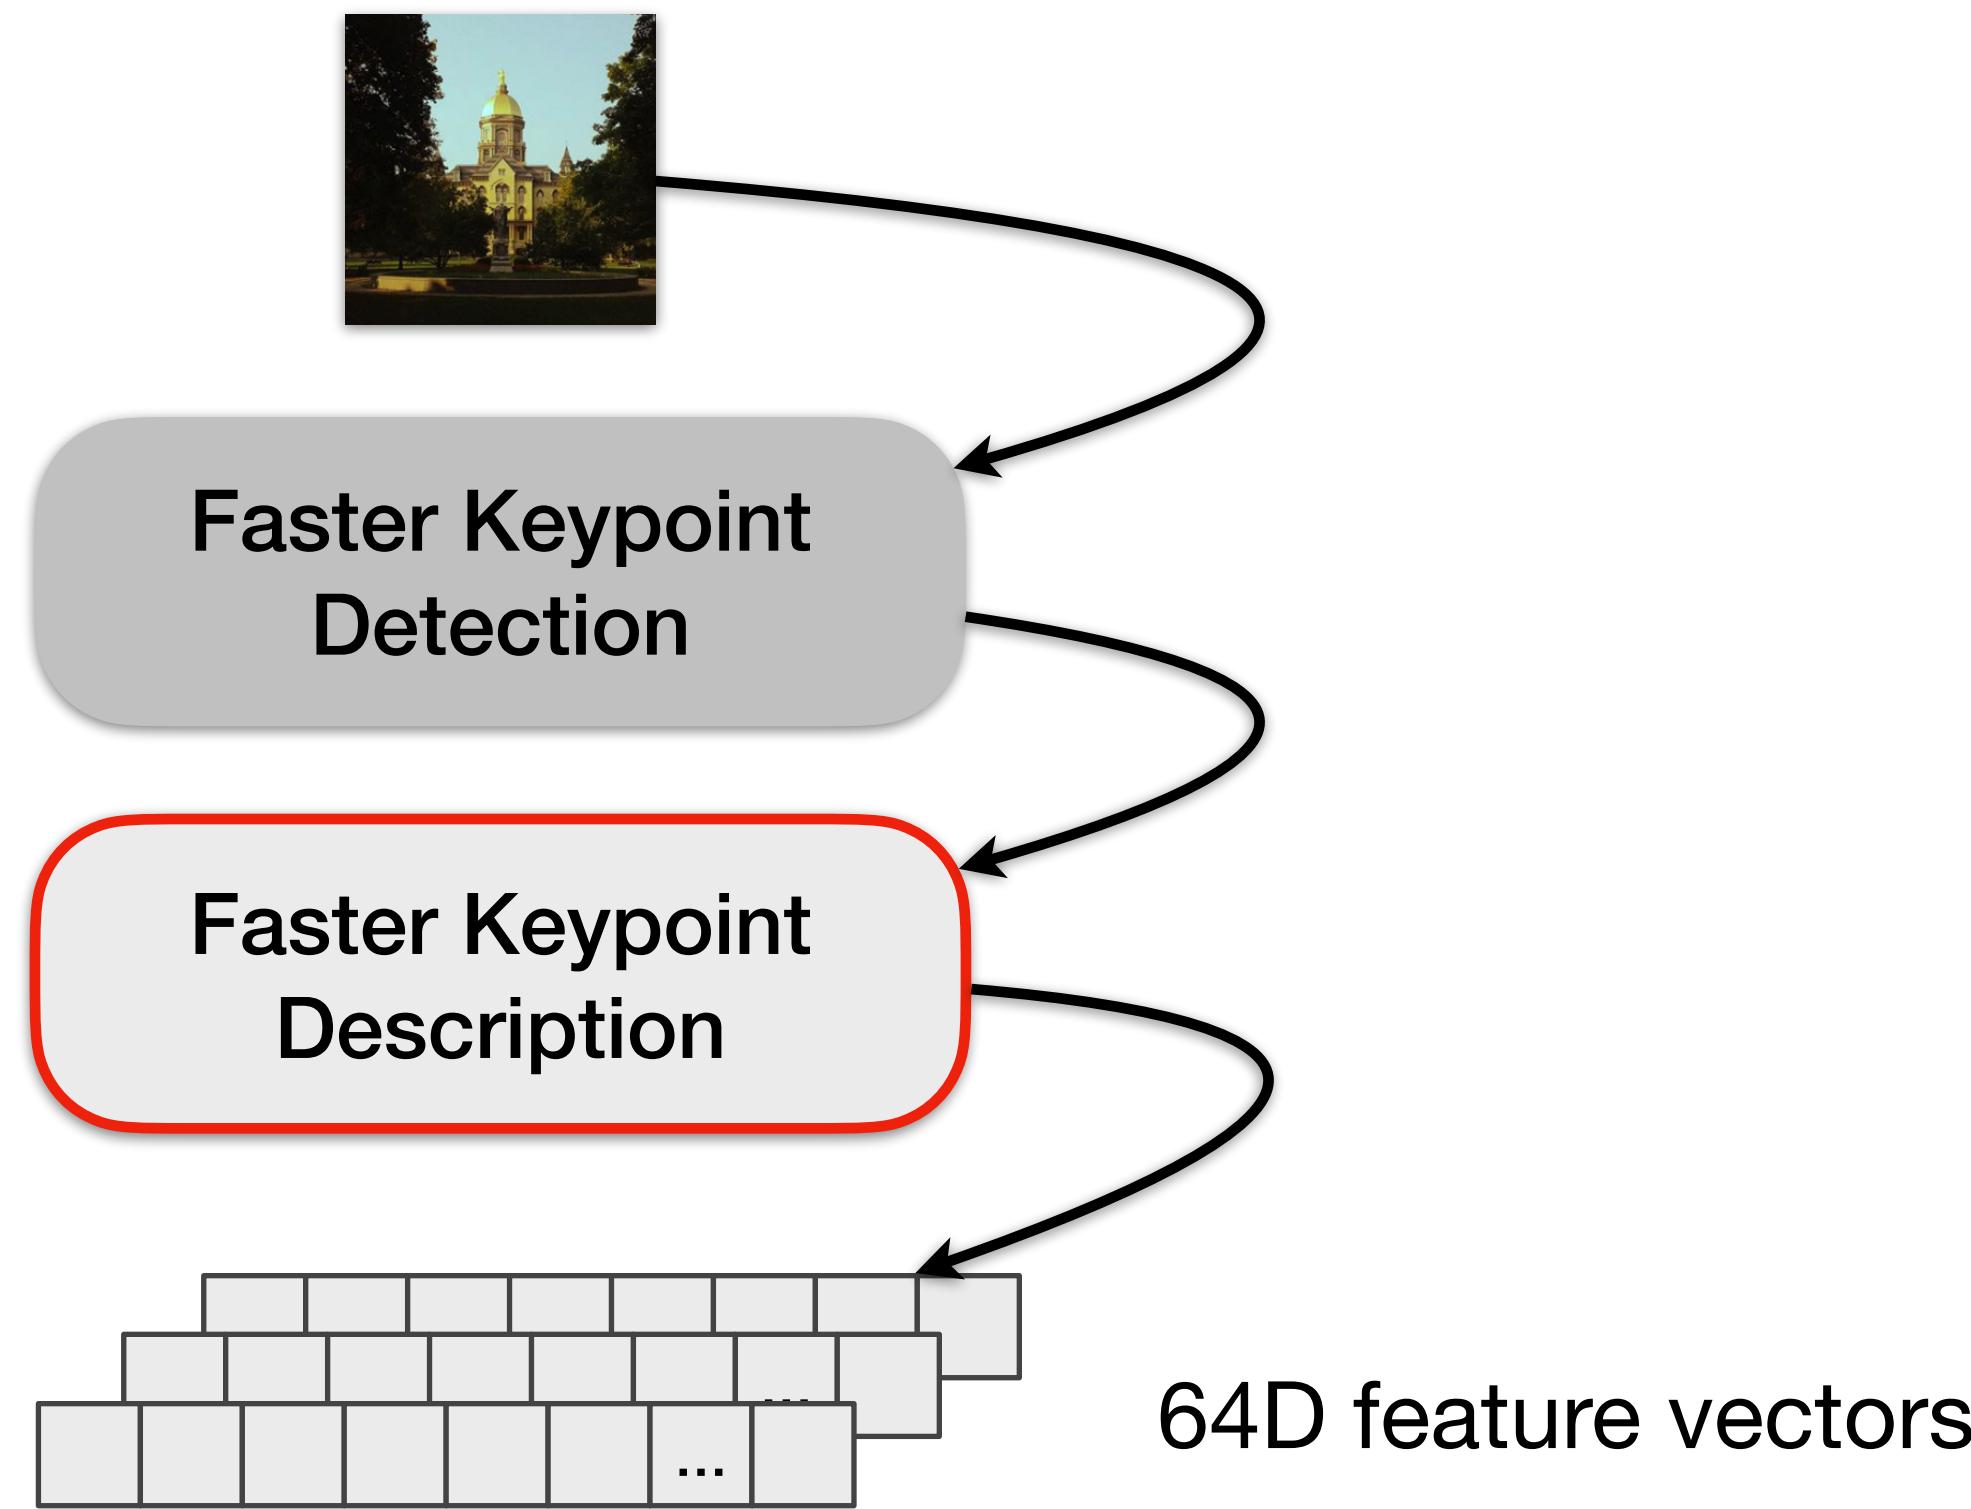
1. Location (x, y)
2. Scale (σ from Hessian matrix)
3. Orientation (dominant local gradient angle)
4. Strength ($\det(H)$)



https://docs.opencv.org/3.4/d4/dd2/tutorial_py_surf_intro.html

Speeded-Up Robust Features (SURF)

Stages



SURF: Speeded Up Robust Features

Herbert Bay¹, Tinne Tuytelaars², and Luc Van Gool^{1,2}

¹ ETH Zurich

{bay, vangoool}@vision.ee.ethz.ch

² Katholieke Universiteit Leuven

{Tinne.Tuytelaars, Luc.Vangoool}@esat.kuleuven.be

2006



LOYOLA
UNIVERSITY CHICAGO

Speeded-Up Robust Features (SURF)

Faster Keypoint Description

For each keypoint, rotate the image according to the keypoint orientation.

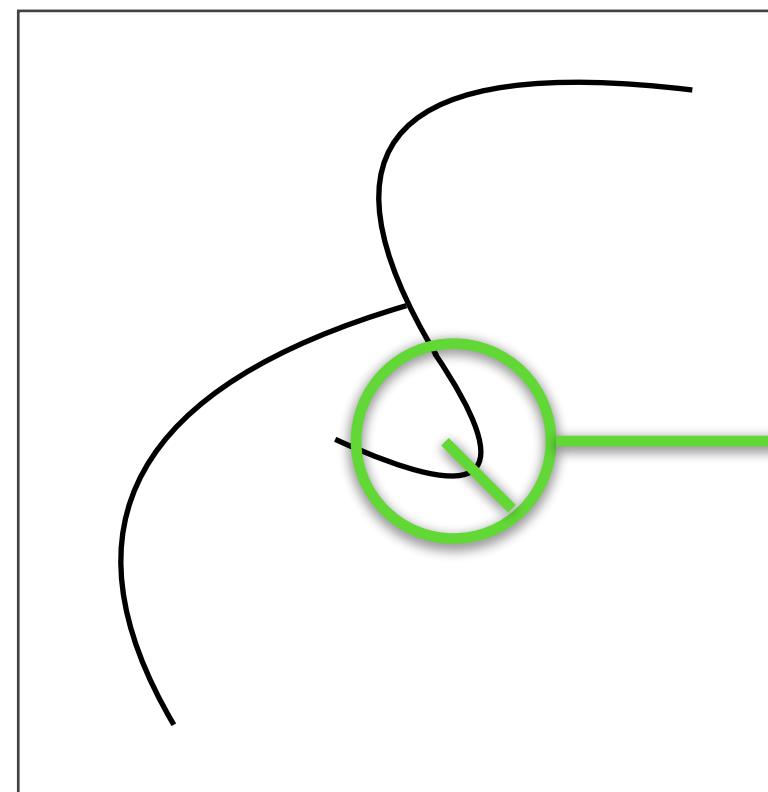


Image 1

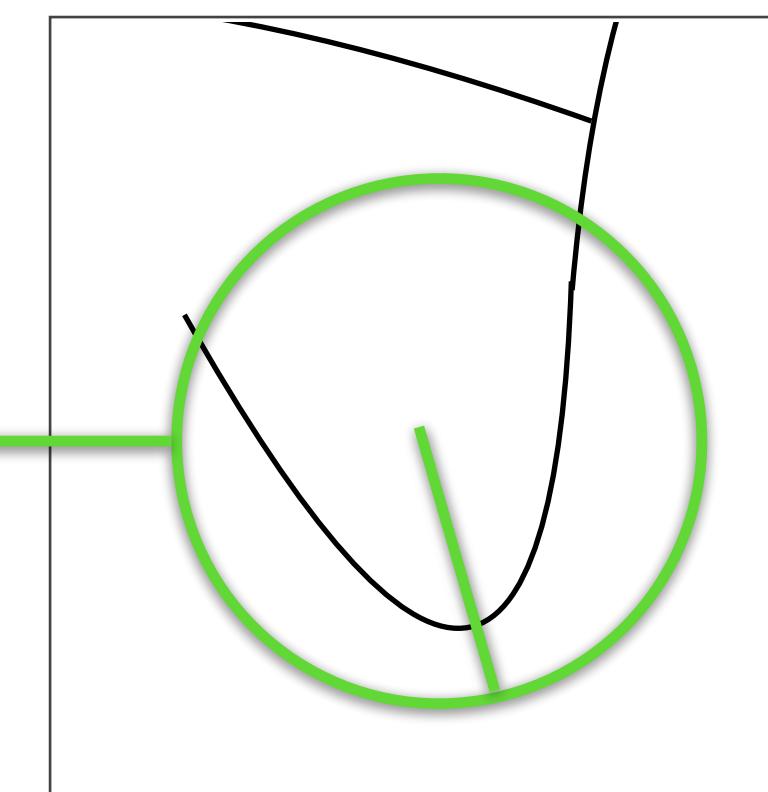


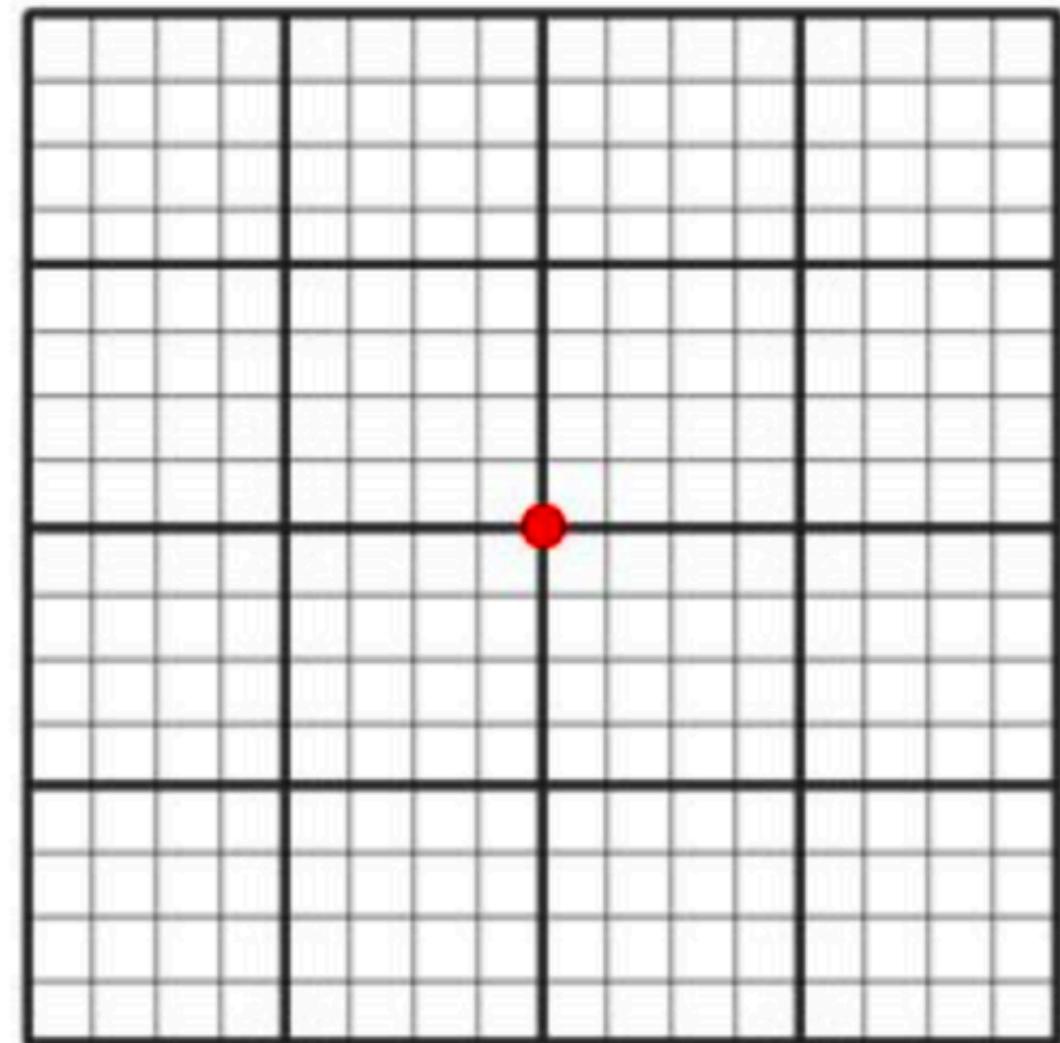
Image 2



Speeded-Up Robust Features (SURF)

Faster Keypoint Description

For each rotated keypoint, sample a 4×4 window on its neighborhood, according to the keypoint scale.



Speeded-Up Robust Features (SURF)

Faster Keypoint Description

For each rotated keypoint, sample a 4×4 window on its neighborhood, according to the keypoint scale.

For each one of the 4×4 cells, compute 4 sums:

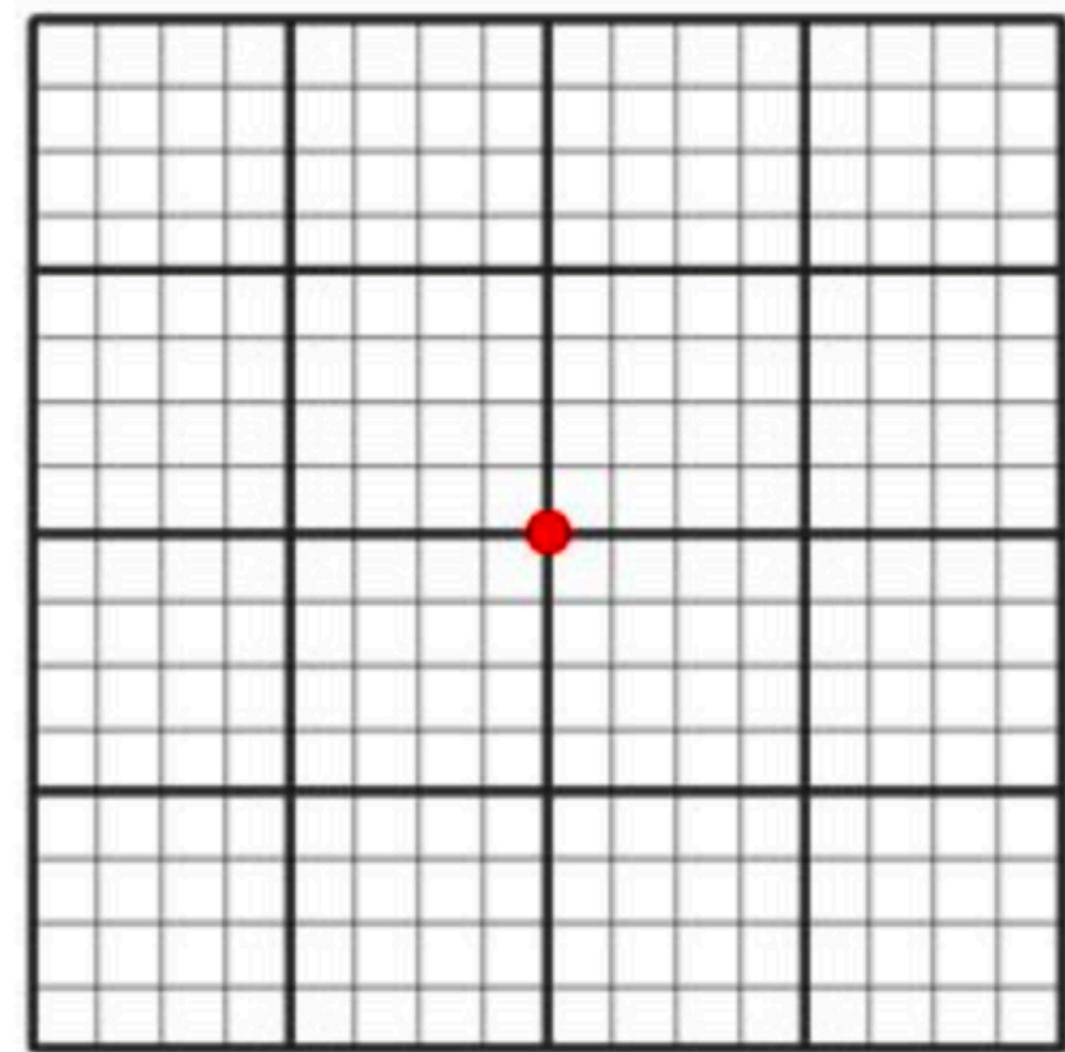
$$(1) \sum d_x, (2) \sum |d_x|, (3) \sum d_y, \text{ and } (4) \sum |d_y|.$$



d_x



d_y



LOYOLA
UNIVERSITY CHICAGO

Speeded-Up Robust Features (SURF)

Faster Keypoint Description

For each rotated keypoint, sample a 4×4 window on its neighborhood, according to the keypoint scale.

For each one of the 4×4 cells, compute 4 sums:

$$(1) \sum d_x, (2) \sum |d_x|, (3) \sum d_y, \text{ and } (4) \sum |d_y|.$$

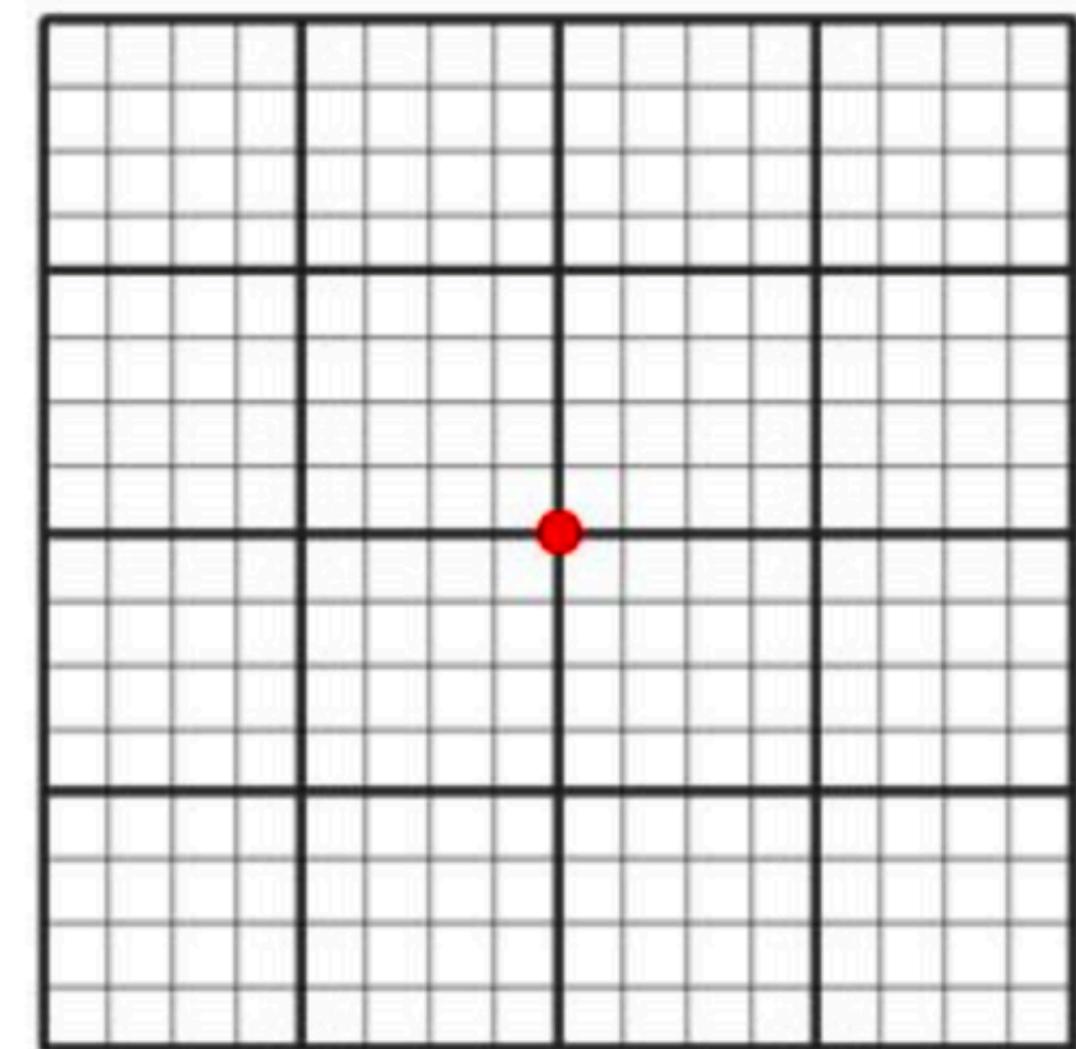
Fill out a feature vector with the $4 \times 4 \times 4 = 64$ values.



d_x

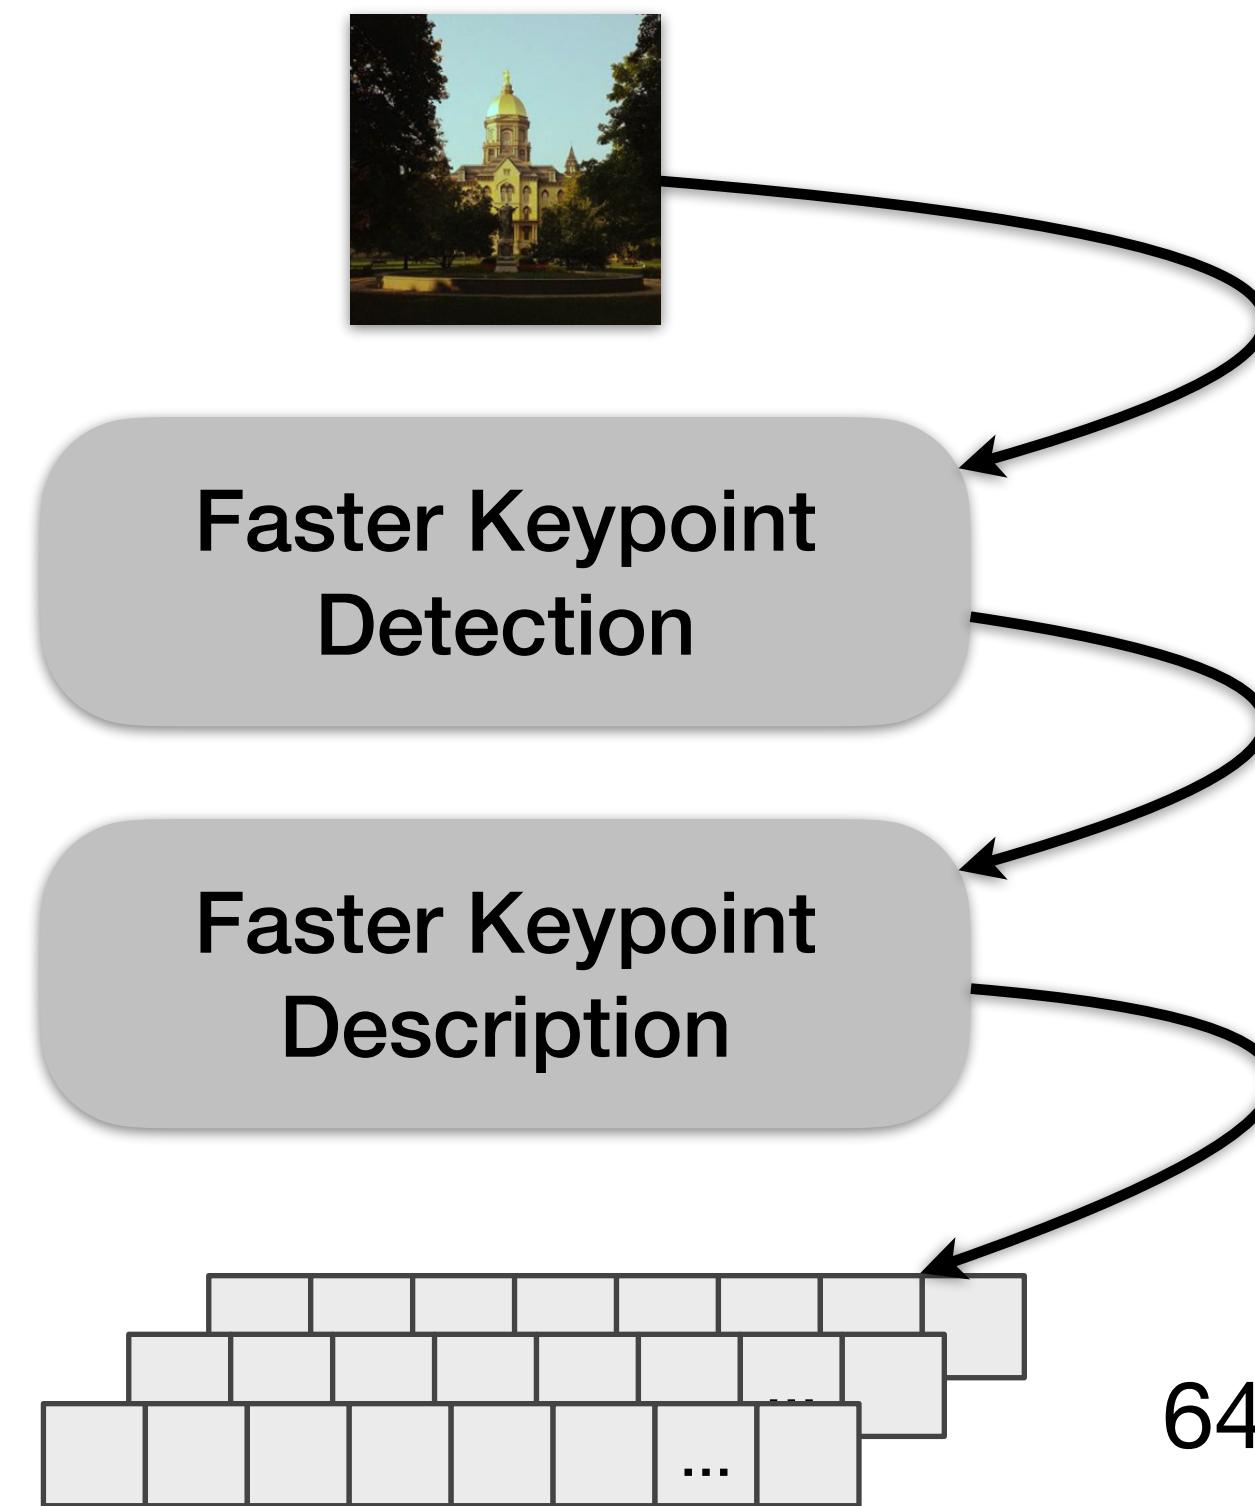


d_y



Speeded-Up Robust Features (SURF)

Stages



64D feature vectors, compare two vectors with L2- or cosine distance.

SURF: Speeded Up Robust Features

Herbert Bay¹, Tinne Tuytelaars², and Luc Van Gool^{1,2}

¹ ETH Zurich

{bay, vangool}@vision.ee.ethz.ch

2006

² Katholieke Universiteit Leuven

{Tinne.Tuytelaars, Luc.Vangool}@esat.kuleuven.be

Speeded-Up Robust Features (SURF)

Example



<https://bit.ly/3qOj3jx>

The screenshot shows a Jupyter Notebook interface. On the left, there is a sidebar titled "Table of contents" with the following structure:

- SURF Example
 - Webcam Usage
 - SURF keypoints detection and description
 - Importing of the necessary libraries.
 - Loading webcam image into memory.
 - Detecting and Describing SURF keypoints
- Section

On the right, the main area shows two sections expanded:

- SURF Example**
- Webcam Usage**

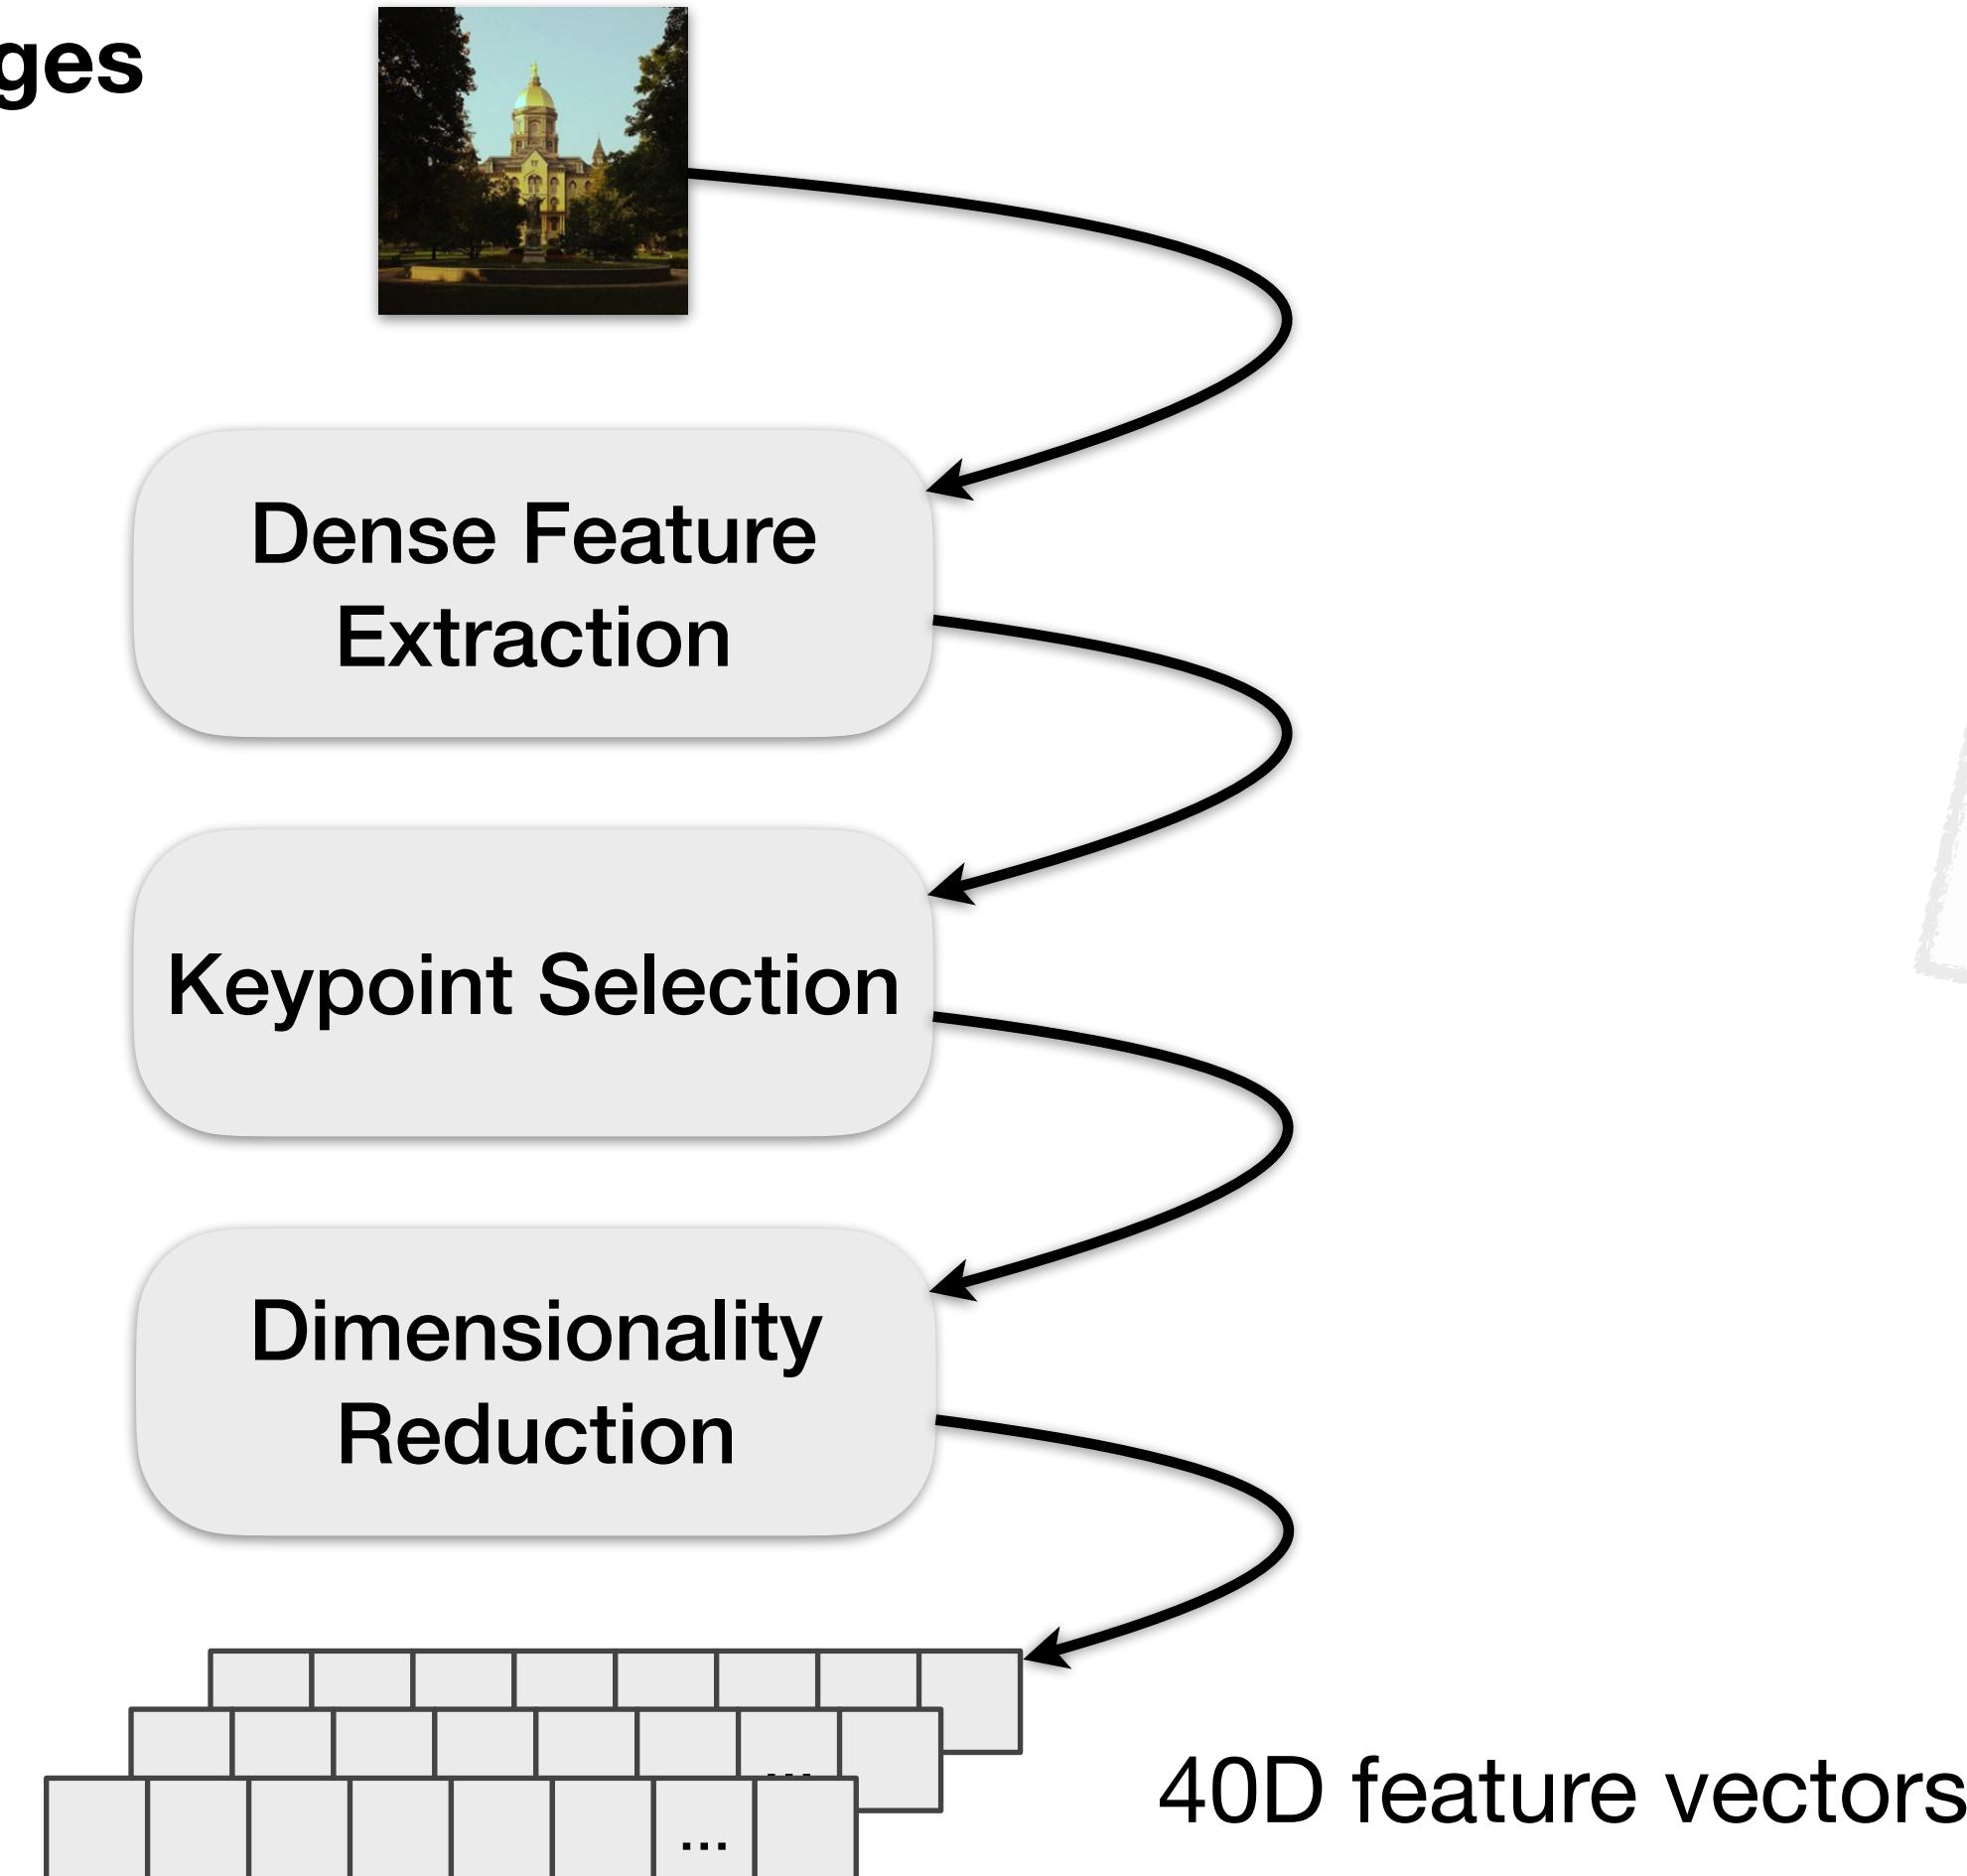
Under "Webcam Usage", there is some Python code:

```
[1] from IPython.display import display, Javascript
from google.colab.output import eval_js
from base64 import b64decode

def take_photo(filename='photo.jpg', quality=0.8):
    js = Javascript('''
        async function takePhoto(quality) {
            const div = document.createElement('div');
            const capture = document.createElement('button');
```

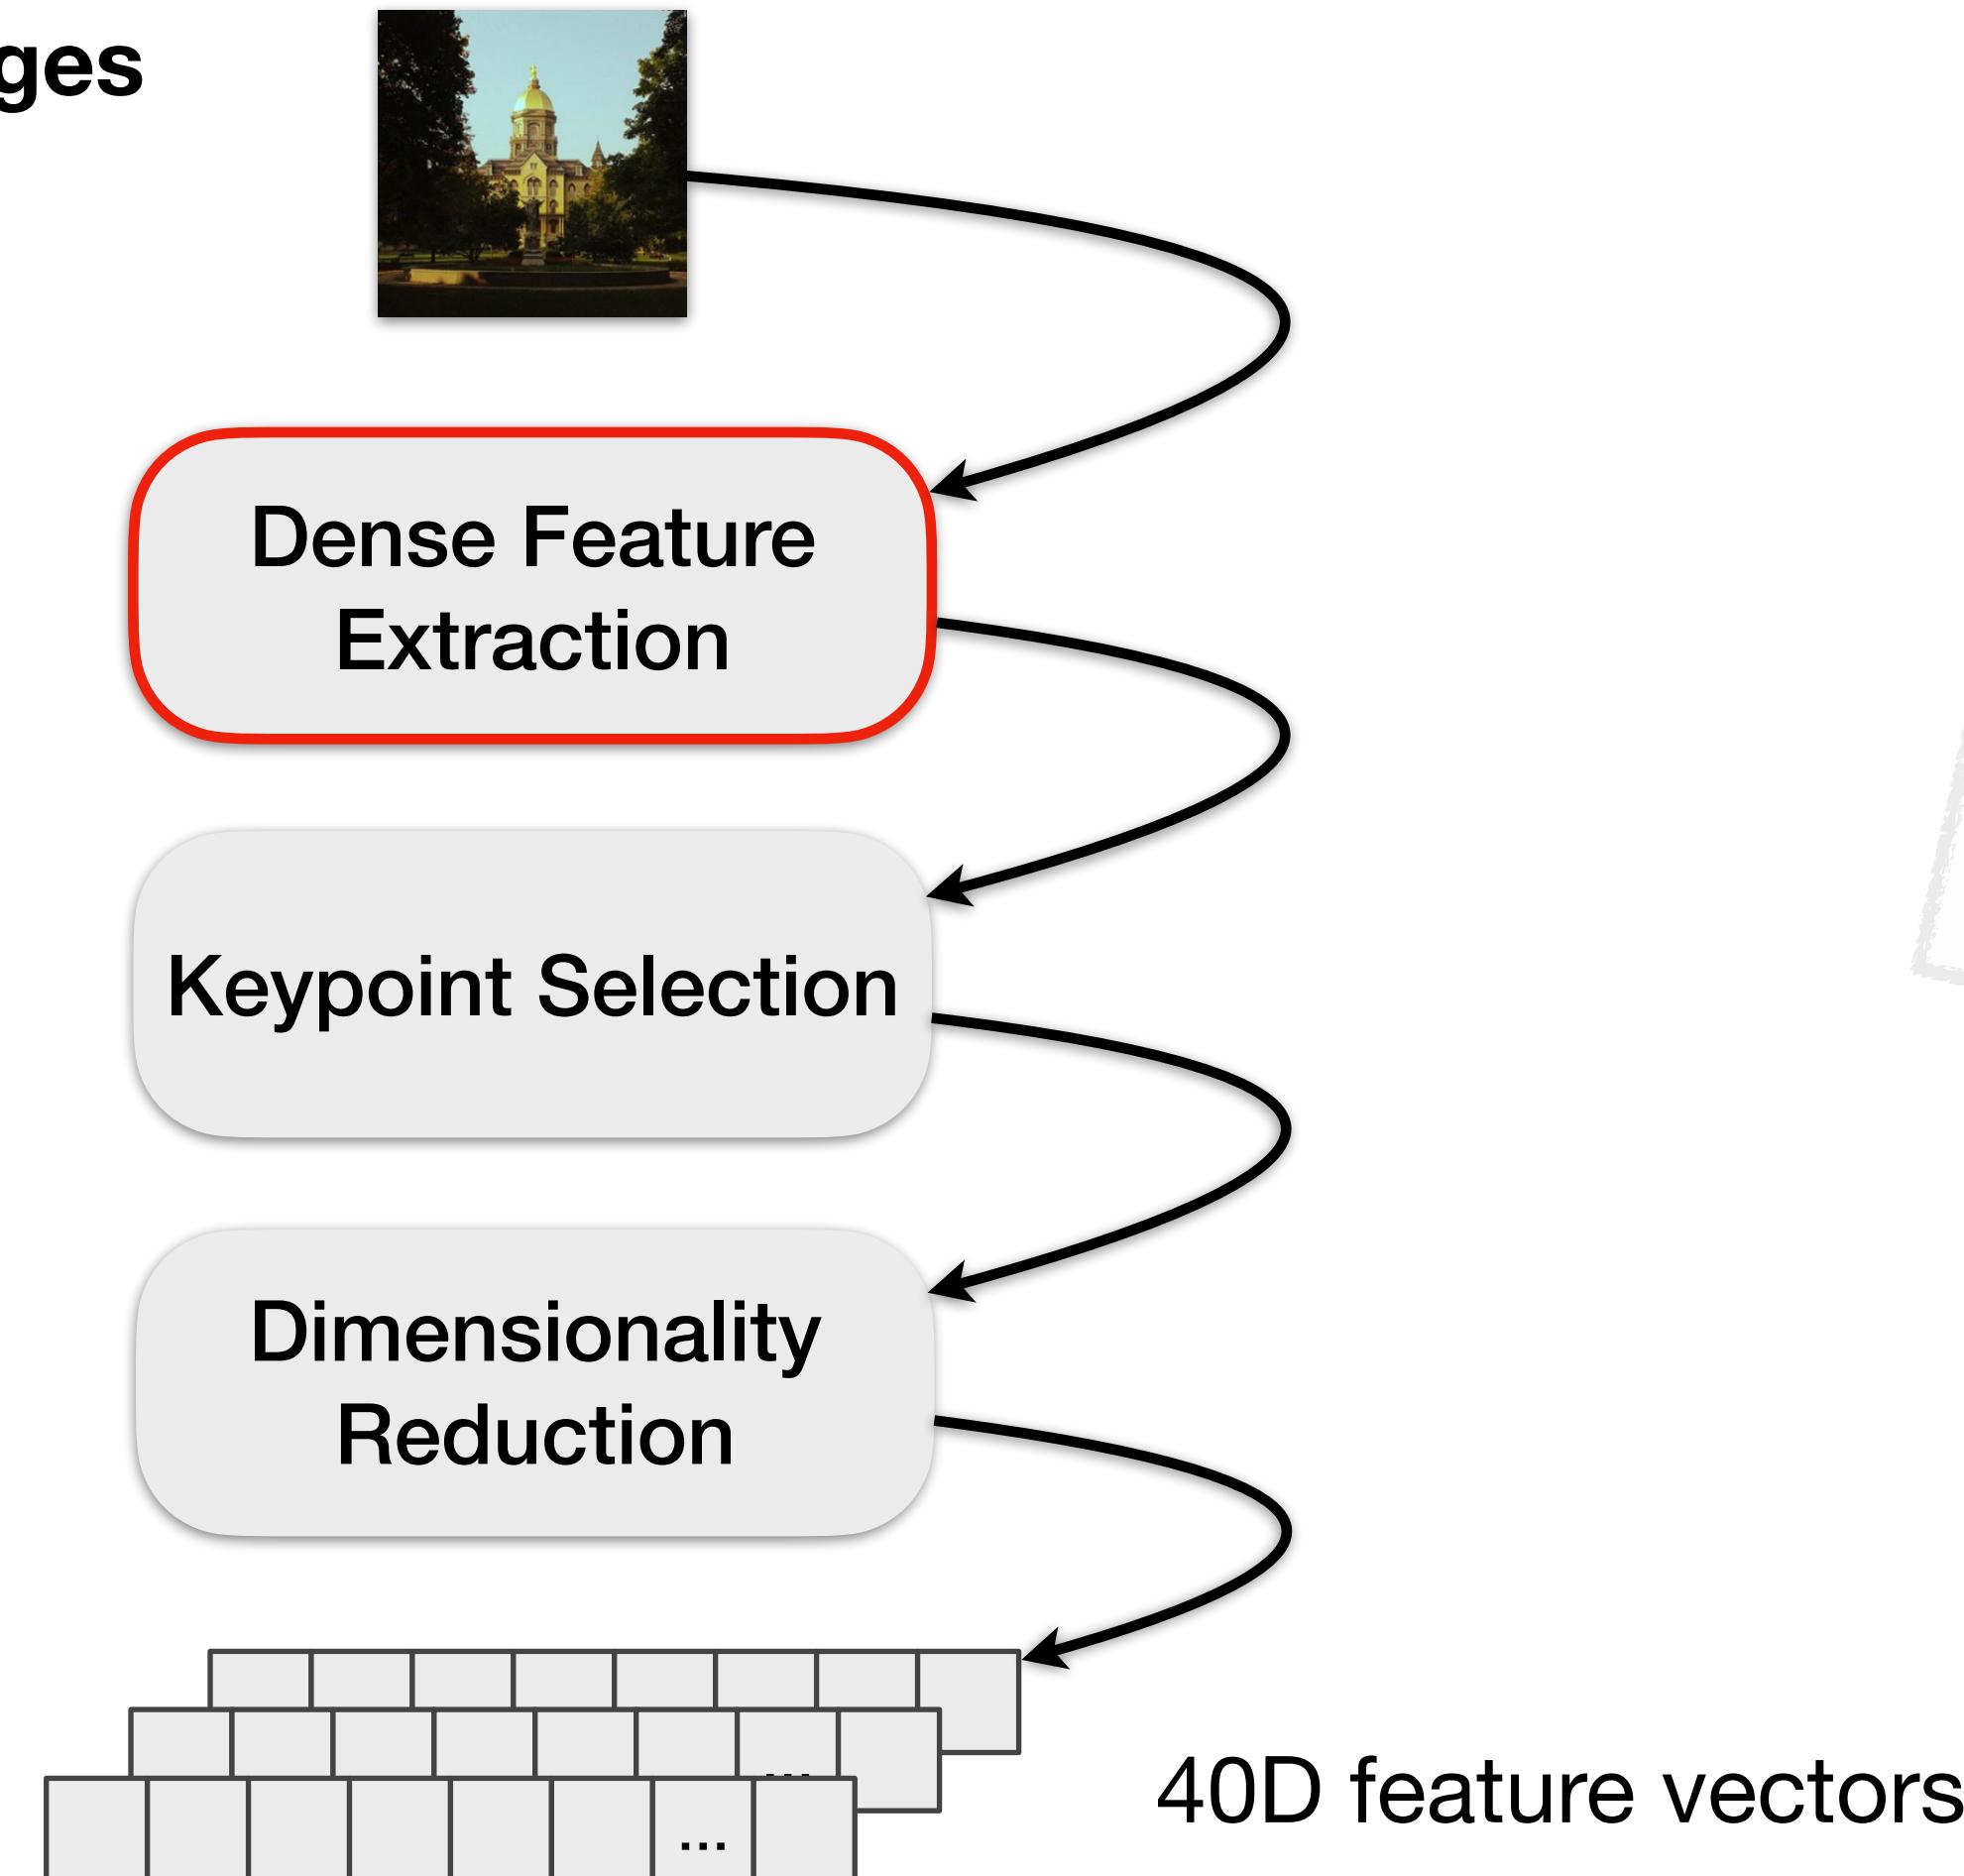
Deep Local Features (DELF)

Stages



Deep Local Features (DELF)

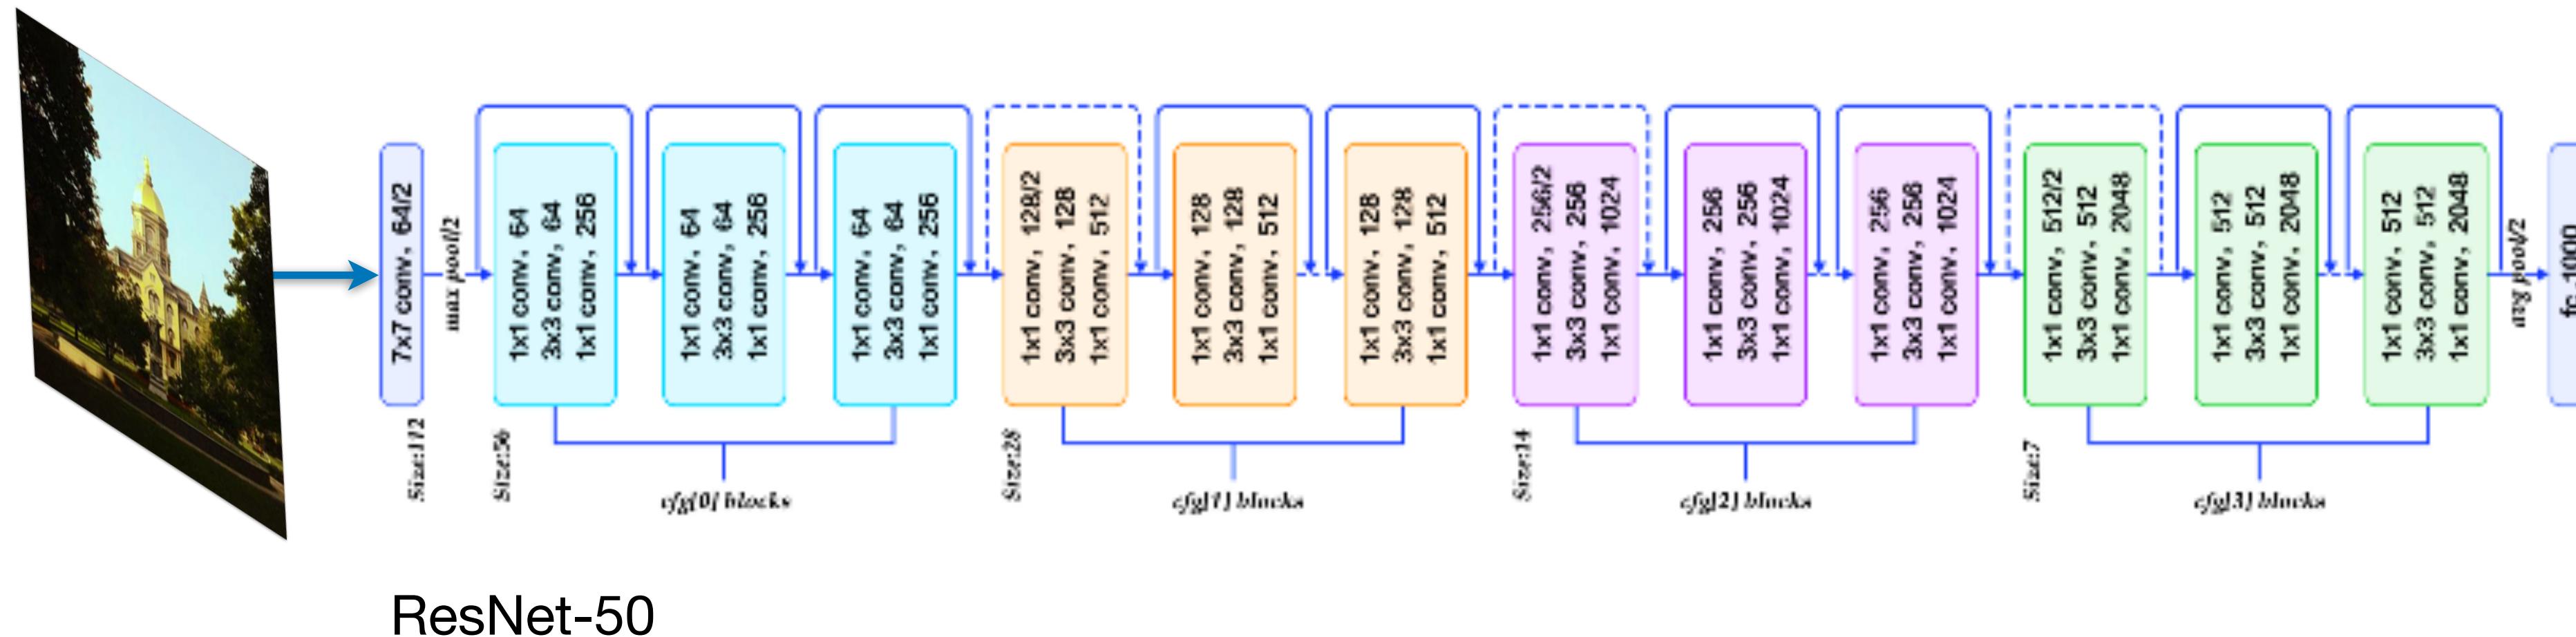
Stages



Deep Local Features (DELF)

Dense Feature Extraction

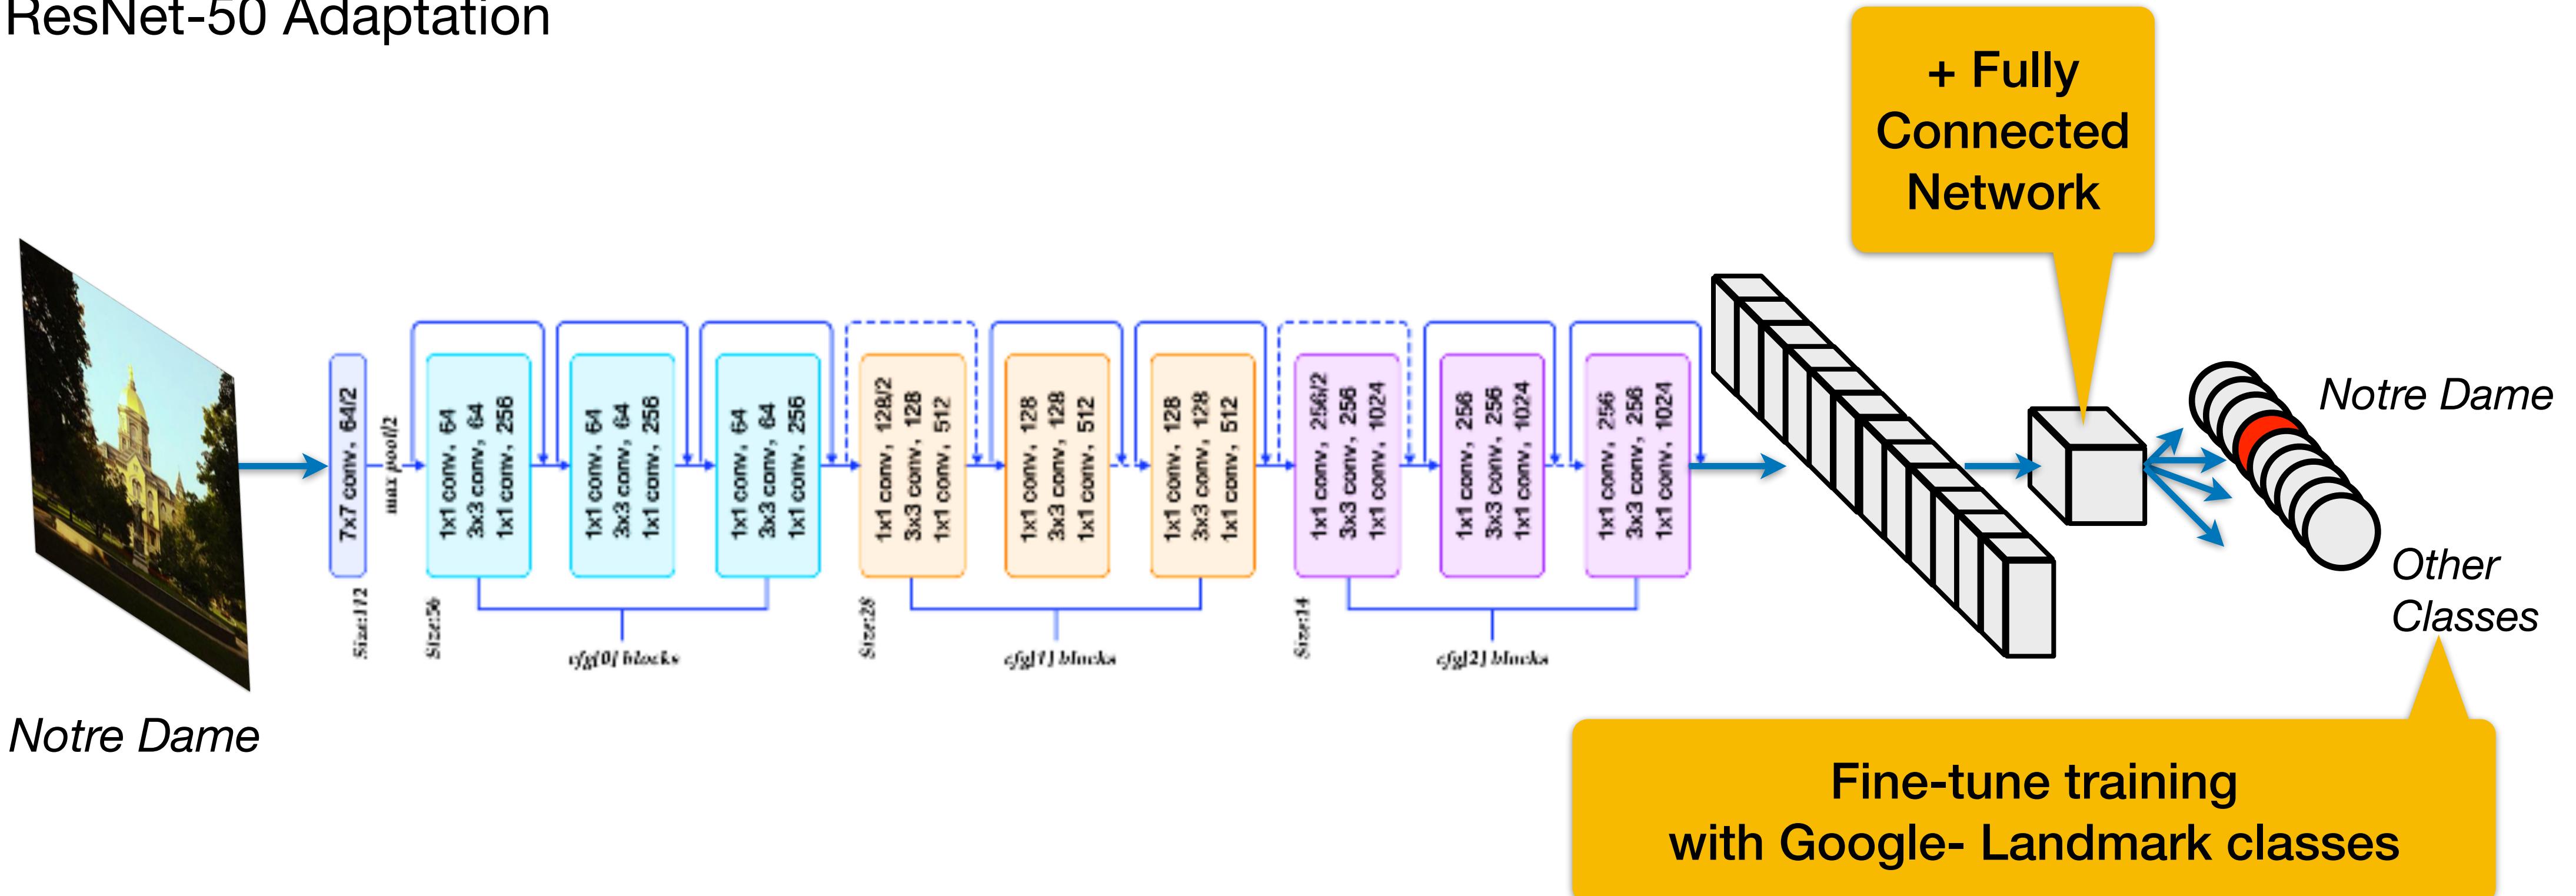
Starting Point: ResNet-50



Deep Local Features (DELF)

Dense Feature Extraction

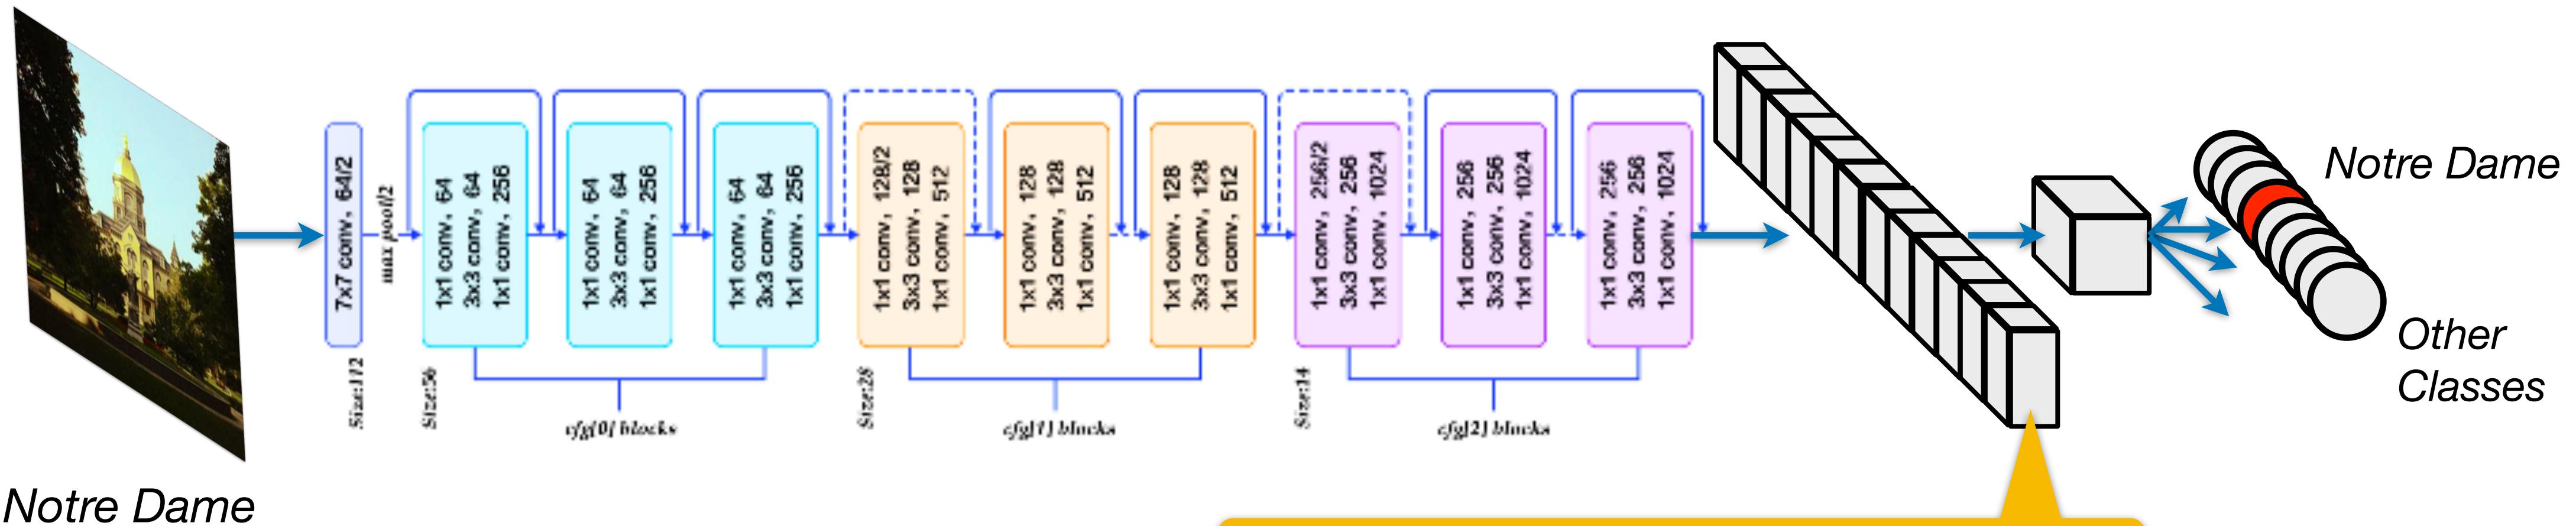
ResNet-50 Adaptation



Deep Local Features (DELF)

Dense Feature Extraction

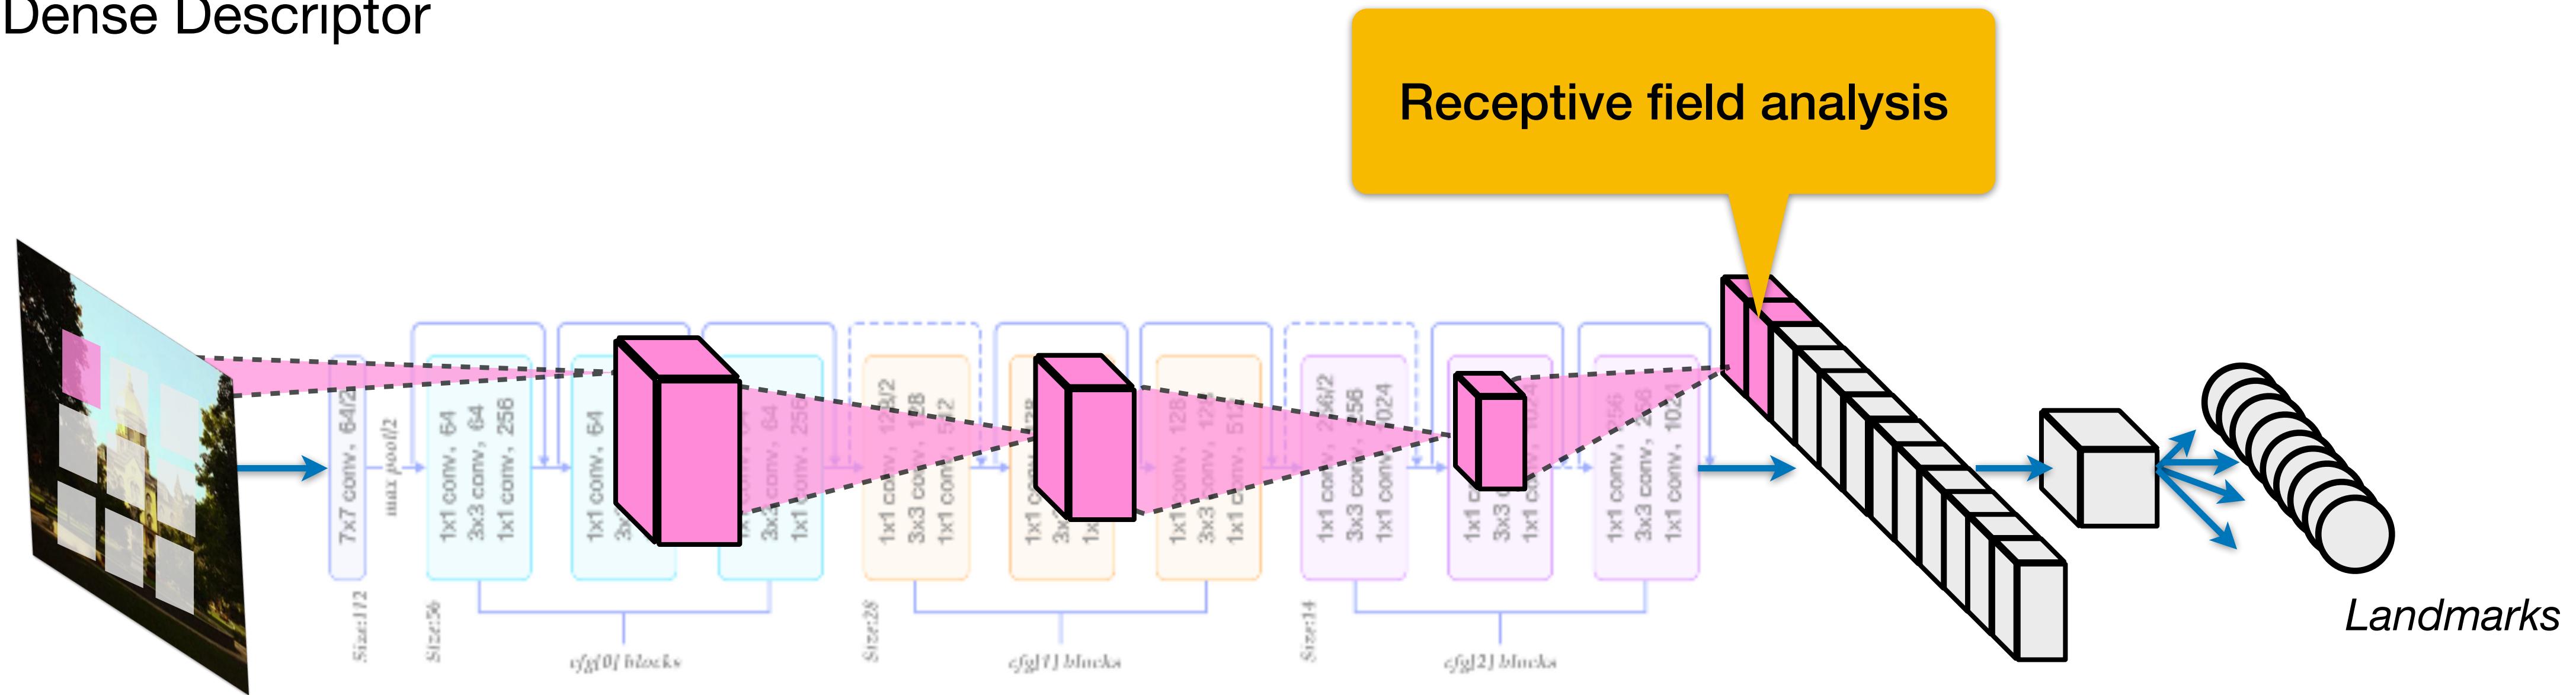
How to convert global description to dense description?



Deep Local Features (DELF)

Dense Feature Extraction

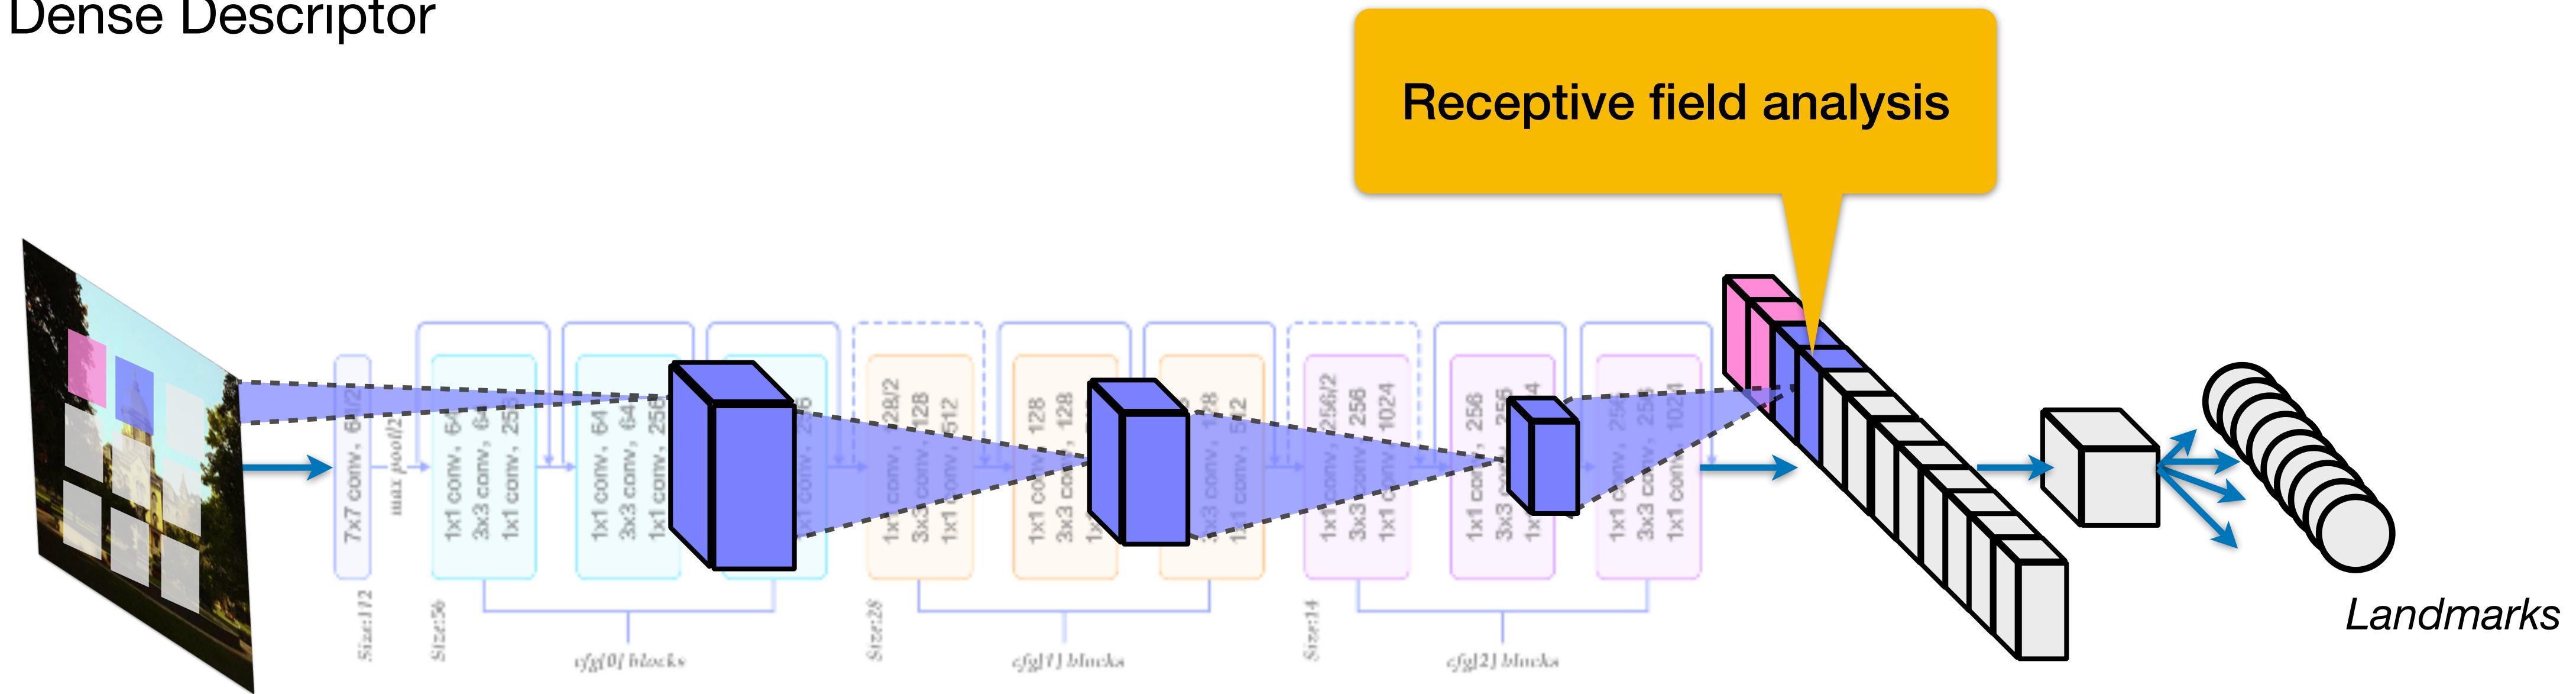
Dense Descriptor



Deep Local Features (DELF)

Dense Feature Extraction

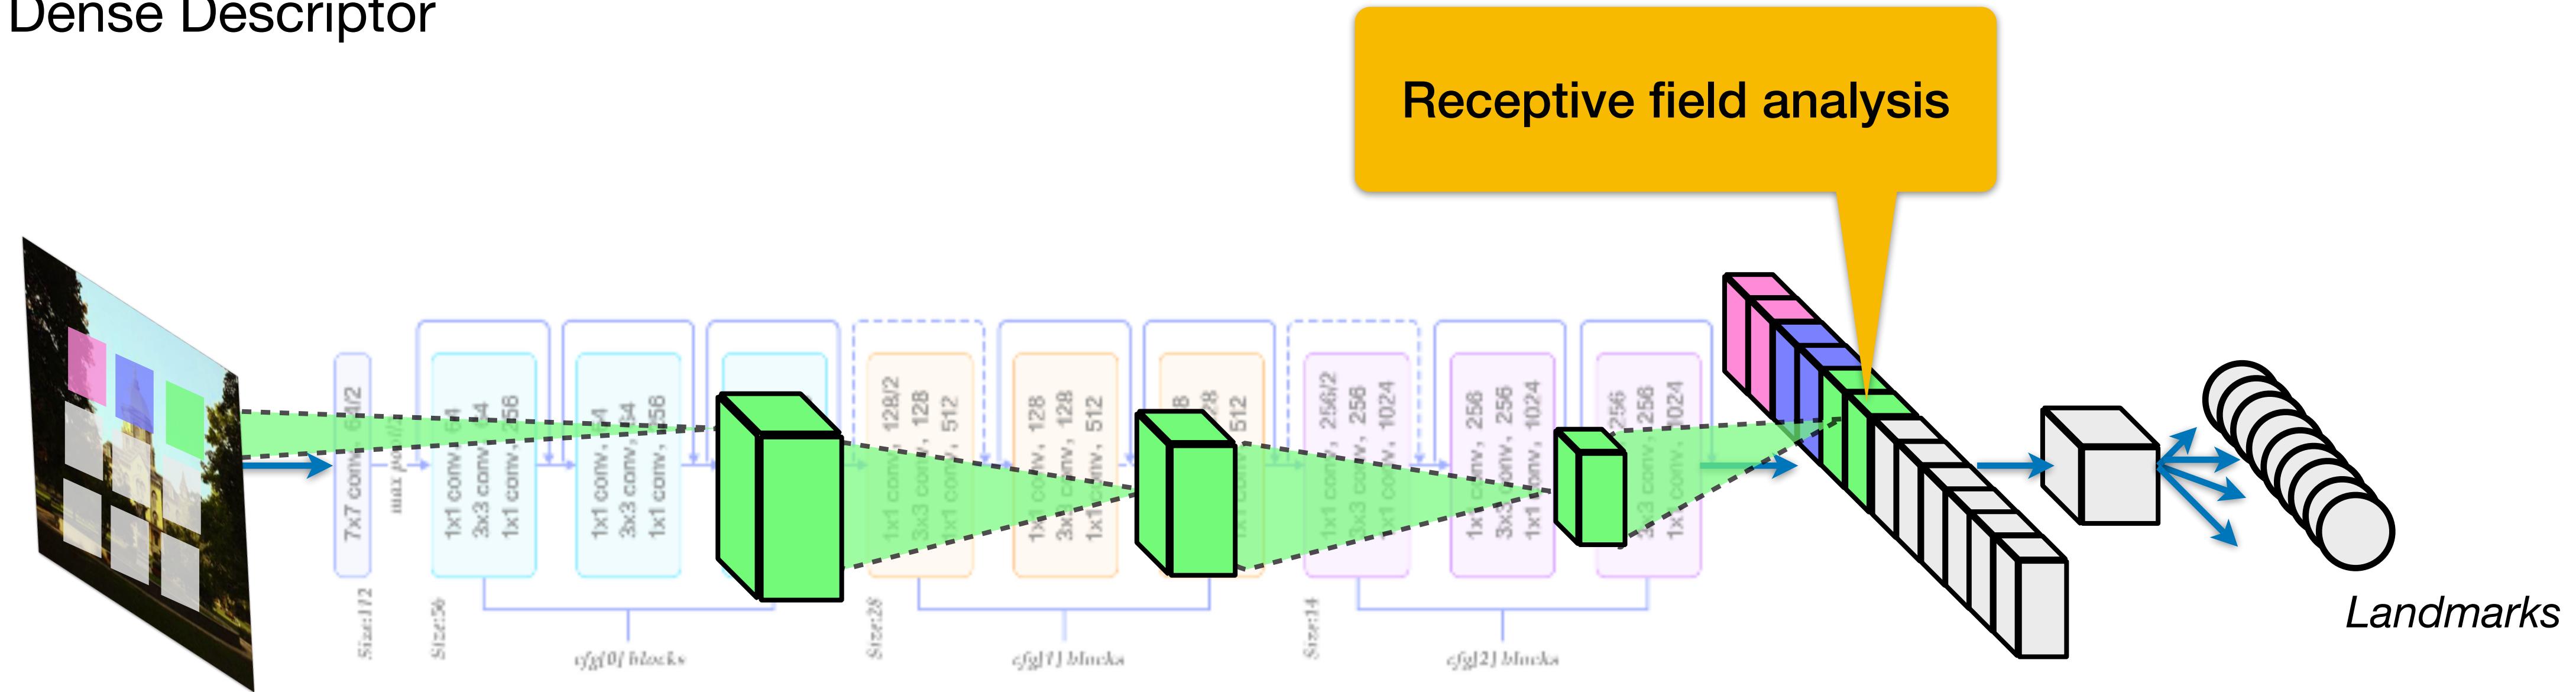
Dense Descriptor



Deep Local Features (DELF)

Dense Feature Extraction

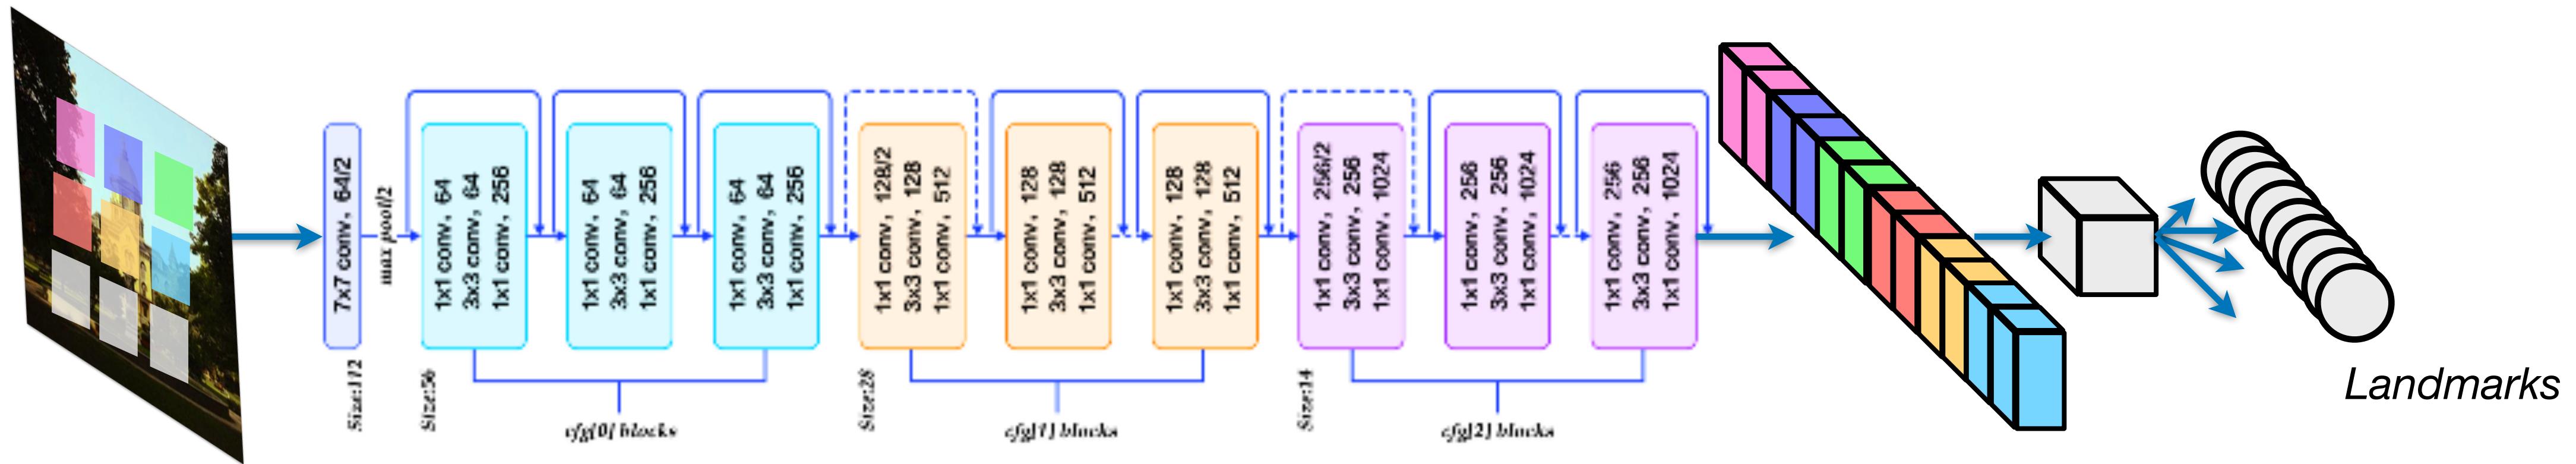
Dense Descriptor



Deep Local Features (DELF)

Dense Feature Extraction

Dense Descriptor

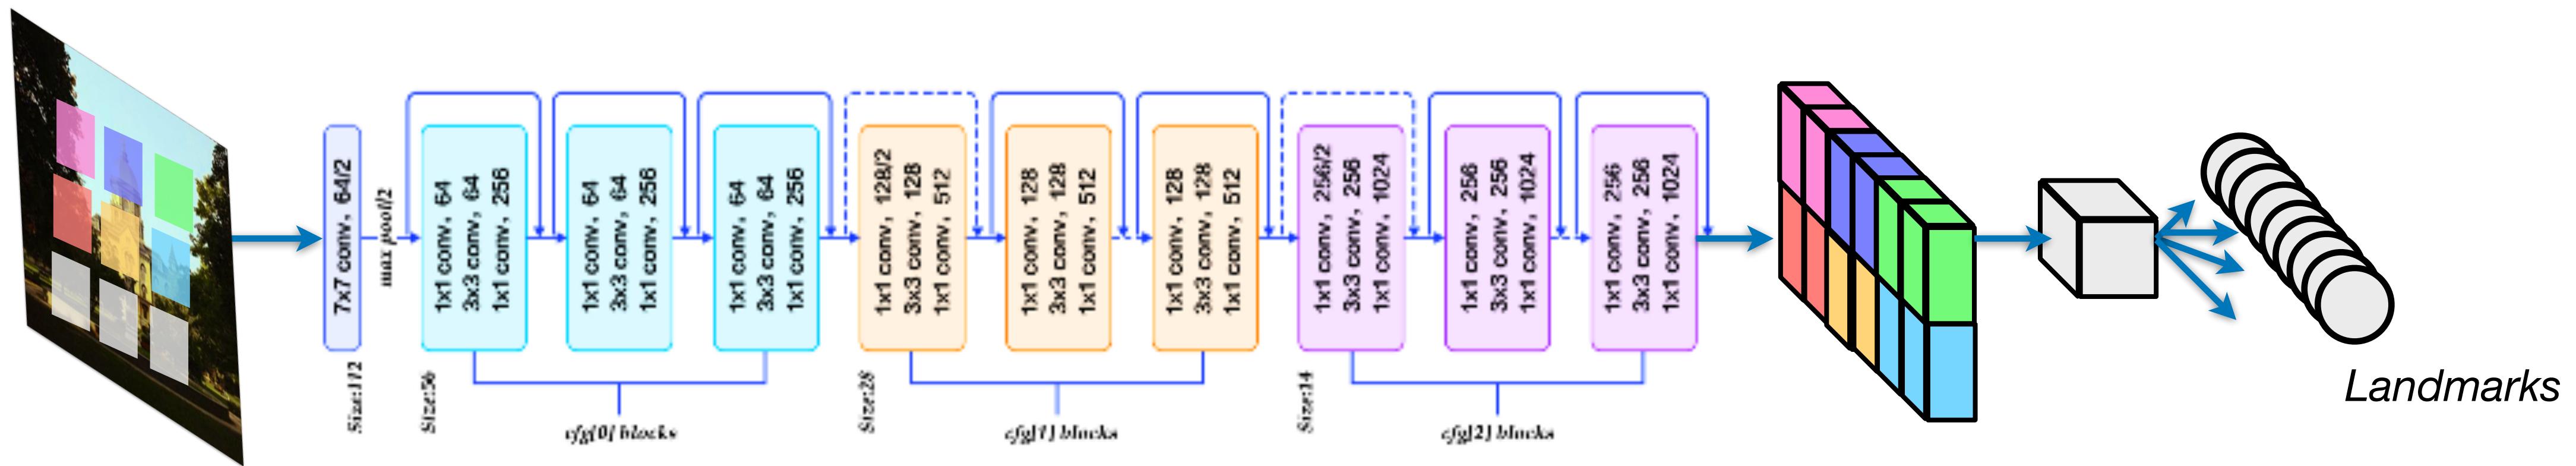


Deep Local Features (DELF)

Dense Feature Extraction

Dense Descriptor

Logical meaning: one feature vector for each densely sampled image region.

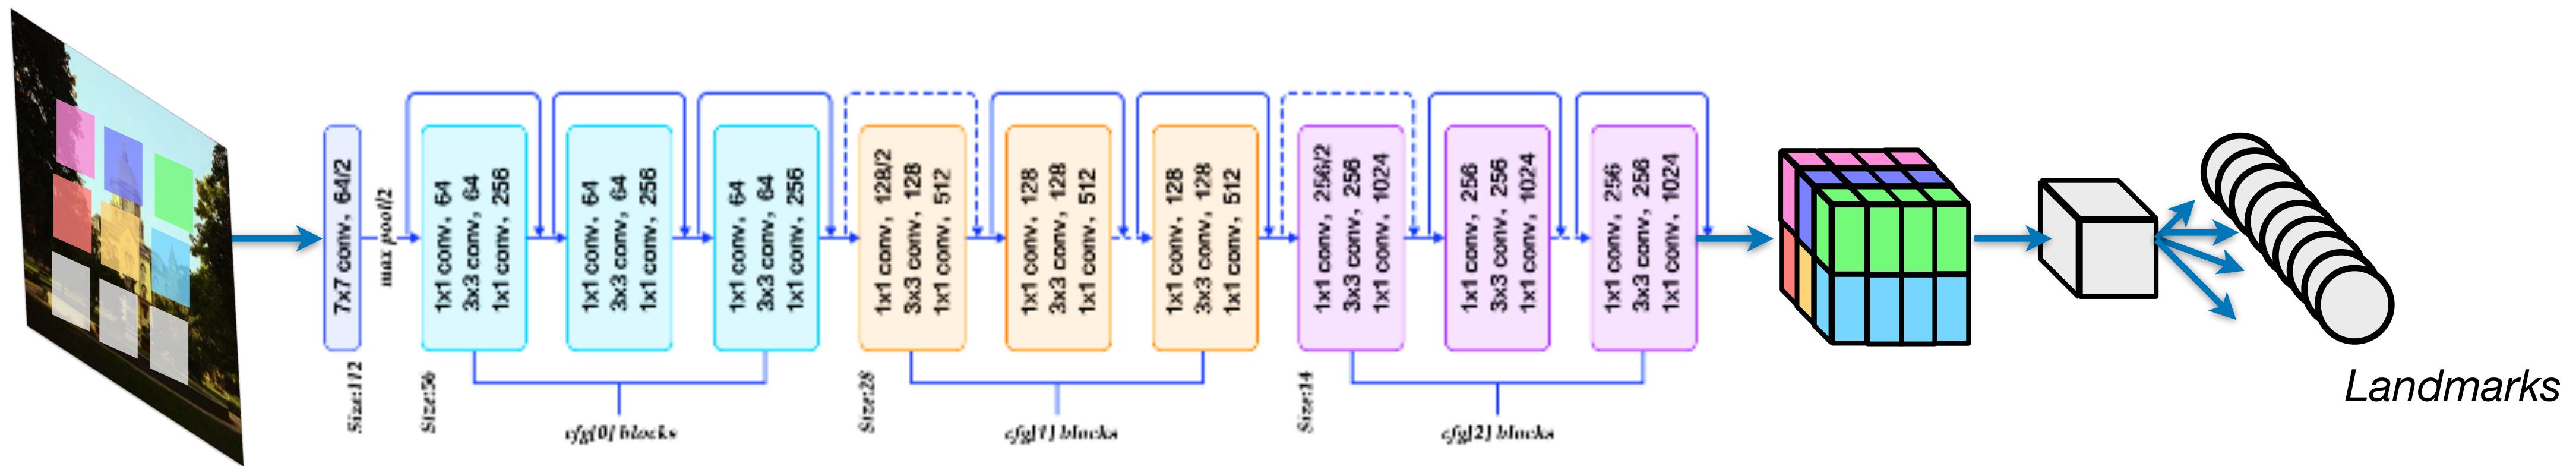


Deep Local Features (DELF)

Dense Feature Extraction

Dense Descriptor

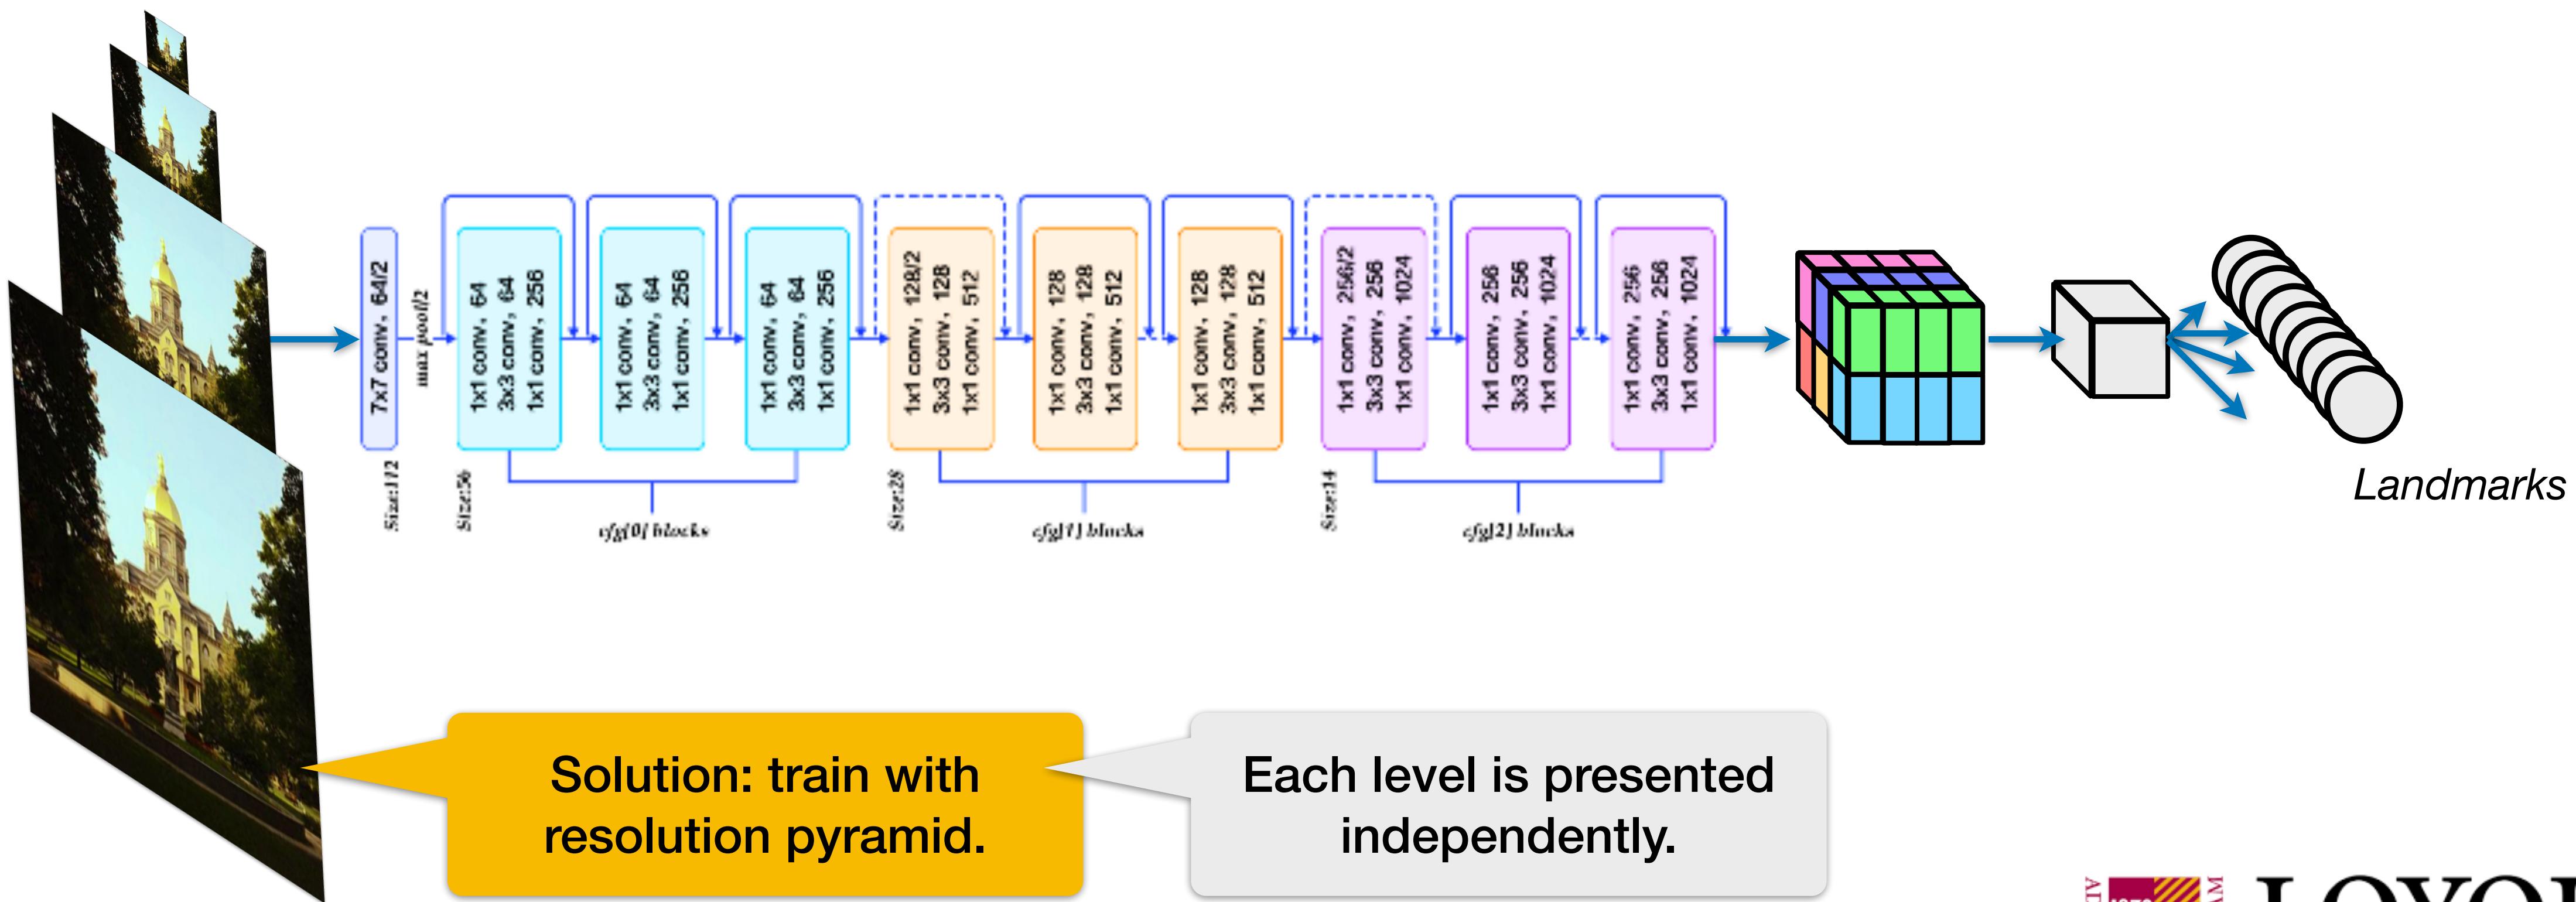
Logical meaning: one feature vector for each densely sampled image region.



Deep Local Features (DELF)

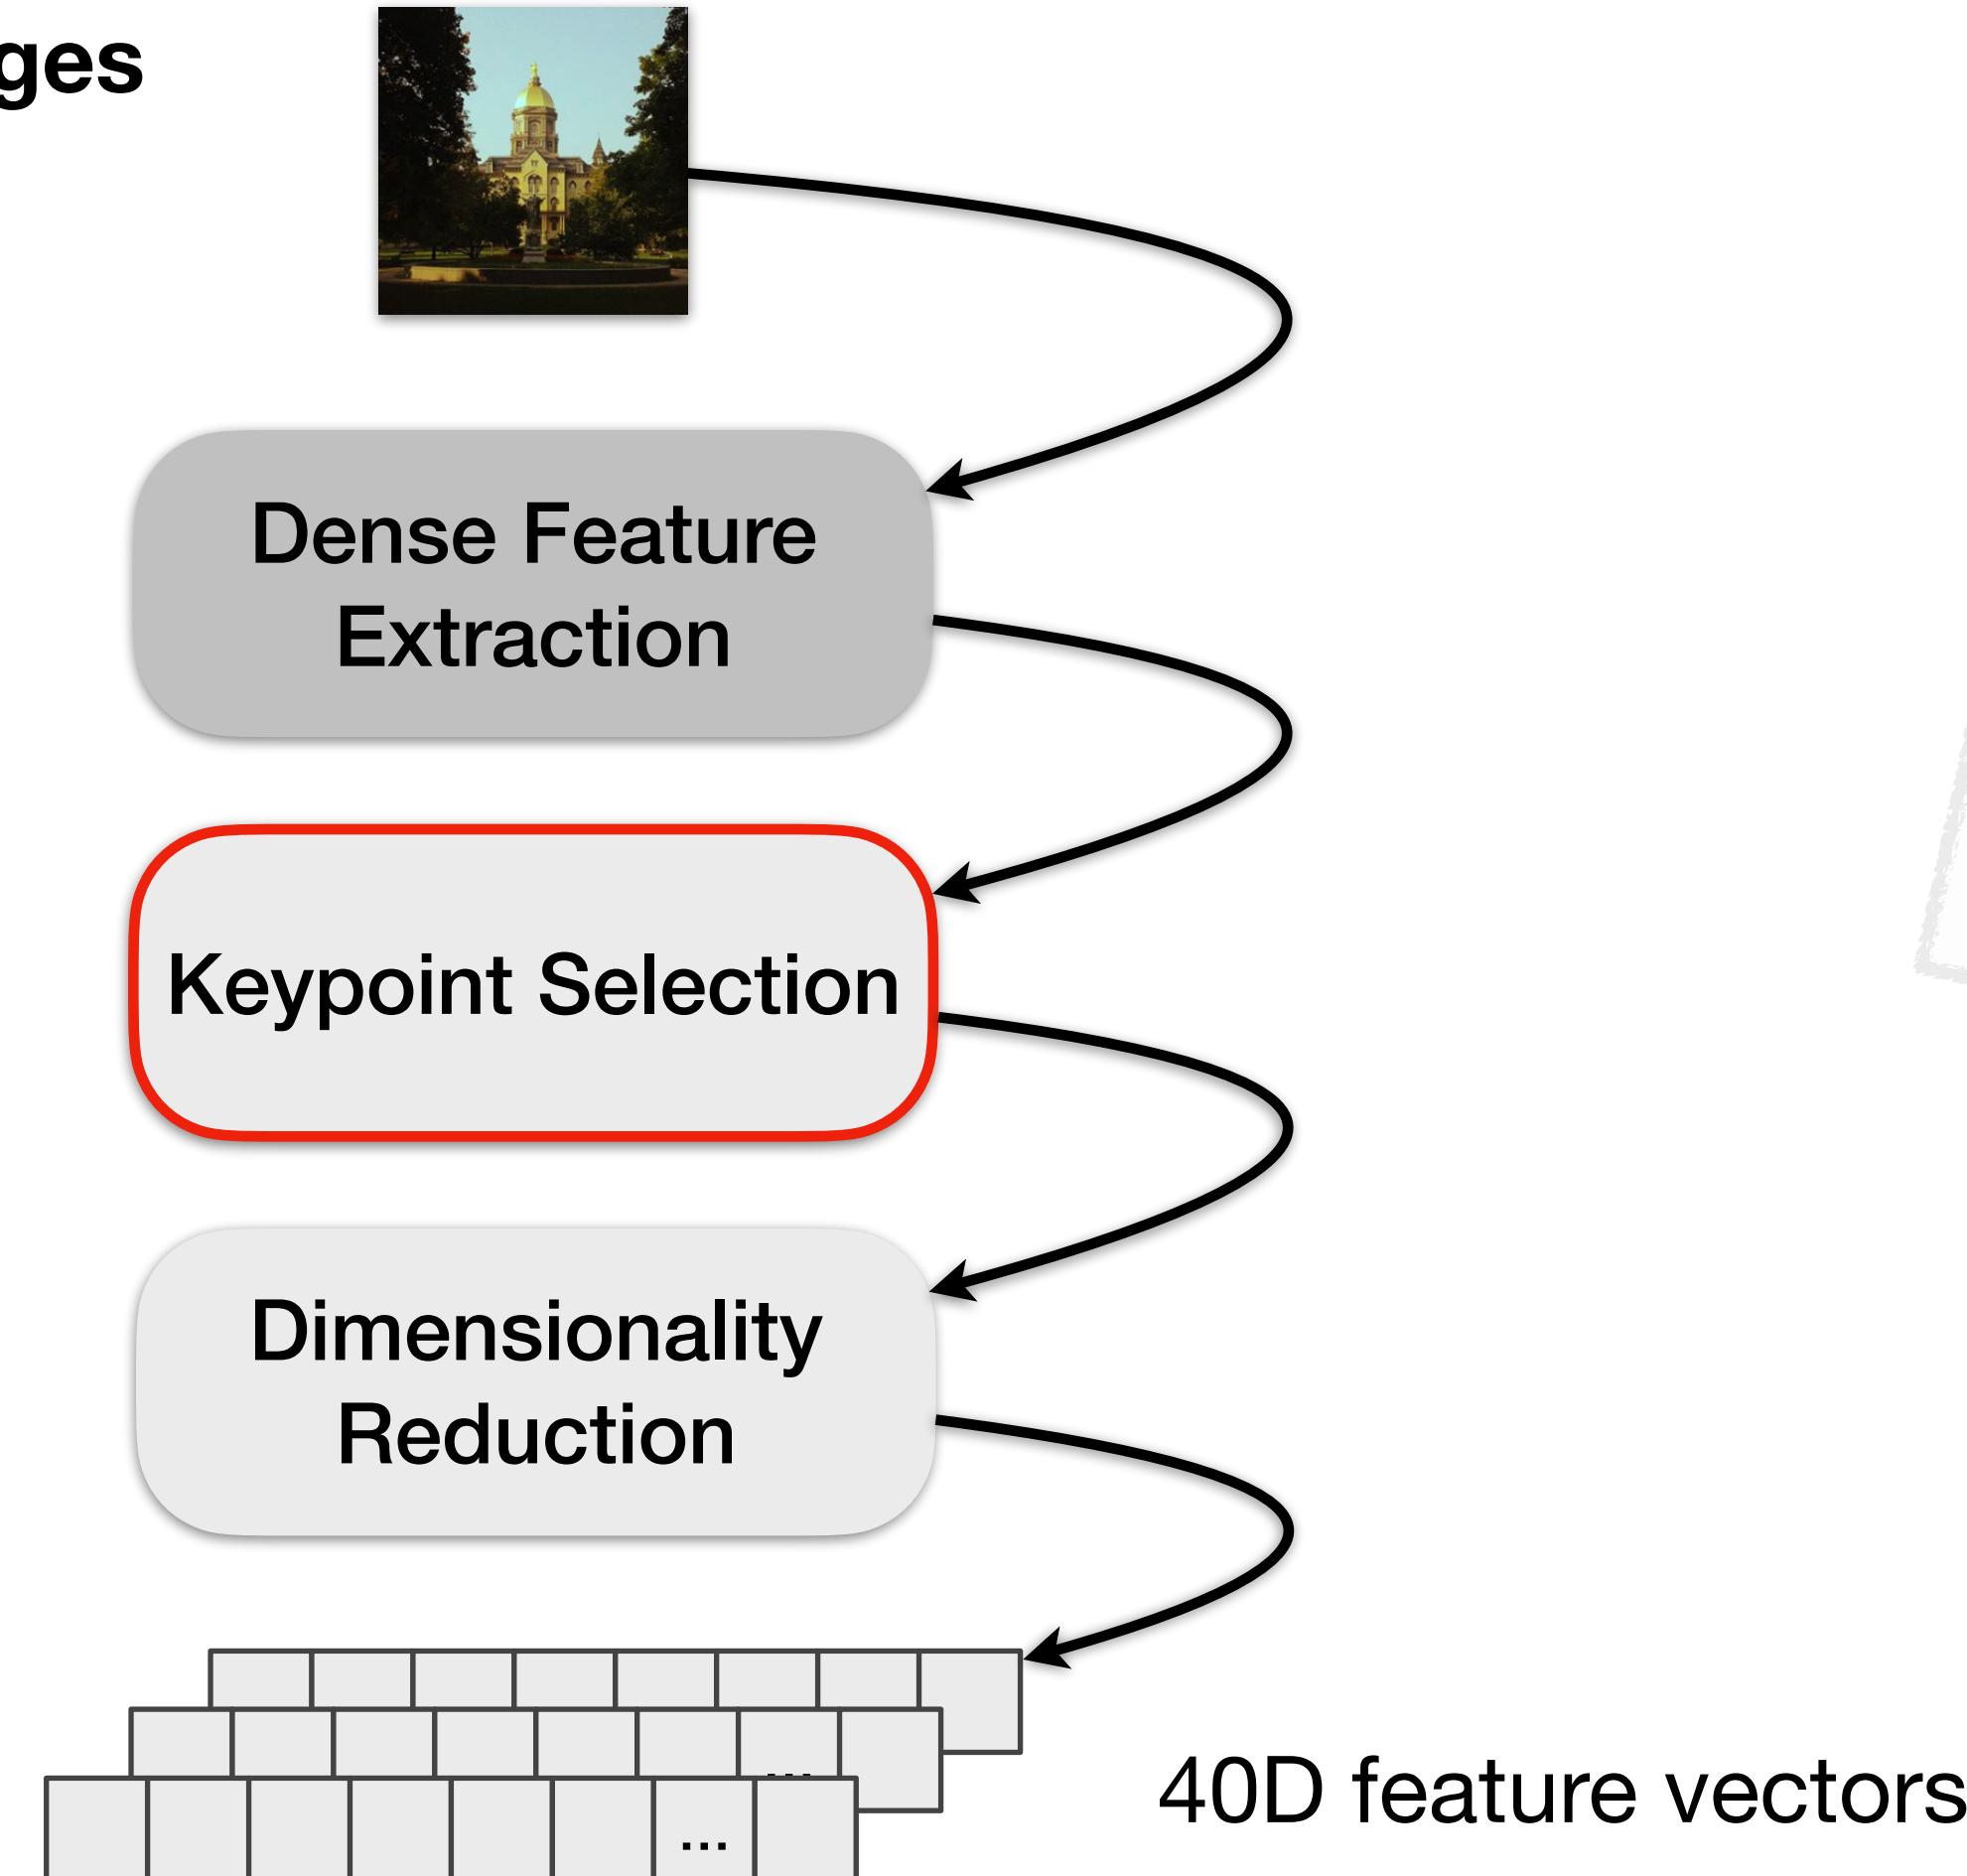
Dense Feature Extraction

How to deal with multiple resolutions?



Deep Local Features (DELF)

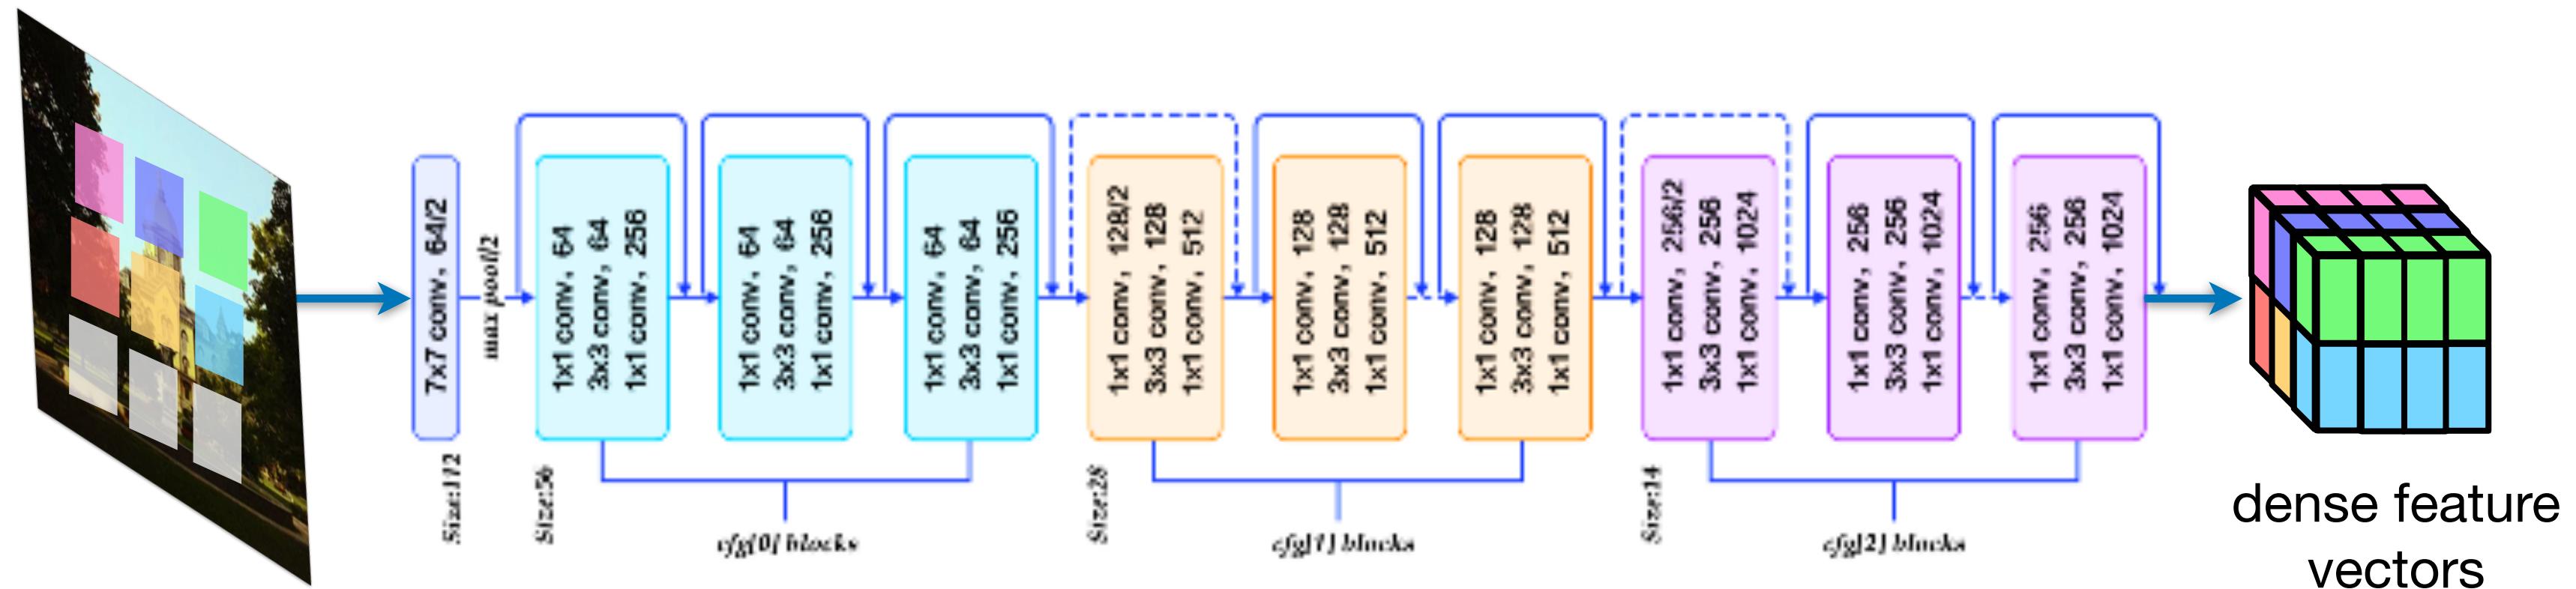
Stages



Deep Local Features (DELF)

Attention-based Keypoint Selection

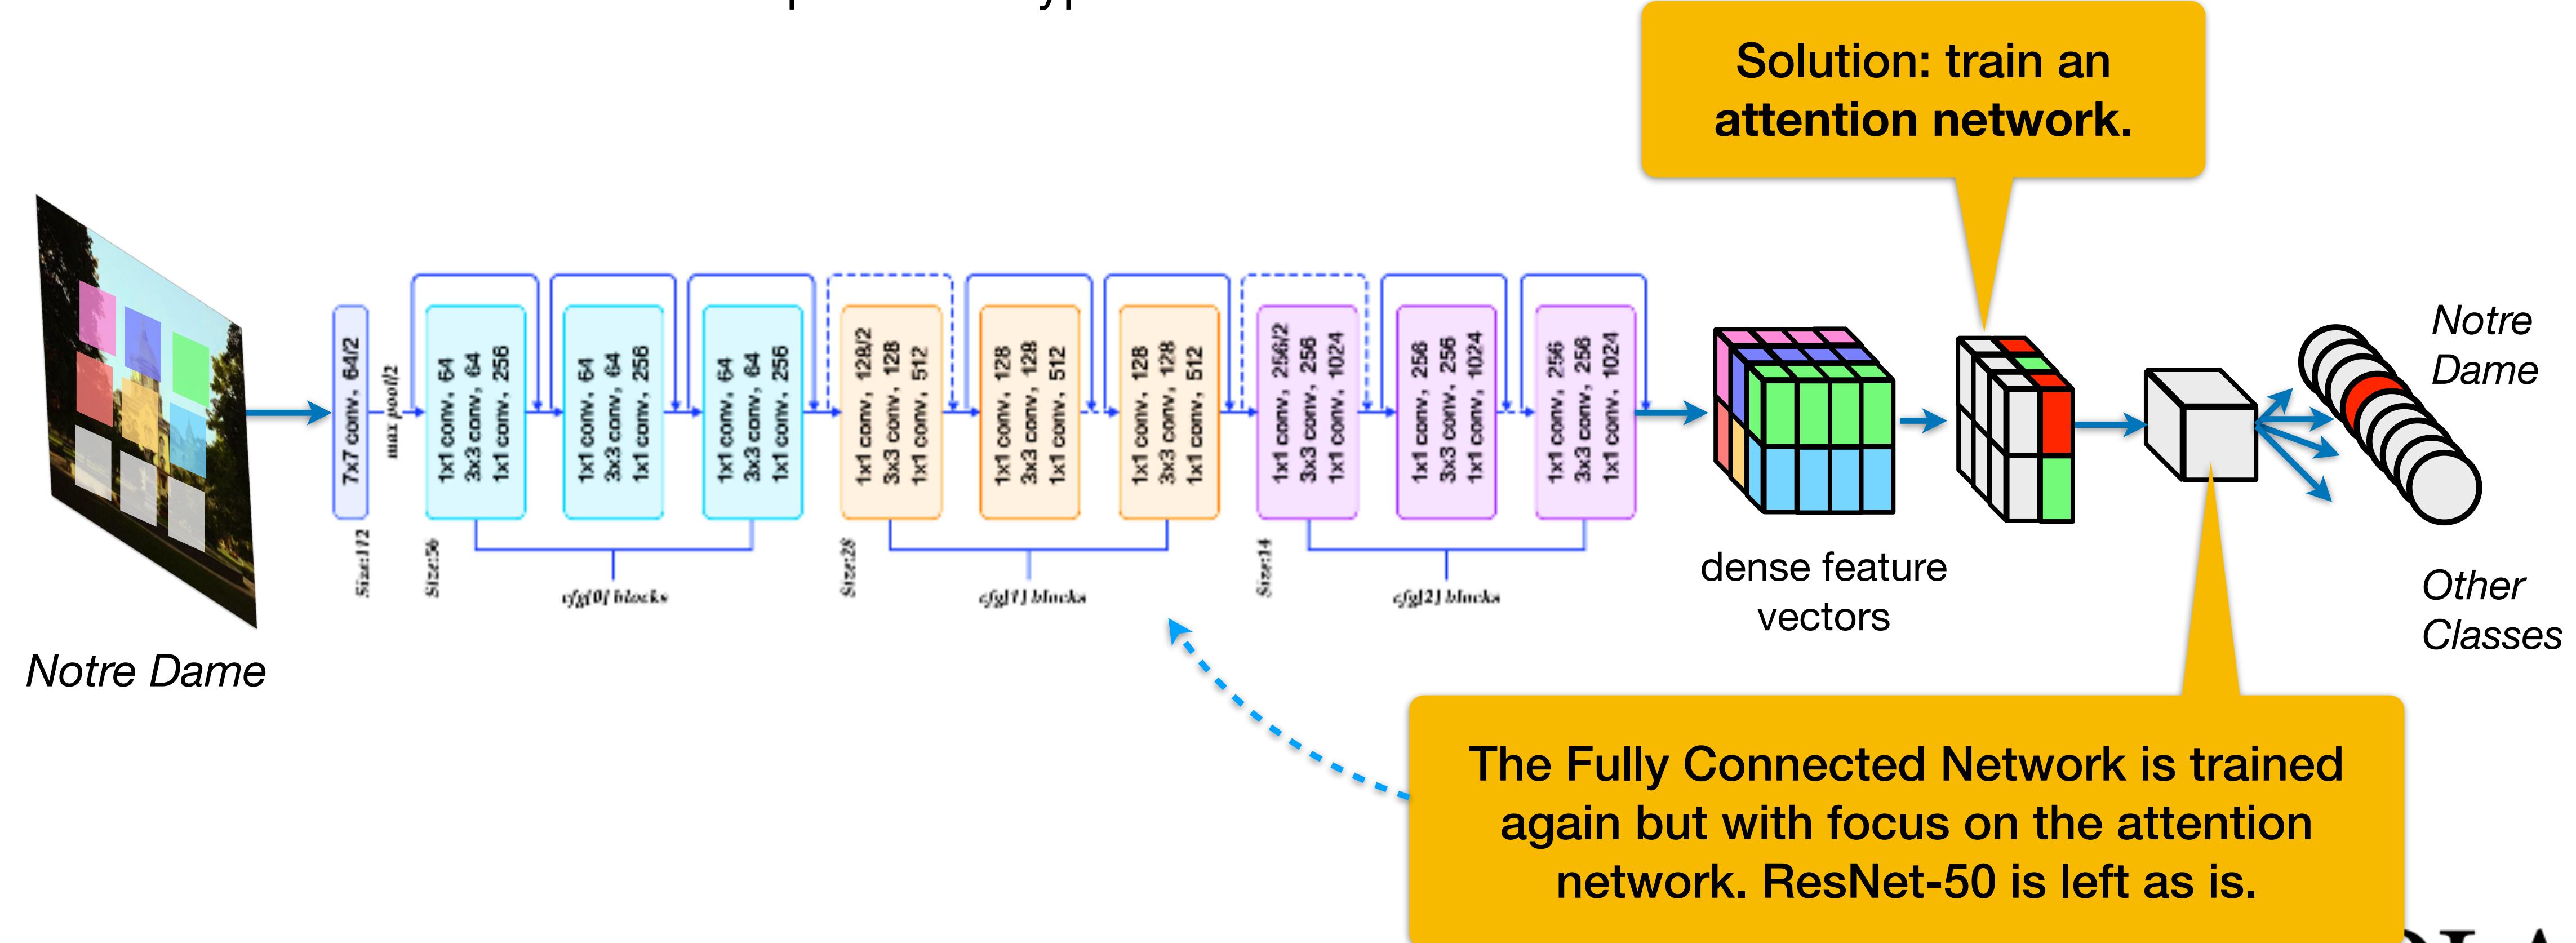
How to convert dense description to keypoint selection?



Deep Local Features (DELF)

Attention-based Keypoint Selection

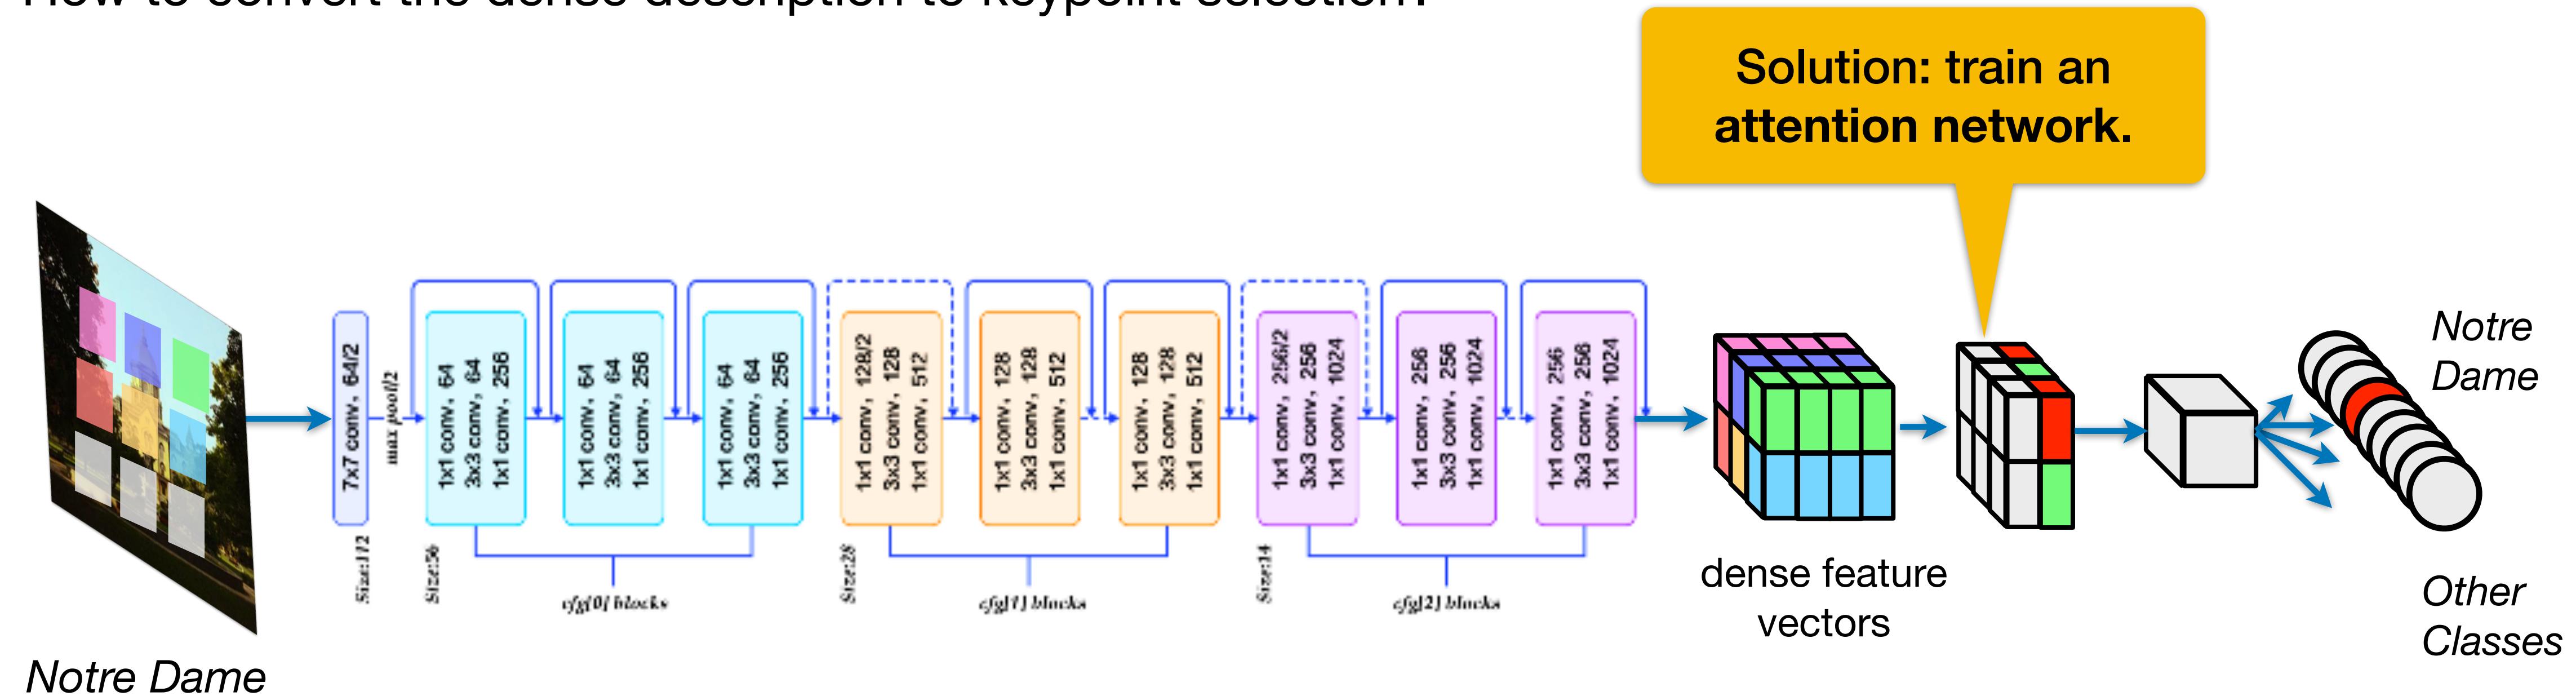
How to convert the dense description to keypoint selection?



Deep Local Features (DELF)

Attention-based Keypoint Selection

How to convert the dense description to keypoint selection?



Notre Dame

Attention Network: Weights the dense feature vectors by either suppressing or inciting them.

Rationale: only the locations helpful for classification are kept.

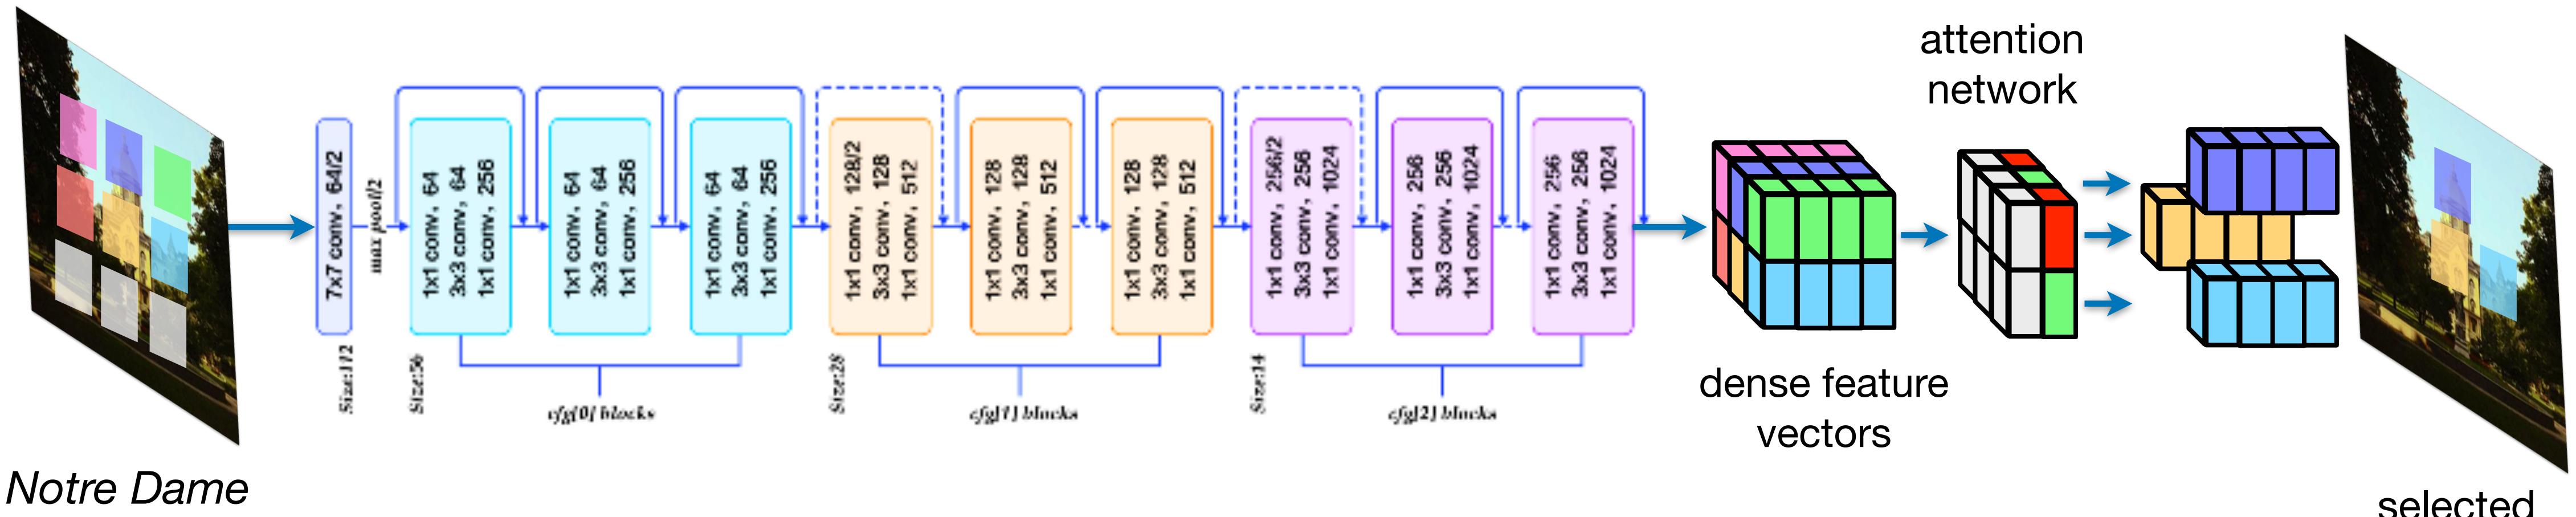


LOYOLA
UNIVERSITY CHICAGO

Deep Local Features (DELF)

Attention-based Keypoint Selection

How to convert the dense description to keypoint selection?



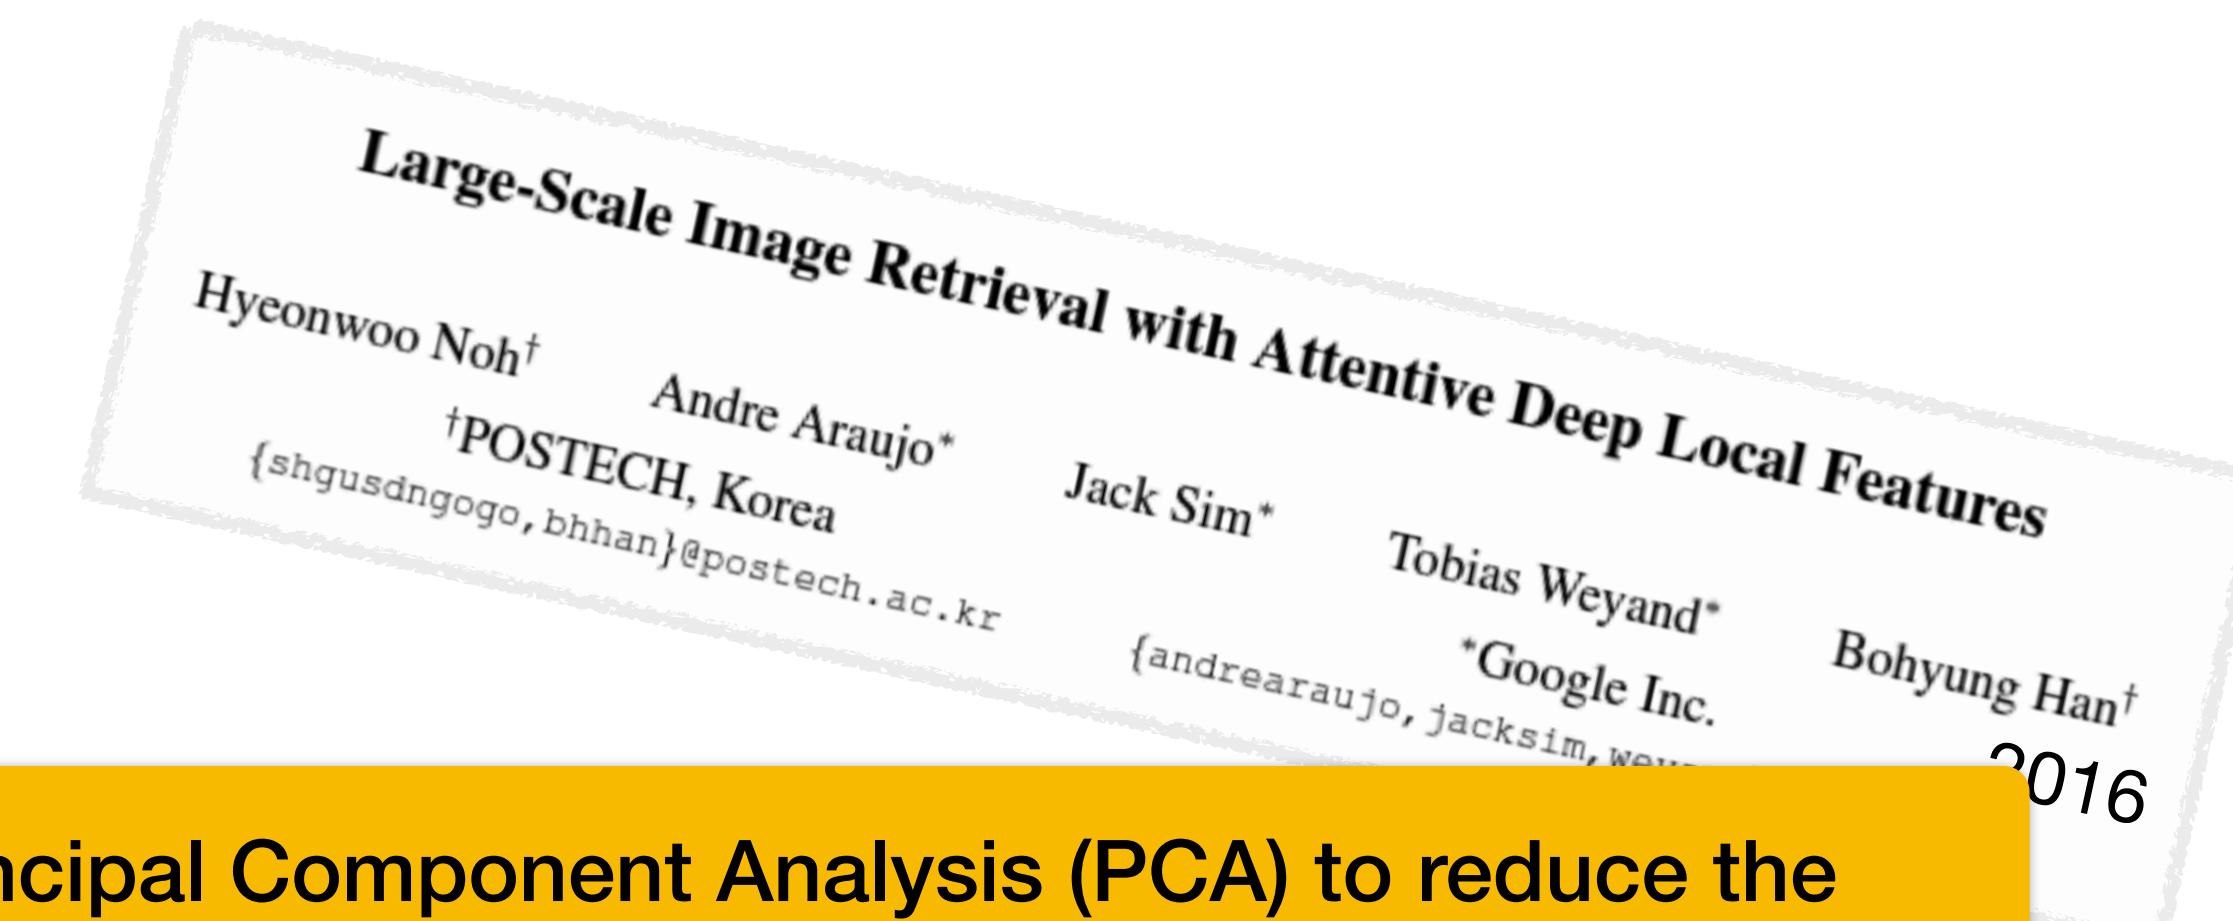
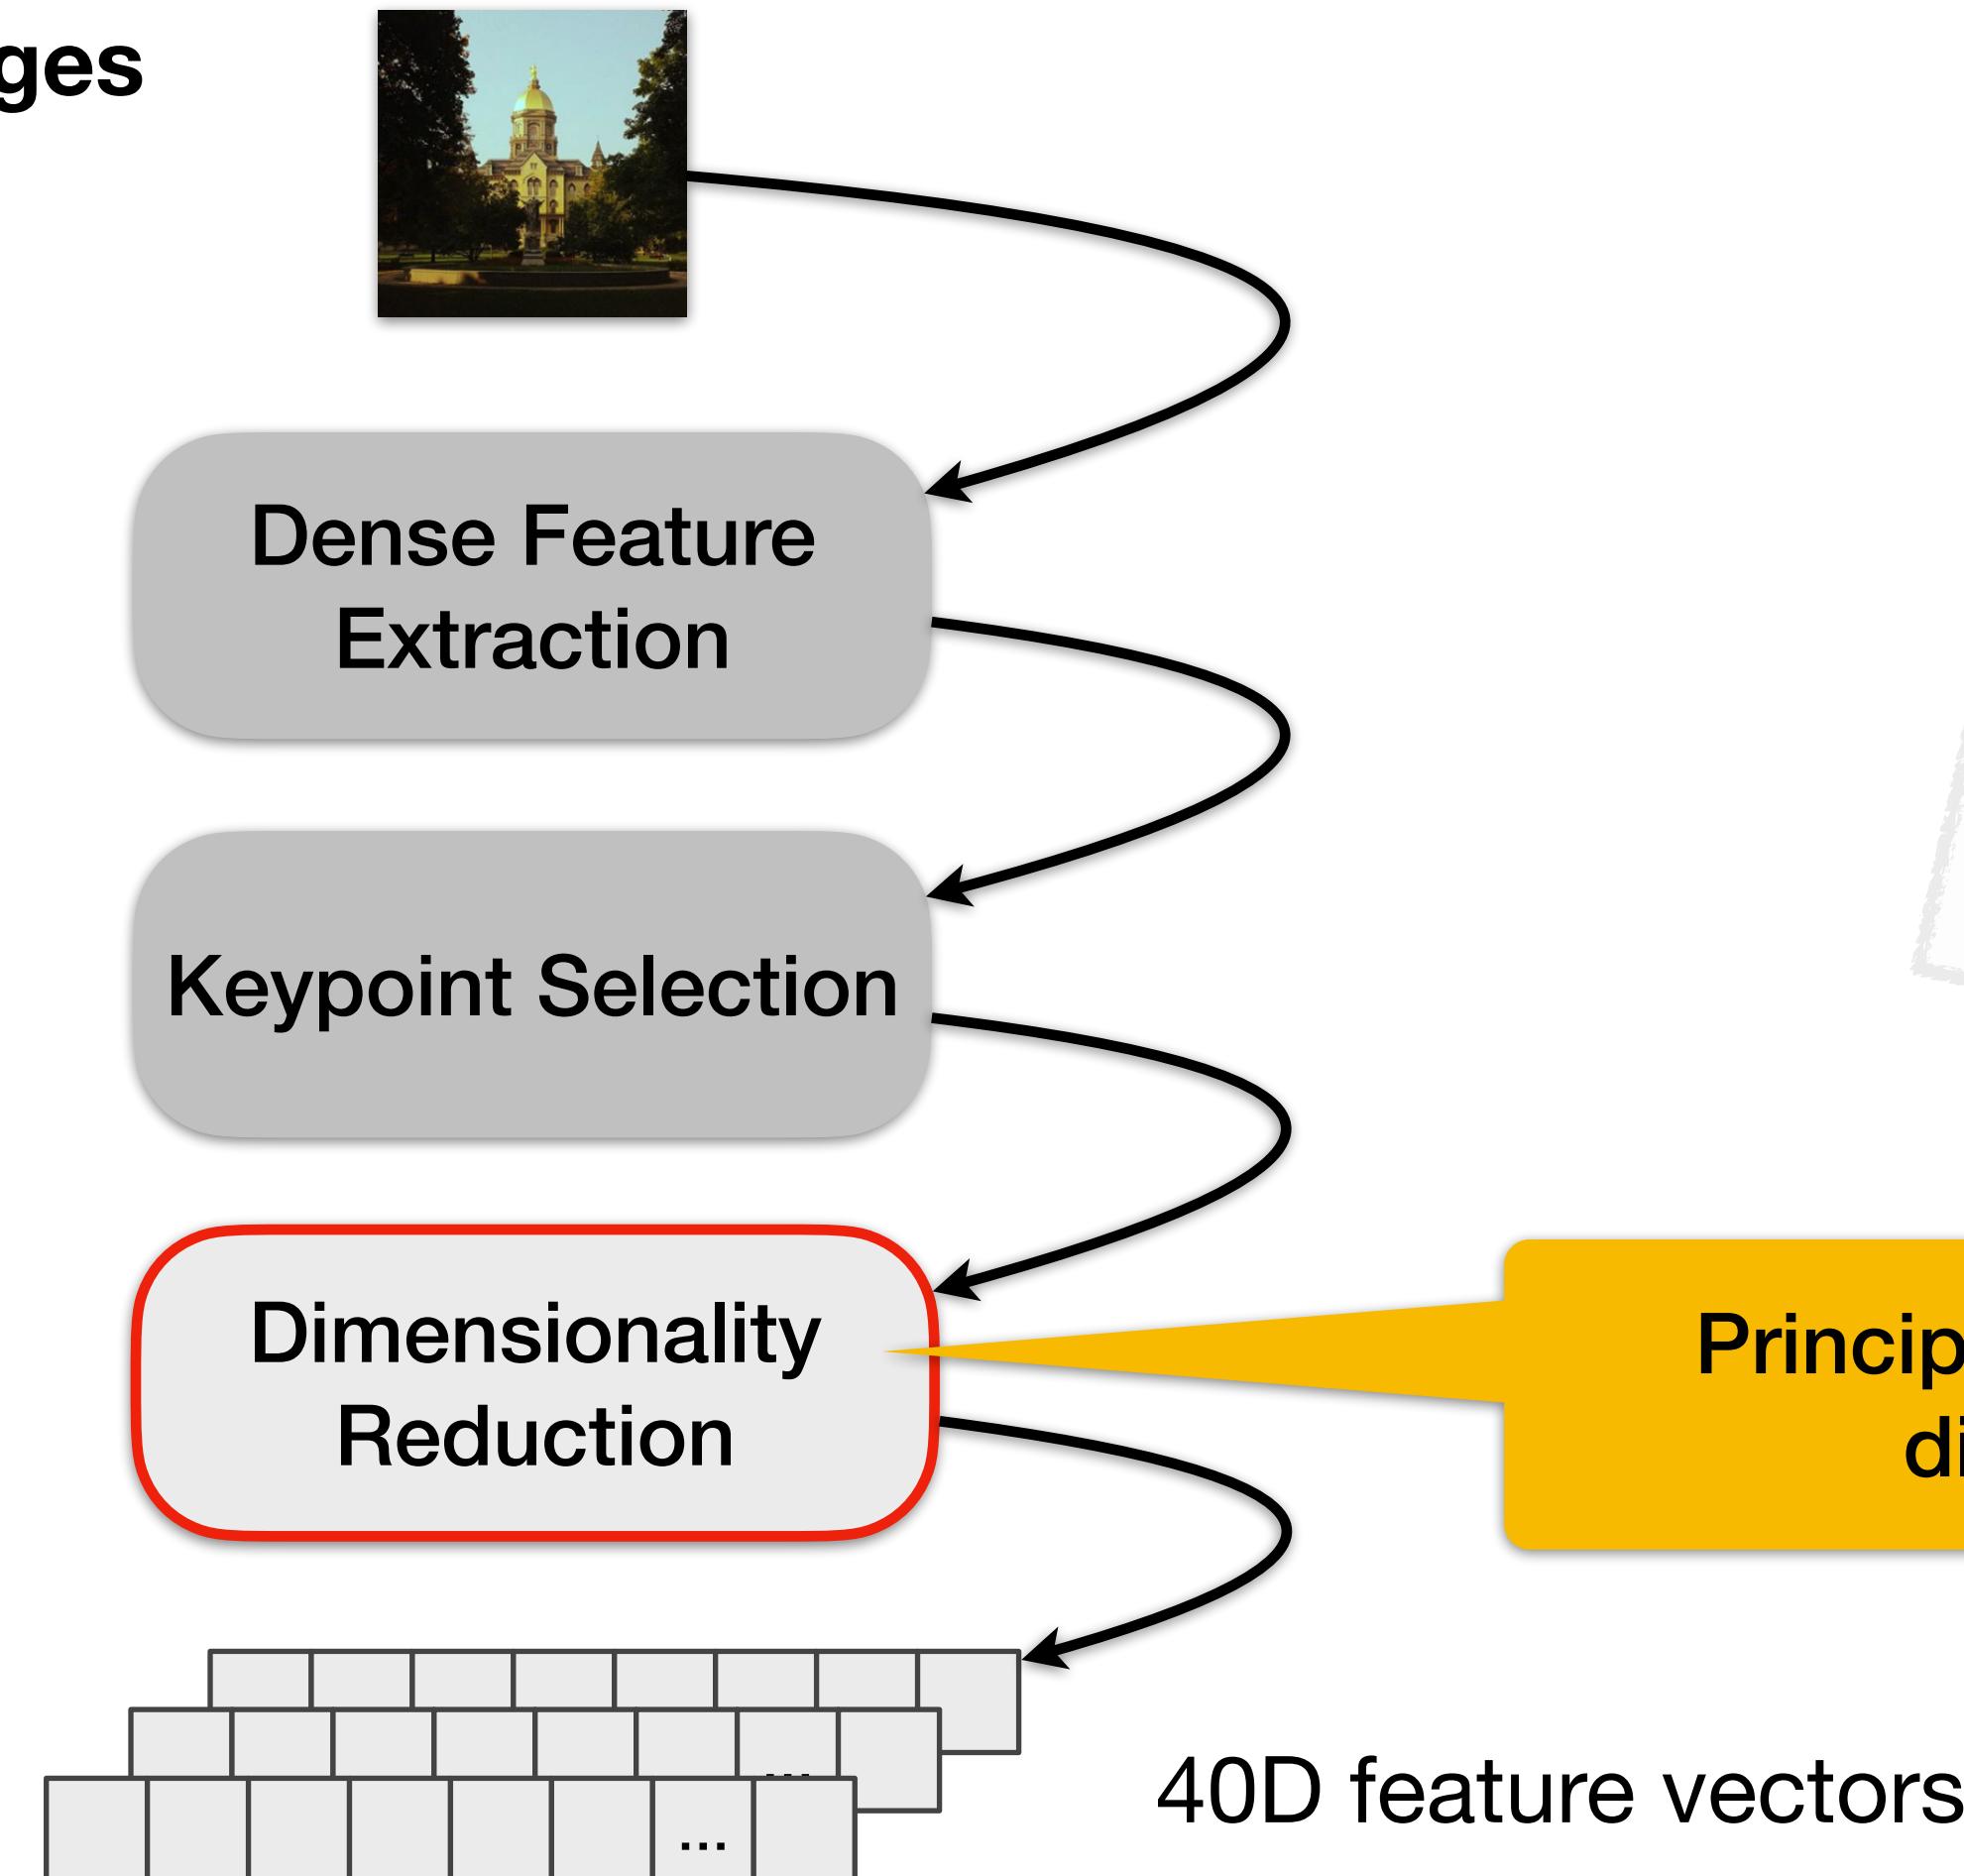
Notre Dame

Attention Network: Weights the dense feature vectors by either suppressing or inciting them.

Rationale: only the locations helpful for classification are kept.

Deep Local Features (DELF)

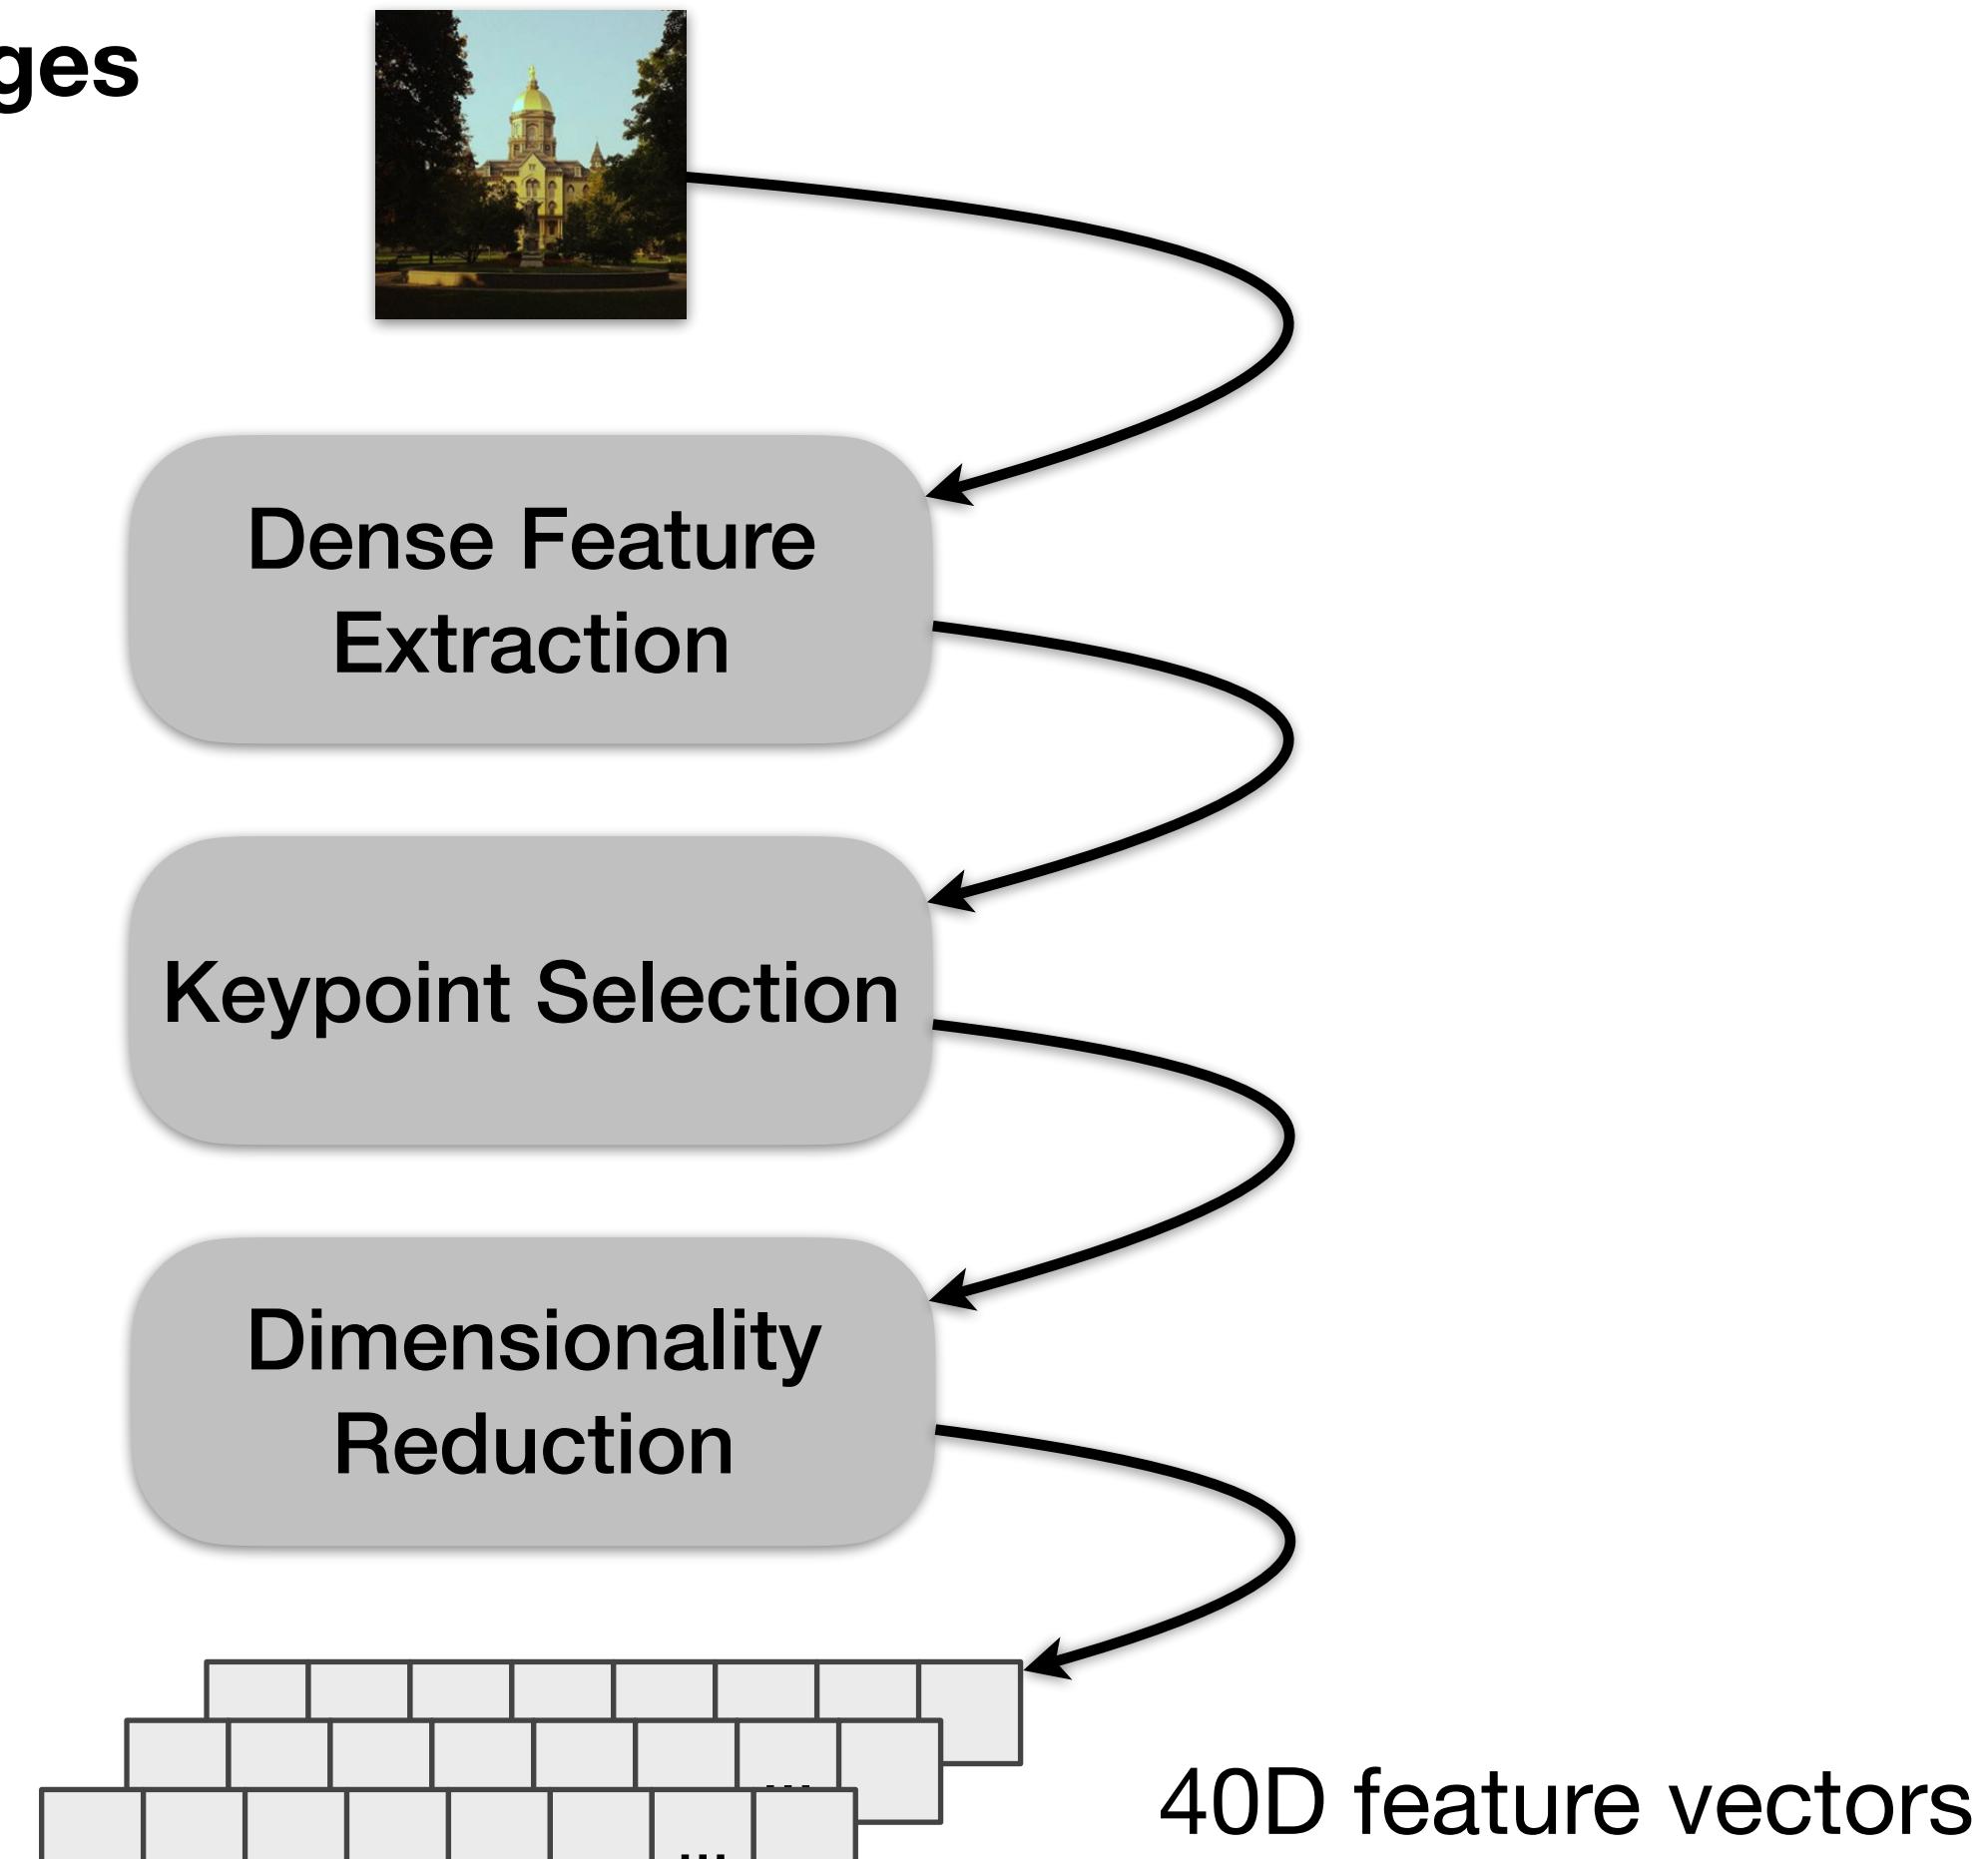
Stages



Principal Component Analysis (PCA) to reduce the dimension of the feature vectors to 40D

Deep Local Features (DELF)

Stages



Pros

Learning-based method
(not handcrafted).

Cons

DELF keypoints have (x, y) position and response (attention network weight).
They have neither scale nor orientation.

What do you think?

Are they scale and rotation invariant?

What's Up Next?

Usage of Feature Vectors

Nick and Jesus will lead the discussion of Image Retrieval.

Deadline of Assignment #2

Tomorrow is the deadline for submitting assignment #2. Let's share our thoughts with Nick and Jesus.

Release of Assignment #3

Topic: Image classification.

